



Matter Specification

Version 1.6

Document:	23-27349-011_Matter-1.6-Core-Specification.pdf June 16, 2026
Sponsored by:	Connectivity Standards Alliance
Accepted by:	This document has been accepted for release by the Connectivity Standards Alliance Board of Directors on June 16, 2026
Abstract:	The Matter specification defines fundamental requirements to enable an interoperable application layer solution for smart home devices over the Internet Protocol.
Keywords:	Referenced in Chapter 1.

Copyright © 2022-2026 Connectivity Standards Alliance, Inc.
508 Second Street, Suite 109B Davis, CA 95616 - USA
www.csa-iot.org
All rights reserved.

Permission is granted to members of the Connectivity Standards Alliance to reproduce this document for their own use or the use of other Connectivity Standards Alliance members only, provided this notice is included. All other rights reserved. Duplication for sale, or for commercial or for-profit use is strictly prohibited without the prior written consent of the Connectivity Standards Alliance.



Matter Specification

Version 1.6, 2026-06-02 10:31:01 -0700: Approved

Copyright Notice, License and Disclaimer

Copyright © Connectivity Standards Alliance (2021-2026). All rights reserved. The information within this document is the property of the Connectivity Standards Alliance and its use and disclosure are restricted, except as expressly set forth herein.

Connectivity Standards Alliance hereby grants you a fully-paid, non-exclusive, nontransferable, worldwide, limited and revocable license (without the right to sublicense), under Connectivity Standards Alliance's applicable copyright rights, to view, download, save, reproduce and use the document solely for your own internal purposes and in accordance with the terms of the license set forth herein. This license does not authorize you to, and you expressly warrant that you shall not: (a) permit others (outside your organization) to use this document; (b) post or publish this document; (c) modify, adapt, translate, or otherwise change this document in any manner or create any derivative work based on this document; (d) remove or modify any notice or label on this document, including this Copyright Notice, License and Disclaimer. The Connectivity Standards Alliance does not grant you any license hereunder other than as expressly stated herein.

Elements of this document may be subject to third party intellectual property rights, including without limitation, patent, copyright or trademark rights, and any such third party may or may not be a member of the Connectivity Standards Alliance. Connectivity Standards Alliance members grant other Connectivity Standards Alliance members certain intellectual property rights as set forth in the Connectivity Standards Alliance IPR Policy. Connectivity Standards Alliance members do not grant you any rights under this license. The Connectivity Standards Alliance is not responsible for, and shall not be held responsible in any manner for, identifying or failing to identify any or all such third party intellectual property rights. Please visit www.csa-iot.org for more information on how to become a member of the Connectivity Standards Alliance.

This document and the information contained herein are provided on an "AS IS" basis and the Connectivity Standards Alliance DISCLAIMS ALL WARRANTIES EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO (A) ANY WARRANTY THAT THE USE OF THE INFORMATION HEREIN WILL NOT INFRINGE ANY RIGHTS OF THIRD PARTIES (INCLUDING WITHOUT LIMITATION ANY INTELLECTUAL PROPERTY RIGHTS INCLUDING PATENT, COPYRIGHT OR TRADEMARK RIGHTS); OR (B) ANY IMPLIED WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, TITLE OR NONINFRINGEMENT. IN NO EVENT WILL THE CONNECTIVITY STANDARDS ALLIANCE BE LIABLE FOR ANY LOSS OF PROFITS, LOSS OF BUSINESS, LOSS OF USE OF DATA, INTERRUPTION OF BUSINESS, OR FOR ANY OTHER DIRECT, INDIRECT, SPECIAL OR EXEMPLARY, INCIDENTAL, PUNITIVE OR CONSEQUENTIAL DAMAGES OF ANY KIND, IN CONTRACT OR IN TORT, IN CONNECTION WITH THIS DOCUMENT OR THE INFORMATION CONTAINED HEREIN, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH LOSS OR DAMAGE.

All company, brand and product names in this document may be trademarks that are the sole property of their respective owners.

This notice and disclaimer must be included on all copies of this document.

Connectivity Standards Alliance
508 Second Street, Suite 206
Davis, CA 95616, USA

Participants

Agrawal, Amit	Alagappan, Vignesh	Alami, Amine	Alexander, Rob
Allegue, Michel	Ananthakrishnan, Krithika	Ashok Kumar, Thivya	Asterjadhi, Alfred
Axelsson, Ulf	Azria, Shana	Bak, Naama	Balducci, Alex
Bao, Yongming	Bartolome, Diego	Bauer-Schwan, Stefan	Beach, Chris
Beck, Austin	Becker, Markus	Beliveau, Louis-Philip	Ben, Thomas
Bhetanabottla, Sriram	Binjola, Sharad	Boehl, Florian	Bollimpalli, Satyanaag
Bonnell, Corey	Borges, Daniel	Boyington, Marcos	Brandbyge, Søren
Bultman, Rob	C, Rajashree	Carlin, Broderick	Carmel-Veilleux, Tennessee
Carter, Harrison	Casallas, Ricardo	Cave, Tony	Chalmers, Andrew
Chan, Osborn	Chandarana, Janak	Chen, Shu	Cheshire, Stuart
Chudinov, Adrian	Chung, Cliff	Chupp, Anton	Coppock, Kevin
Cowan, Michael	Cragie, Robert	Crettenand, Alexander	Cullen, Sam
Cuyckens, Thomas	Damle, Makarand	Darling, Don	De, Pradip
Decenzo, Chris	Delplancke, Julien	Desbiolles, Arnaud	Dhayagude, Hrishikesh
Dieter Van der Meulen	Ding, Li-An	Dok, Hrishikesh	Dolan, David
Dong, Kangping	Drake, Jeff	Duda, Łukasz	Dyck, Nathan
Erickson, Grant	Farnum, Robert	Feraru, Eugen	Finet, Vincent
Fischer, Ingo	Fominaya, Antonio	Foreman, Zack	Franck, Fabien
Freeman, Cecille	Fu, Kenneth Xiaoming	Fuentes, Pedro	Fyall, Ian
Garbus, Mathias Kielgast	Garg, Pankaj	Goto, Kendall	Graf, Tobias
Granbery, Hasty	Griffiths, Andrew	Gucea, Doru	Guiheneuf, Robin-Charles
Guo, Jiacheng	Guo, Song	Haefner, Kyle	Hamilton, Ryan
Hampson, Terence	Hanna, John	Hanna, Steve	Haque, Asad
Harris, Will	Harrow, James	Hartwig, Thomas	Harvey, Gene
Hassanli, Arya	Hazley, Matt	Hegedič, Matjaž	Heide, Janus
Heo, Joonhaeng	Hernandez-Palomares, Martin	Hicklin, William	Hilal, Rawad
Ho, Nguyen	Hoang, Minhhoa	Holbrook, Trevor	Hollebeek, Tim
Houtepen, Rob	Huang, Xiaolong	Hui, Jonathan	Hui, Li
Hui, Yang	Jackson, Donald Clark	Jain, Amit	Jain, Ankur

Jandhyala, Chaitanya	Jayakumar, Liju	Jo, Youngjun	Johns, Jerry
Josefsen, René	KVS, Mohan	KY, Suma	Karabashov, Aziz
Kardous, Mathieu	Kasperczyk, Kamil	Kassel, Gabe	Katira, Utsav
Kavali, Vishnu	Khatrri, Shivam	Kim, Eun Chul	Knörzer, Clemens
Kohr, John	Kommareddi, Naveen	Kontra, Andrew	Koos, Chris
Kovacic, Lazar	Krawetz, Bryan	Królik, Damian	Kumar, Sandeep
Kumar, Saurabh	Lauric, Petru	Lazar, Alin	Le Tutour, Jean
Lea, Thomas	Lee, Byungjoo	Lepage, Marc	Letnick, Chris
Levkov, Stoycho	Liang, Deng	Lindeman, Ryan	Litvin, Andrei
Liu, Eltham	Lorente, Olivier	Lyu, Rashid	Ma, Longfei
Maes, Timothy	Makdissi, Shadi	Mamo, Fesseha	Manley, Tom
Mann, Bryan	Mansour, Peter	Mapp, Chris	Marche, Mikael
Margolis, Evgeni	Marquez, Raul	Martinez, Junior	Matignon, Florent
Matosian, Dan	Meissner, Bryan	Melo, Sara	Menzopol, Andrei
Mikolits, Marc	Moneta, Daniel	Montenegro, Gabriel	Morales, Victor
Morozov, Evgeniy	Morris, Simon	Mouquet, Antoine	Mégevand, Jonathan
Nadathur, Anush	Nagappan, Ramesh	Najaf-Zadeh, Hossein	Naruka, Dev
Neishaboori, Azin	Nicolas, Vivien	Nuyts, Wim	Oleson, Kiel
Olson, Rodney	P, Aswathy	Pacheco, JC	Page, Jason
Pan, Liam	Pan, Shaofeng	Pansy, Jürgen	Parausanu, Dragos
Patil, Shubham	Patwardhan, Aditya	Penven, Jean-Francois	Perumal, Saravana
Po, Kevin	Powell, Ken	Pyasi, Madhur	Rahman, Mo
Rasmussen, Søren Møller	Rempel, David	Rhees, Jon	Richard, Arnaud
Rosenberg, Aron	Rozendaal, Leo	Rupp, Michael	Rush, Matthew
Ryan, Kyle	S, Sowmya	Sallas, Sal	Samad, Abdul
Sambles, Philip	Sandstedt, Michael	Sarkar, Nivi	Sarma, Bhaskar
Scherbakov, Alex	Schiller, Bill	Schlosser, Walter	Schoinas, Yannis
Schultze, Julianne	Seenivasan, Suraj	Segal, Oren	Selvaraj, Vijay
Sena, Joe	She, Chengqiang	Shreve, Erik	Singh, Sumit
Siu, Irene	Smirl, Jon	Smith, Bill	Smith, David
Smith, Matt	Soares, Sergio	Soloway, Alan	Son, Jae
Spade, Lorey	Sperling, Karsten	Stonehouse, Alex	Struchtrup, Sebastian Schulze
Surnin, Dmitry	Swan, James	Szablowski, Michał	Szatmary-Ban, Zoltan

Szczodrak, Marcin	Szewczyk, Robert	Tabassum, Nadira	Taleb, Ali
Terrones, Moises	Trayer, Mark	Truskovsky, Alexander	Tung, Berkat
Turon, Martin	Tursunov, Abdulbois	Uebach, Till	Umesh, Deepakumar
Van der Beken, Peter	Vartanov, Igor	Vauclair, Marc	Verma, Lochan
Wais, Jackie	Wang Qixiang	Wang, David	Wang, Yufeng
Wang, Yunhan	Wechselberger, Fabian	Wei, Qingyun	Weil, Jason
Weinshel, Ben	Weir, Tristan	Williams, Cam	Winnik, Wojciech
Wood, Justin	Wyseur, Brecht	Xu, Yakun	Yang, Carol
Zang, Mingjie	Zbarsky, Boris	Zgrablic, Leonard	Zhang, Xili
Zhao, Betty	Zhao, Ru	Zhodzishsky, Victor	Zijian, Wang

Revision History

Revision	Date	Details	Editor
1	May 11, 2020	Initial draft	Robert Szewczyk
2	September 23, 2022	Version 1.0	Robert Szewczyk
3	May 17, 2023	Version 1.1	Robert Szewczyk
4	October 18, 2023	Version 1.2	Robert Szewczyk
5	April 17, 2024	Version 1.3	Robert Szewczyk
6	November 4, 2024	Version 1.4	Robert Szewczyk
7	March 17, 2025	Version 1.4.1	Robert Szewczyk
8	July 16, 2025	Version 1.4.2	Robert Szewczyk
9	November 10, 2025	Version 1.5	Robert Szewczyk
10	March 16, 2026	Version 1.5.1	Robert Szewczyk
11	June 16, 2026	Version 1.6	Robert Szewczyk

Table of Contents

Copyright Notice, License and Disclaimer	1
Participants	3
Revision History	7
1. Introduction	35
1.1. Scope and Purpose	35
1.2. Acronyms and Abbreviations	35
1.3. Definitions	38
1.4. Standards Terminology Mapping	41
1.5. Conformance Levels	42
1.6. References	42
1.6.1. Connectivity Standards Alliance Reference Documents	43
1.6.2. External Reference Documents	43
1.7. Informative References	50
1.7.1. Connectivity Standards Alliance Reference Documents	50
1.8. Conventions	50
1.8.1. Enumerations and Reserved Values	50
1.8.2. Reserved Bit Fields	51
1.8.3. Number Format	51
1.8.4. Provisional	52
2. Architecture	53
2.1. Overview	53
2.2. Layered Architecture	53
2.3. Network Topology	55
2.3.1. Single network	55
2.3.2. Star network topology	56
2.4. Scoped names	57
2.5. Identifiers	58
2.5.1. Fabric References and Fabric Identifier	58
2.5.2. Vendor Identifier (Vendor ID, VID)	58
2.5.3. Product Identifier (Product ID, PID)	59
2.5.4. Group Identifier (GID)	59
2.5.5. Node Identifier	60
2.5.6. IPv6 Addressing	62
2.6. Device identity	63
2.7. Security	64
2.8. Device Commissioning	64
2.9. Intermittently Connected Device (ICD)	65
2.9.1. Sleepy End Device (SED)	66

2.10. Data Model Root	66
2.11. Stack Limits	66
2.11.1. System Model Limits	66
2.11.2. Interaction Model Limits	67
2.12. Migration from Groups to Groupcast (Matter 1.6.x)	68
2.12.1. Backward compatibility for updated nodes	68
2.12.2. Device-managed migration into Groupcast Membership	68
2.12.3. Administrators and controllers Groupcast migration perspective	70
2.12.4. Resource sizing and limits: legacy Groups vs Groupcast	70
2.13. List of Provisional Items	70
2.13.1. Invoke Wildcard Paths	70
2.13.2. WildcardPathFlags options	70
2.13.3. Tag compression encoding for AttributePathIB, EventPathIB, and AttributeDataIB	70
2.13.4. CacheAndSync Feature on the GroupKeyManagement cluster	70
2.13.5. Asynchronous transfer in Bulk Data Exchange Protocol (BDX)	71
2.13.6. Network Recovery	71
2.13.7. Thread Commissioning Protocol	71
2.13.8. Groupcast Cluster	71
2.13.9. Auxiliary ACL	71
2.13.10. IANA Assigned IPv6 Multicast Address	71
3. Cryptographic Primitives	73
3.1. Deterministic Random Bit Generator (DRBG)	73
3.2. True Random Number Generator (TRNG)	74
3.3. Hash function (Hash)	74
3.4. Keyed-Hash Message Authentication Code (HMAC)	75
3.5. Public Key Cryptography	75
3.5.1. Group	75
3.5.2. Key generation	76
3.5.3. Signature and verification	76
3.5.4. ECDH	78
3.5.5. Certificate validation	78
3.5.6. Time and date considerations for certificate path validation	79
3.6. Data Confidentiality and Integrity	80
3.6.1. Generate and encrypt	81
3.6.2. Decrypt and verify	82
3.7. Message privacy	83
3.7.1. Privacy encryption	83
3.7.2. Privacy decryption	84
3.8. Key Derivation Function (KDF)	84
3.9. Password-Based Key Derivation Function (PBKDF)	86
3.10. Password-Authenticated Key Exchange (PAKE)	87

3.10.1. Computation of pA	89
3.10.2. Computation of pB	89
3.10.3. Computation of transcript TT	90
3.10.4. Computation of cA, cB and Ke	90
4. Secure Channel	93
4.1. General Description	93
4.1.1. Messages	93
4.2. IPv6 Reachability	95
4.2.1. Stub Router Behavior	95
4.2.2. Matter Node Behavior	95
4.3. Discovery	96
4.3.1. Commissionable Node Discovery	98
4.3.2. Operational Discovery	115
4.3.3. Commissioner Discovery	121
4.3.4. Common TXT Key/Value Pairs	124
4.4. Message Frame Format	126
4.4.1. Message Header Field Descriptions	127
4.4.2. Message Footer Field Descriptions	130
4.4.3. Protocol Header Field Descriptions	130
4.4.4. Message Size Requirements	133
4.5. Message Framing Over Stream-Oriented Transports	133
4.5.1. Message Length (16/32 bits)	133
4.6. Message Counters	133
4.6.1. Message Counter Types	134
4.6.2. Secure Session Message Counters	135
4.6.3. Check-In Counter	136
4.6.4. Message Counters as Encryption Nonces	137
4.6.5. Replay Prevention and Duplicate Message Detection	137
4.6.6. Counter Processing of Outgoing Messages	140
4.6.7. Counter Processing of Incoming Messages	140
4.7. Message Processing	141
4.7.1. Message Transmission	141
4.7.2. Message Reception	142
4.8. Message Security	142
4.8.1. Data confidentiality and integrity with data origin authentication parameters	143
4.8.2. Security Processing of Outgoing Messages	143
4.8.3. Security Processing of Incoming Messages	146
4.9. Message Privacy	147
4.9.1. Privacy Key	147
4.9.2. Privacy Nonce	147
4.9.3. Privacy Processing of Outgoing Messages	148

4.9.4. Privacy Processing of Incoming Messages	149
4.10. Message Exchanges	149
4.10.1. Exchange Role	149
4.10.2. Exchange ID	150
4.10.3. Exchange Context	150
4.10.4. Exchange Message Dispatch	150
4.10.5. Exchange Message Processing	151
4.11. Secure Channel Protocol	153
4.11.1. Secure Channel Protocol Messages	153
4.11.2. Parameters and Constants	156
4.12. Message Reliability Protocol (MRP)	157
4.12.1. Reliable Messaging Header Fields	157
4.12.2. Reliable transfer	157
4.12.3. Peer Exchange Management	160
4.12.4. Transport Considerations	160
4.12.5. Reliable Message Processing	160
4.12.6. Reliable Message State	165
4.12.7. MRP Messages	165
4.12.8. Parameters and Constants	166
4.13. Unicast Communication	166
4.13.1. Session Parameters	167
4.13.2. Session Establishment Phase	169
4.13.3. Application Data Phase	171
4.14. Session Establishment	172
4.14.1. Passcode-Authenticated Session Establishment (PASE)	172
4.14.2. Certificate Authenticated Session Establishment (CASE)	179
4.15. Secure Communications over TCP	202
4.15.1. Secure Session Establishment	202
4.15.2. TCP Connection Management	203
4.16. Group Communication	204
4.16.1. Groupcast Session Context	204
4.16.2. Sending a group message	205
4.16.3. Receiving a group message	206
4.17. Group Key Management	206
4.17.1. Operational Groups	206
4.17.2. Operational Group Key Derivation	207
4.17.3. Epoch Keys	209
4.17.4. Distribution of Key Material	213
4.18. Message Counter Synchronization Protocol (MCSP)	214
4.18.1. Message Counter Synchronization Methods	215
4.18.2. Group Peer State	216

4.18.3. MCSP Messages	216
4.18.4. Unsynchronized Message Processing	217
4.18.5. Message Counter Synchronization Exchange	218
4.18.6. Message Counter Synchronization Session Context	221
4.18.7. Sequence Diagram	221
4.19. Bluetooth Transport Protocol (BTP)	223
4.19.1. BTP Session Interface	224
4.19.2. BTP Frame Format	224
4.19.3. BTP Control Frames	225
4.19.4. BTP GATT Service	227
4.19.5. Parameters and Constants	238
4.19.6. Bluetooth Requirements	239
4.20. Public Action Frame Transport Protocol (PAFTP)	239
4.20.1. PAFTP Session Interface	240
4.20.2. PAFTP Frame Format	240
4.20.3. PAFTP Control Frames	242
4.20.4. Parameters and Constants	253
4.21. NFC Transport Layer (NTL)	254
4.21.1. NFC Forum requirements	255
4.21.2. NFC-A technology	255
4.21.3. ISO-DEP	255
4.21.4. APDU layer	257
4.22. Check-In Protocol	261
4.22.1. Requirements	262
4.22.2. Message Content	262
4.22.3. Algorithms	263
4.22.4. Protocol Operation	265
4.23. Thread Commissioning Protocol	268
4.23.1. Becoming a Thread Commissioner	269
4.23.2. Network Discovery and Selection	270
4.23.3. Matter Discovery	271
4.23.4. Message framing and secure session establishment	272
5. Commissioning	275
5.1. Onboarding Payload	275
5.1.1. Onboarding Payload Contents	275
5.1.2. Onboarding Material Representation	276
5.1.3. QR Code	277
5.1.4. Manual Pairing Code	282
5.1.5. TLV Content	285
5.1.6. Concatenation	288
5.1.7. Generation of the Passcode	289

5.1.8. NFC Tag	290
5.2. Initiating Commissioning	291
5.2.1. Purpose and Scope	291
5.2.2. User Journey Details	291
5.3. User Directed Commissioning	297
5.3.1. Overview	297
5.3.2. UDC Protocol Messages	299
5.3.3. Message format	299
5.3.4. Message Exchanges	300
5.3.5. IdentificationDeclaration Message	300
5.3.6. CommissionerDeclaration Message	306
5.3.7. Example Message Exchanges	309
5.4. Device Discovery	312
5.4.1. Purpose and Scope	312
5.4.2. Announcement by Device	314
5.4.3. Discovery by Commissioner	326
5.5. Commissioning Flows	329
5.5.1. Commissioning Flows Error Handling	336
5.5.2. Commissioning Flow Diagram	337
5.6. Administrator Assisted Commissioning Flows	338
5.6.1. Introduction	338
5.6.2. Basic Commissioning Method (BCM)	338
5.6.3. Enhanced Commissioning Method (ECM)	339
5.6.4. Open Commissioning Window	340
5.7. Device Commissioning Flows	341
5.7.1. Standard Commissioning Flow	341
5.7.2. User-Intent Commissioning Flow	342
5.7.3. Custom Commissioning Flow	344
5.7.4. Enhanced Setup Flow (ESF)	351
5.7.5. Commissioning Fallback Mechanism	360
5.7.6. Onboarding Payload Inclusion (Manual Pairing Code, QR Code, NFC Tag)	362
5.8. In-field Upgrade to Matter	364
5.9. Network Recovery	364
5.9.1. Network Recovery channels	365
5.9.2. Implementer Considerations	365
5.9.3. Network Recovery Flow	365
6. Device Attestation and Operational Credentials	369
6.1. Certificate Common Conventions	369
6.1.1. Encoding of Matter-specific RDNs	369
6.1.2. Key Identifier Extension Constraints	371
6.1.3. Certificate Sizes	371

6.1.4. Presentation of example certificates	371
6.2. Device Attestation	372
6.2.1. Introduction	372
6.2.2. Device Attestation Certificate (DAC)	372
6.2.3. Device Attestation Procedure	387
6.2.4. Device attestation revocation	390
6.3. Certification Declaration	395
6.3.1. Certification Declaration (CD) Format	395
6.3.2. Firmware Information	399
6.3.3. Firmware information validation examples	401
6.4. Node Operational Credentials Specification	403
6.4.1. Introduction	403
6.4.2. Node Operational Credentials Management	403
6.4.3. Node Operational Identifier Composition	404
6.4.4. Node Operational Key Pair	404
6.4.5. Node Operational Credentials Certificates	404
6.4.6. Node Operational Credentials Procedure	407
6.4.7. Node Operational Certificate Signing Request (NOCSR)	408
6.4.8. Node Operational Certificate Renewal	409
6.4.9. Node Operational Certificate Revocation	409
6.4.10. Fabric Table Vendor ID Verification Procedure	409
6.4.11. Security Considerations	416
6.5. Operational Certificate Encoding	416
6.5.1. Introduction	416
6.5.2. Matter certificate	417
6.5.3. Version Number	418
6.5.4. Serial Number	418
6.5.5. Signature Algorithm	418
6.5.6. Issuer and Subject	419
6.5.7. Validity	425
6.5.8. Public Key Algorithm	425
6.5.9. EC Curve Identifier	425
6.5.10. Public Key	426
6.5.11. Extensions	426
6.5.12. Matter certificate Extensions Encoding Rules	430
6.5.13. Signature	431
6.5.14. Invalid Matter certificates	431
6.5.15. Examples	432
6.6. Access Control	440
6.6.1. Scope and Purpose	440
6.6.2. Access Control Model	440

6.6.3. Access Control List Examples	446
6.6.4. Access Control Cluster update side-effects	451
6.6.5. Access Control Cluster handling of Access Restrictions	452
6.6.6. Conceptual Access Control Privilege Granting Algorithm	456
6.6.7. Applying Privileges to Action Paths	464
7. Data Model Specification	467
7.1. Practical Information	467
7.1.1. Revision History	467
7.1.2. Scope & Purpose	467
7.1.3. Origin Story	467
7.1.4. Overview	468
7.1.5. Glossary	468
7.1.6. Conventions	468
7.2. Data Qualities	468
7.2.1. Common Data Table Columns	468
7.2.2. Description Section	469
7.2.3. Other Data Table Columns	469
7.3. Conformance	469
7.3.1. Operands in Conformance	471
7.3.2. Feature Code in Conformance	471
7.3.3. Element in Conformance	471
7.3.4. Optional Conformance	471
7.3.5. Provisional Conformance	472
7.3.6. Mandatory Conformance	472
7.3.7. Disallowed Conformance	472
7.3.8. Deprecated Conformance	472
7.3.9. Value Conformance	472
7.3.10. Exclusivity Conformance	474
7.3.11. Otherwise Conformance	474
7.3.12. Quality Conformance	474
7.3.13. Expressions and Optionality	475
7.3.14. Choice Conformance	476
7.3.15. Blank Conformance	479
7.3.16. Feature Conformance	479
7.4. Element	480
7.4.1. Encoded Element Processing	480
7.5. Fabric	481
7.5.1. Accessing Fabric	481
7.5.2. Fabric-Index	481
7.5.3. Fabric-Scoped Data	482
7.5.4. Fabric-Scoped IDs	482

7.6. Access	483
7.6.1. Read Access	484
7.6.2. Write Access	484
7.6.3. Invoke Access	484
7.6.4. Fabric-Scoped Access	485
7.6.5. Fabric-Sensitive Access	485
7.6.6. View Privilege	486
7.6.7. Operate Privilege	486
7.6.8. Manage Privilege	486
7.6.9. Administer Privilege	486
7.6.10. Timed Interaction	486
7.7. Other Qualities	487
7.7.1. Changes Omitted Quality	488
7.7.2. Fixed Quality	488
7.7.3. Singleton Quality	489
7.7.4. Diagnostics Quality	489
7.7.5. Large Message Quality	489
7.7.6. Non-Volatile Quality	489
7.7.7. Reportable Quality	489
7.7.8. Quieter Reporting Quality	489
7.7.9. Scene Quality	490
7.7.10. Nullable Quality	490
7.7.11. Atomic Quality	490
7.8. Node	490
7.9. Endpoint	491
7.10. Cluster	491
7.10.1. Cluster Revision	492
7.10.2. Cluster Optional Features	493
7.10.3. Cluster Data Version	493
7.10.4. New Cluster	493
7.10.5. Cluster Aliasing	494
7.10.6. Cluster Inheritance	494
7.10.7. Status Codes	495
7.10.8. Cluster Classification	496
7.11. Command	497
7.11.1. Command Fields	498
7.12. Attribute	499
7.12.1. Persistence	500
7.13. Global Elements	501
7.13.1. ClusterRevision Attribute	501
7.13.2. FeatureMap Attribute	502

7.13.3. AttributeList Attribute	503
7.13.4. AcceptedCommandList Attribute	503
7.13.5. GeneratedCommandList Attribute	503
7.13.6. FabricIndex Field	503
7.13.7. AtomicRequest Command	503
7.13.8. AtomicResponse Command	504
7.14. Event	504
7.14.1. Event Record	504
7.14.2. Buffering	505
7.14.3. Event Filtering	505
7.14.4. Fabric-Sensitive Event	506
7.15. Atomic Writes	506
7.15.1. Atomic Write Flow	506
7.15.2. Atomic Writer ID	507
7.15.3. Atomic Write State	507
7.15.4. AtomicRequestTypeEnum Type	508
7.15.5. AtomicAttributeStatusStruct Type	508
7.15.6. AtomicRequest Command	508
7.15.7. AtomicResponse Command	512
7.16. Device Type	513
7.16.1. Device Type Revision	514
7.16.2. Device Type Composition	514
7.16.3. Device Type Classification	514
7.16.4. Extra Clusters on an Endpoint	516
7.16.5. Primary Device Type	516
7.17. Non-Standard	516
7.18. Data Field	517
7.18.1. Nullable	517
7.18.2. Optional or Deprecated	518
7.18.3. Constraint & Value	518
7.18.4. Fallback Column	524
7.19. Data Types	525
7.19.1. Base Data Types	526
7.19.2. Derived Data Types	533
7.20. Units	550
7.21. Manufacturer Specific Extensions	551
7.21.1. Manufacturer Extensible Identifiers	551
7.21.2. Manufacturer Extensible Identifier (MEI)	552
7.21.3. Manufacturer Extensions	554
7.21.4. Discoverability	558
8. Interaction Model Specification	561

8.1. Practical Information	561
8.1.1. Revision History	561
8.1.2. Scope & Purpose	561
8.1.3. Origin Story	561
8.1.4. Purpose	562
8.1.5. Glossary	562
8.1.6. Conventions & Conformance	563
8.2. Concepts	563
8.2.1. Path	563
8.2.2. Interaction	567
8.2.3. Transaction	568
8.2.4. Action	569
8.2.5. Common Action Behavior	569
8.3. Status and Interaction	571
8.3.1. Status Response Action	571
8.4. Read Interaction	572
8.4.1. Read Transaction	573
8.4.2. Read Request Action	573
8.4.3. Report Data Action	574
8.5. Subscribe Interaction	578
8.5.1. Subscribe Transaction	580
8.5.2. Subscribe Request Action	580
8.5.3. Subscribe Response Action	582
8.6. Report Transaction	583
8.6.1. Report Transaction Non-Empty	583
8.6.2. Report Transaction Empty	583
8.7. Write Interaction	584
8.7.1. Write Transaction	584
8.7.2. Write Request Action	584
8.7.3. Write Response Action	586
8.7.4. Timed Request Action	588
8.8. Invoke Interaction	589
8.8.1. Invoke Transaction	589
8.8.2. Invoke Request Action	589
8.8.3. Invoke Response Action	593
8.9. Common Action Information Blocks and Paths	595
8.9.1. Path Information	595
8.9.2. Attribute Information Blocks	595
8.9.3. Event Information Blocks and Paths	601
8.9.4. Command Information Blocks and Paths	603
8.9.5. Status Information Blocks and Paths	605

8.10. Status Codes	606
8.10.1. Status Code Table	606
9. System Model Specification	613
9.1. Practical Information	613
9.1.1. Revision History	613
9.1.2. Scope and Purpose	613
9.1.3. Origin Story	613
9.1.4. Overview	613
9.2. Endpoint Composition	614
9.2.1. Endpoint Requirements	615
9.2.2. Leaf Device Types	616
9.2.3. Composed Device Types	616
9.2.4. Component Device Types	617
9.2.5. Compound Devices	617
9.2.6. Asserting Conditions Among Device Types	617
9.2.7. Endpoint Composition Patterns	619
9.2.8. Root Node Endpoint	620
9.2.9. Disambiguation	620
9.2.10. Dynamic Endpoint Allocation	623
9.2.11. Node Configuration Changes	624
9.2.12. Superset Device Types	625
9.3. Interaction Model Relationships	627
9.3.1. Subscription	627
9.4. Binding Relationship	627
9.5. Descriptor Cluster	628
9.5.1. Revision History	628
9.5.2. Classification	629
9.5.3. Cluster ID	629
9.5.4. Features	629
9.5.5. Data Types	629
9.5.6. Attributes	630
9.6. Binding Cluster	631
9.6.1. Binding Mutation	632
9.6.2. Revision History	632
9.6.3. Classification	632
9.6.4. Cluster ID	633
9.6.5. Data Types	633
9.6.6. Attributes	633
9.7. Label Cluster	634
9.7.1. Revision History	634
9.7.2. Classification	634

9.7.3. Cluster ID	635
9.7.4. Data Types	635
9.7.5. Attributes	635
9.8. Fixed Label Cluster	636
9.8.1. Revision History	636
9.8.2. Classification	636
9.8.3. Cluster ID	636
9.8.4. Attributes	636
9.9. User Label Cluster	637
9.9.1. Revision History	637
9.9.2. Classification	637
9.9.3. Cluster ID	637
9.9.4. Attributes	637
9.10. Access Control Cluster	638
9.10.1. Revision History	638
9.10.2. Classification	638
9.10.3. Cluster ID	638
9.10.4. Features	638
9.10.5. Data Types	640
9.10.6. Attributes	650
9.10.7. Error handling	656
9.10.8. Commands	657
9.10.9. Events	658
9.11. Group Relationship	665
9.12. Bridge for non-Matter devices	666
9.12.1. Introduction	666
9.12.2. Exposing functionality and metadata of Bridged Devices	666
9.12.3. Discovery of Bridged Devices	671
9.12.4. Configuration of Bridged Devices	671
9.12.5. New features for Bridged Devices	673
9.12.6. Changes to the set of Bridged Devices	673
9.12.7. Changes to device names and grouping of Bridged Devices	674
9.12.8. Setup flow for a Bridge (plus Bridged Devices)	674
9.12.9. Access Control	674
9.12.10. Software update (OTA)	675
9.12.11. Best practices for Bridge Manufacturers	675
9.12.12. Best practices for Administrators	676
9.13. Bridged Device Basic Information Cluster	676
9.13.1. Revision History	676
9.13.2. Classification	677
9.13.3. Cluster ID	677

9.13.4. Features	677
9.13.5. Attributes	677
9.13.6. Commands	680
9.13.7. Events	681
9.14. Actions Cluster	683
9.14.1. Revision History	683
9.14.2. Classification	684
9.14.3. Cluster ID	684
9.14.4. Data Types	684
9.14.5. Attributes	689
9.14.6. Commands	690
9.14.7. Events	696
9.14.8. Examples	698
9.15. Intermittently Connected Devices Behavior	704
9.15.1. ICD Server Behavior	704
9.15.2. ICD Client Behavior	713
9.16. ICD Management Cluster	721
9.16.1. Revision History	721
9.16.2. Classification	721
9.16.3. Cluster ID	721
9.16.4. Features	721
9.16.5. Data Types	722
9.16.6. Attributes	726
9.16.7. Commands	732
9.17. Ecosystem Information Cluster	737
9.17.1. Revision History	738
9.17.2. Classification	738
9.17.3. Cluster ID	738
9.17.4. Data Types	738
9.17.5. Attributes	741
9.17.6. Examples	742
10. Interaction Model Encoding Specification	745
10.1. Overview	745
10.2. Messages	745
10.2.1. IM Protocol Messages	745
10.2.2. Common Action Information Encoding	745
10.2.3. Chunking	746
10.2.4. Transaction Flows	747
10.3. Data Types	750
10.3.1. Analog - Integer	751
10.3.2. Analog - Floating Point	751

10.3.3. Discrete - Enumeration	751
10.3.4. Discrete - Bitmap	752
10.3.5. Composite - String	752
10.3.6. Composite - Octet String	752
10.3.7. Collection - Struct	752
10.3.8. Collection - List	752
10.3.9. Derived Types	752
10.3.10. Field IDs	752
10.4. Sample Clusters	752
10.4.1. Disco Ball Cluster	752
10.4.2. Super Disco Ball Cluster	762
10.5. Sample Device Types	764
10.5.1. Disco Ball Device Type	764
10.5.2. Super Disco Ball Device Type	766
10.5.3. Disco Spot Device Type	767
10.5.4. Disco Dance System Device Type	768
10.5.5. Weather Station Device Type	769
10.6. Information Blocks	770
10.6.1. Tag Rules	771
10.6.2. AttributePathIB	771
10.6.3. DataVersionFilterIB	774
10.6.4. AttributeDataIB	774
10.6.5. AttributeReportIB	779
10.6.6. EventFilterIB	779
10.6.7. ClusterPathIB	779
10.6.8. EventPathIB	780
10.6.9. EventDataIB	780
10.6.10. EventReportIB	782
10.6.11. CommandPathIB	782
10.6.12. CommandDataIB	783
10.6.13. InvokeResponseIB	784
10.6.14. CommandStatusIB	784
10.6.15. EventStatusIB	784
10.6.16. AttributeStatusIB	785
10.6.17. StatusIB	785
10.7. Message Definitions	785
10.7.1. StatusResponseMessage	785
10.7.2. ReadRequestMessage	785
10.7.3. ReportDataMessage	786
10.7.4. SubscribeRequestMessage	788
10.7.5. SubscribeResponseMessage	789

10.7.6. WriteRequestMessage	789
10.7.7. WriteResponseMessage	790
10.7.8. TimedRequestMessage	790
10.7.9. InvokeRequestMessage	790
10.7.10. InvokeResponseMessage	791
11. Service and Device Management	793
11.1. Basic Information Cluster	793
11.1.1. Revision History	793
11.1.2. Classification	793
11.1.3. Cluster ID	793
11.1.4. Data Types	793
11.1.5. Attributes	797
11.1.6. Events	803
11.2. Group Key Management Cluster	804
11.2.1. Revision History	805
11.2.2. Classification	805
11.2.3. Cluster ID	805
11.2.4. Features	805
11.2.5. Data Types	806
11.2.6. Attributes	810
11.2.7. Commands	812
11.3. Localization Configuration Cluster	815
11.3.1. Revision History	815
11.3.2. Classification	816
11.3.3. Cluster ID	816
11.3.4. Attributes	816
11.4. Time Format Localization Cluster	816
11.4.1. Revision History	817
11.4.2. Classification	817
11.4.3. Cluster ID	817
11.4.4. Features	817
11.4.5. Data Types	817
11.4.6. Attributes	818
11.5. Unit Localization Cluster	819
11.5.1. Revision History	819
11.5.2. Classification	820
11.5.3. Cluster ID	820
11.5.4. Features	820
11.5.5. Data Types	820
11.5.6. Attributes	820
11.6. Power Source Configuration Cluster	821

11.6.1. Revision History	821
11.6.2. Classification	821
11.6.3. Cluster ID	821
11.6.4. Attributes	822
11.7. Power Source Cluster	822
11.7.1. Revision History	822
11.7.2. Classification	822
11.7.3. Cluster ID	823
11.7.4. Features	823
11.7.5. Dependencies	823
11.7.6. Data Types	823
11.7.7. Attributes	833
11.7.8. Events	841
11.7.9. Configuration Examples	843
11.8. Power Topology Cluster	845
11.8.1. Revision History	845
11.8.2. Classification	845
11.8.3. Cluster ID	845
11.8.4. Features	845
11.8.5. Attributes	846
11.9. Network Commissioning Cluster	846
11.9.1. Revision History	846
11.9.2. Classification	847
11.9.3. Cluster ID	847
11.9.4. Features	847
11.9.5. Data Types	847
11.9.6. Attributes	852
11.9.7. Commands	855
11.9.8. Usage of networking configurations	869
11.10. General Commissioning Cluster	871
11.10.1. Revision History	871
11.10.2. Classification	871
11.10.3. Cluster ID	872
11.10.4. Features	872
11.10.5. Data Types	872
11.10.6. Attributes	875
11.10.7. Commands	879
11.11. Diagnostic Logs Cluster	887
11.11.1. Revision History	887
11.11.2. Classification	887
11.11.3. Cluster ID	888

11.11.4. Data Types	888
11.11.5. Commands	890
11.12. General Diagnostics Cluster	892
11.12.1. Revision History	892
11.12.2. Classification	892
11.12.3. Cluster ID	892
11.12.4. Features	893
11.12.5. Data Types	893
11.12.6. Attributes	899
11.12.7. Commands	902
11.12.8. Events	905
11.13. Software Diagnostics Cluster	908
11.13.1. Revision History	908
11.13.2. Classification	908
11.13.3. Cluster ID	908
11.13.4. Features	908
11.13.5. Data Types	908
11.13.6. Attributes	909
11.13.7. Commands	910
11.13.8. Events	911
11.14. Thread Network Diagnostics Cluster	911
11.14.1. Revision History	912
11.14.2. Classification	912
11.14.3. Cluster ID	912
11.14.4. Features	912
11.14.5. Data Types	913
11.14.6. Attributes	920
11.14.7. Commands	932
11.14.8. Events	932
11.15. Wi-Fi Network Diagnostics Cluster	933
11.15.1. Revision History	933
11.15.2. Classification	933
11.15.3. Cluster ID	934
11.15.4. Features	934
11.15.5. Data Types	934
11.15.6. Attributes	936
11.15.7. Commands	938
11.15.8. Events	939
11.16. Ethernet Network Diagnostics Cluster	940
11.16.1. Revision History	940
11.16.2. Classification	940

11.16.3. Cluster ID	940
11.16.4. Features	941
11.16.5. Data Types	941
11.16.6. Attributes	941
11.16.7. Commands	943
11.17. Time Synchronization Cluster	943
11.17.1. Revision History	944
11.17.2. Classification	944
11.17.3. Cluster ID	944
11.17.4. Terminology	944
11.17.5. Features	945
11.17.6. Data Types	945
11.17.7. Status Codes	950
11.17.8. Attributes	951
11.17.9. Commands	954
11.17.10. Events	958
11.17.11. Time Synchronization at Commissioning	959
11.17.12. Time Synchronization during operation	959
11.17.13. Time source prioritization	960
11.17.14. Time synchronization maintenance	960
11.17.15. Acting as an NTP Server	961
11.17.16. Implementation Guidance	961
11.18. Operational Credentials Cluster	963
11.18.1. Revision History	963
11.18.2. Classification	964
11.18.3. Cluster ID	964
11.18.4. Data Types	964
11.18.5. Attributes	971
11.18.6. Commands	973
11.19. Administrator Commissioning Cluster	990
11.19.1. Revision History	991
11.19.2. Classification	991
11.19.3. Cluster ID	991
11.19.4. Features	991
11.19.5. Data Types	991
11.19.6. Status Codes	992
11.19.7. Attributes	992
11.19.8. Commands	993
11.20. Over-the-Air (OTA) Software Update	997
11.20.1. Scope & Purpose	997
11.20.2. Functional overview	998

11.20.3. Software update workflow	999
11.20.4. Security considerations	1017
11.20.5. Some special situations	1018
11.20.6. OTA Software Update Provider Cluster	1019
11.20.7. OTA Software Update Requestor Cluster	1030
11.21. Over-the-Air (OTA) Software Update File Format	1040
11.21.1. Scope & Purpose	1040
11.21.2. General Structure	1040
11.21.3. Security considerations	1044
11.22. Bulk Data Exchange Protocol (BDX)	1044
11.22.1. Overview	1044
11.22.2. Terminology	1044
11.22.3. Protocol Opcodes and Status Report Values	1046
11.22.4. Security and Transport Constraints	1048
11.22.5. Transfer Management Messages	1048
11.22.6. Data Transfer Messages	1058
11.22.7. Synchronous Transfers Message Flows	1062
11.22.8. Asynchronous Transfers Message Flows	1070
11.23. Distributed Compliance Ledger	1072
11.23.1. Scope & Purpose	1072
11.23.2. Schemas	1073
11.23.3. Vendor Schema	1074
11.23.4. Product Attestation Authority and Intermediate Certificate Schema	1076
11.23.5. Operational Trust Anchors Schema	1079
11.23.6. DeviceModel Schema	1081
11.23.7. DeviceSoftwareVersionModel Schema	1088
11.23.8. DeviceSoftwareCompliance / Compliance test result Schema	1091
11.23.9. Device Attestation PKI Revocation Distribution Points Schema	1093
11.23.10. APIs / CLI	1099
11.24. Joint Fabric Datastore Cluster	1099
11.24.1. Revision History	1100
11.24.2. Classification	1100
11.24.3. Cluster ID	1100
11.24.4. Usage Requirements	1100
11.24.5. Data Types	1104
11.24.6. Attributes	1114
11.24.7. Commands	1117
11.25. Joint Fabric Administrator Cluster	1136
11.25.1. Revision History	1136
11.25.2. Classification	1136
11.25.3. Cluster ID	1137

11.25.4. Data Types	1137
11.25.5. Attributes	1138
11.25.6. Commands	1138
11.26. Commissioner Control Cluster	1143
11.26.1. Revision History	1143
11.26.2. Classification	1143
11.26.3. Cluster ID	1143
11.26.4. Data Types	1143
11.26.5. Attributes	1144
11.26.6. Commands	1144
11.26.7. Events	1147
11.27. Groupcast Cluster	1148
11.27.1. Revision History	1148
11.27.2. Classification	1149
11.27.3. Cluster ID	1149
11.27.4. Features	1149
11.27.5. Data Types	1150
11.27.6. Attributes	1152
11.27.7. Commands	1158
11.27.8. Events	1167
11.27.9. Groupcast Key Management	1169
11.27.10. Groupcast Auxiliary ACL Handling	1169
12. Multiple Fabrics	1173
12.1. Introduction	1173
12.2. Joint Fabric	1173
12.2.1. Introduction	1173
12.2.2. Node ID Generation	1173
12.2.3. Anchor ICAC requirements	1174
12.2.4. Joint Fabric ACL Architecture	1174
12.2.5. Joint Commissioning Method (JCM)	1175
12.2.6. Anchor Administrator Selection	1181
12.2.7. Administrator Removal	1183
12.3. User Consent	1184
12.4. Administrator-Assisted Commissioning Method	1184
12.5. Node Behavior	1184
12.6. Fabric Synchronization	1184
12.6.1. Introduction	1184
12.6.2. Terminology	1185
12.6.3. Fabric Synchronization Composition	1185
12.6.4. Preventing Device Duplication	1186
12.6.5. Changes to device and locations of Synchronized Devices	1188

12.6.6. Changes to the set of Synchronized Devices	1188
12.6.7. Fabric Synchronized Relationships	1189
12.6.8. Setup flow for Fabric Synchronization	1189
13. Security Requirements	1195
13.1. Overview	1195
13.2. Device vs. Node	1195
13.3. Commissioning	1195
13.4. Factory Reset	1196
13.5. Firmware	1196
13.6. Security Best Practices	1197
13.6.1. Cryptography	1197
13.6.2. Commissioning and Administration	1197
13.6.3. Firmware	1198
13.6.4. Manufacturing	1198
13.6.5. Resiliency	1198
13.6.6. Battery Powered Devices	1199
13.6.7. Tamper Resistance	1199
13.6.8. Product Security	1199
13.6.9. Bridging	1199
13.6.10. Distributed Compliance Ledger	1199
13.6.11. Safety	1199
13.7. Threats and Countermeasures	1200
14. Transport Layer Security	1221
14.1. Introduction	1221
14.2. Common Conventions	1221
14.2.1. Transport Layer Security (TLS)	1221
14.2.2. TLS Certificates	1221
14.2.3. TLS Endpoint	1222
14.2.4. TLS Connection	1222
14.3. Description	1223
14.3.1. TLS Certificate Management	1223
14.3.2. TLS Client Management	1224
14.3.3. TLS Minimum Requirements	1225
14.4. TLS Certificate Management Cluster	1225
14.4.1. Revision History	1225
14.4.2. Classification	1225
14.4.3. Cluster ID	1225
14.4.4. Data Types	1226
14.4.5. Attributes	1227
14.4.6. Commands	1228
14.5. TLS Client Management Cluster	1241

14.5.1. Revision History	1241
14.5.2. Classification	1241
14.5.3. Cluster ID	1241
14.5.4. Data Types	1241
14.5.5. Status Codes	1242
14.5.6. Attributes	1243
14.5.7. Commands	1243
Appendix A: Tag-length-value (TLV) Encoding Format	1249
A.1. Scope & Purpose	1249
A.2. Tags	1249
A.2.1. Profile-Specific Tags	1249
A.2.2. Context-Specific Tags	1249
A.2.3. Anonymous Tags	1250
A.2.4. Canonical Ordering of Tags	1250
A.3. Lengths	1250
A.4. Primitive Types	1250
A.5. Container Types	1251
A.5.1. Structures	1251
A.5.2. Arrays	1251
A.5.3. Lists	1251
A.6. Element Encoding	1252
A.7. Control Octet Encoding	1252
A.7.1. Element Type Field	1252
A.7.2. Tag Control Field	1254
A.8. Tag Encoding	1254
A.8.1. Fully-Qualified Tag Form	1254
A.8.2. Implicit Profile Tag Form	1255
A.8.3. Common Profile Tag Form	1255
A.8.4. Context-Specific Tag Form	1255
A.8.5. Anonymous Tag Form	1255
A.9. Length Encoding	1255
A.10. End of Container Encoding	1256
A.11. Value Encodings	1256
A.11.1. Integers	1256
A.11.2. UTF-8 and Octet Strings	1256
A.11.3. Booleans	1256
A.11.4. Arrays, Structures and Lists	1257
A.11.5. Floating Point Numbers	1257
A.11.6. Nulls	1257
A.12. TLV Encoding Examples	1257
Appendix B: Tag-length-value (TLV) Schema Definitions	1261

B.1. Introduction	1261
B.1.1. Basic Structure	1261
B.1.2. Keywords	1261
B.1.3. Naming	1261
B.1.4. Namespaces	1262
B.1.5. Qualifiers	1262
B.1.6. Tagging	1263
B.2. Definitions	1263
B.2.1. Type Definition (type-def)	1263
B.2.2. FIELD GROUP Definition (field-group-def)	1264
B.2.3. Namespace Definition (namespace-def)	1266
B.2.4. PROTOCOL Definition (protocol-def)	1267
B.2.5. VENDOR Definition (vendor-def)	1268
B.3. Types	1269
B.3.1. ARRAY / ARRAY OF	1269
B.3.2. BOOLEAN	1271
B.3.3. FLOAT32 / FLOAT64	1271
B.3.4. SIGNED INTEGER / UNSIGNED INTEGER	1272
B.3.5. LIST / LIST OF	1273
B.3.6. OCTET STRING	1274
B.3.7. NULL	1274
B.3.8. STRING	1275
B.3.9. STRUCTURE	1275
B.4. Pseudo-Types	1278
B.4.1. ANY	1278
B.4.2. CHOICE OF	1278
B.5. Qualifiers	1281
B.5.1. any-order / schema-order / tag-order	1281
B.5.2. extensible	1281
B.5.3. id	1282
B.5.4. length	1283
B.5.5. nullable	1283
B.5.6. optional	1284
B.5.7. range	1284
B.5.8. tag	1285
B.5.9. Documentation and Comments	1287
Appendix C: Tag-length-value (TLV) Payload Text Representation Format	1289
C.1. Introduction	1289
C.2. Format Specification	1289
C.2.1. Tag/Value	1289
C.2.2. Context-Specific Tags	1289

C.2.3. Protocol-Specific Tags	1289
C.2.4. Anonymous Tags	1290
C.2.5. Primitive Types	1290
C.2.6. Complex Types: Structure	1291
C.2.7. Complex Types: Arrays	1291
C.2.8. Complex Types: List	1291
C.3. Examples	1291
C.3.1. TLV Schema	1291
C.3.2. TLV Payloads	1292
Appendix D: Status Report Messages	1295
D.1. Overview	1295
D.2. Status Report elements	1295
D.3. Message Format	1295
D.3.1. General status codes (GeneralCode)	1296
D.3.2. Protocol-specific codes (ProtocolId and ProtocolCode)	1296
D.3.3. Protocol-specific data (ProtocolData)	1297
D.4. Presenting StatusReport messages in protocol specifications	1297
Appendix E: Matter-Specific ASN.1 Object Identifiers (OIDs)	1299
Appendix F: Cryptographic test vectors for some procedures	1301
F.1. Certification Declaration CMS test vector	1301
F.2. Device Attestation Response test vector	1304
F.3. Node Operational CSR Response test vector	1308
F.4. Check-In Protocol test vectors	1312
F.5. Fabric Table Vendor ID Verification Procedure Test Vector	1315
Appendix G: Minimal Resource Requirements	1331

Chapter 1. Introduction

The Matter specification defines fundamental requirements to enable an interoperable application layer solution for smart home devices over the Internet Protocol.

1.1. Scope and Purpose

This specification details everything necessary to implement an application and transport layer stack. It is intended to be used by implementers as a complete specification but where necessary other references are noted with details on how these references apply to this specification.

In case of discrepancies between this specification and the [SDK](https://github.com/project-chip/connectedhomeip/) [https://github.com/project-chip/connectedhomeip/], this specification SHALL take precedence.

1.2. Acronyms and Abbreviations

Acronym	Definition
ACL	Access Control List
AGID	Application Group Identifier
AEAD	Authenticated Encryption with Associated Data
AES	Advanced Encryption Standard (from FIPS 197)
AP	Access Point (from IEEE 802.11-2020)
API	Application Programming Interface
ASN.1	Abstract Syntax Notation 1 (from ITU ASN.1)
BLE	Bluetooth Low Energy
BDX	Bulk Data Exchange
BSS	Basic Service Set (from IEEE 802.11-2020)
BSSID	Basic Service Set Identifier (from IEEE 802.11-2020)
BTP	Bluetooth Transport Protocol
CA	Certificate Authority (also known as Certification Authority)
CASE	Certificate Authenticated Session Establishment
CAT	CASE Authenticated Tag
CBC-MAC	Cipher Block Chaining Message Authentication Code
CCM	Counter mode of encryption with CBC-MAC (AEAD mode) (from NIST 800-38C)
CD	Certification Declaration
CMS	Cryptographic Message Syntax
CN	Common Name (from X.520)

Acronym	Definition
CSR	Certificate Signing Request
CTR	Counter Mode (AES block cipher mode) (from NIST 800-38A)
DAC	Device Attestation Certificate
DER	Distinguished Encoding Rule (from X.690)
DN	Distinguished Name (from X.520)
DNS	Domain Name System
DNS-SD	DNS Based Service Discovery (from RFC 6763)
DRBG	Deterministic Random Bit Generator (from NIST 800-90A)
ECC	Elliptic Curve Cryptography (from SEC 1) (also "Error Correction Code")
ECDHE	Elliptic Curve Ephemeral Diffie-Hellman (from SEC 1)
ECDSA	Elliptic Curve Digital Signature Algorithm (from SEC 1)
ESS	Extended Service Set (from IEEE 802.11-2020)
EUI	Extended Unique Identifier
EUI-64	64-bit EUI
GATT	Bluetooth Generic Attribute Profile
GID	Group Identifier (also referred to as "Group ID")
GKH	Group Key Hash
GUA	Global Unicast Address
HMAC	Keyed-Hash Message Authentication Code (from FIPS 198-1)
ICD	Intermittently Connected Device
ID	Identifier
IP	Internet Protocol
IPK	Identity Protection Key (a Universal Group key shared across a Fabric)
KDF	Key Derivation Function (from RFC 5869)
KDM	Key Derivation Method (from RFC 5869)
LLA	Link local address
LLN	Low power and Lossy Network
MAC	Medium Access Control (or "Message Authentication Code")
MCSP	Message Counter Synchronization Protocol
MIC	Message Integrity Code (used as synonym for MAC (Message Authentication Code) to avoid confusion with MAC (Medium Access Control) as used in network addressing contexts)

Acronym	Definition
MRP	Message Reliability Protocol
NFC	Near Field Communication
NOC	Node Operational Certificate
NOCSR	Node Operational Certificate Signing Request
NTL	NFC Transport Layer
OID	Object Identifier (from ITU ASN.1)
OTA	Over-the-air (used mostly in context of "Over-the-air Software Update")
PAA	Product Attestation Authority
PAFTP	Public Action Frame Transport Protocol
PAI	Product Attestation Intermediate
PAKE	Password-Authenticated Key Exchange (from SPAKE2+)
PASE	Passcode-Authenticated Session Establishment
PBKDF	Password-Based Key Derivation Function (from NIST 800-132)
PDU	Protocol Data Unit
PID	Product Identifier (also Product ID)
PIN	Personal Identification Number
PKI	Public Key Infrastructure
PSK	Pre-Shared Key
QR code	Quick Response (code)
SDF	Service Discovery Frame
SDU	Service Data Unit
SED	Sleepy End Device
SHA	Secure Hash Algorithm (from FIPS 180-4)
SNAC	Stub Network Auto-Configuring (from SNAC)
SRP	Service Registration Protocol (from RFC 9665)
STA	(Wireless) Station (from IEEE 802.11-2020)
TCP	Transmission Control Protocol (from RFC 9293)
TFTP	Trivial File Transfer Protocol (from RFC 1350)
TLV	Tag Length Value (refers mostly to Tag-length-value (TLV) Encoding Format)
TRNG	True Random Number Generator (from NIST 800-90B)
TU	Transmission Unit
UDP	User Datagram Protocol

Acronym	Definition
UGID	Universal Group Identifier
ULA	Unique local address
UTC	Universal Time Coordinated
UUID	Universally Unique Identifier
VID	Vendor Identifier (also Vendor ID)
WPA	Wi-Fi Protected Access
ZCL	Zigbee Cluster Library

1.3. Definitions

Term	Definition
Access Control List	A list of entries in the Access Control Cluster expressing individual rules which grant privileges to access cluster elements.
Administrator	A Node having Administer privilege over at least the Access Control Cluster of another Node.
Advertising Data	A data container used in BLE Advertisements to convey a logical grouping of information.
Anchor CA	This RCAC belongs to Anchor Fabric and is trusted by all devices in the Joint Fabric. It also signs ICA belonging to other fabrics that participate in Joint Fabric. Identical with the TRCA of the Joint Fabric.
Anchor Fabric	This is the fabric that “anchors” Joint Fabric because its operational root CA is trusted by all other fabrics in the Joint Fabric.
Anchor Administrator	Administrator that has the ability to interact with the Anchor CA.
Attribute	A data entity which represents a physical quantity or state. This data is communicated to other Nodes using commands.
Binding	A persistent attachment between an instance on one Node to one-or-more corresponding instances on another (or the same) Node.
Border Router	A router, also known as Edge Router, that provides routing services between two IP subnets (typically, between a hub network and a peripheral network).
Bridge	A Node that represents one or more non-Matter devices on the Fabric.
Bridged Device	A non-Matter device that is represented on the Fabric by a Bridge so it can be used by Nodes on the Fabric.
Broadcast	The transmission of a message to every Node in a particular broadcast domain, be it all Nodes on a Ethernet or Wi-Fi link, and/or all Nodes on a Thread mesh.
Certificate Authority (CA)	An entity that issues digital certificates such as a DAC or NOC

Term	Definition
Certification Declaration	A digitally signed token that conveys Matter certification status of a vendor's certified Device.
Client	A Cluster interface that typically sends commands that manipulate the attributes on the corresponding server cluster. A client cluster communicates with a corresponding remote server cluster with the same cluster identifier.
Cluster	A specification defining one or more attributes, commands, behaviors and dependencies, that supports an independent utility or application function. The term may also be used for an implementation or instance of such a specification on an endpoint.
Command	Requests for action on a value with an expected response which may have parameters and a response with a status and parameters.
Commission	To bring a Node into a Fabric.
Commissionable Node	A Node that is able to be commissioned. Specific actions such as a button press may be required to put a Commissionable Node into Commissioning Mode in order for it to allow Commissioning.
Commissionable Node Discovery	Discovery of a Node that is able to be Commissioned, but not necessarily in Commissioning Mode, for the purpose of performing Commissioning. The Node may be brand new, after factory reset, or it may have already been Commissioned.
Commissioner	A Role of a Node that performs Commissioning.
Commissioner Discovery	Discovery of a Commissioner.
Commissionee	An entity that is being Commissioned to become a Node.
Commissioning	Sequence of operations to bring a Node into a Fabric by assigning an Operational Node ID and Node Operational credentials.
Commissioning Channel	A Secure Channel used to perform Commissioning.
Commissioning Mode	The mode of a Node in which it allows Commissioning.
Controller	A Role of a Node that has permissions to enable it to control one or more Nodes.
Controlee	A Role of a Node that has permissions defined to enable it to be controlled by one or more Nodes.
Device	A piece of equipment containing one or more Nodes.
Device Attestation Certificate	An RFC 5280 [https://www.rfc-editor.org/rfc/rfc5280] compliant X.509 v3 document with attestable attributes.
Discriminator	A 12-bit value used to discern between multiple commissionable Matter device advertisements. See Discriminator value .

Term	Definition
Ecosystem	A collaborative set of functional entities (proximal- or cloud-based) that provide user-facing configuration and information tools and services, which is provided by one or more cooperating companies. An “ecosystem” to the user is the vendor of the “functional entities” which include, but are not limited to, a smartphone OS, tablet, smartphone app, tablet app, computer app, voice assistant, TV interface, wall display/switch, cloud service, or an integrated service across hardware and software platforms.
Endpoint	A particular component within a Node that is individually addressable.
Endpoint Address	The address assigned to an Endpoint.
Fabric	A logical collection of communicating Nodes, sharing a common root of trust, and a common distributed configuration state.
Information Element	A Wi-Fi (IEEE 802.11-2020) data container used to convey various information regarding a particular Wi-Fi network’s capabilities and operation.
Joint Fabric	A single fabric with a single root of trust that is administered by one or more Ecosystems which each provide a Node implementing the Joint Fabric Administrator device type.
Key Center	A system component which takes the NOCSR from a Commissioner and allocates an Operational Node ID that is unique to the Fabric, inserts this Operational Node ID as the DN into the NOC, and signs the NOC.
Manual Pairing Code	An 11-digit or 21-digit numeric code that can be manually entered/spoken instead of scanning a QR code, which contains the information needed to commission a Matter device.
Network	A set of nodes that have addressability, connectivity, and reachability to one another via Internet Protocol.
NFC-based Commissioning	Commissioning using NFC as a transport layer (NTL) before joining the operational network.
NFC Onboarding Payload	Onboarding payload stored in an NFC Tag .
Node	An addressable entity which supports the Matter protocol stack and (once Commissioned) has its own Operational Node ID and Node Operational credentials.
Operational Discovery	Discovery of a previously commissioned Node for the purpose of performing operations with that Node.
Onboarding Payload	The information needed to start the process of commissioning a Device.
OTA Provider	A Node implementing the OTA Software Update Provider role (see OTA Software Update Provider Cluster).
OTA Requestor	A Node implementing the OTA Software Update Requestor role (see OTA Software Update Requestor Cluster).

Term	Definition
Product Attestation Authority	An entity which operates a root level Certificate Authority for the purpose of Device Attestation.
Product Attestation Intermediate	An entity which operates an intermediate level Certificate Authority for the purpose of Device Attestation.
Product ID (PID)	A 16-bit number that identifies the type of a Device, uniquely among the product types made by a given vendor. See Product ID .
QR Code	A machine-readable optical label that contains information about the item to which it is attached (see QR Code).
Rendezvous	The process for commissioner and commissionee to establish an initial communication channel.
Role	Some set of (related) behaviors of a Node. Each Node can have multiple roles.
Router	A device that provides routing services in its network in cooperation with other Routers.
Secure Channel	A channel in which messages are encrypted and authenticated. Unicast secure channels also provide authentication of each peer.
Server	A Cluster interface that typically supports all or most of the attributes of the Cluster. A Server Cluster communicates with a corresponding remote Client Cluster with the same Cluster identifier.
Service Discovery	The ability of a Node to locate services of interest.
Setup Code	The low-entropy passcode used to secure commissioning.
Software Image	A data blob, equivalent to a file, utilized by a Node to update its software. For the purposes of OTA Software Update, this further refers to files conforming to the OTA Software Image File Format .
Thread	A low-power IEEE 802.15.4-based IPv6 mesh networking technology (see Thread specification).
Vendor	A party who makes or provides one of the functional entities that implement the Matter standard.
Vendor ID (VID)	A 16-bit number that uniquely identifies the Vendor of the Device. See Vendor ID .

1.4. Standards Terminology Mapping

Matter	HomeKit	Weave	Thread	Zigbee
Administrator	Admin	Fabric provisioner	Commissioner	Coordinator
Attribute	Characteristics	Property		Attribute
Binding	Event subscription	Subscription	Link	Binding
Broadcast			Broadcast	Broadcast
Client		Service client	Client	Client

Matter	HomeKit	Weave	Thread	Zigbee
Cluster	Services	interface		Cluster
Cluster	Trait	Service		Cluster
Command	Command	Command	Command	Command
Commissioning	Pairing	Pairing	Commissioning	Association
Commissioner	Admin	Fabric provisioner	Commissioner	Coordinator
Device	Accessory	Device	Device	Device
End Device			End Device	End Device
Endpoint	Profile	Resource	Interface	Endpoint
Endpoint Address	Device ID	Resource ID	Endpoint Identifier	Endpoint address
Fabric	Network	Fabric	Partition	Network
Network Manager	Device / Controller	Nest Service	Leader	Network manager
Node	Accessory	Node	Node	Node
Router			Router	Router
Server		Service host	Server	Server
Service Discovery		Service directory		Service Discovery

1.5. Conformance Levels

The key words below are usually capitalized in the document to make the requirement clear.

Key Word	Description
MAY	A key word that indicates flexibility of choice with no implied preference.
NOT	A key word that used to describe that the requirement is the inverse of the behavior specified (i.e. SHALL NOT, MAY NOT, etc)
SHALL	A key word indicating a mandatory requirement. Designers are required to implement all such mandatory requirements.
SHOULD	A key word indicating flexibility of choice with a strongly preferred alternative. Equivalent to the phrase is recommended.

1.6. References

The following standards and specifications contain provisions, which through reference in this document constitute provisions of this specification. All the standards and specifications listed are normative references. At the time of publication, the editions indicated were valid. All standards and specifications are subject to revision, and parties to agreements based on this specification are encouraged to investigate the possibility of applying the most recent editions of the standards and specifications indicated below.

1.6.1. Connectivity Standards Alliance Reference Documents

Reference	Reference Location/URL	Description
[Alliance-PNP]	https://groups.csa-iot.org/wg/members/document/21624	Organizational Processes and Procedures, 13-0625, revision 8, November 2021
[AppClusters]	https://groups.csa-iot.org/wg/members-all/document/27350	Application Clusters
[CSA-05-03874]	https://groups.csa-iot.org/wg/members-all/document/10905	Connectivity Standards Alliance Manufacturer Code Database
[DeviceLibrary]	https://groups.csa-iot.org/wg/members-all/document/27351	Device Library
[Matter Brand Guidelines]	https://groups.csa-iot.org/wg/members-all/document/24657	Matter Brand Guidelines
[Standard-Namespaces]	https://groups.csa-iot.org/wg/members-all/document/31936	Standard Namespaces

1.6.2. External Reference Documents

Reference	Reference Location/URL	Description
[AdProx]	https://datatracker.ietf.org/doc/draft-ietf-dnssd-advertising-proxy/04/	Advertising Proxy for DNS-SD Service Registration Protocol
[ANSI C18]	https://ansi.org	ANSI C18 Standards on Portable Cells and Batteries
[BCP47]	https://tools.ietf.org/rfc/bcp/bcp47.txt	Best Current Practice 47
[Bluetooth@CoreSpec]	https://www.bluetooth.com/specifications/specs/core-specification-amended-4-2/	Bluetooth® Core Specification 4.2 (amended)
[Bluetooth@CSS]	https://www.bluetooth.com/specifications/specs/core-specification-supplement/	Bluetooth® Core Specification Supplement 12
[draft-lemon-stub-networks]	https://datatracker.ietf.org/doc/html/draft-lemon-stub-networks-02	Connecting Stub Networks to Existing Infrastructure

Reference	Reference Location/URL	Description
[FIPS 180-4]	https://csrc.nist.gov/publications/detail/fips/180/4/final	NIST FIPS 180-4 Secure Hash Standard (SHS), August 2015
[FIPS 186-5]	https://csrc.nist.gov/publications/detail/fips/186/5/final	NIST FIPS 186-5 Digital Signature Standard (DSS), February 2023
[FIPS 197]	https://doi.org/10.6028/NIST.FIPS.197	NIST FIPS 197 Advanced Encryption Standard (AES), November 2001
[FIPS 198-1]	https://csrc.nist.gov/publications/detail/fips/198/1/final	NIST FIPS 198-1 The Keyed-Hash Message Authentication Code (HMAC), July 2008
[HDMI]	https://hdmiforum.org/hdmi-forum-releases-version-2-1-hdmi-specification/	HDMI CEC specification
[IANA Time Zone Database]	https://www.iana.org/time-zones	IANA Time Zone Database
[IANA IPv6 Multicast Address Space]	https://www.iana.org/assignments/multicast-addresses/multicast-addresses.xhtml	IANA IPv6 Multicast Address Space
[IEC 60086]	https://www.iec.ch	IEC 60086 standard for Primary Batteries
[IEEE 754-2019]	https://ieeexplore.ieee.org/document/8766229	"IEEE Standard for Floating-Point Arithmetic," in IEEE Std 754-2019 (Revision of IEEE 754-2008) July 2019, doi: 10.1109/IEEESTD.2019.8766229.
[IEEE 802.11-2020]	https://standards.ieee.org/standard/802_11-2020.html	IEEE 802.11-2020 - IEEE Standard for Information Technology - Telecommunications and Information Exchange between Systems - Local and Metropolitan Area Networks - Specific Requirements - Part 11: Wireless LAN Medium Access Control (MAC) and Physical Layer (PHY) Specifications
[IEEE 1588-2008]	https://standards.ieee.org/standard/1588-2008.html	IEEE Standard for a Precision Clock Synchronization Protocol for Networked Measurement and Control Systems

Reference	Reference Location/URL	Description
[IoT Device Security Specification]	https://csa-iot.org/all-solutions/product-security/	IoT Device Security Specification, March 2024
[ISO 639]	https://www.iso.org/iso-639-language-codes.html	Language Codes
[ISO 4217]	https://www.iso.org/iso-4217-currency-codes.html	Currency Codes
[ISO/IEC 7816-4]	https://www.iso.org/obp/ui/en/#iso:std:iso-iec:7816:-4:ed-4:v1:en	ISO/IEC 7816-4:2020(E) Identification cards — Integrated circuit cards — Part 4: Organization, security and commands for interchange
[ISO/IEC 18004:2015]	https://www.iso.org/standard/62021.html	Information technology - Automatic identification and data capture techniques - QR Code bar code symbology specification
[ITU ASN.1]	https://www.itu.int/en/ITU-T/asn1/Pages/asn1_project.aspx	ITU ASN.1 Project
[NFCForum-DEVREQS]	https://nfc-forum.org/uploads/Certification-Files/NFCForum-DevicesRequirements-3.3.00.pdf	NFC Forum Devices Requirements, version 3.3.0
[NFCForum-TS-ACTIVITY]	https://nfc-forum.org/build/specifications	NFC Forum Activity Technical Specification, version 2.3, 2023
[NFCForum-TS-DIGITAL]	https://nfc-forum.org/build/specifications	NFC Forum Digital Protocol Technical Specification, version 2.4
[NFCForum-TS-NDEF 1.0]	https://nfc-forum.org/build/specifications	NFC Forum Data Exchange Format (NDEF) Technical Specification
[NFCForum-TS-RTD 1.0]	https://nfc-forum.org/build/specifications	NFC Forum Record Type Definition (RTD) Technical Specification
[NFCForum-TS-RTD URI 1.0]	https://nfc-forum.org/build/specifications	NFC Forum URI Record Type Definition Technical Specification
[NIST 800-38A]	https://nvlpubs.nist.gov/nistpubs/Legacy/SP/nistspecialpublication800-38a.pdf	NIST SP 800-38A Recommendation for Block Cipher Modes of Operation: Methods and Techniques, December 2001

Reference	Reference Location/URL	Description
[NIST 800-38C]	https://nvlpubs.nist.gov/nistpubs/Legacy/SP/nistspecialpublication800-38c.pdf	NIST SP 800-38C Recommendations for Block Cipher Mode of Operation: The CCM Mode for Authentication and Confidentiality, Morris Dworkin, May 2004 (errata update 2007)
[NIST 800-90A]	https://csrc.nist.gov/publications/detail/sp/800-90a/rev-1/final	NIST SP 800-90A Rev. 1 Recommendation for Random Number Generation Using Deterministic Random Bit Generators
[NIST 800-90B]	https://csrc.nist.gov/publications/detail/sp/800-90b/final	NIST SP 800-90B Recommendation for the Entropy Sources Used for Random Bit Generation
[NIST 800-132]	https://nvlpubs.nist.gov/nistpubs/Legacy/SP/nistspecialpublication800-132.pdf	NIST SP 800-132 Recommendation for Password-Based Key Derivation, Part 1: Storage Applications, December 2010
[NIST 800-186]	https://nvlpubs.nist.gov/nistpubs/SpecialPublications/NIST.SP.800-186-draft.pdf	NIST Draft SP 800-186 Recommendation for Discrete Logarithm-Based Cryptography, October 2019
[RFC 1350]	https://www.rfc-editor.org/rfc/rfc1350	The TFTP Protocol (Revision 2)
[RFC 1738]	https://www.rfc-editor.org/rfc/rfc1738	Uniform Resource Locators (URL)
[RFC 2119]	https://www.rfc-editor.org/rfc/rfc2119	Bradner, S., "Key words for use in RFCs to Indicate Requirement Levels", BCP 14, RFC 2119, DOI 10.17487/RFC2119, March 1997
[RFC 2782]	https://www.rfc-editor.org/rfc/rfc2782	A DNS RR for specifying the location of services (DNS SRV)
[RFC 2986]	https://www.rfc-editor.org/rfc/rfc2986	PKCS #10: Certification Request Syntax Specification Version 1.7
[RFC 3306]	https://www.rfc-editor.org/rfc/rfc3306	Unicast-Prefix-based IPv6 Multicast Addresses
[RFC 3587]	https://www.rfc-editor.org/rfc/rfc3587	IPv6 Global Unicast Address Format
[RFC 3986]	https://www.rfc-editor.org/rfc/rfc3986	Uniform Resource Identifier (URI)

Reference	Reference Location/URL	Description
[RFC 4007]	https://www.rfc-editor.org/rfc/rfc4007	IPv6 Scoped Address Architecture
[RFC 4158]	https://www.rfc-editor.org/rfc/rfc4158	Internet X.509 Public Key Infrastructure: Certification Path Building
[RFC 4191]	https://www.rfc-editor.org/rfc/rfc4191	Default Router Preferences and More-Specific Routes
[RFC 4193]	https://www.rfc-editor.org/rfc/rfc4193	Unique Local IPv6 Unicast Addresses (ULA)
[RFC 4291]	https://www.rfc-editor.org/rfc/rfc4291	IPv6 Addressing Architecture
[RFC 4506]	https://www.rfc-editor.org/rfc/rfc4506	XDR: External Data Representation Standard
[RFC 4648]	https://www.rfc-editor.org/rfc/rfc4648	The Base16, Base32, and Base64 Data Encodings
[RFC 4861]	https://www.rfc-editor.org/rfc/rfc4861	Neighbor Discovery for IP version 6 (IPv6)
[RFC 4862]	https://www.rfc-editor.org/rfc/rfc4862	IPv6 Stateless Address Autoconfiguration
[RFC 5280]	https://www.rfc-editor.org/rfc/rfc5280	Internet X.509 Public Key Infrastructure Certificate and Certificate Revocation List (CRL) Profile
[RFC 5505]	https://www.rfc-editor.org/rfc/rfc5505	Principles of Internet Host Configuration
[RFC 5646]	https://tools.ietf.org/html/rfc5646	Tags for Identifying Languages
[RFC 5652]	https://www.rfc-editor.org/rfc/rfc5652	Cryptographic Message Syntax (CMS)
[RFC 5869]	https://www.rfc-editor.org/rfc/rfc5869	HMAC-based Extract-and-Expand Key Derivation Function (HKDF)
[RFC 5905]	https://www.rfc-editor.org/rfc/rfc5905	Network Time Protocol Version 4
[RFC 5952]	https://www.rfc-editor.org/rfc/rfc5952	A Recommendation for IPv6 Address Text Representation
[RFC 6234]	https://www.rfc-editor.org/rfc/rfc6234	US Secure Hash Algorithms

Reference	Reference Location/URL	Description
[RFC 6282]	https://www.rfc-editor.org/rfc/rfc6282	IPv6 Datagrams on IEEE 802.15.4
[RFC 6335]	https://www.rfc-editor.org/rfc/rfc6335	Service Name and Port Number Procedures
[RFC 6760]	https://www.rfc-editor.org/rfc/rfc6760	Replacement of AppleTalk NBP
[RFC 6762]	https://www.rfc-editor.org/rfc/rfc6762	Multicast DNS
[RFC 6763]	https://www.rfc-editor.org/rfc/rfc6763	DNS-Based Service Discovery
[RFC 6920]	https://www.rfc-editor.org/rfc/rfc6920	Naming Things with Hashes
[RFC 6960]	https://www.rfc-editor.org/rfc/rfc6960	X.509 Internet Public Key Infrastructure Online Certificate Status Protocol - OCSP
[RFC 7230]	https://www.rfc-editor.org/rfc/rfc7230	Hypertext Transfer Protocol (HTTP/1.1): Message Syntax and Routing
[RFC 7346]	https://www.rfc-editor.org/rfc/rfc7346	IPv6 Multicast Address Scopes
[RFC 7468]	https://www.rfc-editor.org/rfc/rfc7468	Textual Encodings of PKIX, PKCS, and CMS Structures
[RFC 7558]	https://www.rfc-editor.org/rfc/rfc7558	Scalable DNS-SD Requirements
[RFC 7731]	https://www.rfc-editor.org/rfc/rfc7731	Multicast Protocol for Low-Power and Lossy Networks (MPL)
[RFC 8305]	https://www.rfc-editor.org/rfc/rfc8305	Happy Eyeballs Version 2: Better Connectivity Using Concurrency
[RFC 8446]	https://www.rfc-editor.org/rfc/rfc8446	The Transport Layer Security (TLS) Protocol Version 1.3
[RFC 8490]	https://www.rfc-editor.org/rfc/rfc8490	DNS Stateful Operations
[RFC 8765]	https://www.rfc-editor.org/rfc/rfc8765	DNS Push Notifications
[RFC 8766]	https://www.rfc-editor.org/rfc/rfc8766	Discovery Proxy

Reference	Reference Location/URL	Description
[RFC 8915]	https://www.rfc-editor.org/rfc/rfc8915	Network Time Security for the Network Time Protocol
[RFC 9293]	https://www.rfc-editor.org/rfc/rfc9293	Transmission Control Protocol
[RFC 9665]	https://www.rfc-editor.org/rfc/rfc9665	Service Registration Protocol for DNS-Based Service Discovery
[SEC 1]	https://www.secg.org/sec1-v2.pdf	SEC 1: Elliptic Curve Cryptography, Version 2.0, Certicom Research, May 2009
[SEC 2]	https://secg.org/sec2-v2.pdf	SEC 2: Recommended Elliptic Curve Domain Parameters, Version 2.0, Certicom Research, January 2010
[SIGMA]	https://doi.org/10.1007/978-3-540-45146-4_24	Krawczyk H. (2003) SIGMA: The ‘SIGn-and-MAC’ Approach to Authenticated Diffie-Hellman and Its Use in the IKE Protocols. In: Boneh D. (eds) Advances in Cryptology - CRYPTO 2003. CRYPTO 2003. Lecture Notes in Computer Science, vol 2729. Springer, Berlin, Heidelberg.
[SNAC]	https://datatracker.ietf.org/group/snac/about/	Stub Network Auto-Configuring
[SPAKE2+]	https://tools.ietf.org/pdf/draft-bar-cfrg-spake2plus-02.pdf	SPAKE2+, an Augmented PAKE (Draft 02, 10 December 2020)
[Thread]	https://www.threadgroup.org	Thread 1.3.0 Specification
[Verhoeff]	https://ir.cwi.nl/pub/13045	Verhoeff, J. (1969). Error detecting decimal codes. MC Tracts. Centrum Voor Wiskunde en Informatica.
[WFA Unsynchronized Service Discovery]	https://www.wi-fi.org/discover-wi-fi/wi-fi-aware	Wi-Fi Alliance Wi-Fi Aware Specification.
[X.501]	https://www.itu.int/rec/T-REC-X.501/en	ITU X.501 : Information technology - Open Systems Interconnection - The Directory: Models

Reference	Reference Location/URL	Description
[X.509]	https://www.itu.int/rec/T-REC-X.509/en	ITU X.509 : Information technology - Open Systems Interconnection - The Directory: Public-key and attribute certificate frameworks
[X.520]	https://www.itu.int/rec/T-REC-X.520/en	ITU X.520 : Information technology - Open Systems Interconnection - The Directory: Selected attribute types
[X.680]	https://www.itu.int/rec/T-REC-X.680/en	ITU X.680 : Information technology - Abstract Syntax Notation One (ASN.1): Specification of basic notation
[X.690]	https://www.itu.int/rec/T-REC-X.690/en	ITU X.690 : Information technology - ASN.1 encoding rules: Specification of Basic Encoding Rules (BER), Canonical Encoding Rules (CER) and Distinguished Encoding Rules (DER)

1.7. Informative References

1.7.1. Connectivity Standards Alliance Reference Documents

Reference	Reference Location/URL	Description
[DotdotArch]	https://groups.csa-iot.org/wg/matter-tsg/document/18649	Dotdot Architecture Model, document 13-0589, revision 14, February 2019
[ZCL]	https://groups.csa-iot.org/wg/members-all/document/23019	Zigbee Cluster Library Specification, document 07-5123, revision 8, December 2019

1.8. Conventions

The following conventions are used in this document.

1.8.1. Enumerations and Reserved Values

An undefined value or range of an enumeration, field, or identifier SHALL be considered reserved for future revisions of this standard and SHALL NOT be available for implementation. It is RECOMMENDED that a value stay undefined, rather than defining it as "reserved".

A value or range of an enumeration, field, or identifier that is available for non-standard implementation SHALL be defined as manufacturer specific ("MS").

A value or range of an enumeration, field, or identifier that is available for other parts of this standard SHALL be defined as such.

A value or range of an enumeration, field, or identifier that is deprecated, and not available for implementation, SHALL be defined as deprecated ("D").

1.8.2. Reserved Bit Fields

An undefined bit or bit field SHALL be considered reserved for future revisions of this standard and SHALL NOT be available for implementation.

It is RECOMMENDED that a bit stay undefined, rather than defining it as "reserved".

An implementation of a revision where a bit is reserved SHALL indicate that bit as zero when conveying that bit in an interaction, and ignore that bit when conveyed from another implementation.

1.8.3. Number Format

In this specification, hexadecimal numbers are prefixed with the designation "0x" and binary numbers are prefixed with the designation "0b". All other numbers are assumed to be decimal unless indicated otherwise within the associated text.

When a hexadecimal number is longer than 4 nibbles, the underscore ("_") character SHALL be used to separate the number into groups of 4 nibbles. If the number of nibbles is not evenly divisible by 4, the most significant group SHALL contain less than 4 nibbles and any lesser significant groups SHALL contain 4 nibbles.

Examples of valid hexadecimal groups separated by an underscore:

- "0xAA_33CC"
- "0x1234_ABCD"
- "0x1234_ABCD_526E_6242"

Binary numbers are specified as successive groups of bits, and SHALL be separated by a space (" ") or underscore ("_") character at least every 4 bits from the most significant bit (next to the 0b prefix and leftmost on the page) to the least significant bit (rightmost on the page), e.g. the binary number 0b0000_1111 represents the decimal number 15 and hexadecimal 0x0F.

Where individual bits are indicated (e.g. bit 3) the bit numbers are relative to the least significant bit which is bit 0.

When a digit is specified as having any value in the range of that digit, it is specified with an "x" (this should not be confused with the "x" in the prefix "0x" for hexadecimal notation).

For example:

- "0b0000 0xxx" indicates that the lower 3 bits can take any value but the upper 5 bits must

each be set to 0.

- “0x0000_0xxx” indicates that the lower 3 nibbles can take any value but the upper 5 nibbles must each be set to 0.

Bit ranges are always indicated starting with the most significant bit number. Therefore bits 7..4 are equivalent to the mask 0b1111_0000 or 0xF0.

1.8.4. Provisional

Per [Alliance-PNP], when a specification is completed there may be sections of specification text (or smaller pieces of a section) that are not certifiable at this stage. These sections (or smaller pieces of a section) are marked as provisional prior to publishing the specification. This specification uses well-defined notation to mark [Provisional Conformance](#) or notes a section of text with the term "provisional".

Chapter 2. Architecture

2.1. Overview

Matter aims to build a universal IPv6-based communication protocol for smart home devices. The protocol defines the application layer that will be deployed on devices as well as the different link layers to help maintain interoperability. The following [diagram](#) illustrates the normal operational mode of the stack:




Figure 1. Application and Network Stack

2.2. Layered Architecture

The architecture is divided into layers to help separate the different responsibilities and introduce a good level of encapsulation amongst the various pieces of the protocol stack. The vast majority of interactions flow through the stack captured in the following [Figure](#).




Figure 2. Layered Architecture

The *Application* layer corresponds to the high order business logic of a device. For example, an application that is focused on lighting might contain logic to handle turning on/off a light bulb, as well as its color characteristics.

The *Data Model* layer corresponds to the data and verb elements that help support the functionality of the application. The Application operates on these data structures when there is intent to interact with the device.

The *Interaction Model* layer defines a set of interactions that can be performed between a client and server device. For example, reading or writing attributes on a server device would correspond to application behavior on the device. These interactions operate on the elements defined at the data model layer.

Once an action is constructed using the *Interaction Model*, it is serialized into a prescribed packed binary format to encode for network transmission. This process is handled in the *Action Framing* layer.

An encoded action frame is then processed by the *Security Layer*: the message is encrypted and appended with a message authentication code. These actions ensure the data remain confidential and authentic between sender and receiver of the message.

With an interaction now serialized, encrypted, and signed, the *Message Layer* constructs the payload format with required and optional header fields, which specify properties of the message as well logical routing information.

After the final payload has been constructed by the *Message Layer*, it is sent to the underlying transport protocol (TCP or Matter's [Message Reliability Protocol](#)) for IP management of the data.

Once the data is received on the peer device, it travels up the protocol stack, where the various layers reverse the operations on the data performed by the sender, to finally deliver the message to the Application for consumption.

In addition to the data flows captured above, this specification defines secure session establishment protocols based on operational certificates (see [Section 4.14.2, “Certificate Authenticated Session Establishment \(CASE\)”](#)), or passcodes (see [Section 4.14.1, “Passcode-Authenticated Session Establishment \(PASE\)”](#)), group communication (see [Section 4.16, “Group Communication”](#)), a bulk data transfer protocol (BDX) suitable for sending bulk data between Nodes, and provisions for defining manufacturer-specific protocols.

2.3. Network Topology

In principle, any IPv6-bearing network is suitable for Matter deployment, subject to supporting a few core IPv6 standards. In this version of the specification, we focus on three link layer technologies: Ethernet, Wi-Fi and Thread. We restrict the specification to the above so that the specification can suitably cover provisioning of these link layers, and so that the amount of testing in certification is suitably bounded.

Matter treats networks as shared resources: it makes no stipulation of exclusive network ownership or access. As a result, it is possible to overlay multiple Matter networks over the same set of constituent IP networks.

This protocol may operate in the absence of globally routable IPv6 infrastructure. This requirement enables operation in a network disconnected or firewalled from the global Internet. It also enables deployment in situations where the Internet Service Provider either does not support IPv6 on consumer premises or where the support proves otherwise limiting, for example, if the delegated prefix cannot accommodate all the networks and devices on premises.

This protocol supports local communications spanning one or more IPv6 subnets. Canonical networks supporting a fabric may include a Wi-Fi/Ethernet subnet, or one or more low power and lossy networks (LLN) subnets. In this version of the specification, Thread is the supported LLN standard.

2.3.1. Single network




Figure 3. Single Thread network




Figure 4. Single Wi-Fi network

In the single network topology, all Matter devices are connected to a single logical network. This could be a Thread/802.15.4 network, a Wi-Fi network or an Ethernet network. In the case of Wi-Fi/Ethernet, the network could in fact span multiple Wi-Fi and/or Ethernet segments provided that all the segments are bridged at the link layer. A Node is a single instance of a Matter device within a fabric, operationally available on an IP network.

Each Node in the single-network topology communicates with every other Node in the Fabric via a single network interface.

2.3.2. Star network topology




Figure 5. Star network topology

The star network topology consists of multiple peripheral networks joined together by a central hub network. The hub network will typically be the customer’s home network (Wi-Fi/Ethernet network), while the peripheral networks can be of any supported network type. A peripheral network SHALL always be joined directly to the hub network via one or more Border Routers.

Architecturally, any number of peripheral networks may be present in a single Fabric, including multiple networks of the same type. Nodes MAY have interfaces onto any network (hub or peripheral), and MAY communicate directly to other Nodes on the same network. However, any communication that must cross a network boundary to reach its destination SHALL flow through a Border Router. For additional details see [Section 4.2, “IPv6 Reachability”](#).

This protocol places a set of requirements on the Border Router. These requirements pertain to address assignment, route assignment and advertisement, multicast support (see [Section 4.2.1, “Stub Router Behavior”](#)), and discovery proxying (see [Section 4.3, “Discovery”](#)).

Note that in this version of the specification, Thread is the primary supported LLN. In many cases, customer installations will attempt to maintain a simple network topology, with a single Wi-Fi/Ethernet subnet, and a single Thread network. However, more than one Thread network is possible and supported.

To support home automation interoperability, this protocol supports the concept of [bridging](#) which makes available, through a data model node, devices implementing other home automation technologies, transports and link layers.

2.4. Scoped names

The Matter protocol explicitly supports multiple administrators, unrelated by any common roots of trust (multi-admin). This functionality is addressed via [multiple fabrics](#) and is enabled by the core aspects of name scoping (see below), and key considerations enabling multiple fabrics in [commissioning](#), [secure communication](#), and aspects of the data model (such as [fabric-scoped data](#)).

A Fabric is a collection of Matter devices sharing a trusted root. The root of trust in Matter is the Root CA that issues the [NOCs](#) which underpin node identities. Within the fabric, each node is uniquely identified by a stable Node ID. The scoped selection and allocation of these constructs within Matter ensures the uniqueness of identifiers and gives clear guidance on ownership and management of namespaces.

The operational root of trust—the root certificate authority (CA) as identified by its public key (Root PK)—is responsible for the allocation of correctly scoped fabric identifiers. The security of all public key infrastructures (PKI) depends on the private key of the CA being protected and neither guessable nor obtainable; that property, in turn, implies that the public key is globally unique. Within any root CA, the fabrics—identified by a 64-bit number—are unique. The uniqueness mechanism emerges from the collaboration of the commissioner and the root CA associated with that particular commissioner. Matter wraps the collaboration between the commissioner and its associated root CA and other possible data stores into a concept called the "administrative domain manager" (ADM). The algorithmic details and policies within the administrative domain manager are out of the scope of the specification as long as the allocation of all identifiers obeys the uniqueness and scoping criteria. Fabrics are uniquely identified by the tuple of their root CA's public key and a Fabric ID. Similarly, within each fabric, the administrative domain manager is responsible for assigning a unique and stable Operational Node ID to every node.

The scoping strategy as outlined here ensures that each scoped identifier can be allocated solely by the entities responsible for the scoping, without consideration for collisions or forgeries. For example, two different CAs may allocate the same fabric identifiers and this would not create any prob-

lems for the devices within the network. Scoped delegation of responsibility also provides for clear guidelines for the removal of specific identifiers.

A Matter device may be a member of multiple fabrics and thus have multiple associated node IDs. The scoping strategy also naturally lends itself towards unambiguous resolution of names and credentials and places a clearly defined responsibility for managing the namespaces on each fabric's associated administrative domain manager service.

Prior to the first commissioning, such as in factory-reset state, a typical device contains no pre-allocated operational roots of trust, and no operational identities in the form of fabric IDs and node IDs. Yet, various interactions expect the fabric ID, or a node ID. These identifiers emerge in a number of internal constructs — from address discovery, through identifying secure sessions, to evaluating access control privileges. In order to regularize all interactions with the device and solve the bootstrapping problem, a special primordial fabric ID is reserved, and associates a set of initial access control privileges with any communication that would be associated with the initial commissioning sessions.

2.5. Identifiers

2.5.1. Fabric References and Fabric Identifier

As described above, a Fabric ID is a 64-bit number that uniquely identifies the Fabric within the scope of a particular root CA. Conceptually, the fully qualified fabric reference consists of the tuple containing the public key of the root certificate authority, and the Fabric ID. Because the fully qualified fabric reference is cumbersome to use, a number of mechanisms for compression of the reference are defined. The Fabric reference, in compressed form, is used during [operational discovery](#) to provide operational naming separation, a form of namespacing, between unrelated collections of devices.

Fabric ID 0 is reserved across all fabric root public key scopes. This fabric ID SHALL NOT be used as the identifier of a fabric.

A fabric is defined in the Data Model as a set of nodes that interact by accessing Data Model elements as defined in the Interaction Model (see [Section 7.5, “Fabric”](#)).

The layers below the data model, that convey data model interactions as messages, SHALL always indicate either the fabric associated with the message, or that there is no fabric associated with the message.

For example: A Data Model message that is conveyed over a message channel that uses the reserved fabric ID '0' does not have a fabric associated with it.

2.5.2. Vendor Identifier (Vendor ID, VID)

A Vendor Identifier (Vendor ID or VID) is a 16-bit number that uniquely identifies a particular product manufacturer, vendor, or group thereof. Each Vendor ID is statically allocated by the Connectivity Standards Alliance (see [\[Connectivity Standards Alliance Manufacturer Code Database\]](#)).

The following Vendor IDs are reserved:

Table 1. Vendor ID Allocations

Range	Type
0x0000	Matter Standard
0x0001 - 0xFFFF0	reserved for individual Manufacturer Codes as per [Connectivity Standards Alliance Manufacturer Code Database]
0xFFFF1	Test Vendor #1
0xFFFF2	Test Vendor #2
0xFFFF3	Test Vendor #3
0xFFFF4	Test Vendor #4

All other allocations of Vendor ID are specified in [\[Connectivity Standards Alliance Manufacturer Code Database\]](#).

NOTE

The *Test Vendor* IDs are reserved for test and development by device manufacturers or hobbyists. Commissioners SHOULD NOT commission devices using one of these VIDs onto an operational Fabric under normal operation unless the user is made fully aware of the security risks of providing an uncertified device with operational and networking credentials.

2.5.3. Product Identifier (Product ID, PID)

A Product Identifier (Product ID or PID) is a 16-bit number that uniquely identifies a product of a vendor. The Product ID is assigned by the vendor and SHALL be unique for each product within a [Vendor ID](#). Once a Product ID has been used, it SHALL NOT be reused by a different product type under the same Vendor ID. These Product IDs SHOULD NOT be specific to a unique physical device; rather they describe the product type, which might have many manufactured instances (e.g. multiple colors of the same product type).

A value of 0x0000 SHALL NOT be assigned to a product since Product ID = 0x0000 is used for these specific cases:

- To announce an anonymized Product ID as part of device discovery (see [Section 5.4.2, “Announcement by Device”](#)).
- To indicate an [OTA software update file](#) applies to multiple Product IDs equally.
- To avoid confusion when presenting the [Onboarding Payload for ECM with multiple nodes](#).

2.5.4. Group Identifier (GID)

A Group Identifier (Group ID or GID) is a 16-bit number that identifies a set of Nodes across a Fabric at the message layer (see [Section 4.17, “Group Key Management”](#)). A Group ID can further be bound to one or more Endpoints within each Node in the group at the interaction layer.

The Group ID space is allocated as described in [Table 2, “Group ID Allocations”](#):

Table 2. Group ID Allocations

Range	Type
0xFF00 - 0xFFFF	Universal Group ID range reserved for static multi-cast and anycast identifiers
0x0001 - 0xFEFF	Application Group ID, assigned by fabric Administrator
0x0000	Null or unspecified Group ID

2.5.4.1. Universal Group ID

A Universal Group ID (UGID) is one that resides in the 16-bit subrange of Group ID that is reserved for groups that are common across this standard. These special multicast, groupcast, or anycast destinations are constant and known to all Nodes on any Fabric. The Universal Group ID space is allocated as described in [Table 3, “Universal Group ID Allocations”](#):

Table 3. Universal Group ID Allocations

Range	Type
0xFFFF	All Nodes
0xFFFFE	All non-ICD Nodes
0xFFFFD	<i>Deprecated</i>
0xFF00-0xFFFC	Reserved for future use

Because the keys and [IPv6 multicast prefixes](#) are different across Fabrics, Universal Groups only enable communication within a specific Fabric.

All Nodes Group

This group is used to message all Nodes in a Fabric.

All non-ICD Nodes Group

This group is used to message all power-capable Nodes in a Fabric. [ICD](#) Nodes SHALL NOT subscribe to this group.

2.5.5. Node Identifier

A Node Identifier (Node ID) is a 64-bit number that uniquely identifies an individual Node or a group of Nodes on a Fabric. The Node Identifier space is allocated as described in [Table 4, “Node Identifier Allocations”](#):

Table 4. Node Identifier Allocations

Range	Type
0xFFFF_FFFF_FFFF_xxxx	Group Node ID

Range	Type
0xFFFF_FFFF_0000_0000 to 0xFFFF_FFF-F_FFFE_FFFF	Reserved for future use
0xFFFF_FFFE_xxxx_xxxx	Temporary Local Node ID
0xFFFF_FFFD_xxxx_xxxx	CASE Authenticated Tag
0xFFFF_FFFC_xxxx_xxxx to 0xFFFF_FF-FC_FFFF_FFFF	Reserved for future use
0xFFFF_FFB_xxxx_xxxx	PAKE key identifiers
0xFFFF_FFF0_0000_0000 to 0xFFFF_FF-FA_FFFF_FFFF	Reserved for future use
0x0000_0000_0000_0001 to 0xFFFF_FFE-F_FFFF_FFFF	Operational Node ID
0x0000_0000_0000_0000	Unspecified Node ID

Node IDs are used for core messaging, within the internal APIs, within the data model, and to resolve the operational IPv6 addresses of Nodes (see [Section 4.3.2, “Operational Discovery”](#)).

The span of Node IDs from 0xFFFF_FFF0_0000_0000 to 0xFFFF_FFFF_FFFF_FFFF, as well as the value 0x0000_0000_0000_0000 are both reserved for special uses.

2.5.5.1. Operational Node ID

An Operational Node ID is a 64-bit number that uniquely identifies an individual Node on a Fabric. All messages must have an Operational Node ID as the source address. All unicast messages must have an Operational Node ID as the destination address.

While source or destination address MAY be elided from a message, it SHALL remain unambiguously derivable from the [Session ID](#).

2.5.5.2. Group Node ID

A Group Node ID is a 64-bit Node ID that contains a particular [Group ID](#) in the lower half of the Node ID.

2.5.5.3. Temporary Local Node ID

A Temporary Local Node ID is a 64-bit Node ID that contains an implementation-dependent value in its lower 32 bits. This allows implementations to keep track of connections or transport-layer links and similar housekeeping internal usage purposes in contexts where an Operational Node ID is unavailable.

2.5.5.4. PAKE key identifiers

This subrange of Node ID is used to assign an access control subject to a particular PAKE key as specified in [Section 6.6.2.1.1, “PASE and Group Subjects”](#). An example usage would be to create an ACL entry to provide administrative access to any commissioner communicating via a PASE session

established with a particular pincode.

2.5.5.5. CASE Authenticated Tag

This subrange of Node ID is used to assign an access control subject to a group of peer nodes that share a single CASE session as specified in [Section 6.6.2.1.2, “Subjects identified by CASE Authenticated Tag”](#).

2.5.5.6. Unspecified Node ID

The Unspecified Node ID (0x0000_0000_0000_0000) is a reserved value that never appears in messages or protocol usage. It exists to mark or detect the presence of uninitialized, missing, or invalid Node IDs.

2.5.6. IPv6 Addressing

This protocol uses IPv6 addressing for its operational communication. Node IDs and Fabric IDs are resolved to various types of IPv6 addresses [[RFC 4291](#)].

2.5.6.1. IPv6 Unicast Address

An IPv6 Unicast Address is one that uniquely identifies and addresses a single Node on an IPv6 network. A primary design goal for this standard is to allow Nodes to leverage native IPv6 technologies. As such, an operational IPv6 Unicast address that provides connectivity and routability between Nodes SHALL be used. This includes a global unicast address (GUA), a link-local address (LLA), or a unique local address (ULA).

2.5.6.2. IPv6 Multicast Address

An IPv6 Multicast Address is formed using Unicast-Prefix-based IPv6 Multicast Addresses [[RFC 3306](#)]:

- The first 12 bits are defined by [[RFC 3306](#)] and are 0xFF3.
- The next 4 bits are "scop" (scope) and set based on [[RFC 7346](#)] Section 2 to:
 - *Site-Local* (0x5) - spans all networks in the Fabric, including Thread, Ethernet, and Wi-Fi.
- The next 8 bits are reserved (0x00).
- The next 8 bits are "plen", and set to 0x40 to indicate a 64-bit long network prefix field.

The network prefix portion of the Multicast Address is the 64-bit bitstring formed by concatenating:

- 0xFD to designate a locally assigned ULA prefix per [[RFC 4193](#)] Section 3.1
- The upper 56-bits of the Fabric ID for the network in big-endian order

The 32-bit group identifier portion of the Multicast Address is the 32-bits formed by:

- The lower 8-bits of the Fabric ID
- 0x00
- The next 16-bits are the Group Identifier for the group, as specified in [Group Identifier](#) in big-

endian order

An example of the site local scoped multicast address for a given <Fabric ID> and <Group ID>:

```
FF35:0040:FD<Fabric ID>00:<Group ID>
```

NOTE

Though *Site-Local* scope is always used, the effective scope MAY be limited by setting the IPv6 hop count. However, hop counts are not deterministic on Thread networks and might cause groupcast communication to fail if the hop limit is set too low. Therefore, the supported hop count of the underlying link technology SHOULD be set to at least 10 for efficient Matter groupcast communication.

The Multicast Address formation ensures a low probability of a node receiving a multicast message it is not interested in. If a collision does occur on the multicast address (which requires two identical 64-bit Fabric IDs and two identical 16-bit Group IDs), processing of the message disambiguates which fabric and group is relevant by checking which operational group key leads to the message's 64-bit MIC.

IANA Assigned IPv6 Multicast Address

Any device supporting the Groupcast cluster, whether as a Listener or Sender, SHALL support the use of the site-scoped IANA-assigned Multicast Address for Matter, see [IANA IPv6 Multicast Address Space](#).

```
FF05::FA
```

When used over Thread, this address benefits from stateless multicast address compression ([RFC 6282](#), Section 3.2.3 - Figure 9). 6LoWPAN layers such as Thread SHALL use stateless multicast address compression.

2.5.6.3. IPv6 Multicast Port Number

The IANA assigned port number is 5540.

2.5.6.4. IPv4 Coexistence

Matter devices SHALL be tolerant of IPv4 addresses and MAY ignore those addresses for the purposes of Matter operations.

2.6. Device identity

Each Matter device holds a number of certificate chains.

A Device Attestation Certificate (DAC) proves the authenticity of the manufacturer and a certification status of the device's hardware and software. The Device Attestation Certificate is used during the commissioning process by the Commissioner to ensure that only trustworthy devices are admitted into a Fabric. The details of the device attestation process are captured in [Section 6.2, "Device](#)

Attestation”.

Each Matter device is issued an [Operational Node ID](#) and a [Node Operational Certificate \(NOC\)](#) for that Operational Node ID. The NOC enables a Node to identify itself within a Fabric by cryptographically binding a unique Node Operational Key Pair to the operational identity of its subject, attestable through the signature of a trusted Certificate Authority (CA). Operational Node IDs are removed on factory reset or removal of Fabrics. A NOC is issued during the commissioning process of a device into a Fabric. These steps help to protect the privacy of the end-user and to adapt to different trust models.

The format of the Node Operational credentials and protocols for generating those credentials are detailed in the following sections: [Section 6.4, “Node Operational Credentials Specification”](#) and [Section 6.5, “Operational Certificate Encoding”](#).

2.7. Security

Matter deploys robust security practices to protect the Fabric. Matter designates a core set of security primitives (detailed in [Chapter 3, *Cryptographic Primitives*](#)) to provide comprehensive protection. Elliptic curve cryptography, based on the NIST P-256 curve (secp256r1) serves as the foundation for public key cryptography and digital signatures. Commonly available AES modes of operation have been selected to provide shared key cryptographic operations. In scenarios involving an out-of-band passcode-based authentication, Matter uses SPAKE2+, an efficient augmented PAKE algorithm.

The core cryptographic primitives form the basis of a number of complementary secure protocols used within Matter. All unicast Node-to-Node messages are secured, authenticated, and provide replay protection. Building on top of IPv6 multicast, Matter also provides secured group messaging facilities, useful for efficiently addressing on an LLN. The group messaging features prioritize low latency of packet processing.

2.8. Device Commissioning

Device commissioning (see [Chapter 5, *Commissioning*](#)) is the process of joining a device to a Fabric. The device being commissioned is referred to as the Commissionee and the device administering commissioning is referred to as the Commissioner. Device commissioning consists of the following steps:

1. Device discovery (see [Section 5.4, “Device Discovery”](#) and see [Section 5.1, “Onboarding Payload”](#)): The Commissioner discovers commissionable devices on network interfaces such as Bluetooth Low Energy, Wi-Fi, or other connected IP network. The Commissioner obtains the out-of-band secret ([Section 5.1.1.6, “Passcode”](#)) from the commissionable device’s [QR Code](#), [Manual Pairing Code](#), [NFC Tag](#) or other means. This secret is used by [Passcode-Authenticated Session Establishment \(PASE\)](#) to establish a secure commissioning session. The order of discovering commissionable devices and obtaining the out-of-band secret from commissionable device is not critical.
2. Security setup with PASE (see [Section 4.14.1, “Passcode-Authenticated Session Establishment \(PASE\)”](#)): Establish encryption keys between the Commissioner and Commissionee using PASE. All messages exchanged between the Commissioner and Commissionee are encrypted using

these PASE-derived keys. The process also establishes an attestation challenge used during the [device attestation procedure](#).

3. Device attestation verification (see [Section 6.2, “Device Attestation”](#)): Commissioner establishes the authenticity of the Commissionee as a certified device, notifying the user if the device is not certified.
4. Information configuration (see [Section 6.4, “Node Operational Credentials Specification”](#), [Section 11.10, “General Commissioning Cluster”](#) and [Section 11.18, “Operational Credentials Cluster”](#)): The Commissioner provides Commissionee with information such as regulatory domain, UTC time, Operational Certificate and network interfaces configuration.
5. Join network (see [Section 11.9, “Network Commissioning Cluster”](#) and [Section 4.3.2, “Operational Discovery”](#)): The Commissioner triggers the Commissionee to connect to the operational network unless the Commissionee is already on the operational network. The Node’s/Commissionee’s IPv6 address is then either used (if already known) or discovered (if not known) by the Commissioner or Administrator.
6. Security setup with CASE (see [Section 4.14.2, “Certificate Authenticated Session Establishment \(CASE\)”](#)): Derive encryption keys used to establish secure communication between the Commissioner or Administrator and Node with CASE. All unicast messages between a Commissioner or Administrator and a Node are encrypted using these CASE-derived keys.
7. Commissioning complete message exchange (see [Section 11.10, “General Commissioning Cluster”](#)): A message exchange encrypted using CASE-derived encryption keys on the operational network that indicates successful completion of commissioning process.

A commissioner can reconfigure the Commissionee multiple times over the operational network after the commissioning is complete or over the commissioning channel after PASE-derived encryption keys are established during commissioning. The commissioning flows are described in [Section 5.5, “Commissioning Flows”](#).

2.9. Intermittently Connected Device (ICD)

One goal of this standard is to provide support for communicating with devices that are intermittently connected to the network or unreachable for periods of time. The Intermittently Connected Device operating mode behavior is specified in the [Section 9.15, “Intermittently Connected Devices Behavior”](#) section. The Intermittently Connected Device (ICD) operating mode is defined to help optimize network traffic, reachability, and power consumption for such nodes. Intermittent connections can be caused by a number of factors, including duty cycling to minimize power consumption, intermittent availability of network connectivity, device mobility, or intermittent availability of power.

The ICD operating mode is designed to be agnostic to underlying network layer mechanisms and may be leveraged over any supported IP interfaces, including Thread and Wi-Fi.

NOTE

Because clients and infrastructure devices may have limited buffering space to cache messages on behalf of intermittently connected devices, ICD communication patterns SHOULD be designed such that the ICD is predominantly the initiator.

2.9.1. Sleepy End Device (SED)

The Sleepy End Device (SED) operating mode is defined by the Thread standard to help extend and optimize battery lifetimes for Thread devices running on limited power sources such as batteries or limited energy scavenging. While the Matter ICD operating mode can leverage Thread Sleepy End Device (SED) behavior for Thread ICD devices, it SHOULD NOT be confused with it.

2.10. Data Model Root

- Endpoint 0 (zero) SHALL be the [root node endpoint](#).
- Endpoint 0 (zero) SHALL support the Root Node device type.

2.11. Stack Limits

2.11.1. System Model Limits

2.11.1.1. Access Control Limits

- A node SHALL guarantee that there are at least four [Access Control Entries](#) available for every fabric supported by the node.

For example: A node that supports 5 fabrics must support at least 20 ACL entries, and if it supports N entries must enforce that any K fabrics together do not use more than $N - 4 \cdot (5 - K)$ entries.

- Device types MAY impose additional constraints on the number of ACL entries that need to be supported.

2.11.1.2. Group Limits

- A node SHALL support at least one group key per fabric for managing the [IPK](#).

When the node implements the [Groupcast Cluster](#), the [Groups cluster mandatory minima](#) are superseded by the [Groupcast cluster minima](#) and do not apply. However, the node SHALL still maintain any existing legacy group configured as described in [Migration from Groups to Groupcast \(Matter 1.6.x\)](#).

NOTE

This backward-compatibility requirement can impact the node's [Groupcast cluster minima](#) (e.g. the configured [MaxMembershipCount](#) value).

Group Limits when using Groups cluster

- If the node implements one or more device types with support for the Groups cluster, the node SHALL additionally support the maximum number of the required groups as specified by all of these implemented device types, without going below the following mandatory minima:
 - The node SHALL support at least three group keys per fabric.

- The node SHALL support at least four group table entries per fabric per endpoint having a Groups cluster instance.
- A node SHALL support one IPv6 multicast group membership for every operational group it supports.
- Each Groups cluster instance SHALL support adding the endpoint to at least four groups.

For example: A node that supports 5 fabrics and has 1 endpoint with a Groups cluster must support at least 20 group table entries. A node that supports 5 fabrics and has 2 endpoints with a Groups cluster must support at least 40 group table entries.

Group Limits when using Groupcast cluster

- If the node implements the [Groupcast Cluster](#), it SHALL support the following mandatory minima:
 - The node SHALL support as many group memberships as the value it reports in the [MaxMembershipCount](#) attribute.
 - All non-Root endpoints supported MAY be added in one or more MembershipStruct entries.
 - For every [MembershipStruct](#) entry, the node SHALL persist all EndpointIds in that entry.
 - The node SHALL support at least four GroupKeySet entries per fabric.

2.11.2. Interaction Model Limits

2.11.2.1. Read Interaction Limits

- A server SHALL ensure that every fabric the node is commissioned into can handle a single [Read Interaction](#) from a client on that fabric containing up to 9 paths.
- A server MAY permit [Read Interactions](#) even when there is no accessing fabric, subject to available resources (e.g over PASE).

2.11.2.2. Subscribe Interaction Limits

- A publisher SHALL ensure that every fabric the node is commissioned into can support at least three [Subscribe Interactions](#) to the publisher and that each subscription SHALL support at least 3 attribute/event paths.
- A server MAY permit [Subscribe Interactions](#) even when there is no accessing fabric, subject to available resources (e.g over PASE).
- Device type specifications MAY require a larger number of supported subscriptions or paths.
- [SUBSCRIPTION_MAX_INTERVAL_PUBLISHER_LIMIT](#) defines the upper limit for the publisher-selected maximum interval for any subscription.
 - If the publisher is an [ICD](#), this SHALL be set to the [Idle Mode Duration](#) or 60 minutes, whichever is greater.
 - Otherwise, this SHALL be set to 60 minutes.

- The minimal supported capabilities, subject to the minimal constraints above, are reported in the CapabilityMinima attribute of the [Basic Information cluster](#).

2.11.2.3. Invoke Interaction Limits

- An [Invoke Request action](#) MAY contain multiple [concrete command paths](#), if supported by the server, but the total number of commands included in a single Invoke Request action is still limited by the maximum message size constraints.

2.12. Migration from Groups to Groupcast (Matter 1.6.x)

Starting with the release of Matter when the [Groupcast Cluster](#) becomes certifiable, the Groupcast cluster is the recommended mechanism for configuring, managing, and using groups.

A grace period is provided where both Groups and Groupcast are supported until the full transition to Groupcast is complete.

This section provides a migration plan for controllers, and for manufacturers updating deployed devices, avoiding burdening devices that never had pre-Groupcast implementation from needing to support prior behaviors.

2.12.1. Backward compatibility for updated nodes

An updated node, that previously supported the Groups cluster, SHALL continue to operate with existing group communications at the message/transport layer (e.g. previously provisioned group keys, and multicast addressing) after the upgrade to the first Matter version where the Groupcast cluster becomes certifiable. This enables deployed clients, bindings, and automations to continue to function while the migration occurs.

NOTE

This deprecation approach allows vendors to update older products to the Groupcast model, without permanently carrying the legacy groups implementation cost, while still maintaining continuity of operation for already-configured groups.

To maintain backwards compatibility while still deprecating legacy group management, device behaviour depends on the device type requirements for the Groups cluster:

- If a device type requires the Groups cluster (mandatory conformance), then the node SHALL keep the Groups cluster instance(s) enabled.

If the node has any pre-existing legacy group configuration from prior firmware, it SHALL be reflected into the [Membership](#) attribute as described below.

2.12.2. Device-managed migration into Groupcast Membership

When a node is updated from a version using the legacy Groups cluster, the existing group configuration, persisted on the device, SHALL be reflected into the [Membership](#) attribute. Since the legacy data is compatible with the [MembershipStruct](#) structure, used by the [Groupcast Cluster](#), the refer-

ence open source SDK's group data provider can interpret the persisted data directly in the [MembershipStruct](#) format, without requiring explicit one-time migration code, which would execute once after a firmware update. Fields introduced by the [Groupcast Cluster](#) that are not present in the legacy data, such as `HasAuxiliaryACL`, are assigned default values (e.g. `HasAuxiliaryACL=false`, `McastAddrPolicy=PerGroup`).

NOTE

Implementations that rely on the reference open source SDK do not need to implement additional migration logic. The mapping described below is handled transparently by that implementation. The following description is primarily relevant for implementations that use a custom data storage layer for group state, which need to ensure that the persisted data is correctly presented through the [Membership](#) attribute after a firmware update.

The legacy group configuration consists of the following elements:

- The per-endpoint group membership state (the set of GroupIDs associated with each endpoint); and
- the fabric-scoped GroupID-to-GroupKeySetID pairs from the [GroupKeyMap](#) in the [Group Key Management Cluster](#).

These inputs are consolidated into one or more [MembershipStruct](#) entries per fabric, ensuring that the resulting (GroupID, Endpoints, KeySetID) relationships preserve the node's pre-update configuration. The data equivalence between the legacy group configuration and the [Membership](#) attribute, per fabric, is described by the following mapping:

1. For each fabric:
 - i. For each GroupID:
 - a. Compute the list of endpoint IDs that are part of that group. Discard Endpoint 0 from the list if it was present in the legacy group configuration.
 - b. Determine the KeySet to use by reading the fabric-scoped [GroupKeyMap](#) list from the [Group Key Management Cluster](#), and selecting the GroupKeySetID associated with that GroupID.
 - c. The corresponding [MembershipStruct](#) entry in the [Membership](#) attribute for that fabric SHALL contain:
 - A. GroupID set to the GroupID from step [i](#).
 - B. Endpoints set to the computed endpoint list from step [a](#).
 - C. KeySetID set to the GroupKeySetID from step [b](#).
 - D. HasAuxiliaryACL set to `false`. Legacy groups are expected to already have functional access control. If desired, an Administrator MAY enable AuxiliaryACL per GroupID, using [ConfigureAuxiliaryACL](#) command.
 - E. McastAddrPolicy set to `PerGroup`.

Additionally, the MaxGroupsPerFabric attribute in the Group Key Management cluster SHALL be set to 0 to indicate that legacy group mapping is deprecated. This will serve to prevent ongoing attempts to allocate groups using the Groups cluster. The equivalent for Groupcast is the MaxMem-

bershipCount attribute in the Groupcast cluster.

2.12.3. Administrators and controllers Groupcast migration perspective

When Groupcast cluster is present, controllers and Administrators SHALL rely on the [Membership](#) attribute to discover the node's current group configuration, and SHALL modify group membership using the [Groupcast Cluster](#) commands (e.g. [JoinGroup](#), [LeaveGroup](#), and [UpdateGroupKey](#)).

Group Key Set definitions MAY still be done using the Group Key Management cluster, in addition to the abilities provided by the [JoinGroup](#) and [UpdateGroupKey](#) commands.

If any key material within a KeySet needs to be modified, an Administrator can do so through the [Group Key Management Cluster](#) using the same GroupKeySetID (e.g. via [KeySetWrite Command](#)), and those keys apply consistently for group communication, independently of whether the group was originally created via the legacy Groups mechanism, or via Groupcast, since the message handling remains the same.

2.12.4. Resource sizing and limits: legacy Groups vs Groupcast

The legacy Groups model imposes per-endpoint group table requirements that can lead to static provisioning expectations (i.e. scaling by fabrics multiplied by endpoints with Groups cluster instances). In contrast, the [Groupcast Cluster](#) makes group participation explicit, via the [Membership](#) list, and exposes a device-defined capacity via [MaxMembershipCount](#), including single-fabric usage limits, defined by the Groupcast cluster. A device migrating support, from existing legacy Groups configuration to Groupcast, SHALL ensure that [MaxMembershipCount](#) is set to a value that can accommodate all migrated memberships; otherwise, some pre-existing groups risk being lost due to insufficient Membership capacity.

2.13. List of Provisional Items

The following is a list of provisional items.

2.13.1. Invoke Wildcard Paths

Support for an Invoke Interaction with [wildcard paths](#) is provisional.

2.13.2. WildcardPathFlags options

Support for [WildcardPathFlags](#) option in wildcard expansion is provisional.

2.13.3. Tag compression encoding for AttributePathIB, EventPathIB, and AttributeDataIB

The [EncodingEnableTagCompression](#) mechanism for encoding the interaction model is provisional.

2.13.4. CacheAndSync Feature on the GroupKeyManagement cluster

Support for the CacheAndSync feature on the GroupKeyManagement cluster is provisional (see [Sec-](#)

tion 11.2.4, “Features”).

2.13.5. Asynchronous transfer in Bulk Data Exchange Protocol (BDX)

Support for the [Asynchronous Mode](#) in the Bulk Data Exchange Protocol (BDX) is provisional.

2.13.6. Network Recovery

Support for the [Network Recovery](#) feature, associated advertising and communication flows, and `RecoveryIdentifier` and `NetworkRecoveryReason` attributes in the General Commissioning Cluster is provisional.

2.13.7. Thread Commissioning Protocol

Support for [Thread Commissioning Protocol](#) is provisional.

2.13.8. Groupcast Cluster

Support for the [Groupcast Cluster](#) is provisional.

2.13.9. Auxiliary ACL

Support for the [Auxiliary Feature](#) and [AuxiliaryACL attribute](#) in the [Access Control Cluster](#), are provisional.

2.13.10. IANA Assigned IPv6 Multicast Address

Support for the [IANA Assigned IPv6 Multicast Address](#) is provisional.

Chapter 3. Cryptographic Primitives

This chapter introduces the various cryptographic primitives, algorithms and protocol building blocks used in this protocol. It introduces for each of them a functional abstraction that can be referred to in the other chapters of this specification. This chapter also maps those cryptographic primitives to specific instances with the corresponding appropriate informative or normative references. Wherever relevant, it also gives necessary or relevant information about the use of these mappings in a specific context to achieve a compliant implementation.

Given a [version](#) of the [Message Format](#), the cryptographic primitives are mapped to specific instances. There is no cryptosuite negotiation in this protocol: one [version](#) of the [Message Format](#) has one cryptosuite as defined in this chapter.

Each section defines cryptographic primitives generically, together with concrete mappings to specific instances of these cryptographic primitives for [version 1.0](#) of the [Message Format](#). This chapter can also be used as guidance about which cryptographic primitives need to be supported by a device, but it must be noted that not all devices will have to support all of them. For example, a device may not require the `Crypto_PBKDF()` primitive, as values based on this operation could in some instances be precomputed and stored during the manufacturing process of the device. The proposed functional mapping in this chapter is normative with respect to the values computed by the functions but informative with respect to the way the functions are interfaced within implementations. For example, a function returning both a `boolean` to indicate success and a value if the operation is successful could also be implemented using exception mechanisms instead of returning a `boolean`.

It must also be noted that not all cryptographic primitives are exposed to the other parts of the specification. For example, the `Crypto_TRNG()` primitive SHALL NOT be called outside of the `Crypto_DRBG()` implementation.

The cryptographic primitives discussed below operate on data local to the host. Where more complex data types are present and their external representation is applicable, the chapter notes the details of the encoding. Simple multi-byte data types without any additional context are assumed to be in host byte order when they are used internally to a procedure, unless otherwise stated.

All octet strings are presented with first octet having index 0, and presented from left to right for indices 0 through N-1 for an octet string of length N.

3.1. Deterministic Random Bit Generator (DRBG)

This protocol relies on random numbers for many security purposes. For example, random numbers are used in generating secret keys, counters, cryptographic signature generation random secrets, etc. Those random numbers SHALL be generated using the `Crypto_DRBG()` function.

Function and description

```
bit[len]  
Crypto_DRBG(int len)
```

Returns an array of **len** random bits.

Mapping (Version 1.0)

Crypto_DRBG() SHALL be implemented with one of the following DRBG algorithms as defined in [NIST 800-90A](#) (the choice of which one is left to the manufacturer because the choice has no impact on the interoperability):

- CTR DRBG (with AES-CTR)
- HMAC DRBG (with SHA-256)
- HMAC DRBG (with SHA-512)
- Hash DRBG (with SHA-256)
- Hash DRBG (with SHA-512)

Crypto_DRBG() SHALL be seeded and reseeded using **Crypto_TRNG()** with at least 256 bits of entropy (see among others Chapter 4, Section 8.4, and Section 8.6.8 of [NIST 800-90A](#)).

3.2. True Random Number Generator (TRNG)

A TRNG (a.k.a. Entropy Source) is required to provide an entropy seed as an input to the DRBG algorithm.

Function and description

```
bit[len]  
Crypto_TRNG(int len)
```

Returns an array of **len** random bits.

Mapping (Version 1.0)

Crypto_TRNG() MAY be implemented according to the [NIST 800-90B](#) implementation guidelines but alternate implementations MAY be used.

In accordance with good security practices, the **Crypto_TRNG()** SHALL never be called directly but rather SHALL be used in the implementation of **Crypto_DRBG()**.

3.3. Hash function (Hash)

Crypto_Hash() computes the cryptographic hash of a message.

Function and description
byte[CRYPTO_HASH_LEN_IN_BYTES] Crypto_Hash(byte[] message) Returns the cryptographic hash digest of the message.
Mapping (Version 1.0)
int CRYPTO_HASH_LEN_BITS := 256 int CRYPTO_HASH_LEN_BYTES := 32 Crypto_Hash(message) := byte[CRYPTO_HASH_LEN_BYTES] SHA-256(M := message) SHA-256() SHALL be computed as defined in Section 6.2 of FIPS 180-4.

3.4. Keyed-Hash Message Authentication Code (HMAC)

Crypto_HMAC() computes the cryptographic keyed-hash message authentication code of a message.

Function and description
byte[CRYPTO_HASH_LEN_BYTES] Crypto_HMAC(byte[] key, byte[] message) Returns the cryptographic keyed-hash message authentication code of a message using the given key.
Mapping (Version 1.0)
Crypto_HMAC(key, message) := byte[CRYPTO_HASH_LEN_BYTES] HMAC(K := key, text := message) HMAC() SHALL be computed as defined in FIPS 198-1 using Crypto_Hash() as the underlying hash function H (this is also referred to as HMAC-SHA256()) and CRYPTO_HASH_LEN_BYTES is defined in Section 3.3, “Hash function (Hash)”.

3.5. Public Key Cryptography

Matter specifies the following scheme and parameters for public key cryptography.

3.5.1. Group

The public key cryptography of Matter relies on the group defined in the following mapping table.

Mapping (Version 1.0)

Matter public key cryptography SHALL be based on Elliptic Curve Cryptography (ECC) with the elliptic curve: `secp256r1` defined in Section 2.4.2 of [SEC 2](#). (This curve is also referred to as `NIST P-256` or `prime256v1` in [FIPS 186-5](#) and [NIST 800-186](#).)

`PrivateKey` is an opaque data type to hold either the private key or any handle or reference that allows other primitives to access the corresponding private key.

`PublicKey` is an opaque data type to hold the public key or any handle or reference that allows other primitives to access the corresponding public key. A public key is a point on the elliptic curve. (At places in the specification where public keys are to be explicitly transmitted, the format in which they are transmitted is specified.)

`int CRYPTO_GROUP_SIZE_BITS := 256`

`int CRYPTO_GROUP_SIZE_BYTES := 32`

`int CRYPTO_PUBLIC_KEY_SIZE_BYTES := (2 * CRYPTO_GROUP_SIZE_BYTES) + 1 = 65` is the size in bytes of the public key representation when encoded using the uncompressed public key format as specified in section 2.3 of [SEC 1](#).

```
struct {
    PublicKey publicKey;
    PrivateKey privateKey;
} KeyPair;
```

3.5.2. Key generation

`Crypto_GenerateKeyPair()` is the function to generate a key pair.

Function and description

KeyPair
Crypto_GenerateKeyPair()

Generates a key pair and returns a `KeyPair`.

Mapping (Version 1.0)

```
Crypto_GenerateKeypair() :=
    KeyPair ECCGenerateKeypair()
```

`ECCGenerateKeypair()` SHALL generate a key pair according to Section 3.2.1 of [SEC 1](#).

3.5.3. Signature and verification

`Crypto_Sign()` is used to sign a message, and `Crypto_Verify()` is used to verify a signature on a message.

These functions either generate or verify a signature of type **Signature** defined by the following mapping.

Mapping (Version 1.0)

```
struct {
    byte[CRYPTO_GROUP_SIZE_BYTES] r,
    byte[CRYPTO_GROUP_SIZE_BYTES] s
} Signature
```

3.5.3.1. Signature

Function and description

```
Signature
Crypto_Sign(
    PrivateKey privateKey,
    byte[] message)
```

Returns the signature of the **message** using the **privateKey**.

Mapping (Version 1.0)

```
Crypto_Sign(privateKey, message) :=
    Signature ECDSASign(dU := privateKey, M := message)
```

ECDSASign() SHALL be the ECDSA signature function as defined in Section 4.1 of [SEC 1](#) using **Crypto_Hash()** as the underlying hash **Hash()** function.

3.5.3.2. Signature verification

Function and description

```
boolean
Crypto_Verify(
    PublicKey publicKey,
    byte[] message,
    Signature signature)
```

Verifies the **signature** of the **message** using the **publicKey**, returns **TRUE** if the verification succeeds, **FALSE** otherwise.

Mapping (Version 1.0)

```
Crypto_Verify(publicKey, message, signature) :=
    boolean ECDSAVerify(QU := publicKey, M := message, S := signature)
```

ECDSAVerify() SHALL be the ECDSA signature verification function as defined in Section 4.1.4 of [SEC 1](#) using **Crypto_Hash()** as the underlying hash function **Hash()**; returns **TRUE** if the verification succeeds and **FALSE** otherwise.

3.5.4. ECDH

Crypto_ECDH() is used to compute a shared secret from the Elliptic Curve Diffie-Hellman (ECDH) protocol.

Function and description

```
byte[CRYPTO_GROUP_SIZE_BYTES]
Crypto_ECDH(
    PrivateKey myPrivateKey,
    PublicKey theirPublicKey)
```

Computes a shared secret using Elliptic Curve Diffie-Hellman.

Mapping (Version 1.0)

```
Crypto_ECDH(myPrivateKey, theirPublicKey) :=
    byte[CRYPTO_GROUP_SIZE_BYTES] ECDH(dU := myPrivateKey, QV := theirPublicKey)
```

The output of **ECDH()** SHALL be the serialization of the x-coordinate of the resultant point as defined in Section 3.3.1 of [SEC 1](#).

3.5.5. Certificate validation

Crypto_VerifyChain() is used to verify [Matter certificates](#).

Crypto_VerifyChainDER() is used to verify public key [X.509 v3](#) certificates in [X.509 v3 DER](#) format.

Function and description

```
boolean
Crypto_VerifyChain(MatterCertificate[] certificates)
```

Given [Matter certificates](#), **Crypto_VerifyChain()** performs all the necessary validity checks on **certificates**, taking in account that the notion of "current time" for the purposes of validation SHALL abide by the rules in [Section 3.5.6, "Time and date considerations for certificate path validation"](#).

Function and description

```
boolean
Crypto_VerifyChainDER(DERCertificate[] certificates)
```

Given a list of DER-encoded certificates in [RFC 5280](#) format, starting at the end-entity (leaf) certificate, and following the chain of trust towards the root, `Crypto_VerifyChainDER()` performs all the necessary validity checks on `certificates`.

The Validity period validation for the root and optional intermediate certificates is performed against the `notBefore` timestamp of the end-entity (leaf certificate) used as value for the current time.

Mapping (Version 1.0)

```
Crypto_VerifyChain(certificates) :=
    boolean verified
```

`verified` is `TRUE` if the `Matter certificates` are verified as prescribed by [RFC 5280](#).

```
Crypto_VerifyChainDER(certificates) :=
    boolean verified
```

`verified` is `TRUE` if the `certificates` are verified as prescribed by [RFC 5280](#).

The primitives as described above verify cryptographic integrity of the certificate chains. This specification imposes a number of additional constraints on certificates discussed below in sections on [Device Attestation Certificates](#), [Node Operational Certificates](#) and [Certificate Common Conventions](#).

3.5.6. Time and date considerations for certificate path validation

The Basic Path Validation algorithm in [RFC 5280](#) mandates the consideration of the "current time" against the validity period (`notBefore`, `notAfter` fields) when validating paths. The usage of "current time" assumes that such a time is available and correct, which is a strong assumption when considering some constrained devices or devices only locally reachable on a network in the absence of infrastructure to synchronize time against a global real-time reference.

When the `Crypto_VerifyChain` primitive is used, rather than overriding the Basic Path Validation algorithm of [RFC 5280](#), Nodes SHALL consider the following definition of "current time" that accounts for the possible lack of a real time reference:

- If a Node has a current real-time clock value which is trusted according to implementation-defined means to be accurate with regard to global real-time, whether using [Time Synchronization](#) features of this specification or other means, then it SHALL use that time;
- Otherwise, the current time SHALL be set to the last-known-good UTC time.

Upon failure to validate a certificate path, where the only reason for failure is an invalid validity period of a path element, a Node MAY apply a policy of its choice to determine whether to ignore

this failure and still consider the path valid.

3.5.6.1. Last Known Good UTC Time

Nodes SHALL maintain a stored Last Known Good UTC Time. This time is used as a fallback for cryptographic credentials expiry enforcement, if all other available time synchronization mechanisms fail.

The last known good UTC time SHALL be updated at commissioning and MAY be updated after a successful time synchronization, or by an embedded time in an OTA. Nodes SHOULD store a Last Known Good UTC Time value to persistent storage at least once a month. A Node's initial out-of-box Last Known Good UTC time SHALL be the compile-time of the firmware.

A Node MAY adjust the Last Known Good UTC Time backwards if it believes the current Last Known Good UTC Time is incorrect and it has a good time value from a trusted source. The Node SHOULD NOT adjust the Last Known Good UTC to a time before the later of:

- The build timestamp of its currently running software image
- The not-before timestamp of any of its operational certificates (see [Section 6.4.5, “Node Operational Credentials Certificates”](#)).

If a Node has used the Last Known Good UTC Time, it SHOULD recheck its security materials and existing connections if it later achieves time synchronization.

3.6. Data Confidentiality and Integrity

Symmetric block ciphers are used to provide message security.

All unicast and multicast messages between Nodes requiring protection for confidentiality and integrity with data origin authentication SHALL use Authenticated Encryption with Associated Data (AEAD) as primitive to protect those messages.

Mapping (Version 1.0)

Data confidentiality and integrity SHALL use the AES-CCM mode as defined in [NIST 800-38C](#) with the following parameters:

- `int CRYPTO_SYMMETRIC_KEY_LENGTH_BITS := 128` (this is the key length of the underlying block cipher in bits)
- `int CRYPTO_SYMMETRIC_KEY_LENGTH_BYTES := 16` (this is the key length of the underlying block cipher in bytes)
- `int CRYPTO_AEAD_MIC_LENGTH_BITS := 128` (this is the MIC length in bits)
- `int CRYPTO_AEAD_MIC_LENGTH_BYTES := 16` (this is the MIC length in bytes)
- `int CRYPTO_AEAD_NONCE_LENGTH_BYTES := 13`
- Key length SHALL be `CRYPTO_SYMMETRIC_KEY_LENGTH_BITS` bits.
- MIC length SHALL be `CRYPTO_AEAD_MIC_LENGTH_BITS` bits.
- The parameter `q` SHALL be 2 (length of encoding of maximum length) as specified in Appendix A.1 of [NIST 800-38C](#).
- The parameter `n` SHALL be `CRYPTO_AEAD_NONCE_LENGTH_BYTES` (length of nonce in bytes) as specified in Appendix A.1 of [NIST 800-38C](#).

`SymmetricKey` is an opaque data type to hold a symmetric block cipher key or any handle or reference that allows other primitives to access the corresponding key.

3.6.1. Generate and encrypt

Function and description

```
byte[lengthInBytes(P) + CRYPTO_AEAD_MIC_LENGTH_BYTES]
Crypto_AEAD_GenerateEncrypt(
    SymmetricKey K,
    byte[lengthInBytes(P)] P,
    byte[] A,
    byte[CRYPTO_AEAD_NONCE_LENGTH_BYTES] N)
```

Performs the generate and encrypt computation on payload `P` and the associated data `A` using the key `K` and a nonce `N`; the output contains the ciphertext and the tag truncated to `Tlen` bits (the encoding of the output depends on the mapping to the specific instance of the cryptographic primitive).

Mapping (Version 1.0)

```
Crypto_AEAD_GenerateEncrypt(K, P, A, N) :=
    byte[lengthInBytes(P) + CRYPTO_AEAD_MIC_LENGTH_BYTES] AES-CCM-GenerateEncrypt(K
:= K, P := P, A := A, N := N, Tlen := CRYPTO_AEAD_MIC_LENGTH_BITS)
```

AES-CCM-GenerateEncrypt() SHALL be the function described in Section 6.1 of [NIST 800-38C](#) with the counter generation function of Appendix A.3 of [NIST 800-38C](#) and the formatting function as defined in Appendix A.2 of [NIST 800-38C](#); returns the encoding of the ciphertext and the tag of length **Tlen** bits, as specified in Section 6.1 of [NIST 800-38C](#) as a byte array.

3.6.2. Decrypt and verify**Function and description**

```
{boolean success, byte[lengthInBytes(P)] payload}
Crypto_AEAD_DecryptVerify(
    SymmetricKey K,
    byte[lengthInBytes(P) + CRYPTO_AEAD_MIC_LENGTH_BYTES] C,
    byte[] A,
    byte[CRYPTO_AEAD_NONCE_LENGTH_BYTES] N)
```

Performs the decrypt and verify computation on the combined ciphertext and tag **C** and the associated data **A** using the key **K** and a nonce **N**. Note that the encoding of **C** depends on the mapping of the specific instance of the cryptography primitive.

This function has two outcomes:

- If tag verification succeeds, the **success** output is **TRUE** and the **payload** array contains the decrypted payload **P**.
- If tag verification fails, the **success** output is **FALSE** and the contents of the **payload** array is undefined.

Mapping (Version 1.0)

```
Crypto_AEAD_DecryptVerify(K, C, A, N) :=
  {boolean, byte[lengthInBytes(P)]} AES-CCM-DecryptVerify(K := K, C := C, A := A,
  N := N, Tlen := CRYPTO_AEAD_MIC_LENGTH_BITS)
```

AES-CCM-DecryptVerify() SHALL be the function described in Section 6.2 of [NIST 800-38C](#) with the counter generation function of Appendix A.3 of [NIST 800-38C](#) and the formatting function as defined in Appendix A.2 of [NIST 800-38C](#) and **C** SHALL be a byte array containing the ciphertext as specified in Section 6.2 of [NIST 800-38C](#).

- If tag verification succeeds, the **success** output is **TRUE** and the **payload** array contains the decrypted payload **P**.
- If tag verification fails, the **success** output is **FALSE** and the contents of the **payload** array is undefined.

3.7. Message privacy

Message privacy is implemented using a block cipher in CTR mode.

Mapping (Version 1.0)

Message privacy SHALL use the AES-CTR mode as defined in [NIST 800-38A](#) with the following parameters:

- **int CRYPTO_PRIVACY_NONCE_LENGTH_BYTES := 13**
- Key length SHALL be **CRYPTO_SYMMETRIC_KEY_LENGTH_BITS** bits.

3.7.1. Privacy encryption

Function and description

```
byte[lengthInBytes(M)]
Crypto_Privacy_Encrypt(
  SymmetricKey K,
  byte[lengthInBytes(M)] M,
  byte[CRYPTO_PRIVACY_NONCE_LENGTH_BYTES] N)
```

Performs the encryption of the message **M** using the key **K** and a nonce **N**; the output contains the data **M** encrypted.

Mapping (Version 1.0)

```
Crypto_Privacy_Encrypt(K, M, N) :=
    byte[lengthInBytes(M)]
    AES-CTR-Encrypt(K := K, P := M, N := N)
```

AES-CTR-Encrypt() SHALL be the encryption function described in Section 6.5 of [NIST 800-38A](#) with the sequence of counters **T** being generated according to the counter generation function of Appendix A.3 of [NIST 800-38C](#) using **N** and the value of **q = 2**; returns the encrypted message as a byte array.

3.7.2. Privacy decryption**Function and description**

```
byte[lengthInBytes(C)]
Crypto_Privacy_Decrypt(
    SymmetricKey K,
    byte[lengthInBytes(C)] C,
    byte[CRYPTO_PRIVACY_NONCE_LENGTH_BYTES] N)
```

Performs the decryption of **C** using the key **K** and a nonce **N**; the output **M** is the decryption of **C**

Mapping (Version 1.0)

```
Crypto_Privacy_Decrypt(K, C, N) :=
    byte[lengthInBytes(C)]
    AES-CTR-Decrypt(K := K, C := C, N := N)
```

AES-CTR-Decrypt() SHALL be the decryption function described in Section 6.5 of [NIST 800-38A](#) with the sequence of counters **T** being generated according to the counter generation function of Appendix A.3 of [NIST 800-38C](#) using **N** and the value of **q = 2**; returns the decrypted message as a byte array.

3.8. Key Derivation Function (KDF)

Matter specifies the following key derivation function to generate encryption keys.

Function and description

```

bit[len]
Crypto_KDF(
    byte[] inputKey,
    byte[] salt,
    byte[] info,
    int len)

```

Returns the key of **len** bits derived from **inputKey** using the **salt** and the **info**; **len** SHALL be a multiple of 8.

Mapping (Version 1.0)

```

Crypto_KDF(inputKey, salt, info, len) :=
    bit[len] HKDF-Expand(
        PRK := HKDF-Extract(salt := salt, IKM := inputKey),
        info := info, L := (len / 8))

```

HKDF-Extract SHALL be the HKDF-Extract primitive from [RFC 5869](#) section 2.2 using **Crypto_HMAC** (i.e. HMAC-SHA256) as the auxiliary **HMAC-Hash** function.

HKDF-Expand SHALL be the HKDF-Expand primitive from [RFC 5869](#) section 2.3 using **Crypto_HMAC** (i.e. HMAC-SHA256) as the auxiliary **HMAC-Hash** function.

The primitive returns a bit array of **len** bits.

When multiple keys of the same length are generated by a single KDF call, the following shorthand notation can be used:

```

Key1 || Key2 || Key3 = Crypto_KDF
(
    inputKey = inputKeyMaterial,
    salt = [],
    info = [],
    // 3 below matches number of keys expressed in concatenated output
    len = 3 * CRYPTO_SYMMETRIC_KEY_LENGTH_BITS
)

```

This is equivalent to the following:

```

Keys = Crypto_KDF
(
    inputKey = inputKeyMaterial,

```

```
    salt = [],  
    info = [],  
    len = 3 * CRYPTO_SYMMETRIC_KEY_LENGTH_BITS  
)
```

1. Set **Key1** to the **CRYPTO_SYMMETRIC_KEY_LENGTH_BITS** most significant bits of **Keys**.
2. Set **Key2** to the next **CRYPTO_SYMMETRIC_KEY_LENGTH_BITS** significant bits of **Keys**.
3. Set **Key3** to the **CRYPTO_SYMMETRIC_KEY_LENGTH_BITS** least significant bits of **Keys**.

3.9. Password-Based Key Derivation Function (PBKDF)

Matter specifies the following password-based key derivation function to compute a derived key from a cryptographically weak password.

Function and description
<pre>bit[len] Crypto_PBKDF(byte[] input, byte[] salt, int iterations, int len)</pre> Returns a value of len bits derived from the input using the salt and iterations iterations.
Type and description
<pre>STRUCTURE Crypto_PBKDFParameterSet</pre> Maintains the set of parameters exchanged between a Commissioner and a Commissionee during their pairing.
Mapping (Version 1.0)
<pre>int CRYPTO_PBKDF_ITERATIONS_MIN := 1000 int CRYPTO_PBKDF_ITERATIONS_MAX := 100000</pre> <pre>Crypto_PBKDF(input, salt, iterations, len) := bit[len] PBKDF2(P := input, S := salt, C := iterations, kLen := len)</pre> PBKDF2() SHALL be the HMAC-based PBKDF function with Crypto_HMAC(key := P, message := U[j-1]) as the auxiliary function HMAC as defined in Section 5.3 of NIST 800-132; it returns a bit array of len bits.

Mapping (Version 1.0)

```
Crypto_PBKDFParameterSet => STRUCTURE [ tag-order ]
{
    iterations [1] : UNSIGNED INTEGER [ range 32-bits ],
    salt [2] : OCTET STRING [ length 16..32 ],
}
```

- **iterations**: An integer value specifying the number of PBKDF2 iterations: `CRYPTO_PBKDF_ITERATIONS_MIN <= iterations <= CRYPTO_PBKDF_ITERATIONS_MAX`.
- **salt**: A random value per device of at least 16 bytes and at most 32 bytes used as the PBKDF2 salt.

3.10. Password-Authenticated Key Exchange (PAKE)

This protocol uses password-authenticated key exchange (PAKE) for the [PASE](#) protocol.

Mapping (Version 1.0)

Matter uses SPAKE2+ as described in [SPAKE2+](#) as PAKE with:

- The SPAKE2+ verifier is the Commissionee/Responder and the SPAKE2+ prover is the Commissioner/Initiator
- `Crypto_PBKDF()` as underlying PBKDF (see [Section 3.9](#), “Password-Based Key Derivation Function (PBKDF)”), with arguments as described in the definition of `Crypto_PAKEValues_Initiator`
- NIST P-256 elliptic curve as underlying group (see [Section 3.5.1](#), “Group”).
 - SPAKE2+ requires two additional points on the curve: **M** and **N**. The values of **M** and **N** are taken from the draft version 2 of the SPAKE2+ specification ([SPAKE2+](#)) and are listed below in compressed format (format defined in section 2.3 of [SEC 1](#)):
 - **M** = 02886e2f97ace46e55ba9dd7242579f2993b64e16ef3dcab95afd497333d8fa12f
 - **N** = 03d8bbd6c639c62937b04d997f38c3770719c629d7014d49a24b4f98baa1292b49
- `Crypto_Hash()` as underlying hash function (see [Section 3.3](#), “Hash function (Hash)”).
- `Crypto_HMAC()` as underlying HMAC function (see [Section 3.4](#), “Keyed-Hash Message Authentication Code (HMAC)”).
- `KDF(info, key, salt) := Crypto_KDF(key, salt, info, CRYPTO_HASH_LEN_BITS)` as underlying KDF function (see [Section 3.8](#), “Key Derivation Function (KDF)”).

Mapping (Version 1.0)

Crypto_PAKEValues_Initiator := (**w0**, **w1**) where **w0** and **w1** SHALL be computed as follows:

```
CRYPTO_W_SIZE_BYTES := CRYPTO_GROUP_SIZE_BYTES + 8
CRYPTO_W_SIZE_BITS := CRYPTO_W_SIZE_BYTES * 8

byte w0s[CRYPTO_W_SIZE_BYTES] || byte w1s[CRYPTO_W_SIZE_BYTES] =
    (byte[2 * CRYPTO_W_SIZE_BYTES])
    bit[2 * CRYPTO_W_SIZE_BITS]
    Crypto_PBKDF(passcode, salt, iterations, 2 * CRYPTO_W_SIZE_BITS)
byte w0[CRYPTO_GROUP_SIZE_BYTES] = w0s mod p
byte w1[CRYPTO_GROUP_SIZE_BYTES] = w1s mod p
```

where:

- **mod** is the mathematical modulo operation and **||** is the string concatenation or split operator.
- **passcode**, is the Passcode defined in [Section 5.1.1.6, “Passcode”](#), serialized as little-endian over 4 octets. For example, passcode **18924017** would be encoded as the octet string **f1:c1:20:01** and the passcode **00000005** would be encoded as the octet string **05:00:00:00**.
- **p** is the order of the underlying elliptic curve.
- Both **w0s** and **w1s** SHALL have a length equal to **(CRYPTO_GROUP_SIZE_BYTES + 8)**.
- **salt** and **iterations** are extracted from the **Crypto_PBKDFParameterSet** values.
- The pair (**w0,w1**) is also referred to as Commissioner PAKE input

Mapping (Version 1.0)

Crypto_PAKEValues_Responder := (**w0**, **L**) where **w0** and **L** SHALL be computed as follows:

```
byte w0s[CRYPTO_W_SIZE_BYTES] || byte w1s[CRYPTO_W_SIZE_BYTES] =
    (byte[2 * CRYPTO_W_SIZE_BYTES])
    bit[2 * CRYPTO_W_SIZE_BITS]
    Crypto_PBKDF(passcode, salt, iterations, 2 * CRYPTO_W_SIZE_BITS)
byte w0[CRYPTO_GROUP_SIZE_BYTES] = w0s mod p
byte w1[CRYPTO_GROUP_SIZE_BYTES] = w1s mod p
byte L[CRYPTO_PUBLIC_KEY_SIZE_BYTES] = w1 * P
```

where:

- **passcode**, is the Passcode defined in [Section 5.1.1.6, “Passcode”](#).
- **p** is the order of the elliptic curve to be used.
- Both **w0s** and **w1s** SHALL have a length equal to (**CRYPTO_GROUP_SIZE_BYTES** + 8).
- **salt** and **iterations** are extracted from the **Crypto_PBKDFParameterSet** values.
- **P** is the generator of the underlying elliptic curve.
- When the computation of **Crypto_PAKEValues_Responder** is done, fields **w0** and **L** SHALL be stored in the Responder and **w1** SHALL NOT be stored in the Responder.
- The pair (**w0**, **L**) is also referred to as Commissionee PAKE input or verification value

3.10.1. Computation of pA**Mapping (Version 1.0)**

```
Crypto_pA(Crypto_PAKEValues_Initiator) :=
    byte pA[CRYPTO_PUBLIC_KEY_SIZE_BYTES]
```

pA is in uncompressed public key format as specified in section 2.3 of [SEC 1](#). **pA** SHALL be computed as specified in [SPAKE2+](#).

3.10.2. Computation of pB**Mapping (Version 1.0)**

```
Crypto_pB(Crypto_PAKEValues_Responder) :=
    byte pB[CRYPTO_PUBLIC_KEY_SIZE_BYTES]
```

pB is in uncompressed public key format as specified in section 2.3 of [SEC 1](#). **pB** SHALL be computed as specified in [SPAKE2+](#).

3.10.3. Computation of transcript TT

Mapping (Version 1.0)

Crypto_Transcript(PBKDFParamRequest, PBKDFParamResponse, pA, pB) :=
byte[] TT

Crypto_Transcript() SHALL compute TT as specified in SPAKE2+ with:

byte ContextPrefixValue [26] = {
0x43, 0x48, 0x49, 0x50, 0x20, 0x50, 0x41, 0x4b, 0x45, 0x20, 0x56, 0x31, 0x20, 0x43,
0x6f, 0x6d,
0x6d, 0x69, 0x73, 0x73, 0x69, 0x6f, 0x6e, 0x69, 0x6e, 0x67
} // "CHIP PAKE V1 Commissioning" - The usage of CHIP here is intentional and due to
implementation in the SDK before the name change, should not be renamed to Matter.

Context := Crypto_Hash(ContextPrefixValue || PBKDFParamRequest || PBKDFParamResponse)

TT :=
lengthInBytes(Context) || Context ||
0x0000000000000000 || 0x0000000000000000 ||
lengthInBytes(M) || M ||
lengthInBytes(N) || N ||
lengthInBytes(pA) || pA ||
lengthInBytes(pB) || pB ||
lengthInBytes(Z) || Z ||
lengthInBytes(V) || V ||
lengthInBytes(w0) || w0 ||

Z and V SHALL be computed from pA and pB as specified in SPAKE2+.

Note the two 0x0000000000000000 null-lengths indicate that no identities are present and each null-lengths is 8 bytes wide since it is specified by the SPAKE2+ specification that lengths are eight-byte little-endian numbers. The SPAKE2+ specification indicates that we **must** include these length fields.

Note in case PBKDFParamRequest and PBKDFParamResponse messages are not exchanged, they SHALL be replaced by empty strings in the Context computation.

3.10.4. Computation of cA, cB and Ke

Mapping (Version 1.0)

```
Crypto_P2(TT, pA, pB) :=  
  {byte cA[CRYPTO_HASH_LEN_BYTES],  
   byte cB[CRYPTO_HASH_LEN_BYTES],  
   byte Ke[CRYPTO_HASH_LEN_BYTES/2]}
```

Crypto_P2() SHALL compute **cA**, **cB** and **Ke** as specified in [SPAKE2+](#) with **cA** := CRYPTO_HMAC(KcA,pB) and **cB** := CRYPTO_HMAC(KcB,pA).

Chapter 4. Secure Channel

4.1. General Description

The Secure Channel and Message Layer provides a consistent networking service substrate to allow Nodes to communicate securely with one another.

During commissioning and unicast communication, a [discovery](#) mechanism is provided to determine peer IPv6 addresses and operational parameters. [Secure session establishment](#) mechanisms are provided using either [certificates \(CASE\)](#) or [shared passcodes \(PASE\)](#).

4.1.1. Messages

Communication is performed using messages. Messages are either secured or unsecured.

Each message has a [Session Type](#) and [Session ID](#) in order to identify whether it is secured and how it is to be decrypted and authenticated if it is. Each message has a [Message Counter](#) field in order to uniquely identify the message for the purposes of security and duplicate detection.

Operational communication is defined as traffic that uses the secured Matter message format between commissioned nodes over an IP transport. All operational communication has [message security](#) enabled. Operational communication between Nodes can be either unicast or multicast.

Unsecured communication is strictly limited to:

- [Discovery](#), which does not use the Matter message format.
- [User Directed Commissioning \(UDC\)](#), which uses unsecured messages to initiate the commissioning process.
- [Session establishment](#), which uses unsecured messages to establish a [CASE](#) or [PASE](#) session.

4.1.1.1. Message Types

Messages are defined as either a *control message* or *data message*. Most messages are data messages. Control messages are reserved for internal protocols such as [MCSP](#) to initialize security. Both message types are identical in format, but use separate message counter domains so they can operate securely over the same security key.

4.1.1.2. Message Transports

Messages are of [finite size](#) and are sent as individual packets over the supported transports:

- UDP transports each message as a separate datagram. Messages support a basic reliability protocol called [MRP](#) for use when the underlying transport (in this case UDP) doesn't provide such features.
- TCP transports each message with a length prepended, performing segmentation and reassembly as needed.
- [BTP](#) transports each message over BLE as a separate SDU, performing segmentation and

- reassembly as needed. [BTP](#) is provided as a transport protocol for commissioning using a BLE interface.
- [PAFTP](#) transports each message over [WFA-USD](#) as a separate SDU, performing segmentation and reassembly as needed. [PAFTP](#) is provided as a transport protocol for commissioning using a Wi-Fi interface.
 - [NFC Transport Layer \(NTL\)](#) transports each message by encapsulating it into one or more chained ISO/IEC 7816-4 command-response pairs. [NFC Transport Layer \(NTL\)](#) is provided as a transport protocol for commissioning using an NFC interface.
 - [Thread Commissioning Protocol](#) transports each message over the native Thread commissioning protocol using the Relay commands. It is used for commissioning using the Thread interface. From the commissionee’s standpoint, it looks exactly like the commissioning protocols implemented using IP. From the commissioner’s side, the messages are encapsulated within the Relay commands as defined by the Thread Specification.

TCP and [MRP \(UDP with added reliability\)](#) are provided as transport protocols for operational messaging.

Transports based on [PAFTP](#), [NTL](#), [Thread Commissioning Protocol](#), and [BTP](#) are provided as temporary transports for the purpose of commissioning devices onto an operational IP network. They MAY additionally be used for other temporary non-operational messaging, i.e., exceptional conditions, maintenance, debugging and efforts focused on getting the Device into normal operation on an IP network. Examples of such non-operational messaging include [Network Recovery](#), [OTA](#) and retrieving [Diagnostic Logs](#). Support of these transports for such temporary non-operational messaging is optional for both Commissioners/Administrators and Devices. These transports are not intended for operational messaging.

4.1.1.3. Message Exchanges

Messages provide an [Exchange Layer](#) to track related messages that make up small, discrete transactions. The Exchange Layer provides this transaction tracking facility to the [Interaction Model Layer](#) above, providing a means to multiplex multiple such concurrent transactions over a given underlying session. The Exchange Layer also integrates the [Message Reliability Protocol \(MRP\)](#) as a service for use over UDP transports. This Message Layer architecture is shown below in [Figure 6](#), “[Message Layer Stack](#)”:




Figure 6. Message Layer Stack

4.2. IPv6 Reachability

This section describes IPv6 network configuration requirements to enable IPv6 reachability between Nodes. As described in [Section 2.3, “Network Topology”](#), a Matter network may be composed of one or more IPv6 networks.

In a single network configuration, all Matter Nodes are attached to the same IPv6 link. A single network configuration may consist of a single bridged Wi-Fi / Ethernet network where all nodes attached to that network are part of the same broadcast domain. When all Matter Nodes are attached to the same Wi-Fi / Ethernet network, link-local IPv6 addressing is sufficient - no additional IPv6 network infrastructure is required.

In a multiple network configuration, a Matter network is composed of (typically one) infrastructure network and one or more stub networks. Unlike an infrastructure network, stub networks do not serve as transit networks. Typically, the infrastructure network is a bridged Wi-Fi / Ethernet network and the Thread networks are stub networks. A stub router connects a stub network to an infrastructure network and provides IPv6 reachability between the two networks.

4.2.1. Stub Router Behavior

A stub router SHALL implement [[draft-lemon-stub-networks](#)]. In a multiple network configuration, both the infrastructure and stub networks require routable IPv6 addresses to communicate across networks. A routable IPv6 address SHALL have global scope (e.g. GUA or ULA) [[RFC 4007](#)] and SHALL be constructed out of a prefix advertised as on-link. If there is no routable prefix on a given network, the stub router SHALL provide its own routable prefix. Note that Thread’s “on-mesh prefix” is equivalent to Wi-Fi / Ethernet’s “on-link prefix”.

Stub routers SHALL advertise reachability to all routable prefixes on the adjacent network. A stub router connecting a Thread network SHALL advertise reachability to all of the Thread network’s routable prefixes to the adjacent infrastructure network using Route Information Options [[RFC 4191](#)] contained in Router Advertisements [[RFC 4861](#)]. That same stub router SHALL also advertise reachability to all of the infrastructure network’s routable prefixes to the adjacent Thread network in the Thread Network Data [[Thread specification](#)].

4.2.2. Matter Node Behavior

Matter Nodes SHALL configure a link-local IPv6 address. In addition, Nodes SHALL support configuration of at least three routable IPv6 addresses (in addition to the link-local and, in the case of Thread, mesh-local addresses). On a Wi-Fi / Ethernet interface, ICMPv6 Router Advertisement (RA) messages advertise prefixes for use on the link [[RFC 4861](#)]. On a Thread interface, the Thread Network Data contains prefixes for use on the link [[Thread specification](#)]. If a Node receives an on-link prefix that allows autonomous configuration on a given interface and the Node has fewer than three routable IPv6 addresses configured, the Node SHALL autonomously configure an IPv6 address out of that prefix.

Matter Nodes SHALL also configure routes to adjacent networks. On Wi-Fi / Ethernet networks, Nodes SHALL process Route Information Options [[RFC 4191](#)] and configure routes to IPv6 prefixes assigned to stub networks via stub routers. Wi-Fi / Ethernet interfaces SHALL support maintaining at least 16 different routes configured using Route Information Options. On Thread networks,

Nodes SHALL route according to routing information provided in the Thread Network Data [[Thread specification](#)]. Thread devices SHALL support as many routes as can be encoded in the Thread Network Data.

Matter Nodes SHALL support a number of IPv6 neighbor cache entries at least as large as the [number of supported CASE sessions](#) plus the number of supported routes.

4.3. Discovery

This section describes Service Advertising and Discovery for Matter.

Service Advertising and Discovery is used within Matter in the following contexts:

- [Commissionable Node Discovery](#)
- [Operational Discovery](#)
- [Commissioner Discovery](#)
- [User Directed Commissioning](#)

Service Advertising and Discovery for Matter uses IETF Standard DNS-Based Service Discovery (DNS-SD) [[RFC 6763](#)]. Matter requires no modifications to IETF Standard DNS-SD.

Using DNS-SD means that both the unicast IPv6 address and port of the offered service are discovered, freeing Matter from requiring a single preallocated fixed port. This also makes it possible to run multiple instances of Matter software on a single device, because each instance has its own dynamically allocated port, instead of conflicting over attempting to use the same preallocated fixed port.

On Wi-Fi and Ethernet networks today, DNS-SD [[RFC 6763](#)] uses Multicast DNS [[RFC 6762](#)] for zero-configuration operation.

Since Matter protocols must support IPv6 at a minimum, Matter software discovering other Matter instances SHALL process DNS AAAA (IPv6 address) records, but also MAY process DNS A (IPv4 address) records.

Because of this, where feasible in the underlying service discovery API, Matter software advertising the availability of a service SHOULD indicate that announcements and answers for this service need include only IPv6 address records, not IPv4 address records. On a general-purpose dual-stack host that supports both IPv4 and IPv6, this can be achieved by having Matter-related SRV records reference a Matter-specific target hostname that has only IPv6 address records attached. This allows a general-purpose dual-stack host to offer discoverable IPv4 addresses for legacy client software that still requires IPv4, while offering optimized IPv6-only address discovery for Matter purposes.

Similarly, since Matter does not use IPv4, Matter software discovering other Matter instances SHOULD NOT expect any IPv4 addresses included in responses.

These two items address the *content* of service discovery messages. When using Multicast DNS similar efficiency questions arise related to the *delivery* of those service discovery messages, sent over IPv4, IPv6, or both.

Where supported in the underlying service discovery API, Matter software using Multicast DNS to advertise the availability of a service SHOULD indicate that announcements and answers for this service need only be performed over IPv6.

Similarly, where supported in the underlying service discovery API, Matter application software using Multicast DNS to issue service discovery queries SHOULD indicate that these queries need only be performed over IPv6.

These optimizations reduce both the size and the number of multicast packets, which is particularly beneficial on Wi-Fi networks. A Matter device that only supports IPv6 gets these optimizations automatically, simply by virtue of not supporting IPv4 at all.

For Thread mesh networks, where excessive use of multicast would be detrimental [RFC 7558], DNS-SD uses Unicast DNS instead, leveraging capabilities of the Thread Service Registry on the Thread Border Router [draft-lemon-stub-networks].

Conceptually, the DNS-SD [RFC 6763] information being communicated is identical to when Multicast DNS [RFC 6762] is being used, except that the information is communicated in unicast packets to and from a designated Service Registry, rather than being communicated in multicast packets to and from every other Node in the same broadcast domain.

Using Service Registration Protocol [RFC 9665] and an Advertising Proxy [AdProx] running on the Thread Border Router, Matter Nodes on a Thread mesh can be discovered by other Matter Nodes on an adjacent Ethernet or Wi-Fi link, without the cost of using multicast on the Thread mesh. All Thread-connected Matter Nodes SHALL implement Service Registration Protocol.

Thread Border Routers advertise available SRP servers in the Thread Network Data [Thread specification]. Thread devices SHALL register their services using an available SRP server [Thread specification].

When Matter Nodes issue short-lived requests to other Matter Nodes, the response is sent back to the source IPv6 address and port of the request. When Matter Nodes issue long-lived requests to other Matter Nodes, by the time the response is generated the requester may have changed IPv6 address or port, so the responder may have to discover the current IPv6 address and port of the initiator in order to send the response.

A Thread Border Router SHALL implement DNS-SD Discovery Proxy [RFC 8766] to enable clients on the Thread mesh (e.g., other Nodes) to discover services (e.g., Matter Nodes) advertised using Multicast DNS on an adjacent Ethernet or Wi-Fi link, also without the cost of using multicast on the Thread mesh [draft-lemon-stub-networks]. For short-lived instantaneous queries, these queries can be performed using unicast DNS over UDP to the DNS-SD Discovery Proxy. For long-lived queries with ongoing change notification, DNS Push Notifications [RFC 8765] with DNS Stateful Operations [RFC 8490] allows clients on the Thread mesh to be notified of changes to the set of discovered services without expensive polling.

In principle, the Thread mesh Service Registry can be run on any capable Node(s) within (or even outside) the Thread mesh, though in practice the Thread Border Router is an attractive candidate to offer the Service Registry. Thread Border Router devices are typically AC-powered, and typically have more capable CPUs with greater flash storage and RAM than more constrained battery-powered Thread Nodes. Matter devices on Thread are dependent on Thread providing reliable service

for those Thread devices on the Thread mesh. This is similar to how Matter devices on Wi-Fi are dependent on the Wi-Fi access point (AP) providing reliable service for those Wi-Fi devices using that Wi-Fi access point.

4.3.1. Commissionable Node Discovery

The Matter protocol family, currently, supports UDP for Matter commissioning of Commissionees already on the customer's IP network (such as Ethernet-connected Nodes, or Wi-Fi Nodes already associated to the Wi-Fi network via other means). Other transport protocols such as TCP MAY be supported in the future.

For these Commissionees, Matter Commissionable Node Discovery is performed using IETF Standard DNS-Based Service Discovery (DNS-SD) [RFC 6763] as described below.

For Matter Commissionable Node Discovery in the *already-on-network* case, the DNS-SD instance name SHALL be a dynamic, pseudo-randomly selected, 64-bit temporary unique identifier, expressed as a fixed-length sixteen-character hexadecimal string, encoded as ASCII (UTF-8) text using capital letters, e.g., `DD200C20D25AE5F7`. A new instance name SHALL be selected when the Node boots. A new instance name SHALL be selected whenever the Node enters Commissioning mode. A new instance name MAY be selected at other times, as long as the instance name does not change while the Node is in commissioning mode.

A commissionable Node that is already connected to an IP-bearing network SHALL only make itself discoverable on the IP network and SHALL use the relevant DNS-SD service (`_matterc._udp`) described below.

If a Node is connected to multiple IP-bearing networks, and it receives either the `OpenCommissioningWindow` or the `OpenBasicCommissioningWindow` command, it MAY make itself discoverable only on the network on which the command was received. Otherwise, a commissionable Node that is connected to multiple IP-bearing networks SHALL make itself discoverable on all of its connected IP-bearing networks.

The Matter Commissionable Node Discovery DNS-SD instance name SHALL be unique within the namespace of the local network. This namespace is the `.local` link-local namespace of the Ethernet and Wi-Fi links [RFC 6762], or the unicast domain selected by the Thread Border Router for devices on the Thread mesh.

In the rare event of a collision in the selection of the 64-bit temporary unique identifier, the existing DNS-SD name conflict detection mechanism will detect this collision, and a new pseudo-randomly selected 64-bit temporary unique identifier SHALL be generated by the Matter Commissionee that is preparing for commissioning. Name conflict detection is described in Section 9 ("Conflict Resolution") of the Multicast DNS specification [RFC 6762] and Section 2.4.3.1 ("Validation of Adds") of the Service Registration Protocol specification [RFC 9665].

The DNS-SD service type [RFC 6335] for Matter Commissionable Node Discovery is `_matterc._udp`.

For link-local Multicast DNS the service domain SHALL be `local`. For Unicast DNS such as used on Thread the service domain SHALL be as configured automatically by the Thread Border Router.

4.3.1.1. Host Name Construction

For DNS-SD a target host name is required, in addition to the instance name. The target host name SHALL be constructed using one of the available link-layer addresses, such as a 48-bit device MAC address (for Ethernet and Wi-Fi) or a 64-bit MAC Extended Address (for Thread) expressed as a fixed-length twelve-character (or sixteen-character) hexadecimal string, encoded as ASCII (UTF-8) text using capital letters, e.g., **B75AFB458ECD**.<domain>. In the event that a device performs MAC address randomization for privacy, then the target host name SHALL use the privacy-preserving randomized version and the hostname SHALL be updated in the record every time the underlying link-layer address rotates. Note that it is legal to reuse the same hostname on more than one interface, even if the underlying link-layer address does not match the hostname on that interface, since the goal of using a link-layer address is to ensure local uniqueness of the generated hostname. If future link layers are supported by Matter that do not use 48-bit MAC addresses or 64-bit MAC Extended Address identifiers, then a similar rule will be defined for those technologies.

4.3.1.2. Extended Discovery

A Matter Commissionee that advertises Commissionable Node Discovery service records is not necessarily in a state that will allow Commissioning (this state is referred to below as "in Commissioning Mode"). While [Section 5.4.2.3, "Announcement Duration"](#) is limited for some forms of device advertisement, a Matter device MAY advertise Matter Commissionable Node Discovery service records for longer periods, possibly permanently. Advertising Commissionable Node Discovery when not in Commissioning Mode is referred to here as Extended Discovery. Extended Discovery is allowed only for DNS-SD advertisements and not for the other forms of [Device Discovery](#) such as [BLE Commissioning Discovery](#).

To protect customer privacy on public networks, a Matter Commissionee SHALL provide a way for the customer to set a timeout on Extended Discovery, or otherwise disable Extended Discovery. The default behavior for Commissionable Node Discovery SHOULD default to limiting announcement as defined in [Section 5.4.2.3, "Announcement Duration"](#) unless the Manufacturer wishes to enable longer periods for specific use cases.

4.3.1.3. Commissioning Subtypes

The following subtypes for Matter Commissionable Node Discovery are defined:

1. **_L<dddd>**, where **<dddd>** provides the full 12-bit [discriminator](#), encoded as a variable-length decimal number in ASCII text, omitting any leading zeroes.
2. **_S<dd>**, where **<dd>** provides the upper 4 bits of the [discriminator](#), encoded as a variable-length decimal number in ASCII text, omitting any leading zeroes.
3. **_V<dddddd>**, where **<dddddd>** provides the 16-bit [Vendor ID](#), encoded as a variable-length decimal number in ASCII text, omitting any leading zeroes.
4. **_T<ddd>**, where **<ddd>** provides the [device type](#) identifier for the device, encoded as a variable-length decimal number in ASCII (UTF-8) text, omitting any leading zeroes. In case the device combines multiple device types, the manufacturer SHOULD choose the device type identifier of the primary function of the device for which the device wishes to be discoverable.
5. **_CM**, which represents "currently in Commissioning Mode" (due to any method, for example, a factory new device that has just been put into commissioning mode by the user, or an already-

commissioned device which has just received the [OpenCommissioningWindow](#) command).

The *long discriminator* subtype (e.g., [_L840](#)) enables filtering of results to find only Commissionees that match the full discriminator code, as provided in the [onboarding payload](#).

The *short discriminator* subtype (e.g., [_S3](#)) enables filtering of results to find only Commissionees that match the upper 4 bits of the discriminator code, as provided in the [manual pairing code](#).

The optional *Vendor ID* subtype (e.g., [_V123](#)) enables a vendor-specific app to achieve filtering of results to find only Nodes that match that Vendor ID.

The *Commissioning Mode* subtype (e.g., [_CM](#)) enables filtering of results to find only Nodes that are currently in Commissioning Mode. Note that the subtype is [_CM](#) regardless of whether the TXT record for commissioning mode is set to 1 ([CM=1](#)), 2 ([CM=2](#)) or 3 ([CM=3](#)). A Commissionee that is not in commissioning mode ([CM=0](#)) SHALL NOT publish this subtype.

The optional *device type* subtype (e.g., [_T10](#)) enables filtering of results to find only Nodes that match the device type, generally used for the [User-Initiated Beacon Detection, Not Yet Commissioned Device](#) and the [User-Initiated Beacon Detection, Already Commissioned Device](#) use cases.

In the event that a vendor-specific app wishes to show the user only some of that vendor's Commissionees awaiting commissioning but not all of them, any desired filtering logic (based upon arbitrary criteria, not only Product ID) MAY be implemented within that vendor's proprietary commissioning app.

4.3.1.4. TXT Records

After discovery, IPv6 addresses are returned in the AAAA records and key/value pairs are returned in the DNS-SD TXT record.

Nodes SHALL publish AAAA records for all available IPv6 addresses upon which they are willing to accept Matter commissioning messages.

TXT records available for Commissionable Node Discovery include the common TXT record key/value pairs defined in [Section 4.3.4, "Common TXT Key/Value Pairs"](#).

Commissioners SHALL silently ignore TXT record keys that they do not recognize. This is to facilitate future evolution of this specification without breaking backwards compatibility with existing Commissioners that do not implement the new functionality.

The following subsections describe key/value pairs that are defined specifically for Commissionable Node discovery.

4.3.1.5. TXT key for discriminator (D)

The key [D](#) SHALL provide the full 12-bit [discriminator](#) for the Commissionable Node and SHALL be present in the DNS-SD TXT record.

The discriminator value SHALL be encoded as a variable-length decimal number in ASCII text, with up to four digits, omitting any leading zeroes.

Any key **D** with a value mismatching the aforementioned format SHALL be silently ignored.

As an example, value **D=840** would indicate that this Commissionable Node has decimal long discriminator **840**. When needed, the upper 4 bits of the discriminator provided by the [manual pairing code](#) can be algorithmically derived from the full discriminator.

4.3.1.6. TXT key for Vendor ID and Product ID (**VP**)

The optional key **VP**, if present, MAY provide [Vendor ID](#) and [Product ID](#) information of the device.

A vendor MAY choose not to include it at all, for privacy reasons.

If the **VP** key is present then it MAY take two forms:

1. **VP=123** gives Vendor ID
2. **VP=123+456** gives Vendor ID + Product ID

The Vendor ID and Product ID SHALL both be expressed as variable-length decimal numbers, encoded as ASCII text, omitting any leading zeroes, and of maximum length of 5 characters each to fit their 16-bit numerical range.

If the Product ID is present, it SHALL be separated from the Vendor ID using a '+' character.

If the **VP** key is present without a Product ID, the value SHALL contain only the Vendor ID alone, with no '+' character.

If the **VP** key is present, the value SHALL contain at least the Vendor ID.

If the **VP** key is present, it SHALL NOT have a missing or empty value.

4.3.1.7. TXT key for commissioning mode (**CM**)

The key **CM** (Commissioning Mode) SHALL indicate whether or not the publisher of the record is currently in Commissioning Mode and available for immediate commissioning. When in commissioning mode, the value associated with the **CM** key indicates the source of the Passcode.

Five situations are legal:

1. The absence of key **CM** SHALL imply a value of 0 (**CM=0**).
2. The key/value pair **CM=0** SHALL indicate that the publisher is not currently in Commissioning Mode.
3. The key/value pair **CM=1** SHALL indicate that the publisher is currently in Commissioning Mode and requires use of a Passcode for commissioning provided by the Commissionee (e.g., embedded in Onboarding Material which is printed on device, on a label or displayed on screen), such as when the device is in a factory-new state or when the [OpenBasicCommissioningWindow](#) command has been used to enter commissioning mode.
4. The key/value pair **CM=2** SHALL indicate that the publisher is currently in Commissioning Mode and requires use of a dynamically generated Passcode for commissioning corresponding to the verifier that was passed to the device using the [OpenCommissioningWindow](#) command.

5. The key/value pair **CM=3** SHALL indicate that the publisher is currently in **Joint Fabric** Commissioning Mode and requires use of a dynamically generated Passcode for commissioning corresponding to the verifier that was passed to the device using the **OpenJointCommissioningWindow** command. The rationale for this value is to allow a Commissioner to distinguish between a Commissionee that is in **Joint Fabric** Commissioning Mode and one that is in Commissioning Mode **CM=2**. **Joint Fabric** commissioning requires the execution of additional mutual verification steps only performed when both the Commissionee and Commissioner are instructed to use the **Joint Commissioning Method**.

A key value of 2 or 3 MAY be used to disambiguate collisions of discriminators between uncommissioned Nodes and commissioned Nodes announcing after a commissioning window was opened. A key value of 2 or 3 serves as a hint to Commissioners to possibly expect multiple Nodes with the same discriminator (see [Section 4.3.1.5, “TXT key for discriminator \(D\)”](#)), and to instruct the user to enter the Onboarding Payload **presented** by another Administrator rather than a code provided by the Commissionee.

Since **Extended Discovery** can be disabled by the customer, a key value of 0 may not ever be returned by a publisher. When Extended Discovery is disabled and the publisher is not in commissioning mode, then the publisher will not respond to Commissionable Node Discovery.

4.3.1.8. TXT key for device type (DT)

The optional key **DT** MAY be used to provide the publisher’s **Primary Device Type**. If present, it SHALL be encoded as a variable-length decimal number in ASCII text, omitting any leading zeroes.

For example, the **DT=10** key/value pair would indicate that the primary device type is 10 (0x000A), which is the device type identifier for a Door Lock.

4.3.1.9. TXT key for device name (DN)

The optional key **DN** MAY provide a device advertisement name. If present, it SHALL be encoded as a valid UTF-8 string with a maximum length of 32 bytes (matching the maximum length of the Node-Label string in the **Basic Information Cluster**).

When provided, the source of this value SHALL be editable by the user with use clearly designated as being for on-network advertising and MAY be the value stored in the NodeLabel attribute of the **Basic Information Cluster** of the Node.

To protect customer privacy on public networks, if a Commissionee supports this key/value pair, then the Commissionee SHALL provide a way for the customer to disable its inclusion.

A Commissionee SHOULD NOT include this field unless doing so for specific use cases which call for it.

For example, the **DN=Living Room** key/value pair indicates that the advertisement name specified by the user is 'Living Room'.

4.3.1.10. TXT key for rotating device identifier (RI)

The optional key **RI** MAY provide a **Rotating Device Identifier**.

If present, the value associated with the **RI** key SHALL contain the octets of the Rotating Device Identifier octet string encoded as the concatenation of each octet's value as a 2-digit uppercase hexadecimal number.

The resulting ASCII string SHALL NOT be longer than 100 characters, which implies a Rotating Device Identifier of at most 50 octets.

4.3.1.11. TXT key for pairing hint (**PH**)

The optional key **PH** MAY provide a *pairing hint*.

If present, it SHALL be encoded as a variable-length decimal number in ASCII text, omitting any leading zeroes.

The pairing hint represents a base-10 numeric value for a bitmap of methods supported by the Commissionee in its current state for putting it in Commissioning Mode.

For example, the **PH=5** key/value pair represents a hint value with bits 0 and 2 are set.

This value MAY change during the lifecycle of the device.

For example, a device may have a value with bit 0 (Power Cycle) set and bit 2 (Administrator app) unset when in a factory reset state, and then have a value with bit 0 unset and bit 2 set after it has been Commissioned.

The bitmap of methods is defined in [Table 5, "Pairing and Reset Hint Values"](#).

If the Commissionee has enabled [Extended Discovery](#), then it SHALL include the key/value pair for **PH** in the DNS-SD TXT record when not in Commissioning Mode (**CM=0**).

This key/value pair MAY be returned when in Commissioning Mode (**CM=1**).

If the Commissioner does not recognize this value, for example, if the value indicates bit indices defined in a newer version of this specification than the version which the Commissioner implements, then the Commissioner MAY utilize the bits that it does understand and MAY utilize additional data sets available for assisting the user. For example, when a Vendor ID and Product ID are available to the Commissioner, the [Section 11.23, "Distributed Compliance Ledger"](#) may also provide a URL for the Device User Guide which can contain additional information to help in Commissioning this Commissionee.

Some of the pairing hints SHALL or MAY require additional information to be encoded for proper expression of their meaning. This is accomplished with the **PI** TXT key, described in a following section. Dependency on usage of the **PI** key is expressed by the **Instruction Dependency** column in the table below. For some entries, the use of the **PI** key is mandatory (indicated as **M**), while for some other entries, the use of the **PI** key is optional (indicated as **O**). For the remaining entries the use of the **PI** key does not apply (indicated as **no**).

The fields in the bitmap are defined for values of the **PH** key in the [Pairing and Reset Hint Values](#) table below. This bitmap describes methods for putting the device into a specific target state. It is referenced here to describe methods for putting the device into Commissioning Mode, and it is referenced elsewhere to describe methods for putting the device into other states, such as the Factory

Reset state in the DCL DeviceModel Schema's [FactoryResetStepsHint](#) field.

Table 5. Pairing and Reset Hint Values

Bit index	Name	Instruction Dependency	Description
0	Power Cycle	no	The Device will automatically enter the target state upon power cycle (unplug/re-plug, remove/re-insert batteries). When used with the target state of Commissioning Mode for a device that is in a factory reset state, this bit SHALL be set to 1 for devices using Standard Commissioning Flow , and set to 0 otherwise.
1	Device Manufacturer URL/App	no	To put the device into the desired state, the user can be sent to the URL specified in the CommissioningCustomFlowUrl of the DeviceModel schema entry indexed by the Vendor ID and Product ID (e.g., as found in the announcement) in the Distributed Compliance Ledger . This URL could lead to a page with instructions, or to an app with instructions. When used with the target state of Commissioning Mode for a device that is in a factory reset state, this bit SHALL be set to 1 for devices requiring Custom Commissioning Flow before they can be available for Commissioning by any Commissioner.
2	Administrator	no	The Device has been commissioned. Any Administrator that commissioned the device provides a user interface that may be used to put the device into Commissioning Mode.
3	Settings menu on the Node	no	The settings menu on the Device provides instructions to put it into the desired state.
4	Custom Instruction	M	The PI key/value pair SHALL describe a custom way to put the Device into the target state. This Custom Instruction option is NOT recommended for use by a Device that does not have knowledge of the user's language preference.
5	Device Manual	no	The Device Manual provides special instructions to put the Device into the target state (see Section 11.23.6.16, "UserManualUrl"). This is a catch-all option to capture user interactions that are not codified by other options in this table.
6	Press Reset Button	no	The Device will enter the target state when reset button is pressed.
7	Press Reset Button with application of power	no	The Device will enter the target state when reset button is pressed when applying power to it.

Bit index	Name	Ins tru c- tio n De pen - den cy	Description
8	Press Reset Button for N seconds	M	The Device will enter the target state when reset button is pressed for N seconds. The exact value of N SHALL be made available via PI key.
9	Press Reset Button until light blinks	O	The Device will enter the target state when reset button is pressed until associated light blinks. Information on color of light MAY be made available via PI key (see Note 1).
10	Press Reset Button for N seconds with application of power	M	The Device will enter the target state when reset button is pressed for N seconds when applying power to it. The exact value of N SHALL be made available via PI key.
11	Press Reset Button until light blinks with application of power	O	The Device will enter the target state when reset button is pressed until associated light blinks when applying power to the Device. Information on color of light MAY be made available via PI key (see Note 1).
12	Press Reset Button N times	M	The Device will enter the target state when reset button is pressed N times with maximum 1 second between each press. The exact value of N SHALL be made available via PI key.
13	Press Setup Button	no	The Device will enter the target state when setup button is pressed.
14	Press Setup Button with application of power	no	The Device will enter the target state when setup button is pressed when applying power to it.
15	Press Setup Button for N seconds	M	The Device will enter the target state when setup button is pressed for N seconds. The exact value of N SHALL be made available via PI key.
16	Press Setup Button until light blinks	O	The Device will enter the target state when setup button is pressed until associated light blinks. Information on color of light MAY be made available via PI key (see Note 1).
17	Press Setup Button for N seconds with application of power	M	The Device will enter the target state when setup button is pressed for N seconds when applying power to it. The exact value of N SHALL be made available via PI key.

Bit index	Name	Instruction Dependency	Description
18	Press Setup Button until light blinks with application of power	O	The Device will enter the target state when setup button is pressed until associated light blinks when applying power to the Device. Information on color of light MAY be made available via PI key (see Note 1).
19	Press Setup Button N times	M	The Device will enter the target state when setup button is pressed N times with maximum 1 second between each press. The exact value of N SHALL be made available via PI key.
20	Power Cycle N times	M	The Device will automatically enter the target state after exactly N power OFF/ON sequences (unplug/re-plug, remove/re-insert batteries) starting from the powered ON state. A non-zero X value provides the minimum number of seconds to keep the device in the ON state before the next iteration. A non-zero Y value specifies the total elapsed time during which exactly N power cycles must be detected in order to enter the desired state. A non-zero Z value means that the device will blink Z times to indicate that the sequence was detected, before entering the desired state. If a variable (X, Y, or Z) is set to 0, then the corresponding aspect of the instruction is not specified. The format of the PI key's value SHALL be N,X,Y,Z where N,X,Y and Z are variable length decimal values (no leading zeros) and there are no spaces or other characters present in the value. For example, 5,1,10,3 indicates that within 10 seconds the device must be powered OFF then ON 5 times, remaining in the ON state each time for at least 1 second, and the device will blink 3 times to indicate success before entering the target state. Alternatively, 8,0,20,0 indicates that within 20 seconds the device must be powered OFF then ON 8 times in order to enter the target state.

Bit index	Name	Instruction Dependency	Description
21	Press Button for N seconds with indication	M	The Device will enter the target state when the button is pressed for N seconds. A non-zero Z value means that the device will blink Z times to indicate that the action was detected, before entering the desired state. A present and non-blank A value can be used to specify the button(s) where buttons are indicated using HDMI CEC Key Codes (see HDMI) in hexadecimal notation separated by operators '&' and ' ' to indicate 'and' (both buttons) and 'or' (any of the buttons) respectively. The format of the PI key's value SHALL be N,Z or N,Z,A, where N and Z are variable length decimal values (no leading zeros), A is blank or a combination of CEC Key Codes separated by '&' or ' ', and there are no spaces or other characters present in the value. For example, 10,3 indicates that any button can be held for 10 seconds, and the device will blink 3 times before entering the target state. Additional examples: 5,,0x46 indicates that the pause button (CEC Key Code 0x46) must be held for 5 seconds, 5,,0x01&0x02 indicates that the up and down buttons (CEC Key Codes 0x01 and 0x02) must be held 5 secs, and 5,,0x01 0x02 indicates that either the up or down buttons must be held 5 secs.

Bit index	Name	Instruction Dependency	Description
22	Power Cycle Until Indication	M	The Device will automatically enter the target state after a variable number of power OFF/ON sequences (unplug/re-plug, remove/re-insert batteries) starting from the powered ON state. A non-zero X value provides the number of seconds to keep the device in the ON state before moving to the OFF state. A non-zero Y value provides the number of seconds to keep the device in the OFF state before moving to the ON state. A non-zero Z value means that the device will blink Z times to indicate that the sequence was detected, before entering the desired state. An M value can be used to indicate the color name of the light blink, where the color name may be one of: white, red, green, blue, orange, yellow, purple, cool white, warm white, sequence (for rainbow sequence). An N value can be used to indicate the initial state (OFF or ON), where the default when missing or left blank is ON. If the variable X, Y, Z, M or N is blank or set to 0, then the corresponding aspect of the instruction is not specified. The format of the PI key's value SHALL be X,Y,Z,M,N where X, Y and Z are variable length decimal values (no leading zeros) and there are no spaces or other characters present in the value, while M is a single color name, and N, when present, is one of either ON or OFF. For example, 1,2,3,yellow,ON indicates that the device must be powered OFF then ON repeatedly, starting in the ON state, remaining in the ON state each time for 1 second, remaining in the OFF state each time for 2 seconds, and the device will blink in yellow 3 times to indicate success before entering the target state. Alternatively, 0,0,0,sequence indicates that the device must be powered OFF then ON repeatedly until it blinks in a sequence of colors to indicate it has entered the target state.

Note 1: When the **PH** key indicates a light to blink (one or more of bits 9, 11, 16 or 18 is set), information on color of light MAY be made available via **PI** key. When using such color indication in **PI** key, only basic primary and secondary colors that could unambiguously be decoded by a commissioner and understood by an end-user, but without worry of localization, SHOULD be used, e.g. white, red, green, blue, orange, yellow, purple, unless otherwise specified.

Note 2: Any undefined values are reserved for future use.

Note 3: A Commissionee can indicate multiple methods of being put into Commissioning Mode by setting multiple bits in the bitmap at the same time. However, only one such method can be speci-

fied which has a mandatory dependency on the PI key (PI Dependency=**M**) at a time.

For example:

- A **PH** value of 33 (bits 0 and 5 are set) indicates that the user can cause the Commissionee to enter Commissioning Mode by either power cycling it or by following special instructions provided in the Device Manual.
- A **PH** value of 9 (bits 0 and 3 are set) indicates that the user can cause the Commissionee to enter Commissioning Mode by either power cycling it or going to the settings menu and following instructions there.
- A **PH** value of 1 (bit 0 is set) indicates that the user can cause the Commissionee to enter Commissioning Mode only by power cycling it.
- A **PH** value of 16 (bit 4 is set) indicates that the user can cause the Commissionee to enter Commissioning Mode following a custom procedure described by the value of the PI key.
- A **PH** value of 256 (bits 8 is set) indicates that the user can cause the Commissionee to enter Commissioning Mode by pressing the reset button for a duration of time in seconds specified via by the value of the PI key.

When the **PH** key is provided, at least one bit in the above bitmap SHALL be set. That is, a **PH** value of 0 is undefined and illegal.

When the **PH** key is provided, the Commissioner SHOULD take its value into account when providing guidance to the user regarding steps required to put the Commissionee into Commissioning Mode.

4.3.1.12. TXT key for pairing instructions (**PI**)

The optional key **PI** MAY give the *pairing instruction*.

If present, the value SHALL be encoded as a valid UTF-8 string with a maximum length of 128 bytes.

The meaning of this key is dependent upon the **PH** key value, see [Table 5, “Pairing and Reset Hint Values”](#).

For example, given **PH=256**, bit 8 is set which indicates "Press Reset Button for N seconds". Therefore, a value of **PI=10** would indicate that N is 10 in that context.

When bit 4 of the value expressed by the **PH** key is set, indicating presence of a custom string, the Commissionee SHALL be responsible for localization (translation to user's preferred language) as required using the Device's currently configured locale. The Custom Instruction option is NOT recommended for use by a Commissionee that does not have knowledge of the user's language preference.

It is RECOMMENDED to keep the length of PI field small and adhere to the guidance given in section 6.2 of [\[RFC 6763\]](#).

This key/value pair SHALL only be returned in the DNS-SD TXT record if the **PH** bitmap value has a bit set which has PI Dependency = **M** or has a bit set which has PI Dependency = **0** and the Commissionee wants to indicate related information as mentioned in Note 1 above. See [Table 5, “Pairing](#)

and [Reset Hint Values](#)” for more information. The **PH** key SHALL NOT have more than one bit set which has a mandatory dependency on the **PI** key (PI Dependency = **M**) to avoid ambiguity in **PI** key usage.

4.3.1.13. TXT key for Joint Fabric

The optional key **JF** indicates the capabilities of an Ecosystem Administrator that can or is participating in the Joint Fabric.

The **JF** key SHALL be present in the DNS-SD TXT record if and only if the Node is capable of being a Joint Fabric Administrator.

The **JF** key SHALL be encoded as a variable-length decimal number in ASCII text, omitting any leading zeroes.

The **JF** key represents a base-10 numeric value for a bitmap of capabilities for a Joint Fabric Administrator device.

The bitmap of the Joint Fabric key values is defined in [Table 6, “Joint Fabric Key Values”](#).

Table 6. Joint Fabric Key Values

Bit index	Capability	Description
0	Available	This device is capable of acting as a Joint Fabric Administrator. (See Note 1)
1	Administrator	This device is acting as a Joint Fabric Administrator. (See Note 2)
2	Anchor	This device is acting as a Joint Fabric Anchor Administrator. (See Note 2)
3	Datastore	This device is acting as a Joint Fabric Datastore. (See Note 3)

Note 1: bit 0 (Available) SHALL be unset for any of bits 1, 2 or 3 to be set. bit 0 (Available) SHALL be set, bits 1, 2 and 3 SHALL be unset as the default value prior to the Administrator Node being commissioned onto the Joint Fabric. Once an Administrator device is commissioned on the Joint Fabric, bit 0 (Available) SHALL be unset.

Note 2: bit 1 (Administrator) SHALL be set for bit 2 (Anchor) to be set. A device SHALL be a Joint Fabric Administrator to be a Joint Fabric Anchor Administrator. A single device for the Ecosystem (Vendor) which is the Joint Fabric Anchor SHALL have bit 3 (Datastore) set. This means one or more devices MAY have bit 1 (Administrator) and bit 2 (Anchor) set and bit 3 (Datastore) cleared, but at most one device SHALL have bit 3 (Datastore) set.

Note 3: bit 1 (Administrator), bit 2 (Anchor), and bit 3 (Datastore) SHALL all be set for the single device which is the Joint Fabric Datastore.

For example:

- A **JF=1** key/value pair (bit 0 is set) indicates that the Administrator Node is capable of being a Joint Fabric Administrator but has not been commissioned onto the Joint Fabric.
- A **JF=14** key/value pair (bits 1, 2, and 3 are set) indicates that the device is a Joint Fabric Administrator, the Joint Fabric Anchor, and a Joint Fabric Datastore device.
- A **JF=2** key/value pair (bit 1 is set) indicates that the device is a Joint Fabric Administrator but NOT the Joint Fabric Anchor or a Joint Fabric Datastore device.

This value MAY change during the lifecycle of the device.

For example, a device might have a value with bit 0 (Available) set, bit 1 (Administrator) unset, bit 2 (Anchor) unset and bit 3 (Datastore) unset prior to being commissioned onto the Joint Fabric, and then have a value with bit 0 (Available) unset and bit 1 (Administrator) set after it has been commissioned onto the Joint Fabric.

The **VP** key SHALL be present in the DNS-SD TXT record if the **JF** key is present and SHALL provide the Vendor ID of the device.

For example:

VP=123 gives the Vendor ID associated with the Joint Fabric Administrator device.

4.3.1.14. Examples

The examples below simulate a Node in commissioning mode advertising its availability for commissioning.

Examples are shown using both the **dns-sd** command-line test tool and the **avahi** command-line test tool. The **dns-sd** command-line test tool is included in all versions of macOS. It is installed as a DOS command with Bonjour for Windows, and is available on Linux by installing the [mDNSResponder package](https://github.com/balaji-reddy/mDNSResponder) [https://github.com/balaji-reddy/mDNSResponder]. The Avahi package of command line tools is available from the [Avahi project](https://github.com/lathiat/avahi) [https://github.com/lathiat/avahi] for most Linux distributions.

These examples are given for illustrative purposes only. Real Matter Commissionees and Commissioners would not use a command-line test tool for advertising and discovery. Real Matter Commissionees and Commissioners would use the appropriate DNS-SD APIs in their respective chosen programming languages.

Consider a device on Wi-Fi using the 48-bit device MAC address of **B75AFB458ECD** as its host name and a value of **DD200C20D25AE5F7** as its commissionable service instance name. DNS-SD records for it can be set up as follows:

```
dns-sd -R DD200C20D25AE5F7 _matterc._udp,_S3,_L840,_CM . 11111 D=840 CM=2
```

or

```
avahi-publish-service --subtype=_S3._sub._matterc._udp
```



```
--subtype=_L840._sub._matterc._udp DD200C20D25AE5F7 --subtype=_CM._sub._matterc._udp
_matterc._udp 11111 D=840 CM=2
```

- Short discriminator is filterable through [_S3](#) subtype and algorithmically through [D=840](#) TXT key.
- Long discriminator is filterable through [_L840](#) subtype and directly through [D=840](#) TXT key.
- The Commissionee is currently in Commissioning Mode after an Administrator having opened a commissioning window (see [Section 4.3.1.7, “TXT key for commissioning mode \(CM\)”](#)), as shown by [CM=2](#) TXT key and availability by browsing the [_CM](#) subtype.
 - Had the Commissionee been discoverable for initial commissioning rather than subsequent additional commissioning, a [CM=1](#) TXT key would have been published instead.
 - Had the Commissionee been in Commissioning Mode due to the [OpenJointCommissioning-Window](#) command, a [CM=3](#) TXT key would have been published instead.

Avahi only sends a single AAAA record. To force the link-local address to be used, use `avahi-publish-address`. For example:

```
avahi-publish-address B75AFB458ECD.local fe80::f515:576f:9783:3f30
```

The DNS-SD service registration commands shown above results in the creation of the following Multicast DNS records:

<code>_matterc._udp.local.</code>	PTR	<code>DD200C20D25AE5F7._matterc._udp.local.</code>
<code>_S3._sub._matterc._udp.local.</code>	PTR	<code>DD200C20D25AE5F7._matterc._udp.local.</code>
<code>_L840._sub._matterc._udp.local.</code>	PTR	<code>DD200C20D25AE5F7._matterc._udp.local.</code>
<code>_CM._sub._matterc._udp.local.</code>	PTR	<code>DD200C20D25AE5F7._matterc._udp.local.</code>
<code>DD200C20D25AE5F7._matterc._udp.local.</code>	SRV	<code>0 0 11111 B75AFB458ECD.local.</code>
<code>DD200C20D25AE5F7._matterc._udp.local.</code>	TXT	<code>"D=840" "CM=2"</code>
<code>B75AFB458ECD.local.</code>	AAAA	<code>fe80::f515:576f:9783:3f30</code>

Consider a device on Wi-Fi using the 48-bit device MAC address of B75AFB458ECD as its host name. DNS-SD records for it can be set up as follows, when it is in Commissionable Node Discovery.

```
dns-sd -R DD200C20D25AE5F7 _matterc._udp,_S3,_L840,_V123,_CM,_T81 . 11111 D=840
VP=123+456 CM=2 DT=266 DN="Kitchen Plug" PH=256 PI=5
```

or

```
avahi-publish-service --subtype=_S3._sub._matterc._udp
--subtype=_L840._sub._matterc._udp --subtype=_V123._sub._matterc._udp
--subtype=_CM._sub._matterc._udp --subtype=_T81._sub._matterc._udp DD200C20D25AE5F7
_matterc._udp 11111 D=840 VP=123+456 CM=2 DT=81 DN="Kitchen Plug" PH=256 PI=5
```

- Short discriminator is [3](#), long discriminator is [840](#).

- Vendor ID is 123, Product ID is 456.
- Commissioning Mode is 2, indicating the Commissionee is currently in Commissioning Mode due to the [OpenCommissioningWindow](#) command.
 - Had the Commissionee been in Commissioning Mode due to the [OpenJointCommissioningWindow](#) command, a CM=3 TXT key would have been published instead.
- Device type is 266 which is a smart plug (On/Off Plugin Unit, Device Type Identifier 0x010A).
- Device name is Kitchen Plug.
- Pairing hint is 256 which indicates that the Commissionee's reset button must be held down for 5 seconds to enter Commissioning Mode where the value 5 is obtained by reading the value of the PI key.
- Pairing instruction is 5.

Avahi only sends a single AAAA record. To force the link-local address to be used, use avahi-publish-address. For example:

```
avahi-publish-address B75AFB458ECD.local fe80::f515:576f:9783:3f30
```

The DNS-SD service registration commands shown above results in the creation of the following Multicast DNS records:

_matterc._udp.local.	PTR	DD200C20D25AE5F7._matterc._udp.local.
_S3._sub._matterc._udp.local.	PTR	DD200C20D25AE5F7._matterc._udp.local.
_L840._sub._matterc._udp.local.	PTR	DD200C20D25AE5F7._matterc._udp.local.
_V123._sub._matterc._udp.local.	PTR	DD200C20D25AE5F7._matterc._udp.local.
_CM._sub._matterc._udp.local.	PTR	DD200C20D25AE5F7._matterc._udp.local.
_T81._sub._matterc._udp.local.	PTR	DD200C20D25AE5F7._matterc._udp.local.
DD200C20D25AE5F7._matterc._udp.local.	TXT	"D=840" "VP=123+456" "CM=1" "DT=81"
"DN=Kitchen Plug" "PH=256" "PI=5"		
DD200C20D25AE5F7._matterc._udp.local.	SRV	0 0 11111 B75AFB458ECD.local.
B75AFB458ECD.local.	AAAA	fe80::f515:576f:9783:3f30

The port number 11111 is given here purely as an example. One of the benefits of using DNS-SD is that services are not constrained to use a single predetermined well-known port. The port, along with the IPv6 address, is discovered by Commissioners at run time.

A Commissioner can discover all available Commissionees awaiting commissioning:

```
dns-sd -B _matterc._udp
```

or

```
avahi-browse _matterc._udp -r
```

A Commissioner can discover Commissionees awaiting commissioning with short discriminator 3:

```
dns-sd -B _matterc._udp,_S3
```

or

```
avahi-browse _S3._sub._matterc._udp -r
```

A Commissioner can discover Commissionees awaiting commissioning with long discriminator 840:

```
dns-sd -B _matterc._udp,_L840
```

or

```
avahi-browse _L840._sub._matterc._udp -r
```

A Commissioner can discover Commissionees awaiting commissioning with Vendor ID 123:

```
dns-sd -B _matterc._udp,_V123
```

or

```
avahi-browse _V123._sub._matterc._udp -r
```

A Commissioner can discover all Commissionees in commissioning mode:

```
dns-sd -B _matterc._udp,_CM
```

or

```
avahi-browse _CM._sub._matterc._udp -r
```

A commissioner can discover Matter Nodes with Device Type 81:

```
dns-sd -B _matterc._udp,_T81
```

or

```
avahi-browse _T81._sub._matterc._udp -r
```

A Commissioner can discover Nodes that are currently in Commissioning Mode as a result of a commissioning window opened by a current Administrator as a result of invoking either the [OpenCommissioningWindow](#) command, the [OpenJointCommissioningWindow](#) command or the [OpenBasicCommissioningWindow](#) command, using the presence of the `_CM` subtype as a browsing filter:

```
dns-sd -B _matterc._udp._CM
```

or

```
avahi-browse _CM._sub._matterc._udp -r
```

4.3.1.15. Efficiency Considerations

Discovering and using an offered service on the network typically involves several steps:

1. Enumeration of instances available on the network ("browsing")
2. Lookup of a selected instance's port number, host name, and other additional information, communicated in DNS-SD using SRV and TXT records ("resolving")
3. Lookup of the IPv6 address(es) associated with the desired target host.
4. Use of IPv6 Neighbor Discovery and/or IPv6 routing to translate from destination IPv6 address to the next-hop link-layer address for that communication.
5. Establishing a cryptographically secure communication channel between the two endpoints, and then engaging in useful communication.

Although the first three steps are exposed in some APIs as separate steps, at a protocol level they usually require only a single network round-trip. When a PTR query is issued to discover service instances, the usual DNS Additional Record mechanism, where packet space permits, automatically places the related SRV, TXT, and address records into the Additional Record section of the reply. These additional records are stored by the client, to enable subsequent steps in the sequence to be performed without additional redundant network operations to learn the same information a second time.

DNS-SD over Multicast DNS works by receiving replies from other Nodes attached to the same local link, Nodes that may have been previously completely unknown to the requester. Because of this, Multicast DNS, like IPv6 Neighbor Discovery, does not have any easy way to distinguish genuine replies from malicious or fraudulent replies. Consequently, application-layer end-to-end security is essential. Should a malicious device on the same local link give deliberately malicious or fraudulent replies, the misbehavior will be detected when the device is unable to establish a cryptographically secure application-layer communication channel. This reduces the threat to a Denial-of-Service attack, which can be remedied by physically removing the offending device.

4.3.2. Operational Discovery

For Matter Nodes that have already been commissioned onto a Matter Fabric, run-time dynamic discovery of operational Matter Nodes is used, rather than assuming a fixed unchanging IPv6

address and port for the lifetime of the product. This is done to allow for greater flexibility, so that the underlying IPv6 network can grow and evolve over time as needed without breaking Matter functionality. This is the same reason that other networked consumer electronics products do not assume a single fixed unchanging IP address for the lifetime of the product [RFC 5505].

4.3.2.1. Operational Instance Name

For Matter operational discovery the DNS-SD instance name is constructed from a 64-bit [compressed Fabric identifier](#), and a 64-bit Node identifier, as assigned by the commissioner, each expressed as a fixed-length sixteen-character hexadecimal string, encoded as ASCII (UTF-8) text using capital letters, separated by a hyphen. For example, a Matter Node with Matter compressed fabric identifier **2906-C908-D115-D362** and Matter Node identifier **8FC7-7724-01CD-0696** has Matter operational discovery DNS-SD instance name **2906C908D115D362-8FC7772401CD0696**.

The Matter operational discovery DNS-SD instance name needs to be unique within the namespace of the local network (the **.local** link-local namespace of the Ethernet and Wi-Fi links [RFC 6762], or the unicast domain selected by the Thread Border Router for devices on the Thread mesh). This uniqueness is assumed to be guaranteed by appropriate selection of a unique Matter fabric identifier and unique Node identifier within that Matter fabric.

4.3.2.2. Compressed Fabric Identifier

In order to reduce the very large size of a full [Fabric Reference](#) which would need to be used as the scoping construct in the [instance name](#), a 64-bit compressed version of the full Fabric Reference SHALL be used. The computation of the Compressed Fabric Identifier SHALL be as follows:

```
byte CompressedFabricInfo[16] = /* "CompressedFabric" */
    {0x43, 0x6f, 0x6d, 0x70, 0x72, 0x65, 0x73, 0x73,
     0x65, 0x64, 0x46, 0x61, 0x62, 0x72, 0x69, 0x63}

CompressedFabricIdentifier =
    Crypto_KDF(
        inputKey := TargetOperationalRootPublicKey,
        salt := TargetOperationalFabricID,
        info := CompressedFabricInfo,
        len := 64)
```

Where:

- **TargetOperationalRootPublicKey** is the raw uncompressed elliptical curve public key of the root certificate for the advertised Node's [Operational Certificate](#) chain, without any format marker prefix byte (i.e. after removing the first byte of the **ec-pub-key** field in the Operational Certificate's root).
- **TargetOperationalFabricID** is the octet string for the Fabric ID as it appears in the advertised Node's [Operational Certificate](#)'s **subject** field, under the [1.3.6.1.4.1.37244.1.5](#) RDN, that is, a 64-bit unsigned integer scalar in big-endian byte order.

For example, if the root public key for a given Operational Certificate chain containing the identity to be advertised were the following:

pub:

```
04:4a:9f:42:b1:ca:48:40:d3:72:92:bb:c7:f6:a7:e1:
1e:22:20:0c:97:6f:c9:00:db:c9:8a:7a:38:3a:64:1c:
b8:25:4a:2e:56:d4:e2:95:a8:47:94:3b:4e:38:97:c4:
a7:73:e9:30:27:7b:4d:9f:be:de:8a:05:26:86:bf:ac:
fa
```

Then the value for **TargetOperationalRootPublicKey** to use in the derivation of the compressed Fabric Identifier would be without the leading **04**:

```
4a:9f:42:b1:ca:48:40:d3:72:92:bb:c7:f6:a7:e1:1e:
22:20:0c:97:6f:c9:00:db:c9:8a:7a:38:3a:64:1c:b8:
25:4a:2e:56:d4:e2:95:a8:47:94:3b:4e:38:97:c4:a7:
73:e9:30:27:7b:4d:9f:be:de:8a:05:26:86:bf:ac:fa
```

If using the above **TargetOperationalRootPublicKey** and a **TargetOperationalFabricID** value of 0x2906_C908_D115_D362 (octet string **29:06:c9:08:d1:15:d3:62** in big-endian), then the **Compressed-FabricIdentifier** to use in advertising would be **87E1B004E235A130** (octet string **87:e1:b0:04:e2:35:a1:30**).

4.3.2.3. Operational Service Type

The DNS-SD service type [RFC 6335] for Matter Operational Discovery is **_matter._tcp**. Note that the string **_tcp** is boilerplate text inherited from the original DNS SRV specification [RFC 2782], and doesn't necessarily mean that the advertised application-layer protocol runs only over TCP. It is merely mnemonic text which is there to help human readers, and in no way affects software advertising or using the application-layer protocol identified by that unique IANA-recorded service type string.

The following subtype is defined:

1. Compressed Fabric Identifier **_I<hhhh>**, where **<hhhh>** is the **Compressed Fabric Identifier** encoded as exactly 16 uppercase hexadecimal characters, for example **_I87E1B004E235A130** for the Compressed Fabric Identifier example of the previous section. This subtype enables filtering of devices per Fabric if service enumeration (browsing) is attempted, to reduce the set of results to Nodes of interest with operational membership in a given Fabric..

4.3.2.4. Operational Service Domain and Host Name

For link-local Multicast DNS the service domain SHALL be **local**. For Unicast DNS such as used on Thread the service domain SHALL be as configured automatically by the Thread Border Router.

For DNS-SD a target host name is required, in addition to the instance name. The target host name SHALL be constructed using one of the available link-layer addresses, such as a 48-bit device MAC

address (for Ethernet and Wi-Fi) or a 64-bit MAC Extended Address (for Thread) expressed as a fixed-length twelve-character (or sixteen-character) hexadecimal string, encoded as ASCII (UTF-8) text using capital letters, e.g., **B75AFB458ECD.<domain>**. In the event that a device performs MAC address randomization for privacy, then the target host name SHALL use the privacy-preserving randomized version and the hostname SHALL be updated in the record every time the underlying link-layer address rotates. Note that it is legal to reuse the same hostname on more than one interface, even if the underlying link-layer address does not match the hostname on that interface, since the goal of using a link-layer address is to ensure local uniqueness of the generated hostname. If future link layers are supported by Matter that do not use 48-bit MAC addresses or 64-bit MAC Extended Address identifiers, then a similar rule will be defined for those technologies.

4.3.2.5. Subtypes

The following subtype for operational discovery service is defined:

- **_IC** (Incomplete Commissioning) allows the administrator who initiated an NFC-based commissioning to filter nodes whose commissioning flow is yet to be completed. A Commissionee using NFC-based commissioning SHALL publish this subtype from the setup phase step 2 up to the step 4 (included), and SHALL withdraw it (using SRP update or DNS-SD with TTL=0) when leaving this step range. A node outside this step range of the commissioning flow SHALL NOT publish this subtype. A Commissionee not using NFC-based commissioning SHALL NOT use this subtype.

4.3.2.6. Operational Discovery Service Records

After discovery, IPv6 addresses are returned in the AAAA records and key/value pairs are returned in the DNS-SD TXT record. The TXT record MAY be omitted if no keys are defined.

Nodes SHALL publish AAAA records for all available IPv6 addresses upon which they are willing to accept operational messages.

TXT records available for Operational Discovery include the common TXT record key/value pairs defined in [Section 4.3.4, “Common TXT Key/Value Pairs”](#).

The section [Section 4.3.2.7, “TXT record key/value for Incomplete Commissioning \(IC\)”](#) defines additional key/value pairs specific to Operational Discovery.

Nodes SHALL silently ignore TXT record keys that they do not recognize.

4.3.2.7. TXT record key/value for Incomplete Commissioning (IC)

The key/value pair **IC=1** indicates that the node is waiting for Commissioning Complete command as defined in [Section 4.3.2.5, “Subtypes” _IC](#). A node publishing **_IC** subtype SHALL include **IC=1** in TXT record of both the subtype and the base services.

The key/value pair **IC=0** indicates that the node is not waiting for Commissioning Complete command as defined in [Section 4.3.2.5, “Subtypes” _IC](#). A Commissionee that is using NFC-based commissioning but is not waiting for the Commissioning Complete command SHALL NOT publish the **_IC** subtype.

The absence of the key **IC** or any undefined value does not provide information on the commissioning step.

When performing Operational Discovery, the Commissioner SHALL identify devices with the 'IC=1' subtype. If it knows how to commission this device, the Commissioner SHALL execute the steps from 3 to 4 to complete the commissioning of the device.

4.3.2.8. Performance Recommendations

To improve overall performance of operational discovery, especially in large installations, the following recommendations SHOULD be taken in account:

1. Nodes SHOULD cache the last-known IPv6 address and port for each peer Node with which they interact from their SRV record resolved using DNS-SD, to save the cost of a run-time network lookup operation when not needed. When the IPv6 address and port for a peer Node is not known, or an attempt to communicate with a peer Node at its last-known IPv6 address and port does not appear to be succeeding within the expected network round-trip time (i.e., the retransmission timeout value for the first message packet) a Node SHOULD then perform a run-time discovery in parallel, to determine whether the desired Node has acquired a new IPv6 address and/or port [RFC 8305].
2. Nodes SHOULD respond to nonspecific service enumeration queries for the generic Matter Operational Discovery service type (**_matter._tcp**), but these queries SHOULD NOT be used in routine operation, and instead it is RECOMMENDED that they only be used for diagnostics purposes or to determine new membership within a fabric. When used, it is RECOMMENDED that service enumeration employ the **_I<HHHH>** Fabric-specific subtype to only enumerate the desired Nodes on the Fabric of interest in the local network. Moreover, Known Answer Suppression [RFC 6762] SHOULD be employed in such cases to further minimize the number of unnecessary responses to such a query.
3. When resolving the operational service record of another Node, a Node SHOULD use an SRV query for the desired operational service instance rather than doing general enumeration of all nodes (e.g. PTR query) followed by filtering for the desired service instance. This recommendation reduces the amount of multicast traffic generated on-link when Multicast DNS is used, and reduces latency to successful service resolution.
4. Since proxied DNS-SD service discovery MAY be in use within a given network, and service record caching is expected of DNS-SD clients, Nodes SHOULD NOT use DNS-SD as an operational liveness determination method. This is because there may be stale records not yet expired after a Node becomes unreachable which may still be available.

4.3.2.9. Operational Discovery DNS-SD Examples

The example below simulates a commissioned Matter Node advertising its availability for control via the Matter protocol.

Examples are shown using both the **dns-sd** command-line test tool and the **avahi** command-line test tool. The **dns-sd** command-line test tool is included in all versions of macOS. It is installed as a DOS command with Bonjour for Windows, and is available on Linux by installing the [mDNSResponder package](https://github.com/balaji-reddy/mDNSResponder) [https://github.com/balaji-reddy/mDNSResponder]. The **avahi** command line-test tool is available from the [Avahi project](https://github.com/lathiat/avahi) [https://github.com/lathiat/avahi] for most Linux distributions.

This example is given for illustrative purposes only. Real Matter Nodes and controllers would not use a command-line test tool for advertising and discovery. Real Matter Nodes and controllers would use the appropriate DNS-SD APIs in their respective chosen programming languages.

Consider a device on Wi-Fi using the 48-bit device MAC address of B75AFB458ECD as its host name. DNS-SD records for can be set up as follows:

```
dns-sd -R 87E1B004E235A130-8FC7772401CD0696 _matter._tcp . 22222
```

or

```
avahi-publish-service 87E1B004E235A130-8FC7772401CD0696 _matter._tcp 22222
```

The port number 22222 is given here purely as an example. One of the benefits of using DNS-SD is that services are not constrained to use a single predetermined well-known port. This means that multiple instances of the Matter Node control service can run on the same device at the same time, listening on different ports [RFC 6760]. The port, along with the IPv6 address, is discovered by the Matter controller at run time.

Avahi only sends a single AAAA record. To force the link-local address to be used, use avahi-publish-address. For example:

```
avahi-publish-address B75AFB458ECD.local fe80::f515:576f:9783:3f30
```

A Matter controller can discover the current IPv6 address and port for a known commissioned Matter Node:

```
dns-sd -L 87E1B004E235A130-8FC7772401CD0696 _matter._tcp  
87E1B004E235A130-8FC7772401CD0696._matter._tcp.local. can be reached at  
B75AFB458ECD.local.:22222
```

```
dns-sd -Gv6 B75AFB458ECD.local  
fe80::f515:576f:9783:3f30
```

or

```
avahi-browse _matter._tcp -r
```

```
hostname = [B75AFB458ECD.local]  
address = [fe80::f515:576f:9783:3f30]  
port = [22222]
```


4.3.3. Commissioner Discovery

A Commissionee MAY initiate the commissioning process by discovering Commissioners on the network (see [Section 5.2.2.4, “Commissioner Discovery, from an On-Network Device”](#)). This MAY be done using *Commissioner Discovery* described in this section.

With Commissioner Discovery, a Commissionee, upon user interaction, MAY discover Commissioners on the network and obtain a list of information for each which may include Vendor ID, Product ID and friendly name. A Commissionee with a user interface, such as a Television, Thermostat or Video Player device, MAY display the list of discovered commissioners to the user for selection. Once selected, the Commissionee MAY use the [User Directed Commissioning](#) protocol with the Commissioner to indicate that the user has selected it for commissioning of the Commissionee. The Commissioner Discovery service records thus enable a form of "door bell" protocol to allow a Commissionee to request Commissioning.

The Commissioner Discovery feature is optional for both the Commissionee and the Commissioner. Any mandatory requirements described in this section SHALL apply only if the Node or the Commissioner supports this feature. To protect customer privacy on public networks, a Matter Commissioner SHALL provide a way for the customer to set a timeout on Commissioner Discovery, or otherwise disable Commissioner Discovery.

For Commissioner Discovery, the DNS-SD instance name is generated the same way it is done for Commissionable Node Discovery and has the same requirements (uniqueness on local network, and collision detection and recovery) as those in Commissionable Node Discovery, but the requirements for when a new instance name is selected from Commissionable Node Discovery do not apply to Commissioner Discovery. The instance name for Commissioner Discovery MAY be selected whenever the Commissioner deems necessary.

The DNS-SD service type [\[RFC 6335\]](#) is `_matterd._udp`.

The port advertised by a `_matterd._udp` service record SHALL be different than any port associated with other advertised `_matterc._udp` or `_matter._tcp` services, in order to ensure that the session-less messaging employed by the [User Directed Commissioning](#) protocol does not cause invalid message handling from fully operational Matter nodes at the same address. In other words, each `_matterd._udp` service instance needs to be independent from other services to ensure unambiguous processing of the incoming User Directed Commissioning messages.

The following subtype is defined:

- `_T<ddd>` where `<ddd>` is the [primary device type](#) identifier, encoded as a variable-length decimal number in ASCII (UTF-8) text, without leading zeroes. This optional Device Type subtype enables filtering of results to find only Commissioners that match a device type, for example, to discover Commissioners of type Video Player (35 is decimal representation for Video Player device type identifier 0x0023). For such a Video Player filter, subtype `_T35` would be used.

For link-local Multicast DNS the service domain SHALL be `local`. For Unicast DNS such as used on Thread the service domain SHALL be as configured automatically by the Thread Border Router.

The target host name is generated the same way it is done for Commissionable Node Discovery (see [Section 4.3.1.1, “Host Name Construction”](#)).

After discovery, IPv6 addresses are returned in the AAAA records and key/value pairs are returned in the DNS-SD TXT record. The TXT record MAY be omitted if no keys are defined.

Nodes SHALL publish AAAA records for all their available IPv6 addresses.

In addition to the common TXT record key/value pairs defined in [Section 4.3.4, “Common TXT Key/Value Pairs”](#), the following key/value pairs are defined specifically for Commissioner discovery:

- The optional key **VP** gives vendor and product information. This key is optional for a vendor to provide, and optional for a commissioner to use. This value takes the same format described for the **VP** key in Commissionable Node Discovery (see [Section 4.3.1.6, “TXT key for Vendor ID and Product ID \(VP\)”](#)). This key/value pair MAY be returned in the DNS-SD TXT record.
- The optional key **DT** gives the [primary device type](#) identifier for the Commissioner. This value takes the same format described for the **DT** key in Commissionable Node Discovery (see [Section 4.3.1.8, “TXT key for device type \(DT\)”](#)). This key/value pair MAY be returned in the DNS-SD TXT record.
- The optional key **DN** gives the device name. This value takes the same format described for the **DN** key in Commissionable Node Discovery (see [Section 4.3.1.9, “TXT key for device name \(DN\)”](#)). This key/value pair MAY be returned in the DNS-SD TXT record. To protect customer privacy on public networks, a Matter Commissioner SHALL provide a way for the customer to disable inclusion of this key.
- The optional key **CP** indicates whether the Commissioner supports the [Commissioner Passcode feature](#). The absence of key **CP** SHALL imply a value of 0 (**CP=0**), and indicates that the publisher does not support the Commissioner Passcode feature. The key/value pair **CP=1** SHALL indicate that the publisher supports the Commissioner Passcode feature.

Commissionees SHALL silently ignore TXT record keys that they do not recognize. This is to facilitate future evolution of the Matter Commissioner Discovery protocol specification without breaking backwards compatibility with existing Commissionees that do not implement the new functionality.

4.3.3.1. Examples

The examples below simulate a Matter Commissioner advertising that it is present on the network.

Examples are shown using both the **dns-sd** command-line test tool and the **avahi** command-line test tool. The **dns-sd** command-line test tool is included in all versions of macOS. It is installed as a DOS command with Bonjour for Windows, and is available on Linux by installing the [mDNSResponder package](https://github.com/balaji-reddy/mDNSResponder) [https://github.com/balaji-reddy/mDNSResponder]. The **avahi** command line-test tool is available from the [Avahi project](https://github.com/lathiat/avahi) [https://github.com/lathiat/avahi] for most Linux distributions.

These examples are given for illustrative purposes only.

Consider a device on Wi-Fi using the 48-bit device MAC address of B75AFB458ECD as its host name. DNS-SD records for can be set up as follows:

```
dns-sd -R DD200C20D25AE5F7 _matterd._udp,_V123,_T35 . 33333 VP=123+456 DT=35  
DN="Living Room TV"
```

or

```
avahi-publish-service --subtype=_V123._sub._matterd._udp DD200C20D25AE5F7
_matterd._udp 33333 VP=123+456 DT=35 DN="Living Room TV"
```

This produces DNS-SD messages with the following characteristics:

- Vendor ID is **123**, Product ID is **456**.
- Device type is **35** which is a Casting Video Player (Device Type Identifier 0x0023).
- Device name is **Living Room TV**.

Avahi only sends a single AAAA record. To force the link-local address to be used, use avahi-publish-address. For example:

```
avahi-publish-address B75AFB458ECD.local fe80::f515:576f:9783:3f30
```

The DNS-SD service registration command shown above results in the creation of the following Multicast DNS records:

_matterd._udp.local.	PTR	DD200C20D25AE5F7._matterd._udp.local.
_V123._sub._matterd._udp.local.	PTR	DD200C20D25AE5F7._matterd._udp.local.
_T35._sub._matterd._udp.local.	PTR	DD200C20D25AE5F7._matterd._udp.local.
DD200C20D25AE5F7._matterd._udp.local.	TXT	"VP=123+456" "DT=35" "DN=Living Room TV"
DD200C20D25AE5F7._matterd._udp.local.	SRV	0 0 33333 B75AFB458ECD.local.
B75AFB458ECD.local.	AAAA	fe80::f515:576f:9783:3f30

The port number 33333 is given here purely as an example.

A Commissionee can discover all Commissioners:

```
dns-sd -B _matterd._udp
```

or

```
avahi-browse _matterd._udp -r
```

A Commissionee can discover Commissioners with device type 35:

```
dns-sd -B _matterd._udp,_T35
```

or

```
avahi-browse _T35._sub._matterd._udp -r
```

A Commissionee can discover Commissioners with Vendor ID 123:

```
dns-sd -B _matterd._udp,_V123
```

or

```
avahi-browse _V123._sub._matterd._udp -r
```

4.3.4. Common TXT Key/Value Pairs

The TXT records provided during Commissionable, Operational and Commissioner discovery MAY contain the following optional key/value pairs which are common to every discovery method:

- The optional key **SII** indicates the [SESSION_IDLE_INTERVAL](#) of the Node. This key defines the [MRP retry interval](#) for a Node that is Idle. This key MAY optionally be provided by a Node to override the default setting. If the key is not included or invalid, the Node querying the service record SHALL use the default MRP parameter value. The **SII** value is an unsigned integer with units of milliseconds and SHALL be encoded as a variable-length decimal number in ASCII encoding, omitting any leading zeros. The **SII** value SHALL NOT exceed 3600000 (1 hour in milliseconds).
 - Example: **SII=5300** to override the initial retry interval value to 5.3 seconds.
- The optional key **SAI** indicates the [SESSION_ACTIVE_INTERVAL](#) of the Node. This key defines the [MRP retry interval](#) for a Node that is Active. This key MAY optionally be provided by a Node to override the default setting. If the key is not included or invalid, the Node querying the service record SHALL use the default MRP parameter value. The **SAI** value is an unsigned integer with units of milliseconds and SHALL be encoded as a variable-length decimal number in ASCII encoding, omitting any leading zeros. The **SAI** value SHALL NOT exceed 3600000 (1 hour in milliseconds).
 - Example: **SAI=1250** to override the active retry interval value to 1.25 seconds.
- The optional key **SAT** indicates the [SESSION_ACTIVE_THRESHOLD](#) of the Node. This key defines the duration of time the Node stays Active after the last network activity. This key MAY optionally be provided by a Node to override the default setting. If the key is not included or invalid, the Node querying the service record SHALL use the default MRP parameter value. The **SAT** value is an unsigned integer with units of milliseconds and SHALL be encoded as a variable-length decimal number in ASCII encoding, omitting any leading zeros. The **SAT** value SHALL NOT exceed 65535 (65.535 seconds).
 - Example: **SAT=1250** to override the active retry interval value to 1.25 seconds.

This key defines the duration of time the Node stays Active after the last network activity.

- The optional key **T** indicates various additional transport protocol modes, apart from [MRP](#) over

UDP, that are supported by a Node. If the key is not included or invalid, the Node querying the service record SHALL assume the default value of **T=0** indicating that only MRP is supported as a transport protocol. The **T** key, if included, SHALL be encoded as a decimal number in ASCII text, omitting any leading zeroes. This key SHALL only be used during operational discovery and not by commissioner/commissionee Nodes.

- It represents a base-10 numeric value for a bitmap of various additional transport protocol modes supported by the advertising Node.
- For example, the **T=6** key/value pair represents a value with bits 1 and 2 set, and bit 0 cleared.
- Bit index 0 is deprecated and SHALL always be set to 0 by the advertising node and ignored by nodes processing this key.
- The fields in the bitmap for values of the **T** key are currently defined in [Table 7, “Supported Transport Mode Values”](#).

Thus, a Node advertising with key **T=6** represents both a TCP client and a TCP server, but a Node advertising with **T=2** can only be a TCP client.

Table 7. Supported Transport Mode Values

Bit index	Name	Description
0	Reserved	This bit index is deprecated and SHALL be set to 0. Clients SHALL silently ignore this bit.
1	TCP Client	The advertising Node implements the TCP Client mode and MAY connect to a peer Node that is a TCP Server.
2	TCP Server	The advertising Node implements the TCP Server mode and SHALL listen for incoming TCP connections.

Because the information carried in DNS-SD records with Matter is not trustworthy (since the source is not authenticated), the value of the T key SHOULD be regarded only as a hint. The only reliable way to determine which transports are supported by a Node is to connect to it using CASE over MRP and get the list of supported transports from the [session parameters](#).

- The optional key **ICD** indicates whether the Node is operating as a Long Idle Time ICD or as a Short Idle Time ICD. The key SHALL NOT be provided by a Node that does not support the ICD Long Idle Time operating mode. If the key is invalid or not included, the Node querying the service record SHALL assume the default value of **ICD=0** indicating that the ICD is not in the Long Idle Time operating mode. If the **ICD** key is included, it SHALL have one of the two valid values:
 - **0** to indicate "Long Idle Time ICD is not operating as a Long Idle Time ICD",
 - **1** to indicate "Long Idle Time ICD is operating as a Long Idle Time ICD".
 - Example: **ICD=1** to announce that the ICD is in the Long Idle Time operating mode.

NOTE

Since the information carried by DNS-SD records are not authenticated in Matter’s usage of this protocol, security-critical decisions SHOULD NOT be made solely on the basis of them. Rather, this information SHOULD be regarded as a hint that needs verification.

4.4. Message Frame Format

This section describes the encoding of the Matter message format. The Matter message format provides flexible support for various communication paradigms, including unicast secure sessions, multicast group messaging, and session establishment itself. The process of encrypting Matter messages is the same in all modes of communication, and assumes symmetric keys are shared between communicating parties. Unencrypted messages are used only for protocols which bootstrap secure messaging, such as session establishments.

Matter messages are used by Matter applications, as well as the Matter protocol stack itself, to convey application-specific data and/or commands. The Protocol portion of a Matter message contains a [Protocol ID](#) and [Protocol Opcode](#) which identify both the semantic meaning of the message as well as the structure of any associated application payload data. Matter messages also convey an [Exchange ID](#), which relates the message to a particular exchange (i.e. conversation) taking place between two nodes. Finally, certain types of Matter messages can convey information that acknowledges the reception of an earlier message. This is used as part of the [Message Reliability Protocol](#) to provide guaranteed delivery of messages over unreliable transports.

All multi-byte integer fields are transmitted in little-endian byte order unless otherwise noted in the field description.

Matter messages are structured as follows:

NOTE

denotes the field is optional.

Table 8. Matter Message format definition

Length	Field
Message Header	
1 byte	Message Flags
2 bytes	Session ID
1 byte	Security Flags
4 bytes	Message Counter
0/8 bytes	[Source Node ID]
0/2/8 bytes	[Destination Node ID]
variable	[Message Extensions . . .]
Message Payload	
variable	[Message Payload . . .] (encrypted)
Message Footer	

Length	Field
variable	[Message Integrity Check]

The Message Payload of a Matter message SHALL contain a Protocol Message with format as follows:

Table 9. Protocol Message format definition

Length	Field
Protocol Header	
1 byte	Exchange Flags
1 byte	Protocol Opcode
2 bytes	Exchange ID
2 bytes	[Protocol Vendor ID]
2 bytes	Protocol ID
4 bytes	[Acknowledged Message Counter]
variable	[Secured Extensions . . .]
Application Payload	
variable	[Application Payload . . .]

4.4.1. Message Header Field Descriptions

4.4.1.1. Message Flags (8 bits)

An unsigned integer bit field containing the following subfields:

Table 10. Message Flags field definition

bit 7	6	5	4	3	2	1	0
Version				-	S	DSIZ	

NOTE

All unused bits in the Message Flags field are reserved and SHALL be set to zero on transmission and SHALL be silently ignored on reception.

Version (4 bits, positions 4-7)

An unsigned integer specifying the version of the Matter Message format used to encode the message. Currently only one version is defined:

- 0 — Matter Message Format version 1.0
- 1-15 — Reserved for future use

Messages with version field set to reserved values SHALL be dropped without sending a message-layer acknowledgement.

NOTE	The Version field conveys information solely about the structure of the Matter message itself, not about the structure of the application payload or the interpretation of the message’s type. Thus, changes to how an application handles or interprets a message do not result in the creation of a new message format version number.
------	--

S Flag (1 bit, position 2)

A single bit field which SHALL be set if and only if the [Source Node ID](#) field is present.

DSIZ Field (2 bits, position 0-1)

This field SHALL indicate the size and meaning of the [Destination Node ID](#) field.

- 0 — Destination Node ID field is not present
- 1 — Destination Node ID field is present as a 64-bit Node ID
- 2 — Destination Node ID field is present as a 16-bit Group ID
- 3 — Reserved for future use

Messages with DSIZ field set to reserved values SHALL be dropped without sending a message-layer acknowledgement.

4.4.1.2. Session ID (16 bits)

An unsigned integer value identifying the session associated with this message. The session identifies the particular key used to encrypt a message out of the set of available keys (either session or group), and the particular encryption/message integrity algorithm to use for the message. The Session ID field is always present. For details on derivation of this field, see respective sections on [unicast](#) and [group](#) session ID derivation.

4.4.1.3. Security Flags (8 bits)

An unsigned integer bit field containing the following subfields:

Table 11. Security Flags field definition

bit 7	6	5	4	3	2	1	0
P	C	MX	Reserved			Session Type	

NOTE	All unused bits in the Security Flags field are reserved and SHALL be set to zero on transmission and SHALL be silently ignored on reception.
------	---

P Flag (1 bit, position 7)

The Privacy (P) flag is a single bit field which, when set, SHALL identify that the message is encoded with privacy enhancements as specified in [Section 4.9.3, “Privacy Processing of Outgoing Messages”](#).

C Flag (1 bit, position 6)

The Control message (C) flag is a single bit field which, when set, SHALL identify that the message is a control message, such as for the [Message Counter Synchronization Protocol](#), and uses the control message counter for the nonce field as specified in [Section 4.8.1.1, “Nonce”](#).

MX Flag (1 bit, position 5)

The Message Extensions (MX) flag is a single bit field which, when set, SHALL indicate that the [Message Extensions](#) portion of the message is present and has non-zero length. Version 1.0 Nodes SHALL set this flag to zero.

Session Type (2 bit, position 0-1)

An unsigned integer specifying the type of session associated with the message. The following values are defined:

- 0 — Unicast Session
- 1 — Group Session
- 2-3 — Reserved for future use

Messages with Session Type set to reserved values are not valid and SHALL be dropped without sending a message-layer acknowledgement.

The Session Type defines how the Session ID is to be interpreted.

The *Unsecured Session* SHALL be indicated when both Session Type and Session ID are set to 0. The *Unsecured Session* SHALL have no encryption, privacy, or message integrity checking.

A *Secure Unicast Session* SHALL be indicated when Session Type is Unicast Session and Session ID is NOT 0.

4.4.1.4. Message Counter (32 bits)

An unsigned integer value uniquely identifying the message from the perspective of the sending node. The message counter is [generated](#) based on the Session Type and increases monotonically for each unique message generated. When messages are retransmitted, using the reliable messaging capabilities, the counter remains the same, as logical retransmission is of a given message as identified by its message counter. Similarly, [acknowledgements](#) refer to values of the message counter being acknowledged.

NOTE

The Message Counter field is scoped to a given Encryption Key. Also, the Message Counter values are independent for control messages and data messages, as indicated by the [C Flag](#). So it is possible to have the same Message Counter for two messages encrypted with different keys, as well as two messages encrypted with the same key but different values of the [C Flag](#).

4.4.1.5. Source Node ID (64 bits)

An optional sequence of 8 bytes containing the unique identifier of the source node. The Source

Node ID field SHALL only be present in a message when the [S Flag](#) in the Message Flags field is set to 1.

4.4.1.6. Destination Node ID

The optional Destination Node ID field contains the unique Node Identifier of the destination Node or group to which the message is being sent. The size and encoding of the Destination Node ID field depends on the value of the [DSIZ](#) field.

4.4.1.7. Message Extensions (variable)

The Message Extensions field is a variable length block of data for providing backwards compatible extensibility. The format of the Message Extensions block is shown in [Table 12](#), “[Message Extensions block format definition](#)”. The Message Extensions block SHALL be present only if the [MX Flag](#) is set to 1 in the Security Flags field.

Table 12. Message Extensions block format definition

Length	Field
2 bytes	Message Extensions Payload Length, in bytes
variable	[Message Extensions Payload]

If the [MX Flag](#) is set to 1, the Message Extensions Payload Length field SHALL be present and SHALL contain the length of the Message Extensions Payload. The Message Extensions Payload Length field SHALL NOT be privacy obfuscated.

Version 1.0 Nodes SHALL ignore the contents of the Message Extensions payload, by skipping it, to access the [Message Payload](#).

4.4.2. Message Footer Field Descriptions

4.4.2.1. Message Integrity Check (variable length)

A sequence of bytes containing the message integrity check value (a.k.a. tag or MIC) for the message. The length and byte order of the field depend on the integrity check algorithm in use as specified in [Section 3.6](#), “[Data Confidentiality and Integrity](#)”.

The Message Integrity Check field SHALL be present for all messages except those of [Unsecured Session Type](#).

The MIC is calculated as described in [Section 4.8.2](#), “[Security Processing of Outgoing Messages](#)”.

4.4.3. Protocol Header Field Descriptions

4.4.3.1. Exchange Flags (8 bits)

An unsigned integer bit field containing the following subfields:

Table 13. Exchange Flags field definition

bit 7	6	5	4	3	2	1	0
-	-	-	V	SX	R	A	I

NOTE

All unused bits in the Exchange Flags field are reserved and SHALL be set to zero on transmission and SHALL be silently ignored on reception.

I Flag (1 bit, position 0)

The Initiator (I) flag is a single bit field which, when set, SHALL indicate that the message was sent by the initiator of the exchange.

A Flag (1 bit, position 1)

The Acknowledgement (A) flag is a single bit field which, when set, SHALL indicate that the message serves as an acknowledgement of a previous message received by the current message sender.

R Flag (1 bit, position 2)

The Reliability (R) flag is a single bit field which, when set, SHALL indicate that the message sender wishes to receive an acknowledgement for the message.

SX Flag (1 bit, position 3)

The Secured Extensions (SX) flag is a single bit field which, when set, SHALL indicate that the [Secured Extensions](#) portion of the message is present and has non-zero length. Version 1.0 Nodes SHALL set this flag to zero.

V Flag (1 bit, position 4)

The Vendor (V) protocol flag is a single bit field which, when set, SHALL indicate whether the Protocol Vendor ID is present.

4.4.3.2. Protocol Opcode (8 bits)

An unsigned integer value identifying the type of the message. The Protocol Opcode is interpreted relative to the Matter protocol specified in the [Protocol ID](#) field.

Opcodes are defined by the corresponding Protocol specification, for example [Secure Channel Protocol](#).

4.4.3.3. Exchange ID (16 bits)

An unsigned integer value identifying the exchange to which the message belongs. An Exchange ID is allocated by the initiator of the exchange, and is unique within the initiator exchange number space as specified in [Section 4.10.2, “Exchange ID”](#).

4.4.3.4. Protocol ID (16 bits)

An unsigned integer value identifying the protocol in which the [Protocol Opcode](#) of the message is defined.

When the [Protocol Vendor ID](#) is the [Matter Standard Vendor ID](#), the Protocol ID SHALL have one of the values specified by [Table 14](#), “Protocol IDs for the Matter Standard Vendor ID”.

Table 14. Protocol IDs for the [Matter Standard Vendor ID](#)

Range	Type	Message Specification
0x0000	PROTOCOL_ID_SECURE_CHANNEL	Section 4.11.1 , “Secure Channel Protocol Messages”
0x0001	PROTOCOL_ID_INTERACTION_MODEL	Section 10.2.1 , “IM Protocol Messages”
0x0002	PROTOCOL_ID_BDX	Section 11.22.3.1 , “BDX Protocol Messages”
0x0003	PROTOCOL_ID_USER_DIRECTED_COMMISSIONING	Section 5.3.2 , “UDC Protocol Messages”
0x0004	PROTOCOL_ID_FOR_TESTING	Reserved for bespoke protocols run in an isolated test environment.
0x0005 - 0xFFFF	reserved	reserved

4.4.3.5. Protocol Vendor ID (16 bits)

An optional, unsigned integer value that contains the Vendor ID namespacing for the [Protocol ID](#) field. This field SHALL only be present when the [V Flag](#) is set; otherwise the default is 0, corresponding to the [Matter Standard Vendor ID](#).

4.4.3.6. Acknowledged Message Counter (32 bits)

An optional, unsigned integer value containing the message counter of a previous message that is being acknowledged by the current message. The Acknowledged Message Counter field is SHALL only be present when the [A Flag](#) in the Exchange Flags field is 1.

4.4.3.7. Secured Extensions (variable)

The Secured Extensions field is a variable length block of data for providing backwards compatible extensibility. The format of the Secured Extensions block is shown in [Table 15](#), “Secured Extensions block format definition”. The Secured Extensions block SHALL be present only if the [SX Flag](#) is set to 1 in the Exchange Flags field.

Version 1.0 Nodes SHALL ignore the contents of the Secured Extensions payload.

The Secured Extensions block SHALL be encrypted and authenticated based on the Security Flags settings.

Table 15. Secured Extensions block format definition

Length	Field
2 bytes	Secured Extensions Payload Length, in bytes.
variable	[Secured Extensions Payload]

4.4.3.8. Application Payload (variable length)

A sequence of zero or more bytes containing the application data conveyed by the message.

4.4.4. Message Size Requirements

Support for IPv6 fragmentation is not mandatory in Matter, and the expected supported MTU is 1280 bytes, the minimum required by IPv6. Therefore, all messages, including transport headers, SHALL fit within that minimal IPv6 MTU. This message size limit SHALL apply to the UDP transport. A message received over UDP that exceeds this message size limit SHALL NOT be processed. Messages sent over TCP, [PAFTP](#), [NTL](#), or [BTP](#) transports MAY exceed the message size limit if both nodes are capable of supporting larger message sizes.

4.5. Message Framing Over Stream-Oriented Transports

When Matter messages are transferred over stream-oriented transport protocols, such as TCP, [PAFTP](#), or [BTP](#), they need to be framed appropriately to enable the receiver to read each message from the stream. To allow that, each Matter Message SHALL be prepended with a [Message Length](#) field. This field SHALL only be present when the message is being transmitted over a stream-oriented channel. When transmitted over a datagram channel, such as UDP or [NTL](#), the message length SHALL be conveyed by the underlying channel. For example, when transmitted over UDP, the message length SHALL be equal to the payload length of the UDP datagram.

4.5.1. Message Length (16/32 bits)

An optional unsigned integer value, in little-endian byte order, specifying the overall length of the message in bytes, not including the size of this field itself. The size of this field, when present, SHALL depend on the transport protocol being used to transfer the message. For example, for TCP, it SHALL be set to 4 bytes to allow for large payloads, whereas for [PAFTP](#) and [BTP](#) it SHALL be set to 2 bytes.

Protocol	Size	Location Description in Stream
TCP	4 bytes	Field present before the Message Header of each Matter message.
BTP	2 bytes	Field present before the segment payload in the beginning BTP Packet PDU for each Matter message/BTP SDU.
PAFTP	2 bytes	Field present before the segment payload in the beginning PAFTP Packet PDU for each Matter message/PAFTP SDU.
MRP	N/A	Field not present.
NTL	N/A	Field not present.

4.6. Message Counters

All messages contain a 32-bit message counter assigned by the sender of the message. Message counters are assigned sequentially, by monotonically increasing the counter value maintained by

the sender of the message. Message counters serve several purposes:

- **Duplicate Message Detection** – Receiving systems use message counters to detect messages that have been retransmitted by the sender, e.g. in response to packet loss in the network.
- **Message Acknowledgement** – In the Message Reliability Protocol (MRP), message counters provide a way for receivers to identify messages for the purpose of acknowledging their receipt.
- **Encryption Nonces** – When encrypted messages are sent, message counters provide an encryption nonce that ensures each message is encrypted in a unique manner.
- **Replay Prevention** – Related to encryption, message counters also provide a means for detecting and preventing the replay of encrypted messages.

4.6.1. Message Counter Types

All Nodes implement three global 32-bit counters to generate message counters for certain types of messages:

- *Global Unencrypted Message Counter*
- *Global Group Encrypted Data Message Counter*
- *Global Group Encrypted Control Message Counter*

Additionally, Nodes implement a separate 32-bit counter for each session as part of secure session state:

- *Secure Session Message Counter*

Additionally, Nodes implement a separate 32-bit counter for each [Check-In Protocol](#) use case:

- *Check-In Counter*

Technical details for how each counter type works are described in the following sections. [Table 16, “Message Counter Type Overview”](#) is provided to summarize higher-level differences between Message Counter Types:

Table 16. Message Counter Type Overview

Message Counter Type	Session Type	Lifetime	Rollover Policy	Nonvolatile
Global Unencrypted	Unsecured	Unlimited	Allowed	Optional
Global Encrypted Data	Group	Operational Group Key	Allowed	Mandatory
Global Encrypted Control	Group	Operational Group Key	Allowed	Mandatory
Secure Session	Unicast	Session Key	Expires	Optional
Check-In Counter	Unsecured	Unlimited	Allowed	Mandatory

4.6.1.1. Message Counter Initialization

All message counters SHALL be initialized with a random value using the `Crypto_DRBG(len = 28) + 1` primitive. Message counters are initialized to a random number to increase the difficulty of traffic analysis attacks by making it harder to determine how long a particular session has been open. The random initializer ranges from 1 to 2^{28} in order to maximize initial entropy while still reserving the vast majority of the range to actual counter values (roughly $2^{32} - 2^{28}$).

4.6.1.2. Global Unencrypted Message Counter

All Nodes SHALL implement an unencrypted message counter, which is used to generate counters for unencrypted messages.

Typically, Nodes store the *Global Unencrypted Message Counter* in RAM. This makes the counter subject to loss whenever the system reboots or otherwise loses its state. This is permissible because retaining the *Global Unencrypted Message Counter* is not essential to the confidentiality or integrity of the message. In the event that the *Global Unencrypted Message Counter* for a Node is lost, Nodes SHALL randomize the initial value of this counter on startup per [Section 4.6.1.1, “Message Counter Initialization”](#).

4.6.1.3. Global Group Encrypted Message Counters

The *Global Group Encrypted Message Counters* are used to generate the counter for messages encrypted using group keys. There are two such counters:

- The *Global Group Encrypted Data Message Counter* is used to encode regular data messages encrypted with a group key.
- The *Global Group Encrypted Control Message Counter* is used to encode [control messages](#) encrypted with a group key.

Some Nodes might not be required to implement communication using group keys, in which case they MAY omit the *Global Group Encrypted Message Counters*. In contrast to the *Global Unencrypted Message Counter*, Nodes are required to persist the *Global Group Encrypted Message Counters* in durable storage. In particular, Nodes are required to ensure that the value of the *Global Group Encrypted Message Counters* never rolls back and that it is monotonic within the bounds of its range for its use for a given group key. A Node SHALL randomize the initial value of this counter on factory reset per [Section 4.6.1.1, “Message Counter Initialization”](#).

While *Global Group Encrypted Message Counters* are shared by many group keys to generate nonces, rollover is not an issue as long as the [Epoch Key](#) that generates each [operational group key](#) rotates frequently enough.

NOTE

If a nonce is duplicated for a given key, the security consequences are isolated only to the specific key with which the duplicate nonce occurred—a key that has not been updated prior to rollover and has been presumably abandoned or aged out.

4.6.2. Secure Session Message Counters

A *Secure Session Message Counter* is a per-session, 32-bit, ephemeral counter that is used by the

encoding of any encrypted messages using an associated session key. Each peer in a [Secure Unicast Session](#) SHALL maintain its own message counters, with independent counters being used for each unique session used. *Session Message Counters* SHALL exist for as long as the associated security session is in effect. A Node SHALL randomize the initial value of this counter on session establishment per [Section 4.6.1.1, “Message Counter Initialization”](#).

The *Secure Session Message Counter* history window SHALL be maintained for the lifetime of the session, and SHALL be deleted at the same time as the session keys, when the session ends.

Sessions SHALL be discarded and re-established before any *Secure Session Message Counter* overflow or repetition occurs.

4.6.3. Check-In Counter

The Check-In Counter is an unsigned 32-bit counter used during the encryption process of the [Check-In Protocol](#). The only purpose of the Check-In Counter is during the encryption process of the [Check-In Protocol](#).

Each Check-In Protocol use-case implemented by a device SHALL be associated with a specific Check-In counter to be used for each node identity (pair of fabric and node identifier). This MAY be a single counter across all node identities, or MAY be one counter per node identity, or anything in between, as long as each node identity uses a single specific instance of the counter. The device SHALL randomize the initial value of the counter on factory reset per [Section 4.6.1.1, “Message Counter Initialization”](#).

The Check-In Counter SHALL be monotonically increased each time a [Check-In message](#) is sent. This monotonicity guarantee SHALL be preserved across idle and active states. The Check-In Counter can roll over to zero when it exceeds the maximum value of the counter (32-bits). Nodes are required to persist the *Check-In Counter* in durable storage.

When a node reboots, the *Check-In Counter* MAY increase by a larger step than 1. Nodes do not need to write every new value of the Check-In Counter to permanent storage each time it is increased (e.g. to prevent flash wear due to many write operations). One example strategy to achieve reduction of non-volatile storage updates is described below:

1. Read the **counter** value at start-up.
2. Before processing the first message after startup, write **counter + N** to storage, where N is a carefully chosen number (e.g. 1000). This number N should be chosen carefully in order not to exhaust the lifetime 32-bit counter space.
3. Process messages normally until the **counter** has a value one less than the **counter** in storage. When this happens, store **counter + N** to storage.

NOTE

The Check-In Counter has an unlimited lifetime until the device is factory reset. To ensure that a nonce is not reused since it is derived from the Check-In Counter and the ICD key, the key needs to be refreshed before using all valid counter values for the key.

4.6.4. Message Counters as Encryption Nonces

In the context of encrypted messages, message counters serve as nonces for the encryption algorithm, ensuring that every message is encrypted in a unique manner. The uniqueness of an encrypted message's counter is vital to the confidentiality of the message, as the encryption algorithm makes it trivial for an eavesdropper to decrypt messages if the attacker can find two different messages with the same message counter that were encrypted using the same key. Specifically, an attacker can XOR the two different messages that share the same key and nonce to obtain a "block key" which can be used to decrypt any message that uses that key and nonce.

Nodes SHOULD rotate their encryption keys on a regular basis, to ensure that old encryption keys are retired before a *Global Group Encrypted Message Counter* has a chance to wrap to a value previously used with the encryption key. In practice, the frequency of message transmission is such that encryption keys generally rotate at a rate that is much faster than the rate at which a *Global Group Encrypted Message Counter* wraps. In the event that a *Global Group Encrypted Message Counter* wraps before the associated keys are rotated, all keys associated with that *Global Group Encrypted Message Counter* are considered exhausted and are no longer valid to use. In such cases, a new unicast session SHALL be established to the Matter Node to rotate such retired keys before secure communication can resume. Given the importance of confidentiality and message integrity, every effort SHOULD be made to ensure that keys are rotated on a regular basis.

4.6.5. Replay Prevention and Duplicate Message Detection

Beyond their role as encryption nonces, message counters also serve as a means to detect repeated reception of the same message. Message duplication may occur for a number of reasons: out-of-order arrival, network latency, malicious attack, or network error. For example, a duplicate can be caused when a sender retransmits a message after failing to receive an acknowledgement, or because a malicious third party attempted to replay an old message to gain some advantage. To detect duplicate messages, Nodes maintain a history window of the message counters they have received from a particular sender (see [Section 4.6.5.1, "Message Reception State"](#)). Whenever a message is received, its message counter is checked against the history window of message counters from that sender to determine whether it is a duplicate. The Message Layer SHALL discard duplicate messages before they reach the application layer.

4.6.5.1. Message Reception State

The state maintained by a Node about the messages it has received from a particular peer is referred to as *message reception state*. Nodes use this state information to determine whether a newly arrived message is a duplicate of a previous message. Message reception state is maintained on a per-peer or per-session basis, depending on the type of message encryption being used.

At a conceptual level, message reception state consists of a set of integers corresponding to the counters of all the messages that have been received from a particular peer. To limit the amount of memory required to store this state, Nodes employ a lossy compression scheme that takes advantage of the fact that message counters are generated sequentially by the sender. The scheme allows for a limited amount of out-of-order message arrivals due to network effects without inducing false detection of duplicates.

In the compressed form, message reception state is structured as a pair of values: a integer repre-

sending the largest valid, or maximum message counter received from the peer (**max_message_counter**), and a bitmap of size **MSG_COUNTER_WINDOW_SIZE** indicating which messages immediately prior to the max message have been received. The offset into the bitmap equates to the difference between the corresponding message counter and the max message counter, i.e. the first bit in the bitmap indicates whether the message with the counter of **max_message_counter - 1** has been received, the second indicates whether message **max_message_counter - 2** has been received, and so on. A message counter is within the range of the bitmap, also known as the message counter window, when the counter value is between $[(\text{max_message_counter} - \text{MSG_COUNTER_WINDOW_SIZE}) \text{ to } (\text{max_message_counter} - 1) \bmod 2^{32}]$. As messages arrive, the message reception state is queried to determine if an arriving message is new or duplicate. If a message is new, the state is then updated to reflect the arrival of the message. When a message arrives with a message counter that is logically greater than the current maximum message counter for that peer, the maximum message counter value for the peer is updated and the bitmap shifted accordingly.




Figure 7. Message Reception State Example

4.6.5.2. Use of Message Reception State for Encrypted Messages

The algorithm for querying and updating message reception state varies slightly depending on whether the system is tracking reception of encrypted messages or unencrypted messages.

Message Counters with maximum

For encrypted messages of **Secure Unicast Session Type**, any arriving message with a counter in the range $[(\text{max_message_counter} + 1) \text{ to } (2^{32} - 1)]$ SHALL be considered new, and cause the **max_message_counter** value to be updated. Message counters within the range of the bitmap SHALL be considered duplicate if the corresponding bit offset is set to true. All other message counters SHALL be

considered duplicate.

Message Counters with rollover

A message counter with rollover is a free running message counter that monotonically increases, but rolls over to zero when it exceeds the maximum value of the counter (32-bits). Group keys are secured by a shared, global message counter with rollover as described in [Section 4.6.1.3, “Global Group Encrypted Message Counters”](#).

For encrypted messages of [Group Session Type](#), any arriving message with a counter in the range $[(\text{max_message_counter} + 1) \text{ to } (\text{max_message_counter} + 2^{31} - 1)] \text{ (modulo } 2^{32})$ SHALL be considered new, and cause the `max_message_counter` value to be updated. Messages with counters from $[(\text{max_message_counter} - 2^{31}) \text{ to } (\text{max_message_counter} - \text{MSG_COUNTER_WINDOW_SIZE} - 1)] \text{ (modulo } 2^{32})$ SHALL be considered duplicate. A message counter equal to `max_message_counter` SHALL be considered duplicate. Message counters within the range of the bitmap SHALL be considered duplicate if the corresponding bit offset is set to true.

This scheme for encrypted messages effectively divides the message counter space in half: those counters that are forward of the max message counter, which are considered new, and those counters that are behind the max message counter, which are considered duplicates unless indicated otherwise by the values in the bitmap.

4.6.5.3. Use of Message Reception State for Unencrypted Messages

For unencrypted messages, the algorithms for tracking messages and detecting duplicates are similar to, but more permissive than for encrypted messages using [Section 4.6.5.2.2, “Message Counters with rollover”](#). This reflects the fact that duplicate detection of unencrypted messages is not done for security reasons, but rather to catch duplicates caused by network errors (e.g. loss of an ack), which are generally more bounded in time. The more relaxed algorithm for unencrypted duplicate detection also relaxes the durability requirement on the sender’s message counter, allowing senders to store the counter in RAM.

For unencrypted messages, any message counter equal to `max_message_counter` or within the message counter window, where the corresponding bit is set to true SHALL be considered duplicate. All other message counters, whether behind the window or ahead of `max_message_counter`, are considered new and SHALL update `max_message_counter` and shift the window accordingly. Messages with a counter behind the window are likely caused by a node rebooting and are thus processed as rolling back the window to the current location. Note that when `max_message_counter` is close to the minimum of the range, the window SHALL roll back to cover message counters near the maximum of the range.

4.6.5.4. Message Reception State Initialization

To initialize [Message Reception State](#) for a given Peer Node ID, initial `max_message_counter`, Message Type (control or data), Encryption Level (encrypted or unencrypted), and Encryption Key (for any Encryption Level except unencrypted):

- The Message Reception State fields SHALL be set as follows:
 - The Peer Node ID SHALL reference the given Peer Node ID.

- The Message Type SHALL be the given Message Type.
- The Encryption Level SHALL be the given Encryption Level.
- If the Encryption Level is NOT unencrypted, the Encryption Key SHALL reference the given key.
- The **max_message_counter** SHALL be set to the given **max_message_counter**.
- The Message Counter bitmap SHALL be set to all **1**, indicating that only new messages with counter greater than **max_message_counter** SHALL be accepted.

4.6.6. Counter Processing of Outgoing Messages

1. Obtain the outgoing message counter of the sending Node for the given Security Flags, Session Id, and Encryption Key:
 - a. A message of **Unsecured Session Type** SHALL use the current *Global Unencrypted Message Counter*.
 - b. A message of **Secure Unicast Session Type** SHALL use the current *Secure Session Message Counter* for the session associated with the Session ID.
 - c. A message of **Group Session Type** SHALL use:
 - i. The *Global Group Encrypted Data Message Counter* if the Security Flags **C Flag** = **0**.
 - ii. The *Global Group Encrypted Control Message Counter* if the Security Flags **C Flag** = **1**.
2. The outgoing message counter from step 1 SHALL be incremented by 1.
3. Store the incremented outgoing message counter in the *OutgoingMessageCounter* element associated with the **Session Context** for the message.
 - a. If the message counter wraps around from 0xFFFF_FFFF to 0x0000_0000 and the message is of **Secure Unicast Session Type**:
 - i. The Encryption Key SHALL be expired in the **Session Context**. The session will need to be renegotiated to resume communication after transmission of this final message.

4.6.7. Counter Processing of Incoming Messages

1. Determine the **Message Reception State** for the sending peer and get the current **max_message_counter**.
 - a. Given a decrypted message of **Unicast Session Type**:
 - i. Get the session-specific **Message Reception State** from the **Secure Unicast Session Context**.
 - b. Given a decrypted message of **Group Session Type**:
 - i. Extract the Source Node ID from the **Message Header**.
 - A. If there is no Source Node ID for the message, drop the message.
 - ii. Get the **Message Reception State** for the **Source Node ID** of the given message:
 - A. If the Security Flags **C Flag** = **0**, get the Data Message Reception State for the peer node.

- B. If the Security Flags **C Flag** = 1, get the Control Message Reception State for the peer node.
 - iii. If there is no **Message Reception State** for the groupcast message, initiate [Section 4.18.4, “Unsynchronized Message Processing”](#).
 - c. Given an unencrypted message:
 - i. Get the **Message Reception State** associated with the **Unsecured Session Context**.
 - ii. If there is no **Message Reception State** for the unencrypted message, **create it** with the information from the given message.
 2. If the Message Counter is outside the valid message counter window, the message SHALL be marked as a duplicate. Note that while messages may be outside of the window for reasons other than being a duplicate, and we always mark them as such.
 3. If the message is a duplicate:
 - a. If the message is marked as encrypted, follow [Section 4.6.5.2, “Use of Message Reception State for Encrypted Messages”](#).
 - b. If the message is marked as unencrypted, follow [Section 4.6.5.3, “Use of Message Reception State for Unencrypted Messages”](#).
 - c. If the message is encrypted and marked as a duplicate, i.e. Message Counter is outside the valid message counter window or marked as previously received in the **Message Reception State**:
 - i. Perform [Section 4.12.5.2, “Reliable Message Processing of Incoming Messages”](#) on the duplicate message.
 - d. Otherwise, update the Message Reception State as detailed in [Section 4.6.5.1, “Message Reception State”](#), and accept the message for further processing.

4.7. Message Processing

This sub-clause describes the fundamental procedures for transmission and reception.

4.7.1. Message Transmission

To prepare a message for transmission with a given Session ID, Destination Node ID (which may be a **group node id** or an **operational node id**) and Security Flags, the following steps SHALL be performed, in order:

1. Process the message as described in [Section 4.6.6, “Counter Processing of Outgoing Messages”](#).
2. If the message’s **Session Type** is a **Unicast Session**:
 - a. Set **SessionTimestamp** to a timestamp from a clock which would allow for the eventual determination of the last session use relative to other sessions.
 - b. Process the message as described in [Section 4.8.2, “Security Processing of Outgoing Messages”](#).
 - c. Process the message as described in [Section 4.9.3, “Privacy Processing of Outgoing Messages”](#).

4.7.2. Message Reception

To process a received message, the following steps SHALL be performed in order:

1. Perform validity checks on the message; if any fail, processing of the message SHALL stop, and a 'message invalid' error SHOULD be indicated to the next higher layer:
 - a. The **Version** field SHALL be 0.
 - b. If the message is of **Secure Unicast Session Type**:
 - i. The **DSIZ** field SHALL NOT indicate a Group ID is present.
 - c. If the message is of **Group Session Type**:
 - i. The **DSIZ** field SHALL NOT be 0.
 - ii. The **S Flag** field SHALL NOT be 0.
2. If the message is NOT of **Unsecured Session Type**:
 - a. Obtain the Privacy and Encryption Keys associated with the given Session ID:
 - i. If no keys are found, security processing SHALL indicate a failure to the next higher layer with a status of 'message security failed' and no further security processing SHALL be done on this message.
 - b. For each Privacy and Encryption Key, of which there may be more than one in the case of group messages:
 - i. If the **P Flag** is set, follow [Section 4.9.4, “Privacy Processing of Incoming Messages”](#) to deobfuscate the message.
 - ii. Follow [Section 4.8.3, “Security Processing of Incoming Messages”](#) to decrypt and authenticate the message.
3. Follow [Section 4.6.7, “Counter Processing of Incoming Messages”](#) to enforce replay protection and duplicate detection.
4. If the message transport is UDP, follow [Section 4.12.5.2, “Reliable Message Processing of Incoming Messages”](#) to process message reliability.
5. If the message's **Session Type** is a **Unicast Session**:
 - a. Set **SessionTimestamp** and **ActiveTimestamp** to a timestamp from a clock which would allow for the eventual determination of the last session use relative to other sessions.
6. The received message is then delivered to [Section 4.10.5, “Exchange Message Processing”](#).

4.8. Message Security

The detailed steps involved in security processing of outgoing and incoming Matter messages are described in [Section 4.8.2, “Security Processing of Outgoing Messages”](#) and [Section 4.8.3, “Security Processing of Incoming Messages”](#) respectively. [Section 4.8.1, “Data confidentiality and integrity with data origin authentication parameters”](#) defines how the cryptographic parameters are set for securing Matter messages.

4.8.1. Data confidentiality and integrity with data origin authentication parameters

This section specifies the parameters to use the data confidentiality and integrity cryptographic primitive as defined in Section 3.6, “Data Confidentiality and Integrity”.

The parameters in this section SHALL apply for all encrypted messages, i.e. all messages except those of Unsecured Session Type.

4.8.1.1. Nonce

The nonce SHALL be formatted as specified in Table 17, “Nonce”.

Table 17. Nonce

Octets: 1	4	8
Security Flags	Message Counter	Source Node ID

The nonce used for the Authenticated Encryption with Additional Data (AEAD) algorithm (see Section 3.6, “Data Confidentiality and Integrity”) for a given message SHALL be defined as the concatenation of the Security Flags, the Message Counter, and the Source Node ID of that message. The scalar fields in the nonce, namely the Message Counter and the Source Node ID SHALL be encoded in little-endian byte order for the purposes of serialization within the nonce, that is, in the same byte ordering as the segment of the message from which its data originates.

The Source Node ID field used in the nonce SHALL be set to the Operational Node ID of the node originating security protection of the message:

- If the message is of Secure Unicast Session Type:
 - For a CASE session, the Nonce Source Node ID SHALL be determined via the Secure Session Context associated with the Session Identifier.
 - For a PASE session, the Nonce Source Node ID SHALL be Unspecified Node ID.
- If the message is of Group Session Type:
 - The S Flag of the message SHALL be 1 and the Nonce Source Node ID SHALL be the Source Node ID of the message.
 - If the S Flag of the message is 0 the message SHALL be dropped.

NOTE

Because PASE negotiates strong one-time keys per session and the I2RKey and R2IKey are distinct for each direction of communication, the use of the Unspecified Node ID as the Nonce Source Node ID remains semantically secure.

4.8.2. Security Processing of Outgoing Messages

The process for encrypting Matter messages is depicted graphically in Figure 8, “Matter Message Encryption” with color code conventions described in Figure 9, “Matter Message Encryption Legend”.




Figure 8. Matter Message Encryption




Figure 9. Matter Message Encryption Legend

To prepare a secure message for transmission with a given Session ID, Destination Node ID (which may be a [group node id](#) or an [operational node id](#)) and Security Flags, the Node SHALL perform the following steps:

1. Obtain the Encryption Key associated with the Session ID and Destination Node ID and the Session Type associated with the Destination Node ID:
 - a. If no key is found for the given Session ID, security processing SHALL fail and no further security processing SHALL be done on this message.
2. Obtain the outgoing message counter of the sending Node as per [Section 4.6.6, “Counter Process-](#)

ing of Outgoing Messages”.

3. Set the Security fields as follows:
 - a. The **Session ID** field SHALL be set to the value provided to step 1.
 - b. The **Security Flags** field SHALL be set to the value provided to step 1 with the following sub-fields updated:
 - i. The **Session Type** field SHALL be set to the value obtained from step 1.
4. Set the **Message Flags**, **Destination Node ID**, and **Source Node ID** fields as follows:
 - a. If the Session Type is a unicast session:
 - i. Set **S Flag** to 0.
 - ii. Set **DSIZ** to 0.
 - iii. Omit both **Destination Node ID**, and **Source Node ID**.
 - b. Else if the **Session Type** is a group session:
 - i. Set **S Flag** to 1.
 - ii. Set **DSIZ** to 2.
 - iii. Set **Source Node ID** field to the **operational node id** of the sending node.
 - iv. Set **Destination Node ID** field to the 16-bit Group ID derived from the Destination Node ID.
5. Set the **Message Counter** field to the outgoing message counter from step 2.
6. Execute the AEAD generate and encrypt operation, as specified in [Section 3.6.1, “Generate and encrypt”](#), with the following instantiations:
 - a. The bit string key *K* SHALL be the Encryption Key obtained from step 1;
 - b. The nonce *N* SHALL be the **CRYPTO_AEAD_NONCE_LENGTH_BYTES**-octet string constructed according to [Table 17, “Nonce”](#);
 - c. The parameter *P* SHALL be the Message Payload;
 - d. The additional data octet string *A* SHALL be the Message Header contents, using little-endian byte order for all scalars, exactly as they appear in the message segments from which they originate:

Message Flags || Session ID || Security Flags || Message Counter

with the optional fields appended according to the Message Flags:

[Source Node ID] || [Destination Node ID] || [Message Extensions]

- e. $C = \text{Crypto_AEAD_GenerateEncrypt}(K, P, A, N)$
7. If the AEAD operation invoked in step 6 results in an error, then security processing SHALL fail and no further security processing SHALL be done on this message.
8. Let *C* be the output from step 6. *C* contains the tag of **CRYPTO_AEAD_MIC_LENGTH_BITS** bits (Message

Integrity Check (MIC)) as specified by [Section 3.6.1, “Generate and encrypt”](#). The secured outgoing message SHALL be:

A || C

4.8.3. Security Processing of Incoming Messages

All incoming message processing SHALL occur in a serialized manner. If an implementation chooses to process messages in a parallel manner, it must ensure that the behavior is opaque-box identical to a serialized processing implementation.

If the transport layer receives a secured message as indicated by the Session ID, it SHALL perform the following steps:

1. Determine the [Session Type](#), [Session ID](#), and [Message Counter](#) from the message header of the received message.
2. Obtain the Encryption Key associated with the [Session Context](#) of the given Session ID and Session Type:
 - a. If no key is found for the given Session ID, security processing SHALL indicate a failure to the next higher layer with a status of 'message security failed' and no further security processing SHALL be done on this message.
3. Execute the AEAD decryption and verification operation as specified in [Section 3.6.2, “Decrypt and verify”](#) with the following instantiations:
 - a. The bit string key *K* SHALL be the Encryption Key obtained from step 2;
 - b. The nonce *N* SHALL be the `CRYPTO_AEAD_NONCE_LENGTH_BYTES`-octet string constructed according to [Table 17, “Nonce”](#);
 - c. The parameter *C* SHALL be the encrypted and authenticated Message Payload;
 - d. The additional data octet string *A* SHALL be the authenticated Message Header;
 - e. `{success, P} = Crypto_AEAD_DecryptVerify(K, C, A, N)`
4. Return the result `{success, P}` of the AEAD operation:
 - a. If the `success` is `FALSE`, security processing SHALL fail and further processing SHALL NOT be performed on this message. An appropriate error SHOULD be raised to the upper layer to indicate the error.
 - b. Otherwise, set the octet string *PlaintextMessage* to the string

A || P

5. *PlaintextMessage* now represents the deciphered, authenticated, received message.
 - a. NOTE: The message has not yet undergone counter processing nor replay detection.
 - b. The *PlaintextMessage* SHALL be marked as successfully security processed and SHALL be released to the next processing layer.

4.9. Message Privacy

Privacy processing of a message describes the obfuscation and deobfuscation of the message header fields after encryption and before decryption.

The detailed steps involved in privacy processing of outgoing and incoming Matter messages are described in [Section 4.9.3, “Privacy Processing of Outgoing Messages”](#) and [Section 4.9.4, “Privacy Processing of Incoming Messages”](#) respectively. They rely on the cryptographic primitives in [Section 3.7, “Message privacy”](#).

4.9.1. Privacy Key

The Privacy Key is a symmetric key specifically used for [Privacy Processing](#) that is derived from the [EncryptionKey](#) used for [Security Processing](#) of a given message. Given a Session ID reference to a specific Encryption Key, the Privacy Key is derived as follows:

```
PrivacyKey =
    Crypto_KDF
    (
        InputKey = EncryptionKey,
        Salt = [],
        Info = "PrivacyKey",
        Length = CRYPTO_SYMMETRIC_KEY_LENGTH_BITS
    )
```

4.9.2. Privacy Nonce

The Privacy Nonce is a nonce specifically used for [Privacy Processing](#) that is derived from the [SessionID](#) and [MIC](#) of the message. The Privacy Nonce SHALL be the [CRYPTO_AEAD_NONCE_LENGTH_BYTES](#)-octet string constructed as the 16-bit Session ID (in big-endian format) concatenated with the lower 11 (i.e. [CRYPTO_AEAD_MIC_LENGTH_BYTES-5](#)) bytes of the MIC:

```
PrivacyNonce = Session ID || MIC[5..15]
```

For example if Session ID is 42 (i.e. 0x002A) and the computed MIC is c5:a0:06:3a:d5:d2:51:81:91:40:0d:d6:8c:5c:16:3b:

```
Session ID = 00:2a
```

```
MIC = c5:a0:06:3a:d5:d2:51:81:91:40:0d:d6:8c:5c:16:3b
```

```
MIC[5..15] = d2:51:81:91:40:0d:d6:8c:5c:16:3b
```

```
PrivacyNonce = SessionID || MIC[5..15] = 00:2a || d2:51:81:91:40:0d:d6:8c:5c:16:3b
```

```
PrivacyNonce = 00:2a:d2:51:81:91:40:0d:d6:8c:5c:16:3b
```

4.9.3. Privacy Processing of Outgoing Messages

The process for privacy encoding Matter message headers is depicted graphically in [Figure 10](#), “Matter Message Privacy”.




Figure 10. Matter Message Privacy

To apply privacy obfuscation to an encrypted message prepared for transmission by [Section 4.7.1](#), “Message Transmission”, apply obfuscation steps as follows:

1. If **P Flag** is not set, do nothing.
2. Obtain the Privacy Key for the Encryption Key used to secure the message.
3. Execute the encryption operation as specified in [Section 3.7.1](#), “Privacy encryption” with the following instantiations:
 - a. The bit string key *K* SHALL be the Privacy Key obtained from step 1;
 - b. The MIC SHALL be the last **CRYPTO_AEAD_MIC_LENGTH_BYTES** bytes of the **C** outcome of the message security protection as specified in [Section 4.8.2](#), “Security Processing of Outgoing Messages” ($MIC = C[(CRYPTO_AEAD_MIC_LENGTH_BYTES-1)..0]$)
 - c. The nonce *N* SHALL be the **PrivacyNonce** of the message.
 - d. The parameter *M* SHALL be the message header fields where optional fields are only present in *M* if they are present in the message:

$M = \text{Message Counter} || [\text{Source ID}] || [\text{Destination ID}]$

- e. $CP = \text{Crypto_Privacy_Encrypt}(K, M, N)$

4. Let *CP* be the obfuscated output from step 2. *CP* SHALL be used in the final private message in place of the message header fields.

4.9.4. Privacy Processing of Incoming Messages

To deobfuscate a private message received by [Section 4.7.2, “Message Reception”](#) with a given Privacy Key, perform security processing as follows:

1. If **P Flag** is not set, do nothing.
2. With the given Privacy Key, execute the decryption as specified in [Section 3.7.2, “Privacy decryption”](#) with the following instantiations:
 - a. The bit string key K SHALL be the Privacy Key obtained from step 1;
 - b. The MIC SHALL be the last **CRYPTO_AEAD_MIC_LENGTH_BYTES** bytes of the **C** outcome of the message security protection as specified in [Section 4.8.3, “Security Processing of Incoming Messages”](#) ($MIC = C[(CRYPTO_AEAD_MIC_LENGTH_BYTES-1)..0]$)
 - c. The nonce N SHALL be the **PrivacyNonce** of the message.
 - d. The parameter CP SHALL be the message header fields where optional fields are only present in CP if they are present in the message:

$CP = \text{Message Counter} || [\text{Source ID}] || [\text{Destination ID}]$

- e. $M = \text{Crypto_Privacy_Decrypt}(K, CP, N)$
3. Let M be the deobfuscated output from step 2.
 - a. M SHALL be used in the final private message in place of the message header fields.
 - b. The first successfully validated message, M , by [Section 4.8.3, “Security Processing of Incoming Messages”](#) SHALL terminate iteration through Privacy Keys in step 2.

4.10. Message Exchanges

An Exchange provides a way to group related messages together, organize communication flows, and enable higher levels of the communication stack to track semantically relevant groupings of messages.

An Exchange SHALL be bound to exactly one underlying session that will transport all associated Exchange messages for the life of that Exchange. The underlying session SHALL be one of the following session types: **secure unicast** (as established by PASE or CASE), **unsecured** (as is used for the initial session establishment phase of a PASE/CASE session), **secure group**, or **MCSP**.

When used with **reliability**, Exchanges assume basic flow control by the upper layer. The Exchange Layer SHALL NOT accept a message from the upper layer when there is an outbound reliable message pending on the same Exchange.

4.10.1. Exchange Role

The first Node to send a message in an Exchange is said to be in the Initiator role, and all the other Nodes that subsequently participate in the Exchange are said to be in a Responder role. An Exchange is always between one Initiator and one or more peer Responder Nodes. An Exchange

does not survive a reboot of one of the participants. Adjacent layers MAY close an Exchange at any time.

4.10.2. Exchange ID

An Exchange of messages is identified by the Exchange ID field described in [Section 4.4.3.3, “Exchange ID \(16 bits\)”](#). The Exchange ID is allocated by the Initiator. The first message the Initiator sends in a new Exchange SHALL contain a fresh value for the Exchange ID field. The Exchange is then identified by the tuple {[Session Context](#), [Exchange ID](#), [Exchange Role](#)} where Session Context is one of an [Unsecured](#), [Secured](#), [Groupcast](#) or [MCSP](#) session context. All messages that are part of a given Exchange, whether they are sent by the Initiator or not, share the same Exchange ID, allowing the Initiator and Responder Nodes to match responses to requests or otherwise group messages together that are part of more complex transactions. The first Exchange ID for a given Initiator Node SHALL be a random integer. All subsequent Exchange IDs created by that Initiator SHALL be the last Exchange ID it created incremented by one. An Exchange ID is an unsigned integer that rolls over to zero when its maximum value is exceeded.

A node SHOULD limit itself to a maximum of 5 concurrent Initiator exchanges over a unicast session. This is to prevent a node from exhausting the [Secure Session Message Counter](#) window of the session with the peer node.

4.10.3. Exchange Context

An Exchange context is the metadata tracked for an Exchange by all exchange participants. An Exchange context tracks the following data:

1. [Exchange ID](#): The Exchange ID assigned by the Initiator
2. [Exchange Role](#): Initiator or Responder
3. [Session Context](#): The underlying [Unsecured](#), [Secured](#), [Groupcast](#) or [MCSP](#) session context
 - Together, Session Context, Exchange ID and Role comprise a unique key allowing participants to identify any exchange.

4.10.3.1. Protocol ID Registration

The Interaction Model layer indicates to the Exchange Layer which [Protocols](#) it will accept. Any message for a [Protocol ID](#) that is not registered with the Exchange Layer SHALL be dropped.

4.10.4. Exchange Message Dispatch

When sending a message to the Exchange Layer, the next higher layer SHALL specify whether the message is part of an existing Exchange, or the first of a new Exchange. For the case of a first message, the Initiator creates a new Exchange. The Node in the Initiator role SHALL always set the [I Flag](#) in the [Exchange Flags](#) of every message it sends in that Exchange.

Each Node in a Responder role for an Exchange SHALL use the Exchange ID received in previous messages for the Exchange. Each Node in the Responder role SHALL NOT set the [I Flag](#) in the [Exchange Flags](#) of every message it sends in that Exchange. Each Node in a Responder role SHALL NOT set the [Destination Node ID](#) field to a value that identifies any Node other than the Node in the

Initiator role for the Exchange.

Processing SHALL then proceed to [Section 4.12.5.1, “Reliable Message Processing of Outgoing Messages”](#).

4.10.5. Exchange Message Processing

After completion of [Section 4.7.2, “Message Reception”](#), if the message [matches an existing Exchange](#), it is dispatched to the appropriate protocol handler in the next higher layer. Messages for an existing Exchange are dispatched to the handler for that Exchange. Otherwise, the [unsolicited message](#) that created the Exchange is dispatched to the unsolicited message handler.

4.10.5.1. Exchange Message Matching

Upon receipt of a message, the Exchange Layer attempts to match the message to an existing [Exchange](#). A given message is part of an Exchange if it satisfies all the following criteria:

1. The message was received over the session associated with the Exchange.
2. The Exchange ID of the message matches the Exchange ID of the Exchange,
3. The message has the [I Flag](#) set and the Exchange Role of the Exchange is Responder,
OR the message does not have the [I Flag](#) set and the Exchange Role of the Exchange is Initiator.

If the message does not match an existing Exchange, the message is considered an [unsolicited message](#).

4.10.5.2. Unsolicited Message Processing

An unsolicited message is processed as follows:

1. If the unsolicited message is not marked as having a duplicate message counter, has a registered [Protocol ID](#), and the [I Flag](#) is set:
 - a. Create a [new exchange](#) from the incoming message.
 - b. The new exchange will be used by the upper layer for generating responses and subsequent processing of the message.
2. Otherwise, if the message has the [R Flag](#) set:
 - a. Create an [ephemeral exchange](#) from the incoming message and send an immediate standalone acknowledgement.
 - b. The message SHALL NOT be forwarded to the upper layer, and excluding the sending of an immediate standalone acknowledgment, SHALL be ignored.
 - c. The *ephemeral exchange* created for such duplicate or unknown messages with [R Flag](#) set is automatically closed in [Section 4.12.5.2.2, “Standalone acknowledgement processing”](#).
3. Otherwise, processing of the message SHALL stop.

Creating an Exchange based on an Incoming Message

The steps to create a new Exchange based on an incoming message are as follows:

1. A new Exchange and [Exchange Context](#) SHALL be created with the following settings:
 - a. The Exchange ID SHALL be set to the Exchange ID of the message.
 - b. The Exchange Role SHALL be set to the inverse of the incoming message [I Flag](#), for example set the Exchange Role to Responder if the message is from an Initiator.
 - c. The Session Context SHALL be set to the Session on which the message was received.

4.10.5.3. Closing an Exchange

An Exchange MAY be closed by the application layer or a fatal connection error from the lower message layer. The process of closing an Exchange follows:

1. Any pending acknowledgements associated with the Exchange SHALL be flushed. If there is a pending acknowledgment in the [acknowledgement table](#) for the Exchange and it has [Stand-aloneAckSent](#) set to false:
 - a. Immediately send a [standalone acknowledgement](#) for the pending acknowledgement.
 - b. Remove the [acknowledgement table](#) entry for the pending acknowledgement.
2. Wait for all pending retransmissions associated with the Exchange to complete.
 - a. If the retransmission list for the Exchange is empty, remove the Exchange.
 - b. Otherwise, leave the Exchange open and only close it once the retransmission list is empty.

These steps are depicted in [Figure 11, “Exchange close flow”](#).




Figure 11. Exchange close flow

4.11. Secure Channel Protocol

This section specifies the formal protocol definition for the Secure Channel Protocol. Secure Channel Protocol defines the control plane for secure channel communication and security.

4.11.1. Secure Channel Protocol Messages

Secure Channel Protocol is composed of a collection of sub-protocols, including:

- Message Counter Synchronization Protocol (MCSP)
- Message Reliability Protocol (MRP)
- Passcode Based Session Establishment (PASE)
- Certificate Based Session Establishment (CASE)

The protocol opcodes for messages within the Secure Channel Protocol are grouped based on the underlying sub-protocol that uses the message type. [Table 18, “Secure Channel Protocol Opcodes”](#) lists the messages defined by Secure Channel Protocol.

Table 18. Secure Channel Protocol Opcodes

Protocol Opcode	Protocol Command Name	Description
Protocol ID = PROTOCOL_ID_SECURE_CHANNEL		
0x00	MsgCounterSyncReq	The Message Counter Synchronization Request message queries the current message counter from a peer to bootstrap replay protection.
0x01	MsgCounterSyncRsp	The Message Counter Synchronization Response message provides the current message counter from a peer to bootstrap replay protection.
0x10	MRP Standalone Acknowledgement	This message is dedicated for the purpose of sending a stand-alone acknowledgement when there is no other data message available to piggyback an acknowledgement on top of.
0x20	PBKDFParamRequest	The request for PBKDF parameters necessary to complete the PASE protocol.
0x21	PBKDFParamResponse	The PBKDF parameters sent in response to PBKDF-ParamRequest during the PASE protocol.
0x22	PASE Pake1	The first PAKE message of the PASE protocol.
0x23	PASE Pake2	The second PAKE message of the PASE protocol.
0x24	PASE Pake3	The third PAKE message of the PASE protocol.
0x30	CASE Sigma1	The first message of the CASE protocol.
0x31	CASE Sigma2	The second message of the CASE protocol.
0x32	CASE Sigma3	The third message of the CASE protocol.

Protocol Opcode	Protocol Command Name	Description
0x33	CASE Sigma2_Resume	The second resumption message of the CASE protocol.
0x40	StatusReport	The Status Report message encodes the result of an operation in the Secure Channel as well as other protocols.
0x50	ICD Check-In message	The Check-in message notifies a client that the ICD is available for communication.

4.11.1.1. Session Establishment - Out of Resources

After a successful session establishment using [CASE](#) or [PASE](#), a responder may not have enough resources to save all of the [session context information](#). To free resources, a responder SHALL evict an existing session using the following procedure:

1. Use the [SessionTimestamp](#) to determine the least-recently used session.
2. Determine the session that was least-recently used then:
 - a. Send a [status report](#): `StatusReport(GeneralCode: SUCCESS, ProtocolId: SECURE_CHANNEL, ProtocolCode: CLOSE_SESSION)` message to the peer node
 - b. Remove all state associated with the session (see [Section 4.13.3.1, “Secure Session Context”](#)). The Node MAY save state necessary to perform [Session Resumption](#), see [Section 4.14.2.2.1, “Session Resumption State”](#) for more details.
3. Respond to the initiator with the appropriate session establishment message

4.11.1.2. Status Report

The Status Report message is sent from protocol handlers to convey the status of an operation using a common format as defined in [Appendix D, Status Report Messages](#). The *StatusReport* message is a part of the Secure Channel protocol, but embeds an additional context-specific ProtocolID field in its message-specific payload. In this way, the *StatusReport* can convey status for any protocol handler.

4.11.1.3. Secure Channel Status Report Messages

Status Reports specific to the Secure Channel are designated by embedding the [PROTOCOL_ID_SECURE_CHANNEL](#) in the ProtocolId field of the *StatusReport* body. All Secure Channel Status Report Messages SHALL use the [PROTOCOL_ID_SECURE_CHANNEL](#) protocol id. For example, a failure to find a common root of trust may be written in the specification as follows: `StatusReport(GeneralCode: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: NO_SHARED_TRUST_ROOTS)`.

There are several cases for which the secure channel layer may emit a status report:

1. To indicate successful session establishment
2. In response to errors during session establishment
3. In response to errors after session establishment
4. To indicate that a Node is terminating a session

For each of these cases, a Secure Channel [Status Report](#) message SHALL be sent with an appropriate [ProtocolCode](#) as detailed below.

The following table describes the Secure Channel Status Report [Protocol Specific](#) codes. Each entry in the list details the appropriate [General Code](#) to be utilized with the message and whether it may be sent unencrypted. Secure Channel Status Report messages which are marked as encrypted below SHALL only be sent encrypted in a session established with [CASE](#) or [PASE](#).

Table 19. Secure Channel Protocol Codes

Protocol Code	Error	General Code	Encrypted	Additional Data	Description
0x0000	SESSION_ESTABLISHMENT_SUCCESS	SUCCESS	N	N	Indication that the last session establishment message was successfully processed.
0x0001	NO_SHARED_TRUST_ROOTS	FAILURE	N	N	Failure to find a common set of shared roots.
0x0002	INVALID_PARAMETER	FAILURE	N	N	Generic failure during session establishment.
0x0003	CLOSE_SESSION	SUCCESS	Y	N	Indication that the sender will close the current session. See Section 4.11.1.4, “CloseSession” for more details.
0x0004	BUSY	BUSY	N	Y	Indication that the sender cannot currently fulfill the request. See Section 4.11.1.5, “Busy” for more details.

4.11.1.4. CloseSession

A node may choose to close a session for a variety of reasons including, but not limited to, the following:

1. The interaction between nodes is complete
2. The node needs to free up resources for a new session
3. Fabric configuration associated with the [CASE](#) session was removed with the [RemoveFabric command](#) invoked by an Administrator while the session was open

The [CloseSession](#) StatusReport SHALL only be sent encrypted within an exchange associated with a [PASE](#) or [CASE](#) session. The [CloseSession](#) StatusReport SHALL be sent within a new exchange and SHALL NOT set the [R Flag](#).

If a Node has either sent or received a [CloseSession](#) StatusReport, that Node SHALL remove all state associated with the session (see [Section 4.13.3.1, “Secure Session Context”](#)). The Node MAY save state necessary to perform [Session Resumption](#), see [Section 4.14.2.2.1, “Session Resumption State”](#) for more details.

4.11.1.5. Busy

When a receiver receives a request to start a new secure session via a [Sigma1](#) or [PBKDFParamRequest](#) message, the receiver MAY respond with the [BUSY StatusReport](#) when it is unable to fulfill the request. The [BUSY StatusReport](#) SHALL:

1. Set the [R Flag](#) to 0
2. Set the [S Flag](#) to 0
3. Set the [StatusReport ProtocolData](#) to a 16-bit (two byte) little-endian value indicating the minimum time in milliseconds to wait before retrying the original request.
4. Set the [Exchange ID](#) to the *Exchange ID* present in the [Sigma1](#) or [PBKDFParamRequest](#) message which triggered this response.

For example, a responder wishing to indicate they are unable to fulfill the request and that the initiator should wait 500 milliseconds before trying again would send `StatusReport(GeneralCode: BUSY, ProtocolId: SECURE_CHANNEL, ProtocolCode: BUSY, ProtocolData: [0xF4, 0x01])`.

The [BUSY StatusReport](#) SHALL NOT be sent in response to any message except for [Sigma1](#) or [PBKDFParamRequest](#).

An initiator receiving a [BUSY StatusReport](#) from a responder SHALL wait for at least a period of *t* milliseconds before retrying the request where *t* is the value obtained from the [Busy StatusReport ProtocolData](#) field.

If the initiator sends a new session establishment request after receiving a [BUSY StatusReport](#), the request SHALL contain new values for all randomized parameters.

4.11.2. Parameters and Constants

Table 20, “Glossary of constants” is a glossary of constants used in the secure channel protocol, along with a brief description and the default for each constant.

Table 20. Glossary of constants

Constant Name	Description	Value
MSG_COUNTER_WINDOW_SIZE	Maximum number of previously processed message counters to accept from a given Node and key.	32
MSG_COUNTER_SYNC_REQ_JITTER	Maximum amount of random delay before sending a <i>MsgCounterSyncReq</i> when the synchronization request is triggered by receipt of a multicast message.	500 milliseconds
MSG_COUNTER_SYNC_TIMEOUT	The maximum amount of time (in milliseconds) which a Node SHALL wait for a <i>MsgCounterSyncRsp</i> after sending a <i>MsgCounterSyncReq</i> .	400 milliseconds

4.12. Message Reliability Protocol (MRP)

The Message Reliability Protocol (MRP) provides confirmation of delivery for messages that require reliability. The protocol is optimized for constrained devices that may not be able to receive a message at the point it is due to be delivered to them. Reliable messaging MAY be enabled on an individual message basis as required by the protocol design of the higher layer application. Reliability is achieved through time-bounded delivery confirmation, ensuring best effort delivery of critical messages over what may be an inherently lossy and unreliable communication medium.

Flow control mechanisms are not incorporated in MRP because it is intended to be used for short interactions with small numbers of messages in them.

4.12.1. Reliable Messaging Header Fields

The following fields are defined in the [Exchange Flags](#) for use exclusively by MRP:

- [R Flag](#)

Indicates a reliable message. This flag SHALL be set by the sender when a message being sent requires the receiver to send back an acknowledgment. To support unreliable messages, this flag bit MAY be clear, so that no acknowledgements are requested from the receiver.

- [A Flag](#)

Indicates the message is acting as an acknowledgement. This flag MAY be set on any message. When set, the *Acknowledged Message Counter* field SHALL be present and valid. This flag SHALL always be set for *MRP Standalone Acknowledgement* messages.

- [Acknowledged Message Counter](#)

This field SHALL be set to the Message Counter of the message that is being acknowledged.

4.12.2. Reliable transfer

When the reliability bit is set, the *reliable message* is transmitted at most [MRP_MAX_TRANSMISSIONS](#) times until an acknowledgement of receipt is received from the peer or a timeout.

4.12.2.1. Retransmissions

Senders provide an automatic retransmission mechanism for reliable messages. In order for the receiver to receive a message reliably, the sender SHALL trigger the automatic retry mechanism after a period of *mrpBackoffTime* milliseconds without receiving an acknowledgement, where *mrpBackoffTime* is calculated according to the formula below. The sender SHALL retry up to a configured maximum number of times (*MRP_MAX_TRANSMISSIONS* - 1) before giving up and notifying the application.

Messages sent to a Node can be lost for various reasons such as lossy network or insufficient buffer space at the receiver. In the case of [Intermittently Connected Devices](#), which are active infrequently to receive messages destined for them, a sender must be aware of the characteristics of the recipient to ensure it does not attempt to send at a rate beyond the recipient's capability. Therefore, the

sender SHALL choose retransmission timeouts based on the [session characteristics](#) of the destination Node exposed via [Section 4.3.2, “Operational Discovery”](#).

At each sender, a retransmission timer is started each time a reliable message is transmitted. The duration of the retransmission timer SHALL be calculated as follows:

```
"mrpBackoffTime" = i * "MRP_BACKOFF_BASE"^(max(0,n-"MRP_BACKOFF_THRESHOLD")) * (1.0 +
"random"(0,1) * "MRP_BACKOFF_JITTER")
```

Where:

```
{:("mrpBackoffTime", =, "the resultant retransmission timeout for this
transmission"),(n, =, "the number of send attempts before the current one for this
message (0 if this is the initial transmission)"),(i, =, "the base retry interval for
the Exchange (either IDLE or ACTIVE)":)}
```

For each unique Exchange, the sender SHALL wait for the acknowledgement message until the retransmission timer, *mrpBackoffTime*, expires.

When waiting for an acknowledgement, an [Intermittently Connected Device](#) will always be in [Active Mode](#) since it will have at least one active exchange. An [Intermittently Connected Device](#) sender SHOULD increase the *mrpBackoffTime* by its [fast polling interval](#) to take into account the delay that might happen in receiving the acknowledgment while in [Active Mode](#).

For the first message of a new exchange, the base interval, *i*, SHALL be set according to the idle state of the peer node as stored in the Session Context of the session (either the [Secure Session Context](#) or the [Unsecured Session Context](#) depending on the Session Type). For all subsequent messages of the exchange, the base interval, *i*, SHOULD be set according to the active state of the peer node as stored in the Session Context of the session (either the [Secure Session Context](#) or the [Unsecured Session Context](#) depending on the Session Type), unless the sender has other means to determine whether the device is active or idle. The backoff base interval SHALL be set to a value at least 10% greater than the idle interval of the destination:

- If [PeerActiveMode](#) in the Session Context is true:
 - *i* = [SESSION_ACTIVE_INTERVAL](#) of the peer
- Else the peer is in idle mode:
 - *i* = [SESSION_IDLE_INTERVAL](#) of the peer
- *i* = [MRP_BACKOFF_MARGIN](#) * *i*

The [MRP_BACKOFF_THRESHOLD](#) parameter creates a two-phase scheme which begins with linear backoff to improve initial latency when congestion is not the cause of packet drops, and then transitions to exponential backoff to provide convergence when the network is congested. If a positive acknowledgment is received before the retransmission timer expires, the retransmission timer is stopped. Otherwise, if the retransmission timer expires, the message is retransmitted and the timer started again.

The following table illustrates minimum, maximum, and cumulative retransmission times using default parameters.

Table 21. Example MRP Retransmission Times

Metric	Transmission Time [ms]				
Min Jitter	330	330	528	845	1352
Max Jitter	413	413	660	1056	1690
Min Total	330	660	1188	2033	3385
Max Total	413	825	1485	2541	4231
Transmission #	0	1	2	3	4

The sender SHOULD initiate [Section 4.3.2, “Operational Discovery”](#) in parallel with the first retry to re-resolve the address of the destination Node if the initial transmission fails after one expected round trip. The sender SHOULD use the latest MRP parameters for the destination that result from subsequent Operational Discovery.

When a client is communicating with an ICD (i.e. the [ICD Management cluster](#) is present), the sender SHALL initiate [Section 4.3.2, “Operational Discovery”](#) in parallel with the first retry to re-resolve the address of the destination Node if the initial transmission fails after one expected round trip. The sender SHALL use the latest MRP parameters for the destination that result from subsequent Operational Discovery.

4.12.2.2. Acknowledgements

A receiver SHALL acknowledge a reliable message by either using a "piggybacked" acknowledgment in the next message destined to the peer, or a standalone acknowledgment, or both.

The acknowledgement message SHALL set the *Acknowledged Message Counter* field to the value of the *Message Counter* of the reliable message to be acknowledged.

Piggybacking Acknowledgments on Responses

Acknowledgements MAY be conveyed at the same time (i.e. piggybacked) as data in a response message. The receiver tries to optimize message transmission by deferring acknowledgments when a reliable message is received (see [Section 4.12.5.2.2, “Standalone acknowledgment processing”](#)) and piggybacking outstanding acknowledgments on messages that it needs to send back (see [Section 4.12.5.1.1, “Piggyback acknowledgment processing”](#) for more details).

Duplicate Message Detection

Since the reliable messaging protocol has a provision for the sender to retransmit messages, there is a significant chance that a duplicate message may arrive at the receiver. The receiver SHALL detect and mark duplicate messages that it receives using the standard authentication and replay protection mechanisms of the secure message layer (see [Section 4.6.5, “Replay Prevention and Duplicate Message Detection”](#)). The receiver SHALL send an acknowledgment message to the sender for each instance of an authenticated, reliable message, including duplicates. The reliability layer SHALL only propagate the first instance of a message to the next higher layer. Any message

marked as a duplicate SHALL be dropped by the reliability layer.

4.12.3. Peer Exchange Management

The Reliable Messaging Protocol operates within the scope of an [Exchange](#) between two Nodes. MRP SHALL support one pending acknowledgement and one pending retransmission per [Exchange](#).

MRP control parameters, detailed in [Table 22, “Glossary of MRP parameters”](#), are computed outside of the [Exchange communication](#) itself; instead, they are valid for the duration of a secure session. The [SESSION_ACTIVE_INTERVAL](#) and [SESSION_IDLE_INTERVAL](#), used in computation of MRP control parameters, are determined during [Operational Discovery](#) or [Section 4.3.1, “Commissionable Node Discovery”](#). Additionally, the initiator of a secure session MAY provide these parameters in the initial [CASE Sigma1](#) or [PASE PBKDFParamRequest](#) messages, and the responder MAY provide its parameters in the corresponding protocol messages [CASE Sigma2](#) or [PBKDFParamResponse](#).

4.12.4. Transport Considerations

When the upper layer requests a reliable message over a UDP transport, the [R Flag](#) SHALL be set on that message indicating that MRP SHALL be used. Reliable messages sent over TCP, [PAFTP](#), or BTP SHALL utilize the underlying reliability mechanisms of those transports and SHOULD NOT set the [R Flag](#). Reliable messages sent over [NTL](#) SHALL utilize the underlying reliability mechanisms and SHALL not set the [R Flag](#).

4.12.5. Reliable Message Processing

4.12.5.1. Reliable Message Processing of Outgoing Messages

To prepare a given Protocol Message for transmission, the message SHALL be processed as follows:

1. Proceed to [Section 4.12.5.1.1, “Piggyback acknowledgment processing”](#).

Piggyback acknowledgment processing

1. Determine if there is a matching pending acknowledgement in the [acknowledgement table](#) for the given message by checking all of the following conditions:
 - a. If the Destination Node Id and Exchange Id of the given message and pending acknowledgement are the same
 - b. AND either
 - i. the Session Id and underlying Session Credentials of the given message and pending acknowledgement are the same
 - ii. OR both the given message and pending acknowledgement are of [Unsecured Session Type](#).
2. If there is a matching pending acknowledgement, the [A Flag](#) SHALL be set on the outbound message so it will serve as a piggybacked acknowledgement.
 - a. For such a piggybacked acknowledgement, the *Acknowledgment Message Counter* field SHALL be set to the message counter of the received message for which an acknowledgement

ment was pending.

- b. If the message being prepared is not a [standalone acknowledgement](#), remove the matching entry from the [acknowledgement table](#).
- c. If the message being prepared is a [standalone acknowledgement](#), set the **StandaloneAckSent** field of the matching entry in the [acknowledgement table](#) to true.

Message retransmission processing

1. If the outbound message is marked to be delivered reliably over a UDP transport, the **R Flag** SHALL be set on the given message to request an acknowledgement from the peer upon receipt.
 - a. Any message flagged for reliable delivery (**R Flag** set) SHALL be stored in the [retransmission table](#) to track the message until it has been successfully acknowledged by the peer.
2. Perform [Section 4.7.1, “Message Transmission”](#) processing step on the message to send the message to the peer:
 - a. The same Session ID, Destination Node ID, Security Flags, and transport as were used for the initial message transmission SHALL be used.
3. If the transport interface returns an error on the send attempt, the error is assessed to determine whether the message can be retried.
 - a. If the error is fatal, the application is notified and the message removed from the [retransmission table](#).
 - b. If there is no error, or a non-fatal error such as no memory, the message is resent
 - i. Update the [retransmission table](#) to reflect the send count.
 - ii. Start a retransmission timer to track the maximum time to wait before attempting another retransmission.
 - iii. For each retry, the [retransmission table](#) is updated to track the number of retries until the maximum number is attempted, at which point the message is evicted from the [retransmission table](#).

Send flow state diagram

The MRP send flow described above is depicted in the control flow diagram [Figure 12, “MRP send flow”](#).




Figure 12. MRP send flow

4.12.5.2. Reliable Message Processing of Incoming Messages

A message received from [Section 4.7.2, “Message Reception”](#) for reliability processing SHALL be processed as follows:

1. Verify the message has a legal combination of reliability flags:
 - a. If the [R Flag](#) is set:
 - i. If [Group Session Type](#) AND [C Flag](#) = 0, drop the message.
 - b. If the [A Flag](#) is set:
 - i. If [Group Session Type](#) AND [C Flag](#) = 0, drop the message.
2. Proceed to [Section 4.10.5.1, “Exchange Message Matching”](#).
3. Proceed to [Section 4.12.5.2.1, “Received acknowledgement processing”](#).

Received acknowledgement processing

1. If the [A Flag](#) is set:
 - a. Query the [retransmission table](#) for the *Acknowledgement Message Counter* contained in the received message.
 - i. If there is a match:
 - A. Remove the entry from the [retransmission table](#).
 - B. Stop the retransmission timer for that entry.
 - ii. If there is no match, it indicates that this is either a duplicate acknowledgment or the [Exchange context](#) does not exist.
2. Proceed to [Section 4.12.5.2.2, “Standalone acknowledgement processing”](#).

Standalone acknowledgement processing

1. If the [R Flag](#) is set, the received message is requesting an acknowledgement be sent back:
 - a. If the message is marked as a duplicate:
 - i. Immediately send a [standalone acknowledgment](#).
 - ii. If the Exchange is marked as an [ephemeral exchange](#) the Exchange SHALL be closed.
 - iii. Drop the message.
 - b. Otherwise, instead of sending an acknowledgement immediately upon the receipt of a reliable message from a peer, the receiver SHOULD wait for a time no longer than [MRP_STAND-ALONE_ACK_TIMEOUT](#) before sending a [standalone acknowledgment](#):
 - i. Add the message counter of the received message to the [acknowledgement table](#) to signal that an outbound acknowledgement is pending. There can be only one outstanding acknowledgement at a time on a single [Exchange](#). If a pending acknowledgement already exists for the Exchange, and it has [StandaloneAckSent](#) set to false, a [standalone acknowledgment](#) SHALL be sent immediately for that pending message counter, and the [acknowledgement table](#) entry SHALL be replaced for the new message.
 - ii. Start the acknowledgement timer for the Exchange.

A. If the timer triggers before being cancelled, a **standalone acknowledgment** SHALL be sent to the source of the message. Sending this standalone acknowledgment SHALL NOT remove the **acknowledgement table** entry and SHALL set the **StandaloneAckSent** field of the entry to true.

2. The received message is then delivered to the next processing step of [Section 4.7.2, “Message Reception”](#).

Receive flow state diagram

The MRP receive flow described above is depicted in [Figure 13, “MRP receive flow”](#).




Figure 13. MRP receive flow

4.12.6. Reliable Message State

4.12.6.1. Retransmission Table

For retransmissions, the sender maintains a *retransmission table* of context records containing information on all reliable messages sent that have acknowledgments still pending. Each such reliable message context record includes the following fields:

- Reference to Exchange Context
- Message Counter
- Reference to fully formed, encoded and encrypted message buffer
- Send count
- Retransmission timeout counter

Each time a message that requires acknowledgment is sent, a new retransmission context record is inserted into the *retransmission table* or an existing record is updated to increment its send count. The message is sent a configurable maximum number of times (*MRP_MAX_TRANSMISSIONS*) and, if still undelivered, the application is notified of the failure.

4.12.6.2. Acknowledgement Table

The receiver maintains an *acknowledgement table* of context records containing information on each reliable message for which an acknowledgment SHALL be sent. Each such reliable message context record includes the following fields:

- Reference to Exchange Context
- Message Counter
- A boolean, **StandaloneAckSent**, indicating whether a **standalone acknowledgement** has been sent for this message counter. Initially false.

An entry SHALL remain in the table until one of the following things happens:

1. The exchange associated with the entry is closed. See [Section 4.10.5.3, “Closing an Exchange”](#).
2. The exchange associated with the entry has switched to track a pending acknowledgement for a new message counter value. See [Section 4.12.5.2.2, “Standalone acknowledgement processing”](#).
3. A message that is not a **standalone acknowledgement** is sent which serves as an acknowledgment for the entry. See [Section 4.12.5.1.1, “Piggyback acknowledgment processing”](#).

4.12.7. MRP Messages

4.12.7.1. MRP Standalone Acknowledgement

The *MRP Standalone Acknowledgement* message SHALL be formed as follows:

- The application payload SHALL be empty.
- The **A Flag** SHALL be set to **1**.

- The [Acknowledged Message Counter](#) SHALL be included in the header.
- The [Protocol ID](#) SHALL be set to `PROTOCOL_ID_SECURE_CHANNEL`.
- The [Protocol Opcode](#) SHALL be set to [MRP Standalone Acknowledgement](#).

The rules for when to send this message are detailed in [Section 4.12.5.2.2, “Standalone acknowledgement processing”](#).

4.12.8. Parameters and Constants

This is a glossary of the MRP parameters used in this chapter with a brief description for each parameter. A Node SHALL use the provided default value for each parameter unless the message recipient Node advertises an alternate value for the parameter via Operational Discovery.

Table 22. Glossary of MRP parameters

Parameter Name	Description	Default Value
<i>MRP_MAX_TRANSMISSIONS</i>	The maximum number of transmission attempts for a given reliable message. The sender MAY choose this value as it sees fit.	5
<i>MRP_BACKOFF_BASE</i>	The base number for the exponential backoff equation.	1.6
<i>MRP_BACKOFF_JITTER</i>	The scaler for random jitter in the backoff equation.	0.25
<i>MRP_BACKOFF_MARGIN</i>	The scaler margin increase to backoff over the peer idle interval.	1.1
<i>MRP_BACKOFF_THRESHOLD</i>	The number of retransmissions before transitioning from linear to exponential backoff.	1
<i>MRP_STANDALONE_ACK_TIMEOUT</i>	Amount of time to wait for an opportunity to piggyback an acknowledgement on an outbound message before falling back to sending a standalone acknowledgement .	200 milliseconds

4.13. Unicast Communication

This section specifies the semantics of establishing a unicast session and the lifecycle of a unicast session.

Unicast sessions exist in one of two phases:

1. Session Establishment Phase: A series of well-defined unencrypted messages that aim to establish a shared key.
2. Application Data Phase: A series of ad-hoc encrypted messages exchanging interaction model protocol actions, application data, etc.

4.13.1. Session Parameters

This is a glossary of the session parameters used during session establishment with a brief description for each parameter. A Node SHALL use the provided default value for each parameter unless the message recipient Node advertises an alternate value for the parameter via Operational Discovery.

Table 23. Glossary of Session parameters

Parameter Name	Description	Default Value
<i>SESSION_IDLE_INTERVAL</i>	Minimum amount of time between sender retries when the destination node is Idle. This SHALL be greater than or equal to the maximum amount of time a node may be non-responsive to incoming messages when Idle.	500 milliseconds
<i>SESSION_ACTIVE_INTERVAL</i>	Minimum amount of time between sender retries when the destination node is Active. This SHALL be greater than or equal to the maximum amount of time a node may be non-responsive to incoming messages when Active.	300 milliseconds
<i>SESSION_ACTIVE_THRESHOLD</i>	Minimum amount of time the node SHOULD stay active after network activity.	4000 milliseconds
<i>DATA_MODEL_REVISION</i>	Version of Data Model for the Session parameters side where it appears.	See Section 11.1.5.1, “Data-ModelRevision” .
<i>INTERACTION_MODEL_REVISION</i>	Version of Interaction Model for the Session parameters side where it appears.	See Section 8.1.1, “Revision History” .
<i>SPECIFICATION_VERSION</i>	Version of Specification for the Session parameters side where it appears.	See Section 11.1.5.22, “Specification-Version” .
<i>MAX_PATHS_PER_INVOKE</i>	The maximum number of elements in the InvokeRequests list that the Node is able to process	See Section 11.1.5.23, “MaxPathsPer-Invoke” .
<i>SUPPORTED_TRANSPORTS</i>	A bitmap of the supported transport protocols in addition to MRP .	See Table 7, “Supported Transport Mode Values” .

Parameter Name	Description	Default Value
<i>MAX_TCP_MESSAGE_SIZE</i>	Maximum size of the message carried over TCP, excluding the framing message length field, that the node is capable of receiving from its peer.	64000 bytes.

These parameters are encoded in the following TLV format when included in CASE / PASE session establishment:

```
session-parameter-struct => STRUCTURE [ tag-order ]
{
  SESSION_IDLE_INTERVAL      [1, optional] : UNSIGNED INTEGER [ range 32-bits ],
  SESSION_ACTIVE_INTERVAL    [2, optional] : UNSIGNED INTEGER [ range 32-bits ],
  SESSION_ACTIVE_THRESHOLD   [3, optional] : UNSIGNED INTEGER [ range 16-bits ],
  DATA_MODEL_REVISION        [4]           : UNSIGNED INTEGER [ range 16-bits ],
  INTERACTION_MODEL_REVISION [5]           : UNSIGNED INTEGER [ range 16-bits ],
  SPECIFICATION_VERSION       [6]           : UNSIGNED INTEGER [ range 32-bits ],
  MAX_PATHS_PER_INVOKE        [7]           : UNSIGNED INTEGER [ range 16-bits ],
  SUPPORTED_TRANSPORTS        [8]           : UNSIGNED INTEGER [ range 16-bits ],
  MAX_TCP_MESSAGE_SIZE        [9, optional] : UNSIGNED INTEGER [ range 32-bits ],
}
```

For backwards compatibility, if any tag after tag 2 (*SESSION_ACTIVE_INTERVAL*) is present, then the *SESSION_ACTIVE_INTERVAL* SHALL also be present.

NOTE

That means that *SESSION_ACTIVE_INTERVAL* is always present if the sender of the CASE/PASE session establishment message supports *DATA_MODEL_REVISION*, *INTERACTION_MODEL_REVISION*, and *SPECIFICATION_VERSION*, but recipients SHALL NOT assume that it is present.

For backwards compatibility, if the *DATA_MODEL_REVISION* field is missing, it implies a DataModelRevision value of either 16 or 17.

For backwards compatibility, if the *INTERACTION_MODEL_REVISION* field is missing, it implies a value of either 10 or 11.

For backwards compatibility, if the *SPECIFICATION_VERSION* field is missing, it implies a SpecificationVersion value strictly smaller than 0x01030000.

For backwards compatibility, if the *MAX_PATHS_PER_INVOKE* field is missing, it implies a MaxPathsPerInvoke set to 1.

For backwards compatibility, if the *SUPPORTED_TRANSPORTS* field is missing, it implies that the node only supports [MRP](#). When present with a value of 0, it indicates that the node only supports [MRP](#).

The *MAX_TCP_MESSAGE_SIZE* field SHALL only be present if the *SUPPORTED_TRANSPORTS* field indicates that TCP is supported.

4.13.2. Session Establishment Phase

Session establishment uses either the *CASE* or *PASE* protocol.

CASE SHALL be used as a session establishment mechanism for **all** sessions except:

1. Communication for the purpose of commissioning when *NOC* has not yet been installed

PASE SHALL only be used for session establishment mechanism during device commissioning. *PASE* SHALL NOT be used as a session establishment mechanism for any other session. *BTP*, *PAFTP* and *NTL* MAY be used as the transport for device commissioning. *BTP*, *PAFTP* and *NTL* SHALL NOT be used as a transport for operational purposes.

Unless otherwise specified, the *CASE*, *PASE*, *User-Directed Commissioning* protocol, and *Secure Channel Status Report* messages SHALL be the only allowed **unencrypted** messages.

This phase aims to:

1. Authenticate peers (CASE-based sessions only).
2. Derive shared secrets to encrypt subsequent session data.
3. Choose session identifiers to identify the subsequent session.

4.13.2.1. Unsecured Session Context

The following session context data SHALL be utilized to associate messages to a particular peer and recover context during unencrypted sessions:

1. **Session Role**: Records whether the node is the session initiator or responder.
2. **Ephemeral Initiator Node ID**: Randomly selected for each session by the initiator from the *Operational Node ID* range and enclosed by initiator as *Source Node ID* and responder as *Destination Node ID*.
 - Initiators SHALL select a new random ephemeral node ID for each unsecured session, and SHALL select an ID that does not conflict with any ephemeral node IDs for any other ongoing unsecured sessions opened by the initiator.
3. **Message Reception State**: Provides tracking for the *Unencrypted Message Counter* of the remote peer.

Matching and responder creation of **Unsecured Session Contexts** SHALL be as follows:

1. Given an incoming unencrypted message
 - a. Locate any **Unsecured Session Context** with matching **Ephemeral Initiator Node ID**
 - i. If any is located, the incoming message SHALL be assumed to be associated with this **Unsecured Session Context**
 - b. Else if the message carries a *Source Node ID*

- i. Create a new **Unsecured Session Context**
 - ii. Set **Session Role** to **responder**
 - iii. Record the incoming message's **Source Node ID** as **Ephemeral Initiator Node ID**
- c. Else discard the message

Initiator creation of **Unsecured Session Contexts** SHALL be as follows:

1. Given the first outgoing message of an unencrypted exchange
 - a. Create a new **Unsecured Session Context**
 - b. Set **Session Role** to **initiator**
 - c. Randomly select a node ID from the **Operational Node ID** range that does not collide with any ephemeral node IDs for any other ongoing unsecured sessions opened by the initiator and record this as **Ephemeral Initiator Node ID**

4.13.2.2. Session Establishment over IP

When establishing a session over IP, the initiator MAY use TCP when both of the following are true:

1. The initiator supports TCP Client as defined in the **Supported Transport Mode Values** table in the **T record**.
2. The responder supports TCP Server as defined in the **Supported Transport Mode Values** table in the **T record**.

Otherwise, the initiator SHALL use MRP to establish the session.

The transport used during the session establishment phase SHALL also be used for the subsequent transport of messages over the established session.

4.13.2.3. Shared Secrets

Both **CASE** and **PASE** produce two shared keys: **I2RKey** and **R2IKey**. These keys will be saved to the session's context and used to encrypt and decrypt messages during the Session Data Phase.

Nodes that support the **CASE session resumption** SHALL also save to the session's context the **SharedSecret** computed during the **CASE** protocol execution.

4.13.2.4. Choosing Secure Unicast Session Identifiers

Both **CASE** and **PASE** allow each participant the ability to choose a unicast session identifier for the subsequent encrypted session. The session identifier SHALL be used to look up the relevant encryption keys and any other metadata for a particular session.

Messages using a unicast session identifier SHALL set the **Session Type** field to **0**. Each peer SHALL specify a **Session Identifier** unique in reference to **their own** active sessions. There SHALL NOT be overlap between the Session ID values allocated for **PASE** and **CASE** sessions, as the **Session Identifier** space is shared across both session establishment methods.

For example, if the initiator has two active sessions with session identifiers 0x0001 and 0x0002, it

could choose any non-zero session identifier besides 0x0001 and 0x0002.

If there are no available session identifiers (i.e. the participant has 65,535 open sessions), the Node SHALL terminate an existing session to free a session identifier.

4.13.3. Application Data Phase

When the last [CASE](#) or [PASE](#) protocol message is sent or received and successfully processed, session establishment has completed.

4.13.3.1. Secure Session Context

During the Application Data Phase, the following conceptual session context data SHALL be utilized to securely process subsequent messages:

1. **Session Type**: Records whether the session was established using [CASE](#) or [PASE](#).
2. **Session Role**: Records whether the node is the session initiator or responder.
3. **Local Session Identifier**: Individually selected by each participant in secure unicast communication during session establishment and used as a unique identifier to recover encryption keys, authenticate incoming messages and associate them to existing sessions.
 - On a given Node, this is the identifier that SHALL be used to map from an incoming message's Session ID field to the session context data.
4. **Peer Session Identifier**: Assigned by the peer during session establishment.
 - On a given Node, this is the identifier that SHALL be used in the [Session ID](#) field of every outgoing message associated with the session, so that it can be interpreted as the **Local Session Identifier** by the remote peer.
5. **I2RKey**: Encrypts data in messages sent from the initiator of session establishment to the responder.
6. **R2IKey**: Encrypts data in messages sent from the session establishment responder to the initiator.
7. **SharedSecret**: Computed during the [CASE](#) protocol execution and re-used when [CASE session resumption](#) is implemented.
8. **Local Message Counter**: [Secure Session Message Counter](#) for outbound messages.
 - At successful session establishment, the **Local Message Counter** SHALL be initialized per [Section 4.6.1.1, "Message Counter Initialization"](#).
9. **Message Reception State**: Provides tracking for the [Secure Session Message Counter](#) of the remote peer.
10. **Local Fabric Index**: Records the local [Index](#) for the session's Fabric, which MAY be used to look up Fabric metadata related to the Fabric for which this session context applies.
 - This field SHALL contain the "no Fabric" value of 0 when the **SessionType** is [PASE](#) and successful invocation of the [AddNOC](#) command has not yet occurred during commissioning.
11. **Peer Node ID**: Records the authenticated node ID of the remote peer, when available.
 - This field SHALL contain the "Unspecified Node ID" value of 0 when the **SessionType** is [PASE](#).

12. **Resumption ID**: The ID used when resuming a session between the local and remote peer.
13. **SessionTimestamp**: A timestamp indicating the time at which the last message was sent or received. This timestamp SHALL be initialized with the time the session was created. See [Section 4.11.1.1, “Session Establishment - Out of Resources”](#) for more information.
14. **ActiveTimestamp**: A timestamp indicating the time at which the last message was received. This timestamp SHALL be initialized with the time the session was created.
15. The following Session parameters (see [Table 23, “Glossary of Session parameters”](#)):
 - a. *SESSION_IDLE_INTERVAL*
 - b. *SESSION_ACTIVE_INTERVAL*
 - c. *SESSION_ACTIVE_THRESHOLD*
 - d. **PeerActiveMode**: A boolean that tracks whether the peer node is in Active or Idle mode. **PeerActiveMode** is set as follows:

```
"PeerActiveMode" = ("now()" - "ActiveTimestamp") < "SESSION_ACTIVE_THRESHOLD"
```

Note that the **Local Fabric Index** and **Peer Fabric Index** reported in the [NOC Response](#) MAY differ in value, while still referring to the same Fabric, since for a given complete [Fabric Reference](#), the short [Fabric Index](#) allocated during commissioning of the respective Nodes on the same Fabric MAY be different. This possible difference is due to the order in which the Fabric in question was joined in the lifecycle of the respective Nodes. See the section on [AddNOC command behavior](#) for details on Fabric Index allocation behavior over time.

There SHALL also be reservation of storage to support [CASE Authenticated Tag \(CAT\)](#) fields. The **CAT** fields are 32-bit values that MAY have been present in RDN [case-authenticated-tag](#) of the remote peer's operational certificate, during [CASE](#).

The **CAT** fields are used to cache Operational Certificate data so that it can be used by the [ACL processing logic](#) to support [CASE Authenticated Tags](#).

Since these fields MAY be omitted from NOCs, they MAY be marked as absent in the context, such that they are not taken into account when missing. When present, they SHALL be stored. Maximum up to 3 **CAT** fields SHALL be supported.

Their value is unused in [PASE](#) session contexts.

4.14. Session Establishment

4.14.1. Passcode-Authenticated Session Establishment (PASE)

This section describes session establishment using a shared [passcode](#) together with an [augmented Password-Authenticated Key Exchange \(PAKE\)](#), in which only one party knows the passcode beforehand, to generate shared keys. This protocol is only used when commissioning a Node (i.e. the Commissionee).

4.14.1.1. Protocol Overview

The Passcode-Authenticated Session Establishment (PASE) protocol aims to establish the first session between a Commissioner and a Commissionee using a known passcode provided out-of-band. The pairing is performed using [Section 3.10, “Password-Authenticated Key Exchange \(PAKE\)”](#) and relies on a [Password-Based Key Derivation Function \(PBKDF\)](#) where the passcode is used as password.

This session establishment protocol provides a means to:

1. Communicate PBKDF parameters.
2. Derive PAKE bidirectional secrets.




Figure 14. Overview of the PASE Protocol

The Commissioner is the Initiator and the Commissionee is the Responder.

It is assumed that the initiator has somehow obtained the passcode and that the responder has the relevant `Crypto_PAKEValues_Responder` corresponding to the passcode before starting a PASE session establishment protocol.

4.14.1.2. Protocol Details

Message format

All PASE messages SHALL be structured as specified in [Section 4.4, “Message Frame Format”](#).

All PASE messages are sent using an [Unsecured Session](#):

- The [Session ID](#) field SHALL be set to **0**.
- The [Session Type](#) bits of the [Security Flags](#) SHALL be set to **0**.
- In the PASE messages from the initiator, [S Flag](#) SHALL be set to **1** and [DSIZ](#) SHALL be set to **0**.
- In the PASE messages from the responder, [S Flag](#) SHALL be set to **0** and [DSIZ](#) SHALL be set to **1**.

For each PASE message, the [application payload](#) is the TLV encoding of the message structure as defined below:

Table 24. PASE Messages

Message Name	Payload TLV Encoding
PBKDFParamRequest	pbkdfparamreq-struct
PBKDFParamResponse	pbkdfparamresp-struct
Pake1	pake-1-struct
Pake2	pake-2-struct
Pake3	pake-3-struct
PakeFinished	N/A (encoded via StatusReport)

The other fields of the [Message format](#) are not specific to the PASE messages.

For all TLV-encoded PASE messages, any context-specific tags not listed in the associated TLV schemas SHALL be reserved for future use, and SHALL be silently ignored if seen by a recipient which cannot understand them.

Message Exchange

The [PBKDFParamRequest](#), [PBKDFParamResponse](#), [Pake1](#), [Pake2](#), [Pake3](#), and [PakeFinished](#) of a distinct session establishment are part of the same [message exchange](#). The initiator and responder SHALL NOT send encrypted application data in the newly established session until [PakeFinished](#) is received by the initiator within the unencrypted session used for establishment.

Each message SHALL use [PROTOCOL_ID_SECURE_CHANNEL](#) as [Protocol ID](#) and the corresponding [Protocol Opcode](#) as defined in [Table 18](#), “Secure Channel Protocol Opcodes”.

The flags of the [Exchange Flags](#) of the [Protocol Header](#) are defined as follows per PASE message:

Message	I Flag
PBKDFParamRequest	1
PBKDFParamResponse	0
Pake1	1
Pake2	0
Pake3	1

All PASE messages SHALL be sent reliably. This may be implicit (e.g. TCP) or explicit (e.g. [MRP](#) reliable messaging) in the underlying transport.

The other fields of the [Protocol Header](#) are not specific to the PASE messages.

PBKDFParamRequest

This message serves to request the PBKDF parameters, with a payload that follows this TLV schema:

```
pbkdfparamreq-struct => STRUCTURE [ tag-order ]
{
```

```

    initiatorRandom      [1] : OCTET STRING [ length 32 ],
    initiatorSessionId   [2] : UNSIGNED INTEGER [ range 16-bits ],
    passcodeId           [3] : UNSIGNED INTEGER [ length 16-bits ],
    hasPBKDFParameters   [4] : BOOLEAN,
    initiatorSessionParams [5, optional] : session-parameter-struct
}

```

1. The initiator SHALL generate a random number `InitiatorRandom = Crypto_DRBG(len = 32 * 8)`.
2. The initiator SHALL generate a session identifier (`InitiatorSessionId`) for subsequent identification of this session. The `InitiatorSessionId` field SHALL NOT overlap with any other existing PASE or CASE session identifier in use by the initiator. See Section 4.13.2.4, “Choosing Secure Unicast Session Identifiers” for more details. The initiator SHALL set the Local Session Identifier in the Session Context to the value `InitiatorSessionId`.
3. The initiator SHALL choose a passcode identifier (`PasscodeId`) corresponding to a particular PAKE passcode verifier installed on the responder. A value of 0 for the `passcodeId` SHALL correspond to the PAKE passcode verifier for the currently-open commissioning window, if any. Non-zero values are reserved for future use. For example, for initial commissioning, the verifier would be the built-in verifier matching the Onboarding Payload's passcode or, equivalently, the multi-fabric Basic Commissioning Method passcode if that method is supported. For the multi-fabric Enhanced Commissioning Method, the verifier would match the verifier provided through the `OpenCommissioningWindow` command.
4. The initiator SHALL indicate whether the PBKDF parameters (salt and iterations) are known for the particular `passcodeId` (for example from the QR code) by setting `HasPBKDFParameters`. If `HasPBKDFParameters` is set to `True`, the responder SHALL NOT return the PBKDF parameters. If `HasPBKDFParameters` is set to `False`, the responder SHALL return the PBKDF parameters.
5. The initiator SHALL send a message with the appropriate Protocol Id and Protocol Opcode from Table 18, “Secure Channel Protocol Opcodes” whose payload is the TLV-encoded `pbkdf-paramreq-struct PBKDFParamRequest` with an anonymous tag for the outermost struct.

```

PBKDFParamRequest =
{
    initiatorRandom (1) = InitiatorRandom,
    initiatorSessionId (2) = InitiatorSessionId,
    passcodeId (3) = PasscodeId,
    hasPBKDFParameters (4) = HasPBKDFParameters,
}

```

PBKDFParamResponse

```

pbkdfparamresp-struct => STRUCTURE [ tag-order ]
{
    initiatorRandom      [1] : OCTET STRING [ length 32 ],
    responderRandom      [2] : OCTET STRING [ length 32 ],
}

```



```

responderSessionId      [3] : UNSIGNED INTEGER [ range 16-bits ],
pbkdf_parameters        [4] : Crypto_PBKDFParameterSet,
responderSessionParams  [5, optional] : session-parameter-struct
}

```

On receipt of **PBKDFParamRequest**, the responder SHALL:

1. Verify **passcodeID** is set to 0. If verification fails, the responder SHALL send a **status report: StatusReport(GeneralCode: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: INVALID_PARAMETER)** and perform no further processing.
2. **Generate a random number** **ResponderRandom** = **Crypto_DRBG(len = 32 * 8)**.
3. Generate a **session identifier** (**ResponderSessionId**) for subsequent identification of this session. The **ResponderSessionId** field SHALL NOT overlap with any other existing PASE or CASE **session identifier** in use by the responder. See **Section 4.13.2.4, “Choosing Secure Unicast Session Identifiers”** for more details. The responder SHALL set the **Local Session Identifier** in the **Session Context** to the value **ResponderSessionId**.
4. Set the **Peer Session Identifier** in the **Session Context** to the value **PBKDFParamRequest.initiatorSessionId**.
5. Construct the appropriate **Crypto_PBKDFParameterSet** (**PBKDFParameters**). If **PBKDFParamRequest.hasPBKDFParameters** is **True** the responder SHALL NOT include the PBKDF parameters (i.e. salt and iteration count) in the **Crypto_PBKDFParameterSet**. If **Msg1.hasPBKDFParameters** is **False** the responder SHALL include the PBKDF parameters (i.e. salt and iteration count) in the **Crypto_PBKDFParameterSet**.
6. Send a message with the appropriate **Protocol Id** and **Protocol Opcode** from **Table 18, “Secure Channel Protocol Opcodes”** whose payload is the TLV-encoded **pbkdfparamresp-struct** **PBKDFParamResponse** with an anonymous tag for the outermost struct.

```

PBKDFParamResponse =
{
    initiatorRandom (1) = PBKDFParamRequest.initiatorRandom,
    responderRandom (2) = ResponderRandom,
    responderSessionId (3) = ResponderSessionId,
    pbkdf_parameters (4) = PBKDFParameters
}

```

Pake1

```

pake-1-struct => STRUCTURE [ tag-order ]
{
    pA [1] : OCTET STRING [ length CRYPTO_PUBLIC_KEY_SIZE_BYTES ],
}

```

On receipt of **PBKDFParamResponse**, the initiator SHALL:

1. Set the **Peer Session Identifier** in the **Session Context** to the value **PBKDFParamResponse.responderSessionId**.
2. Generate the **Crypto_PAKEValues_Initiator** according to the **PBKDFParamResponse.pbkdf_parameters**
3. Using **Crypto_PAKEValues_Initiator**, generate **pA := Crypto_pA(Crypto_PAKEValues_Initiator)**
4. Send a message with the appropriate **Protocol Id** and **Protocol Opcode** from Table 18, “Secure Channel Protocol Opcodes” whose payload is the TLV-encoded **pake-1-struct Pake1** with an anonymous tag for the outermost struct.

```
Pake1 =
{
  pA (1) = pA,
}
```

Pake2

```
pake-2-struct => STRUCTURE [ tag-order ]
{
  pB [1] : OCTET STRING [ length CRYPTO_PUBLIC_KEY_SIZE_BYTES ],
  cB [2] : OCTET STRING [ length CRYPTO_HASH_LEN_BYTES],
}
```

On receipt of **Pake1**, the responder SHALL:

1. Compute **pB := Crypto_pB(Crypto_PAKEValues_Responder)** using the passcode verifier indicated in **PBKDFParamRequest**
2. Compute **TT := Crypto_Transcript(PBKDFParamRequest, PBKDFParamResponse, Pake1.pA, pB)**
3. Compute **(cA, cB, Ke) := Crypto_P2(TT, Pake1.pA, pB)**
4. Send a message with the appropriate **Protocol Id** and **Protocol Opcode** from Table 18, “Secure Channel Protocol Opcodes” whose payload is the TLV-encoded **pake-2-struct Pake2** with an anonymous tag for the outermost struct.

```
Pake2 =
{
  pB (1) = pB,
  cB (2) = cB,
}
```

Pake3

```
pake-3-struct => STRUCTURE [ tag-order ]
{
```

```
cA [1] : OCTET STRING [length CRYPTO_HASH_LEN_BYTES],
}
```

On receipt of **Pake2**, the initiator SHALL:

1. Compute **TT** := **Crypto_Transcript**(**PBKDFParamRequest**, **PBKDFParamResponse**, **Pake1.pA**, **Pake2.pB**)
2. Compute (**cA**, **cB**, **Ke**) := **Crypto_P2**(**TT**, **Pake1.pA**, **Pake2.pB**)
3. Verify **Pake2.cB** against **cB**. If verification fails, the initiator SHALL send a status report: **StatusReport**(**GeneralCode**: **FAILURE**, **ProtocolId**: **SECURE_CHANNEL**, **ProtocolCode**: **INVALID_PARAMETER**) and perform no further processing.
4. Send a message with the appropriate **Protocol Id** and **Protocol Opcode** from Table 18, “Secure Channel Protocol Opcodes” whose payload is the TLV-encoded **pake-3-struct Pake3** with an anonymous tag for the outermost struct.

```
Pake3 =
{
  cA (1) = cA,
}
```

5. The initiator SHALL NOT send any encrypted application data until it receives **PakeFinished** from the responder.

On reception of **Pake3**, the responder SHALL:

1. Verify **Pake3.cA** against **cA**. If verification fails, the responder SHALL send a status report: **StatusReport**(**GeneralCode**: **FAILURE**, **ProtocolId**: **SECURE_CHANNEL**, **ProtocolCode**: **INVALID_PARAMETER**) and perform no further processing.
2. The responder SHALL set **SessionTimestamp** to a timestamp from a clock which would allow for the eventual determination of the last session use relative to other sessions.
3. The responder SHALL encode and send **PakeFinished**.

PakeFinished

To indicate the successful completion of the protocol, the responder SHALL send a status report: **StatusReport**(**GeneralCode**: **SUCCESS**, **ProtocolId**: **SECURE_CHANNEL**, **ProtocolCode**: **SESSION_ESTABLISHMENT_SUCCESS**).

The initiator SHALL set **SessionTimestamp** to a timestamp from a clock which would allow for the eventual determination of the last session use relative to other sessions.

Session Encryption Keys

After verification of **Pake3**, each party can compute their sending and receiving session keys as described below:

```
byte SEKeys_Info[] = /* "SessionKeys" */
```

```
{0x53, 0x65, 0x73, 0x73, 0x69, 0x6f, 0x6e, 0x4b,
  0x65, 0x79, 0x73}
```

```
I2RKey || R2IKey || AttestationChallenge =
  Crypto_KDF
  (
    inputKey = Ke,
    salt = [],
    info = SEKeys_Info,
    len = 3 * CRYPTO_SYMMETRIC_KEY_LENGTH_BITS
  )
```

1. Each key is exactly **CRYPTO_SYMMETRIC_KEY_LENGTH_BITS** bits.
2. The initiator SHALL use **I2RKey** to encrypt and integrity protect messages and the `R2IKey' to decrypt and verify messages.
3. The responder SHALL use **R2IKey** to encrypt and integrity protect messages and the `I2RKey' to decrypt and verify messages.
4. The **AttestationChallenge** SHALL only be used as a challenge during device attestation. See [Section 6.2.3, “Device Attestation Procedure”](#) for more details.

Upon initial installation of the new PASE Session Keys:

1. The Node SHALL initialize its [Local Message Counter](#) in the [Session Context](#) per [Section 4.6.1.1, “Message Counter Initialization”](#).
2. The Node SHALL initialize the [Message Reception State](#) in the [Session Context`](#) and set the synchronized **max_message_counter** of the peer to 0.

where **||** indicates message concatenation and **[]** a zero-length array.

4.14.2. Certificate Authenticated Session Establishment (CASE)

This section describes a certificate-authenticated session establishment (CASE) protocol using [Node Operational credentials](#). This session establishment mechanism provides an authenticated key exchange between exactly two peers while maintaining privacy of each peer. A resumption mechanism allows bootstrapping a new session from a previous one, dramatically reducing the computation required as well as reducing the number of messages exchanged.

4.14.2.1. Protocol Overview

This session establishment protocol provides a means to:

1. Mutually authenticate both peer Nodes
2. Generate cryptographic keys to secure subsequent communication within a session
3. Exchange operational parameters for the session, such as [Session Identifier](#) and [MRP](#) param-

ters

The cryptographic protocol mirrors the [SIGMA] protocol and uses the [Identity Protection Key \(IPK\)](#) to provide better identity protection. Briefly, the protocol will:

1. Exchange ephemeral elliptic curve public keys ([Sigma1.initiatorEphPubKey](#) and [Sigma2.responderEphPubKey](#)) to generate a shared secret
2. Exchange certificates to prove identities ([Sigma2.encrypted2.responderNOC](#) and [Sigma3.encrypted3.initiatorNOC](#))
3. Prove possession of the NOC private key by signing the ephemeral keys and NOC ([sigma-2-tbs-data](#) and [sigma-3-tbsdata](#))

The basic protocol can be achieved within 2 round trips as shown below:




Figure 15. Basic Session Establishment

4.14.2.2. Session Resumption

The protocol also provides a means to quickly resume a session using a previously established session. Resumption does **not** require expensive signature creation and verification which significantly reduces the computation time. Because of this, resumption is favoured for low-powered devices when applicable. Session resumption SHOULD be used by initiators when the necessary [state](#) is known to the initiator.

The nomenclature [Sigma1 with Resumption](#) in the following subsections implies a [Sigma1](#) message with both the optional [resumptionID](#) and [initiatorResumeMIC](#) fields populated in [sigma-1-struct](#).




Figure 16. Session Resumption

In the case where a Responder is not able to resume a session as requested by a [Sigma1 with Resumption](#), the information included in the [Sigma1 with Resumption](#) message SHALL be processed as a [Sig-](#)

ma1 message without any resumption fields to construct a Sigma2 message and continue the standard session establishment protocol without resumption.

To make the resumption succeed, both the Initiator and the Responder SHALL have remembered the SharedSecret they have computed during the previous execution of the CASE session establishment. It SHALL be that SharedSecret that is used to compute the resumption ID.

Session Resumption State

To perform session resumption, the following state from the previous session context must be known to the initiator and responder:

- 1. SharedSecret
- 2. Local Fabric Index
- 3. Peer Node ID
- 4. Peer CASE Authenticated Tags
- 5. ResumptionID

4.14.2.3. Protocol Details

Message format

All CASE messages SHALL be structured as specified in Section 4.4, “Message Frame Format”.

All CASE messages are sent using an Unsecured Session:

- The Session ID field SHALL be set to 0.
- The Session Type bits of the Security Flags SHALL be set to 0.
- In the CASE messages from the initiator, S Flag SHALL be set to 1 and DSIZ SHALL be set to 0.
- In the CASE messages from the responder, S Flag SHALL be set to 0 and DSIZ SHALL be set to 1.

For each CASE message, the application payload is the TLV encoding of the message structure as defined below:

Table 25. CASE Messages

Message Name	Payload TLV Encoding
Sigma1	sigma-1-struct
Sigma2	sigma-2-struct,
Sigma3	sigma-3-struct,
Sigma2_Resume	sigma-2-resume-struct,
SigmaFinished	N/A (encoded via StatusReport)

The other fields of the Message format are not specific to the CASE messages.

Message Exchange

The **Sigma1**, **Sigma2**, **Sigma3**, and **SigmaFinished** of a distinct session establishment are part of the same **message exchange**. The **Sigma1 with resumption**, **Sigma2_Resume** and **SigmaFinished** of a distinct session resumption are part of the same **message exchange**. The **Sigma1 with resumption**, **Sigma2**, **Sigma3** and **SigmaFinished** of a distinct session resumption that failed to perform the resumption are part of the same **message exchange**.

Each message SHALL use **PROTOCOL_ID_SECURE_CHANNEL** as **Protocol ID** and the corresponding **Protocol Opcode** as defined in Table 18, “Secure Channel Protocol Opcodes”.

The **Exchange Flags** of the **Protocol Header** are defined as follows per CASE message:

Message	I Flag
CASE Sigma1	1
CASE Sigma2	0
CASE Sigma3	1
CASE Sigma2_Resume	0

For the **SigmaFinished** message the value of the **I Flag** is set depending on whether the status message is sent by the Initiator or the Responder.

All CASE messages SHALL be sent reliably. This may be implicit (e.g. TCP) or explicit (e.g. **MRP** reliable messaging) in the underlying transport.

The other fields of the **Exchange format** are not specific to the CASE messages.

Generate and Send Sigma1

The initiator encodes and sends a **Sigma1** message, with a payload that follows this TLV schema:

```
sigma-1-struct => STRUCTURE [ tag-order ]
{
    initiatorRandom          [1] : OCTET STRING [ length 32 ],
    initiatorSessionId       [2] : UNSIGNED INTEGER [ range 16-bits ],
    destinationId            [3] : destination-identifier,
    initiatorEphPubKey       [4] : ec-pub-key,
    initiatorSessionParams   [5, optional] : session-parameter-struct,
    resumptionID             [6, optional] : OCTET STRING [ length 16 ],
    initiatorResumeMIC       [7, optional] : OCTET STRING [ length
CRYPTO_AEAD_MIC_LENGTH_BYTES ]
}
```

1. The initiator SHALL **generate a random number** **InitiatorRandom** = **Crypto_DRBG**(Len = 32 * 8).
2. The initiator SHALL generate a **session identifier** (**InitiatorSessionId**) for subsequent identification of this session. The **InitiatorSessionId** field SHALL NOT overlap with any other existing

PASE or CASE [session identifier](#) in use by the initiator. See [Section 4.13.2.4, “Choosing Secure Unicast Session Identifiers”](#) for more details.

3. The initiator SHALL generate a destination identifier ([DestinationId](#)) according to [Destination Identifier](#) to enable the responder to properly select a mutual Fabric and trusted root for the secure session.
4. The initiator SHALL [generate an ephemeral key pair](#) `InitiatorEphKeyPair = Crypto_GenerateKeyPair()`.
5. The initiator MAY encode any relevant [MRP parameters](#).
6. Any context-specific tags not listed in the above TLV schemas SHALL be reserved for future use, and SHALL be silently ignored if seen by a responder which cannot understand them.
7. If the initiator is resuming a session from a previous execution of the CASE with the same peer, the initiator SHALL:
 - a. Note the [ResumptionID](#) of the previous session.
 - b. Generate the [S1RK key](#).
 - c. Generate the [initiatorResumeMIC](#) using the [SharedSecret](#) from the previous session:

```
byte Resume1MIC_P[] = {}
byte Resume1MIC_A[] = {}
byte Resume1MIC_Nonce[13] = /* "NCASE_SigmaS1" */
    {0x4e, 0x43, 0x41, 0x53, 0x45, 0x5f, 0x53, 0x69,
     0x67, 0x6d, 0x61, 0x53, 0x31}

InitiatorResume1MIC = Crypto_AEAD_GenerateEncrypt(
    K = S1RK,
    P = Resume1MIC_P,
    A = Resume1MIC_A,
    N = Resume1MIC_Nonce
)
```

8. The initiator SHALL send a message with [Secure Channel Protocol ID](#) and [Sigma1 Protocol Opcode](#) from [Table 18, “Secure Channel Protocol Opcodes”](#) whose payload is the TLV-encoded [Sigma1 Msg1](#) with an anonymous tag for the outermost struct.

```
Msg1 =
{
    initiatorRandom      (1) = InitiatorRandom,
    initiatorSessionId   (2) = InitiatorSessionId,
    destinationId        (3) = DestinationId,
    initiatorEphPubKey    (4) = InitiatorEphKeyPair.publicKey
    initiatorSessionParams (5) = session-parameter-struct (optional),
    resumptionID          (6) = ResumptionID (optional, only present if performing
```

```

resumption),
  initiatorResumeMIC      (7) = InitiatorResumeMIC (optional, only present if
performing resumption)
}

```

Validate Sigma1

On receipt of **Msg1**, the responder SHALL perform the following:

1. If **Msg1** contains either a **resumptionID** or an **initiatorResumeMIC** field **but not both**, the responder SHALL send a **status report**: **StatusReport**(GeneralCode: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: INVALID_PARAMETER) and perform no further processing.
2. Set the **Peer Session Identifier** in the **Session Context** to the value **Msg1.initiatorSessionId**.
3. If **Msg1** contains both the **resumptionID** and **initiatorResumeMIC** fields, the responder SHALL search for an existing session that has a **Resumption ID** equal to the incoming **resumptionID**. If a single such session exists, the responder SHALL follow the steps in [Section 4.14.2.3.10](#), “**Validate Sigma1 with Resumption**” rather than continue the steps outlined in [Section 4.14.2.3.5](#), “**Validate Sigma1 Destination ID**”.
4. If **Msg1** does not contain a **resumptionID** and **initiatorResumeMIC** field, the responder SHALL continue the steps in [Section 4.14.2.3.5](#), “**Validate Sigma1 Destination ID**”.

Validate Sigma1 Destination ID

1. The responder SHALL validate the incoming **destinationId**:
 - a. The responder SHALL traverse all its installed **Node Operational Certificates (NOC)**, gathering the associated trusted roots' public keys from the associated chains and SHALL generate a **candidateDestinationId** based on the procedure in [Section 4.14.2.4.1](#), “**Destination Identifier**” for that tuple of <Root Public Key, Fabric ID, Node ID>.
 - b. The responder SHALL verify that the incoming **destinationId** matches one of the **candidateDestinationId** generated above. Upon such a match, the associated trusted root, Fabric ID, Node ID and IPK SHALL be recorded for subsequent use.
 - c. Note that at the initiator, only the current **Epoch Key** for the IPK will have been used. At the receiver, several IPK Epoch Keys may be installed, requiring several **candidateDestinationId** to be computed, one per available IPK Operational Key, per NOC.
2. If there is no **candidateDestinationId** matching the incoming **destinationId**, the responder SHALL send a **status report**: **StatusReport**(GeneralCode: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: NO_SHARED_TRUST_ROOTS) and perform no further processing.
3. Otherwise, if a match was found for the **destinationId**, the matched **NOC**, **ICAC** (if present), and associated trusted root SHALL be used for selection of the **responderNOC** and **responderICAC** in the steps for **Sigma2**.

Generate and Send Sigma2

If **validation** is successful, the responder encodes and sends a **Sigma2** message.


```

sigma-2-tbsdata => STRUCTURE [ tag-order ]
{
    responderNOC          [1] : OCTET STRING,
    responderICAC         [2, optional] : OCTET STRING,
    responderEphPubKey    [3] : ec-pub-key,
    initiatorEphPubKey    [4] : ec-pub-key,
}

sigma-2-tbdata => STRUCTURE [ tag-order ]
{
    responderNOC [1] : OCTET STRING,
    responderICAC [2, optional] : OCTET STRING,
    signature     [3] : ec-signature,
    resumptionID  [4] : OCTET STRING [ length 16 ],
}

sigma-2-struct => STRUCTURE [ tag-order ]
{
    responderRandom      [1] : OCTET STRING [ length 32 ],
    responderSessionId   [2] : UNSIGNED INTEGER [ range 16-bits ],
    responderEphPubKey   [3] : ec-pub-key,
    encrypted2           [4] : OCTET STRING,
    responderSessionParams [5, optional] : session-parameter-struct
}

```

NOTE

[sigma-2-tbsdata](#) is **NOT** transmitted but is instead signed; the signature will be encrypted and transmitted.

1. The responder SHALL [generate a random resumption ID](#) `ResumptionID = Crypto_DRBG(len = 16 * 8)`.
 - a. The responder SHALL set the [Resumption ID](#) in the [Session Context](#) to the value `ResumptionID`.
2. The responder SHALL use the Node Operational Key Pair `ResponderNOKeyPair`, `responderNOC`, and `responderICAC` (if present) corresponding to the NOC obtained in [Section 4.14.2.3.4, “Validate Sigma1”](#).
3. The responder SHALL [generate an ephemeral key pair](#) `ResponderEphKeyPair = Crypto_GenerateKeyPair()`.
4. The responder SHALL [generate a shared secret](#):

```

SharedSecret = Crypto_ECDH(
    privateKey = ResponderEphKeyPair.privateKey,
    publicKey  = Msg1.initiatorEphPubKey,

```

)

5. The responder SHALL encode the following items as a [sigma-2-tbsdata](#) with an [anonymous tag](#):
 - a. [responderNOC](#) as a [matter-certificate](#)
 - b. [responderICAC](#) (if present) as a [matter-certificate](#)
 - c. [ResponderEphKeyPair.publicKey](#)
 - d. [Msg1.initiatorEphPubKey](#)
6. The responder SHALL [generate a signature](#):

```
TBSDData2Signature = Crypto_Sign(
    message = sigma-2-tbsdata,
    privateKey = ResponderNOKeyPair.privateKey
)
```

7. The responder SHALL encode the following items as a [sigma-2-tbedata](#), where the encoding of [responderNOC](#) and [responderICAC](#) items SHALL be byte-for-byte identical to the encoding in [sigma-2-tbsdata](#):
 - a. [responderNOC](#) as a [matter-certificate](#). This encoding SHALL be byte-for-byte identical to the encoding in [sigma-2-tbsdata](#).
 - b. [responderICAC](#) (if present) as a [matter-certificate](#). This encoding SHALL be byte-for-byte identical to the encoding in [sigma-2-tbsdata](#).
 - c. [TBSDData2Signature](#)
 - d. [ResumptionID](#)
8. The responder SHALL [generate a random number](#) [Random](#) = [Crypto_DRBG\(len = 32 * 8\)](#).
9. The responder SHALL [generate the S2K key](#) using [Random](#) as Responder Random and [ResponderEphKeyPair.publicKey](#) as Responder Ephemeral Public Key.
10. The responder SHALL [generate the encrypted and integrity protected data](#):

```
byte TBEDData2_A[] = {}
byte TBEDData2_Nonce[13] = /* "NCASE_Sigma2N" */
    {0x4e, 0x43, 0x41, 0x53, 0x45, 0x5f, 0x53, 0x69,
     0x67, 0x6d, 0x61, 0x32, 0x4e}

TBEDData2Encrypted = Crypto_AEAD_GenerateEncrypt(
    K = S2K,
    P = TBEDData2,
    A = TBEDData2_A,
    N = TBEDData2_Nonce
)
```

11. The responder SHALL generate a **session identifier** (**ResponderSessionId**) for subsequent identification of this secured session. The **ResponderSessionId** field SHALL NOT overlap with any other existing PASE or CASE **session identifier** in use by the responder. See [Section 4.13.2.4, “Choosing Secure Unicast Session Identifiers”](#) for more details. The responder SHALL set the **Local Session Identifier** in the **Session Context** to the value **ResponderSessionId**.
12. The responder SHALL use the Fabric IPK configured as described in [Section 4.14.2.6.1, “Identity Protection Key \(IPK\)”](#).
13. Any context-specific tags not listed in the above TLV schemas SHALL be reserved for future use, and SHALL be silently ignored if seen by an initiator which cannot understand them.
14. The responder SHALL send a message with **Secure Channel Protocol ID** and **Sigma2 Protocol Opcode** from [Table 18, “Secure Channel Protocol Opcodes”](#) whose payload is the TLV-encoded **Sigma2 Msg2** with an anonymous tag for the outermost struct.

```
Msg2 =
{
  responderRandom      (1) = Random,
  responderSessionId   (2) = ResponderSessionId,
  responderEphPubKey   (3) = ResponderEphKeyPair.publicKey,
  encrypted2           (4) = TBEData2Encrypted,
  responderSessionParams (5) = session-parameter-struct (optional)
}
```

Validate Sigma2

On receipt of **Msg2**, the initiator SHALL perform the following:

1. The initiator SHALL [generate a shared secret](#):

```
SharedSecret = Crypto_ECDH(
  privateKey = InitiatorEphKeyPair.privateKey,
  publicKey = Msg2.responderEphPubKey,
)
```

2. The initiator SHALL [generate the S2K key](#) using **Msg2.responderRandom** as Responder Random and **Msg2.responderEphPubKey** as Responder Ephemeral Public Key.
3. The initiator SHALL [generate the decrypted data](#):

```
byte TBEData2_A[] = {}
byte TBEData2_Nonce[13] = /* "NCASE_Sigma2N" */
    {0x4e, 0x43, 0x41, 0x53, 0x45, 0x5f, 0x53, 0x69,
     0x67, 0x6d, 0x61, 0x32, 0x4e}

Success, TBEData2 = Crypto_AEAD_DecryptVerify(
```

```

    K = S2K,
    C = Msg2.encrypted2,
    A = TBEData2_A,
    N = TBEData2_Nonce
)

```

4. If the value of **Success** is **FALSE**, the initiator SHALL send a **status report**: **StatusReport(GeneralCode: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: INVALID_PARAMETER)** and perform no further processing.
5. The initiator SHALL verify that the **NOC** in **TBEData2.responderNOC** and **ICAC** in **TBEData2.responderICAC** (if present) fulfills the following constraints:
 - a. The Fabric ID and Node ID SHALL match the intended identity of the receiver Node, as included in the computation of the **Destination Identifier** when generating Sigma1.
 - b. If an **ICAC** is present, and it contains a Fabric ID in its subject, then it SHALL match the FabricID in the NOC leaf certificate.
 - c. The certificate chain SHALL chain back to the Trusted Root CA Certificate **TrustedRCAC** whose public key was used in the computation of the **Destination Identifier** when generating Sigma1.
 - d. All the elements in the certificate chain SHALL respect the **Matter Certificate DN Encoding Rules**, including range checks for identifiers such as Fabric ID and Node ID.
6. If any of the validations from the previous step fail, the initiator SHALL send a **status report**: **StatusReport(GeneralCode: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: INVALID_PARAMETER)** and perform no further processing.
7. The initiator SHALL verify **TBEData2.responderNOC** using:
 - a. **Success = Crypto_VerifyChain(certificates = [TBEData2.responderNOC, TBEData2.responderICAC, TrustedRCAC])**, when **TBEData2.responderICAC** is present, or
 - b. **Success = Crypto_VerifyChain(certificates = [TBEData2.responderNOC, TrustedRCAC])**, when **TBEData2.responderICAC** is not present.
8. If the value of **Success** is **FALSE**, the initiator SHALL send a **status report**: **StatusReport(GeneralCode: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: INVALID_PARAMETER)** and perform no further processing.
9. The initiator SHALL encode the following items as a **sigma-2-tbsdata** with an **anonymous tag**:
 - a. **responderNOC** as copied from **TBEData2**
 - b. **responderICAC** (if present) as copied from **TBEData2**
 - c. **Msg2.responderEphPubKey**
 - d. **InitiatorEphKeyPair.publicKey**
10. The initiator SHALL **verify** **TBEData2.signature** (see **RFC 5280**):

```

Success = Crypto_Verify(
    publicKey = Public key obtained from responderNOC,

```

```

    message = sigma-2-tbsdata,
    signature = TBEData2.signature
)

```

11. If the value of **Success** is **FALSE**, the initiator SHALL send a **status report**: **StatusReport(General-Code: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: INVALID_PARAMETER)** and perform no further processing.
12. Set the **Resumption ID** in the **Session Context** to the value **TBEData2.resumptionID**.
13. Set the **Peer Session Identifier** in the **Session Context** to the value **Msg2.responderSessionId**.

Generate and Send Sigma3

If **validation** is successful, the initiator encodes and sends a **Sigma3** message.

```

sigma-3-tbsdata => STRUCTURE [ tag-order ]
{
    initiatorNOC      [1] : OCTET STRING,
    initiatorICAC     [2, optional] : OCTET STRING,
    initiatorEphPubKey [3] : ec-pub-key,
    responderEphPubKey [4] : ec-pub-key,
}

sigma-3-tbdata => STRUCTURE [ tag-order ]
{
    initiatorNOC [1] : OCTET STRING,
    initiatorICAC [2, optional] : OCTET STRING,
    signature     [3] : ec-signature,
}

sigma-3-struct => STRUCTURE [ tag-order ]
{
    encrypted3 [1] : OCTET STRING,
}

```

NOTE

sigma-3-tbsdata is **NOT** transmitted but is instead signed; the signature will be encrypted and transmitted.

1. The initiator SHALL select its Node Operational Key Pair **InitiatorNOKeyPair**, Node Operational Certificates **initiatorNOC** and **initiatorICAC** (if present), and Trusted Root CA Certificate **Truste-dRCAC** corresponding to the chosen Fabric as determined by the **Destination Identifier** from Sigma1.
2. The initiator SHALL encode the following items as a **sigma-3-tbsdata** with an **anonymous tag**:
 - a. **initiatorNOC** as a **matter-certificate**

- b. `initiatorICAC` (if present) as a `matter-certificate`
- c. `InitiatorEphKeyPair.publicKey`
- d. `Msg2.responderEphPubKey`

3. The initiator SHALL `generate a signature`:

```
TBData3Signature = Crypto_Sign(
    message = sigma-3-tbsdata,
    privateKey = InitiatorNOKeyPair.privateKey
)
```

4. The initiator SHALL encode the following items as a `sigma-3-tbedata`:

- a. `initiatorNOC` as a `matter-certificate`. This encoding SHALL be byte-for-byte identical to the encoding in `sigma-3-tbsdata`.
- b. `initiatorICAC` (if present) as a `matter-certificate`. This encoding SHALL be byte-for-byte identical to the encoding in `sigma-3-tbsdata`.
- c. `TBData3Signature`

5. The initiator SHALL `generate the S3K key`.

6. The initiator SHALL `generate the encrypted and integrity protected data`:

```
byte TBData3_A[] = {}
byte TBData3_Nonce[13] = /* "NCASE_Sigma3N" */
    {0x4e, 0x43, 0x41, 0x53, 0x45, 0x5f, 0x53, 0x69,
     0x67, 0x6d, 0x61, 0x33, 0x4e}

TBData3Encrypted = Crypto_AEAD_GenerateEncrypt(
    K = S3K,
    P = TBData3,
    A = TBData3_A,
    N = TBData3_Nonce
)
```

- 7. Any context-specific tags not listed in the above TLV schemas SHALL be reserved for future use, and SHALL be silently ignored if seen by a responder which cannot understand them.
- 8. The initiator SHALL send a message with `Secure Channel Protocol ID` and `Sigma3 Protocol Opcode` from Table 18, “`Secure Channel Protocol Opcodes`” whose payload is the TLV-encoded `Sigma3 Msg3 = { encrypted3 (1) = TBData3Encrypted }` with an anonymous tag for the outermost struct.
- 9. The initiator SHALL generate the session encryption keys using the method described in Section 4.14.2.6.6, “`Session Encryption Keys`”.

Validate Sigma3

On receipt of **Msg3**, the responder SHALL perform the following:

1. The responder SHALL generate the **S3K** key.
2. The responder SHALL generate the decrypted data:

```
byte TBEData3_A[] = {}
byte TBEData3_Nonce[13] = /* "NCASE_Sigma3N" */
    {0x4e, 0x43, 0x41, 0x53, 0x45, 0x5f, 0x53, 0x69,
     0x67, 0x6d, 0x61, 0x33, 0x4e}

Success, TBEData3 = Crypto_AEAD_DecryptVerify(
    K = S3K,
    C = Msg3.encrypted3,
    A = TBEData3_A,
    N = TBEData3_Nonce
)
```

3. If the value of **Success** is **FALSE**, the responder SHALL send a **status report**: **StatusReport(GeneralCode: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: INVALID_PARAMETER)** and perform no further processing.
4. The responder SHALL verify that the **NOC** in **TBEData3.initiatorNOC** and the **ICAC** in **TBEData3.initiatorICAC** fulfill the following constraints:
 - a. The Fabric ID SHALL match the Fabric ID matched during processing of the **Destination Identifier** after receiving Sigma1.
 - b. If an **ICAC** is present, and it contains a Fabric ID in its subject, then it SHALL match the FabricID in the NOC leaf certificate.
 - c. The certificate chain SHALL chain back to the Trusted Root CA Certificate **TrustedRCAC** whose public key was matched during processing of the **Destination Identifier** after receiving Sigma1.
 - d. All the elements in the certificate chain SHALL respect the **Matter Certificate DN Encoding Rules**, including range checks for identifiers such as Fabric ID and Node ID.
5. If any of the validations from the previous step fail, the responder SHALL send a **status report**: **StatusReport(GeneralCode: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: INVALID_PARAMETER)** and perform no further processing.
6. The responder SHALL verify **TBEData3.initiatorNOC** using:
 - a. **Success = Crypto_VerifyChain(certificates = [TBEData3.initiatorNOC, TBEData3.initiatorICAC, TrustedRCAC])**, when **TBEData3.initiatorICAC** is present, or
 - b. **Success = Crypto_VerifyChain(certificates = [TBEData3.initiatorNOC, TrustedRCAC])**, when **TBEData3.initiatorICAC** is not present.
7. If the value of **Success** is **FALSE**, the responder SHALL send a **status report**: **StatusReport(General-**

Code: **FAILURE**, ProtocolId: **SECURE_CHANNEL**, ProtocolCode: **INVALID_PARAMETER_**) and perform no further processing.

8. The responder SHALL encode the following items as a **sigma-3-tbsdata** with an **anonymous** tag:
 - a. **initiatorNOC** as copied from **TBData3**
 - b. **initiatorICAC** (if present) as copied from **TBData3**
 - c. **Msg1.initiatorEphPubKey**
 - d. **ResponderEphKeyPair.publicKey**
9. The responder SHALL **verify TBData3.signature** (see **RFC 5280**):

```
Success = Crypto_Verify(
    publicKey= public key obtained from initiatorNOC,
    message = sigma-3-tbsdata,
    signature = TBData3.signature
)
```

10. If the value of **Success** is **FALSE**, the responder SHALL send a **status report**: **StatusReport(General-Code: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: INVALID_PARAMETER)** and perform no further processing.
11. The responder SHALL generate the session keys as described in **Section 4.14.2.6.6, "Session Encryption Keys"**.
12. The responder SHALL initialize its **Local Message Counter** in the **Session Context** per **Section 4.6.1.1, "Message Counter Initialization"**.
13. The responder SHALL initialize the **Message Reception State** in the **Session Context** and set the synchronized **max_message_counter** of the peer to **0**.
14. The responder SHALL set **SessionTimestamp** to a timestamp from a clock which would allow for the eventual determination of the last session use relative to other sessions.
15. The responder SHALL encode and send **SigmaFinished**.

Validate Sigma1 with Resumption

The responder SHALL continue validating the **Sigma1** message **Msg1** as follows:

1. Obtain the **SharedSecret** from the **Section 4.13.3.1, "Secure Session Context"** of the resumed session.
2. Generate the **S1RK** key.
3. Verify the **Resume1MIC** by **decrypting** the following values:

```
byte Resume1MIC_A[] = {}
byte Resume1MIC_Nonce[13] = /* "NCASE_SigmaR1" */
    {0x4e, 0x43, 0x41, 0x53, 0x45, 0x5f, 0x53, 0x69,
     0x67, 0x6d, 0x61, 0x53, 0x31}
```



```

Success, Resume1MICPayload = Crypto_AEAD_DecryptVerify(
    K = S1RK,
    C = Msg1.sigma1ResumeMIC,
    A = Resume1MIC_A,
    N = Resume1MIC_Nonce
)

```

4. If the value of **Success** is FALSE, the responder SHALL continue processing **Sigma1** as if it didn't include any resumption information by continuing the steps in [Section 4.14.2.3.5, "Validate Sigma1 Destination ID"](#).
5. If the value of **Success** is TRUE, the responder SHALL:
 - a. Set the **Peer Session Identifier** in the **Session Context** to the value **Msg1.initiatorSessionId**.
 - b. Send a **Sigma2_Resume** message.

Generate and Send Sigma2_Resume

The responder SHALL encode and send a **Sigma2_Resume** message in response to a valid **Sigma1** with response.

```

sigma-2-resume-struct => STRUCTURE [ tag-order ]
{
    resumptionID           [1] : OCTET STRING [ length 16 ],
    sigma2ResumeMIC        [2] : OCTET STRING [ length 16 ],
    responderSessionID     [3] : UNSIGNED INTEGER [ range 16-bits ],
    responderSessionParams [4, optional] : session-parameter-struct
}

```

1. The responder SHALL generate a new resumption ID **ResumptionID** = **Crypto_DRBG(len = 128)**.
2. The responder SHALL generate a **session identifier** (**ResponderSessionId**) for subsequent identification of this session. The **ResponderSessionId** field SHALL NOT overlap with any other existing PASE or CASE **session identifier** in use by the responder. See [Section 4.13.2.4, "Choosing Secure Unicast Session Identifiers"](#) for more details. The responder SHALL set the **Local Session Identifier** in the **Session Context** to the value **ResponderSessionId**.
3. The responder SHALL generate the **S2RK** key.
4. The responder SHALL generate a resumption MIC:

```

byte Resume2MIC_P[] = {}
byte Resume2MIC_A[] = {}
byte Resume2MIC_Nonce[13] = /* "NCASE_SigmaS2" */
    {0x4e, 0x43, 0x41, 0x53, 0x45, 0x5f, 0x53, 0x69,

```

```
0x67, 0x6d, 0x61, 0x53, 0x32}
```

```
Resume2MIC = Crypto_AEAD_GenerateEncrypt(
    K = S2RK,
    P = Resume2MIC_P,
    A = Resume2MIC_A,
    N = Resume2MIC_Nonce
)
```

5. Any context-specific tags not listed in the above TLV schemas SHALL be reserved for future use, and SHALL be silently ignored if seen by an initiator which cannot understand them.
6. The responder SHALL send a message with the **Secure Channel Protocol ID** and **Sigma2Resume Protocol Opcode** from Table 18, “Secure Channel Protocol Opcodes” whose payload is the TLV-encoded **Sigma2_Resume ResumeMsg2** with an anonymous tag for the outermost struct.

```
ResumeMsg2 =
{
    resumptionID          (1) = ResumptionID,
    sigma2ResumeMIC       (2) = ResumeMIC2,
    responderSessionID    (3) = ResponderSessionId,
    responderSessionParams (4) = session-parameter-struct (optional)
}
```

7. The responder SHALL generate the session keys as described in Section 4.14.2.6.7, “Resumption Session Encryption Keys”.

Validate Sigma2_Resume

On receipt of **ResumeMsg2**, the initiator SHALL perform the following:

1. The initiator SHALL **generate the S2RK key**.
2. The initiator SHALL verify the **Resume2MIC** by **decrypting** the following values:

```
byte Resume2MIC_A[] = {}
byte Resume2MIC_Nonce[13] = /* "NCASE_SigmaR2" */
    {0x4e, 0x43, 0x41, 0x53, 0x45, 0x5f, 0x53, 0x69,
     0x67, 0x6d, 0x61, 0x53, 0x32}
```

```
Success, Resume2MICPayload = Crypto_AEAD_DecryptVerify(
    K = S2RK,
    C = ResumeMsg2.sigma2ResumeMIC,
    A = Resume2MIC_A,
    N = Resume2MIC_Nonce
```

)

3. If **Success** is FALSE, the initiator SHALL send a **status report**: **StatusReport**(GeneralCode: FAILURE, ProtocolId: SECURE_CHANNEL, ProtocolCode: INVALID_PARAMETER_) and perform no further processing.
4. The initiator SHALL set the **Resumption ID** in the **Session Context** to the value **Resume2Msg.resumptionID**.
5. The initiator SHALL generate the session keys as described in [Section 4.14.2.6.7](#), “**Resumption Session Encryption Keys**”.
6. The initiator SHALL reset its **Local Message Counter** in the **Session Context** per [Section 4.6.1.1](#), “**Message Counter Initialization**”.
7. The initiator SHALL reset the **Message Reception State** of the **Session Context** and set the synchronized **max_message_counter** of the peer to 0.
8. The initiator SHALL set **SessionTimestamp** to a timestamp from a clock which would allow for the eventual determination of the last session use relative to other sessions.
9. The initiator SHALL set the **Peer Session Identifier** in the **Session Context** to the value **ResumeMsg2.responderSessionId**.
10. The initiator SHALL send [Section 4.14.2.3.13](#), “**SigmaFinished**”.

SigmaFinished

To indicate the successful completion of the protocol, the Node receiving Sigma3 (if a new session is being established) or Sigma2_Resume (if a session is being resumed) SHALL send a **status report**: **StatusReport**(GeneralCode: SUCCESS, ProtocolId: SECURE_CHANNEL, ProtocolCode: SESSION_ESTABLISHMENT_SUCCESS).

On successful receipt of **SigmaFinished**:

1. The receiving node SHALL initialize the **Local Message Counter** according to [Section 4.6.1.1](#), “**Message Counter Initialization**” for the newly established secure session whose success is acknowledged by this message.
2. The receiving node SHALL set **SessionTimestamp** to a timestamp from a clock which would allow for the eventual determination of the last session usage relative to other sessions.

If this message is received out-of-order or unexpectedly, then it SHALL be ignored.

4.14.2.4. Field Descriptions

Destination Identifier

destination-identifier => OCTET STRING [length CRYPTO_HASH_LEN_BYTES]

The Destination Identifier field enables the initiator of the **Sigma1** message to unambiguously express the following, in a privacy-preserving manner:

- Which shared Trusted Root to select
- Which Fabric ID to use for validation of initiator and responder operational certificates
- Which Node ID is targeted in the given Fabric

This serves several purposes:

1. It requires an initiator to have knowledge of both the [IPK](#) and one of the full identities of the responder Node before it forces the responder node to generate a costly [Sigma2](#) message
 - a. Note that the replay of previously recorded initiator messages is possible, and therefore a Node MAY choose to keep memory of some prior destination identifiers that were successfully processed which it would later reject if seen again, for additional replay protection
2. It ensures that there is no ambiguity on the responder as to which Fabric was selected for communication
3. It hides which Fabric was chosen by the initiator

A destination identifier is generated by:

1. Concatenating the following octet strings for subsequent usage as a [destinationMessage](#):
 - [initiatorRandom](#): The value of [initiatorRandom](#) that will be used in the same message as the Destination Identifier
 - [rootPublicKey](#): The public key of the root of trust of the desired fabric, from the [ec-pub-key](#) field of the [Matter Certificate](#) of that root, as an uncompressed elliptic curve point as defined in section 2.3.3 of [SEC 1](#)
 - [fabricId](#): The Fabric ID of the destination, matching the [matter-fabric-id](#) field of the [Matter Certificate](#) of the desired destination's [NOC](#), and encoding the 64-bit scalar as a little-endian byte order octet string
 - [nodeId](#): The Node ID of the destination, matching the [matter-node-id](#) field of the [Matter Certificate](#) of the desired destination's [NOC](#), and encoding the 64-bit scalar as a little-endian byte order octet string
2. Obtaining the appropriate [Identity Protection Key \(IPK\) Operational Group Key](#) for the associated [Fabric](#) under [Group Key Set index](#) 0 within the [Group Key Management Cluster](#).
3. Computing an identifier [destinationIdentifier](#) of length [CRYPTO_HASH_LEN_BYTES](#) using [Crypto_HMAC\(\)](#) with the [IPK](#) as the [key](#) and [destinationMessage](#) as the [message](#)

The above steps can be summarized as:

```
destinationMessage = initiatorRandom || rootPublicKey || fabricId || nodeId
destinationIdentifier = Crypto_HMAC(key=IPK, message=destinationMessage)
```

For example, given the following:

- Root public key for the common Fabric, in uncompressed elliptical curve point form:

```
RootPublicKey := // Raw uncompressed point form
04:4a:9f:42:b1:ca:48:40:d3:72:92:bb:c7:f6:a7:e1:
1e:22:20:0c:97:6f:c9:00:db:c9:8a:7a:38:3a:64:1c:
b8:25:4a:2e:56:d4:e2:95:a8:47:94:3b:4e:38:97:c4:
a7:73:e9:30:27:7b:4d:9f:be:de:8a:05:26:86:bf:ac:
fa
```

- Common Fabric ID of 0x2906_C908_D115_D362 scalar (octets "62:d3:15:d1:08:c9:06:29" in little-endian)
- Desired Destination Node ID of 0xCD55_44AA_7B13_EF14 (octets "14:ef:13:7b:aa:44:55:cd" in little-endian)
- [Identity Protection Key](#) Epoch Key value of:

```
IPKEpochKey := 4a:71:cd:d7:b2:a3:ca:90:24:f9:6f:3c:96:a1:9d:ee
```

- Note that this is the octet string of a group Epoch Key as would be provided in the [IPKValue](#) field of the [AddNOC](#) command in the [Operational Credentials Cluster](#), or in one of the [EpochKey](#) fields of the [KeySetWrite](#) command in the [Group Key Management Cluster](#).
- The [derived Operational Group Key](#) to be used for computation of a destination identifier, given the above values of root public key, Fabric ID and Identity Protection Key Epoch Key, would be:

```
IPK := 9b:c6:1c:d9:c6:2a:2d:f6:d6:4d:fc:aa:9d:c4:72:d4
```

- Initiator Random value of:

```
7e:17:12:31:56:8d:fa:17:20:6b:3a:cc:f8:fa:ec:2f:
4d:21:b5:80:11:31:96:f4:7c:7c:4d:eb:81:0a:73:dc
```

Then, using the above procedure would yield the following:

- DestinationMessage octets:

```
7e:17:12:31:56:8d:fa:17:20:6b:3a:cc:f8:fa:ec:2f:
4d:21:b5:80:11:31:96:f4:7c:7c:4d:eb:81:0a:73:dc:
04:4a:9f:42:b1:ca:48:40:d3:72:92:bb:c7:f6:a7:e1:
1e:22:20:0c:97:6f:c9:00:db:c9:8a:7a:38:3a:64:1c:
b8:25:4a:2e:56:d4:e2:95:a8:47:94:3b:4e:38:97:c4:
a7:73:e9:30:27:7b:4d:9f:be:de:8a:05:26:86:bf:ac:
fa:62:d3:15:d1:08:c9:06:29:14:ef:13:7b:aa:44:55:
cd
```

- DestinationIdentifier octets:

dc:35:dd:5f:c9:13:4c:c5:54:45:38:c9:c3:fc:42:97:
c1:ec:33:70:c8:39:13:6a:80:e1:07:96:45:1d:4c:53

Public Key

ec-pub-key => OCTET STRING [length CRYPTO_PUBLIC_KEY_SIZE_BYTES]

A public key **ec-pub-key** is the byte string representation of an uncompressed elliptic curve point as defined in section 2.3.3 of [SEC 1](#).

4.14.2.5. Signature

An **ec-signature** is the encoding of the signature as defined in [Section 3.5.3, “Signature and verification”](#).

ec-signature => OCTET STRING [length (CRYPTO_GROUP_SIZE_BYTES * 2)]

4.14.2.6. Cryptographic Keys

Identity Protection Key (IPK)

The Identity Protection Key (IPK) SHALL be the operational group key under [GroupKeySetID](#) of **0** for the fabric associated with the originator’s chosen destination.

The IPK SHALL be exclusively used for [Certificate Authenticated Session Establishment](#). The IPK SHALL NOT be used for [operational group communication](#).

For the generation of the [Destination Identifier](#), the originator SHALL use the operational group key with the second newest EpochStartTime, if one exists, otherwise it SHALL use the single operational group key available.

The operational group key index to use to follow the "second newest EpochStartTime" rule is illustrated below:

Number of keys in Group Key Set	Operational key index	Epoch Key
1	0	EpochKey0
2	0	EpochKey0
3	1	EpochKey1

Sigma2 Key (S2K)

1. A transcript hash SHALL be [generated](#):

```
TranscriptHash = Crypto_Hash(message = Msg1)
```

2. The Sigma2 key SHALL be [generated](#):

```
byte S2K_Info[] = /* "Sigma2" */
    {0x53, 0x69, 0x67, 0x6d, 0x61, 0x32}

S2K = Crypto_KDF(
    inputKey = SharedSecret,
    salt = IPK || Responder Random || Responder Ephemeral Public Key ||
    TranscriptHash,
    info = S2K_Info,
    len = CRYPTO_SYMMETRIC_KEY_LENGTH_BITS
)
```

where `||` indicates message concatenation and **IPK** is generated according to [Section 4.14.2.6.1](#), “Identity Protection Key (IPK)”.

Sigma3 Key (S3K)

1. A transcript hash SHALL be [generated](#):

```
TranscriptHash = Crypto_Hash(message = Msg1 || Msg2)
```

2. The Sigma3 key SHALL be [generated](#):

```
byte S3K_Info[] = /* "Sigma3" */
    {0x53, 0x69, 0x67, 0x6d, 0x61, 0x33}

S3K = Crypto_KDF(
    inputKey = SharedSecret,
    salt = IPK || TranscriptHash,
    info = S3K_Info,
    len = CRYPTO_SYMMETRIC_KEY_LENGTH_BITS
)
```

where `||` indicates message concatenation and **IPK** is generated according to [Section 4.14.2.6.1](#), “Identity Protection Key (IPK)”.

Sigma1 Resumption Key

The Sigma1 resumption key SHALL be [generated](#):

```

byte S1RK_Info[] = /* "Sigma1_Resume" */
    {0x53, 0x69, 0x67, 0x6d, 0x61, 0x31, 0x5f,
     0x52, 0x65, 0x73, 0x75, 0x6d, 0x65}

S3K_Info
S1RK = Crypto_KDF(
    inputKey = SharedSecret,
    salt = Sigma1.initiatorRandom || ResumptionID,
    info = S1RK_Info,
    len = CRYPTO_SYMMETRIC_KEY_LENGTH_BITS
)

```

where `||` indicates message concatenation and `ResumptionID` is the identifier for the previous session.

Sigma2 Resumption Key

The Sigma2 resumption key SHALL be [generated](#):

```

byte S2RK_Info[] = /* "Sigma2_Resume" */
    {0x53, 0x69, 0x67, 0x6d, 0x61, 0x32, 0x5f,
     0x52, 0x65, 0x73, 0x75, 0x6d, 0x65}

S2RK = Crypto_KDF(
    inputKey = SharedSecret,
    salt = Sigma1.initiatorRandom || ResumptionID,
    info = S2RK_Info,
    len = CRYPTO_SYMMETRIC_KEY_LENGTH_BITS
)

```

where `||` indicates message concatenation and `ResumptionID` is the new identifier for the this session.

Session Encryption Keys

1. A transcript hash SHALL be [generated](#):

```
TranscriptHash = Crypto_Hash(message = Msg1 || Msg2 || Msg3)
```

2. The session encryption keys SHALL be [generated](#):

```

byte SEKeys_Info[] = /* "SessionKeys" */
    {0x53, 0x65, 0x73, 0x73, 0x69, 0x6f, 0x6e, 0x4b,

```



```

    0x65, 0x79, 0x73}

I2RKey || R2IKey || AttestationChallenge = Crypto_KDF(
    inputKey = SharedSecret,
    salt = IPK || TranscriptHash,
    info = SEKeys_Info,
    len = 3 * CRYPTO_SYMMETRIC_KEY_LENGTH_BITS
)

```

- Each key is exactly **CRYPTO_SYMMETRIC_KEY_LENGTH_BITS** bits.
- The initiator SHALL use **I2RKey** to encrypt and integrity protect messages and the `R2IKey' to decrypt and verify messages.
- The responder SHALL use **R2IKey** to encrypt and integrity protect messages and the `I2RKey' to decrypt and verify messages.
- The **AttestationChallenge** SHALL only be used as a challenge during device attestation. See [Section 6.2.3, “Device Attestation Procedure”](#) for more details.

where `||` indicates message concatenation.

Resumption Session Encryption Keys

- The resumption session encryption keys SHALL be [generated](#):

```

byte RSEKeys_Info[] = /* "SessionResumptionKeys" */
    {0x53, 0x65, 0x73, 0x73, 0x69, 0x6f, 0x6e, 0x52,
     0x65, 0x73, 0x75, 0x6d, 0x70, 0x74, 0x69, 0x6f,
     0x6e, 0x4b, 0x65, 0x79, 0x73}

I2RKey || R2IKey || AttestationChallenge = Crypto_KDF(
    inputKey = SharedSecret,
    salt = Sigmal.initiatorRandom || ResumptionID,
    info = RSEKeys_Info,
    len = 3 * CRYPTO_SYMMETRIC_KEY_LENGTH_BITS
)

```

- Each key is exactly **CRYPTO_SYMMETRIC_KEY_LENGTH_BITS** bits.
- The initiator SHALL use **I2RKey** to encrypt and integrity protect messages and the `R2IKey' to decrypt and verify messages.
- The responder SHALL use **R2IKey** to encrypt and integrity protect messages and the `I2RKey' to decrypt and verify messages.
- The **AttestationChallenge** SHALL only be used as a challenge during device attestation. See [Section 6.2.3, “Device Attestation Procedure”](#) for more details.

where `||` indicates message concatenation and `ResumptionID` is the new identifier for the this session.

4.14.2.7. Session Context Storage

After the session is established successfully at both peers, some fields SHALL be recorded in the secure session context for later use (see [Section 4.13.3, “Application Data Phase”](#)), in addition to the [Session Encryption Keys](#):

- The peer NOC's `matter-node-id` ([1.3.6.1.4.1.37244.1.1](#)) from the subject field
- The [Section 2.5.1, “Fabric References and Fabric Identifier”](#) for the Fabric within which this secure session is being established
- All peer NOC's `case-authenticated-tag` ([1.3.6.1.4.1.37244.1.6](#)) from the subject field, if present

These fields MAY be recorded at any opportune point during the protocol, but SHALL only be committed to the secure session context once the session is established successfully at both peers.

4.14.2.8. Minimal Number of CASE Sessions

A node SHALL support at least 3 CASE session contexts per fabric. Device type specifications MAY require a larger minimum. Unless the device type specification says otherwise, a minimum number it defines is a per-fabric minimum.

The minimal supported capabilities, subject to the minimal constraints above, are reported in the [CapabilityMinima](#) of the [Basic Information cluster](#).

- Example: If a device type requires at least 4 CASE session contexts, and a node supports 7 fabrics, the node would support at least 28 CASE session contexts, and ensure that each fabric could use at least 4 of them.

4.15. Secure Communications over TCP

Matter messages over [MRP](#) need to be sized so that each UDP datagram fits within an IPv6 MTU limit of 1280 bytes, to avoid IPv6 fragmentation and reassembly. Thus, larger messages cannot be transferred over MRP unless the application explicitly performs fragmentation and reassembly. Therefore, nodes that need to send [large command messages](#) must use an alternative transport protocol, such as TCP [[RFC 9293](#)]. TCP support is communicated by nodes through the [TXT record key T](#) in their DNS-SD operational service records.

Since TCP already provides message transmission reliability, a node that is using TCP as the underlying transport protocol SHALL NOT use MRP reliability semantics on its message exchanges.

4.15.1. Secure Session Establishment

When a pair of nodes that both support TCP intend to establish a secure session between themselves, they MAY choose TCP as the underlying transport protocol. With a given peer, a node SHOULD, typically, either have a secure session over MRP or one over TCP, but MAY also have both. For example, a node might be using MRP for all its interactions with a controller, but might set up a session over TCP with the same for specific data-heavy operations such as OTA.

The underlying TCP connection MAY be long-lived or short-lived depending on the use case. Unlike MRP, which does not rely on an underlying connection, a secure session over TCP is unusable when its connection is broken or is closed. Nodes MAY choose to remove the secure session when the connection goes down. If the session is retained after the connection goes away, then the session SHALL be marked appropriately so that the underlying connection is re-established before the session can be used again. Moreover, the [session resumption](#) state MAY be retained to expedite session establishment when the connection is re-established with the corresponding peer.

4.15.2. TCP Connection Management

A TCP connection between a pair of nodes requires appropriate configurations and active management in order to meet the responsiveness demands of the application.

4.15.2.1. TCP Connection Configuration

The TCP connection configuration parameters MAY include the following:

1. **Connection Establishment Timeout:** A node MAY use this option to configure the amount of time that it waits for an initiated connection establishment attempt to succeed, before giving up and notifying the application.
2. **TCP Keep Alive Messages For Connection Liveness:** TCP Keep Alive messages MAY be used to maintain liveness for long-lived connections. They MAY be used at both the client and server to configure the liveness for each half of the connection. The configurable parameters SHALL be:
 - a. *TCP_KEEP_ALIVE_TIME* : The interval between the last data packet sent and the first keep-alive probe.
 - b. *TCP_KEEP_ALIVE_INTERVAL* : The interval between subsequent keep-alive probes.
 - c. *TCP_KEEP_ALIVE_PROBES* : The number of unacknowledged probes to send before considering the connection dead and notifying the application.
3. **TCP User Timeout:** The TCP User Timeout option specifies the amount of time that transmitted data may remain unacknowledged before the TCP connection is forcibly closed. This option MAY be used by a node to limit the time the connection stays open when the data flow is not progressing.

The total **TCP Keep Alive Timeout** is given by:

$$\text{TCP Keep Alive Timeout} = \text{TCP_KEEP_ALIVE_TIME} + \text{TCP_KEEP_ALIVE_INTERVAL} * \text{TCP_KEEP_ALIVE_PROBES}$$

4.15.2.2. TCP Connection Closures And Re-establishment

The TCP connection MAY be closed by either side of the connection. A node MAY proactively close its connection when instructed by the upper layer. It MAY also close the connection based on its current state.

For example, it MAY close its connection under the following circumstances:

1. The **TCP Keep Alive Timeout** expires on an idle connection.

2. Sent data on the connection is not acknowledged at the TCP layer, and the configured **TCP User Timeout** expires.

When the TCP layer of a node gets notified that the peer has closed the connection, it SHALL close its end of the connection as well, and notify the application. Subsequently, all active **Exchanges** over that connection SHOULD also be closed as they would be unusable over a closed connection.

Either side MAY choose to re-establish the connection when it is closed or determined to be broken. A node SHOULD back-off (typically, based on a Fibonacci back-off sequence) a random amount of time after the connection closure, before attempting to establish the connection again.

A node, upon receipt of a connection request from a peer, SHOULD discard its own backed-off connection retry timer to the same peer, if one is active.

This random back-off mechanism SHOULD prevent connection races between peers for most, if not all scenarios. In addition, nodes SHOULD try to reap old unused connections as much as possible to conserve resources.

4.15.2.3. Memory Requirements for Large messages

Since the TCP transport is designed for large messages, the system platform is expected to have memory available for storing the received messages. The system platform MAY configure a **Maximum Message Size** for the payload that it is capable of receiving over the TCP connection. If a node receives a message header that indicates that the message is larger than the **Maximum Message Size** that it supports, then it SHALL close the connection, and SHOULD send a **Status Report** error message with a status code set to **MESSAGE_TOO_LARGE** back to the sender, before closing the connection.

Given that a node MAY be using both MRP and TCP for sending/receiving messages with different peers, and these transports have different **Maximum Message Size** configurations, the system SHOULD be able to dynamically change this configuration based on what transport the current Exchange is using.

4.16. Group Communication

This section specifies the semantics of sending and receiving multicast group messages and the lifecycle of such groupcast sessions. Multicast addressing is accomplished using the 16-bit **Group ID** field as the destination address. A multicast group is a collection of Nodes, all registered under the same multicast Group ID. A multicast message is sent to a particular destination group and is received by all members of that group.

4.16.1. Groupcast Session Context

Groupcast sessions are conceptually long-running, lasting the duration of a node's membership in a group. Group membership is tracked in the **Group Key Management Cluster**. However, on ingress of each groupcast message, the following ephemeral context SHALL be constructed to inform upper layers of groupcast message provenance:

1. **Fabric Index**: Records the local **Fabric Index** for the Fabric to which an incoming message's group is scoped.

2. **Group ID**: Captures the **Group ID** to which a groupcast message was sent.
3. **Source Node ID**: The **Source Node ID** enclosed by the sender of a groupcast message.
 - Together, **Fabric Index**, **Group ID** and **Source Node ID** comprise a unique identifier that upper layers may use to understand the source and destination of groupcast messages.
4. **Source IP Address**: The unicast source IP address for the originator of the message.
5. **Source Port**: The source port for the originator of the message.
 - The source IP address and port MAY be used for unicast responses to group communication peers, as are required for the **Message Counter Synchronization Protocol**.
6. **Operational Group Key**: The **Operational Group Key** that was used to encrypt the incoming group message.
7. **Group Session ID**: Records the **Group Session ID** derived from the **Operational Group Key** used to encrypt the message.

Once a Groupcast Session Context with **trust-first** policy is created to track authenticated messages from a given Source Node ID, that record SHALL NOT be deleted or recycled until the node reboots. This is to prevent replay attacks that first exhaust the memory allocated to group session counter tracking and then inject older messages as valid, and enforces that trust-first authentication works as intended within the full duration of a boot cycle. Any message from a source that cannot be tracked SHALL be dropped.

4.16.2. Sending a group message

To prepare a multicast message to a Group ID with a given **GroupKeySetID** and IPv6 hop count parameter, the Node SHALL:

1. Obtain, for the given **GroupKeySetID**, the current **Operational Group Key** as the Encryption Key, and the associated **Group Session ID**.
 - a. If no key is found for the given **GroupKeySetID**, security processing SHALL fail and no further security processing SHALL be done on this message.
2. Perform **Section 4.7.1, “Message Transmission”** processing steps on the message with the following arguments:
 - a. The Destination Node Id argument SHALL be the **Group Node Id** corresponding to the given Group ID.
 - b. The Session Id argument SHALL be the Group Session ID from step 1.
 - c. The Security Flags SHALL have only the **P Flag** set.
 - d. The transport SHALL be a site-local routable IPv6 interface.

Next, prepare the message as an IPv6 packet:

1. Set the private, secured message from step 2 above as the IPv6 payload.
2. Set the IPv6 hop count to the value given.
3. Set the IPv6 destination to the **Section 2.5.6.2, “IPv6 Multicast Address”** based on the provided destination Group ID, Fabric ID and the address policy of that group. Specifically:

- a. Look up the group data record based on the Fabric ID and Group ID.
 - b. If the group data record is sourced by the Groupcast cluster, set the IPv6 address according to the [Section 11.27.5.4.5, “McastAddrPolicy”](#) field.
 - c. Otherwise set the IPv6 address according to the PerGroup Multicast Address rules.
4. Set the IPv6 source to an operational [IPv6 Unicast Address](#) of the sending Node.
 5. Set the IPv6 UDP port number to the [Matter IPv6 multicast port](#).

4.16.3. Receiving a group message

All Nodes supporting groups SHALL register to receive on the associated [IPv6 multicast address](#), at the [Matter IPv6 multicast port](#), for each group of which they are a member.

Upon receiving an IPv6 message addressed to one of these Multicast Addresses the Node is registered for, the Node SHALL:

1. Extract the message from the IPv6 payload.
2. Perform [Section 4.7.2, “Message Reception”](#) processing steps on the message.

4.17. Group Key Management

This section describes operational group keys, a mechanism for generating, disseminating and managing shared symmetric keys across a group of Nodes in a Fabric. Operational group keys are made available to applications for the purpose of authenticating peers, securing group communications, and encrypting data. These keys allow Nodes to:

- Prove to each other that they are members of the associated group
- Exchange messages confidentially, and without fear of manipulation by non-members of the group
- Encrypt data in such a way that it can only be decrypted by other members of the group

A central feature of operational group keys is the ability to limit access to keys to a trusted set of Nodes. In particular, credentials required to generate operational group keys SHALL only be accessible to Nodes with a certain level of privilege — those deemed a member of the group. Barring software error or compromise of a privileged Node, access to shared keys SHALL be computationally infeasible for non-trusted parties.

Operational group keys are shareable across all types and combinations of Nodes as determined by the Administrator managing the group. Given all Nodes in possession of the current epoch keys for the group can communicate with other Nodes in the group, it is the responsibility of the Administrator managing the group to only compose groups of Nodes where communication is appropriate for the given application and security requirements.

4.17.1. Operational Groups

An operational group is a logical collection of Nodes that are running one or more common application clusters and share a common security domain in the form of a shared, symmetric group key.

For example, a set of Nodes running a lighting application can form an operational group by sharing a common operational group key derived from the mechanisms described here. Subgroups can be formed within the operational group by defining distinct Group Identifiers for each set of Nodes, while sharing a common operational group key.

Membership in the security domain of an operational group is determined by a Node's possession of all the epoch keys required to generate the current, valid *operational group key* for the group. Individual Nodes can be members of multiple operational groups simultaneously. The set of groups to which a Node belongs can change over time as dictated by application requirements and policies. Groups MAY be introduced or withdrawn over time as need arises.

4.17.1.1. Operational Group Ids

Operational groups are identified uniquely within a Fabric by a [Group Identifier](#). Administrators SHALL assign Group Ids such that no two operational groups within a Fabric have the same Group ID. It is assumed a given Administrator has sufficient access to centralized knowledge, so as to allocate unique Group Ids under a given Fabric such that there are no collisions.

Multiple operational groups MAY share the same operational group key, and thus be used to create logical subgroups over a shared security domain. Operational groups which do not share related functionality, such as a lighting group and a sprinkler group, SHOULD NOT share the same operational key. As an example policy, a lighting application could have all lighting Nodes share a single group key, while organizing lighting subgroups for various rooms or spaces within the structure by assigning a different Group ID to each such subgroup.

4.17.2. Operational Group Key Derivation

An *operational group key* is a symmetric key used as the Encryption Key during [Message Processing](#) for group communication. An *operational group key* is produced by applying a key derivation function with an *epoch key* and salt value as inputs as follows:

```
OperationalGroupKey =  
    Crypto_KDF  
    (  
        InputKey = Epoch Key,  
        Salt = CompressedFabricIdentifier,  
        Info = "GroupKey v1.0",  
        Length = CRYPTO_SYMMETRIC_KEY_LENGTH_BITS  
    )
```

where `[]` denotes a zero-length array.

The **Info** portion of the key derivation is specified in [Section 4.17.2.1, “Group Security Info”](#). The **Salt** portion of the key derivation is specified in [Section 4.3.2.2, “Compressed Fabric Identifier”](#).

For example, given:

- An Epoch Key value of: `23:5b:f7:e6:28:23:d3:58:dc:a4:ba:50:b1:53:5f:4b`

- A CompressedFabricIdentifier value of: 87:e1:b0:04:e2:35:a1:30

After the above derivation following the definition of **Crypto_KDF** in [Section 3.8, “Key Derivation Function \(KDF\)”](#), the resulting operational group key would be: a6:f5:30:6b:af:6d:05:0a:f2:3b:a4:bd:6b:9d:d9:60.

Group membership is enforced by limiting access to the epoch keys. Only Nodes that possess the input epoch key can derive a given operational key. Lack of possession of a particular epoch key restricts access, based on the distribution policy of the epoch keys.

The following diagram shows the process by which operational keys are derived from the epoch key material:




Figure 17. Group Key Derivation

4.17.2.1. Group Security Info

A hard-coded *group security info* is used to diversify the set of operational group keys. This value is hard-coded into the standard’s implementation, and thus is distributed with the associated code. Should the standard’s security mechanisms need to evolve (e.g. to upgrade encryption from AES-128 to AES-256), the *group security info* can be changed to ensure that new keys will be derived for use in the new algorithm. The *group security info* SHALL be the byte stream "GroupKey v1.0", i.e. 0x47 0x72 0x6f 0x75 0x70 0x4b 0x65 0x79 0x20 0x76 0x31 0x2e 0x30.

With the exception of the *group security info*, all input key material SHALL be maintained on a per-Fabric basis.

4.17.2.2. Group Key Set

A *group key set* limits the key derivation process to Nodes within the respective operational groups. Access to a *group key set* is limited based on the functionality provided by a Node and/or the privilege afforded to it. For example, certain home security devices, such as a security system or door lock, may have access to a "Physical Access" *group key set*, while devices such as light bulbs or window coverings would not.

Operational group key lifetime is limited by assigning an expiration time to each *epoch key* in a given *group key set*. By constraining the validity of a given epoch key to an epoch, the ability for members to derive and operate with an operational group key can be constrained to particular periods of time. Epoch keys may be rotated on a periodic basis, and denying access to updated versions of these keys serves as a means to eject group members.

4.17.3. Epoch Keys

Epoch keys provide a means for limiting the lifetime of derived operational group keys. They also provide a way for an Administrator to revoke access to Nodes that have been explicitly excluded from an operational group (albeit after a period of time).

Epoch keys are generated, managed, and stored by an Administrator on a per-Fabric basis. Each key SHALL be a [random](#) value of length [CRYPTO_SYMMETRIC_KEY_LENGTH_BITS](#) bits.

```
EpochKey = Crypto_DRBG(len = CRYPTO_SYMMETRIC_KEY_LENGTH_BITS)
```

Each epoch key has associated with it a start time that denotes the time at which the key becomes active for use by transmitting Nodes. Epoch key start times are absolute UTC time in microseconds encoded using the [epoch-us](#) representation.

4.17.3.1. Using Epoch Keys

Nodes sending group messages SHALL use operational group keys that are derived from the *current epoch key* (specifically, the epoch key with the *latest* start time that is not in the future). Nodes that cannot reliably keep track of time calculate the *current epoch key* as described in [Section 4.17.3.4, “Epoch Key Rotation without Time Synchronization”](#).

Nodes receiving group messages SHALL accept the use of any key derived from one of the currently installed epoch keys. This requirement holds regardless of whether the start time for the key is in the future or the past. This means Nodes continue to accept communication secured under an epoch key until that key is withdrawn by explicitly deleting the key from a Node’s group state by the key distribution Administrator.

Note that there is no end time associated with an epoch key. An epoch key marked with the maximum start time SHALL be disabled and render the corresponding epoch key slot unused.

4.17.3.2. Managing Epoch Keys

The epoch keys are managed using the [Group Key Management Cluster](#). For every group key set published by the key distribution Administrator, there SHALL be at least 1 and at most 3 epoch keys in rotation. Key additions or updates are done using the [KeySetWrite](#) command.

Key updates are idempotent operations to ensure the Administrator is always the source of truth. An epoch key update SHALL order the keys from oldest to newest.

Any epoch key update MAY deliver a partial key set but SHALL include [EpochKey0](#) and MAY include [EpochKey1](#) and [EpochKey2](#). Any update of the key set, including a partial update, SHALL remove all previous keys in the set, however many were defined.

An Administrator MAY completely remove a group key set from a Node using the [KeySetRemove](#) command.

4.17.3.3. Epoch Key Rotation

The key distribution Administrator generates new epoch keys on a regular basis, giving each a unique id and adding them to the list of existing epoch keys within a group. The start time for each new epoch key is scheduled to occur after a configurable *key propagation interval*. The propagation interval is set sufficiently large such that the Administrator can synchronize all Nodes in the operational group with the new epoch key list within that time.

The rotation rate for epoch keys is expected to be on the order of days to weeks for typical applications, but the rate is configurable as required by the key distribution Administrator. Because of the relatively long rotation interval, and the overlap of active epoch keys, local clock drift within Nodes is generally not a concern.

4.17.3.4. Epoch Key Rotation without Time Synchronization

Although epoch keys are distributed with an associated start time, it is nonetheless possible for Nodes that do not maintain a synchronized clock to participate in key rotation. Specifically, upon receiving a new epoch key list from the key distribution Administrator, such a Node can note which of the keys is the current epoch key by comparing their relative start times and using the current epoch key which has the second newest time. It can then use the current key for all locally initiated security interactions until such time as it makes contact with the distribution Administrator again.

This scheme requires the Node to receive epoch keys from the key distribution Administrator at a rate that is at least as fast as the configured *key propagation interval*. The Administrator SHOULD provide a sufficient set of epoch keys to Nodes that do not maintain synchronized time so that they can maintain communication with other group members while a key update is in progress. The key distribution Administrator SHOULD update all Nodes without time, such as ICDs, before the *new epoch key* is activated, and then let Nodes with time all roll to the *new epoch key* at the synchronized start time.

4.17.3.5. Epoch Key Rotation logic

The full life-cycle of Epoch Key rotation is shown in [Figure 18, “Epoch Key Rotation”](#). For each epoch key slot, the start time of the key is shown as one of the following values:

- **New** - a key with a start time in the future.
- **Current** - the active key with the newest start time.
- **Previous** - the active key with the second newest start time.
- **Old** - an active key with the third newest start time.

The diagram shows two types of state transitions:

1. **Admin** - an update of an old key by the Administrator. Changes made during this state transition are indicated in green.
2. **Epoch Activate** - activation of an epoch key due to system time becoming greater than the start time. Changes during this state transition are indicated in yellow.




Figure 18. Epoch Key Rotation

The **Admin Refresh** state begins when an entire *group key set* is freshly written to a Node during commissioning or administration, such as when a new group is added. The **Epoch Activate** state is entered when time progresses to activate a fresh *current epoch key*, aging out the other epoch key slots. The **Admin Update** state is entered when an Administrator updates an *old epoch key* with a *new epoch key*. When in steady state, the **Admin Refresh** state MAY be entered in place of an **Admin Update** state transition to update additional keys to the required ones or to completely reset the group security.

Note that in the above diagram, only an example key distribution scheme is illustrated. It is also possible to devise key distribution algorithms that do not rely on time synchronization, or group configurations that never rotate keys in favor of configuring new groups and removing groups and group key sets with expired keys.

Group Key Set ID

The Group Key Set ID is a 16-bit field that uniquely identifies the set of epoch keys used for derivation of an operational group key. Each Group Key Set ID is scoped to a particular Fabric and assigned by an Administrator so as to be unique within a Fabric.

The Group Key Set ID of 0 SHALL be reserved for managing the **Identity Protection Key (IPK)** on a given Fabric. It SHALL NOT be possible to remove the IPK Key Set if it exists.

4.17.3.6. Group Session ID

A *Group Session ID* is a special case of a 16-bit **Session ID** that is specifically used for group communication. When **Session Type** is 1, denoting a group session, the **Session ID** SHALL be a *Group Session ID* as defined here. The *Group Session ID* identifies probable operational group keys across a Fabric. The *Group Session ID* for a given *operational group key* is derived by treating the output of a **Crypto_KDF** against the associated Operational Group Key as a big-endian representation of a 16-bit

integer, as follows:

```
GroupKeyHash =
    Crypto_KDF
    (
        InputKey = OperationalGroupKey,
        Salt = [],
        Info = "GroupKeyHash",
        Length = 16 // Bits
    )

// GroupKeyHash is an array of 2 bytes (16 bits) per Crypto_KDF

// GroupSessionId is computed by considering the GroupKeyHash as a Big-Endian
// value. GroupSessionId is a scalar. Its use in fields within messages may cause a
// re-serialization into a different byte order than the one used for initial
// generation.
GroupSessionId = (GroupKeyHash[0] << 8) | (GroupKeyHash[1] << 0)
```

where `[]` denotes a zero-length array.

For example, given:

- An Operational Group Key value of: `a6:f5:30:6b:af:6d:05:0a:f2:3b:a4:bd:6b:9d:d9:60`

After the above derivation following the definition of `Crypto_KDF` in [Section 3.8, “Key Derivation Function \(KDF\)”](#), the resulting Group Session ID data would be:

- Raw output of GroupKeyHash: `b9:f7`
- Group Session ID scalar value to be used for further processing: `0xB9F7` (47607 decimal)

The *Group Session ID* MAY help receiving nodes efficiently locate the *Operational Group Key* used to encrypt an incoming groupcast message. It SHALL NOT be used as the sole means to locate the associated *Operational Group Key*, since it MAY collide within the fabric. Instead, the *Group Session ID* provides receiving nodes a means to identify *Operational Group Key* candidates without the need to first attempt to decrypt groupcast messages using all available keys.

On receipt of a message of Group Session Type, all valid, installed, operational group key candidates referenced by the given Group Session ID SHALL be attempted until authentication is passed or there are no more operational group keys to try. This is done because the same Group Session ID might arise from different keys. The chance of a Group Session ID collision is 2^{-16} but the chance of both a Group Session ID collision and the message MIC matching two different operational group keys is 2^{-80} .

Group Session Ids are sized to fit within the context of the Session Identifier field of a message. When used in this context, the Group Session ID value allows a receiving Node to identify the appropriate message encryption key to use from the set of active operational keys it has currently

installed.

4.17.4. Distribution of Key Material

The operational group keys mechanism relies on a key distribution Administrator to reliably distribute select epoch key material to appropriate participants. It is assumed the key distribution Administrator is in possession of all epoch keys, has knowledge of the set of group security domain members which require access to those keys, and is responsible for pushing updates of these credentials to all authorized Nodes in those groups it manages.

Key material is distributed to key holders using the [Group Key Management Cluster](#). In general, the key material of a Node is exposed via Attributes with ACL entries that only allow access by the key distribution Administrator. The information exposed in the [Section 11.2, “Group Key Management Cluster”](#) includes the group epoch keys and associated group session identifiers.




Figure 19. Group Key Distribution

4.17.4.1. Installing a Group onto a Newly Commissioned Node

This section provides an example of the operations required to install a group onto a newly commissioned node.




Figure 20. Installing a group onto a new node

Sequence:

- *Admin:*
 - Generate fabric-unique group ID **GID** and random key **key0** for group key set ID **K**.
 - Write the group key set **K** to the remote node, GroupMember, using **KeySetWrite** command.
 - Associate group ID **GID** with key set **K** by writing an entry to the GroupKeyMap list attribute.
- *GroupMember:*
 - Node subscribes to the IPv6 multicast address generated from the fabric ID and group ID.
- *Admin:*
 - Associate endpoint with group ID **GID** by sending the **Groups** cluster's **AddGroup** command to endpoint.

4.18. Message Counter Synchronization Protocol (MCSP)

This section describes the protocol used to securely synchronize the message counter used by a

sender of messages encrypted with a symmetric group key.

Message counter synchronization is an essential part of enabling secure messaging between members of an operational group. Specifically, it protects against replay attacks, where an attacker replays older messages, which may result in unexpected behavior if accepted and processed by the receiver.

The recipient of a message encrypted with a group key can trust and process that message only if it has kept the last message counter received from a given sender using that key.

Underlying MCSP is a simple request / response protocol. When a multicast message with unknown counter is received, synchronization via MCSP begins by sending a synchronization request via unicast UDP to the multicast message originator's unicast IPv6 address. That originator then sends a synchronization response to the unsynchronized node via unicast UDP. After cryptographic verification, the formerly unsynchronized node is now synchronized with the originator's message counter and can trust the original and subsequent messages from the originator node.

4.18.1. Message Counter Synchronization Methods

There are two methods for synchronizing the message counter of a peer node, which are configurable per-group-key based on the [GroupKeySecurityPolicy](#) field of a given group key set (see [Section 11.2.5.4, "GroupKeySetStruct Type"](#)).

4.18.1.1. Trust-first policy

The first authenticated message counter from an unsynchronized peer is trusted, and its message counter is used to configure message-counter-based replay protection on future messages from that node. All control messages (any message with [C Flag](#) set) use the control message counter and SHALL use Trust-first for synchronization. Note that MCSP is not used for Trust-first synchronization.

This policy provides lower latency for less security-sensitive applications such as lighting.

WARNING

Trust-first synchronization is susceptible to accepting a replayed message after a Node has been rebooted.

4.18.1.2. Cache-and-sync policy

The message that triggers message counter synchronization is stored, a [message counter synchronization exchange](#) is initiated, and only when the synchronization is completed is the original message processed. Cache-and-sync provides replay protection even in the case where a Node has been rebooted, at the expense of higher latency.

Support for the cache-and-sync policy and [MCSP](#) is optional. A node indicates its ability to support this feature via the [Group Key Management](#) cluster [FeatureMap](#).

WARNING

Large groups may cause significant message synchronization burden on the sender and message queueing pressure or latency on the receiver when message counter synchronization state is out of sync, such as after a brown-out, on

first message to a new large group, or similar large scale sets of synchronizations being required at a single point in time.

4.18.2. Group Peer State

The *Group Peer State Table* stores information about every peer with which the node had a group message exchange. For every peer node id the following information is available in the table:

- Peer's Encrypted Group Data Message Counter Status:
 - *Synchronized Data Message Counter* - the largest encrypted data message counter received from the peer, if available.
 - Flag to indicate whether this counter value is valid and synchronized.
 - The *message reception state* bitmap tracking the recent window of data message counters received from the peer.
- Peer's Encrypted Group Control Message Counter Status:
 - *Synchronized Control Message Counter* - the largest encrypted control message counter received from the peer, if available.
 - Flag to indicate whether this counter value is valid and synchronized.
 - The *message reception state* bitmap tracking the recent window of control message counters received from the peer.

There are three scenarios where the receiver of an encrypted message does not know the sender's last message counter:

- The encrypted message is the first received from a peer.
- The device rebooted without persisting the *Group Peer State Table* content.

NOTE It is not required to persist the Peer State Table.

- The entry for the Peer in the *Group Peer State Table* was expunged due to the table being full.

There SHALL be at least 10 entries per supported fabric for Peer Encrypted Group data Message Status in the Group Peer State table. This number SHOULD NOT be less than 5 entries per fabric per instance of Groups cluster on an endpoint. The number of entries is a compromise between the number of peers that can send a group message to a receiver while the table is maintained, and the associated memory usage.

There SHALL be at least 2 entries per supported fabric for Peer Encrypted Group Control Message Status in the Group Peer State table.

The next sections describe the functional protocol used to request message counter synchronization with a peer and form responses to such message counter synchronization requests.

4.18.3. MCSP Messages

4.18.3.1. **MsgCounterSyncReq** - Message Counter Synchronization Request

The Message Counter Synchronization Request (*MsgCounterSyncReq*) message is sent when a message was received from a peer whose current message counter is unknown.

A *MsgCounterSyncReq* message SHALL set the **C Flag** in the message header. The control message counter SHALL be used for message protection.

A *MsgCounterSyncReq* message SHALL be secured with the group key for which counter synchronization is requested and SHALL set the **Session Type** to **1**, indicating a group session as per the rules outline in [Section 4.18.5, “Message Counter Synchronization Exchange”](#).

The payload of the *MsgCounterSyncReq* message takes the format defined in [Table 26, “Message Counter Sync Request”](#):

Table 26. Message Counter Sync Request

Field Size	Message Field	Description
8 bytes	Challenge	The Challenge is a 64-bit random number generated using the DRBG by the initiator of a <i>MsgCounterSyncReq</i> to uniquely identify the synchronization request cryptographically.

4.18.3.2. **MsgCounterSyncRsp** - Message Counter Synchronization Response

The Message Counter Synchronization Response (*MsgCounterSyncRsp*) message is sent in response to a *MsgCounterSyncReq*.

A *MsgCounterSyncRsp* message SHALL set the **C Flag** in the message header. The control message counter SHALL be used for message protection.

The *MsgCounterSyncRsp* message has the format defined in [Table 27, “Message Counter Sync Response”](#):

Table 27. Message Counter Sync Response

Field Size	Message Field	Description
4 bytes	Synchronized Counter	The current data message counter for the node sending the <i>MsgCounterSyncRsp</i> message.
8 bytes	Response	The Response SHALL be the same as the 64-bit value sent in the Challenge field of the corresponding <i>MsgCounterSyncReq</i> .

4.18.4. Unsynchronized Message Processing

An unsynchronized message is one that is cryptographically verified from a node whose message counter is unknown. Upon receipt of an unsynchronized message process the message as follows:

1. The message SHALL be of **Group Session Type**, otherwise discard the message.
2. If **C Flag** is set to **1**:
 - a. Create a **Message Reception State** and set its **max_message_counter** to the message counter of the given message, i.e. **trust-first**.
 - b. Accept the message for further processing.
3. If **C Flag** is set to **0**:
 - a. Determine the **GroupKeySecurityPolicy** (as set by the **Section 11.2, “Group Key Management Cluster”**) of the operational group key used to authenticate the message.
 - b. If the key has a *trust-first* security policy, the receiver SHALL:
 - i. Set the peer’s group key data message counter to *Message Counter* of the message.
 - A. Clear the **Message Reception State** bitmap for the group session from the peer.
 - B. Mark the peer 's group key data message counter as synchronized.
 - ii. Process the message.
 - c. If the key has a *cache-and-sync* security policy, the receiver SHALL:
 - i. If MCSP is not in progress for the given peer Node ID and group key:
 - A. Store the message for later processing.
 - B. Proceed to **Section 4.18.5, “Message Counter Synchronization Exchange”**.
 - ii. Otherwise, do not process the message.
 - A. An implementation MAY queue the message for later processing after MCSP completes if resources allow.

For each peer Node ID and group key pair there SHALL be at most one synchronization exchange outstanding at a time.

4.18.5. Message Counter Synchronization Exchange

A message synchronization exchange starts by sending the **MsgCounterSyncReq** message to the peer Node that sent the message with unknown message counter. When a synchronization request is triggered by an incoming multicast message, the Node SHALL first wait for a uniformly **random** amount of time between **0** and **MSG_COUNTER_SYNC_REQ_JITTER**.

The sender of the *MsgCounterSyncReq* message SHALL:

1. Set the Message Header fields as follows:
 - a. The **S Flag** SHALL be set to **1**.
 - b. The **DSIZ** field SHALL be set to **1**.
 - c. The **P Flag** SHALL be set to **1**.
 - d. The **C Flag** SHALL be set to **1**.
 - e. The **Session Type** field SHALL be set to **1**.
 - f. The **Session ID** field SHALL be set to the **Group Session Id** for the operational group key

being synchronized.

- g. The **Source Node ID** field SHALL be set to the Node ID of the sender Node.
- h. The **Destination Node ID** field SHALL be set to the **Source Node ID** of the message that triggered the synchronization attempt.
2. Create a new synchronization **Exchange**.
 - a. The **Exchange ID** of the message SHALL be set to match the new Exchange.
 - b. The **I Flag** SHALL be set to **1**.
 - c. The **A Flag** SHALL be set to **0**.
 - d. The **R Flag** SHALL be set to **1**.
3. Set the *Challenge* field to the value returned by **Crypto_DRBG(len = 8 * 8)** and store that value to resolve synchronization responses from the destination peer.
4. Arm a timer to enforce that a synchronization response is received before **MSG_COUNTER_SYNC_TIMEOUT**.
 - a. Upon firing of the timer:
 - i. The synchronization exchange SHALL be closed.
 - ii. Any message waiting on synchronization associated with the exchange SHALL be discarded.
 - b. The timer SHALL be stopped upon receipt of an authenticated *MsgCounterSyncRsp* message that matches:
 - i. The **Source Node ID** field matches the **Destination Node ID** field of the most recent *MsgCounterSyncReq* message generated for that Node.
 - ii. The *Response* field corresponds to the *Challenge* field of the *MsgCounterSyncReq* message.
5. Invoke **Section 4.7.1, “Message Transmission”** using parameters from step 1 to encrypt and then send the request message over UDP to the IPv6 unicast address of the destination.
 - a. The request message SHALL use the same operational group key as the message which triggered synchronization.
 - b. The group key SHALL be known/derivable by both parties (sender and receiver).

The receiver of *MsgCounterSyncReq* SHALL:

1. Verify the **Destination Node ID** field SHALL match the Node ID of the receiver, otherwise discard the message.
2. Respond with *MsgCounterSyncRsp*.

The sender of the *MsgCounterSyncRsp* response message SHALL:

1. Set the Message Header fields as follows:
 - a. The **S Flag** SHALL be set to **1**.
 - b. The **DSIZ** field SHALL be set to **1**.

- c. The **P Flag** SHALL be set to **1**.
 - d. The **C Flag** SHALL be set to **1**.
 - e. The **Session Type** field SHALL be set to **1**.
 - f. The **Session ID** field SHALL be set to the **Group Session Id** for the operational group key being synchronized.
 - g. The **Source Node ID** field SHALL be set to the Node ID of the sender Node.
 - h. The **Destination Node ID** field SHALL be set to the **Source Node ID** of the *MsgCounterSyncReq*.
2. Set the **MsgCounterSyncRsp** payload fields as follows:
 - a. The *Response* field SHALL be set to the value of the *Challenge* field from the *MsgCounterSyncReq*.
 - b. The *Synchronized Counter* field SHALL be set to the current **Global Group Encrypted Data Message Counter** of the sender.
 3. Use the same **exchange context** as the *MsgCounterSyncReq* being responded to.
 - a. The **Exchange ID** of the message SHALL be set to the *Exchange ID* of the *MsgCounterSyncReq*.
 - b. The **I Flag** SHALL be set to **0**.
 - c. The **A Flag** SHALL be set to **1**.
 - d. The **R Flag** SHALL be set to **1**.
 4. Invoke **Section 4.7.1, “Message Transmission”** using parameters from step 1 to encrypt and then send the response message over UDP to the IPv6 unicast address of the destination.

The receiver of the *MsgCounterSyncRsp* message SHALL:

1. Verify the *MsgCounterSyncRsp* matches a previously sent *MsgCounterSyncReq*:
 - a. An active synchronization exchange SHALL exist with the source node.
 - b. The *Exchange ID* field SHALL match the *Exchange ID* used for the original *MsgCounterSyncReq* message.
 - c. The *Response* field SHALL match the *Challenge* field used for the original *MsgCounterSyncReq* message.
 - d. The **Destination Node ID** field SHALL match the **Source Node ID** of the original *MsgCounterSyncReq* message.
 - e. The **Source Node ID** field SHALL match the **Destination Node ID** of the original *MsgCounterSyncReq* message.
2. On verification failure:
 - a. Silently ignore the *MsgCounterSyncRsp* message.
3. On verification success:
 - a. Set the peer's group key data message counter to *Synchronized Counter*.
 - b. Clear the **Section 4.6.5.1, “Message Reception State”** bitmap for the peer.
 - c. Mark the peer's group key data message counter as synchronized.

- d. Resume processing of any queued message that triggered synchronization according to [Section 4.6.7, “Counter Processing of Incoming Messages”](#).
 - i. If more than one message is queued from the synchronized peer, using the same operational group key, the messages SHALL be processed in the order received.
- e. Close the synchronization exchange created for the original *MsgCounterSyncReq* message.

4.18.6. Message Counter Synchronization Session Context

While conducting Message Counter Synchronization with a peer, nodes SHALL maintain the following session context. For nodes initiating message counter synchronization, this context SHALL be maintained throughout the full exchange of *MsgCounterSyncReq* and *MsgCounterSyncRsp* messages. For nodes responding to *MsgCounterSyncReq* messages, the context SHALL only be maintained long enough to generate and successfully transmit the *MsgCounterSyncRsp* message.

1. **Fabric Index**: Records the [Index](#) for the Fabric within which nodes are conducting message counter synchronization.
 - Fabric Index is derived by identification of an Operational Group Key associated with the fabric through successful decryption with that key and verification of the [Message Integrity Check](#). For nodes initiating counter synchronization, this occurs at decryption of an inbound groupcast message. For nodes in the responder role, this occurs at decryption of an inbound *MsgCounterSyncReq* message.
2. **Peer Node ID**: Records the node ID of the peer with which message counter synchronization is being conducted.
 - For nodes initiating message counter synchronization, this is the node ID of the responder. For nodes responding to message counter synchronization, this is the node ID of the initiator.
3. **Role**: Records whether the node is the initiator of or responder to message counter synchronization.
 - Together, *Fabric Index*, *Peer Node ID* and *Role* comprise a unique key that can be used to match incoming messages to ongoing MCSP exchanges.
4. **Message Reception State**: Provides tracking for the [Control Message Counter](#) of the remote peer.
5. **Peer IP Address**: The unicast IP address of the peer.
6. **Peer Port**: The receiving port of the peer.
7. **Operational Group Key**: The [Operational Group Key](#) that is being used to encrypt messages within the counter synchronization exchange.
8. **Group Session ID**: Records the [Group Session ID](#) derived from the [Operational Group Key](#) that is being used to encrypt messages within the counter synchronization exchange.

4.18.7. Sequence Diagram

The following sequence diagram shows an example of how message counter synchronization behaves in the most common scenario.

4.18.7.1. Scenario 1 — Multicast Receiver Initiated

Assumptions:

- *Sender* transmits a multicast message.
- *Receiver* does not know *Sender* 's message counter.




Figure 21. Multicast Receiver Initiated Message Counter Synchronization

Sequence:

- *Sender*:
 - Generates, encrypts and sends *Msg1* as a multicast message. *Msg1* could be:
 - Regular message that starts encrypted group communication with receivers *Receiver1* and *Receiver2*.
- Receivers *Receiver1* and *Receiver2* each:
 - Receive, decrypt, authenticate, and cache *Msg1* message for later processing.
 - Generate, encrypt, and send *MsgCounterSyncReq* message.
- *Sender*:

- Receives *MsgCounterSyncReq* message.
- Generates, encrypts and sends *MsgCounterSyncRsp* message.
- R1 and R2 each:
 - Receive, decrypt, process, and verify *MsgCounterSyncRsp* message from *Sender*.
 - On verification success: marks *Sender* 's group key message counter as synchronized.
 - Processes cached *Msg1* message.

4.19. Bluetooth Transport Protocol (BTP)

The Bluetooth Transport Protocol (BTP) provides a TCP-like layer of reliable, connection-oriented data transfer on top of GATT. BTP splits individual Service Data Unit (SDU) messages into multiple BTP segments, which are each transmitted via a single GATT write or indication (as shown in [Figure 22, “MATTERoBLE: Matter Message / BTP layering”](#)).

While BTP is a generic protocol, Matter specifically uses BTP to define a Matter-over-Bluetooth Low Energy (MATTERoBLE) Interface. A MATTERoBLE Interface SHALL implement BTP as a universally compatible transport mode. A MATTERoBLE Interface SHALL only be used to transport Matter messages as the BTP SDU.




Figure 22. MATTERoBLE: Matter Message / BTP layering

The BTP session handshake allows devices to check BTP protocol version compatibility and exchange other data before a BTP session is established. Once established, this session is used to send and receive BTP SDUs (such as Matter messages) as BTP Message Segments. A BTP session MAY open and close with no effect on the state of the underlying Bluetooth LE connection, except in the case where a BTP session is closed by the Peripheral Device. Either the Peripheral or the Central MAY signal the end of a BTP session by closing the underlying BLE connection.

Due chiefly to constraints put on design by resource-limited BLE chipsets, BTP defines a receive window for each side of a session in units of GATT PDUs. Each GATT Write Characteristic Value (ATT_WRITE_REQ) PDU or Indication (ATT_HANDLE_VALUE_IND) PDU is sent with a sequence number which the receiver uses to acknowledge receipt of each packet at the BTP layer and open its receive window from the sender’s perspective.

4.19.1. BTP Session Interface

Conceptually, an open BTP session is exposed to the next-higher session layer as a full-duplex message stream.

4.19.2. BTP Frame Format

A BTP Frame consists of an 8-bit header followed by one or more optional fields, as detailed below. BTP uses little endian encoding for any header fields larger than one byte in length.

Table 28, “BTP Packet PDU format” defines the BTP Packet PDU format.

Unused fields SHALL be set to '0'.

Table 28. BTP Packet PDU format

bit 0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Control Flags								[Management Opcode]							
[Ack Number]								[Sequence Number]							
[Message Length]															
[Segment Payload]...															
...															

4.19.2.1. Control Flags

Table 29. BTP Control Flags

bit 7	6	5	4	3	2	1	0
-	H	M	-	A	E	C	B

H (Handshake) bit

Set to '0' for normal BTP packets. When set, this bit indicates a BTP handshake packet for session establishment and has a different packet format described below.

M (Management Message) bit

Indicates the presence ('1') or absence ('0') of the Management Opcode field. All segments of a message SHALL set this bit to the same value.

A (Acknowledgement) bit

Indicates the presence of the Ack Number field.

E (Ending Segment) bit

Set to '1' on the last segment of a BTP SDU and set to '0' for all other segments of the same BTP SDU. A segment MAY have both the Beginning and Ending bits set indicating that a full BTP SDU is included in the message. When set, the segment payload length is equal to the total remaining unre-

ceived message data. When not set, the segment payload length is equal to the maximum allowable BTP session packet size minus header overhead.

C (Continuing Segment) bit

Set to '0' on the first segment of a BTP SDU and set to '1' for all remaining segments of the same BTP SDU.

B (Beginning Segment) bit

Set to '1' on the first segment of a BTP SDU and set to '0' for all remaining segments of the same BTP SDU. It indicates the presence of the Message Length field.

4.19.2.2. Ack Number

Optional field specified in [Section 4.19.4.8, “Packet Acknowledgements”](#).

4.19.2.3. Sequence Number

Mandatory field for regular data messages specified in [Section 4.19.4.6, “Sequence Numbers”](#).

4.19.2.4. Message Length

Optional field present in Beginning Segment only. Value indicates the length in bytes of the full message buffer to be transmitted. None of the BTP Packet PDU fields is included in the Message Length.

4.19.2.5. Segment Payload

Optional field containing a segment of the Service Data Unit (SDU) message in transmission to the receiver.

4.19.3. BTP Control Frames

BTP defines different control frame formats depending on the Management Opcode that is in the BTP Packet PDU header. Valid Management Opcodes for BTP Control Frames are defined in [Table 30, “BTP Control codes”](#).

Table 30. BTP Control codes

Management Opcode	Name	Description
0x6C	Handshake	Request and response for BTP session establishment

4.19.3.1. BTP Handshake Request

Table 31. BTP Handshake Request format

bit 0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Control Flags = 0x65								Management Opcode = 0x6C							
Ver[0]				Ver[1]				Ver[2]				Ver[3]			

bit 0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Ver[4]				Ver[5]				Ver[6]				Ver[7]			
Requested ATT_MTU															
Client Window Size															

H (Handshake) bit

Set to '1' for connection handshake messages.

M (Management) bit

Set to '1' for connection handshake messages.

Ver[i] (Version) nibbles

Used to negotiate the highest version capability between a Device pair. Supported versions are listed once each, newest first, in descending order. Unused version fields are filled with '0'.

The following values are defined:

- 0 — Unused field
- 4 — BTP as defined by Matter v1.0

Requested ATT_MTU

Requested ATT_MTU is a 16-bit unsigned integer field containing the size of the GATT PDU (ATT_MTU) that can be received by the sender minus the size of the GATT header (3). This value is obtained via the standard ATT MTU exchange procedure (see [Bluetooth® Core Specification 4.2 \(amended\)](#) Vol 3, Part F, Section 3.4.2 "MTU Exchange") and is used to validate that both sides of the BLE connection are using the common minimum value. If BTP is not aware of the negotiated GATT MTU, the value SHALL be set to '23', indicating the minimum ATT_MTU defined by GATT. If the client has no preference, the value may be set to '0'.

NOTE

Each GATT PDU used by the BTP protocol introduces 3 byte header overhead making the maximum BTP Segment Size for the session equal to negotiated ATT_MTU-3.

Client Window Size

Value of the maximum receive window size supported by the client, specified in units of BTP packets where each packet may be up to 244 bytes in length. This maximum was chosen so a single BTP segment can fit into a single 255 byte BLE link layer PDU, including all headers from the link layer, L2CAP, GATT, and BTP.

4.19.3.2. BTP Handshake Response

Table 32. BTP Handshake Response format

bit 0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Control Flags = 0x65								Management Opcode = 0x6C							

bit 0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Final Protocol Version				Reserved				Selected ATT_MTU (low byte)...							
...Selected ATT_MTU (high byte)								Selected Window Size							

H (Handshake) bit

Set to '1' for session handshake messages.

M (Management) bit

Set to '1' for session handshake messages.

Final Protocol Version

Value of the BTP protocol version selected by the server.

Reserved

Must be set to '0'.

Selected ATT_MTU

Value of the maximum ATT_MTU for the connection selected by the server as a 16-bit unsigned integer.

Selected Window Size

Value of the maximum receive window size supported by the server, specified in units of BTP packets where each packet may be up to the selected segment size in length.

4.19.4. BTP GATT Service

4.19.4.1. BTP Channelization

Bluetooth Transport Protocol provides a packetized stream interface over GATT but says nothing about the actual contents of the data packets it transports. The BTP Service UUID is used to specify the actual contents of the packets (see [Table 33, “BTP Service UUID”](#)).

Table 33. BTP Service UUID

BTP Datagram Contents	BTP Service UUID
Matter Message frames	MATTER_BLE_SERVICE_UUID

NOTE

See [Section 4.19.6, “Bluetooth Requirements”](#) for terms of use.

While a single BTP connection may exist between a BTP Client and BTP Server, multiple BTP sessions may be established between various peers.

4.19.4.2. BTP GATT Service

The BTP GATT service is composed of a service with three characteristics—C1, C2 and C3 (see [Table 34, “BTP GATT service”](#)). The client SHALL exclusively use C1 to initiate BTP sessions by sending BTP handshake requests and send data to the server via GATT ATT_WRITE_REQ PDUs. While a client is subscribed to allow indications over C2, the server SHALL exclusively use C2 to respond to BTP handshake requests and send data to the client via GATT ATT_HANDLE_VALUE_IND PDUs.

Table 34. BTP GATT service

Attribute	Description
BTP Service	UUID = MATTER_BLE_SERVICE_UUID
C1 Characteristic (Client TX Buffer)	UUID = 18EE2EF5-263D-4559-959F-4F9C429F9D11 Characteristic Properties = Write max length = 244 bytes
C2 Characteristic (Client RX Buffer)	UUID = 18EE2EF5-263D-4559-959F-4F9C429F9D12 Characteristic Properties = Indication max length = 244 bytes
C3 Characteristic (Additional commissioning-related data)	UUID = 64630238-8772-45F2-B87D-748A83218F04 Characteristic Properties = Read max length = 512 bytes

For all messages sent from the BTP Client to BTP Server, BTP SHALL use the GATT Write Characteristic Value sub-procedure. For all messages sent from the BTP Server to BTP Client, BTP SHALL use the GATT Indications sub-procedure.

The values of C1 and C2 SHALL both be limited to a maximum length of 244 bytes. This constraint is imposed to align with maximum PDU size when LE Data Packet Length Extensions (DPLE) is enabled on Bluetooth 4.2 hardware. Per [Bluetooth® Core Specification 4.2 \(amended\)](#) Vol 3, Part F, Section 3.2.9 "Long Attribute Values", the maximum characteristic value length is 512 bytes. The maximum lengths of C1 and C2 may increase in a future version of the BTP specification to allow higher throughput on BLE connections whose ATT_MTU exceeds 247 bytes.

C3 is an optional characteristic that the server MAY use to include additional data required during the provisioning as defined in [Section 5.4.2.5.7, “GATT-based Additional Data”](#).

BTP Clients SHALL perform certain GATT operations synchronously, including GATT discovery, subscribe, and unsubscribe operations. GATT discovery, subscribe, or unsubscribe SHALL NOT be initiated while the result of a previous operation remains outstanding. This requirement is imposed by GATT protocol.

4.19.4.3. Session Establishment

Before a BTP session can be initiated, a central SHALL first establish a BLE connection to a peripheral. Once a BLE connection has been formed, centrals SHALL assume the GATT client role for BTP session establishment and data transfer, and peripherals SHALL assume the GATT server role. If peripheral supports LE Data Packet Length Extension (DPLE) feature it SHOULD perform data length update procedure before establishing a GATT connection.

Before establishing a BTP session, the GATT client SHOULD perform a GATT Exchange MTU procedure.

After that the BTP Client SHALL execute the GATT discovery procedure. The GATT discovery procedure starts with primary service discovery using either the GATT Discover All Primary Services sub-procedure or the GATT Discover Primary Services by Service UUID sub-procedure.

The BTP Client SHALL perform either the GATT Discover All Characteristics of a Service sub-procedure or the GATT Discover Characteristics by UUID sub-procedure in order to discover the C1 and C2 characteristics.

The BTP Client SHALL perform the GATT Discover All Characteristic Descriptors sub-procedure in order to discover the Client Characteristic Configuration descriptor (CCCD) of C2 characteristic.

To initiate a BTP session, a BTP Client SHALL send a BTP handshake request packet to the BTP Server via a ATT_WRITE_REQ PDU on characteristic C1 of the BTP Service. The handshake request packet SHALL include, a list of BTP protocol versions supported by the client, the client's GATT ATT_MTU, and the client's maximum receive window size. The list of supported protocol versions SHALL be sorted in descending numerical order. If the client cannot determine the BLE connection's ATT_MTU, it SHALL specify a value of '23' (the minimum ATT_MTU supported by GATT) for this field in the handshake request. For a detailed specification of the handshake request binary data format, see [Section 4.19.3.1, "BTP Handshake Request"](#).

After the BTP Client acknowledges delivery of the handshake request packet, upon receipt of a GATT Write Response, the BTP Client SHALL enable indications over C2 characteristics by writing value 0x01 to C2's Client Characteristic Configuration descriptor as described in [Bluetooth® Core Specification 4.2 \(amended\)](#) Vol 3, Part C, Section 10.3.1.1 "Handling of GATT Indications and Notifications".

Once the GATT server has received a client's BTP handshake request and confirmed the client's subscription to C2, it SHALL send a BTP handshake response to the client via a ATT_HANDLE_VALUE_IND PDU on C2. This response SHALL contain the window size, maximum BTP packet size, and BTP protocol version selected by the server. For a detailed specification of the handshake response binary data format, see [Section 4.19.3.2, "BTP Handshake Response"](#).

The server SHALL select a window size equal to the minimum of its and the client's maximum window sizes. Likewise, the server SHALL select a maximum BTP Segment Size for the BLE connection by taking the minimum of 244 bytes (the maximum characteristic value length of C1 and C2), server's ATT_MTU-3 and ATT_MTU-3 as declared by the client.

The server SHALL select a BTP protocol version that is the newest which it and the client both support, where newer protocol version numbers are strictly larger than those of older versions. The version number returned in the handshake response SHALL determine the version of the BTP protocol used by client and server for the duration of the session.

If the server determines that it and the client do not share a supported BTP protocol version, the server SHALL close its BLE connection to the client. When a client sends a handshake request, it SHALL start a timer with a globally-defined duration of BTP_CONN_RSP_TIMEOUT seconds. If this timer expires before the client receives a handshake response from the server, the client SHALL close the BTP session and report an error to the application. Likewise, a server SHALL start a timer

with the same duration when it receives a handshake request from a client. If this timer expires before the server receives a subscription request on C2, the server SHALL close the BTP session and report an error to the application. The state machine in [Figure 23, “BTP session handshake”](#) illustrates the function of these timers as part of the BTP session establishment procedure.




Figure 23. BTP session handshake

4.19.4.4. Data Transmission

Once a BTP session has been established, the next-higher-layer application on both peers may use this BLE connection to send and receive data via GATT writes or indications, according to a peer's GATT role. Clients SHALL exclusively use GATT Write Characteristic Value sub-procedure to send data to servers and servers SHALL exclusively use GATT Indication sub-procedure to send data to clients.

All BTP packets sent on an open BLE connection SHALL adhere to the BTP Packet PDU binary data format specified in BTP Frame Formats. All BTP packets SHALL include a header flags byte and an 8-bit unsigned sequence number. All other packet fields are optional. These optional fields include an 8-bit unsigned received packet acknowledgement number, 16-bit unsigned buffer length indication, and variable-length buffer segment payload.

This section is defined entirely within the scope of a single BTP session. Concurrent BTP sessions between the same peer and multiple remote hosts SHALL maintain separate and independent acknowledgement timers, sequence numbers, and receive windows.

4.19.4.5. Message Segmentation and Reassembly

When the session layer (that is, MATTERoBLE) sends a Matter Message as a BTP SDU over a BTP session, that BTP SDU SHALL be split into ordered, non-overlapping BTP segments so the set of all BTP segments may be reassembled into the original BTP SDU (see [Figure 22, “MATTERoBLE: Matter Mes-](#)

sage / BTP layering”). Each BTP segment SHALL be prepended with a BTP packet header and sent as the payload of a single GATT PDU. If a BTP SDU is split into more than one BTP segment, the BTP segments SHALL be sent in order of their position in the original BTP SDU, starting with the BTP segment at the buffer’s head.

At any point in time, only one BTP SDU may be transmitted in each direction over a BTP session. The transmission of BTP segments of any two BTP SDUs SHALL NOT overlap. If the application attempts to send one BTP SDU while transmission of another BTP SDU is in progress, the new BTP SDU SHALL be appended to a first-in, first-out queue. The next BTP SDU SHALL be dequeued from this queue and transmitted once transmission of the current BTP SDU completes, that is, once all BTP segments of the current BTP SDU have been transmitted and received by the peer via GATT.

The first BTP segment of a BTP SDU sent over a BTP session SHALL have its Beginning Segment header flag set to indicate the beginning of a new BTP SDU (see [Table 28, “BTP Packet PDU format”](#)). The presence of this flag SHALL indicate the further presence of a 16-bit unsigned integer field, the Message Length, that provides the receiver with the total length of the BTP SDU. The last BTP segment for a given BTP SDU SHALL have its Ending Segment flag set to indicate the end of the transmitted BTP SDU. A BTP packet that bears an unsegmented BTP SDU—that is, a BTP SDU small enough to fit into a single BTP segment—SHALL have both its Beginning Segment and Ending Segment flags set.

The size of a single BTP SDU sent via BTP is limited to 64KB, that is, the maximum size of the Message Length field in the BTP packet header. The number of segments used to send a buffer is unlimited and delimited by the Beginning Segment and Ending Segment bits in the BTP packet header. The upper layer imposes more stringent requirements over the maximum SDU size, such as [Section 4.4.4, “Message Size Requirements”](#).

The length of the data payload in each BTP segment whose Ending Segment bit is not set SHALL be equal to the session’s maximum BTP packet size minus the size of that packet’s header. If a packet’s Ending Segment bit is set, the length of its BTP segment data payload SHALL equal the size of the original BTP SDU minus the total size of all previously transmitted BTP segments of that BTP SDU. In this way, the length of a SDU’s last BTP segment is implied by its size.

Once a peer receives a complete set of BTP segments, it SHALL reassemble them in the order received, and verify that the reassembled BTP SDU’s total length matches that specified by the Beginning Segment’s Message Length value. If they match, the receiver SHALL pass the reassembled BTP SDU up to the next-higher-layer. If the reassembled buffer’s length does not match that specified by the sender, or if received BTP segment payload size would exceed the maximum BTP packet size, or receiver receives an Ending Segment without the presence of a previous Beginning Segment, or a Beginning Segment when another BTP SDU’s transmission is already in progress, the receiver SHALL close the BTP session and report an error to the application.

4.19.4.6. Sequence Numbers

All BTP packets SHALL be sent with sequence numbers, regardless of whether they contain SDU segments (for example, a packet acknowledgement with no attached segment payload). The purpose of sequence numbers is to facilitate the BTP receive window. A BTP sequence number SHALL be defined as an unsigned 8-bit integer value that monotonically increments by 1 with each packet sent by a given peer. A sequence number incremented past 255 SHALL wrap to zero.

Sequence numbers SHALL be separately defined for either direction of a BTP session. The sequence number of the first packet sent by the client after completion of the BTP session handshake SHALL be zero. The server's BTP handshake response bears an implied sequence number of zero because it occupies a slot in the client's receive window. The client acknowledges the server's BTP handshake response with an acknowledgement sequence of zero. For this reason, the sequence number of the first data packet sent by the server after completion of the BTP session handshake SHALL be 1.

Peers SHALL check to ensure that all received BTP packets properly increment the sender's previous sequence number by 1. If this check fails, the peer SHALL close the BTP session and report an error to the application.

4.19.4.7. Receive Windows

The purpose of the receive window is to enable flow control at the GATT session layer between BTP peers.

Flow control is required at the GATT transport layer for embedded platforms that use "minimal" BLE chipsets. These platforms may have limited space on the host processor to receive packets from their BLE chipsets. In the case of some dual-chip architectures, writes and indications are received and confirmed by the BLE chip with no input from the host processor. When the BLE chip sends the result of a received GATT PDU to the host processor, that payload and the corresponding BTP packet will be permanently lost if the host does not have enough space to receive it. For this reason, knowledge of a remote host's ability to reliably receive GATT PDUs is presented at the transport layer in the form of the BTP receive window.

Both peers in a BTP session SHALL define a receive window, where the window's size indicates the number of GATT PDUs (that is, BTP segments) a peer can reliably receive and store without session-layer acknowledgment. A maximum window size SHALL be established for both peers as part of the BTP session handshake. To prevent sequence number wrap-around, the largest maximum window size any peer may support is 255.

Both peers SHALL maintain a counter to reflect the current size of the remote peer's receive window. Each peer SHALL decrement this counter when it sends a packet via GATT write or indication and increment this counter when a sent packet is acknowledged.

If a local peer's counter for a remote peer's receive window is zero, the window SHALL be considered closed, and the local peer SHALL NOT send packets until the window reopens (is incremented above zero). When a closed window reopens, a local peer SHALL immediately resume any pending BTP packet transmission.

A local peer SHALL also not send packets if the remote peer's receive window has one slot open and the local peer does not have a pending packet acknowledgement. This is to avoid the situation where the receive windows of both peers are full and neither can send an acknowledgement to reopen its window for the other. Because the server's handshake response bears an implicit BTP sequence number of zero, a server SHALL initialize its counter for the client's receive window size at (negotiated maximum window size - 1). A client SHALL initialize its counter for the server's receive window at the negotiated maximum window size.

Both peers SHALL also keep a counter of their own receive window size based on the sequence

number difference between the last packet they received and the last packet they acknowledged. This counter is used to proactively send early packet acknowledgements when a peer's own receive window is about to close. See [Section 4.19.4.8, “Packet Acknowledgements”](#) for details.

An example scenario involving BTP receive windows is depicted in [Figure 24, “Example receive window scenario”](#), complete with packet acknowledgements as specified in [Section 4.19.4.8, “Packet Acknowledgements”](#). In this scenario, the client transmits a three-segment buffer to the server once it receives the server's handshake request. The handshake request occupies one slot in the client's initial receive window. The server's initial receive window is empty. Both client and server have a maximum window size of 4.




Figure 24. Example receive window scenario

4.19.4.8. Packet Acknowledgements

The purpose of sequence numbers and packet receipt acknowledgements is to support the BTP receive window and provide a keep-alive signal when a session is idle to affirm the health and continued operation of a remote BTP stack.

Per BTP Frame Formats, BTP packet receipt acknowledgements SHALL be received as unsigned 8-bit integer values in the header of a BTP packet. The value of this field SHALL indicate the sequence number of the acknowledged packet.

Acknowledgement of a sequence number indicates acknowledgement of the previous sequence number, if it too is unacknowledged. By induction, acknowledgement of a given packet implies acknowledgement of all packets received on the same BTP session prior to the acknowledged packet.

An acknowledgement is invalid if the acknowledged sequence number does not correspond to an outstanding, unacknowledged BTP packet sequence number. In contrast to TCP, BTP acks are not "free." A stand-alone ack—that is, a BTP packet that contains a packet receipt acknowledgement value but no buffer segment payload—consumes a slot in a remote peer's window just like any other packet. Stand-alone acknowledgement packets SHALL be acknowledged by a remote peer. The implications of this are examined in [Section 4.19.4.9, "Idle Connection State"](#).

Each peer SHALL maintain an acknowledgement-received timer. When a peer sends any BTP packet, it SHALL start this timer if it is not already running. The timer's duration SHALL be globally defined as BTP_ACK_TIMEOUT seconds, referred to as the acknowledgement timeout interval.

A peer SHALL restart its acknowledgement-received timer when a valid acknowledgement is received for any but its most recently sent unacknowledged packet. A peer SHALL stop its acknowledgement-received timer if it receives an acknowledgement for its most recently sent unacknowledged packet. If a peer's acknowledgement-received timer expires, or if a peer receives an invalid acknowledgement, the peer SHALL close the BTP session and report an error to the application.

Because the server's handshake response bears an implicit BTP sequence number of zero, a server SHALL start its acknowledgement-received timer when it sends a handshake response.

Each peer SHALL also maintain a send-acknowledgement timer. When it receives any BTP packet, a peer SHALL record the packet's sequence number as the corresponding BTP session's pending acknowledgement value and start the send-acknowledgement timer if it is not already running. The timer's duration SHALL be defined as any value less than one-half the acknowledgement timeout interval. This ensures that on a healthy BLE connection, a peer will always receive acknowledgements for sent packets before its acknowledgement-received timer expires.

A peer SHALL stop its send-acknowledgement timer when any pending acknowledgement is sent, either as a stand-alone BTP packet or piggybacked onto an outgoing buffer segment. If this timer expires and the peer has a pending acknowledgement, the peer SHALL immediately send that acknowledgement. If the peer sends any packet before this timer expires, it SHALL piggyback any pending acknowledgement on the transmitted packet and stop the send-acknowledgement timer.

Because the server's handshake response bears an implicit BTP sequence number of zero, a client SHALL set its pending acknowledgement value to zero and start its send-acknowledgement timer when it receives the server's a handshake response. Operation of the send-acknowledgement and acknowledgement-received timers is illustrated in [Figure 26, "BTP session lifecycle for Central acting as GATT Client"](#) in [Section 4.19.4.11, "Protocol State Diagrams"](#).

If a peer detects that its receive window has shrunk to two or fewer free slots, it SHALL immediately send any pending acknowledgement as a stand-alone BTP packet. This prevents the session from stalling in the interval between when a peer's receive window becomes empty and when its send-acknowledgement timer would normally fire.

4.19.4.9. Idle Connection State

When neither side of a BTP session has data to send, BTP packets will still be exchanged every send-acknowledgement interval due to acknowledgements generated by the receipt of previous data or stand-alone acknowledgement packets, as discussed in [Section 4.19.4.8, "Packet Acknowledgements"](#). The behavior of the acknowledgement-received timer in this scenario doubles as a keep-

alive mechanism, as it will cause a peer to close a BLE connection automatically if the remote BTP stack crashes or becomes unresponsive. This scenario is illustrated in [Figure 25, “Idle connection scenario”](#).




Figure 25. Idle connection scenario

4.19.4.10. Connection Shutdown

To close a BTP session, a GATT client SHALL unsubscribe from characteristic C2. The GATT server SHALL take this action to indicate closure of any BTP session open to the client.

If a BTP Server needs to close the BTP session, it SHALL terminate its BLE connection to the client.

4.19.4.11. Protocol State Diagrams

[Figure 26, “BTP session lifecycle for Central acting as GATT Client”](#) shows the state machine for BTP session management of a BTP Client Device.




Figure 26. BTP session lifecycle for Central acting as GATT Client

Figure 27, “BTP session lifecycle for Peripheral acting as GATT Server” shows the state machine for BTP session management of a BTP Server Device.




Figure 27. BTP session lifecycle for Peripheral acting as GATT Server

Note that in Figure 27, “BTP session lifecycle for Peripheral acting as GATT Server”, the state machine is identical for GATT clients and servers with the distinction that clients send data to servers via confirmed writes, and servers send data to clients via indications.

Figure 28, “State diagram for BTP session post-establishment” shows the state machine for BTP session maintenance at the protocol level, including liveness enforcement through keep alive messages and automatic teardown if acknowledgements are received before the timeout.




Figure 28. State diagram for BTP session post-establishment

4.19.5. Parameters and Constants

Table 35, “Glossary of constants” is a glossary of constants used in this chapter, along with a brief description and the default for each constant.

Table 35. Glossary of constants

Constant Name	Description	Default
BTP_CONN_RSP_TIMEOUT	The maximum amount of time after sending a BTP Session Handshake request to wait for a BTP Session Handshake response before closing the connection.	5 seconds
BTP_ACK_TIMEOUT	The maximum amount of time after receipt of a segment before a stand-alone ACK must be sent.	15 seconds
BTP_CONN_IDLE_TIMEOUT	The maximum amount of time no unique data has been sent over a BTP session before the Central Device must close the BTP session.	30 seconds

4.19.6. Bluetooth Requirements

The UUID is provided by Bluetooth SIG, Inc. and may only be used by its members in compliance with all terms and conditions of use issued by the Bluetooth SIG, Inc. For more information, visit <https://www.bluetooth.com/specifications/assigned-numbers/16-bit-uuids-for-sdos>.

Use of the Bluetooth extensions feature of this specification and specifically the MATTER_BLE_SERVICE_UUID is strictly prohibited unless the product is certified by both the Bluetooth SIG and the Connectivity Standards Alliance by a member of good standing of both organizations.

Table 36. SIG UUID assignment

Constant Name	Description	Value
MATTER_BLE_SERVICE_UUID	The UUID for the Matter-over-BLE service as assigned by the Bluetooth SIG.	0xFFFF6

4.20. Public Action Frame Transport Protocol (PAFTP)

The Public Action Frame Transport Protocol (PAFTP) is akin to BTP. PAFTP is used when commissioning over Wi-Fi Public Action Frame while BTP is used when commissioning over Bluetooth.

Just like BTP, PAFTP provides a TCP-like layer of reliable, connection-oriented data transfer of Matter commissioning messages in the Service Specific Info field inside the Service Descriptor Extension Attribute inside [WFA-USD](#) Service Discovery Frame (SDF) Follow-up messages. The PAFTP splits Matter message SDU into multiple PAFTP segments, which are each transmitted via a single [WFA-USD](#) SDF Follow-up message (as shown in [Figure 29, “MATTERoPAF: Matter Message / PAFTP layering”](#)). Each [WFA-USD](#) SDF Follow-up message is transmitted via a single Wi-Fi Public Action Frame.

Matter uses PAFTP to define a Matter-over-PAF (MATTERoPAF) Interface. A MATTERoPAF Interface

SHALL implement PAFTP. A MATTERoPAF Interface SHALL only be used to transport Matter messages as PAFTP SDU.




Figure 29. MATTERoPAF: Matter Message / PAFTP layering

The PAFTP session handshake allows devices to check PAFTP protocol version compatibility and exchange other data before a PAFTP session is established. Once established, this session is used to send and receive PAFTP SDUs (such as Matter messages) as PAFTP Message Segments. A PAFTP session MAY open and close with no effect on the state of the underlying WFA-USD connection, except in the case where a PAFTP session is closed by the Commissionable Device. Either the Commissionable Device or the Commissioner MAY signal the end of a PAFTP session by closing the underlying WFA-USD connection.

Due chiefly to constraints put on design by resource-limited Wi-Fi chipsets supporting WFA-USD, PAFTP defines a receive window for each side of a session in units of WFA-USD Follow-up message PDUs. Each WFA-USD Follow-up message PDU is sent with a sequence number which the receiver uses to acknowledge receipt of each packet at the PAFTP layer and open its receive window from the sender’s perspective.

4.20.1. PAFTP Session Interface

Conceptually, an open PAFTP session is exposed to the next-higher session layer as a full-duplex message stream.

4.20.2. PAFTP Frame Format

The PAFTP frame format is identical to BTP frame format as shown in Table 28, “BTP Packet PDU format”. A PAFTP Frame consists of an 8-bit header followed by one or more optional fields. PAFTP

uses little endian encoding for any header fields larger than one byte in length.

Table 37, “PAFTP Packet PDU format” defines the PAFTP Packet PDU format.

Unused fields SHALL be set to '0'.

Table 37. PAFTP Packet PDU format

bit 0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Control Flags								[Management Opcode]							
[Ack Number]								[Sequence Number]							
[Message Length]															
[Segment Payload]...															
...															

4.20.2.1. Control Flags

Table 38. PAFTP Control Flags

bit 7	6	5	4	3	2	1	0
-	H	M	-	A	E	C	B

H (Handshake) bit

Set to '0' for normal PAFTP packets. When set, this bit indicates a PAFTP handshake packet for session establishment and has a different packet format described below.

M (Management Message) bit

Indicates the presence ('1') or absence ('0') of the Management Opcode field. All segments of a message SHALL set this bit to the same value.

A (Acknowledgement) bit

Indicates the presence of the Ack Number field.

E (Ending Segment) bit

Set to '1' on the last segment of a PAFTP SDU and set to '0' for all other segments of the same PAFTP SDU. A segment MAY have both the Beginning and Ending bits set indicating that a full PAFTP SDU is included in the message. When set, the segment payload length is equal to the total remaining unreceived message data. When not set, the segment payload length is equal to the maximum allowable PAFTP session packet size minus header overhead.

C (Continuing Segment) bit

Set to '0' on the first segment of a PAFTP SDU and set to '1' for all remaining segments of the same PAFTP SDU.

B (Beginning Segment) bit

Set to '1' on the first segment of a PAFTP SDU and set to '0' for all remaining segments of the same PAFTP SDU. It indicates the presence of the Message Length field.

4.20.2.2. Ack Number

Optional field specified in [Section 4.20.3.8, “Packet Acknowledgements”](#).

4.20.2.3. Sequence Number

Mandatory field for regular data messages specified in [Section 4.20.3.6, “Sequence Numbers”](#).

4.20.2.4. Message Length

Optional field present in Beginning Segment only. Value indicates the length in bytes of the full message buffer to be transmitted. None of the PAFTP Packet PDU fields is included in the Message Length.

4.20.2.5. Segment Payload

Optional field containing a segment of the Service Data Unit (SDU) message in transmission to the receiver.

4.20.3. PAFTP Control Frames

PAFTP defines different control frame formats depending on the Management Opcode that is in the PAFTP Packet PDU header. Valid Management Opcodes for PAFTP Control Frames are defined in [Table 39, “PAFTP Control codes”](#).

Table 39. PAFTP Control codes

Management Opcode	Name	Description
0x6C	Handshake	Request and response for PAFTP session establishment

4.20.3.1. PAFTP Handshake Request

Table 40. PAFTP Handshake Request format

bit 0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Control Flags = 0x65								Management Opcode = 0x6C							
Ver[0]				Ver[1]				Ver[2]				Ver[3]			
Ver[4]				Ver[5]				Ver[6]				Ver[7]			
Supported Maximum Service Specific Info length															
Commissioner Window Size															

H (Handshake) bit

Set to '1' for connection handshake messages.

M (Management) bit

Set to '1' for connection handshake messages.

Ver[i] (Version) nibbles

Used to negotiate the highest version capability between a Device pair. Supported versions are listed once each, newest first, in descending order. Unused version fields are filled with '0'.

The following values are defined:

- 0 — Unused field
- 4 — PAFTP as defined by Matter v1.5

Supported Maximum Service Specific Info Length

Supported Maximum Service Specific Info Length is a 16-bit unsigned integer field containing the maximum size of the service information that can be transmitted and received by the sender in Service Specific Info field in Service Descriptor Extension attribute in [WFA-USD](#) SDF Follow-up message. The PAFTP can query [WFA-USD](#) layer locally within a Device to determine Supported Maximum Service Specific Info Length. If PAFTP is not aware of the Supported Maximum Service Specific Info Length, the value SHALL be set to '350'.

Commissioner Window Size

Value of the maximum receive window size supported by the Commissioner, specified in units of PAFTP packets where each packet length may be up to Supported Maximum Service Specific Info Length bytes minus PAFTP header. This maximum was chosen so a single PAFTP segment can fit into a single [WFA-USD](#) SDF Follow-up message PDU.

4.20.3.2. PAFTP Handshake Response

Table 41. PAFTP Handshake Response format

bit 0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
Control Flags = 0x65								Management Opcode = 0x6C							
Final Protocol Version				Reserved				Selected Maximum Service Specific Info Length (low byte)...							
...Selected Maximum Service Specific Info Length (high byte)															
								Selected Window Size							

H (Handshake) bit

Set to '1' for session handshake messages.

M (Management) bit

Set to '1' for session handshake messages.

Reserved

Must be set to '0'.

Final Protocol Version

Value of the PAFTP protocol version selected by the Commissionable Device.

Selected Maximum Service Specific Info Length

Value of the maximum size of the service information that can be transmitted and received in Service Specific Info field in Service Descriptor Extension attribute in [WFA-USD](#) SDF Follow-up message for this [WFA-USD](#) connection, selected by the Commissionable Device as a 16-bit unsigned integer. The PAFTP can query [WFA-USD](#) layer locally within a Device to determine Supported Maximum Service Specific Info Length. If PAFTP is not aware of the Supported Maximum Service Specific Info Length, the value SHALL be set to '350'. Selected Maximum Service Specific Info Length is equal to minimum of Supported Maximum Service Specific Info Length received in PAFTP Handshake Request and locally Supported Maximum Service Specific Info Length by the Commissionable Device.

Selected Window Size

Value of the maximum receive window size supported by the Commissionable Device, specified in units of PAFTP packets where each packet may be up to Selected Maximum Service Specific Info Length bytes minus PAFTP header.

4.20.3.3. Session Establishment

Before a PAFTP session can be initiated, a Commissioner SHALL first establish a [WFA-USD](#) connection to a Commissionable Device.

To initiate a PAFTP session, a Commissioner SHALL send a PAFTP handshake request packet to the Commissionable Device. The handshake request packet SHALL include, a list of PAFTP protocol versions supported by the Commissioner, the Commissioner's Supported Maximum Service Specific Info Length, and the Commissioner's maximum receive window size. The list of supported protocol versions SHALL be sorted in descending numerical order. For a detailed specification of the handshake request binary data format, see [Section 4.20.3.1, "PAFTP Handshake Request"](#).

Once the Commissionable Device has received a PAFTP handshake request, it SHALL send a PAFTP handshake response to the Commissioner. This response SHALL contain the window size, maximum selected PAFTP packet size, and PAFTP protocol version selected by the Commissionable Device. For a detailed specification of the handshake response binary data format, see [Section 4.20.3.2, "PAFTP Handshake Response"](#).

The Commissionable Device SHALL select a window size equal to the minimum of its and the Commissioner's maximum window sizes. Likewise, the Commissionable Device SHALL select a maximum PAFTP Segment Size including PAFTP header for the [WFA-USD](#) connection by taking the mini-

imum of Supported Maximum Service Specific Info Length received in PAFTP Handshake Request and locally Supported Maximum Service Specific Info Length by the Commissionable Device.

The Commissionable Device SHALL select a PAFTP protocol version that is the newest which it and the Commissioner both support, where newer protocol version numbers are strictly larger than those of older versions. The version number returned in the handshake response SHALL determine the version of the PAFTP protocol used by Commissionable Device and Commissioner for the duration of the session.

If the Commissionable Device determines that it and the Commissioner do not share a supported PAFTP protocol version, the Commissionable Device SHALL close its [WFA-USD](#) connection to the Commissioner. When a Commissioner sends a handshake request, it SHALL start a timer with a globally-defined duration of PAFTP_CONN_RSP_TIMEOUT seconds. If this timer expires before the Commissioner receives a handshake response from the Commissionable Device, the Commissioner SHALL close the PAFTP session and report an error to the application.




Figure 30. PAFTP session handshake

4.20.3.4. Data Transmission

Once a PAFTP session has been established, the next-higher-layer application on both peers may use this [WFA-USD](#) connection to send and receive Matter commissioning messages via [WFA-USD](#) SDF Follow-up messages.

All PAFTP packets sent on a [WFA-USD](#) connection SHALL adhere to the PAFTP Packet PDU binary data format specified in PAFTP Frame Formats. All PAFTP packets SHALL include a header flags byte and an 8-bit unsigned sequence number. All other packet fields are optional. These optional fields include an 8-bit unsigned received packet acknowledgement number, 16-bit unsigned buffer length indication, and variable-length buffer segment payload.

While a single PAFTP connection may exist between a PAFTP pair, multiple PAFTP sessions may be established between various peers. This section is defined entirely within the scope of a single PAFTP session. Concurrent PAFTP sessions between the same peer and multiple peers SHALL maintain separate and independent acknowledgement timers, sequence numbers, and receive windows.

4.20.3.5. Message Segmentation and Reassembly

When the session layer (that is, MATTERoPAF) sends a Matter Message as a PAFTP SDU over a

PAFTP session, that PAFTP SDU SHALL be split into ordered, non-overlapping PAFTP segments so the set of all PAFTP segments may be reassembled into the original PAFTP SDU (see [Figure 29, “MATTERoPAF: Matter Message / PAFTP layering”](#)). Each PAFTP segment SHALL be prepended with a PAFTP packet header and sent in the Service Specific Info field inside the Service Descriptor Extension Attribute inside of a single [WFA-USD](#) SDF Follow-up message. If a PAFTP SDU is split into more than one PAFTP segment, the PAFTP segments SHALL be sent in order of their position in the original PAFTP SDU, starting with the PAFTP segment at the buffer’s head.

At any point in time, only one PAFTP SDU may be transmitted in each direction over a PAFTP session. The transmission of PAFTP segments of any two PAFTP SDUs SHALL NOT overlap. If the application attempts to send one PAFTP SDU while transmission of another PAFTP SDU is in progress, the new PAFTP SDU SHALL be appended to a first-in, first-out queue. The next PAFTP SDU SHALL be dequeued from this queue and transmitted once transmission of the current PAFTP SDU completes, that is, once all PAFTP segments of the current PAFTP SDU have been transmitted and received by the peer.

The first PAFTP segment of a PAFTP SDU sent over a PAFTP session SHALL have its Beginning Segment header flag set to indicate the beginning of a new PAFTP SDU (see [Table 37, “PAFTP Packet PDU format”](#)). The presence of this flag SHALL indicate the further presence of a 16-bit unsigned integer field, the Message Length, that provides the receiver with the total length of the PAFTP SDU. The last PAFTP segment for a given PAFTP SDU SHALL have its Ending Segment flag set to indicate the end of the transmitted PAFTP SDU. A PAFTP packet that bears an unsegmented PAFTP SDU—that is, a PAFTP SDU small enough to fit into a single PAFTP segment—SHALL have both its Beginning Segment and Ending Segment flags set.

The size of a single PAFTP SDU sent via PAFTP is limited to 64KB, that is, the maximum size of the Message Length field in the PAFTP packet header. The number of segments used to send a buffer is unlimited and delimited by the Beginning Segment and Ending Segment bits in the PAFTP packet header. The upper layer imposes more stringent requirements over the maximum SDU size, such as [Section 4.4.4, “Message Size Requirements”](#).

The length of the data payload in each PAFTP segment whose Ending Segment bit is not set SHALL be equal to the session’s maximum PAFTP packet size minus the size of that PAFTP packet’s header. If a packet’s Ending Segment bit is set, the length of its PAFTP segment data payload SHALL equal the size of the original PAFTP SDU minus the total size of all previously transmitted PAFTP segments of that PAFTP SDU. In this way, the length of a SDU’s last PAFTP segment is implied by its size.

A peer SHALL reassemble PAFTP segments in the ascending order of the segment sequence numbers. Once a peer receives a complete set of PAFTP segments, it SHALL verify that the reassembled PAFTP SDU’s total length matches that specified by the Beginning Segment’s Message Length value. If they match, the receiver SHALL pass the reassembled PAFTP SDU up to the next-higher-layer. If the reassembled buffer’s length does not match that specified by the sender, or if received PAFTP segment payload size would exceed the maximum PAFTP packet size, or receiver receives an Ending Segment without the presence of a previous Beginning Segment, or a Beginning Segment when another PAFTP SDU’s transmission is already in progress, the receiver PAFTP SHALL close the PAFTP session and report an error to the application.

4.20.3.6. Sequence Numbers

All PAFTP packets SHALL be sent with sequence numbers, regardless of whether they contain SDU segments (for example, a packet acknowledgement with no attached segment payload). The purpose of sequence numbers is to facilitate the PAFTP receive window. A PAFTP sequence number SHALL be defined as an unsigned 8-bit integer value that monotonically increments by 1 with each packet sent by a given peer. A sequence number incremented past 255 SHALL wrap to zero.

Sequence numbers SHALL be separately defined for either direction of a PAFTP session. The sequence number of the first packet sent by the Commissioner after completion of the PAFTP session handshake SHALL be zero. The Commissionable Device PAFTP handshake response bears an implied sequence number of zero because it occupies a slot in the Commissioner's receive window. The Commissioner acknowledges the Commissionable Device's PAFTP handshake response with an acknowledgement sequence of zero. For this reason, the sequence number of the first data packet sent by the Commissionable Device after completion of the PAFTP session handshake SHALL be 1.

Peers SHALL check to ensure that all received PAFTP packets properly increment the sender's previous sequence number by 1. If this check fails, the peer SHALL close the PAFTP session and report an error to the application.

4.20.3.7. Receive Windows

The purpose of the receive window is to enable flow control at the [WFA-USD](#) layer between PAFTP peers.

Flow control is required at the [WFA-USD](#) transport layer for embedded platforms that use "minimal" Wi-Fi chipsets supporting [WFA-USD](#). These platforms may have limited space on the host processor to receive packets from their Wi-Fi chipsets. When the Wi-Fi chip sends the result of a received [WFA-USD](#) SDF Follow-up message PDU to the host processor, that payload and the corresponding PAFTP packet will be permanently lost if the host does not have enough space to receive it. For this reason, knowledge of a remote host's ability to reliably receive [WFA-USD](#) SDF Follow message PDUs is presented at the transport layer in the form of the PAFTP receive window.

Both peers in a PAFTP session SHALL define a receive window, where the window's size indicates the number of [WFA-USD](#) SDF Follow-up message PDUs (that is, PAFTP segments) a peer can reliably receive and store without session-layer acknowledgment. A maximum window size SHALL be established for both peers as part of the PAFTP session handshake. To prevent sequence number wrap-around, the largest maximum window size any peer may support is 255.

Both peers SHALL maintain a counter to reflect the current size of the remote peer's receive window. Each peer SHALL decrement this counter when it sends a packet and increment this counter when a sent packet is acknowledged.

If a local peer's counter for a remote peer's receive window is zero, the window SHALL be considered closed, and the local peer SHALL NOT send packets until the window reopens (is incremented above zero). When a closed window reopens, a local peer SHALL immediately resume any pending PAFTP packet transmission.

A local peer SHALL also not send packets if the remote peer's receive window has one slot open and the local peer does not have a pending packet acknowledgement. This is to avoid the situation

where the receive windows of both peers are full and neither can send an acknowledgement to reopen its window for the other. Because the Commissionable Device’s handshake response bears an implicit PAFTP sequence number of zero, a Commissionable Device SHALL initialize its counter for the Commissioner’s receive window size at (negotiated maximum window size - 1). A Commissioner SHALL initialize its counter for the server’s receive window at the negotiated maximum window size.

Both peers SHALL also keep a counter of their own receive window size based on the sequence number difference between the last packet they received and the last packet they acknowledged. This counter is used to proactively send early packet acknowledgements when a peer’s own receive window is about to close. See [Section 4.20.3.8, “Packet Acknowledgements”](#) for details.

An example scenario involving PAFTP receive windows is depicted in [Figure 31, “Example receive window scenario”](#), complete with packet acknowledgements as specified in [Section 4.20.3.8, “Packet Acknowledgements”](#). In this scenario, the Commissioner transmits a three-segment buffer to the Commissionable Device once it receives the server’s handshake response. The Commissionable Device’s initial receive window is empty. Both Commissioner and Commissionable Device have a maximum window size of 4.




Figure 31. Example receive window scenario

4.20.3.8. Packet Acknowledgements

The purpose of sequence numbers and packet receipt acknowledgements is to support the PAFTP receive window and its reordering, and provide a keep-alive signal when a session is idle to affirm

the health and continued operation of a remote PAFTP stack.

Per PAFTP Frame Formats, PAFTP packet receipt acknowledgements SHALL be received as unsigned 8-bit integer values in the header of a PAFTP packet. The value of this field SHALL indicate the sequence number of the acknowledged packet.

Acknowledgement of a sequence number indicates acknowledgement of the previous sequence number, if it too is unacknowledged. By induction, acknowledgement of a given packet implies acknowledgement of all packets received on the same PAFTP session prior to the acknowledged packet.

An acknowledgement is invalid if the acknowledged sequence number does not correspond to an outstanding, unacknowledged PAFTP packet sequence number. In contrast to TCP, PAFTP acks are not "free." A stand-alone ack—that is, a PAFTP packet that contains a packet receipt acknowledgement value but no buffer segment payload—consumes a slot in a remote peer's window just like any other packet. Stand-alone acknowledgement packets SHALL be acknowledged by a remote peer. The implications of this are examined in [Section 4.20.3.9, "Idle Connection State"](#).

Each peer SHALL maintain an acknowledgement-received timer. When a peer sends any PAFTP packet, it SHALL start this timer if it is not already running. The timer's duration SHALL be globally defined as PAFTP_ACK_TIMEOUT seconds, referred to as the acknowledgement timeout interval.

A peer SHALL restart its acknowledgement-received timer when a valid acknowledgement is received for any but its most recently sent unacknowledged packet. A peer SHALL stop its acknowledgement-received timer if it receives an acknowledgement for its most recently sent unacknowledged packet. If a peer's acknowledgement-received timer expires, or if a peer receives an invalid acknowledgement, the peer SHALL close the PAFTP session and report an error to the application.

Because the server's handshake response bears an implicit PAFTP sequence number of zero, a server SHALL start its acknowledgement-received timer when it sends a handshake response.

Each peer SHALL also maintain a send-acknowledgement timer. When it receives any PAFTP packet, a peer SHALL record the packet's sequence number as the corresponding PAFTP session's pending acknowledgement value and start the send-acknowledgement timer if it is not already running. The timer's duration SHALL be defined as any value less than one-half the acknowledgement timeout interval. This ensures that on a healthy [WFA-USD](#) connection, a peer will always receive acknowledgements for sent packets before its acknowledgement-received timer expires.

A peer SHALL restart its send-acknowledgement timer when a pending acknowledgement is sent, either as a stand-alone PAFTP packet or piggybacked onto an outgoing buffer segment, for any but its received unacknowledged packet with the largest packet sequence number. A peer SHALL stop its send-acknowledgement timer when a pending acknowledgement is sent, either as a stand-alone PAFTP packet or piggybacked onto an outgoing buffer segment, for its received unacknowledged packet with the largest packet sequence number. If this timer expires and the peer has a pending acknowledgement, the peer SHALL immediately send that acknowledgement, and acknowledgement number SHALL be the PAFTP packet sequence number of which all previous contiguous PAFTP packet sequence numbers are already received. If the peer sends any packet before this timer expires, it SHALL piggyback any pending acknowledgement on the transmitted packet, and acknowledgement number SHALL be the PAFTP packet sequence number of which all previous

contiguous PAFTP packet sequence numbers are already received. If this timer is not expired, the peer MAY send a pending acknowledgement as a stand-alone PAFTP packet, and acknowledgement number SHALL be the PAFTP packet sequence number of which all previous contiguous PAFTP packet sequence numbers are already received.

Because the server’s handshake response bears an implicit PAFTP sequence number of zero, a client SHALL set its pending acknowledgement value to zero and start its send-acknowledgement timer when it receives the server’s a handshake response. Operation of the send-acknowledgement and acknowledgement-received timers is illustrated in [Figure 33, “PAFTP session lifecycle for Commissioner”](#) in [Section 4.20.3.11, “Protocol State Diagrams”](#).

If a peer detects that its receive window has shrunk to two or fewer free slots, it SHALL immediately send any pending acknowledgement as a stand-alone PAFTP packet, and acknowledgement number SHALL be the PAFTP packet sequence number of which all previous contiguous PAFTP packet sequence numbers are already received. This prevents the session from stalling in the interval between when a peer’s receive window becomes empty and when its send-acknowledgement timer would normally fire.

4.20.3.9. Idle Connection State

When neither side of a PAFTP session has data to send, PAFTP packets will still be exchanged every send-acknowledgement interval due to acknowledgements generated by the receipt of previous data or stand-alone acknowledgement packets, as discussed in [Section 4.20.3.8, “Packet Acknowledgements”](#). The behavior of the acknowledgement-received timer in this scenario doubles as a keep-alive mechanism, as it will cause a peer to close a [WFA-USD](#) connection automatically if the remote PAFTP stack crashes or becomes unresponsive. This scenario is illustrated in [Figure 32, “Idle connection scenario”](#).




Figure 32. Idle connection scenario

4.20.3.10. Connection Shutdown

A Commissioner MAY terminate the subscribe instance, to close a PAFTP session. A Commissionable Device MAY terminate the publish instance, to close a PAFTP session.

4.20.3.11. Protocol State Diagrams

Figure 33, “PAFTP session lifecycle for Commissioner” shows the state machine for PAFTP session management of a PAFTP Commissioner Device.




Figure 33. PAFTP session lifecycle for Commissioner

Figure 34, “PAFTP session lifecycle for Commissionable Device” shows the state machine for PAFTP session management of a Commissionable Device.




Figure 34. PAFTP session lifecycle for Commissionable Device

Figure 35, “State diagram for PAFTP session post-establishment” shows the state machine for PAFTP session maintenance at the protocol level, including liveness enforcement through keep alive messages and automatic teardown if acknowledgements are received before the timeout.




Figure 35. State diagram for PAFTP session post-establishment

4.20.4. Parameters and Constants

Table 42, “Glossary of constants” is a glossary of constants used in this chapter, along with a brief description and the default for each constant.

Table 42. Glossary of constants

Constant Name	Description	Default
PAFTP_CONN_RSP_TIMEOUT	The maximum amount of time after sending a PAFTP Session Handshake request to wait for a PAFTP Session Handshake response before closing the WFA-USD connection.	5 seconds

Constant Name	Description	Default
PAFTP_ACK_TIMEOUT	The maximum amount of time after receipt of a segment before a stand-alone ACK must be sent.	15 seconds
PAFTP_CONN_IDLE_TIMEOUT	The maximum amount of time no unique data has been sent over a PAFTP session before the Commissioner must close the PAFTP session.	30 seconds

4.21. NFC Transport Layer (NTL)

The NFC Transport Layer (NTL) defines how to transfer Matter commissioning messages over an NFC interface.

NOTE

NTL is only used when performing device commissioning over an NFC interface (referred as NFC-based Commissioning). NTL is not used for the reading of the Onboarding Payload from an NFC Tag.

Messages are embedded into Application Programming Data Units (APDU) in compliance with [ISO/IEC 7816-4](#). These APDUs are transferred in blocks according to the ISO-DEP protocol as defined in [NFC Forum Digital Specification](#). Each block is transmitted wirelessly using NFC-A technology as defined in [NFC Forum Digital Specification](#).

The layering inside NTL is shown in [NFC Transport Layer stack](#), and described in following sections.




Figure 36. NFC Transport Layer stack

NTL provides a reliable, datagram-oriented, transport interface with asymmetric roles: one end always transmits first, and the other end always responds. When NTL is used for Matter commissioning, the Commissioner always transmits first and the Commissionee responds.

Devices using NTL SHALL also include NFC onboarding payload as defined in [NFC Tag](#), ensuring the NFC interface provides a consistent commissioning solution.

Additionally, devices supporting NTL SHALL support an alternate commissioning channel (e.g., BLE, Wi-Fi PAF, Ethernet). This guarantees that commissioning is always possible, even if NFC-based

commissioning is unavailable on the Commissioner.

4.21.1. NFC Forum requirements

Because NTL relies on correct implementation of NFC Forum specification:

- products implementing NTL as a Commissioner SHALL comply with the NFC Forum requirements for an NFC Forum Reader/Writer Module supporting Type 4A Tag Operation,
- products implementing NTL as a Commissionee SHALL comply with the NFC Forum requirements for either NFC Forum Type 4A Tag Module or NFC Forum Type 4A Tag Platform Module (Card Emulation).

NOTE

Corresponding NFC Forum Devices eligible to receive the NFC Forum Certification Mark are defined in [NFC Forum Devices Requirements](#).

4.21.2. NFC-A technology

Commissioners supporting NTL SHALL act as an NFC Forum Type 4A Tag Platform in Poll Mode as defined in [NFC Forum Digital Specification](#).

Commissionees supporting NTL SHALL act as an NFC Forum Type 4A Tag Platform in Listen Mode as defined in [NFC Forum Digital Specification](#).

4.21.3. ISO-DEP

ISO-DEP is a half-duplex block transmission protocol providing retransmission and chaining features.

The full ISO-DEP protocol SHALL be implemented in compliance with [NFC Forum Digital Specification](#).

This section highlights some features of the ISO-DEP protocol in the context of Matter communication for information.

During NFC protocol activation, the NFC Reader Device sends a RATS frame in which the maximum frame size it can receive (FSD) is coded, and the NFC Tag Device sends an ATS frame in which the maximum frame size it can received (FSC) is coded.




Figure 37. ISO-DEP protocol activation

This information is used by the ISO-DEP layer to segment the APDU to transmit, utilizing the ISO-DEP chaining feature. On the receiver side, the ISO-DEP layer takes care of re-assembling the

chained segments.




Figure 38. ISO-DEP chaining from initiator to responder




Figure 39. ISO-DEP chaining from responder to initiator

A default maximum response timing is announced by the responder in the RATS frame, but when the responder requires more time to respond, it can extend the initiator waiting time by sending appropriate S-block as shown in [ISO-DEP Waiting time extension](#)




Figure 40. ISO-DEP Waiting time extension

The ISO-DEP protocol also features retransmission when a frame is not received or incorrectly received, making it a reliable transport protocol.

4.21.4. APDU layer

The APDU layer encapsulates Matter messages into APDUs conforming to [ISO/IEC 7816-4](#). It provides a datagram-oriented transport interface tunneled over encapsulation in a series of command requests and their associated responses.

For this purpose, the following commands are used:

Table 43. APDU commands used in NTL

Class	Commands
Interindustry	SELECT, GET RESPONSE
Proprietary	TRANSPORT

The commands (C-APDU) are always issued by the NFC Reader/Writer which is the Matter commissioner, and the responses (R-APDU) are always issued by the NFC Listener which is the Matter commissionee.

R-APDU always embeds a status code, and optionally a response payload.

Some smartphone NFC Reader/Writer implementations are limited to a maximum 256-byte APDU payload, which is insufficient to transfer Matter messages required for commissioning. To guarantee interoperability with all NFC Readers/Writers while limiting the number of cases to support, both the NFC Reader/Writer and NFC listener SHALL always use short field coding (aka short length field) of APDUs.

Both commissionee and commissioner SHALL support C-APDU and R-APDU chaining specified in [ISO/IEC 7816-4](#).

In case the size of the Matter message to transmit in the **TRANSPORT** command APDU is bigger than the maximum size that can be transmitted by this APDU, the APDU chaining procedure specified in [ISO/IEC 7816-4](#) SHALL be used, even though the **TRANSPORT** command belongs to proprietary class. Thanks to the APDU chaining procedure, the protocol can handle Matter messages up to 65535 bytes. However, the maximum size of Matter Message than can be actually transferred is limited by memory constraints of the implementation and [Message Size Requirements](#).

4.21.4.1. SELECT command

Matter commissioning SHALL be initiated by the NFC Reader/Writer by issuing the 7816-4 SELECT command with the Application Identifier (AID) 'A0 00 00 09 09 8A 77 E4 01' as shown in [Table 44](#).

Table 44. NTL SELECT command APDU

CLA	INS	P1	P2	Lc	Data (hexadecimal bytes)	Le
0x00	0xA4	0x04	0x0C	0x09	A0:00:00:09:09:8A:77:E4:01	0x00

When in commissioning mode, a commissionee SHALL answer to a correct SELECT command with a successful response APDU as shown in [Table 45](#).

Table 45. NTL SELECT Successful response APDU

Data							SW1	SW2
Version (8 bits)	0x00 (8 bits)	0x0 (4 bits)	Discriminator (12 bits)	Vendor ID (16 bits)	Product ID (16 bits)	Extended Data (0 or more bits)	0x90	0x00

Data fields:

- **Version** is a uint8 that SHALL encode the NTL protocol version supported by the commissionee. The version SHALL be 0x01. Other values SHALL be reserved for future use. The commissioner SHALL use the corresponding NTL protocol to communicate with the commissionee.
- The next 8 bits, cleared, are undefined (reserved).
- The next 4 bits, cleared, encode the format of the following fields.
- The remaining fields are described in [Section 5.4.2.4, “Discovery Information”](#).
 - Devices MAY choose not to advertise either the Vendor ID and Product ID, or only the Product ID due to privacy or other considerations. When choosing not to advertise both Vendor ID and Product ID, the device SHALL set both Vendor ID and Product ID fields to 0. When choosing not to advertise only the Product ID, the device SHALL set the Product ID field to 0. A device SHALL NOT set the Vendor ID to 0 when providing a non-zero Product ID.
 - **Extended Data** MAY be omitted.

When not in commissioning mode, a commissionee SHALL either ignore the command (no response) or answer with an error response APDU as shown in [Table 46](#) indicating that conditions of use are not satisfied.

Table 46. NTL SELECT Error response APDU

Data	SW1	SW2
Version	0x69	0x85

In other error cases, a commissionee MAY answer with another error response as described in [ISO/IEC 7816-4](#).

4.21.4.2. TRANSPORT command

Once commissioning has been successfully initiated with the SELECT command, Matter messages SHALL be exchanged using the proprietary TRANSPORT command. This command is defined by instruction (INS) code 0x20 with parameters P1 and P2 encoding the length of the message to transmit.

Table 47. TRANSPORT Command APDU

CLA	INS	P1	P2	Lc	Data	Le
0x80 (unchained) 0x90 (chained)	0x20	...	...	...	...	...

- The **Lc** single-octet field SHALL encode the length in octets of the payload's **Data** field in compliance with [ISO/IEC 7816-4](#). To optimize underlying ISO-DEP chaining, **Lc** SHOULD be less or equal than $\min(255, FSC-10)$.
- **P1** and **P2** SHALL encode the number of octets of the full message to transmit. It is encoded as a 2-bytes integer with **P1** being the most significant byte. In case the full message fits in a single command, the value is equal to **Lc**. In case the message requires chaining, the value is bigger than **Lc**. The same value SHALL be used in all chained commands.
- The **Data** field SHALL contain a fragment of the message to transmit, possibly the only one.
- The **Le** single-octet field SHALL encode the maximum length in octets that the reader/writer can receive in the response APDU in compliance with [ISO/IEC 7816-4](#). To optimize underlying ISO-DEP chaining, **Le** SHOULD encode the value equal to $\min(256, FSD-6)$.

The response APDU allows the responder to transmit a message to the initiator. In case the full message fits within the number of bytes encoded by **Le** in C-APDU, a successful response SHALL be indicated by the **SW1** and **SW2** values in [Table 48](#).

Table 48. TRANSPORT successful Response APDU - Complete

Data	SW1	SW2
Message or last fragment of message to transmit	0x90	0x00

In case the full message does not fit within the number of bytes encoded by **Le** in C-APDU, a successful but incomplete response SHALL be indicated by the **SW1** value in [Table 49](#), and **SW2** SHALL encode the number of bytes of message to be sent in the next **GET RESPONSE** R-APDU.

Table 49. TRANSPORT successful Response APDU - Incomplete message

Data	SW1	SW2
Fragment of message to transmit	0x61	...

In case the chained message exceeds the maximum supported message size, an error response conforming to [Table 50](#) SHALL be issued, indicating 'Not enough memory space'.

Table 50. TRANSPORT error Response APDU - Not enough memory space

SW1	SW2
0x6A	0x84

4.21.4.3. GET RESPONSE command

This command SHALL be issued following the reception of an incomplete successful R-APDU, in compliance with [ISO/IEC 7816-4](#).

Table 51. GET RESPONSE command APDU

CLA	INS	P1	P2	Le
0x00	0xC0	0x00	0x00	...

- The **Le** single-octet field SHALL encode the maximum length in octets that the reader/writer can receive in the response APDU in compliance with [ISO/IEC 7816-4](#). To optimize underlying ISO-DEP chaining, **Le** SHOULD encode the value equal to $\min(256, \text{FSD}-6)$.

Successful response APDU is in following tables:

Table 52. GET RESPONSE successful Response APDU - Complete message

Data	SW1	SW2
Message or last fragment of message to transmit	0x90	0x00

Table 53. GET RESPONSE successful Response APDU - Incomplete message

Data	SW1	SW2
Fragment of message to transmit	0x61	...

In case GET RESPONSE is received, but not following a [TRANSPORT successful Response APDU - Incomplete message](#), the commissionee SHALL answer with an error indicating the condition of use is not satisfied, as shown in [Table 54](#)

Table 54. GET RESPONSE error Response APDU - Condition of use not satisfied

SW1	SW2
0x69	0x85

4.21.4.4. APDU sequence diagram

A sequence diagram showing the APDU chaining feature is shown in [NTL APDU sequence diagram](#).




Figure 41. NTL APDU sequence diagram

4.22. Check-In Protocol

The goal of the Check-In Protocol is to provide a way for a server to notify a client of an event or state outside of a secure session in a private and secure fashion.

Some of the events or states to be shared with a client are defined by different Check-in Protocol Use Cases. The current list includes:

- An [ICD](#) is available for communication.

The Check-In Protocol is a set of requirements and processes that can be re-used by multiple Check-In Protocol use cases. Multiple Check-In Protocol use cases can coexist on a single device. The

Check-In Protocol sends a sessionless message (Check-In message) that relies on a key that has been given to the server during the registration of the client.

4.22.1. Requirements

4.22.1.1. Server Requirements

A Server use case of the Check-In Protocol SHALL implement these requirements.

- A persistent data store for Keys.
- An implementation of the [Check-In Counter](#).

Persistent Data Store

For the Check-In Protocol to be effective, the use case of the protocol SHALL provide a means to store and persist client registration information. The minimum information that SHALL be stored and persisted is the key to use during the encryption process of the Check-In message.

4.22.1.2. Client Requirements

The server does not store the starting value of the [Check-In Counter](#) for a given key. Therefore, it is the client’s responsibility to do so when the client registers with the server. A Client use case of the Check-In Protocol SHALL implement a means to store sets containing a key, the starting value of [Check-In Counter](#) and the last known valid offset from the starting value of the [Check-In Counter](#). A Key that has been used to register for Check-In messages SHALL always be associated with a starting [Check-In Counter](#) value and the last known valid offset from the starting [Check-In Counter](#) value.

During the decryption process, the client uses the key, its associated starting [Check-In Counter](#) value and offset to decrypt and authenticate a received Check-In message.

4.22.2. Message Content

Table 55. Check-In Payload

Field	Type	Description
nonce	byte[CRYPTO_AEAD_NON-CE_LENGTH_BYTES]	First plaintext field
Check-In Counter	uint32	Encrypted field - Counter used to generate the nonce
Application Data	byte[]	Encrypted field - Application Data that can be added to the Check-In message
MIC	byte[CRYPTO_AEAD-_MIC_LENGTH_BYTES]	Message integrity Check

4.22.2.1. Nonce

The nonce used to encrypt the encrypted fields of the payload SHALL be the first [CRYPTO_AEAD_NONCE_LENGTH_BYTES](#) bytes of the message payload. The nonce is included in plaintext.

4.22.2.2. Encrypted Section

The encrypted section SHALL be a concatenation of the following, in order:

1. The [Check-In Counter](#) value used in the encryption process, represented as a 4-byte little-endian value.
2. The Application Data, if used by the use case.

Check-In Counter

The [Check-In Counter](#) value SHALL be used to generate the nonce. This value SHALL be provided by the Check-In Protocol use case.

Application Data

The application data is defined and provided by the Check-In Protocol use case. The Check-In Protocol use case is not required to use application data. If application data is not used for a specific use-case, creation of a Check-In message SHALL use a zero-length byte array for the application data.

4.22.2.3. MIC

Message integrity check generated by the AEAD is appended after the encrypted section. It SHALL be the last [CRYPTO_AEAD_MIC_LENGTH_BYTES](#) bytes.

4.22.3. Algorithms

The Check-In message is a sessionless message, since one of the goals of the protocol is to provide a means to recover a secure session that was lost. Since no session keys are available to encrypt the entire message, the encryption of the Check-In message is limited to part of the payload.

The inputs provided to the encryption process by the Check-In protocol use case are:

1. The value of [Check-In Counter](#) represented as 4-byte little-endian value. In the following steps, we assume the counter value is in the correct format.
2. The application data bytes represented as a byte array. In the following steps, we assume the application data is in the correct format.
3. A symmetric encryption key represented as a [CRYPTO_SYMMETRIC_KEY_LENGTH_BYTES](#)-byte array. In the following steps, we assume the key is in the correct format.

4.22.3.1. Nonce Generation

To generate the nonce, the Check-In message SHALL use the [Keyed-Hash Message Authentication Code](#) algorithm. The inputs of the algorithm SHALL be the key and the [Check-In Counter](#) value. Both will be provided by the Check-In Protocol use case. From the generated hash, the nonce SHALL be the first [CRYPTO_AEAD_NONCE_LENGTH_BYTES](#) bytes to match the AEAD requirements.

```
nonce = Crypto_HMAC(key = key, message = counter)[0..(CRYPTO_AEAD_NONCE_LENGTH_BYTES-1)]
```

4.22.3.2. Authenticated encryption

The payload of the Check-In message SHALL be encrypted using the [AEAD](#) function. The inputs of the algorithm SHALL be the generated [nonce](#), the key and the plaintext to encrypt. The additional data field SHALL NOT be used.

```
Ciphertext = Crypto_AEAD_GenerateEncrypt(  
    K = key  
    P = plaintext,  
    A = "",  
    N = nonce)
```

4.22.3.3. Client Counter Validation

When receiving a Check-In message, the client SHALL verify the received counter by executing the following steps.

1. Subtract the stored starting value of the counter from the value provided in the Check-In message. The subtraction SHALL be done as unsigned integers mod 2^{32} . This ensures that even if the subtraction rolls over, it will still produce the correct offset.
2. If the stored last known offset from the stored starting value of the counter is smaller than the value from step 1, the verification SHALL return TRUE.
3. Otherwise the verification SHALL return FALSE.

4.22.3.4. Key Refreshing

A given key is only valid for one set of values of the [Check-In Counter](#). A key becomes invalid once all the [Check-In Counter](#) values for that key have been used. The tracking of the used Check-In counter values is done as part of the [decryption procedure](#).

To avoid reuse of a nonce value with a given key, the key (shared secret) needs to be refreshed before all the valid [Check-In Counter](#) values for that key have been used. If a [Check-In Counter](#) value is reused with the same key, the Check-In messages reusing the counter will be discarded. To avoid the situation where the relationship is broken because of a rollover, we are prescribing when a client should refresh its entry.

Key Refresh Validation

When a client receives a Check-In message with a [Check-In Counter](#) value indicating that 2^{31} counter values have been used, the client SHALL refresh its entry with a new key. The client SHALL determine if a key refresh is necessary by executing the following steps.

1. Check whether the stored counter offset is greater than or equal to 2^{31} .

2. If the stored counter offset is greater than or equal to 2^{31} , a key refresh is required.
3. Otherwise, a key refresh is allowed, but not required.

4.22.4. Protocol Operation

4.22.4.1. Encryption Procedure

When sending a Check-In message, the server SHALL encrypt the application payload (i.e. [Check-In Counter](#) and application data provided) following these steps. If the encryption procedure fails to generate the Check-In message, the server SHALL NOT send the associated Check-In message.

1. The server SHALL generate the nonce as described in [Section 4.22.3.1, “Nonce Generation”](#).
 - a. If the generation failed, the encryption procedure SHALL return FAILURE.
 - b. If the generation succeeded, the server SHALL continue from step 2.
2. The server SHALL encrypt the plaintext as described in [Section 4.22.3.2, “Authenticated encryption”](#).
 - a. The plaintext SHALL be a concatenation of the following, in order :
 - i. The [Check-In Counter](#) value used in the encryption process and
 - ii. the Application Data, if used by the use case.
 - b. If the encryption fails, the encryption procedure SHALL return FAILURE.
 - c. If the encryption succeeds, the server SHALL continue from step 3.
3. The server SHALL create the payload string by concatenating the nonce and the output of the AEAD function.

payload = nonce || Ciphertext

Check-In Protocol Encryption Diagram




Figure 42. Check-In Protocol Encryption Legend

4.22.4.2. Decryption Procedure

When a client receives a Check-In message, the client SHALL decrypt and process the message following these steps:

- 1. The client SHALL break down the message into the nonce and the ciphertext/MIC.

```

nonce           = payload[0:CRYPTO_AEAD_NONCE_LENGTH_BYTES]
ciphertext/MIC = payload[CRYPTO_AEAD_NONCE_LENGTH_BYTES:]

```

2. The client SHALL try to decrypt the message with the key associated to the server being tested. When a client has multiple associated servers, the client iterates through this process for the keys associated to each potential server until the keys associated to all possible servers fail or a successful key is found.

```

Success, Plaintext = Crypto_AEAD_DecryptVerify(
    K = Key,
    C = ciphertext/MIC,
    A = "",
    N = nonce)

```

- a. If the decryption succeeds, the device associated with the key is identified as the server checking in. The client SHALL continue from step 3.
 - b. If the decryption fails, the client SHALL move on to the next candidate server key per step 2, or discard the message if there are no remaining registered server keys to try.
3. From the plaintext, the client SHALL retrieve the [Check-In Counter](#) and the application data if present.
4. The client SHALL verify that the received nonce is valid by calculating it as described in [Section 4.22.3.1, “Nonce Generation”](#) from the received [Check-In Counter](#) and the key. The counter value retrieved in the previous step is already in the correct encoding.
 - a. If the calculated nonce matches the received nonce, the client SHALL continue from step 5.
 - b. If the calculated nonce doesn't match the received nonce, the client SHALL discard the message.
5. The client SHALL verify that the received [Check-In Counter](#) is valid as described in [Section 4.22.3.3, “Client Counter Validation”](#). The [Check-In Counter](#) is converted into a unsigned 32-bit integer.
 - a. If the received counter value is valid, the client SHALL continue from step 6.
 - b. If the received counter value is invalid, the client SHALL discard the message.
6. The client SHALL store the offset between the received [Check-In Counter](#) value and the stored starting value of the [Check-In Counter](#). The subtraction SHALL be done as unsigned integers mod 2^{32} . The client SHALL continue from step 7.
7. The client SHALL verify if a key refresh is necessary as described in [Section 4.22.3.4.1, “Key Refresh Validation”](#).
 - a. If a key refresh is necessary, the client SHALL trigger an attempt to [refresh the key](#), unless the client is already in the process of refreshing this specific key for this specific server. The client SHALL continue from step 8 whether the [key refresh](#) succeeds or not.
 - b. If a key refresh is not necessary, the client SHALL continue from step 8.

8. The client SHALL notify the application that the server has checked-in. What this behavior implies is application specific.

4.23. Thread Commissioning Protocol

The Thread Commissioning Protocol is used for commissioning over Thread. From the commissionee's standpoint, it's akin to commissioning over the operational network. From the commissioner's standpoint, the commissioning messages are encapsulated within the Relay commands as specified by the [Thread specification](#). Although the [Thread specification](#) discusses DTLS payloads being exchanged between the Joiner and the Commissioner, the payloads being sent are treated as opaque UDP payloads by the Joiner Routers and by Border Agents. Consequently, when Thread Commissioning Protocol is being used by Matter, the sequence diagrams differ slightly from those in the [Thread specification](#).

NOTE Support for Thread Commissioning Protocol is provisional.

The overall protocol is can be partitioned into several distinct phases:

- Becoming a Thread Commissioner (Matter Commissioner)
- Network Discovery and Selection
- Matter Discovery
- Message framing and secure session establishment
- Finalization

[Figure 43, “Advertising over Thread Commissioning Protocol”](#) shows the combined process for Thread Network Discovery and Selection, and for Matter Discovery.




Figure 43. Advertising over Thread Commissioning Protocol

4.23.1. Becoming a Thread Commissioner

In order to be able to discover the Matter Commissionees, the Matter Commissioner SHALL become the Thread Commissioner. This process is documented in the [Thread specification](#) (Section 8.4.1. and Section 8.4.2.):

- the Commissioner discovers the Thread network using mDNS,
- the Commissioner authenticates with the Thread network
- the Commissioner petitions the Thread network to become the Thread Commissioner

Once the Matter Commissioner has successfully become a Thread Commissioner, it SHALL set the Commissioner Dataset as outlined in [Section 4.23.1.1, “Thread Commissioning Protocol Steering Data”](#).

4.23.1.1. Thread Commissioning Protocol Steering Data

Steering Data SHALL be constructed using the discriminators from the onboarding payload. The [Thread specification](#) Section 8.10.1.9 specifies that mapping for IEEE EUI-64 identifiers. Matter commissioners SHALL construct an octet string of length 8, and use that octet string to populate the Bloom filter. The octet string SHALL be constructed as detailed in [Table 56, “Mapping the 12-bit Matter Discriminator”](#) and [Table 57, “Mapping the 4-bit Matter Discriminator”](#).

Table 56. Mapping the 12-bit Matter Discriminator

Byte 0	Byte 1	Byte 2	Byte 3	Byte 4	Byte 5	Byte 6	Byte 7
0x4D	0x54	0x00	0x00	0x00	0x00	Discriminator & 0xFF	(Discriminator >> 8) & 0x0F

Table 57. Mapping the 4-bit Matter Discriminator

Byte 0	Byte 1	Byte 2	Byte 3	Byte 4	Byte 5	Byte 6	Byte 7
0x4D	0x54	0x00	0x00	0x00	0x00	0x00	Discriminator & 0x0F

4.23.2. Network Discovery and Selection

4.23.2.1. Device Role

Commissionable devices SHALL implement Joiner Protocol as specified in [Thread specification](#) sections 8.4.4.1 through 8.4.4.4. The sections below detail the enhancements and clarifications against the Thread Joiner protocol.

4.23.2.2. MAC Extended Address

Commissionable device SHALL randomly generate its MAC extended address for its Thread network interface used for commissioning. This address SHOULD be obtained through [Crypto_DRBG\(len=64\)](#). It SHALL be generated on power on and SHALL be durable until commissioning is fully completed. To be more specific, for concurrent connection commissioning, the MAC extended address SHALL NOT change even after switching to the operational network. For non-concurrent connection commissioning, the MAC extended address MAY change after the commissioning connection is closed.

Note that Matter commissioning does not use the concept of Thread Joiner ID, which is a fixed identifier derived from the node's EUI-64 for a Thread commissioner to do matching. Matter commissioner does the matching in the upper layer. This avoids exposing that fixed identifier.

4.23.2.3. Joiner Router Selection

The Commissionee sends the Discovery Request ([Thread specification](#), Section 8.4.4.1.1.1) with the Destination PAN ID in the IEEE 802.15.4 MAC header set to the Broadcast PAN ID (0xFFFF) and the Source PAN ID set to a new randomly generated value. Using a specific Destination PAN ID in this context SHALL NOT be permitted. The Discovery Request SHALL omit the Discovery Request TLVs.

In processing the Discovery Response ([Thread specification](#), Section 8.4.4.1.1.2), the Commissionee SHALL only attempt to advertise on networks that allow joining. The Commissionee SHALL attempt to join on networks where the Steering Data in the Discovery Response matches (in order of priority):

1. the 12-bit discriminator from the onboarding payload

2. the upper 4 bits of the discriminator for cases where the Manual Pairing Code is being used by the commissioner
3. all joiners are permitted to join

The mapping of the discriminator to the bloom filter is described in [Section 4.23.1.1, “Thread Commissioning Protocol Steering Data”](#).

When the entire 12-bit discriminator matches the bloom filter in the Commissioner Dataset and the bloom filter is not set to all ones, the Commissionee SHALL stay provisionally joined to the Joiner Router for a period of not less than 3 minutes.

When only the 4-bit discriminator matches the bloom filter in the Commissioner Dataset, and the bloom filter is not set to all ones, the Commissionee SHALL stay provisionally joined to the Joiner Router for a period of not less than 30 seconds.

When the bloom filter in the Commissioner Dataset is set to all 1s, the Commissionee SHALL stay provisionally joined to the Joiner Router for a period of no less than 3750 milliseconds.

On timeout of provisional joining, the Commissionee SHALL attempt to discover Thread networks on all channels permitted by the regulatory domain by sending the Discovery Request on each channel at least twice. Given the relatively low frequency agility of the Thread networks, the Commissionee MAY build up a shorter list of channels with active Thread networks, and limit the subsequent Discovery Requests only to those channels.

The Commissionee SHALL attempt to join different candidate Thread networks for a period of no less than 15 minutes.

4.23.3. Matter Discovery

[Figure 43, “Advertising over Thread Commissioning Protocol”](#) shows the sequence diagram for Matter Discovery.

4.23.3.1. Advertisement from Commissionee

Upon selecting a Joiner Router, the Commissionee SHALL use a DNS Announcement to make itself discoverable. The Commissionee SHALL send an unsolicited DNS Query Response message with Transaction ID 0. This message SHALL contain information as specified in [Section 4.3.1, “Commissionable Node Discovery”](#). The DNS Announcement SHALL be sent in-line with the restrictions imposed by the Thread commissioning protocol, i.e., it shall be sent to the link-local address of the Joiner Router, with the destination UDP port MJ as specified in the Discovery Response message. No response to this message is expected from the Commissioner.

After sending the DNS Announcement, the Commissionee SHALL await an incoming commissioning message (see [Section 4.14.1, “Passcode-Authenticated Session Establishment \(PASE\)”](#)) for no less than 5000 milliseconds. If the commissioning session does not begin to be established within that time window, the device SHOULD retransmit the DNS Announcement. The device SHOULD retransmit the DNS Announcement up to 3 times with a 1250 millisecond interval between retransmissions.

If after the retransmissions the commissioning session still does not begin to be established, the

device SHOULD attempt to find a different Joiner Router using the procedure described in the subsection above.

NOTE

The timeouts and retransmission counts selected above are subject to fine-tuning. The 1250 millisecond interval is chosen to align with typical network discovery retransmission patterns. Per [Thread specification](#), the Commissionee will spend at least 300ms on each channel.

4.23.3.2. Discovery by Commissioner

The Commissioner receives the DNS Announcement encapsulated within the Relay notification (**RLY_RX.ntf**). If the contents of the DNS Announcement match the intents of the Commissioner, the Commissioner SHALL proceed to establish the secure session. If the contents of the DNS Announcement do not match the intents of the Commissioner, the Commissioner SHALL ignore the DNS Announcement and SHALL NOT initiate commissioning with that Commissionee. Since the DNS Announcement is an unsolicited response with Transaction ID 0, the Commissioner SHALL NOT send any DNS response back to the Commissionee.

4.23.4. Message framing and secure session establishment

[Figure 44](#), “[Example message exchange: initial PASE messages](#)” shows the sequence diagram and encapsulation of example Matter messages over the Thread Commissioning Protocol.




Figure 44. Example message exchange: initial PASE messages

From the commissionee's point of view, message framing over Thread Commissioning Protocol is the same as message framing over UDP/MRP; this ensures that the protocol can function within the constraints of the existing [Thread specification](#). Each outgoing message SHALL be sized so as to fit in the Joiner DTLS Encapsulation TLV of the RLY_RX.ntf message. The Commissionee SHALL use the MRP state machine to provide reliable delivery.

From the commissioner's point of view, each message is communicated as a payload of a CoAP message: it is received within the RLY_RX.ntf message and transmitted within the RLY_TX.ntf message. The respective CoAP messages are transmitted without reliability; the commissioner SHALL apply the MRP protocol within those messages to provide reliable delivery. Each outgoing message SHALL be less than TBD bytes so as to fit in the Joiner DTLS Encapsulation TLV of the RLY_TX.ntf message.

4.23.4.1. Address and Port Assignments

The Thread commissioning protocol provides the Thread Joiner with a link local address and an

ability to send UDP packets addressed to the link local address of the Joiner Router at the Commissioner UDP Port TLV as communicated in the DiscoveryResponse. Consequently, the Commissionee SHALL NOT be able to perform any commissioner discovery.

Commissionee SHALL accept incoming messages on the UDP port it chose to advertise in the DNS Announcement as described in [Figure 43, “Advertising over Thread Commissioning Protocol”](#) and in [Section 4.3.1, “Commissionable Node Discovery”](#).

Following the Matter Discovery phase, the Commissioner SHALL address the outgoing messages to the port indicated by the Commissionee by appropriately setting the value of Joiner UDP Port TLV in the outgoing RLY_TX.ntf message.

Chapter 5. Commissioning

5.1. Onboarding Payload

The purpose of this section is to define the contents of the Onboarding Payload needed to allow commissioning a device into a Matter network. It also specifies the representation and encoding of said payload as a QR Code, as a manually entered code, and as content in an NFC tag.

5.1.1. Onboarding Payload Contents

The Onboarding Payload is composed of required and optional information which will be used by the Commissioner to ensure interoperability between commissioners and devices and provide a consistent user experience. Some or all of this content will be encoded into different formats, some human-readable (such as numeric string) and machine-readable (such as QR code and NFC Tag) formats for printing or display on or integration into the device. The following are the elements that may be used in an Onboarding Payload for a Matter device.

5.1.1.1. Version

A version indication provides versioning of the payload and SHALL be included. Version for machine-readable formats is 3 bits with an initial version of **0b000**. Version for Manual Pairing Code is 1 bit with an initial version of **0b0**.

Rationale: This allows a way to introduce changes to the payload as needed going into the future.

5.1.1.2. Vendor ID and Product ID

Vendor ID and **Product ID**, each a 16-bit value, SHALL be included in machine-readable formats and MAY be included in the Manual Pairing Code.

Rationale: This allows a way to identify the make and model of the device, which is used further during the commissioning flow, such as during the [Device Attestation procedure](#). These unique identifiers also help to retrieve device model metadata like product name, product description, and firmware update URL from the [Distributed Compliance Ledger](#), as well as information relevant to the commissioning flow (see [Section 5.7, “Device Commissioning Flows”](#)).

5.1.1.3. Custom Flow

A 2-bit unsigned enumeration specifying the [Device Commissioning Flow](#) SHALL be included in machine-readable formats. For the encoding of Custom Flow in the Manual Pairing Code, see [Section 5.1.4.1.2, “Custom Flow for Manual Pairing Code”](#).

Rationale: This guides the Commissioner as to whether steps are needed before commissioning can take place.

- A value of **0** indicates that no steps are needed (apart from powering the device).
- A value of **1** indicates that user interaction with the device (pressing a button, for example) is required before commissioning can take place. The specific steps required can be found in the [CommissioningModeInitialStepsHint](#) field of the [Distributed Compliance Ledger](#) for the given Vendor ID and Product ID.

- A value of **2** indicates that additional interactions described in [Section 5.7.3, “Custom Commissioning Flow”](#) are required for initial device setup in order to be successfully commissioned by any Commissioner.

5.1.1.4. Discovery Capabilities Bitmask

An 8-bit capabilities bitmask SHALL be included in machine-readable formats.

Rationale: The Discovery Capabilities Bitmask (see [Table 60, “Discovery Capabilities Bitmask”](#)) contains information about the device’s available technologies for device discovery (see [Section 5.4, “Device Discovery”](#)).

5.1.1.5. Discriminator value

A Discriminator SHALL be included as a 12-bit unsigned integer, which SHALL match the value which a device advertises during commissioning. Manufacturers SHALL NOT use a common discriminator value for all devices, to allow a commissioner to distinguish between multiple advertising devices. It is RECOMMENDED that manufacturers use sequential generation with rollover for devices that might be sold or purchased in a multi-pack in order to reduce the chance of discriminator collision between devices in the same pack.

For machine-readable formats, the full 12-bit Discriminator is used. For the Manual Pairing Code, only the upper 4 bits out of the 12-bit Discriminator are used.

Rationale: The Discriminator value helps to further identify potential devices during the setup process and helps to improve the speed and robustness of the setup experience for the user.

5.1.1.6. Passcode

A **Passcode** SHALL be included as a 27-bit unsigned integer, which serves as proof of possession during commissioning. The 27-bit unsigned integer encodes an 8-digit decimal numeric value, and therefore SHALL be restricted to the values **0x0000001** to **0x5F5E0FE** (**00000001** to **99999998** in decimal), excluding the [invalid Passcode](#) values.

Rationale: The Passcode establishes proof of possession and is also used as the shared secret in setting up the initial secure channel over which further commissioning steps take place.

5.1.1.7. TLV Data

Variable-length TLV data using the TLV format MAY be included in machine-readable formats providing optional information. More details about the TLV can be found in [Section 5.1.5, “TLV Content”](#).

5.1.2. Onboarding Material Representation

In order for the users of Matter products to recognize the onboarding material, and be able to use it easily, it is important to keep the representations of the onboarding material unified and of certain minimum size. To support this the [Matter Brand Guidelines](#) specify the characteristics like composition, colors, font, font size, QR Code size and digit-grouping of the Manual Pairing code.

When the onboarding material is printed on product or packaging material it SHALL follow the [Matter Brand Guidelines](#).

Other representations (product display, app, etc) of the onboarding material SHOULD follow the [Matter Brand Guidelines](#).

5.1.3. QR Code

The Onboarding Payload MAY be included on (or with) a device in the form of a QR code. The following sections detail the content, encoding, and formatting of the QR code.

5.1.3.1. Payload

The content of the QR code consists of the concatenation of a prefix string and a Base-38-encoded string containing the required and optional TLV content:

QR code string := <prefix><base-38-content>

Prefix String

The prefix string consists of the three-character string:

MT:

Base-38 Content

The required content of the QR code is composed of a Packed Binary Data Structure containing elements of the Onboarding Payload as detailed below. The resulting data is Base-38 encoded (with a specific alphabet) to form a string compatible with alphanumeric QR encoding.

Packed Binary Data Structure

Individual data elements SHALL be placed into the structure in the order detailed in the table below. Elements being packed are not necessarily byte- or word-aligned. The resulting packed structure is presented to the encoder as a multi-byte array (see [Section 5.1.3.1.5, “Encoding”](#)), which SHALL be padded with '0' bits at the end of the structure to the nearest byte boundary.

The bits of each fixed-size value are placed in the packed binary data structure in order from least significant to most significant. If TLV Data is included, it is appended to the end of the packed binary data.

For example, the first elements in the table below SHALL be packed into the first bytes of the data array as pictured:

Table 58. Packing of the onboarding payload

lsb		Byte 0					msb		Byte 1					Byte 2					Byte 3					...		
0							7	0							7	0							7	0		
version			Vendor ID												Product ID											
0		2	0											15	0											15

Table 59. Packed Binary Data Structure for Onboarding Payload

Onboarding Payload Element	Size (bits)	Required	Notes
Version	3	Yes	3-bit value specifying the QR code payload version. SHALL be 000.
Vendor ID	16	Yes	
Product ID	16	Yes	
Custom Flow	2	Yes	Device Commissioning Flow 0: Standard commissioning flow: such a device, when uncommissioned, always enters commissioning mode upon power-up, subject to the rules in Section 5.4.2.2, “Announcement Commencement”. 1: User-intent commissioning flow: user action required to enter commissioning mode. 2: Custom commissioning flow: interaction with a vendor-specified means is needed before commissioning. 3: Reserved
Discovery Capabilities Bitmask	8	Yes	Defined in table below.
Discriminator	12	Yes	12-bit as defined in Discriminator
Passcode	27	Yes	See Section 5.1.7, “Generation of the Passcode”
Padding	4	Yes	Bit-padding of '0's to expand to the nearest byte boundary, thus byte-aligning any TLV Data that follows.
TLV Data	Variable	No	Variable length TLV data. Zero length if TLV is not included. This data is byte-aligned. See Section 5.1.3.1.4, “TLV Data” for detail.

Table 60. Discovery Capabilities Bitmask

Bit	Size (bits)	Description
0 (lsb)	1	Reserved (SHALL be 0)
1	1	BLE: 0: Device does not support BLE for discovery or is currently commissioned into one or more fabrics. 1: Device supports BLE for discovery when not commissioned.
2	1	On IP network: 1: Device is already on the IP network

Bit	Size (bits)	Description
3	1	Wi-Fi Public Action Frame: 0: Device does not support Wi-Fi Public Action Frame for discovery or is currently commissioned into one or more fabrics. 1: Device supports Wi-Fi Public Action Frame for discovery when not commissioned.
4	1	NFC Transport Layer: 0: Device does not support NFC Transport Layer (NTL) for commissioning or is currently commissioned into one or more fabrics. 1: Device supports NFC Transport Layer (NTL) for commissioning when not commissioned.
5	1	Thread Commissioning Protocol: 0: Device does not support Thread Commissioning Protocol for commissioning or is currently commissioned into one or more fabrics. 1: Device supports Thread Commissioning Protocol for commissioning when not commissioned.
7..6	2	Reserved (SHALL be 0)

NOTE

Support for Thread Commissioning Protocol is provisional. Until it becomes certifiable, Bit 5 should be treated as reserved.

TLV Data

The TLV data is an optional, variable-length payload. The payload is composed of one or more TLV-encoded elements as defined in detail below in the [TLV Content](#) section.

Encoding

The Packed Binary Data Structure is Base-38 encoded (with a specific alphabet) to produce an alphanumeric string suitable for use as a QR code payload.

Alphabet

The Base-38 alphabet to be employed is composed of a subset of the 45 available characters ([A-Z0-9\\$%*+./ :-](#)) in the QR code for alphanumeric encoding as defined by [ISO/IEC 18004:2015](#), with characters [\\$, %, *, +, /, <space>](#), and [:](#) removed.

Table 61. Alphabet for Onboard Payload Encoding

Code	Character	Code	Character	Code	Character	Code	Character	Code	Character
00	0	09	9	18	I	27	R	36	-
01	1	10	A	19	J	28	S	37	.
02	2	11	B	20	K	29	T		
03	3	12	C	21	L	30	U		
04	4	13	D	22	M	31	V		

05	5	14	E	23	N	32	W		
06	6	15	F	24	O	33	X		
07	7	16	G	25	P	34	Y		
08	8	17	H	26	Q	35	Z		

Method

Base-38 encoding is achieved by employing a simplified strategy where every 3 bytes (24 bits) of binary source data are encoded to 5 characters of the Base-38 alphabet.

Data from the Packed Binary Data Structure are encoded starting with the first byte of the structure. Three-byte chunks are formed into a 24-bit unsigned integer for encoding as follows:

$$\text{UINT}_{24} = (\text{BYTE}_{N+2} \ll 16) \mid (\text{BYTE}_{N+1} \ll 8) \mid (\text{BYTE}_N \ll 0)$$

The 24-bit value is subsequently converted to Base-38 radix using the alphabet above to produce a 5-character substring, with the least-significant character appearing first (little-endian).

If a single byte of binary source data remains, it SHALL be converted to Base-38 radix using the alphabet above to produce a 2-character substring, with the least-significant character appearing first.

If two bytes of binary source data remains, they SHALL be formed into a 16-bit unsigned integer for encoding as follows:

$$\text{UINT}_{16} = (\text{BYTE}_{N+1} \ll 8) \mid (\text{BYTE}_N \ll 0)$$

This 16-bit value is subsequently converted to Base-38 radix using the alphabet above to produce a 4-character substring, with the least-significant character appearing first.

The final encoded string is a result of concatenation of all substrings, with the first-encoded substring appearing at the beginning of the concatenated string.

5.1.3.2. QR Code Format

The format selection, which includes the QR code Version and ECC levels as well as size and color, MAY be tailored to the requirements of the manufacturer and their respective product, provided it meets the following requirements:

QR code Version and Encoding

The QR code generated, as defined in [ISO/IEC 18004:2015](#), SHALL be of Version 1 or higher, using alphanumeric encoding. The size of the payload implies a minimum Version, though a higher Version may be needed to allow a higher ECC level. For example, a minimum payload of 22 alphanumeric characters (19 base-38-encoded characters from the packed binary structure plus 3 prefix characters) can be fit into a Version 1 with ECC=L, but for ECC=M, Q or H, the same payload

requires a Version 2 QR code. This allows the Manufacturer to balance between ECC, pixel size and overall size.

Example QR Code Sizes and Payloads

QR Code Version	Module Size	ECC Level	Alphanumeric capacity (chars)	Total available payload, excluding prefix (bits)	Available payload for TLV data (bits)
1	21x21	L	25	104	16
2	25x25	L	47	208	120
		M	38	168	80
		Q	29	120	32
3	29x29	L	77	352	264
		M	61	272	184
		Q	47	208	120
		H	35	152	64
4	33x33	L	114	528	440
		M	90	416	328
		Q	67	304	216
		H	50	224	136
5	37x37	L	154	720	632
		M	122	568	480
		Q	87	400	312
		H	64	288	200

NOTE

Version 1 codes with ECC levels M, Q, and H and version 2 codes with ECC level H have insufficient capacity

NOTE

Total available payload, excluding prefix = $(\text{trunc}((N-3) / 5) * 24)$ where **N** is the number of alphanumeric characters which fit in the QR code. This formula uses **N-3** to account for the prefix characters, and then determines how many groups of 5 base-38-encoded characters can fit; each such group carrying 24 bits of payload. This formula fills groups of 5 characters after the **MT:** prefix. If there are 2,3 or 4 characters left after these groups, an additional 8 bits (for 2,3 characters) or 16 bits (for 4 characters) of TLV data can be accommodated. So the entries in the table take this into account.

Available payload for TLV data = $(\text{Total available payload, excluding prefix} - 88)$ since the minimum payload for the Packed Binary Data Structure is 84 bits before padding, or 88 bits with padding.

Payload Limits

The number of alphanumeric characters in the QR code to represent the Onboarding Payload for a single product SHALL NOT exceed 255 characters. Using this maximum size would yield a total available payload of 1208 bits, and hence 1120 bits (140 octets) of TLV payload.

NOTE

These size limits refer to the length of the QR code for a single product. In case of [concatenation](#), the total size of the concatenated QR code can exceed this size, but still needs to fit within a valid combination of Version and Quality for the total length of characters needed.

For QR codes, [ISO/IEC 18004:2015](#) currently allows for a maximum payload size of 4296 characters using Level 40 and L ECC Level encoding.

For NFC Tags, the actual maximum size is implementation dependent.

ECC Level

The QR code SHOULD employ level M or higher ECC.

NOTE

A higher level ECC does not help against typical 'reading' issues like shiny surfaces, bad contrast or issues with camera resolution/focus, and lack of camera-app processing dedicated for QR codes. Therefore, in certain situations ECC=L MAY be used as well (e.g. to prevent having to move to a higher Version to fit the payload).

5.1.4. Manual Pairing Code

This section describes the content and format of the Manual Pairing Code, which can be used in certain situations next to, or instead of, the QR code or NFC Tag described above.

5.1.4.1. Content

Payload

The payload of the Manual Pairing Code consists of the following required and optional data elements.

Table 62. Manual Pairing Code Elements

Element	Size (bits)	Required	Notes
VERSION	1	Yes	Shall be 0 Version is encoded as part of first digit of the Manual Pairing Code. A value of 1 is reserved for future extension of the specification.
VID_PID_PRESENT	1	Yes	0: no Vendor ID and Product ID present in Manual Pairing Code 1: Vendor ID and Product ID data included
DISCRIMINATOR	4	Yes	4 Most-significant bits of the 12-bits Discriminator described above

Element	Size (bits)	Required	Notes
PASSCODE	27	Yes	Same as 27-bit Passcode described above
VENDOR_ID	16	No	Needed only to support devices that need a user-intent or vendor specific flow before commissioning (i.e. a non-zero Custom Flow value). If an accompanying QR code is present on the device with the Custom Flow field set to a non-zero value, or if the device requires Custom commissioning flow , this element SHALL be included.
PRODUCT_ID	16	No*	* This element SHALL be included <i>if and only if</i> the VENDOR_ID element is present.

The Vendor ID and Product ID elements are optional. Including these may provide additional information for the setup flow at the expense of a substantially longer Manual Pairing Code.

Custom Flow for Manual Pairing Code

The encoding for Manual Pairing Code does not have a dedicated field for Custom Flow, as exists in the [Packed Binary Data Structure](#). Instead, this information is encoded in the following way:

- For [Standard commissioning flow](#), the variant of Manual Pairing Code *without* Vendor ID and Product ID SHALL be used. A commissioner encountering such Manual Pairing Code SHALL assume it is a "standard flow" device.
- For [User-intent commissioning flow](#) and [Custom Commissioning flow](#), the variant of Manual Pairing Code *with* Vendor ID and Product ID SHALL be used. For this case, a commissioner SHOULD use Vendor ID and Product ID to lookup the [CommissioningCustomFlow](#) field in the [Distributed Compliance Ledger](#) to determine which of these values applies for this Vendor ID and Product ID combination.

Encoding

The required and optional elements, along with a check digit, are encoded into either an 11-digit or 21-digit decimal numeric string, depending on whether the optional Vendor and Product ID information is included.

Method

Each group of digits in the Pairing Code SHALL be encoded as described in the table below. The left-most digit of the entire string SHALL be represented by *DIGIT[1]*. Groups of multiple digits SHALL be encoded such that the most-significant digit appears first (left-most).

Table 63. Encoding Method without Vendor and Product ID's (VID_PID_Present == 0)

Digit	Contents	Encoding	Notes
1 (left-most)	- Version 0 - VID_PID present flag - 2 ms-bits of discriminator	DIGIT[1] := (VID_PID_PRESENT << 2) (DISCRIMINATOR >> 10)	Allows first digit typed/spoken to determine version and VID/PID present. Yields a decimal number from 0..7 (0..3 if VID,PID not present). First digit of '8' or '9' would be invalid for v1 and would indicate new format (e.g. version 2)
2..6	- 3rd and 4th ms-bits of Discriminator - 14 ls-bits of PASSCODE	DIGIT[2..6] := ((DISCRIMINATOR & 0x300) << 6) (PASSCODE & 0x3FFF)	Yields a 5-digit decimal number from 00000 to 65535 (0xFFFF/16 bits)
7..10	- 13 ms-bits of PASSCODE	DIGIT[7..10] := (PASSCODE >> 14)	Yields a 4-digit decimal number from 0000 to 8191 (0x1FFF/13 bits)
11	- Check Digit	DIGIT[11] := (CHECK_DIGIT)	See Section 5.1.4.1.5, “Check Digit” for encoding

Table 64. Encoding Method with Vendor and Product ID's included (VID_PID_Present == 1)

Digit	Contents	Encoding	Notes
1 (left-most)	- Version 0 - VID_PID present flag - 2 ms-bits of Discriminator	DIGIT[1] := (VID_PID_PRESENT << 2) (DISCRIMINATOR >> 10)	Allows first digit typed/spoken to determine version and VID/PID present. Yields a decimal number from 0..7 (4..7 if VID,PID present). First digit of '8' or '9' would be invalid for v1 and would indicate new format (e.g. version 2)
2..6	- 3rd and 4th ms-bits of Discriminator - 14 ls-bits of PASSCODE	DIGIT[2..6] := ((DISCRIMINATOR & 0x300) << 6) (PASSCODE & 0x3FFF)	Yields a 5-digit decimal number from 00000 to 65535 (0xFFFF/16 bits)
7..10	- 13 ms-bits of PASSCODE	DIGIT[7..10] := (PASSCODE >> 14)	Yields a 4-digit decimal number from 0000 to 8191 (0x1FFF/13 bits)
11..15	- Vendor ID	DIGIT[11..15] := (VENDOR_ID)	Yields a 5-digit decimal number from 00000 to 65535 (0xFFFF/16 bits)

Digit	Contents	Encoding	Notes
16..20	- Product ID	DIGIT[16..20] := (PRODUCT_ID)	Yields a 5-digit decimal number from 00000 to 65535 (0xFFFF/16 bits)
21	- Check Digit	DIGIT[21] := (CHECK_DIGIT)	See Section 5.1.4.1.5, “Check Digit” for encoding

Check Digit

The *CHECK_DIGIT* element is a single decimal digit computed across all of the preceding digits of the Pairing Code using the [Verhoeff](#) algorithm.

5.1.4.2. Copying between applications

When the Manual Pairing Code is presented in an application within a multi-function device, such as an application on a smartphone, it SHOULD provide a mechanism such as a copy button to allow easy conveyance of the information to other commissioners on the same device.

When a Commissioner is implemented as an application within a multi-function device, such as an application on a smartphone, it SHOULD provide a mechanism such as a paste button to allow easy conveyance of the information from an administrator on the same device.

The receiving application SHALL be robust against characters like dashes and spaces that may be included in the string, such as those added for readability by the user, when passed between the applications. For example, a receiving application seeing the code "1234-567-8910" would need to interpret it as "12345678910".

5.1.5. TLV Content

A variable-length [TLV Data](#) section MAY be encoded into the Packed Binary Data Structure. The TLV section MAY consist of manufacturer-specific information elements and/or elements common to Matter, encoded using TLV. All elements SHALL be housed within an anonymous top-level structure container.

5.1.5.1. Manufacturer-specific Elements

Manufacturer-specific elements SHALL be tagged with context-specific tags that have semantics which are defined by the vendor for use in the products using their Vendor ID, and SHALL use tag numbers 0x80 to 0xFF.

Tag numbers 0x00 to 0x7F are reserved to indicate Matter-common elements.

Manufacturer-specific elements inherit the context of the Vendor ID and Product ID provided in the Packed Binary Data Structure described above. All elements SHALL follow the constraints outlined in [Appendix A, Tag-length-value \(TLV\) Encoding Format](#).

5.1.5.2. Matter-common Elements

All elements common to Matter SHALL use tag numbers in the range 0x00 to 0x7F, as defined in the following section.

Vendors are encouraged to use Matter-common elements where applicable.

Table 65. Matter-common Reserved Tags

Tag	Value	Description	Type(s)
<i>kTag_Serial-Number</i>	0x00	Device Serial #	UTF-8 String (length = 1..32 bytes) Unsigned Integer, up to 8-byte value (has room to represent a 19-digit decimal number)
<i>PBKDFIterations</i> *	0x01	PBKDFParameterSet Iterations	Unsigned Integer (range = CRYPTO_PBKDF_ITERATIONS_MIN.. CRYPTO_PBKDF_ITERATIONS_MAX)
<i>PBKDFSalt</i> *	0x02	PBKDFParameterSet Salt	Octet String (length = 16..32 bytes)
<i>kTag_NumberOfDevices</i>	0x03	Number of devices that are expected to be commissioned using this payload when using the Enhanced Commissioning Method	Unsigned Integer, range 1 to 255
<i>kTag_CommissioningTime-out</i>	0x04	Time, in seconds, during which the device(s) are expected to be commissionable using the Enhanced Commissioning Method	Unsigned Integer, see Section 5.4.2.3, “Announcement Duration”
reserved	0x05.. 0x7F	reserved for future use	

NOTE

PQC Phase 1 will likely soon replace this *AttestationCryptoInfo* mechanism in the Onboarding Payload with an equivalent mechanism using the Operational Credentials Cluster.

* If the PBKDF parameters are to be included in the TLV section, both the **PBKDFSalt** and **PBKDFIterations** SHALL be encoded.

5.1.5.3. TLV Examples

Manufacturer-specific and Matter-common elements

```
{
  vendorTag01 (0x81) = "Vendor",
  kTag_SerialNumber(0) = "1234567890"
}
```

The above notation encodes to the following bytes:

0x15 0x2C 0x81 0x06 0x56 0x65 0x6E 0x64 0x6F 0x72 0x2C 0x00 0x0A 0x31 0x32 0x33
0x34 0x35 0x36 0x37 0x38 0x39 0x30 0x18

Data	Comments
=====	=====
0x15	Control Byte for outermost container (structure) - Tag control 000xxxxxb: Anonymous tag - Elem type xxx10101b: Structure
0x2C	Control Byte for next TLV - Tag control 001xxxxxb: Context-specific tag - Elem type xxx01100b: UTF-8 String, 1-byte length
0x81	Context-specific vendor tag - Matter-common versus vendor tag 1xxxxxxb: Vendor tag - Tag number x0000001b: Vendor tag #1 ----- 10000001b = 0x81
0x06	Length of vendor string (e.g. 6 bytes)
0x56 0x65 0x6E 0x64 0x6F 0x72	UTF-8 encoded vendor string "Vendor"
0x2C	Control byte for next TLV - Tag control 001xxxxxb: Context-specific tag - Elem type xxx01100b: UTF-8 String, 1-byte length ----- 00101100b = 0x2C
0x00	Context-specific Matter-common Serial Number tag - Matter-common versus vendor tag 0xxxxxxb: Matter-common tag - Tag number x0000000b: kTag_SerialNumber ----- 00000000b = 0x00
0x0A	Length of Serial Number string (10 bytes)
0x31 0x32 0x33 0x34 0x35 0x36 0x37 0x38 0x39 0x30	UTF-8 encoded Serial Number string "1234567890"
0x18	End of container

Matter-common elements only

```
{
    kTag_SerialNumber (0) = "1234567890"
```

```
}
```

The above notation encodes to the following bytes:

0x15 0x2C 0x00 0x0A 0x31 0x32 0x33 0x34 0x35 0x36 0x37 0x38 0x39 0x30 0x18	
Data	Comments
=====	
0x15	Control Byte for outermost container (structure) - Tag control 000xxxxb: Anonymous tag - Elem type xxx10101b: Structure
0x2C	Control Byte for next TLV - Tag control 001xxxxb: Context-specific tag - Elem type xxx01100b: UTF-8 String, 1-byte length ----- 00101100b = 0x2C
0x00	Context-specific Matter-common Serial Number tag - Matter-common versus vendor tag 0xxxxxxb: Matter-common tag - Tag number x0000000b: kTag_SerialNumber ----- 00000000b = 0x00
0x0A	Length of Serial Number string (10 bytes)
0x31 0x32 0x33 0x34 0x35 0x36 0x37 0x38 0x39 0x30	UTF-8 encoded Serial Number string "1234567890"
0x18	End of container

5.1.6. Concatenation

The Onboarding Payload MAY be concatenated with additional Onboarding Payloads to be placed in a single QR Code or NFC Tag:

```
QR code string := MT:<onboarding-base-38-content>*<onboarding-base-38-content2>
```

Where ***** is used as the delimiter.

Concatenation of multiple Matter Onboarding Payloads allows a single QR code (or NFC Tag) to provide the onboarding payload for a number of devices. Each <onboarding-base-38-content> represents the Onboarding Payload for a specific device, and can include its own TLV section (for the TLV content specific to that device). Each <onboarding-base-38-content> can have a length as described in [Section 5.1.3.2.3, “Payload Limits”](#), so the overall length of the resulting concatenated QR code can be larger than such limits.

Example use case for this concatenation:

- Easy commissioning for multi-device packaging, e.g. for a package of light bulbs containing four separate bulbs. Each bulb will have its own Onboarding Payload code(s) printed on the bulb itself. The Manufacturer MAY include a leaflet in the box with a larger QR code containing the concatenation of the four individual Onboarding Payloads. The user can then scan this combined QR code (one step for the user) which would give the Commissioner the Onboarding Payload for all four bulbs in one operation, and it can proceed to commission the four bulbs.

All Commissioners SHALL recognize the * separator from the QR code as indication concatenation is used.

A Commissioner which does not support such concatenated Matter Onboarding Payloads SHOULD indicate to the user the need to commission devices one by one by scanning their individual QR codes or entering the individual Manual Pairing Codes.

The Commissioner SHOULD commission the devices in the order as they are provided in the concatenated code. (This ordering is particularly relevant in case of combi-packs where one of the devices needs to be commissioned first, e.g. a Thread Router first, then one or more Thread-connected bulbs).

Example of concatenated Onboarding Payloads:

```
MT:<onboarding-base-38-content_bulb1>*<onboarding-base-38-content_bulb2>*<onboarding-  
base-38-content_bulb3>*<onboarding-base-38-content_bulb4>
```

5.1.7. Generation of the Passcode

A device can support either dynamic or static passcodes for purposes of establishing the shared secret for the initial secure channel over which further commissioning steps take place.

All devices SHALL conform to the following rules for passcodes:

- Passcodes SHALL NOT be derived from public information, such as a serial number, manufacturer date, MAC address, region of origin, etc.
- The Passcode generation process SHALL use a cryptographically secure random number generator.

If a device generates a dynamic Passcode, then it SHALL conform to the following additional rule:

- Passcodes SHALL be accessible to commissioner only during the commissioning process.

If a device cannot generate a dynamic Passcode, then the static Passcode SHALL conform to the following additional rules:

- A random passcode with 27 bits of entropy SHALL be generated and used for each individual device. Because of the disallowed numbers, the entropy remaining in the actual passcode will be somewhere between 26 bits and 27 bits but the initial value before rejecting disallowed numbers SHALL have 27 bits of entropy.
- The device SHALL be supplied with the [PAKE](#) verifier in its internal storage.

- If the static passcode is also supplied to the device, the static passcode SHALL NOT be accessible during operational mode using any data model attributes or commands.
- If the static passcode is supplied to the device, its storage location SHALL be physically or logically isolated from the location where the [PAKE](#) verifier is stored and SHALL NOT be accessible to the executing unit handling the [PAKE](#) verifier.

5.1.7.1. Invalid Passcodes

The following Passcodes SHALL NOT be used for the [PASE](#) protocol due to their trivial, insecure nature:

- 00000000
- 11111111
- 22222222
- 33333333
- 44444444
- 55555555
- 66666666
- 77777777
- 88888888
- 99999999
- 12345678
- 87654321

5.1.8. NFC Tag

A Commissioner MAY use, in addition to the [QR Code Format](#) and [Manual Pairing Code](#) as described above, an NFC tag associated with a Commissionee to retrieve the Onboarding Payload. When an NFC tag is used the following requirements are applicable.

- The data contained in the NFC tag SHALL be formatted as specified in [QR Code Format](#).
- The NFC tag SHALL be one of the types as defined by [NFC Forum Devices Requirements](#).
- The NFC tag SHALL use the NFC Data Exchange Format (NDEF) as defined by [NFC NDEF 1.0](#).
- The NFC tag SHALL use NDEF messages as defined by [NFC RTD 1.0](#).
- The Onboarding Payload for the NFC tag SHALL use NDEF URI Record Type Definition as defined by [NFC RTD URI 1.0](#) and as specified in the following table.

Table 66. NFC NDEF Representation

Offset	Content	Description
0	0xD1	TNF=0x01, SR=1, MB=1, ME=1
1	0x01	Length of Record Type

Offset	Content	Description
2	URI payload size in bytes	Length of payload
3	0x55	Record Name ("U")
4	0x00	URI Identifier Code: No URI abbreviation
5	URI data	MT:<base-38-content>

5.2. Initiating Commissioning

5.2.1. Purpose and Scope

The process of Matter commissioning can be initiated by the User in a number of ways. This section describes different user journeys supported by Matter. For each, a rationale is provided along with a high-level flow description, up until the point where a commissioning secure session is established. References to sections describing dependent functionality in more detail are provided.

The purpose of this section is to connect features provided in other sections to the user journeys for which they are designed.

WARNING

The list of user journeys provided here is not meant to be exhaustive; there may be other journeys not listed here which can be realized using Matter.

This section provides rationales for Matter functionality and does NOT contain normative requirements for Matter.

The following User Journeys are described in this section:

- [Section 5.2.2.1, “Commissioner Onboarding Material Entry, Not Yet Commissioned Device”](#). "Launch Commissioner, Enter Code"
- [Section 5.2.2.2, “User-Initiated Beacon Detection, Not Yet Commissioned Device”](#). "Launch Commissioner, Discover New Devices"
- [Section 5.2.2.3, “User-Initiated Beacon Detection, Already Commissioned Device”](#). "Launch Commissioner, Discover My devices"
- [Section 5.2.2.4, “Commissioner Discovery, from an On-Network Device”](#). "Launch Device User Interface, Discover Commissioners"

5.2.2. User Journey Details

5.2.2.1. Commissioner Onboarding Material Entry, Not Yet Commissioned Device

"Launch Commissioner, Enter Code"

In the Onboarding Material Entry for a Not Yet Commissioned Device use case, the User first initiates an interaction with a Commissioner, and then provides the necessary Onboarding Material from the Commissionee, by scanning an Onboarding Payload (e.g. QR Code or NFC Tag) or other-

wise inputting the Manual Pairing Code through an input method supported by the commissioner.

5.2.2.1.1. Rationale

In this use case, the user will often have the device in-hand, have immediate access to the onboarding payload, and have immediate access to the desired Commissioner.

5.2.2.1.2. High Level Flow

1. User initiates an interaction with a Commissioner.
2. User inputs the onboarding payload from the Commissionee.
3. Commissioner determines which technologies to use for [Device Discovery](#). When attempting to locate the device on IP-bearing networks, the [Commissionable Node Discovery](#) method is used and typically the DNS-SD service subtypes for long or short discriminator, and commissioning mode (see [Section 4.3.1.3, “Commissioning Subtypes”](#)) are specified to filter the results to Commissionees that match the discriminator in the onboarding payload and that are in Commissioning Mode. When attempting to locate the device via BLE advertisements, the discriminator will typically be used to filter the results. When using the NFC interface, the physical NFC tap performed by the user between the commissioner and commissionee allows device-to-device technology detection.
4. Commissioner begins the Commissioning process (see [Section 5.5, “Commissioning Flows”](#)). If more than one Commissionee is discovered, the Commissioner may further refine the results using any additional information such as a Vendor ID or Product ID that may be available in the onboarding payload. If there is still more than one discovered Commissionee, the Commissioner will typically attempt to establish a [PASE](#) secure commissioning session with each.

5.2.2.1.3. Misuse Considerations

When a device has a static onboarding payload, and the value is physically affixed to the product, it is possible for an attacker with one-time physical access to the device to obtain the onboarding payload and use it to compromise the security of the device in the future. For example, if the device is commissioned again using the same onboarding payload (for example, after a reset), then the attacker may be able to perform a person-in-the-middle attack which could result in a compromise of sensitive user data such as network credentials if passed to the device.

When a device includes device-specific information such as Vendor ID and Product ID in advertisements, then a malicious actor within advertisement range can detect this information and potentially associate it with the location of the device (and potentially, additional information about the location, such as its residents) in ways that the user did not intend.

5.2.2.2. User-Initiated Beacon Detection, Not Yet Commissioned Device

"Launch Commissioner, Discover New Devices"

In the User-Initiated Beacon Detection for a Not Yet Commissioned Device use case, the User first initiates an interaction with a Commissioner, and then indicates an intention to commission devices without providing additional information about them (no onboarding payload, etc).

5.2.2.2.1. Rationale

In this use case, the user has immediate access to a Commissioner. However, the user may not know how to locate the onboarding payload (it may be hidden behind a panel, pin-protected in a settings menu, or inaccessible on a device already physically installed).

Example User interactions with the Commissioner include pushing a "Discover New Devices" button, or speaking to a voice agent "Agent, discover new devices".

5.2.2.2.2. High Level Flow

1. User initiates an interaction with a Commissioner.
2. User indicates an intention to commission devices without providing additional information about them.
3. Commissioner determines which technologies to use for [Device Discovery](#). When attempting to locate the device on IP-bearing networks, the [Commissionable Node Discovery](#) method is used.
4. Commissioner constructs a list of Commissionees discovered, using as much information as possible from the Commissionee advertisement. When a Vendor ID and Product ID is provided in the advertisement, the Commissioner may obtain human readable descriptions of the Vendor and Product in order to assist the user with selection by using fields such as [ProductName](#) and [ProductLabel](#) from the [Distributed Compliance Ledger](#) or any other data set available to it. The ledger entry may also include additional URLs which the Commissioner can offer to the user to help in locating the Onboarding Material or otherwise assist in setting up the device such as the [UserManualUrl](#), [SupportUrl](#), and [ProductUrl](#). The Commissioner may have additional data sets available for assisting the user.
5. User selects Commissionee from list.
6. Commissioner instructs the user to locate and input the onboarding payload.
7. Commissioner begins the Commissioning process (see [Section 5.5, "Commissioning Flows"](#)).

5.2.2.2.3. Variation - Filter by Device Type

The user may indicate the type of device to the Commissioner when initiating this flow. For example, the user might speak the following to a voice agent: "Agent, Discover TVs".

When discovering TVs or any other specific device type on the IP network, this flow will be the same except that a subtype which specifies the device type identifier (see [Section 7.16.5, "Primary Device"](#)) is passed to the [Commissionable Node Discovery](#) method (see [Section 4.3.1.3, "Commissioning Subtypes"](#)).

5.2.2.2.4. Misuse Considerations

In addition to the Misuse Considerations described in [Section 5.2.2.2, "User-Initiated Beacon Detection, Not Yet Commissioned Device"](#), a Commissioner which performs [Device Discovery](#) without knowledge of the Onboarding Payload may discover advertisements from devices that the user did not intend to commission with the given Commissioner. This additional information collected by the Commissioner can be associated with the user in ways that the user did not intend.

5.2.2.3. User-Initiated Beacon Detection, Already Commissioned Device

"Launch Commissioner, Discover My Devices"

In the User-Initiated Beacon Detection for an Already Commissioned Device use case, the User first initiates an interaction with a Commissioner, and then indicates an intention to commission devices already on the IP network without providing additional information about them.

5.2.2.3.1. Rationale

A Device may choose to be discoverable by entities on the local IP network, even when not in Commissioning Mode, in order to satisfy specific user journeys. For example, a TV, Network Infrastructure Manager, or Bridge device may choose to be discoverable in order to facilitate connectivity with other Smart Home systems. One way for a device to be discoverable when not in Commissioning Mode is to implement [Extended Discovery](#).

Example User interactions with the Commissioner include pushing a "Discover My Devices" button, or speaking to a voice agent "Agent, discover my devices".

5.2.2.3.2. High Level Flow

1. User initiates an interaction with a Commissioner.
2. User indicates an intention to commission existing devices on the IP network. This may or may not include additional information about them, such as a device type.
3. Commissioner performs [Commissionable Node Discovery](#) via DNS-SD. The subtype for device type (see [Section 4.3.1.3, "Commissioning Subtypes"](#)) can be specified in order to filter the results to Commissionees that have a specific primary function device type.
4. Commissioner constructs a list of Commissionees discovered, using as much information as possible from the Commissionee advertisement. When a Vendor ID and Product ID is provided (see [Section 4.3.1.6, "TXT key for Vendor ID and Product ID \(VP\)"](#)), the Commissioner may obtain human readable descriptions of the Vendor and Product in order to assist the user with selection by using fields such as [ProductName](#) and [ProductLabel](#) from the [Distributed Compliance Ledger](#) or any other data set available to it. The ledger entry may also include additional URLs which the Commissioner can offer to the user to help in locating the Onboarding Material or otherwise assist in setting up the device such as the [UserManualUrl](#), [SupportUrl](#), and [ProductUrl](#). The Commissioner may have additional data sets available for assisting the user. When the Device Type (see [Section 4.3.1.8, "TXT key for device type \(DT\)"](#)) and/or the Device Name (see [Section 4.3.1.9, "TXT key for device name \(DN\)"](#)) values are provided, then the Commissioner may provide this information to the user in order to assist with Commissionee selection.
5. User selects Commissionee from list.
6. The [Commissionable Node Discovery](#) DNS-SD TXT record for the selected Commissionee includes key/value pairs that can help the Commissioner to guide the user through the next steps of the commissioning process.
7. If the Commissioning Mode value (see [Section 4.3.1.7, "TXT key for commissioning mode \(CM\)"](#)) is set to 0, then the Commissionee is not yet in Commissioning Mode and the Commissioner can guide the user through the steps needed to put the Commissionee into Commissioning Mode. The Pairing Hint (see: [Commissioning Pairing Hint](#)) and the Pairing Instruction (see: [Commis-](#)

sioning Pairing Instruction) fields would then indicate the steps that can be followed by the user to put the device into Commissioning Mode. When the Pairing Hint is present and its value has bit 1 set (Device Manufacturer URL), then the Commissioner can utilize the URL specified in the `CommissioningCustomFlowUrl` of the `DeviceModel` schema entry indexed by the Vendor ID and Product ID in the `Distributed Compliance Ledger` and utilize flows described in `Custom Commissioning Flow` to redirect the user to a custom app or website specified by the device vendor, and receive the user back following the callback flow which contains the onboarding payload. The `Custom Commissioning Flow` can provide a deterministic user experience for putting the device into commissioning mode and transferring the onboarding payload to the Commissioner. The Commissioner then verifies the new Commissioning Mode state using `Commissionable Node Discovery`.

8. Commissioner instructs the user to locate and input the onboarding payload if needed. When a Vendor ID and Product ID is available to the Commissioner, the `Distributed Compliance Ledger` may also provide a URL for the Device User Guide which can contain additional information to help in locating the onboarding payload. The Commissioner may have additional data sets available for assisting the user.
9. Commissioner begins the Commissioning process (see [Section 5.5, “Commissioning Flows”](#)).

5.2.2.3.3. Variation - Filter by Device Type

A variation on this use case is when a user action causes the Commissioner to ask the user for permission to discover devices or to discover a specific type of device on the network, such as a TV for remote controlling playback or a Home Router Access Point for managing network credentials.

The user may indicate the type of device to the Commissioner when initiating this flow. For example, the user might speak the following to a voice agent: "Agent, Discover TVs".

Example User interactions which are tied to discovery of specific device types include:

1. When a user indicates their desire to use a specific Matter Admin/Commissioner, the Commissioner might attempt to discover devices of type Network Infrastructure Manager on the network, and prompt the user with instructions for connecting to it.
2. When a user indicates their desire to use a specific Matter Video Remote Control, the Video Remote Control might look for Video Players on the network, and prompt the user with instructions for connecting to it.

When discovering TVs or any other specific device type on the IP network, this flow will be the same except that a subtype which specifies the device type identifier is passed to the `Commissionable Node Discovery` method (see [Section 4.3.1.3, “Commissioning Subtypes”](#)).

5.2.2.3.4. Misuse Considerations

When a Device implements `Commissionable Node Discovery` while not in Commissioning Mode, the time period during which it may unintentionally provide information to a malicious actor on the network is longer than it otherwise would be. This additional information could potentially be associated with the user in ways that the user did not intend. See [Commissionable Node Discovery Privacy Considerations](#) for device requirements relating to this risk.

When a device includes device-specific information such as Vendor ID, Product ID and Device Type,

or user-generated data such as Device Name, in the DNS-SD TXT record, then a malicious actor on the network can detect this information and potentially associate it with the user in ways that the user did not intend.

A Commissioner which performs [Device Discovery](#) without knowledge of the Onboarding Payload may discover devices on the network that the user did not intend to commission with the given Commissioner. This additional information collected by the Commissioner can be associated with the user in ways that the user did not intend.

5.2.2.4. Commissioner Discovery, from an On-Network Device

"Launch Device User Interface, Discover Commissioners"

In the Commissioner Discovery use case for a Device already on the IP network, the User first initiates an interaction with the Device via a display or other user interface, and indicates the intention to have this device commissioned by a Commissioner on the network. The Device might already have been commissioned into one or more Fabrics or it might not yet have been commissioned. Upon this user interaction, the Device discovers candidate Commissioners and allows the user to select one. The Device then requests from that Commissioner to be commissioned.

5.2.2.4.1. Rationale

In this use case, a Device (Commissionee) with a user interface, such as a TV or Thermostat, initiates the commissioning process. For example, this might be done from within a settings menu for Smart Home control. The Device discovers Commissioners on the IP-bearing network, presents the resulting list to the User for selection. Once selected, the Device indicates to the selected Commissioner that it has been selected by the User, the Device enters Commissioning Mode and provides the onboarding payload to the User.

Another example for this use case is a Device or Node (Commissionee) with a user interface, such as a Content Provider Device or Application, that initiates the commissioning process. This might be done from a program guide or while watching a video when the user indicates a desire to play the selected content on a nearby device. The Device discovers Commissioners on the IP-bearing network, presents the resulting list to the User for selection. Once selected, the Commissionee indicates to the selected Commissioner that it has been selected by the User (see [Section 5.3, "User Directed Commissioning"](#)), the Commissionee enters Commissioning Mode and provides the onboarding payload to the User.

5.2.2.4.2. High Level Flow

1. User initiates an interaction with the Device.
2. User indicates a desire to connect this Device with a Commissioner on the network.
3. Device uses [Commissioner Discovery](#) over DNS-SD on the IP bearing network.
4. Device collects candidates from DNS-SD service records found.
5. Device displays list of Commissioners discovered, including as much information as possible from the DNS-SD TXT record. When a Vendor ID and Product ID is provided (see [Section 4.3.1.6, "TXT key for Vendor ID and Product ID \(VP\)"](#)), the Device may obtain human readable descriptions of the Vendor and Product in order to assist the user with selection by using fields such as

ProductName and **ProductLabel** from the [Distributed Compliance Ledger](#) or any other data set available to it. The Device may have additional data sets available for assisting the user. When the Device Type (see [Section 4.3.1.8, “TXT key for device type \(DT\)”](#)) and/or the Device Name (see [Section 4.3.1.9, “TXT key for device name \(DN\)”](#)) values are provided in the DNS-SD TXT record, then the Device may provide this information to the user in order to assist with Commissioner selection.

6. User selects an entry from the list.
7. Device enters Commissioning Mode.
8. Device displays onboarding payload to the user.
9. Device initiates a [Section 5.3, “User Directed Commissioning”](#) session with the selected Commissioner, which includes in the DNS-SD service name of the Device.
10. Commissioner prompts user to confirm intention to commission this device and asks for onboarding payload.
11. User enters onboarding payload into Commissioner UX.
12. Commissioner begins the commissioning process (see [Section 5.5, “Commissioning Flows”](#)).

For additional variations relating to the display and input of the onboarding payload, see [Section 5.3, “User Directed Commissioning”](#).

5.2.2.4.3. Misuse Considerations

In addition to the Misuse Considerations described in [Section 5.2.2.3, “User-Initiated Beacon Detection, Already Commissioned Device”](#), a Commissionee which performs [Commissioner Discovery](#) may discover Commissioners on the network that the user did not intend to be discovered by the given Commissionee. This additional information collected by the Commissionee can be associated with the user in ways that the user did not intend. See [Commissioner Discovery Privacy Considerations](#) for Commissioner requirements relating to this risk.

Since there are no trust mechanisms employed for Commissioners advertising themselves, Commissionees may provide Commissioner selection choices to the User that are from malicious entities masquerading as commissioners.

When a Commissioner includes device-specific information such as Vendor ID, Product ID and Device Type, or user-generated data such as Device Name, in the DNS-SD TXT record, then a malicious actor on the network can detect this information and potentially associate it with the user in ways that the user did not intend.

5.3. User Directed Commissioning

5.3.1. Overview

In User Directed Commissioning (UDC), the Commissionee sends a message to the Commissioner in order to initiate the commissioning process (see [Section 5.5, “Commissioning Flows”](#)). In addition, the Commissionee can specify its UX preferences for the commissioning session using optional parameters in the message. Similarly, the Commissioner can send a message to the Commissionee to indicate its pre-commissioning state in order to help the Commissionee simplify the experience

for the user.

The availability of the UDC protocol is advertised through Commissioner Discovery service records of DNS-SD service type `_matterd._udp` (see [Section 4.3.3, “Commissioner Discovery”](#)).

Overall, the UDC protocol is a lightweight "door bell" message sent by a Commissionee. There is no state or session associated with UDC protocol messages, allowing the recipient to choose to support messages received from different sources at the same time, or choose to ignore messages received while the processing of an earlier message has not completed. The UDC protocol consists of an **Identification Declaration** which provides the Commissionee's `_matterc._udp` DNS-SD service instance name. In addition, **Identification Declaration** MAY include optional information relating to the requested commissioning session.

Upon receiving this message, the Commissioner MAY query the DNS-SD service instance indicated in the **Identification Declaration**, including TXT records, in order to obtain additional information about the Commissionee, MAY obtain the corresponding Onboarding Payload or Passcode from the user for this Commissionee, and MAY initiate the commissioning process with it.

The Commissioner MAY send a **CommissionerDeclaration** to the **Identification Declaration** message sender's source IP and Port. The **CommissionerDeclaration** provides information to a Commissionee indicating the Commissioner's pre-commissioning state, which MAY be used by the Commissionee to simplify the Passcode entry user experience.

Upon receiving this message, the Commissionee MAY provide instructions to the user about how to proceed to the next steps in the commissioning process, for example, by displaying a Passcode to the user, and instructing them to provide it to the Commissioner.

While a sequence of these messages between Commissionee and Commissioner can be used to achieve a specific user experience, as the examples below illustrate, the processing of each message is expected to stand alone and not be dependent upon other messages in the sequence. For this reason there is no use of session or exchange identifiers, and no encryption of the messages. Consequently, there are no guarantees of behavior or acknowledgements, information contained in messages SHOULD be assumed to be visible to the network, and messages SHALL NOT contain private or sensitive information.

The Commissionee and Commissioner SHALL enforce rate limiting of these messages and provide the option for the user to disable handling of them (either temporarily or until manually re-enabled). For example, a rate limiting algorithm of accepting no more than 20 messages (not including retries) in a 10 minute period is recommended.

One possible user journey for this feature is described in [Commissioner Discovery from an Existing Device](#).



Figure 45. Overview of the UDC Protocol Identification Declaration message

The Commissioner is the Initiator and the Commissionee is the Recipient of the IdentificationDeclaration message.




Figure 46. Overview of the UDC Protocol Commissioner Declaration message

The Commissioner is the Initiator and the Commissionee is the Recipient of the CommissionerDeclaration message.

It is assumed that the user has directed the Initiator to send the IdentificationDeclaration message to the Recipient. Upon receipt and before starting a PASE session with the Initiator, it is assumed that the Recipient will obtain DNS-SD record information (including TXT records) for the Initiator, which was either provided in the IdentificationDeclaration message or can be obtained by querying the DNS-SD records, and then prompt the user for approval and to enter its Onboarding Payload or Passcode.

5.3.2. UDC Protocol Messages

Table 67. User Directed Commissioning Protocol

Protocol Opcode	Protocol Command Name	Description
Protocol ID = PROTOCOL_ID_USER_DIRECTED_COMMISSIONING		
0x00	IdentificationDeclaration	The Identification Declaration message provides the DNS-SD Instance Name of the commissionee requesting commissioning to the commissioner selected by the user. It can also include information relating to the requested commissioning session.
0x01	CommissionerDeclaration	The Commissioner Declaration message provides information to a commissionee from the commissioner indicating its pre-commissioning state. This information can be used by the commissionee to simplify the Passcode entry flow for the user.

5.3.3. Message format

All UDC messages SHALL be structured as specified in Section 4.4, “Message Frame Format”.

All UDC messages are unsecured at the message layer:

- The Session ID field SHALL be set to 0.
- The Session Type bits of the Security Flags SHALL be set to 0.
- The S Flag and DSIZ fields of the Message Flags SHALL be set to 0.

The **R Flag** of the **Exchange Flags** for the UDC messages SHALL be set to **0**.

For each UDC message, the **application payload** is the TLV encoding of the message structure as defined below:

Table 68. UDC Messages

Message Name	Payload TLV Encoding
IdentificationDeclaration	IdentificationDeclaration-struct
CommissionerDeclaration	CommissionerDeclaration-struct

The other fields of the **Message format** are not specific to the UDC messages.

5.3.4. Message Exchanges

The flags of the **Exchange Flags** of the **Protocol Header** are defined as follows per UDC message:

Message	I Flag
IdentificationDeclaration	1
CommissionerDeclaration	1

All UDC IdentificationDeclaration messages SHALL be sent without acknowledgement (e.g., unreliably) using UDP, to an IP address found in a A or AAAA record associated with the Commissioner Discovery (**_matterd._udp**) service, using UDP with a destination port as found in the **_matterd._udp** SRV record.

All UDC CommissionerDeclaration messages SHALL be sent without acknowledgement (e.g., unreliably) using UDP, to the source IP address of the sender of IdentificationDeclaration message, and the port specified in the IdentificationDeclaration message payload.

The Initiator MAY send up to 4 retries. Each retransmission SHALL be delayed by at least 100 ms from the previous transmission.

The other fields of the **Protocol Header** are not specific to the UDC messages.

5.3.5. IdentificationDeclaration Message

This message serves to identify the Commissionee. It can also specify information relating to the requested commissioning session, such as whether to display a Passcode entry dialog. It is sent by the Commissionee to the Commissioner.

In addition to a mandatory **instanceName** field, which can be used by the Commissioner to perform **Commissionable Node Discovery** in order to determine the service IP and Port used to establish a PASE session with the Commissionee, there are a number of optional fields in the IdentificationDeclaration which can be used by the Commissionee to improve performance and usability of the user experience.

The payload for the IdentificationDeclaration message begins with 17 bytes which consist of the DNS-SD instance name defined in **Commissionable Node Discovery** encoded as 16 characters of

ascii text, followed by a null termination character.

Following the 17 bytes is the encoded IdentificationDeclaration-struct TLV.

```
instanceName          : OCTET STRING [ length 17 ]
IdentificationDeclaration-struct => STRUCTURE [ tag-order ]
{
  VendorId             [1, optional] : UNSIGNED INTEGER [ range 16-bits ],
  ProductId            [2, optional] : UNSIGNED INTEGER [ range 16-bits ],
  DeviceName           [3, optional] : STRING [ length 0..32 ],
  DeviceType           [4, optional] : UNSIGNED INTEGER [ range 32-bits ],
  PairingInstruction   [5, optional] : STRING [ length 0..32 ],
  PairingHint          [6, optional] : UNSIGNED INTEGER [ range 32-bits ],
  RotatingDeviceId     [7, optional] : STRING [ length 0..100 ].
  Port                 [8, optional] : UNSIGNED INTEGER [ range 16-bits ],
  TargetAppList        [9, optional] : ARRAY OF
  {
    TargetApp          [10, optional] : STRUCTURE [ tag-order ]
    {
      AppVendorId      [11, optional] : UNSIGNED INTEGER [ range 16-bits ],
      AppProductId     [12, optional] : UNSIGNED INTEGER [ range 16-bits ],
    },
  },
  NoPasscode           [13, optional] : BOOLEAN,
  CdUponPasscodeDialog [14, optional] : BOOLEAN,
  CommissionerPasscode [15, optional] : BOOLEAN,
  CommissionerPasscodeReady [16, optional] : BOOLEAN,
  CancelPasscode       [17, optional] : BOOLEAN,
  PasscodeLength       [18, optional] : UNSIGNED INTEGER [ range 8-bits ]
}
```

5.3.5.1. Features

5.3.5.1.1. Bypass Commissionable Node Discovery

In practice, the Commissionable Node Discovery step performed by the Commissioner in response to receiving an IdentificationDeclaration can incur 1 second or more of latency. A Commissionee can provide the same information that a Commissioner would obtain through Commissionable Node Discovery by providing this data in the IdentificationDeclaration message.

The Commissioner MAY bypass the Commissionable Node Discovery step of the [Commissioner Discovery from an Existing Device](#) process when it receives the information it needs in the IdentificationDeclaration message, namely, the **Port** field. The Commissionee MAY also include **VendorId**, **ProductId**, **DeviceName**, **DeviceType**, **PairingInstruction**, **PairingHint**, and **RotatingDeviceId** fields in the IdentificationDeclaration message, similar to the existence of these fields in the DNS-SD Commis-

sionable Node Discovery response.

See [Section 5.3.7.2, “UDC with Commissionable Node Discovery bypass”](#).

5.3.5.1.2. Target Content Application

In some cases, the Commissionee prefers to determine in advance of commissioning if their Content Application is present on the Commissioner (see Device Library, section 10 *Media Device Types*). This allows the Commissionee to warn the user when an experience with reduced functionality will result from commissioning, and may cause the Commissionee or the user to choose not to proceed with the user-directed-commissioning flow.

Some Commissionees will choose not to proceed with the user-directed-commissioning flow if their app is not present for cases where the user experience requires their app, such as when a casting video client app is designed for communication with a specific content app on a casting video player.

The **NoPasscode** field allows the Commissionee to determine if their app is present but has a different account active in it. Some Commissionees do not support guest mode, or multiple simultaneous accounts in their app. In this scenario, the Commissionee might ask the user to switch accounts on the Commissioner in order to proceed.

When the Commissionee provides the **VidList**, the Commissioner SHALL apply this list as a further restriction upon the set of Content Application Vendor IDs that can be used for authentication (see Account Login cluster).

See [Section 5.3.7.4, “UDC with no Passcode prompt \(targeted app selection\)”](#).

5.3.5.1.3. Coordinate Passcode Dialogs

In some cases, the Passcode needed for commissioning is obtained from a Content App running on the Commissioner by way of the Account Login cluster. When the Passcode is obtained in this fashion, no Passcode display on the Commissionee is required, and displaying a Passcode can be confusing for the user.

To address this, the Commissionee can request to be told by the Commissioner when it is requesting a Passcode from the user (so that the Commissionee can display the Passcode) by using the **CdUpon-PasscodeDialog** field in the IdentificationDeclaration message. The Commissioner would then send a CommissionerDeclaration message to indicate that a Passcode is needed.

The Commissionee can also inform the Commissioner that the user has declined or exited the commissioning process by using the **CancelPasscode** field in the IdentificationDeclaration message.

The Commissionee can also inform the Commissioner of the number of digits in the passcode displayed to the user by using the **PasscodeLength** field in the IdentificationDeclaration message.

See [Section 5.3.7.3, “UDC with coordinated Passcode prompt”](#).

5.3.5.1.4. Commissioner-Generated Passcode

In some scenarios, it is simpler for the user to enter a Passcode into the Commissionee UX than it is

to enter a Passcode into a Commissioner UX. For example, when the Commissionee is a mobile phone app and the Commissioner is a TV, it can be easier to enter a Passcode on a mobile phone than a TV remote control (which might not have number buttons).

In this scenario, the Commissioner generates a Passcode and displays it for the user. The user enters the Passcode into the Commissionee UX. The Commissionee UX generates a PASE verifier based upon this Passcode and enters commissioning mode using this PASE verifier. The Commissionee then notifies the Commissioner that it is in commissioning mode (using the new PASE verifier) by sending a IdentificationDeclaration with the **CommissionerPasscodeReady** field set to true.

If, in addition to dynamic passcodes, the Commissioner provides a fixed passcode option, it SHALL be user-updatable and the Commissioner SHALL provide an option to the user to reset the fixed passcode frequently, and SHALL provide an option to disable the fixed passcode feature.

A Commissioner indicates that it supports this feature via the **CP** txt record in [Commissionable Node Discovery](#).

A Commissionee can utilize this feature via the **CommissionerPasscode** and **CommissionerPasscodeReady** fields.

The Commissionee can also inform the Commissioner that the user has declined or exited the commissioning process by using the **CancelPasscode** field in the IdentificationDeclaration message.

See [Section 5.3.7.5, “UDC with Commissioner-generated Passcode”](#).

5.3.5.2. Parameters

This following table lists the parameters used in IdentificationDeclaration with a brief description for each parameter and the feature that uses it.

Table 69. List of IdentificationDeclaration parameters

Parameter Name	Feature	Description
instanceName	Always Required	The DNS-SD instance name defined in Commissionable Node Discovery .
VendorId	Bypass Commissionable Node Discovery	The Vendor ID provided as input to the mdns TXT records (see Section 4.3.1.6, “TXT key for Vendor ID and Product ID (VP)”).
ProductId	Bypass Commissionable Node Discovery	The Product ID provided as input to the mdns TXT records (see Section 4.3.1.6, “TXT key for Vendor ID and Product ID (VP)”).
DeviceName	Bypass Commissionable Node Discovery	The device name as described in the mdns TXT records (see Section 4.3.1.9, “TXT key for device name (DN)”).
DeviceType	Bypass Commissionable Node Discovery	The device type as described in the mdns TXT records (see Section 4.3.1.8, “TXT key for device type (DT)”).

Parameter Name	Feature	Description
PairingInstruction	Bypass Commissionable Node Discovery	The pairing instruction as described in the mdns TXT records (see Section 4.3.1.12 , “TXT key for pairing instructions (PI)”).
PairingHint	Bypass Commissionable Node Discovery	The pairing hint as described in the mdns TXT records (see Section 4.3.1.11 , “TXT key for pairing hint (PH)”).
RotatingDeviceId	Bypass Commissionable Node Discovery	The Rotating Device Identifier as described in the mdns txt records (see Section 4.3.1.10 , “TXT key for rotating device identifier (RI)”).
Port	Bypass Commissionable Node Discovery	The service port for PASE as found in the SRV record of the DNS-SD Commissionable Node Discovery response.
TargetAppList	Target Content Application	The set of content app Vendor IDs (and optionally, Product IDs) that can be used for authentication.
TargetApp	Target Content Application	An entry in the TargetAppList which contains a TargetVendorId and an optional TargetProductId.
AppVendorId	Target Content Application	The content app Vendor IDs that can be used for authentication.
AppProductId	Target Content Application	The content app Product ID associated with the specified Vendor ID that can be used for authentication.
NoPasscode	Target Content Application	Flag to instruct the Commissioner not to display a Passcode input dialog, and instead send a CommissionerDeclaration message if a commissioning Passcode is needed.
CdUponPasscodeDialog	Coordinate Passcode Dialogs	Flag to instruct the Commissioner to send a CommissionerDeclaration message when the Passcode input dialog on the Commissioner has been shown to the user.
Commissioner-Passcode	Commissioner-Generated Passcode	Flag to instruct the Commissioner to use the Commissioner-generated Passcode for commissioning.
Commissioner-PasscodeReady	Commissioner-Generated Passcode	Flag to indicate whether or not the Commissionee has obtained the Commissioner Passcode from the user and is therefore ready for commissioning.

Parameter Name	Feature	Description
CancelPass-code	Coordinate Passcode Dialogs	Flag to indicate when the Commissionee user has decided to exit the commissioning process.
Passcode-Length	Coordinate Passcode Dialogs	The length of the displayed passcode to inform the Commissioner of the number of digits to request from user.

To send a IdentificationDeclaration message, the Commissionee SHALL:

1. Construct the **instanceName** based upon the DNS-SD instance name defined in [Commissionable Node Discovery](#).
2. To optionally allow the Commissioner to bypass the Commissionable Node Discovery request:
 - a. Construct the **Port** from the service port for PASE as found in the SRV record of the DNS-SD [Commissionable Node Discovery](#) response.
 - b. Optionally, construct the **VendorId** value as described as input to the [TXT key for Vendor ID and Product ID](#).
 - c. Optionally, construct the **ProductId** value as described as input to the [TXT key for Vendor ID and Product ID](#).
 - d. Optionally, construct the **DeviceName** value as described in [TXT key for device name](#).
 - e. Optionally, construct the **DeviceType** value as described in [TXT key for device type](#).
 - f. Optionally, construct the **PairingInstruction** value as described in [TXT key for pairing instruction](#).
 - g. Optionally, construct the **PairingHint** value as described in [TXT key for pairing hint](#).
 - h. Optionally, construct the **RotatingDeviceId** value as described in [TXT key for rotating device identifier](#).
3. Optionally, construct the **TargetApplList** based upon the set of content app vid/pids that can be used for authentication. This is a subset of vid/pid selected by the Commissioner based upon the vid/pid of the client.
4. To optionally prevent the Commissioner from automatically showing a Passcode dialog:
 - a. Construct **NoPasscode** and set to true to indicate that this message SHOULD only trigger a commissioning session if it can be done without asking the user to input a Passcode (e.g. the Passcode was provided by external means). When this field is set to true in the IdentificationDeclaration message, the Commissioner SHALL NOT show a Passcode dialog to the user. Instead, the Commissioner SHOULD send a CommissionerDeclaration message with the NeedsPasscodeDialog field set to true (and potentially NoAppsFound to indicate the reason). The Commissionee can then choose whether to prompt a Passcode dialog (by re-sending the IdentificationDeclaration message without this field) or to abandon this commissioning attempt.
5. To optionally request that the Commissioner send a CommissionerDeclaration message to serve as an indication that the Commissioner has presented a Passcode dialog to the user:

- a. Construct the **CdUponPasscodeDialog** and set to true. When this field is set to true in the IdentificationDeclaration message, the Commissioner SHOULD send a CommissionerDeclaration message with the **PasscodeDialogDisplayed** set to true upon showing the Passcode input dialog to the user. This field SHALL NOT be set when the CommissionerPasscode field is set since these two functionalities are mutually exclusive.
6. When using a Commissionee-generated Passcode, to indicate the length of the Passcode displayed to the user by the Commissionee:
 - a. Construct the **PasscodeLength** and set to the length of the displayed Passcode.
7. To optionally request a Commissioner-generated Passcode from a Commissioner which indicates support for this feature via the **CP TXT record** in **Commissionable Node Discovery**:
 - a. Construct the **CommissionerPasscode** and set to true. When this field is set to true in the IdentificationDeclaration message, the Commissioner SHOULD display the commissioning Passcode to the user, and wait to initiate commissioning until an IdentificationDeclaration message is received with the **CommissionerPasscodeReady** field set to true.
 - b. Construct the **CommissionerPasscodeReady** and set to true to indicate the Commissionee has obtained the Commissioner Passcode from the user and is ready for commissioning, or false to indicate that the user has not entered the Passcode yet, and therefore, the Commissioner SHOULD NOT commence commissioning.
8. To optionally indicate that the Commissionee user has cancelled the commissioning process:
 - a. Construct the **CancelPasscode** and set to true. This indicates that the Commissioner can dismiss any dialogs corresponding to commissioning, such as a Passcode input dialog or a Passcode display dialog.
9. Construct and send **IdentificationDeclaration**.

5.3.6. CommissionerDeclaration Message

This message is sent by a Commissioner to a Commissionee in response to a IdentificationDeclaration message. The CommissionerDeclaration provides information indicating the Commissioner's pre-commissioning state.

```
CommissionerDeclaration-struct => STRUCTURE [ tag-order ]
{
  ErrorCode                [1, optional] : UNSIGNED INTEGER [ range 16-bits ],
  NeedsPasscode            [2, optional] : BOOLEAN,
  NoAppsFound              [3, optional] : BOOLEAN,
  PasscodeDialogDisplayed  [4, optional] : BOOLEAN,
  CommissionerPasscode     [5, optional] : BOOLEAN,
  QRCodeDisplayed          [6, optional] : BOOLEAN,
  CancelPasscode           [7, optional] : BOOLEAN,
  PasscodeLength           [8, optional] : UNSIGNED INTEGER [ range 8-bits ]
}
```

5.3.6.1. Parameters

This following table lists the parameters used in IdentificationDeclaration with a brief description for each parameter and the feature that uses it.

Table 70. List of IdentificationDeclaration parameters

Parameter Name	Feature	Description
ErrorCode	All	Indicates errors incurred during commissioning.
NeedsPasscode	Coordinate PIN Dialogs	When NoPasscode field set to true, and the Commissioner determines that a Passcode code will be needed for commissioning.
NoAppsFound	Target Content Application	No apps with Account Login cluster implementation were found for the last IdentificationDeclaration request. Only apps which provide access to the vendor id of the Commissionee will be considered.
PasscodeDialogDisplayed	Coordinate PIN Dialogs	A Passcode input dialog is now displayed for the user on the Commissioner.
Commissioner- Passcode	Commissioner-Generated Passcode	A Passcode is now displayed for the user by the Commissioner.
QRCodeDisplayed	Commissioner-Generated Passcode	The user experience conveying a Passcode to the user also displays a QR code.
CancelPasscode	Coordinate Passcode Dialogs	Flag to indicate whether the Commissioner user has decided to exit the commissioning process.
Passcode- Length	Coordinate Passcode Dialogs	The length of the displayed passcode to inform the Commissionee of the number of digits to request from user.

The allowed values for the ErrorCode field are the following

Code	Description
0	No error
1	Commissionable Node discovery failed
2	PASE connection failed
3	PASE authentication failed (bad Passcode)
4	DAC validation failed
5	Already on fabric
6	Operational Node discovery failed
7	CASE connection failed

Code	Description
8	CASE authentication failed
9	Configuration failed
10	Binding Configuration failed
11	Commissioner Passcode not supported
12	Invalid UDC parameters
13	App Install Consent Pending
14	App Installing
15	App Install Failed
16	App Installed Retry Needed

To send a CommissionerDeclaration message, the Commissioner SHALL:

1. Construct the **ErrorCode** field if an error condition has been encountered during commissioning by the Commissioner.
2. When the IdentificationDeclaration message includes the NoPasscode field set to true, and the Commissioner determines that a Passcode code will be needed for commissioning:
 - a. Construct the **NeedsPasscode** and set to true.
3. Optionally, construct the **NoAppsFound** and set to true to indicate that no apps with Account Login cluster implementation were found for the last IdentificationDeclaration request. Only apps which provide access to the vendor id of the Commissionee will be considered.
4. When the IdentificationDeclaration message includes the CdcForPasscode field set to true:
 - a. Construct the **PasscodeDialogDisplayed** and set to true to indicates that a Passcode input dialog is now displayed for the user on the Commissioner. Upon receipt, the Commissionee SHOULD display its Passcode to the user. This field SHALL NOT be set when the CommissionerPasscode field is set because these fields are mutually exclusive.
5. When the IdentificationDeclaration message includes the CommissionerPasscode field set to true:
 - a. Construct the **CommissionerPasscode** and set to true to indicates that a Passcode is now displayed for the user by the Commissioner. Upon receipt, the Commissionee SHOULD ask the user to input this Passcode into the Commissionee, the Commissionee SHOULD then update its PAKE Verifier so that it uses this Passcode, and then SHOULD send an IdentificationDeclaration message with the **CommissionerPasscodeReady** field set to true.
 - b. Construct the **PasscodeLength** and set to the length of the Passcode displayed by the Commissioner.
 - c. Optionally, construct **QRCodeDisplayed** and set to true to indicates that the Passcode displayed on the Commissioner includes a QR code option. Clients have the option of letting the user scan via the phone camera.
6. Construct and send **CommissionerDeclaration**.

5.3.7. Example Message Exchanges

5.3.7.1. Regular UDC

For reference.

1. UDC with IdentificationDeclaration sent to Commissioner
2. Commissioner reads instanceName, performs Commissionable Node Discovery to determine vid, pid, rotating ID, Device name, PASE IP, PASE Port
3. Commissioner prompts user for permission to commission Device name
4. User approves
5. Commissioner checks for apps granting access to the vid of the Commissionee, applies a filter on this list if the VidList in the IdentificationDeclaration is populated, and attempts to use Account Login on any resulting apps in order to obtain Passcode using rotatingID
6. If no Passcode is obtained, Commissioner prompts user to input Passcode from Commissionee, taking into account PasscodeLength field when provided in the IdentificationDeclaration
7. Commissioner attempts to commission Commissionee

5.3.7.2. UDC with Commissionable Node Discovery bypass

This saves about 1 second of latency in practice due to skipping second step (Commissionable Node Discovery) in Regular UDC.

1. UDC with IdentificationDeclaration containing VendorId, ProductId, RotatingDeviceId and Port sent to Commissioner
2. Commissioner prompts user for permission to commission Device name
3. User approves
4. Commissioner checks for apps granting access to the vid of the Commissionee, applies a filter on this list if the VidList in the IdentificationDeclaration is populated, and attempts to use Account Login on any resulting apps in order to obtain Passcode using rotatingID
5. If no Passcode is obtained, Commissioner prompts user to input Passcode from Commissionee, taking into account PasscodeLength field when provided in the IdentificationDeclaration
6. Commissioner attempts to commission Commissionee

5.3.7.3. UDC with coordinated Passcode prompt

This allows the Commissionee to wait to display the Passcode until it receives a signal from the Commissioner that a Passcode is needed. When the same user account is active on both the Commissionee (often, in a phone app) and the Commissioner (often, in a TV app), then the commissioning process can often complete without requiring the Commissionee to display a Passcode. When no Passcode is required, it can be confusing for the user if the Commissionee shows a Passcode. To address this, the Commissionee can request to be told by the Commissioner when it's Passcode input dialog is displayed to the user by sending the IdentificationDeclaration message with the CdUponPasscodeDialog set to true. Then, when a Passcode is needed, the Commissioner sends a CommissionerDeclaration with the NeedsPasscode set to true.

5.3.7.3.1. UDC with coordinated Passcode prompt, no Passcode needed

1. UDC with IdentificationDeclaration sent to Commissioner containing CdUponPasscodeDialog set to true
2. If Port not provided, then Commissioner reads instanceName, performs Commissionable Node Discovery to determine vid, pid, rotating ID, Device name, PASE IP, PASE Port
3. Commissioner prompts user for permission to commission Device name
4. User approves
5. Commissioner checks for apps granting access to the vid of the Commissionee, applies a filter on this list if the VidList in the IdentificationDeclaration is populated, and attempts to use Account Login on any resulting apps in order to obtain Passcode using rotatingID
6. Passcode is obtained, no need for Commissioner to prompt user to input Passcode from Commissionee
7. Commissioner attempts to commission Commissionee

5.3.7.3.2. UDC with coordinated Passcode prompt, Passcode is needed

1. UDC with IdentificationDeclaration sent to Commissioner containing CdUponPasscodeDialog set to true
2. If Port not provided, then Commissioner reads instanceName, performs Commissionable Node Discovery to determine vid, pid, rotating ID, Device name, PASE IP, PASE Port
3. Commissioner prompts user for permission to commission Device name
4. User approves
5. Commissioner checks for apps granting access to the vid of the Commissionee, applies a filter on this list if the VidList in the IdentificationDeclaration is populated, and attempts to use Account Login on any resulting apps in order to obtain Passcode using rotatingID
6. No Passcode is obtained, Commissioner prompts user to input Passcode from Commissionee
7. Commissioner sends CommissionerDeclaration with the PasscodeDialogDisplayed set to true
8. Commissionee displays Passcode code for user
9. User enters Passcode code into Commissioner, taking into account PasscodeLength field when provided in the IdentificationDeclaration
10. Commissioner attempts to commission Commissionee

5.3.7.4. UDC with no Passcode prompt (targeted app selection)

This allows the Commissionee to determine if their app is present on the Commissioner. Some Commissionees will choose not to proceed with commissioning if their app is not present for cases where the user experience requires their app.

This also allows the Commissionee to determine if their app is present but has a different account active in it. Some Commissionees do not support guest mode, or multiple simultaneous accounts in their app. In this scenario, the Commissionee might ask the user to switch accounts on the Commissioner in order to proceed.

5.3.7.4.1. UDC with no app present

1. UDC with IdentificationDeclaration sent to Commissioner containing NoPasscode set to true
2. If Port not provided, then Commissioner reads instanceName, performs Commissionable Node Discovery to determine vid, pid, rotating ID, Device name, PASE IP, PASE Port
3. Commissioner prompts user for permission to commission Device name
4. User approves
5. Commissioner checks for apps granting access to the vid of the Commissionee, applies a filter on this list if the VidList in the IdentificationDeclaration is populated, and attempts to use Account Login on any resulting apps in order to obtain Passcode using rotatingID
6. No app is found
7. Commissioner sends CommissionerDeclaration with the NoAppsFound and NeedsPasscode fields set to true
8. Commissionee informs user that given Commissioner does not support its app. Provides additional instructions as needed.

5.3.7.4.2. UDC with mismatched account

1. UDC with IdentificationDeclaration sent to Commissioner containing NoPasscode set to true
2. If Port not provided, then Commissioner reads instanceName, performs Commissionable Node Discovery to determine vid, pid, rotating ID, Device name, PASE IP, PASE Port
3. Commissioner prompts user for permission to commission Device name
4. User approves
5. Commissioner checks for apps granting access to the vid of the Commissionee, applies a filter on this list if the VidList in the IdentificationDeclaration is populated, and attempts to use Account Login on any resulting apps in order to obtain Passcode using rotatingID
6. App is found but does not return Passcode
7. Commissioner sends CommissionerDeclaration with the NeedsPasscode field set to true
8. Commissionee informs user that its app has a different active account. Provides additional instructions as needed.

5.3.7.5. UDC with Commissioner-generated Passcode

This allows the Commissionee to utilize a Passcode code generated and displayed on the Commissioner, and input by the user into the Commissionee, rather than the Regular UDC flow where the Passcode code is generated and displayed on the Commissionee, and input by the user into the Commissioner. This flow can simplify the user experience when the Commissionee is a mobile phone app and the Commissioner is a TV since it can be easier to enter a Passcode on a mobile phone than a TV.

1. Commissionee performs Commissioner Discovery, prompts user to select a Commissioner.
2. Selected Commissioner indicates support for [Commissioner Passcode feature](#) (CP=1) in its DNS-SD TXT record.

3. UDC with IdentificationDeclaration sent to Commissioner containing CommissionerPasscode set to true
4. If Port not provided, then Commissioner reads instanceName, performs Commissionable Node Discovery to determine vid, pid, rotating ID, Device name, PASE IP, PASE Port
5. Commissioner prompts user for permission to commission Device name
6. User approves
7. Commissioner checks for apps granting access to the vid of the client, applies a filter on this list if the VidList in the IdentificationDeclaration is populated, and attempts to use Account Login on any resulting apps in order to obtain Passcode using rotatingID
8. If no Passcode is obtained, Commissioner displays a Passcode
9. Commissioner sends CommissionerDeclaration with PasscodeDialogDisplayed and CommissionerPasscode set to true, and PasscodeLength set to the number of digits in the displayed Passcode. If a QR code is also displayed, then QRCodeDisplayed is also set to true.
10. Commissionee prompts user to input Passcode from Commissioner, taking into account PasscodeLength field when provided in the CommissionerDeclaration
11. Commissionee generates a PAKE verifier using input Passcode, enters commissioning mode using this PAKE verifier
12. Commissionee sends IdentificationDeclaration to Commissioner containing CommissionerPasscode and CommissionerPasscodeReady set to true
13. Commissioner attempts to commission Commissionee using Passcode

5.4. Device Discovery

5.4.1. Purpose and Scope

The purpose of this section is to describe the process by which a device is discovered in order to commission it onto an operational Fabric or to fix its operational network configuration. A device is discovered when the commissioner has received information about the device which can be used to trigger the establishment of a point-to-point communication channel with the device which can be used to exchange commissioning messages. For NFC-based commissioning, when the user brings the NFC reader in close proximity to the NFC tag, an activity defined by the NFC Forum will allow the NFC reader to discover the NFC tag.

Depending on the communication interfaces supported by a device and a commissioner, discovering and commissioning a device onto an operational Fabric are possible using Bluetooth Low Energy (BLE), Wi-Fi ([IEEE 802.11-2020](#)), NFC, or over IP if a device is already on an IP network.

Devices that utilize Thread networking technology SHALL support BLE and MAY additionally support NFC for the purpose of discovery and commissioning. Thread network commissioning procedures as defined in [Thread](#) are not used to perform Matter discovery and commissioning.

5.4.1.1. Using BLE

BLE-based commissioning utilizes the Generic Access Profile (GAP) for device discovery, and [Blue-](#)

[tooth Transport Protocol \(BTP\)](#) for message exchange during commissioning.

5.4.1.2. Using Wi-Fi Public Action Frames

Commissioning using Wi-Fi Public Action Frames utilizes Wi-Fi Alliance Unsynchronized Service Discovery ([WFA-USD](#)) for device discovery and communication during commissioning. [WFA-USD](#) is a mechanism for link layer discovery and direct communication between the Commissioner and the device without an IEEE 802.11 association between them. It uses the publish and subscribe public action frames for device discovery and communication during commissioning. The publish and subscribe public action frames are defined by the Wi-Fi Alliance.

After device discovery, the Commissioner and the device use [Public Action Frame Transport Protocol \(PAFTP\)](#) for message exchange during commissioning.

Devices that support commissioning over Wi-Fi Public Action Frame SHALL support such commissioning on all Wi-Fi bands they support.

NOTE

Since support for commissioning via Wi-Fi Public Action Frame is a new feature, it might not be supported by the user's choice of Commissioner. Therefore, devices whose only Matter-defined commissioning method (apart from the mandatory commissioning over existing IP-based network) uses Wi-Fi Public Action Frame, and which are being certified in the first 2 years (August 2025-August 2027) following the release of the Matter 1.4.2 Specification, SHALL support the [Commissioning Fall-back Mechanism](#). This period can be extended beyond August 2027 until there is sufficient support by commissioners and the platforms used by popular commissioners, covering a sufficient proportion of consumer-facing commissioner devices, so that consumers can count on the high likelihood of the feature being available to them using their devices.

5.4.1.3. Using NFC Transport Layer

NFC-based Commissioning utilizes the physical tap performed by the user between the commissioner and commissionee for technology detection, and [NFC Transport Layer \(NTL\)](#) for message exchange during commissioning.

5.4.1.4. Using IP-based network

If a device already has IP-bearing network connectivity (over Wi-Fi, Ethernet, or otherwise), a Commissioner SHALL discover such a device using DNS-based Service Discovery (DNS-SD), and then convey commissioning messages to the device over IP.

5.4.1.5. Using Thread's Commissioning Protocol

Commissioning using Thread's commissioning protocol relies on the Matter Commissioner becoming a Thread Commissioner. The Matter Commissioner must first petition the Thread network of choice to become the Commissioner on that network. Upon successful completion of that petition, the Matter Commissioner configures the Thread network's steering data to indicate a set of Matter devices that it is looking to commission. That indication is based on the set of Discriminators.

The commissionee first scans the available Thread networks, and matches the steering data against its Discriminator. Upon finding such a network, the commissionee sends an DNS Announcement message to convey all information permitted in the Announcement; from the commissionee’s point, this step is very much the same as announcement using IP-based network.

After the device discovery, the commissioning protocol is the same as using IP-based network from the commissionee’s point of view. From the Matter Commissioner point of view, the protocol is sent over Thread’s Relay Commands.

The protocol and its message flows are described in [Section 4.23, “Thread Commissioning Protocol”](#).

5.4.2. Announcement by Device

This section describes how devices announce their commissionable status to allow a Commissioner to discover the device to be commissioned and how devices announce network recovery requests towards their administrators.

5.4.2.1. Announcement and Discovery Using Multiple Technologies

A device SHALL announce on all of the networking technologies it supports as indicated in the Discovery Capability Bitmask (see [Table 60, “Discovery Capabilities Bitmask”](#)), except for NFC (as there is no announcement over NFC interface).

5.4.2.2. Announcement Commencement

A device which is not yet commissioned into a Matter fabric SHALL commence announcing its ability to be commissioned depending on its primary device function and manufacturer-chosen [Device Commissioning Flow](#), per the table below. Nodes that implement [Network Recovery](#) support SHALL announce their need for network recovery on supported commissioning channels, such as Bluetooth LE or Wi-Fi Public Action Frame. Otherwise, Nodes already commissioned into one or more Matter fabrics or already connected to an IP-bearing network that wish to announce their commissioning requests SHALL ONLY do so using DNS-SD over their operational network (see [Section 5.4.2.7, “Using Existing IP-bearing Network”](#) and [Section 4.3, “Discovery”](#)). In the interest of privacy, an already-commissioned Node SHALL NOT commence commissioning announcement using Bluetooth LE or Wi-Fi Public Action Frame.

Primary Device Function	Announcement
Most control originates from a Fabric (excluding Locks and Barrier Access Devices)	SHALL start announcing automatically upon application of power when using Standard commissioning flow . When using User-intent commissioning flow or Custom Commissioning flow , it SHALL NOT start announcing automatically upon application of power.
Most control does not originate from a Fabric (e.g., dishwasher, coffee maker, refrigerator)	SHALL NOT start announcing automatically upon application of power. User-intent commissioning flow or Custom Commissioning flow is required.

Primary Device Function	Announcement
Locks and Barrier Access Devices	SHOULD NOT start announcing automatically upon application of power. If Standard commissioning flow is used, additional security considerations are required as listed in Security Best Practices . User-intent commissioning flow or Custom Commissioning flow SHOULD be used. Due to the complexity of installing these device types, if Custom Commissioning flow is used the device is allowed to start announcing automatically upon application of power. If the device starts announcing automatically upon application of power then all counter measures for Standard commissioning flow SHALL apply.

Note that the above guidelines are in place to avoid unnecessary pollution of the 2.4 GHz spectrum and as a mitigation of the privacy threat created due to unnecessary transmissions by a commissionable device.

If announcement has ceased (see [Section 5.4.2.3, “Announcement Duration”](#)), it may be re-initiated via a device-specific user interaction such as a button press or other action defined by the manufacturer and indicated by the methods specified in [Section 5.7, “Device Commissioning Flows”](#).

5.4.2.3. Announcement Duration

5.4.2.3.1. Commissionable Announcement Duration

In order to minimize unnecessary pollution of the 2.4 GHz and 5 GHz shared wireless spectrum, especially with device discovery, a commissionable device SHALL NOT announce with a rapid interval for a duration longer than 15 minutes after announcement commences. This duration was chosen to capture the primary case of a user setting up immediately after powering on for a range of devices, including time to download, install and launch applications, transit rooms within a home, etc.

Note that devices MAY choose to announce for less time in order to conserve battery life or for other device-specific reasons. Note that an announcement duration that is too short may result in a poor setup experience for users. Shorter announcement intervals SHOULD only be employed to meet otherwise unattainable device functionality/requirements. To help strike a balance between a good setup experience and conserving battery life, a device SHALL NOT announce for a duration of less than 3 minutes after announcement commences.

A failed attempt to commission does not restart or delay the timeout. Moreover, this timeout applies only to cessation of announcements and not to abortion of connections, *i.e.*, a connection SHOULD NOT abort prematurely upon expiration of the announcement duration.

5.4.2.3.2. Network Recovery Announcement Duration

The Recovery Nodes SHALL announce for duration of 48 hours, in line with the limits set forth by [Extended Announcement](#). In order to minimize the chance of advertising storms and to minimize unnecessary pollution of the 2.4 GHz and 5 GHz shared wireless spectrum, devices in Network Recovery state SHALL use advertising rates mandated for the Extended Announcement rate for the

entire Announcement Duration. The specific intervals are called out in the specification of each commissioning channel, for example, see [Section 5.4.2.5.3, “Interval”](#) for BLE advertising interval.

A failed attempt to restore device’s network connectivity does not restart or delay the timeout. Moreover, this timeout applies only to cessation of announcements and not to abortion of connections, *i.e.*, a connection **SHOULD NOT** abort prematurely upon expiration of the announcement duration.

5.4.2.3.3. Extended Announcement

A device **MAY** announce for a longer period, up to 48 hours in total, known as Extended Announcement. For setup cases, this enhances the likelihood of success in cases where a user needs more than 15 minutes after first powering an uncommissioned device (e.g. physical setup takes >15 minutes, user must leave and return later). For network recovery cases, devices **SHALL** employ Extended Announcement to allow sufficient time for corrective action, except in cases where its use would negatively affect battery life.

If a device opts to use Extended Announcement, it **SHALL** set both the PID and VID to 0 in the announcement and **SHALL** elide any Extended Data. This is to preserve user privacy for scenarios where a device may remain uncommissioned for extended periods.

5.4.2.4. Discovery Information

This section details the information advertised by a commissionable Node.

Field	Length	Is Required?
Discriminator	12-bit	Yes
Vendor ID	16-bit	No
Product ID	16-bit	No
Extended Data	Variable	No

5.4.2.4.1. Discriminator

A 12-bit value matching the field of the same name in the Onboarding Payload.

5.4.2.4.2. Vendor ID

A 16-bit value identifying the device manufacturer (see [Section 2.5.2, “Vendor Identifier \(Vendor ID, VID\)”](#)).

5.4.2.4.3. Product ID

A 16-bit value identifying the product (see [Section 2.5.3, “Product Identifier \(Product ID, PID\)”](#)).

5.4.2.4.4. Extended Data

Extended Data **MAY** be made available by commissionable Nodes. This data **SHALL** be encoded using a standard TLV encoding defined in this section. The location of this data varies based on the Node’s commissioning networking technology.

This extended data SHALL be encoded as a [TLV structure](#) tagged with an anonymous tag.

The members of this structure SHALL use [context-specific tags](#) with the values and meanings shown in the table below.

Tag	Value	Member type	Member Description
RotatingIdTag	0x00	octet string	Rotating Device Identifier

5.4.2.4.4.1. Rotating Device Identifier

Some device makers need a way to uniquely identify a device before it has been commissioned for vendor-specific customer support purposes. For example, the device maker may need this to identify factory software version and related features, manufacturing date, or to assist in recovery when the Onboarding Material has been lost or damaged. In order to avoid privacy issues associated with exposing a fixed unique identifier, devices MAY utilize a Rotating Device Identifier for identification purposes. A Rotating Device Identifier is similar to a serial number but rotates at pre-defined moments.

The Rotating Device Identifier provides a non-trackable identifier which is unique per-device and that can be used in one or more of the following ways:

- Provided to the vendor's customer support for help in commissioning or establishing Node provenance;
- Used programmatically to obtain a Node's [Passcode](#) or other information in order to provide a simplified setup flow. Note that the mechanism(s) by which the [Passcode](#) may be obtained is outside of this specification. If the Rotating Device Identifier is to be used for this purpose, the system implementing this feature SHALL require proof of possession by the user at least once before providing the [Passcode](#). The mechanism for this proof of possession, and validation of it, is outside of this specification.

The Rotating Device Identifier is an optional feature for a Node to implement and an optional feature for a Commissioner to utilize. The algorithm used for generating a Rotating Device Identifier SHALL meet the following security and privacy requirements:

1. It SHALL be irreversible in such a way that:
 - a. It SHALL prevent recovery of a unique identifier for the device by entities that do not already have access to the set of possible unique identifiers.
 - b. Leaking of a common key or equivalent could not be used to recover a unique identifier for all devices sharing the common key.
2. It SHALL protect against long-term tracking by rotating upon each commencement of advertising.
3. It SHALL have a total of at least 64 bits of entropy and SHOULD preferably have more, up to 256 bits.
4. It SHALL NOT contain a fixed identifier such as a serial number.

The [Rotating Device Identifier Algorithm](#) below meets these requirements. A Node that implements the Rotating Device Identifier SHALL use either the Rotating Device Identifier Algorithm or a differ-

ent algorithm which has been approved and verified by the Connectivity Standards Alliance for this purpose and which meets the same set of security and privacy requirements listed above.

The Rotating Device Identifier Algorithm employs a key derivation algorithm that combines a monotonically increasing lifetime counter with a persistent unique per-device identifier.

This persistent unique identifier SHALL consist of a randomly-generated 128-bit or longer octet string which SHALL be programmed during factory provisioning or delivered to the device by the vendor using secure means after a software update.

This persistent unique identifier SHALL be protected against reading or writing over the air after initial introduction into the device, and stay fixed during the lifetime of the device.

NOTE

This persistent unique identifier SHALL NOT be the same as the UniqueID attribute exposed in the Basic Information cluster.

The lifetime counter SHALL be an integer at least 16 bits in size, incremented upon each commencement of advertising, and wrapping when the maximum value is reached.

The Rotating Device Identifier Algorithm is defined as follows:

```
Rotating Device ID = Rotation Counter || Crypto\_KDF(
                                inputKey := Persistent Unique ID,
                                salt:= Rotation Counter,
                                info := "RotatingDeviceID",
                                len := 128)
```

(where || is the concatenation operation)

The rotation counter is encoded as 2 bytes using little-endian encoding in the above algorithm, everywhere it appears.

The Rotating Device ID is the concatenation of the current rotation counter and the 16 bytes of the [Crypto_KDF](#) result.

5.4.2.4.4.2. TLV Example

Extended data containing just a Rotating Device Identifier would be encoded as the following bytes:

Offset	Data	Comments
0x00	0x15	Control byte for structure with anonymous tag
0x01	0x30	Control byte for octet string with 1-byte length and a context-specific tag
0x02	0x00	Context-specific tag for Rotating Device Identifier
0x03	0x12	Length of Rotating Device Identifier (e.g. 18 bytes)
0x04	0xXX..0xXX	Rotating Device Identifier

0x16	0x18	End of container
------	------	------------------

5.4.2.5. Using BLE

This section provides details of how a device announces its commissionable status or how a device requests fixes to the operational network configuration using BLE technology (see [Bluetooth® Core Specification 4.2 \(amended\)](#) and [Bluetooth® Core Specification Supplement 12](#)). As required in [Section 5.4.2.2, “Announcement Commencement”](#), Nodes currently commissioned into one or more fabrics or already connected to an IP-bearing network SHALL NOT employ this method for commissioning.

5.4.2.5.1. Device Role

Commissionable devices SHALL implement the role of a Generic Access Profile (GAP) Peripheral.

5.4.2.5.2. Channels

There are three advertising channels used by BLE. All three channels SHOULD be used by commissionable devices for BLE advertising.

5.4.2.5.3. Interval

Commissionable devices SHOULD use an Advertising Interval between 20 ms and 60 ms for the first 30 seconds and a value between 150 ms to 1285 ms for the rest of the Announcement duration. Shorter intervals typically result in shorter discovery times.

If a device opts to use [Extended Announcement](#), it SHALL switch to using an Advertising Interval larger or equal to 1200 ms and SHOULD use a nominal Advertising Interval of 1285 ms. When using Extended Announcement, the device SHALL set the Extended Announcement Flag in the Matter Service Data in the BLE Advertisement (see [Table 72, “Matter BLE Service Data payload format — Commissionable OpCode”](#)).

5.4.2.5.4. Advertising Mode

Commissionable devices SHALL use the GAP General Discoverable mode, sending connectable undirected advertising events.

5.4.2.5.5. Advertising Address

To ensure privacy, commissionable devices SHALL use LE Random Device Address (see [Bluetooth® Core Specification 4.2 \(amended\)](#) Vol 6, Part B, Section 1.3.2.1 "Static device address") for BLE Advertising and SHALL change it at least on every boot.

5.4.2.5.6. Advertising Data

In order to reduce 2.4 GHz spectrum congestion due to active BLE scanning, and to extend battery life in battery-powered devices, all critical data used for device discovery is contained in the Advertising Data rather than the Scan Response Data. This allows a BLE Commissioner to passively scan (i.e., not issue Scan Requests upon receiving scannable advertisements) and still be able to receive all information needed to commission a device.

Note that if additional vendor-specific information is to be conveyed and does not fit within the Advertising Data, it may be included in the Scan Response Data. See [Section 5.4.2.9, “Manufacturer-specific data”](#) for details on including vendor-specific information.

Advertising data for Matter discovery uses "Service Data - 16 bit UUID" advertisement data type (see [Bluetooth® Core Specification Supplement 12](#) Section 1.11 "Service Data"), with 16-bit UUID value of 0xFFFD (see [Table 36, “SIG UUID assignment”](#)).

All multi-byte values are encoded in little-endian byte order within the service data payload.

The advertising format depends on the value of the Matter BLE Device OpCode. The Device OpCode is always embedded in the first byte of the service data. The following Device OpCode values are defined:

Table 71. Matter BLE Device Opcodes

Value	Description
0x00	Commissionable
0x01	Network Recovery
0x02 - 0xFF	Reserved

When the Device OpCode is set to 0x00 (Commissionable), the following fields are defined within the Matter service data:

Table 72. Matter BLE Service Data payload format — Commissionable OpCode

Bytes	Type	Description
0	uint8	Matter BLE Device OpCode. Value == 0x00 (Commissionable)
1-2	uint16	Bits[15:12] == 0x0 (Advertisement version) Bits[11:0] == 12-bit Discriminator (see Section 5.4.2.4.1, “Discriminator”)
3-4	uint16	16-bit Vendor ID (see Section 5.4.2.4.2, “Vendor ID”) Set to 0, if elided
5-6	uint16	16-bit Product ID (see Section 5.4.2.4.3, “Product ID”) Set to 0, if elided
7	uint8	Bit[0] == Additional Data Flag (see Section 5.4.2.5.7, “GATT-based Additional Data”) Bit[1] == Using Extended Announcement Flag (see Section 5.4.2.3.3, “Extended Announcement”) Bits[7:2] are reserved for future use and SHALL be clear (set to 0)

Devices MAY choose not to advertise either the VID and PID, or only the PID due to privacy or other considerations. When choosing not to advertise both VID and PID, the device SHALL set both VID and PID fields to 0. When choosing not to advertise only the PID, the device SHALL set the PID field to 0. A device SHALL NOT set the VID to 0 when providing a non-zero PID.

The Using Extended Announcement flag SHALL be set while the device is in the [Extended](#)

Announcement period and SHALL NOT be set during the initial **Announcement Duration**.

The following table details the contents of an exemplary Advertising PDU payload with a Commissionable OpCode for Vendor ID 0xFFF1, Product ID 0x8000, Discriminator 0x3AB and with **GATT-based Additional Data** marked as present:

Table 73. Exemplary Matter BLE commissionable advertisement

Byte	Value	Description
0	0x02	AD[0] Length == 2 bytes
1	0x01	AD[0] Type == 1 (Flags)
2	0x06	Bit 0 (LE Limited Discoverable Mode) set to 0 Bit 1 (LE General Discoverable Mode) set to 1 Bit 2 (BR/EDR Not supported) set to 1
3	0x0B	AD[1] Length == 11 bytes
4	0x16	AD[1] Type == 0x16 (Service Data - 16-bit UUID)
5-6	0xF6, 0xFF	16-bit Matter UUID assigned by Bluetooth SIG = 0xFFF6
7	0x00	Matter BLE OpCode == 0x00 (Commissionable)
8-9	0xAB, 0x03	Bits[15:12] == 0x0 (Advertisement version) Bits[11:0] == 12-bit Discriminator == 0x3AB
10-11	0xF1, 0xFF	16-bit Vendor ID == 0xFFF1
12-13	0x00, 0x80	16-bit Product ID == 0x8000
14	0x01	Bit[0] == 1: (GATT-based Additional Data present) Bits[7:1] clear (reserved for future use)

Note that the position of the fields within the Advertising PDU (e.g. Flags, 16-bit UUID Service Data, etc.) within a conformant advertisement MAY differ from the example above, since there are no ordering constraints for advertisement fields in the Length-Type-Value format used in BLE advertisements.

When the Device OpCode is set to 0x01 (Network Recovery), the following fields are defined within the Matter Service Data:

Table 74. Matter BLE Service Data payload format — Network Recovery OpCode

Bytes	Type	Description
0	uint8	Matter BLE Device OpCode. Value == 0x01 (Network Recovery)
1	uint8	Bits[7:4] == 0x0 (Advertisement version) Bits [3:0] == 0x0 (reserved for future use)

2-9	uint64	64-bit Recovery ID (see Section 11.10.6.11 , “RecoveryIdentifier”)
10	uint8	Bit[0] == Additional Data Flag (see Section 5.4.2.5.7 , “GATT-based Additional Data”) Bit[7:1] are reserved for future use and SHALL be clear (set to 0)

The following table details the contents of an exemplary Advertising PDU for the Network Recovery OpCode, with Recovery ID set to the octet string **0x11 0x22 0x33 0x44 0x55 0x66 0x77 0x88**

Table 75. Advertising PDU payload when CHIP BLE OpCode indicates Node is in Network Recovery mode

Byte	Value	Description
0	0x02	AD[0] Length == 2 bytes
1	0x01	AD[0] Type == 1 (Flags)
2	0x05	Bit 0 (LE Limited Discoverable Mode) SHOULD be set to 1 Bit 1 (LE General Discoverable Mode) SHOULD be set to 0 Bit 2 (BR/EDR Not Supported) set to 1
3	0x0E	AD[1] Length == 14 bytes
4	0x16	AD[1] Type == 0x16 (Service Data - 16-bit UUID)
5-6	0xF6, 0xFF	16-bit Matter UUID assigned by Bluetooth SIG
7	0x01	Matter BLE OpCode == 0x01 (Network Recovery mode)
8	0x00	Bits[7:4] == 0x0 (Advertisement version) Bits[3:0] == 0x0 (reserved for future use)
9-16	0x11, 0x22, 0x33, 0x44, 0x55, 0x66, 0x77, 0x88	64-bit Recovery ID 11 22 33 44 55 66 77 88
17	0x00	Bit[0] == 0 (GATT-based Additional Data Flag not present) Bits[7:1] clear (reserved for future use)

5.4.2.5.7. GATT-based Additional Data

When the Additional Data Flag is set in the Matter Service Data in the BLE Advertisement, the commissioner MAY access additional commissioning-related data via an unencrypted read-only GATT characteristic C3 (see [Table 34](#), “BTP GATT service”).

The value of the C3 characteristic SHALL be set to the Extended Data payload of the Discovery Information (see [Section 5.4.2.4.4, “Extended Data”](#)).

5.4.2.6. Using Wi-Fi Public Action Frame

This section provides details of how a device announces its commissionable status using Wi-Fi Public Action Frame. As required in [Section 5.4.2.2, “Announcement Commencement”](#), Nodes currently commissioned into one or more fabrics or already connected to an IP-bearing network SHALL NOT employ this method.

5.4.2.6.1. Device Role

Commissionable devices SHALL implement the role of a publisher in Wi-Fi Alliance Unsynchronized Service Discovery ([WFA-USD](#)).

5.4.2.6.2. Channels

The Commissionable device advertises itself on channel 6 in the 2.4 GHz frequency band, which is the Default Publish Channel in [WFA-USD](#). Additionally, the Commissionable device also advertises itself on channels in the Publish Channel List. This Publish Channel List includes all 20 MHz channels in the 2.4 GHz frequency band and 5 GHz band (if supported) that are allowed per regulation in the geographical location.

The Commissionable device iterates between advertising itself on the Default Publish Channel and on channels in the Publish Channel List per behavior described in [WFA-USD](#). For Commissionable devices to advertise on Dynamic Frequency Selection channels within the 5 GHz band, full compliance with the regulatory requirements of the specific geographical location is mandatory.

The Commissionable device can prioritize advertising on channels in the Publish Channel List in which it determines the presence of an IEEE 802.11 Basic Service Set. The Commissionable device can determine presence of an IEEE 802.11 Basic Service Set on a channel by scanning that channel passively (i.e., SHOULD NOT send Probe Requests) in order to reduce unnecessary transmissions in the shared spectrum.

The detection of an IEEE 802.11 beacon on a particular channel serves as a strong indication of the presence of an IEEE 802.11 Basic Service Set operating on that same channel. Moreover, within the IEEE 802.11 beacon, a Commissionable device can find information about the regulatory domain it is currently located in, specifically in the Country Information Element.

5.4.2.6.3. Publisher Operating Parameters

Commissionable devices SHALL create a Publish instance by invoking the Publish Method, using argument values specified in [Table 76, “Publish method arguments values for discovery”](#). Little-endian encoding is used for all argument values.

Table 76. Publish method arguments values for discovery

Argument	Description	Value for discovery
service_name	UTF-8 name string which identifies the service	_matterc._udp

Argument	Description	Value for discovery
matching_filter_tx	Ordered sequence of <length,value> pairs to be included in the Publish message	NULL
matching_filter_rx	Ordered sequence of <length, value> pairs that specify further the response conditions beyond the service name used to filter the Subscribe message	NULL
service_specific_info	Sequence of values that are conveyed in the Publish message	Ordered sequence: <8-bits, Device OpCode >, <16-bits, Device Information >, <16-bits, Vendor ID >, <16-bits, Product ID >, <variable-bits, Extended Data >
configuration_parameters	Miscellaneous configuration parameters	Publish Type = unsolicited and solicited transmissions Discovery range = 0 Solicited transmission type = unicast Announcement period = <i>(see text below)</i> Time to live = <i>(see text below)</i> Event conditions = always Matching filter flag = 1 NAN Ranging Flag = 0 Data Path Flag = 0 Awake DW Interval = Not Applicable Further Service Discovery flag = 1 Further Service Discovery function = 0 NAN Discovery flag = 1 Default Publish Channel = 6 Publish Channel List = [all 20 MHz channels in 2.4 GHz band and 5 GHz band (if supported), per regulation in the geographical region] Nmin = 5 Nmax = 10 Mmin = 5 Mmax = 10

Table 77. Wi-Fi Public Action Frame Device OpCode

Device OpCode Value	Description
0x00	Commissionable
0x01	NetworkRecovery
0x02 - 0xFF	Reserved

Table 78. Wi-Fi Public Action Frame Device Information

Bits	Description
b15:b12	Advertisement version, set to 0x0
b11:b0	12-bit discriminator (see Section 5.4.2.4.1, “Discriminator”)

Some devices MAY choose not to advertise the Vendor ID and/or Product ID due to privacy or other considerations. These devices SHALL set the Vendor ID and Product ID to value 0 in the [service_specific_info](#). Extended Data is optionally present in [service_specific_info](#). The [Time to live](#) value is set according to [Section 5.4.2.3, “Announcement Duration”](#). The recommended value of [Announcement period](#) is 100 TUs. 1 TU is equal to 1024 microseconds.

Single-channel Publish State and Multiple-channels Publish State defined in [WFA-USD](#) govern the advertisement behavior.

1. Single-channel Publish State: The Commissionable device operates in the Default Publish Channel during this state. The dwell period (the time spent on a channel) is randomly determined as $N \times 100$ TU, where N is an integer randomly selected within the range [N_{min} , N_{max}].
2. Multiple-channels Publish State: In this state, the publisher can operate in one or more channels from the Publish Channel List. The dwell period is randomly determined as $M \times 100$ TU, where M is an integer randomly selected within the range [M_{min} , M_{max}]. The implementation of [WFA-USD](#) has flexibility to divide the dwell period across the channels in the Publish Channel List. The Commissionable device MAY publish on a subset of channels within each multiple-channels Publish State, and the minimum time spent on any one channel is 100 ms in [WFA-USD](#).

If a device opts to use [Extended Announcement](#), it SHALL switch to using N_{min} , N_{max} , M_{min} , and M_{max} greater than or equal to 10, 15, 10, 15, respectively, and SHOULD use an [Announcement period](#) of 1000 TUs.

NOTE These recommended values may be updated in a future version of the specification.

5.4.2.7. Using Existing IP-bearing Network

This section details how a device that is already connected to an IP-bearing network advertises its commissionable state. The discovery protocols leverage IETF Standard DNS-based Service Discovery [[RFC 6763](#)]. A device SHALL use multicast DNS [[RFC 6762](#)] on Wi-Fi and Ethernet networks to make itself discoverable. On Thread networks, a device SHALL use the Service Registration Protocol [[RFC 9665](#)] and an Advertising Proxy [[AdProx](#)] running on a Thread Border Router to make itself discoverable. Additional details on application of the above protocols in Matter is found in [Section 4.3, “Discovery”](#). The encoding of the information required for discovery during the commissioning process is covered in [Section 4.3.1, “Commissionable Node Discovery”](#).

5.4.2.8. Using Thread Commissioning Protocol

NOTE Support for Thread Commissioning Protocol is provisional.

As required in [Section 5.4.2.2, “Announcement Commencement”](#), Nodes currently commissioned into one or more fabrics, already connected to a non-Thread IP-bearing network or already con-

nected non-provisionally to a Thread network SHALL NOT employ this method. The overall protocol is described in [Section 4.23, “Thread Commissioning Protocol”](#):

- For overall sequence diagram, see [Figure 43, “Advertising over Thread Commissioning Protocol”](#)
- For the details of selecting the Thread network, see [Section 4.23.2, “Network Discovery and Selection”](#)
- For advertising payloads and timing, see [Section 4.23.3.1, “Advertisement from Commissionee”](#)

5.4.2.9. Manufacturer-specific data

If needed, manufacturer-specific data MAY be advertised by a commissionable device using one of the following mechanisms, based on the supported commissioning technology. Commissioners receiving this data SHOULD treat it as opaque unless they have the need to and possess the ability to correctly interpret the information conveyed.

5.4.2.9.1. Using BLE

Any manufacturer-specific data may be included as a Manufacturer Specific Data AD type in the Advertising Data or in the Scan Response data.

Note that to receive Scan Response data information the Commissioner has to perform BLE active scanning that, in addition to creating additional traffic in the shared 2.4 GHz unlicensed band, can delay device discovery and connection, increasing the overall time required to commission a device.

5.4.2.9.2. Using Wi-Fi Public Action Frame

Any manufacturer-specific data SHOULD be conveyed using the Vendor-specific attributes in [WFA-USD](#).

5.4.3. Discovery by Commissioner

How a Commissioner discovers a commissionable device depends on the communication interfaces that the device and the Commissioner support (see [Section 5.4.2.1, “Announcement and Discovery Using Multiple Technologies”](#)). Though not all communication interfaces must be supported by every device (see [Table 60, “Discovery Capabilities Bitmask”](#)), a Commissioner SHALL support Commissioning (see [Section 5.5, “Commissioning Flows”](#)) using existing IP network and over BLE (if equipped with a BLE interface). Furthermore, commissioners SHOULD support commissioning over Wi-Fi Public Action Frame (if equipped with a Wi-Fi interface), and SHOULD support NFC-based Commissioning (if equipped with an NFC interface).

A Commissioner that is aware of the device’s Discovery Capability Bitmask SHALL initiate Device Discovery on all of the communication interfaces that are supported by both the Commissioner and the device. A Commissioner that is unaware of the device’s Discovery Capability Bitmask SHALL initiate Device Discovery on all of the networking technologies it supports out of BLE, Wi-Fi Public Action Frame, and DNS-SD network discovery on IP networks.

Commissioners SHALL always support discovering a device using DNS-based Service Discovery (DNS-SD) for commissioning, irrespective of the [Discovery Capabilities Bitmask](#) specified in the

Onboarding Payload Contents.

The following sections detail Commissioner behavior for each of these networking technologies. Though a QR code, NFC Tag, or Manual Pairing code may be scanned, tapped or entered prior to discovery, it is not required to do so. However, after scan/tap/entry of the code, the Discriminator, VID and PID elements are available to ensure that the intended device is discovered before proceeding to the connection phase of commissioning.

5.4.3.1. Using BLE

Commissioners SHALL implement the role of a GAP Central. To discover a commissionable device advertising over BLE, a Commissioner SHALL perform a BLE scan across all three advertising channels with a sufficient dwell time, interval, and overall duration of scan. In order to promote quick discovery it is recommended that a Commissioner scan as aggressively as possible within the Commissioner device functionality/UX constraints. In addition, if manufacturer-specific data is not needed, a passive scan (*i.e.*, one that only listens for Advertisement PDUs and does not issue Scan Request PDUs) SHOULD be used.

If discovery procedure is user initiated the scan interval SHOULD be set between 30 ms and 60 ms, and the scan window SHOULD be set to 30 ms. If discovery procedure is not user initiated (*i.e.*, the Commissioner is scanning in the background), the device may use more relaxed scan, for example, the scan interval set to 1.28 seconds and scan window set to 11.25 ms.

NOTE

Recommended values are defined in Appendix A: Timers and Constants of [Blue-tooth® Core Specification 4.2 \(amended\)](#), Vol 3, Part C.

5.4.3.2. Using Wi-Fi Public Action Frame

To discover a commissionable device advertising using Wi-Fi Public Action Frame (see [Section 5.4.2.6, “Using Wi-Fi Public Action Frame”](#)), a Commissioner that supports commissioning using Wi-Fi Public Action Frame SHALL implement the role of a Subscriber in Wi-Fi Alliance Unsynchronized Service Discovery (WFA-USD). The Commissioner SHALL create an active or a passive subscribe instance by invoking the Subscribe Method using the argument values specified [Table 79, “Subscribe method arguments values for discovery”](#) (WFA-USD).

Table 79. Subscribe method arguments values for discovery

Argument	Description	Value for discovery
service_name	UTF-8 name string which identifies the service	_matterc_udp
matching_filter_rx	Ordered sequence of <length, value> pairs used to filter out received publish discovery messages containing the service name	NULL
matching_filter_tx	Ordered sequence of <length,value> pairs to be included in the Subscribe message	NULL

Argument	Description	Value for discovery
ser-vice_speci-fic_info	Sequence of values that are conveyed in the Subscribe message	NULL
configura-tion_para-meters	Miscellaneous configuration parameters	Subscribe Type = Passive or Active Discovery range = 0 Query period = <i>(see text below)</i> Time to live = <i>(see text below)</i> Matching filter flag = 1 NAN Ranging Flag = 0 Awake DW Interval = 1 NAN Discovery flag = 1 or 0 Default Publish Channel = 6 Publish Channel List = [44, 149]

The recommended value of **Query period** is between 100 TUs and 600 TUs when **Subscribe Type = Active**. The **Query period** is not applicable when **Subscribe Type = Passive**. The **Time to live** value is set according to [Section 5.4.2.3, “Announcement Duration”](#).

The Commissioner can create the subscribe instance on Default Publish Channel (channel 6 in the 2.4 GHz band) or any of the 20 MHz channels in the 2.4 GHz frequency band or any of the 20 MHz channels in the 5 GHz frequency band, if supported (see [Section 5.4.2.6.2, “Channels”](#)). The frequency and duration of the Commissioner’s presence on the channel where the subscribe instance is created are based on implementation. For swiftest device discovery in **WFA-USD**, the Commissioner SHOULD follow these guidelines:

1. Whenever possible, the Commissioner SHOULD create a subscribe instance on the Default Publish Channel.
2. If the Commissioner is already associated with an operational network on the 5 GHz band, it can still create a subscribe instance on the same operating channel. However, if Commissioner is not an Access Point and if the operating channel is a Dynamic Frequency Selection channel, then to ensure compliance with the local regulatory requirements, the Commissioner SHOULD choose either channel 44 or channel 149 in the 5 GHz band for the subscribe instance, whichever is permitted with a higher transmission power based on the regulatory domain. Access Points as Commissioners can choose to create subscribe instance on any operating channel, including Dynamic Frequency Selection channels.

When the Commissioner discovers the **_matterc._udp** service name published by a Commissionable device, and the Commissioner uniquely identifies the Commissionable device through **12-bit discriminator** or **Extended Data**, then the Commissioner MAY proceed with the commissioning flow (see [Section 5.5, “Commissioning Flows”](#)). When commissioning over Wi-Fi Public Action Frame, the Commissioner SHALL use concurrent connection commissioning flow. Any Commissionable device supporting Wi-Fi Public Action Frame commissioning SHALL support concurrent connection commissioning flow, indicating this capability in the **SupportsConcurrentConnection** attribute of the **General Commissioning Cluster**.

During commissioning flow, the Commissioner and the Commissionable device SHALL use **WFA-**

USD Service Discovery Frame (SDF) Follow-up messages to transport Matter messages until the Commissionable device joins the operational network. When a **WFA-USD** SDF Follow-up message is used to transport a Matter message, the Service Info field of the Service Descriptor Extension Attribute included in the Wi-Fi USD Follow-up message SHALL be used to encapsulate the Matter message. The value of its Service Protocol Type field SHALL be set to '3' to indicate **Matter**, and its Service Specific Info field SHALL carry the Matter message sent by the Commissionable device or the Commissioner.

5.4.3.3. Using Existing IP-bearing Network

To discover a commissionable device over an existing IP-bearing network connection, the Commissioner SHALL perform service discovery using DNS-SD as detailed in [Section 4.3, “Discovery”](#), and more specifically in [Section 4.3.1, “Commissionable Node Discovery”](#).

5.4.3.4. Using NFC Transport Layer

To discover a commissionable device over NFC interface, the user needs to bring the Commissioner in close proximity to the Commissionee, and the Commissioner SHALL follow the activities defined for POLL Mode in [NFC Forum Activity](#).

5.4.3.5. Using Thread Commissioning Protocol

NOTE | Support for Thread Commissioning Protocol is provisional.

To discover a commissionable device over the Thread commissioning protocol, the Commissioner SHALL first discover, authenticate with and petition the Thread network to become the Commissioner as described in Thread Specification, Sections 8.4.1. and 8.4.2. Once the Matter Commissioner has become the Thread Commissioner, it SHALL configure the Thread Steering Data as outlined in [Section 4.23.1.1, “Thread Commissioning Protocol Steering Data”](#). Additional details on commissioner obligations are captured in [Section 4.23.3.2, “Discovery by Commissioner”](#).

5.5. Commissioning Flows

There are two commissioning flows depending upon the networking capability of the Commissioner and Commissionee, namely concurrent connection commissioning flow and non-concurrent connection commissioning flow.

For additional security requirements related to commissioning flows, refer to [Section 13.6, “Security Best Practices”](#).

A Commissioner and Commissionee with concurrent connection have the ability to maintain two network connections simultaneously. One connection is between the Commissioner (or Commissionee) and the operational network (e.g., home Wi-Fi network or Thread network) that the Commissionee is being programmed to join. The second connection is between the Commissioner and Commissionee for commissioning as is referred to as commissioning channel. A Commissioner and Commissionee with non-concurrent connection capability cannot be simultaneously connected to both the operational network that the Commissionee is being configured to join, and the commissioning channel.

The two connections MAY either be on the same or on different networking interfaces. For example, a Commissioner uses its Wi-Fi interface to connect to the operational network, but use its Bluetooth Low Energy interface for commissioning.

To determine whether a Commissionee has concurrent or non-concurrent connection capability, the Commissioner can use the [Section 11.10.6.5, “SupportsConcurrentConnection”](#) attribute of the [Section 11.10, “General Commissioning Cluster”](#).

The Commissionee communication channel used at the beginning of commissioning MAY be on the same power domain as the operational network interface or MAY be on a different power domain. For example, a Commissionee could use NFC energy harvesting for commissioning through [NFC Transport Layer \(NTL\)](#), but require main power to connect to the operational network.

To determine whether a Commissionee is in a power configuration that allows connecting to the operational network, the Commissioner can use the [IsCommissioningWithoutPower](#) attribute of the [General Commissioning Cluster](#).

Commissioning, including both initial and setup phases, SHALL be a time-bound process that completes before expiration of a fail-safe timer. The fail-safe timer SHALL be set at the beginning of commissioning. If the device is not in a power configuration that allows connecting to the operational network at the end of the initial phase, then the fail-safe timer countdown is paused, and it is resumed during the second phase of commissioning. If the fail-safe timer expires prior to commissioning completion, the Commissioner and Commissionee SHALL behave as specified in [Section 11.10.7.2.2, “Behavior on expiry of Fail-Safe timer”](#). Successful completion of commissioning SHALL disarm the fail-safe timer.

A Commissionee that is ready to be commissioned SHALL accept the request to establish a [PASE](#) session with the first Commissioner that initiates the request. When a Commissioner is either in the process of establishing a [PASE](#) session with the Commissionee or has successfully established a session, the Commissionee SHALL NOT accept any more requests for new PASE sessions until one of the following events occurs:

- session establishment fails,
- the successfully established PASE session is terminated on the commissioning channel (see [Section 4.11.1.4, “CloseSession”](#) in [Section 4.11.1.3, “Secure Channel Status Report Messages”](#)),
- the PASE session is established through [NFC Transport Layer \(NTL\)](#) and Commissionee receives a valid [SELECT command](#)

In the event a [Section 4.11.1.4, “CloseSession”](#) status message is sent or received , or whenever commissioning through [NFC Transport Layer \(NTL\)](#) is interrupted by [SELECT command](#) :

1. If the fail-safe timer is armed, the fail-safe timer SHALL be considered expired and the cleanup steps detailed in [Section 11.10.7.2, “ArmFailSafe”](#) SHALL be executed.
2. If the commissioning window is still open, the Commissionee SHALL continue listening for commissioning requests.

In order to avoid locking out the Commissionee from accepting new PASE session requests indefinitely, a Commissionee SHALL expect a PASE session to be established within 60 seconds of receiving the initial request. This means the Commissionee SHALL expect to receive the [PAKE3](#) message

within 60 seconds after sending a [PBKDFParamResponse](#) in response to a [PBKDFParamRequest](#) message from the Commissioner to establish a PASE session. If the PASE session is not established within the expected time window the Commissionee SHALL terminate the current session establishment using the [INVALID_PARAMETER](#) status code as described in [Section 4.11.1.3, “Secure Channel Status Report Messages”](#).

During commissioning:

- the Commissionee SHALL expose all the supported endpoints and descriptor clusters of the device
- the Commissionee SHALL support the [Access Control cluster](#), the [Basic Information cluster](#), the [General Commissioning cluster](#) and the [Operational Credentials cluster](#)
- if the Commissionee requires network commissioning then it SHALL support the [Network Commissioning cluster](#)
- the Commissionee MAY expose other clusters, but it does not have to support subscriptions for arbitrary clusters
- the Commissionee, when using NFC-based commissioning, MAY reject attribute writes for clusters used during the initial phase of commissioning, therefore the Commissioner SHALL handle these errors

When using NFC-based commissioning, it can be interrupted without waiting for the expiry of the fail-safe timer thanks to the physical tap done by the user on the device. Since the transmission of [SELECT command](#) to interrupt the commissioning is under the control of the commissioner, it can decide whether it should issue it, for example because its internal watchdog fired, or not, filtering out accidental taps.

The commissioning commands and attributes are defined in Clusters (see [Section 11.9, “Network Commissioning Cluster”](#), [Section 11.10, “General Commissioning Cluster”](#), [Section 11.14, “Thread Network Diagnostics Cluster”](#), and [Section 11.15, “Wi-Fi Network Diagnostics Cluster”](#)) and are sent, written, or read using the Interaction Model (see [Chapter 8, *Interaction Model Specification*](#)).

[Figure 47, “Commissioning flow diagram”](#) depicts the commissioning flow between the Commissioner and Commissionee, with optional steps applying if the condition within square brackets is met. The specific steps are described below. Unless indicated otherwise, a commissioner SHALL complete a step, including waiting for any responses to commands it sends in that step, before moving on to the next step.

The initial phase of the Commissioning process consists of the following steps:

1. The Commissioner initiating the commissioning SHALL have regulatory and fabric information available, and SHOULD have accurate date, time and timezone.
2. Commissioner and Commissionee SHALL establish a commissioning channel between each other after discovery or direct physical tap (see [Section 5.4, “Device Discovery”](#)).
3. If the Commissioner and device support [Section 5.7.4.1, “Terms and Conditions \(TC\) Acknowledgement”](#), then the Commissioner SHALL obtain the Terms and Conditions to present to the user from [Section 11.23.6.23, “EnhancedSetupFlowTCUrl”](#).
4. If the Commissioner and device support [Section 5.7.4.1, “Terms and Conditions \(TC\) Acknowledgement”](#), then the Commissioner SHALL obtain the Terms and Conditions to present to the user from [Section 11.23.6.23, “EnhancedSetupFlowTCUrl”](#).

- edgement”, the Commissioner SHALL present the Terms and Conditions to the user following [Section 5.7.4.1, “Terms and Conditions \(TC\) Acknowledgement”](#), unless the Commissioner already has user provided responses that can be used.
5. If the Commissioner and device support [Section 5.7.4.1, “Terms and Conditions \(TC\) Acknowledgement”](#), and Terms and Conditions were presented in step 4, then the Commissioner SHALL receive the user responses to the provided terms for use in step 9.
 6. Commissioner and Commissionee SHALL establish encryption keys with PASE (see [Section 4.14.1, “Passcode-Authenticated Session Establishment \(PASE\)”](#)) on the commissioning channel. All subsequent messages on the commissioning channel are encrypted using PASE-derived encryption keys. Upon completion of PASE session establishment, the Commissionee SHALL autonomously arm the [Fail-safe timer](#) for a timeout of 60 seconds. This is to guard against the Commissioner aborting the Commissioning process without arming the fail-safe, which may leave the device unable to accept additional connections.
 7. Commissioner SHALL re-arm the Fail-safe timer on the Commissionee to the desired commissioning timeout within 60 seconds of the completion of PASE session establishment, using the `ArmFailSafe` command (see [Section 11.10.7.2, “ArmFailSafe”](#)). A Commissioner MAY obtain device information including guidance on the fail-safe value from the Commissionee by reading `BasicCommissioningInfo` attribute (see [Section 11.10.6.2, “BasicCommissioningInfo”](#)) prior to invoking the [Section 11.10.7.2, “ArmFailSafe”](#) command.
 8. If the Commissionee has at least one instance of the [Network Commissioning cluster](#) on any endpoint with either the `WI` (i.e. Wi-Fi) or `TH` (i.e. Thread) [feature flags set](#) in its `FeatureMap`, Commissioner SHALL configure regulatory information in the Commissionee using the [Section 11.10.7.4, “SetRegulatoryConfig”](#) command.
 9. If the [Terms & Conditions Enhanced Setup Flow](#) feature is supported by both the device and by the commissioner, and [Section 11.10.6.9, “TCAcknowledgementsRequired”](#) is `True`, the Commissioner SHALL present them as documented in [Terms and Conditions Acknowledgement](#). The user’s responses SHALL be propagated back to the node with [Section 11.10.7.8, “SetTCAcknowledgements”](#).
 10. Commissioner SHALL establish the authenticity of the Commissionee as a certified Matter device (see [Section 6.2.3, “Device Attestation Procedure”](#)).
 - If the Commissionee fails the Device Attestation Procedure, for any reason, the Commissioner MAY choose to either continue to the Commissioning, or terminate it, depending on implementation-dependent policies.
 - Upon failure of the procedure, the Commissioner SHOULD warn the user that the Commissionee is not a fully trusted device, and MAY give the user the choice to authorize or deny the commissioning. Such a warning enables user choice in Commissionee trust on their Fabric, for development workflows, as well as homebrew device development. Such a warning SHOULD contain as much information as the commissioner can provide about the Commissionee, and SHOULD be adapted to the reason of the failure, for example by being different between the case of an expired certificate versus a revoked PAI certificate.
 - Reasons for failing the Device Attestation procedure MAY include, but are not limited to, the following:
 - The Commissionee being of a device type currently in development or not yet certified (see [certification_type](#) in the [Certification Declaration](#)). Commissioners SHOULD accept a

Certification Declaration with [certification_type=1](#) (provisional) and MAY inform the customer.

- The Commissionee's PAA not being in the Commissioner's trusted set.
 - The Commissioner having obtained knowledge that a [PAA or PAI certificate](#) presented has been revoked, or that the particular [Device Attestation Certificate](#) has been revoked (see [Section 6.2.4, "Device attestation revocation"](#)).
 - The Commissioner cannot obtain a revocation set or cannot obtain updated revocation information from the CA (see [Section 6.2.4.2, "Determining Revocation Status of an Entity"](#)).
 - The Commissionee failing to prove possession of the Device Attestation private key, either by programming error, malicious intent or other reasons.
 - One of the elements of the Commissionee's Device Attestation Certificate chain not meeting the policy validation steps of the Device Attestation Procedure, including errors on validity period.
- If a Commissioner denies commissioning for any reason, it SHOULD notify the user of the reason with sufficient details for the user to understand the reason, so that they could determine if it would be possible to commission the device using a different Commissioner.
11. Following the Device Attestation Procedure yielding a decision to proceed with commissioning, the Commissioner SHALL request operational CSR from Commissionee using the CSRRequest command (see [Section 11.18.6.5, "CSRRequest Command"](#)). The CSRRequest command will cause the generation of a new operational key pair at the Commissionee.
 12. Commissioner SHALL generate or otherwise obtain an Operational Certificate containing Operational ID after receiving the CSRResponse command from the Commissionee (see [Section 11.18.6.5, "CSRRequest Command"](#)), using implementation-specific means.
 13. Commissioner SHALL install operational credentials (see [Figure 58, "Node Operational Credentials flow"](#)) on the Commissionee using the [Section 11.18.6.13, "AddTrustedRootCertificate"](#) and [AddNOC](#) commands, and SHALL use the [Section 11.18.6.11, "UpdateFabricLabel"](#) command to set a string that the user can recognize and relate to this Commissioner/Administrator. The AdminVendorId field of the AddNOC command SHALL be set to a value for which the [Vendor Schema in DCL](#) contains the name and other information of the Commissioner's manufacturer.
 14. Commissioner MAY configure the Access Control List (see [Section 9.10, "Access Control Cluster"](#)) on the Commissionee in any way it sees fit, if the singular entry added by the [AddNOC](#) command in the previous step granting Administer privilege over [CASE](#) authentication type for the Node ID provided with the command is not sufficient to express its desired access control policies.
 - a. The Commissioner MAY read the Commissioning Access Restriction List (see [Section 9.10, "Access Control Cluster"](#)) on the Commissionee in order to identify fabric restrictions. When restrictions exist, the Administrator MAY invoke the ReviewFabricRestrictions command on the same cluster in order to initiate the review process (see [Section 9.10.4.2, "ManagedDevice"](#)), however, the Administrator SHALL wait until after commissioning completes before sending this command.
 15. If the Commissionee both supports it and requires it, the Commissioner SHALL configure the operational network at the Commissionee using commands such as AddOrUpdateWiFiNetwork

(see [Section 11.9.7.3, “AddOrUpdateWiFiNetwork Command”](#)) and `AddOrUpdateThreadNetwork` (see [Section 11.9.7.4, “AddOrUpdateThreadNetwork Command”](#)). A Commissionee requires network commissioning if it is not already on the desired operational network. A Commissionee supports network commissioning if it has any [NetworkCommissioning](#) cluster instances. A Commissioner MAY learn about the networks visible to the Commissionee using `ScanNetworks` command (see [Section 11.9.7.1, “ScanNetworks Command”](#)).

16. The Commissioner SHALL trigger the Commissionee to connect to the operational network using `ConnectNetwork` command (see [Section 11.9.7.8, “ConnectNetwork Command”](#)) unless the Commissionee is already on the desired operational network. In case the device is not fully powered during the [initial phase of commissioning](#), the Commissionee SHALL signal this to the Commissioner via the [IsCommissioningWithoutPower](#) attribute, the fail-safe timer and the Cumulative Fail Safe Context timer countdown SHALL be paused, and the connection to the operational network SHALL be deferred to the setup phase of the Commissioning process.
17. If the [IsCommissioningWithoutPower](#) attribute is true, the Commissioner SHALL wait for an indication that the commissionee is in a power configuration that is capable of joining an operational network before proceeding to the next step.

This concludes the initial phase of the Commissioning process. If the device was fully powered during the initial phase of the Commissioning process then the setup phase SHALL happen immediately after the end of the initial phase, else the setup phase shall start as soon as the device is powered up and running.

The setup phase of the Commissioning process consists of the following steps:

1. If the device was not fully powered during the initial phase of the Commissioning process then it SHALL rearm the failsafe timer and the Cumulative Fail Safe Context timer countdowns with the remaining times that they had when they were paused in the first phase of the Commissioning process, and it SHALL connect to the operational network.
2. The Commissionee SHALL use [Section 4.3.2, “Operational Discovery”](#) to be discoverable on the operational network. An Administrator configured in the ACL of the Commissionee by the Commissioner SHALL use [Section 4.3.2, “Operational Discovery”](#) to discover the Commissionee. This Administrator MAY be the Commissioner itself, or another Node to which the Commissioner has delegated the task. If the device was not fully powered during the initial phase of the Commissioning process, then the Commissionee MAY pause the fail-safe timer and the Cumulative Fail Safe Context timer countdown until the Administrator discovers the Commissionee, to give the Administrator sufficient time to discover the Commissionee.
3. The Administrator SHALL open a CASE (see [Section 4.14.2, “Certificate Authenticated Session Establishment \(CASE\)”](#)) session with the Commissionee over the operational network.
4. The Administrator having established a CASE session with the Commissionee over the operational network in the previous step SHALL invoke the `CommissioningComplete` command (see [Section 11.10.7.6, “CommissioningComplete”](#)). A success response after invocation of the `CommissioningComplete` command ends the commissioning process. In case of NFC-based commissioning, the Commissionee SHALL update the operational discovery service according to the requirement in [Section 4.3.2.5, “Subtypes”](#).

At this point the Commissioner MAY do additional setup like configure UTC time:

1. If the device supports the [Time Synchronization Cluster](#) server:

- The Commissioner SHOULD configure UTC time using the [Section 11.17.9.1, “SetUTCTime”](#) command.
- The Commissioner SHOULD set the time zone using the [Section 11.17.9.3, “SetTimeZone”](#) command, if the [Section 11.17.5.1, “TimeZone”](#) feature is supported.
- The Commissioner SHOULD set the DST offsets using the [Section 11.17.9.5, “SetDSTOffset”](#) command if the [Section 11.17.5.1, “TimeZone”](#) feature is supported, and the [Section 11.17.9.4, “SetTimeZoneResponse”](#) from the Commissionee had the `DSTOffsetRequired` field set to `True`.
- The Commissioner SHOULD set a Default NTP server using the [Section 11.17.9.6, “SetDefaultNTP”](#) command if the [Section 11.17.5.2, “NTPClient”](#) feature is supported and the [Section 11.17.8.5, “DefaultNTP”](#) attribute is null. If the current value is non-null, Commissioners MAY opt to overwrite the current value.
- The Commissioner SHOULD set a trusted time source using the [Section 11.17.9.2, “SetTrustedTimeSource”](#) command if the [Section 11.17.5.4, “TimeSyncClient”](#) feature is supported, the [Section 11.17.8.4, “TrustedTimeSource”](#) attribute is null and there is an available trusted time source on the fabric. The Commissioner SHOULD ensure the ACL on the `TrustedTimeSource` is set to grant the Commissionee View privilege to the Time Synchronization cluster. If the `TrustedTimeSource` is non-null, the Commissioner MAY opt to overwrite the current value with a node from its own fabric by sending the [Section 11.17.9.2, “SetTrustedTimeSource”](#) command.

2. After having used [NFC Transport Layer \(NTL\)](#) as a temporary transport protocol for commissioning, the commissioner SHALL rediscover on the operational network the supported endpoints, clusters, attributes and events of the commissionee. This is necessary since the device capabilities and state exposed in the [initial phase of commissioning](#) (when the device has no operational power) MAY be a subset of the device capabilities in the [setup phase of commissioning](#) and afterwards.

While the Administrator of steps in the [setup phase of commissioning](#) will, in many situations, be the Commissioner Node itself, it MAY be a different Node that was configured by the Commissioner to have Administer privilege against the Commissionee’s [Section 11.10, “General Commissioning Cluster”](#). This is to support flexibility in finalizing the Commissioning. From a Commissionee’s perspective, all Nodes with Administer privilege in the Commissionee’s ACL are equivalent once the Node has a [Node Operational Certificate](#) and associated [Node Operational Identifier](#) on the Fabric into which it was just commissioned.

A Commissioner MAY configure UTC time, Operational ID, and Operational certificates, etc., information over an arbitrary number of interactions at the Commissionee, over the operational network after the commissioning is complete, or over the commissioning channel after PASE-derived encryption keys are established during commissioning.

In concurrent connection commissioning flow the commissioning channel SHALL terminate after successful setup phase step 4 (`CommissioningComplete` command invocation). In non-concurrent connection commissioning flow the commissioning channel SHALL terminate after successful initial phase step 16 (trigger joining of operational network at Commissionee). The PASE-derived encryption keys SHALL be deleted when commissioning channel terminates. The PASE session

SHALL be terminated by both Commissioner and Commissionee once the CommissioningComplete command is received by the Commissionee.

In both concurrent connection commissioning flow and non-concurrent connection commissioning flow, the Commissioner MAY choose to continue commissioning and override the failure in initial phase step 10 (Commissionee attestation).

5.5.1. Commissioning Flows Error Handling

Overall, all Commissioning operations employ actions using cluster attributes and commands that are also, in certain cases, available during normal steady-state operation once commissioned.

If a Commissionee requires network commissioning, the Commissioner SHOULD attempt to configure the [primary network interface](#) on the Root Node endpoint initially.

If the initial attempt fails for networking-related reasons, then the Commissioner SHOULD attempt to configure [secondary network interfaces](#) via additional endpoints that have a server instance of the Network Commissioning cluster, if such endpoints exist.

Before attempting to configure any such alternative interface, the commissioner SHALL revert network configuration changes made as part of the preceding unsuccessful configuration attempt, specifically the [Section 11.9.7.6, “RemoveNetwork”](#) command SHALL be used to remove any Thread or Wi-Fi network configurations added by such an unsuccessful attempt. This ensures that when commissioning ultimately succeeds, only the configuration changes relating to the network interface that was successfully configured will be committed by the [Section 11.10.7.6, “Commissioning-Complete”](#) command.

The Commissioner MAY leverage out-of-band information to prioritize network interfaces, considering factors like the availability of Thread Border Routers, Wi-Fi Access Point status, and user preferences.

Whenever the [Fail-Safe timer](#) is armed, Commissioners and Administrators SHALL NOT consider any cluster operation to have timed-out before waiting at least 30 seconds for a valid response from the cluster server. Some commands and attributes with complex side-effects MAY require longer and have specific timing requirements stated in their respective cluster specification.

Some request commands used for Commissioning and administration have a 'Breadcrumb' argument. When set, this argument SHALL be used to update the value of the [Breadcrumb Attribute](#) as a side-effect of successful execution of those commands. On command failures, the [Breadcrumb Attribute](#) SHALL remain unchanged.

In concurrent connection commissioning flow, the failure of any of the steps from initial phase step 2 through 14 SHALL result in the Commissioner and Commissionee returning to initial phase step 2 (device discovery and commissioning channel establishment) and repeating each step. The failure of any of the steps from initial phase step 15 through setup phase step 4 in concurrent connection commissioning flow SHALL result in the Commissioner and Commissionee returning to initial phase step 15 (configuration of operational network information), and MAY result in an attempt to configure secondary network interfaces. In the case of failure of any of the steps from initial phase step 15 through setup phase step 4 in concurrent connection commissioning flow, the Commissioner and Commissionee SHALL reuse the existing PASE-derived encryption keys over the commis-

sioning channel and all steps up to and including initial phase step 14 are considered to have been successfully completed.

In non-concurrent connection commissioning flow, the failure of any of the steps from initial phase step 2 through 17 SHALL result in the Commissioner and Commissionee returning to initial phase step 2 (device discovery and commissioning channel establishment) and repeating each step. In case `IsCommissioningWithoutPower` attribute was false in initial phase step 17 the failure of any of steps from setup phase step 1 through 4 SHALL result in the Commissioner and Commissionee returning to initial phase step 2 (device discovery and commissioning channel establishment) and repeating each step. In case `IsCommissioningWithoutPower` attribute was true in initial phase step 17, the failure of any of the steps from setup phase step 1 through 4 SHALL result in the Commissioner and Commissionee returning to setup phase step 1 and repeating each step.

Commissioners that need to restart from initial phase step 2 MAY immediately expire the fail-safe by invoking the `ArmFailSafe` command with an `ExpiryLengthSeconds` field set to 0. Otherwise, Commissioners will need to wait until the current fail-safe timer has expired for the Commissionee to begin accepting PASE again.

In both concurrent connection commissioning flow and non-concurrent connection commissioning flow, the Commissionee SHALL exit Commissioning Mode after 20 failed attempts.

Once a Commissionee has been successfully commissioned by a Commissioner into its fabric, the commissioned Node SHALL NOT accept any more PASE requests until any one of the following conditions is met:

- Device is factory-reset.
- Device enters commissioning mode.

Ongoing administration of Nodes by Administrators employs many of the same clusters and constraints related to Fail-Safe timer and cluster operation time-outs used for initial or subsequent Commissioning into new Fabrics. The respective cluster specifications for the [Section 11.18, “Operational Credentials Cluster”](#) and the [Section 11.9, “Network Commissioning Cluster”](#) reflect the necessary usage of the `ArmFailSafe` and `CommissioningComplete` commands of the [Section 11.10, “General Commissioning Cluster”](#) to achieve consistent state during administrative operations.

5.5.2. Commissioning Flow Diagram



Figure 47. Commissioning flow diagram

5.6. Administrator Assisted Commissioning Flows

5.6.1. Introduction

In this method, a current Administrator of a Node first sends the [OpenCommissioningWindow](#) command to the Node over a [CASE](#) session. The new Administrator SHALL already have network connectivity and complete commissioning based on the two flows described below.

The commands for these flows are defined in [Section 11.19, “Administrator Commissioning Cluster”](#).

5.6.2. Basic Commissioning Method (BCM)

This method is **OPTIONAL** for Nodes and Administrators/Commissioners to implement. In this method, the current Administrator SHALL send the [OpenBasicCommissioningWindow](#) command to the Node over a [CASE](#) session. The Node SHALL advertise its presence over DNS-SD (see [Section 5.4.2.7, “Using Existing IP-bearing Network”](#) and [Section 4.3.1, “Commissionable Node Discovery”](#)).

after receiving the [OpenBasicCommissioningWindow](#) command.

The new Administrator's Commissioner then completes commissioning with the Node using similar Commissioning flow as it would do for a factory-new device (although note that IP channel is used for discovery). It can either read the QR code format or NFC tag, or use the Manual Pairing Code format of the [Onboarding Payload](#) of the Node.

The following steps describe a possible sequence of events for BCM commissioning:

1. Current Administrator puts the Node in [OpenBasicCommissioningWindow](#) for a specified time window, and receives success response from the Node on the [OpenBasicCommissioningWindow](#) command.
 - a. When the targeted Node is a [ICD](#), the current Administrator can guide the user to perform some action to 'wake' the device from its sleep cycle.
2. New Administrator completes commissioning within the prescribed window using steps outlined in [Figure 47, "Commissioning flow diagram"](#).

5.6.3. Enhanced Commissioning Method (ECM)

This method SHALL be implemented for Nodes and Commissioners/Administrators. When using ECM, the Node's current Administrator instructs the Node over a [CASE](#) session, to go into [OpenCommissioningWindow](#). It SHALL choose a new [RANDOM](#) passcode and SHALL compute and send the corresponding [PAKE](#) passcode verifier to the Node. Actual value of the passcode SHALL NOT be sent to the Node.

The current Administrator then presents the new passcode and discriminator as described [below](#). The Node SHALL advertise its presence over DNS-SD (see [Section 5.4.2.7, "Using Existing IP-bearing Network"](#) and [Section 4.3.1, "Commissionable Node Discovery"](#)) after receiving the [OpenCommissioningWindow](#) command.

5.6.3.1. Presentation of Onboarding Payload for ECM

Presentation of the passcode and other relevant information SHALL be done at least with one or more of the methods below, depending on the capabilities of the first Administrator opening the OCW:

1. If a user interface display is supported, the temporary Onboarding Payload SHALL be displayed using a textual representation of the Manual Pairing Code, using the [11-digit variant](#): it SHALL NOT contain the [VENDOR_ID](#) or [PRODUCT_ID](#) as the commissioning of the node(s) using the ECM cannot be subject to User-Intent or Custom Flows.
2. If a user interface display is supported, the temporary Onboarding Payload SHOULD also be displayed using the definitions included in [Section 5.1.3, "QR Code"](#) subject to the following constraints:
 - a. If only a single Node is being subjected to the ECM, the Vendor ID and Product ID in the [onboarding payload](#) SHALL be the same as those of that Node.
 - b. If multiple Nodes are being subjected to the ECM using the same [onboarding payload](#), the Vendor ID SHALL be set to [0x0000](#) (Matter Standard) and the Product ID SHALL be set to

0x0000 (consistent with the value used for *not* advertising a Product ID in [Device Announcement](#)).

- c. The **Custom Flow** element SHALL be set to 0 to indicate standard flow.
 - d. The **Discovery Capabilities Mask** SHALL have ONLY bit 2 set to indicate the Node is only discoverable on the IP network.
 - e. The **Passcode** element SHALL be set by the existing Administrator to the same value as the passcode chosen for this ECM operation.
 - f. The **Discriminator** element SHALL be set by the existing Administrator to the same value as the **Discriminator** parameter in [Section 11.19.8.1, “OpenCommissioningWindow”](#).
 - g. If multiple Nodes are subjected to ECM, the [Section 5.1.5, “TLV Content”](#) SHALL contain an entry with **kTag_NumberOfDevices** containing the number of devices that are expected to participate in the commissioning with this ECM operation.
 - h. When the Commissioning Timeout parameter of the OCW command is set to less than the allowed maximum (15 minutes), the [Section 5.1.5, “TLV Content”](#) SHALL contain an entry with **kTag_CommissioningTimeout** containing the value of the Commissioning Timeout parameter used for this ECM operation.
3. If only audio output is supported, the temporary Onboarding Payload SHALL be delivered using a voice prompt of the Manual Pairing Code format. A method SHOULD be available for the user to have the pairing code repeated.

Remote UIs, both visual and audio — such as a manufacturer-specific mobile app or a web UI — are expressly permitted in the set of acceptable mechanisms for conveyance of the onboarding material.

This method allows a current Administrator to set multiple Nodes for commissioning with a new administrator with an appropriate **Commissioning Window**, by opening a commissioning window using the [OpenCommissioningWindow](#) command and sending the **PAKE** passcode verifier to a series of Nodes. The new Administrator uses the information in [Manual Pairing Code](#) to discover the Nodes that are in Commissioning mode and commission them using the new passcode.

The following steps describe a possible sequence of events for ECM commissioning:

1. Current Administrator puts the Node(s) in commissioning mode for a specified time window with a new setup passcode, and receives success responses from the involved Node(s) on the [OpenCommissioningWindow](#) command.
 - a. When one or more **ICD** are among the targeted Nodes, the current Administrator can guide the user to perform some action to 'wake' these devices from their sleep cycle.
2. Current Administrator presents Onboarding Payload as described [above](#).
3. New Administrator completes commissioning within the prescribed window using steps outlined in [Figure 47, “Commissioning flow diagram”](#).

5.6.4. Open Commissioning Window

The following sequence diagram shows steps current Administrator takes to enable [OpenCommissioningWindow](#).




Figure 48. Open Commissioning Window (Administrator A)

5.7. Device Commissioning Flows

This section describes the three different flows for out-of-box commissioning that a Matter device manufacturer may select (while adhering to the restrictions listed in the table in [Section 5.4.2.2, “Announcement Commencement”](#)) for a given product. For each flow, a description is provided which includes actions required to place the device into commissioning mode, fields in the [Onboarding Payload](#) which identify the flow selected by the device manufacturer, fields in the [Distributed Compliance Ledger](#) that provide information used for commissioning and help the Commissioner provide the user with appropriate instructions, and requirements for device packaging relating to the Onboarding Payload. The three flows are the following:

- Standard Commissioning Flow
- User-Intent Commissioning Flow
- Custom Commissioning Flow

Matter device manufacturers SHALL use the [Section 11.23, “Distributed Compliance Ledger”](#) to provide commissioners with information and instructions for both initial and secondary commissioning, and SHOULD use this Ledger to provide links to the user guide, a link to a manufacturer app, and other pre-setup information, to enable an optimal commissioning flow without requiring bilateral arrangements between each commissioner manufacturer and each device manufacturer. Some fields in the Ledger SHALL or SHOULD be populated, depending on the type of commissioning flow, as detailed in the text below and in the [Distributed Compliance Ledger](#) section.

5.7.1. Standard Commissioning Flow

- A Standard Commissioning Flow device SHALL be available for initial commissioning by any Matter commissioner.
- A Standard Commissioning Flow device, when in factory-new state, SHALL start advertising automatically upon power on (see [Commencement](#)).
- A Standard Commissioning Flow device SHALL set [Custom Flow bits](#) in the Onboarding Payload to indicate '0 - Standard Flow'.
- A Standard Commissioning Flow device SHALL follow the rules for [Manual Pairing Code and QR Code Inclusion](#).
- For the case where the device has stopped advertising (e.g. user has powered on the device longer ago than the advertisement period), the manufacturer SHOULD provide guidance about how to bring the device back into advertising mode using the `CommissioningModeInitialStepsHint` field from the [Distributed Compliance Ledger](#). Commissioners SHOULD use this infor-

mation to guide the user for this case.

- When commissioning fails, the commissioner MAY also reference [Distributed Compliance Ledger](#) fields such as `CommissioningFallbackUrl` (see [Section 5.7.5, “Commissioning Fallback Mechanism”](#)), `UserManualUrl`, `SupportUrl` and `ProductUrl` to assist the user in further steps to resolve the issue(s).
- The [Section 11.23, “Distributed Compliance Ledger”](#) entries for Standard Commissioning Flow devices SHALL include the [Section 11.23.6.8, “CommissioningCustomFlow”](#) field set to '0 - Standard' and the [Section 11.23.6.10, “CommissioningModeInitialStepsHint”](#) field set to a non-zero integer value, with bit 0 (Power Cycle) being set to 1. The [Section 11.23.6.11, “CommissioningModeInitialStepsInstruction”](#) field SHALL be set when `CommissioningModeInitialStepsHint` has a mandatory Pairing Instruction dependency and MAY be set when `CommissioningModeInitialStepsHint` has an optional Pairing Instruction dependency.

Table 80. Values of Ledger fields to represent Standard Commissioning Flow

Field Name	Value(s)
<code>CommissioningCustomFlow</code>	0 - Standard (Mandatory)
<code>CommissioningModeInitialStepsHint</code>	This field SHALL be set to a non-zero integer value. See Table 5, “Pairing and Reset Hint Values” for a complete list of pairing instructions. Example value: 33 - The following bits are set: 0 (Power Cycle - Mandatory), 5 (Device Manual - Optional). Bit 1 (Device Manufacturer URL) MAY be set.
<code>CommissioningModeInitialStepsInstruction</code>	The field SHALL be set when <code>CommissioningModeInitialStepsHint</code> has a mandatory Pairing Instruction dependency and MAY be set when <code>CommissioningModeInitialStepsHint</code> has an optional Pairing Instruction dependency. See PI Dependency column of Pairing Hint Table to determine which pairing hints have a mandatory or optional Pairing Instruction dependency and therefore require or allow this field to be populated.

5.7.2. User-Intent Commissioning Flow

- A User-Intent Commissioning Flow device SHALL be available for initial commissioning by any Matter commissioner.
- A User-Intent Commissioning Flow device, when in factory-new state, SHALL NOT start advertising automatically upon application of power (see [Section 5.4.2.2, “Announcement Commencement”](#)).
- To place a User-Intent Commissioning Flow device into advertising mode, some form of user interaction with the device beyond application of power is required (see [Table 5, “Pairing and Reset Hint Values”](#)). If a Device Manufacturer setup artifact is required for this, beyond docu-

mentation, then the device is a [Section 5.7.3, “Custom Commissioning Flow”](#) device and not a User-Intent Commissioning Flow device. The documentation MAY be printed or in the form of online documentation (e.g. [Section 11.23.6.16, “UserManualUrl”](#)).

- A User-Intent Commissioning Flow device SHALL follow the rules for [Manual Pairing Code and QR Code Inclusion](#).
- The [Section 11.23, “Distributed Compliance Ledger”](#) entries for User-Intent Commissioning Flow devices SHALL include the [Section 11.23.6.8, “CommissioningCustomFlow”](#) field set to '1 - User Intent' and the [Section 11.23.6.10, “CommissioningModeInitialStepsHint”](#) field set to a non-zero integer value. Bit 0 (Power Cycle) in the [Section 11.23.6.10, “CommissioningModeInitialStepsHint”](#) field SHALL be set to 0. The [Section 11.23.6.11, “CommissioningModeInitialStepsInstruction”](#) field SHALL be set when CommissioningModeInitialStepsHint has a Pairing Instruction dependency and MAY be set when CommissioningModeInitialStepsHint has an optional Pairing Instruction dependency.
- A User-Intent Commissioning Flow device SHALL set [Custom Flow bits](#) in the Onboarding Payload to indicate '1 - User Intent'.
- The commissioner SHOULD reference [Distributed Compliance Ledger](#) fields such as [Section 11.23.6.10, “CommissioningModeInitialStepsHint”](#), [Section 11.23.6.11, “CommissioningModeInitialStepsInstruction”](#), [UserManualUrl](#), and [SupportUrl](#) to assist the user during commissioning, e.g. to explain how to bring the device into commissioning mode.

Table 81. Values of Ledger fields to represent User-Intent Commissioning Flow

Field Name	Value(s)
CommissioningCustomFlow	1 - User Intent (Mandatory)
CommissioningModeInitialStepsHint	This field SHALL be set to a non-zero integer value. See Table 5, “Pairing and Reset Hint Values” for a complete list of pairing instructions. Example value: 96 - The following bits are set: 6 (Press Reset Button - Optional), 5 (Device Manual - Optional). Bit 1 (Device Manufacturer URL) MAY be set.
CommissioningModeInitialStepsInstruction	The field SHALL be set when CommissioningModeInitialStepsHint has a mandatory Pairing Instruction dependency and MAY be set when CommissioningModeInitialStepsHint has an optional Pairing Instruction dependency. See PI Dependency column of Pairing Hint Table to determine which pairing hints have a mandatory or optional Pairing Instruction dependency and therefore require or allow this field to be populated.

5.7.3. Custom Commissioning Flow

- A Custom Commissioning Flow device SHALL require interaction with custom steps, guided by a service provided by the manufacturer for initial device setup, before it can be commissioned by any Matter commissioner.
- A Custom Commissioning Flow device that supports the [Enhanced Setup Flow](#) SHALL include a usable [Onboarding Payload](#) on-device or in packaging, so that a Commissioner which supports the [Enhanced Setup Flow](#) can commission the device using the information provided in the Onboarding Payload.
 - Such a device SHALL follow the rules for [Manual Pairing Code and QR Code Inclusion](#).
- A Custom Commissioning Flow device which does not support [Enhanced Setup Flow](#) MAY include the [Onboarding Payload](#) on-device or in packaging.
 - If the Onboarding Payload is included, then it SHALL follow the rules for [Manual Pairing Code and QR Code Inclusion](#).
 - If the Onboarding Payload is not included, then it SHALL be provided to the user through other means provided by the manufacturer.
- A Custom Commissioning Flow device SHALL set [Custom Flow bits](#) in the Onboarding Payload to indicate '2 - Custom'.
- The [Distributed Compliance Ledger](#) entries for Custom Commissioning Flow devices SHALL include:
 - the [CommissioningCustomFlow](#) field set to '2 - Custom'
 - the [CommissioningModeInitialStepsHint](#) with bit 0 (Power Cycle) set to 0 and bit 1 (Device Manufacturer URL) set to 1
 - the [CommissioningCustomFlowUrl](#) field populated in order to indicate to commissioners that initial commissioning can only be completed by the user visiting the URL contained therein.

This URL will typically lead to a web page with relevant instructions and/or to a server which (e.g. by looking at the User-Agent) redirects the user to allow viewing, downloading, installing or using a manufacturer-provided means for guiding the user through the process and bring the device into a state that it is available for commissioning by any commissioner. Since the URL is retrieved from a DCL entry corresponding to a specific VID and PID combination, the device manufacturer MAY choose to use any constructed URL valid in a HTTP GET request (i.e. dedicated for the product the user wants to commission) such as, for example, <https://domain.example/download-install-app?vid=FFF1&pid=1234>.
- When a Commissioner encounters a device with [Custom Flow](#) field (in Onboarding Payload) or its [CommissioningCustomFlow](#) field (in [Distributed Compliance Ledger](#)) set to '2 - Custom', it SHOULD use the [CommissioningCustomFlowUrl](#) to guide the user on how to proceed, unless it has alternative means to guide the user to successful commissioning (for example, if the commissioner supports the set of [EnhancedSetupFlowOptions](#) required by the device).
 - If a Commissioner follows or launches the [CommissioningCustomFlowUrl](#) after a User request, it SHALL expand it as described in [Section 5.7.3.1, “CommissioningCustomFlowUrl format”](#).
- If [EnhancedSetupFlowOptions](#) indicates to use the Enhanced Setup Flow and the flags indicated

are supported by the Commissioner, then the commissioner SHOULD perform the required steps described in [Enhanced Setup Flow](#) rather than following the [CommissioningCustomFlowUrl](#).

- A manufacturer contemplating using this flow should realize that
 - This flow typically requires internet access to access the URL, so initial commissioning of the device may fail if there is no internet connection at that time/location.
 - If the flow requires an app, it needs to be made available for popular platforms amongst the user population; some of their platforms running a commissioner (e.g. a smart speaker not running a popular mobile OS) may thus not be able to be used for the initial commissioning of such devices.

Table 82. Values of Ledger fields to represent Custom Commissioning Flow

Field Name	Value(s)
CommissioningCustomFlow	2 - Custom (Mandatory)
CommissioningCustomFlowUrl	'URL' - Device Manufacturer URL (Mandatory)
CommissioningModeInitialStepsHint	This field SHALL be set to a non-zero integer value with at least bit 1 set (Device Manufacturer URL). See Table 5, “Pairing and Reset Hint Values” for a complete list of pairing instructions. Example value: 2 - The following bits are set: 1 (Device Manufacturer URL) (Mandatory).
CommissioningModeInitialStepsInstruction	The field SHALL be set when CommissioningModeInitialStepsHint has a mandatory Pairing Instruction dependency and MAY be set when CommissioningModeInitialStepsHint has an optional Pairing Instruction dependency. See PI Dependency column of Pairing Hint Table to determine which pairing hints have a mandatory or optional Pairing Instruction dependency and therefore require or allow this field to be populated.

5.7.3.1. CommissioningCustomFlowUrl format

The [CommissioningCustomFlowUrl](#) MAY contain a query component (see [RFC 3986](#) section 3.4). If a query is present, it SHALL be composed of one or more key-value pairs:

- The query SHALL use the [&](#) delimiter between key/value pairs.
- The key-value pairs SHALL be in the format [name=<value>](#) where [name](#) is the key name, and [<value>](#) is the contents of the value encoded with proper URL-encoded escaping.
- If key [MTcu](#) is present, it SHALL have a value of ["_"](#) (i.e. [MTcu=_](#)). This is the "callback URL ([Call-backUrl](#)) placeholder".

- If key **MTop** is present, it SHALL have a value of "_" (i.e. **MTop=_**). This is the "onboarding payload placeholder".
- Any key whose name begins with **MT** not mentioned in the previous bullets SHALL be reserved for future use by this specification. Manufacturers SHALL NOT include query keys starting with **MT** in either the **CommissioningCustomFlowUrl** or **CallbackUrl** unless they are referenced by a version of this specification.

When the **CommissioningCustomFlowUrl** for a Custom Commissioning Flow device includes the **MTop** key, the **Passcode** embedded in any Onboarding Payload placed on-device or in packaging SHALL NOT be one that can be used for secure channel establishment with the device. This requirement is intended to ensure a shared secret used for proof of possession will not be transferred to a server without user consent. A Custom Commissioning Flow device MAY utilize Onboarding Payload fields such as the Serial Number (see **kTag_SerialNumber**) to pass device identification to the server specified in **CommissioningCustomFlowUrl**, as these fields by themselves could not be used to gain access to the device on their own like the Passcode could.

When the **CommissioningCustomFlowUrl** for a Custom Commissioning Flow device includes the **MTop** key, the **Passcode** embedded in any Onboarding Payload placed on-device or in packaging MAY be set to 0 (one of the **invalid values**) in order to provide a hint to the Commissioner that it is not one that can be used for secure channel establishment with the device. This would allow the Commissioner to avoid attempting to commission the device if an advertisement from it is detected.

Any other element in the **CommissioningCustomFlowUrl** query field not covered by the above rules, as well as the fragment field (if present), SHALL remain as obtained from the **Distributed Compliance Ledger's CommissioningCustomFlowUrl** field, including the order of query key/value pairs present.

5.7.3.1.1. Expansion of CommissioningCustomFlowUrl by Commissioner

Once the URL is obtained, it SHALL be expanded to form a final URL (**ExpandedCommissioningCustomFlowUrl**) by proceeding with the following substitution algorithm on the original **CommissioningCustomFlowUrl**:

1. If key **MTcu** is present, compute the **CallbackUrl** desired (see [Section 5.7.3.2, "CallbackUrl format for Custom Commissioning Flow response"](#)), and substitute the placeholder value "_" (i.e. in **MTcu=_**) in the **CommissioningCustomFlowUrl** with the desired contents, encoded with proper URL-encoded escaping (see [RFC 3986](#) section 2).
2. If key **MTop** is present, substitute the placeholder value "_" (i.e. in **MTop=_**) in the **CommissioningCustomFlowUrl** with either **Manual Pairing Code**, or **QR code body** including the **MT:** prefix and TLV data (if present), encoded with proper URL-encoded escaping (see [RFC 3986](#) section 2). Note that both methods SHOULD be supported by the Manufacturer's custom flow.

A Commissioner SHALL NOT append the **MTop=** query key/value pair unless the key/value pair was already present as **MTop=_** in the **CommissioningCustomFlowUrl** previously obtained. This constraint enables the determination of which products make use of the payload in their Custom Commissioning Flow infrastructure by inspection of the Distributed Compliance Ledger records.

The final URL after expansion (**ExpandedCommissioningCustomFlowUrl**) SHALL be the one to follow per [Section 5.7.3, "Custom Commissioning Flow"](#), rather than the original value obtained from the Distributed Compliance Ledger.

5.7.3.2. CallbackUrl format for Custom Commissioning Flow response

If a **CallbackUrl** field (i.e. **MTcu=**) query field placeholder is present in the **CommissioningCustomFlowUrl**, the Commissioner MAY replace the placeholder value "_" in the **ExpandedCommissioningCustomFlowUrl** with a URL that the manufacturer custom flow can use to make a smooth return to the Commissioner when the device is in a state that it can be commissioned.

This URL field MAY contain a query component (see [RFC 3986](#) section 3.4).

If a query is present, it SHALL be composed of one or more key-value pairs:

- The query SHALL use the **&** delimiter between key/value pairs.
- The key-value pairs SHALL follow the format **name=<value>** where **name** is the key name, and **<value>** is the contents of the value encoded with proper URL-encoded escaping.
- If key **MTrop** is present, it SHALL have a value of "_" (i.e. **MTrop=**_). This is the placeholder for a "returned onboarding payload" provided by the manufacturer flow to the Commissioner.
- Any key whose name begins with **MT** not mentioned in the previous bullets SHALL be reserved for future use by this specification.

Any other element in the **CallbackUrl** query field not covered by the above rules, as well as the fragment field (if present), SHALL remain as provided by the Commissioner through embedding within the **ExpandedCommissioningCustomFlowUrl**, including the order of query key/value pairs present.

5.7.3.2.1. Expansion of CallbackUrl by the manufacturer custom flow

Once the **CallbackUrl** is obtained by the manufacturer flow, it MAY be expanded to form a final **ExpandedCallbackUrl** URL to be used by proceeding with the following substitution algorithm on the provided **CallbackUrl**:

- If key **MTrop** is present, the manufacturer custom flow having received the initial query containing the **CallbackUrl** MAY compute an **Onboarding Payload** in QR code format including **MT:** prefix, and substitute the placeholder value "_" (i.e. in **MTrop=**_) in the **CallbackUrl** with the desired contents, encoded with proper URL-encoded escaping (see [RFC 3986](#) section 2).
 - The contents of the **MTrop=**_ key/value pair in the **ExpandedCallbackUrl** SHALL only be expanded if the manufacturer custom flow, having received the initial query containing the **CallbackUrl**, supports opening a commissioning window on the target device and supports conveying the corresponding onboarding payload to the Commissioner.
 - The return onboarding payload, if provided, SHALL contain an ephemeral **Passcode** and not a permanent code that can be used in a subsequent commissioning window. If the manufacturer wants the **Passcode** embedded in the Onboarding Payload placed on-device or in packaging to be the one used for session establishment with the Commissioner, then the manufacturer SHALL NOT include the **MTop** key in its **CommissioningCustomFlowUrl** and SHALL NOT populate the **MTrop** value in the **CallbackUrl** expansion.
 - The contents of the return onboarding payload, if provided, SHALL be constructed to match the state of the device at the moment the **ExpandedCallbackUrl** is opened. At least one ingredient which needs to be adapted relative to the received Onboarding Payload is the **Custom Flow** field which needs to be **0** for the return onboarding payload.

- The presence of this field is to assist automatically resuming commissioning without additional data entry (QR code or Manual Pairing Code) by the user at the Commissioner that initially triggered the custom flow. The manufacturer custom flow SHOULD provide an alternate means of conveying the onboarding payload, such as a manual pairing code.
- Note that if the information in the initial onboarding payload that caused triggering of a Custom Commissioning Flow was directly usable, it may be used by the Commissioner, either upon being triggered through the `ExpandedCallbackUrl` having been opened, or autonomously as a fallback.
- Commissioners providing a `CallbackUrl` to the manufacturer custom flow through the `ExpandedCommissioningCustomFlowUrl` SHOULD support using the `ExpandedCallbackUrl` to trigger resumption of Commissioning flow if the `ExpandedCallbackUrl` is followed, otherwise the Commissioner SHOULD NOT substitute the `MTcu` query field when expanding the `CommissioningCustomFlowUrl` into the `ExpandedCommissioningCustomFlowUrl`.
- If the manufacturer custom flow failed to make the device commissionable, it SHALL NOT replace the placeholder value "_" of an included `MTrop=_` key/value pair, to avoid a Commissioner attempting to discover or commission a device not made ready by the custom flow.

A manufacturer custom flow having received an `ExpandedCommissioningCustomFlowUrl` SHOULD attempt to open the `ExpandedCallbackUrl`, on completion of the steps, if an `ExpandedCallbackUrl` was computed from the `CallbackUrl` and opening such a URL is supported.

5.7.3.3. Examples of CommissioningCustomFlow URLs

Below are some examples of valid `ExpandedCommissioningCustomFlowUrl` for several valid values of `CommissioningCustomFlowUrl`, as well as some examples of invalid values of `CommissioningCustomFlowUrl`:

- Valid URL with no query string:
 - Before expansion: <https://company.domain.example/matter/custom/flows/vFFF1p1234>
 - After expansion: <https://company.domain.example/matter/custom/flows/vFFF1p1234> (no change)
- Invalid URL with no query string: `http` scheme is not allowed:
 - <http://company.domain.example/matter/custom/flows/vFFF1p1234>
- Valid URL with basic manufacturer-specific scheme for query:
 - Before expansion: <https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234>
 - After expansion: <https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234> (no change)
- Valid URL with `MTop=_` placeholder using QR format onboarding payload embedding:
 - Before expansion: https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234&MTop=_
 - After expansion: <https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234&MTop=MT%3A-MOA57ZU02IT2L2BJ00>

- Onboarding payload QR content **MT: -M0A57ZU02IT2L2BJ00** was embedded within **MTop** key
- Valid URL with **MTop=_** placeholder using Manual Pairing Code onboarding payload embedding:
 - Before expansion: https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234&MTop=_
 - After expansion: <https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234&MTop=610403146665521046600>
 - Onboarding Manual Pairing Code **610403146665521046600** was embedded within **MTop** key
- Valid URL with **MTop=_** placeholder using Manual Pairing Code onboarding payload embedding, using a different order of keys/value pairs than the previous example:
 - Before expansion: https://company.domain.example/matter/custom/flows?pid=1234&MTop=_&vid=FFF1
 - After expansion: <https://company.domain.example/matter/custom/flows?pid=1234&MTop=610403146665521046600&vid=FFF1>
 - Onboarding Manual Pairing Code **610403146665521046600** was embedded within **MTop** key
- Valid URL with onboarding payload elided (because commissioner could not provide it):
 - Before expansion: https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234&MTop=_
 - After expansion: https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234&MTop=_ (no change)
- Valid URL, return onboarding payload and **CallbackUrl** requested:
 - Before expansion:

```
https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234&MTop=_&MTcu=_
```

- After expansion:

```
https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234&MTop=MT%3A-M0A57ZU02IT2L2BJ00&MTcu=https%3A%2F%2Fcommissioner.domain.example%2Fcb%3Ftoken%3DmAsJ6_vqbr-vjDiG_w%253D%253D%26MTrop%3D_
```

- The **ExpandedCommissioningCustomFlow** URL contains:
 - An embedded onboarding payload QR content value of **MT: -M0A57ZU02IT2L2BJ00**
 - A **CallbackUrl** with a Commissioner-provided arbitrary **token=** key/value pair and the **MTrop=** key/value pair place-holder to indicate support for a return onboarding payload: https://commissioner.domain.example/cb?token=mAsJ6_vqbr-vjDiG_w%3D%3D&MTrop=_
 - After expansion of the **CallbackUrl** (**MTcu** key) into an **ExpandedCallbackUrl**, with an example return onboarding payload of **MT: -M0A5.GB00V68T62010**, the **ExpandedCallbackUrl** would be:


```
https://commissioner.domain.example/cb?token=mAsJ6_vqbr-
vjDiG_w%3D%3D&MTrop=MT%3A-MOA5.GB00V68T62010
```

Note that the **MTcu** key/value pair was initially provided URL-encoded within the **ExpandedCommissioningCustomFlow** URL and the **MTrop=_** key/value pair placeholder now contains a substituted returned onboarding payload.

- Invalid URL, due to **MTza=79** key/value pair in reserved **MT**-prefixed keys reserved for future use:
 - https://company.domain.example/matter/custom/flows?vid=FFF1&pid=1234&MTop=_&MTza=79

5.7.3.4. Example Custom Commissioning Flow

An example of this flow is illustrated below. The "DCL info" concept denotes that the Commissioner SHALL collect the information from the DCL via some mechanism, such as a network resource accessible to the Commissioner containing a replicated set of the DCL's content.




Figure 49. Custom Commissioning Flow sequence diagram

In the flow above:

- In the final steps, the User might have to perform the trigger to the first Commissioner, so that it can start or continue the commissioning process.
- If possible, a Commissioner MAY continue to scan for announcements from the device in the background while any manufacturer-specific app is configuring the device to be available for commissioning. The Commissioner might need a new OnboardingPayload provided to the User by the Manufacturer Website or App.
- In order to simplify the flow, the Commissioner MAY:

- Include the onboarding payload obtained from the user (see **MTop** key in [Section 5.7.3.1, “CommissioningCustomFlowUrl format”](#)) within the **CommissioningCustomFlowUrl**.
- Include a callback URL (see **MTcu** key in [Section 5.7.3.1, “CommissioningCustomFlowUrl format”](#)) within the **ExpandedCommissioningCustomFlowUrl**.
- The Manufacturer Website or App MAY utilize the **CallbackUrl** field, if provided in the query string, in order to simplify the process for signaling the completion of the manufacturer-specific part of the flow back to the Commissioner. When doing so, the Manufacturer Website or App SHOULD put the device into Commissioning mode and SHOULD provide the corresponding onboarding payload to the Commissioner using the **MTrop** key/value pair within the **ExpandedCallbackUrl**.

5.7.4. Enhanced Setup Flow (ESF)

The Enhanced Setup Flow feature provides an optional way for a Custom Commissioning Flow device to specify in a standardized way a set of additional commissioning steps that a device requires which are not included in the Standard Commissioning Flow performed by commissioners. For example, some devices require the user to read and accept a set of terms and conditions provided by the manufacturer before the device setup can complete. Since Standard Commissioning Flow does not include this functionality, a device with this requirement is compelled to use Custom Commissioning Flow. The Enhanced Setup Flow feature allows a commissioner which understands how to perform the additional commissioning steps required by a device to replace the Custom Commissioning Flow by performing the additional commissioning steps as part of an enhanced Standard Commissioning Flow.

For the Enhanced Setup Flow feature, additional commissioning steps are defined such as the need for a Terms and Conditions acknowledgement. A Custom Commissioning Flow device can indicate via [Enhanced Setup Flow Options](#) whether it supports the Enhanced Setup Flow feature and if so, which of the defined common additional commissioning steps it requires.

- A manufacturer contemplating using this flow has to consider that:
 - Enhanced Setup Flow needs to be supported by the Commissioner in order to be used during commissioning.
 - If Enhanced Setup Flow is not supported by the commissioner, then the commissioner and device both need to support [Section 5.7.3, “Custom Commissioning Flow”](#) in order for the device to be successfully commissioned.

Table 83. Enhanced Setup Flow Options Values

Bit index	Name	Description
0	TC on initial Commission	A commissioner is required to show Terms and Conditions content if this is the initial commissioning for the device out of a factory-reset state.

Bit index	Name	Description
1	Disallow user TC responses reuse for required TC	The commissioner SHALL NOT reuse user's responses to the required Terms and Conditions from commissioning a previous product (same VID) with the same T&C.
2	Disallow user TC responses reuse for non-required TC	The commissioner SHALL NOT reuse user's responses to the non-required Terms and Conditions from commissioning a previous product (same VID) with the same T&C.
3	Disallow user TC responses reuse between PIDs	The commissioner SHALL NOT reuse user's responses to the Terms and Conditions from commissioning a previous product (same VID, other PID) with the same T&C.

5.7.4.1. Terms and Conditions (TC) Acknowledgement

One additional step defined in the Enhanced Setup Flow feature involves the display and collection of user acknowledgement of a set of Terms and Conditions defined by the device manufacturer. Commissioners that support this feature SHALL obtain the set of Terms and Conditions texts from the URL defined in [Enhanced Setup Flow Url](#). This file SHALL only be used when successfully validated using [Section 11.23.6.25, "EnhancedSetupFlowTCDigest"](#), [Section 11.23.6.26, "EnhancedSetupFlowTCFileSize"](#), and the [TC File Format](#). When the ESF file is not found, or validation fails, the Commissioner SHALL proceed using the Custom Commissioning Flow.

The Commissioner SHALL present the terms in a such a way that the user is given primary focus on the text and the text is not hidden, obscured, or omitted, unless allowed by exceptions indicated in this section. The user SHALL be given an option to either accept or decline each of the provided texts. If the user declines any text that are marked as required, then the commissioner SHALL end commissioning and notify the user that the device cannot be commissioned due to their refusal of the terms.

5.7.4.2. Reuse of Cached Acknowledgements

In certain cases, such as commissioning multiple instances of the same or similar product, the device manufacturer might not require repeated acknowledgements for each instance.

Examples (the corresponding setting of the bits is indicated in the table below)

1. A manufacturer of motion sensors (same PID) requires acknowledgement of the T&C only when commissioning the first motion sensor - but not for a Smoke/CO sensor (other PID) from same manufacturer - not allowing reuse of acceptance between PIDs.

2. A manufacturer of luminaires of different sizes (various PIDs) requires acknowledgement of the T&C only when commissioning the first luminaire - allowing reuse of acceptance between PIDs with same T&C.
3. A manufacturer of a range of cameras has both required and non-required T&C. The non-required T&C concern uploading video stills to a manufacturer's server for e.g. storage and delayed replay. The user's consent to the non-required T&C (for remote data storage) can be different on how/where the camera is used. For example, they might want to allow this for a camera overlooking the front yard, but not for a camera in a bedroom. In this case, the "Disallow user TC responses reuse for non-required TC" bit is set to 1, which would make the commissioner present the required T&C only once, but present the non-required T&C for every camera device being commissioned, so the user can make a decision on remote data storage for each camera separately.

The table below defines when caching MAY be used, based on the bits "Disallow user TC responses reuse for required TC", "Disallow user TC responses reuse for non-required TC" and "Disallow user TC responses reuse between PIDs" of [EnhancedSetupFlowOptions](#). In conditions marked with **yes** in the table, a commissioner MAY cache a user's acknowledgements and use them for future commissioning of products with the same T&C. In conditions marked with **no**, the commissioner SHALL NOT use a cached user's acknowledgements.

Table 84. Enhanced Setup Flow - Use of Cached Acknowledgements

Bit 1	Bit 2	Bit 3	same PID, required T&C	same PID, non-required T&C	other PID, required T&C	other PID, non-required T&C	example
0	0	0	yes	yes	yes	yes	2
0	1	0	yes	no	yes	no	3
0	0	1	yes	yes	no	no	1
0	1	1	yes	no	no	no	
1	x	x	no	no	no	no	

Reuse of cached acknowledgements SHALL NOT be used when the VID for the product currently being commissioned is different from the product for which the user has reviewed and acknowledged the T&C.

If a Commissioner caches Terms and Conditions acknowledgement responses, then it SHOULD also provide the user a way to remove or update the cached responses. For example, an option could be provided to the user for removing all saved responses, ensuring they would be prompted with Terms and Conditions the next time they commissioned a device which previously had cached responses.

5.7.4.2.1. Definition of "same T&C"

The above text uses the phrase "same T&C". Caching SHALL NOT be done when the ESF file presented to the user has changed between when they reviewed and acknowledged it and the current commissioning, i.e. if values of one or both of [Section 11.23.6.25](#), "EnhancedSetupFlowTCDigest"

and [Section 11.23.6.24, “EnhancedSetupFlowTCRevision”](#) have changed.

5.7.4.3. Use of T&C Acknowledgment during the Custom Commissioning Flow

When [Custom Commissioning Flow](#) is used for displaying Terms and Conditions, the manufacturer-provided means for obtaining user consent SHALL ensure that commissioning can be completed by the original commissioner without performing Terms and Conditions steps by setting [TCAcknowledgementsRequired](#) to False, [TCAcceptedVersion](#) to the accepted version of the Terms and Conditions, and [TCAcknowledgements](#) with the responses provided by the user.

5.7.4.4. Update of T&C as part of a product software update

A manufacturer contemplating using Terms and Conditions should realize that extending or adding them to a product using software updates needs to be considered with great care and SHOULD be avoided to prevent consumer frustration. For example:

- Extending the list of required terms (more requirements that the user SHALL agree to) has the risk that the user will not accept the new terms and be confronted with a device which will no longer function as before.
- Adding Terms and Conditions to a product which previously did not have them will likely result in confusion with the user who could be confronted with a Terms and Conditions dialog after the update.

5.7.4.5. Example of Enhanced Setup Flow for Terms and Conditions




Figure 50. Enhanced Setup Flow TC commissioning sequence diagram

5.7.4.6. Presenting Updated Terms and Conditions

When required by the manufacturer, a device MAY update [TCMinRequiredVersion](#) and set [TCAcknowledgementsRequired](#) to True, indicating to an Administrator that updated Terms and Conditions are available for presentation to the user. The Administrator SHOULD subscribe to the [TCAc-](#)

[knowledgeRequired](#) attribute in order to detect this situation. The Administrator SHOULD present the updated Terms and Conditions in a similar manner to how they would be presented during the commissioning flow, as documented in [Terms and Conditions Acknowledgement](#). The user's responses SHALL be propagated back to the node with [SetTCAcknowledgements](#).

Until new acknowledgements have been received, the device MAY operate with limited functionality related to such functionality for which the updated Terms & Conditions are more restrictive than the agreed ones - until new acknowledgements have been received. Since Administrators are recommended but not required to present updated Terms and Conditions to the user, any such limitation in functionality could lead to a potentially confusing and frustrating experience for the user. As a result, this condition SHOULD be avoided by the manufacturer if possible. If a device opts to enforce updated Terms and Conditions in this way, then it SHOULD inform the user of this requirement in the provided Terms and Conditions and use [TCUpdateDeadline](#) to indicate the deadline before any action takes place. If [TCUpdateDeadline](#) is non-null, then the Administrator SHOULD inform the user by when they must accept the updated Terms and Conditions to avoid any impact to device functionality. Any interactions which cannot succeed due to this functionality limitation SHALL return the `TERMS_AND_CONDITIONS_CHANGED` failure status code. An administrator encountering this status code SHOULD direct the user to the URL located at [EnhancedSetupFlowMaintenanceUrl](#), where the user can be directed on required actions for resolution. Upon satisfactory acceptance of the updated Terms and Conditions, the device SHALL operate with the functionality of the related features, taking into account the user response.

Where re-use of T&C acceptance between products is allowed (see [Section 5.7.4.2, “Reuse of Cached Acknowledgements”](#)), upon acceptance of acknowledgements for a new [Section 11.23.6.24, “EnhancedSetupFlowTCRevision”](#) by the user, the Administrator MAY set the latest acknowledgements to all applicable devices using [SetTCAcknowledgements](#).




Figure 51. Presenting updated Enhanced Setup Flow features diagram

5.7.4.7. Terms and Conditions File Format

The file contained at [Section 11.23.6.23, “EnhancedSetupFlowTCUrl”](#) is a JSON file containing the localized version of the texts to be displayed to the user for the Enhanced Setup Flow Terms & Con-

ditions. Encoding of the file SHALL be UTF-8.

5.7.4.7.1. **schemaVersion Field**

This field SHALL indicate the version of the JSON schema format. This helps in managing changes to the structure of the schema over time. Whenever there is a change to the schema (such as adding, removing, or modifying fields), the **schemaVersion** SHALL be incremented. This ensures that the software processing these configuration files can recognize which version of the schema it is dealing with and handle it appropriately.

This field SHALL be present in the ESF file.

5.7.4.7.2. **esfRevision Field**

This field is a 16-bit integer version number of the file that SHALL match the Version at [Section 11.23.6.24, “EnhancedSetupFlowTCRevision”](#).

This field SHALL be present in the ESF file.

5.7.4.7.3. **defaultCountry Field**

This field is an ISO 3166-1 alpha-2 country code string that indicates the default country to be used if the user’s country is not defined in the file.

This field SHALL be present in the ESF file.

5.7.4.7.4. **countryEntries Field**

This field contains a list of terms and condition entries using the associated country as the key.

This field SHALL be present in the ESF file and contain at least one entry.

5.7.4.7.5. **defaultLanguage Field**

This field is an ISO 639-1 language code string that indicates the default language to be used if the user’s language is not defined for a given country.

This field SHALL be present in the ESF file for each entry in countryEntries.

5.7.4.7.6. **languageEntries Field**

This field contains a list of the different texts to be displayed using the associated language as the key.

This field SHALL be present in the ESF file for each entry in countryEntries and contain at least one entry.

5.7.4.7.7. **ordinal Field**

This field SHALL specify which bit number this text maps to in [Section 11.10.6.8, “TCAcknowledgements”](#). Each bit in the map SHALL keep the same purpose through the life of the product to ensure a correct interpretation of a user’s response across device firmware versions.

This field SHALL be present in the ESF file for each entry in countryEntries.

5.7.4.7.8. required Field

This field is a boolean that indicates whether the corresponding text is required for the device's functionality. If set to true, the user must accept this condition to commission the device.

This field SHALL be present in the ESF file for each entry in the list of terms under the language key in languageEntries.

5.7.4.7.9. title Field

This field is a string that represents the title to use for the terms and conditions subsection when presenting the text. It MAY contain the following subset of HTML tags and their corresponding closing tags to allow for basic formatting of the presented title. Other HTML tags SHALL NOT be used. The permitted HTML tags are: `<b>`, `<em>`, `<i>`, `<small>`, `<strong>`, and `<u>`.

This field SHALL be present in the ESF file for each entry in the list of terms under the language key in languageEntries.

5.7.4.7.10. text Field

This field is a string that contains the text to be presented to the user. It MAY contain the following subset of HTML tags and their corresponding closing tags to allow for formatting of the presented terms. Other HTML tags SHALL NOT be used. The permitted HTML tags are: `<b>`, `<br>`, `<em>`, `<h1>`, `<h2>`, `<h3>`, `<h4>`, `<h5>`, `<h6>`, `<hr>`, `<i>`, `<li>`, `<ol>`, `<p>`, `<small>`, `<strong>`, `<u>`, and `<ul>`.

This field SHALL be present in the ESF file for each entry in the list of terms under the language key in languageEntries.

5.7.4.8. Terms and Conditions File Localization

The locale selected for the TC texts SHALL be a [BCP47](#) language tag. It SHALL be selected with the user's preferred language and an ISO 3166-1 alpha-2 country code. The country code SHALL represent the country, dependent territory, or special area of geographic interest in which the Node is located at the time. If the requested locale is not found then the default SHALL be used in its place.

5.7.4.9. Enhanced Setup Flow File Example (Terms and Conditions)

```
{
  "schemaVersion": 1,
  "esfRevision": 1,
  "defaultCountry": "US",
  "countryEntries": {
    "US": {
      "defaultLanguage": "en",
      "languageEntries": {
        "en": [
          {
```

```

        "ordinal": 0,
        "required": true,
        "title": "Terms and Conditions",
        "text": "<p><b>Feature 1 Text</b><br><br>Please accept
these.</p>"
    },
    {
        "ordinal": 1,
        "required": false,
        "title": "Privacy Policy",
        "text": "<p>Feature 2 Text</p>"
    }
],
"es": [
    {
        "ordinal": 0,
        "required": true,
        "title": "Términos y condiciones",
        "text": "<p><b>Característica 1 Texto</b><br><br>Por favor
acéptelos.</p>"
    },
    {
        "ordinal": 1,
        "required": false,
        "title": "Política de privacidad",
        "text": "<p>Característica 2 Texto</p>"
    }
]
},
"MX": {
    "defaultLanguage": "es",
    "languageEntries": {
        "es": [
            {
                "ordinal": 0,
                "required": true,
                "title": "Términos y condiciones",
                "text": "<p><b>Característica 1 Texto</b><br><br>Por favor
acéptelos.</p>"
            }
        ]
    }
}

```

```

    },
    "CN": {
      "defaultLanguage": "zh",
      "languageEntries": {
        "zh": [
          {
            "ordinal": 0,
            "required": true,
            "title": "条款和条件",
            "text": "<p><b>产品1文字</b></p>"
          },
          {
            "ordinal": 1,
            "required": false,
            "title": "隐私条款",
            "text": "<p><b>产品2文字</b></p>"
          }
        ]
      }
    },
    "RU": {
      "defaultLanguage": "ru",
      "languageEntries": {
        "ru": [
          {
            "ordinal": 0,
            "required": true,
            "title": "Условия и положения",
            "text": "<p><b>Текст функции 1</b><br><br>Пожалуйста,
примите эти условия пользования.</p>"
          },
          {
            "ordinal": 1,
            "required": false,
            "title": "Положение о конфиденциальности",
            "text": "<p>Текст функции 2</p>"
          }
        ]
      }
    }
  }
}

```

5.7.5. Commissioning Fallback Mechanism

Commissioning using any of the methods described above can potentially fail for a number of reasons. In some cases this will be a transient failure (e.g., due to wireless interference), while in other cases failure will be systematic (e.g., a mismatch in capabilities of the Commissioner and the Commissionee).

Especially for the latter cases, a fallback mechanism is defined which allows the Commissioner to direct the user to a manufacturer-provided mechanism where the user can be guided towards resolution of the issues since the manufacturer typically has more knowledge about the specific device than the commissioner.

Examples of means that could be employed to attempt to resolve the issues:

- The [Section 11.23.6.14, “CommissioningFallbackUrl”](#) can lead to a web page listing suggestions to resolve the failure (e.g., "restart the device to start another commissioning attempt" which helps in case the 15 minute announcement interval since power-on had expired before commissioning was attempted)
- The [Section 11.23.6.14, “CommissioningFallbackUrl”](#) can lead to a manufacturer-provided mechanism to get the device on the IP-bearing network (which can then be followed by "[commissioning while already on IP-network](#)").

5.7.5.1. Commissioning Fallback flow for Device Manufacturers

For cases where commissioning failure can be anticipated (such as devices which rely on commissioning via a technology which is not mandatory for Commissioners e.g. supporting only commissioning via [Wi-Fi Public Action Frame](#)), the [Section 11.23.6.14, “CommissioningFallbackUrl”](#) and [Section 11.23.6.7, “DiscoveryCapabilitiesBitmask”](#) fields in the [DeviceModel](#) schema in the DCL SHALL be populated. For the [Section 5.1.4, “Manual Pairing Code”](#) for such a product, the [Encoding Method with Vendor and Product ID’s included](#) SHALL be used so the Commissioner has access to the VendorID and ProductID for DCL-lookup of the URL.

For other cases, this field MAY be populated to provide a manufacturer-provided mechanism in case commissioning fails for any (other) reason.

5.7.5.2. Commissioning Fallback flow for Commissioners

When a Commissioner encounters a situation where the commissioning fails systematically (e.g., due to a mismatch in capabilities of the Commissioner and the Commissionee), and it has no other means to guide the user to resolution of the issues, it SHALL use the [Section 11.23.6.14, “CommissioningFallbackUrl”](#) (when provided) in the [DeviceModel](#) schema in the DCL to guide the user to the manufacturer-provided means towards resolution of the commissioning issues.

- In cases where using this URL is technically impossible (e.g. a headless commissioner, or the lack of internet access) or when the Commissioner is performing offline mode commissioning, the Commissioner SHOULD inform the user of the assumed reason for commissioning failure, and MAY provide additional advice on how this could potentially be resolved.

When a Commissioner encounters a situation where the commissioning failure is transient, the Commissioner MAY try again; it also MAY use the [Section 11.23.6.14, “CommissioningFallbackUrl”](#)

(when provided) to guide the user to the manufacturer-provided means towards resolution of the commissioning issues.

5.7.5.3. URL used in Commissioning Fallback Mechanism

The [Section 11.23.6.14, “CommissioningFallbackUrl”](#) MAY be constructed in the same way as the [Section 11.23.6.9, “CommissioningCustomFlowUrl”](#) field, including making it specific for a particular VendorID and/or ProductID (as described in [Section 5.7.3, “Custom Commissioning Flow”](#), 5th bullet). Also, the query components as described in [Section 5.7.3.1, “CommissioningCustomFlowUrl format”](#) can be added to the URL to improve the overall flow. For example, usage of **MTop**, **MTcu** and **MTrop** parameters would be useful to smooth the flow in case the manufacturer-provided mechanism uses a manufacturer-defined mechanism to get the device on the IP network, after which the original commissioner can take care of the Matter commissioning (“[device already on IP-network](#)”).

An example of this flow is illustrated below. The “DCL info” concept denotes that the Commissioner SHALL collect the information from the DCL via some mechanism, such as a network resource accessible to the Commissioner containing a replicated set of the DCL’s content.




Figure 52. Commissioning Fallback Mechanism sequence diagram

In the flow above:

- In the final steps, the User might have to perform the trigger to the Commissioner, so that it can start or continue the commissioning process.
- If possible, a Commissioner MAY continue to scan for announcements from the device in the background while any manufacturer-specific app is configuring the device to be available for commissioning. The Commissioner might need a new OnboardingPayload provided to the User by the Manufacturer Website or App.
- In order to simplify the flow, the Commissioner MAY:
 - Include the onboarding payload obtained from the user (see [Section 5.7.3.1, “CommissioningCustomFlowUrl format”](#) for **MTop** key details) within the **CommissioningFallbackUrl**.

- Include a callback URL (see [Section 5.7.3.1, “CommissioningCustomFlowUrl format”](#) for **MTcu** key details) within the **ExpandedFallbackUrl**.
- The Manufacturer Website or App MAY utilize the **CallbackUrl** field, if provided in the query string, in order to simplify the process for signaling the completion of the manufacturer-specific part of the flow back to the Commissioner. When doing so, the Manufacturer Website or App SHOULD put the device into Commissioning mode and SHOULD provide the corresponding onboarding payload to the Commissioner using the **MTrop** key/value pair within the **Expanded-CallbackUrl**.

5.7.6. Onboarding Payload Inclusion (Manual Pairing Code, QR Code, NFC Tag)

The Onboarding Payload formats (Manual Pairing Code, QR code, or NFC Tag) enable secure commissioning and provide a consistent experience that many users are familiar with. However, because they contain a symmetric security code, it is not appropriate in all circumstances to have them be in a readily accessible location on the device, such as printed on the back.

The following are the requirements and recommendations regarding the QR Code, NFC Tag and Manual Pairing code for Standard and User Intent Commissioning Flow Devices, as well as for Custom Commissioning Flow Devices when they satisfy the conditions described in [Section 5.7.3, “Custom Commissioning Flow”](#).

The term 'on-device' allows for a physical label affixed to the device or printed directly on the device, as well as one that can be displayed on demand through some physical interface properties of the device (e.g. visual or audio).

1. Devices SHALL include the Manual Pairing code on-device or in packaging.
2. Devices SHALL NOT have the QR nor the manual pairing code in an unprotected format on the outer packaging.
3. Devices SHOULD include the QR Code, and SHOULD include it alongside the Manual Pairing Code on-device or in packaging.
 - a. Permutations (e.g. Manual Pairing Code on-device and QR Code in packaging) are allowed.
 - b. However, for devices supporting [NFC Transport Layer](#), the QR Code SHALL be included to allow NFC-based commissioning to be initiated.
4. Manual Pairing Code and QR Code on-device MAY be removable or obscured to allow the owner to prevent commissioning without their consent.
5. Devices MAY include the QR Code and Manual Pairing Code in multiple forms (see below).
6. Devices MAY include an NFC Tag.
7. Devices including an NFC Tag SHOULD include it alongside the mandatory Manual Pairing Code on-device or in packaging.
 - a. Permutations (e.g., an NFC Tag on-device and a Manual Pairing Code in packaging, or vice versa) are allowed.
8. An on-device NFC Tag MAY be removable or disabled to allow the owner to prevent commissioning without their consent.

Presentation of the QR Code and Manual Pairing code on-device can occur in many forms to allow for adherence to device security requirements and manufacturing considerations. For example security devices could limit the access to the QR code or Manual Pairing Code to avoid an unauthorized user obtaining the information by simple inspection, or make the QR code and/or Manual Pairing Code removable.

The following is a list of possible ways that are acceptable to satisfy the requirements of inclusion of the QR code, Manual Pairing Code and/or NFC Tag. An entry in the list should not be interpreted as being mutually exclusive with another entry. A device **SHOULD** include as many of these ways as possible.

- QR and Manual Pairing Code shown via an on-device display (when available).
- QR and Manual Pairing Code printed on-device, with removal/obscuring considerations noted above.
- When available, NFC Tag antenna on the device, readable by NFC-enabled Commissioners (e.g., by tapping with a smartphone), with Manual Pairing Code printed on-device, including the removal/obscuring/disabling considerations noted above.
- Manual Pairing Code presented on-device via audio output (when available).
- QR and Manual Pairing Code printed on in-packaging materials.

The following are examples of QR code, Manual Pairing Code and NFC Tag inclusion.

- QR Code and Manual Pairing Code printed on a Matter wireless shade inside the battery compartment cover, and provided in the packaging.
- QR Code and Manual Pairing code on a Matter Smart Thermostat that can be activated via an on-device User Interface and displayed only on screen.
- QR Code and Manual Pairing code for a security sensor that is provided in the packaging, and on-device hidden behind a tamper-monitored cover.
- NFC Tag and Manual Pairing code for a smoke detector present on the back of the device, facing the mounting plate and thus obscured when mounted. Additionally, the mounting plate is metal which prevents NFC tapping when in packaging or mounted.
- QR code provided on an E12 light bulb, with manual pairing code on a removable label (the area of QR code likely fits better on small form factor bulb than the area for a 13 character string).
- A wearable device with only a Manual Pairing Code printed on the fabric. No QR code is present because of the difficulty in scanning a QR code on an irregular surface.
- A Smart speaker, without printed QR or manual pairing code on the device (but possibly in-packaging), that can be triggered to read out a Manual Pairing Code.

5.7.6.1. Inclusion of Onboarding Payload for Concatenated QR Codes

Concatenated QR codes (or NFC Tags) as defined in [Section 5.1.6, “Concatenation”](#) convey the Onboarding Payload for multiple devices provided together (e.g. multi-packs). Therefore, they **SHALL NOT** be used on-device. Rather, they **SHOULD** be provided in packaging. Other requirements and recommendations from the previous section apply, e.g. they **SHALL NOT** be used on the outer

packaging.

5.8. In-field Upgrade to Matter

This (informative) section discusses the case of a pre-Matter device currently in the user's home which gets software updated to support Matter, and which steps (either Matter-specified or manufacturer specific) would typically be applied to accomplish this goal.

- The initial situation is a device which is connected to the local network, and some manufacturer specific means (e.g. a manufacturer-provided app) is used to provide new firmware (including Matter functionality) to the device, along with the associated [Certification Declaration](#). Also, a unique [Device Attestation Certificate](#) is provided into the device using secure, manufacturer-specific means.
- The device restarts to enable the new firmware, and is now an uncommissioned Matter device.
- The device can be commissioned by any Commissioner; the [Onboarding Payload](#) needs to be provided to that Commissioner (since this information is not provided on or with the device out of the factory).
 - For this, similar mechanisms as discussed as in [Section 5.6.3, “Enhanced Commissioning Method \(ECM\)”](#) can be employed:
 - information equivalent to the parameters of the [OpenCommissioningWindow](#) command is sent to the device using some secure manufacturer-defined means
 - presentation of the Onboarding Payload and other relevant information can be performed using the mechanisms described in [Section 5.6.3.1, “Presentation of Onboarding Payload for ECM”](#).
 - For devices with a means to output the [Onboarding Payload](#) themselves (e.g. device with a display or audio output), alternatively, similar mechanisms as discussed as in [Section 5.6.2, “Basic Commissioning Method \(BCM\)”](#) can be employed:
 - information equivalent to the parameters of the [OpenBasicCommissioningWindow](#) command is sent to the device using some secure manufacturer-defined means
 - the device itself presents [Onboarding Payload](#).

5.9. Network Recovery

The operational network credentials (e.g., network name, password) can be changed after a node has been commissioned (see [Section 5.5, “Commissioning Flows”](#)). If this occurs, a Node might no longer be able to access its operational network. In that situation, the Node and an Administrator can execute a sequence of steps, called Network Recovery, to allow the Node to attempt to re-establish network connectivity.

Support for Network Recovery is optional and is indicated via the Network Recovery feature bit being set in the FeatureMap attribute of the General Commissioning Cluster.

NOTE | Support for Network Recovery flows is provisional.

5.9.1. Network Recovery channels

Conceptually, the Network Recovery process is possible over any commissioning channel. As the Matter specification evolves, the designers SHOULD consider providing a Network Recovery feature for any new commissioning channel. In order to fully specify the Network Recovery procedure, commissioning channel definitions SHALL specify:

- an advertising format that distinguishes a network recovery advertisement from the commissioning advertisement. The advertisement for Network Recovery SHALL contain the value of the [RecoveryIdentifier](#) attribute and SHOULD contain the error codes that led to the need for network recovery.
- provisions for [Extended Announcement](#) of Network Recovery advertisements.
- an ability to establish a [CASE](#) session over the commissioning channel. That CASE session SHALL be used for the recovery of network connectivity and not for steady state operational interaction.

5.9.2. Implementer Considerations

When the network connectivity is lost — for example, resulting from equipment replacement, network reconfiguration, etc. — it is usually desirable to restore connectivity to all affected devices. In some cases, the loss of network connectivity is intentional with the user expecting not to reconnect some of the end devices or some of the controllers to the network. It MAY be impossible for an Administrator to discern the user's intent that led to the loss of connectivity, so Administrators SHALL seek user consent / intent to determine which Nodes would be subject to the Network Recovery process.

5.9.3. Network Recovery Flow

The Network Recovery process consists of the following steps enumerated below. Recovery Node refers specifically to the node that has lost its operational network connectivity. This process assumes that the operational network credentials for one or more Nodes has been changed such that those Nodes are no longer able to communicate with other Nodes, including Administrators. Normative requirements below apply to Nodes that choose to implement the Network Recovery feature.

1. At some time before the operational network credentials change, an Administrator SHOULD retrieve the value of the [RecoveryIdentifier](#) attribute from the GeneralCommissioning cluster on the Recovery Node. This MAY be done during the commissioning process.
2. The Administrator SHALL have the new credentials required to connect to the operational network, and the new operational network SHALL be present and operational.
3. Should operational network connectivity be lost, the Recovery Node SHALL continue to attempt to connect to the operational network for a duration of at least 120 seconds using its existing credentials. Afterwards it MAY choose to enter Network Recovery mode and commence announcement. Recovery Nodes SHALL attempt to reconnect to the existing operational network while in the recovery state. Given that temporary network outages occur with some frequency, waiting before entering Network Recovery mode and commencing announcement limits wireless spectrum pollution/interference and helps to limit attack opportunities.

4. The Recovery Node SHALL perform recovery announcement using any of its available commissioning channels. A Recovery Node utilizing its BLE interface SHALL advertise the Network Recovery payload (see [BLE Service Data payload format](#)). A Recovery Node utilizing its Wi-Fi interface with Wi-Fi Unsynchronized Service Discovery SHALL publish its Recovery ID using Wi-Fi Unsynchronized Service Discovery (see [Table 76, “Publish method arguments values for discovery”](#)). The Administrator uses the Node’s advertised OpCode (see [BLE Device Opcode](#) and [Wi-Fi Public Action Frame Device Opcode](#)) to identify Recovery Nodes in the Network Recovery mode. Recovery Node announcements SHALL follow the announcement duration requirements described in [Announcement Duration](#) and SHALL take advantage of the [Extended Announcement](#) duration.
5. If the [Extended Announcement](#) duration ends and the Recovery Node is still unable to reach its operational network, it SHALL continue to attempt to reconnect to the existing operational network at a manufacturer-defined interval. The Recovery Node SHALL restart the Network Recovery procedure when the steps enumerated in [CommissioningModeInitialStepsHint](#) are performed.
6. The Administrator SHALL obtain user consent before providing new credentials to a Recovery Node. This step need not occur in the order it appears in this list of steps (*i.e.*, it could have been obtained earlier). User consent helps ensure that credentials are not inadvertently provided to imposter devices, erroneously provided during a temporary network outage, etc.
7. If user consent has not been granted, the Administrator terminates its execution of the Network Recovery flow and the following steps SHALL NOT be performed.
8. The Administrator establishes a connection with the Recovery Node over BLE or Wi-Fi as it would to commission an uncommissioned device.
9. Once connected, the Administrator and Recovery Node SHALL establish a secure session using [CASE](#) for purposes of provisioning the Recovery Node with updated network credentials. All messages exchanged over the recovery connection SHALL be encrypted using CASE-derived keys. Network recovery procedure SHALL preempt autonomous network reconnect attempts: once the secure connection is established, any ongoing attempts taken autonomously by the Recovery Node to re-establish connectivity to the previously configured operational network SHALL be suspended and SHALL NOT interfere with the session established in this step. Autonomous attempts to re-establish network connectivity SHALL resume in the event of failure of the network recovery procedure. The remainder of the steps is substantially the same as the network commissioning during a commissioning flow, except that they are transacted over a CASE session instead of a PASE session. In addition to the remaining network recovery steps, Recovery Nodes SHALL allow any and all transactions allowed over a CASE session, such as reading Basic Information cluster or accessing device configuration data and topology, as Administrators might require them to fully restore network connectivity.
10. Upon successful establishment of the [CASE](#) session, the Recovery Node SHALL autonomously arm the [fail-safe timer](#) for a timeout of 60 seconds. This is to guard against the Administrator not proceeding with the rest of the flow in a timely fashion, and is analogous to initial phase step 6 of the standard [commissioning flow](#).
11. Within 60 seconds of establishing the CASE session, the Administrator SHALL [arm the fail-safe timer](#). If the fail-safe could not be successfully armed, the Administrator SHALL terminate the session.
12. The Administrator MAY choose to delete a Recovery Node’s existing network credentials by

sending a [RemoveNetwork](#). The Administrator SHALL provide the new credentials of the operational network to the Recovery Node using the [AddOrUpdateWiFiNetwork](#) or [AddOrUpdateThreadNetwork](#) commands.

13. The Administrator SHALL then instruct the Recovery Node to connect to the operational network using the new credentials by issuing a [ConnectNetwork](#) command.
14. The Recovery Node SHALL attempt to connect to the operational network using the newly-provisioned credentials. Once connected, all messages associated with the remaining Network Recovery steps between the Administrator and Recovery Node SHALL be exchanged over the operational network. These messages SHALL be communicated over a [CASE](#) session.
15. The Administrator and the Recovery Node SHALL discover each other on the operational network using operational discovery (see [Section 4.3.2, “Operational Discovery”](#)).
16. The Administrator and the Recovery Node SHALL indicate successful completion of the Network Recovery flow using the [CommissioningComplete](#) and [CommissioningCompleteResponse](#) commands. These commands SHALL be sent over a CASE session over the operational network. The Recovery Node SHALL reject the [CommissioningComplete](#) that is not received over the operational network.




Figure 53. Network Recovery sequence diagram

Chapter 6. Device Attestation and Operational Credentials

This chapter describes the procedures and cryptographic credentials involved in establishing trust between entities.

The [Device Attestation](#) section provides mechanisms for Commissioners and Administrators to determine whether a Node is a genuine certified product before sharing sensitive information such as keys and other credentials. The Device Attestation feature relies on a [Device Attestation Certificate \(DAC\)](#) chain and on a [Certification Declaration \(CD\)](#).

The [Node Operational Credentials](#) section describes the credentials used by all Nodes to mutually authenticate each other during [Certificate-Authenticated Session Establishment](#), including the [Node Operational Certificate \(NOC\)](#) chain. These credentials form the basis of how Nodes are identified and take part in securing operational [unicast communication](#).

6.1. Certificate Common Conventions

This chapter makes use of digital certificates in several subsections. All certificates within this specification are based on X.509v3-compliant certificates as defined in [RFC 5280](#). The storage format of the certificates depends on application (e.g., DAC or NOC chain), but all certificates are directly compatible with X.509v3 DER representation after suitable loading or decompression.

In order to simplify further exposition, this subsection contains some common normative conventions that SHALL apply to all digital certificates described in this specification.

The following certificate formats are defined within this specification:

- Compressed [Node Operational credentials](#) certificate chain elements in [Matter Operational Certificate Encoding](#) or "Matter Certificate" format:
 - Node Operational Certificate ([NOC](#))
 - Intermediate CA Certificate ([ICAC](#))
 - Root CA Certificate ([RCAC](#))
- Device Attestation certificate chain elements in Standard X.509 DER format:
 - Device Attestation Certificate (DAC): see [Section 6.2.2.3, "Device Attestation Certificate \(DAC\)"](#)
 - Product Attestation Intermediate (PAI): see [Section 6.2.2.4, "Product Attestation Intermediate \(PAI\) Certificate"](#)
 - Product Attestation Authority (PAA): see [Section 6.2.2.5, "Product Attestation Authority \(PAA\) Certificate"](#)

6.1.1. Encoding of Matter-specific RDNs

In addition to the standard DN (Distinguished Names) attribute types that appear in certificate Subject and Issuer fields, there are Matter-specific DN attribute types under the [1.3.6.1.4.1.1.37244](#) pri-

vate arc. These are listed in [Table 85, “Matter-specific DN Object Identifiers”](#). These OID values are assigned by the Connectivity Standards Alliance for use with Matter. All of these Matter-specific RDNs encode values normatively defined as scalars.

When used in Matter Operational Certificate (TLV) format (see [Section 6.5, “Operational Certificate Encoding”](#)), Matter-specific DN attribute types SHALL be encoded in Matter TLV as unsigned integers with the specified length.

When used in X.509 ASN.1 DER format certificate encoding, Matter-specific DN attribute types SHALL have their value encoded as either a [UTF8String](#) or [PrintableString](#) according to the table below. The values SHALL be encoded in network byte order as exactly twice their specified maximum octet length, encoded as uppercase hexadecimal number format without any separators or prefix, and without omitting any leading zeroes.

For example:

- A scalar value [0x0123_4567_89AB_CDEF](#) for [matter-node-id](#):
 - Scalar maximal length: 8 octets (64 bits)
 - Resulting string: "0123456789ABCDEF" (without quotes)
 - Resulting length: 16 characters
- A scalar value [0xAA_33CC](#) for [matter-noc-cat](#):
 - Scalar maximal length: 4 octets (32 bits)
 - Resulting string: "00AA33CC" (without quotes)
 - Resulting length: 8 characters

Table 85. Matter-specific DN Object Identifiers

TLV Tag	Matter name	Length (octets)	String length	ASN.1 OID	Types Allowed in X.509
17	matter-node-id	8	16	1.3.6.1.4.1.3724.4.1.1	UTF8String
18	matter-firmware-signing-id	8	16	1.3.6.1.4.1.3724.4.1.2	UTF8String
19	matter-icac-id	8	16	1.3.6.1.4.1.3724.4.1.3	UTF8String
20	matter-rcac-id	8	16	1.3.6.1.4.1.3724.4.1.4	UTF8String
21	matter-fabric-id	8	16	1.3.6.1.4.1.3724.4.1.5	UTF8String
22	matter-noc-cat	4	8	1.3.6.1.4.1.3724.4.1.6	UTF8String
23	matter-vvs-id	8	16	1.3.6.1.4.1.3724.4.1.7	UTF8String

TLV Tag	Matter name	Length (octets)	String length	ASN.1 OID	Types Allowed in X.509
N/A	matter-oid-vid	2	4	1.3.6.1.4.1.3724.4.2.1	UTF8String, PrintableString
N/A	matter-oid-pid	2	4	1.3.6.1.4.1.3724.4.2.2	UTF8String, PrintableString

6.1.2. Key Identifier Extension Constraints

Whenever an X.509 certificate contains Authority Key Identifier or Subject Key Identifier extensions, the associated Key Identifier SHALL be of a length of 20 octets, consistent with the length of derivation method (1) described in section 4.2.1.2 of [\[RFC 5280\]](#).

Further constraints related to the exact derivation appear in the following subsections:

- Matter Certificates (NOC, ICAC, RCAC) Subject Key Identifier extension: see [Section 6.5.11.4, “Subject Key Identifier Extension”](#)
- Matter Certificates Authority Key Identifier extension: see [Section 6.5.11.5, “Authority Key Identifier Extension”](#)
- Device Attestation Certificate (DAC) extensions: see [Section 6.2.2.3, “Device Attestation Certificate \(DAC\)”](#)
- Product Attestation Intermediate (PAI) Certificate extensions: see [Section 6.2.2.4, “Product Attestation Intermediate \(PAI\) Certificate”](#)
- Product Attestation Authority (PAA) Certificate extensions: see [Section 6.2.2.5, “Product Attestation Authority \(PAA\) Certificate”](#)

6.1.3. Certificate Sizes

All certificates SHALL NOT be longer than 600 bytes in their uncompressed DER format. This constraint SHALL apply to the entire DAC chain (DAC, PAI, PAA) and NOC chain (NOC, ICAC, RCAC).

Wherever [Matter Operational Certificate Encoding](#) representation is used, all certificates SHALL NOT be longer than 400 bytes in their TLV form. This constraint only applies to the NOC chain (NOC, ICAC, RCAC) since the DAC chain (DAC, PAI, PAA) only appears in DER format.

All certificates used within Matter SHOULD be as short as possible.

6.1.4. Presentation of example certificates

Certificate bodies are presented for exemplary purposes in multiple formats within this chapter. Since the translation of an X.509 certificate from ASN.1 DER format to human-readable text format may lose fidelity, especially with regards to equivalent types (e.g., PrintableString versus IA5String versus UTF8String) or serialization when non-standard OIDs are seen, textual examples SHALL NOT be considered to be normative. Only direct encoding of DER encoding, such as PEM blocks, should be used to further study the examples. In case of unforeseen divergence between an example certificate illustration and the normative rules expressed in prose, the normative prose SHALL

take precedence over an ambiguous interpretation of an example.

6.2. Device Attestation

6.2.1. Introduction

Certification of a Device includes configuring the Device with immutable credentials that can be cryptographically verified. Device Attestation is the step of the [Commissioning](#) process whereby a Commissioner cryptographically verifies a Commissionee is in fact a certified Device. This chapter describes the Device Attestation Certificate (DAC) and the systems involved in the verification of a DAC.

The processes used to convey the DAC from a Commissionee to a Commissioner, how to verify that a Commissionee holds the private key corresponding to its DAC, and specifically how the DAC is verified are described in [Section 6.2.3, “Device Attestation Procedure”](#).

This chapter refers to the signature algorithm [ECDSA](#) with [SHA256](#) and to the elliptic curve [secp256r1](#) (a.k.a. [prime256v1](#) and [NIST P-256](#)) in compliance with the mapping for [version 1.0](#) of the [Matter Message Format](#) of the cryptographic primitives as specified in [Chapter 3, Cryptographic Primitives](#). Future versions of this specification might adapt these references accordingly.

6.2.2. Device Attestation Certificate (DAC)

All commissionable Matter Nodes SHALL include a Device Attestation Certificate (DAC) and corresponding private key, unique to that Device. The DAC is used in the Device Attestation process, as part of Commissioning a Commissionee into a Fabric. The DAC SHALL be a DER-encoded X.509v3-compliant certificate as defined in [RFC 5280](#) and SHALL be issued by a Product Attestation Intermediate (PAI) that chains directly to an approved Product Attestation Authority (PAA), and therefore SHALL have a certification path length of 2.

The DAC also SHALL contain specific values of Vendor ID and Product ID (see [Section 6.2.2.2, “Encoding of Vendor ID and Product ID in subject and issuer fields”](#)) in its [subject](#) field to indicate the [Vendor ID](#) and [Product ID](#) provenance of the attestation certificate. See [Section 6.2.3.1, “Attestation Information Validation”](#) for how these are used.

The validity period of a DAC is determined by the vendor and MAY be set to the maximum allowed value of 99991231235959Z [GeneralizedTime](#) to indicate that the DAC has no well-defined expiration date.

The notation used in this section to describe the specifics of the DAC uses the ASN.1 basic notation as defined in [X.680](#). The notation below also leverages types defined in [RFC 5280](#) such as [AlgorithmIdentifier](#), [RelativeDistinguishedName](#), [Validity](#), [Time](#), [UTCTime](#), [GeneralizedTime](#), or permitted extension types. Additionally, the notation below uses the ASN.1 definitions captured in the figure below:

```
-- Matter signatures are ECDSA with SHA256
```

```
MatterSignatureIdentifier ::= SEQUENCE {
    algorithm OBJECT IDENTIFIER(id-x962-ecdsa-with-sha256) }
```

-- Matter Names

-- The second to last RelativeDistinguishedName object in MatterDACName SEQUENCE SHALL contain an

-- attribute with type equal to matter-oid-vid and the last RelativeDistinguishedName object in the

-- SEQUENCE SHALL contain an attribute with type field set to matter-oid-pid

MatterDACName ::= SEQUENCE OF RelativeDistinguishedName

-- There are two acceptable formats for MatterPA name. The first is identical to MatterDACName,

-- i.e. the second to last RelativeDistinguishedName object in MatterPAName SEQUENCE SHALL contain an

-- attribute with type equal to matter-oid-vid and the last RelativeDistinguishedName object in the

-- SEQUENCE SHALL contain an attribute with type field set to matter-oid-pid. In the second acceptable

-- format, the last element of the MatterPAName SEQUENCE SHALL be an RelativeDistinguishedName with

-- an attribute with type field set to matter-oid-vid

MatterPAName ::= SEQUENCE OF RelativeDistinguishedName

-- Object definitions and references

-- X962 OIDs

id-x962 OBJECT IDENTIFIER ::= { iso(1) member-body(2) us(840) ansi-x962(10045) }

id-x962-ecdsa-with-sha256 OBJECT IDENTIFIER ::= { id-x962 signatures(4) ecdsa-with-SHA2(3) ecdsa-with-SHA256(2) }

id-x962-prime256v1 OBJECT IDENTIFIER ::= { id-x962 curves(3) prime(1) prime256v1(7) }

-- Connectivity Standards Alliance and Matter specific OIDs

csa-root OBJECT IDENTIFIER ::= { iso(1) identified-organization(3) dod(6) internet(1) private(4) enterprise(1) zigbee(37244) }

-- root arc for attestation certificates

matter-att-root OBJECT IDENTIFIER ::= { csa-root 2 }

```
-- Matter Device Attestation Certificate DN attribute for the Vendor ID (VID)
matter-oid-vid OBJECT IDENTIFIER ::= { matter-att-root 1 }

-- Matter Device Attestation Certificate DN attribute for the Product ID (PID)
matter-oid-pid OBJECT IDENTIFIER ::= { matter-att-root 2 }
```

6.2.2.1. Device Attestation Public Key Infrastructure (PKI)

The Device Attestation PKI hierarchy consists of the PAA, PAI and individual DAC. The public key from the associated [PAI certificate](#) is used to cryptographically verify the DAC signature. The PAI certificate in turn is signed and attested to by the Product Attestation Authority (PAA) CA. The public key from the associated PAA certificate is used to cryptographically verify the PAI certificate signature. The PAA certificate is an implicitly trusted self-signed root certificate. In this way, the DAC chains up to the PAI certificate, which in turn chains up to the PAA root certificate. A PAI SHALL be assigned to a [Vendor ID](#) value. A PAI MAY further be scoped to a single [ProductID](#) value. If a PAI is used for multiple products, then it cannot be scoped to a ProductID value, otherwise the [Device Attestation Procedure](#) will fail policy validations.

Commissioners SHALL use PAA and PAI certificates to verify the authenticity of a Commissionee before proceeding with the rest of the Commissioning flow.

The [subject](#) of all DAC and PAI certificates SHALL be unique among all those issued by their issuer, as intended by [RFC 5280](#) section 4.1.2.6, through the use of [RelativeDistinguishedName](#)s that ensure the uniqueness, such as for example a unique combination of [commonName](#) (OID 2.5.4.3), [serialNumber](#) (OID 2.5.4.5), [organizationalUnitName](#) (OID 2.5.4.11), etc. The exact additional constraints, including for the subject field, for PAA, PAI and DAC certificates, are presented in the following subsections.

The following figure shows the Device Attestation PKI hierarchy.




Figure 54. Device Attestation PKI hierarchy

6.2.2.2. Encoding of Vendor ID and Product ID in subject and issuer fields

The following subsections contain references to [VendorID](#) and [ProductID](#):

- [Section 6.2.2.3, “Device Attestation Certificate \(DAC\)”](#)
- [Section 6.2.2.4, “Product Attestation Intermediate \(PAI\) Certificate”](#)
- [Section 6.2.2.5, “Product Attestation Authority \(PAA\) Certificate”](#)

The values for VendorID and ProductID, where possible or required in [issuer](#) or [subject](#) fields SHALL be encoded by only one of two methods, without mixing the methods within a given field:

1. The "preferred method", using Matter-specific [RelativeDistinguishedName](#) attributes:
 - a. VendorID encoded as [AttributeTypeAndValue](#) entry with [type](#) equal to [1.3.6.1.4.1.37244.2.1](#), and value respecting encoding specified in [Section 6.1.1, “Encoding of Matter-specific RDNs”](#).
 - b. ProductID encoded as [AttributeTypeAndValue](#) entry with [type](#) equal to [1.3.6.1.4.1.37244.2.2](#), and value respecting encoding specified in [Section 6.1.1, “Encoding of Matter-specific RDNs”](#).
2. A "fallback method" to support certificate authorities that only allow customary [RFC 5280](#) OIDs in the arc [{joint-iso-itu-t\(2\) ds\(5\) attributeType\(4\)}](#) for [type](#) values in [AttributeTypeAndValue](#) entries of [RelativeDistinguishedName](#) elements is to encode them as substrings within the [commonName](#) attribute type ([{joint-iso-itu-t\(2\) ds\(5\) attributeType\(4\) commonName\(3\)}](#)):
 - a. VendorID value encoded with substring [Mvid](#): followed by exactly 4 uppercase hexadecimal characters without elision of leading zeroes, anywhere within the [commonName](#), such as for example:
 - i. VendorID 0xFFFF1 (65521 decimal): [Mvid:FFF1](#)
 - ii. VendorID 0x2A (42 decimal): [Mvid:002A](#)
 - b. ProductID value encoded with substring [Mpid](#): followed by exactly 4 uppercase hexadecimal characters without elision of leading zeroes, anywhere within the [commonName](#), such as for example:
 - i. ProductID 0xC20A (49674 decimal): [Mpid:C20A](#)
 - ii. ProductID 0x3A5 (933 decimal): [Mpid:03A5](#)
 - c. For both VendorID and ProductID, the following additional rules apply to the "fallback method":
 - i. The leftmost match having a correct encoding SHALL be used, with other correct matches discarded.
 - ii. For either of VendorID and ProductID, if the prefix string is found on its own anywhere within the [commonName](#), but there is no fully correct match anywhere in the [commonName](#), the field SHALL be considered incorrectly formatted.

The "preferred method" leaves more space for content in the [commonName](#) attribute type if present. It is also less ambiguous, which might allow simpler processing of certificate issuance policy validations in CAs that support the Matter-specific [RelativeDistinguishedName](#) attributes, and simplify the audit of certificates where Vendor ID and Product ID appear.

The "fallback method" is present to support less flexible CA infrastructure.

Fallback method to encode VendorID and ProductID

The "fallback method" requires exactly 9 characters that are safe to use in both `PrintableString` and `UTF8String` for either VendorID or ProductID encoding. Since these VendorID and ProductID substrings have unambiguous format, they MAY be provided anywhere within a `commonName` value, and therefore separator selection does not need to be considered. Note that the standard [RFC 5280](#) length limitation for `commonName` attribute value is 64 characters in total (see `ub-common-name` in [RFC 5280](#)).

Using the "fallback method" for embedding of VendorID and ProductID in `commonName` in the `subject` field of a [Device Attestation Certificate](#) claiming VendorID 0xFFFF1 and ProductID 0x00B1 can be illustrated with the following valid and invalid examples (without the double quotes):

- "ACME Matter Devel DAC 5CDA9899 Mvid:FFF1 Mpid:00B1": valid and recommended since easily human-readable
- "ACME Matter Devel DAC 5CDA9899 Mpid:00B1 Mvid:FFF1": valid and recommended since easily human-readable
- "Mpid:00B1,ACME Matter Devel DAC 5CDA9899,Mvid:FFF1": valid example showing that order or separators are not considered at all for the overall validity of the embedded fields
- "ACME Matter Devel DAC 5CDA9899 Mvid:FFF1Mpid:00B1": valid, but less readable
- "Mvid:FFF1ACME Matter Devel DAC 5CDAMpid:00B19899": valid, but highly discouraged, since embedding of substrings within other substrings might be confusing to human readers.
- "ACME Matter Devel DAC 5CDA9899 Mvid:FF1 Mpid:00B1": invalid, since substring following `Mvid:` is not exactly 4 uppercase hexadecimal digits
- "ACME Matter Devel DAC 5CDA9899 Mvid:fff1 Mpid:00B1": invalid, since substring following `Mvid:` is not exactly 4 uppercase hexadecimal digits
- "ACME Matter Devel DAC 5CDA9899 Mvid:FFF1 Mpid:B1": invalid, since substring following `Mpid:` is not exactly 4 uppercase hexadecimal digits
- "ACME Matter Devel DAC 5CDA9899 Mpid: Mvid:FFF1": invalid, since the prefix `Mpid:` was found but there is no occurrence of `Mpid:` followed by exactly 4 uppercase hexadecimal digits.

In addition, the following example shows a highly discouraged, though technically valid encoding:

- "Mpid:Mvid:FFF1 Mpid:12cd Matter Test Mpid:FE67": VendorID 0xFFFF1 and ProductID 0xFE67 (leftmost match is the correct match for `Mpid:` prefix).
 - Instances such as this example SHOULD be avoided since they can cause confusion in different implementations, even if the rules are obeyed.

If either the Vendor ID ([1.3.6.1.4.1.37244.2.1](#)) or Product ID ([1.3.6.1.4.1.37244.2.2](#)) Matter-specific OIDs appear in any `RelativeDistinguishedName` in the `subject` or `issuer` fields of a certificate which is part of the Device Attestation Certificate chain path, then that certificate within the chain SHALL NOT have its `commonName`, if present, parsed for the "fallback method", in the rest of the `issuer` or `subject` field where the Matter-specific OIDs appear. In other words, considering a field such as `subject` or `issuer`, the presence of either of these OIDs as the `type` for any `AttributeTypeAndValue` within any `RelativeDistinguishedName` of that field SHALL cause the "fallback method" to be skipped altogether for that field. Otherwise, when the "fallback method" can legally be used, it SHALL only be used

against **AttributeTypeAndValue** sequences where the **type** field is **commonName** (**{joint-iso-itu-t(2) ds(5) attributeType(4) commonName(3)}**) in the **issuer** and **subject** fields, and any mention thereafter of using or matching a "Vendor ID" or "Product ID" with regards to a **Device Attestation Procedure** step SHALL rely on values obtained with that method.

For example, if a given **Product Attestation Intermediate certificate** has a **subject** field employing a particular method of encoding the VendorID and ProductID, either using only Matter-specific OIDs or only the fallback method, then it follows that a **Device Attestation Certificate** issued by the certificate authority of that Product Attestation Intermediate SHALL have the same Distinguished-Name content in its **issuer** field, so that the basic path validation algorithm works. That Device Attestation Certificate MAY however have the "fallback method" used within its **subject** field, if the Product Attestation Intermediate certificate authority is unable to encode/reflect the Matter-specific OIDs in RelativeDistinguishedName attributes within the **subject** field. The rules for whether to consider the canonical or "fallback method" for VendorID and ProductID encoding applies field by field independently for each instance of **subject** or **issuer** field found in certificates within the DAC chain.

6.2.2.3. Device Attestation Certificate (DAC)

The attributes in a DAC include:

```
Certificate ::= SEQUENCE {
    tbsCertificate      DACTBSCertificate,
    signatureAlgorithm  AlgorithmIdentifier,
    signatureValue      BIT STRING }
```

```
DACTBSCertificate ::= SEQUENCE {
    version             INTEGER ( v3(2) ),
    serialNumber        INTEGER,
    signature           MatterSignatureIdentifier,
    issuer              MatterPAName,
    validity            Validity,
    subject             MatterDACName,
    subjectPublicKeyInfo SEQUENCE {
        algorithm        OBJECT IDENTIFIER(id-x962-prime256v1),
        subjectPublicKey  BIT STRING },
    extensions          DACExtensions }
```

```
DACExtensions ::= SEQUENCE {
    basicConstraint Extension({extnID id-ce-basicConstraints, critical TRUE, extnValue
BasicConstraints {cA FALSE} })),
    keyUsage Extension({extnID id-ce-keyUsage, critical TRUE, extnValue
KeyUsage({digitalSignature}})),
    authorityKeyIdentifier Extension({extnID id-ce-authorityKeyIdentifier}),
    subjectKeyIdentifier Extension({extnID id-ce-subjectKeyIdentifier}),
```



```

    extendedKeyUsage Extension({extnID id-ce-extKeyUsage}) OPTIONAL,
    authorityInformationAccess Extension({extnID id-pe-authorityInfoAccess}) OPTIONAL,
    subjectAlternateName Extension({extnID id-ce-subjectAltName}) OPTIONAL
}

```

The DAC certificate SHALL follow the following constraints layered on top of the encoding specified by [RFC 5280](#) within the [TBSCertificate](#) structure:

1. The [version](#) field SHALL be set to [2](#) to indicate v3 certificate.
2. The [signature](#) field SHALL contain the identifier for signatureAlgorithm [ecdsa-with-SHA256](#).
3. The [issuer](#) field SHALL be a sequence of [RelativeDistinguishedName](#) s.
4. The [issuer](#) field SHALL have exactly one [VendorID](#) value present.
5. The [issuer](#) field SHALL have exactly zero or one [ProductID](#) value present.
6. The [issuer](#) field SHALL match, byte-for-byte, the [subject](#) field of the PAI certificate for the PAI that issued this DAC.
7. The [subject](#) field SHALL be a sequence of [RelativeDistinguishedName](#) s.
8. The [subject](#) field SHALL have exactly one [VendorID](#) value present.
 - a. The VendorID value present in the [issuer](#) field SHALL match the VendorID value found in [subject](#) field.
9. The [subject](#) field SHALL have exactly one [ProductID](#) value present.
 - a. If a ProductID value was present in the [issuer](#) field, the ProductID value found in [subject](#) field SHALL match the value found in the [issuer](#) field.
10. The algorithm field in [subjectPublicKeyInfo](#) field SHALL be the object identifier for [prime256v1](#).
11. The certificate SHALL carry the following Extensions:
 - a. Basic Constraint extension SHALL be marked [critical](#) and have the [cA](#) field set to FALSE.
 - b. Key Usage extension SHALL be marked [critical](#)
 - i. The [KeyUsage](#) bitstring SHALL only have the [digitalSignature](#) bit set.
 - ii. Other bits SHALL NOT be set
 - c. Authority Key Identifier
 - d. Subject Key Identifier
12. The certificate MAY also carry the following additional Extensions:
 - a. Extended Key Usage
 - b. Authority Information Access
 - c. Subject Alternate Name
 - d. CRLDistributionPoints
 - e. Any other extension allowed in [RFC 5280](#) where inclusion does not violate size limitations. These extensions insofar not defined in this specification SHALL be ignored by commissioners conforming to this specification, but MAY be present to allow flexibility in CA operation.

Valid example DAC with associated private key, in [X.509 PEM](#) format

```
-----BEGIN CERTIFICATE-----
MIIB6TCCAY+gAwIBAgIIDgY7dCvPvL0wCgYIKoZIzj0EAwIwRjEYMBYGA1UEAwP
TWF0dGVyIFRlc3QgUEFJMRQwEgYKKwYBBAGConwCAQwERkZGMTEUMBIGCisGAQQB
gqJ8AgIMBDgwMDAwIBcNMjEwNjI4MTQyMzQzWhgPOTk50TEyMzEyMzU5NTlaMEsx
HTABBgNVBAMFE1hdHRlcjBUZXN0IERBQyAwMDAxMRQwEgYKKwYBBAGConwCAQwE
RkZGMTEUMBIGCisGAQQBgqJ8AgIMBDgwMDAwTATBgqhkJOPQIBBgqhkJOPQMB
BwNCAATCJYMix9xyc3wzvU1wczeqJIW8Rnk+TVrJp1rXQ1JmyQoCjuyvJLD+cAnv
/K7L6tHyw9EkNd7C6tPZkpW/ztbDo2AwXjAMBGNVHRMBAf8EAjAAMA4GA1UdDwEB
/wQEAWIHgDAdBgNVHQ4EFgQUlsLZJJTqL4XA0WcI44jxwJHqD9UwHwYDVROjBBgw
FoAUR0K3CU3r1RXsbs8zuBEVIl8yUogwCgYIKoZIzj0EAwIDSAAwRQIgX8sppA08
NabozmBlxtCdphc9xbJF7DIEkePTSTK3PhcCIQC0VpkPUgUQBfo4j3V0dxVAoESX
kjGWRV5EDWgl2WEDZA==
-----END CERTIFICATE-----

-----BEGIN EC PRIVATE KEY-----
MHcCAQEEIHtcWp+0aVVH+DAQ38iXpphqmT7LfMnMD4V/kIqszwfu0AoGCCqGSM49
AwEHoUQDQgAEwiWDIsfccnN8M77tcHM3qiSFvEZ5Pk1ayada10NSZskKA07sryZQ
/nAJ7/yuy+rR8sPRJDxewurT2ZKVv87Www==
-----END EC PRIVATE KEY-----
```

Human-readable contents of example DAC X.509 certificate

Certificate:

Data:

Version: 3 (0x2)

Serial Number: 1010560536528535133 (0xe063b742bcfbe5d)

Signature Algorithm: ecdsa-with-SHA256

Issuer: CN = Matter Test PAI, 1.3.6.1.4.1.37244.2.1 = FFF1,
1.3.6.1.4.1.37244.2.2 = 8000

Validity

Not Before: Jun 28 14:23:43 2021 GMT

Not After : Dec 31 23:59:59 9999 GMT

Subject: CN = Matter Test DAC 0001, 1.3.6.1.4.1.37244.2.1 = FFF1,
1.3.6.1.4.1.37244.2.2 = 8000

Subject Public Key Info:

Public Key Algorithm: id-ecPublicKey

Public-Key: (256 bit)

pub:

04:c2:25:83:22:c7:dc:72:73:7c:33:be:ed:70:73:

37:aa:24:85:bc:46:79:3e:4d:5a:c9:a7:5a:d7:43:

52:66:c9:0a:02:8e:ec:af:26:50:fe:70:09:ef:fc:

```

        ae:cb:ea:d1:f2:c3:d1:24:35:de:c2:ea:d3:d9:92:
        95:bf:ce:d6:c3
        ASN1 OID: prime256v1
        NIST CURVE: P-256
    X509v3 extensions:
        X509v3 Basic Constraints: critical
            CA:FALSE
        X509v3 Key Usage: critical
            Digital Signature
        X509v3 Subject Key Identifier:
            96:C2:D9:24:94:EA:97:85:C0:D1:67:08:E3:88:F1:C0:91:EA:0F:D5
        X509v3 Authority Key Identifier:
            keyid:AF:42:B7:09:4D:EB:D5:15:EC:6E:CF:33:B8:11:15:22:5F:32:52:88

    Signature Algorithm: ecdsa-with-SHA256
        30:45:02:20:5f:cb:29:a4:0d:3c:35:a6:e8:ce:60:65:c6:d0:
        9d:a6:17:3d:c5:b2:45:ec:32:04:91:e3:d3:49:32:b7:3e:17:
        02:21:00:b4:56:99:0f:52:05:10:04:5a:38:8f:75:4e:77:15:
        40:a0:44:97:92:31:96:45:5e:44:0d:68:25:d9:61:03:64

```

Valid example DAC with associated private key, in [X.509 PEM](#) format, using "[fallback method](#)" for VendorID and ProductID in Subject

```

-----BEGIN CERTIFICATE-----
MIIBODCCAXegAwIBAgIIbec9lw3wZpAwCgYIKoZIzj0EAwIwRjEYMBYGA1UEAwP
TWF0dGVyIFRlc3QgUEFJMRQwEgYKKwYBBAGConwCAQwERkZGMTEUMBIGCisGAQQB
gqJ8AgIMBDgwMDAwIBcNMjEwNjI4MTQyMzQzWhgPOTk5OTEyMzEyMzU5NTLaMDMx
MTAvBgNVBAMKE1hdHRlcjBUZXN0IERBQyAwMDAxIE12aWQ6RkZGMSBNCGLk0jgw
MDAwWTATBgqhkJOPQIBBgqhkJOPQMBBwNCAATCJYMiX9xyc3wzvu1wczeqJIW8
Rnk+TVrJp1rXQ1JmyQoCjuyvJLD+cAnv/K7L6tHyw9EkNd7C6tPZkpW/ztbDo2Aw
XjAMBGNVHRMBAf8EAjAAMA4GA1UdDwEB/wQEAwIHgDAdBgNVHQ4EFgQUlsLZJJTq
L4XA0WcI44jxwJHqD9UwHwYDVR0jBBgwFoAUR0K3CU3r1RXsbs8zuBEVIL8yUogw
CgYIKoZIzj0EAwIDRwAwRAIgbvYsHaGRTg1JzPTB6TqfVFPABF8LCYKEP1AvV7Ah
yL4CIACKW3A6YixqtqKfkWuvw81mMVymqafU8kx5k1c0zqbe
-----END CERTIFICATE-----

```

Human-readable contents of example DAC X.509 certificate, using "[fallback method](#)" for VendorID and ProductID in Subject

```

Certificate:
    Data:
        Version: 3 (0x2)
        Serial Number: 7919366188737521296 (0x6de73d970df06690)

```

```

Signature Algorithm: ecdsa-with-SHA256
Issuer: CN = Matter Test PAI, 1.3.6.1.4.1.37244.2.1 = FFF1,
1.3.6.1.4.1.37244.2.2 = 8000
Validity
    Not Before: Jun 28 14:23:43 2021 GMT
    Not After : Dec 31 23:59:59 9999 GMT
Subject: CN = Matter Test DAC 0001 Mvid:FFF1 Mpid:8000
Subject Public Key Info:
    Public Key Algorithm: id-ecPublicKey
    Public-Key: (256 bit)
    pub:
        04:c2:25:83:22:c7:dc:72:73:7c:33:be:ed:70:73:
        37:aa:24:85:bc:46:79:3e:4d:5a:c9:a7:5a:d7:43:
        52:66:c9:0a:02:8e:ec:af:26:50:fe:70:09:ef:fc:
        ae:cb:ea:d1:f2:c3:d1:24:35:de:c2:ea:d3:d9:92:
        95:bf:ce:d6:c3
    ASN1 OID: prime256v1
    NIST CURVE: P-256
X509v3 extensions:
    X509v3 Basic Constraints: critical
        CA:FALSE
    X509v3 Key Usage: critical
        Digital Signature
    X509v3 Subject Key Identifier:
        96:C2:D9:24:94:EA:97:85:C0:D1:67:08:E3:88:F1:C0:91:EA:0F:D5
    X509v3 Authority Key Identifier:
        keyid:AF:42:B7:09:4D:EB:D5:15:EC:6E:CF:33:B8:11:15:22:5F:32:52:88

Signature Algorithm: ecdsa-with-SHA256
    30:44:02:20:6e:f6:2c:1d:a1:91:4e:0d:49:cc:f4:c1:e9:3a:
    9f:54:53:c0:04:5f:0b:09:89:04:3f:50:2f:57:b0:21:c8:be:
    02:20:00:8a:5b:70:3a:62:2c:6a:b6:a2:9f:93:0b:af:c3:cd:
    66:31:5c:a6:a9:a7:d4:f2:4c:79:93:57:34:ce:a6:de

```

6.2.2.4. Product Attestation Intermediate (PAI) Certificate

The attributes in a PAI certificate include:

```

Certificate ::= SEQUENCE {
    tbsCertificate      PAITBSCertificate,
    signatureAlgorithm  AlgorithmIdentifier,
    signatureValue      BIT STRING }

```

```

PAITBSCertificate ::= SEQUENCE {
    version                INTEGER ( v3(2) ),
    serialNumber            INTEGER,
    signature               MatterSignatureIdentifier,
    issuer                  Name,
    validity                Validity,
    subject                 MatterPAName,
    subjectPublicKeyInfo    SEQUENCE {
        algorithm            OBJECT IDENTIFIER(id-x962-prime256v1),
        subjectPublicKey     BIT STRING },
    extensions              PAIExtensions }

PAIExtensions ::= SEQUENCE {
    basicConstraint Extension({extnID id-ce-basicConstraints, critical TRUE, extnValue
BasicConstraints {cA TRUE, pathLen 0} })),
    keyUsage Extension({extnID id-ce-keyUsage, critical TRUE, extnValue KeyUsage(<see
text>)})),
    authorityKeyIdentifier Extension({extnID id-ce-authorityKeyIdentifier}),
    subjectKeyIdentifier Extension({extnID id-ce-subjectKeyIdentifier}),
    extendedKeyUsage Extension({extnID id-ce-extKeyUsage}) OPTIONAL
}

```

The PAI certificate SHALL follow the following constraints layered on top of the encoding specified by [RFC 5280](#) within the **TBSCertificate** structure:

1. The **version** field SHALL be set to **2** to indicate v3 certificate.
2. The **signature** field SHALL contain the identifier for signatureAlgorithm **ecdsa-with-SHA256**.
3. The **issuer** field SHALL be a sequence of **RelativeDistinguishedName** s.
4. The **issuer** field SHALL have exactly zero or one **VendorID** value present.
5. The **issuer** field SHALL match, byte-for-byte, the **subject** field of the PAA certificate for the PAA that issued this PAI certificate.
6. The **subject** field SHALL be a sequence of **RelativeDistinguishedName** s.
7. The **subject** field SHALL have exactly one **VendorID** value present.
 - a. If a VendorID value was present in the **issuer** field, the VendorID value found in **subject** field SHALL match the value found in the **issuer** field.
8. The **subject** field SHALL have exactly zero or one **ProductID** value present.
9. The algorithm field in **subjectPublicKeyInfo** field SHALL be the object identifier for **prime256v1**.
10. The certificate SHALL carry the following Extensions:
 - a. Basic Constraint extension SHALL be marked **critical** and have the **ca** field set to TRUE and **pathLen** field set to 0.

- b. Key Usage extension SHALL be marked **critical**
 - i. Both the **keyCertSign** and **cRLSign** bits SHALL be set in the **KeyUsage** bitstring
 - ii. The **digitalSignature** bit MAY be set in the **KeyUsage** bitstring
 - iii. Other bits SHALL NOT be set
 - c. Authority Key Identifier
 - d. Subject Key Identifier
11. The certificate MAY also carry the following additional Extensions:
- a. Extended Key Usage
 - b. CRLDistributionPoints
 - c. Any other extension allowed in [RFC 5280](#) where inclusion does not violate size limitations. These extensions insofar not defined in this specification SHALL be ignored by commissioners conforming to this specification, but MAY be present to allow flexibility in CA operation.

The PAI certificate presented in the following example is for the issuer of the example DAC certificate from the previous section.

Valid example PAI with associated private key, in [X.509 PEM](#) format

```
-----BEGIN CERTIFICATE-----
MIIB1DCCAXqgAwIBAgIIPmzmUJrYQM0wCgYIKoZIzj0EAwIwMDEYMBYGA1UEAwWP
TWF0dGVyIFRlc3QgUEFBMRQwEgYKKwYBBAGConwCAQwERkZGMTAgFw0yMTA2Mjgx
NDIzNDNaGA850Tk5MTIzMTIzNTk1OVowRjEYMBYGA1UEAwWPTWF0dGVyIFRlc3Qg
UEFJMRQwEgYKKwYBBAGConwCAQwERkZGMTAUMBIGIisGAQQBgqJ8AgIMBDgwMDAw
WTATBgqhkJOPQIBBgqhkJOPQMBBwNCAASA3fEbIo8+MfY7z1eY2hRiOuu96C7z
e06tv7GP4av0MdC01LIGBLbMxtm1+rZ0feEMt0vgF8nsFRYFbXDyzQsio2YwZDAS
BgNVHRMBAf8ECDAGAQH/AgEAMA4GA1UdDwEB/wQEAwIBBjAdBgNVHQ4EFgQUr0K3
CU3r1RXsbs8zuBEVIl8yUogwHwYDVROjBBgwFoAUav0idx9RH+y/FkGXZxDc3DGh
cX4wCgYIKoZIzj0EAwIDSAAwRQIhAJbJyM8uAYhgBdj1vHLAe3X9mlpWsSRETET
i+oDPOUDAiAlVJQ75X1T1sR199I+v8/CA2zSm6Y5PsfvrYcUq3GCGQ==
-----END CERTIFICATE-----

-----BEGIN EC PRIVATE KEY-----
MHcCAQEEIEZ7LYpps1z+a9sPw2qBp9jj5F0GLffNuCJY88hAHcMYoAoGCCqGSM49
AwEHoUQDQgAEgN3xGyKPPjH2089XmNoUYjrrvegu83jurb+xj+GrzjHQjtSyBgS2
zMbZtfq2Tn3hDLdL4BfJ7BUWBW1w8s0LIg==
-----END EC PRIVATE KEY-----
```

Human-readable contents of example PAI X.509 certificate

Certificate:

Data:

Version: 3 (0x2)

Serial Number: 4498223361705918669 (0x3e6ce6509ad840cd)

Signature Algorithm: ecdsa-with-SHA256

Issuer: CN = Matter Test PAA, 1.3.6.1.4.1.37244.2.1 = FFF1

Validity

Not Before: Jun 28 14:23:43 2021 GMT

Not After : Dec 31 23:59:59 9999 GMT

Subject: CN = Matter Test PAI, 1.3.6.1.4.1.37244.2.1 = FFF1,

1.3.6.1.4.1.37244.2.2 = 8000

Subject Public Key Info:

Public Key Algorithm: id-ecPublicKey

Public-Key: (256 bit)

pub:

04:80:dd:f1:1b:22:8f:3e:31:f6:3b:cf:57:98:da:

14:62:3a:eb:bd:e8:2e:f3:78:ee:ad:bf:b1:8f:e1:

ab:ce:31:d0:8e:d4:b2:06:04:b6:cc:c6:d9:b5:fa:

b6:4e:7d:e1:0c:b7:4b:e0:17:c9:ec:15:16:05:6d:

70:f2:cd:0b:22

ASN1 OID: prime256v1

NIST CURVE: P-256

X509v3 extensions:

X509v3 Basic Constraints: critical

CA:TRUE, pathlen:0

X509v3 Key Usage: critical

Certificate Sign, CRL Sign

X509v3 Subject Key Identifier:

AF:42:B7:09:4D:EB:D5:15:EC:6E:CF:33:B8:11:15:22:5F:32:52:88

X509v3 Authority Key Identifier:

keyid:6A:FD:22:77:1F:51:1F:EC:BF:16:41:97:67:10:DC:DC:31:A1:71:7E

Signature Algorithm: ecdsa-with-SHA256

30:45:02:21:00:96:c9:c8:cf:2e:01:88:60:05:d8:f5:bc:72:

c0:7b:75:fd:9a:57:69:5a:c4:91:11:31:13:8b:ea:03:3c:e5:

03:02:20:25:54:94:3b:e5:7d:53:d6:c4:75:f7:d2:3e:bf:cf:

c2:03:6c:d2:9b:a6:39:3e:c7:ef:ad:87:14:ab:71:82:19

6.2.2.5. Product Attestation Authority (PAA) Certificate

The attributes in a PAA certificate include:

```
Certificate ::= SEQUENCE {
    tbsCertificate      PAATBSCertificate,
    signatureAlgorithm  AlgorithmIdentifier,
```


signatureValue	BIT STRING }
----------------	--------------

```

PAATBSCertificate ::= SEQUENCE {
    version                INTEGER ( v3(2) ),
    serialNumber            INTEGER,
    signature               MatterSignatureIdentifier,
    issuer                  Name,
    validity                Validity,
    subject                 Name,
    subjectPublicKeyInfo    SEQUENCE {
        algorithm            OBJECT IDENTIFIER(id-x962-prime256v1),
        subjectPublicKey     BIT STRING },
    extensions              PAAExtensions }

PAAExtensions ::= SEQUENCE {
    basicConstraint Extension({extnID id-ce-basicConstraints, critical TRUE, extnValue
BasicConstraints {cA TRUE, pathLen 1} }),
    keyUsage Extension({extnID id-ce-keyUsage, critical TRUE, extnValue KeyUsage(<see
text>)}),
    authorityKeyIdentifier Extension({extnID id-ce-authorityKeyIdentifier}) OPTIONAL,
    subjectKeyIdentifier Extension({extnID id-ce-subjectKeyIdentifier}),
    extendedKeyUsage Extension({extnID id-ce-extKeyUsage}) OPTIONAL
}

```

The PAA certificate SHALL follow the following constraints layered on top of the encoding specified by [RFC 5280](#) within the **TBSCertificate** structure:

1. The **version** field SHALL be set to **2** to indicate v3 certificate.
2. The **signature** field SHALL contain the identifier for signatureAlgorithm **ecdsa-with-SHA256**.
3. The **issuer** field SHALL be a sequence of **RelativeDistinguishedName** s.
4. The **issuer** field SHALL have exactly zero or one **VendorID** value present.
5. The **subject** field SHALL be a sequence of **RelativeDistinguishedName** s.
6. The **subject** field SHALL have exactly zero or one **VendorID** value present.
7. The **issuer** and **subject** fields SHALL match exactly.
8. A **ProductID** value SHALL NOT be present in either the **subject** or **issuer** fields.
9. The algorithm field in **subjectPublicKeyInfo** field SHALL be the object identifier for **prime256v1**.
10. The certificate SHALL carry the following Extensions:
 - a. Basic Constraint extension SHALL be marked **critical** and have the **cA** field set to TRUE. The 'pathLen' field MAY be set and if the 'pathLen' is field is present it SHALL be set to 1.
 - b. Key Usage extension SHALL be marked **critical**.

- i. Both the **keyCertSign** and **cRLSign** bits SHALL be set in the **KeyUsage** bitstring
 - ii. The **digitalSignature** bit MAY be set in the **KeyUsage** bitstring
 - iii. Other bits SHALL NOT be set
 - c. Subject Key Identifier
11. The certificate MAY also carry the following additional Extensions:
- a. Extended Key Usage
 - b. Authority Key Identifier
 - c. Any other extension allowed in [RFC 5280](#) where inclusion does not violate size limitations. These extensions insofar not defined in this specification SHALL be ignored by commissioners conforming to this specification, but MAY be present to allow flexibility in CA operation.

The PAA certificate presented in the following example is for the issuer of the example PAI certificate from the previous section.

Valid example PAA with associated private key, in [X.509 PEM](#) format

```
-----BEGIN CERTIFICATE-----
MIIBvTCCAWsGAwIBAgIITqjoMYLUHBwwCgYIKoZIzj0EAwIwMDEYMBYGA1UEAwP
TWF0dGVyIFRlc3QgUEFBRQwEgYKKwYBBAGConwCAQwERkZGNTAgFw0yMTA2Mjgx
NDIzNDNaGA850Tk5MTIzMTIzNTk1OVowMDEYMBYGA1UEAwPTWF0dGVyIFRlc3Qg
UEFBRQwEgYKKwYBBAGConwCAQwERkZGNTBZMBMGByqGSM49AgEGCCqGSM49AwEH
A0IABLbLY3KI fyko9brIGqnZ0uJDHK2p154kL2UXfvn02TKijs0Duq9qj8oYShpQ
NUKWDUU/MD8fGUIdR6Pjxqam3WjZjBkMBIGA1UdEwEB/wQIMAYBAf8CAQEwDgYD
VR0PAQH/BAQDAgEGMB0GA1UdDgQWBBRq/SJ3H1Ef7L8WQZdnENzcMaFxfjAfbgNV
HSMEGDAwBRq/SJ3H1Ef7L8WQZdnENzcMaFxfjAKBgqhkhjOPQDAGNHADBEAiBQ
qoAC9NkyqaAFOPZTaKOP/8jvu8m+t9pWmDXPmqdRDgIgI7rI/g8j51RFtLM5CBpH
mUkpxyqvChVI1A0DTVFLJd4=
-----END CERTIFICATE-----

-----BEGIN EC PRIVATE KEY-----
MHcCAQEEIGUSyuyuz8VD1gYjFhWXF i8BRoTFZaEpti/SjCerHMxQoAoGCCqGSM49
AwEHoUQDQgAEtstjcoh/KSj1usgaqdk64kMcranXniQvZrD++c7ZMqK0zQ06r2qP
yhhKGLA1QpYNRT8wPx8ZQh11Ho+PGpqbdQ==
-----END EC PRIVATE KEY-----
```

Human-readable contents of example PAA X.509 certificate

Certificate:

Data:

Version: 3 (0x2)

Serial Number: 5668035430391749660 (0x4ea8e83182d41c1c)

Signature Algorithm: ecdsa-with-SHA256

Issuer: CN = Matter Test PAA, 1.3.6.1.4.1.37244.2.1 = FFF1

Validity

Not Before: Jun 28 14:23:43 2021 GMT

Not After : Dec 31 23:59:59 9999 GMT

Subject: CN = Matter Test PAA, 1.3.6.1.4.1.37244.2.1 = FFF1

Subject Public Key Info:

Public Key Algorithm: id-ecPublicKey

Public-Key: (256 bit)

pub:

04:b6:cb:63:72:88:7f:29:28:f5:ba:c8:1a:a9:d9:

3a:e2:43:1c:ad:a9:d7:9e:24:2f:65:17:7e:f9:ce:

d9:32:a2:8e:cd:03:ba:af:6a:8f:ca:18:4a:1a:50:

35:42:96:0d:45:3f:30:3f:1f:19:42:1d:75:1e:8f:

8f:1a:9a:9b:75

ASN1 OID: prime256v1

NIST CURVE: P-256

X509v3 extensions:

X509v3 Basic Constraints: critical

CA:TRUE, pathlen:1

X509v3 Key Usage: critical

Certificate Sign, CRL Sign

X509v3 Subject Key Identifier:

6A:FD:22:77:1F:51:1F:EC:BF:16:41:97:67:10:DC:DC:31:A1:71:7E

X509v3 Authority Key Identifier:

keyid:6A:FD:22:77:1F:51:1F:EC:BF:16:41:97:67:10:DC:DC:31:A1:71:7E

Signature Algorithm: ecdsa-with-SHA256

30:44:02:20:50:aa:80:02:f4:d9:32:a9:a0:05:38:f6:53:68:

ad:0f:ff:c8:ef:bb:c9:be:b7:da:56:98:35:cf:9a:a7:51:0e:

02:20:23:ba:c8:fe:0f:23:e7:54:45:b6:53:39:08:1a:47:99:

49:29:c7:2a:af:0a:15:48:d4:0d:03:4d:51:4b:25:de

6.2.3. Device Attestation Procedure

The device attestation procedure SHALL be executed by Commissioners when commissioning a device. It serves to validate whether a particular device is certified for Matter compliance and that it was legitimately produced by the certified manufacturer. See [Section 5.5, “Commissioning Flows”](#) for the possible outcomes based on whether the Device Attestation Procedure succeeds or fails to attest the device.

The Device Attestation Certificate Chain MAY be read at any time, either prior to or after receipt of the [AttestationResponse](#). The Commissionee SHALL make the Certificate Chain available whenever requested using [CertificateChainRequest](#). If the Commissioner does not already have this information, to proceed with the validation, it SHALL request the Commissionee’s Device Attestation Certificate Chain using [CertificateChainRequest](#).

The procedure is as follows:

1. The Commissioner SHALL generate a random 32 byte attestation nonce using `Crypto_DRBG()`.
2. The Commissioner SHALL send the `AttestationNonce` to the Commissionee and request Attestation Information using `AttestationRequest`.
3. The Commissionee SHALL return the signed `Attestation Information` to the Commissioner using `AttestationResponse`.

After execution of the procedure, the Attestation Information SHOULD be validated using the checks described in [Section 6.2.3.1, “Attestation Information Validation”](#).

6.2.3.1. Attestation Information Validation

A Commissioner validating the Attestation Information SHOULD record sufficient information to provide detailed results of the validation outcome to users. Therefore, prior to validating Attestation Information, a Commissioner SHOULD have previously obtained the Device Attestation Certificate chain for the Commissionee, so that the DAC and PAI necessary for the procedure are available.

In order to consider a Commissionee successfully attested, a Commissioner SHALL have successfully validated at least the following:

- The PAA SHALL be validated for presence in the Commissioner’s trusted root store, which SHOULD include at least the set of globally trusted PAA certificates present in the [Distributed Compliance Ledger](#) at the issuing timestamp (`notBefore`) of the DAC.
- The DAC certificate chain SHALL be validated using the `Crypto_VerifyChainDER()` function, taking into account the mandatory presence of the PAI and of the PAA. It is especially important to ensure the entire chain has a length of exactly 3 elements (PAA certificate, PAI certificate, Device Attestation Certificate) and that the necessary format policies previously exposed are validated, to avoid unauthorized path chaining (e.g., through multiple PAI certificates).
 - Chain validation SHALL be performed with respect to the `notBefore` timestamp of the DAC to ensure that the DAC was valid when it was issued. This way of validating is abided by the `Crypto_VerifyChainDER()` function.
 - Chain validation SHALL include revocation checks of the DAC and PAI, based on the Commissioner’s best understanding of revoked entities. See [Section 6.2.4, “Device attestation revocation”](#) for an interoperable method of obtaining revocation information in the Device Attestation PKI.
- The `VendorID` value found in the subject DN of the DAC SHALL match the `VendorID` value in the subject DN of the PAI certificate.
- If the PAA certificate contains a `VendorID` value in its subject DN, its value SHALL match the `VendorID` value in the subject DN of the PAI certificate.
- The Device Attestation Signature (`attestation_signature`) field from `Attestation Response` SHALL be validated:

```
Success = Crypto_Verify(  
    publicKey = Public key from DAC,
```

```

message = Attestation Information TBS (attestation\_tbs),
signature = Device Attestation Signature (attestation\_signature)
)

```

where the fields are encoded as described in [Section 11.18.4.7, “Attestation Information”](#).

- The [AttestationChallenge](#) SHALL be obtained from a [CASE session](#), [resumed CASE session](#), or [PASE session](#) depending on the method used to establish the secure session within which device attestation is conducted.
- The [AttestationNonce](#) in [Device Attestation elements](#) SHALL match the Commissioner’s provided [AttestationNonce](#).
- The Certification Declaration signature SHALL be validated using the [Crypto_Verify\(\)](#) function and the public key obtained from the Connectivity Standards Alliance’s Certificate Authority Certificate.
- The [Certification Declaration](#) SHALL be validated:
 - The [vendor_id](#) field in the [Certification Declaration](#) SHALL match the [VendorID attribute](#) found in the Basic Information cluster.
 - The [product_id_array](#) field in the [Certification Declaration](#) SHALL contain the value of the [ProductID attribute](#) found in the Basic Information cluster.
 - The Certification Declaration SHALL be considered valid only if it contains both or neither of the [dac_origin_vendor_id](#) and [dac_origin_product_id](#) fields.
 - If the Certification Declaration has both the [dac_origin_vendor_id](#) and the [dac_origin_product_id](#) fields, the following validation SHALL be done:
 - The [VendorID](#) value from the subject DN in the DAC SHALL match the [dac_origin_vendor_id](#) field in the [Certification Declaration](#).
 - The [VendorID](#) value from the subject DN in the PAI SHALL match the [dac_origin_vendor_id](#) field in the [Certification Declaration](#).
 - The [ProductID](#) value from the subject DN in the DAC SHALL match the [dac_origin_product_id](#) field in the [Certification Declaration](#).
 - The [ProductID](#) value from the subject DN in the PAI, if such a ProductID value appears, SHALL match the [dac_origin_product_id](#) field in the [Certification Declaration](#).
 - If the Certification Declaration has neither the [dac_origin_vendor_id](#) nor the [dac_origin_product_id](#) fields, the following validation SHALL be done:
 - The [VendorID](#) value from the subject DN in the DAC SHALL match the [vendor_id](#) field in the [Certification Declaration](#).
 - The [VendorID](#) value from the subject DN in the PAI SHALL match the [vendor_id](#) field in the [Certification Declaration](#).
 - The [ProductID](#) value from the subject DN in the DAC SHALL be present in the [product_id_array](#) field in the [Certification Declaration](#).
 - The [ProductID](#) value from the subject DN in the PAI, if such a Product ID is present, SHALL match one of the values present in the [product_id_array](#) field in the [Certification Declaration](#).

Declaration.

- If the Certification Declaration contains the **authorized_paa_list** field, the following validation SHALL be done:
 - The Subject Key Identifier (SKI) extension value of the PAA certificate, which is the root of trust of the DAC, SHALL be present as one of the values in the **authorized_paa_list** field.
- The **certificate_id** field SHOULD match the **CDCertificateID** field found in the entry of the **DeviceSoftwareCompliance** schema in the **Distributed Compliance Ledger** where the entry's **VendorID**, **ProductID** and **SoftwareVersion** field match the respective **VendorID**, **ProductID** and **SoftwareVersion** attributes values found in the **Basic Information Cluster**. For further clarity, a scenario where a mismatch is most likely to occur is with devices that have received an updated firmware but not an updated CD.
 - If the **certificate_id** does not match the **CDCertificateID** field due to the above point, a Commissioner SHOULD NOT reject the Commissionee. Instead, the Commissioner SHOULD check that there is an entry with the **VendorID**, **ProductID** and **SoftwareVersion** in the **DeviceSoftwareCompliance** from the **Distributed Compliance Ledger** to validate that the current version is certified. For the case of offline commissioning without access to DCL, the Commissioner SHOULD proceed trusting that a valid CD matching the reported **VendorID** and **ProductID** at least ensures the device has been previously certified and SHOULD check the DCL at a later time to ensure the current **SoftwareVersion** is also certified.
- The **firmware_information** field in the **Attestation Information**, if present, SHALL match the content of an entry in the **Distributed Compliance Ledger** for the specific device as explained in **Section 6.3.2, "Firmware Information"**. If the Commissioner does not support Firmware Information validation, it MAY skip checking this match.

The order of execution of the above validation steps MAY be optimized by Commissioners. For example, if some validation steps are deemed by a Commissioner to make the remainder of the steps unnecessary because they have no chance of succeeding, then the validation steps could be ordered such that superfluous steps or rounds trips are omitted.

6.2.4. Device attestation revocation

The **Distributed Compliance Ledger** contains a trusted list of revocation distribution points for the Device Attestation PKI. The schema is documented in **Section 11.23.9, "Device Attestation PKI Revocation Distribution Points Schema"**.

The format of certificates in the Device Attestation PKI is constrained to maximum sizes and a specific format to reduce the complexity and system requirements for Device and Commissioner implementations. Therefore, there are no requirements for common revocation information, such as **cRLDistributionPoints** (see **RFC 5280** section 4.2.1.13) to be included. Furthermore, OCSP (see **RFC 6960**) is not used in order to allow Commissioner implementations which rely on previously cached revocation information, or which do not have concurrent access both to the public Internet and to the Commissionee when commissioning. The **Device Attestation PKI Revocation Distribution Points Schema** in the Distributed Compliance Ledger is thus the main method of delivering interoperable and universally accessible revocation information for Matter Commissioners and Adminis-

trators.

The revocation information MAY be partitioned over multiple [Device Attestation PKI Revocation Distribution Points Schema](#) entries, as the provider of such information deems fit. In other words, there might be more than one record to consult to determine revocation status of a given certificate.

This might occur when PAAs or PAIs use partitioned CRL, in which case multiple CRL partitions will be created by a PAA or PAI, and multiple CRL URI entries will be added to the DCL for a particular PAA or PAI. Providers need to ensure URIs for all partitions have been uploaded to the DCL for a given PAA or PAI. The algorithm for revocation set construction will check the CRL's Issuing Distribution Point (IDP) extension and ensure it matches the URI stored in the DCL that was used to fetch the CRL.

If a DAC, PAI or PAA certificate contains the [cRLDistributionPoints](#) (see [RFC 5280](#) section 4.2.1.13) extension, it SHALL be ignored, since the [Device Attestation PKI Revocation Distribution Points Schema](#) in the Distributed Compliance Ledger is the only supported method of determining up-to-date revocation status.

CRL delegation (also known as "indirect CRLs") is supported for both PAI and PAA, allowing certificate authorities to delegate CRL signing to other entities. This provides flexibility in managing revocation lists while maintaining security through proper validation chains and reduces operational overhead for certificate authorities.

6.2.4.1. Conceptual algorithm for revocation set construction

To build the list of revoked entities for a given PAI or PAA certificate authority, the table described in the [Device Attestation PKI Revocation Distribution Points Schema](#) can be used with a conceptual algorithm, which can then be used at time of revocation checks.

To build an oracle of revoked entities, the conceptual algorithm is as follows:

- Inputs:
 - List of all entries in the [Device Attestation PKI Revocation Distribution Points](#) table of the Distributed Compliance Ledger.
- Outputs:
 - For each tuple of ([CertificateAuthorityKeyIdentifier](#), [CertificateAuthorityName](#)), a set of [RevokedEntitySerialNumber](#) and an associated tuple of ([CRLSignerCertificate](#), [CRLSignerDelegator](#)) certificates.
 - [CertificateAuthorityKeyIdentifier](#) is the Subject Key Identifier of a PAA or PAI as an octet string (e.g. ASN.1 [OCTET STRING](#), or protobuf [bytes](#)).
 - [CertificateAuthorityName](#) is the X.500 Name associated with the PAA or PAI whose revocation information is being aggregated, as it appears in the [Issuer](#) field of a certificate it issues.
 - [RevokedEntitySerialNumber](#) is a serial number containing a large integer matchable to an ASN.1 [INTEGER](#) as seen in the [CertificateSerialNumber](#) field of an [RFC 5280](#)-compliant certificate.

- **CRLSignerCertificate** is the certificate present in the CRLSignerCertificate field from the Device Attestation PKI Revocation Distribution Points table entry processed.
- **CRLSignerDelegator** is the certificate present in the CRLSignerDelegator field from the Device Attestation PKI Revocation Distribution Points table entry processed, if present.

NOTE

While the algorithm here does not mention the recording of any other metadata per revoked entity to simplify illustration, implementers MAY do so to allow richer diagnostics information to be provided to end-users if a revoked entity is found.

- Algorithm:

- For each **entry** in the **Device Attestation PKI Revocation Distribution Points**
 1. If **entry.RevocationType** is not equal to 1 (not of type **RFC 5280 CRL**), stop further processing of the entry, and continue to next entry, as it is a format not defined in this Specification.
 2. Parse the certificate found **entry.CRLSignerCertificate** from its PEM encoding, keeping it as **CRLSignerCertificate**.
 3. If **entry.IsPAA** is TRUE:
 - a. If **CRLSignerCertificate** encoded a VendorID, validate that it matches **entry.VendorID**. In case of mismatch, stop further processing of the entry, and continue to next entry, if any.
 4. Else if **entry.IsPAA** is FALSE:
 - a. If the **CRLSignerCertificate** is a PAI certificate (rather than a PAI-delegated certificate), validate that the **CRLSignerCertificate** 's VendorID, matches **entry.VendorID**. Else validate that the 'CRLSignerDelegator' 's VendorID, matches **entry.VendorID**. In case of mismatch, stop further processing of the entry, and continue to next entry, if any.
 - b. If the **CRLSignerCertificate** is a PAI certificate (rather than a PAI-delegated certificate), validate that the **CRLSignerCertificate** 's ProductID matches **entry.ProductID**, if present in the certificate's subject. Else, validate that the 'CRLSignerDelegator' 's ProductID matches **entry.ProductID**, if present in the certificate's subject. In case of mismatch, stop further processing of the entry, and continue to next entry, if any.
 5. Validate the certification path containing **CRLSignerCertificate** and 'CRLSignerDelegator' (when present), using the approved PAAs currently in force in the Distributed Compliance Ledger as the trust anchors. In case of failure, stop further processing of the entry, and continue to next entry, if any.
 6. Obtain the CRL at **entry.DataURL**, keeping it as **CRLFile**.
 7. Perform **CRLFile** validation:
 - a. Validate that the **crlExtensions** field of **CRLFile** has an **Authority Key Identifier** extension matching the **CRLSignerCertificate** 's **Subject Key Identifier**. In case of mismatch, stop further processing of the entry, and continue to next entry, if any.
 - b. If more than one entry exists in the Distributed Compliance Ledger where the **entry.VendorID** and **entry.IssuerSubjectKeyID** are matching, then:

- i. Validate that the **CRLFile** has the Issuing Distribution Point critical CRL extension present (see [RFC 5280](#) section 5.2.5), and then:
 - A. Validate that the **distributionPoint** field of the extension has a single **GeneralName** of type **uniformResourceIdentifier**.
 - B. Validate that the **uniformResourceIdentifier** in the **GeneralName** field matches exactly, byte-for-byte, the value found in the **entry.DataUrl** field.
 - c. If any of above validations fail, stop further processing of the entry, and continue to next entry, if any.
8. Validate the CRL according to Section 6.3 of [RFC 5280](#), using the certificate specified in **CRLSignerCertificate** as the trust anchor. Delta CRLs are not supported, so the "use-deltas" input variable SHALL be false. Note that since the conceptual algorithm in this section constructs a revocation set in advance without access to the candidate certificate, validations from Section 6.3 of [RFC 5280](#) SHALL be adapted by the implementer, and implementations MAY use any algorithm that conforms to the algorithm equivalence clause of Section 6.3. In case of failure, stop further processing of the entry, and continue to next entry, if any.
 9. If **entry.IsPAA** is true and **entry.CRLSignerCertificate** is not self-signed, then assign the subject field of the PAA Certificate that signed **entry.CRLSignerCertificate** to **CertificateAuthorityName**. Otherwise, if **entry.CRLSignerDelegator** is present and non-empty, then assign its subject field to **CertificateAuthorityName**. Otherwise assign the subject field of **entry.CRLSignerCertificate** to **CertificateAuthorityName**. In each case, similarly assign the value of the corresponding subject key identifier extension to **CertificateAuthorityKeyIdentifier**. The latter value should match **entry.IssuerSubjectKeyID**. In case of mismatch, stop further processing of the entry, and continue to next entry, if any.
 10. For each entry in the **CRLFile** 's **CertificateList.tbsCertList.revokedCertificates** sequence:
 - a. If the entry in **revokedCertificates** has the **certificateIssuer** entry extension (or if that extension appears in a previous entry in the sequence but also applies to this entry as specified in section 5.3.3 of [RFC 5280](#)) and the **CertificateAuthorityName** variable does not match the distinguished name found in the **certificateIssuer** extension, ignore this entry.
 - b. Access the correct revocation set by the tuple of (**CertificateAuthorityKeyIdentifier**, **CertificateAuthorityName**). If the set is missing, create it and assign the (**CRLSignerCertificate**, **CRLSignerDelegator**) tuple with the set.
 - c. Set the **RevokedEntitySerialNumber** to the revoked certificate's **userCertificate** field (which is a serial number).
 - d. Insert the **RevokedEntitySerialNumber** entry in the revocation set found earlier.
 - Optionally record metadata to record when each revocation distribution point requires an update.

NOTE

Validation failures during this algorithm are likely to indicate deliberate compromise or accidental failure of the CRL generation or distribution systems, which might result in accepting revoked certificates. Thus, such failures SHOULD be noted

and investigated.

After execution of this algorithm, an oracle exists which MAY then be used to determine the revocation of any PAI or DAC certificate.

6.2.4.2. Determining Revocation Status of an Entity

During the device attestation procedure, a Commissioner SHALL use a revocation set it maintains to determine whether the PAI and DAC certificates are revoked, unless the Commissioner does not have access to the revocation set due to a transient lack of access to necessary resources it uses to maintain the revocation set. If the revocation set was unavailable, the Commissioner SHOULD notify the user of the fact that commissioning could succeed for some non-genuine devices, due to lack of access to some of the necessary information.

If the revocation set is available but no updated CRL is available beyond the timestamp in the "nextUpdate" field in the PAI or the PAA CRL, the Commissioner MAY assume that no updated revocation information is being provided by the CA. The Commissioner SHOULD notify the user of the fact that commissioning could succeed for some non-genuine devices, due to lack of access to some of the necessary information.

The revocation set SHOULD at least include the [Device Attestation PKI Revocation Distribution Points Schema](#) of the [Distributed Compliance Ledger](#).

Since there might be multiple formats for revocation information developed over time, the algorithm in this section is purposefully described in generic terms that would allow expansion to future formats.

The following conceptual algorithm is an example of how to determine whether a DAC or PAI certificate is "revoked" or "not revoked":

- Inputs:
 - The **AuthorityKeyIdentifier** extension of the certificate being queried.
 - The **issuer** name field of the certificate being queried.
 - The **serialNumber** field of the certificate being queried.
 - A conceptual **RevocationDataSet** indexed by (**CertificateAuthorityKeyIdentifier**, **CertificateAuthorityName**), yielding a set of **RevokedEntitySerialNumber**, and an associated tuple of (**CRLSignerCertificate**, **CRLSignerDelegator**) values.
 - **CertificateAuthorityKeyIdentifier** is the Subject Key Identifier of a PAA or PAI as an octet string (e.g. ASN.1 **OCTET STRING**, or protobuf **bytes**).
 - **CertificateAuthorityName** is the X.500 Name associated with the PAA or PAI whose revocation information is being looked-up.
 - **RevokedEntitySerialNumber** is a serial number containing as an ASN.1 **INTEGER** as seen in the **CertificateSerialNumber** field of an [RFC 5280](#)-conformant certificate.
 - **CRLSignerCertificate** is the certificate that signed the CRL.
 - **CRLSignerDelegator** if present, is the certificate that signed the CRLSignerCertificate.

- See [Section 6.2.4.1, “Conceptual algorithm for revocation set construction”](#) for a usable method to generate this dataset.
- Outputs:
 - A boolean value **IsRevoked**, where true indicates that the entity is revoked, and false indicates that it is not revoked.
- Algorithm:
 1. Initialize **IsRevoked** to FALSE.
 2. Obtain the correct **RevocationSet** by indexing the **RevocationDataSet** by the **AuthorityKeyIdentifier** and **issuer** of the certificate whose status is being determined.
 3. If no matching **RevocationSet** is found, return immediately.
 4. If a matching **RevocationSet** is found, determined if **serialNumber** is a member of the set.
 - a. If the entity type whose revocation status is being verified is a PAI, then the subject PAI’s issuer (a PAA) SHALL fulfill one of the following two cases, otherwise return immediately:
 - i. The Subject and Subject Key of the PAI certificate’s issuer matches exactly the CRLSignerCertificate’s subject (i.e. the CRLSignerCertificate is a PAA).
 - ii. The Subject and Subject Key of the PAI certificate’s issuer matches exactly the PAA which is the issuer of the CRLSignerCertificate (i.e. the CRLSignerCertificate is a CRL signer delegated by a PAA).
 - b. If the entity type whose revocation status is being verified is a DAC, then:
 - i. If the **CRLSignerDelegator** is present, then the subject DAC’s issuer (a PAI) SHALL match the **CRLSignerDelegator** in both Subject Key and Subject, otherwise return immediately.
 - ii. If the **CRLSignerDelegator** is absent, then the subject DAC’s issuer (a PAI) SHALL match the **CRLSignerCertificate** in both Subject Key and Subject, otherwise return immediately.
 - c. If the **serialNumber** is part of the set, set **IsRevoked** to TRUE.

The certificate in question is deemed revoked if **IsRevoked** was TRUE after execution of the above algorithm.

6.3. Certification Declaration

A Certification Declaration (CD) is a cryptographic document that allows a Matter device to assert its protocol compliance. It is encoded in a **CMS** format described in [RFC 5652](#). Upon successful completion of certification by a device type, Connectivity Standards Alliance creates the CD for that device type so that it can be included in the device firmware by the manufacturer.

6.3.1. Certification Declaration (CD) Format

The Certification Declaration is a **CMS**-encoded SignedData payload (see [RFC 5652](#) section 5) whose EncapsulatedContentInfo **eContent** field contains a TLV-encoded **certification-elements** structure

with an anonymous tag:

Certification Elements TLV structure

```
certification-elements => STRUCTURE [ tag-order ]
{
    format_version [0]      : UNSIGNED INTEGER [ range 16-bits ]
    vendor_id [1]           : UNSIGNED INTEGER [ range 16-bits ]
    product_id_array [2]    : ARRAY [ length 1..100 ] OF UNSIGNED INTEGER [ range 16-
bits ]
    device_type_id [3]      : UNSIGNED INTEGER [ range 32-bits ]
    certificate_id [4]      : STRING [ length 19 ]
    security_level [5]      : UNSIGNED INTEGER [ range 8-bits ]
    security_information [6] : UNSIGNED INTEGER [ range 16-bits ]
    version_number [7]      : UNSIGNED INTEGER [ range 16-bits ]
    certification_type [8]  : UNSIGNED INTEGER [ range 8-bits]
    dac_origin_vendor_id [9, optional] : UNSIGNED INTEGER [ range 16-bits ]
    dac_origin_product_id [10, optional] : UNSIGNED INTEGER [ range 16-bits ]
    authorized_paa_list [11, optional] : ARRAY [ length 1..10 ] OF OCTET STRING [
length 20 ]
}
```

The Certification Elements TLV is encoded with data to form a **cd_content** message to be signed.

```
cd_content =
{
    format_version (0)      = 1,
    vendor_id (1)           = <vendor_id>,
    product_id_array (2)    = <array of product_id values>,
    device_type_id (3)      = <primary device type identifier>,
    certificate_id (4)      = <globally unique certificate ID issued by
Connectivity Standards Alliance>,
    security_level (5)      = 0,
    security_information (6) = 0,
    version_number (7)      = <version_number>,
    certification_type (8)  = <certification_type>,
    dac_origin_vendor_id (9) = <Vendor ID associated with the DAC, optional>,
    dac_origin_product_id (10) = <Product ID associated with the DAC, optional>,
    authorized_paa_list (11) = <array of PAA SKIs, optional>
}
```

The **format_version** field SHALL contain the value 1.

The **vendor_id** field SHALL contain the **Vendor ID** associated with the Certification Declaration.

The **product_id_array** field SHALL contain an array of a number of **Product IDs** which are covered by the same certification (e.g. certification by similarity). All other fields of a Certification Declaration apply to all products in this array.

The **device_type_id** field SHALL contain the **device type** identifier for the primary function of the device. For example, if **device_type_id** is 10 (0x000A), it would indicate that the device has a primary function of a Door Lock device type. See also the **_T** subtype in [Section 4.3.1.3, “Commissioning Subtypes”](#).

The **device_type_id** field in a given Certification Declaration indicates the **primary device type** and SHOULD match the **device_type_id** value in the **DCL entries** associated with the VendorID and ProductID combinations present in that Certification Declaration.

The **certificate_id** field SHALL contain a globally unique serial number allocated by the Connectivity Standards Alliance for this Certification Declaration.

The **security_level** and **security_information** fields are reserved for future use and SHALL be ignored at read time, and set to zero at issuance time.

The **version_number** field SHALL contain a version number for the CD itself, assigned by the Connectivity Standards Alliance, that matches the Vendor ID and Product ID used in a **DeviceSoftwareVersionModel** entry in the **Distributed Compliance Ledger** matching the certification record associated with the product presenting this CD.

The value of the **version_number** field is not meant to be interpreted by commissioners and SHALL be recorded as assigned.

The value of the **version_number** field is not related to any of the associated products' firmware or software version number, but is rather a numerical version number for the CD itself.

The **certification_type** field SHALL contain the type of certification for this CD, interpreted according to the following table:

certification_type	meaning
0	used for development and test purposes
1	provisional - used to allow production and distribution to occur in parallel with certification (with potential software fixes yielding a higher SoftwareVersion which gets certification)
2	official - allocated after passing certification
other values	reserved

For details about the usage of the **certification_type** field in the Device Attestation Procedure, see [failure of Device Attestation Procedure](#).

A product MAY include a Certification Declaration with **certification_type**=1 (provisional) in factory firmware if the Certification Declaration can be replaced via software update.

The **dac_origin_vendor_id** field, if present, SHALL contain the **Vendor ID** value expected to be found

in the Device Attestation Certificate's subject DN.

The `dac_origin_product_id` field, if present, SHALL contain the [Product ID](#) value expected to be found in the Device Attestation Certificate's subject DN.

The `dac_origin_vendor_id` and `dac_origin_product_id` SHALL only be present together.

The use of the `dac_origin_vendor_id` and `dac_origin_product_id` fields allows for a target of the device attestation procedure to have a manufacturing provenance which differs from the entity that obtains the ultimate certification. If present, they tie a given Certification Declaration to an original manufacturer's device attestation chain of trust, so that DACs MAY be issued at manufacturing time without *a priori* knowledge of the ultimate vendor.

The optional `authorized_paa_list` field, if present, SHALL contain a list of one or more Product Attestation Authority (PAA) which is/are authorized (by the device manufacturer) to sign the Product Attestation Intermediate (PAI) Certificate which signs the Device Attestation Certificate for a product carrying this Certification Declaration. Each such PAA is identified by the Subject Key Identifier (SKI) extension value of its certificate.

Any context-specific tags not listed in the above schema for Certification Elements SHALL be reserved for future use, and SHALL be silently ignored if seen by a Commissioner which cannot understand them.

See [Section 6.2.3, "Device Attestation Procedure"](#) for more details about usage of the Certification Declaration fields.

Certification Declaration CMS ASN.1 Encoding Format

```

CertificationDeclaration ::= SEQUENCE {
    version                INTEGER ( v3(3) ),
    digestAlgorithm         OBJECT IDENTIFIER sha256 (2.16.840.1.101.3.4.2.1),
    encapContentInfo        EncapsulatedContentInfo,
    signerInfo              SignerInfo }

EncapsulatedContentInfo ::= SEQUENCE {
    eContentType            OBJECT IDENTIFIER pkcs7-data (1.2.840.113549.1.7.1),
    eContent                OCTET STRING cd_content }

SignerInfo ::= SEQUENCE {
    version                INTEGER ( v3(3) ),
    subjectKeyIdentifier    OCTET STRING,
    digestAlgorithm         OBJECT IDENTIFIER sha256 (2.16.840.1.101.3.4.2.1),
    signatureAlgorithm      OBJECT IDENTIFIER ecdsa-with-SHA256 (1.2.840.10045.4.3.2),
    signature               OCTET STRING }

```

The Certification Declaration encoding rules:

1. The format SHALL only support CMS `version` v3.

2. The `digestAlgorithm` SHALL use the `sha256` algorithm.
3. The `signatureAlgorithm` SHALL use the `ecdsa-with-SHA256` (ECDSA with SHA256) and `secp256r1` curve, as defined in Section 2.4.2 of [SEC 2](#).
4. The `eContentType` SHALL use the `pkcs7-data` type.
5. The `subjectKeyIdentifier` SHALL contain the subject key identifier (SKI) of a well-known Connectivity Standards Alliance certificate, that was used to generate the `signature`. The format of the key identifiers supported is available as part of the Certification Policy. A product MAY contain a Certification Declaration signed by a test certificate with the `certification_type` flag set to 0 for certification testing. Certification Declarations used by products in production environments SHALL be signed by a valid Connectivity Standards Alliance certificate.
6. The `eContent` field SHALL contain the TLV-encoded Certification Elements TLV structure (see [Certification Elements TLV structure](#)).

Note that Certification Declarations SHALL NOT be generated by any Node, but rather, they SHALL be stored and transmitted to a Commissioner by a Commissionee during the conveyance of the [Attestation Information](#) in response to an [Attestation Request command](#).

See [Section F.1, “Certification Declaration CMS test vector”](#) for a complete example of generating a Certification Declaration.

6.3.2. Firmware Information

Firmware Information is an optional component of the [Device Attestation Information](#) (see [Section 6.2.3, “Device Attestation Procedure”](#)).

Firmware Information MAY contain one or more Firmware Digests that correspond to the components in the firmware that have been measured and recorded during the boot process (e.g., bootloader, kernel, root filesystem, etc), and MAY contain other metadata. A Firmware Digest SHALL either represent a hash of the corresponding firmware layer or a hash of the signed manifest that was used to validate the corresponding firmware layer during secure boot. An implementation MAY choose to hash the measurements of all components into a single hash and include only that hash in Firmware Information.

A device MAY report Firmware Information containing its firmware digests only if it implements a secure subsystem that protects the device attestation private key and is able to securely collect and report firmware digests as shown in [Figure 55, “Illustration of the measured boot process”](#). This process is known as "measured boot".

Ideally, the measured boot process SHOULD be rooted in silicon such as a boot ROM, similar to the secure boot process found in many systems-on-chip (SoCs). Since many SoCs and microcontrollers are unable to perform measured boot in hardware, the process SHOULD start at the earliest firmware component possible (for example, at the bootloader shown in the figure below). In this case, this firmware component is not measured and in fact, it is the root for measurement. Therefore, it SHALL be resistant to attacks compromising subsequent firmware components (e.g., the ROM must verify its authenticity (secure boot) or it may be placed in a locked partition at the factory that cannot be updated by software in the field).




Figure 55. Illustration of the measured boot process

The device secure subsystem SHALL use the device attestation private key to sign [attestation-elements](#) and [NOCSR-elements](#). The device secure subsystem SHALL fill the [attestation-elements](#) fields using information compiled into its image or generated during the measured boot process. The device secure subsystem SHALL validate all signing requests so that if the device software, but not its secure subsystem, gets compromised it cannot act as a signing oracle to sign Attestation Information Responses with fake Firmware Digests.

The [firmware_information](#) field in [attestation-elements](#) SHALL NOT be generated by devices that do not implement a separate secure subsystem, in software or hardware, which maintains and controls the use of the device attestation private key.

For devices that support secure boot, it is straightforward to add support for measured boot. Specifically, the hashes of the different firmware components that are already generated and verified sequentially during secure boot SHALL be collected and stored for reporting. Devices that do not support secure boot MAY implement measured boot by generating the hashes in software during the boot process implementing the root for measurement in the earliest firmware component.

If a device that chooses to send Firmware Digests and which supports an industry-standard measured boot architecture and which can generate signed firmware attestation reports, the secure subsystem in the device MAY validate the firmware attestation reports locally and SHALL report the raw firmware digests in [attestation-elements](#) so that the [firmware_information](#) field in [attestation-elements](#) has the same values in all devices of the same model that run the specific Software Image.

Firmware Digests SHALL NOT be reported by devices that implement a single firmware component in the boot chain, because there is nothing to measure and report subsequently, unless they have support for measured boot built in the device's boot ROM.

Commissioners MAY use the reported firmware information to confirm that the firmware version is authorized to run on the device, that it has not been revoked, or that it does not contain known vulnerabilities. Commissioners and Administrators that choose to verify this information SHOULD refer to canonical databases, such as the Distributed Compliance Ledger (see [Section 11.23](#), “Distrib-

uted Compliance Ledger”) to validate that the reported firmware information matches what is expected for an authorized Software Image associated with a given Certification Declaration. The firmware information, when validated, SHALL be validated as an opaque well-known octet string. Internal semantic validation MAY be applied for error-reporting, but the exact format is out of the scope of this specification.

In cases where a Commissioner or Administrator detects such an invalid or problematic firmware version, Commissioners and Administrators MAY, after consultation with the user, refuse to commission the device, provide it with operational credentials, or otherwise operate it, until the firmware has been updated, to avoid putting the user at risk from compromised software.

6.3.3. Firmware information validation examples

Below is an illustrative example of the Commissioner actions to validate the firmware information.

1. Retrieve the `firmware_information` field from `attestation-elements`
2. Retrieve all `Distributed Compliance Ledger DeviceSoftwareVersionModel` entries for the Commissionee’s Vendor ID and Product ID.
3. Verify that there is a valid, non-revoked, entry where the `FirmwareInformation` field exactly matches the `firmware_information` field in `attestation-elements`.
4. If verification fails, report error to the user
5. If verification succeeds, proceed with device commissioning

Below is an example of the corresponding Device actions. For illustrative purposes, it is assumed that the device implements a secure subsystem that maintains the private device attestation key and signs `attestation-elements` using this key but it does not have direct hardware support for measure boot. This is expected to be the common case for many devices covered by this version of the specification. Consequently, the measurement process can only start from the bootloader shown in the figure above.

1. The device bootloader produces a measurement of the OS kernel using a supported hash algorithm from RFC 5912 and delivers it to the secure subsystem.
2. The secure subsystem receives the measurement and stores in a location inaccessible to the OS.
3. The OS kernel produces a hash of the root filesystem and delivers the measurement to the secure subsystem.
4. When the secure subsystem is asked to sign an `attestation-elements` structure using its private device attestation key, it generates two `FirmwareDigests` or one combined `FirmwareDigest` from these measurements, fills the `firmware_information` field in `attestation-elements` using these measurements, fills the CD blob compiled into the secure environment and signs the `attestation-elements` structure.

The Device Vendor is responsible to provide the `FirmwareInformation` field when a new Software Image entry is reported in the corresponding `Distributed Compliance Ledger` entry.

Below is an exemplary ASN.1 schema for an encoding scheme that could be used to encode firmware information.

Firmware Information encoding example

```

HashAlgorithm ::= SEQUENCE {
    id OBJECT IDENTIFIER,
    params ANY OPTIONAL
}

FirmwareDigest ::= SEQUENCE {
    digestAlgorithm HashAlgorithm,
    digestHash      OCTET STRING
}

FirmwareInformation ::= SEQUENCE {
    firmwareDigests SEQUENCE OF FirmwareDigest
}

-- Example HashAlgorithm id
id-sha256 OBJECT IDENTIFIER ::= {
    joint-iso-itu-t(2) country(16) us(840) organization(1) gov(101)
    csor(3) nistAlgorithms(4) hashalgs(2) 1
}

-- Below is an example value for the above exemplary FirmwareInformation

firmwareInformation FirmwareInformation ::= {
    -- The firmwareDigests contain two values, for two separate components.
    firmwareDigests {
        {
            digestAlgorithm {
                id id-sha256
            },
            digestHash '00112233445566778899AABBCCDDEEFF00112233445566778899AABBCCDDEEFF'H
        },
        {
            digestAlgorithm {
                id id-sha256
            },
            digestHash '101112131415161718191A1B1C1D1E1F101112131415161718191A1B1C1D1E1F'H
        }
    }
}

```

The above example would yield the following DER-encoded octet string:

```
30663031 300d0609 60864801 65030402 01050004 20001122 33445566 778899AA
BBCCDDEE FF001122 33445566 778899AA BBCCDDEE FF303130 0D060960 86480165
03040201 05000420 10111213 14151617 18191A1B 1C1D1E1F 10111213 14151617
18191A1B 1C1D1E1F
```

6.4. Node Operational Credentials Specification

6.4.1. Introduction

The Node Operational credentials are a collection of credentials to enable a Node to identify itself within a Fabric. The Node Operational credentials are distinct from the Device Attestation credentials. The Node Operational credentials are installed during [Commissioning](#).

The Node Operational credentials include the following items:

- [Node Operational Key Pair](#)
- [Node Operational Certificate \(NOC\)](#)
- [Intermediate Certificate Authority \(ICA\) Certificate \(optional\)](#)
- [Trusted Root Certificate Authority \(CA\) Certificate\(s\)](#)

Each Node in a Fabric is identified with a [Node Operational Identifier](#). In order to securely identify the Node, the Node Operational Identifier is bound to the Node Operational Public Key as both are contained within the signed NOC. The Node Operational Identifier is a constituent part of the subject field of the NOC, according to the rules described in [Matter DN Encoding Rules](#). A connecting Node can attest to the validity of the Node Operational Public Key and the Node Operational Identifier in a received NOC because the NOC is signed by a CA that the connecting Node trusts. Used with [Certificate Authenticated Session Establishment \(CASE\)](#), these data provide the basis for secure communications on the Fabric.

6.4.2. Node Operational Credentials Management

Commands from the [Operational Credentials Cluster](#) are used to install and update Node Operational credentials.

A Node receives its initial set of Node Operational credentials through the [AddNOC](#) command when it is commissioned onto a Fabric by a Commissioner.

Once installed, Node Operational credentials MAY be updated by an Administrator with the appropriate privileges using the [UpdateNOC](#) command.

Once installed, Node Operational credentials MAY be removed by an Administrator with the appropriate privileges using the [RemoveFabric](#) command. The removal uses [RemoveFabric](#), since the Fabric association for the given Node Operational credentials may underpin a variety of bindings and other [fabric-scoped](#) configuration, which would remain in an inconsistent state if the Node Operational credentials alone were removed, as opposed to the entire associated Fabric and data.

6.4.3. Node Operational Identifier Composition

The Node Operational Identifier is used for Node discovery and network address resolution within a network segment. The **FabricID** portion of the Node Operational Identifier serves a scoping purpose to identify disjoint operational Fabrics within a given network segment. The **NodeID** portion of the Node Operational Identifier is the logical addressing identifier used:

- within Message-layer messages for logical addressing (see [Section 4.4, “Message Frame Format”](#))
- within Data Model bindings to express data subscription relationships between Nodes (see [Chapter 9, System Model Specification](#))
- within Access Control List Entries to refer to individual Nodes as access control grantees (subjects) when **CASE** sessions are used for communication (see [Section 9.10, “Access Control Cluster”](#))

In addition to the **FabricID** and **NodeID**, a Node Operational Identifier MAY include at most three 32-bit **CASEAuthenticatedTag** ([1.3.6.1.4.1.37244.1.6](#)) attributes used to tag the operational identifier to implement access control based on [CASE Authenticated Tags](#).

The Fabric ID is a 64-bit value that identifies the Fabric and is scoped to a particular Root CA. For example, two fabrics with the same Fabric ID are not equivalent unless their Root CA are the same. The Fabric ID MAY be chosen randomly or algorithmically but it SHALL be allocated uniquely within the set of all possible Fabric IDs for which a given Root CA will sign operational certificates. Before allocating the Fabric ID, the Commissioner SHOULD attempt to ensure that an existing Fabric is reused and joined, if any is applicable from the perspective of the Commissioner in the current commissioning context. The method used for determining local Fabric ID existence is vendor-specific.

The Node ID is a 64-bit value that identifies a Node within a Fabric. The Node ID MAY be chosen randomly or algorithmically but it SHALL be allocated uniquely within the Fabric before it is given to the Node or otherwise used. The Node ID SHALL be chosen, by a Commissioner, at the time of Node commissioning.

The uniqueness constraint for Fabric ID is only required to be ensured within the scope of the Root CA serving the Commissioner.

6.4.4. Node Operational Key Pair

A Node Operational Key Pair, comprised of a Node Operational Public Key and a Node Operational Private Key, is created using the **Crypto_GenerateKeypair** function. A new Node Operational Key Pair is generated for each Commissioning Session in accordance with [security requirements](#).

6.4.5. Node Operational Credentials Certificates

All certificates in the Node Operational credentials are X.509v3 certificates compliant with [RFC 5280](#), encoded in such a way that they respect the constraints in the [Operational_Certificate](#) section. They may be encoded as [X.509v3](#) certificates or [Matter Operational Certificates](#) (“Matter Certificates” thereafter). The signature field of a certificate SHALL be calculated using the X.509v3 encoding of the certificate.

6.4.5.1. Node Operational Certificate (NOC)

The NOC SHALL be issued by either a Root CA trusted within the Fabric or by an Intermediate Certificate Authority (ICA) whose ICA certificate is directly issued by such a Root CA. The NOC is bound to the Node Operational Key Pair through the [Node Operational Credential Signing Request \(NOCSR\)](#).

The validity period specifies the time period for which a NOC is valid. For constrained or sleepy devices that lack accurate time, enforcement of an NOC's validity period MAY be omitted.

6.4.5.2. Intermediate CA (ICA) Certificate

In the case where an intermediate CA (ICA) issues the NOC, the ICA certificate is used to attest to the validity of the NOC. The Root CA certificate associated with the issuer of the ICA certificate is used in turn to attest to the validity of the ICA certificate.

6.4.5.3. Trusted Root CA Certificates

Each Node has one or more trusted Root CA certificates in its Node Operational credentials that it uses to verify ICA certificates and Node Operational Certificates presented by other Nodes, treating them as trust anchors as described in [RFC 5280](#). A Root CA certificate is self-signed. They are not verified but rather trusted because they were provisioned by a trusted Commissioner.

In the case where a Root CA issues the NOC, the Root CA certificate is used to attest to the validity of the NOC.

The trusted Root CA certificates that a Device trusts when the Device is verifying operational certificates are those stored in the [TrustedRootCertificates attribute](#) of that Device's [Operational Credentials cluster](#).

A device MAY have Root CA certificates that it trusts for purposes other than for operational credential verification. These certificates SHALL NOT appear in any Node's [TrustedRootCertificates attribute](#) of the [Operational Credentials cluster](#). The certificates configured in that cluster SHALL only be added during the commissioning process by the Commissioner, or during root rotation operations by an Administrator already trusted by the Node. Nodes SHALL NOT modify the [Trusted-RootCertificates attribute](#) outside of the processing of [Operational Credentials cluster](#) commands.

The figures below show the Node Operational Certificate hierarchies, with and without optional ICAC.




Figure 56. Node Operational Certificate PKI hierarchy with optional ICAC




Figure 57. Node Operational Certificate PKI hierarchy without optional ICAC

6.4.5.4. Vendor ID Verification Signer Certificate (VVSC)

The [Fabric Table Vendor ID Verification Procedure](#) MAY rely on a Vendor Verification Signer Certificate (VVSC) in some cases.

These certificates appear in the [Operational Trust Anchors Schema](#) of the [Distributed Compliance Ledger](#).

The VVSC allows a fabric operator to cryptographically associate its Vendor ID with a fabric, even if the root of trust of the fabric is locally unique, and not shared with anyone else outside the

premises or operational network on which it operates.

6.4.6. Node Operational Credentials Procedure

The following procedure is used by a Node to obtain an Operational Credential. This procedure is part of [Commissioning](#).

6.4.6.1. Node Operational Certificate Signing Request (NOCSR) Procedure

After the Commissioner validates Device Attestation Information, the following procedure is used to generate a Node Operational Key Pair and obtain the NOCSR.

1. The Commissioner SHALL generate a random 32 byte nonce named **CSRNonce** using **Crypto_DRBG()**.
2. The Commissioner SHALL send the CSRNonce to the Node and request NOCSR Information using the **CSRRequest Command**.
 - a. The Node SHALL create a new candidate Node Operational Key Pair, using **Crypto_GenerateKeyPair()**, valid for the duration of the **Fail-Safe Context** currently in progress.
 - b. The Node SHOULD verify that the newly generated candidate Node Operational Key Pair does not match any other existing Node Operational Key Pair on the device. If such a key collision was to be found, it would indicate a key pair that was not properly randomly generated. The procedure SHALL fail if such a collision is detected. See [Section 11.18.6.5, “CSRRequest Command”](#) for the error generated in that situation.
 - c. The candidate Node Operational Key Pair SHALL only be committed to persistent storage upon successful execution of the next **AddNOC Command** executed with a Node Operational Certificate whose public key matches the candidate key.
 - d. The Node SHALL create a Certificate Signing Request (CSR) by following the format and procedure in **PKCS #10**, which includes a signature using the Node Operational Private Key (see [RFC 2986](#) section 4.2).
 - e. The CSR's subject MAY be any value and the device SHOULD NOT expect the final operational certificate to contain any of the CSR's subject DN attributes.
3. The Node SHALL generate and return the NOCSR Information (see [Section 11.18.4.9, “NOCSR Information”](#) for encoding) to the Commissioner using the **CSRResponse Command**. The NOCSR Information includes a signature using the Device Attestation Private Key.

Node Operational CSR Information Validation

1. The Commissioner SHALL validate the Device Attestation Signature (**attestation_signature**) field from **CSRResponse Command**:

```
Success = Crypto_Verify(  
    publicKey = Public key from DAC,  
    message = NOCSR Information TBS (nocsr_tbs),  
    signature = Device Attestation Signature (attestation_signature)  
)
```

where the fields are encoded as described in [Section 11.18.4.9, “NOCSR Information”](#).

- The **AttestationChallenge** SHALL be obtained from a **CASE session**, **resumed CASE session**, or **PASE session** depending on the method used to establish the secure session within which device attestation is conducted.
 - The CSR Nonce in **NOCSR Information** SHALL match the Commissioner’s CSR Nonce.
2. The inner signature in the PKCS#10 **csr** sub-field of the **CSRResponse Command**’s **NOCSRElements** field SHALL be verified, per the definition of CSR signatures in **PKCS #10**.




Figure 58. Node Operational Credentials flow

6.4.7. Node Operational Certificate Signing Request (NOCSR)

A Node creates a NOCSR in response to the Commissioner, so that the Commissioner can request a **NOC** on the Node’s behalf from its trusted Certificate Authority. The CSR itself SHALL follow the encoding and rules from **PKCS #10**, with the minimum attributes shown in the example below.

Note that the subject field MAY be any value.

*NOC*SR

Certificate Request:

Data:

Version: 1 (0x0)

Subject:

Subject Public Key Info:

Public Key Algorithm: id-ecPublicKey

Public-Key: (256 bit)

pub:

04:12:3b:90:f5:.....

ASN1 OID: prime256v1

NIST CURVE: P-256

Attributes:

Requested Extensions:

Signature Algorithm: ecdsa-with-SHA256

30:46:02:21:00:95:ff:.....

6.4.8. Node Operational Certificate Renewal

A [NOC](#) can be renewed by an Administrator (a Node that has Administer privileges on the Node to be updated). The Administrator triggers the process by sending an [CSRRequest Command](#).

6.4.9. Node Operational Certificate Revocation

A Node's access to other Nodes can be revoked by removing the associated Node ID from [Access Control Entry subjects](#) where it appears. This action is taken by an Administrator which has the privilege to update the Access Control Cluster for its Nodes.

6.4.10. Fabric Table Vendor ID Verification Procedure

NOTE

Vendors operating operational CAs need to consider whether they are operationally sharing a root with other vendors. This could cause the potential for that other vendor to associate NOC chains with a different vendor, since validation of chaining is done only against the trust anchor (root).

The [Fabrics attribute](#) of the [Operational Credentials Cluster](#) contains a table of all fabrics on which a node is operating, recorded using [FabricDescriptorStructs](#).

The VendorID field of the Fabric attribute MAY be used to identify which organization among the members of the Connectivity Standards Alliance has created the fabric. This information MAY assist in expressing to end-users which "ecosystem" or manufacturer has control of the node via that fabric.

Since it is critical to avoid claiming that a fabric is operated by an entity that the user trusts when it's not actually the case, there is a requirement to be able to authenticate the statement of associa-

tion of a fabric to a given Vendor ID.

The Fabric Table Vendor ID Verification Procedure aims to verify the authenticity of a Vendor ID bound to a given [Fabric Reference](#) composed of a root public key and Fabric ID.

The procedure is uniform across fabrics and MAY be used even to validate the Vendor ID association of fabrics other than the accessing fabric executing the procedure. This is especially important when providing information to an end-user or other client about all fabrics to which a device is joined.

The procedure supports all types of operational credentials, whether involving a global operational root of trust (i.e. a large number of fabrics share a same root of trust), or not.

NOTE

Nodes implementing a version pre-dating version 1.4.2 of this specification will not have the necessary commands to execute this procedure. This is a known deficiency that will resolve itself with time, since the procedure was introduced in 1.4.2, along with revision 2 of the [Operational Credentials Cluster](#). Commissioners and administrators SHALL NOT assume that the Vendor ID association of a fabric in the fabric table is genuine unless this procedure is executed. This includes the case of all devices that pre-date this feature, which means that older implementations will simply appear with all fabrics being possibly not genuine.

There are several components to the procedure:

- A [vendor_fabric_binding_message](#) binding together the fabric reference and the VendorID of its administrator.
 - This message, along with a fresh challenge, will be signed using the NOC operational key by the challenged fabric member, whenever the procedure is performed.
- An optional [vid_verification_statement](#) authenticating a fabric reference as being associated with a particular VendorID, if the root of trust associated with the NOC is not global.
 - In some operational credential chain cases, such as when the root of trust of the fabric is retained by local equipment only (e.g. one root per home), the [vid_verification_statement](#) allows delegation to a separate signer whose public key is globally trusted. This avoids needing to record every root of trust in a global trust store.
- Execution of the procedure by a client invoking a command on a node's Operational Credentials cluster, followed by cross-validation of the response against other data from the node and from the [Distributed Compliance Ledger](#), to determine if the fabric truly is associated with the VendorID.

The procedure is constructed such that if the association of Vendor ID with the fabric is validated, the node that generated the response is also validated to definitely be a trusted part of that fabric.

6.4.10.1. Algorithm

There are two parties:

1. a "verifier" client;

2. a candidate node whose membership in the fabric, and associated VendorID for the fabric, will be validated ("candidate" thereafter).

The algorithm of the procedure at the verifier is as follows:

1. Read the **NOCs** attribute of the Operational Credentials Cluster using a non-fabric-filtered read.
2. Read the **Fabrics** attribute of the Operational Credentials Cluster using a non-fabric-filtered read.
3. Read the **TrustedRootCertificates** attribute of the Operational Credentials Cluster using a non-fabric-filtered read.
4. Select a fabric listed in the **Fabrics** attribute whose **VendorID** field is to be verified.
5. Generate a 32-octet random challenge **client_challenge** := **Crypto_DRBG**(32 * 8)
6. Invoke the **SignVIDVerificationRequest** command of the Operational Credential Cluster against the candidate using the **FabricIndex** of the entry in **Fabrics** whose **VendorID** is being verified, and the **client_challenge** generated in a previous step. This yields **SignVIDVerificationResponse**(**fabric_index**, **fabric_binding_version**, **signature**).
7. Obtain the **attestation_challenge** from **CASE session**, **resumed CASE session** or **PASE session**, depending on the secure session context of the session over which the **SignVIDVerificationRequest** command was sent.
8. Locally generate the **vendor_fabric_binding_message**, following the **fabric_binding_version** obtained in the **SignVIDVerificationResponse**.
 - For a **fabric_binding_version** version of 0x01, generate **vendor_fabric_binding_message** := **fabric_binding_version** (0x01) || **root_public_key** || **fabric_id** || **vendor_id**
 - **root_public_key** is the entire octet string of the **RootPublicKey** field of the entry in **Fabrics** whose **VendorID** is being verified.
 - **fabric_id** is the 64-bit **FabricID** field of the entry in **Fabrics** whose **VendorID** is being verified, encoded as a big-endian value without omitting leading zeroes.
 - **vendor_id** is the 16-bit **VendorID** field of the entry in **Fabrics** whose **VendorID** is being verified, encoded as a big-endian value without omitting leading zeroes.
9. Locally generate the **vendor_id_verification_tbs** message
 - **vendor_id_verification_tbs** := **fabric_binding_version** || **client_challenge** || **attestation_challenge** || **fabric_index** || **vendor_fabric_binding_message** || <**vid_verification_statement**>
 - **fabric_binding_version** is the value from the **SignVIDVerificationResponse**.
 - **client_challenge** is the 32-octet string previously randomly generated.
 - **attestation_challenge** is the attestation challenge value previously obtained from the secure session context.
 - **fabric_index** is the 1-octet value from the **SignVIDVerificationResponse**
 - **vendor_fabric_binding_message** is the octet string generated in a previous step.
 - **vid_verification_statement** is the 85-octet (for cryptographic primitives mapping 1.0) value from the **VIDVerificationStatement** field of the entry in **Fabrics** whose **VendorID** is

being verified, if it was present.

10. Given the subject public key associated with the fabric being verified, validate that `Crypto_Verify(noc_public_key, vendor_id_verification_tbs, signature)` succeeds, otherwise the procedure terminates as failed.
 - `noc_public_key` value is the subject key extracted from the `NOCValue` field of the entry in `NOCs` whose `FabricIndex` field matches the fabric being verified.
 - `vendor_id_verification_tbs` is the value generated in the previous step.
 - `signature` is the value from `SignVIDVerificationResponse` received previously.
11. Verify that the NOC chain is valid using `Crypto_VerifyChain()` against the entire chain reconstructed from both the `NOCs` and `TrustedRootCertificates` attribute entries whose `FabricIndex` field matches the fabric being verified. If the chain is not valid, the procedure terminates as failed.
12. Proceed with cross-referencing device data with the Distributed Compliance Ledger's data, according to the following cases:
 - a. If the `VIDVerificationStatement` field was present in the entry in `Fabrics` whose `VendorID` is being verified:
 - i. Find the `VIDVerificationStatement` signer according to the following steps:
 - A. Validate the `statement_version` octet of the `VIDVerificationStatement`, ensuring it is the known value `0x21`, otherwise the procedure terminates as failed.
 - B. Extract the `vid_verification_signer_skid` field from the `VIDVerificationStatement`.
 - C. If the `VVSC` field of the `NOCs` attribute entry whose `FabricIndex` field matches the fabric being verified is present, then extract that certificate as the `VIDVerificationSignerCertificate`. Otherwise, look up the entry in the `Operational Trust Anchors Schema` under the expected `VendorID` (`VID` field) being verified whose `IsVIDVerificationSigner` field is true and whose `SubjectKeyID` matches the `vid_verification_signer_skid` field value as the `VIDVerificationSignerCertificate`. If `VIDVerificationSignerCertificate` is not found or empty after this step, then the procedure terminates as failed.
 - ii. Validate that the `SubjectKeyID` of the `VIDVerificationSignerCertificate` from the previous step matches the value of the `vid_verification_signer_skid`. If it does not match, the procedure terminates as failed.
 - iii. Validate the path of `VIDVerificationSignerCertificate` using the certificates in the `Operational Trust Anchors Schema` under the expected `VendorID` (`VID` field) being verified. If the path is invalid (e.g. any element of the chain missing, any certificate malformed, any signature verification fails, etc), then the procedure terminates as failed. Only the self-signed certificates among the set SHALL be considered as trust anchors during certificate path validation. The path length SHALL NOT be longer than 3. The verifier MAY apply knowledge of revoked entities obtained from outside the Distributed Compliance Ledger.
 - iv. Locally generate the `vid_verification_statement_tbs` message as described in [Section 6.4.10.3, "Generation of the vid_verification_statement_tbs message"](#). The `version` to use SHALL be the value from the `VIDVerificationStatement`. The `vid_verification_signer_skid` to use SHALL be the one from the `VIDVerificationStatement`.

- v. Validate that `Crypto_Verify(vid_verification_signer_pk, vendor_id_verification_tbs, signature)` succeeds, otherwise the procedure terminates as failed.
 - `vid_verification_signer_pk` is the subject public key extracted from the `VIDVerificationSignerCertificate`.
 - `vendor_id_verification_tbs` is the octet string having the value of `vid_verification_statement_tbs`.
 - `signature` is the signature portion of the `VIDVerificationStatement`.
 - b. Otherwise (i.e. `VIDVerificationStatement` not used):
 - i. Look-up the entry in the `Operational Trust Anchors Schema` under the expected `VendorID` (`VID field`) being verified whose `IsRoot` field is true and whose `SubjectKeyID` matches the `SubjectKeyID` field value of the `TrustedRootCertificates` attribute entry whose `FabricIndex` field matches the fabric being verified (i.e. the RCAC of the candidate fabric). If a matching RCAC certificate is not found or empty after this step, then the procedure terminates as failed. If found, save the certificate as `GloballyTrustedRoot`
 - ii. Verify that the NOC chain is valid using `Crypto_VerifyChain()` against the entire chain reconstructed from the `NOCs` attribute entry whose `FabricIndex` field matches the fabric being verified, but populating the trusted root with the `GloballyTrustedRoot` certificate rather than the value in `TrustedRootCertificates` associated with the candidate fabric. If the chain is not valid, the procedure terminates as failed.
13. Given that all prior steps succeeded, the candidate fabric's VendorID SHALL be deemed authentic.

The figure below illustrates the procedure with a sequence diagram.




Figure 59. Sequence diagram of Fabric Table Vendor ID Verification Procedure

6.4.10.2. Generation of the VIDVerificationStatement message

The **VIDVerificationStatement** message allows a globally trusted signer whose public certificate chains back to entities in the **Operational Trust Anchors Schema** under the expected **VendorID** to authenticate the association of a given VendorID with a fabric reference.

This allows verifying a large number of fabrics, with a large associated number of RCAC with a small number of vendor-bound signers, thus reducing the number of RCAC to store in the Distributed Compliance Ledger associated with each manufacturer or vendor.

The **VIDVerificationStatement** is computed as follows:

1. Locally generate the `vid_verification_statement_tbs` message using the procedure in [Section 6.4.10.3](#), “Generation of the `vid_verification_statement_tbs` message”.
2. Select an issuer (signer) for the `VIDVerificationStatement`. Store the 20-octet SubjectKeyID extension value of the associated issuer public certificate as `vid_verification_signer_skid`.
 - The certificate chain for the issuer SHALL chain back to elements in the [Operational Trust Anchors Schema](#) under the expected `VendorID` that was used to generate `vid_verification_statement_tbs`.
3. Generate `vid_verification_statement_signature` := `Crypto_Sign(signer.private_key, vid_verification_statement_tbs)`.
 - Only the [Cryptographic Primitives](#) in force in this specification SHALL be used for signing.
4. Generate `VIDVerificationStatement` := `statement_version || vid_verification_signer_skid || vid_verification_statement_signature`.
 - The `statement_version` SHALL be the same as the one used in the `vid_verification_statement_tbs`.

The **VIDVerificationStatement** message value can now be installed on a device using the **SetVIDVerificationStatement** command of the Operational Credentials Cluster.

The construction of this payload is illustrated below:




Figure 60. Construction of `VIDVerificationStatement`

6.4.10.3. Generation of the `vid_verification_statement_tbs` message

The `vid_verification_statement_tbs` is message that can always be reconstructed given the information associated with a `vendor_fabric_binding_message` and the identity of an issuer (signer) of a `VIDVerificationStatement`. It takes part in procedures in [Section 6.4.10.1, “Algorithm”](#) and [Section 6.4.10.2, “Generation of the VIDVerificationStatement message”](#).

The `vid_verification_statement_tbs` is computed as follows:

- `vid_verification_statement_tbs := statement_version || vendor_fabric_binding_message || vid_verification_signer_skid`
 - `statement_version` is the 1-octet value `0x21` indicating an original version associated with version 1.0 of the Matter [Cryptographic Primitives](#).
 - `vendor_fabric_binding_message` is the octet string binding a fabric reference to a vendor, as described in [Section 6.4.10.1, “Algorithm”](#).
 - `vid_verification_signer_skid` is the 20-octet string of the Subject Key ID extension value of the associated issuer public certificate which will sign the statement, see [Section 6.4.10.2, “Generation of the VIDVerificationStatement message”](#).

6.4.11. Security Considerations

A `NOC` is a Node’s credential to operate on a Fabric. It SHALL be protected against the following threats:

1. The Node Operational Private Key SHALL be protected from unauthorized access.
2. The Node Operational Private Key SHOULD never leave the device.
3. The `NOC` SHALL NOT contain information that may violate the user’s privacy.
4. The `NOC` chain and associated data (e.g. operational private key, Vendor ID Verification Statement, Vendor ID Verification Signing Certificate, etc) SHALL be wiped if the Node is factory reset.

6.5. Operational Certificate Encoding

6.5.1. Introduction

This section details the **Matter certificate data structure** (hereafter “Matter certificate”), a specific encoding that is sometimes used as a compact alternative to the standard X.509 certificate format [RFC 5280] for bandwidth-efficient transmission. A [Node Operational Certificate \(NOC\)](#), [Intermediate CA certificate](#), [Root CA certificate](#) or [Vendor Verification Signer Certificate \(VVSC\)](#) MAY all be encoded as a Matter certificate.

To compress the structure more efficiently than an X.509 certificate, a Matter certificate SHALL be encoded with the Matter TLV structured data interchange language [[Appendix A, Tag-length-value \(TLV\) Encoding Format](#)] instead of the ASN.1 Distinguished Encoding Rules (DER) [X.690].

This section provides a technical specification of the structure of data comprising a Matter certificate with accompanying requirements for their semantic validation, and their conversion to and

from X.509 certificates. In some cases, as noted, the limitations on the semantic interpretation of parts of a Matter certificate follow from limitations applied by [\[RFC 5280\]](#).

A certificate comprises a record of the following conceptual fields:

- Certificate Text
 - Version Number
 - Serial Number
 - Signature Algorithm ID
 - Issuer Name
 - Validity period
 - Not Before
 - Not After
 - Subject name
 - Subject Public Key Info
 - Public Key Algorithm
 - Subject Public Key
 - Issuer Unique Identifier
 - Subject Unique Identifier
 - Extensions
- Certificate Signature Algorithm
- Certificate Signature

6.5.1.1. ASN.1 Object Identifiers (OID)

Several important components of X.509 certificates follow the pattern commonly used in ASN.1 data models where some types are constructed with an ASN.1 object identifier (OID) to identify each variant. For example, the cryptographic algorithm used in the digital signature is identified by its OID.

Matter certificates do not use ASN.1 OIDs. Instead, each valid ASN.1 OID SHALL be mapped to a Matter TLV tag within its reference category. Each reference category defines the context of the Matter tag, and tag values are assigned to the reference categories according to the type of fields where they can appear in X.509 certificates.

6.5.2. Matter certificate

A Matter certificate encodes a subset of the object identifiers (OIDs) specified in X.509. Only some attribute types for relative distinguished names are valid, only certain cryptographic algorithms (corresponding to the algorithms as defined in [Chapter 3, Cryptographic Primitives](#)) are used, and only a limited set of extensions are used. Therefore, every Matter certificate can be represented as a corresponding X.509 certificate. However, the converse is not true; not every X.509 certificate can be represented as a Matter certificate.

The signature included in a Matter certificate is the **signatureValue** of the corresponding X.509 certificate, not a signature of the preceding Matter TLV data in the Matter certificate structure. Accord-

ingly, validating the signature in a Matter certificate entails its logical conversion to the corresponding X.509 certificate to recover the original **tbsCertificate** of the basic syntax signed by the Certificate Authority (CA).

```
matter-certificate [anonymous] => STRUCTURE [tag-order]
{
    serial-num [1]      : OCTET STRING [ length 0..20 ],
    sig-algo [2]        : signature-algorithm,
    issuer [3]          : LIST [ length 1.. ] OF dn-attribute,
    not-before [4]      : UNSIGNED INTEGER [ range 32-bits ],
    not-after [5]       : UNSIGNED INTEGER [ range 32-bits ],
    subject [6]         : LIST [ length 1.. ] OF dn-attribute,
    pub-key-algo [7]    : public-key-algorithm,
    ec-curve-id [8]     : elliptic-curve-id,
    ec-pub-key [9]      : OCTET STRING,
    extensions [10]     : LIST [ length 1.. ] OF extension,
    signature [11]      : ec-signature,
}
```

6.5.3. Version Number

Matter certificates SHALL only support version X.509 v3. This field is not encoded in the Matter certificate structure.

6.5.4. Serial Number

The context-specific tag **serial-num [1]** SHALL be used to identify the serial number field in the Matter certificate structure.

A Matter certificate follows the same limitation on admissible serial numbers as in [\[RFC 5280\]](#), i.e., that implementations SHALL admit serial numbers up to 20 octets in length, and certificate authorities SHALL NOT use serial numbers longer than 20 octets in length.

6.5.5. Signature Algorithm

Like an X.509 certificate, a Matter certificate SHALL include a digital signature in its signature component. The signature algorithm component of a Matter certificate specifies the cryptographic algorithm used for composing and validating the signature embedded in the signature component of the certificate. The signature algorithm SHALL match the algorithm in [Section 3.5.3, “Signature and verification”](#).

The context-specific tag **sig-algo [2]** SHALL be used to identify the signature algorithm field in the Matter certificate structure.

```
signature-algorithm => UNSIGNED INTEGER [ range 8-bits ]
```

```
{
    ecdsa-with-sha256    = 1,
}
```

The following values SHALL be defined for **signature-algorithm**:

Table 86. Signature Algorithm Object Identifiers

Value	ASN.1 OID
1	iso(1) member-body(2) us(840) ansi-x962(10045) signatures(4) ecdsa-with-SHA2(3) ecdsa-with-SHA256(2)

6.5.6. Issuer and Subject

The context-specific tags **issuer** [3] and **subject** [6] SHALL be used to identify the issuer and the subject DN fields in the Matter certificate structure. The entries in the lists SHALL be Distinguished Names (DNs), which are described in [Section 6.5.6.1, “X.501 Distinguished Names”](#).

6.5.6.1. X.501 Distinguished Names

The Issuer Name and Subject Name components of an X.509 certificate contain DN as defined in [\[RFC 5280\]](#). The ASN.1 format of a DN is a sequence of Relative Distinguished Names (RDNs). Two distinguished names DN1 and DN2 match if they have the same number of RDNs, for each RDN in DN1 there is a matching RDN in DN2, and the matching RDNs appear in the same order in both DNs.

The RDN in an X.509 certificate may be encoded as a set of one or more DN attributes, although in practice it is usually a single DN attribute. The RDN in a Matter certificate SHALL be always a single DN attribute. Two relative distinguished names RDN1 and RDN2 match if the attribute in RDN1 matches the attribute in RDN2.

```
dn-attribute => CHOICE OF
{
    // Standard and Matter-specific DN attributes.
    // Of these, all are encoded as UTF8String except domain-component,
    // which is encoded as IA5String in X.509 form.
    common-name [1]           : STRING,
    surname [2]               : STRING,
    serial-num [3]            : STRING,
    country-name [4]          : STRING,
    locality-name [5]         : STRING,
    state-or-province-name [6] : STRING,
    org-name [7]              : STRING,
    org-unit-name [8]         : STRING,
    title [9]                 : STRING,
    name [10]                 : STRING,
    given-name [11]          : STRING,
```



```

initials [12]           : STRING,
gen-qualifier [13]      : STRING,
dn-qualifier [14]       : STRING,
pseudonym [15]          : STRING,
domain-component [16]   : STRING,
matter-node-id [17]      : UNSIGNED INTEGER,
matter-firmware-signing-id [18] : UNSIGNED INTEGER,
matter-icac-id [19]      : UNSIGNED INTEGER,
matter-rcac-id [20]      : UNSIGNED INTEGER,
matter-fabric-id [21]    : UNSIGNED INTEGER,
matter-noc-cat [22]      : UNSIGNED INTEGER,
matter-vvs-id [23]       : UNSIGNED INTEGER,

// Standard DN attributes when encoded as PrintableString in X.509 form
// NOTE: The tags for these SHALL be the base tags + 0x80.
common-name-ps [129]     : STRING,
surname-ps [130]         : STRING,
serial-num-ps [131]      : STRING,
country-name-ps [132]    : STRING,
locality-name-ps [133]   : STRING,
state-or-province-name-ps [134] : STRING,
org-name-ps [135]        : STRING,
org-unit-name-ps [136]   : STRING,
title-ps [137]           : STRING,
name-ps [138]            : STRING,
given-name-ps [139]      : STRING,
initials-ps [140]        : STRING,
gen-qualifier-ps [141]   : STRING,
dn-qualifier-ps [142]    : STRING,
pseudonym-ps [143]       : STRING,
}

```

Table 87, “Standard DN Object Identifiers” lists the context-specific tags defined for the standard DN attribute types used in Matter that can be encoded in X.509 certificates as either **UTF8String** or as **PrintableString** format. In Matter certificates, the context-specific tag is logically-ORed with **0x80** (and its name given a corresponding **-ps** suffix) to indicate that the corresponding X.509 encoding of the attribute uses the **PrintableString** format instead of **UTF8String**.

Table 87. Standard DN Object Identifiers

Tag base	Matter name base	ASN.1 OID
1	common-name	joint_iso_ccitt(2) ds(5) attributeType(4) commonName(3)
2	surname	joint_iso_ccitt(2) ds(5) attributeType(4) surname(4)

Tag base	Matter name base	ASN.1 OID
3	serial-num	joint_iso_ccitt(2) ds(5) attributeType(4) serialNumber(5)
4	country-name	joint_iso_ccitt(2) ds(5) attributeType(4) countryName(6)
5	locality-name	joint_iso_ccitt(2) ds(5) attributeType(4) localityName(7)
6	state-or-province-name	joint_iso_ccitt(2) ds(5) attributeType(4) stateOrProvinceName(8)
7	org-name	joint_iso_ccitt(2) ds(5) attributeType(4) organizationName(10)
8	org-unit-name	joint_iso_ccitt(2) ds(5) attributeType(4) organizationalUnit-Name(11)
9	title	joint_iso_ccitt(2) ds(5) attributeType(4) title(12)
10	name	joint_iso_ccitt(2) ds(5) attributeType(4) name(41)
11	given-name	joint_iso_ccitt(2) ds(5) attributeType(4) givenName(42)
12	initials	joint_iso_ccitt(2) ds(5) attributeType(4) initials(43)
13	gen-qualifier	joint_iso_ccitt(2) ds(5) attributeType(4) generationQualifier(44)
14	dn-qualifier	joint_iso_ccitt(2) ds(5) attributeType(4) dnQualifier(46)
15	pseudonym	joint_iso_ccitt(2) ds(5) attributeType(4) pseudonym(65)

Table 88, “Standard DN Domain Component Object Identifier” lists the context-specific tag defined for the standard DN attribute type used in Matter that is encoded in X.509 certificates as **IA5String**.

Table 88. Standard DN Domain Component Object Identifier

Tag	Matter name	ASN.1 OID
16	domain-component	itu_t(0) data(9) pss(2342) ucl(19200300) pilot(100) pilotAttribute-Type(1) domainComponent(25)

In addition to the standard DN attribute types, there are Matter-specific DN attribute types under the [1.3.6.1.4.1.1.37244](#) private arc. See [Section 6.1.1](#), “Encoding of Matter-specific RDNs” for constraints and examples related to usage of Matter-specific DN attribute types.

6.5.6.2. Matter Certificate Types

The Matter-specific DN attribute types convey information about Matter-specific certificate types as listed in [Table 89](#), “Matter Certificate Types”.

Table 89. Matter Certificate Types

Matter name	Description
matter-node-id	Certifies the identity of a Matter Node Operational Certificate (NOC) .
matter-firmware-signing-id	Certifies the identity of a firmware signing certificate.
matter-icac-id	Certifies the identity of a Matter Intermediate CA Certificate (ICAC) .
matter-rcac-id	Certifies the identity of a Matter Root CA Certificate (RCAC) .

Matter name	Description
matter-vvs-id	Certifies the identity of a Vendor Verification Signer .

The value of **matter-icac-id**, **matter-rcac-id** and **matter-vvs-id** DN attribute types MAY be any 64-bit identifier desired by the certificate's issuer. Apart from marking what type of certificates are involved, they MAY be used for debugging purposes to determine the specific CA in use, for example if different production tiers or regions are used.

6.5.6.3. Matter DN Encoding Rules

The rules that SHALL be followed for Matter-specific attribute types when encoding the subject DN are:

- For a [Matter Node Operational Certificate \(NOC\)](#):
 - The subject DN SHALL encode exactly one **matter-node-id** attribute.
 - The **matter-node-id** attribute's value SHALL be in the Operational Node ID range (0x0000_0000_0000_0001 to 0xFFFF_FFEF_FFFF_FFFF), see [Table 4, “Node Identifier Allocations”](#).
 - The subject DN SHALL encode exactly one **matter-fabric-id** attribute.
 - The **matter-fabric-id** attribute's value SHALL NOT be 0 (see [Section 2.5.1, “Fabric References and Fabric Identifier”](#)).
 - The subject DN SHALL NOT encode any **matter-icac-id** attribute.
 - The subject DN SHALL NOT encode any **matter-rcac-id** attribute.
 - The subject DN MAY encode at most three **matter-noc-cat** attributes.
 - Each **matter-noc-cat** attribute present, if any, SHALL encode a different [CASE Authenticated Tag identifier](#) (upper 16 bits of value) than is used by other **matter-noc-cat** attributes (CATs).
 - The subject DN SHALL NOT encode any **matter-vvs-id** attribute.
- For a [Matter ICA Certificate \(ICAC\)](#):
 - The subject DN SHALL NOT encode any **matter-node-id** attribute.
 - The subject DN MAY encode at most one **matter-fabric-id** attribute.
 - If present, the **matter-fabric-id** attribute's value SHALL NOT be 0 (see [Section 2.5.1, “Fabric References and Fabric Identifier”](#)).
 - The subject DN SHALL encode exactly one **matter-icac-id** attribute.
 - The subject DN SHALL NOT encode any **matter-rcac-id** attribute.
 - The subject DN SHALL NOT encode any **matter-noc-cat** attribute.
 - The subject DN SHALL NOT encode any **matter-vvs-id** attribute.
- For a [Matter Root CA Certificate \(RCAC\)](#):
 - The subject DN SHALL NOT encode any **matter-node-id** attribute.
 - The subject DN MAY encode at most one **matter-fabric-id** attribute.

- If present, the **matter-fabric-id** attribute's value SHALL NOT be 0 (see [Section 2.5.1, “Fabric References and Fabric Identifier”](#)).
- The subject DN SHALL NOT encode any **matter-icac-id** attribute.
- The subject DN SHALL encode exactly one **matter-rcac-id** attribute.
- The subject DN SHALL NOT encode any **matter-noc-cat** attribute.
- The subject DN SHALL NOT encode any **matter-vvs-id** attribute.
- For a [Vendor Verification Signer Certificate \(VVSC\)](#):
 - The subject DN SHALL encode exactly one **matter-vvs-id** attribute.
 - The subject DN SHALL NOT encode any **matter-node-id** attribute.
 - The subject DN SHALL NOT encode any **matter-fabric-id** attribute.
 - The subject DN SHALL NOT encode any **matter-noc-cat** attribute.
- The attributes SHALL appear in the same order in the Matter certificate and in the corresponding X.509 certificates.
- When any **matter-fabric-id** attributes are present in either the [Matter Root CA Certificate](#) or the [Matter ICA Certificate](#), the value SHALL match the one present in the [Matter Node Operational Certificate \(NOC\)](#) within the same certificate chain.
- The order of the attributes can be issuer-specific and is not enforced by Matter specifications.
- All implementations SHALL accept, parse, and handle Matter certificates with up to 5 RDNs in a single DN.
- All implementations SHALL reject Matter certificates with more than 5 RDNs in a single DN.

In addition to the above rules, the encoding constraints in [Section 6.1.1, “Encoding of Matter-specific RDNs”](#) SHALL be followed.

6.5.6.4. Matter DN Examples

In all the examples below, the long hexadecimal numbers are actually encoded as hexadecimal UTF8String in the X.509 certificates, using the scheme in [Section 6.1.1, “Encoding of Matter-specific RDNs”](#).

The following is an example of subject DN encoding for a [Matter Node Operational Certificate \(NOC\)](#). Typically, it is a list of two RDN attributes:

```
subject = [[
    matter-node-id    = 0x0102030405060708,
    matter-fabric-id = 0xFAB000000000001D
]]
```

In addition to the mandatory attributes, it may also encode other supported RDN attributes such as **common-name** and [CASE Authenticated Tags](#) as presented below:

```
subject = [[
    common-name      = "NOC Example",
    matter-node-id    = 0x0102030405060708,
    matter-fabric-id = 0xFAB0000000000001D,
    matter-noc-cat    = 0xABCD0002
]]
```

The following subject DN example illustrates that multiple RDN attributes of the same type can be encoded. The specific order of attributes is not enforced. Note that number of RDN attributes in the **subject** field SHALL NOT exceed five:

```
subject = [[
    matter-noc-cat    = 0xABCD0004,
    matter-node-id    = 0x0102030405060708,
    matter-noc-cat    = 0xABCE0018,
    matter-fabric-id = 0xFAB0000000000001D,
    matter-noc-cat    = 0xABCF0002
]]
```

The following example illustrates an illegal subject DN due to the presence of the same **CASE Authenticated Tag** value with two different version numbers.

```
subject = [[
    matter-node-id    = 0x0102030405060708,
    matter-fabric-id = 0xFAB0000000000001D,
    matter-noc-cat    = 0xABCD0004,      # <-- Value 0xABCD, Version 0x0004
    matter-noc-cat    = 0xABCD0002,      # <-- Value 0xABCD, Version 0x0002
]]
```

The following is an example of subject DN encoding for a **Matter Root CA certificate**. In this case, the Matter Root CA certificate is not associated with a specific Matter fabric:

```
subject = [[
    matter-rcac-id = 0xCA00000000000001D
]]
```

The following is another example of subject DN encoding for a **Matter Root CA certificate**. In this case, the Matter Root CA certificate is associated with a specific Matter fabric. This DN also encodes an issuer-specific **common-name** RDN attribute:

```
subject = [[
    matter-rcac-id    = 0xCA00000000000001D,
```

```

matter-fabric-id = 0xFAB000000000001D,
common-name      = "ROOT CA HOME 3"
]]

```

The following is an example of subject DN encoding of a [Vendor Verification Signer certificate](#).

```

subject = [[
  common-name = "Acme Vendor Verification",
  matter-vvs-id = 0xEFFF001199887766
]]

```

6.5.7. Validity

The context-specific tags **not-before** [4] and **not-after** [5] SHALL be used to identify the not-before and not-after fields in the Matter certificate structure, which indicate the period of validity for the certificate. These two fields SHALL be encoded as unsigned integers. The value of these fields SHALL be encoded as a UTC time of type **epoch-s** (Epoch Time in Seconds).

Special value 0, when encoded in the **not-after** field, corresponds to the X.509/[RFC 5280](#) defined special time value 99991231235959Z meaning no well-defined expiration date.

6.5.8. Public Key Algorithm

The context-specific tag **pub-key-algo** [7] SHALL be used to identify the public key algorithm field in the Matter certificate structure.

```

public-key-algorithm => UNSIGNED INTEGER [ range 8-bits ]
{
  ec-pub-key = 1,
}

```

The following values SHALL be defined for **public-key-algorithm**:

Table 90. Public Key Algorithm Object Identifiers

Value	ASN.1 OID
1	iso(1) member-body(2) us(840) ansi-x962(10045) keyType(2) ecPublicKey(1)

6.5.9. EC Curve Identifier

The context-specific tag **ec-curve-id** [8] SHALL be used to identify the elliptic curve field in the Matter certificate structure.

```

elliptic-curve-id => UNSIGNED INTEGER [ range 8-bits ]
{

```

```
prime256v1 = 1,
}
```

The following values SHALL be defined for **elliptic-curve-id**:

Table 91. Elliptic Curve Object Identifiers

Value	ASN.1 OID
1	iso(1) member_body(2) us(840) ansi-x962(10045) curves(3) prime(1) prime256v1(7)

6.5.10. Public Key

The context-specific tag **ec-pub-key** [9] SHALL be used to identify the elliptic curve public key material field in the Matter certificate structure. The public key SHALL be a byte string representation of an uncompressed elliptic curve point as defined in section 2.3.3 of [SEC 1](#).

6.5.11. Extensions

The context-specific tag **extensions** [10] SHALL be used to identify the extensions field in the Matter certificate structure. The **extensions** list SHALL NOT contain more than one instance of a particular extension. The following table summarizes context-specific tags defined for the certificate extension types used in Matter.

extension => CHOICE OF
{
 basic-cnstr [1] : basic-constraints,
 key-usage [2] : UNSIGNED INTEGER [range 16-bits],
 extended-key-usage [3] : ARRAY [length 1..] OF key-purpose-id,
 subject-key-id [4] : OCTET STRING [length 20],
 authority-key-id [5] : OCTET STRING [length 20],
 future-extension [6] : OCTET STRING,
}

Table 92. Extensions Object Identifiers

Tag	Matter Name	ASN.1 OID
1	basic-cnstr	joint-iso-itu-t(2) ds(5) certificateExtension(29) basic-Constraints(19)
2	key-usage	joint-iso-itu-t(2) ds(5) certificateExtension(29) keyUsage(15)
3	extended-key-usage	joint-iso-itu-t(2) ds(5) certificateExtension(29) extended-KeyUsage(37)
4	subject-key-id	joint-iso-itu-t(2) ds(5) certificateExtension(29) subjectKeyIdentifier(14)
5	authority-key-id	joint-iso-itu-t(2) ds(5) certificateExtension(29) authorityKeyIdentifier(35)

Tag	Matter Name	ASN.1 OID
6	future-extension	any valid ASN.1 OID (future extension)

These context-specific tags identify the extension entries in the **extensions** list. The type of each extension is further described in the subsections below.

6.5.11.1. Basic Constraints Extension

The basic constraints extension identifies whether the subject of the certificate is a CA and the maximum depth of valid certification paths that include this certificate.

When present, the basic constraints extension SHALL be treated as **critical** and it SHALL be marked as **critical** in the corresponding X.509 certificate. The **critical** field SHALL NOT be encoded in the Matter certificate structure.

The context-specific tag **basic-cnstr** [1] SHALL be used to identify a basic constraints extension entry in the Matter certificate **extensions** list.

```
basic-constraints => STRUCTURE [tag-order]
{
    is-ca [1] : BOOLEAN,
    path-len-constraint [2, optional] : UNSIGNED INTEGER [ range 8-bits ],
}
```

The **is-ca** field SHALL be encoded regardless of the value (**true** or **false**). The **path-len-constraint** MAY be present only when **is-ca** == **true**.

6.5.11.2. Key Usage Extension

The key usage extension defines the purpose of the key contained in the certificate.

When present, the key usage extension SHALL be treated as **critical** and it SHALL be marked as **critical** in the corresponding X.509 certificate. The **critical** field SHALL NOT be encoded in the Matter certificate structure.

The context-specific tag number **key-usage** [2] SHALL be used to identify a key usage extension entry in the Matter certificate **extensions** list.

The **key-usage** field is derived as a logical OR of all **key-usage-flag** values that apply to the corresponding public key:

```
key-usage-flag => UNSIGNED INTEGER [ range 16-bits ]
{
    digitalSignature      = 0x0001,
    nonRepudiation        = 0x0002,
    keyEncipherment       = 0x0004,
    dataEncipherment      = 0x0008,
```

```
keyAgreement      = 0x0010,
keyCertSign       = 0x0020,
CRLSign           = 0x0040,
encipherOnly      = 0x0080,
decipherOnly      = 0x0100,
}
```

6.5.11.3. Extended Key Usage Extension

The extended key usage extension indicates one or more purposes for which the certified public key may be used, in addition to or in place of the basic purposes indicated in the key usage extension.

When present, the extended key usage extension SHALL be treated as **critical** and it SHALL be marked as **critical** in the corresponding X.509 certificate. The **critical** field SHALL NOT be encoded in the Matter certificate structure.

The context-specific tag number **extended-key-usage** [3] SHALL be used to identify an extended key usage extension entry in the Matter certificate **extensions** list.

The **extended-key-usage** field SHALL be encoded as an array of **key-purpose-id** values, where each **key-purpose-id** value SHALL be encoded as 8-bit unsigned integer:

```
key-purpose-id => UNSIGNED INTEGER [ range 8-bits ]
```

The following values SHALL be defined for **key-purpose-id**:

Table 93. Key Purpose Object Identifiers

Value	ASN.1 OID
1	iso(1), organization(3), dod(6), internet(1), security(5), mechanisms(5), pkix(7), 3, server-Auth(1)
2	iso(1), organization(3), dod(6), internet(1), security(5), mechanisms(5), pkix(7), 3, clientAuth(2)
3	iso(1), organization(3), dod(6), internet(1), security(5), mechanisms(5), pkix(7), 3, code-Signing(3)
4	iso(1), organization(3), dod(6), internet(1), security(5), mechanisms(5), pkix(7), 3, email-Protection(4)
5	iso(1), organization(3), dod(6), internet(1), security(5), mechanisms(5), pkix(7), 3, time-Stamping(8)
6	iso(1), organization(3), dod(6), internet(1), security(5), mechanisms(5), pkix(7), 3, OCSP-Signing(9)

The **key-purpose-id** values in the **extended-key-usage** array SHALL be encoded in the same order as they appeared in the corresponding X.509 certificate.

6.5.11.4. Subject Key Identifier Extension

The subject key identifier extension provides a means of identifying Matter certificates that contain a particular public key.

When present, the subject key identifier extension SHALL be treated as **non-critical** and it SHALL be marked as **non-critical** in the corresponding X.509 certificate. The **critical** field SHALL NOT be encoded in the Matter certificate structure.

The context-specific tag number **subject-key-id** [4] SHALL be used to identify a subject key identifier extension entry in the Matter certificate **extensions** list.

The Subject Key Identifier field SHALL be derived from the public key using method (1) described in section 4.2.1.2 of [RFC 5280]. Thus, the **subject-key-id** SHALL be composed of the 160-bit SHA-1 hash of the certificate's subject public key value.

See [Section 6.1.2, “Key Identifier Extension Constraints”](#) for additional constraints.

6.5.11.5. Authority Key Identifier Extension

The authority key identifier extension provides a means of identifying the public key corresponding to the private key used to sign a Matter certificate.

When present, the authority key identifier extension SHALL be treated as **non-critical** and it SHALL be marked as **non-critical** in the corresponding X.509 certificate. The **critical** field SHALL NOT be encoded in the Matter certificate structure.

The context-specific tag number **authority-key-id** [5] SHALL be used to identify an authority key identifier extension entry in the Matter certificate **extensions** list.

Note that the authority key identifier extension field in an X.509 certificate may optionally include issuer and serial number fields, which are not supported by Matter certificates.

The Authority Key Identifier field SHALL be derived from the public key using method (1) described in section 4.2.1.2 of [RFC 5280]. Thus, the **authority-key-id** SHALL be composed of the 160-bit SHA-1 hash of the public key used to verify the certificate's signature.

See [Section 6.1.2, “Key Identifier Extension Constraints”](#) for additional constraints.

6.5.11.6. Future Extension

The Matter certificate is designed with extensibility in mind and this field is added to support arbitrary certificate extension in the future.

Note that implementations that do not support specific future extension will ignore it but will be able to use it for the Matter certificate signature validation. If ignored extension is marked as critical then validation of the corresponding Matter certificate SHALL fail.

The context-specific tag number **future-extension** [6] SHALL be used to identify all future extension entries in the Matter certificate **extensions** list. There MAY be more than one future extension field.

The **future-extension** field SHALL be encoded as OCTET STRING and it SHALL be an exact copy of the DER encoded extension field (including the DER encoded ASN.1 OID of the extension) in the corresponding X.509 certificate. These extension fields in a Matter certificate SHALL be encoded in the same order as they appeared in the original X.509 certificate.

6.5.12. Matter certificate Extensions Encoding Rules

The rules that SHALL be followed when encoding the Matter certificate are:

- For a **Matter Node Operational Certificate (NOC)**:
 - The basic constraints extension SHALL be encoded with **is-ca** set to **false**.
 - The key usage extension SHALL be encoded with exactly one flag: **digitalSignature**.
 - The extended key usage extension SHALL be encoded with exactly two **key-purpose-id** values: **serverAuth** and **clientAuth**.
 - The subject key identifier extension SHALL be present.
 - The authority key identifier extension SHALL be present.
- For **Matter ICA Certificate** and **Matter Root CA Certificate**:
 - The basic constraints extension SHALL be encoded with **is-ca** set to **true**.
 - The key usage extension SHALL be encoded with at least the two flags **keyCertSign** and **CRL-Sign** included, and MAY include a third flag: **digitalSignature**. No additional flags are permitted.
 - The extended key usage extension SHALL NOT be present.
 - The subject key identifier extension SHALL be present.
 - The authority key identifier extension SHALL be present.
 - For the **Matter Root CA Certificate** the authority key identifier extension SHALL be equal to the subject key identifier extension.
- For **Vendor ID Verification Signer Certificate**:
 - The basic constraints extension SHALL be encoded with **is-ca** set to **false**.
 - The key usage extension SHALL be encoded with exactly one flag: **digitalSignature**.
 - The subject key identifier extension SHALL be present.
 - The authority key identifier extension SHALL be present.
- The extensions SHALL appear in the same order in the Matter certificate and in the corresponding X.509 certificates.
- For a Matter Firmware Signing Certificate these rules SHALL be followed:
 - The basic constraints extension SHALL be encoded with **is-ca** set to **false**.
 - The key usage extension SHALL be encoded with exactly one flag: **digitalSignature**.
 - The extended key usage extension SHALL be encoded with exactly one **key-purpose-id** values: **codeSigning**.
 - The subject key identifier extension SHALL be present.

- The authority key identifier extension SHALL be present.
- The extensions SHALL appear in the same order in the Matter certificate and in the corresponding X.509 certificates.

Note that Matter doesn't specify how firmware images are signed and implementation of firmware image signing is manufacturer-specific. However, since firmware image signing is a common feature, the format for Matter TLV certificates has affordances for encoding firmware signing certificates.

6.5.13. Signature

The context-specific tag **signature** [11] SHALL be used to identify the signature field in the Matter certificate structure.

An **ec-signature** is the encoding of the signature as defined in [Section 3.5.3, “Signature and verification”](#).

```
ec-signature => OCTET STRING [ length (CRYPTO_GROUP_SIZE_BYTES * 2) ]
```

6.5.14. Invalid Matter certificates

The Matter certificate is considered invalid if it violates Matter certificate encoding rules defined in this section. The processing of invalid Matter certificate SHOULD fail and an error SHOULD be reported to the application. Here is a non-exhaustive list of errors that may invalidate the certificate:

- Matter certificate structure includes elements that are not defined in this section.
- Matter certificate elements are encoded in a wrong order.
- Matter certificate element has wrong type.
- Length of an element is outside of its valid range, for example:
 - **serial-num** field is longer than 20 octets.
 - **not-before** or **not-after** field is longer than 32-bits.
 - **subject-key-id** or **authority-key-id** field is different from 20 octets.
- Matter OID values encoded in Matter certificate are not defined in this section, for example:
 - **sig-algo** field encodes value, which is not defined in [Table 86, “Signature Algorithm Object Identifiers”](#).
 - **pub-key-algo** field encodes value, which is not defined in [Table 90, “Public Key Algorithm Object Identifiers”](#).
 - **ec-curve-id** field encodes value, which is not defined in [Table 91, “Elliptic Curve Object Identifiers”](#).
 - **key-purpose-id** field of the Extended Key Usage Extension encodes value, which is not defined in [Table 93, “Key Purpose Object Identifiers”](#).
- Invalid Matter Distinguished Names encoding for Issuer and Subject DNs. Refer to [Section](#)

6.5.6.3, “Matter DN Encoding Rules” for more details. For example:

- Node Subject DN doesn’t include Matter specific **matter-node-id** or **matter-fabric-id** attribute.
- Firmware signing Subject DN doesn’t include Matter specific **matter-firmware-signing-id** attribute.
- Intermediate CA Subject DN doesn’t include Matter specific **matter-icac-id** attribute.
- Root CA Subject DN doesn’t include Matter specific **matter-rcac-id** attribute.
- Certificate Subject DN encodes **matter-node-id** and **matter-rcac-id**, which contradict each other.
- Multiple **matter-cat-id** with the same identifier value and different version numbers, or any **matter-cat-id** with a version number of 0.
- CA certificate doesn’t have Basic Constraints Extension **is-ca** field set to 'true'.
- **key-usage** field of the Key Usage Extension has undefined flags.
- Certificate extension that SHALL be marked as **critical** is marked as **non-critical** in the X.509 representation and vice versa.

6.5.15. Examples

6.5.15.1. Example of Operational Root CA Certificate (RCAC)

The same RCAC in X.509 and Matter TLV formats is presented in this section.

RCAC with Corresponding Private Key in an X.509 PEM Format

```
-----BEGIN CERTIFICATE-----
MIIBnTCCAU0gAwIBAgIIWeqmMpR/VBwwCgYIKoZIzj0EAwIwIjEgMB4GCi sGAQQB
gqJ8AQQMEENBQ0FDQUNBMDAwMDAwMDEwHhcNMjAxMDE1MTQyMzQzWhcNNDExMDE1
MTQyMzQyWjAiMSAwHgYKKwYBBAGConwBBAwQQ0FDQUNBQ0EwMDAwMDAwMTBZMBMG
ByqGSM49AgEGCCqGSM49AwEHA0IABBNT07PvHacIxJCASAF0QH1ZkM4ivE6zPppa
yyWoVgPrptzYITZmpORPWsoT63Z/r6fc3dwzQR+CowtUPdHSS6ijYzBhMA8GA1Ud
EwEB/wQFMAMBAf8wDgYDVROPAQH/BAQDAgEGMBOGA1UdDgQWBBQTr4GrNzdLLtKp
ZJsSt60kKH4VHTAfBgNVHSMEGDAWgBQTr4GrNzdLLtKpZJsSt60kKH4VHTAKBgq
hkjOPQQDAgNIADBFAiBFgWRGbI8ZWrwKu3xstaJ6g/QdN/jV0+7FIKvSoNoFCQIh
ALinwLwELjDPZNww/jNOEgAZZk5RUEkTT1eBI4RE/HUx
-----END CERTIFICATE-----

-----BEGIN EC PRIVATE KEY-----
MHcCAQEEIH1zW+/pFqHAYgL4ypiB5CZjqqaucQzsom+JnAQdXQaoAoGCCqGSM49
AwEHoUQDQgAEE10js+8dpwjEkIBIAU5AfVmQziK8TrM+mLrLJahWA+um3NghNmak
5E9ayhPrdn+vp9zd3DNBH4KjC1Q90dJLqA==
-----END EC PRIVATE KEY-----
```

RCAC Printed in an X.509 DER Format

Certificate:

Data:

Version: 3 (0x2)

Serial Number: 6479173750095827996 (0x59eaa632947f541c)

Signature Algorithm: ecdsa-with-SHA256

Issuer: 1.3.6.1.4.1.37244.1.4 = CACACACA00000001

Validity

Not Before: Oct 15 14:23:43 2020 GMT

Not After : Oct 15 14:23:42 2040 GMT

Subject: 1.3.6.1.4.1.37244.1.4 = CACACACA00000001

Subject Public Key Info:

Public Key Algorithm: id-ecPublicKey

Public-Key: (256 bit)

pub:

04:13:53:a3:b3:ef:1d:a7:08:c4:90:80:48:01:4e:

40:7d:59:90:ce:22:bc:4e:b3:3e:9a:5a:cb:25:a8:

56:03:eb:a6:dc:d8:21:36:66:a4:e4:4f:5a:ca:13:

eb:76:7f:af:a7:dc:dd:dc:33:41:1f:82:a3:0b:54:

3d:d1:d2:4b:a8

ASN1 OID: prime256v1

NIST CURVE: P-256

X509v3 extensions:

X509v3 Basic Constraints: critical

CA:TRUE

X509v3 Key Usage: critical

Certificate Sign, CRL Sign

X509v3 Subject Key Identifier:

13:AF:81:AB:37:37:4B:2E:D2:A9:64:9B:12:B7:A3:A4:28:7E:15:1D

X509v3 Authority Key Identifier:

keyid:13:AF:81:AB:37:37:4B:2E:D2:A9:64:9B:12:B7:A3:A4:28:7E:15:1D

Signature Algorithm: ecdsa-with-SHA256

30:45:02:20:45:81:64:46:6c:8f:19:5a:bc:0a:bb:7c:6c:b5:

a2:7a:83:f4:1d:37:f8:d5:3b:ee:c5:20:ab:d2:a0:da:05:09:

02:21:00:b8:a7:c2:5c:04:2e:30:cf:64:dc:30:fe:33:4e:12:

00:19:66:4e:51:50:49:13:4f:57:81:23:84:44:fc:75:31

RCAC in Matter TLV Format

```

15 30 01 08 59 ea a6 32 94 7f 54 1c 24 02 01 37 03 27 14 01 00 00 00 ca
ca ca ca 18 26 04 ef 17 1b 27 26 05 6e b5 b9 4c 37 06 27 14 01 00 00 00

```



```

ca ca ca ca 18 24 07 01 24 08 01 30 09 41 04 13 53 a3 b3 ef 1d a7 08 c4
90 80 48 01 4e 40 7d 59 90 ce 22 bc 4e b3 3e 9a 5a cb 25 a8 56 03 eb a6
dc d8 21 36 66 a4 e4 4f 5a ca 13 eb 76 7f af a7 dc dd dc 33 41 1f 82 a3
0b 54 3d d1 d2 4b a8 37 0a 35 01 29 01 18 24 02 60 30 04 14 13 af 81 ab
37 37 4b 2e d2 a9 64 9b 12 b7 a3 a4 28 7e 15 1d 30 05 14 13 af 81 ab 37
37 4b 2e d2 a9 64 9b 12 b7 a3 a4 28 7e 15 1d 18 30 0b 40 45 81 64 46 6c
8f 19 5a bc 0a bb 7c 6c b5 a2 7a 83 f4 1d 37 f8 d5 3b ee c5 20 ab d2 a0
da 05 09 b8 a7 c2 5c 04 2e 30 cf 64 dc 30 fe 33 4e 12 00 19 66 4e 51 50
49 13 4f 57 81 23 84 44 fc 75 31 18

```

RCAC Printed in Matter TLV Schema Format

```

matter-certificate = {
  serial-num       = 59 EA A6 32 94 7F 54 1C,
  sig-algo         = 0x01U,
  issuer           = [[ matter-rcac-id = 0xCACACACA00000001U ]],
  not-before       = 0x271B17EFU,
  not-after        = 0x4CB9B56EU,
  subject          = [[ matter-rcac-id = 0xCACACACA00000001U ]],
  pub-key-algo     = 0x01U,
  ec-curve-id      = 0x01U,
  ec-pub-key       = 04 13 53 A3 B3 EF 1D A7 08 C4 90 80 48 01 4E 40
                    7D 59 90 CE 22 BC 4E B3 3E 9A 5A CB 25 A8 56 03
                    EB A6 DC D8 21 36 66 A4 E4 4F 5A CA 13 EB 76 7F
                    AF A7 DC DD DC 33 41 1F 82 A3 0B 54 3D D1 D2 4B
                    A8,
  extensions       = [[
    basic-constraints = {
      is-ca = true
    },
    key-usage         = 0x60U,
    subject-key-id     = 13 AF 81 AB 37 37 4B 2E D2 A9 64 9B 12 B7 A3 A4
                      28 7E 15 1D,
    authority-key-id   = 13 AF 81 AB 37 37 4B 2E D2 A9 64 9B 12 B7 A3 A4
                      28 7E 15 1D,
  ]],
  signature         = 45 81 64 46 6C 8F 19 5A BC 0A BB 7C 6C B5 A2 7A
                    83 F4 1D 37 F8 D5 3B EE C5 20 AB D2 A0 DA 05 09
                    B8 A7 C2 5C 04 2E 30 CF 64 DC 30 FE 33 4E 12 00
                    19 66 4E 51 50 49 13 4F 57 81 23 84 44 FC 75 31
}

```

6.5.15.2. Example of Operational Intermediate CA Certificate (ICAC)

The same ICAC in X.509 and Matter TLV formats is presented in this section. An issuer of this ICAC is RCAC from previous section.

ICAC with Corresponding Private Key in an X.509 PEM Format

```
-----BEGIN CERTIFICATE-----
MIIBnTCCAUOgAwIBAgIILbREhVZBrT8wCgYIKoZIzj0EAwIwIjEgMB4GCisGAQQB
gqJ8AQQMEENBQ0FDQUNBMDAwMDAwMDEwHhcNMjAxMDE1MTQyMzQzWhcNNDExMDE1
MTQyMzQyWjAiMSAwHgYKKwYBBAGConwBAwwQQ0FDQUNBQ0EwMDAwMDAwMzBZMBMG
ByqGSM49AgEGCCqGSM49AwEHA0IABMXQhhu4+QxAXBIXTkxevuqTn3J3S8wzI54v
Wfb0avjcfUaCoOPMxkbm3ynqhr9WKucgqJgzfTg/MsCgnkFgGeqjYzBhMA8GA1Ud
EwEB/wQFMAMBAf8wDgYDVR0PAQH/BAQDAgEGMB0GA1UdDgQWBRTUtcFnpwVpQiQ
aGKGSAGinx9B0zAfBgNVHSMEGDAwgbQTr4GrNzdLLtKpZJsSt60kKH4VHTAKBgqq
hkjOPQQDAgNIADBFAiEAhBoG1Dten+zSToexJE61HGos8g2bXmugfxHmAC9+DKMC
IE4ypgLDYJOAktNIvb0ZihFGRr1BzxA3g2Qa4l4/I/Om
-----END CERTIFICATE-----

-----BEGIN EC PRIVATE KEY-----
MHcCAQEEIIBGE09zwrSBtsQJRpu2sWB11+ZL8tSJ1KiFs15xxdUapoAoGCCqGSM49
AwEHoUQDQgAExdCGG7j5DEBcEjFOTF6+6p0fcndLzDMjni9Z9vRq+Nx9RoKg48zG
RubfKqGv1Yq5yComDN90D8ywKCeQWAZ6g==
-----END EC PRIVATE KEY-----
```

ICAC Printed in an X.509 DER Format

Certificate:

Data:

Version: 3 (0x2)

Serial Number: 3293332566983159519 (0xdb444855641aedef)

Signature Algorithm: ecdsa-with-SHA256

Issuer: 1.3.6.1.4.1.37244.1.4 = CACACACA00000001

Validity

Not Before: Oct 15 14:23:43 2020 GMT

Not After : Oct 15 14:23:42 2040 GMT

Subject: 1.3.6.1.4.1.37244.1.3 = CACACACA00000003

Subject Public Key Info:

Public Key Algorithm: id-ecPublicKey

Public-Key: (256 bit)

pub:

04:c5:d0:86:1b:b8:f9:0c:40:5c:12:31:4e:4c:5e:

be:ea:93:9f:72:77:4b:cc:33:23:9e:2f:59:f6:f4:

6a:f8:dc:7d:46:82:a0:e3:cc:c6:46:e6:df:29:ea:

```

      86:bf:56:2a:e7:20:a8:98:33:7d:38:3f:32:c0:a0:
      9e:41:60:19:ea
      ASN1 OID: prime256v1
      NIST CURVE: P-256
    X509v3 extensions:
      X509v3 Basic Constraints: critical
      CA:TRUE
      X509v3 Key Usage: critical
      Certificate Sign, CRL Sign
      X509v3 Subject Key Identifier:
        53:52:D7:05:9E:9C:15:A5:08:90:68:62:86:48:01:A2:9F:1F:41:D3
      X509v3 Authority Key Identifier:
        keyid:13:AF:81:AB:37:37:4B:2E:D2:A9:64:9B:12:B7:A3:A4:28:7E:15:1D

    Signature Algorithm: ecdsa-with-SHA256
      30:45:02:21:00:84:1a:06:d4:3b:5e:9f:ec:d2:4e:87:b1:24:
      4e:b5:1c:6a:2c:f2:0d:9b:5e:6b:a0:7f:11:e6:00:2f:7e:0c:
      a3:02:20:4e:32:a6:02:c3:60:9d:00:92:d3:48:bd:bd:19:8a:
      11:46:46:bd:41:cf:10:37:83:64:1a:e2:5e:3f:23:fd:26

```

ICAC in Matter TLV Format

```

15 30 01 08 2d b4 44 85 56 41 ae df 24 02 01 37 03 27 14 01 00 00 00 ca
ca ca ca 18 26 04 ef 17 1b 27 26 05 6e b5 b9 4c 37 06 27 13 03 00 00 00
ca ca ca ca 18 24 07 01 24 08 01 30 09 41 04 c5 d0 86 1b b8 f9 0c 40 5c
12 31 4e 4c 5e be ea 93 9f 72 77 4b cc 33 23 9e 2f 59 f6 f4 6a f8 dc 7d
46 82 a0 e3 cc c6 46 e6 df 29 ea 86 bf 56 2a e7 20 a8 98 33 7d 38 3f 32
c0 a0 9e 41 60 19 ea 37 0a 35 01 29 01 18 24 02 60 30 04 14 53 52 d7 05
9e 9c 15 a5 08 90 68 62 86 48 01 a2 9f 1f 41 d3 30 05 14 13 af 81 ab 37
37 4b 2e d2 a9 64 9b 12 b7 a3 a4 28 7e 15 1d 18 30 0b 40 84 1a 06 d4 3b
5e 9f ec d2 4e 87 b1 24 4e b5 1c 6a 2c f2 0d 9b 5e 6b a0 7f 11 e6 00 2f
7e 0c a3 4e 32 a6 02 c3 60 9d 00 92 d3 48 bd bd 19 8a 11 46 46 bd 41 cf
10 37 83 64 1a e2 5e 3f 23 fd 26 18

```

ICAC Printed in Matter TLV Schema Format

```

matter-certificate = {
  serial-num       = 2D B4 44 85 56 41 AE DF,
  sig-algo         = 0x01U,
  issuer           = [[ matter-rcac-id = 0xCACACACA00000001U ]],
  not-before       = 0x271B17EFU,
  not-after        = 0x4CB9B56EU,
  subject          = [[ matter-icac-id = 0xCACACACA00000003U ]],

```

```

pub-key-algo      = 0x01U,
ec-curve-id       = 0x01U,
ec-pub-key        = 04 C5 D0 86 1B B8 F9 0C 40 5C 12 31 4E 4C 5E BE
                   EA 93 9F 72 77 4B CC 33 23 9E 2F 59 F6 F4 6A F8
                   DC 7D 46 82 A0 E3 CC C6 46 E6 DF 29 EA 86 BF 56
                   2A E7 20 A8 98 33 7D 38 3F 32 C0 A0 9E 41 60 19
                   EA,
extensions        = [[
    basic-constraints = {
        is-ca = true
    },
    key-usage          = 0x60U,
    subject-key-id      = 53 52 D7 05 9E 9C 15 A5 08 90 68 62 86 48 01 A2
                       9F 1F 41 D3,
    authority-key-id    = 13 AF 81 AB 37 37 4B 2E D2 A9 64 9B 12 B7 A3 A4
                       28 7E 15 1D,
]],
signature          = 84 1A 06 D4 3B 5E 9F EC D2 4E 87 B1 24 4E B5 1C
                   6A 2C F2 0D 9B 5E 6B A0 7F 11 E6 00 2F 7E 0C A3
                   4E 32 A6 02 C3 60 9D 00 92 D3 48 BD BD 19 8A 11
                   46 46 BD 41 CF 10 37 83 64 1A E2 5E 3F 23 FD 26
}

```

6.5.15.3. Example of Node Operational Certificate (NOC)

The same NOC in X.509 and Matter TLV formats is presented in this section. An issuer of this NOC is ICAC from previous section.

NOC with Corresponding Private Key in an X.509 PEM Format

```

-----BEGIN CERTIFICATE-----
MIIB4DCCAYagAwIBAgIIPvz/FwK5oXowCgYIKoZIzj0EAwIwIjEgMB4GCi sGAQQB
gqJ8AQMMEENBQ0FDQUNBMDAwMDAwMDMwHhcNMjAxMDE1MTQyMzQzWhcNNDAxMDE1
MTQyMzQyWjBEMSAwHgYKKwYBBAGConwBAQwQREVERURFREUwMDAxMDAwMTEgMB4G
Ci sGAQQBgqJ8AQUMEEZBQjAwMDAwMDAwMDAwMUQwWTATBgcqhkjOPQIBBggqhkjO
PQMBBwNCAASaKi Fvs53WtvohG4NciePmr7ZsFPdYMZVPn/T3o/ARLIoNjq8pxlMp
TUju4HCKAyzKOTk8OntG8YGuoHj+rYODo4GDMIGAMAwGA1UdEwEB/wQCMAAwDgYD
VR0PAQH/BAQDAgeAMCAGA1UdJQEB/wQWMBQGCCsGAQUFBwMCBggrBgEFBQcDATA d
BgNVHQ4EFgQUn1Wia35DA+YIig+kTv5T0+14qYWEwHwYDVROjBBgwFoAUU1LXBZ6c
FaUIkGhihkgBop8fQdMwCgYIKoZIzj0EAwIDSAAwRQIgeVXCAMLS6TVkSumMi/f
KPie3+WvnA5XK9ihSq7TRICIQC4PKF8ewX7Fkt315xSLhMxa8/ReJXksqTyQEUY
FzJxWQ==
-----END CERTIFICATE-----

```

```

-----BEGIN EC PRIVATE KEY-----
MHcCAQEEIKVls/ooq01qdPtvD/ik00DZ4a6Y8h36HwpZp0oCGhYnoAoGCCqGSM49
AwEHoUQDQgAEmiohb70d1rb6IRuDXInj5q+2bBT3WDGVT5/096PwESyKDY6vKcZT
KU1I7uBwigMsyjk5PDp7RvGBrqb4/q2Dgw==
-----END EC PRIVATE KEY-----

```

NOC Printed in an X.509 DER Format

Certificate:

Data:

Version: 3 (0x2)

Serial Number: 4538782998777667962 (0x3efcff1702b9a17a)

Signature Algorithm: ecdsa-with-SHA256

Issuer: 1.3.6.1.4.1.37244.1.3 = CACACACA00000003

Validity

Not Before: Oct 15 14:23:43 2020 GMT

Not After : Oct 15 14:23:42 2040 GMT

Subject: 1.3.6.1.4.1.37244.1.1 = DEDEDEDE00010001, 1.3.6.1.4.1.37244.1.5 = FAB000000000001D

Subject Public Key Info:

Public Key Algorithm: id-ecPublicKey

Public-Key: (256 bit)

pub:

04:9a:2a:21:6f:b3:9d:d6:b6:fa:21:1b:83:5c:89:

e3:e6:af:b6:6c:14:f7:58:31:95:4f:9f:f4:f7:a3:

f0:11:2c:8a:0d:8e:af:29:c6:53:29:4d:48:ee:e0:

70:8a:03:2c:ca:39:39:3c:3a:7b:46:f1:81:ae:a0:

78:fe:ad:83:83

ASN1 OID: prime256v1

NIST CURVE: P-256

X509v3 extensions:

X509v3 Basic Constraints: critical

CA:FALSE

X509v3 Key Usage: critical

Digital Signature

X509v3 Extended Key Usage: critical

TLS Web Client Authentication, TLS Web Server Authentication

X509v3 Subject Key Identifier:

9F:55:A2:6B:7E:43:03:E6:08:83:E9:13:BF:94:F4:FB:5E:2A:61:61

X509v3 Authority Key Identifier:

keyid:53:52:D7:05:9E:9C:15:A5:08:90:68:62:86:48:01:A2:9F:1F:41:D3

Signature Algorithm: ecdsa-with-SHA256

```

30:45:02:20:79:55:c2:02:63:0b:4b:a4:d5:91:25:26:32:2f:
df:28:f8:9e:df:e5:af:9c:0e:57:2b:d8:a1:4a:aa:bb:4d:12:
02:21:00:b8:3c:a1:7c:7b:05:fb:16:4b:77:d7:9c:52:96:13:
31:6b:cf:d1:78:95:e4:b2:a4:f2:40:4b:98:17:32:71:59

```

NOC in Matter TLV Format

```

15 30 01 08 3e fc ff 17 02 b9 a1 7a 24 02 01 37 03 27 13 03 00 00 00 ca
ca ca ca 18 26 04 ef 17 1b 27 26 05 6e b5 b9 4c 37 06 27 11 01 00 01 00
de de de de 27 15 1d 00 00 00 00 00 b0 fa 18 24 07 01 24 08 01 30 09 41
04 9a 2a 21 6f b3 9d d6 b6 fa 21 1b 83 5c 89 e3 e6 af b6 6c 14 f7 58 31
95 4f 9f f4 f7 a3 f0 11 2c 8a 0d 8e af 29 c6 53 29 4d 48 ee e0 70 8a 03
2c ca 39 39 3c 3a 7b 46 f1 81 ae a0 78 fe ad 83 83 37 0a 35 01 28 01 18
24 02 01 36 03 04 02 04 01 18 30 04 14 9f 55 a2 6b 7e 43 03 e6 08 83 e9
13 bf 94 f4 fb 5e 2a 61 61 30 05 14 53 52 d7 05 9e 9c 15 a5 08 90 68 62
86 48 01 a2 9f 1f 41 d3 18 30 0b 40 79 55 c2 02 63 0b 4b a4 d5 91 25 26
32 2f df 28 f8 9e df e5 af 9c 0e 57 2b d8 a1 4a aa bb 4d 12 b8 3c a1 7c
7b 05 fb 16 4b 77 d7 9c 52 96 13 31 6b cf d1 78 95 e4 b2 a4 f2 40 4b 98
17 32 71 59 18

```

NOC Printed in Matter TLV Schema Format

```

matter-certificate = {
    serial-num      = 3E FC FF 17 02 B9 A1 7A,
    sig-algo        = 0x01U,
    issuer          = [[ matter-icac-id = 0xCACACACA00000003U ]],
    not-before      = 0x271B17EFU,
    not-after       = 0x4CB9B56EU,
    subject         = [[ matter-node-id = 0xDEDEDEDE00010001U,
                        matter-fabric-id = 0xFAB0000000000001DU ]],
    pub-key-algo    = 0x01U,
    ec-curve-id     = 0x01U,
    ec-pub-key      = 04 9A 2A 21 6F B3 9D D6 B6 FA 21 1B 83 5C 89 E3
                      E6 AF B6 6C 14 F7 58 31 95 4F 9F F4 F7 A3 F0 11
                      2C 8A 0D 8E AF 29 C6 53 29 4D 48 EE E0 70 8A 03
                      2C CA 39 39 3C 3A 7B 46 F1 81 AE A0 78 FE AD 83
                      83,
    extensions      = [[
        basic-constraints = {
            is-ca = false
        },
        key-usage          = 0x01U,
        extended-key-usage = [ 0x02U, 0x01U ],
    ]],
}

```

```

    subject-key-id    = 9F 55 A2 6B 7E 43 03 E6 08 83 E9 13 BF 94 F4 FB
                        5E 2A 61 61,
    authority-key-id  = 53 52 D7 05 9E 9C 15 A5 08 90 68 62 86 48 01 A2
                        9F 1F 41 D3,

  ]],
  signature          = 79 55 C2 02 63 0B 4B A4 D5 91 25 26 32 2F DF 28
                        F8 9E DF E5 AF 9C 0E 57 2B D8 A1 4A AA BB 4D 12
                        B8 3C A1 7C 7B 05 FB 16 4B 77 D7 9C 52 96 13 31
                        6B CF D1 78 95 E4 B2 A4 F2 40 4B 98 17 32 71 59
}

```

6.6. Access Control

6.6.1. Scope and Purpose

This section specifies the features related to controlling access to a Node's Endpoint Clusters ("Targets" hereafter) from other Nodes. The overall features are collectively named "Access Control" hereafter.

The Access Control features aim to ensure that only authorized Nodes are permitted access to given application-layer functionality exposed by the Data Model, through the Interaction Model. Access Control is the fundamental link between the Secure Channel and the Interaction Model.

In order to implement a policy of Access Control, Administrators on the fabric create and maintain a consistent distributed configuration of Access Control Lists (ACLs) across all Nodes. Each Node has an ACL containing Access Control Entries which codify the policy. The Access Control Cluster exposes a data model view of a Node's ACL which enables its maintenance. In addition to the ACL, a per-fabric [Access Restriction List](#) (ARL), which is set by the device, MAY exist. The ARL contains Access Restriction Entries, which identify the attributes, commands and events on specific endpoint clusters which are not accessible on a given fabric. The Access Control Cluster exposes a data model view of a Node's ARL, which enables its discovery.

6.6.2. Access Control Model

The Access Control system is rule-based with no implicit access permitted by default. Access to a Node's Targets is denied unless the Access Control system grants the required privilege level to a given Subject to interact with given Targets on that Node. Initial Access Control privileges are [bootstrapped](#) during the commissioning phase, and maintained thereafter by Administrators.

All access privileges, from the [AccessControlEntryPrivilegeEnum](#) enumeration, SHALL be granted by Access Control. Thus, the Initiator ("Subject" hereafter) of any Interaction Model action SHALL NOT succeed in executing a given action on a Node's Target unless that Node's Access Control system explicitly grants the required privilege to that Subject for that particular action.

The Access Control system grants privileges by checking and verifying all attempted access against rules explicitly codified in Access Control Entries within the Node's Access Control List and its Access Restriction List. Additionally, Access Control implicitly grants administrative access privi-

leges to an Administrative Subject during a Node's [commissioning phase](#), with the exception of commands, attributes and events identified by the Node's Commissioning Access Restriction List.

Access Control Entries contain:

- A [FabricIndex](#) scoping the entry to the [Associated Fabric](#).
- A Privilege level granted by the entry (see [Section 9.10.5.2, "AccessControlEntryPrivilegeEnum"](#))
 - *View*: reading or subscribing to data
 - *Operate*: writing operational data and invoking operational commands
 - *Manage*: writing configuration data and invoking configuration commands (for example, [Binding](#) and [Group](#) Clusters access)
 - *Administer*: writing administrative data and invoking administrative commands (for example, [Access Control](#) and [Commissioning](#) Clusters access)
- A list of target Clusters to which the entry applies
- A list of source Subjects to which the entry applies
- An Authentication Mode that describes the type of secure channel authentication method to which the entry's subjects apply

Access Restriction Entries contain:

- A [FabricIndex](#) scoping the entry to the [Associated Fabric](#) (entries applied during commissioning do not include a Fabric Index)
- An endpoint to which the entry applies
- A cluster to which the entry applies
- A list of restricted data model elements (attributes, commands and events) for the cluster

6.6.2.1. Access Control Subjects

The meaning of a "Subject" is primarily that of describing the source for an action, using a given authentication method provided by the Secure Channel architecture. A subject SHALL be one of:

- A passcode, identified by a Passcode ID, authenticated locally by a [PASE](#) session, during the commissioning phase.
- Note that any Passcode ID other than 0, which is the default commissioning passcode, is reserved for future use.
- Furthermore, ACL entries with a PASE authentication mode SHALL NOT be explicitly added to the Access Control List, since there is an invisible implicit administrative entry (see [Section 6.6.2.9, "Bootstrapping of the Access Control Cluster"](#)) always equivalently present on the Commissionee (but not the Commissioner) during PASE sessions.
- A source node, authenticated by a [CASE](#) session using its Operational Certificate, during the operational phase. The source node can be identified by its [Node ID](#) and/or by [CASE Authenticated Tags](#).
- A destination group, identified by a destination Group ID in the message, authenticated by an

Operational Group Key from the Group Key Management Cluster, during the operational phase.

6.6.2.1.1. PASE and Group Subjects

Note that the subject is not considered to be an individual Node when the authentication is via pass-code or group symmetric key; in these cases, the administrative root of trust is conditional only upon bearing the correct passcode during session establishment, or bearing the Operational Group Key when constructing a group message.

6.6.2.1.2. Subjects identified by CASE Authenticated Tag

In contrast, a CASE Authenticated Tag (CAT) is a [special subject distinguished name](#) within the Operational Certificate shared during CASE session establishment that functions as a group-like tag. Such a tag can be applied to several Nodes, thereby facilitating management of Access Control Entries that use the same set of Nodes as subjects. Because these tags are authenticated within the CASE session context, the administrative root of trust does chain back through the individual source Node to the Fabric’s [trusted root](#). This makes CATs suitable for group-like use while maintaining secure authentication and attribution ability.

Each CAT is 32-bit and equally divided into identifier value and its corresponding version:

- The upper 16-bits are allocated to identifier value
- The lower 16-bits are allocated to the version number

The identifier value uniquely identifies a CAT.

The CAT identifier value space is allocated as described in the [CAT Identifier Space](#). This table reserves space for the use of CAT identifiers that SHOULD be used only for the purposes specified. [CAT Identifier Space](#) lists the allocation rules of CAT identifiers that SHALL be used for the [Joint Fabric](#).

Table 94. CAT Identifier Value Allocations

Range	Type
0xFFFF	Administrator identifier
0xFFFE	Anchor identifier
0xF000 to 0xFFFD	Reserved for future use
0x0000 to 0xEFFF	Generally available

The [CAT Identifier Space](#) reserves the following identifiers for special use within the context of the [Joint Fabric](#) and SHALL NOT apply to fabrics other than the [Joint Fabric](#).

- 0xFFFF SHALL represent the [Administrator CAT](#) identifier. The Administrator CAT SHALL be used to grant **Administer** privilege to a group of Administrator Nodes participating in the context of the [Joint Fabric](#). See [Section 12.2.4.1, “Administrator CAT”](#) for more details. The version number of the Administrator CAT SHALL be initially set to 0x0001.
- 0xFFFE SHALL represent the [Anchor CAT](#) identifier. The Anchor CAT SHALL be used to validate a Node as a [Joint Fabric](#) Anchor. See [Section 12.2.4.2, “Anchor CAT”](#) for more details. The version

number of the Anchor CAT SHALL be initially set to 0x0001.

The version number represents the current version of specific identifier value. An Administrator MAY increment the version number on any changes in the set of Nodes sharing the given tag. Version number is a monotonically increasing natural number in the range of 1 to 65535. A version number of 0 is invalid and SHALL NOT be used. On reaching the maximum value (65535), wrap-around is not supported and the tag identifier SHOULD be retired by an Administrator since version increase will no longer be possible.

When a CAT appears in the **Subjects** list of an Access Control Entry, it SHALL be encoded within the CASE Authenticated Tag sub-space of **Node Identifiers**, with the upper 32 bits set to 0xFFFF_FFFD. Note that this encoding cannot appear as an operational Node ID. It is merely a sub-encoding allowing the 64-bit scalars in an Access Control Entry's **Subjects** list to represent both Node IDs and CATs.

Example CASE Authenticated Tags:

- Joint Fabric Administrator Tag identifier 0xFFFF, Version 0x0001
 - CAT value: 0xFFFF_0001
 - Appears in Access Control Entry **Subjects** list as 0xFFFF_FFFD_FFFF_0001
 - Appears in **Node Operational Certificate** subject under OID 1.3.6.1.4.1.37244.1.6 with value 0xFFFF0001
 - Would appear in X.509 certificate subject under OID 1.3.6.1.4.1.37244.1.6 as UTF8String FFFF0001 for signature validation purposes.
- Tag identifier 0xAB12, Version 0x0003
 - CAT value: 0xAB12_0003
 - Appears in Access Control Entry **Subjects** list as 0xFFFF_FFFD_AB12_0003
 - Appears in **Node Operational Certificate** subject under OID 1.3.6.1.4.1.37244.1.6 with value 0xAB120003
 - Would appear in X.509 certificate subject under OID 1.3.6.1.4.1.37244.1.6 as UTF8String AB120003 for signature validation purposes.
- Tag identifier 0x071C, Version 0x1074
 - CAT value: 0x071C_1074
 - Appears in Access Control Entry **Subjects** list as 0xFFFF_FFFD_071C_1074
 - Appears in **Node Operational Certificate** subject under OID 1.3.6.1.4.1.37244.1.6 with value 0x071C1074
 - Would appear in X.509 certificate subject under OID 1.3.6.1.4.1.37244.1.6 as UTF8String 071C1074 for signature validation purposes.

6.6.2.2. Wildcards

The Subjects list of an Access Control Entry MAY grant a given privilege to more than one Subject, if the Authentication Mode allows it, such as in the case of the **CASE** and Group Authentication Modes. An empty Subjects list SHALL mean that every possible Subject employing the stated Authentication Mode is granted the entry's privilege over the Targets.

The Targets list of an Access Control Entry MAY grant a given privilege to more than one Target. An empty Targets list SHALL mean that every Cluster on every Endpoint exposed by the Node is accessible using the granted privilege to any matching Subject. Each Target in the Targets list SHALL specify Cluster instances directly by Cluster ID (on any Endpoint, or limited to particular Endpoints), indirectly by Endpoint ID (all Cluster instances on that Endpoint), or indirectly by Device Type ID (all Cluster instances on all Endpoints containing that Device Type).

For both the Subjects list and Targets list of an Access Control Entry, empty lists permit a rudimentary form of "wildcard" behavior, which is especially useful for codifying policies providing common view/read/discover access to a given subset of Nodes based on Authentication Mode.

CAUTION

Given that "wildcard" (that is, **any** subject/target) granting is possible with an empty Subjects list or an empty Targets list, it follows that care must be taken by Administrators generating and distributing Access Control Lists to ensure unintended access does not arise. It is RECOMMENDED to avoid updating Access Control Entries in such a way as to remove Subjects or Targets one by one, which may result in a wildcard after individual actions. Rather, entire Targets/Subjects lists SHOULD be written atomically in a single action, to ensure a complete final state is achieved, with either wildcard or not, and that no accidental wildcards arise. Furthermore, such ACL entries with wildcard subject should be deployed with care as they grant the named privilege level to potentially many senders, especially when Group Authentication Mode is specified.

6.6.2.3. Subjects do not correspond to users

The Subjects for an Access Control Entry are logical subjects, configured through policy by an Administrator, including possibly a Commissioner during the commissioning phase. A given implementation of administrative logic MAY assign authentication identities to Nodes directly associated with physical end-users (for example, a mobile device of a given end-user). However, since Nodes are logical networking entities, the specific policy of how Node identities are mapped to physical end-users and physical devices is implementation-specific. Therefore, the access granted by a given Node's Access Control system should not be construed as having any particular meaning in regards to physical end-users other than the fact that a given set of Administrators computed a consistent set of Access Control Lists to effect a desired system functionality across all Nodes they administer and end-users they represent.

6.6.2.4. Implementation-specific Rules

Since the target of a given Access Control Entry is a list of Targets, and since Targets (that is, Clusters on Endpoints) are Interaction Model constructs, it should be assumed that access control functionality as described within this model is constrained to the interaction model layer. However, constraints on incoming session establishment requests MAY be affected by the Access Control system, based on implementation-defined rules. For example, a Node MAY deny [CASE](#) session establishment from an initiator whose identity doesn't match any Access Control Entry. These types of rules are implementation-specific and SHOULD be carefully considered, if applied at all. For example, due to the richness of Access Control Entry encoding for Subjects, significant care has to be taken to avoid incorrectly rejecting an incoming [CASE](#) session establishment that could be valid. Rejecting valid connections could cause a Node to become unreachable. Any constraints on transport-level and net-

work-level functionality, including but not limited to the availability of commissioning-mode connectivity, are out of Access Control scope.

6.6.2.5. Incoming Subject Descriptors

The Message Layer SHALL provide sufficient metadata (e.g. Authentication Mode, source identity) about incoming messages to the Interaction Model protocol layer in order to properly enforce Access Control.

An Incoming Subject Descriptor (ISD) is a mapping from the security layer fields of an incoming Message to a tuple of `<AuthMode, SubjectDescriptor>` that can map unambiguously to an Access Control Entry's `Subjects` and `AuthMode` fields. See [Section 6.6.6.1.3, "Incoming Subject Descriptor \(ISD\) Structure"](#) for further details.

6.6.2.6. Access Control Extensions

An implementation MAY use [Access Control Extensions](#) to extend the base Access Control model. Since all extensions are installed by Administrators for a fabric, it is expected that only extensions that would improve overall security will be applied. Since every Vendor MAY implement extensions as they see fit, it SHOULD NOT be expected that an extension will be supported by every Node. It is therefore RECOMMENDED that careful consideration of interoperability concerns be given when implementing Access Control Extensions. A fabric's Administrators MAY always read a given Node's Access Control Entries and Access Control Extensions pertaining to the fabric. Therefore, Administrators MAY use extensions to record auditing metadata about Access Control Entries which are not for operational use by the Node.

A Node SHALL preserve every field of the installed Access Control Cluster, including extensions when present, without internally-initiated modifications, so that they may be read-back verbatim upon receiving an appropriate request from an Administrator.

6.6.2.7. Application-level Permissions

The Access Control Cluster SHALL NOT be used to encode application-level permissions and configurations such as smart lock PIN codes or similar user-facing security functionality. Application-level security is best served by finer-grained capabilities described and addressed by application-domain-specific clusters.

6.6.2.8. Access Restrictions

An Access Control Entry in the ACL describes permissions granted by the fabric Administrator. In contrast, an Access Restriction Entry in the ARL describes restrictions imposed by the device. The ARL restrictions override permissions granted in the ACL.

6.6.2.9. Bootstrapping of the Access Control Cluster

Updates to the Access Control List through Access Control Cluster attributes and commands SHALL be restricted by the same Access Control mechanisms as all other clusters on the Node, and therefore require a grant of `Administer` privilege. Administrators are able to bootstrap a Node's Access Control List during the [commissioning phase](#) due to the [Access Control Privilege Granting algorithm](#) implicitly granting the `Administer` privilege to Administrative Subject Nodes over a PASE commis-

sioning channel; this implicit privilege grant applies for the Commissioner to administer the Commissionee, but not in the opposite direction.

6.6.2.10. Action attribution

The recording of a given Interaction Model Action's attribution to a source entity is distinct from the contents of an Access Control Entry. Action Attribution SHALL be recorded against the Incoming Subject Descriptor (see [Section 6.6.6.1.3, "Incoming Subject Descriptor \(ISD\) Structure"](#)) rather than against any matched Access Control Entry's contents.

6.6.2.11. Restrictions on Administer Level Privilege Grant

Since the **Administer** privilege level grants wide access to a Node for a given Subject, it SHALL NOT be valid to have an **Administer** privilege set on an Access Control Entry, unless **AuthMode** is "CASE". For example, an **AuthMode** of "Group", which admits no source Node authentication and reduced attribution ability, SHALL NOT be used to grant **Administer** privilege.

6.6.3. Access Control List Examples

The following Access Control Lists illustrate the flexibility of codifying access control policy using concrete examples.

Upon Factory Data Reset, the Access Control Cluster is empty, having an Access Control List with no entries.

```
Access Control Cluster: {
  ACL: [],          // empty
  Extension: [],    // empty (omitted hereafter)
  ARL: []           // empty for this example (omitted hereafter)
}
```

However, the Access Control Privilege Granting algorithm behaves as if, over a PASE commissioning channel during the [commissioning phase](#), the following implicit Access Control Entry were present on the Commissionee (but not the Commissioner) to grant **Administer** privilege for the entire Node.

```
Access Control Cluster: {
  ACL: [
    0: {              // implicit entry only; does not explicitly exist!
      FabricIndex: 0, // not fabric-specific
      Privilege: Administer,
      AuthMode: PASE,
      Subjects: [],
      Targets: []     // entire node
    }
  ]
}
```


}

During the commissioning phase, the [AddNOC](#) command automatically creates an Access Control Entry granting [Administer](#) privilege for the appropriate [CaseAdminSubject](#).

In the scenario the node is being commissioned onto a separately-administered fabric, [Administer](#) privilege is granted for the entire Node, using the appropriate CASE authenticated Subject (in this case, Node ID 0xAAAA_AAAA_AAAA_AAAA) on the appropriate Fabric (in this case, Fabric 0xFAB0_0000_0000_001D as Fabric index 1).

```
Access Control Cluster: {
  ACL: [
    0: {
      FabricIndex: 1,
      Privilege: Administer,
      AuthMode: CASE,
      Subjects: [ 0xAAAA_AAAA_AAAA_AAAA ],
      Targets: []
    }
  ]
}
```

In the scenario the node is being commissioned onto the Joint Fabric the [Administer](#) privilege is granted for all Nodes with the Administrator CAT as the [CaseAdminSubject](#) (in this case, 0xFFFF_FFFD_FFFF_0001) on the appropriate Fabric (in this case, Fabric 0xFAB0_0000_0000_001D as Fabric index 1).

```
Access Control Cluster: {
  ACL: [
    0: {
      FabricIndex: 1,
      Privilege: Administer,
      AuthMode: CASE,
      Subjects: [ 0xFFFF_FFFD_FFFF_0001 ],
      Targets: []
    }
  ]
}
```

An Administrator adds an Access Control Entry which grants [View](#) privilege, for the entire Node, to all CASE authenticated Nodes.

```
Access Control Cluster: {
```



```

ACL: [
  ...
  1: {
    FabricIndex: 1,
    Privilege: View,
    AuthMode: CASE,
    Subjects: [],
    Targets: []
  }
]
}

```

An Administrator adds an Access Control Entry which grants **Manage** privilege, for endpoints 1 and 3, to any Nodes which can authenticate as members of Group 1.

```

Access Control Cluster: {
  ACL: [
    ...
    2: {
      FabricIndex: 1,
      Privilege: Manage,
      AuthMode: Group,
      Subjects: [ 0x0000_0000_0000_0001 ],
      Targets: [ { Endpoint: 1 }, { Endpoint: 3 } ]
    }
  ]
}

```

An Administrator revises this Access Control Entry to grant the same privilege, for only the pump configuration and control cluster (0x0202) on endpoint 3, and for any door lock cluster (0x0101) on the entire Node, to the same Nodes.

```

Access Control Cluster: {
  ACL: [
    ...
    2: {
      FabricIndex: 1,
      Privilege: Manage,
      AuthMode: Group,
      Subjects: [ 0x0000_0000_0000_0001 ],
      Targets: [ { Endpoint: 1 }, { Endpoint: 3, Cluster: 0x0000_0202 }, { Cluster:
0x0000_0101 } ]
    }
  ]
}

```

```
]
}
```

An Administrator adds an Access Control Entry which grants **Operate** privilege, for all endpoints containing the extended color light device (0x010D) on the entire Node, to CASE authenticated Nodes 0x1111_1111_1111_1111 and 0x2222_2222_2222_2222.

```
Access Control Cluster: {
  ACL: [
    ...
    3: {
      FabricIndex: 1,
      Privilege: Operate,
      AuthMode: CASE,
      Subjects: [ 0x1111_1111_1111_1111, 0x2222_2222_2222_2222 ],
      Targets: [ { DeviceType: 0x0000_010D } ]
    }
  ]
}
```

A Commissioner adds four more Nodes into an existing fabric. These new Nodes have Node IDs 0x3333_3333_3333_3333, 0x4444_4444_4444_4444, 0x5555_5555_5555_5555 and 0x6666_6666_6666_6666 respectively. The Fabric Administration policy requires associating these four nodes and an existing node (0x2222_2222_2222_2222) into a CAT group.

To achieve this, an Administrator will issue NOCs to all five nodes in this CAT group with a CAT value of 0xABCD_0001, which is tag identifier value 0xABCD and version 1, and encoded as subject value 0xFFFF_FFFD_ABCD_0001. To distribute these CATs, an Administrator obtains NOCs from its certificate authority with the requisite subjects including the desired CAT. They are either initially provisioned with the **AddNOC** command during initial commissioning (for the new Nodes) or updated with **UpdateNOC** (for existing Nodes).

Then the Administrator grants permissions to the five nodes by updating the ACL of all relevant targets by adding an entry with subject of **CAT** (0xFFFF_FFFD_ABCD_0001). The Administrator may also remove entries where Node 0x2222_2222_2222_2222 appears as an explicit Subject if presentation of the CAT identifier value 0xABCD and version value 0x0001 confers an equivalent privilege. Note that any Node with CAT identifier value of 0xABCD and version value 0x0001 or higher in their NOC will have this privilege.

```
Access Control Cluster: {
  ACL: [
    ...
    3: {
      FabricIndex: 1,
      Privilege: Operate,
```

```

    AuthMode: CASE,
    Subjects: [ 0x1111_1111_1111_1111, 0xFFFF_FFFD_ABCD_0001],
    Targets: [ { DeviceType: 0x0000_010D } ]
  }
]
}

```

An Administrator wants to remove the Node with Node ID 0x3333_3333_3333_3333 from the CAT group defined by CAT identifier value of 0xABCD, as installed in the previous example.

The Administrator could follow the steps outlined below:

1. Administrator will ensure that the removed node having Node ID 0x3333_3333_3333_3333 will not receive a new NOC with an CAT identifier value of 0xABCD. Note that the node removed from the group will continue to hold existing NOC (with CAT identifier of 0xABCD and version 0x0001).
2. An Administrator updates NOCs with CAT identifier value of 0xABCD and version 0x0002 (encoded as subject 0xFFFF_FFFD_ABCD_0002) in all remaining currently reachable Nodes within the CAT group to ensure they continue to have same privilege as before. In this example, Nodes having Node IDs 0x2222_2222_2222_2222, 0x4444_4444_4444_4444 and 0x5555_5555_5555_5555 are currently reachable.
3. Node with Node ID 0x6666_6666_6666_6666 from this CAT group is a valid member but was not reachable by an Administrator at the time of this change. This Node will continue to hold existing NOC with CAT of 0xABCD_0001.
4. After updating NOCs of all reachable Nodes, the Administrator revises the Access Control Entry of all reachable nodes who have the previous CAT (encoded as subject 0xFFFF_FFFD_ABCD_0001) in an ACL entry, to remove privilege from the Node no longer in the grouping (i.e. those with version 0x0001) by increasing trusted version value to be higher than 0x0001. The Administrator decides to increment version value by one to set the new version value to be 0x0002.
5. Once ACL changes are propagated to all controlled nodes, they will no longer allow access privileges to any Node with older version (i.e. value less than 0x0002) of CAT identifier value 0xABCD. Hence, the node removed from the group, having Node ID 0x3333_3333_3333_3333 and CAT with identifier 0xABCD and version of 0x0001, can no longer access any of the controlled nodes whose ACL entries were updated to have a subject of 0xFFFF_FFFD_ABCD_0002 (CAT identifier value 0xABCD, version 0x0002).
6. Node having Node ID of 0x6666_6666_6666_6666 will not be able to access any Nodes by relying on CAT, since it does not have an NOC with latest CAT (with version 0x0002). However, it can still access Nodes that list it as a subject Node ID explicitly. When an Administrator eventually establishes connection to this Node, the Administrator updates the NOC to the latest version, with CAT set to 0xABCD_0002. After having its NOC updated to have the newest version of the CAT, the Node with Node ID 0x6666_6666_6666_6666 will again have access to Nodes that list subject 0xFFFF_FFFD_ABCD_0002 (CAT identifier value 0xABCD, version 0x0002), with no further updates to ACL entries of existing Nodes.
7. Any controlled Node which previously held an ACL Entry with prior version of the updated CAT

(subject 0xFFFF_FFFD_ABCD_0001) but was not reachable by an Administrator at the time of update, will continue to hold the previous Access Control Entry with a subject allowing CAT with identifier of 0xABCD and version 0x0001 or higher. Thus, these Nodes will grant privileges to any Node from the original CAT group (including Node ID 0x3333_3333_3333_3333). When an Administrator eventually establishes connection to this Node with older ACL entry, the Administrator updates it with the latest value, so that Node ID 0x3333_3333_3333_3333 no longer has privileges.

Note that in the above example, the CAT identifier value remained the same (0xABCD) in NOCs and ACL entries throughout these steps. Only the version portion was updated to effect changes to the meaning of the CAT.

As can be seen in the example above, there are multiple steps involving updates to NOCs and ACL entries to affect CAT-based grouping and aliasing policies. It is therefore possible that some Nodes may not receive these changes immediately, due to network reachability issues, such as being powered down for an extended period, and thus have ACL entries or NOCs that grant temporarily obsolete privileges. This is true as well with direct Node ID subjects, in general.

Administrators aim for best-effort eventual consistency while executing the steps outlined above.

```
Access Control Cluster: {
  ACL: [
    ...
    3: {
      FabricIndex: 1,
      Privilege: Operate,
      AuthMode: CASE,
      Subjects: [ 0x1111_1111_1111_1111, 0xFFFF_FFFD_ABCD_0002 ],
      Targets: [ { DeviceType: 0x0000_010D } ]
    }
  ]
}
```

6.6.4. Access Control Cluster update side-effects

Updates to the Access Control Cluster SHALL take immediate effect in the Access Control system. For example, given an Interaction Model action message containing the following actions, the Access Control Privilege Granting algorithm would grant a privilege of **None** for the second action, since the first action would take effect immediately beforehand.

- Pre-conditions:
- Access Control List has single entry: [{Privilege: Administer, Authmode: CASE (2), Subjects: [0x0011223344556677], Targets: []}]
- Node ID 0x0011223344556677 over CASE is allowed Administer privilege for all targets
- Incoming message Source is Node ID 0x0011223344556677 over CASE: matches Access Control

Entry subject

- Actions:
 1. Action: Write (1st path)
 1. Path: Endpoint[0]/Cluster[AccessControl]/Attribute[ACL]/ListIndex[0]/Field[Targets]
 2. Value: `[{Endpoint: 2}]`
 - Single entry updated to target only Endpoint 2
 3. **Granted:** Administer privilege granted, due to Access Control Entry match
 2. Action: Write (2nd path)
 1. Path: Endpoint[1]/Cluster[OnOff]/Attribute[OnTime]
 2. **Denied:** No privilege granted, because prior action **in the same message** had updated Access Control List to only allow access to Endpoint 2, and this action targets Endpoint 1
- Post-conditions:
 - Access Control List has single entry, updated by first path of Write Action: `[{Privilege: Administer, Authmode: CASE (2), Subjects: [0x0011223344556677], Targets: [{Endpoint: 2}]}]`
 - Node ID 0x0011223344556677 over CASE is allowed **Administer** privilege for only Endpoint 2 target

Note that in this example, the Node has inadvertently lost its ability to update the Access Control Cluster by limiting its Administer privilege to Endpoint 2.

6.6.5. Access Control Cluster handling of Access Restrictions

In the example above, the ARL was empty, meaning that there were no access restrictions imposed upon the fabric by the device. However, the following Access Control Cluster examples illustrate how access restrictions would impact access granted to a given client.

Upon Factory Data Reset, the Access Control Cluster has an Access Control List with no entries, an Access Restriction List with no entries, and a CommissioningARL set by the device which applies during commissioning sessions.

```
Access Control Cluster: {
  ACL: [],           // empty
  Extension: [],     // empty (omitted hereafter)
  CommissioningARL: [
    0: {
      Endpoint: 1,      // application endpoint 1
      Cluster: 0x0453,  // application cluster 0x0453
      Restrictions: [
        0: {
          Type: AttributeAccessForbidden, // access limitation for attribute
          Id: 0x0000                      // attribute 0x0000
        },
      ],
    },
  ],
}
```

```

    1: {
        Type: CommandForbidden, // command limitation
        Id: NULL                // all commands
    },
    2: {
        Type: EventForbidden, // event limitation
        Id: NULL              // all events
    }
]
}
]
ARL: [] // empty
}

```

Upon allocation of a new FabricIndex during commissioning (see [Section 11.18.6.8, “AddNOC Command”](#)), one or more ARL entries with the new FabricIndex MAY be added to the ARL attribute at the time that the initial ACL entry for the new FabricIndex is added. Until CommissioningComplete command is invoked in the GeneralCommissioning cluster, the CommissioningARL continues to apply.

Example:

```

Access Control Cluster: {
    ACL: [
        0: {
            FabricIndex: 1,
            Privilege: Administer,
            AuthMode: CASE,
            Subjects: [ 0xAAAA_AAAA_AAAA_AAAA ],
            Targets: []
        }
    ],
    Extension: [], // empty (omitted hereafter)
    CommissioningARL: [
        0: {
            Endpoint: 1, // application endpoint 1
            Cluster: 0x0453, // application cluster 0x0453
            Restrictions: [
                0: {
                    Type: AttributeAccessForbidden, // access limitation for attribute
                    Id: 0x0000 // attribute 0x0000
                },
                1: {

```

```

        Type: CommandForbidden, // command limitation
        Id: NULL                // all commands
    },
    2: {
        Type: EventForbidden, // event limitation
        Id: NULL              // all events
    }
]
}
]
ARL: [
    0: {
        FabricIndex: 1,
        Endpoint: 1,          // application endpoint 1
        Cluster: 0x0453,      // application cluster 0x0453
        Restrictions: [
            0: {
                Type: AttributeAccessForbidden, // access limitation for attribute
                Id: 0x0000                      // attribute 0x0000
            },
            1: {
                Type: CommandForbidden, // command limitation
                Id: NULL                // all commands
            },
            2: {
                Type: EventForbidden, // event limitation
                Id: NULL              // all events
            }
        ]
    }
]
}

```

In this example, the restrictions during commissioning, expressed by the CommissioningARL attribute, include 3 limitations on Cluster 0x0453 of Endpoint 1: access to attribute 0x0000, all commands, all events. The restrictions for the fabric at FabricIndex 1, expressed by the ARL attribute, include the same 3 limitations on Cluster 0x0453 of Endpoint 1.

The CommissioningARL and ARL attributes SHALL NOT include limitations that would prevent commissioning (see [Section 9.10.4.2.1, “Managed Device Feature Usage Restrictions”](#)). If a wildcard restriction appears in the ARL that would restrict access in such a case, the wildcard MAY appear but the restriction SHALL NOT take effect.

Superimposing this ARL value upon the example earlier in this section, when an Administrator

adds, for example, an Access Control Entry which grants **Operate** privilege for the entire Node, to all CASE-authenticated subjects), this grant MAY be further limited by restrictions imposed by the ARL. In other words, even though the ACL entry grants **Operate** privilege to all data model elements, attempts to read or write attribute 0x0000, or to invoke commands upon Cluster 0x0453 of Endpoint 1 would result in an error of ACCESS_RESTRICTED, since the Access Restriction List is a subsequent overriding of an initial privilege granted.

```

Access Control Cluster: {
  ACL: [
    ...
    1: {
      FabricIndex: 1,
      Privilege: View,
      AuthMode: CASE,
      Subjects: [],
      Targets: []
    }
  ],
  CommissioningARL [ ... ], // (omitted hereafter)
  ARL: [
    0: {
      FabricIndex: 1,
      Endpoint: 1,           // application endpoint 1
      Cluster: 0x0453,       // application cluster 0x0453
      Restrictions: [
        0: {
          Type: AttributeAccessForbidden, // access limitation for attribute
          Id: 0x0000                      // attribute 0x0000
        },
        1: {
          Type: CommandForbidden, // command limitation
          Id: NULL                // all commands
        },
        2: {
          Type: EventForbidden, // event limitation
          Id: NULL              // all events
        }
      ]
    }
  ]
}

```

Updates to the ARL attribute can happen at any time, for example, as a result of configuration

changes made by the user within an external management user interface, out of band of the Matter protocol. In particular, the fabric Administrator can request that the device initiates a review of fabric restrictions using the `ReviewFabricRestrictions` command on the Access Control Cluster. This command returns a `ReviewFabricRestrictionsResponse` containing a `Token` which can be used to correlate with a `FabricRestrictionReviewUpdate` event, indicating completion of the fabric restriction review. Such a review can result in an asynchronous update to the ARL attribute and an `AccessRestrictionEntryChanged` event.

In this example, after the fabric restriction review, the ARL value is updated: the result is a single limitation on Cluster 0x0453 of Endpoint 1: access to command 0x0000. In other words, restrictions on attributes and events on Cluster 0x0453 of Endpoint 1 have been removed, and all command restrictions on the cluster except command 0x0000 have been removed.

```
Access Control Cluster: {
  ACL: [
    ...
    1: {
      FabricIndex: 1,
      Privilege: View,
      AuthMode: CASE,
      Subjects: [],
      Targets: []
    }
  ],
  ARL: [
    ...
    1: {
      FabricIndex: 1,
      Endpoint: 1,           // application endpoint 1
      Cluster: 0x0453,       // application cluster 0x0453
      Restrictions: [
        0: {
          Type: CommandForbidden, // command limitation
          Id: 0x0000              // command 0x0000
        }
      ]
    }
  ]
}
```

6.6.6. Conceptual Access Control Privilege Granting Algorithm

This section describes an overall Conceptual Access Control Privilege Granting algorithm. Implementations of this algorithm SHALL have an identical outcome to the output of this conceptual

algorithm described below.

The Interaction Model protocol, through its message handling, SHALL determine the privilege level granted per Target, on every instance where a Target is referenced for use.

6.6.6.1. Necessary Data Structures

6.6.6.1.1. Access Control List

The [Access Control List](#) contains several [Access Control Entries](#), previously described in [Section 6.6.2, “Access Control Model”](#).

The entry fields are:

- Subjects List ([subject-id\[\]](#) [Subjects](#))
- Targets List ([AccessControlTargetStruct\[\]](#) [Targets](#))
- Authentication Mode value ([AccessControlEntryAuthModeEnum](#) [AuthMode](#))
- Privilege value ([AccessControlEntryPrivilegeEnum](#) [Privilege](#))

6.6.6.1.2. Access Restriction List

The [Access Restriction List](#) contains several [Access Restriction Entries](#), previously described in [Section 6.6.2, “Access Control Model”](#).

The entry fields are:

- Endpoint ([endpoint-no](#) [Endpoint](#))
- Cluster ([cluster-id](#) [Cluster](#))
- Restrictions List ([AccessRestrictionStruct\[\]](#) [Restrictions](#))

6.6.6.1.3. Incoming Subject Descriptor (ISD) Structure

Each incoming message has a unique [AuthMode](#), [SubjectDescriptor](#) applicable to it, whose derivation is deterministic based on both incoming message fields and session metadata fields. For example, if a message arrives that matches a given CASE Session ID, then the metadata for that CASE session would be used.

Computation of the ISD is described in [Section 6.6.6.3, “Derivation of ISD from Incoming Message”](#).

The ISD fields are as follows:

- Commissioning Flag ([bool](#) [IsCommissioning](#)), whether the authentication is over a commissioning channel for an ongoing commissioning.
- Authentication Mode ([AuthModeEnum](#) [AuthMode](#)), mapping to an authentication mode, directly comparable to [Access Control Entry AuthMode](#).
- Subjects List ([list<SubjectID>](#) [Subjects](#)), mapping incoming message source to a type of subject, such as a CASE session Source Node ID.
- Fabric Index ([FabricIndex](#) [FabricIndex](#)), mapping to a fabric reference.

6.6.6.2. Overall Algorithm

The algorithm takes as input:

- the Access Control List (**acl**) containing the logical set of all entries from the ACL and Auxiliary-ACL (if present) attributes from the Access Control Cluster
- the Privilege (**request_privilege**) required for the given request
- the ISD of Incoming Message (**subject_desc**)
- the Endpoint ID (**endpoint_id**) for which the querier requires a Privilege level
- the Cluster ID (**cluster_id**) for which the querier requires a Privilege level
- the Cluster Element (**cluster_element**) for which the querier requires a Privilege level
- the Request Type (**request_type**) for which the querier requires a Privilege level (ReadAttribute, ReadEvent, WriteAttribute, InvokeCommand)
- the ARL attribute (**arl_attr**) from the Access Control Cluster
- the CommissioningARL attribute (**commissioning_arl_attr**) from the Access Control Cluster

NOTE

Logically speaking, the total ACL is the union of two attributes (ACL and, if present, AuxiliaryACL) in the Access Control Cluster. Implementations are free to manage the storage of the entries in the way that is most appropriate without loss of generality to the Algorithm presented herein.

The output of the algorithm is:

- The access control status of the request, which is one of {AccessGranted, AccessRestricted, AccessDenied}.

The algorithm first determines the set of privileges granted for the Action Path, which is a subset of {View, Operate, Manage, Administer} as described in [Section 9.10.5.2, “AccessControlEntryPrivilegeEnum”](#). If the Privilege required for the request is contained within this set, and the request is not restricted by the Access Restriction List, then the AccessGranted status is returned. If the Privilege required for the request is contained within this set, and the request is restricted by the Access Restriction List, then the AccessRestricted status is returned. Otherwise, the AccessDenied status is returned.

The computation of the ISD is a pre-condition to the algorithm and is described in [Section 6.6.6.3, “Derivation of ISD from Incoming Message”](#).

The goal is to find the complete set of privileges granted given the input. The principle of least privilege is respected by virtue of the entire Access Control List having been computed with rules such that the least privilege is granted to all subjects. Therefore, any Access Control Entry granting the **required privilege** to the subject for a given target is sufficient to determine whether access is allowed.

The algorithm SHALL function as follows:

```
def subject_matches(acl_subject, isd_subject):
```

```
# Subjects must match exactly, or both are CAT
# with matching CAT ID and acceptable CAT version
return (acl_subject == isd_subject) or
    (is_cat(acl_subject) and is_cat(isd_subject) and
    (get_cat_id(acl_subject) == get_cat_id(isd_subject)) and
    (get_cat_version(isd_subject) >= get_cat_version(acl_subject)))

def add_granted_privilege(granted_privileges, privilege):
    # Add the new privilege to the granted privileges set
    granted_privileges.add(privilege)
    # Also add any privileges subsumed by the new privilege
    if (privilege == PrivilegeEnum.Operate):
        granted_privileges.add(PrivilegeEnum.View)
    elif (privilege == PrivilegeEnum.Manage):
        granted_privileges.add(PrivilegeEnum.Operate)
        granted_privileges.add(PrivilegeEnum.View)
    elif (privilege == PrivilegeEnum.Administer):
        granted_privileges.add(PrivilegeEnum.Manage)
        granted_privileges.add(PrivilegeEnum.Operate)
        granted_privileges.add(PrivilegeEnum.View)

def get_granted_privileges(acl, acl_feature_map, subject_desc, endpoint_id,
cluster_id) -> set[Privilege]:
    # Granted privileges set is initially empty
    granted_privileges = set()

    # PASE commissioning channel implicitly grants administer privilege to commissioner
    if subject_desc.AuthMode == AuthModeEnum.PASE and
subject_desc.IsCommissioneeDuringCommissioning():
        add_granted_privilege(granted_privileges, PrivilegeEnum.Administer)

    for acl_entry in acl:
        # End checking if highest privilege is granted
        if PrivilegeEnum.Administer in granted_privileges: break

        # FabricIndex must match, there are no valid entries with FabricIndex == 0
        # other than the implicit PASE entry, which we will not see explicitly in the
        # access control list
        if acl_entry.FabricIndex == 0: continue
        if acl_entry.FabricIndex != subject_desc.FabricIndex: continue

        # Auth mode must match
        if acl_entry.AuthMode != subject_desc.AuthMode: continue
```

```
# Subject must match, or be "wildcard"
if is_empty(acl_entry.Subjects):
    # Precondition: only CASE and Group auth can have empty subjects
    assert(acl_entry.AuthMode in [AuthModeEnum.CASE, AuthModeEnum.Group])

    # Empty is wildcard, no match required
else:
    # Non-empty requires a match
    matched_subject = False
    for acl_subject in acl_entry.Subjects:
        for isd_subject in subject_desc.Subjects:
            if subject_matches(acl_subject, isd_subject):
                matched_subject = True
                break
        if matched_subject: break
    if not matched_subject: continue

# Target must match, or be "wildcard"
if is_empty(acl_entry.Targets):
    # Group subjects cannot grant access to Endpoint 0 when Auxiliary feature is
    # enabled.
    if (acl_entry.AuthMode == AuthModeEnum.Group) and (acl_feature_map & Auxiliary
    != 0) and (endpoint_id == 0):
        continue

    # Otherwise, empty is wildcard, no match required
else:
    # Non-empty requires a match
    matched_target = False
    for target in acl_entry.Targets:
        # Precondition: target cannot be empty
        assert(target.Cluster != null or target.Endpoint != null or target.DeviceType
        != null)

        # Precondition: target cannot specify both endpoint and device type
        assert(target.Endpoint == null or target.DeviceType == null)

        # Cluster must match, or be wildcard
        if target.Cluster != null and target.Cluster != cluster_id:
            continue

        # Endpoint must match, or be wildcard
        if target.Endpoint != null and target.Endpoint != endpoint_id:
            continue

        # Endpoint may be specified indirectly via device type
```

```
        if target.DeviceType != null and not
endpoint_contains_device_type(endpoint_id, target.DeviceType):
            continue
        matched_target = True
        break
    if not matched_target: continue

    # Extensions processing must not fail
    if not extensions_are_valid(acl, acl_entry, subject_desc, endpoint_id,
cluster_id): continue

    # All checks have passed, add privilege to granted privilege set
    add_granted_privilege(granted_privileges, acl_entry.privilege)

    # Should never grant Administer privilege to a Group.
    if subject_desc.AuthMode == AuthModeEnum.Group:
        assert (PrivilegeEnum.Administer not in granted_privileges)

    return granted_privileges

# Note that `arl` is either the `arl_attr` or `commissioning_arl_attr`.
def is_request_restricted(arl, subject_desc, request_type, endpoint_id, cluster_id,
cluster_element):
    # Handles ensuring that only endpoints and clusters for which access restrictions
are allowed
    # are subject to the restriction. Only mandatory clusters on endpoints with device
type allowing
    # Managed Device ACL cluster feature can be restricted. Some specific clusters and
elements can
    # never be restricted. See the Managed Device feature section of the ACL cluster for
details.
    if are_restrictions_disallowed(request_type, endpoint_id, cluster_id,
cluster_element):
        return False

    for arl_entry in arl:
        if not subject_desc.IsCommissioning:
            if arl_entry.FabricIndex != subject_desc.FabricIndex: continue

        if arl_entry.Endpoint != endpoint_id: continue

        if arl_entry.Cluster != cluster_id: continue

        for arl_restriction in arl_entry.Restrictions:
```



```
# if the restriction is attribute access
if arl_restriction.Type == AccessRestrictionTypeEnum.AttributeAccessForbidden
and request_type in [RequestTypeEnum.AttributeRead, RequestTypeEnum.AttributeWrite]:
    # if the attribute id is NULL (wildcard) or matches the cluster_element
    if arl_restriction.ID == NULL or arl_restriction.ID == cluster_element:
        # then no granted privileges
        return True

if arl_restriction.Type == AccessRestrictionTypeEnum.AttributeWriteForbidden and
request_type == RequestTypeEnum.AttributeWrite:
    # if the attribute id is NULL (wildcard) or matches the cluster_element
    if arl_restriction.ID == NULL or arl_restriction.ID == cluster_element:
        # then no granted privileges
        return True

if arl_restriction.Type == AccessRestrictionTypeEnum.CommandForbidden and
request_type == RequestTypeEnum.CommandInvoke:
    # if the command id is NULL (wildcard) or matches the cluster_element
    if arl_restriction.ID == NULL or arl_restriction.ID == cluster_element:
        # then no granted privileges
        return True

if arl_restriction.Type == AccessRestrictionTypeEnum.EventAccessForbidden and
request_type == RequestTypeEnum.EventRead:
    # if the event id is NULL (wildcard) or matches the cluster_element
    if arl_restriction.ID == NULL or arl_restriction.ID == cluster_element:
        # then no granted privileges
        return True

return false

def get_access_status(
    acl, acl_feature_map, arl_attr, commissioning_arl_attr,
    request_privilege, subject_desc, request_type, endpoint_id, cluster_id,
    cluster_element) -> access_status:
    granted_privileges = get_granted_privileges(acl, acl_feature_map, subject_desc,
    endpoint_id, cluster_id)

    if request_privilege in granted_privileges:
        # Select the correct source of ARL depending on context.
        arl = commissioning_arl_attr if subject_desc.IsCommissioning else arl_attr
        if is_request_restricted(arl, subject_desc, request_type, endpoint_id, cluster_id,
        cluster_element):
            return AccessStatusEnum.AccessRestricted
```

```

else:
    return AccessStatusEnum.AccessGranted

return AccessStatusEnum.AccessDenied

```

6.6.6.3. Derivation of ISD from Incoming Message

The algorithm to derive the ISD from an incoming message takes as input:

- The incoming message (**message**)
- The Session ID of the incoming message (**session_id**)
- A conceptual Sessions Metadata database (**sessions_metadata**)
- The Group Key Management Cluster (**group_key_management_cluster**)

The output of the algorithm is the **SubjectDescriptor** structure below:

```

DEFAULT_COMMISSIONING_PASSCODE = 0

enum AuthModeEnum {
    None = 0, # conceptual "no auth" value
    PASE = 1,
    CASE = 2,
    Group = 3
}

struct SubjectDescriptor {
    bool IsCommissioning;
    AuthModeEnum AuthMode;
    list<SubjectID> Subjects; # max 3 items
    FabricIndex FabricIndex;
}

```

The algorithm SHALL function as follows:

```

def get_isd_from_message(message) -> SubjectDescriptor:
    isd = {
        IsCommissioning: False,
        AuthMode: AuthModeEnum.None,
        Subjects: [],
        FabricIndex: 0
    }

    session_id = message.get_session_id()

```

```

if sessions_metadata.get_auth_mode(session_id) == AuthModeEnum.PASE:
    isd.AuthMode = AuthModeEnum.PASE
    isd.IsCommissioning = True
    isd.Subjects.append(DEFAULT_COMMISSIONING_PASSCODE)
    isd.FabricIndex = sessions_metadata.get_fabric_index(session_id) # may be zero
else if sessions_metadata.get_auth_mode(session_id) == AuthModeEnum.CASE:
    isd.IsCommissioning = sessions_metadata.get_fabric_index(session_id) ==
sessions_metadata.get_pending_fabric_index()
    isd.AuthMode = AuthModeEnum.CASE
    isd.Subjects.append(sessions_metadata.get_src_node_id(session_id))
    # CASE session may contain CATs which also serve as subjects
    # Append all CATs if present (can be up to 3)
    if sessions_metadata.has_src_case_authenticated_tags(session_id):

isd.Subjects.append(sessions_metadata.get_src_case_authenticated_tags(session_id))
    isd.FabricIndex = sessions_metadata.get_fabric_index(session_id)
    assert(isd.FabricIndex != 0) # cannot be zero
else if sessions_metadata.get_auth_mode(session_id) == AuthModeEnum.Group:
    # Message is assumed to have been decrypted and matched properly prior to
    # this procedure occurring.
    group_id = message.get_dst_group_id()
    group_key_id = sessions_metadata.get_group_key_id(message)
    # Group membership must be verified against Group Key Management Cluster
    if group_key_management_cluster.group_key_map_has_mapping(group_id, group_key_id):
        isd.AuthMode = AuthModeEnum.Group
        isd.Subjects.append(group_id)
        isd.FabricIndex = sessions_metadata.get_fabric_index(message)
        assert(isd.FabricIndex != 0) # cannot be zero
else:
    # Do nothing on error, ISD remains unchanged
    assert(isd.IsCommissioning == False)
    assert(isd.AuthMode == AuthModeEnum.None)
    assert(is_empty(isd.Subjects))
    assert(isd.FabricIndex == 0)

return isd

```

6.6.7. Applying Privileges to Action Paths

The Data Model specifies which privilege is required for each data element, via its access qualities.

The Interaction Model specifies how each action is processed, for both its request and its response. This includes details on how the Interaction Model uses Access Control to determine whether to

allow the request (i.e. continue processing), or to deny the request (and whether/how that is indicated in the response).

Determining whether to allow or deny an action for a request path entails:

- Determining the [required privilege](#) for the action, given the action, request path and type of access requested;
- Determining the [set of granted privileges](#) for the action, given the request path and requesting subject;
- Checking whether the required privilege is present in the set of granted privileges:
 - If present, the action is allowed;
 - If not present, the action is denied.

Note that the Interaction Model may allow the action for some request paths while denying it for other request paths in the same action. Also, note that Access Control is merely one of the checks used by the Interaction Model, and an action that is allowed by Access Control may fail for other reasons.

Chapter 7. Data Model Specification

7.1. Practical Information

7.1.1. Revision History

Please note that Matter revision 1 SHALL be considered equivalent to revision 16, Subsequent releases SHALL start at revision 17.

Revision	Description
1	Equivalent to revision 16 (1st Matter release of the Data Model)
2-11	Approved specifications, but not released for implementation.
12-15	Released as revisions of the Dotdot Architecture Model specification
16	Initial release of this specification
17	Added tag data type
18	Added the LocationDescriptorStruct data type
19	Changed F quality description
20	Removed P quality
21	Added Revision conformance

7.1.2. Scope & Purpose

This is part of a package of Data Model specifications that are agnostic to underlying layers (encoding, message, network, transport, etc.). Each specification below may be independently maintained. This package, as a whole, SHALL be independently maintained as agnostic and decoupled from lower layers. This package may be referenced by inclusion in vertical protocol stack specifications.

Data Model	Defines first order elements and namespace for endpoints, clusters, attributes, data types, etc.
Interaction Model	Defines interactions, transactions and actions between nodes.
System Model	Defines relationships that are managed between endpoints and clusters.
Cluster Library	Reference library of cluster specifications.
Device Library	Reference library of device type specifications.

7.1.3. Origin Story

The origin of this section is the Dotdot Architecture Model [[Dotdot Architecture](#)] and parts of Chap-

ter 2 of the Zigbee Cluster Library specification [ZCL] that define the data model.

This document defines the specific model used in Matter, which has diverged from Zigbee.

7.1.4. Overview

This document defines first order elements and namespace of the data model and can also be called the meta-model (of the data model). This document is the "read me first" specification in the data model. This data model is ultimately implemented in the application layer of a communication stack.

This specification does not define how data is stored, encoded or what interactions are allowed on the data.

7.1.5. Glossary

Term	Description
MS	Manufacturer or Vendor Specific
N/A	not applicable
desc	see detailed description section

7.1.6. Conventions

See the [Conventions](#) section for conventions used in the Data Model specifications.

7.2. Data Qualities

[Cluster](#) specifications and [device type](#) definitions have tables that define the qualities of elements that make up the cluster or device type. Not all elements have all qualities. For example, a command does not have the read quality. Some elements have intrinsic qualities, that are not listed. For example, an event always has the read quality. Qualities SHALL be defined in columns in tables that describe data model elements.

7.2.1. Common Data Table Columns

The following columns are common across tables describing attributes, commands, events and structs:

- ID**
Defines an identifier for the data model element that is unique at its context.
- Name**
Defines a CamelCase name of the element to be used in specification text, not the protocol. Text usage SHALL always be followed with the element name (e.g. CurrentLevel attribute, Stopped event, or Left field).

Field

Same as **Name**. Other headers like "Field", "Bit Field", and "Command Field", are deprecated. Use **Name**.

Conformance

Defines dependencies on whether an element is optional or mandatory.

Access

Defines how an element is accessed (e.g. read or write) and what privileges are required to access the data.

Summary

A short summary of the element with no normative text.

7.2.2. Description Section

A separate section, describing a table row element, is required to include normative text, such as behavior associated with the element.

7.2.3. Other Data Table Columns

Other columns specific to the element:

Data Type

A **Data Field** requires this column for attribute, event or command data.

Quality

This is a catchall column for uncategorized qualities.

Fallback

This defines a fallback value for data fields.

Response

Cluster command tables have this column.

Direction

Cluster command tables have this column.

Priority

Event tables have this column.

Value

Enumerations use this column instead of the **ID** column.

7.3. Conformance

A Conformance column defines optionality and dependency for any data model element or set of elements. This column is valid for attributes, commands, events, enumerations, and fields of com-

mands, events or structures.

Conformance Column	Name	Summary
M	Mandatory	This is part of the base mandatory feature set and is mandatory for current revision.
O	Optional	This is a purely optional element with no dependencies, except the M set.
P	Provisional	See Provisional below.
D	Deprecated	This is a deprecated item that MAY occur in legacy implementations, but not in the current revision.
X	Disallowed	This is disallowed for the cluster derivation.
<i>AB</i>	Mandatory	This SHALL be supported if <i>AB</i> is true. <i>AB</i> is a boolean expression
[<i>AB</i>]	Optional	This MAY be supported only if <i>AB</i> is true. Brackets also act as parentheses.
<i>EF</i>	Operand	True if <i>EF</i> feature or element is supported.
<i>EF</i> == <i>x</i>	Equal	True if <i>EF</i> is equal to the non-changing value <i>x</i> .
<i>EF</i> != <i>x</i>	Unequal	True if <i>EF</i> does not equal the non-changing value <i>x</i> .
<i>x</i> <= <i>EF</i> & <i>EF</i> <= <i>y</i>	Range	True if <i>EF</i> is with in the range defined by the non-changing values <i>x</i> and <i>y</i> (inclusive).
<i>AB</i> <i>CD</i>	Or	True if either <i>AB</i> or <i>CD</i> is true.
<i>AB</i> ^ <i>CD</i>	Xor	True if only one of <i>AB</i> or <i>CD</i> is true, not both.
<i>AB</i> & <i>CD</i>	And	True if both <i>AB</i> and <i>CD</i> are true.
! <i>AB</i>	Exclusivity	True if <i>AB</i> is false.
(<i>AB</i> & <i>CD</i>)	Parentheses	Parentheses can be put around any conformance expression to combine.

Conformance Column	Name	Summary
<i>C1, C2...</i>	Otherwise	A conformance that is a list of boolean expressions, processed left to right.
<i>C.an</i>	Choice	Exclusive choices between a number of elements with the same conformance

7.3.1. Operands in Conformance

Supported operands in a conformance expression are:

- feature code
- element in the same table (e.g. attribute, field or value)
- element in the attribute table when annotated according to [Section 7.3.13, “Expressions and Optionality”](#) (e.g. attribute, field or value)
- element in the request command definition table when used as conformance for a response command definition table, when annotated according to [Section 7.3.13, “Expressions and Optionality”](#) (e.g. field or value).

7.3.2. Feature Code in Conformance

A feature code evaluates to TRUE, if the feature is supported, otherwise it SHALL evaluate to FALSE.

7.3.3. Element in Conformance

An element name, by itself, SHALL evaluate to TRUE if the element exists, otherwise it SHALL evaluate to FALSE.

NOTE

Command field as conformance SHALL only be used as element conformance when used in the context of command (request or response).

An element name MAY be used in a value conformance expression. See [Value Conformance](#) for further details.

7.3.4. Optional Conformance

Optionality with "[]" SHALL be defined for an entire expression and not for parts of an expression. Individual conformance list entries MAY define optionality with "[]", but not the entire list.

For example: The expression "[AA] & BB" is illegal, however, "[AA & BB]" or "AA & BB" is legal.

For example: The expression "AA | [BB]" is illegal, however, "[AA | BB]" or "AA, [BB]" is legal.

The tag "O" SHALL define the element as optional for the revision. "O" SHALL only be used by itself, without operators, to mean optional without dependencies, or SHALL be used in a conformance list ending in ", O", to mean otherwise optional.

7.3.5. Provisional Conformance

The tag "P" defines the element as provisional. "P" may be used in a list, where the intended conformance or list follows the "P". If the intended conformance has not been determined, then nothing appears after the "P".

It is recommended that the intended future conformance be noted, so that when the provisional marking is removed, the intended conformance becomes the current conformance.

For example: "P, M" means provisional, but mandatory, when not provisional in the future.

For example: "P, [AA & BB]" means provisional for now, but optional if AA & BB are true, when not provisional in the future.

For example: "[AA], P" means optional for AA and provisional otherwise, where the future conformance is unknown at this time.

Each provisional element SHALL be listed in a higher level specification, that includes the data model, and data model derived specifications that use this notation.

7.3.6. Mandatory Conformance

The tag "M" SHALL define the element as mandatory. "M" SHALL only be used by itself, without operators, to mean mandatory without dependencies, or SHALL be used in a conformance list ending in ", M", to mean otherwise mandatory.

7.3.7. Disallowed Conformance

The tag "X" is used when a derived cluster removes support for some elements. The tag "X" SHALL define the element as disallowed for the revision of the derivation. "X" SHALL only be used by itself, without operators.

7.3.8. Deprecated Conformance

The tag "D" SHALL define the element as disallowed for the revision. Previous revisions MAY support this element and conformance is defined in such previous revisions. "D" SHALL only be used by itself, without operators, to mean disallowed for the current revision, or SHALL be used in a conformance list ending in ", D", to mean that now it is still allowed (based on the conformance preceding the ", D"), but it is planned to be deprecated in the future after some grace period.

7.3.9. Value Conformance

Value conformance is when the conformance of an element is defined in relation to the value of another non-changing value, which MAY be the value of another element. The following operators are used to define such value relationships, and SHALL be used when the name of another element is involved, to differentiate from [Element in Conformance](#):

Operator	Name
==	Equal

Operator	Name
!=	Unequal
<	Less
<=	Equal or less
>	Larger
>=	Equal or larger

7.3.9.1. Range

The applicable operators can be used to define a range of values to be used for conformance if the value is within or outside the defined range. Such as $x \leq EF \ \& \ EF \leq y$ to indicate that the element is mandatory if the value of *EF* is between the values of *x* and *y* (both included) or $EF > x$ to indicate that the element is mandatory if the value of *EF* is larger than *x*.

A range conformance SHALL only be used when an actual range is defined; in cases where $x \leq EF \ \& \ EF \leq y$ is used and *x* and *y* are equal, this notation is NOT allowed and the equal operator SHALL be used instead.

7.3.9.2. Element

An element name SHALL only support evaluation as a value, when the element name identifies one of the following:

- an attribute with the Fixed quality, or
- a field of a struct when defined as the data type for an attribute with the Fixed quality, or
- a field in a command when referenced in a command payload.

NOTE

Using the name of a command field as a value SHALL only be used in the context of command (request or response).

For example: An attribute conformance of "AttributeX == TRUE", indicates mandatory if the value of AttributeX is TRUE.

For example: An attribute conformance of "AttributeX == 10", indicates mandatory if the AttributeX has the value 10.

For example: An attribute conformance of "AttributeX.MinValue == 10", indicates mandatory if the MinValue field of AttributeX has the value 10.

For example: A command field conformance, for an element in a response, of "ReferenceToCommandRequest. InputCommandField == EnumValue", indicates the field is mandatory in the response if the request has the InputCommandField set to the specific "EnumValue".

7.3.9.3. Revision

The revision of a cluster or device type MAY be used in a conformance expression, by using the term "Rev" in the conformance statement. The following rules applies when using revision as con-

formance:

- SHALL only occur as a value conformance,
- SHALL only occur with a "v" prefix before the cluster revision value, e.g. v2,
- SHALL only use the larger or equal ($\geq$) operator,
- SHALL NOT use a value of v1 (The initial revision).

The use of Revision conformance serves as a hint to clients as to exactly when an element was added. Elements with this conformance SHALL be treated by the server as either M or O (when the Revision conformance is True) as illustrated in the examples below; when the Revision conformance is False, this SHALL be treated as disallowed. The conformance can be combined, e.g. when an element that was optional in one revision and was changed to mandatory in a later revision.

For example: An element with conformance of "Rev $\geq$ v3", indicates mandatory if the revision is equal or higher than 3.

For example: An element with conformance of "[Rev $\geq$ v2]", indicates optional if the revision is equal or higher than 2.

For example: An element with conformance of "Rev $\geq$ v4, [Rev $\geq$ v2]", indicates optional if the revision is equal or higher than 2 and less than 4, and changed to mandatory if the Cluster Revision of the cluster is equal or higher than 4.

7.3.10. Exclusivity Conformance

Exclusivity occurs when the entire expression only excludes.

7.3.11. Otherwise Conformance

Otherwise conformance is a list evaluated from top to bottom (see [Quality Conformance](#)), and left to right, where an expression is mutually exclusive to the previous expressions (above and to the left). It is a shorthand that allows defining conformance which depends on the previously evaluated true expression.

Some examples:

- "AB, O" means mandatory for AB and optional otherwise.
- "AB, [CD]" means mandatory for AB, optional if CD is true and AB is false, otherwise not allowed.
- "!AB, O" means mandatory if AB is false, otherwise optional (if AB is true).
- "[AB], M" is the equivalent to "!AB, O", and a clearer way to define the conformance.
- "[AA], [BB], [CC]" is the equivalent to "[AA | BB | CC]".

7.3.12. Quality Conformance

To support differing quality conformance for an element, the element row is duplicated with the qualities that differ and the conformance for those qualities. Duplicating the rows requires top to bottom evaluation, similar to left to right evaluation for [Otherwise Conformance](#). However, if 2

adjacent rows evaluate to the same conformance (e.g. "AB, O" and "O" when AB is false), then the choice of qualities is optional, if the row element is implemented.

For the example below, the MinLevel element is mandatory for LT with the minimum of 1 (not zero), mandatory for AB with the minimum of 0 (zero), otherwise the element is optional with a minimum of 0 (zero).

ID	Name	Type	Constraint	Quality	Access	Fallback	Conformance
43	MinLevel	uint8	1 to 100	F	R V	1	LT
43	MinLevel	uint8	0 to 100	F	R V	0	AB, O

For the example below, the MaxLevel element is fixed, read only and mandatory for LT, writeable and mandatory for AB, otherwise the element is optional and access may be read only or writable.

ID	Name	Type	Constraint	Quality	Access	Fallback	Conformance
44	MaxLevel	uint8	1 to 100	F	R V	100	LT
44	MaxLevel	uint8	0 to 100		RW VM	100	AB, O
44	MaxLevel	uint8	0 to 100	F	R V	100	O

7.3.13. Expressions and Optionality

These rules apply to the conformance tags for a cluster, when used to indicate conformance for elements (e.g. features, attributes, fields) within a cluster specification:

- A conformance tag for a cluster MAY be the [ClusterRevision](#) evaluated as described in [Value Conformance](#).
- A conformance tag for a cluster MAY be the name of a cluster feature (see [FeatureMap](#)).
- A conformance tag for a cluster MAY be the name of an element in the Name column of the same table.
- A conformance tag for a cluster MAY be the struct field name of an element, when annotated with the name in the Name column, e.g. "<ElementName>.<FieldName>".
- A conformance tag for a cluster MAY be the name of an attribute in the attribute table, when annotated with a reference to the element.
- A conformance tag for a command field of response commands in a cluster MAY be the name of a command field in the associated request command, when annotated with a reference to the element.

These rules apply to the conformance tags for a device type, when used to indicate conformance for a cluster, cluster requirement, etc. within a device type specification:

- A conformance tag for a device type definition MAY be a condition of the node.
- A conformance tag for a device type definition MAY be the [Device Type Revision](#) evaluated as

described in [Value Conformance](#).

This rule generally apply for conformance tags:

- A conformance expression supports conformance tags as operands.

The annotations MAY be combined, e.g. "<ReferenceToElement>.<StructFieldName>" to use element conformance of a struct field from a Fixed attribute in another table.

Expressions SHALL represent a dependency with boolean logic using:

- the NOT operator such as "!AA"
- the OR operator such as "AA | BB"
- the AND operator such as "AA & BB"
- the XOR (exclusive or) operator such as "AA ^ BB"
- the equal operator such as "AA==10"
- the not equal operator such as "AA!=10"

Equality operators require that a value can be resolved for the left operand. If the left operand is not supported, the [fallback](#) is the value. "NULL" is a valid value for the equality operator.

Simple dependencies MAY also be defined in conformance. If a Max attribute has a dependency on a Min attribute, then the conformance for Max is "Min". Exclusive logic also applies. For example, if the Absolute attribute is mutually exclusive to the Percentage attribute, and one of the two must be supported, then conformance for Absolute would be "!Percentage", and conformance for Percentage would be "!Absolute".

Unless overridden with parentheses, the order of operations is:

- NOT operator "!"
- AND operator "&", OR operator "|", XOR operator "^"
- equal "==", not equal "!="

If an expression is false, there SHALL be no assumption of general optionality. If the conformance expression evaluates to false but optional, the expression is false.

Some examples of conformance definitions:

- "[AB]" means the element is not allowed, if AB evaluates to false.
- "[AA & BB]" means optional if AA and BB are both true, but excluded (and not optional) otherwise.
- "[!AA & BB]" means only optional if AA is false and BB is true.

7.3.14. Choice Conformance

Choice conformance defines requirements for implementing a set of elements at the same level. This set of elements is called the choice set.

Choice parameters determine the requirements for a choice set and are added to a conformance expression following a period, for example *C.an+*.

C is any logical optional *[AB]* conformance expression or optional conformance list, including "O" (optional), but not "M", "X", "D", or "P". If *C* is a conformance list, then the list SHALL be surrounded by parentheses. A conformance list may also include optional choice expressions among others in the list that do not support a choice.

It is invalid to use mandatory conformance for choice such as "M.a", "AB.a+" or "!AB.a-". This gives the reader the false impression that the individual element is mandatory.

The choice parameters are *a*, *n*, "+" (plus sign) and "-" (minus sign):

- *a* is a lower case letter identifying 2 or more elements in the choice set, that SHALL be at the same scope (in the same table).
- *n* determines the number from the choice set that SHALL be supported after evaluation of the conformance *C* and subject to the conditions described below. *n* SHALL only be present as a suffix to a choice set.
- + and - serve to limit the number of elements that may be included from the choice set in an implementation and SHALL only be present as a suffix to *n* or to the choice set if *n* is omitted. These choice parameters SHALL NOT appear together in the same conformance designation.

The *n*, +, and - choice parameters, if present, are interpreted as follows:

- If *n* is omitted, then *n* is considered to be one.
- If *n* appears without a plus ("+") or minus ("-") sign choice parameter, *n* represents the exact number of elements from the choice set that SHALL be included in an implementation. *n* SHALL be between 1 and the number of elements in the choice set minus 1.
- If *n* appears with a plus ("+") sign choice parameter, *n* represents that at least *n* elements of the choice set SHALL be included in an implementation. *n* SHALL be between 1 and the number of elements in the choice set minus 1.
- If *n* appears with a minus ("-") sign choice parameter, *n* represents that at most *n* elements of the choice set SHALL be included in an implementation. *n* SHALL be between 1 and the number of elements in the choice set minus 1.
- *n* can appear as a range, e.g. 2-4, without a trailing plus ("+") or minus ("-") sign, to indicate that the number of elements included in an implementation SHALL be within the indicated range inclusive. The values in the range SHALL be between 1 and the number of elements in the choice set minus 1. If a range is used, trailing choice parameters plus ("+") and minus ("-") SHALL NOT be used.

When choice conformance notation is used, each element in the choice set SHALL have identical choice parameters.

Invalid example: It is illegal to use "O.a" for one element and "O.a2" for another.

Choice sets SHALL NOT be shared between tables within the same device type or cluster specification. For example, if choice set **a** is used in the Attributes table and a second choice set is used in

another table within the cluster specification, the second set must use a different choice set identifier, such as **b**.

Invalid example: It is illegal to have one attribute marked as O.a in the attributes table and one command marked as O.a in the commands table.

Valid examples:

- "O.a" means exactly one of the "a" elements SHALL be supported.
- "O.a2" means exactly two of the "a" elements SHALL be supported.
- "O.a2+" means at least two of the "a" elements SHALL be supported.
- "O.a+" means at least one of the "a" elements SHALL be supported.
- "O.a2-" means at most two of the "a" elements SHALL be supported.
- "O.a-" means at most one of the "a" elements SHALL be supported.
- "O.a2-4" means 2, 3 or 4 of the "a" elements SHALL be supported.

Although individual element conformance such as "O.a", "O.a+", "[AB].a+" or "[AB].a2-4" allows an individual element to be optional, the restrictions on the choice set defined by the choice set parameters must be satisfied once the conformance expression is evaluated and an implementation SHALL support a number of elements as allowed by the choice set parameters. This may be a specific number of elements, or one of a range of possible numbers of elements, depending on the choice set parameters.

For example: The following is valid, because the *n* value of 1 or more is always supported.

If AB is true and CD is false then AbsoluteValue must be supported and optionally PercentValue. If AB is false and CD is true then PercentValue must be supported and optionally AbsoluteValue. If AB is true and CD is true then both elements must be supported. If AB is false and CD is false then either one or both elements must be supported.

ID	Name	Type	Constraint	Fallback	Conformance
43	AbsoluteValue	uint16	0 to 5000	0	AB, O.a1+
44	PercentValue	uint8	0 to 100	0	CD, O.a1+

Different expressions for *C* are allowed, which may limit the choices, based on conformance, to greater than zero, but less than *n*. In such cases, consideration is needed to define the choices and conformance so that the *n* value is satisfied in every valid combination.

For example: The following is invalid if conformance allows AB to be false and CD to be true, because the *n* value of 2 is not satisfied.

ID	Name	Type	Constraint	Fallback	Conformance
43	AbsoluteCm	uint16	0 to 5000	0	[AB & CM].a2
44	AbsoluteIn	uint16	0 to 2000	0	[AB & IN].a2
45	Percent	uint8	0 to 100	0	[CD].a2

7.3.15. Blank Conformance

If an element does not have a designated conformance (the column is blank or omitted), then it SHALL inherit conformance from the next highest element in the model hierarchy. For example: A data field in a struct attribute inherits its conformance from the attribute.

7.3.16. Feature Conformance

The above examples use (all capitalized) cluster feature codes in conformance expressions.

If a feature table has no conformance column then each feature is considered to be optional for that revision and therefore either true (feature supported) or false (feature unsupported).

If a feature table has a Conformance column then these are the allowed conformance:

Conformance	Feature	Meaning
M	true	This means the feature is now mandatory, and the conformance expression is true, and the feature supported (true).
O	<i>implementation</i>	This is the default conformance (if blank), which means the feature support is indicated by the implementation (true or false).
D	false	The conformance expression is false, and the feature is unsupported (false). A legacy implementation may still support the feature with a true value.
X	false	This is used in a device type that overrides cluster specification conformance, which means the conformance expression is false, and the feature is unsupported (false).
P	false	This means the conformance expression is false, and the feature is provisional and unsupported.
<i>false mandatory expression</i>	false	Conformance expressions are used to define more complex feature dependencies. A list of expressions evaluates to a single expression. Choice options only support optional expressions.
<i>false optional expression</i>	false	
<i>true mandatory expression</i>	true	
<i>true optional expression</i>	<i>implementation</i>	

- If the feature indicator (bit) is unsupported, then the feature indicator SHALL be false and the feature SHALL be unsupported (false).
- Otherwise, if the conformance expression evaluates to false, then the feature indicator SHALL be false and the feature SHALL be unsupported (false).
- Otherwise, if the conformance expression evaluates to mandatory, then the feature indicator and feature support SHALL be the value of the conformance expression.
- Otherwise:
 - the feature support (true or false) SHALL be the value or the feature indicator (bit) as implemented.
 - If the conformance expression evaluates to a [choice conformance](#), then the indicators for the choice set SHALL also conform to the choice parameters.

7.4. Element

An element of the data model is a data construct that supports an instance of data. Listed below are the elements of the data model.

First order elements

fabric, node, endpoint, cluster

Cluster first order elements

command, event, attribute

Nested elements

command field, event field, struct field

Dynamic element

list entry

Semantic elements

device type, data type

Attribute data

elements (above) that are part of an attribute

Data field

attribute, field element, or list entry (see [Data](#))

7.4.1. Encoded Element Processing

When parsing or processing encoded payloads of elements as represented by an encoding layer, such as [TLV format](#), the following general rules apply:

- Unknown elements SHALL be ignored and skipped. This provides forward compatibility with future elements.
- Elements SHALL be present when conveyed according to the element's conformance.

- Elements that are present and conformant SHALL be processed.

7.5. Fabric

A fabric is a set of nodes that interact by accessing data model elements as defined in the [Interaction Model](#). A fabric is a security domain that allows a set of nodes to be identified and communicate within the context of the domain. A node is considered to be 'on' a fabric, when it can be identified and interact in the context of that fabric. An interaction is considered to occur 'on' a fabric, when the interaction occurs in the context of that fabric (see [Accessing Fabric](#)). Each interaction occurs either on a single fabric, or without a fabric context (see [Accessing Fabric](#)).

A node MAY be identified and interact on one or more fabrics.

How a fabric is established and how a node comes to be on a fabric is not defined here and left to the lower layers.

7.5.1. Accessing Fabric

If an interaction is associated with a particular fabric, that fabric is called the "accessing fabric".

An Accessing Fabric is available for an interaction in 3 cases:

- The session is a secure session that was established via [Certificate Authenticated Session Establishment \(CASE\)](#), which causes every interaction over that session to be associated with the fabric identified by the Local Fabric Index in the [Secure Session Context](#).
- After the [AddNOC](#) command is successfully executed within a secure session established over [Passcode-Authenticated Session Establishment \(PASE\)](#), which causes every interaction over that session to thereafter be associated with the fabric created by the command.
- On receipt of a valid groupcast message, the associated fabric is the one identified by the Fabric Index stored in the [Groupcast Session Context](#).

If the interaction is not associated with a fabric, the accessing fabric does not exist. In this case any comparison of the accessing fabric to any existing fabric SHALL consider them not equal.

7.5.2. Fabric-Index

Each fabric supported on a node is referenced by [fabric-index](#) that is unique on the node. This fabric-index enables the look-up of the full fabric information from the fabric-index. A fabric-index of 0 (zero) or NULL SHALL indicate that there is no fabric associated with the context in which the fabric-index is being used. If fabric-index is used in a context that is exclusively associated with a fabric, such as fabric-scoped data model elements, then the fabric-index values SHALL NOT include 0 (zero) or NULL.

The fabric-index corresponding to the [accessing fabric](#) is called the "accessing fabric-index". If the accessing fabric does not exist, the accessing fabric-index SHALL indicate no fabric with a fabric-index of 0.

7.5.3. Fabric-Scoped Data

Most cluster data instances are accessible regardless of the accessing fabric. However, data that is exclusively associated with a particular fabric SHALL be defined as being fabric-scoped. Fabric-scoped data SHALL be defined with [fabric-scoped access](#).

The fabric associated with fabric-scoped data is called the "associated fabric".

Fabric-scoped data allows multiple accessing fabrics to manipulate a list of data items without interfering with each other. See [Section 7.19.1.8.2, "Fabric-Filtered List"](#).

Fabric-scoped data SHALL be limited to the following:

- [list of fabric-scoped structs](#), which MAY include fabric-sensitive fields
- [fabric-sensitive event](#)

A fabric-scoped data instance is always a composite struct-like data instance, with multiple fields.

Fabric-scoped data SHALL always include the [FabricIndex field](#) to indicate the associated fabric. The [FabricIndex field](#) for fabric-scoped data SHALL NOT be 0 or NULL.

Any interaction, including cluster commands, SHALL NOT cause modification of fabric-scoped data, directly or indirectly, if the interaction has an [accessing fabric](#) different than the [associated fabric](#) for the data, except in the case of a cluster command that explicitly defines an exception to this rule.

Data fields in a fabric-scoped struct MAY also have the [fabric-sensitive access modifier](#).

7.5.4. Fabric-Scoped IDs

Some data types are fabric-scoped IDs, including, but not limited to, node ID and group ID.

A fabric-scoped ID MAY require the presence of a [fabric-index](#) data type field within the same nesting scope to indicate the fabric associated with the ID in these cases:

- If the fabric-scoped ID is not part of fabric-scoped data.
- If the fabric-scoped ID is part of fabric-scoped data with an associated fabric that is not the fabric associated with the ID.

Fabric-scoped IDs SHALL only be indicated in these elements:

- [structs](#)
- [events](#)
- [commands](#)

Where necessary, specification text SHOULD define the data to which the [fabric-index](#) applies.

7.6. Access

Data model elements have access qualities. Some elements have intrinsic access or access limitations. For example: Cluster commands or command fields are not writable.

An Access column defines access to a data model element or set of elements. This column is valid for attributes, commands, events, and nested attribute data fields.

Access Column	Description
R	Read Access
W	Write Access
R[W]	Read Access and optionally Write Access
R*W	Deprecated: use R[W]
Privileges - separate with space	
V	Read Access or Invoke Access requires View privilege
O	Read Access , Write Access , or Invoke Access requires Operate privilege
M	Read Access , Write Access , or Invoke Access requires Manage privilege
A	Read Access , Write Access , or Invoke Access requires Administer privilege
Fabric - separate with space	
F	Fabric-Scoped Access
S	Fabric-Sensitive Access
Timed - separate with space	
T	Write Access or Invoke Access with timed interaction only

Attributes, commands, and events SHALL define their access, and SHALL include privileges in their access definition. For example: An attribute defines whether it is readable or writable, and what privileges are required to do so.

Attributes, commands, and events that do not define any privileges as access qualities SHALL be deemed to have the following:

- View privilege required for Read access,
- Operate privilege required for Write access,
- Operate privilege required for Invoke access for request commands.
- No privileges defined for response commands.

For example: An event with implicit read access or explicit 'R' access defaults to access 'R V'. An

attribute with access 'RW' defaults to access 'RW VO'. A request command with implicit invoke access defaults to privilege 'O'.

Nested elements MAY define their access, but SHALL NOT include privileges in their access definition. Nested elements SHALL inherit their privileges from the next highest element in the model hierarchy. Nested elements that do not define their access SHALL inherit their access from the next highest element in the model hierarchy. For example: A data field in a struct attribute inherits its access qualities from attribute.

Elements SHALL only include the lowest required privilege for a type of access.

That means:

- An event SHALL define the single privilege required for Read access.
- A command SHALL define the single privilege required for Invoke access.
- A readable attribute SHALL define the single privilege required for Read access.
- A writable (but not readable) attribute SHALL define the single privilege (that is not View) required for Write access.
- A readable and writable attribute MAY define a single privilege (that is not View) required for both Read and Write access.
- A readable and writable attribute MAY define the View privilege as required for Read access and one other privilege (that is not View) as required for Write access.

7.6.1. Read Access

Read access means that a request for data values associated with an element SHALL be supported. This access modifier SHALL only be defined for cluster event and attribute data definitions. This access modifier SHALL NOT be defined for cluster command definitions.

This is implicitly defined for cluster events and does not need to be stipulated explicitly.

7.6.2. Write Access

Write access means that a request to modify attribute data values SHALL be supported. This access modifier SHALL only be defined for cluster attribute data definitions. This access modifier SHALL NOT be defined for cluster event and command definitions.

A cluster specification SHALL define the conditions when write access attribute data is not writable, and SHALL define normative or recommended behavior to follow when this occurs.

An implementation that does not support write access for a field with optional write access SHALL have this declared in its product Declaration of Conformity.

7.6.3. Invoke Access

Invoke access means that a request to execute a command SHALL be supported. This access modifier SHALL only be defined for cluster command definitions, by defining an appropriate privilege level for the command. This access modifier SHALL NOT be defined for cluster event and attribute

data definitions.

7.6.4. Fabric-Scoped Access

This access modifier defines that a data model element is scoped to an [associated fabric](#).

When applied to a data element this access modifier defines that element as [fabric-scoped data](#).

This access modifier MAY be applied to the following data model [elements](#):

- Attributes
- Structs
- Commands

7.6.4.1. Attributes

- When this access modifier is applied to an attribute, the attribute SHALL be defined as a [fabric-scoped list](#).

7.6.4.2. Structs

- When this access modifier is applied to a struct, the struct SHALL be a [fabric-scoped struct](#).

7.6.4.3. Commands

- When this access modifier is applied to a command, it requires the command to have an associated fabric to succeed.

7.6.5. Fabric-Sensitive Access

This access modifier further restricts access to data that is sensitive to the [associated fabric](#).

This access modifier acts as an additional constraint over those imposed by the fabric-scoped access, namely:

- Fabric-sensitive data SHALL NOT be readable unless the [accessing fabric](#) is the [associated fabric](#) of the data. See [Section 7.5.3, “Fabric-Scoped Data”](#).
- A cluster command SHALL NOT return fabric-sensitive data if the [accessing fabric](#) is not the [associated fabric](#) of the data.

Data that is fabric-sensitive SHALL always be fabric-scoped.

This access modifier MAY be applied to the following data model [elements](#):

- Attributes that are lists of [fabric-scoped structs](#)
- Fields of a [fabric-scoped struct](#)
- Events

7.6.5.1. Attributes

- When applied to an attribute, this access modifier defines the attribute as a [fabric-scoped list](#) that is always read as a [fabric-filtered list](#).

7.6.5.2. Struct fields

- When applied to a field of a struct, this access modifier indicates that the field's value is only available to the struct's [associated fabric](#).

7.6.5.3. Events

- When applied to an event, the event is defined as a [fabric-sensitive event](#).

7.6.6. View Privilege

An element with the View privilege SHALL support Read (if readable) and Invoke (if invocable) access if the source of the request is granted the View privilege.

A command with the View privilege defined SHALL NOT alter data that is part of its function (e.g. settings, configuration), but MAY alter data that is internal or diagnostic in nature (e.g. usage statistics).

7.6.7. Operate Privilege

An element with the Operate privilege defined SHALL support Read (if readable), Write (if writable), and Invoke (if invocable) access if the source of the request is granted the Operate privilege.

7.6.8. Manage Privilege

An element with the Manage privilege defined SHALL support Read (if readable), Write (if writable), and Invoke (if invocable) access if the source of the request is granted the Manage privilege.

7.6.9. Administer Privilege

An element with the Administer privilege defined SHALL support Read (if readable), Write (if writable), and Invoke (if invocable) access if the source of the request is granted the Administer privilege.

7.6.10. Timed Interaction

This access modifier requires the use of a timed interaction.

Timed interactions are used to limit the amount of time an action message is valid and can interact with a node. They are used to prevent a timing attack on the system. For example, a malicious attacker could perform an "intercept, interfere, and replay" procedure whereby a legitimate message is intercepted, receipt by the intended destination is jammed, and the attacker sends the message at a later time to cause a malicious action such as unlocking a door at an unintended time.

While the practical difficulties of such an attack are high, and the malicious eavesdropper cannot decrypt the action message, the timed interaction provides further mitigation of risk for critical actions.

The timed interaction can be thought of as a 2-phase commit. A precursor action ([Timed Request Action](#)) is sent to indicate the valid time window for arrival of some subsequent, primary action. Since the timed request requires a response, an attacker cannot do the store-and-forward timing attack anymore. The lack of an authenticated response from the intended destination will prevent the subsequent primary action from being sent.

A command with this modifier SHALL require a timed invoke interaction. A writable attribute with this modifier SHALL require a timed write interaction.

An attempted untimed write interaction to a writable attribute with this modifier SHALL generate an error response.

An untimed invoke interaction for a request command with this modifier SHALL generate an error response.

7.7. Other Qualities

A Quality column defines other qualities not covered in other columns. Some qualities are limited to a specific set of elements. If an element does not have designated qualities, then it SHALL inherit qualities from the next highest element in the model hierarchy. For example: A data field in a struct attribute inherits its access qualities from attribute.

Quality Column	Name	Elements	Description
C	Changes Omitted	attribute data	Fast changing data or data where deltas are meaningless to report, and which will not cause delta changes on subscriptions
F	Fixed	attribute data	The read only value is a value that changes infrequently (see description)
I	Singleton	cluster	The cluster is a singleton on the node for the device type
K	Diagnostics	cluster	The cluster is a verbose diagnostics cluster, which could be omitted from wildcard expansion

Quality Column	Name	Elements	Description
L	Large Message	command	The command payload can be large, resulting in the message exceeding the minimum IPv6 MTU of 1280 bytes.
N	Non-Volatile	attribute data	The attribute data value is non-volatile and is persistent across restarts
Q	Quieter Reporting	attribute data	Data with fluctuating data rate or where some deltas are meaningless or undesirable to report.
S	Scene	attribute	The attribute is part of a scene
T	Atomic	attribute	The attribute requires writes to be part of an atomic interaction
X	Nullable	data fields	The data type of the data field is nullable

7.7.1. Changes Omitted Quality

This quality MAY be given to attribute data that is deemed to have a high rate of change or where changes are not meaningful or too large to convey as part of [Subscribe interaction](#).

Attribute data with this quality SHALL support [Read Access](#), but SHALL NOT have delta changes published as part of a [Subscribe interaction](#).

7.7.2. Fixed Quality

Data with this quality is read-only and SHALL only change in the following cases:

- The program image changes.
- The functionality changes due to installation or configuration changes of the device.

A [Node Configuration Change](#) that affects one or multiple attributes (on any endpoint of the node) marked with this quality, SHALL cause an increase of the value of the [ConfigurationVersion](#) attribute in the [Basic Information Cluster](#) cluster when that attribute is supported.

A [Node Configuration Change](#) that affects one or multiple attributes (on any endpoint) of a **bridged** node which are marked with this quality, SHALL cause an increase of the value of the [ConfigurationVersion](#) attribute in the [Bridged Device Basic Information Cluster](#) cluster when that attribute is supported, and SHALL also cause an increase of the value of the [ConfigurationVersion](#) attribute in

the [Basic Information Cluster](#) cluster of the bridge when that attribute is supported.

Additional considerations regarding the update of ConfigurationVersion are:

- If a configuration change on a node affects multiple attributes with this quality, the change SHOULD be considered as a single change, by only causing the [ConfigurationVersion](#) attribute to increase once.
- If the configuration change includes the addition of a new bridged node or removal of an existing bridged node, it is considered as a configuration change on the bridge itself and the value of the [ConfigurationVersion](#) attribute in [Basic Information Cluster](#) cluster, for the bridge, SHALL be increased.

7.7.3. Singleton Quality

This quality indicates that a cluster is a singleton on the node, representing the entire node.

7.7.4. Diagnostics Quality

This quality is given to Clusters which contain very large amounts of seldom-consulted attributes.

Such a cluster SHALL be omitted from [Request Path Expansion](#) when the [WildcardSkipDiagnostic-clusters](#) bit in [WildcardPathFlags](#) is set.

7.7.5. Large Message Quality

Command messages with this quality are typically larger than the minimum IPv6 MTU of 1280 bytes and cannot be transported over [MRP](#). They SHALL require [TCP](#) for communication. Both the request command and the response command SHALL be marked with this quality if at least one of them is expected to be large. This is because they both are part of the same [Exchange](#) that is bound to one underlying session.

7.7.6. Non-Volatile Quality

See [Persistence](#).

7.7.7. Reportable Quality

The [Subscribe interaction](#) supports all attribute data. This quality is supported by other interactions that only require attribute data with this quality to support interval or change reporting.

7.7.8. Quieter Reporting Quality

This quality MAY be given to attribute data that changes in a way where some of the changes are considered to have a high rate of change or where changes are not meaningful or desirable to report.

This quality can be used for fast changing attributes, where the majority of reports would have little to no value for the client (e.g. the continuous countdown of a timer in the server where the client can have a local digital twin) but where some reports may be important to synchronize the

local twin, (e.g. when a countdown starts, stops or is changed). Another case where this quality can be used is for attributes that can change value for a long period (e.g. the `CurrentLevel` of a light performing a ramp with a long transition time) or where the measured value can change quite often. This prevents over-reporting of such an attribute to many subscribed clients which could lead to link degradation on bandwidth-constrained links.

Attributes with this quality SHALL specify, in the attribute description, the details of the conditions under which changes to the attribute value SHALL, SHOULD or MAY be reported.

Changes to the value under conditions other than those specified in the attribute description SHOULD NOT be reported. To ensure clients can understand the parameters related to specifics of that attribute's reporting, the description MAY include details regarding when the attribute value might change without being reported.

Attributes MAY get this quality assigned in a revision of a cluster, to limit the amount of reporting, therefore Clients SHALL generally be resilient to more reporting than constrained by the attribute description, from older revisions of the cluster.

Attribute data with this quality SHALL support [Read Access](#).

7.7.9. Scene Quality

This quality is supported and described in the Scenes Management cluster. It relates to the ability of a field to be settable by Scenes. This quality SHALL only apply to attributes having unsigned integer or boolean data types of size at most 4 bytes from the Base Data Types list (e.g. `bool`, `uint8`, `uint16`, `uint32`), or derived types thereof (e.g. `enum8`, `map8`). This quality has no effect in contexts unrelated to the Scenes Management cluster.

7.7.10. Nullable Quality

See [Nullable](#).

7.7.11. Atomic Quality

This quality indicates that an attribute requires atomic writes. See [Atomic Writes](#).

An attribute with this quality SHALL support [Write Access](#).

7.8. Node

A Node is an addressable entity which supports the Matter protocol stack and (once Commissioned) has its own Operational Node ID and Node Operational credentials. It is a unique resource on the network that has a set of functions and capabilities that a user may recognize distinctly as a functional whole.

This distinction is usually physical, such as the physical device itself, or a logical instance of a physical device.

A node is the highest or outermost first order element in the data model. A node is the outermost

unique addressable element of the data model.

A node MAY have multiple Operational Node IDs and Node Operational Credentials, with each ID scoped to a particular fabric. When an Operational Node ID is used as the target address of an interaction, the fabric under which the ID is scoped is the [accessing fabric](#) for the interaction.

The lower layers in a communication stack supporting this data model SHALL support interactions between nodes on a logical inter-network of nodes. See [Interaction Model Specification](#) and [System Model Specification](#) that describe relationships and interactions between nodes and data model elements on each node.

It is possible for parts of a node to reside on different processors (e.g. separate application and network processors).

A physical device MAY host multiple nodes.

7.9. Endpoint

A node is composed of one or more endpoints. An endpoint is an instance of something that could be a service or virtual device as indicated by a [device type](#). An endpoint is a particular component within a Node that is individually addressable.

Each endpoint conforms to one or more [device type](#) definitions that define the clusters supported on the endpoint. Clusters are object classes that are instantiated on an endpoint.

The word 'device', depending on the context, may be used as shorthand to denote the device type definition as represented by a device type ID, a device type implementation, or an endpoint (device type instance).

There are also many examples in specification text where 'device' is used, when it would be better, and more accurate to use 'node', 'physical device', or 'product'.

The word 'device' may also be used in cluster specifications to describe application software that is supporting an instance of a cluster server or client. In this case, it would be better, and more accurate to use either 'client' or 'server'.

One must be careful to make sure there is no ambiguity when using the word 'device' in specification text, or better yet, use another word.

7.10. Cluster

Clusters are the functional building block elements of the data model. A cluster specification defines both a client and server side that correspond with each other through interactions. A cluster may be considered an interface, service, or object class and is the lowest independent functional element in the data model. Each cluster is defined by a cluster specification that defines elements of a cluster including attributes, events, commands, as well as behavior associated with interactions with these elements. Cluster attributes, events, commands and behaviors are mandatory or optional depending on the definition of the cluster. Optional items may have dependencies.

A cluster specification SHALL list one or more Cluster Identifiers. A Cluster Identifier SHALL reference a single cluster specification and SHALL define conformance to that specification. A cluster instance SHALL be indicated and discovered by a Cluster Identifier on an endpoint. A Cluster Identifier also defines the purpose of the instance.

The server cluster supports attribute data, events and cluster commands. The client cluster initiates interactions, including invocation of cluster commands.

7.10.1. Cluster Revision

The revision of a cluster is to enforce backward and forward compatibility, but still allow clusters to be enhanced, fixed, or updated, without changing the cluster's basic function.

A cluster revision SHALL be associated with an approved revision and release of a cluster specification. The revision of an instance of a cluster SHALL be represented by the global, mandatory, and read only [ClusterRevision attribute](#).

Changes to a cluster specification SHALL only augment, not modify the primary function of the cluster. Changes to a cluster specification SHALL be represented by incrementing the cluster revision. New revisions of a client cluster SHALL interoperate with older revisions of the server cluster and vice versa. Interoperability between corresponding cluster instances MAY require reading the cluster revision.

For example: If a new product client application supporting revision 3 of cluster X wishes to take advantage of the new behavior that is mandated by revision 3, then the application can read the revision of the corresponding server cluster X in each remote endpoint. If a corresponding cluster X supports revision 3 or greater, then the behavior is supported.

Examples of changes to a cluster that require incrementing the revision:

- Changing the behavior of the cluster
- Changing a read only attribute to become writable
- Adding new attributes (e.g. min and max of an existing attribute value)
- Adding new commands, actions, or behavior
- Adding one or more fields to an existing command
- Adding a new enumerated value to an attribute
- Adding qualities to an existing attribute (e.g. Q)
- Changing anything that is optional to mandatory
- Changing dependencies of optional items
- Deprecating parts of the cluster specification
- Any non-editorial specification text change

7.10.2. Cluster Optional Features

In general, as the number of optional elements in a cluster specification increases, the number of possible combinations increases, which could decrease the interoperability of that cluster.

Each cluster has a mandatory feature set that consists of mandatory elements such as attributes, commands, fields, values, dependencies, behavior, etc.

A cluster specification MAY have optional feature sets, each supported by a set of elements (see [Section 7.13.2, “FeatureMap”](#)).

There is no requirement that each cluster instance supports the same set of optional elements.

If an application knows the [ClusterRevision](#) and [FeatureMap](#) supported by a cluster instance, then it knows the exact specification text required to be implemented by that instance.

7.10.3. Cluster Data Version

A cluster data version is a metadata increment-only counter value, maintained for each cluster instance. A cluster data version represents an exact & coherent state of cluster attribute data at present. An application may externally hold a data version (called a held data version) published by a cluster instance which then represents a cluster instance state at some time in the past. An application may use a held data version to optimize future interactions, by indicating the held data version. A cluster data version is surfaced in the [Interaction Model](#) when data is requested. It is used to optimize data [read](#) transactions by reducing the need to send the same data. Write interactions may also be qualified with a held data version to disallow changes, unless the cluster instance has the same data version (see [Chapter 8, Interaction Model Specification](#)). A cluster data version is published as information in some interactions (see [Chapter 8, Interaction Model Specification](#)). An externally held data version may be included as information in some interactions (see [Chapter 8, Interaction Model Specification](#)).

A cluster data version SHALL increment or be set (wrap) to zero if incrementing would exceed its maximum value. A cluster data version SHALL be maintained for each cluster instance. A cluster data version SHALL be initialized randomly when it is first published. A cluster data version SHALL be incremented if any attribute data changes.

7.10.4. New Cluster

When considering the creation of a new cluster specification, it is recommended to consider reusing and extending an existing cluster specification. These are the mechanisms to consider, in order, to extend a cluster:

1. Optional elements: attribute data, commands, events, enumerations, etc.
2. Optional feature(s) in the [FeatureMap](#) attribute for a set of elements (see 1)
3. [Cluster Aliasing](#) to reuse a cluster specification as a whole, but with a different semantic
4. [Cluster Inheritance](#)
5. A new cluster specification

7.10.5. Cluster Aliasing

Cluster aliasing allows the reuse of approved and validated specifications and derived documents, such as test plans, scripts, etc.

- More than one Cluster Identifier, each with unique purpose and semantic content, MAY map to a single cluster specification.

For example: A Concentration Measurement cluster specification may be quite abstract but have many mapped Cluster Identifiers each with a more concrete purpose, such as CO₂ or O₂ concentration measurement.

7.10.6. Cluster Inheritance

Cluster inheritance allows the reuse of approved and validated specifications and derived documents, such as test plans, scripts, etc. This allows a new cluster specification to be defined as extending or reducing the requirements of an existing cluster specification, called the base cluster. This also allows an existing cluster specification to be defined as a derived cluster, by creating a new base cluster that is more generic, allowing potential new clusters to be derived from the new base cluster.

- A derived cluster specification MAY have mandatory requirements that are optional in the base specification.
- A derived cluster specification MAY remove requirements that are optional in the base specification.
- A derived cluster specification MAY remove or make optional a requirement that is mandatory in the base specification, if the resulting specification is deemed useful in its reduced form, and logically a subset of the base clusters.

For example: The Bridged Device Basic Information cluster is derived as a reduced form of the base Basic Information cluster, where many informational attributes are not mandatory, because the information is not available from devices behind the bridge. However, the derived cluster provides the same, but reduced, function as the base cluster.

- It is RECOMMENDED that a derived cluster makes changes to the base cluster by extending or reducing one or more features or independent elements, and not by modifying features, elements or cluster behavior.

Such modifications SHOULD be defined as a feature, which can be used to extend or reduce.

- All new features, elements or behavior introduced by the derived cluster that may be useful across other derivations SHOULD be defined in the base cluster specification and made optional (in that base cluster specification), to maintain the entire set of requirements and identifier namespace in one place.
- When a derived cluster is created, the base and derived cluster SHALL define a limited range of

element ID values allocated for the base and derived cluster, to ensure no overlap will occur. A limited range allows for sufficient use of the ID space for most use cases while allowing additional ranges to be defined in the future should these ranges become exhausted. The range SHALL consider any existing element IDs or bits in the existing clusters, but if possible the following ranges are recommended as the default values reserved for a base cluster:

- Feature bits: 0-19
- Attribute IDs: 0x0000-0x007F
- Command IDs: 0x00-0x7F
- Event IDs: 0x00-0x7F

NOTE

Readers should be aware that the Color Control, Door Lock, Level Control and OnOff clusters do not follow these recommended ranges, due to being created before these recommended ranges were defined.

- A derived cluster specification SHALL define its own revision ([ClusterRevision attribute](#)) that is independent of the base specification, with the single exception that if the base specification is modified/updated in a way that requires its revision to be increased, as detailed in [Section 7.10.1, “Cluster Revision”](#), then the revision of any derived cluster specifications SHALL also be increased.
- A base cluster specification MAY be created from an original base cluster, which then becomes a derived cluster to the newly created base cluster.

NOTE

If a cluster is derived from an already derived cluster, or a new base cluster is created for an existing cluster, it is important to ensure there is no overlap in the defined ranges of the base and derived clusters. It also means that a derived or base cluster MAY contain elements which are outside of the proposed ranges. In that case, the clusters SHALL define custom ranges which fit into the already defined elements.

If an endpoint supports multiple server clusters that derive or map to the same base cluster specification, then each SHALL represent a single implementation and operate as a single entity or instance. This makes it possible to deploy a new device endpoint with both a base and a derived cluster identifier, which SHALL remain backward compatible to legacy devices that support only the original cluster identifier. Cluster identifiers that are mapped to a single base cluster specification, but are defined for distinctly different purposes, MAY exist together on a device endpoint. If there is no base cluster identifier defined, or no base cluster identifier exists on the same endpoint, then each cluster identifier SHALL represent a separate instance.

It is a good practice to explore the possibility of either deriving a cluster from an existing cluster or creating a base cluster to map or derive new and existing cluster identifiers. See [New Cluster](#) for other options.

7.10.7. Status Codes

A cluster specification defines status code responses to actions depending on the cluster instance state. A status code is either a [global Interaction Model status code](#), or a cluster specific status code

that is unique to the cluster specification. A global status code is either scoped to the entire action, or to a cluster request path. A cluster specific status code scoped to a cluster instance is indicated by a cluster path. When an interaction defines a Status Response response, the responder SHALL return a global Interaction Model status code. When an interaction response needs to communicate a cluster specific status code, the responder SHALL return the path to the cluster instance, the global status code SUCCESS or FAILURE, and the cluster specific status code. Each cluster specific status code SHALL be associated with either SUCCESS or FAILURE, not both. A cluster specific status code SHALL be, by default, associated with FAILURE unless it is defined as associated with SUCCESS. The global SUCCESS status code means the action was executed for the request path; the global FAILURE status code means that it was not executed.

- Cluster-specific status code SHALL be defined using the [status](#) type.
- Cluster-specific status codes MAY have the same numeric values as global status codes. Interaction model messages SHALL make it clear whether a particular message field is a global status code or a cluster-specific status code.
- Cluster-specific status codes SHALL communicate more information than just a generic success or failure condition. Global status codes SHALL be used to communicate such conditions.
- A server cluster SHALL NOT return a cluster-specific code from another cluster.

7.10.8. Cluster Classification

A cluster SHALL be defined as either a utility cluster or an application cluster.

7.10.8.1. Utility Cluster

A utility cluster is not part of the primary application operation of an endpoint. It may be used for configuration, discovery, addressing, diagnostics, monitoring device health, software update, etc. It may have a temporary relationship with its cluster counterpart.

Utility cluster examples scoped to an endpoint: Identify, Descriptor, Binding, Groups, etc. Utility cluster examples scoped to the node: Basic Information, various Diagnostics clusters, OTA update related clusters, etc.

7.10.8.2. Application Cluster

An application cluster supports the primary operation of the endpoint. An application cluster supports one or more persistent application interactions between client and server.

Example application cluster transactions:

- On/Off cluster - client sends command to server
- Temperature Measurement cluster - server reports data to client

An application cluster is not a utility cluster even though it may support utility functions for itself, such as calibration, modes of operation, etc. An application cluster specification SHALL be agnostic

about layers and processes outside of its application domain.

Example: A particular temperature measurement cluster may exist on different devices, or in different networks, each with different security & commissioning mechanisms and/or policies.

Example: A commissioning cluster's domain is commissioning, but not temperature measurement.

7.11. Command

A cluster command is a set of [data fields](#), each of a [data type](#) that is conveyed between client and server cluster instances to invoke a behavior on the receiver of the command.

Each command SHALL be listed in a table with [data quality](#) columns: ID, Name, Direction, Response, Access, Conformance.

- The command table SHALL define the direction of the command as either client to server or server to client.
 - A request command is always from client to server.
 - A response command is always from server back to client, in response to a request command.
- The command table SHALL define the access privileges for each request command or omit the privileges for the default (see [default access privileges](#)).
- The command table SHALL NOT define privileges for a response command.
- The command table SHALL define a possible response to the command, if any.
- The command table SHALL define conformance for each command.

Conformance for a client to server command means the server SHALL recognize and support the client to server command and generate responses as defined. Conformance for a server to client command means the server SHALL send the command as cluster behavior defines, such as in response to a client to server command. Conformance for a command can depend on supported server features. A client SHALL NOT be required to support optional commands or commands depending on an optional feature.

A command description SHALL define when a command is generated. A command description SHALL define the effect upon receipt of a command which may be:

- a response command
- a success status response
- an error status response
- no response

A command definition SHALL clearly define any side-effects on fabric-scoped data, if applicable.

A command is identified and indicated with a command ID that SHALL be unique to the cluster^{*}.

NOTE

Some legacy clusters have reused the same command ID twice to indicate one command from the client and another from the server. Moving forward, command IDs SHALL NOT be reused in that fashion.

A cluster command table SHALL have a Response column with the following values:

Response Column	Description
N	no response is returned for this command
Y	status only is returned for this command
<i>command</i>	name of the response command to this command

A cluster command table SHALL have a Direction column with the following values:

Direction Column	Description
client ⇒ server	command is conveyed from the client to the server cluster
client ⇐ server	command is conveyed from the server to the client cluster

Each command SHALL be described in its own section with a table defining command fields (if any).

7.11.1. Command Fields

A command MAY indicate zero or more fields that are defined in a table. Each command field is defined as a row in the table with these columns:

Column	Description
ID	This is the unique field ID of the field
Name	This is the unique name of the field
Type	This is the data type of the field
Constraint	See Section 7.18.3.3, “Constraint”
Quality	See Section 7.7, “Other Qualities”
Fallback	See Section 7.18.4, “Fallback Column”
Conformance	See Section 7.3, “Conformance”

Command field conformance defines the sender requirements to include the field in a well-formed command for the revision of the cluster. A new command field or a newly made-mandatory command field in a newly revised cluster specification may be omitted by a legacy sender. The cluster

specification SHALL define clear behavior upon receipt of any possible well-formed command with fields that are not present. The cluster specification SHALL take into consideration the revision history of possible well-formed commands from legacy implementations. To allow deprecation, it is recommended that command fields have a well-defined default value (such as NULL), and associated default behavior, that is equivalent to omitting the field. Well-defined behavior, for a field that is not present, may be no behavior at all.

For example

A newly revised Noise cluster adds a new mandatory Volume field to the MakeNoise command. Legacy receivers will ignore the Volume field, and legacy senders will not include the field.

Another example

The Volume field is mandatory for the original cluster and there is a proposal to make it optional. The Volume field NULL value has the semantic of ignoring the field, so instead of making it optional, the default value is used. This would make the receiver logic simpler.

7.12. Attribute

An attribute is cluster data. Each attribute SHALL be listed in a table with [data quality](#) columns: ID, Name, (Data) Type, Constraint, other Quality, Access, Default (value), and Conformance. An attribute SHALL also define its associated semantics and behavior. Attributes reflect queryable/settable state, configuration and capabilities of a device. If no privileges are explicitly defined for an attribute, then [default access privileges](#) take effect. Attribute data MAY also have these other qualities:

Quality	Short	Description
Scene	S	indicates that the data is part of a scene
Persistent	N	indicates that the data value is persistent across restarts
Fixed	F	indicates that the read only value is a value that changes infrequently (see description)
Nullable	X	indicates that the data may have a value of NULL
Atomic	T	indicates that the attribute may only be written to in the context of an atomic write

Fabric-scoped attribute data SHALL be defined as a [fabric-scoped list](#).

7.12.1. Persistence

Persistent data retains its value across a restart.

A restart is:

- a program restart or reboot
- power cycle reboot
- user-initiated reboot
- reboot initiated from a program image upgrade

A factory reset is not such a restart. A factory reset is a deliberate behavior to reset persistent data back to its original state when the product left the factory.

Cluster attributes that represent configuration data SHALL be persistent data unless otherwise specified.

For example: a writable attribute that persistently changes the behavior (or mode) of the cluster.

Examples of non-configuration data: device state data, data that is calculated or comes from an external source, such as a sensor value, a time value, etc.

Many clusters define persistent data that is not surfaced as attributes, but is managed by commands. Commissioning or configuration data that is created to allow the cluster to perform its function is persistent data. A group table entry and binding entries are both persistent data across a restart.

When a persistent cluster attribute represents controlled state of the device, the device SHALL restore the attribute value to the value before the restart was initiated, and put the device in the state that is represented by the restored value.

For example: After an OTA cluster restart, clusters that have visible state attributes, such as the state of a light, or a window shade SHALL be persistent and define these attributes as persistent.

Some cluster specifications add a dependency with a persistent configuration attribute A that contains a value to use to restore persistent state attribute B after a restart. This is perfectly valid but cluster specific.

Cluster state data that is not controlled, such as sensor data, is not considered persistent.

The cluster specification may put dependencies and limitations on persistent data.

7.13. Global Elements

Below is a list of global elements. These are used for self-description of the server.

Table 95. Global Attributes

ID	Name	Element	Type	Con- straint	Quality	Access	Fallback	Confor- mance
0xFFFD	Cluster-Revision	attribute	uint16	min 1	F	R V		M
0xFFFC	FeatureMap	attribute	map32		F	R V	0	M
0xFFFB	AttributeList	attribute	list[attrib- id]		F	R V		M
0xFFFA	EventList	attribute						D
0xFFF9	AcceptedCom- mandList	attribute	list[com- mand-id]		F	R V		M
0xFFF8	GeneratedCom- mandList	attribute	list[com- mand-id]		F	R V		M

Table 96. Global Fields

ID	Name	Element	Type	Con- straint	Quality	Access	Fallback	Confor- mance
0xFE	FabricIndex	struct or event field	fabric-idx	1 to 254		R V F		fabric- scoped

Table 97. Global Commands

ID	Name	Direction	Response	Access	Conformance
0xFE	AtomicRe- quest	client ⇒ server	AtomicRe- sponse	O	desc
0xFD	AtomicRe- sponse	client ⇐ server	N		AtomicRequest

7.13.1. ClusterRevision Attribute

The ClusterRevision attribute indicates the revision of the server cluster specification supported by the cluster instance. An implementation of a cluster specification before the ClusterRevision attribute was added SHALL have an assumed cluster revision of 0 (zero). For a new cluster specification, the initial value for the ClusterRevision attribute SHALL be 1 (not zero).

A history of revision numbers for a cluster specification release is listed in the Revision History sec-

tion for a cluster specification. Each new revision of a cluster specification SHALL specify a new revision number incremented (by 1) from the last. The highest revision number in a cluster specification's Revision History is the revision number for the cluster specification. Therefore, a ClusterRevision attribute value SHALL be the (highest) revision number of the cluster specification that has been implemented.

7.13.2. FeatureMap Attribute

Each instance of a cluster SHALL support this attribute.

The FeatureMap attribute SHALL indicate whether the server supports zero or more optional cluster features. A cluster feature is a set of cluster elements that are mandatory or optional for a defined feature of the cluster. If a cluster feature is supported by the cluster instance, then the corresponding bit SHALL be set to 1, otherwise the bit SHALL be set to 0 (zero). All undefined bits in this attribute SHALL be set to 0 (zero).

The set of cluster elements that are designated as mandatory (M) are implicitly part of the mandatory cluster feature set, and do not have a bit in the FeatureMap attribute.

A cluster specification SHALL support this attribute if the cluster supports features. If a cluster specification is revised to support features (and so this attribute), each bit in the FeatureMap attribute SHALL have a defined default value (1 or 0), to conform with the previous revision of the cluster specification, that did not support the FeatureMap attribute. The value of 1 means the feature elements were mandatory (M) in the previous revision. The value of 0 (zero) means the elements were optional in the previous revision.

Any newly created feature set of a cluster SHALL be dependent on that cluster.

Feature sets are revision controlled as part of a cluster using the [ClusterRevision attribute](#). The cluster specification is the independent element that is revision controlled. A remote application reading the [FeatureMap](#) and [ClusterRevision attribute](#) will then know exactly what features are supported in the cluster instance.

Each feature set SHALL be well defined within the cluster specification. Each feature SHALL be mapped to a short capitalized code name for the feature set to be referenced as a [conformance tag](#) in the cluster specification text, including the Conformance columns defining the elements of the cluster.

If a cluster defines more than 32 feature sets, then it will be necessary to add another feature bitmap attribute. Any client trying to reference the new feature set will know about the new bitmap, because it knows about the new feature set(s). Legacy products will not know about the new feature set nor the new bitmap.

For a cluster whose definition which does not define a FeatureMap, the server SHALL set this attribute to 0 (zero).

See [Section 7.3.16, "Feature Conformance"](#) for details on conformance.

7.13.3. AttributeList Attribute

Each instance of a cluster SHALL support this attribute. This attribute SHALL be a list of the attribute IDs of the attributes supported by the cluster instance.

7.13.4. AcceptedCommandList Attribute

This attribute is a list of client generated commands which are supported by this cluster server instance.

Each instance of a cluster SHALL support this attribute.

This attribute SHALL be a list of the command IDs for client generated commands that are supported and processed by the server.

For each client request command in this list that mandates a response from the server, the response command SHALL be indicated in the GeneratedCommandList attribute.

If any attribute on a server supports atomic writes, this attribute SHALL contain the command ID for AtomicRequest.

7.13.5. GeneratedCommandList Attribute

This attribute is a list of server generated commands. A server generated command is a server to client command.

Each instance of a cluster SHALL support this attribute.

This attribute SHALL be a list of the command IDs for server generated commands.

For each command in this list that is a response to a client request command, the request command SHALL be indicated in the AcceptedCommandList attribute.

If any attribute on a server supports atomic writes, this attribute SHALL contain the command ID for AtomicResponse.

7.13.6. FabricIndex Field

This field SHALL be present for [fabric-scoped data](#). This field does not have to be defined explicitly in the field table for fabric-scoped data.

This field SHALL NOT be present in a write interaction. For a write interaction, the server SHALL provide the value of the [accessing fabric-index](#) as the FabricIndex field value to processing logic, after receipt of the interaction. For a [read](#) interaction this field SHALL be included in all reported data that is defined as fabric-scoped.

7.13.7. AtomicRequest Command

This command begins, commits or rolls back atomic writes. See [Section 7.15.6, “AtomicRequest”](#).

This command SHALL be supported by any cluster instance which has attributes with the [Atomic](#)

Quality.

7.13.8. AtomicResponse Command

This command is returned in response to an AtomicRequest command. See [Section 7.15.7, “AtomicResponse”](#).

This command SHALL be supported by any cluster instance which has attributes with the [Atomic Quality](#).

7.14. Event

An event defines a record of something that occurred in the past. In this regard, an event record can be thought of as a log entry, with an event record stream providing a chronological view of the events on the node.

Unlike attributes, which do not provide any edge-preserving capabilities (i.e. no guarantees that every attribute change will be conveyed to observers), events permit capturing every single edge or change and conveying it reliably to an observer. This is critical for safety and security applications that rely upon such guarantees for correct behavior.

Each cluster event is listed in a table that defines: [ID](#), [Priority](#), [Access](#), [Quality](#), [Conformance](#).

Absence of a [Quality](#) column implies that the event does not support any qualities.

Event records are [readable](#), and do not require the read access modifier to be explicitly defined.

Each generated event record SHALL have an event priority that MAY override the defined priority for that event.

Each event SHALL be described in a section that defines the purpose of the event and the data fields of the event (if any).

7.14.1. Event Record

An event record is created by the node at the time the event happens. That record SHALL have the following data fields associated with it that are common to all events:

Name	Type	Conformance
Number	event-no	M
SystemTimeStamp	stime-ms	O.a
EpochTimeStamp	posix-ms	O.a
Priority	priority	M
Data	struct	O

7.14.1.1. Number Field

This is an [event number](#) value that is scoped to the node. This number SHALL be monotonically increasing for the life of the node. This monotonicity guarantee SHALL be preserved across [restarts](#).

Between restarts, each event record SHALL be assigned a number that is exactly 1 greater than the last created event record on that Node.

When a node restarts, the event number MAY increase by a larger step than 1. *Rationale:* Nodes do not need to write every new value of the event number counter to permanent storage each time it is increased (e.g. to prevent flash wear due to many write operations). One example strategy to achieve reduction of non-volatile storage updates is described below:

1. Read the [counter](#) value at start-up.
2. Before processing any message, write [counter](#) + N to storage, where N is a carefully chosen number (e.g. 1000). This number N should be chosen carefully in order not to exhaust the lifetime 64-bit counter space.
3. Process messages normally until the [counter](#) has a value one less than the [counter](#) in storage. When this happens, store [counter](#) + N to storage.

7.14.1.2. SystemTimeStamp and EpochTimeStamp Fields

Each event record SHALL have a timestamp at the time it was created (and not when it is reported to a client). This timestamp SHALL either be [System Time in milliseconds](#) or [POSIX Time in milliseconds](#).

7.14.1.3. Priority Field

Each event record has an associated [priority](#). This priority describes the usage semantics of the event.

7.14.1.4. Data Field

Event data fields SHALL be defined in the form of a [struct](#) in a table with the following columns for each field: ID, Name, Type, Constraint, Quality, Default, Conformance.

7.14.2. Buffering

Event records SHALL be buffered on the Node, with priority given to events of a higher priority level over a lower priority level. Within a priority level, newer event records SHALL overwrite older event records. The Node SHOULD only overwrite older events if there are newer events created and there is insufficient space to retain both.

7.14.3. Event Filtering

Interactions that report event records MAY be filtered by [event ID](#) and/or [event number](#).

7.14.4. Fabric-Sensitive Event

An entire event MAY be defined as having [fabric-sensitive access](#); otherwise, it SHALL NOT be associated with a fabric.

A [read](#) interaction SHALL NOT filter event records, based on fabric, for event records that are not associated with a fabric.

A [read](#) interaction SHALL NOT report fabric-sensitive event records that are associated with a fabric different than the [accessing fabric](#).

A [fabric-sensitive](#) event SHALL include the global [FabricIndex field](#). For a [fabric-sensitive](#) event it is not required to define the [FabricIndex field](#) in the event field table.

7.15. Atomic Writes

Atomic writes allow a client to make multiple writes to certain sets of attributes atomically, such that changes are either applied entirely or none at all.

7.15.1. Atomic Write Flow

Stage	Action		Action Flow	Description
1	Invoke AtomicRequest		Client ⇒ Server	Begins atomic write
	Return AtomicResponse		Client ⇐ Server	Acknowledges beginning of atomic write or returns error
2	Repeated	Read/Write Request	Client ⇒ Server	Reads or writes to attributes
		Read/Write Response	Client ⇐ Server	Acknowledges write or errors
3	Invoke AtomicRequest		Client ⇒ Server	Commits or cancels atomic write
	Return AtomicResponse		Client ⇐ Server	Acknowledges commit of atomic write or returns error

Atomic writes have three stages:

1. The client requests to begin an atomic write by providing a list of attributes with the Atomic Write quality and a timeout for completing its expected writes
 - a. The server responds with a list of [attribute statuses](#), and possibly an alternative timeout
 - b. If the client does not accept the statuses or the alternative timeout, then the client SHALL cancel the atomic write using a [Rollback Write](#) request.
2. Otherwise, the client makes a series of write requests to the attributes
 - a. The server pends these writes in an internal buffer, potentially returning errors if a write

can not possibly succeed during a commit

3. If the client no longer wants to continue, then it requests to roll back the atomic write; otherwise, the client requests to commit the atomic write
 - a. The server evaluates all the writes, and either applies all of them or returns an error to the client, discarding the pending writes

7.15.2. Atomic Writer ID

An Atomic Writer ID is an [Operational Node ID](#) which identifies the client within the [accessing fabric](#).

If the session context for the transaction is a [Secure Session Context](#), the Atomic Writer ID SHALL be the Peer Node ID stored in the context.

If the session context for the transaction is a [Groupcast Session Context](#), the Atomic Writer ID SHALL be invalid.

An Atomic Writer ID that is not a valid [Operational Node ID](#) SHALL be invalid.

7.15.3. Atomic Write State

An Atomic Write State associates a client with a set of atomic attributes on a server. It is the combination of:

1. An endpoint-no
2. A cluster-id
3. An [Atomic Writer ID](#)
4. The [accessing fabric](#)
5. A set of attrib-ids

An attribute on a given server cluster instance is associated with an Atomic Write State if the state's set of attrib-ids contains the ID of the attribute, the cluster-id matches the ID of the cluster instance and its endpoint-no matches the endpoint of cluster instance.

A client is associated with an Atomic Write State if both its [Atomic Writer ID](#) and [accessing fabric](#) match the [Atomic Writer ID](#) and [accessing fabric](#) of the client.

If a server receives a Write Request for an attribute that is not associated with an Atomic Write State that is also associated with the client making the request, the server SHALL return the error code INVALID_IN_STATE.

Reading an attribute with the Atomic quality from a client with an associated Atomic Write State SHALL return the current value of the attribute if the client has not successfully written to the attribute; if the client has written to the attribute while having an associated Atomic Write State, the server SHALL return the updated value of the attribute.

When writing to any attribute associated with an Atomic Write State, any integrity checks SHALL be evaluated in the context of the pending values of all the attributes associated with the Atomic

Write State.

Any writes to attributes with the Atomic quality from a client associated with an Atomic Write State SHALL only be visible to that client until committed, and SHALL NOT have any effect on the operation of the server.

If a server receives a command which modifies an attribute with the Atomic quality while a client has an Atomic Write State containing the attrib-id of that attribute, it SHALL apply these changes as if there were no atomic write; this MAY cause the atomic write to fail when the client attempts to commit (see [Commit Write](#)).

7.15.4. AtomicRequestTypeEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	BeginWrite	Begin an atomic write	M
1	CommitWrite	Commit an atomic write	M
2	RollbackWrite	Rollback an atomic write, discarding any pending changes	M

7.15.5. AtomicAttributeStatusStruct Type

This struct indicates the status of an attribute during an atomic write.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	AttributeID	attrib-id					M
1	StatusCode	status					M

7.15.5.1. AttributeID Field

This field SHALL indicate the ID of the attribute with the associated [StatusCode](#).

7.15.5.2. StatusCode Field

This field SHALL indicate the atomic status of an attribute.

7.15.6. AtomicRequest Command

This command is used to begin, commit or cancel an atomic write.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Request-Type	AtomicRequestTypeEnum				M
1	AttributeRequests	list[attrib-id]				M
2	Timeout	uint16				desc

7.15.6.1. RequestType Field

This field SHALL indicate the type of atomic request being made by the client.

7.15.6.2. AttributeRequests Field

This field SHALL indicate a list of attribute IDs the client wishes to write to during an atomic request.

7.15.6.3. Timeout Field

This field SHALL indicate an interval in milliseconds in which the client will either commit or roll-back the atomic write. If the value of the RequestType field is BeginWrite, this field SHALL be included, otherwise it SHALL be omitted.

7.15.6.4. Effect on Receipt

If an AtomicRequest command is received without an accessing fabric, the server SHALL return a status code of INVALID_COMMAND.

If an AtomicRequest command is received with an [Atomic Writer ID](#) which is either unavailable or not a valid [Operational Node ID](#), the server SHALL return a status code of INVALID_COMMAND.

For AtomicRequests with any RequestType:

1. If the [AttributeRequests](#) field is empty, the server SHALL return an error code of INVALID_COMMAND.
2. If the [AttributeRequests](#) field contains duplicate values, the server SHALL return an error code of INVALID_COMMAND.
3. If the [AttributeRequests](#) contains attrib-ids for attributes not supported by the server, the server SHALL return an error code of INVALID_COMMAND.

Begin Write

If the RequestType of the AtomicRequest is BeginWrite, the server SHALL attempt to begin an atomic write and return an AtomicResponse:

1. If the client is already associated with an [Atomic Write State](#) that is associated with any attributes on the same cluster and endpoint, the server SHALL return a status code of INVALID_IN_S-

TATE.

2. If the request does not have a [Timeout](#), the server SHALL return an error code of INVALID_COMMAND.
3. Otherwise, the server SHALL return an [AtomicResponse](#):
 - a. For each attrib-id in [AttributeRequests](#), the server SHALL return a matching [AtomicAttributeStatusStruct](#) in [AttributeStatus](#):
 - i. If the client does not have sufficient privilege to write to the attribute, the [status](#) SHALL be UNSUPPORTED_ACCESS.
 - ii. If the specified attribute does not support atomic writes, the [status](#) SHALL be INVALID_COMMAND.
 - iii. If the specified attribute is currently being used by a different atomic write, the [status](#) SHALL be BUSY.
 - iv. If the server does not have enough resources to pend writes to the attribute, the [status](#) SHALL be RESOURCE_EXHAUSTED.
 - v. Otherwise, the [status](#) SHALL be SUCCESS to indicate that the attribute is available for writing.
 - b. For internal reasons, a server MAY implement atomic writes for a set of attributes with a shared resource, such that beginning an atomic write to any of the attributes in the set requires including all of the attributes in the set. If attributes other than the requested attributes are included in an atomic write, the server SHALL return statuses for those attributes as if they had been requested.
 - c. Subsequent invocations of [AtomicRequest](#) of type [Commit Write](#) or [Rollback Write](#) SHALL provide the attrib-ids returned in [AttributeStatus](#), not the attrib-ids requested in [AttributeRequests](#).
 - d. If any of the requested attributes have a [status code](#) other than SUCCESS, the [StatusCode](#) of the [AtomicResponse](#) SHALL be FAILURE.
 - i. The server SHALL NOT associate an [Atomic Write State](#) with the client.
 - ii. Depending on the [status codes](#) received for individual attributes, the client MAY choose to try again; e.g. if the attribute the client was attempting to write to has status BUSY, the client could attempt to write to it again after a pause.
 - e. Otherwise, the [StatusCode](#) SHALL be SUCCESS.
 - i. The server SHALL create an [Atomic Write State](#) with the endpoint number and ID of the cluster instance handling the [AtomicRequest](#) command, the [accessing fabric](#), the [Atomic Writer ID](#), and the set of attrib-ids returned.
 - ii. If the requested timeout is longer than the maximum timeout the server allows for the requested set of attributes, the server MAY return the maximum timeout for the requested set of attributes in the [Timeout](#) field. Otherwise, the server SHALL return the requested timeout in the [Timeout](#) field.
 - iii. If the server does not receive a matching [AtomicRequest](#) with a [RequestType](#) of [CommitWrite](#) from the associated client before the timeout provided by the server in the [AtomicResponse](#), the server SHALL roll back any pending writes and discard the atomic

write.

Commit Write

If the RequestType of the AtomicRequest is CommitWrite, the server SHALL attempt to commit any pending writes and return an AtomicResponse:

1. If the client is not associated with an [Atomic Write State](#) whose endpoint-no and cluster-id match the endpoint number and ID of the cluster instance handling the request, and whose set of attrib-ids match the attrib-ids requested, irrespective of order, the server SHALL return an error code of INVALID_IN_STATE.
2. The server SHALL process all pending writes to the requested atomic attributes as if they had arrived in a single message.
3. The server SHALL create a list of [attribute statuses](#) for each attribute that would be modified by the processing of the atomic write:
 - a. If the server is able to determine that the atomic write will not succeed, the [status code](#) SHALL indicate the error code that would have been received on the first failed write to that attribute.
 - i. For example, if a write is an attempt to replace a list, but the data version has changed since the atomic write began due to a command invocation, the [status code](#) would be DATA_VERSION_MISMATCH.
 - b. Otherwise, the [status code](#) SHALL be SUCCESS.
4. If the [status code](#) of any [AtomicAttributeStatusStruct](#) in the list is not SUCCESS, then the server SHALL:
 - a. Discard all pending writes to the attributes associated with the [Atomic Write State](#) associated with the client
 - b. Discard the [Atomic Write State](#) associated with the client
 - c. Return an [AtomicResponse](#) with the list of [attribute statuses](#) and an [atomic status code](#) of FAILURE.
5. Otherwise, if the writes would collectively violate a constraint, then the server SHALL:
 - a. Discard all pending writes to the attributes associated with the [Atomic Write State](#) associated with the client
 - b. Discard the [Atomic Write State](#) associated with the client
 - c. Return an [AtomicResponse](#) with the list of [attribute statuses](#) and an [atomic status code](#) of CONSTRAINT_ERROR.
6. Otherwise, the server SHALL:
 - a. Create a list of [attribute statuses](#) for each attribute referenced by a pending write
 - b. Attempt to apply all pending writes to the attributes in the [Atomic Write State](#) associated with the client:
 - i. For each pending write, write the pending value to the attribute, and store the returned status code in the [status code](#) field on the [attribute status](#) whose [AttributeID](#) matches the ID of the attribute.

- ii. If any pending write fails, return an **AtomicResponse** with an **atomic status code** of FAILURE and the list of **attribute statuses**.
 - iii. Otherwise, return an **AtomicResponse** with an **atomic status code** of SUCCESS and an empty list of **attribute statuses**.
- c. Discard the **Atomic Write State** associated with the client

Rollback Write

If the RequestType of the AtomicRequest is RollbackWrite, the server SHALL attempt to rollback any pending writes and return an AtomicResponse.

1. If the client is not associated with an **Atomic Write State** whose endpoint-no and cluster-id match the endpoint number and ID of the cluster instance handling the request, and whose set of attrib-ids match the attrib-ids requested, irrespective of order, the server SHALL return an error code of INVALID_IN_STATE.
2. Otherwise, the server SHALL:
 - a. Discard all pending writes to the attributes associated with the **Atomic Write State** associated with the client
 - b. Discard the **Atomic Write State** associated with the client
 - c. Return an **AtomicResponse** with an **atomic status code** of SUCCESS.

7.15.7. AtomicResponse Command

The AtomicResponse command is sent by a server in response to an AtomicRequest command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	StatusCode	status				M
1	AttributeStatus	list[AtomicAttributeStatusStruct]				M
2	Timeout	uint16				desc

7.15.7.1. StatusCode Field

This field SHALL indicate the status of the atomic write. See [Section 7.15.6, “AtomicRequest”](#).

7.15.7.2. AttributeStatus Field

This field SHALL indicate a list of statuses for the attributes involved in the atomic write. This list SHALL contain a status for each of the attribute IDs provided in the AtomicRequest. See [Section 7.15.6, “AtomicRequest”](#).

7.15.7.3. Timeout Field

This field SHALL indicate an effective timeout provided by the server. If the value of the Request-

Type field in the AtomicRequest is BeginWrite, this field SHALL be included, otherwise it SHALL be omitted. See [Section 7.15.6, “AtomicRequest”](#).

7.16. Device Type

In this architecture model, a device type is the highest semantic element. A device type defines conformance for a set of one or more endpoints. A device type defines a set of requirements for the node or endpoint in the market.

A device type SHALL define the cluster support for an endpoint. A composed device type MAY define one or more other device types as part of the composed device type.

A device type definition MAY define or use predefined conditions from requirements, limitations and/or capabilities of the node. A device type definition MAY define or use predefined conditions on one or more underlying stack standard(s).

A device type MAY assert conditions defined in itself or in other device types. See [Section 9.2.6, “Asserting Conditions Among Device Types”](#) for details.

A device type MAY define support of a cluster as dependent upon a condition. A device type definition MAY specify optional clusters that are recommended as enhancements.

A device type definition MAY refine cluster conformance:

- Support of optional cluster elements or features MAY be changed to mandatory depending on device type conditions.
- Support of optional cluster elements or features MAY depend on device type conditions.

A device type definition SHALL NOT alter the choice sets of cluster elements in a way that would be incompatible with the cluster definition.

Examples of choice set changes:

Cluster Choice Set	Change by Device Type Def.	Allowed?	Reason
O.a	O.a+	No	Client expects mutual exclusion
O.a+	O.a	Yes	Subset of client expectations
O.a3+	O.a4+	No	Superset of client expectations
O.a4+	O.a3+	Yes	Subset of client expectations

A device type definition SHALL specify a device type ID, device revision, and a set of one or more mandatory clusters including each cluster’s minimum revision.

A device type definition MAY be generic and allow many similar clusters, where at least one

instance SHALL be required.

For example: a simple sensor device.

If all sensor devices are common in cluster requirements (except the clusters that perform the sensing), then there is no reason to create a device type for each sensor cluster.

A device type definition MAY be very specific and list particular clusters as mandatory.

For example: a door lock device or thermostat.

7.16.1. Device Type Revision

A device type revision is an unsigned integer that is associated with an approved revision and release of a device type definition. The initial value for a device type revision SHALL be 1. The initial revision (1) of a device type definition SHALL require the latest (at the time of definition the cluster) certifiable revisions of the clusters it mandates. Any mandatory changes to the device type definition SHALL only augment, not modify, the function of the device. Any changes SHALL increment the version of the device. Newer versions of the device SHALL interoperate with older revisions at the older revision's level of functionality.

Examples of changes to a device type definition that require incrementing the revision:

- Mandating a higher revision of one or more mandatory clusters
- Changing an item from optional to mandatory
- Deprecating parts of the device type definition

7.16.2. Device Type Composition

A device type definition MAY be a composed device type and therefore require other device types for its composition. A device type instance MAY be composed of other endpoints that support extra cluster instances. See [Chapter 9, System Model Specification](#) for more details.

7.16.3. Device Type Classification

Each device type definition SHALL specify the endpoint as being either a Utility, or Application. Each device type definition SHALL specify the scope as either endpoint or node. Each Application device type definition SHALL specify the endpoint as being either Simple or Dynamic.

7.16.3.1. Utility Device Type

A Utility device type supports configuration and settings. A utility device type definition SHALL define requirements for utility clusters. A utility device type MAY also represent the physical device or product. There MAY be more than one endpoint supporting a utility device type on a node.

Examples of utility device types: Root Node, Power Source, Bridged Node.

7.16.3.2. Application Device Type

Application devices types are typically the most common endpoints on a node and in the network. An endpoint supporting an application device type is an application endpoint. An Application device type SHALL be scoped to the endpoint. An application endpoint SHALL support clusters the primary application function of the endpoint.

Examples of application device type categories: HVAC, lighting, sensing, etc.

7.16.3.3. Simple Device Type

A Simple device type supports local control that is persistent, independent, and unsupervised. A Simple device type is an Application device type. Simple devices types are typically the most common endpoints in the network. Simple device type examples: sensors, actuators, lights, on/off switches, on/off power outlets, etc. Simple endpoints support independent operation without central control or gateways. An endpoint supporting a simple device type is a simple endpoint. Simple endpoints SHALL support relationships through bindings.

7.16.3.4. Dynamic Device Type

A Dynamic device type supports intelligent and supervisory services, such as commissioning, monitoring, trend analysis, scheduling and central management. A dynamic device type is an application device type. An endpoint supporting a dynamic device type is a dynamic endpoint. A dynamic endpoint is typically found on a central controller where there exists an intelligent supervisory application that manages simple device control applications. Typically, a dynamic endpoint supports client clusters for central control, management, monitoring or supervisory functions. Typically, the product supporting a dynamic endpoint has visibility into the entire network (or part thereof) of simple endpoints.

A dynamic endpoint client cluster instance MAY be used to multiplex transactions to or from multiple simple device server clusters in the network. A dynamic endpoint client cluster MAY initiate interactions to many server clusters in the network. A dynamic endpoint client cluster MAY receive data from many server clusters in the network. Dynamic endpoints MAY support relationships through bindings. A dynamic device endpoint MAY support one or more external agents, outside the node stack, that manage relationships. External agents include, but are not limited to, a cloud application, a smartphone, an in-home display, or a configuration tool.

7.16.3.5. Device Type Scope

A node device type is a utility device type scoped to a node. A node device type definition SHALL support clusters that represent the entire node. An endpoint supporting a node device type is a node endpoint. A node endpoint MAY also represent the physical device or product. There MAY be more than one node endpoint on a node.

Other classes of device types are endpoint scoped device types.

7.16.4. Extra Clusters on an Endpoint

An endpoint MAY support extra clusters not mandated by the device type definition. An endpoint MAY support optional features or cluster items (attributes, commands, events, etc.), that are not mandated by the device type definition. Extra clusters, features, or cluster items, SHALL only augment, not modify, the function of the device type or clusters.

A device type definition MAY list clusters as optional, which implies they can augment the device type. Clients MAY use those clusters to augment the operation of the other clusters for the device type.

In case of clusters appearing on an endpoint which are not listed in the Cluster Requirements table of a device type definition, there may not be a well-defined standard interpretation with respect to the context of that device type, so clients MAY ambiguously interpret how, or when, the extra clusters are intended to be used. Clients MAY ignore any clusters which are not listed in the Cluster Requirements table of a device type definition when interacting with the endpoint. This is true even for clusters that would otherwise appear to a client's or server's implementer as "obvious".

7.16.5. Primary Device Type

The Primary Device Type is selected according to the following rules:

- When a device combines multiple application device types, the manufacturer SHALL choose the application device type identifier of the primary function of the device as the Primary Device Type.
- When a device has a single application device type, the manufacturer SHALL use that as the Primary Device Type.
- For devices exposing multiple application device types for which the [superset device type](#) condition holds, the superset device type of the set of device types SHALL be used as Primary Device Type.

The Primary Device Type is used in the DCL ([field DeviceTypeID](#)), in the [Certification Declaration](#) (field `device_type_id`) and in announcements ([TXT key for device type](#)). Since the DCL field [DeviceTypeID](#) is not mutable, the Primary Device Type cannot change over the lifecycle of a product (e.g. via a software update).

Example: A desk lamp contains a light (Extended Color Light) and a switch (Generic Switch) to control the light. The Extended Color Light matches the primary function of the device, so this is selected as Primary Device Type.

Example: A dimmable light bulb exposes Dimmable Light and On/Off Light device types. Since Dimmable Light is the superset device type, this is selected as Primary Device Type.

7.17. Non-Standard

This architecture model provides mechanisms for non-standard or manufacturer specific items such as clusters, commands, events, attributes and attribute data fields. These items MAY be supported on a certified product. Such vendor specific items SHALL NOT change the standard behavior

of the standard items. The specific function of a vendor specific item cannot be tested as part of certification. They can only be tested to verify that they do no harm, and conform to proper behavior with regard to identification, discovery, error processing, etc. A non-standard item **SHOULD NOT** take the place of a standard item that provides the same function. It is up to the certification authority to make a judgment call that is in keeping with the spirit of these requirements. Implementers are encouraged to develop and certify standard items, not non-standard items.

7.18. Data Field

A data field is any attribute, field or entry that is not a collection data type, or data that is not surfaced as an attribute, but defined in a cluster specification.

Optional attribute data **MAY** be referenced as data fields in other attribute specifications within the same cluster specification. Cluster specifications also define data fields that are not surfaced, such as temporary calculated values, or persistent state values. Any defined data value in a cluster specification is a data field.

Each cluster data field **SHALL** be defined with a table including these columns for data qualities:

- [Data Type](#)
- [Constraint](#)
- [Quality](#)
- [Access](#)
- [Fallback](#)
- [Conformance](#)

A data field **SHALL** inherit (if possible) the qualities from the [cluster first-order element](#) of which it is part, unless overridden. It **SHOULD** be rare to override inherited qualities.

For example: If an attribute is a struct data type, that is readable and writable, then all fields of the struct are readable and writable.

New or optional data fields might not be recognized by a receiver, such as a legacy receiver. The data field description **SHALL** define default behavior (such as absence of behavior) when a new or optional data field is not present. It is recommended to define or use a feature when adding new or optional data fields, to better indicate conformance. It is recommended to define a [fallback value](#), such as a NULL value, that indicates such default behavior.

7.18.1. Nullable

When a data field value is required to designate an unknown, invalid, or undefined data value, and there is no obvious data value (e.g. zero), that is within the valid range to indicate this, the data field **MAY** be designated as nullable, so that an implemented instance of data **MAY** have the value of NULL.

In this context, these wordings have the same meaning:

- The data field has the value of null.
- The data field has the null value.
- The data field is the null value.
- The data field is null.
- The data field has the value of NULL.
- The data field has the NULL value.
- The data field is the NULL value.
- The data field is NULL.

Representation of NULL for the implementation of the Data Model is a consideration of the underlying encoding specification. The encoding layer SHALL have the capability to indicate NULL for any nullable data field. How the encoding layer indicates NULL is outside the scope of the Data Model specification.

All data fields MAY be defined to be nullable, regardless of data type.

A cluster specification SHALL define whether a data field is nullable. A cluster specification SHALL define the meaning of the NULL value.

Composite data types that have a length (i.e. octet string and list), and derived types that have those as the base type, SHALL NOT differentiate semantically between the NULL value and the empty (zero length) value. In particular a zero-length value SHALL be allowed for nullable values of these types no matter what other length constraints are imposed on the value, and SHALL have the same semantics as the NULL value.

A fabric-scoped list SHALL NOT be nullable.

7.18.2. Optional or Deprecated

An optional or deprecated data field that is not implemented, and therefore does not exist, SHALL NOT be indicated as the NULL value. How the encoding layer encodes non-existent data is outside the scope of the Data Model specification.

The Conformance column SHALL define if a data field is optional or deprecated. To manage the data identifier namespace, a deprecated data field SHALL NOT be removed from text that lists its identifier and default value. The description text of a deprecated data field SHALL be removed for new revisions of specification text.

If the specification text of a cluster depends on the value of an optional or deprecated data field of the same cluster, then the data field SHALL have a well-defined default value that SHALL be used when the data field is not implemented.

7.18.3. Constraint & Value

The tables below describe the nomenclature for describing constraints and default data values. This nomenclature is used in the cluster specifications for data value constraints, defaults, and other definitions.

7.18.3.1. Common Literal Values

These values are commonly used in cluster text and Default columns in cluster data definition tables.

Value	Description
0	The numeral zero is used to indicate the zero value for analog data. This is equivalent to the boolean value FALSE.
1	This is used for any analog data type to mean that the value is 1. This is equivalent to the boolean value TRUE.
FALSE	"FALSE", "false" or "False" is a boolean value and is equivalent to 0 (zero).
TRUE	"TRUE", "true" or "True" is a boolean value and is equivalent to 1.
NaN	Not a Number defined for any floating point values.
NULL	This indicates the value of NULL, and is the preferred presentation of the NULL value in this specification.
null	This indicates the value of NULL. It is the equivalent of the NULL value but is not the preferred presentation of the NULL value. Note that this definition of null will be removed from a future version of this specification.
empty	This indicates empty list or string data.
min	The minimum possible data value for the data type.
max	The maximum possible data value for the data type.
<i>numeric units</i>	Some number in some well-defined units as described in the data type (e.g. 100° C)

7.18.3.2. Common Expressions

The following expressions can be used to define values used for constraints. The result of the expression is a single numerical value and is be used to define constraints as described in [Constraint](#).

Expression	Description
<code>minOf(x, y, ...)*</code>	Use the smallest value of x, y or other provided values in the dataset.

Expression	Description
<code>maxOf(x, y, ...)*</code>	Use the largest value of x, y or other provided values in the dataset.

* x, y, or other elements used in the expressions SHALL be a numerical value.

7.18.3.3. Constraint

The Constraint column is valid for any composed device type, attribute or data field of an attribute, event, command or struct. It is RECOMMENDED to always define a constraint for any data field.

Constraint	Description
desc	Defines the constraint is defined in the description section
Composed Device Type Constraints	
<code>x*</code>	Defines the exact number of endpoints, greater than zero, supporting this device type as a part of the composed device type
<code>x to y</code>	Defines the allowed number of endpoints, greater than zero, supporting this device type as a part of the composed device type
<code>max y</code>	Defines the maximum number of endpoints, greater than zero, supporting this device type as a part of the composed device type
<code>min x</code>	Defines the minimum number of endpoints, greater than zero, supporting this device type as a part of the composed device type
Numeric Data Type Constraints	
<code>x</code>	Defines a value that is supported.
<code>x to y</code>	Defines a supported value range.
<code>max y</code>	Defines the value range from min to y
<code>min x</code>	Defines the value range from x to max.
<code>all</code>	Defines that all values are supported. Same as "min to max".
<code>constraint, constraint...</code>	Defines support of a union of two or more value and value range constraints
Octet String Data Type Constraints	
<code>x*</code>	Defines the size in bytes to be exactly x
<code>x to y</code>	Defines that the size range supported is from x to y bytes, inclusive

Constraint	Description
min <i>x</i>	Defines that the size limit supported is a minimum of <i>x</i> bytes, inclusive
max <i>y</i>	Defines the size limit supported is a maximum of <i>x</i> bytes, inclusive
all	Defines no constraint on size of the string. Same as "min to max".
<i>constraint, constraint...</i>	Defines support of a union of two or more size ranges or size limit constraints
List Data Type Constraints	
<i>x</i> *	Defines the number of entries to be exactly <i>x</i>
<i>x</i> to <i>y</i>	Defines the range for the number of entries supported is from <i>x</i> to <i>y</i> , inclusive
min <i>x</i>	Defines the limit for the number of entries supported is a minimum of <i>x</i> entries, inclusive
max <i>y</i>	Defines the limit for the number of entries supported is a maximum of <i>x</i> entries, inclusive
all	Defines no constraint on the number of entries in the list. Same as "min to max".
<i>constraint, constraint...</i>	Defines support of a union of two or more range or limit constraints on the number of entries.
<i>list_constraint</i> [<i>entry_constraint</i>]	Defines <i>list_constraint</i> as a list constraint and <i>entry_constraint</i> as a constraint on the entry data type. See also list entry qualities
Character String Data Type Constraints	
<i>char_constraint</i> { <i>z</i> }	Defines <i>char_constraint</i> as the string constraint in bytes and <i>z</i> as the maximum number of Unicode codepoints.
Enumeration Data Type Constraints	
all	Defines that all enumeration values that are defined for the enumeration type definition are supported for interactions on the element having this constraint.
Boolean Data Type Constraints	
all	Defines that both logical values are supported for interactions on the element having this constraint.

* *x*, *y*, or *z* are literal values of the data type or from the [Common Literal Values](#).

Valid examples of constraints with expressions are:

- "minOf(5, elementValue)" defines the minimum of either the value 5 or the value of the element named elementValue will be used within the constraint expression.
- "min maxOf(10, elementValue)" defines the maximum of either the value 10 or the value of the element named elementValue will be used to indicate the minimum value for the constraint.
- "1 to maxOf(15, elementValue)" defines the supported range for the element to be between the value of 1 and the maximum value of either the value 15 or the value of the element named elementValue.

7.18.3.4. Numeric Constraint

The constraints of a Numeric Data Type can be defined by the value of other elements, e.g. an element constrained by minimum and maximum values defined by the associated min/max elements (attribute or data field).

These rules apply to the constraint when elements are used as constraints:

- A constraint for an element MAY use the name of an element in the Name column of the same table.
- A constraint for an element MAY use the struct field name of an element of the same table, when annotated with the name of the constraining element from the Name column, e.g. "<ConstrainingElementName>.<Field>".
- A constraint for an element MAY use the name of an attribute in the attribute table, when annotated with a reference to the element.
- A constraint for an element for command field of response commands in a cluster MAY use the name of a command field in the associated request command, when annotated with a reference to the element.

7.18.3.5. List and String Constraints

The constraints for Lists and Strings SHALL indicate the capability of the server with respect to the number of entries of the associated List or size of the associated String.

The minimum number of entries for a list or the size of a string SHALL be 0 (zero), unless redefined using the above notation.

A **range** constraint for a List or String defines the number of entries or size, inclusive, for the element that SHALL be supported by the server. Clients writing to or otherwise affecting changes in elements with a range constraint MAY supply a list of any length or string of any size provided that the length of the list or size of the string are between x and y inclusive.

A **limit** constraint for a List or String defines a single-ended limit to the number of entries or size, inclusive, for the element that SHALL be supported by the server.

Notes on clients writing to or affecting changes in elements with a limit constraint:

- *min*: Clients writing to or otherwise affecting changes in elements with this constraint MAY supply fewer elements than the constraint value. For example, a client writing to the a list attribute with a constraint of *min 4* MAY write a list with the number of items being between 0 (zero) and

4 inclusive, or more than 4 provided that (the client has knowledge that) the server supports more than 4 items.

- *max*: Clients writing to or otherwise affecting changes in elements with this constraint MAY supply fewer elements than the constraint value. For example, a client writing to a string attribute with a constraint of *max 16* MAY write a string of any size between 0 (zero) and 16.

A **union** constraint for a List or String is a comma-delimited set of constraints. A union constraint SHALL only have one minimum (min x) constraint and one maximum (max y) constraint. A union constraint SHALL NOT define a range below the minimum constraint or range greater than the maximum constraint, including the defined [minimum](#) (min) and [maximum](#) (max) for the data type.

A constraint on a list or string data means that the data SHALL always be indicated within that constraint. A constraint on a [writeable](#) list or string data means that the data SHALL support writing within the constraint, and SHALL NOT support writing outside the constraint.

The actual number of list entries or string size for an element MAY differ from the capability as determined by the constraints but SHALL NOT violate the constraints for the element.

For example, a string attribute may be defined with a constraint of 5 to 7:

- A value of "bear" would violate the constraints as the size of 4 is less than the lower end of the constraint range, which is 5.
- A value of "bison" would be acceptable as the string size of 5 is the lower end of the range of the constraint.
- A value of "beaver" would be acceptable as the string size of 6 is between the lower and upper range of the constraint.
- A value of "buffalo" would be acceptable as the string size of 7 is the upper end of the range of the constraint.
- A value of "bontebok" would violate the constraints as the size of 8 is more than the upper end of the constraint range, which is 7.

7.18.3.6. Effective Maximum for Character String Data Type

A server SHALL support up to the maximum in [char_constraint](#) for a character string data type. The character string data SHALL NOT contain more than z Unicode codepoints.

Example: A string with a constraint of "max 128{32}" dictates that the server provide for a 128 byte string, but the string may contain up to 32 Unicode codepoints.

7.18.3.7. Nullable in Range

If data is [nullable](#) then NULL SHALL be a valid value.

If the data type is a list or derived from a list, and the list is nullable, then a length of 0 (zero) SHALL be supported, and defined in the constraint column.

If the data type is an octet string, or derived from an octet string (e.g. character string), and the data is nullable, then a length of 0 (zero) SHALL be supported, and defined in the constraint column.

7.18.3.8. Enumeration Constraint

When the constraint for an element is 'all', if any attempt is made in an interaction to use a value that is not defined in the data type definition associated with the element, the result for the interaction SHALL be a constraint violation with the response defined by the actual interaction, e.g. an attribute write or a cluster command.

NOTE

Values that are well-defined could still result in a constraint error or other violation due to features unsupported by a server, the current system state, or other reasons associated with the particular device.

7.18.4. Fallback Column

A fallback value defined in this column is not meant to be the value used when the server returns to factory default settings. Specified conformance for data fields may be optional or change over time. A fallback value is defined to complete dependencies when the actual data field value is not present.

A data field SHALL have a defined fallback value when:

- the data field is new, and a fallback is required for backwards compatibility with legacy instances
- the data field is optional, deprecated, or obsolete and therefore is not always present
- an initial value is needed before the application starts
- the value cannot be determined by the application for the instance
- there is a dependency on the attribute value to formulate other data or affect behavior

If a fallback value is not defined for a data field, the fallback value is determined by the following conditions:

- If the data field is nullable then the fallback value SHALL be NULL
- Else the fallback value SHALL be one of the following, depending on type:
 - Boolean: false
 - Analog: 0 or 0.0, depending on range
 - Bitmaps: 0
 - Enumeration: MS
 - Composite:
 - String: empty
 - List: empty
 - Struct: fallback is recursively composited from the fallback of its member fields

- Derived types: use the fallback value of the base type

These are the options for the Fallback column used for attributes or attribute, command or event data:

Fallback Column	Description
x	a literal value x of the data type, or as defined in Common Literal Values
MS	a manufacturer or implementation specific value

If the fallback value of a data field is specified as manufacturer specific, then there SHALL be no defined fallback value and the application SHALL support a manufacturer specific value that is in the valid range.

7.19. Data Types

Each data field in a cluster specification SHALL have a well-defined data type. Each attribute in a cluster specification SHALL map to a single data type.

The table indicates for each data type whether it defines an analog or discrete value. Values of analog types MAY be added to or subtracted from other values of the same type and are typically used to measure the value of physical properties that can vary continuously over a range. Values of discrete data types only have meaning as individual values and SHALL NOT be added or subtracted.

Some data types specify bit-widths for future potential growth in range (analog) or number of values (discrete).

Cluster specifications SHALL use the data type short name to reduce the text size of the specification.

Data types can be defined in three locations:

1. Cluster specification.
2. Cluster category sections.
3. Data Model.

The data types defined in the Data Model are the shareable data types and MAY be used within any cluster specification. The data types defined in the cluster category sections are the shareable data types within that specific cluster category and SHALL only be used by clusters within that cluster category. The data types defined in a cluster specification only apply to that specific cluster and SHALL NOT be reused in other cluster specifications.

The data types defined in the Data Model and cluster category section SHALL NOT have naming conflicts. A data type defined in the cluster specification MAY have the same name as a data type defined in the Data Model or cluster category section.

To determine which data type is used by an element, the following rules apply, in order from top to

bottom:

- If the data type is defined in the cluster specification, this is used.
- If the data type is defined in the cluster category section and not in the cluster specification within that cluster category section, this is used.
- If the data type is not defined in the cluster specification or the cluster category section, the data type defined in the Data Model is used.

7.19.1. Base Data Types

Class	Data Type	Short	Size
Discrete	Boolean		
	Boolean	bool	1 byte
	Bitmap		
	8-bit bitmap	map8	1 byte
	16-bit bitmap	map16	2 bytes
	32-bit bitmap	map32	4 bytes
	64-bit bitmap	map64	8 bytes
Analog	Unsigned Integer		
	Unsigned 8-bit integer	uint8	1 byte
	Unsigned 16-bit integer	uint16	2 bytes
	Unsigned 24-bit integer	uint24	3 bytes
	Unsigned 32-bit integer	uint32	4 bytes
	Unsigned 40-bit integer	uint40	5 bytes
	Unsigned 48-bit integer	uint48	6 bytes
	Unsigned 56-bit integer	uint56	7 bytes
	Unsigned 64-bit integer	uint64	8 bytes
	Signed Integer		
	Signed 8-bit integer	int8	1 byte
	Signed 16-bit integer	int16	2 bytes
	Signed 24-bit integer	int24	3 bytes
	Signed 32-bit integer	int32	4 bytes
	Signed 40-bit integer	int40	5 bytes
	Signed 48-bit integer	int48	6 bytes
	Signed 56-bit integer	int56	7 bytes
	Signed 64-bit integer	int64	8 bytes

Class	Data Type	Short	Size
Analog	Floating Point		
	Single precision	single	4 bytes
	Double precision	double	8 bytes
Composite	String		
	Octet string	octstr	desc
	Collection		
	List	list	desc
	Struct	struct	desc

7.19.1.1. Boolean Type

The Boolean type represents a logical value, either FALSE or TRUE.

- FALSE SHALL be equivalent to the value 0 (zero).
- TRUE SHALL be equivalent to the value 1 (one).

7.19.1.2. Bitmap Type (8, 16, 32 and 64-bit)

This data type is typically used to represent simple cluster settings or state that are treated as whole.

The [Reserved Bit Fields](#) conventions define reserved bitmap data.

- It is RECOMMENDED to define more bits than initially needed to be able to support more values for later revisions.
- The Bitmap type MAY be used to support up to 8, 16, 32 or 64 boolean values.
- Bits MAY be combined to enumerate other values.
- Bits SHOULD be combined as contiguous bit fields.
- Future revisions MAY require non-contiguous bit fields.
- The conformance for a bit in a bitmap SHALL be mandatory or dependent upon an existing discoverable element, and therefore SHALL NOT be purely optional.

Allowable Conformance for a bit in a bitmap:

- Mandatory
- Dependent upon a Feature supported in the FeatureMap attribute.
- Dependent upon the support of an attribute.

A nullable bitmap SHALL NOT permit use of the most significant bit.

7.19.1.3. Unsigned Integer (8, 16, 24, 32, 40, 48, 56 and 64-bit)

This type represents an unsigned integer with length of N bits and a usable range of:

- $[0..2^N-1]$ if not nullable OR
- $[0..2^N-2]$ if nullable.

The following table presents the representable values following the above rules:

Width N (bits)	Minimum value	Maximum value if nullable	Maximum value if not nullable
8	0 (0x00)	254 (0xFE)	255 (0xFF)
16	0 (0x0000)	65534 (0xFFFE)	65535 (0xFFFF)
24	0 (0x000000)	16777214 (0xFFFFFE)	16777215 (0xFFFFF)
32	0 (0x00000000)	4294967294 (0xFFFFFEE)	4294967295 (0xFFFFFFF)
40	0 (0x0000000000)	1099511627774 (0xFFFFFFFEE)	1099511627775 (0xFFFFFFF)
48	0 (0x000000000000)	281474976710654 (0xFFFFFFFEE)	281474976710655 (0xFFFFFFF)
56	0 (0x00000000000000)	72057594037927934 (0xFFFFFFFEE)	72057594037927935 (0xFFFFFFF)
64	0 (0x0000000000000000)	18446744073709551614 (0xFFFFFFFEE)	18446744073709551615 (0xFFFFFFF)

7.19.1.4. Signed Integer Type (8, 16, 24, 32, 40, 48, 56 and 64-bit)

This type represents an signed integer with length of N bits and a usable range of:

- $[-(2^{(N-1)})..2^{(N-1)}-1]$ if not nullable OR
- $[-(2^{(N-1)}-1)..2^{(N-1)}-1]$ if nullable.

Whether to use two's complement or another representation for the implementation of the Data Model is a consideration of the underlying encoding specification.

The following table presents the representable values in base-10 following the above rules:

Width N (bits)	Minimum value if nullable	Minimum value if not nullable	Maximum value
8	-127	-128	127
16	-32767	-32768	32767

Width N (bits)	Minimum value if nullable	Minimum value if not nullable	Maximum value
24	-8388607	-8388608	8388607
32	-2147483647	-2147483648	2147483647
40	-549755813887	-549755813888	549755813887
48	-140737488355327	-140737488355328	140737488355327
56	-36028797018963967	-36028797018963968	36028797018963967
64	-9223372036854775807	-9223372036854775808	9223372036854775807

7.19.1.5. Single-Precision Type

The single precision number format is based on the [IEEE 754-2019](#) single precision (32-bit) format for binary floating-point arithmetic.

See [IEEE 754-2019](#) for more details on the representable values.

7.19.1.6. Double Precision Type

The double precision number format is based on the [IEEE 754-2019](#) double precision (64-bit) format for binary floating-point arithmetic.

The format and interpretation of values of this data type follow the same rules as given for the single precision data type, but with wider mantissa and exponent ranges.

See [IEEE 754-2019](#) for more details on the representable values.

7.19.1.7. Octet String Type

The octet string data type defines a sequence of octets with a finite octet count from 0 to 65534. It is RECOMMENDED to define a [constraint](#) on the maximum possible count.

7.19.1.8. List

A list is defined as a collection of entries of the same data type, with a finite count from 0 to 65534. A cluster specification may define further constraints on the maximum possible count. The list entry data type SHALL be any defined data type, except a [list](#) data type, or any data type derived from a list.

The quality columns for a list definition are for the list.

The list entries are indicated with an index that is an unsigned integer starting at 0 (zero). The maintained order of entries, by index, is defined in the cluster specification, or undefined. Data that is defined as a list is indicated with "list[X]" where X is the entry type. The data type of the list entry has its own qualities, constraints, and conformance.

To define qualities for the list entry data type, make the list entry data type a defined local derived data type, with a table including the columns required to define and constrain the data type.

For example: *Derived data types defined here:*

Name	Type	Constraint	Quality	...
Month- NameString	string	3	F	...
MonthNumber	uint8	1 to 12		...

SummerStruct defined here:

ID	Name	Type	Constraint	Quality	...
0	Year	int16	-1000 to 3000		...
1	Summer- Months	list[Month- Number]	max 12	N	...

Used Here:

ID	Name	Type	Constraint	Quality	...
0	MonthNames	list[Month- NameString]	12	N	...
1	SummerYears	list[Summer- Struct]	max 50		...

There is an inline shortcut to define the list entry data type constraints. See [List Constraints](#).

For example:

ID	Name	Type	Constraint	Quality	...
0	MonthNames	list[string]	12[3]	N	...

It is RECOMMENDED to put a maximum [constraint](#) on the list and list entry data types.

It is RECOMMENDED that a list entry data type be a struct, to enable the addition of new fields to the list's entries in the future.

- The cluster data version SHALL be incremented when the list order or entries change.
- An entry SHALL NOT be NULL.
- The list SHALL support reading and reporting all entries.
- The list SHALL support reporting, updates, and/or deletion of one or more entries.
- If the list is writable, it SHALL support writing or deleting the entire list.
- If the list is writable, it SHALL support updating one or more individual entries by indicating an index per updated entry.

- If the list is writable, it SHALL support deleting one or more individual entries by indicating an index per deleted entry.
- If the list is writable, it SHALL support adding one or more individual entries.
- A list MAY define an entry that is a struct that is fabric-scoped (see [Section 7.6.4, “Fabric-Scoped Access”](#)).

Fabric-Scoped List

- A fabric-scoped list SHALL define an entry data type that is a struct, which SHALL also be fabric-scoped (see [Fabric-Scoped Struct](#)).

Each entry in a fabric-scoped list SHALL be fabric-scoped to a particular fabric or no fabric.

Fabric-Filtered List

A fabric-scoped list supports a fabric-filter that filters the view of the list for [read](#) and [write](#) interactions. This filter simplifies client side logic that does not want to read or write fabric data that is not associated with the [accessing fabric](#).

- An interaction upon a list with fabric-filtering SHALL only indicate and access entries where the [associated fabric](#) matches the [accessing fabric](#), and all other entries SHALL be ignored.
- Fabric-filtered list entries SHALL be in the same order as the full list.
- Fabric-filtered list entries SHALL be indexed from 0 with no gaps, as if the other entries did not exist.
- For a [write](#) interaction, fabric-filtering SHALL be enabled.
- When writing to a fabric-scoped list, the write interaction SHALL be on an [accessing fabric](#), otherwise, the write interaction SHALL fail (see [Interaction Model Specification](#)).
- For a [read](#) interaction on a list, fabric-filtering MAY be enabled.
- For a [read](#) interaction on a list, with fabric-filtering disabled, the list SHALL be reported as a full list with all entries.

For example: A fabric-scoped full list with each entry having an associated FabricIndex and Value field:

```
List = [ { FabricIndex = A, Value = 20 },
         { FabricIndex = B, Value = 30 },
         { FabricIndex = A, Value = 40 },
         { FabricIndex = B, Value = 50 },
         { FabricIndex = B, Value = 60 } ]
```

would be a fabric-filtered list when accessed with fabric B:

```
List = [ { FabricIndex = B, Value = 30 },
         { FabricIndex = B, Value = 50 },
```



```
{ FabricIndex = B, Value = 60 } ]
```

Reading a fabric-filtered list entry index 2 accessed with fabric B reports:

```
list[2] = [ { FabricIndex = B, Value = 60 } ]
```

Writing fabric-filtered list entry index 1 when accessed with fabric B:

```
list[1] = [ { FabricIndex = B, Value = 55 } ]
```

changes the full list to:

```
list = [ { FabricIndex = A, Value = 20 },
         { FabricIndex = B, Value = 30 },
         { FabricIndex = A, Value = 40 },
         { FabricIndex = B, Value = 55 },
         { FabricIndex = B, Value = 60 } ]
```

7.19.1.9. Struct Type

A struct is a sequence of fields of any data type. Individual fields are identified by a field ID of unsigned integer, starting at 0 (zero), for the first field.

- A struct itself SHALL have no constraint qualities.
- Each struct field SHALL have its own qualities.
- Access, conformance and persistence qualities, when not explicitly defined, SHALL be inherited from the instance of the struct itself.
- Struct fields MAY have optional conformance.
- A struct SHALL support reading and reporting of all fields.
- A struct SHALL support reporting changes to one or more fields.
- If the struct is writable, it SHALL support writing the entire struct.
- If a field of the struct is writable, the struct SHALL support updating the field.
- Because of optional struct field conformance, instances of the same struct MAY support multiple 'flavors' of the same struct data type, but with a different set of optional fields.

Fabric-Scoped Struct

- A fabric-scoped struct SHALL only be defined and occur as an entry in a fabric-scoped list.
- A fabric-scoped struct SHALL support the [global FabricIndex field](#) of type [fabric-index](#), which indicates the [associated fabric](#) of the struct, or indicates that there is no associated fabric.
- The table that defines fields of a fabric-scoped struct SHALL NOT list the [global FabricIndex](#)

field, which is a global field and defined implicitly.

- The [global FabricIndex field](#) of a fabric-scoped struct SHOULD NOT be indicated in a write interaction.
- The [global FabricIndex field](#) of a fabric-scoped struct SHALL be ignored in a write interaction.
- The [global FabricIndex field](#) SHOULD NOT be indicated on a fabric-scoped struct contained in the payload of a request command.
- The [global FabricIndex field](#) SHALL be ignored on a fabric-scoped struct contained in the payload of a request command.
- When a [write](#) interaction creates a fabric-scoped struct entry (in a fabric-scoped list), the server SHALL implicitly load the [accessing fabric-index](#) into the [global FabricIndex field](#) of the struct.
- When the payload of a request command contains a fabric-scoped struct, the server SHALL implicitly load the [accessing fabric-index](#) into the [global FabricIndex field](#) of the struct.
- A fabric-scoped struct MAY be defined with some fields that are [fabric-sensitive](#).
- For interactions on a fabric-scoped struct that report back data, fabric-sensitive struct fields SHALL be omitted when reporting data back to the client, when the struct has an [associated fabric](#), and it is not the [accessing fabric](#).

7.19.2. Derived Data Types

These data types are commonly used and derived from the base data types. If a data type is used by more than one cluster specification, then it SHALL be listed here as a derived data type. Such common data types can then be reused instead of redefined in each cluster specification.

Class	Data Type	Short	Base Type	Size
Analog	Relative			
	Percentage units 1%	percent	uint8	1 bytes
	Percentage units 0.01%	percent100ths	uint16	2 bytes
	Time			
	Epoch Time in microseconds	epoch-us	uint64	8 bytes
	Epoch Time in sec- onds	epoch-s	uint32	4 bytes
	UTC Time	utc	same as Epoch Time in Seconds but Deprecated	
	POSIX Time in mil- liseconds	posix-ms	uint64	8 bytes
	System Time in microseconds	systime-us	uint64	8 bytes
	System Time in milliseconds	systime-ms	uint64	8 bytes
	Elapsed Time in seconds	elapsed-s	uint32	4 bytes
	Physical Quantities			
	Temperature	temperature	int16	2 bytes
	Power	power-mW	int64	8 bytes
	Amperage	amperage-mA	int64	8 bytes
	Voltage	voltage-mV	int64	8 bytes
	Energy	energy-mWh	int64	8 bytes
	Apparent Power	power-mVA	int64	8 bytes
	Apparent Energy	energy-mVAh	int64	8 bytes
	Reactive Power	power-mVAR	int64	8 bytes
	Reactive Energy	energy-mVARh	int64	8 bytes
	Financial			
	Money	money	int64	8 bytes

Class	Data Type	Short	Base Type	Size
Discrete	Enumeration			
	8-bit enumeration	enum8	uint8	1 byte
	16-bit enumeration	enum16	uint16	2 bytes
	Priority	priority	enum8	1 byte
	Status Code	status	enum8	1 byte
	Identifier			
	Group ID	group-id	uint16	2 bytes
	Endpoint Number	endpoint-no	uint16	2 bytes
	Vendor ID	vendor-id	uint16	2 bytes
	Device Type ID	devtype-id	uint32	4 bytes
	Fabric ID	fabric-id	uint64	8 bytes
	Fabric Index	fabric-idx	uint8	1 byte
	Cluster ID	cluster-id	uint32	4 bytes
	Attribute ID	attrib-id	uint32	4 bytes
	Field ID	field-id	uint32	4 bytes
	Event ID	event-id	uint32	4 bytes
	Command ID	command-id	uint32	4 bytes
	Action ID	action-id	uint8	1 bytes
	Transaction ID	trans-id	uint32	4 bytes
	Node ID	node-id	uint64	8 bytes
	Subject ID	subject-id	uint64	8 bytes
	Index			
	Entry Index	entry-idx	uint16	2 bytes
	Counter			
	Data Version	data-ver	uint32	4 bytes
	Event Number	event-no	uint64	8 bytes
	Tag			
	Namespace	namespace	enum8	1 byte
	Tag	tag	enum8	1 byte
	Three-Level Auto Setting			
	Three-Level Auto Setting	ThreeLevelAutoEnum	enum8	1 byte

Class	Data Type	Short	Base Type	Size
Composite	String			
	Character String	string	octstr	desc
	Address			
	IP Address	ipadr	octstr	4 or 16 bytes
	IPv4 Address	ipv4adr	octstr	4 bytes
	IPv6 Address	ipv6adr	octstr	16 bytes
	IPv6 Prefix	ipv6pre	octstr	1 to 17 bytes
	Hardware Address	hwadr	octstr	6 or 8 bytes
	Tag			
	Semantic Tag	semtag	struct	4 bytes
	Location			
	Location Descriptor	locationdesc	struct	variable
	Financial			
	Currency	currency	struct	variable
	Price	price	struct	variable

7.19.2.1. Percentage Type

A Percentage is an unsigned 8-bit value representing percent with a resolution of 1%. The range is from 0 (0%) to 100 (100%).

7.19.2.2. Percentage100ths Type

A Percentage 100ths is an unsigned 16-bit value representing percent with a resolution of 0.01%. The range is from 0 (0.00%) to 10000 (100.00%).

7.19.2.3. Epoch Time in Microseconds Type

This type represents an offset, in microseconds, from 0 hours, 0 minutes, 0 seconds, on the 1st of January, 2000 UTC (the Epoch), encoded as an unsigned 64-bit scalar value.

This offset is the sum of two parts: time elapsed, not counting leap-seconds, and a local time offset. The local time offset MAY include a timezone offset and a MAY include a DST offset.

Any use of this type SHALL indicate how the associated local time offset is determined in the specific context of that use. This MAY be done, for example, by simply saying the time is a UTC time, in which case the local time offset is 0.

A given Epoch Time value MAY be interpreted in at least two ways:

1. The value can be converted to a local clock date/time (year, month, day, hours, minutes, sec-

onds, microseconds) by treating the local time offset as 0 and finding the UTC (year, month, day, hours, minutes, seconds, microseconds) tuple that corresponds to an elapsed time since the epoch time equal to the given value. The value then represents that tuple, but interpreted in the specific timezone and DST situation associated with the value. This procedure does not require knowing the local time offset of the value.

2. The value can be converted to a UTC time by subtracting the associated local time offset from the Epoch Time value and then treating the resulting value as an elapsed count of microseconds since the epoch time.

For example, an Epoch Time value of 0x0000_0BF1_B7E1_0000 corresponds to an offset of exactly 152 days. This can be interpreted as "00:00:00 on June 1, 2000" in whatever local time zone is associated with the value. That corresponds to the following times in ISO 8601 notation:

- 2000-06-01T00:00Z if the associated local time offset is 0 (i.e. the value is in UTC).
- 2000-05-31T23:00Z if the associated local time offset is +1 hour (e.g. the CET timezone, without daylight savings).
- 2000-06-01T00:00+02 if the associated local time offset is +1 hour.
- 2000-06-01T04:00Z if the associated local time offset is -4 hours (e.g. the EDT time zone, which includes daylight savings).
- 2000-06-01T00:00-04 if the associated local time offset is -4 hours.

Conversion from NTP timestamps

Timestamps from NTP also do not count leap seconds, but have a different epoch. NTP 128-bit timestamps consist of a 64-bit seconds portion (NTP(s)) and a 64-bit fractional seconds portion (NTP(frac)). NTP(s) at 00:00:00 can be calculated from the Modified Julian Day (MJD) as follows:

$$\text{NTP(s)} = (\text{MJD} - 15020) * (24 * 60 * 60)$$

where 15020 is the MJD on January 1, 1900 (the NTP epoch)

NTP(s) on January 1, 2000 00:00:00 UTC (MJD = 51544) is 3155673600 (0xBC17C200)

[Epoch Time](#) has a microsecond precision, and this precision can be achieved by using the most significant 32 bits of the fractional portion (NTP(frac32)). Conversion between the 128-bit NTP timestamps and a UTC [Epoch Time in Microseconds](#) is as follows:

UTC [Epoch Time](#) = (NTP(s) - 0xBC17C200)*10⁶ + ((NTP(frac32)*10⁶) / 2³²) where all numbers are treated as unsigned 64-bit integers and the division is integer division.

7.19.2.4. Epoch Time in Seconds Type

This type represents an offset, in seconds, from 0 hours, 0 minutes, 0 seconds, on the 1st of January, 2000 UTC (the Epoch), encoded as an unsigned 32-bit scalar value. Other than that, this type has the same semantics as [Epoch Time in Microseconds](#).

This type is employed where compactness of representation is important and where the resolution of seconds is still satisfactory.

7.19.2.5. POSIX Time in milliseconds Type

This type represents an offset, in milliseconds, from the UNIX epoch (1970-01-01 00:00:00 UTC), encoded as an unsigned 64-bit scalar value.

This type is employed for compatibility reasons.

7.19.2.6. System Time in microseconds Type

System time in microseconds is an unsigned 64-bit value representing the number of microseconds since boot.

7.19.2.7. System Time in milliseconds Type

System time in milliseconds is an unsigned 64-bit value representing the number of milliseconds since boot.

This type is employed for compatibility reasons.

7.19.2.8. Elapsed Time in seconds Type

Elapsed time in seconds is an unsigned 32-bit value representing the time that has elapsed for an operation or other activity, as determined by the definition of the attribute using this type.

7.19.2.9. Temperature Type

This type, derived from int16, represents a temperature on the Celsius scale with a resolution of 0.01°C.

• $value = (temperature\ in\ ^\circ C) \times 100$

- -4°C ⇒ -400
- 123.45°C ⇒ 12345

The range is constrained by absolute zero: -273.15°C to 327.67°C.

Conversion of Temperature Values for Display

When converting temperature values for display manufacturers SHOULD ensure that calculations **round** to the nearest representable value. Particular care is needed when using integer arithmetic.

For example, assuming a display resolution of 0.5:

Attribute value	Temperature		Display as	
°C/100	°C	°F	°C	°F
1965	19.65	67.37	19.5	67.5
-545	-5.45	22.19	-5.5	22

Attribute value	Temperature		Display as	
-1823	-18.23	-0.81	-18	-1

Sample Conversion Code

Sample code provided to ensure consistent Fahrenheit to Celsius and vice-versa conversion between devices and across vendors.

For degF: the value is a int8u representing 2x temperature value in Fahrenheit (to get 0.5 resolution).

For degC: the value is a int16s representing Celsius in 0.01 resolution as expected by the Matter format.

```

/*
 * Function : translateMatterTempDataType()
 * Description : Converts the temperature setpoints in Matter
 * to the half degF format.
 * The half degF format is 8-bit unsigned,
 * and represents 2x temperature value in
 * Fahrenheit (to get 0.5 resolution).
 * The format used in Matter is 16-bit signed
 * in Celsius and multiplied by 100
 * to get 0.01 resolution.
 * e.g. 2500 (25.00 deg C) ---> 0x9A (77 deg F)
 * Input Para : Temperature in Matter (degC) format
 * Output Para: Temperature in half DegF format
 */
int8u translateMatterTempDataType(int16s temperature)
{
    int32s x = temperature;
    // rearrangement of
    // = (x * (9/5) / 100 + 32) * 2;
    // the added 250 is for proper rounding.
    // a rounding technique that only works
    // with positive numbers

    return (int8u) ((x*9*2 + 250)/ (5*100) + 64);
}

/*
 * Function : translateDegFTemp
 * Description : Converts the temperature in DegF

```

```

* protocol to the format
* expected by the cluster attribute
* Measured Value in the
* Temperature Measurement cluster.
* The half deg F format is 8-bit
* unsigned, and represents
* 2x temperature value in
* Fahrenheit (to get 0.5 resolution).
* The format expected by cluster
* is 16-bit signed in Celsius and
* multiplied by 100 to get
* 0.01 resolution.
* e.g. 0x9A (77 deg F) ---> 2500 (25.00 deg C)
* Input Para : temperature in DegF format
* Output Para: temperature in Matter format
*/
int16s translateDegFTemp(int8u temperature)
{
    int32s x = temperature;

    // rearrangement of
    // = 100 * (x/2 - 32) * 5/9
    // *1000 (should be 100), +90, then /10,
    // is for rounding.

    return (int16s) (((x - 64)*5*1000 + 90) / (10*2*9));
}

```

7.19.2.10. Power Type

This type, derived from int64, represents power measured in milliwatts.

7.19.2.11. Amperage Type

This type, derived from int64, represents amperage measured in milliamps.

7.19.2.12. Voltage Type

This type, derived from int64, represents voltage measured in millivolts.

7.19.2.13. Energy Type

This type, derived from int64, represents energy measured in milliwatt-hours.

7.19.2.14. Apparent Power Type

This type, derived from int64, represents apparent power measured in millivolt-amps.

7.19.2.15. Apparent Energy Type

This type, derived from int64, represents apparent energy measured in millivolt-amp-hours.

7.19.2.16. Reactive Power Type

This type, derived from int64, represents reactive power measured in millivolt-amps reactive.

7.19.2.17. Reactive Energy Type

This type, derived from int64, represents reactive energy measured in millivolt-amp-hours reactive.

7.19.2.18. Money Type

This type, derived from int64, represents an amount of money.

The money type SHALL be used in concert with a [CurrencyStruct](#) to indicate the currency in use, as well as the number of digits after the decimal point in the amount of that currency.

For example, if a given money value is 1015, the value of the associated [Currency](#) field is 840 (USD), and the value of the associated [DecimalPoints](#) field is 4, then the represented amount would be 0.1015 dollars, or 10.15 cents.

Similarly, if a given money value is 1015, the value of the associated [Currency](#) field is 978 (EUR), and the value of the associated [DecimalPoints](#) field is 2, then the represented amount would be 10.15 euros.

Similarly, if a given money value is 1015, the value of the associated [Currency](#) field is 392 (JPY), and the value of the associated [DecimalPoints](#) field is 0, then the represented amount would be 1015 yen.

Operations involving money types SHALL be done using fixed point arithmetic. The results of arithmetic operations involving money SHALL be rescaled in accordance with the associated [Currency](#) structure rather than result in altering the [DecimalPoints](#).

7.19.2.19. Enumeration Type (8-bit, 16-bit)

This data type employs scalars to represent context-specific values available from an enumerated set. This data type is nullable.

A data field of this type SHALL have well-defined values with well-defined conformance for implementation. A base cluster that requires derivation MAY defer the definition of all or some enumerated values to the derivation that can be implemented. Such a base cluster MAY also define constraints, such as ranges and semantics of those ranges for the enumeration.

Values for an enumeration MAY be defined directly in a cluster specification, in a derived cluster specification, or in a separate namespace defined in the cluster specification. For examples, please

see uses of the [SemanticTagStruct](#) data type that supports many (and any) standard namespaces.

Undefined values or ranges SHALL be considered reserved and SHALL NOT be implemented.

External standards may be referenced as well as listing the values for the external standard. If the external standard adds values after a specification is adopted, those new values are allowed, but optional.

Enumeration values are defined in a table with a Conformance column. When the definition of an enumeration is missing a Conformance column, all values SHALL be considered to have mandatory conformance.

All mandatory readable enumeration values SHALL be understood by the client. All mandatory writable enumeration values SHALL be understood by the server.

If a client indicates an enumeration value to the server, that is not supported by the server, because it is optional, deprecated, or a new value unrecognized by a legacy server, then the server SHALL respond with the status code `CONSTRAINT_ERROR`, unless the cluster defines alternate behavior, such as:

- convert the value to a mandatory value
- ignore the value
- generate a cluster specific error

With regard to revising a cluster specification:

- It is RECOMMENDED that a client be as strict as possible by indicating only values that a server supports.
- It is RECOMMENDED that the server be as forgiving as possible when processing unsupported values.

Note that indicated enumerations MAY comprise only a strict subset of the required enumerations.

For example: If a server implementation can never enter an enumerated state *XYZ*, then the value *XYZ* would never be indicated, therefore the server would not have to support *XYZ*.

8-bit Enumeration

This type is used for enumeration types that need no more than 256 possible values.

16-bit Enumeration

This type is used for enumeration types that need more than 256 (but less than 65,536) possible values.

7.19.2.20. Priority Type

This is an enumeration of priority used to tag events and possibly other data. The data type does not define any particular ordering among the values. Specific uses of the data type may assign

semantics to the values that imply an ordering relationship.

Value	Priority	Description
0	DEBUG	Information for engineering debugging/troubleshooting
1	INFO	Information that either drives customer facing features or provides insights into device functions that are used to drive analytics use cases
2	CRITICAL	Information or notification that impacts safety, a critical function, or ongoing reliable operation of the node or application supported on an endpoint.

7.19.2.21. Status Code Type

An enumeration value that means a success or error status. A status code is indicated as a response to an action in an interaction (see [Interaction Model Specification](#)).

A status code SHALL be one of:

- a common status code from the set defined in the [Interaction Model status code table](#);
- a cluster status code that is scoped to a particular cluster.

The following table defines the enumeration ranges for status codes.

Status Code Range	Description
0x00	common status code: SUCCESS
0x01	common status code: FAILURE
0x02 to 0x10	cluster scoped status codes
0x70 to 0xCF	other common status codes defined in Interaction Model Status Code Table .

Status codes in an undefined range, or status codes undefined within a range are reserved and SHALL NOT be indicated.

7.19.2.22. Fabric ID Type

A value to identify a fabric.

7.19.2.23. Fabric Index Type

This is an index that maps to a particular fabric on the node, see [Fabric-Index](#). It is used for:

- the [accessing fabric index](#) of an interaction
- the [FabricIndex global field](#) in [fabric-scoped data](#)

7.19.2.24. Node ID Type

A [64-bit ID for a node](#) scoped and unique to a particular fabric as indicated by an accompanying [fabric-index](#) adjacent instantiation.

7.19.2.25. Subject ID Type

A [64-bit integer](#) that identifies the source of an action, referencing an entity that is authenticated via a method provided by the secure channel architecture.

Any use of a Subject ID needs to be accompanied by a method to determine which type of subject is being represented. There are several types of subjects that may need to be disambiguated, depending on the authentication mode used:

- PASE: Lower 16 bits represent the Passcode ID, upper 48 bits are clear.
- CASE: 64 bits represent either the Node ID or a CASE Authenticated Tag.
- Group: Lower 16 bits represent the Group ID, upper 48 bits are clear.

7.19.2.26. Group ID Type

A [16-bit ID for a group](#) scoped to a particular fabric as indicated by an accompanying [fabric index](#) adjacent instantiation.

7.19.2.27. Endpoint Number Type

An unsigned number that indicates an instance of a [device type](#). Endpoint numbers SHALL NOT be 0xFFFF, to allow all endpoint number values to be expressible in nullable endpoint-no fields.

7.19.2.28. Vendor ID Type

A [Vendor ID](#).

Vendor IDs MAY be used as a prefix in a [Manufacturer Extensible Identifier](#) format.

7.19.2.29. Device Type ID Type

An identifier that indicates conformance to a [device type](#).

Device Type IDs SHALL be a [Manufacturer Extensible Identifier](#). The specifics of its representation are described in [Data Model Types](#).

7.19.2.30. Cluster ID Type

An identifier that indicates conformance to a cluster specification.

Cluster IDs SHALL be a [Manufacturer Extensible Identifier](#). The specifics of its representation are described in [Data Model Types](#).

7.19.2.31. Attribute ID Type

An identifier that indicates an attribute defined in a cluster specification.

Attribute IDs SHALL be a [Manufacturer Extensible Identifier](#). The specifics of its representation are described in [Data Model Types](#).

7.19.2.32. Field ID Type

An identifier that indicates a field defined in a struct.

Field IDs SHALL be a [Manufacturer Extensible Identifier](#). The specifics of its representation are described in [Data Model Types](#).

7.19.2.33. Event ID Type

An identifier that indicates an [Event](#) defined in a cluster specification.

Event IDs SHALL be a [Manufacturer Extensible Identifier](#). The specifics of its representation are described in [Data Model Types](#).

7.19.2.34. Command ID Type

An identifier that indicates a [command](#) defined in a cluster specification.

Command IDs SHALL be a [Manufacturer Extensible Identifier](#). The specifics of its representation are described in [Data Model Types](#).

7.19.2.35. Action ID Type

An identifier that indicates an action as defined in the [Interaction Model](#) specification.

7.19.2.36. Transaction ID Type

An identifier for a transaction as defined in the [Interaction Model](#) specification, see [Transaction ID](#).

7.19.2.37. Entry Index Type

This is an index for a [list data type](#).

7.19.2.38. Data Version Type

An unsigned number that indicates a [Data Version](#).

7.19.2.39. Event Number Type

An unsigned number that indicates an [Event](#) instance.

7.19.2.40. Character String Type

The character string data type is derived from an octet string. The octets SHALL be characters with UTF-8 encoding. An instance of this data type SHALL NOT contain truncated code points.

Note that the character string type is a bounded sequence of characters whose size bound format is not specified in the data model, but rather a property of the underlying encoding. Therefore, no assumptions are to be made about the presence or absence of a length prefix or NULL-terminator byte, or other implementation considerations.

It is RECOMMENDED to define constraints on the maximum possible string length.

If at least one of the code points within the string has value 31 (0x1F), which is Unicode **INFORMATION SEPARATOR 1** and ASCII **Unit Separator**, then any client making use of the string SHALL only consider the code points that appear before the first **INFORMATION SEPARATOR 1** as being the textual information carried by the string. Any comparison between such a string and other strings SHALL use the textual component before the first **INFORMATION SEPARATOR 1**. The remainder of the character string after a first **INFORMATION SEPARATOR 1** is reserved for future use by this specification. Implementations of this version of the specification SHALL NOT produce character strings containing **INFORMATION SEPARATOR 1**.

7.19.2.41. IP Address Types

Either an IPv4 or an IPv6 address as defined below.

7.19.2.42. IPv4 Address Type

The IPv4 address data type is derived from an octet string. The octets SHALL correspond to the four octets in network byte order that comprise an IPv4 address represented utilizing quad-dotted notation.

Examples of encoding:

- Address 192.168.2.235 → C0A802EB
- Address 10.4.200.75 → 0A04C84B

7.19.2.43. IPv6 Address

The IPv6 address data type is derived from an octet string. The octets SHALL correspond to the full 16 octets that comprise an IPv6 address as defined by [RFC 4291](#). The octets SHALL be presented in network byte order.

Examples of encoding:

- Address 2001:DB8:0:0:8:800:200C:417A → 20010DB8000000000000080800200C417A
- Address 2001:0DB8:1122:3344:5566:7788:99AA:BBCC → 20010DB8112233445566778899AABBCC

7.19.2.44. IPv6 Prefix

The IPv6 prefix data type is derived from an octet string. The octets SHALL be encoded such that:

- The first octet SHALL encode the prefix length, in bits, in the range of 0 to 128.
 - A value of 0 indicates an absent/invalid prefix.
- The subsequent octets SHALL encode the contiguous leftmost bits of the prefix, in network byte

order, with left justification, such that the first bit of the prefix is in the most significant bit of the first octet. Encoding SHOULD use the least number of bytes to encode the prefix but MAY include unused trailing zeroes.

Examples of encoding:

- Preferred minimal encoding: Prefix **2001:0DB8:0:CD30::/60** → 9 octets → **3C20010DB80000CD30**
- Preferred minimal encoding: Prefix **2001:0DB8:BB00::/40** → 6 octets → **2820010DB8BB**
- Allowed non-minimal encoding: Prefix **2001:0DB8:BB00::/40** → 7 octets → **2820010DB8BB00**

7.19.2.45. Hardware Address Type

The Hardware Address data type SHALL be either a 48-bit IEEE MAC Address or a 64-bit IEEE MAC Address (e.g. EUI-64). The order of bytes is Big-Endian or display mode, where the first byte in the string is the left most or highest order byte.

7.19.2.46. SemanticTagStruct Type

This data type SHALL be represented by the following structure:

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	MfgCode	vendor-id		X	NULL		M
1	Name-spaceID	namespace					M
2	Tag	tag					M
3	Label	string	max 64	X	NULL		MfgCode != NULL, O

MfgCode Field

If the MfgCode field is not NULL, it SHALL be the [Vendor ID](#) of the manufacturer who has defined a certain namespace and the NamespaceID field SHALL be the ID of a namespace defined by the manufacturer identified in the MfgCode field.

If a manufacturer specific Tag field is indicated in a list of SemanticTagStruct entries, the list SHALL include at least one standard tag which is not from any manufacturer's namespace. A standard tag is a tag from a common namespace, a derived cluster namespace, or an applicable device-specific namespace.

If MfgCode is NULL, the NamespaceID field SHALL indicate a standard namespace.

NamespaceID Field

The NamespaceID field SHALL identify a namespace.

The common and device-specific semantic tag namespaces are listed in [\[StandardNamespaces\]](#).

Tag Field

The Tag field SHALL be the ID of a semantic tag located within the namespace indicated by NamespaceID.

A device MAY expose tags from the common or device-specific namespaces and from manufacturer-specific namespaces in a single [Section 9.5.6.5, “TagList”](#).

Label Field

The Label field, if present, SHALL contain human-readable text suitable for display on a client. The content of the Label field is defined by the manufacturer.

This field SHALL be present when the MfgCode is not NULL. This field SHOULD NOT be used if the Tag is from a standard namespace, unless the Tag requires further qualification. For example: A Tag that has the meaning of "room" in a location namespace, would require the a label string to qualify the type of room, such as "1", "2b", "Bathroom", etc.

7.19.2.47. Namespace Type

The Namespace type identifies the namespace used for a semantic tag.

7.19.2.48. Tag Type

The Tag type SHALL identify a semantic tag located within a namespace.

7.19.2.49. LocationDescriptorStruct Type

This data type SHALL be represented by the following structure:

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Location-Name	string	max 128				M
1	Floor-Number	int16		X			M
2	AreaType	tag		X			M

LocationName Field

This field SHALL indicate the name of the location. For example, "blue room".

If the location name is not user provided, the logic that generates it (clients, devices etc.) SHOULD utilize synthesized user-friendly, understandable, names for the location, rather than opaque values such as "private" or "2fe7c241-a50a-4863-896e-c5878da5ed68".

FloorNumber Field

This field SHALL indicate the level number. Negative values correspond to basement levels.

Value zero indicates this is the main floor, which typically includes the main entrance to the user's home. For a building with multiple levels, it is the client's responsibility to map each level to/from a FloorNumber tag value, using the level numbering convention of the region where the client operates. For example, if the client operates in Europe, building level 1, which is one level up from the street level, SHOULD be mapped to FloorNumber tag value 0x1. If the client operates in North America, building level 1, which is at street level, SHOULD be mapped to FloorNumber tag value 0x0.

A NULL value indicates that this information is not available.

When the clients present the level information for user selection, they SHOULD use the operating region to determine how to render and map this data. For example, if the client operates in North America it SHOULD present the user a list that includes entries labeled "basement", "first", "second", and internally mapped to floor numbers -1, 0, and 1. If operating in Europe, the client SHOULD present a list that includes entries labeled "basement", "ground", "first", internally mapped to floor numbers -1, 0, and 1.

The floor number information is expected to be mostly useful to the clients, rather than the devices, such as for grouping devices that are located on the same level. For example, an automation may be defined for all devices located at the basement level (floor number -1).

NOTE

Handling complex level situations, such as half levels (side split houses), or the levels from an apartment building, is up to the client and/or user.

AreaType Field

This field SHALL be the ID of an area semantic tag, located within the [Common Area Namespace](#). For example, this tag may indicate that the location refers to a bedroom.

If this field is NULL, that indicates that the area type information is not available.

NOTE

This field only indicates the type of the area. Multiple areas of the same type, such as bedrooms, MAY exist in a user's home.

7.19.2.50. ThreeLevelAutoEnum Type

This data type is derived from enum8. It is used for a three-level and Auto setting of an attribute on a device. This data type has four values as enumerated below, and cannot be expanded.

Value	Name	Summary	Conformance
0	Auto	Automatic Level	M
1	Low	Low Level	M
2	Medium	Medium Level	M
3	High	High Level	M

7.19.2.51. CurrencyStruct Type

This data type represents a currency with an associated number of decimal points.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Currency	uint16	max 999				M
1	Decimal-Points	uint8	all				M

7.19.2.52. Currency Field

This field SHALL indicate the numeric currency code. The value of the currency field SHALL match one of the numeric currency code values defined by [ISO 4217](#).

7.19.2.53. DecimalPoints Field

This field SHALL indicate the number of digits to the right of the decimal point in values of the specified currency.

7.19.2.54. PriceStruct Type

This data type represents an amount of money in a given currency.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Amount	money	all				M
1	Currency	CurrencyStruct					M

7.19.2.55. Amount Field

This field SHALL indicate the amount of money as specified by the [currency code](#) and [number of decimal places](#) in the [Section 7.19.2.56, “Currency”](#) field. See [Section 7.19.2.18, “Money”](#).

7.19.2.56. Currency Field

This field SHALL indicate the currency for the value in the [Section 7.19.2.55, “Amount”](#) field.

7.20. Units

In order to align units across clusters, the following base units are defined to be used when adding elements that use one of the units listed in the table below. These base units are based on the SI-units, with an exception for Temperature where Celsius is more widely used in day-to-day measurements. Clusters created before these base units were defined MAY diverge from the use of these units, for legacy elements only. When adding new elements of one of the types defined below, the elements SHALL use the associated units.

Be aware that this section does not define any data types, it defines the base units to be used by data types or cluster elements. Any scaling needed is to be either defined in the data type description or the element description.

Quantity	Base Unit
Length	Meter
Mass	Kilogram
Time	Second
Electrical voltage	Volt
Electrical current	Ampere
Temperature	Degree Celsius
Pressure	Pascal

7.21. Manufacturer Specific Extensions

This section covers Manufacturer Specific (MS) extensions and how they are supported by identifiers, paths, wildcards, discoverability, etc.

7.21.1. Manufacturer Extensible Identifiers

A Manufacturer Extensible Context (MEC) contains a collection of items which MAY be extended by manufacturers. Each item in a MEC has a source which is either Standard, Scoped or a particular Manufacturer Code (MC).

- A Standard source references definitions described in Matter standard clusters.
- A Scoped source adopts the same source as that of the cluster that contains its definition.
- An MC-based source references manufacturer-specific definitions.

Table 98. MEC Example

Context	Source	Items
MEC	Standard	Item 0
		Item 1
		Item 2
	Scoped	Item 0
		Item 1
		Item 2
	MC 1	Item 0
		Item 1
		Item 2
	MC 2	Item 0
		Item 1
		Item 2

A Manufacturer Extensible Identifier (MEI) identifies an item in an MEC and has no meaning beyond the context of that MEC.

7.21.2. Manufacturer Extensible Identifier (MEI)

An MEI has the following format:

Table 99. MEI Format

Field	Prefix	Suffix
Description	Encodes a source (standard, scoped or a particular MC)	Encodes an item's key (in context of MEC + source)
Width	16-bit	16-bit
Bit Positions	31..16	15..0

The MEI permits encoding of ~65K keys in the suffix.

A specific MEI MAY only permit certain combinations of the above.

7.21.2.1. Encoding

The MEI prefix encodes the source [Vendor ID](#) and follows the same rules as outlined in [Table 1, “Vendor ID Allocations”](#), with the exception that a Scoped source is encoded using the same prefix as a Standard source. Consequently, a given MEI SHALL NOT permit both Standard and Scoped source types given the ambiguity in telling them apart.

Given the above, the encoding is as follows:

Table 100. MEI Prefix

Prefix	Source
0x0000	Standard OR Scoped
0x0001 - 0xFFFF0	Manufacturer Code as per [Connectivity Standards Alliance Manufacturer Code Database]
0xFFFF1 - 0xFFFF4	Test Vendor MC

The MEI suffix encodes a key as follows:

Table 101. MEI Suffix

Suffix	Item
0x0000 - 0xFFFFE	Item 0 to 65534

7.21.2.2. Data Model Types

The following data model types SHALL be represented as MEIs:

Table 102. MEI Suffix

Type	Permitted Source Types	Suffix Range
Device Type ID	Standard or MC	0x0000 - 0xBFFF
Cluster ID	Standard or MC	Standard Cluster: 0x0000 - 0x7FFF
		Manufacturer-Specific Cluster: 0xFC00 - 0xFFFFE
Attribute ID (Global)	Standard	0xF000 - 0xFFFFE
Attribute ID (Non-Global)	Scoped or MC	0x0000 - 0x4FFF
Event ID	Scoped or MC	0x00 - 0xFF
Command ID (Global)	Standard	0xE0 - 0xFF
Command ID (Non-Global)	Scoped	0x00 - 0xDF
Command ID	MC	0x00 - 0xFF
Field ID (Global)	Standard	0xE0 - 0xFE
Field ID (Non-Global)	Scoped or MC	0x00 - 0xDF

For those types which are listed with an 8 bit suffix range in this table (i.e. Event ID, Command ID and Field ID), the 16-bit MEI Suffix SHALL be constructed by padding the 8-bit data with a most significant byte set to zero.

For example:

Table 103. MEI Decoding Example

MEI	Description
0x0000_0000	Standard/Scoped item 0

MEI	Description
0x0000_0001	Standard/Scoped item 1
0x0000_0002	Standard/Scoped item 2
0x0000_FFFE	Standard/Scoped item 65534
0x0001_0000	MC 1 item 0
0x0001_0001	MC 1 item 1
0x0001_0002	MC 1 item 2
0x0001_FFFE	MC 1 item 65534
0x0002_0000	MC 2 item 0
0x0002_0001	MC 2 item 1
0x0002_0002	MC 2 item 2
0x0002_FFFE	MC 2 item 65534
0xFFFF_0000	Invalid

7.21.3. Manufacturer Extensions

A manufacturer extensible context MAY be extended with items from any manufacturer. Such extensions SHALL be identified using an MEI with prefix for that particular manufacturer, and SHALL NOT use a standard/scoped prefix.

There are further constraints:

- MS extensions SHALL only be permitted on standard clusters or another existing MS extension of a standard cluster from another manufacturer.
- An extended cluster MAY instantiate a struct definition defined in the standard cluster.
- A struct that has been extended with new fields SHALL have the same definition in all instances of that struct within a given cluster definition.
- A defined element (struct, command, event) SHALL NOT be re-used or instantiable in a different cluster (except in extended clusters)

This is illustrated by the following hypothetical scenario.

Suppose the standard provides cluster **ABCD** which contains related counters and their recent statistics. The counter values are available as attributes **1** and **2**, which are reset daily. The statistics are grouped into a **SummarizedStats** struct, available as attributes **3** and **4**, and track summary statistics for each counter over a recent period (last month). Each instance of the statistics struct has fields **1**, **2**, and **3**, for minimum, maximum, and mean values for that period.

Suppose manufacturer A extends the standard cluster with additional statistics (marked with ¹ below). A adds lifetime counts as attributes **0x000A_0001** and **0x000A_0002**, which are never reset. A also adds quartiles Q1, Q2, and Q3 to the standard **SummarizedStats** struct, as fields **0x000A_0001**, **0x000A_0002**, and **0x000A_0003**. These quartiles are available for all existing instances of the standard struct, such as standard attributes **3** and **4**.

Suppose manufacturer B, a partner of manufacturer A, extends the standard cluster further (marked with ² below). B wishes to add instances of the standard statistics struct, as attributes **0x000B_0001** and **0x000B_0002**, to track summary statistics for each counter, over a different recent period (last year instead of last month). Since manufacturer A had already extended the standard statistics struct, the instantiation of that struct will contain both standard and A's fields. If B desires to create a new version of that statistics struct without A's changes, it would have to declare a new definition of that struct with new fields in it.

Suppose manufacturer C, a partner of manufacturers B and A, adds a MS cluster **0x000C_FC01** (marked with ³ below) that doesn't extend an existing standard cluster. This cluster has a sensor available as attribute **0x0000_0001** of type **SensorStats**, which has fields **0x0000_0001**, **0x0000_0002**, and **0x0000_0003**, for the sensor's value, precision, and accuracy. Since Attribute and Field IDs are defined using the 'Scoped' source type, the prefix of '0000' implicitly equates to the same source as the cluster it is defined in, i.e. Manufacturer C. C also wishes to add an instance of the **SummarizedStats** struct as attribute **0x000C_0002**, to track summary statistics for the sensor over a recent period (last hour). Since this cluster does not extend any previous cluster, it cannot instantiate any of the extended versions of the **SummarizedStats** struct as defined previously. Instead, C will have re-define that structure definition within its cluster definition and use it.

Table 104. Hypothetical Standard Cluster

Endpoint	Cluster	Attribute	Attribute Description	Struct Field	Field Description
0x0001	0x0000_ABCD	0x0000_0001	Counter 1 current value (reset daily)	-	-
		0x0000_0002	Counter 2 current value (reset daily)	-	-
		0x0000_0003	Counter 1 (period = month)	0x0000_0001	Min count
				0x0000_0002	Max count
				0x0000_0003	Mean count
		0x0000_0004	Counter 2 (period = month)	0x0000_0001	Min count
				0x0000_0002	Max count
				0x0000_0003	Mean count

Table 105. Hypothetical Manufacturer A Extension Scenario

Endpoint	Cluster	Attribute	Attribute Description	Struct Field	Field Description
0x0001	0x0000_ABCD	0x0000_0001	Counter 1 current value (reset daily)	-	-
		0x0000_0002	Counter 2 current value (reset daily)	-	-
		0x0000_0003	Counter 1 (period = month)	0x0000_0001	Min count
				0x0000_0002	Max count
				0x0000_0003	Mean count
				0x000A_0001 ¹	Q1 count ¹
				0x000A_0002 ¹	Q2 count ¹
				0x000A_0003 ¹	Q3 count ¹
		0x0000_0004	Counter 2 (period = month)	0x0000_0001	Min count
				0x0000_0002	Max count
				0x0000_0003	Mean count
				0x000A_0001 ¹	Q1 count ¹
				0x000A_0002 ¹	Q2 count ¹
				0x000A_0003 ¹	Q3 count ¹

Table 106. Hypothetical Manufacturer B Extension of A Scenario

Endpoint	Cluster	Attribute	Attribute Description	Struct Field ^	Field Description
0x0001	0x0000_ABCD	0x0000_0001	Counter 1 current value (reset daily)	-	-
		0x0000_0002	Counter 2 current value (reset daily)	-	-
		0x0000_0003	Counter 1 (period = month)	0x0000_0001	Min count
				0x0000_0002	Max count
				0x0000_0003	Mean count
				0x000A_0001 ¹	Q1 count ¹
				0x000A_0002 ¹	Q2 count ¹
				0x000A_0003 ¹	Q3 count ¹
		0x0000_0004	Counter 2 (period = month)	0x0000_0001	Min count
				0x0000_0002	Max count
				0x0000_0003	Mean count
				0x000A_0001 ¹	Q1 count ¹
				0x000A_0002 ¹	Q2 count ¹
				0x000A_0003 ¹	Q3 count ¹
		0x000B_0001 ²	Counter 1 (period = month) ²	0x0000_0001	Min count
				0x0000_0002	Max count
				0x0000_0003	Mean count
				0x000A_0001 ¹	Q1 count ¹
				0x000A_0002 ¹	Q2 count ¹
				0x000A_0003 ¹	Q3 count ¹
		0x000B_0002 ²	Counter 2 (period = month) ²	0x0000_0001	Min count
				0x0000_0002	Max count
				0x0000_0003	Mean count
				0x000A_0001 ¹	Q1 count ¹
				0x000A_0002 ¹	Q2 count ¹
				0x000A_0003 ¹	Q3 count ¹

Table 107. Hypothetical Manufacturer C Custom Cluster

Endpoint	Cluster	Attribute	Attribute Description	Struct Field	Field Description
0x0001	0x000C_FC01 ³	0x0000_0001 ³	Sensor 1 Stats ³	0x0000_0001 ³	Value ³
				0x0000_0002 ³	Precision ³
				0x0000_0003 ³	Accuracy ³
		0x0000_0002 ³	Counter 1 (period = hour) ³	0x0000_0001 ³	Min daily count ³
				0x0000_0002 ³	Max daily count ³
				0x0000_0003 ³	Mean daily count ³

Note the following potential combinations of path components:

Table 108. Hypothetical Manufacturer Extension Path Examples

	Description
0x0001/0x0000_ABCD/0x0000_0003/0x0000_0001	Cluster ID = Standard, Attribute ID = Scoped, Field ID = Scoped: Counter 1 min value
0x0001/0x0000_ABCD/0x0000_0003/0x000A_0001	Cluster ID = Standard, Attribute ID = Scoped, Field ID = MS(A): Counter 1 Q1 daily count over last month
0x0001/0x0000_ABCD/0x000B_0001/0x0000_0001	Cluster ID = Standard, Attribute ID = MS(B), Field ID = Scoped: Counter 1 Q1 daily count over last year
0x0001/0x0000_ABCD/0x000B_0001/0x000A_0001	Cluster ID = Standard, Attribute ID = MS(B), Field ID = MS(A): Counter 1 Q1 daily count over last year
0x0001/0x000C_FC01/0x0000_0001/0x0000_0002	Cluster ID = MS(C), Attribute ID = Scoped, Field ID = Scoped: Sensor 1 precision
0x0001/0x000C_FC01/0x0000_0002/0x0000_0003	Cluster ID = MS(C), Attribute ID = Scoped, Field ID = Scoped: Counter 1 (period = hour) Mean
0x0001/0x000C_FC01/0x0000_FFFD	Cluster ID = MS(C), Attribute ID = Standard: Cluster revision

7.21.4. Discoverability

The Descriptor Cluster reports the device types and clusters on a node's endpoints, whether they are standard or from a particular manufacturer.

For example, if a node supports cluster 0x000C_ABCD on endpoints 1 and 2, that information is available in the Descriptor Cluster.

The Read Interaction provides a means to read the contents of all or part of a cluster.

For example, reading cluster `0x0000_ABCD` on endpoint 1 might return mandatory attribute `0x0000_0001`, optional attribute `0x0000_0009`, and MS attributes `0x000A_0001` and `0x000A_0002`.

Chapter 8. Interaction Model Specification

8.1. Practical Information

8.1.1. Revision History

Please note that Matter revision 1.0 SHALL be considered equivalent to revision 10.

Revision	Description
1	Equivalent to revision 10 (1 st Matter release of the Interaction Model)
2-9	Released as versions of the Zigbee Cluster Library chapter 2 which combined the interaction model with encoding
10	Initial release of this specification (Matter 1.0)
11	Updated Event error processing (Matter 1.2)
12	Added WildcardPathFlags to AttributeDataIB, and associated behavior; added support for multiple CommandDataIB in one Invoke Request
13	Added WildcardFilterConfigurationVersion (Matter 1.3)

8.1.2. Scope & Purpose

This is part of a package of Data Model specifications that are agnostic to underlying layers (encoding, message, network, transport, etc.). Each specification below may be independently maintained. This package, as a whole, SHALL be independently maintained as agnostic and decoupled from lower layers. This package may be referenced by inclusion in vertical protocol stack specifications.

Data Model	Defines first order elements and namespace for endpoints, clusters, attributes, data types, etc.
Interaction Model	Defines interactions, transactions and actions between nodes.
System Model	Defines relationships that are managed between endpoints and clusters.
Cluster Library	Reference library of cluster specifications.
Device Library	Reference library of devices type definitions.

8.1.3. Origin Story

The original baseline for this section comes from the Zigbee Cluster Library [ZCL] Chapter 2 relating to ZCL commands and interactions. This specification addresses these gaps determined by the Data Model Tiger Team:

- Multi-Element Message support
- Synchronized Reporting
- Reduce message types (commands & actions)
- Complex data type support in all messages
- Events
- Interception attack

8.1.4. Purpose

The purpose is to define a layer that abstracts interactions from other layers, including security, transport, message format & encoding. The intent is that this document will align with current cluster specifications in the ZCL (revision 8 at this time), and still support cluster evolution over time.

8.1.5. Glossary

Term	Short	Spec Details	Description
Wildcardable	A	<i>this spec</i>	able to indicate all current instances of the data
Optional	O	Data Model	only required for some action behavior
Quality	Qual, Q	Data Model	quality of information in an information block
Action Flow		<i>this spec</i>	direction flow of actions
Path		<i>this spec</i>	a path to an element (see Path)
Group Path		<i>this spec</i>	a path with a group ID instead of node ID and endpoint number (see Group Path)
Wildcard Path		<i>this spec</i>	a path with one or more elements that are wildcards (see Wildcard Path)
Attribute Path		<i>this spec</i>	a path to an attribute data field path for attribute data (see Attribute Path)
Request Path		<i>this spec</i>	a path that may be a group or wildcard path (see Request Path)

Term	Short	Spec Details	Description
Concrete Path		<i>this spec</i>	a path that is not a group or wildcard path (see Concrete Path)
Existent Path		<i>this spec</i>	a concrete path that exists on a server cluster (see Existent Path)
Supported		Data Model	the indicated element is supported by the implemented instance
Unsupported		Data Model	the indicated element is not supported by the implemented instance

8.1.6. Conventions & Conformance

See the [Data Model](#) specification.

8.2. Concepts

Relationships between devices are established using data model elements and interactions defined here. See the [System Model](#) specification for more information.

An interaction is a sequence of transactions. A transaction is a sequence of actions.

An action is a single logical communication from a source node to one or more destination nodes. An action is conveyed by one or more messages.

The actual construction and encoding of messages is left to the message layer, which is the layer below this layer.

- The protocol layers below this layer MAY have constraints that only support a subset of the functionality described here.

Examples:

- A client may choose Read interactions instead of Subscribe interactions.
- A client may choose to not Write or Invoke commands.

8.2.1. Path

A path is used to indicate one or more element instances in the data model. The path has the form as described in Augmented Backus–Naur:

```
<path> ::= <target> <cluster> <cluster element>
```

```

<target> ::= <group target> | <endpoint target>
<group target> ::= <group ID>
<endpoint target> ::= <node> <endpoint>
<endpoint> ::= <wildcard endpoint> | <concrete endpoint>
<cluster> ::= <wildcard cluster> | <concrete cluster>
<cluster element> ::= <attribute> | <event> | <command>

```

An [Attribute Path](#) is a path indicating an [<attribute>](#).

A [Command Path](#) is a path indicating a [<command>](#).

An [Event Path](#) is a path indicating an [<event>](#).

8.2.1.1. Concrete Path

- A concrete path SHALL NOT have group IDs or wildcards.
- A concrete path SHALL indicate a single element instance that is either:
 - an event with the path ending in an event ID
 - a command with the path ending in a command ID
 - an attribute with the path ending in an attribute ID
 - a struct field with the path ending in a field ID
 - a list entry with the path ending in a list entry index.

8.2.1.2. Existent Path

- An existent path is a concrete path that indicates a single existing instance on the node indicated in the path.

8.2.1.3. Group Path

A group path is a path that targets endpoints that are members of a group, using group ID, instead of indicating a node and endpoint.

- A group path SHALL resolve into zero or more paths.
- A group path SHALL include a group ID that indicates zero or more endpoints that are members of the group.
- A group path MAY include a wildcard cluster indication and therefore also be a [Wildcard Path](#).

8.2.1.4. Wildcard Path

A wildcard path is a path with a wildcard endpoint indication and/or wildcard cluster indication.

- A wildcard path SHALL resolve into zero or more paths.
- A wildcard path SHALL indicate zero or more element instances.
- A wildcard path MAY include a group ID and therefore also be a [Group Path](#).

8.2.1.5. Request Path

A request path is used in actions that request data model elements.

- A request path SHALL be either a concrete path, a group path or a wildcard path.

8.2.1.6. Request Path Expansion

Many actions specify this process step to expand a request path into a list of existent paths. This process does not check access qualities, such as read or write access, privileges, or fabric qualities.

- If the path is a [Group Path](#), it SHALL be replaced with a list of paths, one for each endpoint that is a member of the group on the target node.
 - All concrete paths that are not existent paths in the list generated by the above-mentioned group-to-endpoint path expansion SHALL be removed.
- Else the list SHALL be the path.
- Each path in the list that is a [Wildcard Path](#) SHALL be expanded into a complete list of existent paths. This is done by generating all permutations where the wildcarded elements are replaced with existent elements.
- When this expansion is performed for an Attribute read or subscription use case, any paths that must be omitted due to the processing of the Attribute Wildcard Path Flags SHALL be excluded from the resulting list. For other interactions, the Attribute Wildcard Path Flags SHALL be ignored.

This process produces zero or more existent paths.

8.2.1.7. Attribute Path

An attribute path is used to indicate all or part of a cluster attribute. An attribute path may indicate deeper parts of [collection type data](#).

Associated Information Block: [AttributePathIB](#)

If the attribute data type is a [collection data type](#), such as a struct or list, then the path may indicate deeper nested parts of the data.

The nesting of [collection data](#) is conceptually unlimited, but the actual structure of the data is well-defined in the cluster specification. Attribute data structures are similar to data structures supported in a programming language (see [Data Types](#) in the Data Model specification). An attribute path is conceptually similar to the path or dot notation used to reference programming language data structures.

A field ID for structure data or an entry index for list data are currently the only options in an attribute path, after the attribute ID itself.

- The **<attribute>** component of an attribute path SHALL have the following form:

```
<attribute> ::= <attribute ID> <nesting level>*
```

<nesting level> ::= <struct field ID> | <list entry index>

* <nesting level> occurs zero or more times as defined in a cluster specification.

The endpoint component is subject to wildcard expansion, as constrained in particular actions and contexts.

Attribute Wildcard Path Flags

Processing wildcard expansion of an Attribute Path MAY be affected by the associated WildcardPathFlags and WildcardFilterConfigurationVersion fields.

When expanding a wildcard path indicated in an [AttributePathIB](#) into one or more concrete paths, paths SHALL be omitted from the result set if the [WildcardFilterConfigurationVersion](#) is greater than or equal to the current [ConfigurationVersion](#) in the Basic Information cluster and they satisfy any of the conditions below, each of which depends on the value of [WildcardPathFlags](#) and the path being considered:

- if the [WildcardSkipRootNode](#) bit is set and the concrete path targets endpoint 0;
- if the [WildcardSkipGlobalAttributes](#) bit is set and the concrete path refers to one of the following attribute IDs:
 - GeneratedCommandList/0xFFFF8
 - AcceptedCommandList/0xFFFF9
 - AttributeList/0xFFFFB
- if the [WildcardSkipAttributeList](#) bit is set and the concrete path refers to the AttributeList (ID 0xFFFFB) attribute;
- if the [WildcardSkipCommandLists](#) bit is set and the concrete path refers to either the GeneratedCommandList (ID 0xFFFF8) or AcceptedCommandList (ID 0xFFFF9) attributes;
- if the [WildcardSkipCustomElements](#) bit is set, and the concrete path targets:
 - any element whose Cluster ID has an [MEI prefix](#) not equal to zero;
 - any attribute whose ID has an [MEI prefix](#) not equal to zero, even if included in a standard cluster;
- if the [WildcardSkipFixedAttributes](#) bit is set and the concrete path refers to an attribute with the Fixed (F) quality;
- if the [WildcardSkipChangesOmittedAttributes](#) bit is set and the concrete path refers to an attribute with the Changes Omitted (C) quality;
- if the [WildcardSkipDiagnosticsClusters](#) bit is set and the concrete path refers to a cluster with the [Diagnostics \(K\) quality](#).

WildcardPathFlags SHALL ONLY be used for either Read or Subscribe interactions.

If the client is including WildcardPathFlags in an AttributePathIB destined to a server that pre-dates the introduction of WildcardPathFlags (currently provisional), the client SHALL tolerate the inclusion of reports for paths that would otherwise be omitted by servers compliant with this feature.

NOTE

Support for WildcardPathFlags and WildcardFilterConfigurationVersion is provisional.

8.2.1.8. Command Path

A command path is used to indicate a cluster command.

Associated Information Block: [CommandPathIB](#)

- The **<command>** component of a command path SHALL have the following form:

```
<command> ::= <command ID>
```

- The endpoint field is wildcardable, though this may be disallowed in the various uses of the Command Path in different actions and contexts.

8.2.1.9. Event Path

A event path is used to indicate a cluster event.

Associated Information Block: [EventPathIB](#)

- The **<event>** component of an event path SHALL have the following form:

```
<event> ::= <event ID>
```

See [Section 7.14, “Event”](#) for a description of a cluster event and event data fields.

The endpoint, cluster and event ID fields are wildcardable. These are further constrained in the various uses of the Event Path in different actions and contexts.

- An event path SHALL NOT be a group path.

8.2.2. Interaction

An interaction is a sequence of one or more transactions between nodes, that occurs in the context of an [accessing fabric](#), or no fabric.

How a fabric, or no fabric, context is established for an interaction, is not defined here.

The first transaction (of an interaction) starts with the first action from the node called the **initiator**. The first action destination is called the **target**, which is either a node or group. For the remainder of the interaction, the initiator remains the same.

An interaction may be a single transaction (e.g. Read). An interaction may be an unbounded number of transactions (e.g. Subscribe).

Interaction	Transactions	Description
Read Interaction	Read	This interaction is a request for cluster attributes and/or event data.
Subscribe Interaction	Subscribe, Report	This interaction subscribes to cluster attributes and/or event data.
Write Interaction	Write	This interaction modifies cluster attributes.
Invoke Interaction	Invoke	This interaction invokes cluster commands.

8.2.3. Transaction

A transaction is either the whole, or part of an interaction. A transaction is a sequence of one or more actions. Actions in a transaction are defined as first or following, to better describe dependencies in this specification.

- The first action of a transaction SHALL be initiated by a single node.
- An action in a transaction SHALL have a target destination that is either a single node, called a unicast action or a group of nodes, called groupcast action.

8.2.3.1. Transaction ID

The transaction ID is a field present in all actions (see [Common Action Information](#)) that indicates the logical grouping of those actions.

- All following actions in a transaction SHALL have the same transaction ID as the first action.
- A groupcast action SHALL end a transaction and any subsequent action in the interaction SHALL NOT use the same transaction ID.

The table below lists all transactions.

Transaction	Description
Read Transaction	This transaction is a request for cluster attribute and/or event data.
Subscribe Transaction	This transaction creates a subscription to cluster attributes and/or events.
Report Transaction	This transaction maintains a subscription for the Subscribe interaction.
Write Transaction	This transaction modifies cluster attributes.
Invoke Transaction	This transaction invokes cluster commands.

8.2.4. Action

The table below lists all actions.

Action	Description	Outgoing Message
Status Response Action	This action is a success or error response.	Unicast
Read Request Action	This action is a request for cluster attribute data and/or events.	Unicast
Report Data Action	This action responds to a Read Request Action or Subscribe Request Action .	Unicast
Subscribe Request Action	This action is a request for a subscription to cluster attribute data and/or events.	Unicast
Subscribe Response Action	This action is a response to a Subscribe Request Action .	Unicast
Write Request Action	This action is a request to modify cluster attribute data.	Unicast Groupcast
Write Response Action	This action responds to a Write Request Action .	Unicast
Invoke Request Action	This action executes a cluster command.	Unicast Groupcast
Invoke Response Action	This action is used to respond to an Invoke Request Action with cluster defined responses.	Unicast
Timed Request Action	This action indicates that another action will take place within a Timed interval.	Unicast

8.2.5. Common Action Behavior

The message layer below this interaction layer encodes an action into one or more messages and delivers the messages to a destination. This interaction layer delivers action information to the message layer by passing action information, through some interface (not defined here). The message layer delivers action information, from an incoming message, to this interaction layer.

In all action descriptions in this specification, action information (or information blocks), refers to the information that is transferred to and from the message layer.

There is no designation of mandatory or optional for such information because the implementation is undefined. However, some information fields may be omitted, meaning the information may not be needed for all actions.

8.2.5.1. Common Action Information

Action Field	Type	Conformance	Description
InteractionModelRevision	uint8	M	the revision number of the implemented Interaction Model specification under which the sending node was certified
Action	action-id	M	the action
TransactionID	trans-id	M	the transaction ID
FabricIndex	fabric-idx	M	the accessing fabric index , based on the session used to deliver the action
SourceNode	node-id	M	the node ID of the node that generates the action
DestinationNode	node-id	O.a	the node ID of the destination where the action is sent
DestinationGroup	group-id	O.a	the group ID of the destination where the action is sent
<i>action specific</i>	<i>variable</i>	M	specific action information described in each action section

8.2.5.2. Outgoing Action

- Each generated action SHALL provide the action information above to the message layer.
- If the action is the first action of a transaction, the TransactionID SHALL be a value that uniquely identifies the transaction on the source of the action.
- If the action is a following action, the TransactionID SHALL be the same as the TransactionID in the first action of the transaction.
- If the action is a unicast following action the DestinationNode SHALL be the SourceNode of the previous action in the transaction.
- The generated action information SHALL be submitted to the message layer.
 - Upon receipt of this action information, the message layer SHALL construct and convey one or more messages for this action to the target.
 - If the message layer encounters an error that prevents the complete construction, encoding and/or conveyance of the action, then the message layer SHALL inform this layer of the error.

- c. If the action is not completely conveyed, the action, with the associated transaction and interaction, SHALL terminate.
 - i. If the failed action is NOT a Status Response action, this layer SHOULD, if possible, submit a Status Response action to the message layer, with a status code of FAILURE and the same TransactionID.

8.2.5.3. Incoming Action

1. If the message layer receives a valid message for an action, it SHALL be delivered to this layer with the action information above.
2. If this layer receives a message for an action that is not expected semantically, has invalid action information, or has an error not described in this specification, a Status Response action with an INVALID_ACTION Status Code SHALL be generated as defined in [Status Response Action](#), and the associated transaction and interaction SHALL terminate.
3. If during the receipt and decoding of messages for an action by the message layer, an error occurs that prevents a complete receipt of a valid action, then the message layer SHALL inform this layer of the error.
 - a. When informed of an error from a message layer, the action, with the associated transaction and interaction, SHALL terminate.
4. If the action is not able to be executed due to insufficient resources, a Status Response SHALL be sent to the initiator with a status code of either:
 - a. PATHS_EXHAUSTED if there are not enough resources to support the number of paths in the action information,
 - b. and the number of paths in the action exceeds the number of paths that is guaranteed to be supported for the action (see [Interaction Model Limits](#)),
 - c. BUSY in all other recoverable resource exhausted situations (e.g. if too many Read interactions are already in progress),
 - d. or RESOURCE_EXHAUSTED for any other resource insufficiency,
 - e. and the interaction SHALL be terminated.

It is implementation specific whether the message layer submits logical parts of an action to this layer as it receives and processes each message. The only requirement above is that all the information, or an error, be submitted to this layer.

Global common interaction Status Codes are defined in this document in [Status Codes](#). Cluster specific Status Codes are defined in each cluster specification.

8.3. Status and Interaction

There is no Status interaction, but an error status may be generated as part of any interaction.

8.3.1. Status Response Action

This action is defined as a following action for some actions, or is generated when there is an

unspecified transaction or interaction error. This action conveys status to this layer or conveys status from this layer to another node. The status indicates success or an error as part of a transaction or interaction.

See [Section 8.2.5, “Common Action Behavior”](#) for behavior common to all actions. The specific action information for this action is shown below.

8.3.1.1. Status Response Information

Action Field	Type	Conformance	Description
Status	status	M	a status code (see Status Codes)

8.3.1.2. Outgoing Status Response Action

- This action SHALL be unicast.
- This action SHALL NOT be generated in response to a groupcast.
- This action SHALL be generated as specified in interactions defined here.
- If this action is generated with an error Status, the current transaction and interaction SHALL be terminated.
- This action SHALL only be generated with an error Status when an error occurs as a result of the immediate previous received action in the current transaction.
- This action’s DestinationNode field SHALL be the immediate previous received action’s SourceNode.
- This action’s TransactionID field SHALL be the immediate previous received action’s TransactionID.
- If there is no well-defined Status Code for an error or exception, the Status Code of FAILURE SHALL be used.

8.3.1.3. Incoming Status Response Action

- Upon receipt of this action with a success Status Code, this layer SHALL consume the status and continue the current transaction and interaction.
- Upon receipt of this action with an error Status, this layer SHALL terminate the current transaction and interaction.
- Upon receipt of this action with an error Status, this layer SHALL submit the error to the layer above.

8.4. Read Interaction

This interaction is generated when an initiator wishes to determine the value of one or more attributes or events located on a node.

8.4.1. Read Transaction

Action	Action Flow	Description
Read Request	Initiator ⇒ Target	data report request
Report Data	Initiator ⇐ Target	response report with data

8.4.2. Read Request Action

Read Request action is the first action of a Read transaction (and interaction). See [Section 8.2.5, “Common Action Behavior”](#) for behavior common to all actions. The specific action information for this action is shown below.

8.4.2.1. Read Request Action Information

Action Field	Type	Conformance	Description
AttributeRequests	list[AttributePathIB]	M	a list of zero or more request paths to cluster attribute data
DataVersionFilters	list[DataVersionFilterIB]	AttributeRequests	a list of zero or more cluster instance data versions
EventRequests	list[EventPathIB]	M	a list of zero or more request paths to cluster events
EventFilters	list[EventFilterIB]	EventRequests	a list of zero or more minimum event numbers per specific node
FabricFiltered	bool	M	limits the data read within fabric-scoped lists to the accessing fabric

8.4.2.2. Outgoing Read Request Action

- This action SHALL be unicast.
- This action SHALL be generated as the first action in a Read transaction.
- A valid [AttributePathIB](#) for attribute data SHALL be one in the table [Valid Read Attribute Paths](#).
- A valid [EventPathIB](#) for an event SHALL be one in the table [Valid Event Paths](#).
- A path indicated in AttributeRequests or EventRequests SHALL NOT target a group.

8.4.2.3. Incoming Read Request Action

- Upon receipt of this action, this layer SHALL generate a Report Data action to the subscriber, as defined in [Incoming Read Request and Subscribe Request Action Processing](#).

- If the Report Data was generated successfully, it SHALL be submitted to the message layer.

8.4.3. Report Data Action

This action is either a first action in a Report transaction (as part of a Subscribe interaction), or a following action to a Read Request action or Subscribe Request action.

See [Section 8.2.5, “Common Action Behavior”](#) for behavior common to all actions. The specific action information for this action is shown below.

8.4.3.1. Report Data Action Information

Action Field	Type	Conformance	Description
SuppressResponse	bool	M	do not send a response to this action
SubscriptionId	uint32	O	a SubscriptionId only used in a Subscribe interaction
AttributeReports	list[AttributeReportIB]	O	a list of zero or more attribute data reports
EventReports	list[EventReportIB]	O	a list of zero or more event reports

8.4.3.2. Incoming Read Request and Subscribe Request Action Processing

- Each path indicated by the Report Data action SHALL be a [Concrete Path](#).

Upon receipt of a Read Request action or Subscribe Request action, this process SHALL be followed:

1. Each request path in the AttributeRequests field SHALL be processed as follows:
 - a. If the path does not conform to [Valid Read Attribute Paths](#) then:
 - i. a Status Response with the INVALID_ACTION Status Code SHALL be generated as defined in [Status Response Action](#), a Report Data action SHALL NOT be generated, and this interaction and process SHALL terminate.
 - b. Else if the path is a concrete path, the path SHALL be processed as follows, terminating further processing of the concrete path immediately if a failure [AttributeStatusIB](#) is generated at any point.
 - i. Execute the [ACL Access Granting Algorithm](#) against the concrete path, assuming the **required_privilege** for the element is View, to determine whether the subject would have had at least some access against the concrete path.
 - A. If the outcome is AccessDenied, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_ACCESS Status Code.
 - B. Else if the outcome is AccessRestricted, an [AttributeStatusIB](#) SHALL be generated with the ACCESS_RESTRICTED Status Code.
 - C. Otherwise, continue processing.

- ii. Determine element existence from the concrete path, and apply the following checks:
 - A. If the path indicates a node that is unsupported, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_NODE Status Code.
 - B. Else if the path indicates an endpoint that is unsupported, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_ENDPOINT Status Code.
 - C. Else if the path indicates a cluster that is unsupported, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_CLUSTER Status Code.
 - D. Else if the path indicates an attribute or attribute data field that is unsupported, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_ATTRIBUTE Status Code with the Path field indicating the first unsupported data field (not the entire attribute data path).
 - E. Else if the path indicates an attribute that is not readable, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_READ Status Code.
 - iii. Execute the [ACL Access Granting Algorithm](#) against the concrete path a second time, using the actual [required_privilege](#) for the attribute referenced in the path.
 - A. If the outcome is AccessDenied, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_ACCESS Status Code.
 - B. Else if the outcome is AccessRestricted, an [AttributeStatusIB](#) SHALL be generated with the ACCESS_RESTRICTED Status Code.
 - C. Otherwise, continue processing.
 - iv. If an [AttributeStatusIB](#) was generated, the path SHALL be discarded.
 - c. Else perform [Request Path Expansion](#) and process each expanded existent path as follows:
 - i. If the path indicates attribute data that is not readable, then the path SHALL be discarded.
 - ii. Execute the [ACL Access Granting Algorithm](#) against the path, using the actual [required_privilege](#) for the attribute referenced in the path.
 - A. If the outcome is either AccessDenied or AccessRestricted, then the path SHALL be discarded.
 - B. Otherwise, continue processing.
2. If no error-free existent paths remain, then AttributeRequests are considered empty.
 3. Else each remaining error-free existent path is processed as follows:
 - a. If the DataVersionFilters field indicates [DataVersionFilterIB](#) entries with a Path field that matches the path, where all matching entries have a DataVersion field that matches the data version of the cluster instance in the path, then the path SHALL be ignored
 - b. Else an [AttributeDataIB](#) SHALL be generated with the Data and Path as indicated by the path being processed. If the attribute indicated by the path or any struct field nested inside that attribute (at any level of nesting) is a fabric-scoped list, then any such list within the generated Data SHALL be processed as follows:
 - i. If the FabricFiltered parameter is true, or the attribute indicated by the path is [fabric-sensitive](#), any such list SHALL be generated as a [fabric-filtered list](#) of entries.

- ii. Else, any such list SHALL be generated as an unfiltered list of entries, with each entry indicated as a [fabric-sensitive struct](#).
 - c. Each [AttributeDataIB](#) or [AttributeStatusIB](#) generated from processing [AttributeRequests](#) SHALL be added to the [AttributeReports](#) action field in the Report Data action.
4. Each request path in the [EventRequests](#) field SHALL be processed as follows:
- a. If the path is a concrete path, the path SHALL be processed as follows, terminating further processing of the concrete path immediately if a failure [EventStatusIB](#) is generated at any point.
 - i. Execute the [ACL Access Granting Algorithm](#) against the concrete path, assuming the [required_privilege](#) for the element is View, to determine whether the subject would have had at least some access against the concrete path.
 - A. If the outcome is [AccessDenied](#), an [EventStatusIB](#) SHALL be generated with the [UNSUPPORTED_ACCESS](#) Status Code.
 - B. Else if the outcome is [AccessRestricted](#), an [EventStatusIB](#) SHALL be generated with the [ACCESS_RESTRICTED](#) Status Code.
 - C. Otherwise, continue processing.
 - ii. Determine element existence from the concrete path, and apply the following checks:
 - A. If the path indicates a node that is unsupported, an [EventStatusIB](#) SHALL be generated with the [UNSUPPORTED_NODE](#) Status Code.
 - B. Else if the path indicates an endpoint that is unsupported, an [EventStatusIB](#) SHALL be generated with the [UNSUPPORTED_ENDPOINT](#) Status Code.
 - C. Else if the path indicates a cluster that is unsupported, an [EventStatusIB](#) SHALL be generated with the [UNSUPPORTED_CLUSTER](#) Status Code.
 - D. Else if the path indicates a cluster event that is unsupported, an [EventStatusIB](#) SHALL be generated with the [UNSUPPORTED_EVENT](#) Status Code.
 - iii. Execute the [ACL Access Granting Algorithm](#) against the concrete path a second time, using the actual [required_privilege](#) for the event referenced in the path.
 - A. If the outcome is [AccessDenied](#), an [EventStatusIB](#) SHALL be generated with the [UNSUPPORTED_ACCESS](#) Status Code.
 - B. Else if the outcome is [AccessRestricted](#), an [EventStatusIB](#) SHALL be generated with the [ACCESS_RESTRICTED](#) Status Code.
 - C. Otherwise, continue processing.
 - iv. If an [EventStatusIB](#) was generated, the path SHALL be discarded.
 - b. Else if the path does not conform to [Valid Event Paths](#) then:
 - i. a Status Response with the [INVALID_ACTION](#) Status Code SHALL be generated as defined in [Status Response Action](#),
 - ii. a Report Data action SHALL NOT be generated,
 - iii. and this interaction and process SHALL terminate.
 - c. Else perform [Request Path Expansion](#) and process each expanded existent path as follows:

- i. Execute the [ACL Access Granting Algorithm](#) against the path, using the actual [required_privilege](#) for the event referenced in the path.
 - A. If the outcome is either AccessDenied or AccessRestricted, then the path SHALL be discarded.
 - B. Otherwise, continue processing.
5. If no error free existent paths remain, then EventRequests are considered empty.
6. Else for each unique node indicated in the remaining existent paths:
 - a. Each event record currently queued in the node, in order from lowest to highest event number, SHALL generate an [EventDataIB](#) except for any of the following:
 - i. If the node indicated matches the Node information field of an [EventFilterIB](#) from EventFilters, and the event number is less than the EventMin field in the [EventFilterIB](#).
 - ii. If the event record path does not match a path in the remaining existent event paths.
 - iii. If the event record path is [fabric-sensitive](#), and the [associated fabric](#) does not match the [accessing fabric](#).
 - b. Each information block generated from processing EventRequests SHALL be added to the EventReports action field in the Report Data action.
7. If this action is a Subscribe Request action,
 - a. If both AttributeRequests and EventRequests are empty
 - i. a [Status Response Action](#) with the INVALID_ACTION Status Code SHALL be sent to the initiator,
 - ii. a Report Data action SHALL NOT be generated,
 - iii. and the interaction and process SHALL terminate.
 - b. Else if either MinIntervalFloor or MaxIntervalCeiling is missing, or MinIntervalFloor is greater than MaxIntervalCeiling
 - i. a [Status Response Action](#) with the INVALID_ACTION Status Code SHALL be sent to the initiator,
 - ii. a Report Data action SHALL NOT be generated,
 - iii. and the interaction and process SHALL terminate.
 - c. Else a SubscriptionId which uniquely identifies this subscription on the publisher SHALL be indicated in the Report Data action
8. Else the SubscriptionId SHALL be omitted.

8.4.3.3. Outgoing Report Data Action

1. This action SHALL be unicast.
2. This action MAY have an empty list of AttributeReports and/or EventReports.
3. This action SHALL NOT include any nested attribute data field or nested event data field that is defined as fabric-sensitive, if the [associated fabric](#) for that field does not match the [accessing fabric](#) for the interaction.

4. SuppressResponse MAY be set to TRUE for a Report Data action that initiates a Report transaction that conveys an empty list of AttributeReports and EventReports, otherwise:
 - a. SuppressResponse SHALL be set to TRUE for a Report Data action that is part of a Read transaction.
 - b. SuppressResponse SHALL be set to FALSE for a Report Data action that is part of a Subscribe transaction.
5. This action SHALL be generated as either:
 - a. part of a Read transaction in direct response to a Read Request action.
 - b. part of a Subscribe transaction in direct response to a Subscribe Request action.
 - c. part of a Subscribe interaction as the first action of each synchronized [Report Transaction](#).

8.4.3.4. Incoming Report Data Action

1. Upon receipt of this action, if SuppressResponse is TRUE, a response SHALL NOT be generated;
2. Otherwise a [Status Response Action](#) SHALL be generated with a status code of
 - a. SUCCESS to continue the interaction,
 - b. INVALID_SUBSCRIPTION if the action is part of a Subscribe interaction and the SubscriptionID is invalid,
 - c. FAILURE to terminate the interaction,
 - d. The [Status Response Action](#) SHALL be submitted to the message layer to deliver to the source of this action.

8.5. Subscribe Interaction

The Subscribe interaction is composed of these transactions:

Transaction	Description
Subscribe	start and prime a reporting session
Report	synchronized Report transaction
<i>more reports</i>	continuous Report transactions for the life of the subscription

This interaction allows a subscriber to create a subscription with a publisher on another node for the purposes of receiving data reports from that publisher thereafter, for the duration of the subscription. This allows the subscriber to maintain a coherent snapshot, or twin, of the subscription data as it currently exists on the publisher. The session itself is kept synchronized on both sides through the receipt of timely data reports with the intervals defined by a negotiated maximum interval subscription parameter.

This interaction is started when the initiator (or subscriber), wishes to subscribe to one or more attributes or events located on a target node (the publisher). The attribute data and events requested in the [Subscribe transaction](#) are the subscription data.

This interaction starts by creating a subscription with a [Subscribe transaction](#), which primes the subscriber with initial subscription data. The rest of the subscription is a sequence of [Report transactions](#) initiated by the publisher as defined by parameters of the subscription. Each [Report transaction](#) in a subscription reports changes to the subscription data.

To keep the subscription alive, a [Report transaction](#) is sent from the publisher every maximum interval, or possibly more frequently.

[Report transactions](#) from the publisher are rate limited by the minimum interval subscription parameter, as negotiated between the subscriber and publisher.

The [Subscribe Request](#) action provides boundary values (floor or ceiling) for the publisher to determine the final minimum and maximum interval parameters of the subscription. The time units for these intervals are seconds.

Each [Subscribe](#) interaction is a subscription that is identified by a Subscription ID as generated by the publisher.

- The [Subscribe](#) interaction SHALL start with one [Subscribe transaction](#) followed by a periodic sequence of [Report transactions](#) (see [Report Transaction](#)).
- A [Report transaction](#) SHALL be initiated by a [Report Data](#) action as part of an active subscription for a [Subscribe](#) interaction.
- All [Report Data](#) actions in a [Subscribe](#) interaction SHALL have the same [SubscriptionId](#) parameter value that uniquely identifies the interaction among all subscriptions on the publisher.
- Each [Report transaction](#) in a subscription SHALL report the path for each delta change in the subscription data, including the attribute data that has changed and/or the event that has occurred, since the last [Report transaction](#), with the exception of attribute data with the [Changes Omitted \(C\)](#) quality.
- Each [Report transaction](#) initiated by the publisher SHALL complete successfully before another [Report transaction](#) is initiated by the publisher.
- Each [Report transaction](#) SHALL NOT be initiated by the publisher until the minimum interval has expired since the last [Report transaction](#) in the subscription.
- Attribute changes SHALL be delivered as soon as possible, taking into account the minimum interval.
- Events SHALL always be queued and buffered. Each [Report](#) containing events SHALL deliver queued events without reordering the queue. Queued events MAY be opportunistically delivered whenever some other activity triggers a [Report transaction](#). Absent any such triggers, queued events SHALL be delivered in a [Report transaction](#) generated at the maximum interval.
- When the [IsUrgent](#) flag is FALSE or absent for a subscription's event path in the [EventPathIB](#), event queueing does not automatically trigger a [Report transaction](#).
- When the [IsUrgent](#) flag is TRUE for a subscription's event path in the [EventPathIB](#), the queueing of such an event SHALL trigger a [Report transaction](#) for the subscription, subject to all [Report transaction](#) rules. This [Report transaction](#) will report the events that have been queued by the time the [Report transaction](#) happens.
- If the subscriber does not receive a [Report transaction](#) within the maximum interval from the

last Report Data, the subscriber SHALL terminate the Subscribe interaction.

- If a node receives a Report Data action with an inactive SubscriptionId, a Status Response action SHALL be sent with an INVALID_SUBSCRIPTION Status Code.
- If, in response to a Report Data action, the publisher receives a Status Response action with a status code that is not SUCCESS, the publisher SHALL terminate the Subscribe interaction.
- If the publisher does not receive a Status Response action in response to a Report Data action with SuppressResponse set to FALSE, the publisher MAY terminate the Subscribe interaction or SHALL re-synchronize the subscription in the next Report Data transaction by:
 - Including all subscription data to re-prime the subscription, or
 - Including all deltas since the last successful Report Data transaction.
- The subscriber MAY terminate the subscription and interaction by responding with a Status Response action with an INVALID_SUBSCRIPTION Status Code.
- The publisher MAY terminate the subscription and interaction by not generating a [Report transaction](#) within the maximum interval.
- When a Subscribe interaction is terminated on the publisher or subscriber, the subscription, identified by a SubscriptionId, SHALL also be terminated.

8.5.1. Subscribe Transaction

Action	Action Flow	Description
Subscribe Request	Initiator ⇒ Target	list of event and attribute data identifiers supported on a server cluster
Report Data	Initiator ⇐ Target	primed published data
Status Response	Initiator ⇒ Target	success, or otherwise an error to terminate the subscription
Subscribe Response	Initiator ⇐ Target	provides subscription parameters

8.5.2. Subscribe Request Action

Subscribe Request action is a first action. See [Section 8.2.5, “Common Action Behavior”](#) for behavior common to all actions. The specific action information for this action is shown below.

8.5.2.1. Subscribe Request Action Information

Action Field	Type	Conformance	Description
KeepSubscriptions	bool	M	false to terminate existing subscriptions from initiator

Action Field	Type	Conformance	Description
MinIntervalFloor	uint16	M	the requested minimum interval boundary floor in seconds
MaxIntervalCeiling	uint16	M	the requested maximum interval boundary ceiling in seconds
AttributeRequests	list[AttributePathIB]	O	a list of zero or more request paths to cluster attribute data
DataVersionFilters	list[DataVersionFilterIB]	AttributeRequests	a list of zero or more cluster instance data versions
EventRequests	list[EventPathIB]	O	a list of zero or more request paths to cluster events
EventFilters	list[EventFilterIB]	EventRequests	a list of zero or more minimum event numbers per specific node
FabricFiltered	bool	M	limits the data read within fabric-scoped lists to the accessing fabric

8.5.2.2. Outgoing Subscribe Request Action

- This action SHALL initiate a Subscribe interaction.
- A Subscribe Request action SHALL be unicast from the subscriber to the publisher.
- This action SHALL be generated to initiate a Subscribe interaction (see [Section 8.5, “Subscribe Interaction”](#)).
- This action SHALL include a requested ceiling (highest) maximum interval value as MaxIntervalCeiling.
- This action SHALL include a requested floor (lowest) minimum interval value as MinIntervalFloor.

NOTE

If the publisher is an [intermittently connected device](#), the MinIntervalFloor SHOULD be 0. To avoid needing to switch to Active Mode to process the event and a second time to send the information, an [ICD](#) MAY try to synchronize its updates towards various subscribers so that it only needs to go to Active Mode once and can stay idle for a full idle interval. If the MinIntervalFloor is too high, this limits the freedom for an ICD to synchronize its updates. The higher the MinIntervalFloor is, the higher the likelihood an ICD will not be able to send updates to various subscribers in a synchronized pattern, which might negatively impact energy usage

(e.g. battery life).

- At least one attribute or event SHALL be indicated in the action.
- A valid [AttributePathIB](#) SHALL be one in the table [Valid Read Attribute Paths](#).
- A valid [EventPathIB](#) SHALL be one in the table [Valid Event Paths](#).
- A path indicated in AttributeRequests or EventRequests SHALL NOT target a group.

8.5.2.3. Incoming Subscribe Request Action

- If KeepSubscriptions is FALSE, all existing or pending subscriptions on the publisher for this subscriber SHALL be terminated.
- This layer SHALL process the Subscribe Request action as defined in [Incoming Read Request and Subscribe Request Action Processing](#).

8.5.3. Subscribe Response Action

Subscribe Response action is the last following action in a [Subscribe Transaction](#). This action activates the subscription. See [Section 8.2.5, “Common Action Behavior”](#) for behavior common to all actions. The specific action information for this action is shown below (see [Section 8.5, “Subscribe Interaction”](#)).

8.5.3.1. Subscribe Response Action Information

Action Field	Type	Conformance	Description
SubscriptionId	uint32	M	identifies the subscription
MaxInterval	uint16	M	the final maximum interval for the subscription in seconds

8.5.3.2. Outgoing Subscribe Response Action

- Upon receipt of a successful Status Response action from the subscriber for the Report Data action that primes the subscription, this action SHALL be generated and submitted to the message layer to send to the subscriber.
- This action SHALL be unicast.
- The SubscriptionId value SHALL be the same as the one used in Report Data generated to prime this subscription.
- The publisher SHALL compute an appropriate value for the MaxInterval field in the action. This SHALL respect the following constraint: $\text{MinIntervalFloor} \leq \text{MaxInterval} \leq \text{MAX}(\text{SUBSCRIPTION_MAX_INTERVAL_PUBLISHER_LIMIT}, \text{MaxIntervalCeiling})$
- Upon sending a Subscribe Response action, the subscription, as indicated by the SubscriptionId, SHALL become active on the publisher with a min interval equal to the requested MinIntervalFloor and a max interval equal to the MaxInterval field in the response.

8.5.3.3. Incoming Subscribe Response Action

- Upon receipt of a Subscribe Response action, the subscription, as indicated by the SubscriptionId, SHALL become active to the subscriber.

8.5.3.4. Subscription Activation

- The paths to the subscription data SHALL only be error free existent paths generated from processing the Subscribe Request.

The EventFilters and DataVersionFilters fields in the Subscribe Request are one time parameters for the priming of the subscription.

- Subsequent ReportData actions, as part of the subscription, SHALL include the latest:
 - EventNo associated with each node generating new events.
 - DataVersion associated with each cluster where there are data changes.
- The FabricFiltered parameter from the Subscribe Request SHALL remain in effect for all data reported during the interaction.
- Upon subscription activation, the minimum and maximum interval parameters SHALL take effect to determine the timing and expectation of subsequent Report transactions.

8.6. Report Transaction

There is no Report interaction. A Report transaction is part of a Subscribe interaction. See [Section 8.5, “Subscribe Interaction”](#) for details.

The valid Report transactions are:

8.6.1. Report Transaction Non-Empty

Action	Action Flow	Description
Report Data	Initiator $\Leftarrow$ Target	report of data and/or events with SuppressResponse set to FALSE
Status Response	Initiator $\Rightarrow$ Target	an error ends the interaction

8.6.2. Report Transaction Empty

Action	Action Flow	Description
Report Data	Initiator $\Leftarrow$ Target	report with no data or events with SuppressResponse set to TRUE

8.7. Write Interaction

This interaction is started when an initiator wishes to modify the values of one or more attributes located on one or more nodes. An optional Timed Request action defines a Timeout interval that starts at the sending of the Status Response action.

8.7.1. Write Transaction

- A Write interaction SHALL consist of one of the transactions shown below.

8.7.1.1. Timed Write Transaction

Action	Action Flow	Description
Timed Request	Initiator $\Rightarrow$ Target	time interval defined to send Write Request action
Status Response	Initiator $\Leftarrow$ Target	confirmation
Write Request	Initiator $\Rightarrow$ Target	data to modify
Write Response	Initiator $\Leftarrow$ Target	with errors or success from Write Request action

8.7.1.2. Untimed Write Transaction

Action	Action Flow	Description
Write Request	Initiator $\Rightarrow$ Target	data to modify
Write Response	Initiator $\Leftarrow$ Target	with errors or success from Write Request action

- If there is a preceding successful Timed Request action, the following Write Request action SHALL be received before the end of the Timeout interval.
- If there is a preceding successful Timed Request action, the Timeout interval SHALL start when the Status Response action acknowledging the Timed Request action with a success code is sent.
- If there is a preceding successful Timed Request action, the Write Request action SHALL be unicast.
- If there is not a preceding successful Timed Request action, the Write Request action MAY be groupcast.
- A client MAY choose to use a Timed Write transaction even if the attribute does not have the Timed Interaction quality.
- The server SHALL support a Timed Write transaction for all writeable attributes.

8.7.2. Write Request Action

This action is either the first action of the Write transaction or it follows a successful Timed Request action. See [Section 8.2.5, “Common Action Behavior”](#) for behavior common to all actions. The specific action information for this action is shown below.

8.7.2.1. Write Request Action Information

Action Field	Type	Conformance	Description
SuppressResponse	bool	M	do not send a response to this action
TimedRequest	bool	M	flag action as part of a timed write transaction
WriteRequests	list[AttributeDataIB]	M	a list of one or more path and data tuples.

8.7.2.2. Outgoing Write Request Action

- This action SHALL be generated as the first action in a Write transaction, or following a Timed Request action and successful Status Response action.
- If this action is part of a Timed Write transaction, TimedRequest SHALL be TRUE, else FALSE.
- If not part of a Timed Write transaction, this action MAY be groupcast.
- If this action is groupcast, SuppressResponse SHALL be TRUE.

8.7.2.3. Incoming Write Request Action

- If this action is not able to be executed because the maximum supported number of Write interactions is already in progress, then a Status Response action with the BUSY Status Code SHALL be submitted to the message layer and this interaction SHALL terminate.
- If this action is part of a Timed Write transaction, and the Timeout has expired from the preceding Timed Request action, then a Status Response action with the TIMEOUT Status Code SHALL be submitted to the message layer and this interaction SHALL terminate.

NOTE Devices certified prior to 1.4 MAY return UNSUPPORTED_ACCESS for this condition.

- If this action is part of a Timed Write transaction, and this action has TimedRequest set to FALSE, then a Status Response action with the TIMED_REQUEST_MISMATCH Status Code SHALL be submitted to the message layer and this interaction SHALL terminate.

NOTE Devices certified prior to 1.4 MAY return UNSUPPORTED_ACCESS for this condition.

- If this action is marked with TimedRequest as TRUE but this action is not part of a Timed Write transaction (i.e. there was no corresponding Timed Request action prior to it matching the same TransactionID), then a Status Response action with the TIMED_REQUEST_MISMATCH Status Code SHALL be submitted to the message layer and this interaction SHALL terminate.

NOTE Devices certified prior to 1.4 MAY return UNSUPPORTED_ACCESS for this condition.

See [Outgoing Write Response Action](#) for building a Write Response action and executing the Write Request action.

- If this action was unicast and SuppressResponse is FALSE, a Write Response action SHALL be

generated and submitted to the message layer to send to the initiator, otherwise no Write Response SHALL be sent.

8.7.3. Write Response Action

This action is a following action for a Write Request action. See [Section 8.2.5, “Common Action Behavior”](#) for behavior common to all actions. The specific action information for this action is shown below.

8.7.3.1. Write Response Action Information

Action Field	Type	Conformance	Description
WriteResponses	list[AttributeStatusIB]	O	a list of zero or more concrete paths indicating errors or successes

8.7.3.2. Outgoing Write Response Action

1. This action SHALL be unicast.
2. Each request path in the WriteRequests field of a Write Request SHALL be processed as follows:
 - a. If the path does not conform to [Valid Write Attribute Paths](#) then:
 - i. a Status Response with the INVALID_ACTION Status Code SHALL be generated as defined in [Status Response Action](#), a Write Response action SHALL NOT be generated, and this interaction and process SHALL terminate.
 - b. Else if the path is a concrete path, the path SHALL be processed as follows, terminating further processing of the concrete path immediately if a failure [AttributeStatusIB](#) is generated at any point (one that is not a status of SUCCESS).
 - i. Execute the [ACL Access Granting Algorithm](#) against the concrete path, assuming the **required_privilege** for the element is View, to determine whether the subject would have had at least some access against the concrete path.
 - A. If the outcome is AccessDenied, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_ACCESS Status Code.
 - B. Else if the outcome is AccessRestricted, an [AttributeStatusIB](#) SHALL be generated with the ACCESS_RESTRICTED Status Code.
 - C. Otherwise, continue processing.
 - ii. Determine element existence from the concrete path, and apply the following checks:
 - A. If the path indicates a node that is unsupported, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_NODE Status Code.
 - B. Else if the path indicates an endpoint that is unsupported, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_ENDPOINT Status Code.
 - C. Else if the path indicates a cluster that is unsupported, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_CLUSTER Status Code.
 - D. Else if the path indicates an attribute or attribute data field that is unsupported, an

- [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_ATTRIBUTE Status Code with the Path field indicating only the path to the first unsupported data field (not the entire attribute data path).
- E. Else if the path indicates an attribute that is not writable, an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_WRITE Status Code.
- iii. Execute the [ACL Access Granting Algorithm](#) against the concrete path a second time, using the actual [required_privilege](#) for the attribute referenced in the path.
- A. If the outcome is AccessDenied, a [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_ACCESS Status Code.
- B. Else if the outcome is AccessRestricted, a [AttributeStatusIB](#) SHALL be generated with the ACCESS_RESTRICTED Status Code.
- C. Otherwise, continue processing.
- iv. If the path indicates specific attribute data that requires a Timed Write transaction to write and this action is not part of a Timed Write transaction, an [AttributeStatusIB](#) SHALL be generated with the NEEDS_TIMED_INTERACTION Status Code.
- v. Else if the attribute in the path indicates a fabric-scoped list and there is no [accessing fabric](#), an [AttributeStatusIB](#) SHALL be generated with the UNSUPPORTED_ACCESS Status Code, with the Path field indicating only the path to the attribute.
- vi. Else if the DataVersion field of the [AttributeDataIB](#) is present and does not match the data version of the indicated cluster instance, an [AttributeStatusIB](#) SHALL be generated with the DATA_VERSION_MISMATCH Status Code.
- vii. If the above processing generated an [AttributeStatusIB](#), the path SHALL be discarded.
- viii. Else perform [Write Path Data Process](#)
- c. Else perform [Request Path Expansion](#) and process each expanded existent path as follows:
- i. If the path indicates attribute data that is not writable, then the path SHALL be discarded.
- ii. Execute the [ACL Access Granting Algorithm](#) against the path, using the actual [required_privilege](#) for the attribute referenced in the path.
- A. If the outcome is either AccessDenied or AccessRestricted, then the path SHALL be discarded.
- B. Otherwise, continue processing.
- iii. Else if the path indicates specific attribute data that requires a Timed Write transaction to write and this action is not part of a Timed Write transaction, then the path SHALL be discarded.
- iv. Else perform [Write Path Data Process](#)

8.7.3.3. Write Path Data Process

1. If the path indicates a fabric-scoped list or list entry, it SHALL be processed as a [fabric-filtered list of fabric-scoped structs](#).
2. Each data field indicated by the path, SHALL be processed in the order conveyed as follows:

- a. If a data field is not within the constraints defined by the cluster specification, an **AttributeStatusIB** SHALL be generated with the CONSTRAINT_ERROR Status Code, with the Path field duplicating the path.
- b. Otherwise perform the following:
 - i. The data field SHALL be changed to the data indicated with the path.
 - ii. An **AttributeStatusIB** SHALL be generated with the SUCCESS Status Code, with the Path field duplicating the path.

8.7.3.4. Write Response Generation

- Each **AttributeStatusIB** generated from processing the WriteRequests field of the Write Request action, SHALL be added to the WriteResponses action field of the Write Response action.

NOTE

An empty WriteResponses would occur if all paths were wildcard or group paths that expand to non-accessible data.

8.7.3.5. Incoming Write Response Action

Upon receipt of this action, the action information SHALL be submitted to the layer above.

8.7.4. Timed Request Action

This action is a first action of a transaction. The specific action information for this action is shown below.

This action informs the receiver that another action will be sent in the same direction, within the same transaction, and within a specified Timeout interval. The Timeout interval SHALL start when the Status Response action acknowledging the Timed Request action with a success code is sent.

8.7.4.1. Timed Request Action Information

Action Field	Type	Conformance	Description
Timeout	uint16	M	an interval, in milliseconds, to expect a following action

8.7.4.2. Outgoing Timed Request Action

- This action SHALL be generated as an optional first action in a Write or Invoke transaction.
- This action SHALL be unicast.

8.7.4.3. Incoming Timed Request Action

- Upon receipt of this action, this layer SHALL construct and send a Status Response action with SUCCESS to the initiator.
- This layer SHALL expect a Write Request or Invoke Request action within Timeout milliseconds of sending the Status Response action.

8.8. Invoke Interaction

This interaction is generated when a device wishes to invoke one or more cluster specific commands on one or more nodes. Cluster commands are defined as either client to server or server to client. Invoke Request action SHALL support [group paths](#) and SHOULD support [wildcard paths](#). Invoke Response action does not support [wildcard paths](#).

8.8.1. Invoke Transaction

- The Invoke interaction SHALL consist of one of the Invoke transactions shown below

8.8.1.1. Timed Invoke Transaction

Action	Action Flow	Description
Timed Request	Initiator ⇒ Target	time interval defined to send Invoke Request action
Status Response	Initiator ⇐ Target	confirmation
Invoke Request	Initiator ⇒ Target	commands to invoke
Invoke Response	Initiator ⇐ Target	response(s) defined by the cluster specification

8.8.1.2. Untimed Invoke Transaction

Action	Action Flow	Description
Invoke Request	Initiator ⇒ Target	commands to invoke
Invoke Response	Initiator ⇐ Target	response(s) defined by the cluster specification

- A client MAY choose to use a Timed Invoke transaction even if the command does not have the Timed Interaction quality.
- The server SHALL support a Timed Invoke transaction for all commands.

8.8.2. Invoke Request Action

- This action SHALL be generated as the first action in an Invoke transaction, or following a Timed Request action and successful Status Response action.
- If there is a preceding successful Timed Request action, the following Invoke Request action SHALL be received before the end of the Timeout interval.
- If there is a preceding successful Timed Request action, the Timeout interval SHALL start when the Status Response action acknowledging the Timed Request action with a success code is sent.
- Each Invoke Request and Invoke Response action in a Timed Invoke transaction SHALL be unicast.
- If not part of a Timed Invoke transaction, this action MAY be groupcast.

- If a cluster command is defined to be invoked as the result of a groupcast, the command SHALL be invoked with an Invoke Request action which SHALL start a new transaction.
- Each path in an Invoke Request or Invoke Response action SHALL indicate a server cluster.

Invoke Request action is either the first action of the Invoke transaction or it follows a successful Timed Request action. See [Section 8.2.5, “Common Action Behavior”](#) for behavior common to all actions. The specific action information for this action is shown below.

8.8.2.1. Invoke Request Action Information

Action Field	Type	Constraint	Conformance	Description
SuppressResponse	bool		M	do not send a response to this action
TimedRequest	bool		M	flag action as part of a timed invoke transaction
InvokeRequests	list[CommandDataIB]	desc	M	cluster command(s)

8.8.2.2. Outgoing Invoke Request Action

1. This action SHALL be generated as the first action in an Invoke transaction, or following a Timed Request action and successful Status Response action.
2. If this action is part of a Timed Invoke transaction, TimedRequest SHALL be TRUE, else FALSE.
3. A valid [CommandDataIB](#) SHALL be one in the table [Valid Command Paths](#).
4. If InvokeRequests has 1 CommandDataIB,
 - a. the path indicated in that CommandDataIB MAY target a group or a wildcard.
5. Else (InvokeRequests has more than 1 CommandDataIB),
 - a. each path indicated in InvokeRequests SHALL be a concrete (non-wildcard) path and SHALL NOT target a group, and
 - b. each path indicated in InvokeRequests SHALL be unique.
 - c. The total number of CommandDataIB in InvokeRequests SHALL NOT exceed what the server indicates it can handle (See [Section 11.1.5.23, “MaxPathsPerInvoke”](#)).
 - d. Care must be taken to avoid issuing a set of commands with some data dependencies between them, since there is no guarantee of "happens-after" in the ordering.

NOTE

For InteractionModelRevision $\leq$ 11, having more than one command in InvokeRequests was marked as provisional.

8.8.2.3. Incoming Invoke Request Action

1. If this action is part of a Timed Invoke transaction, and the Timeout has expired from the pre-

ceding Timed Request action, then a Status Response action with the TIMEOUT Status Code SHALL be submitted to the message layer and this interaction SHALL terminate.

NOTE Devices certified prior to 1.4 MAY return UNSUPPORTED_ACCESS for this condition.

2. If this action is part of a Timed Invoke transaction, and this action has TimedRequest set to FALSE, then a Status Response action with the TIMED_REQUEST_MISMATCH Status Code SHALL be submitted to the message layer and this interaction SHALL terminate.
3. If this action is marked with TimedRequest as TRUE, but this action is not part of a Timed Invoke transaction (i.e. there was no immediately previous Timed Invoke action), then a Status Response action with the TIMED_REQUEST_MISMATCH Status Code SHALL be submitted to the message layer and this interaction SHALL terminate.
4. If this action contains more [CommandDataIB](#) elements in the InvokeRequests list than are supported by the device (see [Section 11.1.5.23, “MaxPathsPerInvoke”](#)), then a Status Response action with the INVALID_ACTION Status Code SHALL be submitted to the message layer and this interaction SHALL terminate.
5. Each request path in the InvokeRequests field SHALL be processed as follows:
 - a. If the path does not conform to [Valid Command Paths](#) then:
 - i. a Status Response with the INVALID_ACTION Status Code SHALL be generated as defined in [Status Response Action](#), an Invoke Response action SHALL NOT be generated, and this interaction and process SHALL terminate.
 - b. Else if the path is a concrete path, the path SHALL be processed as follows, terminating further processing of the concrete path immediately if a failure [CommandStatusIB](#) is generated at any point (one that is not a status of SUCCESS). If a failure [CommandStatusIB](#) is generated, its CommandPath field SHALL be a duplicate of the concrete path processed, including the command ID of the original concrete path.
 - i. Execute the [ACL Access Granting Algorithm](#) against the concrete path, assuming the [required_privilege](#) for the element is Operate, to determine whether the subject would have had at least some access against the concrete path.
 - A. If the outcome is AccessDenied, a [CommandStatusIB](#) SHALL be generated with the UNSUPPORTED_ACCESS Status Code.
 - B. Else if the outcome is AccessRestricted, a [CommandStatusIB](#) SHALL be generated with the ACCESS_RESTRICTED Status Code.
 - C. Otherwise, continue processing.
 - ii. Determine element existence from the concrete path, and apply the following checks:
 - A. If the path indicates a node that is unsupported, a [CommandStatusIB](#) SHALL be generated with the UNSUPPORTED_NODE Status Code.
 - B. Else if the path indicates an endpoint that is unsupported, a [CommandStatusIB](#) SHALL be generated with the UNSUPPORTED_ENDPOINT Status Code.
 - C. Else if the path indicates a cluster that is unsupported, a [CommandStatusIB](#) SHALL be generated with the UNSUPPORTED_CLUSTER Status Code.
 - D. Else if the path indicates a command that is unsupported, a [CommandStatusIB](#)

SHALL be generated with the UNSUPPORTED_COMMAND Status Code.

- iii. Execute the [ACL Access Granting Algorithm](#) against the concrete path a second time, using the actual [required_privilege](#) for the command referenced in the path.
 - A. If the outcome is AccessDenied, a [CommandStatusIB](#) SHALL be generated with the UNSUPPORTED_ACCESS Status Code.
 - B. Else if the outcome is AccessRestricted, a [CommandStatusIB](#) SHALL be generated with the ACCESS_RESTRICTED Status Code.
 - C. Otherwise, continue processing.
 - iv. If the command in the path has the [Large Message Quality](#) and was received on a transport that is not capable of transporting Large Messages, a [CommandStatusIB](#) SHALL be generated with the INVALID_TRANSPORT_TYPE Status Code.
 - v. Else if the command in the path is fabric-scoped and there is no accessing fabric, a [CommandStatusIB](#) SHALL be generated with the UNSUPPORTED_ACCESS Status Code.
 - vi. Else if the command in the path requires a Timed Invoke transaction to invoke and this action is not part of a Timed Invoke transaction, a [CommandStatusIB](#) SHALL be generated with the NEEDS_TIMED_INTERACTION Status Code.
 - c. Else perform [Request Path Expansion](#) and process each expanded existent path as follows:
 - i. Execute the [ACL Access Granting Algorithm](#) against the path, using the actual [required_privilege](#) for the command referenced in the path.
 - A. If the outcome is either AccessDenied or AccessRestricted, then the path SHALL be discarded.
 - B. Otherwise, continue processing.
 - ii. Else if the command in the path requires a [Large Message](#) transport and this action was not received on a transport capable of Large Messages, then the path SHALL be discarded.
 - iii. Else if the command in the path is fabric-scoped and there is no accessing fabric, then the path SHALL be discarded.
 - iv. Else if the command in the path requires a Timed Invoke transaction to invoke and this action is not part of a Timed Invoke transaction, then the path SHALL be discarded.
6. Each command in the remaining and error-free concrete command paths SHALL be executed, as defined in [Invoke Execution](#).

Invoke Execution

1. The command SHALL be executed as defined in the cluster specification, and the following applies:
 - a. For each data field in CommandFields:
 - i. If a mandatory data field is missing, or incoming data cannot be mapped to the expected data type for a field, a [CommandStatusIB](#) SHALL be generated with an error status of INVALID_COMMAND, even if the cluster defines another type of response.
 - ii. If a data field violates expected constraints, a [CommandStatusIB](#) SHOULD be generated

with an error status of `CONSTRAINT_ERROR`.

- b. If the cluster specification defines a following command in response to the command, a [CommandDataIB](#) SHALL be generated for the following command with these fields:
 - i. CommandFields defined for the following command
 - ii. ClusterPath field of the CommandPath field that is a duplicate of the command path processed, up to the cluster ID
 - iii. Command field of the CommandPath field that is the command ID of the following command
 - c. Else if the cluster specification defines a success or error status as a response (sometimes specified as a Default Response), a [CommandStatusIB](#) SHALL be generated with these fields:
 - i. CommandPath that is a duplicate of the concrete command path processed
 - ii. Status as defined in the cluster specification
2. If the InvokeRequest contains more than 1 [CommandDataIB](#):
- a. If the InvokeRequest contains certain commands and endpoints for which the server has an efficient way to process those commands:
 - i. The server MAY do so, even if this would deviate from the general concept of "execution started in order" for the commands which are grouped.
 - b. Else:
 - i. The Invoke Execution SHALL be started for each command in the same order as conveyed and expanded.
 - c. Actions MAY run concurrently, and the protocol offers no guarantees about completion of one action preceding the start of another.
 - d. In case one or more of these [CommandDataIB](#) triggers an asynchronous action, any given [Invoke Execution](#) may not complete in the same order as the corresponding [CommandDataIB](#) in the request. As a consequence, the order of generated [CommandDataIB](#) and/or [CommandStatusIB](#) MAY be different from the order of the corresponding [CommandDataIB](#) in the InvokeRequest.
3. An [Invoke Response Generation](#) SHALL be performed on completion of all valid commands. A partial [Invoke Response Generation](#) MAY be performed at any time when processing multiple commands, provided it carries at least a single response.

8.8.3. Invoke Response Action

Invoke Response action is a following action for an Invoke Request action. See [Section 8.2.5, "Common Action Behavior"](#) for behavior common to all actions. The specific action information for this action is shown below.

8.8.3.1. Invoke Response Action Information

Action Field	Type	Conformance	Description
SuppressResponse	bool	M	Ignored by client

Action Field	Type	Conformance	Description
InvokeResponses	list[InvokeResponseIB]	M	command response or status

The SuppressResponse field is mandatory to maintain backwards compatibility with older clients. However, as Matter does not support responses to InvokeResponse actions, this field has no effect.

8.8.3.2. Outgoing Invoke Response Action

- This action SHALL be generated in response to an Invoke Request action, after all valid commands are executed.
- A valid [InvokeResponseIB](#) SHALL only indicate a concrete path.

Invoke Response Generation

- If the previous action in this transaction action is groupcast, this process and transaction terminate with no response.
- Else if SuppressResponse is FALSE, or a [CommandDataIB](#) was generated, an Invoke Response action SHALL be generated as follows:
 - Each generated [CommandStatusIB](#) and [CommandDataIB](#) SHALL be included in an [InvokeResponseIB](#) in the InvokeResponses action field.
 - An Invoke Response action SHALL be submitted to the message layer to send to the source of the previous action.

8.8.3.3. Incoming Invoke Response Action

- Upon receipt of this action, each entry in InvokeResponses SHALL be processed, in order, as follows:
 - If the entry is a [CommandStatusIB](#), it SHALL be submitted to the layer above.
 - Else if the response is a [CommandDataIB](#):
 - If the ClusterPath field of the CommandPath field does not duplicate or match a wildcard of one of the paths in the previous action of the interaction, the entry SHALL be discarded.
 - Else if the command ID value of the Command field of the CommandPath field is not expected by the cluster instance, the entry SHALL be discarded.
 - Else it SHALL be submitted to the layer above.
- For each entry in the initial Invoke Request action that triggered this incoming Invoke Response action, if there isn't a corresponding entry in InvokeResponses, a [CommandStatusIB](#) with the NO_COMMAND_RESPONSE status SHALL be submitted to the layer above.

NOTE

Sending the NO_COMMAND_RESPONSE should only be done after receiving all encoded/transported messages, which indicates that the interaction is completed and the response is missing for those requests.

8.9. Common Action Information Blocks and Paths

Shown below are common information blocks used to submit actions to, and receive actions from, the message layer.

8.9.1. Path Information

- Logically for this specification, the Node or Group SHALL be present in all paths.
- When an outgoing path indicates the same Node or Group as the action target, the message layer MAY optimize the path by removing Node or Group.
- When an incoming path does not indicate a Node or Group, the message layer SHALL add the action target (Node or Group) to the path.

See [Path](#) for more information.

8.9.1.1. ClusterPathIB

Path Field	Type	Quality	Conformance
Group	group-id		!Endpoint
Node	node-id		!Group
Endpoint	endpoint-no	A	!Group
Cluster	cluster-id	A	M

8.9.2. Attribute Information Blocks

8.9.2.1. AttributePathIB

An attribute path supports nesting levels deeper than the attribute with the NestedPath field below. The NestedPath indicates zero or more nesting levels to establish the nesting depth of the attribute path.

- An attribute path MAY indicate data at any nesting depth in the attribute data.
- If the Attribute is indicated as a wildcard then the path SHALL indicate all attributes of the cluster.
- If the Attribute is indicated as a wildcard, then NestedPath SHALL be empty.
- If the final nesting level in an attribute path indicates a [collection data type](#), then the path SHALL indicate all collection data, plus any deeper nested data as part of the collection.
- If the WildcardPathFlags field is absent, it SHALL be considered to be an empty map with all flags unset.

Path Field	Type	Quality	Conformance
Cluster	ClusterPathIB		M
Attribute	attrib-id	A	M

Path Field	Type	Quality	Conformance
NestedPath	list[NestingLevelIB]		Attribute != <i>wildcard</i>
WildcardPathFlags	WildcardPathFlags- Bitmap		P, O
WildcardFilterConfigurationVersion	uint32		P, O

8.9.2.2. NestingLevelIB

This defines a deeper nesting level of an attribute path to [collection data](#), such as a struct field, or list entry.

Path Field	Type	Quality	Conformance
FieldID	field-id	A	O.a
ListIndex	entry-idx	A	O.a

8.9.2.3. WildcardPathFlags

This field is used in conjunction with the [WildcardFilterConfigurationVersion](#) to suppress reports for specified attributes.

- The meaning of the flags present in this field is defined in [WildcardPathFlagsBitmap](#).
- The absence of the [WildcardPathFlags](#) field SHALL indicate that all flags are false.
- If the [WildcardPathFlags](#) value is zero, it SHOULD be omitted by clients.
- If the [WildcardPathFlags](#) value is specified, a [WildcardFilterConfigurationVersion](#) SHOULD also be specified.
- The effect of WildcardPathFlags and the [WildcardFilterConfigurationVersion](#) on AttributePathIB processing is described in [Attribute Wildcard Path Flags](#).

8.9.2.4. WildcardFilterConfigurationVersion

This field is used in conjunction with the [WildcardPathFlags](#) to suppress reports for specified attributes. Its value limits when certain flags in the WildcardPathFlags filter apply:

- If the [ConfigurationVersion](#) in the Basic Information cluster is greater than the value specified in this field, then the device configuration has changed, and certain [WildcardPathFlags](#) do not apply. The effect of WildcardPathFlags and the [WildcardFilterConfigurationVersion](#) on AttributePathIB processing is described in detail in [Attribute Wildcard Path Flags](#).
- If the WildcardFilterConfigurationVersion field is omitted or the value is 0xFFFF_FFFF, the value of the WildcardFilterConfigurationVersion SHALL be considered to be the same value as the current [ConfigurationVersion](#) attribute.
- If the [WildcardFilterConfigurationVersion](#) value is specified, a [WildcardPathFlags](#) SHOULD also be specified.

8.9.2.5. WildcardPathFlagsBitmap

This data type is derived from map32.

The WildcardPathFlagsBitmap indicates flags that apply to the path, affecting wildcard expansion.

The following flags are defined:

Bit	Name	Summary
0	WildcardSkipRootNode	Skip the Root Node endpoint (endpoint 0) during wildcard expansion.
1	WildcardSkipGlobalAttributes	Skip several large global attributes during wildcard expansion.
2	WildcardSkipAttributeList	Skip the AttributeList global attribute during wildcard expansion.
3	DoNotUse	
4	WildcardSkipCommandLists	Skip the AcceptedCommandList and GeneratedCommandList global attributes during wildcard expansion.
5	WildcardSkipCustomElements	Skip any manufacturer-specific clusters or attributes during wildcard expansion.
6	WildcardSkipFixedAttributes	Skip any Fixed (F) quality attributes during wildcard expansion.
7	WildcardSkipChangesOmittedAttributes	Skip any Changes Omitted (C) quality attributes during wildcard expansion.
8	WildcardSkipDiagnosticsClusters	Skip all clusters with the Diagnostics (K) quality during wildcard expansion.

8.9.2.6. Valid Read Attribute Paths

This table is valid for [AttributePathIB](#) for a Read Request action or a Subscribe Request action. See [AttributePathIB](#) for more conformance restrictions.

The table defines valid attribute data path with wildcard and non-wildcard combinations.

Node	Endpoint	Cluster	Attribute	Attribute Data Requested
Specific	Wildcard	Wildcard	Wildcard	all attribute data from all clusters on all endpoints on a specific node
Specific	Wildcard	Wildcard	Specific	specific global attribute data or field for all clusters on all endpoints on a specific node (e.g. ClusterRevision)
Specific	Wildcard	Specific	Wildcard	all attribute data from a specific cluster on all endpoints on a specific node (e.g. Descriptor cluster)
Specific	Wildcard	Specific	Specific	a specific attribute data or field from a specific cluster on all endpoints on a specific node (e.g. Descriptor cluster)
Specific	Specific	Wildcard	Wildcard	all attribute data from all clusters on a specific endpoint on a specific node
Specific	Specific	Wildcard	Specific	a specific global attribute data or field for all clusters on a specific endpoint on a specific node (e.g. ClusterRevision)
Specific	Specific	Specific	Wildcard	all attribute data from a specific cluster on a specific endpoint a specific node

Node	Endpoint	Cluster	Attribute	Attribute Data Requested
Specific	Specific	Specific	Specific	a specific attribute data or field from a specific cluster on a specific endpoint on a specific node

8.9.2.7. Valid Write Attribute Paths

This table is valid for Path field of a [AttributeDataIB](#) for a Write Request action. See [AttributeDataIB](#) and [AttributePathIB](#) for more conformance restrictions.

The table defines valid attribute data paths including combinations with wildcards, non-group (node & endpoint), and group paths. A blank entry means the element does not exist in the path.

Node	Endpoint	Group	Cluster	Attribute	Attribute Data Provided
Specific	Wildcard		Specific	Specific	a specific attribute data or field from a specific cluster on all endpoints on a specific node
Specific	Specific		Specific	Specific	a specific attribute data or field from a specific cluster on a specific endpoint on a specific node
		Specific	Specific	Specific	a specific attribute data or field from a specific cluster on a specific group of zero or more endpoints on each node

8.9.2.8. AttributeDataIB

This is used in Write Request and Report Data actions.

Info Field	Type	Quality	Conformance
DataVersion	data-ver		desc
Change	enum8		desc
Path	AttributePathIB		M
Data	desc		M

DataVersion

- This field SHALL NOT be present for a group or wildcard path.
- This field MAY be present for a path in a Write Request action.
- This field SHALL be present for a path in a Report Data action.

Change

This field indicates a change to a list for a write interaction, or as reported for a read or subscribe interaction.

- This field SHALL only be present if the data type is a [list entry index](#) or [list data type](#).
- The list change SHALL be in the context of a [fabric-filtered list](#), if fabric-filtering is enabled or required.
- This field SHALL be one of the values defined in the table below.

Name	Description
REPLACE	This indicates the entire list changing to another list.
ADD	This adds one or more entries without imposing a specific order.

To delete or modify single entries in the list, the full list needs to be replaced.

REPLACE

- This Change value SHALL indicate replacing the entire list.
- This Change value SHALL only be used when the last nesting level of the Path field indicates a list (Attribute or FieldID).
- The data type of the [AttributeDataIB](#) Data field SHALL be [list](#).
- The Data field MAY be an empty list which effectively deletes all entries from the list.

ADD

- This Change value SHALL indicate adding one entry to the list.
- This Change value SHALL only be used when the last nesting level of the Path field indicates a list (Attribute or FieldID).

- The data type of the [AttributeDataIB](#) Data field SHALL be the data type of the list's entries.

Path

See [AttributePathIB](#).

Data

- The data type of this field SHALL be the [data type](#) of the data indicated by the Path field.
- If the final nesting level indicated by the Path is an [list entry index](#) of null, the data type SHALL be the data type of the list's entries.

8.9.2.9. AttributeReportIB

Info Field	Type	Quality	Conformance
AttributeStatus	AttributeStatusIB		O.a
AttributeData	AttributeDataIB		O.a

- Only one of the above two fields SHALL be present.

8.9.2.10. DataVersionFilterIB

Info Field	Type	Quality	Conformance
Path	ClusterPathIB		M
DataVersion	data-ver		M

- The Path field SHALL indicate a concrete path.

8.9.3. Event Information Blocks and Paths

8.9.3.1. EventFilterIB

Info Field	Type	Quality	Conformance
Node	node-id		M
EventMin	event-no		M

8.9.3.2. EventPathIB

Path Field	Type	Quality	Conformance
Path	ClusterPathIB		M
Event	event-id		O
IsUrgent	bool		O

8.9.3.3. Valid Event Paths

This table is valid for [EventPathIB](#). The table defines valid event paths including combinations with wildcards, but not group paths.

Node	Endpoint	Cluster	Event	Event(s) Requested
Specific	Wildcard	Wildcard	Wildcard	all events from all clusters on all endpoints on a specific node
Specific	Wildcard	Specific	Wildcard	all events from a specific cluster on all endpoints on a specific node (e.g. Descriptor cluster)
Specific	Wildcard	Specific	Specific	a specific event from a specific cluster on all endpoints on a specific node (e.g. Descriptor cluster)
Specific	Specific	Wildcard	Wildcard	all events from all clusters on a specific endpoint on a specific node
Specific	Specific	Specific	Wildcard	all events from a specific cluster on a specific endpoint on a specific node
Specific	Specific	Specific	Specific	a specific event from a specific cluster on a specific endpoint on a specific node

8.9.3.4. EventDataIB

Info Field	Type	Quality	Conformance
Path	EventPathIB		M
EventNumber	event-no		M
Priority	priority		M
EpochTimeStamp	posix-ms		O.a
SystemTimeStamp	system-ms		O.a

Info Field	Type	Quality	Conformance
Data	<i>variable</i>		M

- The Path field SHALL indicate an existent path.
- The Path field SHALL NOT have an IsUrgent field present.

Data Field

The Data field SHALL contain the cluster-specific payload of the Event.

8.9.3.5. EventReportIB

Info Field	Type	Quality	Conformance
EventStatus	EventStatusIB		O.a
EventData	EventDataIB		O.a

- Only one of the above fields SHALL be present.

8.9.4. Command Information Blocks and Paths

8.9.4.1. CommandPathIB

Path Field	Type	Quality	Conformance
ClusterPath	ClusterPathIB		M
Command	command-id		M

8.9.4.2. Valid Command Paths

This table is valid for [CommandPathIB](#). The table defines valid command paths including combinations with wildcards, non-group (node & endpoint), and group paths. A blank entry means the element does not exist in the path.

Node	Endpoint	Group	Cluster	Command	Description	Conformance
Specific	Wildcard		Specific	Specific	a specific cluster command to all endpoints of a node	P, M
Specific	Specific		Specific	Specific	a specific cluster command to a specific endpoint on a node	M

Node	Endpoint	Group	Cluster	Command	Description	Conformance
		Specific	Specific	Specific	a specific cluster command to a group of endpoints	M

NOTE Support for commands paths involving wildcards is provisional.

8.9.4.3. CommandDataIB

Info Field	Type	Quality	Conformance
CommandPath	CommandPathIB		M
CommandFields	<i>variable</i>		M
CommandRef	uint16		O

CommandRef field

This field is used to provide a reference for a command.

1. When used in an Invoke Request Action:
 - a. This field MAY be omitted if InvokeRequests contains only a single CommandDataIB.
 - b. When InvokeRequests contains multiple CommandDataIBs, the CommandRef field SHALL contain a value unique within InvokeRequests.
 - i. Implementations MAY choose to use the index of the CommandDataIB within the InvokeRequests field as the value of this field.
2. When used in an Invoke Response Action:
 - a. CommandRef MAY be omitted from the CommandDataIB in InvokeResponses if the InvokeRequests being responded to contains only a single CommandDataIB, even if CommandRef was included in that CommandDataIB.
 - b. CommandRef SHALL be included in the CommandDataIB in InvokeResponses if InvokeRequests contained more than one CommandDataIB.
 - c. When a CommandDataIB in InvokeResponses has a CommandRef field, its value SHALL be equal to the CommandRef field from the CommandDataIB (in the InvokeRequests) that generated this CommandDataIB.
 - i. If InvokeRequests contained only a single CommandDataIB, and CommandRef was not included in that CommandDataIB, then CommandRef for this CommandDataIB in InvokeResponses SHALL NOT be included.

8.9.4.4. InvokeResponseIB

This is used in response to an invoked command.

Info Field	Type	Quality	Conformance
Command	CommandDataIB		O.a
Status	CommandStatusIB		O.a

- Only one of the above fields SHALL be present.
- The Path field in [CommandDataIB](#) and [CommandStatusIB](#) SHALL indicate a concrete path.

8.9.5. Status Information Blocks and Paths

8.9.5.1. CommandStatusIB

This is used to indicate a command response that represents a status: an invalid command, an error accessing the command, successful [execution](#) of a command without a following command, and so on.

Info Field	Type	Quality	Conformance
CommandPath	CommandPathIB		M
Status	StatusIB		M
CommandRef	uint16		O

- The Path field SHALL indicate a concrete path.

CommandRef field

This field is used to provide a reference for a command.

- CommandRef MAY be omitted from the CommandStatusIB if InvokeRequests contains only a single CommandDataIB, even if CommandRef was included in that CommandDataIB.
- CommandRef SHALL be included in the CommandStatusIB in InvokeResponses if InvokeRequests contained more than one CommandDataIB.
- When a CommandStatusIB in InvokeResponses has a CommandRef field, its value SHALL be equal to the CommandRef field from the CommandDataIB (in InvokeRequests) that generated this CommandStatusIB.
 - If InvokeRequests contained only a single CommandDataIB, and CommandRef was not included in that CommandDataIB, then CommandRef for this CommandStatusIB in InvokeResponses SHALL NOT be included.

8.9.5.2. EventStatusIB

This is used to indicate an invalid event, or an error accessing the event.

Info Field	Type	Quality	Conformance
Path	EventPathIB		M
Status	StatusIB		M

- The Path field SHALL indicate a concrete path.

8.9.5.3. AttributeStatusIB

This is used to indicate an invalid attribute data request path, or an error accessing the data, or success writing to an attribute path.

Info Field	Type	Quality	Conformance
Path	AttributePathIB		M
Status	StatusIB		M

- The Path field SHALL indicate a concrete path.

8.9.5.4. StatusIB

This is used to respond with errors or success to actions. A success Status field is valid for every layer. A non-success Status field is either defined in this layer, or generated and recognized by another layer. ClusterStatus is defined in a cluster specification.

See [Section 8.10, “Status Codes”](#) for valid values for Status.

Info Field	Type	Quality	Conformance
Status	status		M
ClusterStatus	status		O

Status Field

This is the one of the common status code values defined in [Status Code Table](#).

ClusterStatus Field

ClusterStatus values are defined in a cluster specification.

8.10. Status Codes

The table below lists global status codes for the Interaction Model. These MAY be used by interaction model processing of actions and as common status codes for cluster specifications. All values not defined here SHALL be reserved (per general conventions). Cluster specifications that wish to communicate a status not defined in this table MAY use a cluster-specific status code as described in [Status Codes](#).

8.10.1. Status Code Table

Status Code	Value	Summary
0x00	SUCCESS	Operation was successful.
0x01	FAILURE	Operation was not successful.

Status Code	Value	Summary
0x7D	INVALID_SUBSCRIPTION	Subscription ID is not active.
0x7E	UNSUPPORTED_ACCESS NOT_AUTHORIZED	The sender of the action or command does not have authorization or access. NOT_AUTHORIZED is an obsolete name of this error code.
0x7F	UNSUPPORTED_ENDPOINT	The endpoint indicated is unsupported on the node.
0x80	INVALID_ACTION	The action is malformed, has missing fields, or fields with invalid values. Action not carried out.
0x81	UNSUPPORTED_COMMAND UNSUP_COMMAND	The indicated command ID is not supported on the cluster instance. Command not carried out. UNSUP_COMMAND is an obsolete name for this error code.
0x82	<i>reserved</i>	Deprecated: use UNSUPPORTED_COMMAND
0x83	<i>reserved</i>	Deprecated: use UNSUPPORTED_COMMAND
0x84	<i>reserved</i>	Deprecated: use UNSUPPORTED_COMMAND
0x85	INVALID_COMMAND INVALID_FIELD	The cluster command is malformed, has missing fields, or fields with invalid values. Command not carried out. INVALID_FIELD is an obsolete name for this error code.
0x86	UNSUPPORTED_ATTRIBUTE	The indicated attribute ID, field ID or list entry does not exist for an attribute path.

Status Code	Value	Summary
0x87	CONSTRAINT_ERROR INVALID_VALUE	Out of range error or set to a reserved value. Attribute keeps its old value. Note that an attribute value may be out of range if an attribute is related to another, e.g. with minimum and maximum attributes. See the individual attribute descriptions for specific details. INVALID_VALUE is an obsolete name for this error code.
0x88	UNSUPPORTED_WRITE READ_ONLY	Attempt to write a read-only attribute. READ_ONLY is an obsolete name for this error code.
0x89	RESOURCE_EXHAUSTED INSUFFICIENT_SPACE	An action or operation failed due to insufficient available resources. INSUFFICIENT_SPACE is an obsolete name for this error code.
0x8A	<i>reserved</i>	Legacy cluster specification error status code: use SUCCESS
0x8B	NOT_FOUND	The indicated data field or entry could not be found.
0x8C	UNREPORTABLE_ATTRIBUTE	Reports cannot be issued for this attribute.
0x8D	INVALID_DATA_TYPE	The data type indicated is undefined or invalid for the indicated data field. Command or action not carried out.
0x8E	<i>reserved</i>	Legacy cluster specification error status code: use UNSUPPORTED_ATTRIBUTE.
0x8F	UNSUPPORTED_READ	Attempt to read a write-only attribute.
0x90	<i>reserved</i>	Deprecated: use FAILURE
0x91	<i>reserved</i>	Deprecated: use FAILURE
0x92	DATA_VERSION_MISMATCH	Cluster instance data version did not match request path

Status Code	Value	Summary
0x93	<i>reserved</i>	Legacy cluster specification error status code: use FAILURE
0x94	TIMEOUT	The transaction was aborted due to time being exceeded.
0x95	<i>reserved</i>	ZCL OTA Upgrade cluster specific error status code
0x96	<i>reserved</i>	ZCL OTA Upgrade cluster specific error status code.
0x97	<i>reserved</i>	ZCL OTA Upgrade cluster specific error status code.
0x98	<i>reserved</i>	ZCL OTA Upgrade cluster specific error status code.
0x99	<i>reserved</i>	ZCL OTA Upgrade cluster specific error status code.
0x9A	<i>reserved</i>	ZCL OTA Upgrade cluster specific error status code.
0x9B	UNSUPPORTED_NODE	The node ID indicated is not supported on the node.
0x9C	BUSY	The receiver is busy processing another action that prevents the execution of the incoming action.
0x9D	ACCESS_RESTRICTED	The access to the action or command by the sender is permitted by the ACL but restricted by the ARL.
0xC0	<i>reserved</i>	Deprecated: use FAILURE
0xC1	<i>reserved</i>	Deprecated: use FAILURE
0xC2	<i>reserved</i>	Deprecated: use FAILURE
0xC3	UNSUPPORTED_CLUSTER	The cluster indicated is not supported on the endpoint.
0xC4	<i>reserved</i>	Deprecated: use SUCCESS
0xC5	NO_UPSTREAM_SUBSCRIPTION	Used by proxies to convey to clients the lack of an upstream subscription to a source.

Status Code	Value	Summary
0xC6	NEEDS_TIMED_INTERACTION	A Untimed Write or Untimed Invoke interaction was used for an attribute or command that requires a Timed Write or Timed Invoke.
0xC7	UNSUPPORTED_EVENT	The indicated event ID is not supported on the cluster instance.
0xC8	PATHS_EXHAUSTED	The receiver has insufficient resources to support the specified number of paths in the request
0xC9	TIMED_REQUEST_MISMATCH	A request with TimedRequest field set to TRUE was issued outside a Timed transaction or a request with TimedRequest set to FALSE was issued inside a Timed transaction.
0xCA	FAILSAFE_REQUIRED	A request requiring a Fail-safe context was invoked without the Fail-Safe context.
0xCB	INVALID_IN_STATE	The received request cannot be handled due to the current operational state of the device
0xCC	NO_COMMAND_RESPONSE	A CommandDataIB is missing a response in the InvokeResponses of an Invoke Response action.
0xCD	TERMS_AND_CONDITIONS_CHANGED	The node requires updated TC acceptance. The user MAY be directed to visit the Enhanced-SetupFlowMaintenanceUrl to complete this.
0xCE	MAINTENANCE_REQUIRED	The node requires the user to visit the EnhancedSetupFlow-MaintenanceUrl for instructions on further action.
0xCF	DYNAMIC_CONSTRAINT_ERROR	The value for the data type was not accepted due to runtime validation issues. Command or action not carried out.

Status Code	Value	Summary
0xD0	ALREADY_EXISTS	Attempt to create an entity that already exists or create an entity with an identifier that is already in use. Command or action not carried out.
0xD1	INVALID_TRANSPORT_TYPE	Attempt to process on a transport type not valid for this element. Command or action not carried out.

Chapter 9. System Model Specification

9.1. Practical Information

9.1.1. Revision History

Revision	Description
1	Initial release of this specification
2	Added usage description for TagList
3	Added description related to Node reconfiguration

9.1.2. Scope and Purpose

This is part of a package of Data Model specifications that are agnostic to underlying layers (encoding, message, network, transport, etc.). Each specification below may be independently maintained. This package, as a whole, SHALL be independently maintained as agnostic and decoupled from lower layers. This package may be referenced by inclusion in a set of protocol stack specifications for a complete vertical standard.

Data Model	Defines first order elements and namespace for endpoints, clusters, attributes, data types, etc.
Interaction Model	Defines interactions, transactions and actions between nodes.
System Model	Defines relationships that are managed between endpoints and clusters.
Cluster Library	Reference library of cluster specifications.
Device Library	Reference library of devices type definitions.

9.1.3. Origin Story

The origin of this section is the [Dotdot Architecture Model](#) and parts of Chapter 2 of the [Zigbee Cluster Library specification](#) that define the data model.

This document defines the specific model used in Matter, which has diverged from Zigbee.

9.1.4. Overview

This specification defines persistent relationships between local and remote instances of data model elements, that support a system of operational communication between such instances. A system is a set of nodes and persistent relationships that automate data flow and control based on local or external stimulus.

9.2. Endpoint Composition

Endpoint composition SHALL be indicated by these [Descriptor](#) cluster attributes:

- DeviceTypeList SHALL list the device type(s) that the endpoint supports.
- PartsList SHALL indicate endpoints as required for each device type in the DeviceTypeList.

The population of the PartsList of the various endpoints allows for the indication of a hierarchy (composition) for which the following terms are defined:

- "Descendants" of an endpoint E are endpoints that are:
 - in the PartsList of E, or
 - descendants of an endpoint in the PartsList of E.
- "Children" of an endpoint E are endpoints that are descendants of E but not descendants of any endpoint in the PartsList of E. E is the "Parent" endpoint of these Children endpoints. A "Child" endpoint is one of these Children endpoints.
- A "Leaf" endpoint is an endpoint that has an empty PartsList.
- "Sibling endpoints" are endpoints which are children of the same endpoint.




Figure 61. Simplified relationship diagram

An example of composition:

In the [refrigerator composition figure](#) below:

- Endpoint 0 (Root Node, thus using full-family pattern) has a PartsList containing endpoints 1 through 5.
- Endpoint 1 has a PartsList containing endpoints 2 and 4.

- Endpoint 2 has a PartsList containing endpoint 3.
- Endpoint 4 has a PartsList containing endpoint 5.

Therefore, endpoints 2, 3, 4 and 5 are not children of endpoint 0 and thus do *not* qualify as siblings of endpoint 1 (which is the only child of endpoint 0).

Also:

- Endpoint 1 has a PartsList containing endpoints 2 and 4.
- Endpoints 2 and 4 do not occur in any other PartsList of a descendant of endpoint 1.
- Endpoints 2 and 4 are children of endpoint 1.

Therefore, they are sibling endpoints.

Also:

- Endpoint 2 has a PartsList containing only endpoint 3.

Therefore, endpoint 3 is a single child and does not have any siblings.

Also:

- Endpoint 4 has a PartsList containing only endpoint 5.

Therefore, endpoint 5 is a single child and does not have any siblings.

Also:

- Endpoints 3 and 5 are not appearing together in any other PartsList (except that of endpoint 0).

Therefore, they are not siblings.

9.2.1. Endpoint Requirements

Each [Simple endpoint](#) SHALL support exactly one [Application device type](#) with the exception that, the endpoint MAY support additional device types which are subsets of the [Application device type](#) (the superset). See [Section 9.2.12, “Superset Device Types”](#) for cluster requirements of superset devices.

The endpoint MAY support additional [utility device types](#).

If a device type requires endpoint composition, the composition SHALL be defined in the device type specification.

For example: A Color Temperature Light device type may support device type IDs for both a Dimmable Light and On/Off Light, because those are subsets of a Color Temperature Light (the superset).

For example: A Room Temperature Sensor device type and a Room Humidity Sensor device type must be on separate endpoints because they are both Simple device types and neither is a subset of the other.

For example: The Bridged Node device type (a utility device type) may be added to an endpoint, which means this endpoint represents a node behind a bridge, and requires one or more extra clusters.

9.2.2. Leaf Device Types

A leaf device type is not composed of other application device types (see above bullet on subsets).

For example: The Temperature Sensor device type mandates the Temperature Measurement server cluster, and does not require additional device types or endpoints.

Most device types define leaf endpoints without the need for composition.

For example: A Dimmer Switch device type mandates these clusters: On/Off, Level Control, Identify, and does not require additional device types or endpoints.

There can be situations where leaf endpoints have no application device types, and one or more utility device types.

For example: a [light with mains power and battery back-up](#). This needs two endpoints for the two Power Source device types, each of which has a Power Source cluster, to describe both power sources, and only one of those endpoints can host the lighting application device type. The other endpoint will only have the Power Source cluster, and no application device type.

9.2.3. Composed Device Types

A composed device type contains one or more other device types, referred to as [component device types](#), which are listed in the PartsList of the endpoint of the composed device type. The component device types exist on their own endpoints and may also themselves be composed device types or be device types which are not composed.

A composed device type is identified in this specification by the existence of a Device Type Requirements section within the device type definition. This section contains a table that lists the component device types that could be included in a composition along with constraints for each component device type.

For example: An endpoint supporting a Temperature Controlled Cabinet device type inside a Refrigerator device type which has 2 component temperature sensors (one for freezer temperature, and one for ice tray temperature) would have a PartsList containing 2 component temperature sensor leaf endpoints (one for each of the sensors). Those component leaf end-

points would indicate and conform to the temperature sensor device type, and have unique TagList attributes to disambiguate them.

Clusters associated with the component device types MAY have their conformance altered in the Cluster Requirements on Component Device Types table in the composed device type definition.

Elements associated with the component device types MAY have their conformance altered in the Element Requirements on Component Device Types table in the composed device type definition.

Tags associated with the component device types MAY have their conformance altered in the Semantic Tag Requirements on Component Device Types table in the composed device type definition.

For example: Device type A is a composed device type and includes component device type B in its Device Type Requirements table. The definition for device type B includes cluster C in its Cluster Requirements table with optional (O) conformance. Composed device type A lists component device type B and cluster C in its Cluster Requirements on Component Device Types table and alters the conformance for cluster C in device type B to mandatory (M).

9.2.4. Component Device Types

A component device type is a device type that is included in the Device Type Requirements table of a [composed device type](#).

9.2.5. Compound Devices

A compound device is a single physical entity containing multiple device types existing on separate endpoints but for which a composed device type is not defined. A compound device allows a manufacturer to create a product containing multiple device types as seen fit by that manufacturer.

For example: A compound device supporting 2 lighting device endpoints, an "up light" on endpoint 1 and a "down light" on endpoint 2. The root node PartsList contains endpoint 1 and 2. The lighting device types are not defined as composed devices, and no device is defined having such a construction, therefore this device containing multiple lighting device endpoints is a compound device.

For example: A PIR occupancy sensor which also contains a light level sensor and a temperature sensor, would have 3 device endpoints with the corresponding device types (Occupancy Sensor, Light Sensor and Temperature Sensor).

9.2.6. Asserting Conditions Among Device Types

A device type MAY assert conditions defined in itself or other device types which appear on the same or other endpoints of the node, subject to the requirements below.

Definitions:

Term	Description
Assert/Asserting	Asserting is the act of one device type causing a condition defined in itself or some other device type to be true.
Asserting device type	A device type asserting a condition defined in itself or another device type is the asserting device type.
Asserted	A condition is asserted when it is caused to be true.

The following are the device types which can have conditions asserted:

Device Type	Description
Root Node	The closest root or root-like endpoint to the asserting device type
Bridged Node	
Self	The device type itself
Child	See Section 9.2, “Endpoint Composition” for the definition of Child

The conditions which can be asserted by an asserting device type are listed in the Condition Requirements table in the definition of the asserting device type. The contents of this table reference conditions defined in the device type itself or in other device types.

The table contains the following columns:

Column	Description
Location	Indicates where this condition is being asserted and is one of Root, Self, or Descendant
Device Type ID	The device type ID of the device where the condition being asserted is defined
Device Type Name	The device type name associated with the device type ID above
Condition	The condition being asserted
Constraint	Any numerical constraints on this condition
Conformance	The conformance for the asserted condition

For example, the Meter Reference Point device type asserts the TimeSyncCond condition as mandatory on the Root Node device type, where it is defined as optional in the Root Node definition. The Meter Reference Point device type is the asserting device type and the TimeSyncCond condition is the asserted condition.

The conformance for a condition MAY contain any valid conformance expression as allowed by this

specification. A conformance listed as purely optional (O) MAY be used to signal that a condition could be expected to be asserted on the node, which could be helpful to clients.

For details regarding valid conformance designations, see [Section 7.3, “Conformance”](#).

9.2.7. Endpoint Composition Patterns

This specification defines two patterns for endpoint composition supported by a device type:

- Tree pattern of endpoints directly below the composed endpoint.
- Full-family list of all descendant endpoints below the composed endpoint.

The Bridged Node device type has specific requirements on endpoint composition as detailed in the definition of this device type in the Device Library.

9.2.7.1. Tree Pattern

The tree pattern supports a general tree of endpoints, each of which supports the pattern defined by the endpoint's device type. The tree pattern is commonly used for composed application device types. The tree pattern is commonly used for device types that support physical device composition (e.g. a Refrigerator).

- The tree pattern of composition SHALL be defined in the device type specification, by a table of device types, from which the specified device type is composed.
 - Each device type from the table that is included in an implementation SHALL be represented by one or more endpoints in the PartsList attribute of the Descriptor cluster on the endpoint of the parent device type.
 - An implementation MAY include other device types not listed in the table of device types in its composition and each such additional device SHALL be represented by an endpoint in the PartsList attribute of the Descriptor cluster on the endpoint of the parent device type.

For example: A Refrigerator/Freezer endpoint has a PartsList with the Refrigerator and Freezer endpoint, but not the parts of the Refrigerator or Freezer. The Refrigerator and Freezer endpoints each have a PartsList that includes temperature sensor and control endpoints.

9.2.7.2. Full-Family Pattern

The full-family pattern is a list of all descendant endpoints, with no imposed hierarchy.

- The full-family pattern of endpoint composition SHALL be defined in the device type specification by stating that the device type supports the full-family pattern.
- The PartsList attribute of the Descriptor cluster on the endpoint SHALL contain all descendant endpoints.

Example: The Root Node and Aggregator device types use the full-family pattern, as defined

in their device type specification.

9.2.7.3. PartsList Requirements for all Patterns

- An endpoint SHALL NOT include itself in its PartsList attribute.
- Cycles where an endpoint is a (direct or indirect) ancestor of itself SHALL NOT occur.
- An endpoint (except endpoint 0) SHALL have a single parent, i.e. there SHALL be a single path from the Root Node to each endpoint.
- The PartsList SHALL NOT include endpoints that do not exist on the device.
- The PartsList SHALL NOT include duplicate entries as it represents a set.

9.2.8. Root Node Endpoint

A root node device type is a device type that is specified to be supported by all nodes in a fabric. Its main function is to describe the node and what endpoints the node supports.

The following requirements apply to root node endpoints:

- There SHALL be only one root node endpoint for the node.
- A root node device type SHALL be a [singleton](#) (representing the entire node) on the root node endpoint.
- The root node endpoint MAY contain additional [utility device types](#).
- The PartsList of the Descriptor cluster on the root node endpoint SHALL list all endpoints on the node, except the root node endpoint.
- The root node endpoint SHALL NOT exist in any other endpoint's PartsList on the node.
- The root node endpoint requirements are defined by a node scoped device type.

NOTE

The root node endpoint is currently under-specified for nodes which are not commissionable. As a result, the root node endpoint on non-commissionable nodes MAY be absent or otherwise limited in functionality. This specification gap is expected to be addressed in a future revision.

9.2.9. Disambiguation

Using the terms and composition patterns described above, some endpoints will have sibling endpoints with an overlap in application device type(s). For this case, the **Duplicate** condition has been defined (see the Base Device Type in the Device Library).

- A device type definition MAY define a method for disambiguation when the Duplicate condition applies to either the endpoint with that device type or to its children.
- If no disambiguation method is defined in the relevant device type definitions, an endpoint to which the Duplicate condition applies SHALL have a TagList attribute in its Descriptor cluster and its set of tags SHALL be distinct from the set of tags of any sibling endpoint with which it has an overlap in application device type(s). When comparing sets of tags, ordering of the list

SHALL be ignored.

- The TagList feature and attribute MAY also be used in other cases to provide guidance to a client and/or a user.
- When the Duplicate condition applies, clients SHOULD disambiguate the relevant sibling endpoints in their user interfaces and client logic using the method defined in the device type or the TagList attribute values if no other method is defined.

For example: If the Duplicate condition applies to child endpoints of an Aggregator endpoint that represent multiple independent bridged devices, the endpoints make available metadata to allow a client to disambiguate distinct bridged devices with an overlap in application device types. Typically this is done using the NodeLabel attribute of the Bridged Device Basic Information cluster - thus reusing the naming information which the bridge already has to allow disambiguation to the user when using a direct user interface to the bridge.

For instance: The Aggregator in [this example](#), exposes several Color Temperature Lights (endpoints 13 and 22) which are disambiguated with their NodeLabel. Note that the compound device at endpoints 24, 25 and 26 also uses a TagList since, for this case, the bridge knows the lighting direction of both elements of the compound device.

Another example: The following diagram (copied from Chapter 10 of the Device Library) shows a Video Player device containing 3 separate Content Apps on endpoints 21, 31 and 41 which are disambiguated using the ApplicationName attribute of the Application Basic cluster on these endpoints.




Figure 62. Endpoint Composition for Video Player Device

Note that these requirements related to disambiguation hold for non-composed as well as composed device types, as illustrated in the examples below.

Example of using TagList for disambiguation for a non-composed device type: A switch device with two buttons.

- Endpoint 0 has device type Root Node with a PartsList with endpoints 11 and 12 (these are sibling endpoints per the definition above)
- Endpoint 11 has device type Generic Switch and TagList containing two tags: Position.Left and Number.One
- Endpoint 12 has device type Generic Switch and TagList containing two tags: Position.Right and Number.Two

Example of using TagList for disambiguation of a composed device type: a refrigerator with two cabinets (fresh food and freezer), as illustrated in the figure below:

- Endpoint 1 has the composed device type (Refrigerator) and the PartsList contains endpoints 2 and 4 (but not 3 and 5).
- Endpoints 2 and 4 are sibling endpoints which each have a device type Temperature Con-

trolled Cabinet, with a TagList to disambiguate.

- Endpoint 2 has a PartsList with endpoint 3 to indicate that the temperature sensor on endpoint 3 measures the temperature in the refrigerator (fresh food) compartment. Endpoint 3 is a child of endpoint 2.
- Similarly, endpoint 4 has a PartsList with endpoint 5 to indicate that the temperature sensor on endpoint 5 measures the temperature in the freezer compartment. Endpoint 5 is a child of endpoint 4.
- Endpoints 3 and 5 each use a TagList to indicate to the client details about the placement of the temperature sensor inside the Temperature Controlled Cabinet of endpoint 2 and 4. Note that these TagLists in this example are identical (in the product, both sensors are located in the top left position of their respective compartment), which is allowed since endpoints 3 and 5 are not siblings, hence the "Duplicate" condition does not apply. So these endpoints are using a TagList to provide information rather than to disambiguate.




Figure 63. Disambiguation and information via semantic tags

9.2.10. Dynamic Endpoint Allocation

Some nodes MAY require a dynamic number of endpoints, since the functionality they expose can change at run-time, e.g.

- a [Bridge](#) on which Bridged Devices are added or deleted.
- a Casting Video Player in which Content App endpoints are added or deleted (see Device Library, section 10).

Such dynamic entities which need to be exposed with an endpoint, will be referred to as an "exposed entity" in the following description.

For any node which changes its endpoint composition, the endpoint addresses SHALL be allocated

according to the following rules:

- When such exposed entity is exposed for the first time, it SHALL be allocated a *new* endpoint address (or set of endpoint addresses), incrementing from the highest previously (ever) used endpoint address.
 - For the situation where a node following these mechanisms has exhausted all available remaining endpoint addresses for exposed entities, it MAY wrap around to the lowest unused endpoint address, starting at 1, since the RootNode endpoint (endpoint 0) cannot be used to host application device types and is therefore never free.
- For existing exposed entities, the endpoint addresses SHALL NOT be changed.
 - This persistency requirement also holds for the case of restart/reboot of the device.

As a result of these mechanisms, endpoint addresses that were used for exposed entities that were once exposed but now have been removed will *not* be reused in the future (apart from the exceptional wrap-around corner case mentioned above), in order to avoid the possibility of confusing other nodes by re-assigning (reusing) an endpoint address for a *different* exposed entity. Other nodes using the exposed entities from this node SHOULD remove information related to exposed entities no longer being exposed.

Other nodes that wish to be notified of changes of the exposed entities SHOULD monitor changes of the PartsList attribute in the Descriptor cluster on the [root node endpoint](#).

9.2.11. Node Configuration Changes

A node MAY be able to support different capabilities depending on the installation or configuration of the device. Changes in capabilities MAY result in the node gaining and/or losing Matter-related functionality which could impact the representation and use of the node within the system.

Changes in the configuration of the node could impact a controller's ability to understand the capabilities of that node and the overall system, such as configured bindings with other nodes or automations in a controller/ecosystem, which could break due to the functionality no longer being present on the node.

Nodes SHALL indicate changes in their configuration that impacts the representation of the node by increasing the [ConfigurationVersion](#) attribute in the [Basic Information Cluster](#). The [ConfigurationVersion](#) SHALL also be increased in the case a firmware update adds or removes functionality on the node. For bridged nodes, this attribute is located in the corresponding [Bridged Device Basic Information Cluster](#); when such a configuration change occurs, the [ConfigurationVersion](#) attribute in the [Basic Information Cluster](#) of the bridge SHALL also be increased.

A client SHOULD detect a change of [ConfigurationVersion](#) on a node (or bridged node) and SHOULD take the necessary steps to reevaluate the capabilities of the node—these steps are manufacturer specific. The controller MAY attempt to reconfigure e.g. bindings, if required, in order to re-establish the functionality of the system. The steps required to do so are manufacturer specific.

Servers SHALL report any changes to fixed attributes in the same Report Data Action as the corresponding change to the [ConfigurationVersion](#) attribute.

Clients SHOULD wait until they have received an entire Report Data Action before acting on any changes to fixed attributes, or the [ConfigurationVersion](#) attribute.

On any changes seen either to a fixed attribute or the ConfigurationVersion attribute of the Basic Information cluster or the Bridged Basic Information cluster, clients SHOULD hold-off for some time before taking into consideration the re-computation of device topology or configuration, if their design requires it, so as to provide enough time for the server to have reported all dependent changes associated with the reconfiguration and avoid re-computing configuration changes too frequently. For example, it is not expected for a Thermostat configuration to change every second, but rather at most a few times during the installed lifetime. It would be appropriate for a client to wait for several minutes before acting on changes to fixed attributes, to ensure that the entire reconfiguration is done and avoid doing complex re-computations or data migrations of settings or bindings until it's very unlikely that further configuration changes will occur.

9.2.12. Superset Device Types

A superset device type is a device type that is functionally similar or identical to a subset device type and has an overlap in its cluster requirements with a subset device type, subject to the following constraints:

- All clusters in the subset device type with mandatory conformance SHALL be supported in the superset device type with mandatory conformance.
- Any clusters in the subset device type with optional conformance MAY be supported in the superset device type.
- Any clusters in the subset device type with optional conformance that are also supported in the superset device type MAY have either optional or mandatory conformance in the superset device type.
- Element requirements of overlapping clusters in the superset device type SHALL be the same as in the subset device type.
- An endpoint implementing a superset device type MAY include the subset device type(s) in the DeviceTypeList of the descriptor cluster of the endpoint.
- There is a use case (rationale) for exposing the relationship between the superset and subset device type.

When these constraints apply, the subset device type SHALL be indicated in the "Superset Of" column of the table in the "Classification" section of superset device type, and guidance for server and client makers MAY be added.

9.2.12.1. Superset Device Types with identical cluster requirements

A special case of Superset/Subset device types is where the cluster requirements (as expressed in the constraints above) of two device types are identical. This typically occurs in cases where a new device type is introduced for a specific use case which is different from the use case of the existing device type, e.g. mounted on/off control is added for the in-wall use case versus the existing on/off plug-in unit. If there is a rationale to define one of these device types as the superset of the other, this SHALL be indicated in the "Classification" section of superset device type, and guidance for server and client makers SHOULD be added - also in the definition of the subset device type to hint

readers to the superset device type with identical cluster requirements.

This guidance to implementers of servers and clients on how to use one or both device types SHOULD help to realize optimum representation in both 'new' clients, which understand both the existing and the new device type, and 'old' clients, which understand only the existing device type.

9.2.12.2. Superset Device Types - usage in Server

The endpoint exposing the server side of the clusters of the superset device type SHALL include the superset device type and MAY include the subset device type(s) in its DeviceTypeList attribute.

Example for Dimmable Light and Extended Color Light device types (Extended Color Light is a superset of Dimmable Light): An endpoint implementing device type Extended Color Light SHALL include device type Extended Color Light and MAY include device type Dimmable Light in its DeviceTypeList.

The endpoint exposing the server side of the clusters of the subset device type SHALL include the subset device type and SHALL NOT include the superset device type(s) in its DeviceTypeList attribute.

Example for Dimmable Light and Extended Color Light device types (Extended Color Light is a superset of Dimmable Light): An endpoint implementing device type Dimmable Light SHALL include device type Dimmable Light and SHALL NOT include Extended Color Light in its DeviceTypeList.

If an endpoint lists both the superset device type and the subset device type(s) in its DeviceTypeList attribute, a client that is capable of only recognizing and controlling the subset device type(s) would also be able to control the aspects of the superset device type that are in common with the subset device type.

9.2.12.3. Superset Device Types - Primary Device Type indication

For the [Primary Device](#) of a device, the superset device type SHALL be used - NOT any of the subset device type(s).

9.2.12.4. Superset Device Types used in Composed Device Types

A composed device type SHALL indicate that superset device types are allowed in an implementation by adding a + sign after the ID and the name of the subset device type having a subset/superset relationship in the device type requirements table of the composed device type.

An endpoint SHALL NOT implement a subset of a device type listed in the device type requirements table of a composed device type.

For example, the following table (when used in the Device Type Requirements section of the definition of a device type) indicates that an endpoint implementing this device type MAY use a dimmable light or any superset of a dimmable light (such as color temperature light and extended color

light), but it may not use an on/off light as that is a subset of a dimmable light:

ID	Name	Constraint	Conformance
0x0101+	Dimmable Light+		O

9.3. Interaction Model Relationships

This section define some of the system behaviors and their constraints as they apply to interactions specified in the [Interaction Model](#).

9.3.1. Subscription

9.3.1.1. Persistency

A subscription is an ephemeral 'session' between a subscriber and a publisher. A subscription can lose synchronization for a variety of reasons, including (but not limited to):

- Inability to send reports due to network connectivity issues
- Factory-reset of the publisher
- Reboot of either the subscriber or publisher
- Decision by either publisher or subscriber to tear down the subscription to reclaim resources

In all cases, the subscriber can eventually discover the loss of synchronization by not receiving a sync report or data report in the agreed upon sync interval, or through some other failure to communicate with the publisher.

- When a subscriber discovers the loss of synchronization, it MAY then initiate a re-subscription to resume the subscription.
- An implementation MAY choose to persist the details of a subscription across reboots, but it is not necessary.

In all cases, the publisher eventually discovers the loss of synchronization by not receiving a Status Response to a Report Data message that expects a response, or by receiving an error Status Response.

9.4. Binding Relationship

This relationship occurs because a simple client endpoint does not know which endpoints will be the target for the client generated actions, on one or more remote nodes.

For example: A light switch that controls one or more light bulbs, does not know the remote node endpoints of the bulbs.

For example: A thermostat that subscribes to an occupancy sensor, does not know the remote node endpoint of the sensor.

In such cases, a binding is used to direct the local endpoint to the remote endpoint. The existence of the Binding cluster on an endpoint, allows a director to create one or more binding entries (bindings) in the Binding cluster. A director is a remote client that is given access to create such bindings.

Each binding indicates either a remote node’s endpoint or cluster on a remote node’s endpoint. Multiple bindings are allowed, depending on the interaction. A binding is either a unicast binding, where the binding target is a remote endpoint, or a groupcast binding, where the binding target is a group of remote endpoints.

See the Binding Cluster specification for more specification detail.

9.5. Descriptor Cluster

NOTE

The Descriptor cluster is meant to replace the support from the Zigbee Device Object (ZDO) for describing a node, its endpoints and clusters.

This cluster describes an endpoint instance on the node, independently from other endpoints, but also allows composition of endpoints to conform to complex device type patterns.

This cluster supports a list of one or more device type identifiers that represent conformance to device type specifications.

For Example: An Extended Color Light device type may support device type IDs for both a Dimmable Light and On/Off Light, because those are subsets of an Extended Color Light (the superset).

The cluster supports a PartsList attribute that is a list of zero or more endpoints to support compound devices or composed device types.

For Example: A Refrigerator/Freezer appliance device type may be defined as being composed of multiple Temperature Sensor endpoints, a Metering endpoint, and two Thermostat endpoints.

9.5.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Semantic tag list; TagList feature
3	Add EndpointUniqueID attribute

9.5.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Endpoint	DESC

9.5.3. Cluster ID

ID	Name
0x001D	Descriptor

9.5.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Conformance	Summary
0	TAGLIST	TagList	desc	The TagList attribute is present

9.5.4.1. TagList Feature

See the [Disambiguation section in the System Model spec](#) for conformance requirements for this feature and the corresponding attribute.

9.5.5. Data Types

9.5.5.1. DeviceTypeStruct Type

The device type and revision define endpoint conformance to a release of a device type definition. See the Data Model specification for more information.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Device-Type	devtype-id					M
1	Revision	uint16	min 1				M

DeviceType Field

This SHALL indicate the device type definition. The endpoint SHALL conform to the device type definition and cluster specifications required by the device type.

Revision Field

This is the implemented revision of the device type definition. The endpoint SHALL conform to this revision of the device type.

9.5.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Device-TypeList	list[Device-Type-Struct]	min 1	F	desc	R V	M
0x0001	ServerList	list[cluster-id]		F	empty	R V	M
0x0002	ClientList	list[cluster-id]		F	empty	R V	M
0x0003	PartsList	list[end-point-no]			empty	R V	M
0x0004	TagList	list[Semantic-TagStruct]	1 to 6	F	MS	R V	TAGLIST
0x0005	End-pointUniqueID	string	max 32	F		R V	O

9.5.6.1. DeviceTypeList Attribute

This is a list of device types and corresponding revisions declaring endpoint conformance (see [Section 9.5.5.1, “DeviceTypeStruct Type”](#)). At least one device type entry SHALL be present.

An endpoint SHALL conform to all device types listed in the DeviceTypeList. A cluster instance that is in common for more than one device type in the DeviceTypeList SHALL be supported as a shared cluster instance on the endpoint.

9.5.6.2. ServerList Attribute

This attribute SHALL list each cluster ID for the server clusters present on the endpoint instance.

9.5.6.3. ClientList Attribute

This attribute SHALL list each cluster ID for the client clusters present on the endpoint instance.

9.5.6.4. PartsList Attribute

This attribute indicates composition of the device type instance. Device type instance composition SHALL include the endpoints in this list.

See [Endpoint Composition](#) for more information about which endpoints to include in this list.

9.5.6.5. TagList Attribute

This attribute SHALL be used to disambiguate sibling endpoints in certain situations, as defined in

the [Disambiguation section in the System Model specification](#). An example of such a situation might be a device with two buttons, with this attribute being used to indicate which of the two endpoints corresponds to the button on the left side.

It MAY also be used to provide information about an endpoint (e.g. the relative location of a Temperature sensor in a Temperature Controlled Cabinet).

- A client SHOULD use these tags to convey disambiguation information and other relevant information to the user (e.g. showing it in a user interface), as appropriate.
- A client SHOULD use these tags in its logic to make decisions, as appropriate.

For example, a client MAY identify which endpoint maps to a certain function, orientation or labeling.

A client MAY use the Label field of each [SemanticTagStruct](#), if present in each structure, to indicate characteristics of an endpoint, or to augment what is provided in the TagID field of the same structure.

9.5.6.6. EndpointUniqueID Attribute

This attribute SHALL indicate an identifier which allows to uniquely identify the functionality exposed on an endpoint, and therefore SHALL be unique within the device. It is constructed in a manufacturer specific manner.

- If a globally unique identifier is used, the same [rules](#) as defined for the UniqueID attribute in the Basic Information cluster apply.
- If the identifier is only unique in the scope of the device, and cannot be used to track the device, then it MAY remain unchanged at factory reset.

The value does not need to be human readable, since it is intended for machine to machine (M2M) communication.

9.6. Binding Cluster

NOTE

This scope of this document is the Binding cluster as part of the Cluster Library. The Binding cluster is meant to replace the support from the Zigbee Device Object (ZDO) for supporting the binding table.

A binding represents a persistent relationship between an endpoint and one or more other local or remote endpoints. A binding does not require that the relationship exists. It is up to the node application to set up the relationship.

A binding is used to inform a client endpoint of one or more targets for a potential interaction. For example: a light switch that controls one or more light bulbs, needs to be told the nodes and endpoints of the bulbs, or told a group in which the bulbs are members. For example: A client that needs to subscribe to an occupancy sensor, needs to know the node and endpoint of the sensor.

In such cases, a binding is used to direct a local endpoint to a target. The existence of the Binding

cluster on the client endpoint, allows the creation of one or more binding entries (bindings) in the Binding cluster.

Each binding indicates another endpoint or cluster on another endpoint. Multiple bindings are allowed, depending on the interaction.

A binding is either a unicast binding, where the target is a single endpoint on a single node, or a groupcast binding, where the target is a group, which may indicate multiple endpoints on multiple nodes. The binding may also target a single cluster on the target endpoint(s).

When a client cluster requires a target for an interaction, the Binding cluster SHALL exist on the same endpoint.

Once a binding entry is created on the Binding cluster, the client endpoint MAY initiate interactions to the binding target.

9.6.1. Binding Mutation

If, during the creation of multiple bindings, there are no available resources to create an entry, or to establish a binding relationship, the client SHALL respond with a status of RESOURCE_EXHAUSTED, and the binding SHALL NOT be created.

The number of bindings supported is left to the implementation. However, a device type definition MAY prescribe the minimum number of bindings supported on an endpoint. In this case, the number prescribed by the device type SHALL be supported for each fabric the node supports, unless the device type specifies otherwise. The total number of bindings supported SHALL be the sum of the requirements for each endpoint, unless the device types involved specify otherwise.

- For example, if a node supports 6 fabrics and a device type specifies at least 3 bindings must be supported, the node would need to support at least 18 bindings and ensure that at least 3 were available to every fabric.

A binding SHALL only be created with the Cluster field if the indicated client cluster exists on the endpoint.

When a binding is removed, the client endpoint SHALL end the binding relationship with the removed binding target.

9.6.2. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

9.6.3. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Endpoint	BIND

9.6.4. Cluster ID

ID	Name
0x001E	Binding

9.6.5. Data Types

9.6.5.1. TargetStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	Node	node-id	all				Endpoint
2	Group	group-id	min 1				!Endpoint
3	Endpoint	endpoint-no	all				!Group
4	Cluster	cluster-id	all				O

Node Field

This field is the remote target node ID. If the Endpoint field is present, this field SHALL be present.

Group Field

This field is the target group ID that represents remote endpoints. If the Endpoint field is present, this field SHALL NOT be present.

Endpoint Field

This field is the remote endpoint that the local endpoint is bound to. If the Group field is present, this field SHALL NOT be present.

Cluster Field

This field is the cluster ID (client & server) on the local and target endpoint(s). If this field is present, the client cluster SHALL also exist on this endpoint (with this Binding cluster). If this field is present, the target SHALL be this cluster on the target endpoint(s).

9.6.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Binding	list[Target-Struct]	desc	N	[]	RW VM F	M

9.6.6.1. Binding Attribute

Each entry SHALL represent a binding.

Here are some examples:

Endpoint Device Type	Node	Group	Cluster	Endpoint	Example Description
Light Switch Client	omit	1234	omit	omit	switch endpoint sends On/Off, Level Control & Color Control cluster commands to group 1234
Temp Sensor Client	456789	omit	1026	3	temperature sensor client subscribes to node 456789 endpoint 3 temperature measurement cluster

9.7. Label Cluster

This cluster provides a feature to tag an endpoint with zero or more labels. This is a base cluster that requires a derived cluster to create an instance.

9.7.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

9.7.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Endpoint	LABEL

9.7.3. Cluster ID

This is a base cluster with no cluster ID. See derived clusters for more information.

ID	Name
n/a	Label

9.7.4. Data Types

9.7.4.1. LabelStruct Type

This is a string tuple with strings that are user defined.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Label	string	max 16		empty		M
1	Value	string	max 16		empty		M

Label Field

The Label or Value semantic is not defined here.

Label examples: "room", "zone", "group", "direction".

Value Field

The Label or Value semantic is not defined here. The Value is a discriminator for a Label that may have multiple instances.

Label:Value examples: "room":"bedroom 2", "orientation":"North", "floor":"2", "direction":"up"

9.7.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	LabelList	list[LabelStruct]					M

9.7.5.1. LabelList Attribute

This is a list of string tuples. Each entry is a [LabelStruct](#).

9.8. Fixed Label Cluster

This cluster is derived from the Label cluster and provides a feature for the device to tag an endpoint with zero or more read-only labels.

Examples:

- A bridge can use this to indicate grouping of bridged devices. For example: All bridged devices whose endpoints have an entry in their LabelList "room":"bedroom 2" are in the same (bed)room.
- A manufacturer can use this to identify a characteristic of an endpoint. For example to identify the endpoints of a luminaire, one pointing up, the other pointing down, one of the endpoints would have a LabelList entry "orientation":"up" while the other would have "orientation":"down". Using such indication, the user interface of a Node controlling this luminaire knows which of the endpoints is which of the lights.

Note that the [TagList](#) in the Descriptor cluster provides an alternative mechanism for such self-description using standardized tags rather than manufacturer-selected strings, yielding a standardized mechanism for features defined in the various [namespaces](#). The second example above can be implemented using semantic tags `Direction.Upward` and `Direction.Downward` instead of (or in addition to) the Fixed Label cluster.

9.8.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

9.8.2. Classification

Hierarchy	Role	Scope	PICS Code
Label	Utility	Endpoint	FLABEL

9.8.3. Cluster ID

ID	Name
0x0040	Fixed Label

9.8.4. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	LabelList	list[Label-Struct]		N	empty	R V	M

9.9. User Label Cluster

This cluster is derived from the Label cluster and provides a feature to tag an endpoint with zero or more writable labels.

9.9.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

9.9.2. Classification

Hierarchy	Role	Scope	PICS Code
Label	Utility	Endpoint	ULABEL

9.9.3. Cluster ID

ID	Name
0x0041	User Label

9.9.4. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	LabelList	list[Label-Struct]	desc	N	empty	RW VM	M

9.9.4.1. LabelList Attribute

The server SHALL support the storage of up to 4 list entries in this attribute. The server MAY support the storage of more than 4 entries in this attribute.

When reading from this attribute, the server SHALL respond with the actual contents of the attribute which MAY contain any number of entries (possibly more than 4),

When writing to this attribute, a client MAY include any number of entries to be written, or none at all.

If an attempt is made to write to this attribute with a list length that is not supported by the server, the server SHALL respond with RESOURCE_EXHAUSTED.

9.10. Access Control Cluster

The Access Control Cluster exposes a data model view of a Node's Access Control List (ACL), which codifies the rules used to manage and enforce Access Control for the Node's endpoints and their associated cluster instances. Access to this Access Control Cluster itself requires a special Administrator privilege level, such that only Nodes granted such privilege (hereafter termed "Administrators") can manage the Access Control Cluster.

The Access Control Cluster SHALL be present on the [root node endpoint](#) of each Node, and SHALL NOT be present on any other Endpoint of any Node.

9.10.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Added Managed Device feature, Extension feature, fixed conformance
3	Added Auxiliary feature and AuxiliaryACL attribute

9.10.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	ACL

9.10.3. Cluster ID

Cluster ID	Cluster Name
0x001F	AccessControl

9.10.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Conformance	Summary
0	EXTS	Extension	O	Device provides ACL Extension attribute

Bit	Code	Feature	Conformance	Summary
1	MNGD	ManagedDevice	desc	Device is managed
2	AUX	Auxiliary	P, desc	Some auxiliary ACL entries may be auto-generated by other features

9.10.4.1. Extension Feature

This feature indicates the device supports ACL Extension attribute.

9.10.4.2. ManagedDevice Feature

This feature is for a device that is managed by a service associated with the device vendor and which imposes default access restrictions upon each new fabric added to it. This could arise, for example, if the device is managed by a service provider under contract to an end-user, in such a way that the manager of the device does not unconditionally grant universal access to all of a device's functionality, even for fabric administrators. For example, many Home Routers are managed by an Internet Service Provider (a service), and these services often have a policy that requires them to obtain user consent before certain administrative functions can be delegated to a third party (e.g., a fabric Administrator). These restrictions are expressed using an Access Restriction List (ARL).

The purpose of this feature on the Access Control cluster is to indicate to a fabric Administrator that access by it to specific attributes, commands and/or events for specific clusters is currently prohibited. Attempts to access these restricted data model elements SHALL result in an error of ACCESS_RESTRICTED.

A device that implements this feature SHALL have a mechanism to honor the ReviewFabricRestrictions command, such as user interfaces or service interactions associated with a service provider or the device manufacturer, which allows the owner (or subscriber) to manage access restrictions for each fabric. The user interface design, which includes the way restrictions are organized and presented to the user, is not specified, but SHOULD be usable by non-expert end-users from common mobile devices, personal computers, or an on-device user interface.

Controllers and clients SHOULD incorporate generic handling of the ACCESS_RESTRICTED error code, when it appears in allowed contexts, in order to gracefully handle situations where this feature is encountered. Device vendors that adopt this feature SHOULD be judicious in its use given the risk of unexpected behavior in controllers and clients.

For certification testing, a device that implements this feature SHALL provide a way for all restrictions to be removed.

The ARL attribute provides the set of restrictions currently applied to this fabric.

The ReviewFabricRestrictions command provides a way for the fabric Administrator to request that the server triggers a review of the current fabric restrictions, by involving external entities such as end-users, or other services associated with the manager of the device hosting the server. This review process MAY involve communication between external services and the user, and MAY take

an unpredictable amount of time to complete since an end-user may need to visit some resources, such as a mobile application or web site. A FabricRestrictionReviewUpdate event will be generated by the device within a predictable time period of the ReviewFabricRestrictionsResponse (see [Section 9.10.8.1, “ReviewFabricRestrictions”](#) for specification of this time period), and this event can be correlated with the ReviewFabricRestrictionsResponse using a token provided in both. The device MAY provide instructions or a Redirect URL in the FabricRestrictionReviewUpdate event in order to help the user access the features required for managing per-fabric restrictions.

See [Section 6.6.2, “Access Control Model”](#) for a description of how access control is impacted by the ARL attribute.

Managed Device Feature Usage Restrictions

Use of this feature SHALL be limited to the mandatory clusters of endpoints having a device type that explicitly permits its use in the Device Library Specification. As a reminder, the device types associated with an endpoint are listed in the Descriptor cluster of the endpoint.

In addition, use of this feature SHALL NOT restrict the following clusters on any endpoint:

- 1. the Descriptor Cluster (0x001D)
- 2. the Binding Cluster (0x001E)
- 3. the Network Commissioning Cluster (0x0031)
- 4. the Identify Cluster (0x0003)
- 5. the Groups Cluster (0x0004)

In addition, use of this feature SHALL NOT restrict the [global attributes](#) of any cluster.

Because ARLs cannot be used to restrict root node access or access to any clusters required for commissioning, administrators MAY determine the current restrictions of the ARL at any point, including during commissioning after joining the fabric.

9.10.4.3. Auxiliary Feature

This feature indicates that there MAY be entries in the AuxiliaryACL attribute which indicate synthesized ACL entries. For example, when this feature is supported, the configuration of groups via the [Groupcast cluster](#) MAY lead, under some circumstances, to some access being granted via auxiliary means (not explicit ACL attribute list writes) to some subjects by virtue of group membership. See the [AuxiliaryACL](#) attribute.

9.10.5. Data Types

9.10.5.1. ChangeTypeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Changed	Entry or extension was changed	M

Value	Name	Summary	Conformance
1	Added	Entry or extension was added	M
2	Removed	Entry or extension was removed	M

9.10.5.2. AccessControlEntryPrivilegeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
1	View	Can read and observe all (except Access Control Cluster)	M
2	Proxy View		D
3	Operate	View privileges, and can perform the primary function of this Node (except Access Control Cluster)	M
4	Manage	Operate privileges, and can modify persistent configuration of this Node (except Access Control Cluster)	M
5	Administer	Manage privileges, and can observe and modify the Access Control Cluster	M

Operate Value

This value implicitly grants View privileges

Manage Value

This value implicitly grants Operate & View privileges

Administer Value

This value implicitly grants Manage, Operate & View privileges

9.10.5.3. AccessRestrictionTypeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	AttributeAccessForbidden	Clients on this fabric are currently forbidden from reading and writing an attribute	M
1	AttributeWriteForbidden	Clients on this fabric are currently forbidden from writing an attribute	M
2	CommandForbidden	Clients on this fabric are currently forbidden from invoking a command	M
3	EventForbidden	Clients on this fabric are currently forbidden from reading an event	M

9.10.5.4. AccessControlEntryAuthModeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
1	PASE	Passcode authenticated session	M
2	CASE	Certificate authenticated session	M
3	Group	Group authenticated session	M

9.10.5.5. AccessControlAuxiliaryTypeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	System	This ACL entry exists because of some system reason and is likely non-revocable	M
1	Groupcast	Synthesized via Groupcast Cluster administrator-configured group membership	M

9.10.5.6. AccessControlTargetStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Cluster	cluster-id	all	X			M
1	Endpoint	endpoint-no	all	X			M
2	Device-Type	devtype-id	all	X			M

9.10.5.7. AccessControlEntryStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	Privilege	AccessControlEntryPrivilegeEnum	all			S	M
2	AuthMode	AccessControlEntryAuthModeEnum	all			S	M
3	Subjects	list[subject-id]	max SubjectsPerAccessControlEntry	X		S	M
4	Targets	list[AccessControlTargetStruct]	max TargetsPerAccessControlEntry	X		S	M
5	Auxiliary-Type	AccessControlAuxiliaryTypeEnum	all			S	P, O

Privilege Field

The privilege field SHALL specify the level of privilege granted by this Access Control Entry.

Each privilege builds upon its predecessor, expanding the set of actions that can be performed upon a Node. Administer is the highest privilege, and is special as it pertains to the administration of privileges itself, via the Access Control Cluster.

When a Node is granted a particular privilege, it is also implicitly granted all logically lower privilege levels as well. The following diagram illustrates how the higher privilege levels subsume the lower privilege levels:




Figure 64. Access Control Privilege Levels

Individual clusters SHALL define whether attributes are readable, writable, or both readable and writable. Clusters also SHALL define which privilege is minimally required to be able to perform a particular read or write action on those attributes, or invoke particular commands. Device type specifications MAY further restrict the privilege required.

The Access Control Cluster SHALL require the Administer privilege to observe and modify the Access Control Cluster itself. The Administer privilege SHALL NOT be used on Access Control Entries which use the Group auth mode.

E.g. A Fan Control Cluster may require Operate privilege to write to a level attribute (low/medium/high), and to configure each level's RPM setting via a command. The Fan Control Cluster may also expose a current RPM attribute, which requires only View privilege to read. Clients granted Operate privilege will be able to both change the level, and configure each level's RPM. Clients granted View privilege will be able to read the current RPM, but will not be granted sufficient privilege to change the level or configure each level's RPM.

E.g. A Fan Control Cluster may be included in a more industrial device type. To ensure proper operation, this device type may restrict configuration of fan level RPM settings to require Manage privilege. Clients granted Manage privilege will have sufficient privilege to configure each level's RPM; clients granted Operate privilege will not be able to perform such configuration, but will still be able to change the level. This additional restriction would apply only to the Fan Control Cluster as included in this particular device type; a client granted Operate privilege may still be able to perform configuration in Fan Control Clusters included in other

device types on the same Node.

AuthMode Field

The AuthMode field SHALL specify the authentication mode required by this Access Control Entry.

Subjects Field

The subjects field SHALL specify a list of [Subject IDs](#), to which this Access Control Entry grants access.

Device types MAY impose additional constraints on the minimum number of subjects per Access Control Entry.

An attempt to create an entry with more subjects than the node can support SHALL result in a RESOURCE_EXHAUSTED error and the entry SHALL NOT be created.

Subject ID SHALL be of type uint64 with semantics depending on the entry’s AuthMode as follows:

Subject Semantics

AuthMode	Subject
PASE	Lower 16-bits → Passcode ID Upper 48-bits → all bits clear
CASE	64-bits → Node ID or CASE Authenticated Tag
Group	Lower 16-bits → Group ID Upper 48-bits → all bits clear

An empty subjects list indicates a wildcard; that is, this entry SHALL grant access to any Node that successfully authenticates via AuthMode. The subjects list SHALL NOT be empty if the entry’s AuthMode is PASE.

The PASE AuthMode is reserved for future use (see [Section 6.6.2.9, “Bootstrapping of the Access Control Cluster”](#)). An attempt to write an entry with AuthMode set to PASE SHALL fail with a status code of CONSTRAINT_ERROR.

For PASE authentication, the Passcode ID identifies the required passcode verifier, and SHALL be 0 for the default commissioning passcode.

For CASE authentication, the Subject ID is a distinguished name within the Operational Certificate shared during CASE session establishment, the type of which is [determined by its range](#) to be one of:

- a Node ID, which identifies the required source node directly (by ID)
- a [CASE Authenticated Tag](#), which identifies the required source node indirectly (by tag)

E.g. an ACL entry with CASE AuthMode that grants privileges to Subject IDs [0x0000_0000_1111_1111, 0x0000_0000_2222_2222, 0x0000_0000_3333_3333] (which are Node

IDs) will grant access to Nodes with Node ID 0x0000_0000_1111_1111, 0x0000_0000_2222_2222, or 0x0000_0000_3333_3333, but will not grant access to Nodes with Node ID 0x0000_0000_4444_4444 or 0x0000_0000_5555_5555.

E.g. an ACL entry with CASE AuthMode that grants privileges to Subject IDs [0x0000_0000_6666_6666, 0xFFFF_FFFD_ABCD_0002] (which are a Node ID and a CASE Authenticated Tag) will grant access to the Node with Node ID 0x0000_0000_6666_6666 and any Nodes with CAT identifier value 0xABCD if the CAT's version is 0x0002 or higher. It will not grant access to Nodes with other CAT values such as 0x9999_9999. Any node with CAT identifier value of 0xABCD but version less than 0x0002 (for example: 0xFFFF_FFFD_ABCD_0001) will not be granted access.

For Group authentication, the Group ID identifies the required group, as defined in the Group Key Management Cluster.

E.g. an entry with Group AuthMode that grants privileges to Subject IDs [0x0000_0000_0000_1111, 0x0000_0000_0000_2222] (which are Group IDs) will grant access to Nodes in Group 0x1111 or 0x2222, but will not grant access to Nodes in Group 0x3333, even if they share Operational Group Keys.

Targets Field

The targets field SHALL specify a list of [AccessControlTargetStruct](#), which define the clusters on this Node to which this Access Control Entry grants access.

Device types MAY impose additional constraints on the minimum number of targets per Access Control Entry.

An attempt to create an entry with more targets than the node can support SHALL result in a RESOURCE_EXHAUSTED error and the entry SHALL NOT be created.

A single target SHALL contain at least one field (Cluster, Endpoint, or DeviceType), and SHALL NOT contain both an Endpoint field and a DeviceType field.

A target grants access based on the presence of fields as follows:

Target Semantics

Cluster	Endpoint	DeviceType	Target
			Invalid
		D	All clusters on any end-point with Descriptor containing device type D
	E		All clusters on only endpoint E

Cluster	Endpoint	DeviceType	Target
	E	D	Invalid
C			Only cluster C on all endpoints
C		D	Only cluster C on any endpoint with Descriptor containing device type D
C	E		Only cluster C on only endpoint E
C	E	D	Invalid

An empty targets list indicates a wildcard:

- If the subject's AuthMode is Group:
 - If the Auxiliary feature map bit is not set, this entry SHALL grant access to all cluster instances on all endpoints on this Node.
 - Else, when the Auxiliary feature map bit is set, this entry SHALL grant access to all cluster instances on all endpoints on this Node, besides the RootNode endpoint (endpoint 0).
- Else, this entry SHALL grant access to all cluster instances on all endpoints on this Node.

E.g. an entry that grants privileges to the Color Light Bulb Device Type will grant privileges to any cluster on any endpoint that contains the Color Light Bulb device type (whether that cluster is in the Color Light Bulb device type or not), and will not grant privileges to any other cluster on any other endpoint.

E.g. an entry that grants privileges to Endpoint 1 will grant privileges to any cluster on Endpoint 1, and will not grant privileges to any other cluster on any other endpoint.

E.g. an entry that grants privileges to the On/Off Cluster on any endpoint will not grant privileges to any other cluster on any endpoint.

E.g. an entry that grants privileges to the On/Off Cluster with Color Light Bulb Device Type will grant privileges to just the On/Off Cluster on any endpoint that contains the Color Light Bulb device type, and will not grant privileges to any other cluster on any other endpoint (including other clusters in the Color Light Bulb device type, or the On/Off cluster on endpoints that do not contain the Color Light Bulb device type).

E.g. an entry that grants privileges to the On/Off Cluster on Endpoint 1 will not grant privileges to any other cluster on Endpoint 1, or to any other cluster (including the On/Off cluster) on any other endpoint.

E.g. an entry that grants privileges to the whole Node (empty Targets field, and AuthMode set to Group), on a Node with endpoints IDs 0, 1 and 2, would not grant privileges to Endpoint 0 during group message processing if the Auxiliary feature map bit is set, but would grant

access on endpoints 1 and 2 (if those endpoints are joined to the subject’s group ID), while the same request would grant access on endpoints 0, 1 and 2, if the Auxiliary feature map bit is not set.

AuxiliaryType Field

This field, if present, SHALL indicate the type of [AuxiliaryACL](#) entry stored. For example, for [Group-cast cluster](#) synthesized entries, this would be set to [Groupcast](#). An administrator MAY use the information provided by this field to determine the feature that caused the auxiliary ACL entry to be created, so that access can be managed via that other feature’s operations.

This field SHALL NOT be present in entries in the ACL attribute.

9.10.5.8. AccessControlExtensionStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	Data	octstr	max 128			S	M

Data Field

This field MAY be used by manufacturers to store arbitrary [TLV-encoded](#) data related to a fabric’s Access Control Entries.

The contents SHALL consist of a [top-level anonymous list](#); each list element SHALL include a [profile-specific](#) tag encoded in [fully-qualified](#) form.

Administrators MAY iterate over this list of elements, and interpret selected elements at their discretion. The content of each element is not specified, but MAY be coordinated among manufacturers at their discretion.

E.g. a manufacturer could use this field to store structured data, including various metadata and cryptographic signatures. The manufacturer could then verify a fabric’s Access Control List by generating a canonical bytestream from the Access Control Entries for the fabric, then verifying the signature against it. Such a canonical bytestream could be generated by encoding specific entry fields and sub-fields (such as lists) in specific order and specific format (e.g. TLV).

9.10.5.9. AccessRestrictionStruct Type

This structure describes an access restriction that would be applied to a specific data model element on a given endpoint/cluster pair (see [Section 9.10.5.10, “AccessRestrictionEntryStruct”](#)).

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Type	AccessRestrictionTypeEnum	all				M
1	ID	uint32	all	X			M

Type Field

This field SHALL indicate the type of restriction, for example, AttributeAccessForbidden.

ID Field

This field SHALL indicate the element [Manufacturer Extensible Identifier \(MEI\)](#) associated with the element type subject to the access restriction, based upon the [AccessRestrictionTypeEnum](#). When the Type is AttributeAccessForbidden or AttributeWriteForbidden, this value SHALL be considered of type attrib-id (i.e. an attribute identifier). When the Type is CommandForbidden, this value SHALL be considered of type command-id (i.e. an attribute identifier). When the Type is EventForbidden, this value SHALL be considered of type event-id (i.e. an event identifier).

A null value SHALL indicate the wildcard value for the given value of Type (i.e. all elements associated with the Type under the associated endpoint and cluster for the containing AccessRestrictionEntryStruct).

9.10.5.10. AccessRestrictionEntryStruct Type

This structure describes a current access restriction on the fabric.

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Endpoint	endpoint-no	all			S	M
1	Cluster	cluster-id	all			S	M
2	Restrictions	list[AccessRestrictionStruct]	min 1			S	M

Endpoint Field

This field SHALL indicate the endpoint having associated access restrictions scoped to the associated fabric of the list containing the entry.

Cluster Field

This field SHALL indicate the cluster having associated access restrictions under the entry's Endpoint, scoped to the associated fabric of the list containing the entry.

Restrictions Field

This field SHALL indicate the set of restrictions applying to the Cluster under the given Endpoint, scoped to the associated fabric of the list containing the entry.

This list SHALL NOT be empty.

9.10.5.11. CommissioningAccessRestrictionEntryStruct Type

This structure describes a current access restriction when there is no accessing fabric.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Endpoint	endpoint-no	all				M
1	Cluster	cluster-id	all				M
2	Restrictions	list[AccessRestrictionStruct]	min 1				M

Endpoint Field

This field SHALL indicate the endpoint having associated access restrictions scoped to the associated fabric of the list containing the entry.

Cluster Field

This field SHALL indicate the cluster having associated access restrictions under the entry's Endpoint, scoped to the associated fabric of the list containing the entry.

Restrictions Field

This field SHALL indicate the set of restrictions applying to the Cluster under the given Endpoint, scoped to the associated fabric of the list containing the entry.

This list SHALL NOT be empty.

9.10.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	ACL	list[AccessControlEntryStruct]	desc			RW A F	M
0x0001	Extension	list[AccessControlExtensionStruct]	desc			RW A F	EXTS

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0002	SubjectsPerAccessControlEntry	uint16	4 to 65534	F		R V	M
0x0003	TargetsPerAccessControlEntry	uint16	3 to 65534	F		R V	M
0x0004	AccessControlEntriesPerFabric	uint16	4 to 65534	F		R V	M
0x0005	CommissioningARL	list[CommissioningAccessRestrictionEntryStruct]	desc	F		R V	MNGD
0x0006	ARL	list[AccessRestrictionEntryStruct]	desc			R V F	MNGD
0x0007	Auxiliary-ACL	list[AccessControlEntryStruct]	max 2000	C		R A F	P, AUX

9.10.6.1. Default Value

The default values of the Access Control Cluster ACL and Extension attributes are empty; that is, devoid of contents, with each list attribute containing zero elements.

The Access Control List is able to have an initial entry added because the [Access Control Privilege Granting algorithm](#) behaves as if, over a PASE commissioning channel during the [commissioning phase](#), the following implicit Access Control Entry were present on the Commissionee (but not on the Commissioner):

```
Access Control Cluster: {
  ACL: [
    0: {                                // implicit entry only; does not explicitly exist!
      FabricIndex: 0,                  // not fabric-specific
    }
  ]
}
```

```

    Privilege: Administer,
    AuthMode: PASE,
    Subjects: [],
    Targets: []           // entire node
  }
],
Extension: []
}

```

When the [ManagedDevice](#) feature is present, the value of the Access Control Cluster CommissioningARL attribute contains the set of access restrictions applied during commissioning, before a fabric has been assigned. Upon allocation of a new FabricIndex during commissioning (see [Section 11.18.6.8, “AddNOC Command”](#)), one or more ARL entries with the new FabricIndex SHALL be added to the ARL attribute.

In the following example of the CommissioningARL and ARL attributes, the restrictions applied during commissioning consist of a single entry with 3 limitations on Cluster 0x0453 of Endpoint 1: writing to attribute 0x0000, all commands, all events. In contrast, the example also includes a restriction on the fabric with index 1 which consists of 1 limitation on Cluster 0x0453 of Endpoint 1: writing to attribute 0x0000.

```

Access Control Cluster: {
  CommissioningARL: [
    0: {
      Endpoint: 1,           // application endpoint 1
      Cluster: 0x0453,       // application cluster 0x0453
      Restrictions: [
        0: {
          Type: AttributeWriteForbidden, // write limitation for attribute
          Id: 0x0000                     // attribute 0x0000
        },
        1: {
          Type: CommandForbidden, // command limitation
          Id: NULL                // all commands
        },
        2: {
          Type: EventForbidden, // event limitation
          Id: NULL              // all events
        }
      ]
    }
  ]
}
ARL: [
  0: {

```

```

FabricIndex: 1,          // ARL is fabric-scoped
Endpoint: 1,            // application endpoint 1
Cluster: 0x0453,        // application cluster 0x0453
Restrictions: [
  0: {
    Type: AttributeWriteForbidden, // write limitation for attribute
    Id: 0x0000                     // attribute 0x0000
  }
]
}

```

9.10.6.2. Administration Guidelines

The Access Control Cluster is passive in nature and is only responsible for maintaining entries in the Access Control List. It is the responsibility of Administrators to create and maintain Access Control policy by managing the list and its entries. The Access Control Cluster SHALL NOT change or remove Access Control Entries of its own volition. A device that needs to impose access restrictions upon a fabric MAY do so by adding Access Restriction Entries to the CommissioningARL and ARL attributes.

Administrators SHOULD strive to minimize resource usage by combining entries where appropriate. For example, if an Administrator is responsible for an entry that grants privilege to subject Node A, and wishes to grant the same privilege to Node B under the same conditions, then that Administrator SHOULD update the existing entry to apply to subject Node B as well as Node A, instead of creating a new entry. If a set of Nodes is commonly used in entries, then an Administrator MAY consider using [CASE Authenticated Tags](#) (CATs) for those entries, especially if membership in the set of Nodes changes over time.

Administrators SHOULD carefully consider the effect of Access Control Entries they manage, particularly when targeting entire Endpoints (either directly, or indirectly by Device Type), to ensure that granted privileges are appropriate for the set of Clusters that may entail. Administrators SHOULD be mindful that targeting by Device Type grants privileges to all Clusters on all Endpoints with Descriptor containing that Device Type (thereby including Clusters not part of that specified Device Type), now and in the future. Administrators SHOULD consider whether targeting specific Endpoints, or Clusters, or both, might be a more appropriate policy for the given Subjects; studying the Descriptor Cluster for affected Endpoints may help in this determination.

Administrators SHOULD be careful to avoid inadvertently removing their own administrative access. For example, an Administrator SHOULD change its own administrative access entry by updating the existing entry or by creating a new entry before removing the old entry, and SHOULD NOT remove the old entry before creating any new entry.

9.10.6.3. ACL Attribute

An attempt to add an Access Control Entry when no more entries are available SHALL result in a RESOURCE_EXHAUSTED error being reported and the ACL attribute SHALL NOT have the entry

added to it. See [Section 2.11.1.1, “Access Control Limits”](#).

See the [AccessControlEntriesPerFabric](#) attribute for the actual value of the number of entries per fabric supported by the server.

Each Access Control Entry codifies a single grant of privilege on this Node, and is used by the [Access Control Privilege Granting algorithm](#) to determine if a subject has privilege to interact with targets on the Node.

Every AccessControlEntryStruct in the ACL attribute SHALL NOT include the AuxiliaryType field.

9.10.6.4. Extension Attribute

If present, the Access Control Extensions MAY be used by Administrators to store arbitrary [data](#) related to fabric’s Access Control Entries.

The Access Control Extension list SHALL support a single extension entry per supported fabric.

9.10.6.5. SubjectsPerAccessControlEntry Attribute

This attribute SHALL provide the minimum number of [Subjects](#) per entry that are supported by this server.

Since reducing this value over time may invalidate ACL entries already written, this value SHALL NOT decrease across time as software updates occur that could impact this value. If this is a concern for a given implementation, it is recommended to only use the minimum value required and avoid reporting a higher value than the required minimum.

9.10.6.6. TargetsPerAccessControlEntry Attribute

This attribute SHALL provide the minimum number of [Targets](#) per entry that are supported by this server.

Since reducing this value over time may invalidate ACL entries already written, this value SHALL NOT decrease across time as software updates occur that could impact this value. If this is a concern for a given implementation, it is recommended to only use the minimum value required and avoid reporting a higher value than the required minimum.

9.10.6.7. AccessControlEntriesPerFabric Attribute

This attribute SHALL provide the minimum number of [ACL Entries](#) per fabric that are supported by this server.

Since reducing this value over time may invalidate ACL entries already written, this value SHALL NOT decrease across time as software updates occur that could impact this value. If this is a concern for a given implementation, it is recommended to only use the minimum value required and avoid reporting a higher value than the required minimum.

9.10.6.8. CommissioningARL Attribute

This attribute SHALL provide the set of [CommissioningAccessRestrictionEntryStruct](#) applied during

commissioning on a managed device.

When present, the CommissioningARL attribute SHALL indicate the access restrictions applying during commissioning.

Attempts to access data model elements described by an entry in the CommissioningARL attribute during commissioning SHALL result in an error of ACCESS_RESTRICTED. See [Section 6.6.2, “Access Control Model”](#) for more information about the features related to controlling access to a Node’s Endpoint Clusters (“Targets” hereafter) from other Nodes.

See [Section 9.10.4.2.1, “Managed Device Feature Usage Restrictions”](#) for limitations on the use of access restrictions.

9.10.6.9. ARL Attribute

This attribute SHALL provide the set of [AccessRestrictionEntryStruct](#) applied to the associated fabric on a managed device.

When present, the ARL attribute SHALL indicate the access restrictions applying to the accessing fabric. In contrast, the [CommissioningARL](#) attribute indicates the accessing restrictions that apply when there is no accessing fabric, such as during commissioning.

The access restrictions are externally added/removed based on the particular relationship the device hosting this server has with external entities such as its owner, external service provider, or end-user.

Attempts to access data model elements described by an entry in the ARL attribute for the accessing fabric SHALL result in an error of ACCESS_RESTRICTED. See [Section 6.6.2, “Access Control Model”](#) for more information about the features related to controlling access to a Node’s Endpoint Clusters (“Targets” hereafter) from other Nodes.

See [Section 9.10.4.2.1, “Managed Device Feature Usage Restrictions”](#) for limitations on the use of access restrictions.

9.10.6.10. AuxiliaryACL Attribute

This attribute, when present, SHALL provide a list of additional ACL entries generated by the server that SHALL be applied during the [Access Control Privilege Granting algorithm](#). These ACL entries are not writable or removable using this attribute; they reflect the underlying feature that generated them.

Features such as the Groupcast cluster MAY cause entries to appear in this attribute as a result of internal configurations that grant access to certain subjects by mechanisms other than explicit management of the [ACL](#) list.

The purpose of exposing these entries in the data model is to ensure transparency in the access granted to subjects.

Changes to this attribute resulting from reconfiguration of the underlying features SHALL cause the AuxiliaryAccessUpdated event to be generated when they occur. See the [AuxiliaryAccessUpdated Event](#) for additional constraints.

Servers SHOULD attempt to make the list as concise as possible while respecting the limits such as SubjectsPerAccessControlEntry and TargetsPerAccessControlEntry.

The entries in this list SHALL NOT count towards the AccessControlEntriesPerFabric attribute limit, which applies only to the writable entries in the [ACL](#) attribute.

Every AccessControlEntryStruct in the AuxiliaryACL SHALL have the AuxiliaryType field set to a value indicating the reason for existence of the entry in the list.

Each feature which would cause entries to appear in this attribute SHALL individually specify the format of entries that are auto-generated, and under which circumstances they are generated.

9.10.7. Error handling

Administrators SHALL use regular actions to administer the Access Control Cluster (by reading and writing entries in the list). Administrators SHOULD take care to use DataVersion conditional writes when administering the list or its contents.

The Access Control Cluster SHALL fail to write, and return an appropriate error, if an attempt is made to create or update an Access Control Entry or Access Control Extension such that it would have invalid contents.

For example, the following Access Control Entry conditions will result in an error of CONSTRAINT_ERROR:

- Privilege enum value invalid
- AuthMode enum value invalid
- AuthMode is PASE (reserved for future use)
- Subjects element invalid
 - e.g. illegal CAT with CASE AuthMode
- Targets element invalid
 - e.g. no field present
 - e.g. both Endpoint and DeviceType fields present
- Field combinations invalid
 - e.g. Administer Privilege with Group AuthMode

For example, the following Access Control Extension conditions will result in an error of CONSTRAINT_ERROR:

- There is an attempt to add more than 1 entry associated with the given [accessing fabric index](#) in the extension list
- Data value exceeds max length
- Data value not valid [TLV-encoded](#) data

The Access Control Cluster MAY fail to write, and return a RESOURCE_EXHAUSTED error, if an attempt is made to create or update an entry or extension such that storage is exhausted.

9.10.8. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	ReviewFabricRestrictions	client ⇒ server	ReviewFabricRestrictionsResponse	A F	MNGD
0x01	ReviewFabricRestrictionsResponse	client ⇐ server	N		MNGD

9.10.8.1. ReviewFabricRestrictions Command

This command signals to the service associated with the device vendor that the fabric administrator would like a review of the current restrictions on the accessing fabric. This command includes an optional list of ARL entries that the fabric administrator would like removed.

In response, a ReviewFabricRestrictionsResponse is sent which contains a token that can be used to correlate a review request with a FabricRestrictionReviewUpdate event.

Within 1 hour of the ReviewFabricRestrictionsResponse, the FabricRestrictionReviewUpdate event SHALL be generated, in order to indicate completion of the review and any additional steps required by the user for the review.

A review MAY include obtaining consent from the user, which can take time. For example, the user may need to respond to an email or a push notification.

The ARL attribute may change at any time due to actions taken by the user, or the service associated with the device vendor.

NOTE

A malicious fabric administrator could attempt to create a fabric with a VID that it does not own in order to gain enhanced access. For this reason the device SHOULD perform the [Fabric Table Vendor ID Verification Procedure](#).

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ARL	list[CommissioningAccessRestrictionEntryStruct]	desc			M

ARL Field

When the ARL field is provided, it indicates the specific restrictions that are requested for review. An empty list represents a generic request for review of all restrictions.

9.10.8.2. ReviewFabricRestrictionsResponse Command

Returns the review token for the request, which can be used to correlate with a FabricRestriction-ReviewUpdate event.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Token	uint64	all			M

Token Field

This field SHALL specify a Token that can be used to correlate a ReviewFabricRestrictionsResponse with a FabricRestrictionReviewUpdate event.

9.10.9. Events

ID	Name	Priority	Quality	Access	Conformance
0x00	AccessControlEntryChanged	INFO		A S	M
0x01	AccessControlExtensionChanged	INFO		A S	EXTS
0x02	FabricRestrictionReviewUpdate	INFO		A S	MNGD
0x03	AuxiliaryAccessUpdated	INFO		A S	P, AUX

9.10.9.1. AccessControlEntryChanged Event

The server SHALL generate AccessControlEntryChanged events whenever its [ACL](#) attribute data is changed by an Administrator.

- Each added entry SHALL generate an event with ChangeType Added.
- Each changed entry SHALL generate an event with ChangeType Changed.
- Each removed entry SHALL generate an event with ChangeType Removed.

Access Modifier: Fabric-Sensitive						
ID	Name	Type	Constraint	Quality	Fallback	Conformance
1	AdminNodeID	node-id	desc	X		M
2	AdminPasscodeID	uint16	desc	X		M

Access Modifier: Fabric-Sensitive						
3	ChangeType	ChangeType-Enum	all			M
4	LatestValue	AccessControlEntryStruct	all	X		M

AdminNodeID Field

The Node ID of the Administrator that made the change, if the change occurred via a CASE session.

Exactly one of AdminNodeID and AdminPasscodeID SHALL be set, depending on whether the change occurred via a CASE or PASE session; the other SHALL be null.

AdminPasscodeID Field

The Passcode ID of the Administrator that made the change, if the change occurred via a PASE session. Non-zero values are reserved for future use (see [PasscodeId generation in PBKDFParamRequest](#)).

Exactly one of AdminNodeID and AdminPasscodeID SHALL be set, depending on whether the change occurred via a CASE or PASE session; the other SHALL be null.

ChangeType Field

The [type of change](#) as appropriate.

LatestValue Field

The latest value of the changed [entry](#).

This field SHOULD be set if resources are adequate for it; otherwise it SHALL be set to NULL if resources are scarce.

9.10.9.2. AccessControlExtensionChanged Event

The server SHALL generate AccessControlExtensionChanged events whenever its [extension](#) attribute data is changed by an Administrator.

- Each added extension SHALL generate an event with ChangeType Added.
- Each changed extension SHALL generate an event with ChangeType Changed.
- Each removed extension SHALL generate an event with ChangeType Removed.

Access Modifier: Fabric-Sensitive						
ID	Name	Type	Constraint	Quality	Fallback	Conformance
1	AdminNodeID	node-id	desc	X		M

Access Modifier: Fabric-Sensitive						
2	AdminPass-codeID	uint16	desc	X		M
3	ChangeType	ChangeType-Enum	all			M
4	LatestValue	AccessControlExtensionStruct	all	X		M

AdminNodeID Field

The Node ID of the Administrator that made the change, if the change occurred via a CASE session.

Exactly one of AdminNodeID and AdminPasscodeID SHALL be set, depending on whether the change occurred via a CASE or PASE session; the other SHALL be null.

AdminPasscodeID Field

The Passcode ID of the Administrator that made the change, if the change occurred via a PASE session. Non-zero values are reserved for future use (see [PasscodeId generation in PBKDFParamRequest](#)).

Exactly one of AdminNodeID and AdminPasscodeID SHALL be set, depending on whether the change occurred via a CASE or PASE session; the other SHALL be null.

ChangeType Field

The [type of change](#) as appropriate.

LatestValue Field

The latest value of the changed [extension](#).

This field SHOULD be set if resources are adequate for it; otherwise it SHALL be set to NULL if resources are scarce.

9.10.9.3. FabricRestrictionReviewUpdate Event

The server SHALL generate a FabricRestrictionReviewUpdate event to indicate completion of a fabric restriction review. Due to the requirement to generate this event within a bound time frame of successful receipt of the [ReviewFabricRestrictions](#) command, this event MAY include additional steps that the client MAY present to the user in order to help the user locate the user interface for the Managed Device feature.

Access Modifier: Fabric-Sensitive						
ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Token	uint64	all			M

Access Modifier: Fabric-Sensitive						
1	Instruction	string	max 512			O
2	ARLRequest- FlowUrl	string	max 256			O

Token Field

This field SHALL indicate the Token that can be used to correlate a ReviewFabricRestrictionsResponse with a FabricRestrictionReviewUpdate event.

Instruction Field

This field SHALL provide human readable text that MAY be displayed to the user to help them locate the user interface for managing access restrictions for each fabric.

A device SHOULD implement the [Localization Configuration Cluster](#) when it has no other means to determine the locale to use for this text.

Examples include "Please try again and immediately access device display for further instructions." or "Please check email associated with your Acme account."

ARLRequestFlowUrl Field

This field SHALL indicate the URL for the service associated with the device maker which the user can visit to manage fabric limitations. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme for internet-hosted URLs.

- The URL MAY embed the token, fabric index, fabric vendor, or other information transparently in order to pass context about the originating [ReviewFabricRestrictions](#) command to the service associated with the URL. The service associated with the device vendor MAY perform vendor ID verification on the fabric from which the [ReviewFabricRestrictions](#) command originated.
- If the device grants the request, the ARL attribute in the [Access Control Cluster](#) SHALL be updated to reflect the new access rights and a successful response SHALL be returned to the device making the request using the [MTaer](#) field of the [callbackUrl](#). If the request is denied, the ARL attribute SHALL remain unchanged and a failure response SHALL be returned to the device making the request using the [MTaer](#) field of the [callbackUrl](#).
- The device using this mechanism SHALL provide a service at the URL that can accept requests for additional access and return responses indicating whether the requests were granted or denied.
- This URL will typically lead to a server which (e.g. by looking at the User-Agent) redirects the user to allow viewing, downloading, installing or using a manufacturer-provided means for guiding the user through the process to review and approve or deny the request. The device manufacturer MAY choose to use a constructed URL which is valid in a HTTP GET request (i.e. dedicated for the product) such as, for example, <https://domain.example/arl-app?vid=FFF1&pid=1234>. If a client follows or launches the ARLRequestFlowUrl, it SHALL expand it as described in [Section 9.10.9.3.4, "ARLRequestFlowUrl format"](#).

- A manufacturer contemplating using this flow should realize that
 - This flow typically requires internet access to access the URL, and access extension may fail when internet connectivity is not available.
 - If the flow prefers to redirect the user to an app which is available on popular platforms, it SHOULD also provide a fallback option such as a web browser interface to ensure users can complete access extension.
 - A malicious Administrator could tamper with the URL (including any parameters it contains) in order to reduce restrictions for another Fabric. A well-implemented web service or app SHOULD validate that the VID in flow URL invocations match recent requests for review by comparing with VID and token pairs from recent [ReviewFabricRestrictions](#) requests.
- An Administrator supporting this flow should realize that if the device serving this cluster is malicious, it could send dangerous URLs to the client which could take the user to malicious sites. The Administrator SHOULD provide details, such as the VID, to the user about the device sending the event so that the user can make an informed decision about whether to trust the URL.

ARLRequestFlowUrl format

The **ARLRequestFlowUrl** SHALL contain a query component (see [RFC 3986](#) section 3.4) composed of one or more key-value pairs:

- The query SHALL use the **&** delimiter between key/value pairs.
- The key-value pairs SHALL in the format **name=<value>** where **name** is the key name, and **<value>** is the contents of the value encoded with proper URL-encoded escaping.
- If key **MTcu** is present, it SHALL have a value of **"_"** (i.e. **MTcu=_**). This is the "callback URL (**CallbackUrl**) placeholder".
- Any key whose name begins with **MT** not mentioned in the previous bullets SHALL be reserved for future use by this specification. Manufacturers SHALL NOT include query keys starting with **MT** in the **ARLRequestFlowUrl** unless they are referenced by a version of this specification.

Any other element in the **ARLRequestFlowUrl** query field not covered by the above rules, as well as the fragment field (if present), SHALL remain including the order of query key/value pairs present.

Expansion of ARLRequestFlowUrl by client

Once the URL is obtained, it SHALL be expanded to form a final URL (**ExpandedARLRequestFlowUrl**) by proceeding with the following substitution algorithm on the original **ARLRequestFlowUrl**:

1. If key **MTcu** is present, compute the **CallbackUrl** desired (see [Section 9.10.9.3.5, "CallbackUrl format for ARL Request Flow response"](#)), and substitute the placeholder value **"_"** (i.e. in **MTcu=_**) in the **ARLRequestFlowUrl** with the desired contents, encoded with proper URL-encoded escaping (see [RFC 3986](#) section 2).

The final URL after expansion (**ExpandedARLRequestFlowUrl**) SHALL be the one to follow, rather than the original value obtained from the **FabricRestrictionReviewUpdate** event.

CallbackUrl format for ARL Request Flow response

If a **CallbackUrl** field (i.e. **MTcu=**) query field placeholder is present in the **ARLRequestFlowUrl**, the client MAY replace the placeholder value "_" in the **ExpandedARLRequestFlowUrl** with a URL that the manufacturer flow can use to make a smooth return to the client when the ARL flow has terminated.

This URL field MAY contain a query component (see [RFC 3986](#) section 3.4).

If a query is present, it SHALL be composed of one or more key-value pairs:

- The query SHALL use the **&** delimiter between key/value pairs.
- The key-value pairs SHALL follow the format **name=<value>** where **name** is the key name, and **<value>** is the contents of the value encoded with proper URL-encoded escaping.
- If key **MTaer** is present, it SHALL have a value of "_" (i.e. **MTaer=**_). This is the placeholder for a "access extension response" provided by the manufacturer flow to the client. The manufacturer flow SHALL replace this placeholder with the final status of the access extension request, which SHALL be formatted following [Expansion of CallbackUrl by the manufacturer custom flow](#) and encoded with proper URL-encoded escaping.
- Any key whose name begins with **MT** not mentioned in the previous bullets SHALL be reserved for future use by this specification.

Any other element in the **CallbackUrl** query field not covered by the above rules, as well as the fragment field (if present), SHALL remain as provided by the client through embedding within the **ExpandedARLRequestFlowUrl**, including the order of query key/value pairs present.

Expansion of CallbackUrl by the manufacturer custom flow

Once the **CallbackUrl** is obtained by the manufacturer flow, it MAY be expanded to form a final **ExpandedARLRequestCallbackUrl** URL to be used by proceeding with the following substitution algorithm on the provided **CallbackUrl**:

- If key **MTaer** is present, the manufacturer custom flow having received the initial query containing the **CallbackUrl** SHALL substitute the placeholder value "_" (i.e. in **MTaer=**_) in the **CallbackUrl** with the final status of the access extension request flow which SHALL be one of the following. Any value returned in the **MTaer** field not listed above SHALL be considered an error and SHALL be treated as **GeneralFailure**.
 - **Success** - The flow completed successfully and the ARL attribute was updated. The client MAY now read the ARL attribute to determine the new access restrictions.
 - **NoChange** - The ARL attribute was already listing minimum restrictions for the requesting fabric.
 - **GeneralFailure** - The flow failed for an unspecified reason.
 - **FlowAuthFailure** - The user failed to authenticate to the flow.
 - **NotFound** - Access extension failed because the target fabric was not found.

A manufacturer custom flow having received an **ExpandedARLRequestFlowUrl** SHOULD attempt to open the **ExpandedARLRequestCallbackUrl**, on completion of the request, if an **ExpandedARLRequest-**

CallbackUrl was computed from the **CallbackUrl** and opening such a URL is supported.

Examples of ARLRequestFlowUrl URLs

Below are some examples of valid **ExpandedARLRequestFlowUrl** for several valid values of **ARLRequestFlowUrl**, as well as some examples of invalid values of **ARLRequestFlowUrl**:

- Invalid URL with no query string: **http** scheme is not allowed:

- <http://company.domain.example/matter/arl/vFFF1p1234>

- Valid URL :

- <https://company.domain.example/matter/arl/vFFF1p1234>

- Valid URL, **CallbackUrl** requested:

- Before expansion:

```
https://company.domain.example/matter/arl?vid=FFF1&pid=1234&MTcu=_
```

- After expansion:

```
https://company.domain.example/matter/arl?vid=FFF1&pid=1234&MTcu=https%3A%2F%2Fclient.domain.example%2Fcb%3Ftoken%3DmAsJ6_vqbr-vjDiG_w%253D%253D%26MTaer%3D_
```

- The **ExpandedARLRequestFlowUrl** URL contains:

- A **CallbackUrl** with a client-provided arbitrary **token=** key/value pair and the **MTaer=** key/value pair place-holder to indicate support for a return access extension completion status: https://client.domain.example/cb?token=mAsJ6_vqbr-vjDiG_w%3D%3D&MTaer=_
- After expansion of the **CallbackUrl** (**MTcu** key) into an **ExpandedCallbackUrl**, with an example return access extension completion status of **Success**, the **ExpandedARLRequestCallbackUrl** would be:

```
https://client.domain.example/cb?token=mAsJ6_vqbr-vjDiG_w%3D%3D&MTaer=Success
```

NOTE

The **MTcu** key/value pair was initially provided URL-encoded within the **ExpandedARLRequestFlowUrl** URL and the **MTaer=_** key/value pair place-holder now contains a substituted returned completion status.

- Invalid URL, due to **MTza=79** key/value pair in reserved **MT**-prefixed keys reserved for future use:

- https://company.domain.example/matter/arl?vid=FFF1&pid=1234&MTop=_&MTza=79

9.10.9.4. AuxiliaryAccessUpdated Event

The server SHALL generate AuxiliaryAccessUpdated events whenever its **AuxiliaryACL** attribute data changes due to system management operations.

This event **SHOULD** be generated only once per batch of changes, if many changes are done together.

Administrators **MAY** use this event to trigger re-evaluation of overall access control currently granted to the node.

Leaving a Fabric **SHALL NOT** generate this event.

Access Modifier: Fabric-Sensitive						
ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	AdminNodeID	node-id	all	X		M

AdminNodeID Field

The AdminNodeID field **SHALL** contain the NodeID of the Administrator that caused the action which led to changes to the AuxiliaryACL. If no information is available, such as when a change is internally initiated, this field **SHALL** be null.

9.11. Group Relationship

A group is a collection of one or more endpoints on one or more nodes. A group is identified by a unique group ID. If the network supports fabrics, each group, its group ID, and nodes that are members of the group, are all scoped to a single fabric.

Conceptually, there is a Group Table on each node that represents endpoint group membership. Each Group Table entry maps a group ID to one or more endpoints on that node, and any endpoint on a node **MAY** be assigned to one or more groups.

A group relationship, that is contained in the Group Table, is managed through the endpoints using the Groups cluster.

The Interaction Model allows a group identifier to be used as the destination of a message or action. If a message received by a node has a group destination, the Group Table is checked to see which endpoints on the node are members of the group. Then, the message will be delivered to those endpoints.

Note that there is a risk that multiple clients allocate the same group identifier for their own purpose. This likely leads to undesired behavior. For this reason, a client **SHOULD** discover the uniqueness of their 'candidate' group ID.

Also note that groupcast relies on its support by the underlying network layer. Depending on this network layer, groupcast may not work to "sleepy" devices that have their radio turned off when idle to preserve battery lifetime.

9.12. Bridge for non-Matter devices

9.12.1. Introduction

A Bridge serves to allow the use of non-Matter IoT devices (e.g. devices on a Zigbee or Z-Wave network, or any other non-Matter connectivity technology) in a Matter Fabric, with the goal to enable the consumer to keep using these non-Matter devices together with their Matter devices.

This is illustrated in the figure below: the non-Matter devices are exposed as Bridged Devices to Nodes on the Fabric. The Matter Nodes can communicate with both the (native) Matter devices as well as the Bridged Devices (by virtue of the Bridge which performs the translation between Matter and the other protocol).




Figure 65. principle of bridging

NOTE

The bridging-concept described here is NOT intended to be used to expose Matter Nodes (which are not on the Fabric) to a Fabric, since direct connection of those Matter Nodes to the Fabric would provide a better experience.

This section will describe how the Data Model concepts can be used by a Bridge to expose the Bridged Devices in such a way that they can be discovered and used by Nodes on the Matter Fabric.

9.12.2. Exposing functionality and metadata of Bridged Devices

After Commissioning, the Bridge SHALL expose (at least) one Node to the Fabric. The device implementing the Bridge MAY have more than one Node. This, however, is orthogonal to the bridging concept and will not be discussed further here.

- On this Node, the Bridge SHALL expose a set of endpoints representing the various Bridged Devices on the non-Matter side of the Bridge.

- Additionally, it SHALL expose an endpoint with the device type **Aggregator** which has a **Descriptor** cluster with a **PartsList** attribute containing all the endpoints representing those Bridged Devices.
- Each Bridged Device is represented by one or more endpoints listed in this **PartsList** (see [Section 9.12.2.3, “Information about Bridged Devices”](#) and examples below). The **Descriptor** cluster on the corresponding endpoint provides information about the particular Bridged Device, such as its device type(s).




Figure 66. example of endpoints representing Bridged Devices

In case the Bridge is bridging to/from multiple technologies (or has some other logical distinction between groups of bridged devices), it MAY expose such groups as two or more such hierarchical trees each with their **Aggregator** device type (e.g. one for each technology, see figure below); it MAY also combine all bridged devices in one hierarchical tree (see figure above). Both figures have the same set of bridged devices - the difference is in how the bridge manufacturer decides to expose them as one or multiple sets.




Figure 67. example of endpoints representing Bridged Devices from two technologies

9.12.2.1. Topology or logical grouping

A Bridge typically has information on topology or logical grouping of the Bridged Devices, which can be of use to Nodes on the Matter Fabric.

- For example, consider a Bridge with 50 lights. If this exposure of grouping, and their naming, would not be present, the user would just see a list of 50 lights on their controller and would not know which of those physical lights would be in which location/group.

If a Bridge has such information on topology or logical grouping, it SHOULD expose such information in the EndpointLists attribute of an **Actions** cluster (the ActionLists of which MAY be empty if no actions are exposed). A Bridge MAY make it possible (e.g., through a Bridge Manufacturer’s app) for its users to restrict whether all or some of such information is exposed to the Fabric. The Node on the Fabric using the Bridged Devices which is interested in using such topology or logical grouping (e.g. to show the grouping of lights per room in an overview to the user), SHOULD derive such grouping (and associated naming) from this EndpointLists attribute.

In the example below, the devices are split over two rooms, as exposed in the EndpointLists attribute. This example also illustrates a compound endpoint for a compound Bridged Device, in this case a lighting device (Bridged Device at EP 24,25,26) which has an up- and downlighter which can be controlled separately, and which have their own set of lighting-related clusters on an individual endpoint (EP 25,26). Note that the **Bridged Device Basic Information** cluster is at the top of the hierarchy for this compound device (EP 24), while the application device types and application clusters are at the leaf endpoints (EP 25,26).

Since EP 25,26 are listed in the PartsList of EP 24, they 'inherit' the Bridged Node device type and information in the **Bridged Device Basic Information** cluster of EP 24.




Figure 68. example of endpoints representing Bridged Devices using grouping




Figure 69. impression of app UI indicating information for the Bridged Devices

9.12.2.2. Native Matter functionality in Bridge

The Bridge MAY also contain native Matter functionality, i.e. non-bridged functionality, such as in the example below, which shows a smart speaker device having, in addition to a Wi-Fi connection, also a Zigbee connection towards a number of Zigbee lights. The speaker functionality (EP 31) is native Matter functionality (and could have a Controller role to allow sending Matter commands upon receiving voice commands), while the remainder of the non-zero endpoints represent the Bridged Devices.




Figure 70. example of Bridge with native Matter functionality

9.12.2.3. Information about Bridged Devices

For each Bridged Device, the Bridge SHALL include a **Bridged Device Basic Information** cluster on the endpoint which represents this Bridged Device. In case a Bridged Device is represented by multiple endpoints, the Bridged Device Basic Information cluster SHALL only be present on the endpoint which is the top level of the hierarchy representing this Bridged Device (example: endpoint 24 in Figure 68, “example of endpoints representing Bridged Devices using grouping”).

On this endpoint with the Bridged Device Basic Information cluster, the Bridge SHALL also include a **Descriptor** cluster with

- a DeviceTypeList attribute containing device type Bridged Node plus the device type(s) of the Bridged Device, and
- a PartsList attribute listing any other endpoints which jointly expose the functionality of this Bridged Device.

Information about power sources for Bridged Devices

In case the Bridged Device contains a power source such as battery or mains power feed, *and* information about the state of that power source is available to the Bridge, the Bridge SHALL also include one or more **Power Source** cluster(s), where their EndpointList attribute SHALL contain the endpoint(s) of the Bridged Device that are powered by this power source. Each endpoint with a Power Source cluster SHALL have the related Power Source device type in its DeviceTypeList.

For Bridged Devices which are represented by a single endpoint:

- The **Power Source** cluster SHALL be included on that single endpoint. An example of this is shown for the battery-powered temperature sensor on endpoint 23 in [Figure 68, “example of endpoints representing Bridged Devices using grouping”](#).

For Bridged Devices which are represented by multiple endpoints:

- If the Bridged Device contains only one power source, the **Power Source** cluster SHALL be present on an endpoint which corresponds to the part of the Bridged Device being powered by the power source.
 - In case this power source provides power to the entire Bridged Device, the power source cluster SHALL be on the endpoint where the Bridged Node device type is located, and contain an EndpointList attribute containing all the endpoints constituting the Bridged Device.

In case a Bridged Device does not contain a power source such as battery or mains power feed, *or* information about the state of that power source is not available to the Bridge, the Bridge SHALL NOT include a **Power Source** cluster on the endpoint(s) representing the Bridged Device.

9.12.2.4. Clusters for Bridged Device functionality (device types)

For each Bridged Device, the Bridge SHALL expose the clusters required for a device of the indicated Matter device type(s).

This allows the Matter Nodes to recognize the device type of the Bridged Device and interact with its clusters in the same manner as with a native Matter Node of that device type.

9.12.2.5. Identify for Bridge and Bridged Devices

If the bridge has a mechanism to identify itself to the user (e.g. blinking LED on the bridge itself), the associated Identify cluster SHOULD be used on the endpoint with the Aggregator device type.

For the identification function of individual bridged devices, the conformance for the associated Identify cluster is determined by the application device types on the endpoint(s) representing the bridged device. For example, in the composed lighting device in [Figure 68, “example of endpoints representing Bridged Devices using grouping”](#), each of the endpoints 25 and 26 would have an identify cluster to allow individual identification of each of those lights.

9.12.3. Discovery of Bridged Devices

A Node which discovers another Node with device type **Aggregator** on one of its endpoints SHOULD walk the entire tree of endpoints via the PartsList attributes and endpoints to discover the list of Bridged Devices, including their device types and other attributes, as well as any native Matter functionality potentially present on the Node.

Each endpoint found containing a **Bridged Node** device type represents a Bridged Device of the device type(s) specified at this endpoint, or one of the endpoints found via its PartsList. If the discovering Node supports this device type, it MAY add this Bridged Device to the list of devices which it could interact with, or could set up configuration for.

This discovering Node SHALL use the acquired information on the available Bridged Devices to set up (configure) (likely with input from the user) how the Bridged Devices can be used with the Matter Nodes (e.g. which switch controls which light, or how to control a light from an app on the user's phone).

Since the Bridge may expose a large number of Bridged Devices, the discovering Node SHALL use the **NodeLabel** attribute in the **Bridged Device Basic Information** cluster of each of the Bridged Devices to allow the user to easily identify and recognize the various Bridged Devices, and expedite the setup/configuration process, rather than present the user with an unannotated list of, for example, 20 lights, 3 sensors and 4 switches. These labels have likely been populated by the user when interacting previously with the Bridge e.g. through means provided by the Bridge Manufacturer, such as a Bridge Manufacturer app.

If power source-related information regarding the Bridged Device is provided in the **Power Source** cluster on the associated endpoint, the discovering Node SHOULD use this information in a similar manner as power source-related information acquired from a Matter Node's **Power Source** cluster. Such information can then be used to inform the user about the state of the power source (e.g. warn about low batteries) in a Bridged Device in a similar manner as done for Matter Nodes.

9.12.4. Configuration of Bridged Devices

For configuration of the discovered Bridged Devices, two basic archetypes are described in the following sections: one for actuators and one for sensors/switches.

Since a Bridged Device of a certain device type has the same set of application clusters as a native Matter device of the same device type, this process is similar to configuring a native Matter device of that device type.

9.12.4.1. Sending commands from a Matter controller to a Bridged Device

For Bridged Devices that are actuators and hence have a Controlee role, a Controller Node on the Fabric MAY send commands to the associated clusters on one or more endpoints on the Bridge's Node, such as an **On** command to the **On/Off** cluster of a Bridged Device. The Bridge SHALL forward this command to the relevant Bridged Device after conversion between the Matter protocol and the non-Matter device's native protocol.

Example:

A Controller creates a Group containing some Matter lights as well as some non-Matter lights, by sending an **AddGroup** command to the instances of the **Group** cluster on both the endpoints of the Matter lights as well as on the Bridge's Node endpoints representing the targeted bridged lights.

Similarly, the Controller creates one or more Scenes using the instances of the **Scenes Management** cluster on these endpoints.

The Controller then sends a (single) **On** command (**On/Off** cluster) to this group to switch on all these lights in a single operation. This (single) multicast message will be received (and interpreted) by the Matter lights which are part of this group as well as by the Bridge, which will forward it (after appropriate protocol conversion) to the relevant bridged lights.

Similarly, the Controller sends a (single) **MoveToLevel** (**Level Control** cluster) or sends a (single) **RecallScene** (**Scenes Management** cluster) to this group, to set the brightness resp. recall a scene contents on all these lights in a single operation.




Figure 71. example of bridging lights

9.12.4.2. Receiving status/events/commands from a Bridged Device

For Bridged Devices like sensors and switches, the Bridge will receive value updates (e.g. Zigbee attribute reports), events and/or commands from those devices, and SHALL (after conversion from the native protocol of the non-Matter devices towards Matter protocol) represent those as attributes, events and/or commands in the appropriate clusters on the associated endpoints of the Bridge. Interactions with those attributes/events/commands on the Matter side (e.g., towards a Controller using the sensor/switch data) SHOULD be identical to interactions with corresponding attributes/events/commands in native Matter sensors and switches (e.g., attribute readout and subscription, proxying and eventing).

Examples:

- A temperature sensor sends a status report to the Bridge over a non-Matter interface. The logic in the Bridge processes this as an update to the **Measured Value** attribute of the **Temperature Measurement** cluster on the endpoint associated with this bridged sensor. Nodes on the Fabric which have an interest in this value can read the updated attribute value, and can configure a subscription on this attribute. This is identical to reading an attribute value or setting up an attribute subscription on a native-Matter temperature sensor Node.
- A user presses a button on a (push-button) switch device. The switch device sends a message to

the Bridge over a non-Matter interface. The logic in the Bridge processes this to generate an **InitialPress** event (**Switch** cluster) on the endpoint representing the switch.

Nodes on the Fabric which have an interest in the switch operation can setup eventing from this cluster.




Figure 72. example of bridging switches and sensors

9.12.5. New features for Bridged Devices

Bridged Devices can have their software updated independently of the Bridge, through Bridge Manufacturer-specific means. These updates MAY result in one or more changes to their capabilities, such as supported clusters and/or attributes, for an endpoint. Like every Matter Node, every endpoint on the Bridge's Node contains a **Descriptor** cluster that contains attributes for the device types (DeviceTypeList), endpoints (PartsList) and supported clusters (ServerList and ClientList). The **ConfigurationVersion** attribute in the [Section 9.13, "Bridged Device Basic Information Cluster"](#) SHALL be increased as needed for such changes, along with an increase of the [Section 11.1.5.24, "ConfigurationVersion"](#) attribute in the [Section 11.1, "Basic Information Cluster"](#) of the bridge. Nodes that wish to be notified of such changes SHOULD monitor changes of these attributes.

9.12.6. Changes to the set of Bridged Devices

Bridged Devices can be added to or removed from the Bridge through Bridge-specific means. For example, the user can use a Manufacturer-provided app to add/remove Zigbee devices to/from their Matter-Zigbee Bridge.

When an update to the set of Bridged Devices (which are exposed according to the [Section 9.12.11, "Best practices for Bridge Manufacturers"](#)) occurs, the Bridge SHALL

- on the **Descriptor** clusters of the **root node endpoint** and of the endpoint which holds the **Aggregator** device type: update the **PartsList** attribute (add/remove entries from this list)
- update the exposed endpoints and their descriptors according to the new set of Bridged Devices
- increase the [Section 11.1.5.24, "ConfigurationVersion"](#) attribute in the **Basic Information Cluster**

of the bridge itself, if that attribute is supported.

Nodes that wish to be notified of added/removed devices SHOULD monitor changes of the PartsList attributes in the Descriptor cluster on the [root node endpoint](#) and on the endpoint which holds the [Aggregator](#) device type, as well as changes of the ConfigurationVersion attribute in the BasicInformation cluster on the root node endpoint.

Allocation of endpoints for Bridged Devices SHALL be performed as described in [Dynamic Endpoint allocation](#).

9.12.7. Changes to device names and grouping of Bridged Devices

Typically, the user has some means (e.g. a Manufacturer-provided app) to assign names to the Bridged Devices, or names could be assigned automatically by the Bridge. The Bridge SHALL expose such names in the NodeLabel attribute of the Bridged Device Basic Information cluster on the applicable endpoint.

Similarly, the user typically has some means to group the Bridged Devices (e.g. via a room/zone-concept) and provide names to such groups, or grouping could be applied automatically by the Bridge. The Bridge SHOULD expose such grouping using the EndpointLists attribute of the Actions cluster as described above.

For such exposed information, when there is a change in naming/grouping (e.g. the user makes changes via a Manufacturer-provided app), the Bridge SHALL update the appropriate attributes.

Nodes that wish to be notified of a change in such a name or grouping SHOULD monitor changes of this attribute or cluster.

9.12.8. Setup flow for a Bridge (plus Bridged Devices)

As described above, the Bridge together with its Bridged Devices is exposed as a single Node with a list of endpoints. Consequently, a single Node ID and a single Operational Certificate is assigned during Commissioning and a single pass through the commissioning flow is required to bring the Bridge (along with its Bridged Devices) onto a Fabric. This provides for a simple user experience, since the user only needs to go through the commissioning flow for the Bridge, and not separately for each of the Bridged Devices.

9.12.9. Access Control

The Bridge is a Matter node, therefore it has a single [Section 9.10, “Access Control Cluster”](#) for the entire Node, like every other Matter Node. This cluster contains all [Access Control Entries](#) for each of its endpoints, including for all Bridged Devices and other native Matter functionality exposed by the Bridge Node.

A typical setup of Access Control would reflect which privilege level a Matter Controller needs to have for one or more Bridged Devices. This overall access set MAY be a subset of all the Bridged Devices on the Bridge, rather than all endpoints on a Bridge. This can be accomplished by setting an Access Control Entry containing as targets a list of those endpoints representing a Bridged Device or a set of Bridged Devices. As defined in [the ACL model](#), it is also possible to specify access towards specific [Targets](#), for example all Bridged Devices of device type [Extended Color Light](#).

9.12.10. Software update (OTA)

The Bridge is a Matter device and its Matter-related functionality MAY be updated using the mechanism described in [Section 11.20, “Over-the-Air \(OTA\) Software Update”](#).

The Bridged Devices, on the other hand, are not native Matter devices, do not have a Product ID, and are not listed in the [Section 11.23, “Distributed Compliance Ledger”](#). They are typically updated using a mechanism defined and deployed by the Bridge Manufacturer. That same mechanism is typically used to update the parts of the Bridge which are not related to Matter, which is particularly relevant to allow synchronization of updates to the non-Matter part of the Bridge with updates to the Bridged Devices. Obviously, such mechanism MAY also be employed to update the entire Bridge firmware, including the Matter-related functionality.

9.12.11. Best practices for Bridge Manufacturers

This section summarizes (in order of priority) the process to determine which non-Matter devices the Bridge exposes to the Fabric.

1. Choice of supported device types

- A Manufacturer MAY choose which of the Matter device types they can or want to support in the Bridge. After implementation of support for bridging of those device types, they SHALL certify the Bridge for those device types.
- By default, a Bridge SHOULD expose to the Fabric all its connected non-Matter devices which can be mapped to a Matter device type for which that Bridge is certified.
E.g., if a Bridge is certified for Matter light bulbs, it SHOULD NOT hide any light bulb on the non-Matter side from the Fabric by default (some situations where the Bridge MAY deviate from this recommendation are in the following text).
- Given the wide variety of device types on a wide variety of standards, there may be device types on the non-Matter side that do not have a corresponding Matter device type. Such devices cannot be bridged to a Matter device type. The Manufacturer MAY choose to not expose such devices with the Bridge or MAY expose them with a manufacturer-specific device type and/or manufacturer-specific clusters.

2. Compatibility issues

- For the device types for which a Bridge is certified, a Bridge Manufacturer MAY decide to not expose certain devices based on any reason, including compatibility and interoperability reasons, or to expose them in a 'best-effort' manner as needed.
 - The Bridge Manufacturer may choose to not expose a device that does not support certain functions or features which are mandatory for a Matter device type, but which are defined as optional, or not defined at all, in the specification for the corresponding device type on the non-Matter side of the bridge. Such a Bridge would expose to the Fabric only Bridged Devices of device types which support those particular control functions or features which are required.
 - The Bridge Manufacturer may also choose to map or emulate such features which are not available in the Bridged Device; example: for a bridged colored light connected via a protocol which does not support scenes, the Bridge could emulate the scene function by storing the scenes in the Bridge and sending corresponding brightness and color com-

mands to the light when a Scene Recall command is received from a Matter Node.

3. User choice

- A Bridge MAY make it possible (e.g., through a Bridge Manufacturer's app) for its users to further restrict which devices are exposed to the Fabric.

For example, a user may decide to prevent exposure to the Fabric of certain Devices Types, such as all occupancy sensors, or of only certain devices of a certain device type, such as only their bedroom occupancy sensor.

9.12.12. Best practices for Administrators

An Administrator MAY indicate to users which devices are native Matter and which ones are Bridged Devices, as determined using the presence of a **Bridged Node** device type on the endpoint, in order to ensure the user does not make assumptions about the Bridged Devices having the same security requirements as native Matter devices.

9.13. Bridged Device Basic Information Cluster

This cluster provides attributes and events for determining basic information about Bridged Nodes.

This cluster is derived from the Basic Information cluster and serves two purposes towards a Node communicating with a Bridge:

- Indicate that the functionality on the Endpoint where it is placed (and its Parts) is bridged, and
- Provide a centralized collection of attributes that the Node MAY collect to aid in conveying information regarding the Bridged Device to a user, such as the vendor name, the model name, or user-assigned name.

This cluster SHALL be exposed by a Bridge on the Endpoint representing each Bridged Device. When the functionality of a Bridged Device is represented using a set of Endpoints, this cluster SHALL only be exposed on the Endpoint which is at the top of the hierarchy for the functionality of that Bridged Device.

This cluster SHALL NOT be used on an endpoint that is not in the Descriptor cluster PartsList of an endpoint with an Aggregator device type.

This cluster has been derived from the **Basic Information** Cluster, and provides generic information about the Bridged Device. Not all of the attributes in the Basic Information Cluster are relevant for a Bridged Device (e.g. ProductID since it is not a Matter device). For other attributes, the information which is listed as Mandatory for the Basic Information Cluster, may not be available when the Bridged Device does not provide it to the Bridge, and the Bridge has no other means to determine it. For such cases where the information for a particular attribute is not available, the Bridge SHOULD NOT include the attribute in the cluster for this Bridged Device. See below for Conformance details.

9.13.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Added ProductAppearance attribute
3	Added SpecificationVersion and MaxPathsPerInvoke attributes
4	Updated conformance for UniqueID to mandatory, ProductID to optional when bridging Matter devices, add the BridgedICDSupport feature
5	Added ConfigurationVersion attribute
6	Adopt latest version of base cluster (Basic Information).

9.13.2. Classification

Hierarchy	Role	Scope	PICS Code
Basic Information	Utility	Endpoint	BRBINFO

9.13.3. Cluster ID

ID	Name
0x0039	Bridged Device Basic Information

9.13.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Conformance	Summary
20	BIS	BridgedICDSupport	[Rev >= v4]	Support bridged ICDs.

9.13.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	DataModelRevision						X
0x0001	VendorName						desc
0x0002	VendorID						desc
0x0003	ProductName						O

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0004	ProductID						desc
0x0005	NodeLabel						O
0x0006	Location						X
0x0007	Hardware-Version						O
0x0008	Hardware-Version-String						O
0x0009	Software-Version						O
0x000A	Software-Version-String						O
0x000B	Manufacturing-Date						O
0x000C	PartNumber						O
0x000D	ProductURL						O
0x000E	Product-Label						O
0x000F	Serial-Number						O
0x0010	LocalConfigDisabled						X
0x0011	Reachable						M
0x0012	UniqueID						Rev >= v4, O
0x0013	CapabilityMinima						X
0x0014	ProductAppearance						[Rev >= v2]
0x0015	SpecificationVersion						X

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0016	Max-PathsPer-Invoke						X
0x0018	ConfigurationVersion						[Rev >= v6]

Since this cluster has been derived from the [Basic Information](#) Cluster, the identifiers of the attributes, their range, quality and default characteristics and their descriptions correspond to those in that Cluster and those descriptions are not repeated here. Several attributes from the Basic Information Cluster which are *not* relevant or applicable for a Bridged Device have been marked with X in column Conformance and SHALL NOT be used in the Bridged Device Basic Information Cluster. The Conformance characteristics of several attributes in this cluster have changed from M to O compared to their Conformance in the Basic Information Cluster, and SHALL be used according to the table above.

9.13.5.1. VendorName, VendorID & ProductID Attributes

The Bridge SHOULD fill these attributes with the available information, which could e.g. come from the Bridged Device provided to the Bridge over the non-Matter interface (e.g. VendorName and VendorID) or could have been provided by the user (e.g. assigned name of a device for NodeLabel).

If the manufacturer of a Bridged Device is known to the Bridge, the Bridge SHALL provide this name (in attribute VendorName), otherwise it SHALL NOT include this attribute.

If the manufacturer of a Bridged Device and the associated Alliance-assigned [Vendor ID](#) are known to the Bridge (e.g. by copying the Manufacturer Code from the Node Descriptor of a Zigbee device or the Basic Information cluster of a Matter device), the Bridge SHALL provide this identifier (in attribute VendorID), otherwise it SHALL NOT include this attribute.

If the associated [Product ID](#) of a bridged Matter device is known to the Bridge (e.g. from the Basic Information cluster of the bridged Matter device), the Bridge SHALL provide this identifier (in attribute ProductID), otherwise it SHALL NOT include this attribute. For bridged non-Matter devices, the ProductID attribute SHALL NOT be included.

9.13.5.2. Reachable Attribute

This attribute SHALL be used to indicate whether the bridged device is reachable by the bridge, so a Matter Node which wants to communicate with a bridged device can get an indication that this might fail (when the attribute is False). Determination of reachability might not be perfect (e.g. depending on technology employed), so the Matter Node SHOULD be aware of the risk of false positives and negatives on reachability determination. For example, a bridged device MAY be marked as unreachable while it could actually be reached, and vice-versa. Also, detection (and indication) that a bridged device is not longer reachable may be delayed due to the technique employed (e.g. detecting that a number of expected messages from the bridged device did not arrive). Also see event ReachableChanged below.

9.13.5.3. UniqueID Attribute

This attribute SHALL, for a Bridged Device, be updated when the Bridge is factory reset. If the bridged device does not provide some unique id (e.g. in the case of bridging from non-Matter devices, or in case of bridging Matter devices from an earlier revision which were not required to provide a UniqueID attribute), the bridge SHALL generate a unique id on behalf of the bridged device.

NOTE

The UniqueID attribute was optional in cluster revisions prior to revision 4.

9.13.5.4. ConfigurationVersion Attribute

This attribute SHALL contain the current version number for the configuration of the bridged device. A larger value of ConfigurationVersion SHALL indicate a newer configuration than a lower value.

If the bridge detects a change on a bridged device, which it deems as a change in the configuration of the bridged device, it SHALL increase this attribute (and the corresponding attribute on the bridge itself) as described in [Section 9.2.11, “Node Configuration Changes”](#).

The ability and the method used to detect such a change on a bridged device is manufacturer specific.

9.13.6. Commands

ID	Name	Direction	Response	Access	Conformance
0x80	KeepActive	client ⇒ server	Y	O	BIS

9.13.6.1. KeepActive Command

Upon receipt, the server SHALL attempt to keep the bridged device active for the duration specified by the command, when the device is next active.

The implementation of this is best-effort since it may interact with non-native protocols. However, several specific protocol requirements are:

- If the bridged device is a Matter Intermittently Connected Device, then the server SHALL send a [StayActiveRequest](#) command with the StayActiveDuration field set to value of the StayActiveDuration field in the received command to the bridged device when the bridged device next sends a checks-in message or subscription report. See [Intermittently Connected Devices Behavior](#) for details on ICD state management.

When the bridge detects that the bridged device goes into an active state, an ActiveChanged event SHALL be generated.

In order to avoid unnecessary power consumption in the bridged device:

- The server SHALL enter a "pending active" state for the associated device when the KeepActive command is received. The server "pending active" state SHALL expire after the amount of time

defined by the TimeoutMs field, in milliseconds, if no subsequent KeepActive command is received. When a KeepActive command is received, the "pending active" state is set, the StayActiveDuration is updated to the greater of the new value and the previously stored value, and the TimeoutMs is updated to the greater of the new value and the remaining time until the prior "pending active" state expires.

- The server SHALL only keep the bridged device active once for a request. (The server SHALL only consider the operation performed if an associated ActiveChanged event was generated.)

This command SHALL have the following data fields:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	StayActiveDuration	uint32	all			M
1	TimeoutMs	uint32	30000 to 3600000			M

StayActiveDuration Field

This field SHALL indicate the duration, in milliseconds, that the device is requested to remain active, once the device becomes active again.

The value of this field MAY be longer than the value supported by the bridged device and would, typically, be used by the client to request the server of the bridged device to stay active and responsive for this period to allow a sequence of message exchanges during that period.

The client MAY slightly overestimate the duration it wants the bridged device to be active for, in order to account for network delays.

TimeoutMs Field

This field SHALL indicate the period, in milliseconds, that the server will wait before the "pending active" state expires. See the [KeepActive Command](#) description for details.

NOTE

TimeoutMs is a timeout for the request, NOT the time the device will be awake for. The server will wait for up to TimeoutMs for the device. If after TimeoutMs the ICD device does NOT check-in, the server will not perform any actions.

9.13.7. Events

ID	Name	Priority	Quality	Access	Conformance
0x00	StartUp				O
0x01	ShutDown				O
0x02	Leave				O
0x03	ReachableChanged				M

ID	Name	Priority	Quality	Access	Conformance
0x80	ActiveChanged	INFO		V	BIS

9.13.7.1. Leave Event

The Leave event SHOULD be generated by the bridge when it detects that the associated device has left the non-Matter network.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	FabricIndex					X

NOTE

The **FabricIndex** field has the **X** conformance, indicating it SHALL NOT be present. This event, in the context of Bridged Device Basic Information cluster, has no usable fields, but the original Basic Information cluster's field definition is kept for completeness.

9.13.7.2. ReachableChanged Event

This event SHALL be generated when there is a change in the Reachable attribute. Its purpose is to provide an indication towards interested parties that the reachability of a bridged device has changed over its native connectivity technology, so they MAY take appropriate action.

After (re)start of a bridge this event MAY be generated.

9.13.7.3. ActiveChanged Event

This event (when supported) SHALL be generated the next time a bridged device becomes active after a KeepActive command is received.

See [KeepActive](#) for more details.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	PromisedActiveDuration	uint32	desc			M

PromisedActiveDuration Field

This field SHALL indicate the minimum duration, in milliseconds, that the bridged device will remain active after receiving the initial request from the [KeepActive](#) processing steps.

If the bridged device is a Matter Intermittently Connected Device, PromisedActiveDuration SHALL be set to the PromisedActiveDuration value returned in the [StayActiveResponse](#) command.

If the bridged device is not a Matter Intermittently Connected Device, the implementation of this is best-effort since it may interact with non-native protocol.

9.14. Actions Cluster

This cluster provides a standardized way for a Node (typically a Bridge, but could be any Node) to expose logical grouping and actions.

Specifically this cluster provides:

- Information about logical grouping of endpoints on the Node (example: lights in a room)
- Information about named actions that can be performed on such a group of endpoints (example: recall a scene for a group of lights by its name)
- Commands to trigger such actions
- Events to receive feedback on the state of such actions.

The information on grouping and available actions is typically provided by the user or Bridge manufacturer via some means not defined in Matter, and therefore provided as read-only to Nodes. For example: a manufacturer-provided app allows a user to set up logical grouping and create/assign scene for such groups.

Using this cluster, a Node can learn about such logical grouping, provided actions, and trigger such actions.

While the origin of this cluster stems from use cases with a Bridge, its server side may also be implemented on any Node which can expose certain grouping, actions or automations to other users.

After defining the attributes, commands and events for this cluster, and the associated data types, several [examples](#) are provided to illustrate the capabilities of this cluster.

Actions can be defined in a flexible manner to suit the needs of the various nodes implementing this cluster. For each action, the commands available for that particular action are defined.

This cluster can be used to expose only the grouping of endpoints without any actions defined by populating the EndpointList attribute accordingly and providing an empty list for ActionList.

The term 'action' in the description of this cluster should not be confused with the term '[action](#)' as used in the Interaction Model.

9.14.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

9.14.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Application	Node	ACT

9.14.3. Cluster ID

ID	Name
0x0025	Actions

9.14.4. Data Types

9.14.4.1. CommandBits Type

This data type is derived from [map16](#).

Bit	Name	Summary
0	InstantAction	Indicate support for InstantAction command
1	InstantActionWithTransition	Indicate support for InstantActionWithTransition command
2	StartAction	Indicate support for StartAction command
3	StartActionWithDuration	Indicate support for StartActionWithDuration command
4	StopAction	Indicate support for StopAction command
5	PauseAction	Indicate support for PauseAction command
6	PauseActionWithDuration	Indicate support for PauseActionWithDuration command
7	ResumeAction	Indicate support for ResumeAction command
8	EnableAction	Indicate support for EnableAction command
9	EnableActionWithDuration	Indicate support for EnableActionWithDuration command
10	DisableAction	Indicate support for DisableAction command
11	DisableActionWithDuration	Indicate support for DisableActionWithDuration command

Note - The bit allocation of this bitmap SHALL follow the ID's of the Commands of this cluster.

9.14.4.2. ActionTypeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Other	Use this only when none of the other values applies	M
1	Scene	Bring the endpoints into a certain state	M
2	Sequence	A sequence of states with a certain time pattern	M
3	Automation	Control an automation (e.g. motion sensor controlling lights)	M
4	Exception	Sequence that will run when something doesn't happen	M
5	Notification	Use the endpoints to send a message to user	M
6	Alarm	Higher priority notification	M

Scene Value

Can be used to set a static state of the associated endpoints (typically using `InstantAction` or `InstantActionWithTransition`), or to bring these endpoints into a more dynamic state (typically using `StartAction`), where the endpoints would e.g. gradually cycle through certain colors for a pleasing effect. A voice controller could use "set" (to map to `InstantAction`) or "play" (to map to `StartAction`) to trigger such actions.

Example: see [Section 9.14.8.1, "Example 1: Scene recall"](#) and [Section 9.14.8.2, "Example 2: Set dynamic light effect"](#).

Sequence Value

Indicates an action which involves a sequence of events/states of the associated endpoints, such as a wake-up experience.

Example: see [Section 9.14.8.4, "Example 4: Wake-up routine"](#).

Automation Value

Indications an automation (e.g. a motion sensor controlling lights, an alarm system) which can be e.g. started, stopped, paused, resumed.

Example: see [Section 9.14.8.3, "Example 3: Pause sensor automation"](#).

Exception Value

Indicates some action which the server will execute when a certain condition (which normally does not happen) is not met.

Example: lock the doors when the server's system has detected no one is at home while the doors are in the 'unlocked' state.

Notification Value

Indicates an action that can be triggered (e.g. by `InstantAction`) to notify the user.

Example: play a pattern on the lights in the living room if there is someone in the garden in the evening.

Alarm Value

Similar to Notification but with a higher priority (and might override other endpoint states which `Type=Notification` would *not* override).

Example: flash all lights in the house when CO sensor triggers.

9.14.4.3. ActionStateEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Inactive	The action is not active	M
1	Active	The action is active	M
2	Paused	The action has been paused	M
3	Disabled	The action has been disabled	M

Note that some of these states are applicable only for certain actions, as determined by their `SupportedCommands`.

9.14.4.4. ActionErrorEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unknown	Other reason not listed in the row(s) below	M
1	Interrupted	The action was interrupted by another command or interaction	M

9.14.4.5. EndpointListTypeEnum Type

This data type is derived from `enum8` and has its values listed below.

Value	Name	Summary	Conformance
0	Other	Another group of endpoints	M
1	Room	User-configured group of endpoints where an endpoint can be in only one room	M
2	Zone	User-configured group of endpoints where an endpoint can be in any number of zones	M

The Room and Zone values are provided for the cases where a user (or the system on behalf of the user) has created logical grouping of the endpoints (e.g. bridged devices) based on location.

Other Value

This value is provided for the case of an endpoint list which is tied specifically to this action i.e. not independently created by the user. For Type=Other the Name MAY be empty. A Matter controller would typically not use this for anything else than just to know which endpoints would be affected by the action.

Room Value

Is used for the situation where an endpoint can only be part of one such rooms (e.g. physical mapping). Using these exposed logical groups, a Matter controller who has a similar grouping concept can use it to place each endpoint (bridged device) in the right room automatically, without user having to redo that setup for each device in each system - both at first contact and upon later updates to the endpoints (e.g. user adds a bridged device or creates a new room).

Zone Value

Is a more general concept where an endpoint can be part of multiple zones, e.g. a light in the living room can be part of the "reading corner" zone (subset of the lights in the living room) but also part of the "downstairs" zone which contains all the lights on a floor, e.g. combining living room, kitchen and hallway. This indicates that a user has defined this list of endpoints as something they logically would like to control as a group, so Matter controllers could provide the user with a way to do as such.

9.14.4.6. ActionStruct Type

This data type holds the details of a single action, and contains the data fields below.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	ActionID	uint16	all				M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	Name	string	max 128{32}				M
2	Type	Action-TypeEnum	all				M
3	EndPointListID	uint16	all				M
4	SupportedCommands	Command-Bits	0 to 0xFFFF				M
5	State	ActionStateEnum	all				M

ActionID Field

This field SHALL provide a unique identifier used to identify an action.

Name Field

This field SHALL indicate the name (as assigned by the user or automatically by the server) associated with this action. This can be used for identifying the action to the user by the client. Example: "my colorful scene".

Type Field

This field SHALL indicate the type of action. The value of Type of an action, along with its SupportedCommands can be used by the client in its UX or logic to determine how to present or use such action. See [ActionTypeEnum](#) for details and examples.

EndPointListID Field

This field SHALL provide a reference to the associated endpoint list, which specifies the endpoints on this Node which will be impacted by this ActionID.

SupportedCommands Field

This field is a bitmap which SHALL be used to indicate which of the cluster's commands are supported for this particular action, with a bit set to 1 for each supported command according to the table below. Other bits SHALL be set to 0.

State Field

This field SHALL indicate the current state of this action.

9.14.4.7. EndpointListStruct Type

This data type holds the details of a single endpoint list, which relates to a set of endpoints that have some logical relation, and contains the data fields below.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	End-pointListID	uint16	all			R	M
1	Name	string	max 128{32}			R	M
2	Type	EndpointListTypeEnum	all			R	M
3	Endpoints	list[endpoint-no]	max 256			R	M

EndPointListID Field

This field SHALL provide an unique identifier used to identify the endpoint list.

Name Field

This field SHALL indicate the name (as assigned by the user or automatically by the server) associated with the set of endpoints in this list. This can be used for identifying the action to the user by the client. Example: "living room".

Type Field

This field SHALL indicate the type of endpoint list, see [EndpointListTypeEnum](#).

Endpoints Field

This field SHALL provide a list of endpoint numbers.

9.14.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	ActionList	list[ActionStruct]	max 256		empty	R V	M
0x0001	EndPointLists	list[EndpointListStruct]	max 256		empty	R V	M
0x0002	SetupURL	string	max 512		empty	R V	O

9.14.5.1. ActionList Attribute

The ActionList attribute holds the list of actions. Each entry SHALL have an unique ActionID, and its EndpointListID SHALL exist in the EndpointLists attribute.

9.14.5.2. EndpointLists Attribute

The EndpointLists attribute holds the list of endpoint lists. Each entry SHALL have an unique EndpointListID.

9.14.5.3. SetupURL Attribute

The SetupURL attribute (when provided) SHALL indicate a URL; its syntax SHALL follow the syntax as specified in [RFC 1738](#), max. 512 ASCII characters and SHALL use the [https](#) scheme. The location referenced by this URL SHALL provide additional information for the actions provided:

- When used without suffix, it SHALL provide information about the various actions which the cluster provides.
 - Example: SetupURL could take the value of `example://Actions` or <https://domain.example/Matter/bridgev1/Actions> for this generic case (access generic info how to use actions provided by this cluster).
- When used with a suffix of `"/?a="` and the decimal value of ActionID for one of the actions, it MAY provide information about that particular action. This could be a deeplink to manufacturer-app/website (associated somehow to the server node) with the information/edit-screen for this action so that the user can view and update details of the action, e.g. edit the scene, or change the wake-up experience time period.
 - Example of SetupURL with suffix added: `example://Actions/?a=12345` or <https://domain.example/Matter/bridgev1/Actions/?a=12345> for linking to specific info/editing of the action with ActionID 0x3039.

9.14.6. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	InstantAction	client ⇒ server	Y	O	desc
0x01	InstantAction- WithTransi- tion	client ⇒ server	Y	O	desc
0x02	StartAction	client ⇒ server	Y	O	desc
0x03	StartAction- WithDuration	client ⇒ server	Y	O	desc
0x04	StopAction	client ⇒ server	Y	O	desc
0x05	PauseAction	client ⇒ server	Y	O	desc
0x06	PauseAction- WithDuration	client ⇒ server	Y	O	desc

ID	Name	Direction	Response	Access	Conformance
0x07	ResumeAction	client ⇒ server	Y	O	desc
0x08	EnableAction	client ⇒ server	Y	O	desc
0x09	EnableAction- WithDuration	client ⇒ server	Y	O	desc
0x0A	DisableAction	client ⇒ server	Y	O	desc
0x0B	DisableAction- WithDuration	client ⇒ server	Y	O	desc

Conformance: a server SHALL support a command when it is listed in the SupportedCommands data field of one or more actions listed in the ActionList provided by the server.

Some general requirements and data fields for all the commands:

- The ActionID data field SHALL indicate the action identifier. If there is no entry in the ActionList holding the same action identifier, a response SHALL be generated with the StatusCode NOT_FOUND.
- The InvokeID data field MAY be provided by the client when invoking a command. When this value is provided, the server SHALL generate a [StateChanged](#) event when the action changes to a new state or an [ActionFailed](#) event when execution of the action fails; in the data of such events, the value of the InvokeID data field from the invoke will be provided, so that the client can relate the event back to the corresponding command. It is up to the client to determine which value is provided in InvokeID. When sending multiple commands for the same action, with different InvokeID, the server SHALL provide in the event the InvokeID value from the most recent command for a particular ActionID.
- If the command refers to an action which currently is not in a state where the command applies, a response SHALL be generated with the StatusCode INVALID_COMMAND.
- Actions are typically mapped to state changes of the involved endpoints. Those endpoints can also be controlled with commands from other clusters, controlled by other nodes and impacted by non-Matter interactions (e.g. local controls). Such other interactions can cause the state of the endpoints to differ from the results of the command which triggered the action. A client interested in such changes can use the InvokeID data field (see above) to receive events StateChanged and ActionFailed for feedback for such cases.

9.14.6.1. InstantAction Command

This command is used to trigger an instantaneous action.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O

This command triggers an action (state change) on the involved endpoints, in a "fire and forget" manner. Afterwards, the action's state SHALL be Inactive.

Example: recall a scene on a number of lights.

9.14.6.2. InstantActionWithTransition Command

This command is used to trigger an instantaneous action with a transition over a given time.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O
2	Transition-Time	uint16	all		MS	M

It is recommended that, where possible (e.g., it is not possible for attributes with Boolean data type), a gradual transition SHOULD take place from the old to the new state over this time period. However, the exact transition is manufacturer dependent.

This command triggers an action (state change) on the involved endpoints, with a specified time to transition from the current state to the new state. During the transition, the action's state SHALL be Active. Afterwards, the action's state SHALL be Inactive.

Example: recall a scene on a number of lights, with a specified transition time.

TransitionTime Field

This field SHALL indicate the transition time in 1/10th of seconds.

9.14.6.3. StartAction Command

This command is used to trigger the commencement of an action.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O

This command triggers the commencement of an action on the involved endpoints. Afterwards, the action's state SHALL be Active.

Example: start a dynamic lighting pattern (such as gradually rotating the colors around the set-points of the scene) on a set of lights.

Example: start a sequence of events such as a wake-up experience involving lights moving through several brightness/color combinations and the window covering gradually opening.

9.14.6.4. StartActionWithDuration Command

This command is used to trigger the commencement of an action with a duration.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O
2	Duration	uint32	all		MS	M

This command triggers the commencement of an action on the involved endpoints, and SHALL change the action's state to Active. After the specified Duration, the action will stop, and the action's state SHALL change to Inactive.

Example: start a dynamic lighting pattern (such as gradually rotating the colors around the set-points of the scene) on a set of lights for 1 hour (Duration=3600).

Duration Field

This field SHALL indicate the requested duration in seconds.

9.14.6.5. StopAction Command

This command is used to stop an action.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O

This command stops the ongoing action on the involved endpoints. Afterwards, the action's state SHALL be Inactive.

Example: stop a dynamic lighting pattern which was previously started with StartAction.

9.14.6.6. PauseAction Command

This command is used to pause an action.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O

This command pauses an ongoing action, and SHALL change the action's state to Paused.

Example: pause a dynamic lighting effect (the lights stay at their current color) which was previously started with StartAction.

9.14.6.7. PauseActionWithDuration Command

This command is used to pause an action with a duration.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O
2	Duration	uint32	all		MS	M

This command pauses an ongoing action, and SHALL change the action's state to Paused. After the specified Duration, the ongoing action will be automatically resumed, which SHALL change the action's state to Active.

Example: pause a dynamic lighting effect (the lights stay at their current color) for 10 minutes (Duration=600).

The difference between Pause/Resume and Disable/Enable is on the one hand semantic (the former is more of a transitional nature while the latter is more permanent) and on the other hand these can be implemented slightly differently in the implementation of the action (e.g. a Pause would be automatically resumed after some hours or during a nightly reset, while an Disable would remain in effect until explicitly enabled again).

Duration Field

This field SHALL indicate the requested duration in seconds.

9.14.6.8. ResumeAction Command

This command is used to resume an action.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O

This command resumes a previously paused action, and SHALL change the action's state to Active.

The difference between ResumeAction and StartAction is that ResumeAction will continue the action from the state where it was paused, while StartAction will start the action from the beginning.

Example: resume a dynamic lighting effect (the lights' colors will change gradually, continuing from the point they were paused).

9.14.6.9. EnableAction Command

This command is used to enable an action.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O

This command enables a certain action or automation. Afterwards, the action's state SHALL be Active.

Example: enable a motion sensor to control the lights in an area.

9.14.6.10. EnableActionWithDuration Command

This command is used to enable an action with a duration.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O
2	Duration	uint32	all		MS	M

This command enables a certain action or automation, and SHALL change the action's state to be Active. After the specified Duration, the action or automation will stop, and the action's state SHALL change to Disabled.

Example: enable a "presence mimicking" behavior for the lights in your home during a vacation; the Duration field is used to indicated the length of your absence from home. After that period, the presence mimicking behavior will no longer control these lights.

Duration Field

This field SHALL indicate the requested duration in seconds.

9.14.6.11. DisableAction Command

This command is used to disable an action.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O

This command disables a certain action or automation, and SHALL change the action's state to Inactive.

Example: disable a motion sensor to no longer control the lights in an area.

9.14.6.12. DisableActionWithDuration Command

This command is used to disable an action with a duration.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16	all			M
1	InvokeID	uint32	all			O
2	Duration	uint32	all		MS	M

This command disables a certain action or automation, and SHALL change the action's state to Disabled. After the specified Duration, the action or automation will re-start, and the action's state SHALL change to either Inactive or Active, depending on the actions (see [Section 9.14.8.4, "Example 4: Wake-up routine"](#) and [Section 9.14.8.6, "Example 6: Alarm system"](#)).

Example: disable a "wakeup" experience for a period of 1 week when going on holiday (to prevent them from turning on in the morning while you're not at home). After this period, the wakeup experience will control the lights as before.

Duration Field

This field SHALL indicate the requested duration in seconds.

9.14.7. Events

ID	Name	Priority	Access	Conformance
0x00	StateChanged	INFO	V	M
0x01	ActionFailed	INFO	V	M

9.14.7.1. StateChanged Event

This event SHALL be generated when there is a change in the State of an ActionID during the execution of an action *and* the most recent command using that ActionID used an InvokeID data field.

It provides feedback to the client about the progress of the action.

Example: When InstantActionWithTransition is invoked (with an InvokeID data field), two StateChanged events will be generated:

- one when the transition starts (NewState=Active)
- one when the transition completed (NewState=Inactive)

This event SHALL have the following data fields:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16				M

ID	Name	Type	Constraint	Quality	Fallback	Conformance
1	InvokeID	uint32				M
2	NewState	ActionStateEnum				M

ActionID Field

This field SHALL be set to the ActionID of the action which has changed state.

InvokeID Field

This field SHALL be set to the InvokeID which was provided to the most recent command referencing this ActionID.

NewState Field

This field SHALL be set to state that the action has changed to.

9.14.7.2. ActionFailed Event

This event SHALL be generated when there is some error which prevents the action from its normal planned execution *and* the most recent command using that ActionID used an InvokeID data field.

It provides feedback to the client about the non-successful progress of the action.

Example: When InstantActionWithTransition is invoked (with an InvokeID data field), and another controller changes the state of one or more of the involved endpoints during the transition, thus interrupting the transition triggered by the action, two events would be generated:

- StateChanged when the transition starts (NewState=Active)
- ActionFailed when the interrupting command occurs (NewState=Inactive, Error=interrupted)

Example: When InstantActionWithTransition is invoked (with an InvokeID data field = 1), and the same client invokes an InstantAction with (the same or another ActionId and) InvokeID = 2, and this second command interrupts the transition triggered by the first command, these events would be generated:

- StateChanged (InvokeID=1, NewState=Active) when the transition starts
- ActionFailed (InvokeID=2, NewState=Inactive, Error=interrupted) when the second command interrupts the transition
- StateChanged (InvokeID=2, NewState=Inactive) upon the execution of the action for the second command

This event SHALL have the following data fields:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ActionID	uint16				M
1	InvokeID	uint32				M
2	NewState	ActionStateEnum				M
3	Error	Action-ErrorEnum				M

ActionID Field

This field SHALL be set to the ActionID of the action which encountered an error.

InvokeID Field

This field SHALL be set to the InvokeID which was provided to the most recent command referencing this ActionID.

NewState Field

This field SHALL be set to state that the action is in at the time of generating the event.

Error Field

This field SHALL be set to indicate the reason for non-successful progress of the action.

9.14.8. Examples

This section provides some examples how the attributes and commands in this cluster can be used in practical situations. Details of the behavior typically depend on the details of the logic built into the server.

9.14.8.1. Example 1: Scene recall

User has defined a scene on a number of lights. The corresponding action would have these data fields describing it:

- Name="sunset"
- Type=Scene
- EndpointListID references a struct referencing the set of involved endpoints
 - Name="living room"
 - Type=Room
 - Endpoints=list of the endpoints of the lights in this room
- SupportedCommands: InstantAction, InstantActionWithTransition

The InstantAction command (e.g. triggered by a voice command "set sunset in living room") will

trigger the server to activate that scene on those lights.

When a slow fade-in is preferred, the `InstantActionWithTransition` can be used, with a `Transition-Time` parameter of e.g. 50 (denoting 5 s).




Figure 73. State diagram for example 'scene recall'

9.14.8.2. Example 2: Set dynamic light effect

User has defined a scene on a number of lights. The corresponding action would have these data fields describing it:

- Name="sunset"
- Type=Scene
- EndpointListID references a struct referencing the set of involved endpoints (same as in Example 1)
- SupportedCommands: StartAction, StartActionWithDuration, StopAction

The `StartActionWithDuration` command (e.g. triggered by a voice command "play sunset in living room for 1 hour") will trigger the server to activate a dynamic pattern with colors inspired by sunset on the associated lights. At any time, the `StopAction` can be used to stop the effect.

Please note that the most of the data fields in the `ActionStruct` for this example are identical to those in example 1 - except for the `SupportedCommands`. The different sets of supported commands indicate whether it is an instant scene recall (example 1) vs. a long-term dynamic effort (example 2). A server could expose this action also as a single action with the combined set of supported commands.




Figure 74. State diagram for example 'dynamic light effect'

9.14.8.3. Example 3: Pause sensor automation

User has defined an automation: a motion sensor controls the lights in a certain room. Sometimes, they want to override that automatic behavior, e.g. when having a party.

The action for this example would refer to such automation, which is typically active, but can be paused (=temporarily disabled).

The corresponding action would have these data fields describing it:

- Name="motion sensor"
- Type=Automation
- EndpointListID references a struct referencing the set of involved endpoints (same as in Example 1)
- SupportedCommands: EnableAction, DisableAction, PauseAction, PauseActionWithDuration, ResumeAction

Typically, the action has been started when the user defines the motion sensor behavior, so without a Matter command the action's state would be 'Active'. The `PauseActionWithDuration` command (e.g. triggered by a voice command "disable the motion sensor in living room for 2 hours") will trigger the server to disable the behavior associated with the motion sensor for the specified time. A `ResumeAction` command would make this behavior active again. The automation could also have internal logic to abort the disabling after several hours or during the night-time reset, to accommodate for the case the user 'paused' and forgot to 'resume'.




Figure 75. State diagram for example 'pause sensor automation'

9.14.8.4. Example 4: Wake-up routine

User has defined a wake-up routine: the lights in the bedroom will gradually become brighter and change in color temperature to simulate a sunrise. The sequence lasts for e.g. 30 minutes. Near the end of the sequence, the window coverings would be (partially) opened as well.

The corresponding action would have these data fields describing it:

- Name="wakeup"
- Type=Sequence
- EndpointListID references a struct referencing the set of involved endpoints (lights and window coverings in the bedroom)
- SupportedCommands: EnableAction, DisableAction, DisableActionWithDuration, PauseAction, PauseActionWithDuration

When the user wants to snooze some more, he can use a voice command to trigger the PauseAction command (which could automatically timeout after e.g. 10 minutes). The StopAction command could similarly be used to cancel the remainder of the whole sequence. The DisableActionWithDuration (with parameter 172,800 = 2*24*60*60 s) can be used on Friday evening to disable the sequence for the weekend.

When such action has been defined, a Matter node which is aware of the user's calendar for the day, can use the StartAction command to trigger this sequence of events.




Figure 76. State diagram for example 'wake-up experience'

9.14.8.5. Example 5: Scheduled scene recall

User has setup a scene to be recalled at a certain time of day (e.g. colorful garden lighting to switch on around sunset) and switching off those lights (e.g. at midnight). The server's automation system takes care of this. On certain occasions (e.g. garden party), the user wants to override this behavior (i.e. the scene should not be recalled at sunset because another scene has been set for the party).

The corresponding action would have these data fields describing it:

- Name="colorful evening garden"
- Type=Automation
- EndpointListID references a struct referencing the set of involved endpoints (lights in the garden)
- SupportedCommands: EnableAction, DisableAction, DisableActionWithDuration

After installation, this action is in Inactive state. At the scheduled "on" time, the colorful garden lighting scene is activated and the action's state changes to the Active state. At the scheduled "off" time, the lights are switched off, and the action's state changes to the Inactive state. Using the DisableAction, the user can prevent these automated steps to occur - the action's state changes to the Disabled state. Using the DisableActionWithDuration, the user can do similarly, but also indicate an automatic re-enabling after the specified time period. Using the EnableAction, the user can re-enable the automation.




Figure 77. State diagram for example 'scheduled scene recall'

9.14.8.6. Example 6: Alarm system

User has an alarm system which exposes this cluster, with an action that allows to arm/disarm the system by voice commands from a Matter node which is a client to this cluster.

The corresponding action would have these data fields describing it:

- Name="alarm system"
- Type=Automation
- EndpointListID references a struct referencing the set of involved endpoints (elements of the alarm system)
- SupportedCommands: EnableAction, DisableAction, DisableActionWithDuration

After installation, this action could be in Inactive state (assume user is at home installing so system is not armed). Using the EnableAction, the alarm system would be armed. Using the DisableAction, the alarm system would be disarmed (or disarmed for a period with DisableWithDuration).




Figure 78. State diagram for example 'alarm system'

9.15. Intermittently Connected Devices Behavior

9.15.1. ICD Server Behavior

9.15.1.1. Operational States

Idle Mode

Idle Mode defines the operational state of an ICD where the node is unreachable due to it not being able to receive messages for a certain period of time. The ICD will remain in Idle Mode for a maximum of one [Idle Mode Duration](#) as defined in the [ICD Management cluster](#).

Slow Polling

When in **Idle Mode**, the ICD SHALL configure its network interface in *Slow Polling*. *Slow Polling* defines the fastest frequency at which the device will typically receive messages in Idle Mode. The [Slow Polling](#) interval MAY be the same as the [Idle Mode Duration](#).

Active Mode

Active Mode defines the operational state of an ICD where it is reachable on the network and will answer in a timely manner. The minimum amount of time a device will typically remain in Active Mode is [Active Mode Duration](#) as defined in the [ICD Management cluster](#).

Fast Polling

When in **Active Mode**, the ICD SHALL configure its network interface in *Fast Polling*. *Fast Polling* defines the fastest frequency at which the device can receive messages in Active Mode.

Operational State Switch

Some of the reasons why an ICD can switch Operational states from Idle Mode to Active Mode are:

- The ICD has remained in Idle Mode for a full [Idle Mode Duration](#).
- An attribute change needs to be reported to a subscriber.
- When the maximum reporting interval has been reached for an active subscription.
- A generated event needs to be reported to a subscriber.

When an ICD switches from Idle Mode to Active Mode, it MAY send all pending messages. These messages could be:

- Sending subscription reports and events.
- Sending Check-In messages.
- Sending Interaction Messages.

A node typically determines whether it is in Active or Idle mode based on whether it has any outstanding [Exchanges](#) in the Message Layer. While there are Exchanges active, a node typically will remain in Active mode. Once all Exchanges are closed and the node has no other outstanding reasons for staying active, a node SHOULD transition into Idle mode. An ICD node MAY use additional modes and logic for switching between them.

9.15.1.2. ICD Session Configurations

Due to the ICD operating modes, there may be an additional delay before the ICD receives a message. This causes the travel time of a message to be affected by the ICD's configuration and current operating mode.

When in [Idle Mode](#), the device will typically not receive messages at a faster frequency than the [Slow Polling Interval](#). As such, a client SHOULD NOT try to send the same message at a frequency higher than the [Slow Polling Interval](#) when the ICD is in Idle Mode. The ICD SHALL set its [SESSION_IDLE_INTERVAL](#) to a value greater than or equal to the [Slow Polling Interval](#).

When in [Active Mode](#), the device will typically not receive messages at a faster frequency than the [Fast Polling Interval](#). As such, a client SHOULD NOT try to send the same message at a frequency higher than the [Fast Polling Interval](#) when the ICD is in Active Mode. The ICD SHALL set its [SESSION_ACTIVE_INTERVAL](#) to a value greater than or equal to the [Fast Polling Interval](#).

The ICD SHALL set its [SESSION_ACTIVE_THRESHOLD](#) to the [Active Mode Threshold](#) stored in the [ICD Management cluster](#).

9.15.1.3. Check-In Protocol Support

This section describes the Check-In Protocol use case used by Intermittently Connected Devices (ICDs) to maintain a known relationship in case subscriptions with clients are lost. This includes how a client shares an [ICDToken](#) with the ICD, when Check-In messages are sent and how the Check-In Protocol requirements are respected.

The Check-In Protocol is a fail-safe mechanism which allows an Intermittently Connected Device (ICD) to notify a registered client that it is available for communication when all subscriptions between the client and ICD are lost. A subscription can be lost for several reasons, such as:

- The ICD might not have full RAM retention when it is in an idle state.
- When the ICD is powered off to change the battery.
- Power or network outage causing the connection between the client and the ICD to be interrupted.
- The client is unavailable for any reason (e.g. during a software update or hosted on a mobile device that is sometimes out-of-home).

The Check-In message is sessionless and relies on a shared secret that has been given to the ICD during the registration of the client using the [ICD Management cluster](#).

Use Case Requirements

The ICD Check-In Protocol representation of the [Check-In Counter](#) SHALL be the [ICD Counter](#). It SHALL follow all the requirements of the [Check-In Counter](#).

The ICD Check-In Protocol representation of the Check-In Protocol [persistent data store](#) SHALL be the [RegisteredClients](#) attribute in the [ICD Management cluster](#). The attribute SHALL store all the [client registration information](#).

The [Application Data](#) for the ICD Check-In Protocol use case SHALL be the [Active Mode Threshold](#) attribute in the [ICD Management cluster](#). Before being used in the [encryption process](#), the [Active Mode Threshold](#) SHALL be converted into a 2-byte little-endian value.

Protocol State Machine

The state machine diagram represents the states an ICD can be in. The diagram represents the state of the ICD for one client or [MonitoredSubject](#). The ICD can be in a different state for each client simultaneously.




Figure 79. ICD State Machine

Uncommissioned State

An ICD is in the *Uncommissioned* state when it has not been provided with any credentials to join a fabric. In this state, the ICD is passively waiting for incoming messages during the [open commissioning window](#).

Unregistered State

An ICD is in the *Unregistered* state with a given [MonitoredSubject](#) when no established relationship with the client exists, meaning the ICD has neither a registration entry in the [ICD Management cluster](#) nor an active subscription associated with the client. In this state, the ICD is passively waiting for incoming messages during its [Active Mode Duration](#).

Check-in State

An ICD is in the *Check-in* state when a registration entry is present for the client in the [ICD Management cluster](#), but no active subscription exists for that client. In this state, the ICD sends Check-In messages instead of subscription reports/events. The Check-In message notifies the client that the ICD is available for communication.

When in the Check-in state, each time an ICD transitions from [Idle Mode](#) to [Active Mode](#), the ICD SHALL attempt to send a Check-In message to each registered ([MonitoredSubject](#)) client that are [resident](#) clients without an active subscription, effectively attempting to send a Check-In message every [Idle Mode Duration](#). The ICD SHOULD NOT attempt to send Check-In messages to clients that have been registered as [ephemeral](#) clients.

The ICD MAY back-off (i.e. skip some of these Check-In messages) when it does not experience interactions from the registered client (e.g. if that client has left the home network) after several consecutive Check-In messages. The time between Check-In messages SHALL NOT be larger than the duration (in seconds) specified in [MaximumCheckInBackoff](#). The number of consecutive missed Check-In messages required to start backing-off is manufacturer specific. The missed Check-In message counters SHALL be reset to 0 for all registered clients when a [user active mode trigger](#) is executed.

NOTE

A commissioned and registered ICD that loses connectivity or goes offline (e.g., due to device mobility or loss of power) SHALL NOT transition to an uncommissioned or unregistered state solely because it failed to send periodic check-in messages for longer than [MaximumCheckInBackoff](#). Any such state change requires explicit user action.

Subscribed State

An ICD is in the *Subscribed* state with a given [MonitoredSubject](#) when it has at least one active subscription with the client. It is not necessary to have an entry in the [ICD Management cluster](#) to be in the subscribed state. The ICD MAY be in the subscribed state most of its operating time. In this state, the ICD's goal is to notify the client when it is available for communication with subscription reports.

Message Format

All Check-In messages SHALL be structured as specified in [Message Format](#).

All Check-In messages are unsecured at the message layer:

- The [Session ID](#) field SHALL be set to **0**.
- The [Session Type](#) bits of the [Security Flags](#) SHALL be set to **0**.
- The [S Flag](#) field of the [Message Flags](#) SHALL be set to **0**.
- The [Message Counter](#) value SHALL be the current value of the [global unencrypted message counter](#).
- The [DSIZ](#) field of the [Message Flags](#) SHALL be set to **0**.

Protocol Header Field

The Check-In message will have the following values for the protocol header field :

- The [Initiator flag](#) field SHALL be set to **1**.
- The [Reliability flag](#) field SHALL be set to **0**.
- The [SX flag](#) field SHALL be set to **0**.
- The [Vendor Flag](#) field SHALL be set to **0**.

Protocol Id and Opcode

The ICD Check-In message SHALL have [secure channel](#) protocol ID. The ICD Check-In message SHALL have the [ICD Check-In message](#) protocol opcode.

Sending Check-In Message Steps

When sending an ICD Check-In message, the server SHALL respect the following requirements:

1. The server SHALL NOT send more than one Check-In message with a given [ICD counter value](#) to the same [RegisteredClient CheckInNodeID](#).
2. The server SHALL follow the [encryption process](#) of the Check-In Protocol.
3. The server SHALL configure the Matter header as defined in [message format](#) section.

The server MAY try to resolve the IP address and the port of the [CheckInNodeID](#) each time the server needs to send a Check-In message to a [RegisteredClient CheckInNodeID](#).

9.15.1.4. User Active Mode Trigger

Since ICDs are not immediately responsive, they require a means to render them available for communication for user initiated use cases. Some of the user initiated use cases are:

- Opening a new commissioning window to add another administrator.
- Reconfiguration of an existing fabric (e.g. IPKs, NOC rotation, ACL changes).
- Reconfiguration of cluster functionality (e.g. ICD Management, Bindings, Groups, Scenes).
- Removal of a device from a fabric.
- Changes to the device's settings.

An ICD that requires user initiated use cases SHALL provide an Active Mode Trigger as a way for a user to put the device into active mode and make it responsive to controller communication. The [UserActiveModeTrigger](#) feature in the [ICD Management cluster](#) indicates whether a particular device implements an active mode trigger.

9.15.1.5. Short Idle Time (SIT) ICD

ICDs with [Slow Polling Interval](#) shorter than or equal to 15 seconds SHOULD be configured as a Short Idle Time ICD. For example, this is a typical scenario for door locks and window coverings, where commands need to be sent to the ICD with a use case imposed latency requirement. Typically, devices that are Short Idle Time ICDs are not initiators in the communication flow.

Requirements

A SIT ICD SHALL advertise its [SESSION_IDLE_INTERVAL](#) per [SII](#). It SHALL also advertise its [SESSION_ACTIVE_INTERVAL](#) per [SAI](#). The device MAY have a [user action](#) that triggers the ICD to transition to Active mode.

Having a long [ActiveModeThreshold](#) makes the ICD more responsive for subsequent messages at the expense of increased power consumption. The ICD stays in [Active Mode](#) for the entire duration of the [ActiveModeThreshold](#).

On the other hand, having an [ActiveModeThreshold](#) of 0 means that the device might enter [Idle Mode](#) immediately after the last active [exchange](#) is closed. One good example is if a client reads five attributes in a row from an ICD and the [ActiveModeThreshold](#) is 0 then it MAY have to wait up to one [Idle Mode Duration](#) after each attribute read.

In order to get the benefits of the [ActiveModeThreshold](#) attributes while keeping the power consumption low, the ICD MAY ignore the [Active Mode Threshold](#) and immediately return to [IdleMode](#) upon the completion of an exchange if the ICD was the Initiator of the exchange being closed and there are no other reasons keeping it in [active mode](#).

9.15.1.6. Long Idle Time (LIT) ICD

ICDs with a [Slow Polling Interval](#) longer than 15 seconds SHOULD be configured as Long Idle Time ICDs. For example, this is a typical scenario for sensors and light switches, where data reports are initiated by the ICD for an updated sensed value or a switch event. Typically, devices that are Long Idle Time ICDs are initiators in the communication flow.

Overview

The high-level steady state behavior of a LIT ICD is to alternate between long IDLE mode periods and short ACTIVE mode periods, where the ACTIVE mode period begins with a transmission to registered clients. A graphical depiction of ICD steady state behavior is shown in [Figure 80, “ICD Steady State Behavior Diagram”](#).




Figure 80. ICD Steady State Behavior Diagram

The ICD stays active a minimum amount of time each active mode period, and that time can be extended based on received network activity from registered clients. A graphical depiction of ICD behavior when in ACTIVE mode is shown in Figure 81, “ICD Active Mode Behavior Diagram”. The **ActiveModeDuration** starts immediately upon transitioning out of IDLE mode into ACTIVE mode. The **ActiveModeThreshold** starts immediately upon sending the Check-In message or any other transmission. The Exchange Context ends upon receipt of a response acknowledgement or timeout. Receipt of a response acknowledgement starts a new **ActiveModeThreshold** and keeps the ICD ACTIVE for longer. If further messages are received during the **ActiveModeThreshold**, a new Exchange can be created to keep the ICD ACTIVE for even longer.




Figure 81. ICD Active Mode Behavior Diagram

Requirements

The device SHALL indicate it is a LIT ICD by setting the **LongIdleTime** feature to 1 in the **ICD Management cluster**. When the device is operating as a LIT ICD, it SHALL advertise 1 with the **ICD discovery TXT key**. This informs clients that they SHOULD wait for a notification from the LIT ICD before sending any messages to the device. When the device is operating as a SIT ICD, it SHALL advertise 0 with the **ICD discovery TXT key**. This informs clients that they SHOULD NOT wait for a notification from the LIT ICD before sending any messages to the device. It SHALL also advertise its **SESSION_ACTIVE_INTERVAL** per **SAI**. A LIT ICD SHOULD NOT advertise its **SESSION_IDLE_INTERVAL** with **SII**. The device SHALL support the **ICD Check-In use case** requirements. The device SHALL have a **user action** that triggers the ICD to transition to Active mode. The **Fast Polling** interval SHALL be smaller than the **Active Mode Duration**. The **ActiveModeThreshold** SHALL NOT be smaller than 5 seconds.

Client Registration Requirement

For LIT ICDs to be fully functional, they require clients to support and use the ICD client behavior.

The main client requirements are the client registration process with the [ICD Management cluster](#), and the processing of incoming Check-In messages received from the ICD.

A LIT ICD MAY operate as a LIT ICD when at least one client has successfully registered.

A commissioned LIT ICD SHALL operate as a [SIT ICD](#) if it doesn't have at least one registration with any client on any fabric in the [ICD Management cluster](#). In case this situation prolongs for an extended period of time, *P1* (ICD_UNREGISTERED_THRESHOLD_PERIOD), without communication with clients, a manufacturer MAY implement mechanisms to preserve battery life e.g. take the device offline for some period of time, *P2* (ICD_UNREGISTERED_OFFLINE_PERIOD), before returning to an active period of SIT operation.

- Such mechanism can be selected to balance user convenience vs. battery life, and likely will depend on the device type, functionality and use case.
- When a [User Active Mode Trigger](#) occurs, the device SHALL return to operate as a [SIT ICD](#), and restart tracking the *P1* period.
- Likewise any asynchronous operational action on the device such as a user pressing a switch button or a sensor sampling relevant data SHALL return the device to a SIT mode of operation.
- The period *P2* MAY be adaptive (e.g. gradually increasing if subsequent attempts fail) but SHALL be no more than [Idle Mode Duration](#).
- Detection of such extended period *P1* is manufacturer specific, and MAY differentiate between these failure cases:
 - network operational but no communication with clients
 - network lost (e.g. Thread/Wi-Fi sensor device when the mains-powered Thread devices/Wi-Fi access point have been switched off)
 - In this situation, the LIT ICD cannot communicate with the network

When operating as a [SIT ICD](#), a LIT ICD:

1. SHALL NOT set its Slow Polling Interval to a value higher than the [Short Idle Time ICD](#) maximum.
2. SHALL advertise its [SESSION_IDLE_INTERVAL](#) using the [SII discovery TXT key](#).
3. SHALL advertise 0 with the [ICD discovery TXT key](#) to inform clients that it can be reached asynchronously.
4. MAY ignore the [Active Mode Threshold](#) as described in the [SIT ICD Requirements](#) section.

After a client has successfully registered with [ICD Management cluster](#), a LIT ICD MAY operate as one. When an ICD transitions from operating as a SIT to a LIT ICD, the LIT ICD SHALL respect all the LIT requirements.

Runtime Operating Mode Switching

Due to the registration requirement, LIT ICDs SHALL have the capacity to switch between LIT and SIT operating modes. This ability can be leveraged by devices that have a use case where the device MAY want to remain in the SIT operating mode even if a client has registered. A good example is a Smoke & CO alarm device. While the device is line-powered, it can be very responsive to clients and

remain in the SIT operating mode. If it loses line-power, it can transition to LIT operating mode if it has at least one registration and leverage all the LIT ICD features and power efficiency. If the device supports switching between the SIT and LIT operation modes even with a registered client, it SHALL set the [DynamicSitLitSupport](#) feature in the [ICD Management cluster](#).

Each time a LIT ICD switches between operating modes, both the [mDNS ICD key](#) and [OperatingMode](#) attribute of the ICDM cluster must be updated:

- [ICD=0](#) and [OperatingMode=SIT](#) if "Long Idle Time ICD is operating as a Short Idle Time ICD",
- [ICD=1](#) and [OperatingMode=LIT](#) if "Long Idle Time ICD is operating as a Long Idle Time ICD".

Runtime Polling Interval Changes

In addition to being able to dynamically switch between LIT and SIT operating modes, an ICD can also dynamically change its [Fast Polling Interval](#), [Slow Polling Interval](#) and its [Active Mode Threshold](#). A change in [Active Mode Threshold](#) value SHALL trigger a DNS-SD update of the [SAT](#) TXT key. Changes of the [SESSION_ACTIVE_INTERVAL](#) or [SESSION_IDLE_INTERVAL](#) SHALL also trigger a DNS-SD update of the corresponding [SAI](#) or [SII](#) TXT keys. Note that an update of the polling intervals MAY NOT trigger an update of the [SAI](#) or [SII](#) TXT keys if the new polling intervals are still compliant with the rules from [ICD Session Configurations](#) section.

9.15.1.7. Long Idle Time (LIT) ICD Implementation settings (informative)

From LIT-measurements in a realistic environment, the following minimum settings are advised:

- [ActiveModeDuration](#) (SDK [CHIP_CONFIG_ICD_ACTIVE_MODE_DURATION_MS](#)): 0, i.e. it would use Active Mode Threshold as ICD Active Mode Duration.
- [ActiveModeThreshold](#) (SDK [CHIP_CONFIG_ICD_ACTIVE_MODE_THRESHOLD_MS](#)): 25000-30000. Such range (25-30 seconds) allows time for approximately five MRP retries in case of connectivity issues in a congested network (In some situations, queuing in the Thread layer can occur, which might result in significant delays to Matter traffic)
- [IdleModeDuration](#) (SDK [CHIP_CONFIG_ICD_IDLE_MODE_DURATION_SEC](#)): 3600 seconds (1 hour) or some larger value, but no more than the maximum 68400 seconds (19 hours)
- [SAI \(Session Active Interval\)](#): 3 seconds (e.g., set SDK [CHIP_CONFIG_MRP_LOCAL_ACTIVE_RETRY_INTERVAL](#) to 2 seconds and fast polling to 1 second). Since LIT ICDs do not advertise [SII](#), a larger SAI will improve MRP stability.

When a client does not support Long Idle Time (LIT), specific Short Idle Time (SIT) fallback configurations are required. In the current SDK, the maximum subscription interval defaults to the SDK's idle mode duration. To optimize connectivity, please consider the following configuration:

- Short Idle Mode Duration: Set [CHIP_CONFIG_ICD_SHORT_IDLE_MODE_DURATION_SEC](#) to a lower value (e.g., 5 minutes).
- SIT Polling Interval: Set [CHIP_DEVICE_CONFIG_ICD_SIT_POLLING_INTERVAL](#) to less than 15 seconds (e.g., : 8 seconds) to increase the likelihood of connection during idle mode.

9.15.2. ICD Client Behavior

This section provides a description of the client requirements to support [Long Idle Time ICDs](#). A client does not require a specific feature-set to support [Short Idle Time ICDs](#).

9.15.2.1. Check-In Protocol Support

A client that supports [Long Idle Time ICDs](#) SHALL implement the Check-In Protocol client [requirements](#).

9.15.2.2. Registration Processes

In order to maximize the reliability of communications with [Long Idle Time ICDs](#), which can be unreachable for extended periods of time, the [ICD Management cluster](#) provides a mechanism for clients to register themselves or other monitoring nodes to receive Check-In messages from an ICD. A client that needs to be able to interact with a [Long Idle Time ICDs](#) SHALL invoke the [Register-Client](#) command against it, typically during commissioning, or have the commissioner or some other existing client send the RegisterClient command on its behalf.

The [RegisterClient](#) command supplies the following parameters to the server:

1. An [ICDToken](#) (Key) which is used to encrypt the Check-In messages sent by the server.
2. A node identifier which specifies the node that will receive the Check-In messages.
3. A subject-id used to determine if a particular client has an active subscription.
4. The type of the client registering, permanent or ephemeral (e.g. mobile).

When a client registers with multiple servers, it SHALL use unique keys since the key is used to identify which server sent the Check-In message. Check-In messages do not have an originator address. Therefore, successfully decrypting a Check-In message with a key and having that key correspond to a specific server is the only way to identify which server sent the Check-In message.

If the commissioner is ephemeral and a separate stationary controller exists in the network, the commissioner SHOULD register the stationary controller in the [ICD Management cluster](#). Because an ICD is only required to provide one guaranteed client slot per fabric, clients SHOULD self-limit commissioning time registration.

Commissioner Self-Registration

A commissioner that is also a controller can register itself with the ICD during the device configuration portion of the commissioning process. This is the simplest client registration scenario, where a commissioner registers itself during commissioning. At the end of the commissioning process, the commissioner and ICD are [coupled](#).

This diagram shows one possible process a commissioner can use to register itself with an ICD. Other processes are also possible such as a Commissioner which registers itself right after the [Add NOC](#) step and before the [Operational Discovery](#) step.




Figure 82. Commissioner Self-Registration to ICD

External Controller Registration

A commissioner can also register an external controller at the time of commissioning. This requires that the commissioner knows the node ID of the controller and that commissioner and controller share an [ICDToken](#) used for the client registration by some out-of-band means.

This diagram shows one possible process a commissioner can use to register the controller with an ICD. Other processes are also possible such as a Commissioner which registers the controller right after the [Add NOC](#) step and before the [Operational Discovery](#) step. The same is true for when the [ICD-Token](#) is shared between devices. Other processes may choose to share the key at a different moment in the process.




Figure 83. External Controller Registration

Delegated Controller Registration

A commissioner can also delegate client registration to some other assisting device or controller. In such a scenario, the commissioner does not register the controller, but instead completes the commissioning process and then hands control to the delegated controller to register itself or other clients with the ICD.

When doing so, a commissioner would need to send a command to keep the ICD awake for some amount of time, in order to ensure that an external controller can reach it.

Alternatively, the infrastructure could also be notified of the newly-commissioned ICD before invoking `CommissioningComplete`, as this would keep the device awake and available for interactions.

This diagram shows one possible process a commissioner can use to register the controller with an ICD.




Figure 84. Delegated Controller Registration

Assisted On Demand Registration

If a new controller is added to a fabric with existing ICDs, it may be necessary that some existing controller in the infrastructure assist with client registration by providing notification to a new controller when an ICD is active and thus available.

In this scenario, an out-of-band method can be used for the existing controller to detect or be informed of the addition of the new controller and signal to it when present ICDs are active. It should be noted that it is also feasible for a node identity to be instantiated on different hardware to suit the needs of the user. This could allow controllers to be replaced in a system, or have controller identities roam to different hardware without the need to re-register in the client.

This diagram shows one possible process a commissioner can use to register the controller with an ICD.




Figure 85. Infrastructure Assisted On-Demand Registration

Assisted Proxy Registration

Alternatively, instead of the new controller receiving notification when an ICD wakes up, an existing controller could be requested to register the new controller when an ICD checks-in. In this way, the existing controller acts as a proxy to register on behalf of the new controller. This requires a queuing or caching mechanism in the existing controller.




Figure 86. Infrastructure Assisted Proxy Registration

Client Registration upon ICD Upgrade

When a device upgrades to become ICD capable per [Section 9.15.2.5.2, “Asynchronous Interactions”](#), a client SHOULD register with the device using procedure described in [Section 9.15.2.2, “Registration Processes”](#).

- After detecting the LIT capability, the client SHOULD interact with the ICD and initiate the registration process within the ICD Management cluster. This enables the client to receive Check-In messages and other notifications from the LIT ICD.
- An ephemeral client SHOULD NOT register itself in the ICD Management cluster. If it does, it SHOULD register as an **ephemeral** client in the ICD Management client. A stationary controller in the network SHOULD be registered instead.




Figure 87. In field upgrade of ICD functionality

9.15.2.3. Unregistration Process

When a commissioner or client wants to unregister itself, it can send an **UnregisterClient** command to the ICD. When a registered client is about to permanently leave the network, it SHOULD send an **UnregisterClient** command to free its slot for another client.

When an ICD receives the **RemoveFabric** command, all registrations associated with the fabric SHALL be removed during the fabric removal process.

9.15.2.4. Post-Commissioning Process

In addition to registering a client with the ICD, the commissioner or client SHOULD ensure that some node that matches the registered **subject-id** establishes and maintains a subscription with the ICD. A client is not required to establish and maintain a subscription with the ICD. In such case, it can rely on Check-In messages as a heartbeat to provide an indication when it can communicate with an ICD.

9.15.2.5. Client and ICD Relationship

This section describes how a client and a LIT ICD arrange to communicate with each other. Because **SIT ICDs** are reachable on the network without special action by clients, clients need not register with SIT ICDs.

Communication with an ICD

LIT ICDs have two operating modes. They can either be operating as LIT or as **SIT ICD**. If the ICD is operating as a **SIT ICD**, it SHALL advertise this by setting the **ICD discovery TXT key** to 0. When the **LIT ICD** is operating as a **SIT ICD**, clients can send their message to the ICD without prior notification. If the ICD is operating as a **LIT ICD**, it SHALL advertise this by setting the **ICD discovery TXT key** to 1. If the ICD is operating as a **LIT ICD**, it will be unreachable for extended periods of time. **Permanent** clients need to wait for prior notification before sending a message to the ICD. Prior

notification typically arrives with these two types of messages :

- Report transaction from the ICD on an ongoing [Subscription](#).
- Check-In message from the ICD.

[Ephemeral](#) clients SHOULD leverage [asynchronous interactions](#) since they might not be present in the network when the ICD notifies it is available for communication.

After sending either of these messages, an ICD will remain available for a period of time to allow clients to send their messages. Clients that support [LIT ICDs](#) SHOULD sequentially send any buffered interactions for a peer ICD immediately after receipt of a notification from that [LIT ICD](#).

Upon receipt of a Check-In message from an ICD, a client that has lost its previous subscription or intends to establish a new subscription with that ICD SHOULD initiate an attempt to subscribe to the ICD. If the client initiates a session establishment with the ICD in response to a Check-In message, then the client SHOULD assume the ICD to be active and use the [SESSION_ACTIVE_INTERVAL](#) for its [MRP](#) retry intervals. If the client chooses not to establish a subscription, it MAY utilize the Check-In message as a means of tracking a regular heartbeat from the given ICD. If the client no longer intends to interact with the ICD, the uninterested client SHOULD unregister itself from the ICD.

When a client wants the ICD to stay in active mode for a longer duration (typically greater than the duration of [ActiveModeThreshold](#)), it can send a [StayActiveRequest](#) command to the ICD server. The StayActiveRequest command SHALL contain a requested time duration (in milliseconds) that the client wants the ICD to stay active for. In response, the ICD SHALL reply with the duration that the ICD SHALL be active for.

Asynchronous Interactions

Consider a sensor [LIT ICD](#), which primarily exposes telemetry data via read-only attributes and allows a few commands to be executed from its application clusters. Such devices must still be administered, and that can only be carried out when the ICD is reachable.

Likely targets for administration are:

- [Opening a new commissioning window](#) to add another administrator.
- Reconfiguration of an existing fabric (e.g. [IPKs](#), [NOC rotation](#), [ACL changes](#)).
- Reconfiguration of cluster functionality (e.g. [ICD Management cluster](#), [Binding cluster](#), Groups cluster, Scenes Management cluster).
- Removal of a device from a [fabric](#).
- Changes to the device's settings.

To enable these user initiated use cases, LIT ICDs expose to clients [instructions](#) on how they can be forced into Active Mode. Clients SHOULD read these instructions from the ICD during the registration process and persist them. Clients SHOULD expose these instructions to users to permit users to trigger Active Mode as needed.

9.15.2.6. In-field upgrade of ICD functionality

This section details the behavior and requirements for an ICD upgrading from a Short Idle Time (SIT) implementation in Matter 1.3 or earlier, to a Long Idle Time (LIT) capable implementation in Matter 1.4 or later. The goal is for an ICD-capable client to detect that the device has been upgraded and is now LIT-capable, and for the client to automatically register with the device and enable the enhanced power savings mode of LIT operation.

Pre-Upgrade State

- The device is running a version of Matter that does not support the LIT ICD feature, either:
 - an older Matter version that did not define the LIT ICD feature yet (Matter 1.3 or earlier)
 - or the device did not implement optional support for the LIT ICD feature (Matter 1.4 or later)
- The device operates as a SIT ICD (i.e., any non-LIT device, including devices running Matter 1.3 or earlier).
- A client (with ICD support) exists on the network, but has not registered with the ICD via the ICD Management cluster.

Post-Upgrade State

- The device upgrades to Matter 1.4 or later and is now LIT capable or otherwise enables LIT support after initial deployment.
- The device now supports the Long Idle Time feature within the ICD Management cluster.

Client Behavior and Discovery

- Following the upgrade, a client will see the ICD Management cluster appear, signaling that the device is now LIT (Long-Idle-Time) capable. This is due to new clusters and attributes becoming available post-upgrade.
- To identify the transition from SIT to LIT, a client SHOULD detect the following:
 - The [ConfigurationVersion](#), when present, increases from its prior value. This SHOULD trigger the client to query for the following device capabilities:
 - The presence of the [ICD Management Cluster](#) either via a read, wildcard subscription, or other mechanisms.
 - The [LongIdleTimeSupport](#) feature within the ICD Management cluster is set to **1**.
 - The ICD discovery TXT key in the DNS-SD record is present.

Handling Subscription Loss in ICD Upgrades

If a client loses its subscription (even with server-side [persistent subscriptions](#)), the client SHOULD:

- Monitor for attribute changes, which the server will report if the subscription is active.
- If no reports are received (indicating a lost subscription), the client SHOULD resubscribe to the ICD Management cluster to detect changes in critical attributes (such as LongIdleTimeSupport, ConfigurationVersion, etc.).

When subscribing, clients SHOULD use wildcard subscriptions or explicitly subscribe to all relevant ICD Management attributes to ensure they receive updates about state transitions (e.g., SIT to LIT capability).

9.16. ICD Management Cluster

ICD Management Cluster enables configuration of the ICD's behavior and ensuring that listed clients can be notified when an intermittently connected device, ICD, is available for communication.

The cluster implements the requirements of the Check-In Protocol that enables the ICD Check-In use case.

9.16.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Addition of LIT ICD support, UserActiveModeTrigger, ActiveModeDuration, StayActiveResponse; removal of field Key in MonitoringRegistrationStruct
3	Addition of the ClientType attribute; addition of the ICD Check-In Backoff support; addition of the DynamicSitLitSupport feature

9.16.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	ICDM

9.16.3. Cluster ID

ID	Name
0x0046	ICDManagement

9.16.4. Features

This cluster SHALL support the FeatureMap global attribute.

Bit	Code	Feature	Conformance	Summary
0	CIP	CheckInProtocol-Support	LITS, O	Device supports attributes and commands for the Check-In Protocol support.
1	UAT	UserActiveMode-Trigger	LITS, O	Device supports the user active mode trigger feature.
2	LITS	LongIdleTimeSupport	O	Device supports operating as a Long Idle Time ICD .
3	DSLS	DynamicSitLitSupport	[LITS]	Device supports dynamic switching from SIT to LIT operating modes .

9.16.4.1. CheckInProtocolSupport Feature

When this feature is supported, the device SHALL support all the associated commands and attributes to properly support the [Check-In Protocol](#).

9.16.4.2. UserActiveModeTrigger Feature

This feature is supported if and only if the device has a [user active mode trigger](#).

9.16.4.3. LongIdleTimeSupport Feature

This feature is supported if and only the device is a [Long Idle Time ICD](#).

9.16.4.4. DynamicSitLitSupport Feature

This feature is supported if and only if the device can switch between SIT and LIT operating modes even if it has a valid registered client. See the [dynamic SIT / LIT operating mode switching](#) for more details.

9.16.5. Data Types

9.16.5.1. UserActiveModeTriggerBitmap Type

This data type is derived from map32. See the [UserActiveModeTriggerHint table](#) for requirements associated to each bit.

Bit	Name	Summary	Conformance
0	PowerCycle	Power Cycle to transition the device to ActiveMode	M
1	SettingsMenu	Settings menu on the device informs how to transition the device to ActiveMode	M
2	CustomInstruction	Custom Instruction on how to transition the device to ActiveMode	M
3	DeviceManual	Device Manual informs how to transition the device to ActiveMode	M
4	ActuateSensor	Actuate Sensor to transition the device to ActiveMode	M
5	ActuateSensorSeconds	Actuate Sensor for N seconds to transition the device to Active-Mode	M
6	ActuateSensorTimes	Actuate Sensor N times to transition the device to ActiveMode	M
7	ActuateSensorLights-Blink	Actuate Sensor until light blinks to transition the device to ActiveMode	M
8	ResetButton	Press Reset Button to transition the device to ActiveMode	M
9	ResetButtonLights-Blink	Press Reset Button until light blinks to transition the device to ActiveMode	M
10	ResetButtonSeconds	Press Reset Button for N seconds to transition the device to Active-Mode	M
11	ResetButtonTimes	Press Reset Button N times to transition the device to ActiveMode	M

Bit	Name	Summary	Conformance
12	SetupButton	Press Setup Button to transition the device to ActiveMode	M
13	SetupButtonSeconds	Press Setup Button for N seconds to transition the device to Active-Mode	M
14	SetupButtonLights-Blink	Press Setup Button until light blinks to transition the device to ActiveMode	M
15	SetupButtonTimes	Press Setup Button N times to transition the device to ActiveMode	M
16	AppDefinedButton	Press the N Button to transition the device to ActiveMode	M

ClientTypeEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	Permanent	The client is typically resident, always-on, fixed infrastructure in the home.	M
1	Ephemeral	The client is mobile or non-resident or not always-on and may not always be available in the home.	M

9.16.5.2. OperatingModeEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	SIT	ICD is operating as a Short Idle Time ICD.	M
1	LIT	ICD is operating as a Long Idle Time ICD.	M

9.16.5.3. MonitoringRegistrationStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	CheckInNodeID	node-id	all	N		S	M
2	MonitoredSubject	subject-id	all	N		S	M
3	Key						D
4	ClientType	ClientTypeEnum	all	N	0	S	M

CheckInNodeID Field

This field SHALL indicate the NodeID of the Node to which Check-In messages will be sent when the MonitoredSubject is not subscribed.

MonitoredSubject Field

This field SHALL indicate the monitored Subject ID. This field SHALL be used to determine if a particular client has an active subscription for the given entry. The MonitoredSubject, when it is a NodeID, MAY be the same as the CheckInNodeID. The MonitoredSubject gives the registering client the flexibility of having a different CheckInNodeID from the MonitoredSubject. A subscription SHALL count as an active subscription for this entry if:

- It is on the associated fabric of this entry, and
- The subject of this entry matches the [ISD](#) of the SubscriptionRequest message that created the subscription. Matching SHALL be determined using the [subject_matches](#) function defined in the [Access Control Privilege Granting Algorithm](#).

For example, if the MonitoredSubject is Node ID 0x1111_2222_3333_AAAA, and one of the subscribers to the server on the entry's associated fabric bears that Node ID, then the entry matches.

Another example is if the MonitoredSubject has the value 0xFFFF_FFFD_AA12_0002, and one of the subscribers to the server on the entry's associated fabric bears the [CASE Authenticated TAG](#) value 0xAA12 and the version 0x0002 or higher within its [NOC](#), then the entry matches.

Key Field

This field is deprecated. Use the [RegisterClient](#) command to set the [ICDToken](#).

ClientType Field

This field SHALL indicate the client's type to inform the ICD of the availability for communication of the client.

9.16.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	IdleModeDuration	uint32	1 to 64800	F	1	R V	M
0x0001	ActiveModeDuration	uint32	all	F	300	R V	M
0x0002	ActiveModeThreshold	uint16	all	F	300	R V	M
0x0003	RegisteredClients	list[MonitoringRegistrationStruct]	desc	N	[]	R A F	CIP
0x0004	ICD-Counter	uint32	all	C N	0	R A	CIP
0x0005	ClientsSupportedPerFabric	uint16	min 1	F	1	R V	CIP
0x0006	UserActiveModeTriggerHint	UserActiveModeTriggerBitmap	desc	F	0	R V	UAT
0x0007	UserActiveModeTriggerInstruction	string	max 128	F	""	R V	desc
0x0008	OperatingMode	OperatingModeEnum	all			R V	LITS
0x0009	MaximumCheckInBackoff	uint32	IdleModeDuration to 64800	F	1	R V	CIP

9.16.6.1. IdleModeDuration Attribute

This attribute SHALL indicate the maximum interval in seconds the server can stay in idle mode. The IdleModeDuration SHALL NOT be smaller than the ActiveModeDuration.

9.16.6.2. ActiveModeDuration Attribute

This attribute SHALL indicate the minimum interval in milliseconds the server typically will stay in active mode after initial transition out of idle mode. The ActiveModeDuration does not include the ActiveModeThreshold.

9.16.6.3. ActiveModeThreshold Attribute

This attribute SHALL indicate the minimum amount of time in milliseconds the server typically will stay active after network activity when in active mode.

9.16.6.4. RegisteredClients Attribute

This attribute SHALL contain all clients registered to receive notification if their subscription is lost. The maximum number of entries that can be in the list SHALL be [ClientsSupportedPerFabric](#) for each fabric supported on the server, as indicated by the value of the [SupportedFabrics](#) attribute in the [Operational Credentials cluster](#).

9.16.6.5. ICDCounter Attribute

This attribute SHALL indicate the value of the ICD Counter. The ICD Counter is used as the [Check-In Counter](#) to encrypt the check-in message for this client.

9.16.6.6. ClientsSupportedPerFabric Attribute

This attribute SHALL indicate the maximum number of entries that the server is able to store for each fabric in the [RegisteredClients](#) attribute.

9.16.6.7. UserActiveModeTriggerHint Attribute

This attribute SHALL indicate which user action(s) will trigger the ICD to switch to Active mode. If the attribute indicates support for a trigger that is dependent on the [UserActiveModeTriggerInstruction](#) in the [UserActiveModeTriggerHint table](#), the [UserActiveModeTriggerInstruction](#) attribute SHALL be implemented and SHALL provide the required information, unless specified otherwise in the requirement column of the [UserActiveModeTriggerHint table](#).

ActuateSensorLightsBlink, ResetButtonLightsBlink and SetupButtonLightsBlink (i.e. bits 7, 9 and 14) have a dependency on the [UserActiveModeTriggerInstruction](#) attribute but do not require the attribute to be present.

Bit index	Name	Requirement	UserActiveModeTriggerInstruction Dependency
0	PowerCycle	The Device will automatically enter Active Mode upon power cycle (e.g. remove/re-insert batteries).	False

Bit index	Name	Requirement	UserActiveModeTriggerInstruction Dependency
1	SettingsMenu	The settings menu on the Device provides instructions to put it into Active Mode.	False
2	CustomInstruction	The UserActiveModeTriggerInstruction key/value pair describes a custom way to put the Device into Active Mode. This Custom Instruction option is NOT recommended for use by a Device that does not have knowledge of the user's language preference.	True
3	DeviceManual	The Device Manual provides special instructions to put the Device into Active Mode (see Section 11.23.6.16 , " UserManualUrl "). This is a catch-all option to capture user interactions that are not codified by other options in this table.	False
4	ActuateSensor	The Device will enter Active Mode when the sensor is actuated, e.g. for a door sensor, open/close the door.	False
5	ActuateSensorSeconds	The Device will enter Active Mode when the sensor is actuated for N seconds. The exact value of N SHALL be made available via the UserActiveModeTriggerInstruction attribute.	True

Bit index	Name	Requirement	UserActiveModeTriggerInstruction Dependency
6	ActuateSensorTimes	The Device will enter Active Mode when reset button is pressed. The exact value of N SHALL be made available via the UserActiveModeTriggerInstruction attribute.	True
7	ActuateSensorLightsBlink	The Device will enter Active Mode when the sensor is actuated until associated light blinks. Information on color of light MAY be made available via the UserActiveModeTriggerInstruction attribute.	True
8	ResetButton	The Device will enter Active Mode when reset button is pressed.	False
9	ResetButtonLightsBlink	The Device will enter Active Mode when reset button is pressed until associated light blinks. Information on color of light MAY be made available via the UserActiveModeTriggerInstruction attribute.	True
10	ResetButtonSeconds	The Device will enter Active Mode when reset button is pressed for N seconds. The exact value of N SHALL be made available via the UserActiveModeTriggerInstruction attribute.	True

Bit index	Name	Requirement	UserActiveModeTriggerInstruction Dependency
11	ResetButtonTimes	The Device will enter Active Mode when reset button is pressed N times with maximum 1 second between each press. The exact value of N SHALL be made available via UserActiveModeTriggerInstruction attribute.	True
12	SetupButton	The Device will enter Active Mode when setup button is pressed.	False
13	SetupButtonSeconds	The Device will enter Active Mode when setup button is pressed for N seconds. The exact value of N SHALL be made available via the UserActiveModeTriggerInstruction attribute.	True
14	SetupButtonLightsBlink	The Device will enter Active Mode when setup button is pressed until associated light blinks. Information on color of light MAY be made available via the UserActiveModeTriggerInstruction attribute.	True
15	SetupButtonTimes	The Device will enter Active Mode when setup button is pressed N times with maximum 1 second between each press. The exact value of N SHALL be made available via the UserActiveModeTriggerInstruction attribute.	True

Bit index	Name	Requirement	UserActiveModeTriggerInstruction Dependency
16	AppDefinedButton	The Device will enter Active Mode when N button is pressed where N identifies which button to press. The value of N SHALL be made available via the UserActiveModeTriggerInstruction attribute	True

An [ICD](#) can indicate multiple ways of being put into Active Mode by setting multiple bits in the bitmap at the same time. However, a device SHALL NOT set more than one bit which has a dependency on the [UserActiveModeTriggerInstruction](#) attribute.

9.16.6.8. UserActiveModeTriggerInstruction Attribute

The meaning of the attribute is dependent upon the [UserActiveModeTriggerHint](#) attribute value, and the conformance is indicated in the "dependency" column in [UserActiveModeTriggerHint table](#). The [UserActiveModeTriggerInstruction](#) attribute MAY give additional information on how to transition the device to Active Mode. If the attribute is present, the value SHALL be encoded as a valid UTF-8 string with a maximum length of 128 bytes. If the [UserActiveModeTriggerHint](#) has the [ActuateSensorSeconds](#), [ActuateSensorTimes](#), [ResetButtonSeconds](#), [ResetButtonTimes](#), [SetupButtonSeconds](#) or [SetupButtonTimes](#) set, the string SHALL consist solely of an encoding of N as a decimal unsigned integer using the ASCII digits 0-9, and without leading zeros.

For example, given [UserActiveModeTriggerHint](#)="1024", [ResetButtonSeconds](#) is set which indicates "Press Reset Button for N seconds". Therefore, a value of [UserActiveModeTriggerInstruction](#)="6" would indicate that N is 6 in that context.

When [CustomInstruction](#) is set by the [UserActiveModeTriggerHint](#) attribute, indicating presence of a custom string, the [ICD](#) SHOULD perform localization (translation to user's preferred language, as indicated in the Device's currently configured locale). The Custom Instruction option SHOULD NOT be used by an [ICD](#) that does not have knowledge of the user's language preference.

When the [UserActiveModeTriggerHint](#) key indicates a light to blink ([ActuateSensorLightsBlink](#), [ResetButtonLightsBlink](#) or [SetupButtonLightsBlink](#)), information on color of light MAY be made available via the [UserActiveModeTriggerInstruction](#) attribute. When using such color indication in the [UserActiveModeTriggerInstruction](#) attribute, the string SHALL consist of exactly 6 hexadecimal digits using the ASCII characters 0-F and encoding the RGB color value as used in HTML encodings.

9.16.6.9. OperatingMode Attribute

This attribute SHALL indicate the operating mode of the ICD as specified in the [OperatingModeEnum](#).

- If the ICD is operating as a LIT ICD, OperatingMode SHALL be LIT.
- If the ICD is operating as a SIT ICD, OperatingMode SHALL be SIT.

9.16.6.10. MaximumCheckInBackoff Attribute

This attribute SHALL indicate the maximum time in seconds between two Check-In messages when [back-off](#) is active. The MaximumCheckInBackoff SHALL NOT be smaller than the [IdleModeDuration](#).

If the MaximumCheckInBackoff is equal to the [IdleModeDuration](#), it means the ICD does not [back-off](#).

9.16.7. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	RegisterClient	client ⇒ server	RegisterClientResponse	M F	CIP
0x01	RegisterClientResponse	client ⇐ server	N		CIP
0x02	UnregisterClient	client ⇒ server	Y	M F	CIP
0x03	StayActiveRequest	client ⇒ server	StayActiveResponse	O	LITS, O
0x04	StayActiveResponse	client ⇐ server	N		LITS, O

9.16.7.1. RegisterClient Command

This command allows a client to register itself with the ICD to be notified when the device is available for communication.

This command SHALL have the following data fields:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CheckInNodeID	node-id	all			M
1	MonitoredSubject	subject-id	all			M
2	Key	octstr	16			M
3	VerificationKey	octstr	16			O
4	ClientType	ClientTypeEnum	all			M

CheckInNodeID Field

This field SHALL provide the node ID to which a Check-In message will be sent if there are no active subscriptions matching MonitoredSubject.

MonitoredSubject Field

This field SHALL provide the monitored subject ID.

Key Field

This field SHALL contain the ICDToken, a 128-bit symmetric key shared by the ICD and the ICD Client, used to encrypt [Check-In](#) messages from this ICD to the MonitoredSubject.

VerificationKey Field

This field SHALL provide the verification key. The verification key represents the key already stored on the server. The verification key provided in this field SHALL be used by the server to guarantee that a client with manage permissions can only modify entries that contain a Key equal to the verification key. The verification key SHALL be provided for clients with manage permissions. The verification key SHOULD NOT be provided by clients with administrator permissions for the server cluster. The verification key SHALL be ignored by the server if it is provided by a client with administrator permissions for the server cluster.

ClientType Field

This field SHALL provide the client type of the client registering.

Effect on Receipt

On receipt of the RegisterClient command, the server SHALL perform the following procedure:

1. The server verifies that an entry for the fabric is available in the server's list of registered clients.
 - a. If one of the entries in storage for the fabric has the same CheckInNodeID as the received CheckInNodeID, the server SHALL continue from step 2.
 - b. If there is an available entry for the fabric, an entry is created for the fabric and the received CheckInNodeID, MonitoredSubject, Key and ClientType are stored. The server SHALL continue from step 5.
 - c. If there are no available entries for the fabric, the status SHALL be RESOURCE_EXHAUSTED and the server SHALL continue from step 6.
2. The server SHALL verify the privileges of the command's [ISD](#).
 - a. If the [ISD](#) of the command has administrator privileges for the server cluster, the server SHALL continue from step 4.
 - b. If the [ISD](#) of the command does not have administrator privileges for the server cluster, the server SHALL continue from step 3.
3. The server SHALL verify that the received verification key is equal to the key previously stored in the list of registered clients with the matching CheckInNodeID.

- a. If the verification key does not have a valid value, the status SHALL be FAILURE. the server SHALL continue from step 6.
 - b. If the verification key is not equal to the Key value stored in the entry, the status SHALL be FAILURE. The server SHALL continue from step 6.
 - c. If the verification key is equal to the Key value stored in the entry, the server SHALL continue from step 4.
4. The entry SHALL be updated with the received CheckInNodeID, MonitoredSubject, Key and ClientType.
 - a. If the update fails, the status SHALL be FAILURE. The server SHALL continue from step 6.
 - b. If the update succeeds, the server SHALL continue from step 5.
 5. The server SHALL persist the client information.
 - a. If the persistence fails, the status SHALL be FAILURE and the server SHALL continue from step 6.
 - b. If the persistence succeeds, the status SHALL be SUCCESS and the server SHALL continue from step 6.
 6. The server SHALL generate a response.
 - a. If the status is SUCCESS, the server SHALL generate a RegisterClientResponse command.
 - b. If the status is not SUCCESS, the server SHALL generate a default response with the Status field set to the evaluated error status.

9.16.7.2. RegisterClientResponse Command

This command SHALL be sent by the ICD Management Cluster server in response to a successful RegisterClient command.

This command SHALL have the following data fields:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ICDCounter	uint32	all			M

When Generated

This command SHALL be generated in response to a successful RegisterClient command. The ICD-Counter field SHALL be set to the ICDCounter attribute of the server.

9.16.7.3. UnregisterClient Command

This command allows a client to unregister itself with the ICD. Example: a client that is leaving the network (e.g. running on a phone which is leaving the home) can (and should) remove its subscriptions and send this UnregisterClient command before leaving to prevent the burden on the ICD of an absent client.

This command SHALL have the following data fields:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CheckInNodeID	node-id	all			M
1	VerificationKey	octstr	16			O

CheckInNodeID Field

This field SHALL provide the registered client node ID to remove from storage.

VerificationKey Field

This field SHALL provide the verification key associated with the CheckInNodeID to remove from storage. The verification key represents the key already stored on the server. The verification key provided in this field SHALL be used by the server to guarantee that a client with manage permissions can only remove entries that contain a Key equal to the stored key. The verification key SHALL be provided for clients with manage permissions. The verification key SHOULD NOT be provided by clients with administrator permissions for the server cluster. The verification key SHALL be ignored by the server if it is provided by a client with administrator permissions for the server cluster.

Effect on Receipt

On receipt of the UnregisterClient command, the server SHALL perform the following procedure:

1. The server SHALL check whether there is a entry stored on the device for the fabric with the same CheckInNodeID.
 - a. If there are no entries stored for the fabric, the status SHALL be NOT_FOUND. The server SHALL continue from step 6.
 - b. If there is an error when reading from storage, the status SHALL be FAILURE. The server SHALL continue from step 6.
 - c. If there is at least one entry stored on the server for the fabric, the server SHALL continue from step 2.
2. The server SHALL verify if one of the entries for the fabric has the corresponding CheckInNodeID received in the command.
 - a. If no entries have the corresponding CheckInNodeID, the status SHALL be NOT_FOUND. The server SHALL continue from step 6.
 - b. If an entry has the corresponding CheckInNodeID, the server SHALL continue to step 3.
3. The server SHALL check whether the **ISD** of the command has administrator permissions for the server cluster.
 - a. If the **ISD** of the command has administrator privileges for the server cluster, the server SHALL continue from step 5.
 - b. If the **ISD** of the command does not have administrator privileges for the server cluster, the server SHALL continue from step 4.

4. The server SHALL verify that the received verification key is equal to the key previously stored in the list of registered clients with the matching CheckInNodeID.

a. If the verification key does not have a valid value, the status SHALL be FAILURE. the server SHALL continue from step 6.

b. If the verification key is not equal to the Key value stored in the entry, the status SHALL be FAILURE. The server SHALL continue from step 6.

c. If the verification key is equal to the Key value stored in the entry, the server SHALL continue from step 5.
5. The server SHALL delete the entry with the matching CheckInNodeID from storage and will persist the change.

a. If the removal of the entry fails, the status SHALL be FAILURE. The server SHALL continue from step 6.

b. If the removal succeeds, the status SHALL be SUCCESS and the server SHALL continue to step 6.
6. The server SHALL generate a response with the Status field set to the evaluated status.

9.16.7.4. StayActiveRequest Command

This command allows a client to request that the server stays in active mode for at least a given time duration (in milliseconds) from when this command is received.

This command SHALL have the following data fields:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	StayActive-Duration	uint32	all			M

This StayActiveDuration MAY be longer than the ActiveModeThreshold value and would, typically, be used by the client to request the server to stay active and responsive for this period to allow a sequence of message exchanges during that period. The client MAY slightly overestimate the duration it wants the ICD to be active for, in order to account for network delays.

Effect on Receipt

When receiving a StayActiveRequest command, the server SHALL calculate the maximum PromisedActiveDuration it can remain active as the greater of the following two values:

- StayActiveDuration: Specified in the received command by the client.

• Remaining Active Time: The server’s planned remaining active time based on the ActiveMode-Threshold and its internal resources and power budget.

A server MAY replace StayActiveDuration with Minimum Active Duration in the above calculation.

PromisedActiveDuration represents the guaranteed minimum time the server will remain active, taking into account both the requested duration and the server’s capabilities.

The ICD SHALL report the calculated **PromisedActiveDuration** in a **StayActiveResponse** message back to the client.

9.16.7.5. StayActiveResponse Command

This message SHALL be sent by the ICD in response to the **StayActiveRequest** command and SHALL contain the computed duration (in milliseconds) that the ICD intends to stay active for.

This command SHALL have the following data fields:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	PromisedActiveDuration	uint32	desc			M

PromisedActiveDuration Field

This field SHALL provide the actual duration that the ICD server can stay active from the time it receives the **StayActiveRequest** command.

Minimum Value for PromisedActiveDuration

The minimum value of the **PromisedActiveDuration** field SHALL be equal to either **30000 milliseconds** or **StayActiveDuration** (from the received **StayActiveRequest** command), whichever is smaller.

Example scenarios:

- A Client requests an ICD to stay awake for 20000 milliseconds in its **StayActiveDuration** field. The ICD responds with 20000 in its **PromisedActiveDuration** if it can stay active for that duration.
- A Client requests an ICD to stay awake for 35000 milliseconds in its **StayActiveDuration** field. The ICD responds with 30000 in its **PromisedActiveDuration** since it can only stay active for that minimal amount.
- A Client requests an ICD to stay awake for 10000 milliseconds in its **StayActiveDuration** field, but the ICD's remaining active time is 20000 milliseconds. The ICD responds with 20000 milliseconds in its **PromisedActiveDuration** field since it intends to stay active that long.

9.17. Ecosystem Information Cluster

The Ecosystem Information Cluster provides extended device information for all the logical devices represented by a Bridged Node. The Ecosystem Information Cluster presents the view of device name and location metadata for presentation by a client of the cluster to a user. This cluster is intended to support Fabric Synchronization and be present on an endpoint with the BridgedNode device type listed in the DeviceTypeList of the endpoint's Descriptor cluster.

This augments the [Bridged Device Basic Information](#) Cluster in the following ways:

- The Ecosystem Information Cluster adds support for providing a name and location for individual endpoints. (The endpoints do not need to be present on the Bridge for their name and location information to be present.)
- The Ecosystem Information Cluster adds metadata to support conflict resolution between multiple sources of the name and location data.
- The Ecosystem Information Cluster supports user control for the presence of the name and location information by specifying more restricted access.

A client SHOULD use the information provided by the Ecosystem Information Cluster to help the user organize and interact with their devices. Some examples may include:

- Directly organizing and labeling the devices in a client's user interface.
- Providing hints in the user interface, which can assist the user in organizing and labeling their devices.

For the purposes of the Ecosystem Information Cluster section, an instance of the Ecosystem Information Cluster will be referred to as an "instance".

9.17.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

9.17.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Endpoint	ECOINFO

9.17.3. Cluster ID

ID	Name
0x0750	Ecosystem Information

9.17.4. Data Types

9.17.4.1. DeviceTypeStruct Type

The device type and revision define endpoint conformance to a release of a device type definition. See the Data Model specification for more information.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Device-Type	devtype-id					M
1	Revision	uint16	min 1				M

DeviceType Field

This SHALL indicate the device type definition.

Revision Field

This is the implemented revision of the device type definition.

9.17.4.2. EcosystemDeviceStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Device-Name	string	max 64		empty	S	O
1	Device-Name-LastEdit	epoch-us	all		0	S	Device-Name
2	Bridge-dEndpoint	endpoint-no	desc			S	desc
3	Original-Endpoint	endpoint-no	desc			S	desc
4	Device-Types	list[Device-Type-Struct]	desc			S	M
5	UniqueLocationIDs	list[string]	max 64 [max 64]			S	M
6	UniqueLocationIDsLastEdit	epoch-us	all		0	S	M

DeviceName Field

This field SHALL indicate the device's name, which is provided externally if the user consents. (For example, provided by the user in an ecosystem specific interface.)

DeviceNameLastEdit Field

This field SHALL indicate the timestamp of when the DeviceName was last modified.

BridgedEndpoint Field

This field SHALL indicate the endpoint this EcosystemDeviceStruct is associated with on this Bridge.

This field SHALL be present if the device is accessible through the bridge. If present, this field SHALL be set to a valid endpoint.

OriginalEndpoint Field

This field SHALL indicate the endpoint this EcosystemDeviceStruct is associated with on the original device represented by this bridge’s Bridged Node. If this bridge is receiving the device from another bridge, then the OriginalEndpoint field value would be the same on both bridges. This field SHALL be present if the original device is a Matter device. If present, this field SHALL be set to a valid endpoint on the original device.

DeviceTypes Field

This field SHALL indicate all of the DeviceTypes within the DeviceTypeList in the Descriptor Cluster associated with this EcosystemDeviceStruct entry.

This field SHALL contain a list of valid device type ids.

UniqueLocationIDs Field

This field SHALL specify the EcosystemLocationStruct entries in the LocationDirectory attribute associated with this EcosystemDeviceStruct.

UniqueLocationIDsLastEdit Field

This field SHALL indicate the timestamp of when the UniqueLocationIDs was last modified.

NOTE

If multiple server instances update the UniqueLocationIDs field at the same time, it is possible one of the updates will be missed. This is considered an acceptable limitation to reduce the complexity of the design. Since this is meant to be provided from user input, it is unlikely these signals would be happening at one time.

9.17.4.3. EcosystemLocationStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	UniqueLocationID	string	max 64			S	M
1	Location-Descriptor	location-desc				S	M
2	Location-DescriptorLastEdit	epoch-us			0	S	M

UniqueLocationID Field

This field SHALL indicate a unique identifier for a specific Ecosystem Information Cluster server instance representing the location independent of its LocationDescriptor field.

UniqueLocationID can be used by the client to determine if the change is a relocation of the device or just a renaming of the location.

No guarantees are given for consistency of the ID between server instances. The same location may be represented by different IDs on different Ecosystem Information Cluster server instances, so only the history from a single server instance should be considered when evaluating a change.

UniqueLocationID SHALL be changed when the LocationDescriptor changes from one existing location to another location as a result of an external interaction. (For example, the user changes the location assignment.)

UniqueLocationID SHALL NOT be changed when the LocationDescriptor changes name, but still represents the same location. (For example, the user renames a room.)

UniqueLocationID SHALL be changed when LocationDescriptor changes as a result of another Ecosystem Information Cluster server instance changing and the UniqueLocationID on the remote server instance also changes.

UniqueLocationID SHALL NOT be changed when LocationDescriptor changes as a result of another Ecosystem Information Cluster server instance changing and the UniqueLocationID on the remote server instance does not change.

LocationDescriptor Field

This field SHALL indicate the location (e.g. living room, driveway) and associated metadata that is provided externally if the user consents. (For example, provided by the user in an ecosystem specific interface.)

"Location" in this context is typically used by the user's grouping into rooms, areas or other logical groupings of how devices are used. So a device might be part of multiple such "Locations"s.

LocationDescriptorLastEdit Field

This field SHALL indicate the timestamp of when the LocationDescriptor was last modified.

9.17.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	DeviceDirectory	list[EcosystemDeviceStruct]		N		R M F	M
0x0001	LocationDirectory	list[EcosystemLocationStruct]		N		R M F	M

NOTE

The DeviceDirectory and LocationDirectory attributes are Fabric Scoped with Fabric Sensitive fields. The user can direct an ecosystem to disable sharing of user generated data with specific fabrics in the DeviceDirectory and LocationDirectory attributes. This can result in a client ecosystem receiving a list of entries scoped only to other fabrics (with only the fabric index visible since all other fields are fabric sensitive) if the user has disabled sharing of user generated data with specific fabrics.

9.17.5.1. DeviceDirectory Attribute

This attribute SHALL contain the list of logical devices represented by a Bridged Node. Most of the time this will contain a single entry, but may grow with more complex device compositions (e.g. another bridge.)

An empty list indicates that the information is not available.

9.17.5.2. LocationDirectory Attribute

This attribute SHALL contain the list of rooms, areas and groups associated with the DeviceDirectory entries, and SHALL NOT contain locations which are dynamically generated and removed by an ecosystem. (E.g. a location that is generated and removed based on the user being home is not permitted. However, an initially generated location name that does not quickly change is acceptable.)

An empty list indicates that the information is not available.

LocationDirectory entries SHALL be removed if there is no DeviceDirectory that references it.

9.17.6. Examples

9.17.6.1. Bridged Node with One Device and Two Locations

In the example below, there is a single ecosystem representation of a device for the Bridged Node. The user has configured the name "Dimmable Light" for that device and configured it into two locations in the ecosystem (the "Living Room" is on the "First Floor" so the device can be in two locations).

This example demonstrates the relationship between DeviceDirectory and LocationDirectory, synchronization data is omitted for clarity.




Figure 88. Ecosystem Information Cluster Examples - One Device, Two Locations

9.17.6.2. Bridged Node with Two Devices and Two Locations

In the example below, there are ecosystem representations of two devices for the Bridged Node. The user has configured them named "Outlet 1" and "Outlet 2". The devices are each configured by the user to be in a single location in the ecosystem.

This example demonstrates the relationship between DeviceDirectory and LocationDirectory, synchronization data is omitted for clarity.




Figure 89. Ecosystem Information Cluster Examples - Two Devices, Two Locations

Chapter 10. Interaction Model Encoding Specification

10.1. Overview

This specification details the encoding of the [Interaction Model](#) (IM) in the [Matter TLV format](#).

Specifically, it details the encoding of the application payload for Matter messages that map to the Interaction Model. The details of the message header are described in [Section 4.4, “Message Frame Format”](#) and out of scope of this document.

10.2. Messages

10.2.1. IM Protocol Messages

Each Action in the IM specification SHALL be mapped to a message with a unique [Protocol Opcode](#), namespaced under the [PROTOCOL_ID_INTERACTION_MODEL Protocol ID](#).

- [Vendor ID](#) = 0x0000 (Matter Common)
- [Protocol ID](#) = [PROTOCOL_ID_INTERACTION_MODEL](#)

Protocol Opcode	Action	Message
0x01	Status Response	StatusResponseMessage
0x02	Read Request	ReadRequestMessage
0x03	Subscribe Request	SubscribeRequestMessage
0x04	Subscribe Response	SubscribeResponseMessage
0x05	Report Data	ReportDataMessage
0x06	Write Request	WriteRequestMessage
0x07	Write Response	WriteResponseMessage
0x08	Invoke Request	InvokeRequestMessage
0x09	Invoke Response	InvokeResponseMessage
0x0A	Timed Request	TimedRequestMessage

10.2.2. Common Action Information Encoding

Every action SHALL encode the fields specified in [Section 8.2.5.1, “Common Action Information”](#). The methods for encoding Common Action Information fields are:

- As a field in the message header
- As a context tagged field in the action payload

For every field appearing in TLV-encoded data described by the schemas of the following sections,

and where a context-specific tag is used, any context-specific tag not listed in a given schema SHALL be reserved for future use and SHALL be silently ignored by clients and servers if seen in a payload.

10.2.2.1. Message Header Encoded Action Information

The following Common Action Information fields are encoded into the message header:

Header Field	Type	Description
Message Exchange ID	16-bit integer	Used to convey the Transaction ID
Source Node ID	64-bit integer	Node ID of the source that generated the transaction
Destination ID	64-bit integer	Either a Node ID or Group ID is encoded in here depending on what the IM indicates
Protocol ID	32 bits	The protocol to which this message belongs; all messages in this spec SHALL use the PROTOCOL_ID_INTERACTION_MODEL Protocol ID
Protocol OpCode	8 bits	The specific message type

10.2.2.2. Context Tag Encoded Action Information

The following Common Action Information fields are encoded as context tagged fields in the action message payload. All action messages defined in [Section 10.7, “Message Definitions”](#) SHALL include these tagged fields:

Common Action Field	Context Tag
InteractionModelRevision	0xFF

10.2.3. Chunking

Chunking is the act of splitting an Action that contains attribute/event data, specifically ReportData and WriteRequest actions, into multiple messages at logical boundaries due to the size limitations imposed by IPv6 for UDP packets (see [Section 4.4.4, “Message Size Requirements”](#) for more details).

Since attribute/event data within Actions are already organized into a series of [AttributeDataIBs](#) (for attributes) and [EventDataIBs](#) for event records, chunking entails maximally packing these information blocks (IBs) into a series of 'data' messages.

To ensure in-order delivery of a chunked set of IBs, each data message requires a response before the next data message can be sent. For [ReportDataMessage](#) and [WriteRequestMessage](#), a [StatusResponseMessage](#) and [WriteResponseMessage](#) are the respective response messages.

A [MoreChunkedMessages](#) flag SHALL be set on every data message except the last to convey to the

receiver possible delivery of more chunked messages within a given Action. This is specified in the [WriteRequestMessage](#) and [ReportDataMessage](#).

While most data types can be easily encoded in this scheme to fit within a message, the fact that lists can be of variable, and arbitrary, length can lead to complications. Specific strategies to encode lists that are chunking friendly are provided in [Section 10.6.4.3.1, “Lists”](#).

10.2.4. Transaction Flows

10.2.4.1. Read (Success)




Figure 90. Read message flow

10.2.4.2. Read (Server Error)




Figure 91. Read message with server-side error flow

10.2.4.3. Read (Client Error)




Figure 92. Read message with client-side error flow

10.2.4.4. Write (Success)




Figure 93. Write message flow

10.2.4.5. Write (Server Error)




Figure 94. Write message with server-side error flow

10.2.4.6. Subscribe (Success)




Figure 95. Subscription flow

10.2.4.7. Subscribe (Server Error)




Figure 96. Subscription with server-side error flow

10.2.4.8. Subscribe (Client Error)




Figure 97. Subscription with client-side error flow

10.3. Data Types

The IM specification defines a number of schema data types that are usable in a given cluster schema definition.

This section will outline their encoding onto TLV wire types, and their specific representations.

Class	Schema Data Type	TLV Type
Analog	uint8, uint16, uint24, uint32, uint40, uint48, uint56, uint64	Unsigned Integer (width is selected automatically depending on data value)
	int8, int16, int24, int32, int40, int48, int56, int64	Signed Integer (width is selected automatically depending on data value)
	float32	Floating Point Number, 4-byte value
	float64	Floating Point Number, 8-byte value

Class	Schema Data Type	TLV Type
Discrete	enum8, enum16	Unsigned Integer (width is selected automatically depending on data value)
	data8, data16, data32, data64	Unsigned Integer (width is selected automatically depending on data value)
	map8, map16, map32, map64	Unsigned Integer (width is selected automatically depending on data value)
	boolean	Boolean
Composite	string	UTF-8 string (length is selected automatically depending on data value)
	octstr	TLV octet string
Collection	list	TLV array
	struct	TLV structure

10.3.1. Analog - Integer

All signed integer schema types SHALL be encoded using the TLV signed integer type. The specific TLV element type (1-byte, 2-byte, 4-byte and 8-byte signed integer types) SHALL be selected automatically at runtime depending on the actual value.

In this regard, the actual width of the over-the-wire type can be narrower than the width specified in schema.

E.g. a 32-bit value defined in schema will be encoded to a 1-byte TLV signed integer type if the value doesn't exceed (-128 to +127).

Similarly, all unsigned integer schema types SHALL be encoded using the TLV unsigned integer type.

10.3.2. Analog - Floating Point

Both single and double precision floating point analog schema types SHALL be encoded using equivalent TLV floating point types as well.

10.3.3. Discrete - Enumeration

Enumerations SHALL be encoded using the TLV unsigned integer type, with the width selected automatically at runtime based on the actual value.

10.3.4. Discrete - Bitmap

Bitmaps SHALL be encoded using the TLV unsigned integer type, with the width selected automatically at runtime based on the actual value.

10.3.5. Composite - String

While strings are a derived data type, they SHALL be encoded using the [TLV UTF-8 string](#) type.

10.3.6. Composite - Octet String

Octet strings SHALL be encoded using TLV Byte Strings.

10.3.7. Collection - Struct

Structure types in schema SHALL be encoded using the TLV structure type.

10.3.8. Collection - List

The entirety of a list SHALL be encoded as a TLV array.

A list index SHALL start at 0.

Lists SHALL have a maximum size of 65535 elements ($2^{16}-1$).

10.3.9. Derived Types

All derived types (with the exception of strings) SHALL be encoded according to their base type.

10.3.10. Field IDs

Field IDs SHALL be encoded as:

- A context tag when the MEI prefix encodes a standard/scoped source.
- A fully-qualified profile-specific tag when the MEI prefix encodes a manufacturer code. The Vendor ID SHALL be set to the manufacturer code, the profile number set to 0 and the tag number set to the MEI suffix.

NOTE Support for encoding Field IDs with an MC source is provisional.

10.4. Sample Clusters

This section defines sample clusters (with attributes, events, and commands) for illustrative purposes; they SHALL NOT be interpreted as real clusters.

10.4.1. Disco Ball Cluster

This example cluster controls an imaginary mirrored disco ball, for the express purpose of disco

dancing.

10.4.1.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Deprecated Off enum value in Rotate attribute
3	Added pattern list feature
4	Added Name attribute
5	Added Reverse feature
6	Updated columns from Data & Interaction Model specs
7	Deprecated Party feature
8	Added WobbleSupport and WobbleSetting
9	Made Pattern attribute non-volatile

10.4.1.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Application	Endpoint	DISCO

10.4.1.3. Cluster ID

ID	Name
0x3456	Disco Ball

10.4.1.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Conformance	Summary
0	PTY	Party	D	Deprecated feature
1	AX	Axis	O	Allows the disco ball rotational axis to change

Bit	Code	Feature	Conformance	Summary
2	WBL	Wobble	O	Allows the disco ball to wobble on its axis as speed (dependent on Axis)
3	PAT	Pattern	O	Supports a list of patterns to cycle through automatically
4	STA	Statistics	O	Supports a request to statistics
5	REV	Reverse	P, O	Supports a Reverse command and counterclockwise rotation

Spec Writer Note

Describe features each in a separate section (if needed).

10.4.1.4.1. Statistics Feature

This feature indicates the ability to collect and read out statistics regarding the usage of the Disco Ball.

10.4.1.5. Data Types**10.4.1.5.1. WobbleBitmap Type**

This data type is derived from [map8](#).

Bit	Name	Summary	Conformance
0	Perpendicular	Indicate wobble perpendicular to axis	O
1	AlongAxis	Indicate wobble along Axis	O
2	Around	Indicate wobble around	O

10.4.1.5.2. RotateEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Off	Deprecated	D*
1	Clockwise	Rotation is currently clockwise	M
2	CounterClockwise	Rotation is currently counterclockwise	REV

***Spec Writer Note**

"D" means deprecated.

Clockwise Value

This value SHALL indicate that the disco ball is rotating in the clockwise direction

CounterClockwise Value

This value SHALL indicate that the disco ball is rotating in the counterclockwise direction

Spec Writer Note

Description sections are optional and should only be present, if additional details are needed for a given enum value

10.4.1.5.3. PatternStruct Type

This indicates a pattern of operation for a running disco ball.

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Duration	uint16	all		0		M
1	Rotate	RotateEnum	desc	X*	null*		M
2	Speed	uint8	max 200	X	null		M
3	Axis	uint8	max 90	X	null		AX WBL
4	Wobble-Speed	uint8	max 200	X	null		WBL**, O
5	Passcode	string	max 6	X	null	S	M

***Spec Writer Note**

Nullable data ("X") means that in some cases, the data can be null. Null data meaning must be defined (as it is below).

****Conformance**
"WBL, O" means mandatory for WBL feature, otherwise optional.

Duration Field

This field SHALL indicate the time in seconds for the disco ball to perform the pattern.

Rotate Field

This field SHALL indicate the rotation direction or null to not change the direction.

Speed Field

This field SHALL indicate the speed of the rotation, or null to not change the speed.

Axis Field

This field SHALL indicate the angle of the axis of rotation, or null to not change the angle.

WobbleSpeed Field

This field SHALL indicate the speed of the axis wobble, or null to not change the speed.

Passcode Field

An optionally specified passcode that if present, needs to always be provided in the Pattern Request command to successfully invoke this pattern.

10.4.1.6. Status Codes

10.4.1.6.1. StatusCodeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary
0x02*	UNSUPPORTED_PATTERN	The movement pattern is unsupported on the device even though all values are within constraints

***Spec Writer Note**
The list contains cluster specific status codes only indicated for a particular instance of this cluster.

The list SHALL start at 2, after the global error status values of 0 for SUCCESS and 1 for FAILURE.

10.4.1.7. Attributes**Spec Writer Note**

Attributes are supported by the server cluster only.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Run	bool	all*		0	R V	M
0x0001	Rotate	RotateEnum	all		0	R V	M
0x0002	Speed	uint8	0 to 200*		0	R V	M
0x0003	Axis	uint8	0 to 90		0	RW VO	AX WBL
0x0004	Wobble-Speed	uint8	0 to 200		0	RW VO	WBL
0x0005	Patterns	list[PatternStruct]	max 16*	N*	0	RW VM T*	PAT
0x0006	Name	string	max 16	N	0	RW VM	P, O
0x0007	Wobble-Support	WobbleBitmap	desc			R V	WobbleSetting
0x0008	Wobble-Setting	WobbleBitmap	desc			RW VM	[WBL]

Spec Writer Note*Constraint**

- "all" means all possible values.
- "desc" means see attribute description for constraints on attribute.
- "X to Y" means a value range from X=minimum to Y=maximum value.
- "max X" means range or maximum number of entries for a list or bytes for a string type derived from octstr.

Quality

- "N" indicates the read only, write only or read & write value is non-volatile across restarts.
- "F" indicates that the read-only value is static (fixed) and will not change in the future (like the ClusterRevision attribute).
- If there is no "N" or "F" then the value is volatile such that the value of the attribute may change at some point in the future.

Access

- Access column indicates R=Read Only, RW=Read Write, R*W=Read [Write], T=Timed Write, View=Read, Operate=Write, Manage=Write, Administer=Write for ACL processing.

Conformance

- Any attribute that is "M" is part of the base mandatory feature set. "O" is purely optional.
- To support the Axis *feature* any attribute with "AX" conformance must be supported (see Data Model).
- To support the Wobble feature any attribute with WBL conformance must be supported (see Data Model).
- "AX | WBL" means either feature mandates this attribute. If neither are true, then the attribute is not allowed.
- "AX & WBL" would require both features supported to mandate.
- "[PAT]" means optional for Pattern feature.
- "P, O" means provisional otherwise optional (after it's no longer provisional).
- The name of another attribute in the conformance column means the attribute is mandatory if the other attribute is supported. Conformance based on the presence of another attribute can be made optional by using []. Conformance MAY be set using any element in the same table.

10.4.1.7.1. Run Attribute

This attribute SHALL indicate if the disco ball is operating. If the Run attribute is 0, then the Speed, Rotate and WobbleSpeed attributes SHALL be 0.

10.4.1.7.2. Rotate Attribute

This attribute SHALL indicate the direction of rotation either clockwise or counterclockwise.

10.4.1.7.3. Speed Attribute

This attribute SHALL indicate the speed of the disco ball in revolutions per minute.

10.4.1.7.4. Axis Attribute

This attribute SHALL indicate the tilt of the axis of the disco ball, in degrees.

10.4.1.7.5. WobbleSpeed Attribute

This attribute SHALL indicate the speed of the wobble rotation in revolutions per minute.

10.4.1.7.6. Pattern Attribute

This attribute SHALL contain an ordered list of pattern entries. This list of patterns SHALL be used to operate the disco ball when the Pattern Request command is invoked.

10.4.1.7.7. Name Attribute

This attribute SHALL indicate a display name.

10.4.1.7.8. WobbleSupport Attribute

This attribute SHALL indicate the bits of the [WobbleBitmap](#) supported by the device.

10.4.1.7.9. WobbleSetting Attribute

This attribute SHALL indicate the selected type of wobble. This attribute is constrained to, in case of a write interaction, only accept the bits indicated in the [WobbleSupport](#) attribute.

10.4.1.8. Commands**Spec Writer Note**

Commands are supported by the client & server, but always initiated by the client.

ID	Name	Direction	Response**	Access	Quality	Conformance
0x00	StartRequest	client ⇒ server	Y	O		M
0x01	StopRequest	client ⇒ server	Y	O		M
0x02	ReverseRequest	client ⇒ server	Y	O		REV
0x03	WobbleRequest	client ⇒ server	Y	O		WBL
0x04	PatternRequest	client ⇒ server	Y	M T*		PAT
0x05	StatsRequest	client ⇒ server	StatsResponse**	O		STA
0x06	StatsResponse	client ⇐ server	N	O		STA

***Spec Writer Note**

Commands are operable (O) with Invoke. Commands are operable with Timed (T) Invoke only.

****Spec Writer Note**

"StatsRequest" command has a "StatsResponse" command. "Y" means that the command requires just a status in the Invoke Response. "N" means no response required (most response commands do not need a response).

10.4.1.8.1. StartRequest Command

Upon receipt, this SHALL start the disco ball rotating using the data as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Speed	uint8	max 200		MS*	M
1	Rotate	RotateEnum	all		Clockwise	O

***Spec Writer Note**

"MS" means Manufacturer Specific.

Speed Field

This field SHALL indicate the rotation speed.

Rotate Field

This field SHALL indicate the rotation direction.

10.4.1.8.2. StopRequest Command

Upon receipt, this SHALL stop the disco ball rotating, and SHALL set the Run, Speed and Rotate attributes to 0.

10.4.1.8.3. ReverseRequest Command

Upon receipt, this SHALL reverse the direction of the disco ball rotation. This command MAY generate an error response of UNSUPPORTED_PATTERN.

10.4.1.8.4. WobbleRequest Command

Upon receipt, this SHALL wobble the disco ball on its axis at the speed in the WobbleSpeed attribute. This command MAY generate an error response of UNSUPPORTED_PATTERN.

10.4.1.8.5. PatternRequest Command

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Passcode	string	max 6		empty	M

Passcode Field

If the passcode field is an empty string, this SHALL start the disco ball rotating using unprotected (i.e patterns that have no passcode) pattern list entries in sequence to control the disco ball. When the final entry in the list is processed the sequence SHALL restart at the first entry.

If the passcode field is not an empty string, only the patterns that correspond to the provided passcode SHALL be invoked.

10.4.1.8.6. StatsRequest Command

Upon receipt, this SHALL generate a StatsResponse command.

10.4.1.8.7. StatsResponse Command

This command SHALL be generated in response to a StatsRequest command. The data for this command SHALL be as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	LastRun	uint32	all		0	M
1	NumPatternsChanged	uint32	all		0	[PAT] *

***Spec Writer Note**

Patterns field is an optional only for the PAT feature.

LastRun Field

This field SHALL indicate the duration in seconds for the last time the disco ball was run.

NumPatternsChanged Field

This field SHALL indicate the number of pattern changes from the Patterns attribute within the last run time.

10.4.1.9. Events**Spec Writer Note**

Events are supported by the server cluster only.

ID	Name	Priority	Access	Conformance
0x00	Started	INFO	V*	M
0x01	Stopped	INFO	V	M
0x02	PatternChange	INFO	V	[PAT]

***Spec Writer Note**

All events are viewable (V).

10.4.1.9.1. Started Event

This event SHALL be generated, when the Run attribute changes from false to true.

10.4.1.9.2. Stopped Event

This event SHALL be generated, when the Run attribute changes from true to false.

10.4.1.9.3. PatternChange Event

This event SHALL be generated when the Rotate, Speed, or WobbleSpeed attributes are written or changed locally as the result of processing the Pattern attribute.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	PrevPattern	Pattern-Struct		X	null	M
1	CurPattern	Pattern-Struct				M
2	NextPattern	Pattern-Struct		X	null	M
3	Label	string	max 32	X	null	O

PrevPattern Field

This field SHALL be the previous pattern run. If there is no previous pattern, then PrevPattern SHALL be null.

CurPattern Field

This field SHALL be the current pattern being run.

NextPattern Field

This field SHALL be the next in the pattern list. If there is no next pattern, the NextPattern event field SHALL be null.

Label Field

This field SHALL indicate additional readable text that describes the reason for the pattern change.

10.4.2. Super Disco Ball Cluster

This is derived* from the Disco Ball cluster, with overrides for qualities and conformance.

***Spec Writer Note**

This is an example of a derived cluster, where only stricter conformance overrides, and the Name attribute gets a longer allowed length. If both the Disco Ball and Super Disco Ball

server clusters are on the same endpoint, they would represent a singleton instance of a cluster. This allows legacy clients implemented before the Super Disco Ball was specified to still interoperate. Blank column entries define no change to the qualities.

10.4.2.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

10.4.2.2. Classification

Hierarchy	Role	Scope	PICS Code
Disco Ball	Application	Endpoint	SUPDISC

10.4.2.3. Cluster ID

ID	Name
0xBBCC	Super Disco Ball

10.4.2.4. Features

Bit	Code	Feature	Conformance	Summary
0	PTY	Party	D	Deprecated feature
1	AX	Axis	M	Allows the disco ball rotational axis to change
2	WBL	Wobble	M	Allows the disco ball to wobble on its axis as speed (dependent on Axis)
3	PAT	Pattern	M	Supports a list of patterns to cycle through automatically
4	STA	Statistics	M	Supports a request to statistics

Bit	Code	Feature	Conformance	Summary
5	REV	Reverse	P, M	Supports a Reverse command and counterclock-wise rotation

10.4.2.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Run						M
0x0001	Rotate						M
0x0002	Speed						M
0x0003	Axis						AX WBL
0x0004	Wobble-Speed						WBL
0x0005	Pattern						PAT
0x0006	Name		max 32				P, M
0x0007	Wobble-Support						[WBL]
0x0008	Wobble-Setting						[WBL]

10.4.2.6. Events

ID	Name	Priority	Access	Conformance
0x00	Started			M
0x01	Stopped			M
0x02	PatternChange			PAT

10.5. Sample Device Types

This section defines sample device types based on the [Section 10.4, “Sample Clusters”](#) for illustrative purposes; they SHALL NOT be interpreted as real device types.

10.5.1. Disco Ball Device Type

This defines conformance to the example Disco Ball device type.

10.5.1.1. Revision History

Revision	Description
1	Initial revision
2	Attribute Name access changed to Administer
3	Initial Matter release

10.5.1.2. Classification

Device Type ID	Device Type Name	Superset Of	Class	Scope
0x5678	Disco Ball		Simple	Endpoint

10.5.1.3. Conditions

The following table lists conditions that MAY be defined for this device type. See Base Type Conditions for additional conformance tags.

Condition	Description
StatsAllowedCond	The node is allowed to collect statistics.

***Spec Writer Note**

If there are no condition defines, then this section SHALL be omitted.

10.5.1.4. Cluster Requirements

Each endpoint supporting this device type SHALL include these clusters based on the conformance defined below.

Cluster ID	Cluster Name	Client/Server	Quality	Conformance
0x0003	Identify	Server		M
0x3456	Disco Ball	Server		M

10.5.1.4.1. Identify Cluster

This is used to identify the endpoint.

Spec Writer Note

Describe each cluster use with normative requirements in a separate section (like this if needed).

10.5.1.4.2. Disco Ball Cluster

This is what controls the disco ball.

10.5.1.5. Element Requirements

The table below lists qualities and conformance that override the cluster specification requirements. A blank table cell means there is no change to that item and the value from the cluster specification applies.

Cluster ID	Cluster Name	Element	Name	Constraint	Access	Conformance
0x3456	Disco Ball	Feature	Pattern			Ethernet Wi-Fi
0x3456	Disco Ball	Feature	Statistics			StatsAllowedCond
0x3456	Disco Ball	Attribute	Name	max 10	VA	M
0x3456	Disco Ball	Event Field	Pattern-Change.Label			M

10.5.2. Super Disco Ball Device Type

This defines conformance to the example Super Disco Ball device type.

10.5.2.1. Revision History

Revision	Description
1	Initial revision

10.5.2.2. Classification

Device Type ID	Device Type Name	Superset Of	Class
0x3457	Super Disco Ball	Disco Ball	Simple

10.5.2.3. Cluster Requirements

Each endpoint supporting this device type SHALL include these clusters based on the conformance defined below.

Cluster ID	Cluster Name	Client/Server	Quality	Conformance
0x0003	Identify	Server		M
0x3456	Super Disco Ball	Server		M

10.5.2.3.1. Identify Cluster

This is used to identify the endpoint.

10.5.2.3.2. Super Disco Ball Cluster

This is what controls the super disco ball.

10.5.2.4. Element Requirements

The table below lists qualities and conformance that override the cluster specification requirements. A blank table cell means there is no change to that item and the value from the cluster specification applies.

Cluster ID	Cluster Name	Element	Name	Constraint	Access	Conformance
0x3456	Disco Ball	Feature	Pattern			Ethernet Wi-Fi
0x3456	Disco Ball	Feature	Statistics			StatsAllowedCond
0x3456	Disco Ball	Attribute	Name	max 10	VA	M
0x3456	Disco Ball	Event Field	Pattern-Change.Label			M

10.5.3. Disco Spot Device Type

This defines a spot light that shines on a disco ball to create a disco dancing experience.

10.5.3.1. Revision History

Revision	Description
1	Initial revision

10.5.3.2. Classification***Spec Writer Note**

When new device types are being created, the ID SHALL be set to ID-TBD until an ID can be assigned to the new device type.

Device Type ID	Device Type Name	Superset Of	Class	Scope
N/A	Disco Spot		Simple	Endpoint

Spec Writer Note

When adding new device types or clusters (non-sample) that do not have an allocated ID, please use **ID-TBD** in the ID column and contact the Matter data model tiger team chair for instructions on how to allocate new IDs.

10.5.3.3. Condition Requirements

The table below lists conditions and conformance values that override the device type definition.

Location	Device Type ID	Device Type Name	Condition	Constraint	Conformance
Root	NULL	Base	TimeLocale		M
Root	0x0016	Root Node	PowerSource-Cond		M

***Spec Writer Note**

If there are no condition overrides, then this section SHALL be omitted.

10.5.3.4. Cluster Requirements

Each endpoint supporting this device type SHALL include these clusters based on the conformance defined below.

***Spec Writer Note**

When new clusters are being created, the ID SHALL be set to ID-TBD until an ID can be assigned to the new cluster.

Cluster ID	Cluster Name	Client/Server	Quality	Conformance
N/A	Light Movement	Server		M
0x010D	Extended Color Light	Server		M

10.5.3.5. Element Requirements

***Spec Writer Note**

If there are no element overrides, then this section SHALL be omitted.

10.5.4. Disco Dance System Device Type

This defines the components needed for a complete disco dancing system.

10.5.4.1. Revision History

This is the revision history for this document.

Revision	Description
1	Initial revision

10.5.4.2. Classification

Device Type ID	Device Type Name	Superset Of	Class	Scope
N/A	Disco Dance System		Simple	Endpoint

10.5.4.3. Device Type Requirements****Spec Writer Note**

This is the section to define that a device type is composed of other endpoints (device types).

Each endpoint supporting this device type SHALL include endpoints with these device types based on the conformance defined below.

Device Type ID	Device Type Name	Constraint	Conformance
N/A	Disco Ball	1*	M
N/A	Disco Spot	min 1*	M

***Spec Writer Note**

Constraints are limitations on how many endpoints (parts) for this device type. In this example, there may be multiple (possibly many) spot lights for a single disco ball.

10.5.4.4. Cluster Requirements***Spec Writer Note**

If there are no cluster requirements, then this section SHALL be omitted.

10.5.4.5. Element Requirements***Spec Writer Note**

If there are no element overrides, then this section SHALL be omitted.

10.5.5. Weather Station Device Type

This defines the components needed for a weather station.

10.5.5.1. Revision History

Revision	Description
1	Initial revision

10.5.5.2. Classification

Device Type ID	Device Type Name	Superset Of	Class	Scope
N/A	Weather Station		Simple	Endpoint

10.5.5.3. Device Type Requirements

Each endpoint supporting this device type SHALL include endpoints with these device types based on the conformance defined below.

Device Type ID	Device Type Name	Constraint	Conformance
0x0302	Temperature Sensor	min 1	M
0x0307	Humidity Sensor	min 1	O
N/A	Wind Speed Sensor	min 1	O
N/A	Wind Direction Sensor	min 1	O
0x0044	Rain Sensor	min 1	O

10.5.5.4. Cluster Requirements

Cluster ID	Cluster Name	Client/Server	Quality	Conformance
0x0003	Identify	Server		M

10.5.5.5. Element Requirements***Spec Writer Note**

If there are no element overrides, then this section SHALL be omitted.

10.5.5.5.1. Semantic Tag Requirements on Component Device Types

The table below lists conformance requirements for semantic tags associated with the device types listed in the table when used in a composed device.

Device Type ID	Device Type Name	Namespace ID	Namespace	Tag ID	Tag	Conformance
0x0302	Temperature Sensor	0x06	Common Location	0x01	Outdoor	M

10.6. Information Blocks

These are elements that may apply to multiple message types, and are defined in a common way to permit re-use as a definition. Unless stated otherwise, these correspond to their identically named counterparts in the Interaction Model Specification.

10.6.1. Tag Rules

Unless otherwise noted, all context tags SHALL be emitted in the order as defined in the appropriate specification. This is done to reduce receiver side complexity in having to deal with arbitrary order tags.

10.6.2. AttributePathIB

TLV Type: List					
Tag	Comments	Tag Type	Tag Number	TLV Type	Range
EnableTagCompression		Context Tag	0	bool	-
Node		Context Tag	1	Unsigned Int	64 bits
Endpoint		Context Tag	2	Unsigned Int	16 bits
Cluster		Context Tag	3	Unsigned Int	32 bits
Attribute		Context Tag	4	Unsigned Int	32 bits
ListIndex		Context Tag	5	Unsigned Int	16 bits, nullable
WildcardPathFlags		Context Tag	6	Unsigned Int	32 bits
WildcardFilterConfigurationVersion		Context Tag	7	Unsigned Int	32 bits

- The contents of [ClusterPathIB](#) in the Interaction Model specification have been expanded here for encoding efficiency.
- The ClusterPathIB Group field is omitted here (see Node field description).
- Maximum nesting is restricted to referencing a list element in an attribute. Consequently, the [NestedPath](#) field is removed and replaced with a single [ListIndex](#) field.

10.6.2.1. EnableTagCompression

This tag is used to select between two different interpretations on the receiver when the [Node](#), [Endpoint](#), [Cluster](#), [Attribute](#) tags are omitted:

- When false or not present, omission of any of the tags in question (with the exception of [Node](#)) indicates wildcard semantics.
- When true, indicates the use of a tag compression scheme. In this case the value for any omitted tag SHALL be set to the value for that tag in the last AttributePathIB that had [EnableTagCompression](#) not present or set to false and was seen in a message that is part of the same interaction model Action as the current message.
 - The AttributePathIB the values end up coming from MAY appear in the same message (but earlier in it) as the current AttributePathIB.
 - The values that come from the previous AttributePathIB MAY still be missing. In that case, with the exception of [Node](#), they indicate wildcard semantics.

NOTE

Support for encoding using the EnableTagCompression tag is provisional.

10.6.2.2. Node

- If the **Group** field is present in the **IM representation**, the Group ID is encoded in the **DST** field in the message header and elided from the encoding here.
- The **Node** tag MAY be omitted if the target node of the path matches the NodeID of the server involved in the interaction.

10.6.2.3. Endpoint, Cluster

- Each of these tags can be omitted. The semantics of such omission depend on the value of **EnableTagCompression**.

10.6.2.4. Attribute, ListIndex

- When **EnableTagCompression** is false or not present, they have the following semantics:

Attribute	ListIndex	Description
Omitted	Omitted	Selects all attributes within the specified Node, Endpoint, Cluster
Present	Omitted	Selects a specific attribute within the specified Node, Endpoint, Cluster.
Present	Present	Selects a specific list item within a top-level attribute of type list.

This does not allow expressing all possible paths defined in the interaction model. Only paths that can be expressed MAY be used.

- The **ListIndex** tag is nullable. The null value SHALL only be used when this AttributePathIB is used in an **AttributeDataIB** and indicates a list append operation. See [Section 10.6.4.3.1, “Lists”](#) for more details.

10.6.2.5. Examples

Select all attributes on a given cluster and endpoint:

```
AttributePath = [[ Endpoint = 10, Cluster = Disco Ball ]]
```

Select all attributes in all clusters on a given endpoint:

```
Path = [[ Endpoint = 10 ]]
```

Select all attributes in all clusters on the node:

```
Path = [[ ]]
```

Select a specific attribute:

```
Path = [[ Endpoint = 10, Cluster = Disco Ball, Attribute = Axis ]]
```

Select a specific item in a top-level list:

```
Path = [[ Endpoint = 10, Cluster = Disco Ball, Attribute = Pattern, ListIndex = 4 ]]
```

Select all attributes in all clusters on a given endpoint on a proxied node:

```
Path = [[ Node = 0x18B430003020203, Endpoint = 10 ]]
```

Tag Compression Example #1:

```
Path1 = [[ Node = 0x18B430003020203, Endpoint = 10, Cluster = Disco Ball, Attribute =  
Pattern, ListIndex = 3 ]] // Start tracking path elements.  
Path2 = [[ EnableTagCompression = true, ListIndex = 4 ]] // Node, Endpoint, Cluster,  
Attribute are re-used from Path1  
Path3 = [[ EnableTagCompression = true, ListIndex = 5 ]] // Node, Endpoint, Cluster,  
Attribute are re-used from Path1  
Path4 = [[ EnableTagCompression = true, Attribute = Axis ]] // Endpoint, Cluster are  
re-used from Path1
```

Tag Compression Example #2:

```
Path1 = [[ Node = 0x18B430003020203, Cluster = Disco Ball, Attribute = Pattern,  
ListIndex = 3 ]] // Endpoint is wildcard, start tracking path elements.  
Path2 = [[ EnableTagCompression = true, ListIndex = 4 ]] // Node, Endpoint (including  
wildcard), Cluster, Attribute are re-used from Path1  
Path3 = [[ EnableTagCompression = true, ListIndex = 5 ]] // Node, Endpoint (including  
wildcard), Cluster, Attribute are re-used from Path1
```

Tag Compression Example #3:

```
Path1 = [[ Node = 0x18B430003020203, Endpoint = 10, Cluster = Disco Ball, Attribute =  
Pattern, ListIndex = 3 ]] // Start tracking path elements.  
Path2 = [[ EnableTagCompression = true, ListIndex = 4 ]] // Node, Endpoint, Cluster,  
Attribute are re-used from Path1  
Path3 = [[ EnableTagCompression = true, ListIndex = 5 ]] // Node, Endpoint, Cluster,
```

```
Attribute are re-used from Path1
Path4 = [[ Node = 0x18B430003020203, Endpoint = 20, Cluster = Disco Ball, Attribute =
Axis ]] // Reset tracker variables
Path5 = [[ EnableTagCompression = true, Attribute = Pattern, ListIndex = 5]] // Node,
Endpoint, Cluster are re-used from Path4.
```

10.6.3. DataVersionFilterIB

TLV Type: Structure					
Element	Comments	Tag Type	Tag Number	Type	Range
Path		Context Tag	0	ClusterPathIB	-
DataVersion		Context Tag	1	Unsigned Int	32 bits

10.6.4. AttributeDataIB

TLV Type: Structure					
Element	Comments	Tag Type	Tag Number	Type	Range
DataVersion		Context Tag	0	Unsigned Int	32 bits
Path		Context Tag	1	AttributePathIB	-
Data		Context Tag	2	Variable (see below)	-

- The **Change** field in the Interaction Model specification is not encoded directly. Instead, it is encoded through the use of special values in the **Path** and **Data** fields (see [Section 10.6.4.3.1, “Lists”](#)).

10.6.4.1. DataVersion

This tag can be omitted if the value of **EnableTagCompression** in the Path field is true. In this case, the value for the omitted tag SHALL be set to the value for that tag (if present) in the last AttributeDataIB that had tag compression disabled (i.e **EnableTagCompression** not present or set to false) and was seen in a message that is part of the same interaction model Action as the current message. If this tag was not present and tag compression was disabled, it SHALL be interpreted as though a data version was not specified in that, or subsequent AttributeDataIBs.

NOTE

Support for encoding using the EnableTagCompression tag is provisional.

10.6.4.2. Path

In addition to the rules specified for [AttributePathIB](#), the **Attribute** and 'Cluster' fields within that element SHALL always be present.

10.6.4.3. Data

Upon [path expansion](#) of the value in **Path**, the hierarchy and structure of the encoded data for each concrete Path SHALL be based on the schema description of the specified attribute within the specified cluster. The TLV encoding of each element in the data SHALL follow the rules of encoding as provided in [Data Types](#).

10.6.4.3.1. Lists

The various values in the **Change** enumeration are realized as follows:

Change Type	Realization
REPLACE	Path SHALL refer to a list with ListIndex omitted and Data SHALL contain new values that will replace the existing contents of the list.
ADD	Path SHALL refer to a list with ListIndex containing a value of null and Data containing the new value of the list item that will be added to the list.

NOTE

ListIndex is currently only allowed to be omitted or null. Any other value SHALL be interpreted as an error.

The below bullets detail the suggested patterns for mutating a list, depending on the encoding of the list.

- A single [AttributeDataIB](#) containing a path to the list itself and Data that contains all items in the list encoded as a TLV array. This option SHOULD be selected if it is possible to encode the entirety of the list in a single [AttributeDataIB](#) that fits in a single message.

This option SHOULD be selected when sending AttributeReports as a part of a ReportDataMessage, when it is possible to encode the entirety of the list in a single [AttributeDataIB](#) that fits in a single message.

This option SHOULD also be selected when modifying the ACL or Extension attribute of the [Access Control Cluster](#) as part of a WriteRequestMessage, when it is possible to encode the entirety of the list in a single [AttributeDataIB](#) that fits in a single message.

This option SHALL NOT be used in a WriteRequestMessage for other list attributes.

- A series of [AttributeDataIBs](#), with the first containing a path to the list itself and Data that is an empty array, which signals clearing the list, and subsequent [AttributeDataIBs](#) each containing a path to each list item with **ListIndex** being null, in order, and Data that contains the value of the list item.
- A series of [AttributeDataIBs](#), with the first containing a path to the list itself and Data that is an array that contains some number of list items, which signals replacing the list with those items, and subsequent [AttributeDataIBs](#) each containing a path to each list item with **ListIndex** being null, in order, and **Data** that contains the value of the list item.

This option SHOULD be selected when sending AttributeReports as a part of a ReportDataMessage, when it is NOT possible to encode the entirety of the list in a single [AttributeDataIB](#) that fits in a single message.

This option SHOULD also be selected when modifying the ACL or Extension attribute of the [Access Control Cluster](#) as part of a WriteRequestMessage, when it is NOT possible to encode the entirety of the list in a single [AttributeDataIB](#) that fits in a single message.

This option SHALL NOT be used in a WriteRequestMessage for other list attributes.

Because lists may be processed in chunks or by item, clients that receive an error status for a write action to a list attribute SHOULD NOT assume that the list contents are unchanged.

10.6.4.4. Examples

10.6.4.4.1. Simple Types

Update a top-level attribute:

```
AttributeDataIB = {  
    DataVersion = 1,  
    Path = [[ Endpoint = 10, Cluster = Disco Ball, FieldID = Axis ]],  
    Data = 90  
}
```

10.6.4.5. Collection Types (List)

Add an item to a list:

```
AttributeDataIB = {  
    DataVersion = 1,  
    Path = [[ Endpoint = 10, Cluster = Disco Ball, FieldID = Pattern, ListIndex =  
null]],  
    Data = {  
        Duration = 100,  
        Rotate = Counterclockwise, // On the wire enum value (2) is used  
        Speed = 12,  
        Axis = 90,  
        // WobbleSpeed omitted; this cluster instance does not support Wobble  
        Passcode = "9876"  
    }  
}
```

Replace a list (Single IB):


```

AttributeDataIB = {
  DataVersion = 1,
  Path = [[ Endpoint = 10, Cluster = Disco Ball, FieldID = Pattern]],
  Data = [[
    {
      Duration = 900,
      Rotate = Clockwise, // On the wire enum value (1) is used
      Speed = 12,
      Axis = 0,
      // WobbleSpeed omitted; this cluster instance does not support Wobble
      Passcode = "1234"
    }
    {
      Duration = 100,
      Rotate = Counterclockwise, // On the wire enum value (2) is used
      Speed = 12,
      Axis = 90,
      // WobbleSpeed omitted; this cluster instance does not support Wobble
      Passcode = "9876"
    }
  ]],
}

```

Replace a list (Multiple IBs):

```

AttributeDataIB1 = {
  DataVersion = 1,
  Path = [[ Endpoint = 10, Cluster = Disco Ball, FieldID = Pattern ]],
  Data = [
  ]
}

AttributeDataIB2 = {
  DataVersion = 1,
  Path = [[ Endpoint = 10, Cluster = Disco Ball, FieldID = Pattern, ListIndex =
null]],
  Data = {
    Duration = 900,
    Rotate = Clockwise, // On the wire enum value (1) is used
    Speed = 12,
    Axis = 0,
    // WobbleSpeed omitted; this cluster instance does not support Wobble

```

```

        Passcode = "1234"
    }
}

AttributeDataIB3 = {
    DataVersion = 1,
    Path = [[ Endpoint = 10, Cluster = Disco Ball, FieldID = Pattern, ListIndex =
null]],
    Data = {
        Duration = 100,
        Rotate = Counterclockwise, // On the wire enum value (2) is used
        Speed = 12,
        Axis = 90,
        // WobbleSpeed omitted; this cluster instance does not support Wobble
        Passcode = "9876"
    }
}

```

Replace a list (Multiple IBs, first one is nonempty):

```

AttributeDataIB1 = {
    DataVersion = 1,
    Path = [[ Endpoint = 10, Cluster = Disco Ball, FieldID = Pattern ]],
    Data = [[
        {
            Duration = 900,
            Rotate = Clockwise, // On the wire enum value (1) is used
            Speed = 12,
            Axis = 0,
            // WobbleSpeed omitted; this cluster instance does not support Wobble
            Passcode = "1234"
        }
    ]]
}

AttributeDataIB2 = {
    DataVersion = 1,
    Path = [[ Endpoint = 10, Cluster = Disco Ball, FieldID = Pattern, ListIndex =
null]],
    Data = {
        Duration = 100,
        Rotate = Counterclockwise, // On the wire enum value (2) is used
        Speed = 12,

```

```

    Axis = 90,
    // WobbleSpeed omitted; this cluster instance does not support Wobble
    Passcode = "9876"
  }
}

```

Modifying or deleting single list items SHALL be done by completely replacing the list with the options provided above.

10.6.5. AttributeReportIB

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	TLV Type	Range
AttributeSta- tus		Context Tag	0	AttributeSta- tusIB	-
AttributeData		Context Tag	1	Attribute- DataIB	-

10.6.6. EventFilterIB

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	TLV Type	Range
Node		Context Tag	0	Unsigned Int	64 bits
EventMin		Context Tag	1	Unsigned Int	64 bits

- The **Node** tag MAY be omitted if the target node of the path matches the NodeID of the server involved in the interaction.

10.6.7. ClusterPathIB

TLV Type: List					
Element	Comments	Tag Type	Tag Number	TLV Type	Range
Node		Context Tag	0	Unsigned Int	64 bits
Endpoint		Context Tag	1	Unsigned Int	16 bits
Cluster		Context Tag	2	Unsigned Int	32 bits

- The **Node** tag MAY be omitted if the target node of the path matches the NodeID of the server involved in the interaction.
- If the **Group** field is present, the Group ID is encoded in the **DST** field in the message header and elided from the encoding here.

10.6.8. EventPathIB

TLV Type: List					
Element	Comments	Tag Type	Tag Number	TLV Type	Range
Node		Context Tag	0	Unsigned Int	64 bits
Endpoint		Context Tag	1	Unsigned Int	16 bits
Cluster		Context Tag	2	Unsigned Int	32 bits
Event		Context Tag	3	Unsigned Int	32 bits
'IsUrgent'		Context Tag	4	Boolean	-

- The contents of [ClusterPathIB](#) have been expanded here to increase encoding efficiency.
- The **Node** tag MAY be omitted if the target node of the path matches the NodeID of the server involved in the interaction.
- Omission of the **Endpoint**, **Cluster** and **Event** tags SHALL have different interpretations depending on where the EventPathIB is used. See [Section 10.7.2.2, “EventRequests”](#), [Section 10.7.4.1, “EventRequests”](#), and [Section 10.7.3.1, “EventReports”](#) for the different contexts.

10.6.8.1. Examples

Select a particular event type:

```
Path = [[ Endpoint = 10, Cluster = Disco Ball, Event = Pattern Change ]]
```

Select all events on a given cluster (used in Read/Subscribe requests):

```
Path = [[ Endpoint = 10, Cluster = Disco Ball ]]
```

Select all events on a given cluster with urgency (used in Read/Subscribe requests):

```
Path = [[ Endpoint = 10, Cluster = Disco Ball, IsUrgent = true ]]
```

10.6.9. EventDataIB

TLV Type: Structure					
Element	Comments	Tag Type	Tag Number	Type	Range
Path		Context Tag	0	EventPathIB	-
EventNumber		Context Tag	1	Unsigned Int	64 bits
Priority		Context Tag	2	Unsigned Int	8 bits
one-of {					

TLV Type: Structure					
→ EpochTime-stamp		Context Tag	3	Unsigned Int	64 bits
→ SystemTime-stamp		Context Tag	4	Unsigned Int	64 bits
→ DeltaEpochTimestamp	Optional	Context Tag	5	Unsigned Int	64 bits
→ DeltaSystemTimestamp	Optional	Context Tag	6	Unsigned Int	64 bits
}					
Data		Context Tag	7	Variable (see below)	-

10.6.9.1. DeltaEpochTimestamp

This tag is present when delta encoding the UTC timestamp relative to a prior event in a given stream of events.

When this tag is present, all other timestamp tags SHALL be omitted.

This SHALL have the same units as EpochTimestamp.

10.6.9.2. DeltaSystemTimestamp

This tag is present when delta encoding the System timestamp relative to a prior event in a given stream of events.

When this tag is present, all other timestamp tags SHALL be omitted.

This SHALL have the same units as SystemTimestamp.

10.6.9.3. Data

This contains the cluster-specific payload of the Event.

The entirety of the Event is represented as a TLV Structure type.

The TLV encoding of each field in the event SHALL follow the rules of encoding as provided in Data Types.

If the cluster does not define any payload for the given event instance, this Data field SHALL be encoded as a struct with no member elements.

10.6.9.4. Examples

Single event:

```
EventDataElement = {
    Path = [[ Endpoint = 10, Cluster = Disco Ball, EventID = Started ]],
```

```
    EventNumber = 1001,
    Priority = INFO,
    EpochTimestamp = 102340234293,
    Data = {
        // Started event contains no data
    }
}
```

10.6.10. EventReportIB

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	TLV Type	Range
EventStatus		Context Tag	0	EventStatusIB	-
EventData		Context Tag	1	EventDataIB	-

10.6.11. CommandPathIB

TLV Type: List					
Element	Comments	Tag Type	Tag Number	TLV Type	Range
Endpoint		Context Tag	0	Unsigned Int	16 bits
Cluster		Context Tag	1	Unsigned Int	32 bits
Command		Context Tag	2	Unsigned Int	32 bits

- The contents of ClusterPathIB have been expanded into the CommandPathIB here to increase encoding efficiency.
- Wildcarding is achieved by omission of the respective tag.
- The Node field in the IM representation is the NodeID of the server involved in the interaction. This is omitted in the encoding here since it is retrievable from the message layer for the message containing this element.
- The Group field in the IM representation is encoded in the DST field in the message header.

10.6.11.1. Examples

Select a particular command:

```
Path = [[ Endpoint = 10, Cluster = Disco Ball, Command = Stop Request ]]
```

Select a particular command (addressed to a group):

```
Path = [[ Cluster = Disco Ball, Command = Stop Request ]]
```

10.6.12. CommandDataIB

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
CommandPath		Context Tag	0	Command-PathIB	-
CommandFields	Optional	Context Tag	1	variable	-
CommandRef	Optional	Context Tag	2	Unsigned Int	16 bits

10.6.12.1. CommandFields

This field SHALL contain the full set of arguments as specified in the description of the command request/response. The arguments SHALL follow the rules of encoding as provided in [Data Types](#).

The entirety of the arguments SHALL be encapsulated in a TLV structure, with each argument encoded appropriately using its field id as the context tag number.

If there are no arguments in the Request or Response, this tag MAY be omitted entirely.

10.6.12.2. CommandRef

This field SHALL contain an unsigned integer used to correlate responses to commands with the requests that led to those responses.

10.6.12.3. Examples

Request + Response:

```
RequestCommandElement = {
    CommandPath = [[ Endpoint = 10, Cluster = Disco Ball, Command = Stats Request ]],
    CommandFields = {} // Empty CommandFields MAY be encoded as an empty container
    CommandRef = 17
}

ResponseCommandElement = {
    CommandPath = [[ Endpoint = 10, Cluster = Disco Ball, Command = Stats Response ]],
    CommandFields = {
        LastRun = 100,
        Patterns = 1
    }
    CommandRef = 17 // the same number as in the request
}
```

Empty request:


```

RequestCommandElement = {
    CommandPath = [[ Endpoint = 10, Cluster = Disco Ball, Command = Stop Request ]]
    // Empty CommandFields MAY also be omitted entirely
    // CommandRef MAY be omitted when the request contains only a single CommandDataIB
}

// No cluster specific response is returned; a status is passed via Invoke Response at
the Interaction layer

```

10.6.13. InvokeResponseIB

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
Command		Context Tag	0	Command-DataIB	-
Status		Context Tag	1	CommandStatusIB	-

10.6.14. CommandStatusIB

TLV Type: Structure					
Element	Comments	Tag Type	Tag Number	Type	Range
CommandPath		Context Tag	0	Command-PathIB	-
Status		Context Tag	1	StatusIB	-
CommandRef	Optional	Context Tag	2	Unsigned Int	16 bits

10.6.14.1. CommandRef

This field SHALL contain an unsigned integer used to correlate responses to commands with the requests that led to those responses.

10.6.15. EventStatusIB

TLV Type: Structure					
Element	Comments	Tag Type	Tag Number	Type	Range
Path		Context Tag	0	EventPathIB	-
Status		Context Tag	1	StatusIB	-

10.6.16. AttributeStatusIB

TLV Type: Structure					
Element	Comments	Tag Type	Tag Number	Type	Range
Path		Context Tag	0	Attribut- ePathIB	-
Status		Context Tag	1	StatusIB	-

10.6.17. StatusIB

TLV Type: Structure					
Element	Comments	Tag Type	Tag Number	Type	Range
Status		Context Tag	0	Unsigned Int	8 bits
ClusterStatus		Context Tag	1	Unsigned Int	8 bits

10.7. Message Definitions

This section contains message definitions that correspond to their equivalent actions in the [Interaction Model Specification](#). Unless specifically indicated, all fields in the ensuing definitions SHALL match their equivalents in the appropriate Actions defined in the [Interaction Model Specification](#).

10.7.1. StatusResponseMessage

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
Status		Context Tag	0	Unsigned Int	8-bits

10.7.2. ReadRequestMessage

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
Attribut- eRequests	Optional	Context Tag	0	Array of Attrib- utePathIB	-
EventRequests	Optional	Context Tag	1	Array of Event- PathIB	-
EventFilters	Optional	Context Tag	2	Array of Event- FilterIB	-
FabricFiltered		Context Tag	3	boolean	-
DataVersion- Filters	Optional	Context Tag	4	Array of DataVersionFil- terIB	-

10.7.2.1. AttributeRequests

- If not present SHALL be treated as an empty list.

10.7.2.2. EventRequests

- If not present SHALL be treated as an empty list.
- Omission of any of the **Endpoint**, **Cluster**, **Event** tags indicates wildcard semantics.

10.7.2.3. EventFilters

- If not present SHALL be treated as an empty list.
- MAY be ignored (i.e. not decoded) if **EventRequests** is empty.

10.7.2.4. DataVersionFilters

- If not present SHALL be treated as an empty list.
- MAY be ignored (i.e. not decoded) if **AttributeRequests** is empty.

10.7.3. ReportDataMessage

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
SubscriptionID	Optional	Context Tag	0	Unsigned Integer	32 bits
AttributeReports	Optional	Context Tag	1	Array of AttributeReportIB	
EventReports	Optional	Context Tag	2	Array of EventReportIB	
MoreChunkedMessages	Can be omitted if false.	Context Tag	3	Boolean	-
SuppressResponse	Omit if 'false'	Context Tag	4	Boolean	-

- Multiple ReportDataMessages SHALL be sent if a Report Data action does not fit into a single message. In this case all the ReportDataMessages except the last one SHALL have **MoreChunkedMessages** set to true.
- For each ReportDataMessage received, a successful StatusResponse message SHALL be sent back to the sender unless **SuppressResponse** is true.
- If **MoreChunkedMessages** is true, **SuppressResponse** SHALL be false.
- If **MoreChunkedMessages** is false, **SuppressResponse** SHALL be set to the SuppressResponse value in the Report Data action being encoded.

10.7.3.1. EventReports

A list of [EventReportIB](#) encoded as a TLV array that have certain compression schemes applied to them to reduce redundant data.

For each EventReportIB in the list:

- The Path tag SHALL utilize the same [tag compression scheme](#) as that utilized by the tags in [AttributePathIB](#). Specifically:
 - The tag compression scheme SHALL only apply to the **Node**, **Endpoint**, **Cluster** and **Event** tags within the [EventPathIB](#) element.
 - The first element within the list SHALL specify all the necessary tags and hence serve as the anchor on which subsequent items MAY rely for compression.
- The EventNumber MAY be omitted if it is exactly one greater than the EventNumber of the previous Event.
- The 'Delta' tags SHALL be used to encode timestamps as deltas from the prior event to improve compression of large timestamps.

NOTE | Support for using the tag compression scheme for the Path is provisional.

10.7.3.1.1. Examples

Event list (highlighting compressions):

```
EventReports = [  
  {  
    Path = [[ Endpoint = 10, Cluster = Disco Ball, EventID = Started ]],  
    ImportanceLevel = INFO,  
    Number = 1001,  
    UTCTimestamp = 102340234293,  
    Data = {  
    }  
  },  
  {  
    Path = [[ EventID = PatternChange]], // same endpoint and cluster but different  
    event type  
    DeltaUTCTimestamp = 1000,  
    Data = {  
      PrevPattern = null,  
      CurPattern = {  
        Duration = 900,  
        Rotate = Clockwise, // On the wire enum value (1) is used  
        Speed = 12,  
        Axis = 0,  

```

```

        // WobbleSpeed omitted; this cluster instance does not support Wobble
        Passcode = "1234"
    },
    NextPattern = {
        Duration = 100,
        Rotate = Counterclockwise, // On the wire enum value (2) is used
        Speed = 12,
        Axis = 90,
        // WobbleSpeed omitted; this cluster instance does not support Wobble
        Passcode = "9876"
    }
}
{
    Path = [[ ]], // same path as the previous path
    DeltaUTCTimestamp = 900000000,
    Data = {
        PrevPattern = {
            Duration = 900,
            Rotate = Clockwise,
            Speed = 12,
            Axis = 0,
            Passcode = "1234"
        },
        CurPattern = {
            Duration = 100,
            Rotate = Counterclockwise,
            Speed = 12,
            Axis = 90,
            Passcode = "9876"
        },
        NextPattern = null
    }
}
]

```

10.7.3.2. MoreChunkedMessages

This flag is set to ‘true’ when there are more chunked messages in a transaction.

10.7.4. SubscribeRequestMessage

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
KeepSubscriptions		Context Tag	0	Boolean	
MinIntervalFloor		Context Tag	1	Unsigned Int	16 bits
MaxIntervalCeiling		Context Tag	2	Unsigned Int	16 bits
AttributeRequests	Optional	Context Tag	3	Array of AttributePathIB	-
EventRequests	Optional	Context Tag	4	Array of EventPathIB	-
EventFilters	Optional. Only present if EventRequests is present.	Context Tag	5	Array of EventFilterIB	-
FabricFiltered		Context Tag	7	boolean	-
DataVersionFilters	Optional. Only present if AttributeRequests is present.	Context Tag	8	Array of DataVersionFilterIB	-

10.7.4.1. EventRequests

- Omission of any of Endpoint, Cluster, Event tags indicates wildcard semantics.

10.7.5. SubscribeResponseMessage

This is sent after all Reports have been sent back to the client. The sole purpose of this is to convey the final set of parameters for the subscription back to the client.

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
SubscriptionID		Context Tag	0	Unsigned Int	32 bits
MaxInterval		Context Tag	2	Unsigned Int	16 bits

10.7.6. WriteRequestMessage

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
SuppressResponse	Omit if 'false'	Context Tag	0	Boolean	-

TLV Type: Structure (Anonymous)

TimedRequest		Context Tag	1	Boolean	-
WriteRequests		Context Tag	2	Array of AttributeDataIB	
MoreChunkedMessages	Can be omitted if false	Context Tag	3	Boolean	-

10.7.6.1. SuppressResponse

- If **MoreChunkedMessages** is true, **SuppressResponse** SHALL be false.
- If **MoreChunkedMessages** is false, **SuppressResponse** SHALL be set to the SuppressResponse value in the Write Request action being encoded.

10.7.6.2. MoreChunkedMessages

- Multiple WriteRequestMessages SHALL be sent in a single transaction if the set of AttributeDataIBs have to be sent across multiple packets. All but the last message SHALL have the **MoreChunkedMessages** flag set to true to indicate this situation. Before sending the next WriteRequestMessage, the sender SHALL await the WriteResponseMessage associated with the previous WriteRequestMessage.
- A Write Request action that is part of a Timed Write Interaction SHALL NOT be chunked.

10.7.7. WriteResponseMessage

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
WriteResponses		Context Tag	0	Array of AttributeStatusIB	-

10.7.8. TimedRequestMessage

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
Timeout		Context Tag	0	Unsigned Int	16 bits

10.7.9. InvokeRequestMessage

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
SuppressResponse		Context Tag	0	Boolean	-
TimedRequest		Context Tag	1	Boolean	-

TLV Type: Structure (Anonymous)					
InvokeRequests		Context Tag	2	Array of CommandDataIB	-

NOTE

In this version of the specification, InvokeRequestMessage contains no provisions for spanning multiple messages (see [Section 4.4.4, “Message Size Requirements”](#)).

10.7.10. InvokeResponseMessage

NOTE

The interaction model is written at the abstract level of actions, and as such, it does not address fragmentation. The interaction model describes a complete response to the invoke interaction as a unit, the chunking protocol described here maintains that illusion as much as possible. The status response messages (Status = SUCCESS) are used for flow control. Status response messages with any other Status indicate the sender of any further chunked InvokeResponseMessages must exit early. In exiting early, the sender MAY send an InvalidAction StatusResponse.

TLV Type: Structure (Anonymous)					
Element	Comments	Tag Type	Tag Number	Type	Range
SuppressResponse		Context Tag	0	Boolean	-
InvokeResponses		Context Tag	1	Array of InvokeResponseIB	-
MoreChunkedMessages		Context Tag	2	Boolean	-

10.7.10.1. SuppressResponse

- If MoreChunkedMessages is true, SuppressResponse SHALL be false.
- If MoreChunkedMessages is false, SuppressResponse SHALL be set to the SuppressResponse value in the Invoke Response action being encoded.

10.7.10.2. InvokeResponses

If the InvokeResponseMessage is being generated in response to an InvokeRequestMessage containing InvokeRequests of length 1, it SHALL contain an InvokeResponses element, and it SHALL NOT contain MoreChunkedMessages.

10.7.10.3. MoreChunkedMessages

Multiple InvokeResponseMessages MAY be sent in a single transaction when multiple CommandDataIBs were present in the Invoke Request action. All but the last message SHALL have the MoreChunkedMessages flag set to true to indicate this situation, and the last message SHALL have the MoreChunkedMessages flag set to false.

Implementations MAY choose to send the InvokeResponseMessage at any granularity, such as:

- Responding with one `InvokeResponseMessage` to each `CommandDataIB` in the corresponding `InvokeRequestMessage`;
- Waiting to batch the responses until a message full boundary is reached;
- Accumulating all responses to all `CommandDataIBs` from the `InvokeRequestMessage`, followed by sending as many messages as needed to fit all the responses.

These different `InvokeMessageResponse` buffering strategies all have different buffering costs to the server, different latency for the client, and different latency for command execution.

Upon reception of any `InvokeResponseMessage` with `MoreChunkedResponses` set to true, the receiver SHALL respond with `StatusResponseMessage`. Before sending the next `InvokeResponseMessage`, the sender SHALL await the `StatusResponseMessage` associated with the previous `InvokeResponseMessage`. If the Status field in the `StatusResponseMessage` is set to SUCCESS, that indicates the receiver is ready to process the next `InvokeResponseMessage` in the chunked series, and in that case the sender SHALL send the next chunk as soon as possible. Any other value of the Status field other than SUCCESS in the `StatusResponseMessage` SHALL cause the sender of the `InvokeResponseMessage` to terminate further transmission of `InvokeResponseMessages`, close the exchange, and consider the Invoke Interaction completed.

Chapter 11. Service and Device Management

11.1. Basic Information Cluster

This cluster provides attributes and events for determining basic information about Nodes, which supports both Commissioning and operational determination of Node characteristics, such as Vendor ID, Product ID and serial number, which apply to the whole Node.

11.1.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Added ProductAppearance attribute
3	Added SpecificationVersion and MaxPathsPerInvoke attributes
4	Updated conformance for UniqueID to mandatory
5	Added ConfigurationVersion attribute as provisional
6	Added extended information fields to CapabilityMinimaStruct and removed provisional status for ConfigurationVersion.

11.1.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	BINFO

11.1.3. Cluster ID

ID	Name
0x0028	Basic Information

11.1.4. Data Types

11.1.4.1. ProductFinishEnum Type

The data type of ProductFinishEnum is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Other	Product has some other finish not listed below.	M
1	Matte	Product has a matte finish.	M
2	Satin	Product has a satin finish.	M
3	Polished	Product has a polished or shiny finish.	M
4	Rugged	Product has a rugged finish.	M
5	Fabric	Product has a fabric finish.	M

11.1.4.2. ColorEnum Type

The data type of ColorEnum is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Black	Approximately RGB #000000.	M
1	Navy	Approximately RGB #000080.	M
2	Green	Approximately RGB #008000.	M
3	Teal	Approximately RGB #008080.	M
4	Maroon	Approximately RGB #800000.	M
5	Purple	Approximately RGB #800080.	M
6	Olive	Approximately RGB #808000.	M
7	Gray	Approximately RGB #808080.	M
8	Blue	Approximately RGB #0000FF.	M
9	Lime	Approximately RGB #00FF00.	M
10	Aqua	Approximately RGB #00FFFF.	M

Value	Name	Summary	Conformance
11	Red	Approximately RGB #FF0000.	M
12	Fuchsia	Approximately RGB #FF00FF.	M
13	Yellow	Approximately RGB #FFFF00.	M
14	White	Approximately RGB #FFFFFF.	M
15	Nickel	Typical hardware "Nickel" color.	M
16	Chrome	Typical hardware "Chrome" color.	M
17	Brass	Typical hardware "Brass" color.	M
18	Copper	Typical hardware "Copper" color.	M
19	Silver	Typical hardware "Silver" color.	M
20	Gold	Typical hardware "Gold" color.	M

11.1.4.3. ProductAppearanceStruct Type

This structure provides a description of the product's appearance.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Finish	ProductFinishEnum	all				M
1	Primary-Color	ColorEnum	all	X			M

Finish Field

This field SHALL indicate the visible finish of the product.

PrimaryColor Field

This field indicates the representative color of the visible parts of the product. If the product has no representative color, the field SHALL be null.

11.1.4.4. CapabilityMinimaStruct Type

This structure provides constant values related to overall global capabilities of this Node, that are not cluster-specific.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	CaseSessionsPerFabric	uint16	3 to 10000		3		M
1	SubscriptionsPerFabric	uint16	3 to 10000		3		M
2	SimultaneousInvocationsSupported	uint16	1 to 10000				Rev >= v6
3	SimultaneousWritesSupported	uint16	1 to 10000				Rev >= v6
4	ReadPathsSupported	uint16	9 to 10000				Rev >= v6
5	SubscribePathsSupported	uint16	3 to 10000				Rev >= v6

CaseSessionsPerFabric Field

This field SHALL indicate the actual minimum number of concurrent [CASE](#) sessions that are supported per fabric.

This value SHALL NOT be smaller than the required minimum indicated in [Section 4.14.2.8, “Minimal Number of CASE Sessions”](#).

SubscriptionsPerFabric Field

This field SHALL indicate the actual minimum number of concurrent [subscriptions](#) supported per fabric.

This value SHALL NOT be smaller than the required minimum indicated in [Section 8.5.1, “Subscribe Transaction”](#).

SimultaneousInvocationsSupported Field

This field SHALL indicate the actual maximum number of concurrent [Invoke interactions](#) that can be processed simultaneously by the node before possibly returning a BUSY status code.

SimultaneousWritesSupported Field

This field SHALL indicate the actual minimum number of concurrent [Write interactions](#) that can be processed simultaneously by the node before possibly returning a BUSY status code.

ReadPathsSupported Field

This field SHALL indicate the actual maximum number of read paths (i.e. the sum of lengths of the lists of [AttributePathIB](#) and [EventPathIB](#) in the action) which a node guarantees being able to process in any [Read Request Action](#).

This is related to [Section 2.11.2.1, “Read Interaction Limits”](#).

SubscribePathsSupported Field

This field SHALL indicate the actual maximum number of subscription paths (i.e. the sum of lengths of the lists of [AttributePathIB](#) and [EventPathIB](#) in the action) which a node guarantees being able to process in any [Subscribe Request Action](#).

This is related to [Section 2.11.2.2, “Subscribe Interaction Limits”](#).

11.1.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	DataModelRevision	uint16	desc	F	MS	R V	M
0x0001	Vendor-Name	string	max 32	F	MS	R V	M
0x0002	VendorID	vendor-id	all	F	MS	R V	M
0x0003	Product-Name	string	max 32	F	MS	R V	M
0x0004	ProductID	uint16	all	F	MS	R V	M
0x0005	NodeLabel	string	max 32	N	""	RW VM	M
0x0006	Location	string	2	N	"XX"	RW VA	M
0x0007	Hardware-Version	uint16	all	F	0	R V	M
0x0008	Hardware-Version-String	string	1 to 64	F	MS	R V	M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0009	Software-Version	uint32	desc	F	0	R V	M
0x000A	Software-Version-String	string	1 to 64	F	MS	R V	M
0x000B	Manufacturing-Date	string	8 to 16	F	MS	R V	O
0x000C	PartNumber	string	max 32	F	MS	R V	O
0x000D	ProductURL	string	max 256	F	MS	R V	O
0x000E	Product-Label	string	max 64	F	MS	R V	O
0x000F	Serial-Number	string	max 32	F	MS	R V	O
0x0010	LocalConfigDisabled	bool	all	N	False	RW VM	O
0x0011	Reachable	bool	all		True	R V	O
0x0012	UniqueID	string	max 32	F	MS	R V	Rev >= v4, O
0x0013	CapabilityMinima	CapabilityMinimaStruct	all	F	MS	R V	M
0x0014	ProductAppearance	ProductAppearanceStruct	all	F	MS	R V	[Rev >= v2]
0x0015	SpecificationVersion	uint32	desc	F	0	R V	Rev >= v3
0x0016	Max-PathsPer-Invoke	uint16	min 1	F	1	R V	Rev >= v3
0x0018	ConfigurationVersion	uint32	min 1	N	1	R V	Rev >= v6

11.1.5.1. DataModelRevision Attribute

This attribute SHALL be set to the revision number of the Data Model against which the Node is certified. The value of this attribute SHALL be one of the valid values listed in [Section 7.1.1, “Revision History”](#).

11.1.5.2. VendorName Attribute

This attribute SHALL specify a human readable (displayable) name of the vendor for the Node.

11.1.5.3. VendorID Attribute

This attribute SHALL specify the [Vendor ID](#).

11.1.5.4. ProductName Attribute

This attribute SHALL specify a human readable (displayable) name of the model for the Node such as the model number (or other identifier) assigned by the vendor.

11.1.5.5. ProductID Attribute

This attribute SHALL specify the [Product ID](#) assigned by the vendor that is unique to the specific product of the Node.

11.1.5.6. NodeLabel Attribute

This attribute SHALL represent a user defined name for the Node. This attribute SHOULD be set during initial commissioning and MAY be updated by further reconfigurations.

11.1.5.7. Location Attribute

This attribute SHALL be an ISO 3166-1 alpha-2 code to represent the country, dependent territory, or special area of geographic interest in which the Node is located at the time of the attribute being set. This attribute SHALL be set during initial commissioning (unless already set) and MAY be updated by further reconfigurations. This attribute MAY affect some regulatory aspects of the Node’s operation, such as radio transmission power levels in given spectrum allocation bands if technologies where this is applicable are used. The Location’s region code SHALL be interpreted in a case-insensitive manner. If the Node cannot understand the location code with which it was configured, or the location code has not yet been configured, it SHALL configure itself in a region-agnostic manner as determined by the vendor, avoiding region-specific assumptions as much as is practical. The special value XX SHALL indicate that region-agnostic mode is used.

11.1.5.8. HardwareVersion Attribute

This attribute SHALL specify the version number of the hardware of the Node. The meaning of its value, and the versioning scheme, are vendor defined.

11.1.5.9. HardwareVersionString Attribute

This attribute SHALL specify the version number of the hardware of the Node. The meaning of its value, and the versioning scheme, are vendor defined. The HardwareVersionString attribute SHALL

be used to provide a more user-friendly value than that represented by the [HardwareVersion](#) attribute.

11.1.5.10. SoftwareVersion Attribute

This attribute SHALL contain the current version number for the software running on this Node. A larger value of SoftwareVersion is newer than a lower value, from the perspective of software updates (see [Section 11.20.3.3, “Availability of Software Images”](#)). Nodes MAY query this field to determine the currently running version of software on another given Node.

11.1.5.11. SoftwareVersionString Attribute

This attribute SHALL contain a current human-readable representation for the software running on the Node. This version information MAY be conveyed to users. The maximum length of the SoftwareVersionString attribute is 64 bytes of UTF-8 characters. The contents SHOULD only use simple 7-bit ASCII alphanumeric and punctuation characters, so as to simplify the conveyance of the value to a variety of cultures.

Examples of version strings include "1.0", "1.2.3456", "1.2-2", "1.0b123", "1.2_3".

11.1.5.12. ManufacturingDate Attribute

This attribute SHALL specify the date that the Node was manufactured. The first 8 characters SHALL specify the date of manufacture of the Node in international date notation according to ISO 8601, i.e., YYYYMMDD, e.g., 20060814. The final 8 characters MAY include country, factory, line, shift or other related information at the option of the vendor. The format of this information is vendor defined.

11.1.5.13. PartNumber Attribute

This attribute SHALL specify a human-readable (displayable) vendor assigned part number for the Node whose meaning and numbering scheme is vendor defined.

Multiple products (and hence PartNumbers) can share a ProductID. For instance, there may be different packaging (with different PartNumbers) for different regions; also different colors of a product might share the ProductID but may have a different PartNumber.

11.1.5.14. ProductURL Attribute

This attribute SHALL specify a link to a product specific web page. The specified URL SHOULD resolve to a maintained web page available for the lifetime of the product. The syntax of this attribute SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme. The maximum length of this attribute is 256 ASCII characters.

11.1.5.15. ProductLabel Attribute

This attribute SHALL specify a vendor specific human readable (displayable) product label. The ProductLabel attribute MAY be used to provide a more user-friendly value than that represented by the [ProductName](#) attribute. The ProductLabel attribute SHOULD NOT include the name of the vendor as defined within the [VendorName](#) attribute.

11.1.5.16. SerialNumber Attribute

This attribute SHALL specify a human readable (displayable) serial number.

11.1.5.17. LocalConfigDisabled Attribute

This attribute SHALL allow a local Node configuration to be disabled. When this attribute is set to True the Node SHALL disable the ability to configure the Node through an on-Node user interface. The value of the LocalConfigDisabled attribute SHALL NOT in any way modify, disable, or otherwise affect the user's ability to trigger a factory reset on the Node.

11.1.5.18. Reachable Attribute

This attribute (when used) SHALL indicate whether the Node can be reached. For a native Node this is implicitly True (and its use is optional).

Its main use case is in the derived [Bridged Device Basic Information](#) cluster where it is used to indicate whether the bridged device is reachable by the bridge over the non-native network.

11.1.5.19. UniqueID Attribute

This attribute SHALL indicate a unique identifier for the device, which is constructed in a manufacturer specific manner.

It MAY be constructed using a permanent device identifier (such as device MAC address) as basis. In order to prevent tracking,

- it SHOULD NOT be identical to (or easily derived from) such permanent device identifier
- it SHALL be updated when the device is factory reset
- it SHALL NOT be identical to the SerialNumber attribute
- it SHALL NOT be printed on the product or delivered with the product

The value does not need to be human readable, since it is intended for machine to machine (M2M) communication.

NOTE

The conformance of the UniqueID attribute was optional in cluster revisions prior to revision 4.

NOTE

This UniqueID attribute SHALL NOT be the same as the Persistent Unique ID which is used in the [Rotating Device Identifier](#) mechanism.

11.1.5.20. CapabilityMinima Attribute

This attribute SHALL provide the minimum guaranteed value for some system-wide resource capabilities that are not otherwise cluster-specific and do not appear elsewhere. This attribute MAY be used by clients to optimize communication with Nodes by allowing them to use more than the strict minimum values required by this specification, wherever available.

The values supported by the server in reality MAY be larger than the values provided in this attribute, such as if a server is not resource-constrained at all. However, clients SHOULD only rely

on the amounts provided in this attribute.

Note that since the fixed values within this attribute MAY change over time, both increasing and decreasing, as software versions change for a given Node, clients SHOULD take care not to assume forever unchanging values and SHOULD NOT cache this value permanently at Commissioning time.

11.1.5.21. ProductAppearance Attribute

This attribute SHALL provide information about the appearance of the product, which could be useful to a user trying to locate or identify the node.

11.1.5.22. SpecificationVersion Attribute

This attribute SHALL contain the current version number for the specification version this Node was certified against. A larger value of SpecificationVersion is newer than a lower value.

Nodes MAY query this field to determine the currently supported version of the specification on another given Node.

The format of this number is segmented as its four component bytes.

Bit positions for the fields are as follows:

Bits	Name	Summary
31 .. 24	Major	Major version of specification.
23 .. 16	Minor	Minor version of specification.
15 .. 8	Dot	Dot version of the specification.
7 .. 0	Reserved1	Future reserved version field 1, set to 0 until defined.

For example, a SpecificationVersion value of 0x01040200 is composed of 4 version components, representing a version 1.4.2.0.

In the example above:

- Major version is the most significant byte (0x01).
- Minor version is the second most significant byte (0x04).
- Dot version is the third most significant byte (0x02).
- Reserved1 value is the least significant byte (0x00).

The initial revision (1.0) of this specification (1.0) was 0x01000000. Matter Spring 2024 release (1.3) was 0x01030000.

If the **SpecificationVersion** is absent or zero, such as in Basic Information cluster revisions prior to Revision 3, the specification version cannot be properly inferred unless other heuristics are employed.

Comparison of **SpecificationVersion** SHALL always include the total value over 32 bits, without

masking reserved parts.

11.1.5.23. MaxPathsPerInvoke Attribute

This attribute SHALL indicate the maximum number of elements in a single InvokeRequests list (see [Section 8.8.2, “Invoke Request Action”](#)) that the Node is able to process. Note that since this attribute MAY change over time, both increasing and decreasing, as software versions change for a given Node, clients SHOULD take care not to assume forever unchanging values and SHOULD NOT cache this value permanently at Commissioning time.

If the **MaxPathsPerInvoke** attribute is absent or zero, such as in Basic Information cluster revisions prior to Revision 3, clients SHALL assume a value of 1.

11.1.5.24. ConfigurationVersion Attribute

This attribute SHALL contain the current version number for the configuration of the Node. A larger value of ConfigurationVersion SHALL indicate a newer configuration than a lower value.

11.1.6. Events

ID	Name	Priority	Quality	Access	Conformance
0x00	StartUp	CRITICAL		V	M
0x01	ShutDown	CRITICAL		V	O
0x02	Leave	INFO		V	O
0x03	Reach-ableChanged	INFO		V	Reachable

11.1.6.1. StartUp Event

The StartUp event SHALL be generated by a Node as soon as reasonable after completing a boot or reboot process. The StartUp event SHOULD be the first Data Model event recorded by the Node after it completes a boot or reboot process.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Software-Version	uint32				M

SoftwareVersion Field

This field SHALL be set to the same value as the one available in the [SoftwareVersion](#) attribute.

11.1.6.2. ShutDown Event

The ShutDown event SHOULD be generated by a Node prior to any orderly shutdown sequence on a best-effort basis. When a ShutDown event is generated, it SHOULD be the last Data Model event

recorded by the Node. This event SHOULD be delivered urgently to current subscribers on a best-effort basis. Any subsequent incoming interactions to the Node MAY be dropped until the completion of a future boot or reboot process.

11.1.6.3. Leave Event

The Leave event SHOULD be generated by a Node prior to permanently leaving a given Fabric, such as when the [RemoveFabric](#) command is invoked for a given fabric, or triggered by factory reset or some other manufacturer specific action to disable or reset the operational data in the Node. When a Leave event is generated, it SHOULD be assumed that the fabric recorded in the event is no longer usable, and subsequent interactions targeting that fabric will most likely fail.

Upon receipt of Leave Event on a subscription, the receiving Node MAY update other nodes in the fabric by removing related bindings, access control list entries and other data referencing the leaving Node.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	FabricIndex	fabric-idx	1 to 254			M

FabricIndex Field

This field SHALL contain the local [fabric-index](#) of the fabric which the node is about to leave.

11.1.6.4. ReachableChanged Event

This event (when supported) SHALL be generated when there is a change in the Reachable attribute.

Its main use case is in the derived [Bridged Device Basic Information](#) cluster.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Reachable- NewValue	bool				M

ReachableNewValue Field

This field SHALL indicate the value of the [Reachable](#) attribute after it was changed.

11.2. Group Key Management Cluster

The Group Key Management cluster manages shared symmetric keys for the node. The cluster is scoped to the node and is a singleton for the node. This cluster maintains a list of groups supported by the node. Each group list entry supports a single group, with a single group ID and single group key. Duplicate groups are not allowed in the list. Additions or removal of a group entry are per-

formed via modifications of the list. Such modifications require Administer privilege.

Each group entry includes a membership list of zero or more endpoints that are members of the group on the node. Modification of this membership list is done via the Groups cluster, which is scoped to an endpoint. See the [Chapter 9, System Model Specification](#) specification for more information on groups.

11.2.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Clarify KeySetWrite validation and behavior on invalid epoch key lengths
3	Refocus on pure key management; introduce Groupcast feature associated with Groupcast cluster

11.2.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	GRPKEY

11.2.3. Cluster ID

ID	Name
0x003F	GroupKeyManagement

11.2.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Conformance	Summary
0	CS	CacheAndSync	P	The ability to support Cache-AndSync security policy and MCSP .
1	GCAST	Groupcast	P, O	The device supports groups using the Groupcast cluster .

11.2.4.1. CacheAndSync Feature

The [CacheAndSync](#) security policy has been provisional since Matter v1.0.

11.2.4.2. Groupcast Feature

When set, group management and group key mapping is done using the [Section 11.27, “Groupcast Cluster”](#).

If the Groupcast cluster is present on the Root Node endpoint, then this feature bit SHALL be set.

When this feature map bit is set, this cluster SHOULD be used solely for key management as the Groupcast cluster offers more direct and long-term supported methods of managing group key mapping.

11.2.5. Data Types

11.2.5.1. GroupKeySecurityPolicyEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	TrustFirst	Message counter synchronization using trust-first	M
1	CacheAndSync	Message counter synchronization using cache-and-sync	P, CS

11.2.5.2. GroupKeyMulticastPolicyEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	PerGroupID	Indicates filtering of multicast messages for a specific Group ID	M
1	AllNodes	Indicates not filtering of multicast messages	M

PerGroupID Value

The 16-bit Group Identifier of the Multicast Address SHALL be the Group ID of the group.

AllNodes Value

The 16-bit Group Identifier of the Multicast Address SHALL be 0xFFFF.

11.2.5.3. GroupKeyMapStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	GroupId	group-id	all				M
2	GroupKey-SetID	uint16	1 to 65535				M

GroupId Field

This field uniquely identifies the group within the scope of the given Fabric.

GroupKeySetID Field

This field references the set of group keys that generate operational group keys for use with this group, as specified in [Section 4.17.3.5.1, “Group Key Set ID”](#).

A GroupKeyMapStruct SHALL NOT accept GroupKeySetID of 0, which is reserved for the [IPK](#).

11.2.5.4. GroupKeySetStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	GroupKey-SetID	uint16	all				M
1	GroupKey-Security-Policy	GroupKey-Security-PolicyEnum	all				M
2	EpochKey 0	octstr	16	X			M
3	EpochStar tTime0	epoch-us	all	X			M
4	EpochKey 1	octstr	16	X			M
5	EpochStar tTime1	epoch-us	all	X			M
6	EpochKey 2	octstr	16	X			M
7	EpochStar tTime2	epoch-us	all	X			M

Access Modifier: Fabric Scoped							
8	GroupKey-Multicast-Policy	GroupKey-Multicast-PolicyEnum	all				D

GroupKeySetID Field

This field SHALL provide the fabric-unique index for the associated group key set, as specified in [Section 4.17.3.5.1, “Group Key Set ID”](#).

GroupKeySecurityPolicy Field

This field SHALL provide the security policy for an operational group key set.

When CacheAndSync is not supported in the [FeatureMap](#) of this cluster, any action attempting to set CacheAndSync in the GroupKeySecurityPolicy field SHALL fail with an INVALID_COMMAND error.

EpochKey0 Field

This field, if not null, SHALL be the InputKey used in the derivation of an OperationalGroupKey for epoch slot 0 of the given group key set. The derived OperationalGroupKey SHALL be persistently stored for the lifetime of the derived key; however, the InputKey itself SHALL NOT be stored. If EpochKey0 is not null, EpochStartTime0 SHALL NOT be null.

EpochStartTime0 Field

This field, if not null, SHALL define when EpochKey0 becomes valid as specified by [Section 4.17.3, “Epoch Keys”](#). Units are absolute UTC time in microseconds encoded using the [epoch-us](#) representation.

EpochKey1 Field

This field, if not null, SHALL be the InputKey used in the derivation of an OperationalGroupKey for epoch slot 1 of the given group key set. The derived OperationalGroupKey SHALL be persistently stored for the lifetime of the derived key; however, the InputKey itself SHALL NOT be stored. If EpochKey1 is not null, EpochStartTime1 SHALL NOT be null.

EpochStartTime1 Field

This field, if not null, SHALL define when EpochKey1 becomes valid as specified by [Section 4.17.3, “Epoch Keys”](#). Units are absolute UTC time in microseconds encoded using the [epoch-us](#) representation.

EpochKey2 Field

If the GCAST feature bit is set in the FeatureMap, this field SHALL be null, unless the GroupKeySetId is 0 (the Identity Protection Key).

This field, if not null, SHALL be the InputKey used in the derivation of an OperationalGroupKey for

epoch slot 2 of the given group key set. The derived `OperationalGroupKey` SHALL be persistently stored for the lifetime of the derived key; however, the `InputKey` itself SHALL NOT be stored. If `EpochKey2` is not null, `EpochStartTime2` SHALL NOT be null.

EpochStartTime2 Field

If the GCAST feature bit is set in the `FeatureMap`, this field SHALL be null, unless the `GroupKeyId` is 0 (the Identity Protection Key).

This field, if not null, SHALL define when `EpochKey2` becomes valid as specified by [Section 4.17.3, “Epoch Keys”](#). Units are absolute UTC time in microseconds encoded using the `epoch-us` representation.

11.2.5.5. GroupInfoMapStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	GroupId	group-id	all				M
2	Endpoints	list[endpoint-no]	min 1				M
3	Group-Name	string	max 16				O

GroupId Field

This field uniquely identifies the group within the scope of the given Fabric.

Endpoints Field

This field provides the list of Endpoint IDs on the Node to which messages to this group SHALL be forwarded.

GroupName Field

This field provides a name for the group. This field SHALL contain the last `GroupName` written for a given `GroupId` on any Endpoint via the Groups cluster.

11.2.5.6. GroupcastAdoptionStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	GroupcastAdopted	bool	all				M

GroupcastAdopted Field

This field SHALL indicate whether Groupcast was adopted by the associated Fabric's administrators.

11.2.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Group-KeyMap	list[GroupKeyMapStruct]	desc	N	empty	RW VM F	M
0x0001	GroupTable	list[GroupInfoMapStruct]	desc		empty	R V F	M
0x0002	Max-GroupsPerFabric	uint16	all	F	0	R V	M
0x0003	Max-GroupKeysPerFabric	uint16	1 to 65535	F	1	R V	M
0x0004	GroupcastAdoption	list[GroupcastAdoptionStruct]	desc	N		RW A F	P, GCAST

11.2.6.1. GroupKeyMap Attribute

If the GCAST feature bit is set in the FeatureMap attribute, the following rules apply to the accessing Fabric:

- When Groupcast is adopted (the **GroupcastAdoption** entry has **GroupcastAdopted** set to true):
 - This attribute SHALL be empty.
 - Any attempt to write to this attribute SHALL fail with an **INVALID_IN_STATE** status code.
- Otherwise (Groupcast is not adopted or the entry is missing):
 - This attribute SHALL contain the Group Key Set mappings derived from the Groupcast cluster's Membership attribute (one mapping per group per fabric).
 - GroupKeyMapStruct entry updates SHALL cause the associated Groupcast cluster's Membership attribute (by GroupID) to be updated with the provided GroupKeySetID. If an entry is missing for a given GroupID in the GroupKeyMap, which exists in the Groupcast cluster's Membership attribute for a given fabric, then the Groupcast cluster's membership attribute SHALL use placeholder value 65535 for the KeySetID. While this KeySetID is technically valid, administrators SHOULD avoid allocating it for actual usage to avoid value aliasing for this field.

This attribute is a list of [GroupKeyMapStruct](#) entries. Each entry associates a logical Group Id with a particular group key set.

11.2.6.2. GroupTable Attribute

If the GCAST feature is set in the FeatureMap:

- If the GroupcastAdoption attribute has an entry for the accessing Fabric and that entry has the GroupcastAdopted field set to true, then this field SHALL be empty.
- Else this attribute SHALL contain the Group mappings computed in equivalence to the Groupcast cluster's Membership attribute (one mapping per group per fabric).

This attribute is a list of GroupInfoMapStruct entries. Each entry provides read-only information about how a given logical Group ID maps to a particular set of endpoints, and a name for the group. The content of this attribute reflects data managed via the Groups cluster (see [AppClusters](#)), and is in general terms referred to as the 'node-wide Group Table'.

The GroupTable SHALL NOT contain any entry whose [GroupInfoMapStruct](#) has an empty Endpoints list. If a RemoveGroup or RemoveAllGroups command causes the removal of a group mapping from its last mapped endpoint, the entire GroupTable entry for that given GroupId SHALL be removed.

11.2.6.3. MaxGroupsPerFabric Attribute

If the Groupcast support is enabled (GCAST feature is set), this SHALL be set to 0 indicating group management is done using the [Groupcast cluster](#) and not the legacy Groups cluster.

This attribute SHALL indicate the maximum number of legacy groups that this node supports per fabric. For legacy usage, the value of this attribute SHALL be set to be no less than the required minimum supported groups as specified in [Section 2.11.1.2, "Group Limits"](#).

The length of the [GroupKeyMap](#) and [GroupTable](#) list attributes SHALL NOT exceed the value of the MaxGroupsPerFabric attribute multiplied by the number of [supported fabrics](#).

11.2.6.4. MaxGroupKeysPerFabric Attribute

This attribute SHALL indicate the maximum number of [group key sets](#) this node supports per fabric. The value of this attribute SHALL be set according to the minimum number of group key sets to support as specified in [Section 2.11.1.2, "Group Limits"](#).

11.2.6.5. GroupcastAdoption Attribute

This attribute SHALL indicate whether the accessing fabric claims to have migrated to Groupcast.

When a Fabric's entry has the GroupcastAdopted field set to true, the behavior of the GroupKeyMap and GroupTable attributes will change (see description of respective attributes).

There SHALL NOT be more than 1 entry per fabric in this attribute.

If a Fabric has not yet written an entry for themselves, the server SHALL act as if that Fabric had written an entry with [GroupcastAdopted](#) set to false, even if not present in the list.

11.2.7. Commands

All commands in this cluster SHALL be [scoped](#) to the accessing fabric.

ID	Name	Direction	Response	Access	Conformance
0x00	KeySetWrite	client ⇒ server	Y	A F	M
0x01	KeySetRead	client ⇒ server	KeySetRead-Response	A F	M
0x02	KeySetRead-Response	client ⇐ server	N		M
0x03	KeySetRemove	client ⇒ server	Y	A F	M
0x04	KeySetReadAllIndices	client ⇒ server	KeySetReadAllIndicesResponse	A F	M
0x05	KeySetReadAllIndicesResponse	client ⇐ server	N		M

11.2.7.1. KeySetWrite Command

This command is used by Administrators to set the state of a given Group Key Set, including atomically updating the state of all epoch keys.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupKey-Set	GroupKeySetStruct				M

Effect on Receipt

The following validations SHALL be done against the content of the GroupKeySet field:

- If the EpochKey0 field is null or its associated EpochStartTime0 field is null, then this command SHALL fail with an INVALID_COMMAND status code responded to the client.
- If the EpochKey0 field's length is not exactly 16 bytes, then this command SHALL fail with a CONSTRAINT_ERROR status code responded to the client.
- If the EpochStartTime0 is set to 0, then this command SHALL fail with an INVALID_COMMAND status code responded to the client. Note that internally, a GroupKeySetStruct's EpochStartTime0 MAY be set to zero, due to the behavior of the AddNOC command which synthesizes a GroupKeySetStruct (see [Section 11.18.6.8.1, "IPKValue Field"](#)). However, the value 0 is illegal in the GroupKeySet field sent by a client.
- If the EpochKey1 field is not null, then the EpochKey0 field SHALL NOT be null. Otherwise this command SHALL fail with an INVALID_COMMAND status code responded to the client.

- If the EpochKey1 field is not null, and the field's length is not exactly 16 bytes, then this command SHALL fail with a CONSTRAINT_ERROR status code responded to the client.
- If the EpochKey1 field is not null, its associated EpochStartTime1 field SHALL NOT be null and SHALL contain a later epoch start time than the epoch start time found in the EpochStartTime0 field. Otherwise this command SHALL fail with an INVALID_COMMAND status code responded to the client.
- If exactly one of the EpochKey1 or EpochStartTime1 is null, rather than both being null, or neither being null, then this command SHALL fail with an INVALID_COMMAND status code responded to the client.
- If the EpochKey2 field is not null, then the EpochKey1 and EpochKey0 fields SHALL NOT be null. Otherwise this command SHALL fail with an INVALID_COMMAND status code responded to the client.
- If the EpochKey2 field is not null, and the field's length is not exactly 16 bytes, then this command SHALL fail with a CONSTRAINT_ERROR status code responded to the client.
- If the EpochKey2 field is not null, its associated EpochStartTime2 field SHALL NOT be null and SHALL contain a later epoch start time than the epoch start time found in the EpochStartTime1 field. Otherwise this command SHALL fail with an INVALID_COMMAND status code responded to the client.
- If exactly one of the EpochKey2 or EpochStartTime2 is null, rather than both being null, or neither being null, then this command SHALL fail with an INVALID_COMMAND status code responded to the client.

If there exists a Group Key Set associated with the accessing fabric which has the same GroupKeySetID as that provided in the GroupKeySet field, then the contents of that group key set SHALL be replaced. A replacement SHALL be done by executing the equivalent of entirely removing the previous Group Key Set with the given GroupKeySetID, followed by an addition of a Group Key Set with the provided configuration. Otherwise, if the GroupKeySetID did not match an existing entry, a new Group Key Set associated with the accessing fabric SHALL be created with the provided data. The Group Key Set SHALL be written to non-volatile storage.

Upon completion, this command SHALL send a status code back to the initiator:

- If the Group Key Set was properly installed or updated on the Node, the status code SHALL be set to SUCCESS.
- If there are insufficient resources on the receiver to store an additional Group Key Set, the status code SHALL be set to RESOURCE_EXHAUSTED (see [Section 2.11.1.2, "Group Limits"](#));
- Otherwise, this status code SHALL be set to FAILURE.

11.2.7.2. KeySetRead Command

This command is used by Administrators to read the state of a given Group Key Set.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupKey-SetID	uint16	all			M

Effect on Receipt

If there exists a Group Key Set associated with the accessing fabric which has the same GroupKey-SetID as that provided in the GroupKeySetID field, then the contents of that Group Key Set SHALL be sent in a [KeySetReadResponse](#) command, but with the EpochKey0, EpochKey1 and EpochKey2 fields replaced by null.

Otherwise, if the GroupKeySetID does not refer to a Group Key Set associated with the accessing fabric, then this command SHALL fail with a NOT_FOUND status code.

11.2.7.3. KeySetReadResponse Command

This command SHALL be generated in response to the [KeySetRead](#) command, if a valid Group Key Set was found. It SHALL contain the configuration of the requested Group Key Set, with the EpochKey0, EpochKey1 and EpochKey2 key contents replaced by null.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupKey-Set	GroupKey-SetStruct				M

11.2.7.4. KeySetRemove Command

This command is used by Administrators to remove all state of a given Group Key Set.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupKey-SetID	uint16	all			M

Effect on Receipt

If there exists a Group Key Set associated with the accessing fabric which has the same GroupKey-SetID as that provided in the GroupKeySetID field, then the contents of that Group Key Set SHALL be removed, including all epoch keys it contains.

If there exist any entries for the accessing fabric within the [GroupKeyMap](#) attribute that refer to the GroupKeySetID just removed, then these entries SHALL be removed from that list.

This command SHALL fail with an INVALID_COMMAND status code back to the initiator if the GroupKeySetID being removed is 0, which is the Key Set associated with the [Identity Protection Key \(IPK\)](#). The only method to remove the IPK is usage of the [RemoveFabric](#) command or any operation which causes the equivalent of a RemoveFabric to occur by side-effect.

This command SHALL send a SUCCESS status code back to the initiator on success, or NOT_FOUND if the GroupKeySetID requested did not exist.

11.2.7.5. KeySetReadAllIndices Command

This command is used by Administrators to query a list of all Group Key Sets associated with the accessing fabric.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	DoNotUse					X

Effect on Receipt

Upon receipt, this command SHALL iterate all stored GroupKeySetStruct associated with the accessing fabric and generate a [KeySetReadAllIndicesResponse](#) command containing the list of GroupKeySetID values from those structs.

11.2.7.6. KeySetReadAllIndicesResponse Command

This command SHALL be generated in response to [KeySetReadAllIndices](#) and it SHALL contain the list of [GroupKeySetID](#) for all Group Key Sets associated with the scoped Fabric.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupKey-SetIDs	list[uint16]				M

GroupKeySetIDs Field

This field references the set of group keys that generate operational group keys for use with the accessing fabric.

Each entry in GroupKeySetIDs is a [GroupKeySetID](#) field.

11.3. Localization Configuration Cluster

Nodes should be expected to be deployed to any and all regions of the world. These global regions may have differing common languages, units of measurements, and numerical formatting standards. As such, Nodes that visually or audibly convey information need a mechanism by which they can be configured to use a user's preferred language, units, etc.

This cluster supports an interface to a Node. It provides attributes for determining and configuring localization information that a Node SHALL utilize when conveying values to a user.

11.3.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table

below.

Revision	Description
1	Initial revision

11.3.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	LCFG

11.3.3. Cluster ID

ID	Name
0x002B	Localization Configuration

11.3.4. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	ActiveLocale	string	max 35	N	MS	RW VM	M
0x0001	SupportedLocales	list[string]	max 32[max 35]	F	MS	R V	M

11.3.4.1. ActiveLocale Attribute

The ActiveLocale attribute SHALL represent the locale that the Node is currently configured to use when conveying information. The ActiveLocale attribute SHALL be a Language Tag as defined by [BCP47](#). The ActiveLocale attribute SHALL have a default value assigned by the Vendor and SHALL be a value contained within the SupportedLocales attribute.

An attempt to write a value to ActiveLocale that is not present in SupportedLocales SHALL result in a CONSTRAINT_ERROR error.

11.3.4.2. SupportedLocales Attribute

The SupportedLocales attribute SHALL represent a list of locale strings that are valid values for the [ActiveLocale](#) attribute. The list SHALL NOT contain any duplicate entries. The ordering of items within the list SHOULD NOT express any meaning.

11.4. Time Format Localization Cluster

Nodes should be expected to be deployed to any and all regions of the world. These global regions may have differing preferences for how dates and times are conveyed. As such, Nodes that visually or audibly convey time information need a mechanism by which they can be configured to use a

user's preferred format.

This cluster supports an interface to a Node. It provides attributes for determining and configuring time and date formatting information that a Node SHALL utilize when conveying values to a user.

11.4.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.4.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	LTIME

11.4.3. Cluster ID

ID	Name
0x002C	Time Format Localization

11.4.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Summary
0	CALFMT	CalendarFormat	The Node can be configured to use different calendar formats when conveying values to a user.

11.4.5. Data Types

11.4.5.1. HourFormatEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	12hr	Time conveyed with a 12-hour clock	M
1	24hr	Time conveyed with a 24-hour clock	M

Value	Name	Summary	Conformance
255	UseActiveLocale	Use active locale clock	M

11.4.5.2. CalendarTypeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Buddhist	Dates conveyed using the Buddhist calendar	O.a+
1	Chinese	Dates conveyed using the Chinese calendar	O.a+
2	Coptic	Dates conveyed using the Coptic calendar	O.a+
3	Ethiopian	Dates conveyed using the Ethiopian calendar	O.a+
4	Gregorian	Dates conveyed using the Gregorian calendar	O.a+
5	Hebrew	Dates conveyed using the Hebrew calendar	O.a+
6	Indian	Dates conveyed using the Indian calendar	O.a+
7	Islamic	Dates conveyed using the Islamic calendar	O.a+
8	Japanese	Dates conveyed using the Japanese calendar	O.a+
9	Korean	Dates conveyed using the Korean calendar	O.a+
10	Persian	Dates conveyed using the Persian calendar	O.a+
11	Taiwanese	Dates conveyed using the Taiwanese calendar	O.a+
255	UseActiveLocale	calendar implied from active locale	O.a+

11.4.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	HourFormat	HourFormatEnum	all	N		RW VM	M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0001	ActiveCalendarType	CalendarTypeEnum	all	N		RW VM	CALFMT
0x0002	SupportedCalendarTypes	list[CalendarTypeEnum]	desc	F	N/A	R V	CALFMT

11.4.6.1. HourFormat Attribute

This attribute SHALL represent the format that the Node is currently configured to use when conveying the hour unit of time.

If not UseActiveLocale, this value SHALL take priority over any unit implied through the [ActiveLocale](#) attribute.

If UseActiveLocale, any unit implied through the [ActiveLocale](#) attribute is used as the hour format, and if [ActiveLocale](#) is not present, the hour format is unknown.

11.4.6.2. ActiveCalendarType Attribute

This attribute SHALL represent the calendar format that the Node is currently configured to use when conveying dates.

If not UseActiveLocale, this value SHALL take priority over any unit implied through the [ActiveLocale](#) attribute.

If UseActiveLocale, any unit implied through the [ActiveLocale](#) attribute is used as the calendar type, and if [ActiveLocale](#) is not present, the calendar type is unknown.

11.4.6.3. SupportedCalendarTypes Attribute

This attribute SHALL represent a list of [CalendarTypeEnum](#) values that are supported by the Node. The list SHALL NOT contain any duplicate entries. The ordering of items within the list SHOULD NOT express any meaning. The maximum length of the SupportedCalendarTypes list SHALL be equivalent to the number of enumerations within [CalendarTypeEnum](#).

11.5. Unit Localization Cluster

Nodes should be expected to be deployed to any and all regions of the world. These global regions may have differing preferences for the units in which values are conveyed in communication to a user. As such, Nodes that visually or audibly convey measurable values to the user need a mechanism by which they can be configured to use a user's preferred unit.

This cluster supports an interface to a Node. It provides attributes for determining and configuring the units that a Node SHALL utilize when conveying values in communication to a user.

11.5.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table

below.

Revision	Description
1	Initial revision
2	Added SupportedTemperatureUnits attribute

11.5.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	LUNIT

11.5.3. Cluster ID

ID	Name
0x002D	Unit Localization

11.5.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Summary
0	TEMP	TemperatureUnit	The Node can be configured to use different units of temperature when conveying values to a user.

11.5.5. Data Types

11.5.5.1. TempUnitEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Fahrenheit	Temperature conveyed in Fahrenheit	O.a2+
1	Celsius	Temperature conveyed in Celsius	O.a2+
2	Kelvin	Temperature conveyed in Kelvin	O.a2+

11.5.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	TemperatureUnit	TempUnitEnum	all	N		RW VM	TEMP
0x0001	SupportedTemperatureUnits	list[TempUnitEnum]	2 to 3	F		R V	TEMP

11.5.6.1. TemperatureUnit Attribute

This attribute SHALL indicate the unit for the Node to use only when conveying temperature in communication to the user, for example such as via a user interface on the device. If provided, this value SHALL take priority over any unit implied through the [ActiveLocale Attribute](#).

An attempt to write to this attribute with a value not included in the SupportedTemperatureUnits attribute list SHALL result in a CONSTRAINT_ERROR.

11.5.6.2. SupportedTemperatureUnits Attribute

This attribute SHALL indicate a list of units supported by the Node to be used when writing the TemperatureUnit attribute of this cluster. Each entry in the list SHALL be unique.

11.6. Power Source Configuration Cluster

This cluster is used to describe the configuration and capabilities of a Device's power system. It provides an ordering overview as well as linking to the one or more endpoints each supporting a Power Source cluster.

11.6.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.6.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	PSCFG

11.6.3. Cluster ID

ID	Name
0x002E	Power Source Configuration

11.6.4. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Sources	list[end-point-no]	max 6	N		R V	M

11.6.4.1. Sources Attribute

This list SHALL contain the set of all power sources capable of participating in the power system of this Node. Each entry in the list SHALL be the endpoint number of an endpoint having a Power Source cluster, which corresponds to a physical power source. The endpoint number SHALL be unique within the list.

The order of power sources on a Node is defined by the [Order](#) attribute of its associated Power Source cluster provided on the endpoint. List entries SHALL be sorted in increasing order, that is, an entry with a lower order SHALL have a lower index than any entry with a higher order. Multiple entries MAY have the same order, there are no restrictions on their relative sorting.

11.7. Power Source Cluster

This cluster is used to describe the configuration and capabilities of a physical power source that provides power to one or more endpoints on a node. In case the node has multiple power sources, each SHALL be described by its own cluster instance. Each instance of this cluster MAY be associated with one or more endpoints or the entire node.

11.7.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Added EndpointList attribute that maps a power source to a list of endpoints
3	Added Q quality (replacing C quality) on BatPercentRemaining, BatTimeRemaining and BatTimeToFullCharge

11.7.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	PS

11.7.3. Cluster ID

ID	Name
0x002F	Power Source

11.7.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Conformance	Summary
0	WIRED	Wired	O.a	A wired power source
1	BAT	Battery	O.a	A battery power source
2	RECHG	Rechargeable	[BAT]	A rechargeable battery power source
3	REPLC	Replaceable	[BAT]	A replaceable battery power source

11.7.5. Dependencies

This cluster SHOULD be on an endpoint that supports a device type that requires this cluster. This cluster SHOULD NOT be just added as an extra cluster to an endpoint to conserve endpoint instances.

11.7.6. Data Types

11.7.6.1. WiredFaultEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	Unspecified	The Node detects an unspecified fault on this wired power source.	M

Value	Name	Summary	Conformance
1	OverVoltage	The Node detects the supplied voltage is above maximum supported value for this wired power source.	M
2	UnderVoltage	The Node detects the supplied voltage is below maximum supported value for this wired power source.	M

11.7.6.2. BatFaultEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	Unspecified	The Node detects an unspecified fault on this battery power source.	M
1	OverTemp	The Node detects the temperature of this battery power source is above ideal operating conditions.	M
2	UnderTemp	The Node detects the temperature of this battery power source is below ideal operating conditions.	M

11.7.6.3. BatChargeFaultEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	Unspecified	The Node detects an unspecified fault on this battery source.	M
1	AmbientTooHot	The Node detects the ambient temperature is above the nominal range for this battery source.	M

Value	Name	Summary	Conformance
2	AmbientTooCold	The Node detects the ambient temperature is below the nominal range for this battery source.	M
3	BatteryTooHot	The Node detects the temperature of this battery source is above the nominal range.	M
4	BatteryTooCold	The Node detects the temperature of this battery source is below the nominal range.	M
5	BatteryAbsent	The Node detects this battery source is not present.	M
6	BatteryOverVoltage	The Node detects this battery source is over voltage.	M
7	BatteryUnderVoltage	The Node detects this battery source is under voltage.	M
8	ChargerOverVoltage	The Node detects the charger for this battery source is over voltage.	M
9	ChargerUnderVoltage	The Node detects the charger for this battery source is under voltage.	M
10	SafetyTimeout	The Node detects a charging safety timeout for this battery source.	M

11.7.6.4. PowerSourceStatusEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	Unspecified	Indicate the source status is not specified	M
1	Active	Indicate the source is available and currently supplying power	M

Value	Name	Summary	Conformance
2	Standby	Indicate the source is available, but is not currently supplying power	M
3	Unavailable	Indicate the source is not currently available to supply power	M

11.7.6.5. WiredCurrentTypeEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	AC	Indicates AC current	M
1	DC	Indicates DC current	M

11.7.6.6. BatChargeLevelEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	OK	Charge level is nominal	M
1	Warning	Charge level is low, intervention may soon be required.	M
2	Critical	Charge level is critical, immediate intervention is required	M

11.7.6.7. BatReplaceabilityEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	Unspecified	The replaceability is unspecified or unknown.	M
1	NotReplaceable	The battery is not replaceable.	M
2	UserReplaceable	The battery is replaceable by the user or customer.	M

Value	Name	Summary	Conformance
3	FactoryReplaceable	The battery is replaceable by an authorized factory technician.	M

11.7.6.8. BatCommonDesignationEnum Type

This data type is derived from enum16.

Value	Name	Summary	Conformance
0	Unspecified	Common type is unknown or unspecified	M
1	AAA	Common type is as specified	M
2	AA	Common type is as specified	M
3	C	Common type is as specified	M
4	D	Common type is as specified	M
5	4v5	Common type is as specified	M
6	6v0	Common type is as specified	M
7	9v0	Common type is as specified	M
8	1_2AA	Common type is as specified	M
9	AAAA	Common type is as specified	M
10	A	Common type is as specified	M
11	B	Common type is as specified	M
12	F	Common type is as specified	M
13	N	Common type is as specified	M
14	No6	Common type is as specified	M

Value	Name	Summary	Conformance
15	SubC	Common type is as specified	M
16	A23	Common type is as specified	M
17	A27	Common type is as specified	M
18	BA5800	Common type is as specified	M
19	Duplex	Common type is as specified	M
20	4SR44	Common type is as specified	M
21	523	Common type is as specified	M
22	531	Common type is as specified	M
23	15v0	Common type is as specified	M
24	22v5	Common type is as specified	M
25	30v0	Common type is as specified	M
26	45v0	Common type is as specified	M
27	67v5	Common type is as specified	M
28	J	Common type is as specified	M
29	CR123A	Common type is as specified	M
30	CR2	Common type is as specified	M
31	2CR5	Common type is as specified	M
32	CR_P2	Common type is as specified	M
33	CR_V3	Common type is as specified	M

Value	Name	Summary	Conformance
34	SR41	Common type is as specified	M
35	SR43	Common type is as specified	M
36	SR44	Common type is as specified	M
37	SR45	Common type is as specified	M
38	SR48	Common type is as specified	M
39	SR54	Common type is as specified	M
40	SR55	Common type is as specified	M
41	SR57	Common type is as specified	M
42	SR58	Common type is as specified	M
43	SR59	Common type is as specified	M
44	SR60	Common type is as specified	M
45	SR63	Common type is as specified	M
46	SR64	Common type is as specified	M
47	SR65	Common type is as specified	M
48	SR66	Common type is as specified	M
49	SR67	Common type is as specified	M
50	SR68	Common type is as specified	M
51	SR69	Common type is as specified	M
52	SR516	Common type is as specified	M

Value	Name	Summary	Conformance
53	SR731	Common type is as specified	M
54	SR712	Common type is as specified	M
55	LR932	Common type is as specified	M
56	A5	Common type is as specified	M
57	A10	Common type is as specified	M
58	A13	Common type is as specified	M
59	A312	Common type is as specified	M
60	A675	Common type is as specified	M
61	AC41E	Common type is as specified	M
62	10180	Common type is as specified	M
63	10280	Common type is as specified	M
64	10440	Common type is as specified	M
65	14250	Common type is as specified	M
66	14430	Common type is as specified	M
67	14500	Common type is as specified	M
68	14650	Common type is as specified	M
69	15270	Common type is as specified	M
70	16340	Common type is as specified	M
71	RCR123A	Common type is as specified	M

Value	Name	Summary	Conformance
72	17500	Common type is as specified	M
73	17670	Common type is as specified	M
74	18350	Common type is as specified	M
75	18500	Common type is as specified	M
76	18650	Common type is as specified	M
77	19670	Common type is as specified	M
78	25500	Common type is as specified	M
79	26650	Common type is as specified	M
80	32600	Common type is as specified	M

11.7.6.9. BatApprovedChemistryEnum Type

This data type is derived from enum16.

Value	Name	Summary	Conformance
0	Unspecified	Cell chemistry is unspecified or unknown	M
1	Alkaline	Cell chemistry is alkaline	M
2	LithiumCarbonFluoride	Cell chemistry is lithium carbon fluoride	M
3	LithiumChromiumOxide	Cell chemistry is lithium chromium oxide	M
4	LithiumCopperOxide	Cell chemistry is lithium copper oxide	M
5	LithiumIronDisulfide	Cell chemistry is lithium iron disulfide	M

Value	Name	Summary	Conformance
6	LithiumManganeseDioxide	Cell chemistry is lithium manganese dioxide	M
7	LithiumThionylChloride	Cell chemistry is lithium thionyl chloride	M
8	Magnesium	Cell chemistry is magnesium	M
9	MercuryOxide	Cell chemistry is mercury oxide	M
10	NickelOxyhydride	Cell chemistry is nickel oxyhydride	M
11	SilverOxide	Cell chemistry is silver oxide	M
12	ZincAir	Cell chemistry is zinc air	M
13	ZincCarbon	Cell chemistry is zinc carbon	M
14	ZincChloride	Cell chemistry is zinc chloride	M
15	ZincManganeseDioxide	Cell chemistry is zinc manganese dioxide	M
16	LeadAcid	Cell chemistry is lead acid	M
17	LithiumCobaltOxide	Cell chemistry is lithium cobalt oxide	M
18	LithiumIon	Cell chemistry is lithium ion	M
19	LithiumIonPolymer	Cell chemistry is lithium ion polymer	M
20	LithiumIronPhosphate	Cell chemistry is lithium iron phosphate	M
21	LithiumSulfur	Cell chemistry is lithium sulfur	M
22	LithiumTitanate	Cell chemistry is lithium titanate	M
23	NickelCadmium	Cell chemistry is nickel cadmium	M
24	NickelHydrogen	Cell chemistry is nickel hydrogen	M

Value	Name	Summary	Conformance
25	NickelIron	Cell chemistry is nickel iron	M
26	NickelMetalHydride	Cell chemistry is nickel metal hydride	M
27	NickelZinc	Cell chemistry is nickel zinc	M
28	SilverZinc	Cell chemistry is silver zinc	M
29	SodiumIon	Cell chemistry is sodium ion	M
30	SodiumSulfur	Cell chemistry is sodium sulfur	M
31	ZincBromide	Cell chemistry is zinc bromide	M
32	ZincCerium	Cell chemistry is zinc cerium	M

11.7.6.10. BatChargeStateEnum Type

This data type is derived from enum8.

Value	Name	Summary	Conformance
0	Unknown	Unable to determine the charging state	M
1	IsCharging	The battery is charging	M
2	IsAtFullCharge	The battery is at full charge	M
3	IsNotCharging	The battery is not charging	M

11.7.7. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Status	Power-SourceStatusEnum	desc			R V	M
0x0001	Order	uint8	all	N		R V	M
0x0002	Description	string	max 60	F		R V	M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0003	WiredAssessedInputVoltage	uint32	all	C X		R V	[WIRED]
0x0004	WiredAssessedInputFrequency	uint16	all	C X		R V	[WIRED]
0x0005	WiredCurrentType	WiredCurrentTypeEnum	desc	F		R V	WIRED
0x0006	WiredAssessedCurrent	uint32	all	C X		R V	[WIRED]
0x0007	Wired-Nominal-Voltage	uint32	all	F		R V	[WIRED]
0x0008	Wired-MaximumCurrent	uint32	all	F		R V	[WIRED]
0x0009	WiredPresent	bool	all			R V	[WIRED]
0x000A	ActiveWiredFaults	list[WiredFaultEnum]	max 8			R V	[WIRED]
0x000B	BatVoltage	uint32	all	C X		R V	[BAT]
0x000C	BatPercentRemaining	uint8	max 200	Q X		R V	[BAT]
0x000D	Bat-TimeRemaining	uint32	all	Q X		R V	[BAT]
0x000E	BatCharge Level	BatChargeLevelEnum	desc			R V	BAT
0x000F	BatReplacement-Needed	bool	all			R V	BAT
0x0010	BatReplaceability	BatReplaceabilityEnum	all	F		R V	BAT

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0011	BatPresent	bool	all			R V	[BAT]
0x0012	ActiveBatFaults	list[BatFaultEnum]	max 8			R V	[BAT]
0x0013	BatReplacementDescription	string	max 60	F		R V	REPLC
0x0014	BatCommonDesignation	BatCommonDesignationEnum	desc	F		R V	[REPLC]
0x0015	BatANSIDesignation	string	max 20	F		R V	[REPLC]
0x0016	BatIECDesignation	string	max 20	F		R V	[REPLC]
0x0017	BatApprovedChemistry	BatApprovedChemistryEnum	desc	F		R V	[REPLC]
0x0018	BatCapacity	uint32	all	F		R V	[REPLC RECHG]
0x0019	BatQuantity	uint8	all	F		R V	REPLC
0x001A	BatChargeState	BatChargeStateEnum	desc			R V	RECHG
0x001B	BatTimeToFullCharge	uint32	all	Q X		R V	[RECHG]
0x001C	BatFunctionalWhileCharging	bool	all			R V	RECHG
0x001D	BatChargingCurrent	uint32	all	C X		R V	[RECHG]

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x001E	Active-BatChargeFaults	list[BatChargeFaultEnum]	max 16			R V	[RECHG]
0x001F	EndpointList	list[endpoint-no]				R V	M

11.7.7.1. Status Attribute

This attribute SHALL indicate the participation of this power source in providing power to the Node as specified in [PowerSourceStatusEnum](#).

11.7.7.2. Order Attribute

This attribute SHALL indicate the relative preference with which the Node will select this source to provide power. A source with a lower order SHALL be selected by the Node to provide power before any other source with a higher order, if the lower order source is available (see [Section 11.7.7.1, “Status Attribute”](#)).

Note, Order is read-only and therefore NOT intended to allow clients control over power source selection.

11.7.7.3. Description Attribute

This attribute SHALL provide a user-facing description of this source, used to distinguish it from other power sources, e.g. "DC Power", "Primary Battery" or "Battery back-up". This attribute SHALL NOT be used to convey information such as battery form factor, or chemistry.

11.7.7.4. WiredAssessedInputVoltage Attribute

This attribute SHALL indicate the assessed RMS or DC voltage currently provided by the hard-wired source, in mV (millivolts). A value of NULL SHALL indicate the Node is currently unable to assess the value. If the wired source is not connected, but the Node is still able to assess a value, then the assessed value MAY be reported.

11.7.7.5. WiredAssessedInputFrequency Attribute

This attribute SHALL indicate the assessed frequency of the voltage, currently provided by the hard-wired source, in Hz. A value of NULL SHALL indicate the Node is currently unable to assess the value. If the wired source is not connected, but the Node is still able to assess a value, then the assessed value MAY be reported.

11.7.7.6. WiredCurrentType Attribute

This attribute SHALL indicate the type of current the Node expects to be provided by the hard-wired source as specified in [WiredCurrentTypeEnum](#).

11.7.7.7. **WiredAssessedCurrent Attribute**

This attribute SHALL indicate the assessed instantaneous current draw of the Node on the hard-wired source, in mA (milliamps). A value of NULL SHALL indicate the Node is currently unable to assess the value. If the wired source is not connected, but the Node is still able to assess a value, then the assessed value MAY be reported.

11.7.7.8. **WiredNominalVoltage Attribute**

This attribute SHALL indicate the nominal voltage, printed as part of the Node's regulatory compliance label in mV (millivolts), expected to be provided by the hard-wired source.

11.7.7.9. **WiredMaximumCurrent Attribute**

This attribute SHALL indicate the maximum current, printed as part of the Node's regulatory compliance label in mA (milliamps), expected to be provided by the hard-wired source.

11.7.7.10. **WiredPresent Attribute**

This attribute SHALL indicate if the Node detects that the hard-wired power source is properly connected.

11.7.7.11. **ActiveWiredFaults Attribute**

This attribute SHALL indicate the set of wired faults currently detected by the Node on this power source. This set is represented as a list of [WiredFaultEnum](#). When the Node detects a fault has been raised, the appropriate [WiredFaultEnum](#) value SHALL be added to this list, provided it is not already present. This list SHALL NOT contain more than one instance of a specific [WiredFaultEnum](#) value. When the Node detects all conditions contributing to a fault have been cleared, the corresponding [WiredFaultEnum](#) value SHALL be removed from this list. An empty list SHALL indicate there are currently no active faults. The order of this list SHOULD have no significance. Clients interested in monitoring changes in active faults MAY subscribe to this attribute, or they MAY subscribe to [WiredFaultChange](#).

11.7.7.12. **BatVoltage Attribute**

This attribute SHALL indicate the currently measured output voltage of the battery in mV (millivolts). A value of NULL SHALL indicate the Node is currently unable to assess the value.

11.7.7.13. **BatPercentRemaining Attribute**

This attribute SHALL indicate the estimated percentage of battery charge remaining until the battery will no longer be able to provide power to the Node. Values are expressed in half percent units, ranging from 0 to 200. E.g. a value of 48 is equivalent to 24%. A value of NULL SHALL indicate the Node is currently unable to assess the value.

Changes to this attribute SHALL only be marked as reportable in the following cases:

- At most once every 10 seconds, or
- When it changes from null to any other value and vice versa.

Since reporting consumes power, devices SHOULD be careful not to over-report.

11.7.7.14. BatTimeRemaining Attribute

This attribute SHALL indicate the estimated time in seconds before the battery will no longer be able to provide power to the Node. A value of NULL SHALL indicate the Node is currently unable to assess the value.

Changes to this attribute SHALL only be marked as reportable in the following cases:

- At most once every 10 seconds, or
- When it changes from null to any other value and vice versa.

Since reporting consumes power, devices SHOULD be careful not to over-report.

11.7.7.15. BatChargeLevel Attribute

This attribute SHALL indicate a coarse ranking of the charge level of the battery, used to indicate when intervention is required as specified in [BatChargeLevelEnum](#).

11.7.7.16. BatReplacementNeeded Attribute

This attribute SHALL indicate if the battery needs to be replaced. Replacement MAY be simple routine maintenance, such as with a single use, non-rechargeable cell. Replacement, however, MAY also indicate end of life, or serious fault with a rechargeable or even non-replaceable cell.

11.7.7.17. BatReplaceability Attribute

This attribute SHALL indicate the replaceability of the battery as specified in [BatReplaceabilityEnum](#).

11.7.7.18. BatPresent Attribute

This attribute SHALL indicate whether the Node detects that the batteries are properly installed.

11.7.7.19. ActiveBatFaults Attribute

This attribute SHALL indicate the set of battery faults currently detected by the Node on this power source. This set is represented as a list of [BatFaultEnum](#). When the Node detects a fault has been raised, the appropriate [BatFaultEnum](#) value SHALL be added to this list, provided it is not already present. This list SHALL NOT contain more than one instance of a specific [BatFaultEnum](#) value. When the Node detects all conditions contributing to a fault have been cleared, the corresponding [BatFaultEnum](#) value SHALL be removed from this list. An empty list SHALL indicate there are currently no active faults. The order of this list SHOULD have no significance. Clients interested in monitoring changes in active faults MAY subscribe to this attribute, or they MAY subscribe to [BatFaultChange](#).

11.7.7.20. BatReplacementDescription Attribute

This attribute SHALL provide a user-facing description of this battery, which SHOULD contain

information required to identify a replacement, such as form factor, chemistry or preferred manufacturer.

11.7.7.21. BatCommonDesignation Attribute

This attribute SHALL indicate the ID of the common or colloquial designation of the battery, as specified in [BatCommonDesignationEnum](#).

11.7.7.22. BatANSIDesignation Attribute

This attribute SHALL indicate the string representing the ANSI designation for the battery as specified in [ANSI C18](#).

11.7.7.23. BatIECDesignation Attribute

This attribute SHALL indicate the string representing the IEC designation for the battery as specified in [IEC 60086](#).

11.7.7.24. BatApprovedChemistry Attribute

This attribute SHALL indicate the ID of the preferred chemistry of the battery source as specified in [BatApprovedChemistryEnum](#).

11.7.7.25. BatCapacity Attribute

This attribute SHALL indicate the preferred minimum charge capacity rating in mAh of individual, user- or factory-serviceable battery cells or packs in the battery source.

11.7.7.26. BatQuantity Attribute

This attribute SHALL indicate the quantity of individual, user- or factory-serviceable battery cells or packs in the battery source.

11.7.7.27. BatChargeState Attribute

This attribute SHALL indicate the current state of the battery source with respect to charging as specified in [BatChargeStateEnum](#).

11.7.7.28. BatTimeToFullCharge Attribute

This attribute SHALL indicate the estimated time in seconds before the battery source will be at full charge. A value of NULL SHALL indicate the Node is currently unable to assess the value.

Changes to this attribute SHALL only be marked as reportable in the following cases:

- At most once every 10 seconds, or
- When it changes from null to any other value and vice versa.

Since reporting consumes power, devices SHOULD be careful not to over-report.

11.7.7.29. BatFunctionalWhileCharging Attribute

This attribute SHALL indicate whether the Node can remain operational while the battery source is charging.

11.7.7.30. BatChargingCurrent Attribute

This attribute SHALL indicate assessed current in mA (milliamps) presently supplied to charge the battery source. A value of NULL SHALL indicate the Node is currently unable to assess the value.

11.7.7.31. ActiveBatChargeFaults Attribute

This attribute SHALL indicate the set of charge faults currently detected by the Node on this power source. This set is represented as a list of [BatChargeFaultEnum](#). When the Node detects a fault has been raised, the appropriate [BatChargeFaultEnum](#) value SHALL be added to this list, provided it is not already present. This list SHALL NOT contain more than one instance of a specific [BatChargeFaultEnum](#) value. When the Node detects all conditions contributing to a fault have been cleared, the corresponding [BatChargeFaultEnum](#) value SHALL be removed from this list. An empty list SHALL indicate there are currently no active faults. The order of this list SHOULD have no significance. Clients interested in monitoring changes in active faults MAY subscribe to this attribute, or they MAY subscribe to the [BatFaultChange event](#).

11.7.7.32. EndpointList Attribute

This attribute SHALL indicate a list of endpoints that are powered by the source defined by this cluster. Multiple instances of this cluster MAY list the same endpoint, because it is possible for power for an endpoint to come from multiple sources. In that case the Order attribute indicates their priority.

For each power source on a node, there SHALL only be one instance of this cluster.

A cluster instance with an empty list SHALL indicate that the power source is for the entire node, which includes all endpoints.

A cluster instance with a non-empty list SHALL include the endpoint, upon which the cluster instance resides.

The above rules allow that some endpoints can have an unknown power source, and therefore would not be indicated by any instance of this cluster.

Legacy Implementations

Legacy implementations of this cluster before revision 2, before this attribute was defined, would have implemented this cluster on an application endpoint without indicating it in EndpointList (since that attribute did not exist in revision 1), because it represented a power source for the endpoint, not the entire node.

For example: Bridge implementations support endpoints for bridged devices that have different power sources.

Such implementations followed device type requirements and semantics outside of this cluster, because this attribute did not exist.

Future updates of such a cluster instance on the same endpoint, would allow that same endpoint to be an entry in the EndpointList attribute. Therefore it is valid to list the endpoint upon which the cluster instance exists.

Empty list examples

Typically, there is one power source for the node. Also common is mains power for the node with battery backup power for the node. In both these common cases, for each cluster instance described, the list is empty.

Populated list example

A node has a mains power source with Order as 0 (zero), but some application endpoints (not all) have a battery back up source with Order as 1, which means this list is empty for the Power Source cluster associated with the mains power, because it indicates the entire node, but the Power Source cluster instance associated with the battery backup would list the endpoints that have a battery backup.

11.7.8. Events

ID	Name	Priority	Quality	Access	Conformance
0x00	WiredFault-Change	INFO		V	[WIRED]
0x01	BatFault-Change	INFO		V	[BAT]
0x02	BatCharge-FaultChange	INFO		V	[RECHG]

11.7.8.1. WiredFaultChange Event

The WiredFaultChange Event SHALL be generated when the set of wired faults currently detected by the Node on this wired power source changes. This event SHALL correspond to a change in value of [ActiveWiredFaults](#).

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Current	list[Wired-FaultEnum]	max 8		empty	M
1	Previous	list[Wired-FaultEnum]	max 8		empty	M

Current Field

This field SHALL represent the set of faults currently detected, as per [ActiveWiredFaults](#).

Previous Field

This field SHALL represent the set of faults detected prior to this change event, as per [ActiveWired-Faults](#).

11.7.8.2. BatFaultChangeEvent

The BatFaultChangeEvent SHALL be generated when the set of battery faults currently detected by the Node on this battery power source changes. This event SHALL correspond to a change in value of [ActiveBatFaults](#).

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Current	list[BatFaultEnum]	max 8		empty	M
1	Previous	list[BatFaultEnum]	max 8		empty	M

Current Field

This field SHALL represent the set of faults currently detected, as per [ActiveBatFaults](#).

Previous Field

This field SHALL represent the set of faults detected prior to this change event, as per [ActiveBatFaults](#).

11.7.8.3. BatChargeFaultChangeEvent

The BatChargeFaultChangeEvent SHALL be generated when the set of charge faults currently detected by the Node on this battery power source changes. This event SHALL correspond to a change in value of [ActiveBatChargeFaults](#).

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Current	list[BatChargeFaultEnum]	max 16		empty	M
1	Previous	list[BatChargeFaultEnum]	max 16		empty	M

Current Field

This field SHALL represent the set of faults currently detected, as per [ActiveBatChargeFaults](#).

Previous Field

This field SHALL represent the set of faults detected prior to this change event, as per [ActiveBatChargeFaults](#).

11.7.9. Configuration Examples

The following examples illustrate use of the Order attribute in the Power Source Cluster.

11.7.9.1. Example: Redundant Mains Power Sources

This example describes a design with symmetric, dual-redundant mains power sources, where the system is powered by either one of the power sources. The Node must define two Power Source Clusters, one for each mains source. The system indicates no preference for either source, so the sources have the same Order.

Power Source (on Endpoint 2)

Description: "Mains A"
Order: 0

Power Source (on Endpoint 5)

Description: "Mains B"
Order: 0

11.7.9.2. Example: Battery with Mains Power Back-up

This example describes a design with a built-in battery as the primary source, where the mains power serves to keep the battery charged and act as back-up if the battery fails. The Node must define two Power Source Clusters, one for the battery and another for the mains. Since the battery is primary, it must have a lower Order than the mains source.

Power Source (on Endpoint 2)

Description: "Primary Battery"
Order: 0

Power Source (on Endpoint 5)

Description: "Mains"
Order: 1

11.7.9.3. Example: Mains Power with Battery Back-up

This example describes a design where the system always runs from the a mains power source and the back-up battery is out-of-circuit until mains power fails at which point the back-up battery powers the system. The Node must define two Power Source Clusters, one for the mains and another for the battery. Since the mains source is primary, it must have a lower Order than the battery source.

Power Source (on Endpoint 2)

Description: "Mains"

Order: 0

Power Source (on Endpoint 5)

Description: "Battery Back-up"
Order: 1

11.7.9.4. Example: Battery with Dual Back-up

This example describes a design with a built-in battery as the primary source, and where two wired sources, USB and a DC adapter, redundantly serve to keep the battery charged and act as back-up if the battery fails. The Node must define three Power Source Clusters, one for each of the battery, the USB source, and the DC adapter. Since the battery is primary, the battery source must have a lower Order than the other sources. This system has no preference between the DC Adapter and USB sources, so these sources will have the same Order.

Power Source (on Endpoint 2)

Description: "Primary Battery"
Order: 0

Power Source (on Endpoint 5)

Description: "DC Adapter"
Order: 1

Power Source (on Endpoint 7)

Description: "USB Power"
Order: 1

11.7.9.5. Example: Mains Power with Battery Powered Peripheral

This example describes a design with a mains powered core and battery powered peripheral. In this example both power sources are required for proper operation. The Node must define two Power Source Clusters, one for the wired source and one for the battery. Since both sources are required, both sources will have the same Order. We will use Endpoint 2 for the mains power and Endpoint 7 for the battery.

Power Source (on Endpoint 2)

Description: "Mains Power"
Order: 0

Power Source (on Endpoint 7)

Description: "Peripheral Battery"

Order: 0

11.8. Power Topology Cluster

The Power Topology Cluster provides a mechanism for expressing how power is flowing between endpoints.

11.8.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.8.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Application	Endpoint	PWRTL

11.8.3. Cluster ID

ID	Name
0x009C	Power Topology

11.8.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Conformance	Summary
0	NODE	NodeTopology	O.a	This endpoint provides or consumes power to/from the entire node
1	TREE	TreeTopology	O.a	This endpoint provides or consumes power to/from itself and its child endpoints
2	SET	SetTopology	O.a	This endpoint provides or consumes power to/from a specified set of endpoints

Bit	Code	Feature	Conformance	Summary
3	DYPF	DynamicPower-Flow	[SET]	The specified set of endpoints may change

11.8.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	AvailableEndpoints	list[end-point-no]	max 20	F		R V	SET
0x0001	ActiveEndpoints	list[end-point-no]	max 20	N		R V	DYPF

11.8.5.1. AvailableEndpoints Attribute

This attribute SHALL indicate the list of endpoints capable of providing power to and/or consuming power from the endpoint hosting this server.

11.8.5.2. ActiveEndpoints Attribute

This attribute SHALL indicate the current list of endpoints currently providing or consuming power to or from the endpoint hosting this server. This list SHALL be a subset of the value of the [AvailableEndpoints](#) attribute.

11.9. Network Commissioning Cluster

Network commissioning is part of the overall Node commissioning. The main goal of Network Commissioning Cluster is to associate a Node with or manage a Node's one or more network interfaces. These network interfaces can include the following types.

- Wi-Fi ([IEEE 802.11-2020](#))
- Ethernet (802.3)
- Thread (802.15.4)

An instance of the Network Commissioning Cluster only applies to a single network interface instance present. An interface, in this context, is a unique entity that can have an IPv6 address assigned to it and ingress and egress IP packets.

11.9.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Support determining capabilities for Wi-Fi and Thread interfaces. Additional Wi-Fi directed scanning requirements.

11.9.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	CNET

11.9.3. Cluster ID

ID	Name
0x0031	Network Commissioning

11.9.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Conformance	Summary
0	WI	WiFiNetworkInterface	O.a	Wi-Fi related features
1	TH	ThreadNetworkInterface	O.a	Thread related features
2	ET	EthernetNetworkInterface	O.a	Ethernet related features

11.9.5. Data Types

11.9.5.1. WiFiSecurityBitmap Type

This data type is derived from [map8](#).

WiFiSecurityBitmap encodes the supported Wi-Fi security types present in the Security field of the [WiFiInterfaceScanResultStruct](#).

Bit	Name	Summary	Conformance
0	Unencrypted	Supports unencrypted Wi-Fi	M
1	WEP	Supports Wi-Fi using WEP security	M

Bit	Name	Summary	Conformance
2	WPA-PERSONAL	Supports Wi-Fi using WPA-Personal security	M
3	WPA2-PERSONAL	Supports Wi-Fi using WPA2-Personal security	M
4	WPA3-PERSONAL	Supports Wi-Fi using WPA3-Personal security	M

11.9.5.2. ThreadCapabilitiesBitmap Type

This data type is derived from [map16](#).

The ThreadCapabilitiesBitmap encodes the supported Thread features and capabilities of a Thread-enabled network interface.

Bit	Name	Summary	Conformance
0	IsBorderRouterCapable	Thread Border Router functionality is present	O
1	IsRouterCapable	Router mode is supported (interface could be in router or REED mode)	O
2	IsSleepyEndDeviceCapable	Sleepy end-device mode is supported	O
3	IsFullThreadDevice	Device is a full Thread device (opposite of Minimal Thread Device)	O
4	IsSynchronizedSleepyEndDeviceCapable	Synchronized sleepy end-device mode is supported	O

NOTE

The valid combinations of capabilities are restricted and dependent on Thread version.

11.9.5.3. WiFiBandEnum Type

This data type is derived from [enum8](#).

WiFiBandEnum encodes a supported Wi-Fi frequency band present in the WiFiBand field of the [WiFiInterfaceScanResultStruct](#).

Value	Name	Summary	Conformance
0	2G4	2.4GHz - 2.401GHz to 2.495GHz (802.11b/g/n/ax)	O.b+
1	3G65	3.65GHz - 3.655GHz to 3.695GHz (802.11y)	O.b+
2	5G	5GHz - 5.150GHz to 5.895GHz (802.11a/n/ac/ax)	O.b+
3	6G	6GHz - 5.925GHz to 7.125GHz (802.11ax / Wi-Fi 6E)	O.b+
4	60G	60GHz - 57.24GHz to 70.20GHz (802.11ad/ay)	O.b+
5	1G	Sub-1GHz - 755MHz to 931MHz (802.11ah)	O.b+

11.9.5.4. NetworkCommissioningStatusEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Success	OK, no error	M
1	OutOfRange	Value Outside Range	M
2	BoundsExceeded	A collection would exceed its size limit	M
3	NetworkIDNotFound	The NetworkID is not among the collection of added networks	M
4	DuplicateNetworkID	The NetworkID is already among the collection of added networks	M
5	NetworkNotFound	Cannot find AP: SSID Not found	M
6	RegulatoryError	Cannot find AP: Mismatch on band/channels/regulatory domain / 2.4GHz vs 5GHz	M
7	AuthFailure	Cannot associate due to authentication failure	M

Value	Name	Summary	Conformance
8	UnsupportedSecurity	Cannot associate due to unsupported security mode	M
9	OtherConnectionFailure	Other association failure	M
10	IPv6Failed	Failure to generate an IPv6 address	M
11	IPBindFailed	Failure to bind Wi-Fi <-> IP interfaces	M
12	UnknownError	Unknown error	M

11.9.5.5. NetworkInfoStruct Type

NetworkInfoStruct struct describes an existing network configuration, as provided in the [Networks](#) attribute.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	NetworkID	octstr	1 to 32				M
1	Connected	bool					M

NetworkID Field

Every network is uniquely identified (for purposes of commissioning) by a NetworkID mapping to the following technology-specific properties:

- SSID for Wi-Fi
- Extended PAN ID for Thread
- Network interface instance name at operating system (or equivalent unique name) for Ethernet.

The semantics of the NetworkID field therefore varies between network types accordingly. It contains SSID for Wi-Fi networks, Extended PAN ID (XPAN ID) for Thread networks and netif name for Ethernet networks.

NOTE

SSID in Wi-Fi is a collection of 1-32 bytes, the text encoding of which is not specified. Implementations must be careful to support reporting byte strings without requiring a particular encoding for transfer. Only the commissioner should try to potentially decode the bytes. The most common encoding is UTF-8, however this is just a convention. Some configurations may use Latin-1 or other character sets. A commissioner MAY decode using UTF-8, replacing encoding errors with "?" at the application level while retaining the underlying representation.

XPAN ID is a big-endian 64-bit unsigned number, represented on the first 8 octets of the octet string.

Connected Field

This field SHALL indicate the connected status of the associated network, where "connected" means currently linked to the network technology (e.g. Associated for a Wi-Fi network, media connected for an Ethernet network).

11.9.5.6. WiFiInterfaceScanResultStruct Type

WiFiInterfaceScanResultStruct represents a single Wi-Fi network scan result.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Security	WiFiSecurityBitmap	all				M
1	SSID	octstr	max 32				M
2	BSSID	octstr	6				M
3	Channel	uint16	all				M
4	WiFiBand	WiFiBandEnum	all				O
5	RSSI	int8	all				O

WiFiBand Field

This field, if present, MAY be used to differentiate overlapping channel number values across different Wi-Fi frequency bands.

RSSI Field

This field, if present, SHALL denote the signal strength in dBm of the associated scan result.

11.9.5.7. ThreadInterfaceScanResultStruct Type

ThreadInterfaceScanResultStruct represents a single Thread network scan result.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	PanId	uint16	max 65534				M
1	Extended-PanId	uint64	all				M
2	Network-Name	string	1 to 16				M
3	Channel	uint16	all				M
4	Version	uint8	all				M
5	Extended Address	hwadr	all				M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
6	RSSI	int8	all				M
7	LQI	uint8	all				M

ExtendedAddress Field

ExtendedAddress stands for an IEEE 802.15.4 Extended Address.

11.9.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	MaxNetworks	uint8	min 1	F		R A	M
0x0001	Networks	list[NetworkInfoStruct]	max MaxNetworks		empty	R A	M
0x0002	ScanMaxTimeSeconds	uint8	desc	F		R V	WI TH
0x0003	ConnectMaxTimeSeconds	uint8	desc	F		R V	WI TH
0x0004	InterfaceEnabled	bool		N	true	RW VA	M
0x0005	LastNetworkingStatus	NetworkCommissioningStatusEnum		X	null	R A	M
0x0006	LastNetworkID	octstr	1 to 32	X	null	R A	M
0x0007	LastConnectErrorValue	int32		X	null	R A	M
0x0008	SupportedWiFiBands	list[WiFiBandEnum]	min 1	F	MS	R V	WI

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0009	SupportedThreadFeatures	ThreadCapabilitiesBitmap		F	MS	R V	TH
0x000A	ThreadVersion	uint16		F	MS	R V	TH

11.9.6.1. MaxNetworks Attribute

This SHALL indicate the maximum number of network configuration entries that can be added, based on available device resources. The length of the [Networks](#) attribute SHALL be less than or equal to this value.

11.9.6.2. Networks Attribute

This attribute SHALL indicate the network configurations that are usable on the network interface represented by this cluster server instance.

The order of configurations in the list reflects precedence. That is, any time the Node attempts to connect to the network it SHALL attempt to do so using the configurations in Networks Attribute in the order as they appear in the list.

The order of list items SHALL only be modified by the AddOrUpdateThreadNetwork, AddOrUpdateWiFiNetwork and ReorderNetwork commands. In other words, the list SHALL be stable over time, unless mutated externally.

Ethernet networks SHALL be automatically populated by the cluster server. Ethernet Network Commissioning Cluster instances SHALL always have exactly one [NetworkInfoStruct](#) instance in their [Networks](#) attribute. There SHALL be no way to add, update or remove Ethernet network configurations to those Cluster instances.

11.9.6.3. ScanMaxTimeSeconds Attribute

This attribute SHALL indicate the maximum duration taken, in seconds, by the network interface on this cluster server instance to provide scan results.

See [Section 11.9.7.1, “ScanNetworks Command”](#) for usage.

11.9.6.4. ConnectMaxTimeSeconds Attribute

This attribute SHALL indicate the maximum duration taken, in seconds, by the network interface on this cluster server instance to report a successful or failed network connection indication. This maximum time SHALL account for all operations needed until a successful network connection is deemed to have occurred, including, for example, obtaining IP addresses, or the execution of necessary internal retries.

11.9.6.5. InterfaceEnabled Attribute

This attribute SHALL indicate whether the associated network interface is enabled or not. By default all network interfaces SHOULD be enabled during initial commissioning (InterfaceEnabled set to true).

It is undefined what happens if InterfaceEnabled is written to false on the same interface as that which is used to write the value. In that case, it is possible that the Administrator would have to await expiry of the fail-safe, and associated recovery of network configuration to prior safe values, before being able to communicate with the node again (see [Section 11.10.7.2, “ArmFailSafe”](#)).

It MAY be possible to disable Ethernet interfaces but it is implementation-defined. If not supported, a write to this attribute with a value of false SHALL fail with a status of INVALID_ACTION. When disabled, an Ethernet interface would no longer employ media detection. That is, a simple unplug and replug of the cable SHALL NOT re-enable the interface.

On Ethernet-only Nodes, there SHALL always be at least one of the Network Commissioning server cluster instances with InterfaceEnabled set to true.

11.9.6.6. LastNetworkingStatus Attribute

This attribute SHALL indicate the status of the last attempt either scan or connect to an operational network, using this interface, whether by invocation of the [ConnectNetwork](#) command or by autonomous connection after loss of connectivity or during initial establishment. If no such attempt was made, or no network configurations exist in the [Networks](#) attribute, then this attribute SHALL be set to null.

This attribute is present to assist with error recovery during Network commissioning and to assist in non-concurrent networking commissioning flows.

11.9.6.7. LastNetworkID Attribute

This attribute SHALL indicate the NetworkID used in the last attempt to connect to an operational network, using this interface, whether by invocation of the [ConnectNetwork](#) command or by autonomous connection after loss of connectivity or during initial establishment. If no such attempt was made, or no network configurations exist in the [Networks](#) attribute, then this attribute SHALL be set to null.

If a network configuration is removed from the [Networks](#) attribute using the RemoveNetwork command after a connection attempt, this field MAY indicate a NetworkID that is no longer configured on the Node.

This attribute is present to assist with error recovery during Network commissioning and to assist in non-concurrent networking commissioning flows.

11.9.6.8. LastConnectErrorValue Attribute

This attribute SHALL indicate the ErrorValue used in the last failed attempt to connect to an operational network, using this interface, whether by invocation of the [ConnectNetwork](#) command or by autonomous connection after loss of connectivity or during initial establishment. If no such attempt

was made, or no network configurations exist in the [Networks](#) attribute, then this attribute SHALL be set to null.

If the last connection succeeded, as indicated by a value of Success in the [LastNetworkingStatus](#) attribute, then this field SHALL be set to null.

This attribute is present to assist with error recovery during Network commissioning and to assist in non-concurrent networking commissioning flows.

11.9.6.9. SupportedWiFiBands Attribute

This attribute SHALL indicate all the frequency bands supported by the Wi-Fi interface configured by the cluster instance.

11.9.6.10. SupportedThreadFeatures Attribute

This attribute SHALL indicate all of the Thread features supported by the Thread interface configured by the cluster instance.

This attribute is primarily used to determine the most important general capabilities of the Thread interface associated with the cluster instance, as opposed to the current runtime dynamic configuration.

NOTE

Most run-time details of the actual Thread interface are found in the [Thread Network Diagnostics cluster](#), if supported.

11.9.6.11. ThreadVersion Attribute

This attribute SHALL indicate the Thread version supported by the Thread interface configured by the cluster instance.

The format SHALL match the value mapping found in the "Version TLV" section of [Thread specification](#). For example, Thread 1.3.0 would have [ThreadVersion](#) set to 4.

11.9.7. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	ScanNetworks	client ⇒ server	ScanNetworksResponse	A	WI TH
0x01	ScanNetworksResponse	client ⇐ server	N		WI TH
0x02	AddOrUpdateWiFiNetwork	client ⇒ server	NetworkConfigResponse	A	WI
0x03	AddOrUpdateThreadNetwork	client ⇒ server	NetworkConfigResponse	A	TH

ID	Name	Direction	Response	Access	Conformance
0x04	RemoveNetwork	client ⇒ server	NetworkConfigResponse	A	WI TH
0x05	NetworkConfigResponse	client ⇐ server	N		WI TH
0x06	ConnectNetwork	client ⇒ server	ConnectNetworkResponse	A	WI TH
0x07	ConnectNetworkResponse	client ⇐ server	N		WI TH
0x08	ReorderNetwork	client ⇒ server	NetworkConfigResponse	A	WI TH

11.9.7.1. ScanNetworks Command

This command is used to scan for available networks on the network interface associated with the cluster instance.

This command SHALL scan on the Cluster instance's associated network interface for either of:

- All available networks (non-directed scanning)
- Specific networks (directed scanning)

Scanning for available networks detects all networks of the type corresponding to the cluster server instance's associated network interface that are possible to join, such as all visible Wi-Fi access points for Wi-Fi cluster server instances, all Thread PANs for Thread cluster server instances, within bounds of the maximum response size.

Scanning for a specific network (i.e. directed scanning) takes place if a network identifier (e.g. Wi-Fi SSID) is provided in the command arguments. Directed scanning SHALL restrict the result set to the specified network only.

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

The client SHALL NOT expect the server to be done scanning and have responded with ScanNetworksResponse before [ScanMaxTimeSeconds](#) seconds have elapsed. Enough transport time affordances for retries SHOULD be expected before a client determines the operation to have timed-out.

This command SHALL fail with a status code of BUSY if the server determines that it will fail to reliably send a response due to changes of networking interface configuration at runtime for the interface over which the command was invoked, or if it is currently unable to proceed with such an operation.

For Wi-Fi-supporting servers (WI feature) the server SHALL always honor directed scans, and attempt to provide all matching BSSID which are reachable on the bands which would otherwise be attempted if a ConnectNetwork having the specified SSID were to take place. This command is use-

ful for clients to determine reachability capabilities as seen by the server's own radios.

For Wi-Fi-supporting servers the server SHALL always scan on all bands supported by the interface associated with the cluster instance on which the command was invoked.

If the command was invoked over the same link whose configuration is managed by a given server cluster instance, there MAY be an impact on other communication from the invoking client, as well as other clients, while the network interface is processing the scan. Clients SHOULD NOT use this command unless actively in the process of re-configuring network connectivity.

The arguments for this command are as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	SSID	octstr	1 to 32	X	null	[WI]
1	Breadcrumb	uint64	all			O

SSID Field

This field, if present, SHALL contain the SSID for a directed scan of that particular Wi-Fi SSID. Otherwise, if the field is absent, or if it is null, this SHALL indicate scanning of all BSSID in range. This field SHALL be ignored for ScanNetworks invocations on non-Wi-Fi server instances.

Breadcrumb Field

The Breadcrumb field, if present, SHALL be used to atomically set the [Breadcrumb](#) attribute in the General Commissioning cluster on success of the associated command. If the command fails, the [Breadcrumb](#) attribute in the General Commissioning cluster SHALL be left unchanged.

11.9.7.2. ScanNetworksResponse Command

This command is used to report the results of a ScanNetworks command.

This command SHALL contain the status of the last ScanNetworks command, and the associated scan results if the operation was successful.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Network- ingStatus	Network- Commission- ingSta- tusEnum	desc			M
1	DebugText	string	max 512			O
2	WiFiScan- Results	list[WiFiIn- terfaceScan- ResultStruct]	desc			WI

ID	Name	Type	Constraint	Quality	Fallback	Conformance
3	Thread-ScanResults	list[ThreadInterfaceScanResultStruct]	desc			TH

Results are valid only if NetworkingStatus is Success.

Before generating a ScanNetworksResponse, the server SHALL set the [LastNetworkingStatus](#) attribute value to the NetworkingStatus matching the response.

NetworkingStatus Field

The NetworkingStatus field SHALL indicate the status of the last scan operation, taking one of these values:

- Success: Scanning succeeded.
- [NetworkNotFound](#): No instance of an explicitly-provided network identifier was found during the scan. This error cannot occur if no network identifier was provided, such as when scanning for all available networks.
- [OutOfRange](#): Network identifier was invalid (e.g. empty, too long, etc).
- [RegulatoryError](#): Could not scan on any bands due to lack of regulatory configuration.
- [UnknownError](#): An internal error occurred during scanning.

DebugText Field

This field, if present and non-empty, MAY contain error information which MAY be communicated to the user in case the NetworkingStatus was not Success. Its purpose is to help developers in troubleshooting errors and MAY go into logs or crash reports.

WiFiScanResults Field

If NetworkingStatus was Success, this field SHALL contain the Wi-Fi network scan results. The list MAY be empty if none were found in range on the bands supported by the interface, or if directed scanning had been used and the desired SSID was not found in range.

The maximum number of results present in the result list supported MAY depend on memory and MAY contain a subset of possibilities, to avoid memory exhaustion on the cluster server and avoid crossing the maximum command response size supported (see [Section 4.4.4, “Message Size Requirements”](#)).

The order in which results are reported is implementation-specific. Results SHOULD be reported in decreasing RSSI order, even if RSSI is not reported in the response, to maximize the likelihood that most likely to be reachable elements are included within the size limits of the response.

ThreadScanResults Field

If NetworkingStatus was Success, this field SHALL contain the Thread network scan results. The list MAY be empty if none were found in range on the bands supported by the interface.

The maximum number of results present in the result list supported MAY depend on memory and MAY contain a subset of possibilities, to avoid memory exhaustion on the cluster server and avoid crossing the maximum command response size supported (see [Section 4.4.4, “Message Size Requirements”](#)).

The order in which results are reported is implementation-specific. Results SHOULD be reported in decreasing LQI order, to maximize the likelihood that most likely to be reachable elements are included within the size limits of the response.

11.9.7.3. AddOrUpdateWiFiNetwork Command

This command is used to add or update a Wi-Fi network configuration.

This command SHALL be used to add or modify Wi-Fi network configurations.

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

The Credentials associated with the network are not readable after execution of this command, as they do not appear in the [Networks](#) attribute, for security reasons.

If this command contains a ClientIdentifier, and the [Networks](#) list does not contain an entry with a matching ClientIdentifier, then this command SHALL fail with a status of NOT_FOUND.

See [Section 11.9.7.5, “Common processing of AddOrUpdateWiFiNetwork and AddOrUpdateThreadNetwork”](#) for behavior of addition/update.

ID	Name	Type	Constraint	Fallback	Conformance
0	SSID	octstr	max 32		M
1	Credentials	octstr	max 64		M
2	Breadcrumb	uint64	all		O

SSID Field

This field SHALL indicate the SSID to which to attempt connection. Specific BSSID selection is not supported by this cluster.

Credentials Field

Credentials is the passphrase or PSK for the network (if any is needed).

Security type, cipher and credential format (passphrase or PSK) SHALL be contextually auto-selected during execution of the [ConnectNetwork](#) command and during subsequent operational state network connections, based on the most secure Wi-Fi security type available within beacons

and probe responses for the set of all discovered BSSIDs for the configured SSID. The type of PSK or passphrase used SHALL be inferred based on the length and contents of the Credentials field provided, matching the security type chosen.

Valid Credentials length are:

- 0 bytes: Unsecured (open) connection
- 5 bytes: WEP-64 passphrase
- 10 hexadecimal ASCII characters: WEP-64 40-bit hex raw PSK
- 13 bytes: WEP-128 passphrase
- 26 hexadecimal ASCII characters: WEP-128 104-bit hex raw PSK
- 8..63 bytes: WPA/WPA2/WPA3 passphrase
- 64 bytes: WPA/WPA2/WPA3 raw hex PSK

These lengths SHALL be contextually interpreted based on the security type of the BSSID where connection will occur.

When the length of Credentials and available set of BSSID admits more than one option, such as the presence of both WPA2 and WPA security type within the result set, WPA2 SHALL be considered more secure.

NOTE	It might occur that a station cannot connect to a particular access point with higher security and selects a lower security connectivity type if the link quality is deemed to be too low to achieve successful operation, or if all retry attempts fail.
------	---

Breadcrumb Field

See [Section 11.9.7.1.2, “Breadcrumb Field”](#) for usage.

11.9.7.4. AddOrUpdateThreadNetwork Command

This command is used to add or update a Thread network configuration.

This command SHALL be used to add or modify Thread network configurations.

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

See [Section 11.9.7.5, “Common processing of AddOrUpdateWiFiNetwork and AddOrUpdateThread-Network”](#) for behavior of addition/update.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Operational-Dataset	octstr	max 254			M

ID	Name	Type	Constraint	Quality	Fallback	Conformance
1	Breadcrumb	uint64	all			O

The XPAN ID in the OperationalDataset serves as the NetworkID for the network configuration to be added or updated.

If the [Networks](#) attribute does not contain an entry with the same NetworkID as the one provided in the OperationalDataset, the operation SHALL be considered an addition, otherwise, it SHALL be considered an update.

OperationalDataset Field

This field SHALL indicate the Thread Network Parameters, including channel, PAN ID, and Extended PAN ID.

The encoding for the OperationalDataset field is defined in the [Thread specification](#).

The client SHALL pass the OperationalDataset as an opaque octet string.

Breadcrumb Field

See [Section 11.9.7.1.2, “Breadcrumb Field”](#) for usage.

11.9.7.5. Common processing of AddOrUpdateWiFiNetwork and AddOrUpdateThreadNetwork

Both AddOrUpdateWiFiNetwork and AddOrUpdateThreadNetwork operate similarly, against different underlying technologies. The processing of these commands in the addition and update case is covered by the following subsections.

Processing an addition

If the [Networks](#) attribute is already full, the command SHALL immediately respond with NetworkConfigResponse having NetworkingStatus status field set to BoundsExceeded.

If any of the parameters in the OperationalDataset are invalid, the command SHALL immediately respond with NetworkConfigResponse having NetworkingStatus status field set to a value different than Success and consistent with the error.

If validation of all parameters has succeeded, this command SHALL append the configuration at the end of the existing list in the [Networks](#) attribute, making this new network the one with least priority.

On success, the NetworkConfigResponse command SHALL have its NetworkIndex field set to the 0-based index of the entry in the [Networks](#) attribute that was just added.

Processing an update

If any of the parameters in the OperationalDataset are invalid, the command SHALL immediately respond with NetworkConfigResponse having NetworkingStatus status field set to a value different

than Success and consistent with the error.

If validation of all parameters has succeeded, this command SHALL update the existing entry indexed by NetworkID in the [Networks](#) attribute, keeping existing position within the list.

On success, the NetworkConfigResponse command SHALL have its NetworkIndex field set to the 0-based index of the entry in the [Networks](#) attribute that was just updated, and a NetworkingStatus status field set to Success.

11.9.7.6. RemoveNetwork Command

This command is used to remove a network configuration on the network interface associated with the cluster instance.

This command SHALL remove the network configuration from the Cluster if there was already a network configuration with the same NetworkID. The relative order of the entries in the [Networks](#) attribute SHALL remain unchanged, except for the removal of the requested network configuration.

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NetworkID	octstr	1 to 32			M
1	Breadcrumb	uint64	all			O

If the [Networks](#) attribute does not contain a matching entry, the command SHALL immediately respond with NetworkConfigResponse having NetworkingStatus status field set to NetworkIdNotFound.

On success, the NetworkConfigResponse command SHALL have its NetworkIndex field set to the 0-based index of the entry in the [Networks](#) attribute that was just removed, and a NetworkingStatus status field set to Success.

NetworkID Field

This field SHALL indicate the NetworkID for the entry to remove: the SSID for Wi-Fi and XPAN ID for Thread.

Breadcrumb Field

See [Section 11.9.7.1.2, “Breadcrumb Field”](#) for usage.

11.9.7.7. NetworkConfigResponse Command

This command is used to report the results of a network configuration command (AddOrUpdateWiFiNetwork, AddOrUpdateThreadNetwork, or RemoveNetwork).

This response command relates status information for some commands which require it as their response command. See each individual cluster server command for the situations that may cause a NetworkingStatus different than Success.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NetworkingStatus	NetworkCommissioningStatusEnum	desc			M
1	DebugText	string	max 512			O
2	NetworkIndex	uint8	max (MaxNetworks - 1)			NetworkingStatus == Success

Before generating a NetworkConfigResponse, the server SHALL set the [LastNetworkingStatus](#) attribute value to the NetworkingStatus matching the response.

Before generating a NetworkConfigResponse, the server SHALL set the [LastNetworkID](#) attribute value to the NetworkID that was used in the command for which an invocation caused the response to be generated.

NetworkingStatus Field

This field SHALL indicate the status of the last operation attempting to modify the [Networks](#) attribute configuration, taking one of these values:

- Success: Operation succeeded.
- [OutOfRange](#): Network identifier was invalid (e.g. empty, too long, etc).
- [BoundsExceeded](#): Adding this network configuration would exceed the limit defined by [MaxNetworks](#).
- [NetworkIdNotFound](#): The network identifier was expected to be found, but was not found among the added network configurations in [Networks](#) attribute.
- [UnknownError](#): An internal error occurred during the operation.

DebugText Field

See [Section 11.9.7.2.2, “DebugText Field”](#) for usage.

NetworkIndex Field

This field SHALL indicate the 0-based index of the entry in the [Networks](#) attribute that was last added, updated or removed successfully by the associated request command.

11.9.7.8. ConnectNetwork Command

This command is used to connect to a network on the network interface associated with the cluster instance.

This command SHALL attempt to connect to a network whose configuration was previously added by either the [AddOrUpdateWiFiNetwork](#) or [AddOrUpdateThreadNetwork](#) commands. Network is identified by its NetworkID.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NetworkID	octstr	1 to 32			M
1	Bread-crumbs	uint64	all			O

This command SHALL fail with a BUSY status code returned to the initiator if the server is currently unable to proceed with such an operation, such as if it is currently attempting to connect in the background, or is already proceeding with a prior ConnectNetwork.

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

Before connecting to the new network, the Node SHALL disconnect the operational network connections managed by any other Network Commissioning cluster instances (whether under the Root Node or a Secondary Network Interface), where those connections are not represented by an entry in the Networks attribute of the corresponding cluster instance. This ensures that an Administrator or Commissioner can reliably reconfigure the operational network connection of a device that has one or more Secondary Network interfaces, for example by removing the active network configuration from one cluster instance, followed by adding a new configuration and calling ConnectNetwork on a different cluster instance.

Success or failure of this command SHALL be communicated by the ConnectNetworkResponse command, unless some data model validations caused a FAILURE status to be sent prior to finishing execution of the command. The ConnectNetworkResponse SHALL indicate the value Success in the NetworkingStatus field on successful connection. On failure to connect, the ConnectNetworkResponse SHALL contain an appropriate NetworkingStatus, DebugText and ErrorValue indicating the reason for failure.

The amount of time needed to determine successful or failing connectivity on the cluster server’s associated interface is provided by the [ConnectMaxTimeSeconds](#) attribute. Clients SHALL NOT consider the connection to have timed-out until at least that duration has taken place. For non-concurrent commissioning situations, the client SHOULD allow additional margin of time to account for its delay in executing [operational discovery](#) of the Node once it is connected to the new network.

On successful connection, the entry associated with the given Network configuration in the [Networks](#) attribute SHALL indicate its Connected field set to true, and all other entries, if any exist, SHALL indicate their Connected field set to false.

On failure to connect, the entry associated with the given Network configuration in the [Networks](#) attribute SHALL indicate its Connected field set to false.

The precedence order of any entry subject to ConnectNetwork SHALL NOT change within the [Networks](#) attribute.

Even after successfully connecting to a network, the configuration SHALL revert to the prior state of configuration if the CommissioningComplete command (see [Section 11.10.7.6, “CommissioningComplete”](#)) is not successfully invoked before expiry of the Fail-Safe timer.

When non-concurrent commissioning is being used by a Commissioner or Administrator, the ConnectNetworkResponse SHALL be sent with the NetworkingStatus field set to Success prior to closing the commissioning channel, even if not yet connected to the operational network, unless the device would be incapable of joining that network, in which case the usual failure path described in the prior paragraphs SHALL be followed. Once the commissioning channel is closed, if the attribute [IsCommissioningWithoutPower](#) is absent or false, the operational channel SHALL start without waiting. However, if [IsCommissioningWithoutPower](#) is true, the operational channel SHALL start automatically as soon as the device is powered up and running. It is possible that the only method to determine success of the operation is operational discovery of the Node on the new operational network. Therefore, before invoking the ConnectNetwork command, the client SHOULD re-invoke the [Arm Fail-Safe](#) command with a duration that meets the following:

1. Sufficient time to meet the minimum required time (see [Section 11.9.6.4, “ConnectMaxTimeSeconds Attribute”](#)) that may be taken by the server to connect to the desired network.
2. Sufficient time to account for possible message-layer retries when a response is requested.
3. Sufficient time to allow [operational discovery](#) on the new network by a Commissioner or Administrator.
4. Sufficient time to establish a [CASE](#) session after operational discovery
5. Not so long that, in error situations, the delay to reverting back to being discoverable for commissioning with a previous configuration would cause significant user-perceived delay.

Note as well that the CommissioningTimeout duration provided in a prior [OpenCommissioningWindow](#) or [OpenBasicCommissioningWindow](#) command may impact the total time available to proceed with error recovery after a connection failure.

The [LastNetworkingStatus](#), [LastNetworkID](#) and [LastConnectErrorValue](#) attributes MAY assist the client in determining the reason for a failure after reconnecting over a Commissioning channel, especially in non-concurrent commissioning situations.

NetworkID Field

This field SHALL indicate the NetworkID for the entry used to configure the connection: the SSID for Wi-Fi and XPAN ID for Thread.

Breadcrumb Field

See [Section 11.9.7.1.2, “Breadcrumb Field”](#) for usage.

11.9.7.9. ConnectNetworkResponse Command

This command is used to report the results of a ConnectNetwork command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NetworkingStatus	NetworkCommissioningStatusEnum	all			M
1	DebugText	string				O
2	ErrorValue	int32	all	X		M

Before generating a ConnectNetworkResponse, the server SHALL:

- Set the [LastNetworkingStatus](#) attribute value to the NetworkingStatus matching the response.
- Set the [LastNetworkID](#) attribute value to the NetworkID that was used in the ConnectNetwork command which caused the response to be generated.
- Set the [LastConnectErrorValue](#) attribute value to the ErrorValue matching the response, including setting it to null if the ErrorValue is not applicable.

NetworkingStatus Field

This field SHALL indicate the status of the last connection attempt, taking one of these values:

- Success: Connection succeeded.
- [NetworkNotFound](#): No instance of an explicitly-provided network identifier was found during the attempt to join the network.
- [OutOfRange](#): Network identifier was invalid (e.g. empty, too long, etc).
- [NetworkIdNotFound](#): The network identifier was not found among the added network configurations in [Networks](#) attribute.
- [RegulatoryError](#): Could not connect to a network due to lack of regulatory configuration.
- [UnknownError](#): An internal error occurred during the operation.
- Association errors (see also description of errors in [NetworkCommissioningStatusEnum](#)): [AuthFailure](#), [UnsupportedSecurity](#), [OtherConnectionFailure](#), [IPV6Failed](#), [IPBindFailed](#)

DebugText Field

See [Section 11.9.7.2.2](#), “[DebugText Field](#)” for usage.

ErrorValue Field

This field SHALL indicate the status of the associated request command.

- ErrorValue interpretation for Wi-Fi association errors:
 - On any association failure during enabling of a network, the ErrorValue field SHALL be set

to the Status Code value that was present in the last frame related to association where Status Code was not equal to zero and which caused the failure of a final retry attempt, if this final failure was due to one of the following Management frames:

- Association Response (Type 0, Subtype 1)
- Reassociation Response (Type 0, Subtype 3)
- Authentication (Type 0, Subtype 11)
- Table 9-50 "Status Codes" in [IEEE 802.11-2020](#) contains a description of all values possible, which can unambiguously be used to determine the cause, such as an invalid security type, unsupported rate, etc.
- Otherwise, the ErrorValue field SHALL contain an implementation-dependent value which MAY be used by a reader of the structure to record, report or diagnose the failure.

11.9.7.10. ReorderNetwork Command

This command is used to re-order the network configuration list.

This command SHALL set the specific order of the network configuration selected by its NetworkID in the [Networks](#) attribute to match the position given by NetworkIndex.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NetworkID	octstr	1 to 32			M
1	NetworkIndex	uint8	desc			M
2	Breadcrumb	uint64	all			O

NetworkID Field

This field SHALL indicate the NetworkID for the entry to reorder: the SSID for Wi-Fi and XPAN ID for Thread.

NetworkIndex Field

This field SHALL indicate the 0-based index of the new desired position of the entry in the [Networks](#) attribute.

Breadcrumb Field

See [Section 11.9.7.1.2, “Breadcrumb Field”](#) for usage.

Effect when received

If the [Networks](#) attribute does not contain a matching entry, the command SHALL immediately respond with NetworkConfigResponse having NetworkingStatus status field set to NetworkIdNotFound.

If the `NetworkIndex` field has a value larger or equal to the current number of entries in the `Networks` attribute, the command SHALL immediately respond with `NetworkConfigResponse` having `NetworkingStatus` status field set to `OutOfRange`.

On success, the `NetworkConfigResponse` command SHALL have its `NetworkIndex` field set to the 0-based index of the entry in the `Networks` attribute that was just updated, matching the incoming `NetworkIndex`, and a `NetworkingStatus` status field set to `Success`.

The entry selected SHALL be inserted at the new position in the list. All other entries, if any exist, SHALL be moved to allow the insertion, in a way that they all retain their existing relative order between each other, with the exception of the newly re-ordered entry.

Re-ordering to the same `NetworkIndex` as the current location SHALL be considered as a success and yield no visible changes of the `Networks` attribute.

Examples of re-ordering

To better illustrate the re-ordering operation, consider this initial state, exemplary of a Wi-Fi device:

Index in list	NetworkID field	Connected field
0	FancyCat	false
1	BlueDolphin	true
2	Home-Guest	false
3	WillowTree	false

On receiving `ReorderNetwork` with:

- `NetworkID` = Home-Guest
- `NetworkIndex` = 0

The outcome, after applying to the **initial state** would be:

Index in list	NetworkID field	Connected field
0	Home-Guest	false
1	FancyCat	false
2	BlueDolphin	true
3	WillowTree	false

In the above outcome, `FancyCat` and `BlueDolphin` moved "down" and `Home-Guest` became the highest priority network in the list.

On receiving `ReorderNetwork` with:

- `NetworkID` = FancyCat
- `NetworkIndex` = 3

The outcome, after applying to the **initial state** would be:

Index in list	NetworkID field	Connected field
0	BlueDolphin	true
1	Home-Guest	false
2	WillowTree	false
3	FancyCat	false

In the above outcome, BlueDolphin, Home-Guest and WillowTree moved "up" and FancyCat became the lowest priority network in the list.

11.9.8. Usage of networking configurations

This section describes how to ensure deterministic and well-behaved network connectivity, both when concurrent and non-concurrent commissioning flows (see [Section 5.5, "Commissioning Flows"](#)) are used.

Operational Networking configuration is managed by the set of Network Commissioning cluster server instances distributed on a Node's Endpoints.

Since Matter employs IPv6 communication and DNS-SD for operational [Discovery](#), there is a fundamental aspect of multi-homing present, where a Node with multiple concurrently operated network interfaces device may be reachable using a variety of addresses on different network technologies. Care SHOULD be taken by Administrators and Commissioners to avoid making strong assumptions about single address reachability. Administrators and Commissioners SHOULD be prepared to attempt reachability tests against specific network technologies if the final desired state of networking requires a specific reachable path.

A Node MAY be configured in such a way that there are no Network Commissioning cluster server instances present, in which case the remainder of this section SHALL NOT apply.

11.9.8.1. Network Interfaces

Primary Network Interface

The primary network interface of a Node SHALL be the one present on the root node endpoint (see [Section 9.2, "Endpoint Composition"](#)). This interface is expected to be the one most likely to yield an operational reachable state if appropriately configured.

Secondary Network Interfaces

The secondary network interfaces SHOULD be additional technologies that MAY increase reachability or support a [stub router feature](#). These interfaces SHALL be hosted by endpoints other than the root node endpoint, which SHALL support the Secondary Network Interface device type.

A Node supporting secondary network interfaces SHALL support the concurrent connection commissioning flow (see [Figure 47, "Commissioning flow diagram"](#)). Therefore, if the primary network configuration fails, the Commissioner can proceed to configure the secondary network interfaces

via the established commissioning channel.

11.9.8.2. Order of connectivity during connection establishment

When at least one Network Commissioning cluster server instance (hereafter, "Network Commissioning cluster" for short), the following behavior SHOULD take place for each interface associated with a Network Commissioning cluster, in increasing order of associated endpoint number:

1. If the Network Commissioning cluster's InterfaceEnabled attribute is set to false, skip the processing the interface altogether.
2. Set all configurations of the [Networks attribute](#) entry's Connected field to false
3. Iterate through all configurations in the [Networks attribute](#)
 - a. If there was a "last known good" network configuration, that is, the one which was both last successfully connected during prior boot and over which at least one secure channel exchange message was received, it MAY be used as the first attempt. Otherwise, iterate through all configurations in the precedence order of the list, starting at index 0.
4. Attempt to connect to the technology, using the current iteration's network configuration
5. On success, set the Connected state of the list entry to true, and stop attempting further connection. Otherwise, on failure, move to the next configuration.

11.9.8.3. Connectivity management during commissioning or administration

When a network interface is configured during commissioning or reconfigured during ongoing administration, behavior is different than for the startup case described previously, since there are tentative attempts being made to make a Node reachable on an operational network.

Network configuration can be seen as:

- A list of existing configurations, reflected by the [Networks attribute](#).
 - The list SHALL be managed by the AddOrUpdateWiFiNetwork, AddOrUpdateThreadNetwork, RemoveNetwork and ReorderNetwork commands.
 - The list SHALL be tentative until committed by successful invocation of the [CommissioningComplete](#) command, or reverted to prior configuration by the expiry of the Fail-Safe timer (see [Section 11.10.7.2, "ArmFailSafe"](#)).
- A current candidate configuration, the subject of the most recent ConnectNetwork command.
 - There SHALL NOT be new connections to any network during the Fail-Safe timer period unless attempted by invocation of the ConnectNetwork command.

Failures of connection during the Fail-Safe timer window SHALL cause the Node to follow the steps in [Section 5.5.1, "Commissioning Flows Error Handling"](#) after recording the cause of the failure in the [LastNetworkingStatus](#), [LastNetworkID](#) and [LastConnectErrorValue](#) attributes.

After Commissioning or reconfiguration ends successfully, because of successful invocation of the [CommissioningComplete](#) command, the cluster server SHOULD NOT attempt to change connected network until connectivity failure or restart occurs, but rather it SHALL commit the tentative configuration to persistent storage so that it is usable the next time connectivity establishment is

needed.

After Commissioning or reconfiguration ends in failure due to expiry of the Fail-Safe timer, the Node SHALL revert to the network configuration present prior to the Fail-Safe timer being armed.

Because it is possible that multiple network configurations being present could successfully result in an established operational network connection, but only some of these configurations actually have the desired reachability by Administrators on certain fabrics, the following precautions SHOULD be taken to avoid a situation where a Node forever dwells on a network with successful connectivity, but no reachable peers:

- Commissioners and Administrators MAY notify users if multiple independent configurations exist that could cause an alternate configuration to make the device unreachable for reconfiguration by Nodes on the current client's fabric in the future.
- Commissioners and Administrators SHOULD avoid configuring Nodes in ways where it may be ambiguous to end-users which final network configuration will take place.
- Cluster servers on devices with no user interface to express current network configuration to an end-user SHOULD be configured to only support a single entry in the [Networks attribute](#).
- Upon discovering that a user is desiring to configure a Network in a way that would change the set of configured networks, and there are multiple fabrics configured in the [Fabrics](#) attribute of the [Operational Credentials](#) cluster, the client SHOULD notify the user that some other Administrators on other fabrics MAY fail to reach the Node and report connectivity failures.

11.10. General Commissioning Cluster

This cluster is used to manage basic commissioning lifecycle.

This cluster also represents responsibilities related to commissioning that don't well fit other commissioning clusters, like [Section 11.9, "Network Commissioning Cluster"](#). It also hosts functionalities those other clusters may depend on.

11.10.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Add Enhanced Setup Flow

11.10.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	CGEN

11.10.3. Cluster ID

ID	Name
0x0030	General Commissioning

11.10.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Conformance	Summary
0	TC	TermsAndConditions	O	Supports Terms & Conditions acknowledgement
1	NR	NetworkRecovery	P, O	Supports Network Recovery

11.10.4.1. Enhanced Setup Flow Terms & Conditions Feature

The node supports commissioning using [Enhanced Setup Flow Terms & Conditions](#). Support for this feature is limited to nodes that use [Commissioning Custom Flow](#).

11.10.5. Data Types

11.10.5.1. CommissioningErrorEnum Type

This data type is derived from [enum8](#).

This enumeration is used by several response commands in this cluster to indicate particular errors.

Value	Name	Summary	Conformance
0	OK	No error	M
1	ValueOutsideRange	Attempting to set regulatory configuration to a region or indoor/outdoor mode for which the server does not have proper configuration.	M
2	InvalidAuthentication	Executed CommissioningComplete outside CASE session.	M

Value	Name	Summary	Conformance
3	NoFailSafe	Executed CommissioningComplete when there was no active Fail-Safe context .	M
4	BusyWithOtherAdmin	Attempting to arm fail-safe or execute CommissioningComplete from a fabric different than the one associated with the current fail-safe context.	M
5	RequiredTCNotAccepted	One or more required TC features from the Enhanced Setup Flow were not accepted.	TC
6	TCAcknowledgementsNotReceived	No or insufficient acknowledgements from the user for the TC features were received.	TC
7	TCMinVersionNotMet	The version of the TC features acknowledged by the user did not meet the minimum required version.	TC

11.10.5.2. RegulatoryLocationTypeEnum Type

This data type is derived from [enum8](#).

This enumeration is used by the RegulatoryConfig and LocationCapability attributes to indicate possible radio usage.

Value	Name	Summary	Conformance
0	Indoor	Indoor only	M
1	Outdoor	Outdoor only	M
2	IndoorOutdoor	Indoor/Outdoor	M

11.10.5.3. NetworkRecoveryReasonEnum Type

This data type is derived from [enum8](#), however, the maximum value of the enumeration SHALL be less than 15.

This enumeration is used by the NetworkRecoveryReason attribute and may be embedded in the

beacons advertising the need for Network Recovery.

Value	Name	Summary	Conformance
0	Unspecified	Unspecified / unknown reason of network failure	P, M
1	Auth	Credentials for the configured operational network are not valid	P, M
2	Visibility	Configured network cannot be found (e.g. the device cannot see the configured Wi-Fi SSID, Thread end-node is unable to find a parent router on the PAN)	P, M
15 - 255	DoNotUse		X

11.10.5.4. BasicCommissioningInfo Type

This structure provides some constant values that MAY be of use to all commissioners.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	FailSafeExpiryLengthSeconds	uint16	all				M
1	MaxCumulativeFail-safeSeconds	uint16	desc				M

FailSafeExpiryLengthSeconds Field

This field SHALL contain a conservative initial duration (in seconds) to set in the FailSafe for the commissioning flow to complete successfully. This may vary depending on the speed or sleepiness of the Commissionee. This value, if used in the [Section 11.10.7.2, “ArmFailSafe”](#) command’s ExpiryLengthSeconds field SHOULD allow a Commissioner to proceed with a nominal commissioning without having to-rearm the fail-safe, with some margin.

MaxCumulativeFailsafeSeconds Field

This field SHALL contain a conservative value in seconds denoting the maximum total duration for which a fail safe timer can be re-armed. See [Section 11.10.7.2.1, “Fail Safe Context”](#).

The value of this field SHALL be greater than or equal to the FailSafeExpiryLengthSeconds. Absent additional guidelines, it is RECOMMENDED that the value of this field be aligned with the initial [Section 5.4.2.3, “Announcement Duration”](#) and default to 900 seconds.

11.10.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Bread-crumbs	uint64	all		0	RW VA	M
0x0001	BasicCommissioning-Info	BasicCommissioningInfo	desc	F		R V	M
0x0002	RegulatoryConfig	RegulatoryLocationTypeEnum	all		Location-Capability	R V	M
0x0003	Location-Capability	RegulatoryLocationTypeEnum	all	F	IndoorOutdoor	R V	M
0x0004	SupportsConcurrentConnection	bool	all	F	true	R V	M
0x0005	TCAcceptedVersion	uint16	all	N		R A	TC
0x0006	TCMinRequiredVersion	uint16	all	N		R A	TC
0x0007	TCAcknowledgements	map16	all	N		R A	TC
0x0008	TCAcknowledgementsRequired	bool	all	N	true	R A	TC
0x0009	TCUpdateDeadline	uint32	all	N X		R A	TC
0x000A	RecoveryIdentifier	octstr	8	N	0	R M	P, NR

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x000B	Net-workRe-cov-eryReason	Net-workRe-cov-eryReasonEnum	all	X	null	R M	P, NR
0x000C	IsCommissioning-Without-Power	bool	all		false	R V	P, O

11.10.6.1. Breadcrumb Attribute

This attribute allows for the storage of a client-provided small payload which Administrators and Commissioners MAY write and then subsequently read, to keep track of their own progress. This MAY be used by the Commissioner to avoid repeating already-executed actions upon re-establishing a commissioning link after an error.

On start/restart of the server, such as when a device is power-cycled, this attribute SHALL be reset to zero.

Some commands related to commissioning also have a side-effect of updating or resetting this attribute and this is specified in their respective functional descriptions.

The format of the value within this attribute is unspecified and its value is not otherwise used by the functioning of any cluster, other than being set as a side-effect of commands where this behavior is described.

11.10.6.2. BasicCommissioningInfo Attribute

This attribute SHALL describe critical parameters needed at the beginning of commissioning flow. See [Section 11.10.5.4, “BasicCommissioningInfo”](#) for more information.

11.10.6.3. RegulatoryConfig Attribute

This attribute SHALL indicate the regulatory configuration for the product.

Note that the country code is part of [Section 11.1, “Basic Information Cluster”](#) and therefore NOT listed on the RegulatoryConfig attribute.

11.10.6.4. LocationCapability Attribute

LocationCapability is statically set by the manufacturer and indicates if this Node needs to be told an exact RegulatoryLocation. For example a Node which is "Indoor Only" would not be certified for outdoor use at all, and thus there is no need for a commissioner to set or ask the user about whether the device will be used inside or outside. However a device which states its capability is "Indoor/Outdoor" means it would like clarification if possible.

For Nodes without radio network interfaces (e.g. Ethernet-only devices), the value IndoorOutdoor

SHALL always be used.

The default value of the [RegulatoryConfig](#) attribute is the value of LocationCapability attribute. This means devices always have a safe default value, and Commissioners which choose to implement smarter handling can.

11.10.6.5. SupportsConcurrentConnection Attribute

This attribute SHALL indicate whether this device supports "concurrent connection flow" commissioning mode (see [Section 5.5, "Commissioning Flows"](#)). If false, the device only supports "non-concurrent connection flow" mode.

11.10.6.6. TCAcceptedVersion Attribute

This attribute SHALL indicate the last version of the T&Cs for which the device received user acknowledgements. On factory reset this field SHALL be reset to 0.

When [Custom Commissioning Flow](#) is used to obtain user consent (e. g. because the Commissioner does not support the TC feature), the manufacturer-provided means for obtaining user consent SHALL ensure that this attribute is set to a value which is greater than or equal to [TCMinRequiredVersion](#) before returning the user back to the originating Commissioner (see [Section 5.7.4, "Enhanced Setup Flow \(ESF\)"](#)).

11.10.6.7. TCMinRequiredVersion Attribute

This attribute SHALL indicate the minimum version of the texts presented by the Enhanced Setup Flow that need to be accepted by the user for this device. This attribute MAY change as the result of an OTA update.

If an event such as a software update causes TCAcceptedVersion to become less than TCMinRequiredVersion, then the device SHALL update [TCAcknowledgementsRequired](#) to True so that an administrator can detect that a newer version of the texts needs to be presented to the user.

11.10.6.8. TCAcknowledgements Attribute

This attribute SHALL indicate the user's response to the presented terms. Each bit position corresponds to a user response for the associated index of matching text, such that bit 0 (bit value 1) is for text index 0. Bit 15 (bit value 0x8000) is for text index 15. A bit value of 1 indicates acceptance and a value of 0 indicates non-acceptance. For example, if there are two texts that were presented where the first (bit 0, value 1) was declined and the second accepted (bit 1, value 2), we would expect the resulting value of the map to be 2.

Whenever a user provides responses to newly presented terms and conditions, this attribute SHALL be updated with the latest responses. This MAY happen in response to updated terms that were presented to the user. On a factory reset this field SHALL be reset with all bits set to 0.

11.10.6.9. TCAcknowledgementsRequired Attribute

This attribute SHALL indicate whether [SetTCAcknowledgements](#) is currently required to be called with the inclusion of mandatory terms accepted.

This attribute MAY be present and False in the case where no terms and conditions are currently mandatory to accept for [CommissioningComplete](#) command to succeed.

This attribute MAY appear, or become True after commissioning (e.g. due to a firmware update) to indicate that new Terms & Conditions are available that the user must accept.

Upon Factory Data Reset, this attribute SHALL be set to a value of True.

When [Section 5.7.3, “Custom Commissioning Flow”](#) is used to obtain user consent (e.g. because the Commissioner does not support the TC feature), the manufacturer-provided means for obtaining user consent SHALL ensure that this attribute is set to False before returning the user back to the original Commissioner (see [Section 5.7.4, “Enhanced Setup Flow \(ESF\)”](#)).

11.10.6.10. TCUpdateDeadline Attribute

This attribute SHALL indicate the System Time in seconds when any functionality limitations will begin due to a lack of acceptance of updated Terms and Conditions, as described in [Section 5.7.4.6, “Presenting Updated Terms and Conditions”](#).

A null value indicates that there is no pending deadline for updated TC acceptance.

11.10.6.11. RecoveryIdentifier Attribute

This attribute SHALL contain the identifier to be included in the advertisements used during the [Network Recovery Flow](#). This identifier is intended to be advertised over the air and used by an Administrator to establish a Node’s identity without revealing its Node ID.

The attribute SHALL contain a random 64-bit value, that value SHALL be reset on factory reset and SHALL remain unchanged until a next factory reset. It is important that this value be selected at random from a 64-bit number space to ensure a high likelihood of uniqueness from values selected by other Nodes. This value SHOULD be obtained through `Crypto_DRBG(len = 64)`.

11.10.6.12. NetworkRecoveryReason Attribute

This attribute SHALL contain the primary reason that triggered the Network Recovery flow and its associated advertisements. This attribute SHALL be null when the Node is not undergoing a Network Recovery flow.

11.10.6.13. IsCommissioningWithoutPower Attribute

The server SHALL set this attribute to true if and only if is currently operating on the commissioning channel but cannot operate on the operational channel because it is not powered.

This may happen during NFC-based commissioning, when the commissioning channel is [NFC Transport Layer \(NTL\)](#), because it can harvest energy from NFC to operate. However, such a Commissionee must be powered on to switch to the operational channel.

This attribute is used by the Commissioner as described in initial phase step 17 of [Section 5.5, “Commissioning Flows”](#). This attribute is linked to the Commissionee behavior after reception of the [ConnectNetwork Command](#).

11.10.7. Commands

For all client-to-server commands in this cluster, if the client deems that it has timed-out in receiving the corresponding response command to any request, the corresponding step in the commissioning flow SHALL be considered to have failed, with the error handled as described in [Section 5.5.1, “Commissioning Flows Error Handling”](#).

ID	Name	Direction	Response	Access	Conformance
0x00	ArmFailSafe	client ⇒ server	ArmFailSafeResponse	A	M
0x01	ArmFailSafeResponse	client ⇐ server	N		M
0x02	SetRegulatoryConfig	client ⇒ server	SetRegulatoryConfigResponse	A	M
0x03	SetRegulatoryConfigResponse	client ⇐ server	N		M
0x04	CommissioningComplete	client ⇒ server	CommissioningCompleteResponse	A F	M
0x05	CommissioningCompleteResponse	client ⇐ server	N		M
0x06	SetTCAcknowledgements	client ⇒ server	SetTCAcknowledgementsResponse	A	TC
0x07	SetTCAcknowledgementsResponse	client ⇐ server	N		TC

11.10.7.1. Common fields in General Commissioning cluster responses

Some response commands have a DebugText argument which SHOULD NOT be presented directly in user interfaces. Its purpose is to help developers in troubleshooting errors. The value MAY go into logs or crash reports.

11.10.7.2. ArmFailSafe Command

This command is used to arm or disarm the fail-safe timer.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ExpiryLengthSeconds	uint16			900	M
1	Breadcrumb	uint64				M

Success or failure of this command SHALL be communicated by the [ArmFailSafeResponse](#) command, unless some data model validations caused a failure status code to be issued during the processing of the command.

If the [fail-safe timer](#) is not currently armed, the commissioning window is open, and the command was received over a [CASE](#) session, the command SHALL leave the current fail-safe state unchanged and immediately respond with an [ArmFailSafeResponse](#) containing an ErrorCode value of [BusyWithOtherAdmin](#). This is done to allow commissioners, which use PASE connections, the opportunity to use the failsafe during the relatively short commissioning window.

Otherwise, the command SHALL arm or re-arm the "fail-safe timer" with an expiry time set for a duration of ExpiryLengthSeconds, or disarm it, depending on the situation:

- If ExpiryLengthSeconds is 0 and the fail-safe timer was already armed and the [accessing fabric](#) matches the Fabric currently associated with the fail-safe context, then the fail-safe timer SHALL be immediately expired (see further below for side-effects of expiration).
- If ExpiryLengthSeconds is 0 and the fail-safe timer was not armed, then this command invocation SHALL lead to a success response with no side-effects against the fail-safe context.
- If ExpiryLengthSeconds is non-zero and the fail-safe timer was not currently armed, then the fail-safe timer SHALL be armed for that duration.
- If ExpiryLengthSeconds is non-zero and the fail-safe timer was currently armed, and the [accessing Fabric](#) matches the fail-safe context's associated Fabric, then the fail-safe timer SHALL be re-armed to expire in ExpiryLengthSeconds.
- Otherwise, the command SHALL leave the current fail-safe state unchanged and immediately respond with [ArmFailSafeResponse](#) containing an ErrorCode value of [BusyWithOtherAdmin](#), indicating a likely conflict between commissioners.

The value of the Breadcrumb field SHALL be written to the [Breadcrumb](#) on successful execution of the command.

If the receiver restarts unexpectedly (e.g., power interruption, software crash, or other reset) the receiver SHALL behave as if the fail-safe timer expired and perform the sequence of clean-up steps listed below.

On successful execution of the command, the ErrorCode field of the [ArmFailSafeResponse](#) SHALL be set to OK.

Fail Safe Context

When first arming the fail-safe timer, a 'Fail Safe Context' SHALL be created on the receiver, to

track the following state information while the fail-safe is armed:

- The fail-safe timer duration.
- Whether the device was fully powered during the [initial phase of commissioning](#).
- The state of all Network Commissioning [Networks](#) attribute configurations, to allow recovery of connectivity after Fail-Safe expiry.
- Whether an [AddNOC](#) command or [UpdateNOC](#) command has taken place.
- A [fabric-index](#) for the fabric-scoping of the context, starting at the [accessing fabric index](#) for the ArmFailSafe command, and updated with the Fabric Index associated with an [AddNOC](#) or an [UpdateNOC](#) command being invoked successfully during the ongoing Fail-Safe timer period.
- The operational credentials associated with any Fabric whose configuration is affected by the [UpdateNOC](#) command.
- Optionally: the previous state of non-fabric-scoped data that is mutated during the fail-safe period.

Note the following to assist in understanding the above state-keeping, which summarizes other normative requirements in the respective sections:

- The [AddNOC](#) command can only be invoked once per contiguous non-expiring fail-safe timer period, and only if no UpdateNOC command was previously processed within the same fail-safe timer period.
- The [UpdateNOC](#) command can only be invoked once per contiguous non-expiring fail-safe timer period, can only be invoked over a [CASE](#) session, and only if no AddNOC command was previously processed in the same fail-safe timer period.

On creation of the Fail Safe Context a second timer SHALL be created to expire at MaxCumulative-FailsafeSeconds as specified in [Section 11.10.5.4, “BasicCommissioningInfo”](#). This Cumulative Fail Safe Context timer (CFSC timer) serves to limit the lifetime of any particular Fail Safe Context; it SHALL NOT be extended or modified on subsequent invocations of ArmFailSafe associated with this Fail Safe Context. Upon expiry of the CFSC timer, the receiver SHALL execute cleanup behavior equivalent to that of fail-safe timer expiration as detailed in [Section 11.10.7.2.2, “Behavior on expiry of Fail-Safe timer”](#). Termination of the session prior to the expiration of that timer for any reason (including a successful end of commissioning or an expiry of a fail-safe timer) SHALL also delete the CFSC timer.

Behavior on expiry of Fail-Safe timer

If the device is fully powered during the [initial phase of commissioning](#), and the fail-safe timer expires anytime before the [CommissioningComplete](#) command is successfully invoked, then the fail-safe cleanup steps SHALL be executed on the receiver, and the Commissioner and Commissionee SHALL terminate commissioning.

If the device is not fully powered during the [initial phase of commissioning](#), and the fail-safe timer expires anytime before initial phase step 17, then the fail-safe cleanup steps SHALL be executed on the receiver, and the Commissioner and Commissionee SHALL terminate commissioning.

If the device was not fully powered during the [initial phase of commissioning](#), and the fail-safe

timer expires anytime during the [setup phase of commissioning](#), then the fail-safe cleanup steps SHALL NOT be executed and it is up to the receiver to decide whether and when to retry the [setup phase of commissioning](#).

The fail-safe cleanup steps are, in order:

1. Disarm the failsafe timer.
2. Terminate any open [PASE](#) secure session by clearing any associated [Section 4.13.3.1, “Secure Session Context”](#) at the Server.
3. Revoke the temporary administrative privileges granted to any open [PASE](#) session (see [Section 6.6.2.9, “Bootstrapping of the Access Control Cluster”](#)) at the Server.
4. If an AddNOC or UpdateNOC command has been successfully invoked, terminate all [CASE](#) sessions associated with the Fabric whose Fabric Index is recorded in the Fail-Safe context (see [Section 11.10.7.2, “ArmFailSafe”](#)) by clearing any associated [Section 4.13.3.1, “Secure Session Context”](#) at the Server.
5. Reset the configuration of all Network Commissioning [Section 11.9.6.2, “Networks”](#) attribute to their state prior to the Fail-Safe being armed.
6. If an UpdateNOC command had been successfully invoked, revert the state of operational key pair, NOC and ICAC for that Fabric to the state prior to the Fail-Safe timer being armed, for the Fabric Index that was the subject of the UpdateNOC command.
7. If an AddNOC command had been successfully invoked, achieve the equivalent effect of invoking the [RemoveFabric command](#) against the [fabric-index](#) stored in the Fail-Safe Context for the Fabric Index that was the subject of the AddNOC command. This SHALL remove all associations to that Fabric including all fabric-scoped data, and MAY possibly factory-reset the device depending on current device state. This SHALL only apply to Fabrics added during the fail-safe period as the result of the [AddNOC command](#).
8. If the [CSRRequest](#) command had been successfully invoked, but no AddNOC or UpdateNOC command had been successfully invoked, then the new operational key pair temporarily generated for the purposes of NOC addition or update (see [Section 6.4.6.1, “Node Operational Certificate Signing Request \(NOCSR\) Procedure”](#)) SHALL be removed as it is no longer needed.
9. Remove any RCACs added by the [AddTrustedRootCertificate command](#) that are not currently referenced by any entry in the [Fabrics attribute](#).
10. Reset the [Section 11.10.6.1, “Breadcrumb”](#) attribute to zero.
11. Optionally: if no factory-reset resulted from the previous steps, it is RECOMMENDED that the Node rollback the state of all non fabric-scoped data present in the Fail-Safe context.

11.10.7.3. ArmFailSafeResponse Command

This command is used to report the result of the ArmFailSafe command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ErrorCode	CommissioningErrorEnum			OK	M
1	DebugText	string	max 128		""	M

ErrorCode Field

This field SHALL contain the result of the operation, based on the behavior specified in the functional description of the ArmFailSafe command.

DebugText Field

See [Section 11.10.7.1, “Common fields in General Commissioning cluster responses”](#).

11.10.7.4. SetRegulatoryConfig Command

This command is used to set the regulatory configuration for the device.

This SHALL add or update the regulatory configuration in [the RegulatoryConfig Attribute](#) to the value provided in the NewRegulatoryConfig field.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NewRegulatoryConfig	RegulatoryLocationTypeEnum				M
1	CountryCode	string	2			M
2	Breadcrumb	uint64				M

Success or failure of this command SHALL be communicated by the [SetRegulatoryConfigResponse](#) command, unless some data model validations caused a failure status code to be issued during the processing of the command.

The CountryCode field SHALL conform to ISO 3166-1 alpha-2 and SHALL be used to set the [Location](#) attribute reflected by the [Basic Information Cluster](#).

If the server limits some of the values (e.g. locked to a particular country, with no regulatory data for others), then setting regulatory information outside a valid country or location SHALL still set the [Location](#) attribute reflected by the [Basic Information Cluster](#) configuration, but the SetRegulatoryConfigResponse replied SHALL have the ErrorCode field set to [ValueOutsideRange](#) error.

If the [LocationCapability](#) attribute is not Indoor/Outdoor and the NewRegulatoryConfig value received does not match either the Indoor or Outdoor fixed value in LocationCapability, then the SetRegulatoryConfigResponse replied SHALL have the ErrorCode field set to [ValueOutsideRange](#)

error and the [RegulatoryConfig](#) attribute and associated internal radio configuration SHALL remain unchanged.

If the [LocationCapability](#) attribute is set to Indoor/Outdoor, then the [RegulatoryConfig](#) attribute SHALL be set to match the NewRegulatoryConfig field.

On successful execution of the command, the ErrorCode field of the [SetRegulatoryConfigResponse](#) SHALL be set to OK.

The Breadcrumb field SHALL be used to atomically set the [Breadcrumb](#) attribute on success of this command, when SetRegulatoryConfigResponse has the ErrorCode field set to OK. If the command fails, the [Breadcrumb](#) attribute SHALL be left unchanged.

11.10.7.5. SetRegulatoryConfigResponse Command

This command is used to report the result of the SetRegulatoryConfig command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ErrorCode	CommissioningErrorEnum			OK	M
1	DebugText	string			""	M

ErrorCode Field

This field SHALL contain the result of the operation, based on the behavior specified in the functional description of the SetRegulatoryConfig command.

DebugText Field

See [Section 11.10.7.1, “Common fields in General Commissioning cluster responses”](#).

11.10.7.6. CommissioningComplete Command

This command is used to indicate that the commissioning process is complete.

Success or failure of this command SHALL be communicated by the [CommissioningCompleteResponse](#) command, unless some data model validations caused a failure status code to be issued during the processing of the command.

This command signals the Server that the Commissioner or Administrator has successfully completed all steps needed during the Fail-Safe period, such as commissioning (see [Section 5.5, “Commissioning Flows”](#)) or other Administrator operations requiring usage of the Fail Safe timer. It ensures that the Server is configured in a state such that it still has all necessary elements to be fully operable within a Fabric, such as ACL entries (see [Section 9.10, “Access Control Cluster”](#)) and operational credentials (see [Section 6.4, “Node Operational Credentials Specification”](#)), and that the Node is reachable using CASE (see [Section 4.14.2, “Certificate Authenticated Session Establishment \(CASE\)”](#)) over an operational network.

An ErrorCode of [NoFailSafe](#) SHALL be responded to the invoker if the CommissioningComplete command was received when no [Fail-Safe context](#) exists.

If [Terms and Conditions](#) are required, then an ErrorCode of [TCAcknowledgementsNotReceived](#) SHALL be responded to the invoker if the user acknowledgements to the required Terms and Conditions have not been provided.

This command is fabric-scoped, so cannot be issued over a session that does not have an associated fabric, i.e. over PASE session prior to an AddNOC command. In addition, this command is only permitted over CASE and must be issued by a node associated with the ongoing Fail-Safe context. An ErrorCode of [InvalidAuthentication](#) SHALL be responded to the invoker if the CommissioningComplete command was received outside a [CASE](#) session (e.g., over Group messaging, or [PASE](#) session after AddNOC), or if the accessing fabric is not the one associated with the ongoing Fail-Safe context.

This command SHALL only result in success with an ErrorCode value of OK in the [CommissioningCompleteResponse](#) if received over a [CASE](#) session and the [accessing fabric index](#) matches the Fabric Index associated with the current Fail-Safe context. In other words:

- If no AddNOC command had been successfully invoked, the CommissioningComplete command must originate from the Fabric that initiated the Fail-Safe context.
- After an AddNOC command has been successfully invoked, the CommissioningComplete command must originate from the Fabric which was joined through the execution of that command, which updated the Fail-Safe context's Fabric Index.

On successful execution of the CommissioningComplete command, where the CommissioningCompleteResponse has an ErrorCode of OK, the following actions SHALL be undertaken on the Server:

1. The [Fail-Safe timer](#) associated with the current Fail-Safe context SHALL be disarmed.
2. The commissioning window at the Server SHALL be closed.
3. Any temporary administrative privileges automatically granted to any open [PASE](#) session SHALL be revoked (see [Section 6.6.2.9, "Bootstrapping of the Access Control Cluster"](#)).
4. The [Secure Session Context](#) of any [PASE](#) session still established at the Server SHALL be cleared.
5. The [Breadcrumb](#) attribute SHALL be reset to zero.

After receipt of a CommissioningCompleteResponse with an ErrorCode value of OK, a client cannot expect any previously established [PASE](#) session to still be usable, due to the server having cleared such sessions.

11.10.7.7. CommissioningCompleteResponse Command

This command is used to report the result of the CommissioningComplete command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ErrorCode	CommissioningErrorEnum			OK	M
1	DebugText	string			""	M

ErrorCode Field

This field SHALL contain the result of the operation, based on the behavior specified in the functional description of the CommissioningComplete command.

DebugText Field

See [Section 11.10.7.1, “Common fields in General Commissioning cluster responses”](#).

11.10.7.8. SetTCAcknowledgements Command

This command is used to set the user acknowledgements received in the Enhanced Setup Flow Terms & Conditions into the node.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	TCVersion	uint16	all			M
1	TCUserResponse	map16	all			M

TCVersion field

This field SHALL contain the version of the Enhanced Setup Flow Terms & Conditions that were presented to the user.

TCUserResponse field

This field SHALL contain the user responses to the Enhanced Setup Flow Terms & Conditions as a map where each bit set in the bitmap corresponds to an accepted term in the file located at [Section 11.23.6.23, “EnhancedSetupFlowTCUrl”](#).

Effect on Receipt

This command SHALL copy the user responses and accepted version to the presented Enhanced Setup Flow Terms & Conditions from the values provided in the TCUserResponse and TCVersion fields to the [TCAcknowledgements Attribute](#) and the [TCAcceptedVersion Attribute](#) fields respectively.

This command SHALL result in success with an ErrorCode value of OK in the [SetTCAcknowledgementsResponse](#) if all required terms were accepted by the user. Specifically, all bits have a value of 1 in TCAcknowledgements whose ordinal is marked as required in the file located at [EnhancedSe-](#)

[tupFlowTCUrl](#).

If the TCVersion field is less than the [TCMinRequiredVersion](#), then the ErrorCode of [TCMinVersion-NotMet](#) SHALL be returned and [TCAcknowledgements](#) SHALL remain unchanged.

If TCVersion is greater than or equal to TCMinRequiredVersion, but the TCUserResponse value indicates that not all required terms were accepted by the user, then the ErrorCode of [RequiredTCNotAccepted](#) SHALL be returned and [TCAcknowledgements](#) SHALL remain unchanged.

11.10.7.9. SetTCAcknowledgementsResponse Command

This command is used to report the result of the SetTCAcknowledgements command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ErrorCode	CommissioningErrorEnum			OK	M

ErrorCode Field

This field SHALL contain the result of the operation, based on the behavior specified in the [functional description of the SetTCAcknowledgements command](#).

11.11. Diagnostic Logs Cluster

This Cluster supports an interface to a Node. It provides commands for retrieving unstructured diagnostic logs from a Node that may be used to aid in diagnostics. It will often be the case that unstructured diagnostic logs will be Node-wide and not specific to any subset of Endpoints. When present, this Cluster SHALL be implemented once for the Node. The Node SHOULD also implement the BDX Initiator and BDX Sender roles as defined in the [BDX Protocol](#).

11.11.1. Revision History

The global ClusterRevision Attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.11.2. Classification

Hierarchy	Role	Scope	PICS Code	Quality
Base	Utility	Node	DLOG	K

11.11.3. Cluster ID

ID	Name
0x0032	Diagnostic Logs

11.11.4. Data Types

11.11.4.1. IntentEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	EndUserSupport	Logs to be used for end-user support	M
1	NetworkDiag	Logs to be used for network diagnostics	M
2	CrashLogs	Obtain crash logs from the Node	M

EndUserSupport Value

SHALL indicate that the purpose of the log request is to retrieve logs for the intention of providing support to an end-user.

NetworkDiag Value

SHALL indicate that the purpose of the log request is to diagnose the network(s) for which the Node is currently commissioned (and/or connected) or has previously been commissioned (and/or connected).

CrashLogs Value

SHALL indicate that the purpose of the log request is to retrieve any crash logs that may be present on a Node.

11.11.4.2. StatusEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Success	Successful transfer of logs	M
1	Exhausted	All logs have been transferred	M

Value	Name	Summary	Conformance
2	NoLogs	No logs of the requested type available	M
3	Busy	Unable to handle request, retry later	M
4	Denied	The request is denied, no logs being transferred	M

Success Value

SHALL be used if diagnostic logs will be or are being transferred.

Exhausted Value

SHALL be used when a BDX session is requested, however, all available logs were provided in a [LogContent](#) field.

NoLogs Value

SHALL be used if the Node does not currently have any diagnostic logs of the requested type (Intent) to transfer.

Busy Value

SHALL be used if the Node is unable to handle the request (e.g. in the process of another transfer) and the Client SHOULD re-attempt the request later.

Denied Value

SHALL be used if the Node is denying the current transfer of diagnostic logs for any reason.

11.11.4.3. TransferProtocolEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	ResponsePayload	Logs to be returned as a response	M
1	BDX	Logs to be returned using BDX	M

ResponsePayload Value

SHALL be used by a Client to request that logs are transferred using the [LogContent](#) attribute of the response

BDX Value

SHALL be used by a Client to request that logs are transferred using BDX as defined in [BDX Protocol](#)

11.11.5. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	RetrieveLogsRequest	client ⇒ server	RetrieveLogsResponse	O	M
0x01	RetrieveLogsResponse	client ⇐ server	N		M

11.11.5.1. RetrieveLogsRequest Command

Reception of this command starts the process of retrieving diagnostic logs from a Node.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Intent	IntentEnum	all			M
1	RequestedProtocol	TransferProtocolEnum	all			M
2	TransferFileDesignator	string	max 32			O

Intent Field

This field SHALL indicate why the diagnostic logs are being retrieved from the Node. A Node MAY utilize this field to selectively determine the logs to transfer.

RequestedProtocol Field

This field SHALL be used to indicate how the log transfer is to be realized. If the field is set to BDX, then if the receiving Node supports BDX it SHALL attempt to use BDX to transfer any potential diagnostic logs; if the receiving Node does not support BDX then the Node SHALL follow the requirements defined for a TransferProtocolEnum of ResponsePayload. If this field is set to ResponsePayload the receiving Node SHALL only utilize the [LogContent](#) field of the RetrieveLogsResponse command to transfer diagnostic log information.

TransferFileDesignator Field

This field SHALL be present if the RequestedProtocol is BDX. The TransferFileDesignator SHALL be set as the [File Designator](#) of the BDX transfer if initiated.

Effect on Receipt

On receipt of this command, the Node SHALL respond with a RetrieveLogsResponse command.

If the RequestedProtocol is set to BDX the Node SHOULD immediately realize the RetrieveLogsResponse command by initiating a [BDX Transfer](#), sending a [BDX SendInit](#) message with the File Designator field of the message set to the value of the TransferFileDesignator field of the RetrieveLogsRequest. On reception of a [BDX SendAccept](#) message the Node SHALL send a RetrieveLogsResponse command with a Status field set to Success and proceed with the log transfer over BDX. If a failure StatusReport is received in response to the SendInit message, the Node SHALL send a RetrieveLogsResponse command with a Status of Denied. In the case where the Node is able to fit the entirety of the requested logs within the [LogContent](#) field, the Status field of the RetrieveLogsResponse SHALL be set to Exhausted and a BDX session SHALL NOT be initiated.

If the RequestedProtocol is set to BDX and either the Node does not support BDX or it is not possible for the Node to establish a BDX session, then the Node SHALL utilize the [LogContent](#) field of the RetrieveLogsResponse command to transfer as much of the current logs as it can fit within the response, and the Status field of the RetrieveLogsResponse SHALL be set to Exhausted.

If the RequestedProtocol is set to ResponsePayload the Node SHALL utilize the [LogContent](#) field of the RetrieveLogsResponse command to transfer as much of the current logs as it can fit within the response, and a BDX session SHALL NOT be initiated.

If the RequestedProtocol is set to BDX and there is no TransferFileDesignator the command SHALL fail with a Status Code of INVALID_COMMAND.

If the Intent and/or the RequestedProtocol arguments contain invalid (out of range) values the command SHALL fail with a Status Code of INVALID_COMMAND.

11.11.5.2. RetrieveLogsResponse Command

This SHALL be generated as a response to the RetrieveLogsRequest.

The data for this command is shown in the following.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Status	StatusEnum	all			M
1	LogContent	octstr	max 1024			M
2	UTCTime-Stamp	epoch-us	all			O
3	TimeSince-Boot	systime-us	all			O

Status Field

This field SHALL indicate the result of an attempt to retrieve diagnostic logs.

LogContent Field

This field SHALL be included in the command if the Status field has a value of Success or Exhausted. A Node SHOULD utilize this field to transfer the newest diagnostic log entries. This field SHALL be empty if BDX is requested and the Status field has a value of Success.

UTCTimeStamp Field

This field SHOULD be included in the command if the Status field has a value of Success and the Node maintains a wall clock. When included, the UTCTimeStamp field SHALL contain the value of the oldest log entry in the diagnostic logs that are being transferred.

TimeSinceBoot Field

This field SHOULD be included in the command if the Status field has a value of Success. When included, the TimeSinceBoot field SHALL contain the time of the oldest log entry in the diagnostic logs that are being transferred represented by the number of microseconds since the last time the Node went through a reboot.

11.12. General Diagnostics Cluster

The General Diagnostics Cluster, along with other diagnostics clusters, provide a means to acquire standardized diagnostics metrics that MAY be used by a Node to assist a user or Administrator in diagnosing potential problems. The General Diagnostics Cluster attempts to centralize all metrics that are broadly relevant to the majority of Nodes.

11.12.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	UpTime attribute now mandatory, and added TimeSnapshot command, added DMTEST feature
3	Added DeviceLoadStatus attribute

11.12.2. Classification

Hierarchy	Role	Scope	PICS Code	Quality
Base	Utility	Node	DGGEN	K

11.12.3. Cluster ID

ID	Name
0x0033	General Diagnostics

11.12.4. Features

The following table indicates the features for this cluster:

Bit	Code	Feature	Conformance	Summary
0	DMTEST	DataModelTest	desc	Support specific testing needs for extended Data Model features

11.12.4.1. DataModelTest Feature

This feature indicates support for extended Data Model testing commands, which are required in some situations.

This feature SHALL be supported if the [MaxPathsPerInvoke](#) attribute of the Basic Information Cluster has a value > 1.

11.12.5. Data Types

11.12.5.1. HardwareFaultEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unspecified	The Node has encountered an unspecified fault.	M
1	Radio	The Node has encountered a fault with at least one of its radios.	O
2	Sensor	The Node has encountered a fault with at least one of its sensors.	O
3	ResettableOverTemp	The Node has encountered an over-temperature fault that is resettable.	O
4	NonResettableOverTemp	The Node has encountered an over-temperature fault that is not resettable.	O

Value	Name	Summary	Conformance
5	PowerSource	The Node has encountered a fault with at least one of its power sources.	O
6	VisualDisplayFault	The Node has encountered a fault with at least one of its visual displays.	O
7	AudioOutputFault	The Node has encountered a fault with at least one of its audio outputs.	O
8	UserInterfaceFault	The Node has encountered a fault with at least one of its user interfaces.	O
9	NonVolatileMemory-Error	The Node has encountered a fault with its non-volatile memory.	O
10	TamperDetected	The Node has encountered disallowed physical tampering.	O

11.12.5.2. RadioFaultEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unspecified	The Node has encountered an unspecified radio fault.	M
1	WiFiFault	The Node has encountered a fault with its Wi-Fi radio.	O
2	CellularFault	The Node has encountered a fault with its cellular radio.	O
3	ThreadFault	The Node has encountered a fault with its 802.15.4 radio.	O

Value	Name	Summary	Conformance
4	NFCFault	The Node has encountered a fault with its NFC radio.	O
5	BLEFault	The Node has encountered a fault with its BLE radio.	O
6	EthernetFault	The Node has encountered a fault with its Ethernet controller.	O

11.12.5.3. NetworkFaultEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unspecified	The Node has encountered an unspecified fault.	M
1	HardwareFailure	The Node has encountered a network fault as a result of a hardware failure.	O
2	NetworkJammed	The Node has encountered a network fault as a result of a jammed network.	O
3	ConnectionFailed	The Node has encountered a network fault as a result of a failure to establish a connection.	O

11.12.5.4. InterfaceTypeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unspecified	Indicates an interface of an unspecified type.	M
1	WiFi	Indicates a Wi-Fi interface.	O
2	Ethernet	Indicates a Ethernet interface.	O

Value	Name	Summary	Conformance
3	Cellular	Indicates a Cellular interface.	O
4	Thread	Indicates a Thread interface.	O

11.12.5.5. BootReasonEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unspecified	The Node is unable to identify the Power-On reason as one of the other provided enumeration values.	M
1	PowerOnReboot	The Node has booted as the result of physical interaction with the device resulting in a reboot.	M
2	BrownOutReset	The Node has rebooted as the result of a brown-out of the Node's power supply.	M
3	SoftwareWatchdogReset	The Node has rebooted as the result of a software watchdog timer.	M
4	HardwareWatchdogReset	The Node has rebooted as the result of a hardware watchdog timer.	M
5	SoftwareUpdateCompleted	The Node has rebooted as the result of a completed software update.	M
6	SoftwareReset	The Node has rebooted as the result of a software initiated reboot.	M

11.12.5.6. NetworkInterface Type

This structure describes a network interface supported by the Node, as provided in the [NetworkInterfaces attribute](#).

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Name	string	max 32				M
1	IsOperational	bool					M
2	Off-PremiseServicesReachableIPv4	bool		X	null		M
3	Off-PremiseServicesReachableIPv6	bool		X	null		M
4	HardwareAddress	hwaddr					M
5	IPv4Addresses	list[ipv4addr]	max 4				M
6	IPv6Addresses	list[ipv6addr]	max 8				M
7	Type	InterfaceTypeEnum					M

Name Field

This field SHALL indicate a human-readable (displayable) name for the network interface, that is different from all other interfaces.

IsOperational Field

This field SHALL indicate if the Node is currently advertising itself operationally on this network interface and is capable of successfully receiving incoming traffic from other Nodes.

OffPremiseServicesReachableIPv4 Field

This field SHALL indicate whether the Node is currently able to reach off-premise services it uses by utilizing IPv4. The value SHALL be null if the Node does not use such services or does not know whether it can reach them.

OffPremiseServicesReachableIPv6 Field

This field SHALL indicate whether the Node is currently able to reach off-premise services it uses by utilizing IPv6. The value SHALL be null if the Node does not use such services or does not know whether it can reach them.

HardwareAddress Field

This field SHALL contain the current link-layer address for a 802.3 or [IEEE 802.11-2020](#) network interface and contain the current extended MAC address for a 802.15.4 interface. The byte order of the [octstr](#) SHALL be in wire byte order. For addresses values less than 64 bits, the first two bytes SHALL be zero.

IPv4Addresses Field

This field SHALL provide a list of the IPv4 addresses that are currently assigned to the network interface.

IPv6Addresses Field

This field SHALL provide a list of the unicast IPv6 addresses that are currently assigned to the network interface. This list SHALL include the Node's link-local address and SHOULD include any assigned GUA and ULA addresses. This list SHALL NOT include any multicast group addresses to which the Node is subscribed.

Type Field

This field SHALL indicate the type of the interface using the [InterfaceTypeEnum](#).

11.12.5.7. DeviceLoadStruct Type

This structure describes some important resource utilization metrics for the Node, as provided in the [DeviceLoadStatus](#) attribute.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	CurrentSubscriptions	uint16	all				M
1	CurrentSubscriptions-ForFabric	uint16	all				M
2	TotalSubscriptions-Established	uint32	all				M
3	TotalInteraction-ModelMessagesSent	uint32	all				M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
4	TotalInteractionModelMessagesReceived	uint32	all				M

NOTE

For all the fields, the value SHALL remain at the maximum representable (clamp to max) if the maximum value is reached.

CurrentSubscriptions Field

This field SHALL indicate the number of currently-active Interaction Model subscriptions across all fabrics on the node.

CurrentSubscriptionsForFabric Field

This field SHALL indicate the number of currently-active Interaction Model subscriptions for the [accessing fabric](#) only. If no accessing fabric is available, this field SHALL be set to zero.

TotalSubscriptionsEstablished Field

This field SHALL indicate the total number of Interaction Model subscriptions successfully established across all fabrics on the node since start-up.

TotalInteractionModelMessagesSent Field

This field SHALL indicate the number of outgoing Interaction Model protocol messages sent since start-up. These are messages that have the [Protocol ID](#) set to [PROTOCOL_ID_INTERACTION_MODEL](#), excluding any retries of such messages.

TotalInteractionModelMessagesReceived Field

This field SHALL indicate the number of incoming Interaction Model protocol messages received since start-up. These are messages that have the [Protocol ID](#) set to [PROTOCOL_ID_INTERACTION_MODEL](#), excluding any retries of such messages.

11.12.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	NetworkInterfaces	list[NetworkInterface]	max 8			R V	M
0x0001	RebootCount	uint16		N		R V	M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0002	UpTime	uint64		C		R V	M
0x0003	TotalOperational-Hours	uint32		C N		R V	O
0x0004	BootReason	BootReasonEnum				R V	O
0x0005	Active-Hardware-Faults	list[HardwareFaultEnum]	max 11			R V	O
0x0006	ActiveRadioFaults	list[RadioFaultEnum]	max 7			R V	O
0x0007	ActiveNetwork-Faults	list[NetworkFaultEnum]	max 4			R V	O
0x0008	TestEvent-TriggersEnabled	bool	all			R V	M
0x0009	DoNotUse						X
0x000A	Device-LoadStatus	Device-LoadStruct	all	C		R V	Rev >= v3

Attribute 0x0009 SHALL NOT be used in any implementation of previous, current or future version of this specification.

11.12.6.1. NetworkInterfaces Attribute

The NetworkInterfaces attribute SHALL be a list of [NetworkInterface](#) structs. Each logical network interface on the Node SHALL be represented by a single entry within the NetworkInterfaces attribute.

11.12.6.2. RebootCount Attribute

The RebootCount attribute SHALL indicate a best-effort count of the number of times the Node has rebooted. The RebootCount attribute SHOULD be incremented each time the Node reboots. The RebootCount attribute SHALL NOT be incremented when a Node wakes from a low-power or sleep state. The RebootCount attribute SHALL only be reset to 0 upon a factory reset of the Node.

11.12.6.3. UpTime Attribute

The UpTime attribute SHALL indicate a best-effort assessment of the length of time, in seconds,

since the Node's last reboot. This attribute SHOULD be incremented to account for the periods of time that a Node is in a low-power or sleep state. This attribute SHALL only be reset upon a device reboot. This attribute SHALL be based on the same System Time source as those used to fulfill any usage of the [systime-us](#) and [systime-ms](#) data types within the server.

11.12.6.4. TotalOperationalHours Attribute

The TotalOperationalHours attribute SHALL indicate a best-effort attempt at tracking the length of time, in hours, that the Node has been operational. The TotalOperationalHours attribute SHOULD be incremented to account for the periods of time that a Node is in a low-power or sleep state. The TotalOperationalHours attribute SHALL only be reset upon a factory reset of the Node.

11.12.6.5. BootReason Attribute

The BootReason attribute SHALL indicate the reason for the Node's most recent boot.

11.12.6.6. ActiveHardwareFaults Attribute

The ActiveHardwareFaults attribute SHALL indicate the set of faults currently detected by the Node. When the Node detects a fault has been raised, the appropriate [HardwareFaultEnum](#) value SHALL be added to this list. This list SHALL NOT contain more than one instance of a specific [HardwareFaultEnum](#) value. When the Node detects that all conditions contributing to a fault has been cleared, the corresponding [HardwareFaultEnum](#) value SHALL be removed from this list. An empty list SHALL indicate there are currently no active faults. The order of this list SHOULD have no significance. Clients interested in monitoring changes in active faults MAY subscribe to this attribute, or they MAY subscribe to [HardwareFaultChange](#).

11.12.6.7. ActiveRadioFaults Attribute

The ActiveRadioFaults attribute SHALL indicate the set of faults currently detected by the Node. When the Node detects a fault has been raised, the appropriate [RadioFaultEnum](#) value SHALL be added to this list. This list SHALL NOT contain more than one instance of a specific [RadioFaultEnum](#) value. When the Node detects that all conditions contributing to a fault has been cleared, the corresponding [RadioFaultEnum](#) value SHALL be removed from this list. An empty list SHALL indicate there are currently no active faults. The order of this list SHOULD have no significance. Clients interested in monitoring changes in active faults MAY subscribe to this attribute, or they MAY subscribe to [RadioFaultChange](#).

11.12.6.8. ActiveNetworkFaults Attribute

The ActiveNetworkFaults attribute SHALL indicate the set of faults currently detected by the Node. When the Node detects a fault has been raised, the appropriate [NetworkFaultEnum](#) value SHALL be added to this list. This list SHALL NOT contain more than one instance of a specific [NetworkFaultEnum](#) value. When the Node detects that all conditions contributing to a fault has been cleared, the corresponding [NetworkFaultEnum](#) value SHALL be removed from this list. An empty list SHALL indicate there are currently no active faults. The order of this list SHOULD have no significance. Clients interested in monitoring changes in active faults MAY subscribe to this attribute, or they MAY subscribe to [NetworkFaultChange](#).

11.12.6.9. TestEventTriggersEnabled Attribute

The TestEventTriggersEnabled attribute SHALL indicate whether the Node has any [TestEventTrigger](#) configured. When this attribute is true, the Node has been configured with one or more test event triggers by virtue of the internally programmed EnableKey value (see [Section 11.12.7.1, “TestEventTrigger Command”](#)) being set to a non-zero value. This attribute can be used by Administrators to detect if a device was inadvertently commissioned with test event trigger mode enabled, and take appropriate action (e.g. warn the user and/or offer to remove all fabrics on the Node).

11.12.6.10. DeviceLoadStatus Attribute

This attribute SHALL indicate the status of various resources used.

11.12.7. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	TestEventTrigger	client ⇒ server	Y	M	M
0x01	TimeSnapshot	client ⇒ server	TimeSnapshotResponse	O	M
0x02	TimeSnapshotResponse	client ⇐ server	N		M
0x03	Payload-TestRequest	client ⇒ server	PayloadTestResponse	M	DMTEST
0x04	Payload-TestResponse	server ⇒ client	N		DMTEST

11.12.7.1. TestEventTrigger Command

This command SHALL be supported to provide a means for certification tests to trigger some test-plan-specific events, necessary to assist in automation of device interactions for some certification test cases. This command SHALL NOT cause any changes to the state of the device that persist after the last fabric is removed.

The fields for the TestEventTrigger command are as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	EnableKey	octstr	16			M
1	EventTrigger	uint64				M

EnableKey Field

The EnableKey is a 128 bit value provided by the client in this command, which needs to match a value chosen by the manufacturer and configured on the server using manufacturer-specific

means, such as pre-provisioning. The value of all zeroes is reserved to indicate that no EnableKey is set. Therefore, if the EnableKey field is received with all zeroes, this command SHALL FAIL with a response status of CONSTRAINT_ERROR.

The EnableKey SHOULD be unique per exact set of devices going to a certification test.

Devices not targeted towards going to a certification test event SHALL NOT have a non-zero EnableKey value configured, so that only devices in test environments are responsive to this command.

In order to prevent unwittingly actuating a particular trigger, this command SHALL respond with a response status of CONSTRAINT_ERROR if the EnableKey field does not match the a-priori value configured on the device.

EventTrigger Field

This field SHALL indicate the test or test mode which the client wants to trigger.

The expected side-effects of EventTrigger values are out of scope of this specification and will be described within appropriate certification test literature provided to manufacturers by the Connectivity Standards Alliance, in conjunction with certification test cases documentation.

Values of EventTrigger in the range 0xFFFF_FFFF_0000_0000 through 0xFFFF_FFFF_FFFF_FFFF are reserved for testing use by manufacturers and will not appear in the Connectivity Standards Alliance certification test literature.

If the value of EventTrigger received is not supported by the receiving Node, this command SHALL fail with a status code of INVALID_COMMAND.

Otherwise, if the EnableKey value matches the configured internal value for a particular Node, and the EventTrigger value matches a supported test event trigger value, the command SHALL succeed and execute the expected trigger action.

If no specific test event triggers are required to be supported by certification test requirements for the features that a given product will be certified against, this command MAY always fail with the INVALID_COMMAND status, equivalent to the situation of receiving an unknown EventTrigger, for all possible EventTrigger values.

11.12.7.2. TimeSnapshot Command

This command MAY be used by a client to obtain a correlated view of both System Time, and, if currently synchronized and supported, "wall clock time" of the server. This can help clients establish time correlation between their concept of time and the server's concept of time. This is especially useful when processing event histories where some events only contain System Time.

Upon command invocation, the server SHALL respond with a TimeSnapshotResponse.

11.12.7.3. TimeSnapshotResponse Command

This command SHALL be generated in response to a TimeSnapshot command.

When generating this response, all fields SHALL be gathered as close together in time as possible, so that the time jitter between the values is minimized.

If the Time Synchronization cluster is supported by the node, the PosixTimeMs field SHALL NOT be null unless the [UTCTime](#) attribute in the [Time Synchronization](#) cluster is also null.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0x00	System-TimeMs	systime-ms	all		MS	M
0x01	Posix-TimeMs	posix-ms	all	X	null	M

SystemTimeMs Field

This SHALL indicate the current System Time in milliseconds (type [systime-ms](#)), with the value taken at the time of processing of the TimeSnapshot command that generated this response.

The value SHALL be taken from the same clock which populates the [Timestamp field](#) in events when using System Time for the field.

PosixTimeMs Field

This SHALL indicate the current time in [POSIX Time in milliseconds](#), with the value taken from the same source that could populate the [Timestamp field of events](#). This value SHALL only be null when any the following are true:

- The node doesn't support the Time Synchronization cluster.
- The node's Time Synchronization cluster instance's UTCTime attribute is null.

11.12.7.4. PayloadTestRequest Command

This command provides a means for certification tests or manufacturer's internal tests to validate particular command handling and encoding constraints by generating a response of a given size.

This command SHALL use the same [EnableKey](#) behavior as the [TestEventTrigger](#) command, whereby processing of the command is only enabled when the TestEventTriggersEnabled field is true, which SHALL NOT be true outside of certification testing or manufacturer's internal tests.

The fields for the PayloadTestRequest command are as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0x00	EnableKey	octstr	16			M
0x01	Value	uint8				M
0x02	Count	uint16	max 2048			M

EnableKey field

This field SHALL have the same meaning and usage as the [TestEventTrigger EnableKey field](#).

Value field

This field SHALL indicate the value to use in every byte of the PayloadTestResponse’s Payload field.

Count field

This field SHALL indicate the number of times to repeat the Value in the PayloadTestResponse’s Payload field.

Effect upon receipt

This command SHALL respond with a response status of CONSTRAINT_ERROR if either:

- The EnableKey field does not match the a-priori value configured on the device.
- The TestEventTriggersEnabled field is currently false.

Otherwise, the server SHALL respond with a PayloadTestResponse command with a Payload field value containing **Count** instances of the **Value** byte. If the response is too large to send, the server SHALL fail the command and respond with a response status of RESOURCE_EXHAUSTED.

For example:

- If Value is 0x55 and the Count is zero, then the PayloadTestResponse would have the Payload field set to an empty octet string.
- If Value is 0xA5 and the Count is 10, the PayloadTestResponse would have the Payload field set to a content whose hexadecimal representation would be **A5A5A5A5A5A5A5A5A5A5**, and base64 representation would be **paWlpWlpWlpQ==**.

11.12.7.5. PayloadTestResponse Command

This command is sent by the server on receipt of the PayloadTestRequest command.

This command SHALL have the following data fields:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0x00	Payload	octstr	max 2048			M

Payload Field

This field SHALL contain the computed response of the PayloadTestRequest command.

11.12.8. Events

ID	Name	Priority	Quality	Access	Conformance
0x00	HardwareFaultChange	CRITICAL		V	O
0x01	RadioFaultChange	CRITICAL		V	O
0x02	NetworkFaultChange	CRITICAL		V	O
0x03	BootReason	CRITICAL		V	M

11.12.8.1. HardwareFaultChange Event

The HardwareFaultChange Event SHALL indicate a change in the set of hardware faults currently detected by the Node.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Current	list[HardwareFaultEnum]	max 11			M
1	Previous	list[HardwareFaultEnum]	max 11			M

Current Field

This field SHALL represent the set of faults currently detected, as per [HardwareFaultEnum](#).

Previous Field

This field SHALL represent the set of faults detected prior to this change event, as per [HardwareFaultEnum](#).

11.12.8.2. RadioFaultChange Event

The RadioFaultChange Event SHALL indicate a change in the set of radio faults currently detected by the Node.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Current	list[RadioFaultEnum]	max 7			M

ID	Name	Type	Constraint	Quality	Fallback	Conformance
1	Previous	list[Radio-FaultEnum]	max 7			M

Current Field

This field SHALL represent the set of faults currently detected, as per [RadioFaultEnum](#).

Previous Field

This field SHALL represent the set of faults detected prior to this change event, as per [RadioFaultEnum](#).

11.12.8.3. NetworkFaultChange Event

The NetworkFaultChange Event SHALL indicate a change in the set of network faults currently detected by the Node.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Current	list[Network-FaultEnum]	max 4			M
1	Previous	list[Network-FaultEnum]	max 4			M

Current Field

This field SHALL represent the set of faults currently detected, as per [NetworkFaultEnum](#).

Previous Field

This field SHALL represent the set of faults detected prior to this change event, as per [NetworkFaultEnum](#).

11.12.8.4. BootReason Event

The BootReason Event SHALL indicate the reason that caused the device to start-up.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	BootReason	BootReasonEnum				M

BootReason Field

This field SHALL contain the reason for this BootReason event.

11.13. Software Diagnostics Cluster

The Software Diagnostics Cluster provides a means to acquire standardized diagnostics metrics that MAY be used by a Node to assist a user or Administrator in diagnosing potential problems. The Software Diagnostics Cluster attempts to centralize all metrics that are relevant to the software that may be running on a Node.

11.13.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.13.2. Classification

Hierarchy	Role	Scope	PICS Code	Quality
Base	Utility	Node	DGSW	K

11.13.3. Cluster ID

ID	Name
0x0034	Software Diagnostics

11.13.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Summary
0	WTRMRK	Watermarks	Node makes available the metrics for high watermark related to memory consumption.

11.13.5. Data Types

11.13.5.1. ThreadMetricsStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	ID	uint64	all				M
1	Name	string	max 8		empty		O
2	Stack-FreeCurrent	uint32	all		MS		O
3	Stack-FreeMinimum	uint32	all		MS		O
4	StackSize	uint32	all		MS		O

ID Field

The Id field SHALL be a server-assigned per-thread unique ID that is constant for the duration of the thread. Efforts SHOULD be made to avoid reusing ID values when possible.

Name Field

The Name field SHALL be set to a vendor defined name or prefix of the software thread that is static for the duration of the thread.

StackFreeCurrent Field

The StackFreeCurrent field SHALL indicate the current amount of stack memory, in bytes, that are not being utilized on the respective thread.

StackFreeMinimum Field

The StackFreeMinimum field SHALL indicate the minimum amount of stack memory, in bytes, that has been available at any point between the current time and this attribute being reset or initialized on the respective thread. This value SHALL only be reset upon a Node reboot or upon receiving of the [ResetWatermarks](#) command.

StackSize Field

The StackSize field SHALL indicate the amount of stack memory, in bytes, that has been allocated for use by the respective thread.

11.13.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Thread-Metrics	list[ThreadMetric-Struct]	max 64	C		R V	O

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0001	CurrentHeapFree	uint64		C		R V	O
0x0002	CurrentHeapUsed	uint64		C		R V	O
0x0003	CurrentHeapHighWatermark	uint64		C		R V	WTRMRK

11.13.6.1. ThreadMetrics Attribute

This attribute SHALL be a list of [ThreadMetricsStruct](#) structs. Each active thread on the Node SHALL be represented by a single entry within the ThreadMetrics attribute.

11.13.6.2. CurrentHeapFree Attribute

This attribute SHALL indicate the current amount of heap memory, in bytes, that are free for allocation. The effective amount MAY be smaller due to heap fragmentation or other reasons.

11.13.6.3. CurrentHeapUsed Attribute

This attribute SHALL indicate the current amount of heap memory, in bytes, that is being used.

11.13.6.4. CurrentHeapHighWatermark Attribute

This attribute SHALL indicate the maximum amount of heap memory, in bytes, that has been used by the Node. This value SHALL only be reset upon a Node reboot or upon receiving of the [ResetWatermarks](#) command.

11.13.7. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	ResetWatermarks	client ⇒ server	Y	M	WTRMRK

11.13.7.1. ResetWatermarks Command

This command is used to reset the high watermarks for heap and stack memory.

Receipt of this command SHALL reset the following values which track high and lower watermarks:

- The StackFreeMinimum field of the [ThreadMetrics](#) attribute
- The [CurrentHeapHighWatermark](#) attribute

Effect on Receipt

On receipt of this command, the Node SHALL make the following modifications to attributes it supports:

If implemented, the server SHALL set the value of the [CurrentHeapHighWatermark](#) attribute to the value of the CurrentHeapUsed attribute.

If implemented, the server SHALL set the value of the [StackFreeMinimum](#) field for every thread to the value of the corresponding thread's [StackFreeCurrent](#) field.

11.13.8. Events

ID	Name	Priority	Quality	Access	Conformance
0x00	SoftwareFault	INFO		V	O

11.13.8.1. SoftwareFault Event

This Event SHALL be generated when a software fault occurs on the Node.

The event's data are as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ID	uint64	all		0	M
1	Name	string	max 8		empty	O
2	Fault-Recording	octstr	max 1024		empty	O

ID Field

This field SHALL be set to the ID of the software thread in which the last software fault occurred.

Name Field

This field SHALL be set to a manufacturer-specified name or prefix of the software thread in which the last software fault occurred.

FaultRecording Field

This field SHALL be a manufacturer-specified payload intended to convey information to assist in further diagnosing or debugging a software fault. The FaultRecording field MAY be used to convey information such as, but not limited to, thread backtraces or register contents.

11.14. Thread Network Diagnostics Cluster

The Thread Network Diagnostics Cluster provides a means to acquire standardized diagnostics metrics that MAY be used by a Node to assist a user or Administrator in diagnosing potential problems.

The Thread Network Diagnostics Cluster attempts to centralize all metrics that are relevant to a potential Thread radio running on a Node.

11.14.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Remove optionality from FrameErrorRate and MessageErrorRate
3	Add ExtAddress and Rloc16 attributes

11.14.2. Classification

Hierarchy	Role	Scope	PICS Code	Quality
Base	Utility	Node	DGTHREAD	K

11.14.3. Cluster ID

ID	Name
0x0035	Thread Network Diagnostics

11.14.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Summary
0	PKTCNT	PacketCounts	Server supports the counts for the number of received and transmitted packets on the Thread interface.
1	ERRCNT	ErrorCounts	Server supports the counts for the number of errors that have occurred during the reception and transmission of packets on the Thread interface.
2	MLECNT	MLECounts	Server supports the counts for various MLE layer happenings.

Bit	Code	Feature	Summary
3	MACCNT	MACCounts	Server supports the counts for various MAC layer happenings.

11.14.5. Data Types

11.14.5.1. NetworkFaultEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unspecified	Indicates an unspecified fault.	M
1	LinkDown	Indicates the Thread link is down.	M
2	HardwareFailure	Indicates there has been Thread hardware failure.	M
3	NetworkJammed	Indicates the Thread network is jammed.	M

11.14.5.2. ConnectionStatusEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Connected	Node is connected	M
1	NotConnected	Node is not connected	M

11.14.5.3. RoutingRoleEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unspecified	Unspecified routing role.	M
1	Unassigned	The Node does not currently have a role as a result of the Thread interface not currently being configured or operational.	M

Value	Name	Summary	Conformance
2	SleepyEndDevice	The Node acts as a Sleepy End Device with RX-off-when-idle sleepy radio behavior.	M
3	EndDevice	The Node acts as an End Device without RX-off-when-idle sleepy radio behavior.	M
4	REED	The Node acts as an Router Eligible End Device.	M
5	Router	The Node acts as a Router Device.	M
6	Leader	The Node acts as a Leader Device.	M

11.14.5.4. NeighborTableStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	ExtAddress	uint64	all				M
1	Age	uint32	all				M
2	Rloc16	uint16	all				M
3	Link-Frame-Counter	uint32	all				M
4	MleFrame-Counter	uint32	all				M
5	LQI	uint8	0 to 255				M
6	AverageRssi	int8	-128 to 0	X	null		M
7	LastRssi	int8	-128 to 0	X	null		M
8	FrameErrorRate	uint8	0 to 100		0		M
9	MessageErrorRate	uint8	0 to 100		0		M
10	RxOn-WhenIdle	bool	all				M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
11	FullThreatDevice	bool	all				M
12	FullNetworkData	bool	all				M
13	IsChild	bool	all				M

ExtAddress Field

This field SHALL specify the IEEE 802.15.4 extended address for the neighboring Node. The uint64 value is composed by taking the 8 octets of the extended address EUI-64 and treating them as a big-endian integer. For example, octet string (in hexadecimal, from first octet to last) **00112233AABBCCDD** would lead to a value of 0x00112233AABBCCDD.

Age Field

This field SHALL specify the duration of time, in seconds, since a frame has been received from the neighboring Node.

Rloc16 Field

This field SHALL specify the RLOC16 of the neighboring Node. The uint16 value is composed by taking the two RLOC16 and treating the octet string as if it was encoding a big-endian integer. For example, octet string (in hexadecimal, from first octet to last) **44AA** would lead to a value of 0x44AA.

LinkFrameCounter Field

This field SHALL specify the number of link layer frames that have been received from the neighboring node. This field SHALL be reset to 0 upon a reboot of the Node.

MleFrameCounter Field

This field SHALL specify the number of Mesh Link Establishment frames that have been received from the neighboring node. This field SHALL be reset to 0 upon a reboot of the Node.

LQI Field

This field SHALL specify the implementation specific mix of IEEE 802.15.4 PDU receive quality indicators, scaled from 0 to 255.

AverageRssi Field

This field SHOULD specify the average RSSI across all received frames from the neighboring Node since the receiving Node's last reboot. If there is no known received frames this field SHOULD have the value of null. This field SHALL have the units of dBm, having the range -128 dBm to 0 dBm.

LastRssi Field

This field SHALL specify the RSSI of the most recently received frame from the neighboring Node. If

there is no known last received frame the LastRssi field SHOULD have the value of null. This field SHALL have the units of dBm, having the range -128 dBm to 0 dBm.

FrameErrorRate Field

This field SHALL specify the percentage of received frames from the neighboring Node that have resulted in errors.

MessageErrorRate Field

This field SHALL specify the percentage of received messages from the neighboring Node that have resulted in errors.

RxOnWhenIdle Field

This field SHALL specify if the neighboring Node is capable of receiving frames while the Node is in an idle state.

FullThreadDevice Field

This field SHALL specify if the neighboring Node is a full Thread device.

FullNetworkData Field

This field SHALL specify if the neighboring Node requires the full Network Data. If set to False, the neighboring Node only requires the stable Network Data.

IsChild Field

This field SHALL specify if the neighboring Node is a direct child of the Node reporting the [NeighborTable](#) attribute.

11.14.5.5. RouteTableStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	ExtAddress	uint64					M
1	Rloc16	uint16					M
2	RouterId	uint8					M
3	NextHop	uint8					M
4	PathCost	uint8					M
5	LQIIn	uint8					M
6	LQIOut	uint8					M
7	Age	uint8					M
8	Allocated	bool					M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
9	LinkEstablished	bool					M

ExtAddress Field

This field SHALL specify the IEEE 802.15.4 extended address for the Node for which this route table entry corresponds. The uint64 value is composed by taking the 8 octets of the extended address EUI-64 and treating them as a big-endian integer. For example, octet string (in hexadecimal, from first octet to last) **00112233AABBCCDD** would lead to a value of 0x00112233AABBCCDD.

Rloc16 Field

This field SHALL specify the RLOC16 for the Node for which this route table entry corresponds. The uint16 value is composed by taking the two RLOC16 and treating the octet string as if it was encoding a big-endian integer. For example, octet string (in hexadecimal, from first octet to last) **44AA** would lead to a value of 0x44AA.

RouterId Field

This field SHALL specify the Router ID for the Node for which this route table entry corresponds.

NextHop Field

This field SHALL specify the Router ID for the next hop in the route to the Node for which this route table entry corresponds.

PathCost Field

This Field SHALL specify the cost of the route to the Node for which this route table entry corresponds.

LQIIn Field

This field SHALL specify the implementation specific mix of IEEE 802.15.4 PDU receive quality indicators, scaled from 0 to 255, from the perspective of the Node reporting the neighbor table.

LQIOut Field

This field SHALL specify the implementation specific mix of IEEE 802.15.4 PDU receive quality indicators, scaled from 0 to 255, from the perspective of the Node specified within the NextHop field.

Age Field

This field SHALL specify the duration of time, in seconds, since a frame has been received from the Node for which this route table entry corresponds.

Allocated Field

This field SHALL specify if the router ID as defined within the RouterId field has been allocated.

LinkEstablished Field

This field SHALL specify if a link has been established to the Node for which this route table entry corresponds.

11.14.5.6. SecurityPolicy Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Rotation-Time	uint16					M
1	Flags	uint16					M

RotationTime Field

This field SHALL specify the interval of time, in hours, that Thread security keys are rotated. This attribute SHALL be null when there is no dataset configured.

Flags Field

This field SHALL specify the flags as specified in Thread 1.3.0 section 8.10.1.15. This attribute SHALL be null when there is no dataset configured.

11.14.5.7. OperationalDatasetComponents Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Active-Time-stampPresent	bool					M
1	Pending-Time-stampPresent	bool					M
2	MasterKeyPresent	bool					M
3	Network-NamePresent	bool					M
4	Extended-PanIdPresent	bool					M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
5	MeshLocalPrefixPresent	bool					M
6	DelayPresent	bool					M
7	PanIdPresent	bool					M
8	ChannelPresent	bool					M
9	PskcPresent	bool					M
10	SecurityPolicyPresent	bool					M
11	ChannelMaskPresent	bool					M

ActiveTimestampPresent Field

This field SHALL be True if the Node has an active timestamp present, else False.

PendingTimestampPresent Field

This field SHALL be True if the Node has a pending timestamp is present, else False.

MasterKeyPresent Field

This field SHALL be True if the Node has the Thread master key, else False.

NetworkNamePresent Field

This field SHALL be True if the Node has the Thread network's name, else False.

ExtendedPanIdPresent Field

This field SHALL be True if the Node has an extended Pan ID, else False.

MeshLocalPrefixPresent Field

This field SHALL be True if the Node has the mesh local prefix, else False.

DelayPresent Field

This field SHALL be True if the Node has the Thread network delay set, else False.

PanIdPresent Field

This field SHALL be True if the Node has a Pan ID, else False.

ChannelPresent Field

This field SHALL be True if the Node has configured an operational channel for the Thread network, else False.

PskcPresent Field

This field SHALL be True if the Node has been configured with the Thread network Pskc, else False.

SecurityPolicyPresent Field

This field SHALL be True if the Node has been configured with the Thread network security policies, else False.

ChannelMaskPresent Field

This field SHALL be True if the Node has available a mask of available channels, else False.

11.14.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Channel	uint16	all	X		R V	M
0x0001	Routing-Role	RoutingRoleEnum	all	X		R V	M
0x0002	Network-Name	string	max 16	X		R V	M
0x0003	PanId	uint16	all	X		R V	M
0x0004	Extended-PanId	uint64	all	X		R V	M
0x0005	MeshLocalPrefix	ipv6pre	all	X		R V	M
0x0006	Overrun-Count	uint64	all	C	0	R V	ERRCNT
0x0007	NeighborTable	list[NeighborTableStruct]	all		[]	R V	M
0x0008	RouteTable	list[RouteTableStruct]	all		[]	R V	M
0x0009	PartitionId	uint32	all	X		R V	M
0x000A	Weighting	uint16	max 255	X		R V	M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x000B	DataVersion	uint16	max 255	X		R V	M
0x000C	Stable-DataVersion	uint16	max 255	X		R V	M
0x000D	Leader-RouterId	uint8	max 62	X		R V	M
0x000E	DetachedRoleCount	uint16	all	C	0	R V	[MLECNT]
0x000F	ChildRoleCount	uint16	all	C	0	R V	[MLECNT]
0x0010	Router-RoleCount	uint16	all	C	0	R V	[MLECNT]
0x0011	Leader-RoleCount	uint16	all	C	0	R V	[MLECNT]
0x0012	AttachAttempt-Count	uint16	all	C	0	R V	[MLECNT]
0x0013	Partition-IdChange-Count	uint16	all	C	0	R V	[MLECNT]
0x0014	BetterPartitionAttachAttempt-Count	uint16	all	C	0	R V	[MLECNT]
0x0015	ParentChange Count	uint16	all	C	0	R V	[MLECNT]
0x0016	TxTotal-Count	uint32	all	C	0	R V	[MACCNT]
0x0017	TxUnicast-Count	uint32	all	C	0	R V	[MACCNT]
0x0018	TxBroadcastCount	uint32	all	C	0	R V	[MACCNT]
0x0019	TxAckRequested-Count	uint32	all	C	0	R V	[MACCNT]

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x001A	TxAcked-Count	uint32	all	C	0	R V	[MACCNT]
0x001B	TxNoAck-Request-Count	uint32	all	C	0	R V	[MACCNT]
0x001C	TxData-Count	uint32	all	C	0	R V	[MACCNT]
0x001D	TxDat-aPoll-Count	uint32	all	C	0	R V	[MACCNT]
0x001E	TxBeacon-Count	uint32	all	C	0	R V	[MACCNT]
0x001F	TxBeacon-Request-Count	uint32	all	C	0	R V	[MACCNT]
0x0020	TxOther-Count	uint32	all	C	0	R V	[MACCNT]
0x0021	TxRetryCount	uint32	all	C	0	R V	[MACCNT]
0x0022	TxDirect-MaxRetry-ExpiryCount	uint32	all	C	0	R V	[MACCNT]
0x0023	TxIndirect-MaxRetry-ExpiryCount	uint32	all	C	0	R V	[MACCNT]
0x0024	TxErrCca-Count	uint32	all	C	0	R V	[MACCNT]
0x0025	TxErrAbortCount	uint32	all	C	0	R V	[MACCNT]
0x0026	TxErrBusy-Channel-Count	uint32	all	C	0	R V	[MACCNT]
0x0027	RxTotal-Count	uint32	all	C	0	R V	[MACCNT]

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0028	RxUnicast-Count	uint32	all	C	0	R V	[MACCNT]
0x0029	RxBroadcastCount	uint32	all	C	0	R V	[MACCNT]
0x002A	RxData-Count	uint32	all	C	0	R V	[MACCNT]
0x002B	RxDatapoll-Count	uint32	all	C	0	R V	[MACCNT]
0x002C	RxBeacon-Count	uint32	all	C	0	R V	[MACCNT]
0x002D	RxBeacon-Request-Count	uint32	all	C	0	R V	[MACCNT]
0x002E	RxOther-Count	uint32	all	C	0	R V	[MACCNT]
0x002F	RxAddressFiltered-Count	uint32	all	C	0	R V	[MACCNT]
0x0030	RxDestinationFiltered-Count	uint32	all	C	0	R V	[MACCNT]
0x0031	RxDuplicated-Count	uint32	all	C	0	R V	[MACCNT]
0x0032	RxErrorNoFrame Count	uint32	all	C	0	R V	[MACCNT]
0x0033	RxErrorUnknown-Nighbor-Count	uint32	all	C	0	R V	[MACCNT]
0x0034	RxErrorInvalidSrcAddr-Count	uint32	all	C	0	R V	[MACCNT]
0x0035	RxErrorSec-Count	uint32	all	C	0	R V	[MACCNT]

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0036	RxErrFcsCount	uint32	all	C	0	R V	[MACCNT]
0x0037	RxEr-rOther-Count	uint32	all	C	0	R V	[MACCNT]
0x0038	Active-Time-stamp	uint64	all	X	0	R V	O
0x0039	Pending-Time-stamp	uint64	all	X	0	R V	O
0x003A	Delay	uint32	all	X	0	R V	O
0x003B	Security-Policy	Security-Policy		X		R V	M
0x003C	Channel-Page0-Mask	octstr	4	X		R V	M
0x003D	Operational-Dataset-Components	Operational-Dataset-Components		X		R V	M
0x003E	ActiveNetwork-FaultsList	list[NetworkFault-Enum]	max 4			R V	M
0x003F	ExtAddress	uint64	all	X		R V	P, M
0x0040	Rloc16	uint16	all	X		R V	P, M

11.14.6.1. Channel Attribute

This attribute SHALL indicate the 802.15.4 channel number configured on the Node's Thread interface (that is, the Active Operational Dataset's current Channel value). A value of null SHALL indicate that the Thread interface is not currently configured or operational.

11.14.6.2. RoutingRole Attribute

This attribute SHALL indicate the role that this Node has within the routing of messages through the Thread network, as defined by [RoutingRoleEnum](#). The potential roles are defined in the following table. A value of null SHALL indicate that the Thread interface is not currently configured or operational.

11.14.6.3. NetworkName Attribute

This attribute SHALL indicate a human-readable (displayable) name for the Thread network that the Node has been configured to join to. A value of null SHALL indicate that the Thread interface is not currently configured or operational.

11.14.6.4. PanId Attribute

This attribute SHALL indicate the 16-bit identifier of the Node on the Thread network. A value of null SHALL indicate that the Thread interface is not currently configured or operational.

11.14.6.5. ExtendedPanId Attribute

This attribute SHALL indicate the unique 64-bit identifier of the Node on the Thread network. A value of null SHALL indicate that the Thread interface is not currently configured or operational.

11.14.6.6. MeshLocalPrefix Attribute

This attribute SHALL indicate the mesh-local IPv6 prefix for the Thread network that the Node has been configured to join to. A value of null SHALL indicate that the Thread interface is not currently configured or operational.

11.14.6.7. OverrunCount Attribute

This attribute SHALL indicate the number of packets dropped either at ingress or egress, due to lack of buffer memory to retain all packets on the ethernet network interface. The OverrunCount attribute SHALL be reset to 0 upon a reboot of the Node.

11.14.6.8. NeighborTable Attribute

This attribute SHALL indicate the current list of Nodes that comprise the neighbor table on the Node.

11.14.6.9. RouteTable Attribute

This attribute SHALL indicate the current list of router capable Nodes for which routes have been established.

11.14.6.10. PartitionId Attribute

This attribute SHALL indicate the Thread Leader Partition Id for the Thread network to which the Node is joined. This attribute SHALL be null if not attached to a Thread network.

11.14.6.11. Weighting Attribute

This attribute SHALL indicate the Thread Leader Weight used when operating in the Leader role. This attribute SHALL be null if not attached to a Thread network.

11.14.6.12. DataVersion Attribute

This attribute SHALL indicate the full Network Data Version the Node currently uses. This attribute

SHALL be null if not attached to a Thread network.

11.14.6.13. StableDataVersion Attribute

This attribute SHALL indicate the Network Data Version for the stable subset of data the Node currently uses. This attribute SHALL be null if not attached to a Thread network.

11.14.6.14. LeaderRouterId Attribute

This attribute SHALL indicate the 8-bit LeaderRouterId the Node SHALL attempt to utilize upon becoming a router or leader on the Thread network. This attribute SHALL be null if not attached to a Thread network.

11.14.6.15. DetachedRoleCount Attribute

This attribute SHALL indicate the number of times the Node entered the OT_DEVICE_ROLE_DETACHED role as specified within the Thread specification. This value SHALL only be reset upon a Node reboot.

11.14.6.16. ChildRoleCount Attribute

This attribute SHALL indicate the number of times the Node entered the OT_DEVICE_ROLE_CHILD role as specified within the Thread specification. This value SHALL only be reset upon a Node reboot.

11.14.6.17. RouterRoleCount Attribute

This attribute SHALL indicate the number of times the Node entered the OT_DEVICE_ROLE_ROUTER role as specified within the Thread specification. This value SHALL only be reset upon a Node reboot.

11.14.6.18. LeaderRoleCount Attribute

This attribute SHALL indicate the number of times the Node entered the OT_DEVICE_ROLE_LEADER role as specified within the Thread specification. This value SHALL only be reset upon a Node reboot.

11.14.6.19. AttachAttemptCount Attribute

This attribute SHALL indicate the number of attempts that have been made to attach to a Thread network while the Node was detached from all Thread networks. This value SHALL only be reset upon a Node reboot.

11.14.6.20. PartitionIdChangeCount Attribute

This attribute SHALL indicate the number of times that the Thread network that the Node is connected to has changed its Partition ID. This value SHALL only be reset upon a Node reboot.

11.14.6.21. BetterPartitionAttachAttemptCount Attribute

This attribute SHALL indicate the number of times a Node has attempted to attach to a different Thread partition that it has determined is better than the partition it is currently attached to. This value SHALL only be reset upon a Node reboot.

11.14.6.22. ParentChangeCount Attribute

This attribute SHALL indicate the number of times a Node has changed its parent. This value SHALL only be reset upon a Node reboot.

11.14.6.23. TxTotalCount Attribute

This attribute SHALL indicate the total number of unique MAC frame transmission requests. The attribute SHALL only be incremented by 1 for each MAC transmission request regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions. This value SHALL only be reset upon a Node reboot.

11.14.6.24. TxUnicastCount Attribute

This attribute SHALL indicate the total number of unique unicast MAC frame transmission requests. The attribute SHALL only be incremented by 1 for each unicast MAC transmission request regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions. This value SHALL only be reset upon a Node reboot.

11.14.6.25. TxBroadcastCount Attribute

This attribute SHALL indicate the total number of unique broadcast MAC frame transmission requests. The attribute SHALL only be incremented by 1 for each broadcast MAC transmission request regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions. This value SHALL only be reset upon a Node reboot.

11.14.6.26. TxAckRequestedCount Attribute

This attribute SHALL indicate the total number of unique MAC frame transmission requests with requested acknowledgment. The attribute SHALL only be incremented by 1 for each MAC transmission request with requested acknowledgment regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions. This value SHALL only be reset upon a Node reboot.

11.14.6.27. TxAckedCount Attribute

This attribute SHALL indicate the total number of unique MAC frame transmission requests that were acked. The attribute SHALL only be incremented by 1 for each MAC transmission request that is acked regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions. This value SHALL only be reset upon a Node reboot.

11.14.6.28. TxNoAckRequestedCount Attribute

This attribute SHALL indicate the total number of unique MAC frame transmission requests without requested acknowledgment. The attribute SHALL only be incremented by 1 for each MAC trans-

mission request that is does not request acknowledgement regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions.

11.14.6.29. TxDataCount Attribute

This attribute SHALL indicate the total number of unique MAC Data frame transmission requests. The attribute SHALL only be incremented by 1 for each MAC Data frame transmission request regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions. This value SHALL only be reset upon a Node reboot.

11.14.6.30. TxDataPollCount Attribute

This attribute SHALL indicate the total number of unique MAC Data Poll frame transmission requests. The attribute SHALL only be incremented by 1 for each MAC Data Poll frame transmission request regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions. This value SHALL only be reset upon a Node reboot.

11.14.6.31. TxBeaconCount Attribute

This attribute SHALL indicate the total number of unique MAC Beacon frame transmission requests. The attribute SHALL only be incremented by 1 for each MAC Beacon frame transmission request regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions.

11.14.6.32. TxBeaconRequestCount Attribute

This attribute SHALL indicate the total number of unique MAC Beacon Request frame transmission requests. The attribute SHALL only be incremented by 1 for each MAC Beacon Request frame transmission request regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions. This value SHALL only be reset upon a Node reboot.

11.14.6.33. TxOtherCount Attribute

This attribute SHALL indicate the total number of unique MAC frame transmission requests that are not counted by any other attribute. The attribute SHALL only be incremented by 1 for each MAC frame transmission request regardless of the amount of CCA failures, CSMA-CA attempts, or retransmissions. This value SHALL only be reset upon a Node reboot.

11.14.6.34. TxRetryCount Attribute

This attribute SHALL indicate the total number of MAC retransmission attempts. The attribute SHALL only be incremented by 1 for each retransmission attempt that may be triggered by lack of acknowledgement, CSMA/CA failure, or other type of transmission error. This value SHALL only be reset upon a Node reboot.

11.14.6.35. TxDirectMaxRetryExpiryCount Attribute

This attribute SHALL indicate the total number of unique MAC transmission packets that meet maximal retry limit for direct packets. The attribute SHALL only be incremented by 1 for each unique MAC transmission packets that meets the maximal retry limit for direct packets. This value SHALL only be reset upon a Node reboot.

11.14.6.36. TxIndirectMaxRetryExpiryCount Attribute

This attribute SHALL indicate the total number of unique MAC transmission packets that meet maximal retry limit for indirect packets. The attribute SHALL only be incremented by 1 for each unique MAC transmission packets that meets the maximal retry limit for indirect packets. This value SHALL only be reset upon a Node reboot.

11.14.6.37. TxErrCcaCount Attribute

This attribute SHALL indicate the total number of CCA failures. The TxErrCcaCount attribute SHALL only be incremented by 1 for each instance of a CCA failure. This value SHALL only be reset upon a Node reboot.

11.14.6.38. TxErrAbortCount Attribute

This attribute SHALL indicate the total number of unique MAC transmission request failures caused by an abort error. The attribute SHALL only be incremented by 1 for each unique MAC transmission request failure caused by an abort error.

11.14.6.39. TxErrBusyChannelCount Attribute

This attribute SHALL indicate the total number of unique MAC transmission request failures caused by an error as the result of a busy channel (a CSMA/CA fail). The attribute SHALL only be incremented by 1 for each unique MAC transmission request failure caused by a busy channel such as a CSMA/CA failure.

11.14.6.40. RxTotalCount Attribute

This attribute SHALL indicate the total number of received unique MAC frames. This value SHALL only be reset upon a Node reboot.

11.14.6.41. RxUnicastCount Attribute

This attribute SHALL indicate the total number of received unique unicast MAC frames. This value SHALL only be reset upon a Node reboot.

11.14.6.42. RxBroadcastCount Attribute

This attribute SHALL indicate the total number of received unique broadcast MAC frames. This value SHALL only be reset upon a Node reboot.

11.14.6.43. RxDataCount Attribute

This attribute SHALL indicate the total number of received unique MAC Data frames. This value SHALL only be reset upon a Node reboot.

11.14.6.44. RxDataPollCount Attribute

This attribute SHALL indicate the total number of received unique MAC Data Poll frames. This value SHALL only be reset upon a Node reboot.

11.14.6.45. RxBeaconCount Attribute

This attribute SHALL indicate the total number of received unique MAC Beacon frames. This value SHALL only be reset upon a Node reboot.

11.14.6.46. RxBeaconRequestCount Attribute

This attribute SHALL indicate the total number of received unique MAC Beacon Request frames. This value SHALL only be reset upon a Node reboot.

11.14.6.47. RxOtherCount Attribute

This attribute SHALL indicate the total number of received unique MAC frame requests that are not counted by any other attribute. This value SHALL only be reset upon a Node reboot.

11.14.6.48. RxAddressFilteredCount Attribute

This attribute SHALL indicate the total number of received unique MAC frame requests that have been dropped as a result of MAC filtering. This value SHALL only be reset upon a Node reboot.

11.14.6.49. RxDestAddrFilteredCount Attribute

This attribute SHALL indicate the total number of received unique MAC frame requests that have been dropped as a result of a destination address check. This value SHALL only be reset upon a Node reboot.

11.14.6.50. RxDuplicatedCount Attribute

This attribute SHALL indicate the total number of received MAC frame requests that have been dropped as a result of being a duplicate of a previously received MAC frame request. This value SHALL only be reset upon a Node reboot.

11.14.6.51. RxErrNoFrameCount Attribute

This attribute SHALL indicate the total number of received unique MAC frame requests that have been dropped as a result of missing or malformed frame contents. This value SHALL only be reset upon a Node reboot.

11.14.6.52. RxErrUnknownNeighborCount Attribute

This attribute SHALL indicate the total number of received unique MAC frame requests that have been dropped as a result of originating from an unknown neighbor device. This value SHALL only be reset upon a Node reboot.

11.14.6.53. RxErrInvalidSrcAddrCount Attribute

This attribute SHALL indicate the total number of received unique MAC frame requests that have been dropped as a result of containing an invalid source address. This value SHALL only be reset upon a Node reboot.

11.14.6.54. RxErrSecCount Attribute

This attribute SHALL indicate the total number of received unique MAC frame requests that have been dropped as a result of an error with the security of the received frame. This value SHALL only be reset upon a Node reboot.

11.14.6.55. RxErrFcsCount Attribute

This attribute SHALL indicate the total number of received unique MAC frame requests that have been dropped as a result of an error with the FCS of the received frame. This value SHALL only be reset upon a Node reboot.

11.14.6.56. RxErrOtherCount Attribute

This attribute SHALL indicate the total number of received unique MAC frame requests that have been dropped as a result of an error that is not counted by any other attribute. This value SHALL only be reset upon a Node reboot.

11.14.6.57. ActiveTimestamp Attribute

This attribute SHALL be null when there is no dataset configured.

11.14.6.58. PendingTimestamp Attribute

This attribute SHALL be null when there is no dataset configured.

11.14.6.59. Delay Attribute

This attribute SHALL be null when there is no dataset configured.

11.14.6.60. SecurityPolicy Attribute

This attribute SHALL indicate the current security policies for the Thread partition to which a Node is connected. This attribute SHALL be null when there is no dataset configured.

11.14.6.61. ChannelPage0Mask Attribute

This attribute SHALL indicate the channels within channel page 0, in the 2.4GHz ISM band. The channels are represented in most significant bit order, with bit value 1 meaning selected, bit value 0 meaning unselected. For example, the most significant bit of the left-most byte indicates channel 0. If channel 0 and channel 10 are selected, the mask would be: 80 20 00 00. This attribute SHALL be null when there is no dataset configured.

11.14.6.62. OperationalDatasetComponents Attribute

This attribute SHALL indicate a collection of flags to indicate the presence of various operationally acquired values.

11.14.6.63. ActiveNetworkFaults Attribute

This attribute SHALL indicate the set of faults currently detected by the Node. When the Node

detects a fault has been raised, the appropriate [NetworkFaultEnum](#) value SHALL be added to this list. This list SHALL NOT contain more than one instance of a specific [NetworkFaultEnum](#) value. When the Node detects that all conditions contributing to a fault has been cleared, the corresponding [NetworkFaultEnum](#) value SHALL be removed from this list. An empty list SHALL indicate there are currently no active faults. The order of this list SHOULD have no significance. Clients interested in monitoring changes in active faults MAY subscribe to this attribute, or they MAY subscribe to [NetworkFaultChange](#)

11.14.6.64. ExtAddress Attribute

This attribute SHALL indicate the IEEE 802.15.4 extended address for the Node. A value of null SHALL indicate that the extended address is not yet known. The uint64 value is composed by taking the 8 octets of the extended address EUI-64 and treating them as a big-endian integer. For example, octet string (in hexadecimal, from first octet to last) **00112233AABBCCDD** would lead to a value of 0x00112233AABBCCDD.

11.14.6.65. Rloc16 Attribute

This attribute SHALL indicate the RLOC16 of the Node. A value of null SHALL indicate that the Thread interface is not currently configured or operational. The uint16 value is composed by taking the two RLOC16 and treating the octet string as if it was encoding a big-endian integer. For example, octet string (in hexadecimal, from first octet to last) **44AA** would lead to a value of 0x44AA.

11.14.7. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	ResetCounts	client ⇒ server	Y	M	ERRCNT

11.14.7.1. ResetCounts Command

This command is used to reset the count attributes.

Reception of this command SHALL reset the following attributes to 0:

- [OverrunCount](#)

Upon completion, this command SHALL send a status code of SUCCESS back to the initiator.

11.14.8. Events

ID	Name	Priority	Access	Conformance
0x00	ConnectionStatus	INFO	V	O
0x01	NetworkFault-Change	INFO	V	O

11.14.8.1. NetworkFaultChange Event

The NetworkFaultChange Event SHALL indicate a change in the set of network faults currently

detected by the Node.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Current	list[Network-FaultEnum]	max 4			M
1	Previous	list[Network-FaultEnum]	max 4			M

Current Field

This field SHALL represent the set of faults currently detected, as per [Section 11.14.5.1, “Network-FaultEnum Type”](#).

Previous Field

This field SHALL represent the set of faults detected prior to this change event, as per [Section 11.14.5.1, “NetworkFaultEnum Type”](#).

11.14.8.2. ConnectionStatus Event

The ConnectionStatus Event SHALL indicate that a Node’s connection status to a Thread network has changed.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Connection-Status	Connection-StatusEnum				M

11.15. Wi-Fi Network Diagnostics Cluster

The Wi-Fi Network Diagnostics Cluster provides a means to acquire standardized diagnostics metrics that MAY be used by a Node to assist a user or Administrator in diagnosing potential problems. The Wi-Fi Network Diagnostics Cluster attempts to centralize all metrics that are relevant to a potential Wi-Fi radio running on a Node.

11.15.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.15.2. Classification

Hierarchy	Role	Scope	PICS Code	Quality
Base	Utility	Node	DGWIFI	K

11.15.3. Cluster ID

ID	Name
0x0036	Wi-Fi Network Diagnostics

11.15.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Summary
0	PKTCNT	PacketCounts	Node makes available the counts for the number of received and transmitted packets on the Wi-Fi interface.
1	ERRCNT	ErrorCounts	Node makes available the counts for the number of errors that have occurred during the reception and transmission of packets on the Wi-Fi interface.

11.15.5. Data Types

11.15.5.1. SecurityTypeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unspecified	Indicate the usage of an unspecified Wi-Fi security type	M
1	None	Indicate the usage of no Wi-Fi security	M
2	WEP	Indicate the usage of WEP Wi-Fi security	M
3	WPA	Indicate the usage of WPA Wi-Fi security	M
4	WPA2	Indicate the usage of WPA2 Wi-Fi security	M

Value	Name	Summary	Conformance
5	WPA3	Indicate the usage of WPA3 Wi-Fi security	M

11.15.5.2. WiFiVersionEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	a	Indicate the network interface is currently using IEEE 802.11a against the wireless access point.	M
1	b	Indicate the network interface is currently using IEEE 802.11b against the wireless access point.	M
2	g	Indicate the network interface is currently using IEEE 802.11g against the wireless access point.	M
3	n	Indicate the network interface is currently using IEEE 802.11n against the wireless access point.	M
4	ac	Indicate the network interface is currently using IEEE 802.11ac against the wireless access point.	M
5	ax	Indicate the network interface is currently using IEEE 802.11ax against the wireless access point.	M
6	ah	Indicate the network interface is currently using IEEE 802.11ah against the wireless access point.	M

11.15.5.3. AssociationFailureCauseEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unknown	The reason for the failure is unknown.	M
1	AssociationFailed	An error occurred during association.	M
2	AuthenticationFailed	An error occurred during authentication.	M
3	SsidNotFound	The specified SSID could not be found.	M

11.15.5.4. ConnectionStatusEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Connected	Indicate the node is connected	M
1	NotConnected	Indicate the node is not connected	M

11.15.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	BSSID	octstr	6	X	null	R V	M
0x0001	Security-Type	Security-TypeEnum	all	X	null	R V	M
0x0002	WiFiVersion	WiFiVersionEnum	all	X	null	R V	M
0x0003	Channel-Number	uint16	all	X	null	R V	M
0x0004	RSSI	int8	-120 to 0	C X	null	R V	M
0x0005	Beacon-LostCount	uint32	all	C X	0	R V	ERRCNT
0x0006	BeaconRx-Count	uint32	all	C X	0	R V	PKTCNT

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0007	Packet-MulticastRxCount	uint32	all	C X	0	R V	PKTCNT
0x0008	Packet-MulticastTxCount	uint32	all	C X	0	R V	PKTCNT
0x0009	PacketUnicastRxCount	uint32	all	C X	0	R V	PKTCNT
0x000A	PacketUnicastTxCount	uint32	all	C X	0	R V	PKTCNT
0x000B	Current-MaxRate	uint64	all	C X	0	R V	O
0x000C	Overrun-Count	uint64	all	C X	0	R V	ERRCNT

For all attributes listed above, a null value SHALL be returned if the interface is not currently configured or operational.

11.15.6.1. BSSID Attribute

This attribute SHALL indicate the BSSID for which the Wi-Fi network the Node is currently connected.

11.15.6.2. SecurityType Attribute

This attribute SHALL indicate the current type of Wi-Fi security used.

11.15.6.3. WiFiVersion Attribute

This attribute SHALL indicate the current [IEEE 802.11](#) standard version in use by the Node, per the table below.

11.15.6.4. ChannelNumber Attribute

This attribute SHALL indicate the channel that Wi-Fi communication is currently operating on.

11.15.6.5. RSSI Attribute

This attribute SHALL indicate the current RSSI of the Node's Wi-Fi radio in dBm.

11.15.6.6. BeaconLostCount Attribute

This attribute SHALL indicate the count of the number of missed beacons the Node has detected. If

the Node does not have an ability to count beacons expected and not received, this value MAY remain set to zero.

11.15.6.7. BeaconRxCount Attribute

This attribute SHALL indicate the count of the number of received beacons. The total number of expected beacons that could have been received during the interval since association SHOULD match the sum of BeaconRxCount and BeaconLostCount. If the Node does not have an ability to report count of beacons received, this value MAY remain set to zero.

11.15.6.8. PacketMulticastRxCount Attribute

This attribute SHALL indicate the number of multicast packets received by the Node.

11.15.6.9. PacketMulticastTxCount Attribute

This attribute SHALL indicate the number of multicast packets transmitted by the Node.

11.15.6.10. PacketUnicastRxCount Attribute

This attribute SHALL indicate the number of unicast packets received by the Node.

11.15.6.11. PacketUnicastTxCount Attribute

This attribute SHALL indicate the number of unicast packets transmitted by the Node.

11.15.6.12. CurrentMaxRate Attribute

This attribute SHALL indicate the current maximum PHY rate of transfer of data in bits-per-second.

11.15.6.13. OverrunCount Attribute

This attribute SHALL indicate the number of packets dropped either at ingress or egress, due to lack of buffer memory to retain all packets on the network interface. The attribute SHALL be reset to 0 upon a reboot of the Node.

11.15.7. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	ResetCounts	client ⇒ server	Y	O	ERRCNT

11.15.7.1. ResetCounts Command

This command is used to reset the count attributes.

Reception of this command SHALL reset the following attributes to 0:

- [BeaconLostCount](#)
- [BeaconRxCount](#)

- [PacketMulticastRxCount](#)
- [PacketMulticastTxCount](#)
- [PacketUnicastRxCount](#)
- [PacketUnicastTxCount](#)

11.15.8. Events

ID	Name	Priority	Quality	Access	Conformance
0x00	Disconnection	info		V	O
0x01	Association-Failure	info		V	O
0x02	Connection-Status	info		V	O

11.15.8.1. Disconnection Event

The Disconnection Event SHALL indicate that a Node's Wi-Fi connection has been disconnected as a result of de-authenticated or dis-association and indicates the reason.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ReasonCode	uint16				M

ReasonCode Field

This field SHALL contain the Reason Code field value for the Disassociation or Deauthentication event that caused the disconnection and the value SHALL align with Table 9-49 "Reason codes" of [IEEE 802.11-2020](#).

11.15.8.2. AssociationFailure Event

The AssociationFailure event SHALL indicate that a Node has attempted to connect, or reconnect, to a Wi-Fi access point, but is unable to successfully associate or authenticate, after exhausting all internal retries of its supplicant.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Association-Failure-Cause	Association-Failure-CauseEnum	all			M
1	Status	uint16	all			M

AssociationFailureCause Field

The Status field SHALL be set to a value from the [AssociationFailureCauseEnum](#).

Status Field

The Status field SHALL be set to the Status Code value that was present in the last frame related to association where Status Code was not equal to zero and which caused the failure of a last trial attempt, if this last failure was due to one of the following Management frames:

- Association Response (Type 0, Subtype 1)
- Reassociation Response (Type 0, Subtype 3)
- Authentication (Type 0, Subtype 11)

Table 9-50 "Status codes" of [IEEE 802.11-2020](#) contains a description of all values possible.

11.15.8.3. ConnectionStatus Event

The ConnectionStatus Event SHALL indicate that a Node's connection status to a Wi-Fi network has changed. Connected, in this context, SHALL mean that a Node acting as a Wi-Fi station is successfully associated to a Wi-Fi Access Point.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Connection-Status	Connection-StatusEnum				M

11.16. Ethernet Network Diagnostics Cluster

The Ethernet Network Diagnostics Cluster provides a means to acquire standardized diagnostics metrics that MAY be used by a Node to assist a user or Administrator in diagnosing potential problems. The Ethernet Network Diagnostics Cluster attempts to centralize all metrics that are relevant to a potential Ethernet connection to a Node.

11.16.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.16.2. Classification

Hierarchy	Role	Scope	PICS Code	Quality
Base	Utility	Node	DGETH	K

11.16.3. Cluster ID

ID	Name
0x0037	Ethernet Network Diagnostics

11.16.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Summary
0	PKTCNT	PacketCounts	Node makes available the counts for the number of received and transmitted packets on the ethernet interface.
1	ERRCNT	ErrorCounts	Node makes available the counts for the number of errors that have occurred during the reception and transmission of packets on the ethernet interface.

11.16.5. Data Types

11.16.5.1. PHYRateEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Rate10M	PHY rate is 10Mbps	M
1	Rate100M	PHY rate is 100Mbps	M
2	Rate1G	PHY rate is 1Gbps	M
3	Rate2_5G	PHY rate is 2.5Gbps	M
4	Rate5G	PHY rate is 5Gbps	M
5	Rate10G	PHY rate is 10Gbps	M
6	Rate40G	PHY rate is 40Gbps	M
7	Rate100G	PHY rate is 100Gbps	M
8	Rate200G	PHY rate is 200Gbps	M
9	Rate400G	PHY rate is 400Gbps	M

11.16.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	PHYRate	PHYRateEnum	all	X	null	R V	O
0x0001	FullDuplex	bool	all	X	null	R V	O
0x0002	PacketRxCount	uint64	all	C	0	R V	PKTCNT
0x0003	PacketTxCount	uint64	all	C	0	R V	PKTCNT
0x0004	TxErrCount	uint64	all	C	0	R V	ERRCNT
0x0005	CollisionCount	uint64	all	C	0	R V	ERRCNT
0x0006	OverrunCount	uint64	all	C	0	R V	ERRCNT
0x0007	CarrierDetect	bool	all	C X	null	R V	O
0x0008	TimeSyncReset	uint64	all	C	0	R V	O

11.16.6.1. PHYRate Attribute

This attribute SHALL indicate the current nominal, usable speed at the top of the physical layer of the Node. A value of null SHALL indicate that the interface is not currently configured or operational.

11.16.6.2. FullDuplex Attribute

This attribute SHALL indicate if the Node is currently utilizing the full-duplex operating mode. A value of null SHALL indicate that the interface is not currently configured or operational.

11.16.6.3. PacketRxCount Attribute

This attribute SHALL indicate the number of packets that have been received on the ethernet network interface. The attribute SHALL be reset to 0 upon a reboot of the Node.

11.16.6.4. PacketTxCount Attribute

This attribute SHALL indicate the number of packets that have been successfully transferred on the ethernet network interface. The attribute SHALL be reset to 0 upon a reboot of the Node.

11.16.6.5. TxErrCount Attribute

This attribute SHALL indicate the number of failed packet transmissions that have occurred on the

ethernet network interface. The attribute SHALL be reset to 0 upon a reboot of the Node.

11.16.6.6. CollisionCount Attribute

This attribute SHALL indicate the number of collisions that have occurred while attempting to transmit a packet on the ethernet network interface. The attribute SHALL be reset to 0 upon a reboot of the Node.

11.16.6.7. OverrunCount Attribute

This attribute SHALL indicate the number of packets dropped either at ingress or egress, due to lack of buffer memory to retain all packets on the ethernet network interface. The attribute SHALL be reset to 0 upon a reboot of the Node.

11.16.6.8. CarrierDetect Attribute

This attribute SHALL indicate the value of the Carrier Detect control signal present on the ethernet network interface. A value of null SHALL indicate that the interface is not currently configured or operational.

11.16.6.9. TimeSinceReset Attribute

This attribute SHALL indicate the duration of time, in minutes, that it has been since the ethernet network interface has reset for any reason.

11.16.7. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	ResetCounts	client ⇒ server	Y	M	PKTCNT ERRCNT

11.16.7.1. ResetCounts Command

This command is used to reset the count attributes.

Reception of this command SHALL reset the following attributes to 0:

- [PacketRxCount](#)
- [PacketTxCount](#)
- [TxErrCount](#)
- [CollisionCount](#)
- [OverrunCount](#)

11.17. Time Synchronization Cluster

Accurate time is required for a number of reasons, including scheduling, display and validating security materials.

This section describes a mechanism for Nodes to achieve and maintain time synchronization. The Time Synchronization cluster provides attributes for reading a Node's current time. It also allows Administrators to set current time, time zone and daylight savings time (DST) settings.

The Time Synchronization cluster MAY be present on the [root node endpoint](#), and SHALL NOT be present on any other Endpoint of any Node.

11.17.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Make TrustedTimeSource fabric-aware, add TSC feature, define list max sizes, change writable attributes to commands, remote port, add attribute for DNS support

11.17.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	TIMESYNC

11.17.3. Cluster ID

ID	Name
0x0038	Time Synchronization

11.17.4. Terminology

Term	Definition
DNS-SD	Domain name service - service discovery RFC 6763
DST	Daylight Savings Time
GNSS	Global Navigation Satellite System. This is a satellite system that provides global geo-spatial positioning. GNSS systems include the NAVSTAR global positioning system (GPS), Galileo, GLONASS and BeiDou
NTP	Network Time Protocol RFC 5905
NTS	Network Time Security RFC 8915
PTP	Precision Time Protocol IEEE 1588-2008

Term	Definition
SNTP	Simple Network Time Protocol. This is a simplified version of the Network Time Protocol. It is also covered by RFC 5905

11.17.5. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Summary
0	TZ	TimeZone	Server supports time zone.
1	NTPC	NTPClient	Server supports an NTP or SNTP client.
2	NTPS	NTPServer	Server supports an NTP server role.
3	TSC	TimeSyncClient	Time synchronization client cluster is present.

11.17.5.1. TimeZone Feature

Allows a server to translate a UTC time to a local time using the time zone and daylight savings time (DST) offsets. If a server supports the TimeZone feature, it SHALL support the [SetTimeZone](#) and [SetDSTOffset](#) commands, and [TimeZone](#) and [DSTOffset](#) attributes, and SHALL expose the local time through the [LocalTime](#) attribute.

11.17.5.2. NTPClient Feature

Allows a node to use NTP/SNTP for time synchronization.

11.17.5.3. NTPServer Feature

Allows a Node to host an NTP server for the network so that other Nodes can achieve a high accuracy time synchronization within the network. See [Section 11.17.15, “Acting as an NTP Server”](#).

11.17.5.4. TimeSyncClient Feature

This node also supports a time synchronization client and can connect to and read time from other nodes.

11.17.6. Data Types

11.17.6.1. GranularityEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	NoTimeGranularity	This indicates that the node is not currently synchronized with a UTC Time source and its clock is based on the Last Known Good UTC Time only.	M
1	MinutesGranularity	This indicates the node was synchronized to an upstream source in the past, but sufficient clock drift has occurred such that the clock error is now > 5 seconds.	M
2	SecondsGranularity	This indicates the node is synchronized to an upstream source using a low resolution protocol. UTC Time is accurate to ± 5 seconds.	M
3	MillisecondsGranularity	This indicates the node is synchronized to an upstream source using high resolution time-synchronization protocol such as NTP, or has built-in GNSS with some amount of jitter applying its GNSS timestamp. UTC Time is accurate to ± 50 ms.	M
4	MicrosecondsGranularity	This indicates the node is synchronized to an upstream source using a highly precise time-synchronization protocol such as PTP, or has built-in GNSS. UTC time is accurate to ± 10 μ s.	M

11.17.6.2. TimeSourceEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	None	Node is not currently synchronized with a UTC Time source.	M
1	Unknown	Node uses an unlisted time source.	M
2	Admin	Node received time from a client using the SetUTCTime Command.	M
3	NodeTimeCluster	Synchronized time by querying the Time Synchronization cluster of another Node.	M
4	NonMatterSNTP	SNTP from a server not in the Matter network. NTS is not used.	M
5	NonMatterNTP	NTP from servers not in the Matter network. None of the servers used NTS.	M
6	MatterSNTP	SNTP from a server within the Matter network. NTS is not used.	M
7	MatterNTP	NTP from servers within the Matter network. None of the servers used NTS.	M
8	MixedNTP	NTP from multiple servers in the Matter network and external. None of the servers used NTS.	M
9	NonMatterSNTPNTS	SNTP from a server not in the Matter network. NTS is used.	M
10	NonMatterNTPNTS	NTP from servers not in the Matter network. NTS is used on at least one server.	M
11	MatterSNTPNTS	SNTP from a server within the Matter network. NTS is used.	M

Value	Name	Summary	Conformance
12	MatterNTPNTS	NTP from a server within the Matter network. NTS is used on at least one server.	M
13	MixedNTPNTS	NTP from multiple servers in the Matter network and external. NTS is used on at least one server.	M
14	CloudSource	Time synchronization comes from a vendor cloud-based source (e.g. "Date" header in authenticated HTTPS connection).	M
15	PTP	Time synchronization comes from PTP.	M
16	GNSS	Time synchronization comes from a GNSS source.	M

11.17.6.3. TimeZoneDatabaseEnum Type

This data type is derived from [enum8](#). It indicates what the device knows about the contents of the [IANA Time Zone Database](#). Partial support on a device MAY be used to omit historical data, less commonly used time zones, and/or time zones not related to the region a product is sold in.

Value	Name	Summary	Conformance
0	Full	Node has a full list of the available time zones	M
1	Partial	Node has a partial list of the available time zones	M
2	None	Node does not have a time zone database	M

11.17.6.4. TrustedTimeSourceStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	FabricIndex	fabric-idx	all				M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	NodeID	node-id	all				M
2	Endpoint	endpoint-no	all				M

FabricIndex Field

The [Fabric Index](#) associated with the Fabric of the client which last set the value of the trusted time source node.

NodeID Field

Node ID of the trusted time source node on the Fabric associated with the entry.

Endpoint Field

Endpoint on the trusted time source node that contains the Time Synchronization cluster server.

11.17.6.5. FabricScopedTrustedTimeSourceStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	NodeID	node-id	all				M
1	Endpoint	endpoint-no	all				M

NodeID Field

Node ID of the trusted time source node on the Fabric of the issuer.

Endpoint Field

Endpoint on the trusted time source node that contains the Time Synchronization cluster server. This is provided to avoid having to do discovery of the location of that endpoint by walking over all endpoints and checking their [Descriptor Cluster](#).

11.17.6.6. TimeZoneStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Offset	int32	-43200 to 50400				M
1	ValidAt	epoch-us	all				M
2	Name	string	0 to 64				O

Offset Field

The time zone offset from UTC in seconds.

ValidAt Field

The UTC time when the offset SHALL be applied.

Name Field

The time zone name SHOULD provide a human-readable time zone name and it SHOULD use the country/city format specified by the [IANA Time Zone Database](#). The Name field MAY be used for display. If the node supports a TimeZoneDatabase it MAY use the Name field to set its own DST offsets if it has database information for the supplied time zone Name and the given Offset matches.

11.17.6.7. DSTOffsetStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Offset	int32	desc				M
1	ValidStarting	epoch-us	all				M
2	ValidUntil	epoch-us	all	X			M

Offset Field

The DST offset in seconds. Normally this is in the range of 0 to 3600 seconds (1 hour), but this field will accept any values in the int32 range to accommodate potential future legislation that does not fit with these assumptions.

ValidStarting Field

The UTC time when the offset SHALL be applied.

ValidUntil Field

The UTC time when the offset SHALL stop being applied. Providing a null value here indicates a permanent DST change. If this value is non-null the value SHALL be larger than the ValidStarting time.

11.17.7. Status Codes

11.17.7.1. StatusCodeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0x02	TimeNotAccepted	Node rejected the attempt to set the UTC time	M

11.17.8. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	UTCTime	epoch-us	all	C X	null	R V	M
0x0001	Granularity	GranularityEnum	desc		NoTime-Granularity	R V	M
0x0002	Time-Source	Time-SourceEnum	desc		None	R V	O
0x0003	Trusted-Time-Source	Trusted-Time-SourceStruct	all	N X	null	R V	TSC
0x0004	Default-NTP	string	max 128	N X	null	R V	NTPC
0x0005	TimeZone	list[TimeZoneStruct]	1 to 2	N		R V	TZ
0x0006	DSTOffset	list[DSTOffsetStruct]		N	[]	R V	TZ
0x0007	LocalTime	epoch-us	all	C X	null	R V	TZ
0x0008	TimeZone-Database	TimeZone-DatabaseEnum	all	F	None	R V	TZ
0x0009	NTPServer Available	bool	all		False	R V	NTPS
0x000A	Time-ZoneList-MaxSize	uint8	1 to 2	F		R V	TZ
0x000B	DSTOffsetList-MaxSize	uint8	min 1	F		R V	TZ
0x000C	Supports-DNSResolve	bool	all	F	False	R V	NTPC

11.17.8.1. UTCTime Attribute

If the node has achieved time synchronization, this attribute SHALL indicate the current time as a UTC [epoch-us](#) (Epoch Time in Microseconds).

If the node has not achieved time synchronization, this attribute SHALL be null. This attribute MAY be set when a [SetUTCTime](#) is received.

11.17.8.2. Granularity Attribute

This attribute SHALL indicate granularity of the error that the node is willing to guarantee on the time synchronization. It is of type [GranularityEnum](#).

This value SHALL be set to NoTimeGranularity if UTCTime is null and SHALL NOT be set to NoTimeGranularity if UTCTime is non-null.

11.17.8.3. TimeSource Attribute

This attribute SHALL indicate the node's time source. This attribute indicates what method the node is using to sync, whether the source uses NTS or not and whether the source is internal or external to the Matter network. This attribute MAY be used by a client to determine its level of trust in the UTCTime. It is of type [TimeSourceEnum](#).

If a node is unsure if the selected NTP server is within the Matter network, it SHOULD select one of the NonMatter* values.

This value SHALL be set to None if UTCTime is null and SHALL NOT be set to None if UTCTime is non-null.

11.17.8.4. TrustedTimeSource Attribute

This attribute SHALL indicate the Node ID, endpoint, and associated [fabric index](#) of a Node that MAY be used as trusted time source. See [Section 11.17.13, "Time source prioritization"](#). This attribute reflects the last value set by an administrator using the [SetTrustedTimeSource](#) command. If the value is null, no trusted time source has yet been set.

11.17.8.5. DefaultNTP Attribute

This attribute SHALL indicate the default NTP server that this Node MAY use if other time sources are unavailable. This attribute is settable by an Administrator using the [SetDefaultNTP](#) command. It SHOULD be set by the Commissioner during commissioning. If no default NTP server is available, the Commissioner MAY set this value to null. The default IANA assigned NTP port of 123 SHALL be used to access the NTP server.

If set, the format of this attribute SHALL be a domain name or a static IPv6 address with no port, in text format, as specified in [RFC 5952](#). The address format SHALL follow the recommendations in Section 4 and SHALL NOT contain a port number.

11.17.8.6. TimeZone Attribute

This attribute SHALL contain a list of time zone offsets from UTC and when they SHALL take effect.

This attribute uses a list of time offset configurations to allow Nodes to handle scheduled regulatory time zone changes. This attribute SHALL NOT be used to indicate daylight savings time changes (see [Section 11.17.8.7, “DSTOffset Attribute”](#) for daylight savings time).

The first entry SHALL have a ValidAt entry of 0. If there is a second entry, it SHALL have a non-zero ValidAt time.

If a node supports a [TimeZoneDatabase](#), and it has data for the given time zone Name and the given Offset matches, the node MAY update its own [DSTOffset](#) attribute to add new DST change times as required, based on the Name fields of the TimeZoneStruct. Administrators MAY add additional entries to the DSTOffset of other Nodes with the same time zone, if required.

If a node does not support a [TimeZoneDatabase](#), the Name field of the [TimeZoneStruct](#) is only applicable for client-side localization. In particular:

- If the node does not support a [TimeZoneDatabase](#), the Name field SHALL NOT be used to calculate the local time.
- If the node does not support a [TimeZoneDatabase](#), the Name field SHALL NOT be used to calculate DST start or end dates.

When time passes, the node SHOULD remove any entries which are no longer active and change the ValidAt time for the currently used TimeZoneStruct list item to zero.

This attribute SHALL have at least one entry. If the node does not have a default time zone and no time zone has been set, it MAY set this value to a list containing a single [TimeZoneStruct](#) with an offset of 0 (UTC) and a ValidAt time of 0.

11.17.8.7. DSTOffset Attribute

This attribute SHALL contain a list of offsets to apply for daylight savings time, and their validity period.

List entries SHALL be sorted by ValidStarting time.

A list entry SHALL NOT have a ValidStarting time that is smaller than the ValidUntil time of the previous entry. There SHALL be at most one list entry with a null ValidUntil time and, if such an entry is present, it SHALL appear last in the list.

Over time, the node SHOULD remove any entries which are no longer active from the list.

Over time, if the node supports a [TimeZoneDatabase](#) and it has information available for the given time zone name, it MAY update its own list to add additional entries.

If a time zone does not use DST, this SHALL be indicated by a single entry with a 0 offset and a null ValidUntil field.

11.17.8.8. LocalTime Attribute

This attribute SHALL indicate the computed current local time of the node as a [epoch-us](#) (Epoch Time in Microseconds). The value of LocalTime SHALL be the sum of the [UTCTime](#), the offset of the currently valid [TimeZoneStruct](#) from the [TimeZone](#) attribute (converted to microseconds), and the

offset of the currently valid [DSTOffsetStruct](#) from the [DSTOffset](#) attribute (converted to microseconds), if such an entry exists.

If the node has not achieved time synchronization, this SHALL be null. If the node has an empty [DSTOffset](#), this SHALL be null.

11.17.8.9. TimeZoneDatabase Attribute

This attribute SHALL indicate whether the node has access to a time zone database. Nodes with a time zone database MAY update their own [DSTOffset](#) attribute to add new entries and MAY push [DSTOffset](#) updates to other Nodes in the same time zone as required.

11.17.8.10. NTPServerAvailable Attribute

This attribute SHALL indicate if the node is running an [RFC 5905](#) NTPv4 compliant server on port 123, this value SHALL be True.

If the node is not currently running an NTP server, this value SHALL be False.

11.17.8.11. TimeZoneListMaxSize Attribute

This attribute SHALL indicate the number of supported list entries in the [TimeZone](#) attribute. This attribute may take the value of 1 or 2, where the optional second list entry MAY be used to handle scheduled regulatory time zone changes.

11.17.8.12. DSTOffsetListMaxSize Attribute

This attribute SHALL indicate the number of supported list entries in [DSTOffset](#) attribute. This value must be at least 1.

11.17.8.13. SupportsDNSResolve Attribute

This attribute SHALL indicate if the node supports resolving a domain name. DefaultNTP Address values for these nodes MAY include domain names. If this is False, the Address for a DefaultNTP SHALL be an IPv6 address.

11.17.9. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	SetUTCTime	client ⇒ server	Y	A	M
0x01	SetTrusted-TimeSource	client ⇒ server	Y	A F	TSC
0x02	SetTimeZone	client ⇒ server	SetTime-ZoneResponse	M	TZ
0x03	SetTime-ZoneResponse	client ⇐ server	N		TZ
0x04	SetDSTOffset	client ⇒ server	Y	M	TZ

ID	Name	Direction	Response	Access	Conformance
0x05	SetDefaultNTP	client ⇒ server	Y	A	NTPC

11.17.9.1. SetUTCTime Command

This command is used to set the UTC time of the node.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	UTCTime	epoch-us	all		0	M
1	Granularity	GranularityEnum	all		NoTime-Granularity	M
2	TimeSource	TimeSourceEnum	all		None	O

This command MAY be issued by Administrator to set the time. If the Commissioner does not have a valid time source, it MAY send a Granularity of NoTimeGranularity.

Upon receipt of this command, the node MAY update its UTCTime attribute to match the time specified in the command, if the stated Granularity and TimeSource are acceptable. The node SHALL update its UTCTime attribute if its current Granularity is NoTimeGranularity.

If the time is updated, the node SHALL also update its Granularity attribute based on the granularity specified in the command and the expected clock drift of the node. This SHOULD normally be one level lower than the stated command Granularity. It SHALL also update its TimeSource attribute to Admin. It SHALL also update its Last Known Good UTC Time as defined in [Section 3.5.6.1, “Last Known Good UTC Time”](#).

If the node updates its UTCTime attribute, it SHALL accept the command with a status code of SUCCESS. If it opts to not update its time, it SHALL fail the command with a cluster specific Status Code of TimeNotAccepted.

UTCTime Field

This field SHALL give the Client’s UTC Time.

Granularity Field

This field SHALL give the Client’s Granularity, as described in [Granularity](#).

TimeSource Field

This field SHALL give the Client’s TimeSource, as described in [TimeSource](#).

11.17.9.2. SetTrustedTimeSource Command

This command is used to set the TrustedTimeSource attribute.

Upon receipt of this command:

- If the TrustedTimeSource field in the command is null, the node SHALL set the TrustedTimeSource attribute to null and SHALL generate a [MissingTrustedTimeSource](#) event.
- Otherwise, the node SHALL set the TrustedTimeSource attribute to a struct which has NodeID and Endpoint fields matching those in the TrustedTimeSource field and has its FabricIndex field set to the command's [accessing fabric index](#).

Access Modifier: Fabric Scoped						
ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Trusted-TimeSource	FabricScope-dTrusted-Time-SourceStruct		X		M

TrustedTimeSource Field

This field contains the Node ID and endpoint of a trusted time source on the accessing fabric.

11.17.9.3. SetTimeZone Command

This command is used to set the time zone of the node. The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	TimeZone	list[TimeZoneStruct]	1 to 2			M

If the given list is larger than the TimeZoneListMaxSize, the node SHALL respond with RESOURCE_EXHAUSTED and the TimeZone attribute SHALL NOT be updated.

If the given list does not conform to the list requirements in [TimeZone](#) attribute the node SHALL respond with a CONSTRAINT_ERROR and the TimeZone attribute SHALL NOT be updated.

If there are no errors in the list, the TimeZone field SHALL be copied to the TimeZone attribute. A [TimeZoneStatus](#) event SHALL be generated with the new time zone information.

If the node supports a time zone database and it has information available for the time zone that will be applied, it MAY set its [DSTOffset](#) attribute, otherwise the [DSTOffset](#) attribute SHALL be set to an empty list. A [DSTTableEmpty](#) event SHALL be generated if the [DSTOffset](#) attribute is empty. A [DSTStatus](#) event SHALL be generated if the node was previously applying a DST offset.

11.17.9.4. SetTimeZoneResponse Command

This command is used to report the result of a SetTimeZone command. This command SHALL be generated in response to a SetTimeZone command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	DSTOffsetRequired	bool	all		true	M

DSTOffsetRequired Field

If the node supports a time zone database with information for the time zone that will be applied, it MAY use this information to set the [DSTOffset](#) attribute. If the node is setting its own DSTOffset attribute, the DSTOffsetRequired field SHALL be set to false, otherwise it SHALL be set to true.

11.17.9.5. SetDSTOffset Command

This command is used to set the DST offsets for a node.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	DSTOffset	list[DSTOffsetStruct]				M

- If the length of DSTOffset is larger than DSTOffsetListMaxSize, the node SHALL respond with RESOURCE_EXHAUSTED.
- Else if the list entries do not conform to the list requirements for [DSTOffset](#) attribute, the node SHALL respond with CONSTRAINT_ERROR.

If there are no errors in the list, the DSTOffset field SHALL be copied to the DSTOffset attribute.

If the DSTOffset attribute change causes a corresponding change to the DST state, a [DSTStatus](#) event SHALL be generated. If the list is empty, the node SHALL generate a [DSTTableEmpty](#) event.

11.17.9.6. SetDefaultNTP Command

This command is used to set the DefaultNTP attribute.

If the DefaultNTP Address field does not conform to the requirements in the [DefaultNTP](#) attribute description, the command SHALL fail with a status code of INVALID_COMMAND. If the node does not support DNS resolution (as specified in [SupportsDNSResolve](#)) and the provided Address is a domain name, the command SHALL fail with a status code of INVALID_COMMAND. Otherwise, the node SHALL set the DefaultNTP attribute to match the DefaultNTP provided in this command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	DefaultNTP	string	max 128	X		M

DefaultNTP Field

This field contains the address of an NTP server than can be used as a fallback for time synchronization. The format of this field SHALL follow the requirements in the [DefaultNTP](#) attribute

description.

11.17.10. Events

ID	Name	Priority	Quality	Access	Conformance
0x00	DST-TableEmpty	INFO		V	TZ
0x01	DSTStatus	INFO		V	TZ
0x02	TimeZoneStatus	INFO		V	TZ
0x03	TimeFailure	INFO		V	M
0x04	MissingTrustedTimeSource	INFO		V	TSC

11.17.10.1. DSTTableEmpty Event

This event SHALL be generated when the node stops applying the current DSTOffset and there are no entries in the list with a larger ValidStarting time, indicating the need to possibly get new DST data. This event SHALL also be generated if the DSTOffset list is cleared either by a [SetTimeZone](#) command, or by a [SetDSTOffset](#) command with an empty list.

The node SHALL generate this event if the node has not generated a DSTTableEmpty event in the last hour, and the DSTOffset list is empty when the node attempts to update its time. DST-TableEmpty events corresponding to a time update SHOULD NOT be generated more often than once per hour.

There is no data for this event.

11.17.10.2. DSTStatus Event

This event SHALL be generated when the node starts or stops applying a DST offset.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	DSTOffsetActive	bool				M

DSTOffsetActive Field

Indicates whether the current DST offset is being applied (i.e, daylight savings time is applied, as opposed to standard time).

11.17.10.3. TimeZoneStatus Event

This event SHALL be generated when the node changes its time zone offset or name. It SHALL NOT be sent for DST changes that are not accompanied by a time zone change.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Offset	int32	-43200 to 50400			M
1	Name	string	0 to 64			O

Offset Field

Current time zone offset from UTC in seconds.

Name Field

Current time zone name. This name SHOULD use the country/city format specified by the [IANA Time Zone Database](#).

11.17.10.4. TimeFailure Event

This event SHALL be generated if the node has not generated a TimeFailure event in the last hour, and the node is unable to get a time from any source. This event SHOULD NOT be generated more often than once per hour.

11.17.10.5. MissingTrustedTimeSource Event

This event SHALL be generated if the TrustedTimeSource is set to null upon fabric removal or by a [SetTrustedTimeSource](#) command.

This event SHALL also be generated if the node has not generated a MissingTrustedTimeSource event in the last hour, and the node fails to update its time from the TrustedTimeSource because the TrustedTimeSource is null or the specified peer cannot be reached. MissingTrustedTimeSource events corresponding to a time update SHOULD NOT be generated more often than once per hour.

11.17.11. Time Synchronization at Commissioning

During commissioning the Commissioner SHOULD set the UTCTime, and set up the TrustedTimeSource, DefaultNTP, TimeZone and DSTOffsets as required. See [Commissioning Flows](#). Note that the Commissionee might not be able to support the [Time Synchronization Cluster](#) server. Note that the commissioner MAY opt to not set the time so the node SHOULD NOT depend on having time during commissioning.

11.17.12. Time Synchronization during operation

Nodes MAY perform time synchronization using a trusted external source, [RFC 5905](#) or by reading the UTC time from the Time Synchronization cluster of another Node with the desired Granularity and TimeSource. When using NTPv4 / SNTPv4, Nodes capable of external communication SHOULD use network time security (NTS) if that is available on the server ([RFC 8915](#)). This specification DOES NOT mandate that Nodes should include a TCP/TLS stack and a set of adequate Root CA certificates solely to support NTS, but that if a Node already has these capabilities, then it SHOULD implement and attempt NTS for NTP.

Nodes SHOULD attempt to perform a time synchronization after a restart or upon any change of Node state where timekeeping was lost. Nodes SHOULD attempt this time synchronization prior to using any security material which may have expired. If a Node is unable to achieve time synchronization using the steps outlined in [Section 11.17.13, “Time source prioritization”](#), the Node MAY retry or fall back to the stored Last Known Good UTC Time ([Section 3.5.6.1, “Last Known Good UTC Time”](#)).

Since Matter does not mandate external network connectivity for use, device manufacturers should be aware that it is possible to be operating in a network where time is not available from any source. Product manufacturers may wish to account for this possibility when specifying their clock drift requirements. The estimated clock drift is reported using the [Granularity](#) attribute.

11.17.13. Time source prioritization

Nodes SHALL prioritize time synchronization sources in the following order:

1. GNSS time source, if supported by the Node.
2. Trusted, high-resolution, external time source (ex. PTP, network NTP, trusted cloud-based source), if supported by the Node and external connectivity is available.
3. If NTPClient (NTPC) feature is supported: NTP server defined in the DHCPv6 NTP server option, if DHCPv6 is supported on the Node.
4. If NTPClient (NTPC) feature is supported: NTP server defined by DHCPv4 if the Node supports IPv4.
5. If NTPClient (NTPC) feature is supported: NTP server identified by a DNS-SD query for `_ntp._udp`. If multiple servers respond, Nodes with full NTP support SHOULD query multiple servers. Nodes using SNTP SHOULD select any server from the list.
6. If TimeSyncClient (TSC) feature is supported: the Time Synchronization cluster of another Node SHOULD be queried in the following order:
 - a. TrustedTimeSource provided by an Administrator.
 - b. Any of the Node’s current peers per any data model binding that support the Time Synchronization cluster and have the desired Granularity and TimeSource.
 - c. Enumerate all Nodes using DNS-SD query for `_matter._tcp` and select a peer that supports the Time Synchronization cluster and has the desired Granularity and TimeSource.
7. If NTPClient (NTPC) feature is supported: Fallback NTP server defined during commissioning.

Nodes that use GNSS or a trusted external source SHOULD check the remaining time synchronization sources to determine if they SHOULD act as an NTP server (see [Section 11.17.15, “Acting as an NTP Server”](#)).

11.17.14. Time synchronization maintenance

Nodes SHALL adjust their Granularity attribute based on their assessed time synchronization error. Nodes running an NTP server ([Section 11.17.15, “Acting as an NTP Server”](#)) SHALL maintain a Granularity of SecondsGranularity or better and SHOULD maintain an accuracy of MillisecondsGranularity or better.

A Node with a Granularity of NoTimeGranularity SHALL attempt to sync its time at least once a day. Nodes SHOULD NOT query the Time Synchronization cluster of another Node more than once per 30 minutes.

11.17.15. Acting as an NTP Server

Any capable Node with always-on power source that has and can maintain a time synchronization Granularity of MillisecondsGranularity or better, SHOULD act as an NTP server. Any capable Node that reaches the final stage in the [Time source prioritization](#) mechanism SHOULD act as an NTP server for the network.

A Node hosting an NTP server SHALL update the NTPServerAvailable attribute and SHOULD advertise an _ntp._udp DNS-SD service. A Node hosting an NTP server SHOULD support network time security (NTS). A Node hosting an NTP server SHOULD implement leap smearing.

11.17.16. Implementation Guidance

This specification offers several options for getting time in order to support Nodes with various capabilities. This section is not intended as a requirement, but is used to illustrate how various Nodes might implement this specification.

11.17.16.1. Example: Constrained Node with no schedule capabilities

This type of Node would not need to support the TimeZone (TZ) feature or the NTPServer (NTPS) feature. A constrained node can choose to get time either by using a time synchronization client cluster to read the time from a trusted node (if the TSC feature is supported), and / or by using the NTPClient (if the NTPC feature is supported). A constrained node supporting NTPC would likely use an SNTP implementation.

A node supporting both the NTPC and TSC features would have the following attributes:

- UTCTime
- Granularity
- TimeSource (optional)
- TrustedTimeSource
- DefaultNTP

To achieve time synchronization, the Node would start at Step 3 of [Time source prioritization](#) (Steps 1 and 2 do not apply).

If the Node opts to support only the TSC feature, it would have the following attributes:

- UTCTime
- Granularity
- TimeSource (optional)
- TrustedTimeSource

To achieve time synchronization, the Node would start at step 6 (querying a Time Synchronization cluster) of [Time source prioritization](#) (other steps do not apply).

If the Node opts to support only the NTPC feature, it would have the following attributes:

- UTCTime
- Granularity
- TimeSource (optional)
- DefaultNTP

To achieve time synchronization, the Node would start at step 3 of [Time source prioritization](#) and would bypass step 6.

In all cases, if the Node wishes to maintain time synchronization by re-querying, it would set its Granularity attribute as appropriate at each update. This would probably be either MillisecondsGranularity or SecondsGranularity depending on the round-trip delay and the frequency of updates. If the Node wishes to use a single time synchronization, but will not continue to synchronize during operation, it may either set the Granularity as appropriate and downgrade as the clock drifts, or may simply opt to set the Granularity to MinutesGranularity. If supported, the TimeSource attribute would be set as appropriate.

11.17.16.2. Example: Intermittently connected device with schedule capability

This type of Node would need to support the TimeZone (TZ) feature to support scheduling, but would not support the NTPServer (NTPS) feature as it is an intermittently connected device and would not make a reliable server. As with the [Constrained Node with no schedule capabilities](#) example, this type of Node can choose to get time either by using the time synchronization client cluster to read the time from a trusted node (if the TSC feature is supported) and / or by using an NTPClient (if the NTPC feature is supported).

This type of Node would support all the attributes described in the [Constrained Node with no schedule capabilities](#) example with the appropriate features. Because it additionally supports the TimeZone (TZ) feature, it would also support the following attributes:

- TimeZoneListMaxSize
- TimeZone
- DSTOffsetListMaxSize
- DSTOffset
- LocalTime
- TimeZoneDatabase

During commissioning, the Node should receive time zone information from the commissioner. If the Node has access to a time zone database, it can opt to fill its own DSTOffset attribute, otherwise it should also receive DSTOffset information from the Commissioner. During operation, this type of Node would use the same methods of getting time as the [Constrained Node with no schedule capabilities](#) example, setting the Granularity and optional TimeSource as appropriate. It would also use the information from the TimeZone and DSTOffset to calculate a local time.

11.17.16.3. Example: Hub-like Node

These types of Nodes are normally relatively capable, high-powered and always-on. These would most likely include time zone support for scheduling and display. Often these will have significant software beyond this specification and are likely to already have built-in mechanisms for time synchronization. Such Nodes SHOULD support an NTP server to serve time to other Nodes in the network. This type of Node may opt to support the TimeSyncClient (TSC) and NTPClient (NTPC) as backup, but can omit these if it has a reliable time source. It would support the following attributes based on the supported features:

- UTCTime
- Granularity
- TimeSource (optional)
- TrustedTimeSource (if TSC supported)
- DefaultNTP (if NTPC supported)
- TimeZoneListMaxSize
- TimeZone
- DSTOffsetListMaxSize
- DSTOffset
- LocalTime
- TimeZoneDatabase
- NTPServerAvailable

To achieve time synchronization, these Nodes would likely use Step 1 or Step 2 of [Time source prioritization](#), using the remaining options as fallback if necessary. The Node would then set its Granularity and TimeSource as appropriate, and would maintain its times, Granularity and TimeSource. The Node SHOULD also start an NTP server, update the NTPServerAvailable attribute and advertise using DNS-SD on _ntp._udp.

If a time-zone database is supported (Node can calculate DST times from TimeZone settings), these Nodes MAY subscribe to the DSTTableEmpty Events of Nodes with no TimeZoneDatabase support. Upon receipt of these events the Node SHOULD calculate and set new DST values for such Nodes by writing the DSTOffset attribute. These nodes MAY also subscribe to the TimeFailure and MissingTrustedTimeSource events for Nodes in its fabric and supply time or time sources as required.

11.18. Operational Credentials Cluster

This cluster is used to add or remove Node Operational credentials on a Commissionee or already-configured Node, as well as manage the associated Fabrics.

11.18.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision
2	Introduction of command, attributes and fields to support the Fabric Table Vendor ID Verification Procedure . The procedure increased flash storage cost by approximately 90 bytes per fabric.

NOTE

The name was re-normalized from "Node Operational Credentials" cluster to "Operational Credentials" cluster in Matter 1.4.2, to match much of the implementation legacy that used the latter name.

11.18.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	OPCREDS

11.18.3. Cluster ID

ID	Name
0x003E	Operational Credentials

11.18.4. Data Types

11.18.4.1. RESP_MAX Constant Type

A RESP_MAX constant appears in the description of some command fields in this cluster and within the description of some associated serialization schemes.

The value RESP_MAX SHALL be 900 bytes.

The current limit of 900 bytes was chosen to accommodate the maximum size of IPv6 frames, transport headers, message layer headers and integrity protection and Interaction Model protocol encoding, while accounting for sufficient remaining space for signatures and to allow extensions to larger key and digest sizes in the future.

11.18.4.2. CertificateChainTypeEnum Type

This data type is derived from [enum8](#).

This enumeration is used by the [CertificateChainRequest](#) command to convey which certificate from the device attestation certificate chain to transmit back to the client.

Value	Name	Summary	Conformance
1	DACCertificate	Request the DER-encoded DAC certificate	M

Value	Name	Summary	Conformance
2	PAICertificate	Request the DER-encoded PAI certificate	M

11.18.4.3. NodeOperationalCertStatusEnum Type

This data type is derived from [enum8](#).

This enumeration is used by the [NOCResponse](#) common response command to convey detailed outcome of several of this cluster's operations.

Value	Name	Summary	Conformance
0	OK	OK, no error	M
1	InvalidPublicKey	Public Key in the NOC does not match the public key in the NOCSR	M
2	InvalidNodeOpId	The Node Operational ID in the NOC is not formatted correctly.	M
3	InvalidNOC	Any other validation error in NOC chain	M
4	MissingCsr	No record of prior CSR for which this NOC could match	M
5	TableFull	NOCs table full, cannot add another one	M
6	InvalidAdminSubject	Invalid CaseAdminSubject field for an AddNOC command.	M
9	FabricConflict	Trying to AddNOC instead of UpdateNOC against an existing Fabric.	M
10	LabelConflict	Label already exists on another Fabric.	M
11	InvalidFabricIndex	FabricIndex argument is invalid.	M

11.18.4.4. NOCStruct Type

This encodes a [NOC](#) chain, underpinning a commissioned Operational Identity for a given Node.

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	NOC	octstr	max 400				M
2	ICAC	octstr	max 400	X			M
3	VVSC	octstr	max 400				desc

NOTE

The VVSC field is mutually exclusive with the ICAC field. If the ICAC field is non-null, the VVSC field SHALL be omitted. If the VVSC field is present in the structure, the ICAC field SHALL be null. The reason for this is to optimize storage usage, as the VID Verification Signer Certificate (VVSC) is a field that is only needed in root-per-fabric situations without ICAC present.

NOTE

The [Trusted Root CA Certificate \(RCAC\)](#) is not included in this structure. The roots are available in the [TrustedRootCertificates](#) attribute under the same [associated fabric](#) as the one for the NOCStruct entry.

NOC Field

This field SHALL contain the [NOC](#) for the struct's [associated fabric](#), encoded using [Matter Certificate Encoding](#).

ICAC Field

This field SHALL contain the [ICAC](#) for the struct's [associated fabric](#), encoded using [Matter Certificate Encoding](#). If no ICAC is present in the chain, this field SHALL be set to null.

VVSC Field

This field SHALL contain the [Vendor Verification Signer Certificate \(VVSC\)](#) for the struct's [associated fabric](#), encoded using [Matter Certificate Encoding](#). If no VVSC is needed, this field SHALL be omitted (in that there SHALL NOT be a value present, not even an empty octet string). If the ICAC field is non-null, this field SHALL NOT be present.

11.18.4.5. FabricDescriptorStruct Type

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	RootPublicKey	octstr	65				M
2	VendorID	vendor-id	desc				M
3	FabricID	fabric-id					M
4	NodeID	node-id					M

Access Modifier: Fabric Scoped							
5	Label	string	max 32		""		M
6	VIDVerification-Statement	octstr	85				O

This structure encodes a [Fabric Reference](#) for a fabric within which a given Node is currently commissioned.

RootPublicKey Field

This field SHALL contain the public key for the trusted root that scopes the fabric referenced by FabricIndex and its associated operational credential (see [Section 6.4.5.3, “Trusted Root CA Certificates”](#)). The format for the key SHALL be the same as that used in the ec-pub-key field of the [Matter Certificate Encoding](#) for the root in the operational certificate chain.

VendorID Field

This field SHALL contain the value of VendorID associated with the fabric.

This value SHALL have been provided by the AdminVendorID value provided in the [AddNOC](#) command that led to the creation of this FabricDescriptorStruct, or the value updated via the [SetVIDVerificationStatement](#) command, whichever was last completed. The set of allowed values is defined in [AdminVendorID](#).

The intent is to provide user transparency about which entities have Administer privileges on the Node.

Clients SHALL consider the VendorID field value to be untrustworthy until the [Fabric Table Vendor ID Verification Procedure](#) has been executed against the fabric entry having this VendorID.

FabricID Field

This field SHALL contain the FabricID allocated to the fabric referenced by FabricIndex. This field SHALL match the value found in the matter-fabric-id field from the operational certificate providing the operational identity under this Fabric.

NodeID Field

This field SHALL contain the NodeID in use within the fabric referenced by FabricIndex. This field SHALL match the value found in the matter-node-id field from the operational certificate providing this operational identity.

Label Field

This field SHALL contain a commissioner-set label for the fabric referenced by FabricIndex. This field is set by the [UpdateFabricLabel](#) command.

VIDVerificationStatement Field

This field, if present, SHALL contain a previously-installed administrator-set `vid_verification_statement` value (see [Section 6.4.10, “Fabric Table Vendor ID Verification Procedure”](#)) for the fabric referenced by `FabricIndex`. This field is set by the `SetVIDVerificationStatement` command.

11.18.4.6. Attestation Elements

The Attestation Elements contain the metadata related to attestation, encoded in [Matter TLV](#).

Attestation Elements TLV

```
attestation-elements => STRUCTURE [ tag-order ]
{
    certification_declaration[1]      : OCTET STRING,
    attestation_nonce[2]              : OCTET STRING [ length 32 ],
    timestamp[3]                     : UNSIGNED INTEGER [ range 32-bits ],
    firmware_information[4, optional] : OCTET STRING,
}
```

Any context-specific tags not listed in the above schema for Attestation Elements SHALL be reserved for future use, and SHALL be silently ignored if seen by a Commissioner which cannot understand them.

11.18.4.7. Attestation Information

The Attestation Information is the combination of a [Matter TLV](#) payload, containing the [Attestation Elements](#), as well as a signature over an `attestation_tbs` message containing the in-band-transmitted `attestation_elements_message` and a secret out-of-band Attestation Challenge.

The Attestation Information SHALL be computed as follows:

1. Encode the [attestation-elements](#) structure with an anonymous tag into an octet string called `attestation_elements_message`.
 - The `firmware_information` field SHALL be an octet string, as described in [Section 6.3.2, “Firmware Information”](#).
 - The `certification_declaration` field SHALL be the DER-encoded octet string representation of a [CMS\(RFC5652\)-encoded](#) certification declaration, as described in [Section 6.3, “Certification Declaration”](#).
 - The `timestamp` field SHALL be in [epoch-s](#).
 - The `attestation_nonce` field SHALL match the `AttestationNonce` field provided in the [AttestationRequest Command](#) that triggered the generation of the Attestation Elements.
 - Vendor specific information, if present, SHALL be encoded using [fully qualified tags](#). Such fields allow the Node taking part in the Device Attestation Procedure to communicate vendor-specific information that MAY aid in device commissioning. Commissioners that do not understand the format of the data MAY ignore them.

- The resulting `attestation_elements_message`, including optional fields, SHALL be no more than `RESP_MAX` bytes once serialized.

```
attestation_elements_message =
{
    certification_declaration(1) = certification_declaration,
    attestation_nonce(2)         = attestation_nonce,
    timestamp(3)                 = timestamp,
    firmware_information(4)       = firmware_information,

    ... optional fields per attestation-elements ..
}
```

2. Obtain the AttestationChallenge from a [CASE session](#), [resumed CASE session](#), or [PASE session](#) depending on the method used to establish the current session. This is an octet string of length `CRYPTO_SYMMETRIC_KEY_LENGTH_BITS`. Save it for the next step as `attestation_challenge`.
3. Locally generate an `attestation_tbs` message as an octet string by concatenating the `attestation_elements_message` and the `attestation_challenge`:
 - `attestation_tbs = attestation_elements_message || attestation_challenge`
4. Compute the [attestation_signature](#) and record it as an [ec-signature](#) octet string:

```
attestation_signature = Crypto_Sign(
    message = attestation_tbs,
    privateKey = Device Attestation Private Key
)
```

5. Fill the AttestationElements field of the [AttestationResponse Command](#) using the contents of the `attestation_elements_message` octet string.
6. Fill the AttestationSignature field of the [AttestationResponse Command](#) using the contents of the `attestation_signature` octet string.
7. Note that the `attestation_challenge` SHALL NOT be included in any of the payloads conveyed as part of the Attestation Information.

See [Section F.2, “Device Attestation Response test vector”](#) for an example computation of the above messages and application payloads.

11.18.4.8. NOCSR Elements

The NOCSR Elements contain the metadata related to NOCSR, encoded in [Matter TLV](#).

NOCSR Elements TLV

```
nocsr-elements => STRUCTURE [ tag-order ]
{
```

```

csr[1] : OCTET STRING,
CSRNonce[2] : OCTET STRING [ length 32 ],
vendor_reserved1[3, optional] : OCTET STRING,
vendor_reserved2[4, optional] : OCTET STRING,
vendor_reserved3[5, optional] : OCTET STRING
}

```

Any context-specific tags not listed in the above schema for NOCSR Elements SHALL be reserved for future use, and SHALL be silently ignored if seen by a Commissioner which cannot understand them.

11.18.4.9. NOCSR Information

The NOCSR Information is the combination of a [Matter TLV](#) payload, containing the [NOCSR Elements](#), as well as a signature over an `nocsr_tbs` message containing the in-band-transmitted `nocsr_elements_message` and a secret out-of-band Attestation Challenge, using the Attestation Private Key that is unique to the device producing the NOCSR Information.

The NOCSR Information SHALL be computed as follows:

1. Encode the [nocsr-elements](#) structure with an anonymous tag into an octet string called `nocsr_elements_message`.
 - The `csr` field SHALL be a DER-encoded octet string of a properly encoded [PKCS #10](#) Certificate Signing Request (CSR), signed with the [Node Operational Private Key](#) associated with the Node Operational Public Key, which is the `subjectPublicKey` field of the CSR. See [Section 6.4.6.1, “Node Operational Certificate Signing Request \(NOCSR\) Procedure”](#) for details about the generation of the Node Operational Key Pair, and the contents of the CSR.
 - The `CSRNonce` field SHALL match the CSR Nonce field in the corresponding [CSRRequest Command](#).
 - The `vendor_reserved1` through `vendor_reserved3` fields allow the Node taking part in the NOCSR Procedure to communicate vendor-specific information that MAY aid in device commissioning. Commissioners that do not understand the format of the data MAY ignore them.
 - The resulting `nocsr_elements_message`, including optional fields, SHALL be no more than [RESP_MAX](#) bytes once serialized.

```

nocsr_elements_message =
{
    csr(1)      = node_operational_csr_der_bytes,
    CSRNonce(2) = CSRNonce,

    ... optional fields per nocsr-elements ..
}

```

2. Obtain the AttestationChallenge from a [CASE session](#), [resumed CASE session](#), or [PASE session](#)

depending on the method used to establish the current session. This is an octet string of length [CRYPTO_SYMMETRIC_KEY_LENGTH_BITS](#). Save it for the next step as `attestation_challenge`.

3. Locally generate an `nocsr_tbs` message as an octet string by concatenating the `nocsr_elements_message` and the `attestation_challenge`:
 - `nocsr_tbs = nocsr_elements_message || attestation_challenge`
4. Compute the [attestation_signature](#) and record it as an [ec-signature](#) octet string:

```
attestation_signature = Crypto_Sign(
    message = nocsr_tbs,
    privateKey = Device Attestation Private Key
)
```

5. Fill the `NOCSRElements` field of the [CSRResponse Command](#) using the contents of the `nocsr_elements_message` octet string.
6. Fill the `AttestationSignature` field of the [CSRResponse Command](#) using the contents of the `attestation_signature` octet string.
7. Note that the `attestation_challenge` SHALL NOT be included in any of the payloads conveyed as part of the NOCSR Information.

See [Section F.3, “Node Operational CSR Response test vector”](#) for an example computation of the above messages and application payloads.

11.18.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	NOCs	list[NOC-Struct]	max SupportedFabrics	C N		R A F	M
0x0001	Fabrics	list[FabricDescriptorStruct]	max SupportedFabrics	N		R V F	M
0x0002	SupportedFabrics	uint8	5 to 254	F		R V	M
0x0003	Commissioned-Fabrics	uint8	max SupportedFabrics	N		R V	M
0x0004	Trusted-RootCertificates	list[octstr]	max SupportedFabrics[max 400]	C N		R V	M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0005	Current-FabricIndex	uint8			0	R V	M

11.18.5.1. NOCs Attribute

This attribute SHALL contain all NOCs applicable to this Node, encoded as a read-only list of [NOC-Struct](#).

Operational Certificates SHALL be added through the [AddNOC](#) command, and SHALL be removed through the [RemoveFabric](#) command.

Upon Factory Data Reset, this attribute SHALL be set to a default value of an empty list.

The number of entries in this list SHALL match the number of entries in the Fabrics attribute.

11.18.5.2. Fabrics Attribute

This attribute SHALL indicate all fabrics to which this Node is commissioned, encoded as a read-only list of [FabricDescriptorStruct](#). This information MAY be computed directly from the NOCs attribute.

The Fabrics attribute is also known as "the fabric table".

Upon Factory Data Reset, this attribute SHALL be set to a default value of an empty list.

The number of entries in this list SHALL match the number of entries in the NOCs attribute.

11.18.5.3. SupportedFabrics Attribute

This attribute SHALL indicate the number of Fabrics that are supported by the device. This value is fixed for a particular device.

11.18.5.4. CommissionedFabrics Attribute

This attribute SHALL indicate the number of Fabrics to which the device is currently commissioned. This attribute SHALL be equal to the following:

- The number of entries in the NOCs attribute.
- The number of entries in the Fabrics attribute.

Upon Factory Data Reset, this attribute SHALL be set to a default value of 0.

11.18.5.5. TrustedRootCertificates Attribute

This attribute SHALL contain the list of [Trusted Root CA Certificates \(RCAC\)](#) installed on the Node, as octet strings containing their [Matter Certificate Encoding](#) representation.

These certificates are installed through the [AddTrustedRootCertificate](#) command.

Depending on the method of storage employed by the server, either shared storage for identical root certificates shared by many fabrics, or individually stored root certificate per fabric, multiple identical root certificates MAY legally appear within the list.

To match a root with a given fabric, the root certificate's subject and subject public key need to be cross-referenced with the NOC or ICAC certificates that appear in the [NOCs](#) attribute for a given fabric.

Upon Factory Data Reset, this attribute SHALL be set to a default value whereby the list is empty.

11.18.5.6. CurrentFabricIndex Attribute

This attribute SHALL indicate the [accessing fabric index](#).

This attribute is useful to contextualize [Fabric-Scoped](#) entries obtained from response commands or attribute reads, since a given Fabric may be referenced by a different Fabric Index locally on a remote Node.

11.18.6. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	AttestationRequest	client ⇒ server	AttestationResponse	A	M
0x01	AttestationResponse	client ⇐ server	N		M
0x02	Certificate-ChainRequest	client ⇒ server	Certificate-ChainResponse	A	M
0x03	Certificate-ChainResponse	client ⇐ server	N		M
0x04	CSRRequest	client ⇒ server	CSRResponse	A	M
0x05	CSRResponse	client ⇐ server	N		M
0x06	AddNOC	client ⇒ server	NOCResponse	A	M
0x07	UpdateNOC	client ⇒ server	NOCResponse	A F	M
0x08	NOCResponse	client ⇐ server	N		M
0x09	UpdateFabricLabel	client ⇒ server	NOCResponse	A F	M
0x0A	RemoveFabric	client ⇒ server	NOCResponse	A	M
0x0B	AddTrusted-RootCertificate	client ⇒ server	Y	A	M

0x0C	SetVIDVerificationStatement	client $\Rightarrow$ server	Y	A F	M
0x0D	SignVIDVerificationRequest	client $\Rightarrow$ server	SignVIDVerificationResponse	A	M
0x0E	SignVIDVerificationResponse	client $\Leftarrow$ server	N		M

NOTE

The SetVIDVerificationStatement, SignVIDVerificationRequest and SignVIDVerificationResponse commands appeared and were made mandatory in cluster revision 2, in specification version 1.4.2. Prior revisions of this cluster's servers will NOT have the necessary commands. Clients SHOULD use the AcceptedCommandList and GeneratedCommandList global attributes, and/or the cluster revision, to determine support of the feature, if they intend to rely on it.

11.18.6.1. AttestationRequest Command

This command is used to perform an attestation request.

This command SHALL be generated to request the [Attestation Information](#), in the form of an [AttestationResponse Command](#). If the AttestationNonce that is provided in the command is malformed, a recipient SHALL fail the command with a Status Code of INVALID_COMMAND. The Attestation-Nonce field SHALL be used in the computation of the Attestation Information.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Attestation-Nonce	octstr	32			M

11.18.6.2. AttestationResponse Command

This command is used to report the results of the AttestationRequest command. This command SHALL be generated in response to an [Attestation Request command](#).

See [Section 11.18.4.7, “Attestation Information”](#) for details about the generation of the fields within this response command.

See [Section F.2, “Device Attestation Response test vector”](#) for an example computation of an AttestationResponse.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	AttestationElements	octstr	max RESP_MAX			M

ID	Name	Type	Constraint	Quality	Fallback	Conformance
1	AttestationSignature	octstr	64			M

AttestationElements Field

This field SHALL contain the octet string of the serialized attestation_elements_message.

AttestationSignature Field

This field SHALL contain the octet string of the necessary attestation_signature as described in [Section 11.18.4.7, “Attestation Information”](#).

11.18.6.3. CertificateChainRequest Command

This command is used to request a certificate from the device attestation certificate chain.

If the CertificateType is not a valid value per [CertificateChainTypeEnum](#) then the command SHALL fail with a Status Code of INVALID_COMMAND.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CertificateType	CertificateChainTypeEnum	desc			M

11.18.6.4. CertificateChainResponse Command

This command is used to report the results of the CertificateChainRequest command. This command SHALL be generated in response to a [CertificateChainRequest command](#).

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Certificate	octstr	max 600			M

Certificate Field

This field SHALL be the DER encoded certificate corresponding to the CertificateType field in the [CertificateChainRequest command](#).

11.18.6.5. CSRRequest Command

This command is used to perform a CSR request.

This command SHALL be generated to execute the [Node Operational CSR Procedure](#) and subsequently return the [NOCSR Information](#), in the form of a [CSRResponse Command](#).

The CSRNonce field SHALL be used in the computation of the NOCSR Information. If the CSRNonce is malformed, then this command SHALL fail with an INVALID_COMMAND status code.

If the IsForUpdateNOC field is present and set to true, but the command was received over a [PASE](#) session, the command SHALL fail with an INVALID_COMMAND status code, as it would never be possible to use a resulting subsequent certificate issued from the CSR with the [UpdateNOC](#) command, which is forbidden over [PASE](#) sessions.

If the IsForUpdateNOC field is present and set to true, the internal state of the CSR associated key pair SHALL be tagged as being for a subsequent UpdateNOC, otherwise the internal state of the CSR SHALL be tagged as being for a subsequent AddNOC. See [Section 11.18.6.8, “AddNOC Command”](#) and [Section 11.18.6.9, “UpdateNOC Command”](#) for details about the processing.

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

If a prior UpdateNOC or AddNOC command was successfully executed within the fail-safe timer period, then this command SHALL fail with a CONSTRAINT_ERROR status code sent back to the initiator.

If the Node Operational Key Pair generated during processing of the [Node Operational CSR Procedure](#) is found to collide with an existing key pair already previously generated and installed, and that check had been executed, then this command SHALL fail with a FAILURE status code sent back to the initiator.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CSRNonce	octstr	32			M
1	IsForUpdateNOC	bool			false	O

11.18.6.6. CSRResponse Command

This command is used to report the results of the CSRRequest command. This command SHALL be generated in response to a [CSRRequest Command](#).

See [Section 11.18.4.9, “NOCSR Information”](#) for details about the generation of the fields within this response command.

See [Section F.3, “Node Operational CSR Response test vector”](#) for an example computation of a CSR-Response.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NOCSRElements	octstr	max RESP_MAX			M

ID	Name	Type	Constraint	Quality	Fallback	Conformance
1	AttestationSignature	octstr	64			M

NOCSRElements Field

This field SHALL contain the octet string of the serialized `nocsr_elements_message`.

AttestationSignature Field

This field SHALL contain the octet string of the necessary `attestation_signature` as described in [Section 11.18.4.9, “NOCSR Information”](#).

11.18.6.7. AddNOC and UpdateNOC Commands Overview

The AddNOC command is used to commission a Node into a Fabric by providing a usable [NOC](#) and [ICAC](#), with associated Node Operational IDs.

The UpdateNOC command is used to update existing credentials within a Fabric, for the purposes of:

- Rotating the Node Operational Key Pair in use
- Updating the contents of the [NOC](#) and optionally the [ICAC](#) certificates (subjects, issuers, keys, etc) under the current root of trust and Fabric

Both of these commands receive an `NOCValue` and optional `ICACValue` fields and require some common validation in addition to their specific behavior.

NOCValue and ICACValue Fields

The `NOCValue` and `ICACValue` fields SHALL be octet strings that represent a certificate encoded using [Matter Certificate Encoding](#).

Upon receipt, the `NOCValue` and `ICACValue` chain SHALL be validated in the following ways:

1. Verify the NOC using:
 - a. [Crypto_VerifyChain\(certificates = \[NOCValue, ICACValue, RootCACertificate\]\)](#) if `ICACValue` is present,
 - b. [Crypto_VerifyChain\(certificates = \[NOCValue, RootCACertificate\]\)](#) if `ICACValue` is not present.

If this verification fails, the error status SHALL be `InvalidNOC`.

2. The public key of the NOC SHALL match the last generated operational public key on this session, and therefore the public key present in the last [CSRResponse](#) provided to the Administrator or Commissioner that sent the [AddNOC](#) or [UpdateNOC](#) command. If this check fails, the error status SHALL be `InvalidPublicKey`.
3. The [DN Encoding Rules](#) SHALL be validated for every certificate in the chain, including valid

value range checks for identifiers such as Fabric ID and Node ID. If this validation fails, the error status SHALL be `InvalidNodeOpId` if the `matter-node-id` attribute in the subject DN of the NOC has a value outside the Operational Node ID range and `InvalidNOC` for all other failures.

If any of the above validation checks fail, the server SHALL immediately respond to the client with an `NOCResponse`. The `StatusCode` field of the `NOCResponse` SHALL be set to the error status value specified in the above validation checks.

These certificate validation steps are performed to ensure that operational credentials installed match an operational key pair generated by the Device and respect the trust model assumptions expressed in [Section 6.4.5.1, “Node Operational Certificate \(NOC\)”](#).

Handling Errors

For any error described in the following subsections, the device SHALL immediately respond to the client with an `NOCResponse` with the prescribed `StatusCode` field, and SHALL leave all non-volatile state of the device untouched, as if the `AddNOC` command had never been received. The information about the last CSR state associated with this session SHALL also be untouched in this case, so that a valid `AddNOC` command MAY still be issued later that would match that CSR state. The `DebugText` field in the `NOCResponse` MAY be filled with debug information.

The following failed preconditions error cases apply to all invocations of `AddNOC`:

- If the device already has the `CommissionedFabrics` attribute equal to the `SupportedFabrics` attribute, then the device’s operational credentials table is considered full and the device SHALL process the error by responding with a `StatusCode` of `TableFull` as described in [Section 11.18.6.7.2, “Handling Errors”](#).
- If no context or memory exists of a prior `CSRRequest` command having been invoked in the **same secure session** as that which is receiving this `AddNOC` or `UpdateNOC` invocation, then the Node SHALL process the error by responding with a `StatusCode` of `MissingCsr` as described in [Section 11.18.6.7.2, “Handling Errors”](#).

11.18.6.8. AddNOC Command

This command is used to add a new NOC to the device.

This command SHALL add a new `NOC` chain to the device and commission a new Fabric association upon successful validation of all arguments and preconditions.

The new value SHALL immediately be reflected in the `NOCs` list attribute.

A Commissioner or Administrator SHALL issue this command after issuing the `CSRRequest` command and receiving its response.

A Commissioner or Administrator SHOULD issue this command after performing the [Attestation Procedure](#).

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NOCValue	octstr	max 400			M
1	ICACValue	octstr	max 400			O
2	IPKValue	octstr	16			M
3	CaseAdmin-Subject	subject-id				M
4	AdminVendorId	vendor-id				M

IPKValue Field

This field SHALL contain the value of the Epoch Key for the Identity Protection Key (IPK) to set for the Fabric which is to be added. This is needed to bootstrap a necessary configuration value for subsequent [CASE](#) to succeed. See [Section 4.14.2.6.1, “Identity Protection Key \(IPK\)”](#) for details.

The IPK SHALL be provided as an octet string of length CRYPTO_SYMMETRIC_KEY_LENGTH_BYTES.

On successful execution of the AddNOC command, the side-effect of having provided this field SHALL be equivalent to having done a [GroupKeyManagement](#) cluster [KeySetWrite](#) command invocation using the newly joined fabric as the [accessing fabric](#) and with the following argument fields (assuming KeySetWrite allowed a GroupKeySetID set to 0):

```

KeySetWrite
(
    GroupKeySetStruct :=
    {
        GroupKeySetID := 0,
        GroupKeySecurityPolicy := 0,
        EpochKey0 := <Contents of IPKValue field>,
        EpochStartTime0 := 0,
        EpochKey1 := null
        EpochStartTime1 := null
        EpochKey2 := null,
        EpochStartTime2 := null
    }
)

```

CaseAdminSubject Field

If the AddNOC command succeeds according to the semantics of the following subsections, then the Access Control [subject-id](#) SHALL be used to atomically add an [Access Control Entry](#) enabling that Subject to subsequently administer the Node whose operational identity is being added by this command.

The format of the new Access Control Entry, created from this, SHALL be:

```
{
  FabricIndex: **FabricIndex derived from current or new Fabric**,
  Privilege: Administer,
  AuthMode: CASE,
  Subjects: [**CaseAdminSubject provided in the AddNOC command**],
  Targets: [],      // entire node
  Extension: []     // empty
}
```

NOTE

Unless such an Access Control Entry is added atomically as described here, there would be no way for the caller on its given Fabric to eventually add another [Access Control Entry](#) for CASE authentication mode, to enable the new Administrator to administer the device, since the [Fabric Scoping](#) of the Access Control List prevents the current Node from being able to write new entries scoped to that Fabric, if the session is established from [CASE](#). While a session established from [PASE](#) does gain Fabric Scope of a newly-joined Fabric, this argument is made mandatory to provide symmetry between both types of session establishment, both of which need to eventually add an "Administer Node over CASE" Access Control Entry to finalize new Fabric configuration and subsequently be able to call the [CommissioningComplete](#) command.

AdminVendorID Field

This field SHALL be set to the [Vendor ID](#) of the entity issuing the AddNOC command. This value SHALL NOT be one of the reserved Vendor ID values defined in [Table 1, "Vendor ID Allocations"](#).

Effect on Receipt

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, "ArmFail-Safe"](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

If a prior UpdateNOC or AddNOC command was successfully executed within the fail-safe timer period, then this command SHALL fail with a CONSTRAINT_ERROR status code sent back to the initiator.

If the prior [CSRRequest](#) state that preceded AddNOC had the IsForUpdateNOC field indicated as true, then this command SHALL fail with a CONSTRAINT_ERROR status code sent back to the initiator.

If no prior [AddTrustedRootCertificate](#) command was successfully executed within the fail-safe timer period, then this command SHALL process an error by responding with a [NOCResponse](#) with a StatusCode of [InvalidNOC](#) as described in [Section 11.18.6.7.2, "Handling Errors"](#). In other words, AddNOC always requires that the client provides the root of trust certificate within the same Fail-Safe context as the rest of the new fabric's operational credentials, even if some other fabric

already uses the exact same root of trust certificate.

If the NOC provided in the NOCValue encodes an [Operational Identifier](#) for a <Root Public Key, FabricID> pair already present on the device, then the device SHALL process the error by responding with a StatusCode of [FabricConflict](#) as described in [Section 11.18.6.7.2, “Handling Errors”](#).

If the device already has the CommissionedFabrics attribute equal to the SupportedFabrics attribute, then the device’s operational credentials table is considered full and the device SHALL process the error by responding with a StatusCode of [TableFull](#) as described in [Section 11.18.6.7.2, “Handling Errors”](#).

If the [CaseAdminSubject](#) field is not a valid [ACL subject](#) in the context of [AuthMode](#) set to CASE, such as not being in either the [Operational](#) or [CASE Authenticated Tag](#) range, then the device SHALL process the error by responding with a StatusCode of [InvalidAdminSubject](#) as described in [Section 11.18.6.7.2, “Handling Errors”](#).

Otherwise, the command is considered an addition of credentials, also known as "joining a fabric", and the following SHALL apply:

1. A new FabricIndex SHALL be allocated, taking the next valid [fabric-index](#) value in monotonically incrementing order, wrapping around from 254 (0xFE) to 1, since value 0 is reserved and using 255 (0xFF) would prevent cluster specifications from using nullable fabric-idx fields.
2. An entry within the [Fabrics attribute](#) table SHALL be added, reflecting the matter-fabric-id RDN within the NOC’s subject, along with the public key of the [trusted root](#) of the chain and the [AdminVendorID](#) field.
3. The operational key pair associated with the incoming NOC from the NOCValue, and generated by the prior CSRRequest command, SHALL be recorded for subsequent use during [CASE](#) within the fail-safe timer period (see [Section 5.5, “Commissioning Flows”](#)).
4. The incoming NOCValue and ICACValue (if present) SHALL be stored under the FabricIndex associated with the new [Fabric Scope](#), along with the RootCACertificate provided with the prior successful AddTrustedRootCertificate command invoked in the same fail-safe period.
 - a. Implementation of certificate chain storage MAY separate or otherwise encode the components of the array in implementation-specific ways, as long as they follow the correct format when being read from the [NOCs](#) list or used within other protocols such as [CASE](#).
5. The [NOCs](#) list SHALL reflect the incoming NOC from the NOCValue field and ICAC from the ICACValue field (if present).
6. The [operational discovery](#) service record SHALL immediately reflect the new Operational Identifier, such that the Node immediately begins to exist within the Fabric and becomes reachable over [CASE](#) under the new operational identity.
7. The receiver SHALL create and add a new Access Control Entry using the [CaseAdminSubject](#) field to grant subsequent Administer access to an Administrator member of the new Fabric. It is RECOMMENDED that the Administrator presented in CaseAdminSubject exist within the same entity that is currently invoking the AddNOC command, within another of the Fabrics of which it is a member.
 - a. If the Managed Device Feature is implemented by the ACL cluster, then one or more ARL entries with the new FabricIndex MAY be added to the [ARL attribute](#).

8. The incoming IPKValue SHALL be stored in the Fabric-scoped slot within the Group Key Management cluster (see [Section 11.2.7.1, “KeySetWrite Command”](#)), for subsequent use during [CASE](#).
9. The Fabric Index associated with the armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)) SHALL be updated to match the Fabric Index just allocated.
10. If the current secure session was established with [PASE](#), the receiver SHALL:

a. Augment the [secure session context](#) with the FabricIndex generated above, such that subsequent interactions have the proper [accessing fabric](#).
11. If the current secure session was established with [CASE](#), subsequent configuration of the newly installed Fabric requires the opening of a new CASE session from the Administrator from the Fabric just installed. This Administrator is the one listed in the [CaseAdminSubject](#) argument.

Thereafter, the Node SHALL respond with an NOCResponse with a [StatusCode](#) of OK and a FabricIndex field matching the FabricIndex under which the new [Node Operational Certificate \(NOC\)](#) is scoped.

11.18.6.9. UpdateNOC Command

This command is used to update an existing NOC on the device.

This command SHALL replace the NOC and optional associated ICAC (if present) scoped under the accessing fabric upon successful validation of all arguments and preconditions. The new value SHALL immediately be reflected in the [NOCs](#) list attribute.

A Commissioner or Administrator SHALL issue this command after issuing the [CSRRequest Command](#) and receiving its response.

A Commissioner or Administrator SHOULD issue this command after performing the [Attestation Procedure](#).

Access Modifier: Fabric Scoped						
ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NOCValue	octstr	max 400			M
1	ICACValue	octstr	max 400			O

Effect on Receipt

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

If a prior UpdateNOC or AddNOC command was successfully executed within the fail-safe timer period, then this command SHALL fail with a CONSTRAINT_ERROR status code sent back to the initiator.

If a prior [AddTrustedRootCertificate](#) command was successfully invoked within the fail-safe timer

period, then this command SHALL fail with a `CONSTRAINT_ERROR` status code sent back to the initiator, since the only valid following logical operation is invoking the [AddNOC](#) command.

If the prior [CSRRequest](#) state that preceded UpdateNOC had the `IsForUpdateNOC` field indicated as false, then this command SHALL fail with a `CONSTRAINT_ERROR` status code sent back to the initiator.

If any of the following conditions arise, the Node SHALL process an error by responding with an [NOCResponse](#) with a `StatusCode` of [InvalidNOC](#) as described in [Section 11.18.6.7.2, “Handling Errors”](#):

- The NOC provided in the `NOCValue` does not refer in its subject to the [FabricID](#) associated with the [accessing fabric](#).
- The ICAC provided in the `ICACValue` (if present) has a `FabricID` in its subject that does not match the [FabricID](#) associated with the [accessing fabric](#).

Otherwise, the command is considered an update of existing credentials for a given Fabric, and the following SHALL apply:

1. The Operational Certificate under the [accessing fabric index](#) in the [NOCs](#) list SHALL be updated to match the incoming `NOCValue` and `ICACValue` (if present), such that the Node's Operational Identifier within the Fabric immediately changes.
 - a. The operational key pair associated with the incoming NOC from the `NOCValue`, and generated by the prior [CSRRequest](#) command, SHALL be committed to permanent storage, for subsequent use during [CASE](#).
 - b. The [operational discovery](#) service record SHALL immediately reflect the new Operational Identifier.
 - c. All internal data reflecting the prior operational identifier of the Node within the Fabric SHALL be revoked and removed, to an outcome equivalent to the disappearance of the prior Node, except for the ongoing [CASE](#) session context, which SHALL temporarily remain valid until the [NOCResponse](#) has been successfully delivered or until the next transport-layer error, so that the response can be received by the Administrator invoking the command.

Thereafter, the Node SHALL respond with an [NOCResponse](#) with a `StatusCode` of OK and a `FabricIndex` field matching the `FabricIndex` under which the updated [NOC](#) is scoped.

11.18.6.10. NOCResponse Command

This command is used to report the results of the [AddNOC](#), [UpdateNOC](#), [UpdateFabricLabel](#) and [RemoveFabric](#) commands.

This command SHALL be generated in response to the following commands:

- [AddNOC](#)
- [UpdateNOC](#)
- [UpdateFabricLabel](#)
- [RemoveFabric](#)

It provides status information about the success or failure of those commands.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	StatusCode	NodeOperationalCert-StatusEnum				M
1	FabricIndex	fabric-idx	1 to 254			StatusCode == OK, O
2	DebugText	string	max 128			O

StatusCode Field

This field SHALL contain an [NOCStatus](#) value representing the status of an operation involving a [NOC](#).

FabricIndex Field

If present, it SHALL contain the Fabric Index of the Fabric last added, removed or updated.

DebugText Field

This field MAY contain debugging textual information from the cluster implementation, which SHOULD NOT be presented to user interfaces in any way. Its purpose is to help developers in troubleshooting errors and the contents MAY go into logs or crash reports.

11.18.6.11. UpdateFabricLabel Command

This command is used to set the user-visible fabric label for a given Fabric.

This command SHALL be used by an Administrator to set the user-visible [Label](#) field for a given Fabric, as reflected by entries in the [Fabrics](#) attribute. An Administrator SHALL use this command to set the Label to a string (possibly selected by the user themselves) that the user can recognize and relate to this Administrator

- during the commissioning process, and
- whenever the user chooses to update this string.

The Label field, along with the VendorID field in the same entry of the Fabrics attribute, SHOULD be used by Administrators to provide additional per-fabric context when operations such as Remove-Fabric are considered or used.

Access Modifier: Fabric Scoped						
ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Label	string	max 32			M

Label Field

This field SHALL contain the label to set for the fabric associated with the current secure session.

Effect on Receipt

If the Label field is identical to a Label already in use by a Fabric within the [Fabrics](#) list that is not the [accessing fabric](#), then an NOCResponse with a StatusCode of [LabelConflict](#) SHALL be returned for the command and there SHALL NOT be any permanent changes to any Fabric data.

Otherwise, the Label field for the accessing fabric SHALL immediately be updated to reflect the Label argument provided. Following the update, an NOCResponse with a StatusCode of OK SHALL be returned.

If the command was invoked within a fail-safe context after a successful [UpdateNOC](#) command, then the label update SHALL apply to the pending update state that will be reverted if fail-safe expires prior to a [CommissioningComplete](#) command. In other words, label updates apply to the state of the [Fabrics Attribute](#) as currently visible, even for an existing fabric currently in process of being updated.

11.18.6.12. RemoveFabric Command

This command is used to remove a Fabric from the device.

This command is used by Administrators to remove a given Fabric and **delete all associated fabric-scoped data**.

If the given Fabric being removed is the last one to reference a given [Trusted Root CA Certificate](#) stored in the [Trusted Root Certificates](#) list, then that Trusted Root Certificate SHALL be removed.

WARNING

This command, if referring to an already existing Fabric not under the control of the invoking Administrator, SHALL ONLY be invoked after obtaining some form of explicit user consent through some method executed by the Administrator or Commissioner. This method of obtaining consent SHOULD employ as much data as possible about the existing Fabric associations within the [Fabrics](#) list, so that likelihood is as small as possible of a user removing a Fabric unwittingly. If a method exists for an Administrator or Commissioner to convey Fabric Removal to an entity related to that Fabric, whether in-band or out-of-band, then this method SHOULD be used to notify the other Administrative Domain's party of the removal. Otherwise, users may only observe the removal of a Fabric association as persistently failing attempts to reach a Node operationally.

NOTE

If the Administrator intends to remove a fabric over a [CASE](#) session, the RevokeCommissioning command of the [AdministratorCommissioning Cluster](#) SHOULD be invoked before removal of the fabric and, if the removal is successful, also after the removal of the fabric. This serves as a security measure to prevent a malicious fabric administrator from re-adding themselves through an open commissioning window after being removed.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	FabricIndex	fabric-idx	1 to 254			M

FabricIndex Field

This field SHALL contain the Fabric Index reference (see [Section 7.19.2.23, “Fabric Index”](#)) associated with the Fabric which is to be removed from the device.

Effect on Receipt

If the FabricIndex field does not match the FabricIndex of any entry within the [Fabrics](#) list, then an NOCResponse with a StatusCode of [InvalidFabricIndex](#) SHALL be returned for the command and there SHALL NOT be any permanent changes to any device data.

Otherwise, one of the following outcomes SHALL occur:

1. If the FabricIndex matches the last remaining entry in the [Fabrics](#) list, then the device SHALL delete all Matter related data on the node which was created since it was commissioned. This includes all Fabric-Scoped data, including Access Control List, Access Restriction List, bindings, scenes, group keys, operational certificates, etc. All Trusted Roots SHALL also be removed. If a time synchronization cluster is present on the Node, the TrustedTimeSource and DefaultNtp SHALL be set to null. Any Matter related data including logs, secure sessions, exchanges and interaction model constructs SHALL also be removed. Since this operation involves the removal of the secure session data that may underpin the current set of exchanges, the Node invoking the command SHOULD NOT expect a response before terminating its secure session with the target.
2. If the FabricIndex does not equal the [accessing fabric index](#), then the device SHALL begin the process of irrevocably deleting all associated Fabric-Scoped data, including Access Control Entries, Access Restriction Entries, bindings, group keys, operational certificates, etc. Any remaining Trusted Roots no longer referenced by any operational certificate SHALL also be removed. If a time synchronization cluster is present on the Node, and the TrustedTimeSource FabricIndex matches the given FabricIndex, the TrustedTimeSource SHALL be set to null. All secure sessions, exchanges and interaction model constructs related to the Operational Identity under the given Fabric SHALL also be removed. Following the removal, an NOCResponse with a StatusCode of OK SHALL be returned.
3. If the FabricIndex equals the [accessing fabric index](#), then the device SHALL begin the process of irrevocably deleting all associated Fabric-Scoped data, including Access Control Entries, Access Restriction Entries, bindings, group keys, operational certificates, etc. Any remaining Trusted Roots no longer referenced by any operational certificate SHALL also be removed. If a time synchronization cluster is present on the Node, and the TrustedTimeSource FabricIndex matches the given FabricIndex, the TrustedTimeSource SHALL be set to null. All secure sessions, exchanges and interaction model constructs related to the Operational Identity under the given Fabric SHALL also be removed. Since this operation involves the removal of the secure session data that may underpin the current set of exchanges, the Node invoking the command SHOULD NOT expect a response before terminating its secure session with the target.

11.18.6.13. AddTrustedRootCertificate Command

This command is used to add a trusted root certificate to the device.

This command SHALL add a [Trusted Root CA Certificate](#), provided as its [Matter Certificate Encoding](#) representation, to the [TrustedRootCertificates Attribute](#) list and SHALL ensure the next [AddNOC](#) command executed uses the provided certificate as its root of trust.

If the certificate from the RootCACertificate field is already installed, based on exact byte-for-byte equality, then this command SHALL succeed with no change to the list.

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

If a prior AddTrustedRootCertificate command was successfully invoked within the fail-safe timer period, which would cause the new invocation to add a second root certificate within a given fail-safe timer period, then this command SHALL fail with a CONSTRAINT_ERROR status code sent back to the initiator.

If a prior UpdateNOC or AddNOC command was successfully executed within the fail-safe timer period, then this command SHALL fail with a CONSTRAINT_ERROR status code sent back to the initiator.

If the certificate from the RootCACertificate field fails any validity checks, not fulfilling all the requirements for a valid [Matter Certificate Encoding](#) representation, including a truncated or over-size value, then this command SHALL fail with an INVALID_COMMAND status code sent back to the initiator.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	RootCACertificate	octstr	max 400			M

Note that the only method of removing a trusted root is by removing the Fabric that uses it as its root of trust using the [RemoveFabric](#) command.

11.18.6.14. SetVIDVerificationStatement Command

This command is used to manage the VendorID and VIDVerificationStatement fields of the Fabrics attribute, and the VVSC field of an entry in the NOCs attribute.

This command SHALL be used to one or more of the following:

- Update the VendorID associated with an entry in the Fabrics attribute.
- Associate or remove a [VIDVerificationStatement](#) associated with an entry in the Fabrics attribute.
- Associate or remove a [VendorVerificationSigningCertificate \(VVSC\)](#) associated with an entry in the NOCs attribute.

This command SHALL only operate against the Fabrics and NOCs attribute entries associated with the [accessing fabric index](#).

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	VendorID	vendor-id				O.a+
1	VIDVerificationStatement	octstr	max 85			O.a+
2	VVSC	octstr	max 400			O.a+

Effect on Receipt

If the **VendorID** field is present, the value of the **VendorID** in the [Fabrics](#) attribute entry associated with the [accessing fabric index](#) SHALL have its value replaced with the value from the command field.

If the **VVSC** field is present, but the **ICAC** field is already present in the [NOCs](#) attribute entry associated with the [accessing fabric index](#), then the command SHALL fail with a status code of INVALID_COMMAND.

If the **VIDVerificationStatement** field is present:

- If the length of the field's value is neither exactly 0 nor exactly 85, then the command SHALL fail with a status code of CONSTRAINT_ERROR.
- If the length of the field's value is exactly 0, then the **VIDVerificationStatement** field in the [Fabrics](#) attribute entry associated with the [accessing fabric index](#) SHALL be erased and the field SHALL disappear from the Fabrics entry.
- If the length of the field's value is exactly 85, then the **VIDVerificationStatement** field in the [Fabrics](#) attribute entry associated with the [accessing fabric index](#) SHALL have its value replaced with the value from the command field.

If the **VVSC** field is present:

- If the length of the field's value is exactly 0, then the **VVSC** field in the [NOCs](#) attribute entry associated with the [accessing fabric index](#) SHALL be erased and the field SHALL disappear from the NOCs entry.
- If the length of the field's value is not 0, then the **VVSC** field in the [NOCs](#) attribute entry associated with the [accessing fabric index](#) SHALL have its value replaced with the value from the command field. The contents of the certificate need not be validated by the server. Clients SHALL validate the contents at time of use.

If the command was invoked within a fail-safe context after a successful [AddNOC](#) or [UpdateNOC](#) command, then the field updates SHALL apply to the pending update state that will be reverted if fail-safe expires prior to a [CommissioningComplete](#) command. In other words, field updates apply to the state of the [Fabrics Attribute](#) as currently visible, even for an existing fabric currently in process of being updated.

11.18.6.15. SignVIDVerificationRequest Command

This command is used to authenticate the fabric associated with the FabricIndex.

This command SHALL be used to request that the server authenticate the fabric associated with the FabricIndex given by generating the response described in [Section 6.4.10, “Fabric Table Vendor ID Verification Procedure”](#).

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	FabricIndex	fabric-idx	1 to 254			M
1	ClientChallenge	octstr	32			M

The FabricIndex field SHALL contain the fabric index being targeted by the request.

The ClientChallenge field SHALL contain a client-provided random challenge to be used during the signature procedure.

Effect on Receipt

If the FabricIndex field contains a fabric index which does not have an associated entry in the [Fabrics](#) attribute, then the command SHALL fail with a status code of CONSTRAINT_ERROR.

Otherwise, if no other errors have occurred, the command SHALL generate a [SignVIDVerificationResponse](#).

11.18.6.16. SignVIDVerificationResponse Command

This command is used to report the results of the SignVIDVerificationRequest command. This command SHALL contain the response of the [SignVIDVerificationRequest](#), computed as described below.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	FabricIndex	fabric-idx	1 to 254			M
1	FabricBindingVersion	uint8	1 to 255			M
2	Signature	octstr	min 1			M

The FabricIndex field SHALL contain the same value of **FabricIndex** as the value from the associated [SignVIDVerificationRequest](#).

The FabricBindingVersion field SHALL contain value 0x01 for version 1.0 of the Matter [Cryptographic Primitives](#).

The Signature field SHALL contain the octet string result of **Crypto_Sign(noc_private_key, vendor_id_verification_tbs)**:

- **noc_private_key** is the operational private key associated with the **Node Operational Key Pair** for the FabricIndex requested in the associated **SignVIDVerificationRequest**.
- **vendor_id_verification_tbs** := **fabric_binding_version** || **client_challenge** || **attestation_challenge** || **fabric_index** || **vendor_fabric_binding_message** || **<vid_verification_statement>**
 - **fabric_binding_version** is the value from the FabricBindingVersion field of this **SignVIDVerificationResponse**.
 - **client_challenge** is the 32-octet **ClientChallenge** from the **SignVIDVerificationRequest**.
 - **attestation_challenge** is the AttestationChallenge from a **CASE session**, **resumed CASE session**, or **PASE session** depending on the method used to establish the current secure session context over which the response will be sent.
 - **fabric_index** is the 1-octet value of **FabricIndex** from the **SignVIDVerificationRequest**.
 - **vendor_fabric_binding_message** is the octet string of the vendor_fabric_binding_message defined in **Section 6.4.10.1, “Algorithm”**.
 - **vid_verification_statement** is the 85-octet (for cryptographic primitives mapping 1.0) value from the **VIDVerificationStatement** field of the entry in the **Fabrics** attribute associated with the **fabric_index**, if present. If there is no such field in the **Fabrics** attribute for the fabric_index specified, this field SHALL be omitted from the **vendor_id_verification_tbs** message.

11.19. Administrator Commissioning Cluster

This cluster is used to trigger a Node to allow a new Administrator to commission it. It defines Attributes, Commands and Responses needed for this purpose.

There are two methods of commissioning, Basic Commissioning which MAY be supported and is described in **Section 5.6.2, “Basic Commissioning Method (BCM)”** and Enhanced Commissioning which SHALL be supported and is described in **Section 5.6.3, “Enhanced Commissioning Method (ECM)”**.

For the management of Operational Credentials and Trusted Root Certificates, the **Operational Credentials cluster** is used.

If the Administrator Commissioning Cluster server instance is present on an endpoint with the Root Node device type in the Descriptor cluster DeviceTypeList, then:

- The Commissioning Window SHALL be opened or closed on the node that the Root Node endpoint is on.
- The attributes SHALL indicate the state of the node that the Root Node endpoint is on.

If the Administrator Commissioning Cluster server instance is present on an endpoint with the Bridged Node device type in the Descriptor cluster DeviceTypeList, then:

- The Commissioning Window SHALL be opened or closed on the node represented by the Bridged Node.
- The attributes SHALL indicate the state of the node that is represented by the Bridged Node.

11.19.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.19.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	CADMIN

11.19.3. Cluster ID

ID	Name
0x003C	Administrator Commissioning

11.19.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Summary
0	BC	Basic	Node supports Basic Commissioning Method.

11.19.5. Data Types

11.19.5.1. CommissioningWindowStatusEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	WindowNotOpen	Commissioning window not open	M
1	EnhancedWindowOpen	An Enhanced Commissioning Method window is open	M
2	BasicWindowOpen	A Basic Commissioning Method window is open	BC

11.19.6. Status Codes

11.19.6.1. StatusCodeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0x02	Busy	Could not be completed because another commissioning is in progress	M
0x03	PAKEParameterError	Provided PAKE parameters were incorrectly formatted or otherwise invalid	M
0x04	WindowNotOpen	No commissioning window was currently open	M

11.19.7. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Window-Status	Commissioning-Window-StatusEnum				R V	M
0x0001	Admin-FabricIndex	fabric-idx		X		R V	M
0x0002	Admin-VendorId	vendor-id		X		R V	M

11.19.7.1. WindowStatus Attribute

This attribute SHALL indicate whether a new Commissioning window has been opened by an Administrator, using either the [OpenCommissioningWindow](#) command or the [OpenBasicCommissioningWindow](#) command.

This attribute SHALL revert to WindowNotOpen upon expiry of a commissioning window.

NOTE

An initial commissioning window is not opened using either the [OpenCommissioningWindow](#) command or the [OpenBasicCommissioningWindow](#) command, and therefore this attribute SHALL be set to WindowNotOpen on initial commissioning.

11.19.7.2. AdminFabricIndex Attribute

When the [WindowStatus](#) attribute is not set to WindowNotOpen, this attribute SHALL indicate the FabricIndex associated with the Fabric scoping of the Administrator that opened the window. This MAY be used to cross-reference in the [Fabrics](#) attribute of the [Operational Credentials](#) cluster.

If, during an open commissioning window, the fabric for the Administrator that opened the window is removed, then this attribute SHALL be set to null.

When the [WindowStatus](#) attribute is set to WindowNotOpen, this attribute SHALL be set to null.

11.19.7.3. AdminVendorId Attribute

When the [WindowStatus](#) attribute is not set to WindowNotOpen, this attribute SHALL indicate the [Vendor ID](#) associated with the Fabric scoping of the Administrator that opened the window. This field SHALL match the [VendorID](#) field of the [Fabrics](#) attribute list entry associated with the Administrator having opened the window, at the time of window opening. If the fabric for the Administrator that opened the window is removed from the node while the commissioning window is still open, this attribute SHALL NOT be updated.

When the [WindowStatus](#) attribute is set to WindowNotOpen, this attribute SHALL be set to null.

11.19.8. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	OpenCommissioningWindow	client ⇒ server	Y	A T	M
0x01	OpenBasicCommissioningWindow	client ⇒ server	Y	A T	BC
0x02	RevokeCommissioning	client ⇒ server	Y	A T	M

Only one commissioning window can be active at a time. If a Node receives another open commissioning command when an Open Commissioning Window is already active, it SHALL return a failure response (see [Section 11.19.6, “Status Codes”](#)).

11.19.8.1. OpenCommissioningWindow Command

This command is used by a current Administrator to instruct a Node to go into commissioning mode. The [Enhanced Commissioning Method](#) specifies a window of time during which an already commissioned Node accepts [PASE](#) sessions. The current Administrator SHALL specify a timeout value for the duration of the OpenCommissioningWindow command.

When the OpenCommissioningWindow command expires or commissioning completes, the Node SHALL remove the Passcode by deleting the [PAKE](#) passcode verifier as well as stop publishing the DNS-SD record corresponding to this command as described in [Section 4.3.1, “Commissionable](#)

Node Discovery". The commissioning into a new Fabric completes when the Node successfully receives a [Section 11.10.7.6, "CommissioningComplete"](#) command, see [Section 5.5, "Commissioning Flows"](#).

The parameters for OpenCommissioningWindow command are as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Commissioning-Timeout	uint16	desc			M
1	PAKEPass-codeVerifier	octstr	97			M
2	Discriminator	uint16	0 to 4095			M
3	Iterations	uint32	1000 to 100000			M
4	Salt	octstr	16 to 32			M

A current Administrator MAY invoke this command to put a node in commissioning mode for the next Administrator. On completion, the command SHALL return a cluster specific status code from the [Section 11.19.6, "Status Codes"](#) below reflecting success or reasons for failure of the operation. The new Administrator SHALL discover the Node on the IP network using DNS-based Service Discovery (DNS-SD) for commissioning.

If any format or validity errors related to the PAKEPasscodeVerifier, Iterations or Salt arguments arise, this command SHALL fail with a cluster specific status code of PAKEParameterError.

If a commissioning window is already currently open, this command SHALL fail with a cluster specific status code of Busy.

If the [fail-safe timer](#) is currently armed, this command SHALL fail with a cluster specific status code of Busy, since it is likely that concurrent commissioning operations from multiple separate Commissioners are about to take place.

In case of any other parameter error, this command SHALL fail with a status code of COMMAND_INVALID.

CommissioningTimeout Field

This field SHALL specify the time in seconds during which commissioning session establishment is allowed by the Node. This timeout value SHALL follow guidance as specified in the initial [Section 5.4.2.3, "Announcement Duration"](#). The CommissioningTimeout applies only to cessation of any announcements and to accepting of new commissioning sessions; it does not apply to abortion of connections, i.e., a commissioning session SHOULD NOT abort prematurely upon expiration of this timeout.

PAKEPasscodeVerifier Field

This field SHALL specify an ephemeral PAKE passcode verifier (see [Section 3.10, “Password-Authenticated Key Exchange \(PAKE\)”](#)) computed by the existing Administrator to be used for this commissioning. The field is concatenation of two values (w0 || L) SHALL be (CRYPTO_GROUP_SIZE_BYTES + CRYPTO_PUBLIC_KEY_SIZE_BYTES)-octets long as detailed in [Crypto_PAKEValues_Responder](#). It SHALL be derived from an ephemeral passcode (See [PAKE](#)). It SHALL be deleted by the Node at the end of commissioning or expiration of the OpenCommissioningWindow command, and SHALL be deleted by the existing Administrator after sending it to the Node(s).

Discriminator Field

This field SHALL be used by the Node as the long discriminator for DNS-SD advertisement (see [Section 4.3.1.5, “TXT key for discriminator \(D\)”](#)) for discovery by the new Administrator. The new Administrator can find and filter DNS-SD records by long discriminator to locate and initiate commissioning with the appropriate Node.

Iterations Field

This field SHALL be used by the Node as the PAKE iteration count associated with the ephemeral PAKE passcode verifier to be used for this commissioning, which SHALL be sent by the Node to the new Administrator’s software as response to the [Section 4.14.1.2.3, “PBKDFParamRequest”](#) during PASE negotiation. The permitted range of values SHALL match the range specified in [Section 3.9, “Password-Based Key Derivation Function \(PBKDF\)”](#), within the definition of the Crypto_PBKDFParameterSet.

Salt Field

This field SHALL be used by the Node as the PAKE Salt associated with the ephemeral PAKE passcode verifier to be used for this commissioning, which SHALL be sent by the Node to the new Administrator’s software as response to the [Section 4.14.1.2.3, “PBKDFParamRequest”](#) during PASE negotiation. The constraints on the value SHALL match those specified in [Section 3.9, “Password-Based Key Derivation Function \(PBKDF\)”](#), within the definition of the Crypto_PBKDFParameterSet.

When a Node receives the [Section 11.19.8.1, “OpenCommissioningWindow”](#) command, it SHALL begin advertising on DNS-SD as described in [Section 4.3.1, “Commissionable Node Discovery”](#) and for a time period as described in [Section 11.19.8.1.1, “CommissioningTimeout”](#).

When the command is received by a ICD, it SHALL enter into [active mode](#). The ICD SHALL remain in Active Mode as long as one of these conditions is met:

- A commissioning window is open.
- There is an armed [fail-safe timer](#).

11.19.8.2. OpenBasicCommissioningWindow Command

This command MAY be used by a current Administrator to instruct a Node to go into commissioning mode, if the node supports the [Basic Commissioning Method](#). The [Basic Commissioning Method](#) specifies a window of time during which an already commissioned Node accepts PASE sessions. The current Administrator SHALL specify a timeout value for the duration of the OpenBasicCommis-

sioningWindow command.

If a commissioning window is already currently open, this command SHALL fail with a cluster specific status code of Busy.

If the [fail-safe timer](#) is currently armed, this command SHALL fail with a cluster specific status code of Busy, since it is likely that concurrent commissioning operations from multiple separate Commissioners are about to take place.

In case of any other parameter error, this command SHALL fail with a status code of COMMAND_INVALID.

The commissioning into a new Fabric completes when the Node successfully receives a [Section 11.10.7.6, “CommissioningComplete”](#) command, see [Section 5.5, “Commissioning Flows”](#). The new Administrator SHALL discover the Node on the IP network using DNS-based Service Discovery (DNS-SD) for commissioning.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Commissioning-Timeout	uint16	desc			M

CommissioningTimeout Field

This field SHALL specify the time in seconds during which commissioning session establishment is allowed by the Node. This timeout SHALL follow guidance as specified in the initial [Section 5.4.2.3, “Announcement Duration”](#).

When a Node receives the [Section 11.19.8.2, “OpenBasicCommissioningWindow”](#) command, it SHALL begin advertising on DNS-SD as described in [Section 4.3.1, “Commissionable Node Discovery”](#) and for a time period as described in [Section 11.19.8.2.1, “CommissioningTimeout”](#). When the command is received by a [ICD](#), it SHALL enter into [active mode](#). The [ICD](#) SHALL remain in Active Mode as long as one of these conditions is met:

- A commissioning window is open.
- There is an armed [fail-safe timer](#).

11.19.8.3. RevokeCommissioning Command

This command is used by a current Administrator to instruct a Node to revoke any active [Section 11.19.8.1, “OpenCommissioningWindow”](#) or [Section 11.19.8.2, “OpenBasicCommissioningWindow”](#) command. This is an idempotent command.

Effect on Receipt

1. The Node SHALL perform the following actions regardless of current commissioning window state:

- a. The Node SHALL (for ECM) delete the temporary PAKEPasscodeVerifier and associated data
 - b. The Node SHALL terminate any open PASE sessions or PASE sessions in the process of being established
 - c. The Node SHALL immediately expire any fail-safe held by an open PASE session and perform the cleanup steps outlined in [Section 11.10.7.2.2, “Behavior on expiry of Fail-Safe timer”](#)
2. If the commissioning window was NOT open at the time of receipt, the Node SHALL return a cluster specific status code of WindowNotOpen.
 3. If the commissioning window was open at the time of receipt, the Node SHALL perform the following additional actions:
 - a. The Node SHALL stop accepting new incoming PASE session establishment messages
 - b. The Node SHALL stop publishing the DNS-SD records associated with the advertising it was doing due to the open commissioning window. See [Section 4.3.1, “Commissionable Node Discovery”](#)
 - c. The Node SHALL expire the commissioning window and set the WindowStatus attribute to WindowNotOpen

NOTE

If this command is issued over the PASE connection, the caller SHALL NOT expect to receive a response due to the termination of the PASE session during command execution.

11.20. Over-the-Air (OTA) Software Update

11.20.1. Scope & Purpose

The majority of IoT devices require security and/or functional feature updates during their lifetime.

This section describes a set of OTA software update capabilities which enable an "OTA Requestor" to be informed of, obtain, and install software updates from a Node fulfilling the role of an "OTA Provider".

An "OTA Requestor" is any Node implementing the OTA Requestor Device Type (0x0012), which fulfills the client role for the [OTA Software Update Provider cluster](#) and the server role for the [OTA Software Update Requestor cluster](#). An "OTA Provider" is any Node implementing the OTA Provider Device Type (0x0014), which fulfills the server role for the [OTA Software Update Provider cluster](#) and the client role for the [OTA Software Update Requestor cluster](#).

The OTA updates capabilities are designed to support:

- A mechanism to inform OTA Requestors about available OTA Providers.
- A mechanism to allow OTA Requestors to acquire information about available OTA Software Images.
- A mechanism to allow constrained OTA Requestors to obtain OTA Software Images through a local proxy, e.g. if they are not able or willing to proceed with a direct download from the Inter-

net.

- A mechanism to allow OTA Requestors supporting legacy, non-native, or out-of-band update methods to notify an OTA Provider of having completed an update out-of-band.
- A mechanism to allow deferred installation of a software update, based on administrative rules.
- A mechanism to allow user consent to be considered before offering Software Images to OTA Requestors.

OTA Requestors wishing to update their software using these capabilities MAY need to use an application bootloader and MAY require sufficient additional storage in order to download an OTA image.

Furthermore, to encourage interoperability and timely software updates, the OTA update mechanisms provide means of obtaining Software Images which can be uniformly implemented across OTA Requestors on devices from a variety of different vendors. The OTA Providers SHOULD provide services to OTA Requestors from vendors other than its own, as long as the location of Software Update images for these vendors is found. The [Distributed Compliance Ledger](#) is one such centralized source of software update image locations that MAY allow OTA Providers to provide OTA Software Update Images generically to devices from many vendors.

11.20.2. Functional overview

An OTA Requestor SHALL query the OTA Provider periodically to determine the availability of new Software Images. The OTA Provider MAY learn, from backend systems inside or outside of Fabric scope, of the availability of a new Software Image for an OTA Requestor.

An OTA Requestor which has been updated using a mechanism beyond this Cluster MAY report to an OTA Provider that a Software Image update has been completed.

The OTA Provider MAY announce its presence to OTA Requestors on the Fabric to assist in discovery of this service (see [Section 11.20.7.6.1, “AnnounceOTAProvider Command”](#)).

Nodes SHALL NOT rely solely on unsolicited OTA Provider announcements to discover available OTA Providers and SHALL instead employ other means such as using OTA Provider records provisioned during Commissioning, or dynamic discovery of OTA Providers.

OTA Requestors SHALL only upgrade to numerically newer versions and OTA Providers SHALL only ever provide a numerically newer Software Image than a Node's current version number (see [Section 11.1.5.10, “SoftwareVersion”](#)). Any functional rollback SHALL be affected by the Vendor creating a Software Image with a newer (higher) version number, but whose binary contents may match prior functionality.

All OTA Requestors SHALL support usage of a polling mechanism to send a query command to the OTA Provider in order to determine if the OTA Provider has any new Software Images for it. Polling simplifies processing for OTA Requestors that MAY need to perform special setup to get ready to receive a Software Images, such as unlocking flash or allocating space for a new Software Images.

It is ideal to have OTA Providers maintain as little state as possible since this will scale better when there are hundreds of OTA Requestors in a given Fabric. The OTA Provider is not required to keep

track of what pieces of an image that a particular OTA Requestor has received.

The flow for querying the availability of a new version is done using commands of the OTA Provider Cluster. In case of a new image available matching an OTA Requestor's request, the response to the [QueryImage](#) command SHALL contain a URI where the given image can be downloaded.

The download protocol is separate from the Cluster commands. All OTA Providers SHALL support the [BDX Protocol](#) to allow for the downloading of OTA images by both sleepy End Devices and more capable devices, without requiring access to the public Internet from the OTA Requestor. OTA Requestors SHOULD support the [BDX Protocol](#).

In order to maximize the interoperable combinations of deployed products and Fabric Administrators, the Connectivity Standards Alliance's Distributed Compliance Ledger (DCL) MAY contain sufficient OTA Software Update information to cover a large number of products, using a federated mechanism of data maintenance. See [Section 11.23, "Distributed Compliance Ledger"](#) for details on the Distributed Compliance Ledger common data schemas. See [Section 11.20.3.3.2, "Conceptual algorithm for matching OTA Software Images applicable to a query"](#) for the conceptual algorithm recommended for implementation by OTA Providers to match records available in the DCL to incoming queries.

11.20.3. Software update workflow

The software update workflow consists of several steps executed in a sequence from an OTA Requestor towards an OTA Provider. When a newer Software Image for an OTA Requestor is available on the OTA Provider this results in an updated Software Image being acquired and applied by said OTA Requestor.

The steps, in order, and assuming each step successfully leads to the next, are the following, with each numbered according to [Figure 98, "Detailed OTA Software Update Workflow"](#):

- [10] OTA Provider optionally announces its presence to nodes (see [Section 11.20.7.6.1, "AnnounceOTAProvider Command"](#)). This MAY be used in addition to other OTA Provider discovery methods.
- [11] OTA Requestor determines OTA Provider to query.
- [11] OTA Requestor queries the OTA Provider for availability of an updated Software Image version.
- [30..34] OTA Provider obtains consent from user to apply the OTA update.
- [40..41] OTA Provider obtains a copy of the new Software Image, either in real time or in a time-deferred manner, to provide to the OTA Requestor over [BDX Protocol](#), or over an alternate supported protocol that both OTA Provider and OTA Requestor support. If the Software Image happens to be already available in the OTA Provider's [cache](#), this step can be skipped.
- [52] OTA Requestor downloads the update, either over [BDX Protocol](#) from OTA Provider acting as a proxy, or over an alternate protocol that both OTA Provider and OTA Requestor support.
- [60] OTA Requestor notifies the OTA Provider that the download is complete and that it is ready to apply the downloaded image.

- [61] OTA Provider responds with an authorization to apply the update, after an optional delay.
- [62] OTA Requestor applies the update and starts executing updated software.
- [63] OTA Requestor notifies the OTA Provider of having successfully applied the update.

In order to illustrate more specifically these steps, [Figure 98, “Detailed OTA Software Update Workflow”](#) below depicts a detailed sequence showing the following illustrative (not normative) aspects:

- [10..21] Determining the availability of an OTA software update.
- [22] Deferral of download by OTA Provider responding with a "Busy" condition, while it obtains user consent and obtains the Software Image from a vendor server based on information in the OTA Provider's OTA Software Update logic.
- Obtaining user consent in one of these ways:
 - [30..31] Out-of-band notification through some externally-provided user interface, such as a mobile device terminal operated by an authorized user, and connected to OTA Provider's logic in some way.
 - [32..34] Reuse of prior user consent, perhaps from a continued but revokable authorization, sent back to the OTA Provider by OTA Software Update logic.
 - Via the OTA Provider delegating to the OTA Requestor Node (see [Section 11.20.3.4.1, “User consent delegation to Nodes”](#)). Note that this case is not illustrated in the sequence diagram.
- [40..41] Downloading and temporarily storing a Software Image by the OTA Provider, from a Vendor's server, over the public Internet, for the purposes of eventual proxied local download by the OTA Requestor.
- [50..51] Responding positively to a subsequent query by the OTA Requestor, since an OTA software update is now definitely available.
- [52] Downloading of the Software Image from the OTA Provider by the OTA Requestor, using the [BDX Protocol](#) against the temporary storage of the OTA Provider.
- [60..63] OTA Requestor performs the update (after permission from OTA Provider)

Detailed OTA Software Update Workflow




Figure 98. Detailed OTA Software Update Workflow

Given that some of the above steps MAY fail to complete, and that some MAY provide a variety of outcomes or replies, the following subsections give the necessary normative details describing the sequence.

11.20.3.1. Determining the OTA Provider to query

The discovery of available OTA Providers is necessary for OTA Requestors to be able to query for new software images. Each OTA Requestor SHALL keep a list of OTA Provider Operational Identifiers (see [Section 2.5.5, “Node Identifier”](#)) that it could query.

A given OTA Requestor SHALL have sufficient storage to maintain one OTA Provider entry per Fabric within the [DefaultOTAProviders](#) default list. This default OTA Provider list MAY be augmented by any means deemed acceptable by a given OTA Requestor, such that the internal list of possible locations to query contains at least the DefaultOTAProviders, but it MAY contain more. For example, it may contain cached locations that arose from the [AnnounceOTAProvider](#) command.

When an OTA Requestor determines that it is time to query an OTA Provider, it MAY use any method of its choosing to determine which OTA Provider to contact for its next query.

An OTA Requestor MAY expunge OTA Providers from its OTA Provider list if it determines that the entry is stale or obsolete.

Discovery of additional OTA Providers MAY be done by handling the arrival of [AnnounceOTAProvider](#) commands invoked by OTA Providers.

Commissioners SHOULD add an entry to the [DefaultOTAProviders](#) list attribute, if an OTA Provider is known at commissioning time, to reduce the delay between commissioning and first [QueryImage](#) command.

Whenever communicating with an OTA Provider location obtained either through the DefaultOTAProviders attribute, or the AnnounceOTAProvider command, an OTA Requestor SHALL target all interactions with that Node by interacting with the given Endpoint on the given ProviderNodeID obtained from these sources.

Discovery of additional OTA Providers MAY be done using runtime service discovery, which is outside the scope of this specification.

Nodes MAY attempt to contact OTA Providers that are known to them in any order if failing to reach a default OTA Provider from an entry in the [DefaultOTAProviders](#) list. This approach would assist in maximizing likelihood of eventual success.

11.20.3.2. Querying the OTA Provider

Query of the OTA Provider SHALL be done using the [QueryImage](#) command. The arguments for this command provide sufficient information to allow the OTA Provider to determine the availability of a new image for the querying OTA Requestor.

An OTA Requestor SHALL NOT query more frequently than once every 120 seconds, unless a Node loses its timekeeping state, due to events such as power loss or restart, that prevent applying such a delay. This reduces the burden on both the OTA Providers providing the service to a large number of nodes and the supporting networking infrastructure. It is recommended for OTA Requestors to attempt a daily [QueryImage](#) command, if capable, to ensure timely access to updated software, including security-critical updates.

The OTA Provider SHALL use an algorithm deemed satisfactory by its implementer to determine

the availability of a newer Software Image in response to a [QueryImage](#) command. This algorithm will be called the "OTA Image Selection Logic" thereafter.

The OTA Image Selection Logic MAY use any data it deems useful, either local to the equipment or Node hosting the OTA Provider, or remote through external networks, to determine whether an updated Software Image is available (see [Section 11.20.3.3, "Availability of Software Images"](#)).

OTA Provider requests are idempotent. In cases where an OTA Requestor is repeating a request it has already done, and the OTA Provider can detect this, it SHALL NOT behave differently on any subsequent attempt compared to the first, unless a new Software Image has become available in the meantime. That is, an OTA Provider SHALL NOT prevent an OTA Requestor from trying to make the same query more than once. This requirement is critical to ensure that OTA Requestors which encounter error conditions during OTA Image Query or OTA Image Transfer can eventually succeed through retrying the same operation more than once.

Upon final determination of the outcome of the [QueryImage](#) command, the OTA Provider SHALL reply with a [QueryImageResponse](#) command.

If an OTA Provider employs synchronous proxying (e.g. proxy-while-downloading) method to reach off-Fabric Software Images and provide them over [BDX Protocol](#) to OTA Requestors, it SHALL respond with a Status of DownloadProtocolNotSupported to an OTA Requestor in the [QueryImageResponse](#) command if all the following conditions apply:

1. A new Software Image is determined to be available.
2. The Software Image to proxy is served by a remote server that does NOT support range-based transfers.
3. The OTA Requestor only supports BDX.
4. The OTA Provider does not support asynchronous proxying (e.g. download-then-proxy).

The fields of the [QueryImageResponse](#) command convey the next steps to take. The primary indication of action to be taken by the OTA Requestor is expressed in the Status field, with the other fields providing the necessary details as described in the following subsections.

Failure to receive an application-layer response from the OTA Provider after invoking the [QueryImage](#) command SHOULD be considered equivalent to having received a [QueryImageResponse](#) command with a Status field containing NotAvailable (see [Section 11.20.3.2.4, "Handling NotAvailable value in Status field"](#)).

Access Control Requirements

Commissioners or Administrators SHOULD install necessary [ACL entries](#) at Commissioning time or later to enable the handling of the [AnnounceOTAProvider](#) commands by OTA Requestors.

Below is an exemplary ACL entry for a Node implementing the OTA Requestor cluster server to support the processing of the [AnnounceOTAProvider](#) command:

```
{  
  FabricIndex: <Fabric Index of the fabric in question>,  
  ...  
}
```

```

Privilege: Operate,
AuthMode: CASE,
Subjects: [ <Node ID of the node implementing OTA Requestor cluster client> ],
Targets: [ <Endpoint hosting the OTA Requestor cluster server> ]
}

```

Commissioner or Administrator SHOULD install necessary ACL Entries at Commissioning time or later to enable processing of [QueryImage commands](#) from OTA Requestors on their Fabric, otherwise that OTA Provider will not be usable by OTA Requestors.

Below is an exemplary ACL entry for a Node implementing the OTA Provider cluster server to support the processing of the QueryImage command:

```

{
  FabricIndex: <Fabric Index of the fabric in question>,
  Privilege: Operate,
  AuthMode: CASE,
  Subjects: [ ], // Empty for "any" Node wildcard
  Targets: [ <Endpoint hosting the OTA Provider cluster server> ]
}

```

Note that there may be a variety of ACL entry configurations that fulfill the necessary goals, including wildcard entries for the Administrators on a given Fabric. The examples above are for illustration purposes only.

Handling UpdateAvailable value in Status field

The UpdateAvailable status indicates that the OTA Provider has an update available.

The remaining fields within the [QueryImageResponse](#) command SHALL contain the information necessary to allow the OTA Requestor to obtain an updated Software Image.

The field [ImageURI](#) SHALL be set to a location from where the image can be downloaded. The URI provided SHALL be for a protocol within the list of supported protocols provided in the request (see [Section 11.20.6.5.1.4, “ProtocolsSupported Field”](#)). Selection of the URI is based on the information available in the OTA Provider’s Software Images data set.

The field [UpdateToken](#) SHALL be populated by the OTA Provider with a value of its choosing, to allow tracking of the flow from a given OTA Requestor when it sends further requests. The valid length of the UpdateToken is between 8 and 32 bytes, inclusively. The token SHALL be recorded by the OTA Requestor, until an OTA Software Update Image is either applied or discarded. This value SHALL be provided to any subsequent ApplyUpdateRequest and NotifyUpdateApplied commands.

The field [SoftwareVersion](#) SHALL be set to the version number matching the new Software Image.

If the field [DelayedActionTime](#) is present, the OTA Requestor SHALL wait at least the indicated time before proceeding with the next step.

Handling of SoftwareVersion and ImageURI fields SHALL follow these rules:

- If the SoftwareVersion field matches the version indicated in the header of a previously downloaded OTA Software Image, one of two cases applies:
 1. Image was fully downloaded and verified: the OTA Requestor SHOULD skip the transfer step (see [Section 11.20.3.5, “Transfer of OTA Software Update images”](#)), and move directly to the apply step (see [Section 11.20.3.6, “Applying a software update”](#)).
 2. Image was partially downloaded: the OTA Requestor SHOULD attempt to continue the transfer from where it left off, if it is capable, otherwise it SHALL start the transfer anew. See [Section 11.20.3.5, “Transfer of OTA Software Update images”](#) for a description of complete and restarted downloads.
- If the SoftwareVersion field indicates a newer (numerically higher) version than the version currently installed on the OTA Requestor, the OTA Requestor SHOULD proceed with OTA Image Transfer (see [Section 11.20.3.5, “Transfer of OTA Software Update images”](#)), after awaiting at least the delay stated by the [DelayedActionTime](#) field, if present.
- If the SoftwareVersion field indicates the same or an older (numerically lower) version, or if the ImageURI field somehow contains information which cannot be used by the OTA Requestor, then the OTA Requestor SHALL go back to awaiting its next OTA Software Update query opportunity, following the rules previously stated in [Section 11.20.3.2, “Querying the OTA Provider”](#). In that case, the OTA Requestor MAY attempt to select a different OTA Provider according to [Section 11.20.3.1, “Determining the OTA Provider to query”](#), which MAY cause the OTA Requestor to immediately try another query, but to a different OTA Provider, thus not violating daily allowance of a given OTA Requestor towards a given OTA Provider.

Handling Busy value in Status field

The Busy status indicates that the OTA Provider is busy for any reason and that it can only provide a definite answer at a later time. This MAY be because the OTA Provider is currently determining whether an update is available for the OTA Requestor that made the query. An OTA Requestor SHOULD attempt to query the same OTA Provider again later at least once more if a Busy response is obtained, rather than querying a different OTA Provider in its list, so that the OTA Provider that replied Busy could have had resources available to determine availability.

After a Busy status, the OTA Requestor SHALL NOT re-query the OTA Provider which was the subject of the command sooner than the longest of either:

- 2 minutes (120 seconds) from the last [QueryImage](#) command;
- the delay stated by the [DelayedActionTime](#) field.

Note that if a Node loses its timekeeping state due to events such as power loss or restart, the above timing constraint MAY be ignored, however, the previously stated overriding constraint of a minimum delay of 120 seconds between queries to any single OTA Provider by a given OTA Requestor has to be respected.

Handling NotAvailable value in Status field

The NotAvailable status indicates that there is definitely no update currently available from the queried OTA Provider.

The OTA Requestor MAY choose to attempt a [QueryImage](#) command on a different OTA Provider in its OTA Provider List to determine if an update is available from that other OTA Provider.

Otherwise, if there are no other OTA Providers available to query, the OTA Requestor SHALL NOT re-query the OTA Provider which was the subject of the command sooner than the longest of either:

- 2 minutes (120 seconds) from the last [QueryImage](#) command;
- the delay stated by the [DelayedActionTime](#) field, if present.

Note that if a Node loses its timekeeping state due to events such as power loss or restart, the above timing constraint MAY be ignored, however, the previously stated overriding constraint of a minimum delay of 120 seconds between queries to any single OTA Provider by a given OTA Requestor have to be respected.

Handling errors from QueryImage Command

If an OTA Requestor hits error conditions of any kind in invoking the QueryImage command, including receiving DownloadProtocolNotSupported in Status, there are two possible outcomes:

1. If the OTA Requestor has a Software Image it had previously successfully downloaded and verified, the OTA Requestor SHOULD skip the Query step, and move directly to the Apply step (see [Section 11.20.3.6, “Applying a software update”](#)). This increases the likelihood that the OTA Requestor will eventually succeed to apply a previously transferred Software Image.
2. If the OTA Requestor is still attempting to discover an OTA Update, it MAY choose to attempt a [QueryImage](#) command on a different OTA Provider in its OTA Provider List, in which case the timing for the query SHALL match the query timing constraints expressed in the previous paragraphs of this section. Otherwise, it SHALL continue to query the same OTA Provider, again following the query timing constraints previously expressed.

11.20.3.3. Availability of Software Images

The algorithm used by the Image Selection Logic to determine availability of a new Software Image SHALL consider all fields provided by the OTA Requestor and attempt to provide the newest matching Software Image. The OTA Image Selection Logic SHALL only provide newer (numerically higher) SoftwareVersion than the SoftwareVersion sent in the query. See [Section 11.20.3.3.2, “Conceptual algorithm for matching OTA Software Images applicable to a query”](#) for more details.

The OTA Provider MAY provide a Software Image that only conveys data for a subset of updateable components within the OTA Requestor’s Node. These cases of partial or componentized software updates are determined purely by the entity generating the OTA Software Image, and the OTA Provider SHALL never mutate the contents of an OTA Software Image.

The original provider of a Software Image SHOULD be able to assume the contents of the Software Image will remain unchanged and signatures would remain valid. Therefore, an OTA Provider SHALL NOT modify the contents of any Software Images other than allowing that OTA Requestor to index into the Software Image using the [BDX Protocol](#) or other supported download protocol, such that the OTA Requestor may obtain only the desired parts of the Software Image.

The OTA Provider SHALL NOT hide or otherwise mask the contents of a Software Image available

for transfer to a requestor.

In order for different vendors to participate as widely as possible in the distribution of Software Images for the widest variety of products without requiring bilateral distribution agreements between each pair, it is RECOMMENDED for vendors to participate in distribution schemes that maximize availability across other vendors and OTA Providers.

Vendors SHOULD build data sets aggregating the metadata and payloads of Software Images to support their OTA Image Selection Logic by any means they deem satisfactory. Vendors SHOULD refer to canonical databases, such as the [Distributed Compliance Ledger](#).

Given that most Fabrics likely will contain a reduced subset of Nodes capable of acting as OTA Providers compared to the larger set of vendors represented in the many deployed Nodes, it is advantageous to end-users that Vendors attempt to cover as many other vendors with their data sets. This will ensure that the majority of Software Image queries can be fulfilled if a vendor has released a newer version than that installed on the querying OTA Requestor.

Download Protocol selection

The OTA Image Selection Logic SHALL consider the OTA Requestor's [supported download protocols](#) to determine whether to respond to a [QueryImage](#) command.

If either the BDXSynchronous or BDXAsynchronous protocols are supported by the OTA Requestor, the OTA Provider SHALL prefer to respond to the OTA Requestor with a BDX protocol URI, as long as it can fulfill the role of BDX server for the OTA Requestor.

Otherwise, if the HTTPS protocol is supported by the OTA Requestor, and the OTA Provider determines that an OTA Software Image is available to fulfill the request from a server supporting HTTPS, it SHOULD respond with the direct source URI, so that the OTA Requestor MAY download it directly.

Otherwise, if the VendorSpecific protocol is supported by the OTA Requestor, and the OTA Provider has sufficient knowledge of the OTA Requestor's capabilities based on the [QueryImage](#) command arguments, it SHOULD respond with a URI which is known to be understood by the OTA Requestor. It is RECOMMENDED to limit usage of this modality and prefer BDXSynchronous and BDXAsynchronous.

Conceptual algorithm for matching OTA Software Images applicable to a query

An OTA Provider MAY use any of the fields of the [QueryImage](#) command in any way it deems applicable to determine whether an appropriate OTA Software image is available for the OTA Requestor.

However, to increase interoperability, OTA Providers which have access to the data present in the [Distributed Compliance Ledger \(DCL\)](#), whether from cached subset or from a full replica, SHOULD employ at least the common conceptual algorithm provided in this section to determine whether an OTA Image is available.

The information to access OTA Software Image locations for certified software versions is available in the [DCL DeviceSoftwareVersionModel Schema](#), which is indexed by VendorID, ProductID and

SoftwareVersion.

The inputs to the conceptual algorithm are:

- A subset of the fields of the [QueryImage](#) command as a structure named **requestor**:
- **VendorID** of the requestor as **vendorID**
- **ProductID** of the requestor as **productID**
- Current **SoftwareVersion** of the requestor as **softwareVersion**
- The list of all current entries for the given VendorID and ProductID from the **DeviceSoftwareVersionModel** schema, ordered by SoftwareVersion, accounting for the following fields, as an array **candidates[]**
 - **SoftwareVersion** of the entry as **softwareVersion**
 - **SoftwareVersionValid** of the entry as **softwareVersionValid**
 - **MinApplicableSoftwareVersion** of the entry as **minApplicableSoftwareVersion**
 - **MaxApplicableSoftwareVersion** of the entry as **maxApplicableSoftwareVersion**

The output of the conceptual algorithm is a tuple of:

- **candidateWasFound**, a boolean predicate indicating whether a newer version candidate was found
- **softwareVersionFound**, the version of the newer version candidate, if **candidateWasFound** was true, 0 otherwise.

The algorithm is as follows:

1. Assume no matching candidate found
 - **candidateWasFound** := False
 - **softwareVersionFound** := 0
2. Obtain candidates from a DCL replica or from a DCL-based dataset for the given **vendorID** and **productID** in the query, keeping only entries where **softwareVersionValid** is true.
 - An OTA Provider MAY as well filter available versions by certification compliance status (see [Section 11.23.8, “DeviceSoftwareCompliance / Compliance test result Schema”](#)).
3. Sort all candidates by ascending **softwareVersion**.
4. Iterate through all candidates to find all positive matches within the sorted candidates. A "positive match" is a candidate which fullfills every condition in the following list:
 - **requestor.softwareVersion** < **candidate.softwareVersion**
 - **requestor.softwareVersion** ≥ **candidate.minApplicableSoftwareVersion**
 - **requestor.softwareVersion** ≤ **candidate.maxApplicableSoftwareVersion**
5. From the positive matches, select the very last one of list, which will be the newest (numerically highest) possible softwareVersion that could be used. If no positive matches were found, no new software version is available.

A pseudocode of the conceptual algorithm is presented below:

```
# Get candidates for VendorID/ProductID from DCL DeviceSoftwareVersionModel
# schema (e.g. from a replica)
candidates = obtain_candidates_from_dcl(requestor.vendorID, requestor.productID)

def find_newer_version(candidates, requestor):
    # Ensure all candidate records are in monotonic increasing numerical order
    sort(candidates, key="softwareVersion", order="ASCENDING")

    currentCandidate = None

    for candidate in candidates:
        # Current version already newer: skip
        if requestor.softwareVersion >= candidate.softwareVersion:
            continue

        # Candidate requires higher version than already installed: skip
        if requestor.softwareVersion < candidate.minApplicableSoftwareVersion:
            continue

        # Candidate requires lower version than already installed: skip
        if requestor.softwareVersion > candidate.maxApplicableSoftwareVersion:
            continue

        # Update potential candidate since it is applicable
        currentCandidate = candidate

    if currentCandidate is not None:
        # Last value of currentCandidate is highest matching value
        candidateWasFound = True
        softwareVersionFound = currentCandidate.softwareVersion
    else:
        candidateWasFound = False
        softwareVersionFound = 0

    return (candidateWasFound, softwareVersionFound)
```

If **candidateWasFound** was true, then a version matching (**softwareVersionFound**) was found and its location and associated metadata can be found in the [DeviceSoftwareVersionModel](#) schema of the DCL.

While the [QueryImage](#) command MAY also contain the Location, HardwareVersion and Meta-

dataForProvider fields, they are optional to use by an OTA provider. These additional fields MAY assist an OTA provider in supporting field trial and development policies. The certified releases present in the DCL, however, are only indexed by VendorID, ProductID and SoftwareVersion, based on the associated certification.

Once an OTA software update file location ([OtaUrl](#)) and digest ([OtaChecksum](#)) are found for the associated version candidate, an OTA provider MAY omit re-downloading the file, and serve a cached copy if a local copy of a file exists which matches all of the following constraints from the candidate:

- It has the same OtaChecksum
- It has the same OtaChecksumType
- It has the same OtaFileSize

11.20.3.4. Obtaining user consent for updating software

In the following subsections, the word "User" SHALL be construed to mean "an entity with sufficient privileges associated with the Fabric". For instance, this could be a home dweller having previously configured Nodes or other services and currently having privilege to further affect such configuration. The exact scope for what an "entity" for such a "User" role may be, and what "sufficient privilege" means, SHALL be implementation-dependent.

An OTA Provider SHOULD obtain some form of "User Consent" prior to responding with a URI for a Software Image or letting an OTA Requestor proceed with applying a previously downloaded update. In the context of the OTA Cluster, "User Consent" SHALL be defined as any signal that an OTA Provider may obtain through its implementation-specific logic, that conveys consent to proceed from a User administratively allowed to give such consent in an informed manner.

Example of "User Consent" include:

- Triggering a notification to an interactive application user interface, where at least one User is notified of the availability of new software for a given node, and where a signal of approval to continue with the downloading and applying of that update can be conveyed back to the OTA Provider.
 - Example: An OTA Provider detects the local presence of a mobile device with an application supporting required User accounts through out-of-band means. The OTA Provider makes an implementation-specific request over a protocol of its choice, in-band or out-of-band of the Fabric, to obtain consent. The User is notified on screen with "An ExampleCompany Light Bulb needs update to version 1.2.3. Tap here for release notes. Do you want to proceed?". The User then selects "Agree to Update" and the signal is relayed back to the OTA Provider, which then proceeds.
- Relaying of a previously stored consent signal, previously provided by a User at some point in the past. The original capture of the stored consent signal should have been made after having provided sufficient information to the User to understand the consequences of such stored consent. Multiple signals, covering different Nodes, Vendors or Device Classes, may be stored independently to affect a variety of deferred consent policies.
 - Example: An OTA Provider contacts an implementation-specific server with metadata about

the OTA Requestor and details for available Software Images, and obtains a consent signal based on an Administrator having previously stated an account preference to "Always apply software updates to light bulbs". The OTA Provider then proceeds further.

User consent delegation to Nodes

Some capable Nodes MAY have sufficient hardware capabilities to request user consent by means such as display or voice, and subsequently recover user consent feedback through input mechanisms. These devices MAY request optional delegation of user consent by the following method:

1. The OTA Requestor SHALL set the [RequestorCanConsent](#) field in the QueryImageRequest to True, indicating ability to obtain consent.
2. The OTA Provider, if it determines that the best way to obtain user consent is to delegate to the OTA Requestor Node, SHALL include the [UserConsentNeeded](#) field, with a value set to True in the QueryImageResponse, indicating that user consent was not previously obtained, and that delegation SHALL occur.
3. The OTA Requestor, upon observing presence of [UserConsentNeeded](#) field set to True and the availability of an image in the QueryImageResponse SHOULD proceed to obtain user consent using its onboard means, prior to transferring the OTA Software Image reported. If the UserConsentNeeded field is set to False or absent, the OTA Requestor SHALL assume that the OTA Provider already obtained sufficient user consent during the querying phase.

The above method of obtaining User Consent at the OTA Requestor level SHALL NOT be used if a Node is configured with the [LocalConfigDisabled](#) attribute set to True as reflected by the Basic Information Cluster.

User consent recommendations

Because of the variety of Vendors and Devices, the concept of "User Consent" will necessarily take many different forms. Therefore, it is RECOMMENDED that every implementer of OTA Provider logic implement a transparent and easy-to-use set of functionality to allow Users to provide or deny consent for software updates, in a way that functionally integrates with their products and respects the general requirements stated above. Implementation of this feature is expected to improve "user experience" and assist with building trust regarding the installation of new software on Nodes.

It is RECOMMENDED that any method of consent that stores consent signals also provide a way to revoke this consent in the future.

It is RECOMMENDED that metadata from Software Images be used to convey as much information as required within the available set, so that a User can make an informed decision based on the nature of the product being updated, the human-readable instance of the new version number (e.g. [SoftwareVersionString](#) in OTA Image Header), the changes made available, and their side-effects on product functionality. Any URL for online contents conveyed during this process SHOULD point to content that can be localized at the time of delivery, whenever possible. The responsibility for the maintenance of such version information is on the Vendor providing the URL and metadata. The OTA Provider and associated implementation-specific logic SHALL allow a User to consent to an update, even if errors occur while trying to provide additional release information, as the metadata within the Software Image should suffice to provide a first-order description of the new version, which could then be researched or cross-referenced by the User.

It is RECOMMENDED that Vendor-provided Software Update metadata, such as release note URLs, be maintained in the long-term with stable locations, preferably in a manner allowing historical caching by common online search engines, where applicable. See [Section 11.23.7.12, “ReleaseNotesUrl”](#) and [Section 11.21.2.4.8, “ReleaseNotesUrl field”](#) for sources of such information.

11.20.3.5. Transfer of OTA Software Update images

Execution of an OTA Software Update image’s transfer depends on the protocol provided in the [ImageURI](#) field of the query response.

The following are OTA Software Image transfer examples:

- An OTA Requestor invokes a [QueryImage](#) command stating only support for BDX in its [ProtocolsSupported](#). The OTA Provider, using its OTA Image Selection Logic, determines that a Software Image is available. There are several cases to be considered:
 1. The Software Image is at a URI referring to a resource on the public Internet (e.g. <https://domain.example/images/software.bin>).
 - a. The OTA Provider MAY completely download the Software Image, temporarily, to local storage. It would then reply to the OTA Requestor with a [locally-accessible BDX URI](#), such as bdx://8899AABBCCDDEEFF/software_8ce40aa1.bin. In that case, the OTA Requestor SHALL proceed with the download from the OTA Provider using the [BDX protocol](#).
 - b. The OTA Provider MAY employ a variety of buffering and proxying schemes of underlying HTTPS transfers to support the OTA Requestor downloading in real-time as a form of direct proxy. It would immediately reply to the OTA Requestor with a [locally-accessible BDX URI](#), such as bdx://8899AABBCCDDEEFF/software_8ce40aa1.bin. In that case, the OTA Requestor SHALL proceed with the download from the OTA Provider using the [BDX protocol](#). The only difference with the previous case is the fact that the transfer uses data directly proxied in real-time, as opposed to the OTA Requestor downloading a pre-stored cached copy of the same Software Image.
 2. The Software Image is already cached on the OTA Provider, either from pre-seeding over some implementation-specific scheme, or from having previously served this software update to another OTA Requestor. In that case, the OTA Provider SHALL reply to the OTA Requestor with a [locally-accessible BDX URI](#), such as bdx://8899AABBCCDDEEFF/software_8ce40aa1.bin. In that case, the OTA Requestor SHALL proceed with the download from the OTA Provider using the [BDX protocol](#).
- An OTA Requestor invokes a [QueryImage](#) command stating support for BDX and HTTPS in its [ProtocolsSupported](#). The OTA Provider, using its OTA Image Selection Logic, determines that a Software Image is available. The Software Image is at a URI referring to a resource on the public Internet (e.g. <https://domain.example/images/software.bin>). In the case of support for both HTTPS and BDX, all of the above cases are applicable, in addition to the following:
 - The OTA Provider knows that the OTA Requestor supports HTTPS from the [QueryImage](#) command. The OTA Provider MAY then respond directly to the OTA Requestor with the <https://domain.example/images/software.bin> URI. The OTA Requestor SHALL proceed to download from the public Internet using the HTTPS protocol.

As described above, an OTA Provider SHALL either proxy synchronously or asynchronously the

actual Software Image data for [BDX Protocol](#) clients, when a Software Image is determined to be available from the public Internet or from local storage.

Upon receipt of a [QueryImageResponse](#) command containing a [DelayedActionTime](#) field, the OTA Requestor SHALL wait for at least the stated delay time before initiating the first part of a file transfer for the URI provided.

The OTA Provider MAY expunge any previously cached Software Image downloaded on behalf of other OTA Requestors, to save storage, at any time, as long as no transfer is currently in active progress. It is RECOMMENDED that OTA Providers on a given Fabric maintain as much Software Image cache as practical, to improve availability of software image and reduce latency between QueryImage requests and availability of the matching QueryImageResponse for a new Software Image ready to be downloaded.

In order to support non-BDX protocols relying on the public Internet, or intranets, the OTA Requestor SHALL only report support for protocols requiring public Internet access if it has determined that it does indeed have access to the necessary network domains beyond the Fabric. The OTA Requestor MAY employ any method it deems satisfactory to determine public Internet reachability. Because of the variety of firewall and security policies on network infrastructure, it is RECOMMENDED that all Nodes whose primary networking interactions lie within protocols in-band of this wider specification support the BDX method, so that even if a Node cannot access the public Internet, it MAY still obtain OTA Software Images by relying on a local OTA Provider which can.

For BDX transfers, the following BDX-specific constraints SHALL be followed:

- Receiver-Drive mode SHALL be used by the OTA Provider for any transfers initiated from a secure channel on non-TCP transport.
- Asynchronous mode MAY be used by the OTA Provider for any transfer initiated from a secure channel on TCP transport. Note that Asynchronous mode is provisional.
- Idle time-out SHALL be no less than 5 minutes for either Receiver-Driver or Asynchronous mode, before aborting a transfer.
- Block sizes constraints SHALL be as follows:
 - Maximum Block Size over all transports SHALL be a power of two if the OTA Requestor requests a value larger than 128 bytes.
 - For an OTA Requestor-requested Maximum Block Size value between 16 and 128, the exact requested value SHALL be used. This constraint allows low-power Nodes to precisely control the block sizes to ensure their power constraints are respected, including enabling single-frame block transfers over communication mediums where MTU is very small.
 - Maximum Block Size requested by OTA Requestors over non-TCP transports SHALL be no larger than 1024 (2^{10}) bytes. OTA Providers SHALL support the Maximum Block Size of at least 1024 bytes in those cases.
 - Maximum Block Size requested by OTA Requestors over TCP transport SHALL be no larger than 8192 (2^{13}) bytes. OTA Providers SHALL support a Maximum Block Size of at least 4096 (2^{12}) bytes in this case and MAY support 8192 bytes.
 - Actual Block Size used over all transports SHALL be the negotiated Maximum Block Size for every block except the last one, which may be of any size less or equal to the Maximum

Block Size (including zero).

- The OTA Requestor SHALL NOT rely on the ReceiveAccept message from the OTA Provider having the RC[DEFLen] bit set (see [Section 11.22.5.4.2.1, “ReceiveAccept RC\[DEFLen\]: definite length present”](#)) and the associated LEN field populated. Instead, OTA Requestors SHALL rely on OTA Software Image metadata to determine the expected size to download.
- The ReceiveInit message from the OTA Requestor MAY have the RC[STARTOFS] bit set and associated STARTOFS field set to indicate the resumption of a transfer previously aborted, or to affect partial windowed access to the portion of a Software Image desired.
- The ReceiveInit message from the OTA Requestor MAY have the RC[DEFLen] bit set and associated DEFLen field set to state the desired maximum size of the transfer.

Since OTA Requestors MAY need to read Software Image in parts, it is RECOMMENDED that OTA Providers maintain cached Software Images for at least 5 minutes after closure of the last OTA Requestor transfer, so that the OTA Requestor MAY come back to read different parts of an OTA file.

In the case of very large Fabrics, it often occurs that there is a large number of the same model of Node within a given location. Because of this, OTA Providers SHOULD avoid downloading the same Software Image repeatedly for proxying if it can determine that multiple OTA Requestors are requesting, or can be expected to request, the same Software Image.

It is RECOMMENDED to keep Software Images cached for as long as practical to reduce having to reach external off-Fabric resources frequently to address the update needs of a large fleet of identical Nodes that could share a single pre-downloaded cached copy in the OTA Provider. This reduces burden on content delivery servers for Software Images and reduces the amount of data transferred by an OTA Provider from external off-Fabric servers to fulfill software update requirements.

It is RECOMMENDED that Sleepy End Devices make their best effort to optimize their sleep intervals during the OTA Software Image transfer process over BDX to ensure that the download completes in a timely manner. However, it is acknowledged that some Sleepy End Devices might not be able to do so, due to limitations related to their batteries or other constrained power sources. Therefore, such devices MAY take much longer to complete the download process.

In the case where a BDX transfer is aborted due to unforeseen circumstances (e.g. power loss, crash, battery drain on either side), the OTA Requestor MAY try to use a partial (i.e. range-based) transfer to recover and continue the download without having to start from the beginning of a given Software Image. An OTA Requestor SHALL abort retrying a transfer after three attempts in a row where each yielded no forward progress.

It is RECOMMENDED for the OTA Provider to validate the length and digest of proxied images whenever possible (see OTA software update file [Header field](#)) to avoid continuing a transfer if the data is obviously corrupted.

In any situation where an OTA Requestor reaches a terminal failure point for a Software Image transfer and all possible retries or alternate OTA Providers have been exhausted, that OTA Requestor SHALL reset its entire software updating state and revert to doing a future query at the next possible scheduled time, so that perhaps a new Software Image may be available again.

The above situation may occur, for example, if an OTA Provider had cleared its Software Update

Image File cache for any reason, or if there is a transient network failure of sufficient duration to prevent a complete transfer to take place.

Once the entirety of a Software Image has been downloaded and is ready to apply, the OTA Requestor SHALL execute the "[Applying a software update](#)" sequence of the next section.

11.20.3.6. Applying a software update

Once a Software Image has been fully downloaded based on a [QueryImageResponse](#) command, the OTA Requestor SHALL proceed with a sequence to determine when to apply the update by invoking the [ApplyUpdateRequest](#) command.

UpdateToken usage

The OTA Requestor SHALL provide an [UpdateToken](#) to the OTA Provider with the [ApplyUpdateRequest](#) command. This token SHALL be a previously provided [UpdateToken](#) from the last [QueryImageResponse](#), unless the token was lost by the OTA Requestor. In case of token loss, the OTA Requestor SHALL use its [Operational Node ID](#) encoded as a 64-bit value in network byte order. This [UpdateToken](#) MAY be used by an OTA Provider to track deferred OTA application or otherwise allow short-term tracking of OTA Requestors for algorithmic bookkeeping. The OTA Provider SHALL NOT consider an invalid [UpdateToken](#) as a reason to continuously deny or delay an OTA Requestor's request to apply a Software Image.

Update application process

On receipt of the [ApplyUpdateRequest](#) command, the OTA Provider SHALL respond with an action to be taken by the OTA Requestor before activating the new version. The method used to determine the [Action](#) field of the response MAY be based on implementation-specific rules and logic.

Note that the [DelayedActionTime](#) field is a relative time delay from the moment of receipt, which needs to be computed by the OTA Provider to reflect the difference between the OTA Provider's current time and the desired time for execution of the Action.

In case of a successful invocation, the following actions to be taken by the OTA Requestor are possible, based on the [Action](#) field in the [ApplyUpdateResponse](#) command:

- Proceed: Apply the update, taking in account the delay time stated in [DelayedActionTime](#).
 - If the [DelayedActionTime](#) is zero, then the OTA Requestor SHALL apply the update without additional delay.
 - If the [DelayedActionTime](#) is non-zero, the OTA Requestor SHALL await at least [DelayedActionTime](#) seconds prior to applying the software update. An example use of this Action by an OTA Provider is to schedule application of a Software Image based on a user's preferred update time for Nodes of a certain type (e.g. light bulbs or window coverings) to occur at a time when the user is not at home, or when the temporary unavailability of the Node during the update would not pose a problem.
 - When the Proceed action is given, the OTA Requestor SHALL NOT invoke the [ApplyUpdateRequest](#) command again, unless the OTA Requestor suffers an error or unexpected condition while proceeding to apply the new Software Image.

- **AwaitNextAction:** Await at least the given delay time in DelayedActionTime before re-invoking an [ApplyUpdateRequest](#) to get a new Action.
 - If the DelayedActionTime is less than 120 seconds (2 minutes), the OTA Requestor SHALL assume a value of 120 seconds.
 - An example use of this Action by an OTA Provider is to schedule application of a Software Image based on non-occupancy of a room, which Nodes collaborating with the OTA Provider may be able to ascertain, but which may require several attempts over time.
 - It is RECOMMENDED to keep usage of this Action to a practical minimum, as it may cause OTA Requestors to be delayed in their application of a Software Image.
 - The AwaitNextAction action SHALL NOT be emitted in such a way as to cause more than 24 hours of delay in applying an available Software Image. It is expected that user consent having been previously granted should satisfy the overall scheduling constraint this imposes.
- **Discontinue:** The OTA Provider is conveying a desire to rescind a previously provided Software Image.
 - The DelayedActionTime SHALL be ignored by the OTA Requestor if present.
 - The OTA Requestor SHOULD clear its previously downloaded and verified Software Image, if it had been obtained from the same OTA Provider as the one providing the Discontinue action.
 - This action SHALL only be used as a stopgap when it is known that a given Software Image previously provided may cause significant negative side-effects to an OTA Requestor, such as unrecoverable loss of functionality, or other damage.
 - In case of receiving this action unexpectedly (e.g. from a different OTA Provider than the one where a Software Image was downloaded), an OTA Requestor MAY ignore it and consider it the same way as if the [ApplyUpdateRequest](#) command had proceeded with an error.

It is RECOMMENDED that for any OTA Requestor invoking the [ApplyUpdateRequest](#) command with an unknown UpdateToken, the OTA Provider SHOULD assume that the OTA Requestor has a Software Image ready to apply and thus respond with the Proceed or Await action, rather than responding with an error or Discontinue action.

In case of time-out or any error in obtaining an answer to the ApplyUpdateRequest command, or in case of a restart or other unrecoverable situation while awaiting the DelayedActionTime for a Proceed or Await action, the OTA Requestor SHOULD retry to execute the "Querying the OTA Provider" flow (see [Section 11.20.3.2, "Querying the OTA Provider"](#)) again, whenever the OTA Requestor deems it possible.

In case of failure of every possible retry mechanism for at least 3 total attempts, or over more than 24 hours, an OTA Requestor having successfully downloaded and verified a Software Image MAY apply the update. This measure of last recourse is to avoid situations where a critical issue affecting a particular software version would prevent an OTA Requestor or OTA Provider from properly executing the "Applying a Software Update" flow, thus leaving an OTA Requestor in reduced or a forever-impaired state that could otherwise be resolved by applying the Software Image it had successfully downloaded and verified.

After completion of an update, an OTA Requestor SHOULD invoke a NotifyUpdateApplied command

to the OTA Provider which provided the initial query response to indicate that the OTA Requestor has successfully applied the OTA Software Update Image. The OTA Requestor SHALL NOT retry at the application level to invoke this command if a response is not received.

In case the device performs a reboot in order to apply the update, the value of the [BootReason](#) attribute in the [General Diagnostics](#) cluster SHALL be set to [SoftwareUpdateCompleted](#).

NOTE

The above requirement was added in Matter 1.4.2. Devices certified before this version MAY return a different boot reason.

11.20.4. Security considerations

Security for the OTA Software Update capabilities encompasses these areas: Software Image verification, Software Image transport, and Software Image encryption. Security mechanisms in given applications dictate the security level of OTA upgrading. For example, an application with strict security policies (such as a smart lock) MAY support Software Image encryption at rest beyond the secure channel data-in-transit encryption, while other applications MAY only support data-in-transit encryption. Each Vendor SHALL decide the list of required security policies for their use of the OTA Software Update capabilities for a particular product.

11.20.4.1. Image Encryption

A Vendor MAY apply at-rest encryption to Software Image bodies, excluding the Software Update Image Header, using any algorithm of its choosing.

It is out of scope of this specification to mandate such means of protecting the confidentiality of the Software Images.

11.20.4.2. Image Verification

Asymmetric Verification of Authenticity and Integrity

The verification of the authenticity and integrity of Software Images by OTA Requestors is mandatory for security reasons. This is most often accomplished through asymmetric encryption technologies where only one entity is able to create a digital signature but many entities are able to verify it. Software Images SHALL be signed by a private key used by the Vendor for software image signing purposes, with that signature attached to the Software Image that is transported to the OTA Requestor. Once the complete Software Image has been received, the signature SHALL be verified using a matching public key known to the OTA Requestor performing the validation. See [Section 13.5, “Firmware”](#) for the associated security requirements. The format and algorithms used for code signatures verification are out of scope of this specification. The OTA Provider SHOULD NOT expect to be able to validate OTA Software Image signatures on its own.

OTA Requestors MAY be pre-installed with the certificate (public key) of the entity that created the signature, or they MAY receive the certificate over-the-air. How the signer's security data is obtained is considered outside the scope of the OTA Software Update Cluster and is Vendor Specific. When signer certificates are sent over-the-air, they SHALL be securely transferred from a trusted source to reduce the chance an attacker MAY inject their own signer certificate into the OTA Requestor.

Software Images with verification mechanisms built in MAY be transported over insecure communication mechanisms while still maintaining their authenticity and integrity. In fact, it is likely that the originator of the Software Image (a Vendor) will not be directly connected to any Fabric and therefore distribute the Software Image across other mediums (such as the Internet) before arriving on the Fabric. Therefore, it is crucial that the Software Image verification be independent of the communication medium. Any attempts to tamper with the signature or the data itself SHALL be detected and SHALL cause the Software Image to be rejected by the target OTA Requestor. A Software Image from an attacker that crafts its own signed image and tries to have it accepted would be rejected since that image will not be signed by the Vendor's signing authority.

Since Software Images MAY contain software for multiple sub-components (e.g. main processor firmware, radio firmware, graphical/audio assets) which MAY each employ different code signing keys, Software Images SHOULD provide at least an overall authenticity and integrity validation for the entire image, regardless of how it is segmented.

Individual Vendors MAY augment the basic security and authenticity schemes provided by the Software Images and provide their own extensions within the payloads. Those extensions are outside the scope of the OTA Software Update Cluster.

Software Images MAY be encrypted with symmetric keys such that only those OTA Requestors that need to decrypt the Software Image have access to the key. This MAY be used to mask or obfuscate the contents of Software Images from intermediate network participants conveying the Software Images, while in transit and at rest. However, the security of this system is dependent on the security of all OTA Requestors that have access to the symmetric key and the method of storage of the final Software Image at rest on a given OTA Requestor. The schemes employed by different Vendors to encrypt the body of Software Images in transit and at rest, is outside of the scope of this specification.

11.20.5. Some special situations

With the mechanisms provided by this Cluster, roll-out of new Software Images applies equally to all OTA Requestors of matching a given <VendorID, ProductID> tuple, uniformly across the world. The subsections below describe two situations where finer-grained roll-out can be achieved for some regions or to distribute special non-standard versions.

11.20.5.1. Regional roll-out of Software Images

Some Manufacturers have a policy of rolling-out by region (i.e. set of countries), to provide worldwide release schedule differentiation, as well as to test roll-outs gradually, among other reasons. These regional roll-outs may only be feasible using manufacturer-specific schemes that are in addition to the common flows described in this cluster. Common recommended behavior (see [Section 11.20.3.3.2, "Conceptual algorithm for matching OTA Software Images applicable to a query"](#)) does not support regional roll-out since there does not exist a location tagging attribute in the [Distributed Compliance Ledger \(DCL\)](#). For regional rolls-outs prior to full roll-out, refer to the overall techniques described in [Section 11.20.5.2, "Roll-out of non-standard Software Images"](#).

11.20.5.2. Roll-out of non-standard Software Images

Many Manufacturers conduct field trials to test different versions of software (e.g. A/B-testing, beta

testing), or provide dedicated Software Images to a subset of Nodes to affect particular diagnosis tasks, etc. The mechanism described in this cluster is not particularly well-suited for such fine-grained deployment (unless the OTA Provider is provided by, or associated with, the Manufacturer).

To achieve more targeted roll-out, Vendors MAY commission a Node on the same Fabric as the devices requiring the special rules, so that it MAY provide OTA Provider capabilities beyond the core interoperable aspects of this Cluster. Finer-grained selection MAY be applied by special OTA Software Image Selection logic in a given OTA Provider, using the [MetadataForProvider](#) field and [MetadataForRequestor](#) fields of the [QueryImage](#) command. Furthermore, such special OTA Provider may identify itself by including the [MetadataForNode](#) field in a given [AnnounceOTAProvider](#) command.

If such a special Software Image is running on an OTA Requestor, the OTA Requestor MAY reject Software Images provided to it by an OTA Provider on the Fabric, to prevent loss of the non-standard Software Image. A Factory Data Reset of the OTA Requestor SHALL remove such override.

11.20.5.3. Other considerations

While it is expected that Nodes fulfilling the role of OTA Provider will likely have high availability within the Fabric, it may be possible some will be battery-operated or be power-cycled frequently. It is RECOMMENDED that an OTA Provider Node that determines it cannot reliably stay available to service OTA Requestors SHOULD avoid responding with an available OTA Software Image when responding to a [QueryImage](#) command (see [Section 11.20.6.5.2, “QueryImageResponse Command”](#)) unless it has sufficient availability to allow a long-running [BDX Protocol](#) transfer to finish. In general, OTA Provider role SHOULD be fulfilled by a Node with a reliable network availability and stable power, especially if it is set as a default in the [DefaultOTAProviders](#) attribute.

11.20.6. OTA Software Update Provider Cluster

This cluster implements the Provider role in the OTA process.

11.20.6.1. Revision History

The global [ClusterRevision](#) attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.20.6.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	OTAP

11.20.6.3. Cluster ID

ID	Name
0x0029	OTA Software Update Provider

11.20.6.4. Data Types

StatusEnum Type

This data type is derived from [enum8](#).

See [Section 11.20.3.2, “Querying the OTA Provider”](#) for the semantics of these values.

Value	Name	Summary	Conformance
0	UpdateAvailable	Indicates that the OTA Provider has an update available.	M
1	Busy	Indicates OTA Provider may have an update, but it is not ready yet.	M
2	NotAvailable	Indicates that there is definitely no update currently available from the OTA Provider.	M
3	DownloadProtocol-NotSupported	Indicates that the requested download protocol is not supported by the OTA Provider.	M

ApplyUpdateActionEnum Type

This data type is derived from [enum8](#).

See [Section 11.20.3.6, “Applying a software update”](#) for the semantics of the values. This enumeration is used in the Action field of the [ApplyUpdateResponse command](#). See ([Action](#)).

Value	Name	Summary	Conformance
0	Proceed	Apply the update.	M
1	AwaitNextAction	Wait at least the given delay time.	M
2	Discontinue	The OTA Provider is conveying a desire to rescind a previously provided Software Image.	M

DownloadProtocolEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	BDXSynchronous	Indicates support for synchronous BDX.	M
1	BDXAsynchronous	Indicates support for asynchronous BDX.	O
2	HTTPS	Indicates support for HTTPS.	O
3	VendorSpecific	Indicates support for vendor specific protocol.	O

Note that only HTTP over TLS (HTTPS) is supported (see [RFC 7230](#)). Using HTTP without TLS SHALL NOT be supported, as there is no way to authenticate the involved participants.

NOTE Support for the asynchronous BDX mode is provisional.

11.20.6.5. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	QueryImage	client ⇒ server	QueryImageResponse	O	M
0x01	QueryImageResponse	client ⇐ server	N		M
0x02	ApplyUpdateRequest	client ⇒ server	ApplyUpdateResponse	O	M
0x03	ApplyUpdateResponse	client ⇐ server	N		M
0x04	NotifyUpdateApplied	client ⇒ server	Y	O	M

QueryImage Command

Upon receipt, this command SHALL trigger an attempt to find an updated Software Image by the OTA Provider to match the OTA Requestor's constraints provided in the payload fields.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	VendorID	vendor-id	all			M
1	ProductID	uint16	all			M

ID	Name	Type	Constraint	Quality	Fallback	Conformance
2	Software-Version	uint32	all			M
3	ProtocolsSupported	list[DownloadProtocolEnum]	max 8			M
4	Hardware-Version	uint16	all			O
5	Location	string	2			O
6	Requestor-CanConsent	bool	all		False	O
7	Meta-dataFor-Provider	octstr	max 512			O

VendorID Field

The value SHALL be the Vendor ID applying to the OTA Requestor's Node and SHALL match the value reported by the [Basic Information Cluster VendorID](#) attribute.

ProductID Field

The value SHALL be the Product ID applying to the OTA Requestor's Node and SHALL match the value reported by the [Basic Information Cluster ProductID](#) attribute.

SoftwareVersion Field

The SoftwareVersion included in the request payload SHALL provide the value representing the current version running on the OTA Requestor invoking the command. This version SHALL be equal to the [Software Version attribute](#) of the [Basic Information Cluster](#).

ProtocolsSupported Field

This field SHALL contain a list of all download protocols supported by the OTA Requestor.

This field SHALL be used by the OTA Provider to generate the correct URI for the location of the Software Image when one is found to be available. BDX Synchronous transfer mode SHALL always be supported by an OTA Provider. Furthermore, OTA Providers with access to external networking SHOULD support the HTTPS protocol. OTA Providers MAY support other protocols.

The algorithm to select the specific protocol to use in a given Software Image URI is implementation-dependent, provided that the rules in [Section 11.20.3.3.1, "Download Protocol selection"](#) are followed.

See [Section 11.20.3.2, "Querying the OTA Provider"](#) and [Section 11.20.3.5, "Transfer of OTA Software Update images"](#) for more details about usage of this field.

HardwareVersion Field

The value of this field, if present, SHALL contain the OTA Requestor's hardware version, and SHALL be equal to the [HardwareVersion attribute](#) of the [Basic Information Cluster](#).

Location Field

The location, if present, SHALL provide the same value as the [Basic Information Cluster Location](#) attribute for the OTA Requestor as configured. This MAY be used by the OTA Provider logic to allow per-region selection of the Software Image.

RequestorCanConsent Field

This field SHALL be set to true by an OTA Requestor that is capable of obtaining user consent for OTA application by virtue of built-in user interface capabilities. Otherwise, it SHALL be false.

See [Section 11.20.3.4, "Obtaining user consent for updating software"](#) for application details about usage.

MetadataForProvider Field

This optional field, if present, SHALL consist of a [top-level anonymous list](#); each list element SHALL have a [profile-specific](#) tag encoded in [fully-qualified](#) form. Each list element SHALL contain a manufacturer-specific payload, which the OTA Requestor invoking this command wants to expose to the receiving OTA Provider. This payload MAY be used for any purpose and SHOULD be as small as practical.

The use of this field SHOULD be restricted to Vendor-specific usage and SHALL NOT be used as a selector required to match for the selection of a Software Image in production environments, unless absolutely necessary, as the interpretation of this field may be ambiguous to OTA Providers implementing the Cluster in a compliant but divergent way from the sender.

An example of usage for this field is for an OTA Requestor to provide specific data about grouping or authentication in field trial environments, where the OTA Provider is likely to understand it and be able to act upon it, either for special selection of image, or recording of activity.

An OTA Provider SHALL report the availability of Software Images, if one is found to be applicable using the other provided fields, even if the MetadataForProvider field is deemed to contain invalid or unknown information. That is, the contents of the MetadataForProvider field SHALL NOT be used to deny a software update to an OTA Requestor, unless both OTA Requestor and OTA Provider have an externally agreed-upon policy whereby strictly correct additional MetadataForProvider is expected to fulfill the OTA Software Update process.

Usage of the QueryImage Command

OTA Requestors SHALL send a QueryImage command to the OTA Provider to determine the availability of a new Software Image.

See [Section 11.20.3.2, "Querying the OTA Provider"](#) for full details about the OTA Software Update Query flow which makes use of this command.

QueryImageResponse Command

This command is sent in response to the [QueryImage](#) command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Status	StatusEnum	all			M
1	DelayedActionTime	uint32	all		0	Status == Busy , 0
2	ImageURI	string	max 256			Status == UpdateAvailable
3	Software-Version	uint32	all			Status == UpdateAvailable
4	Software-Version-String	string	1 to 64			Status == UpdateAvailable
5	UpdateToken	octstr	8 to 32			Status == UpdateAvailable
6	UserConsentNeeded	bool	all		False	0
7	MetadataForRequestor	octstr	max 512			0

Status Field

This field SHALL contain the primary response regarding the availability of a Software Image.

See [Section 11.20.3.2, “Querying the OTA Provider”](#) for details about the possible values for this field and their meaning.

DelayedActionTime Field

This field SHALL be provided when the Status field is set to Busy, and MAY be provided in other cases. This field, if provided, SHALL convey the minimum time to wait, in seconds from the time of this response, before sending another QueryImage command or beginning a download from the OTA Provider. OTA Requestors SHALL respect this minimum delay, unless they had previously restarted and lost track of it. OTA Providers SHOULD expect OTA Requestors to follow this value to their best capability, however, a restarting Node MAY come back sooner, due to having lost track of this state response.

See [Section 11.20.3.2, “Querying the OTA Provider”](#) for details about the rules regarding this field.

ImageURI Field

This field, when present, SHALL contain a URI where the OTA Requestor SHOULD download a Software Image. The syntax of the ImageURI field SHALL follow the URI syntax as specified in [RFC 3986](#).

Beware, this field is conditionally present based on the conformance listed in [Section 11.20.6.5.2, “QueryImageResponse Command”](#).

If the ImageURI specifies a [BDX Protocol bdx:](#) scheme, then the following rules describe the location to be used for download:

1. The URI's [scheme](#) field SHALL be exactly [bdx](#) in lowercase characters.
2. The URI's [authority](#) field SHALL contain only the [host](#) portion and SHALL use string representation of the [Operational Node ID](#) of the Node where to proceed with the download, on the same Fabric on which the OTA Requestor received the QueryImageResponse.
3. The encoding of the Node ID in the [host](#) field SHALL use an uppercase hexadecimal format, using exactly 16 characters to encode the network byte order value of the NodeID, in a similar fashion as the Node Identifier portion of the [Operational Instance Name](#).
 - a. The Operational Node ID in the [host](#) field SHALL match the NodeID of the OTA Provider responding with the QueryImageResponse. The usage of a different Node ID than that of the provider is reserved for future use. This constraint reduces the number of independent [CASE](#) secure channel sessions that have to be maintained to proceed with OTA software updates, thus reducing energy and resource utilization for the software update process.
4. The [user](#) section of the [authority](#) field SHALL be absent, as there are no "users" to be considered.
5. The [port](#) section of the [authority](#) field SHALL be absent, as the port for transport SHALL be determined through [Operational Discovery](#) of the target Node.
6. The URI SHALL NOT contain a [query](#) field.
7. The URI SHALL NOT contain a [fragment](#) field.
8. The [path](#) field SHALL employ absolute path representation and SHALL contain the [file designator](#) of the software image to download at the BDX server. When used with the BDX server, the leading [/](#) separating the URI authority from the path SHALL be omitted. When contacting the BDX server, further processing of the file designator SHALL NOT be done, including handling of URL-encoded escape sequences. Rather, the exact octets of the path, as received SHALL be the values used by both client and server in handling the file designator.
 - a. The [path](#) SHALL only contain valid URI characters.

These rules above for BDX URIs simplify parsing for OTA Requestors receiving Image URIs. The following example procedure shows how the format constraints simplify the extraction of the necessary data to reach the BDX server for download.

1. Verify that the URI is 24 characters or longer, which is the minimum length of a valid BDX URI with all elements present, for example [bdx://00112233AABBCCDD/0](#).
2. Verify the presence of prefix [bdx://](#) indicating a BDX URI.
3. Extract the next 16 characters and convert from uppercase hexadecimal to a 64-bit scalar value,

considering network byte order. This is the destination Node ID.

4. Verify the presence of a path separator / and skip it.
5. Extract the remaining characters of the string as the [file designator](#) to employ when initiating the BDX transfer.

Example ImageURI values are below, and illustrate some but not all of valid and invalid cases:

- Synchronous or Asynchronous [BDX Protocol](#):
 - Valid: `bdx://8899AABBCCDDEEFF/the_file_designator123`
 - Node ID: `0x8899AABBCCDDEEFF`
 - File designator: `the_file_designator123`
 - Valid: `bdx://0099AABBCCDDEE77/the%20file%20designator/some_more`
 - Node ID: `0x0099AABBCCDDEE77`
 - File designator: `the%20file%20designator/some_more`. Note that the %20 are retained and not converted to ASCII 0x20 (space). The file designator is the path as received verbatim, after the first '/' (U+002F / SOLIDUS) following the host.
 - Invalid: `bdx://99AABBCCDDEE77/the_file_designator123`
 - Node ID: Invalid since it is not exactly 16 characters long, due to having omitted leading zeros.
 - Invalid: `bdx://0099aabbccddee77/the_file_designator123`
 - Node ID: Invalid since lowercase hexadecimal was used.
 - Invalid: `bdx:8899AABBCCDDEEFF/the_file_designator123`
 - Invalid since bdx scheme does not contain an authority, that is, it does not have // after the first :.
- HTTP over TLS:
 - Valid: `https://example.domain:8466/software/image.bin`

See [Section 11.20.3.2, “Querying the OTA Provider”](#) for additional details about the flow.

SoftwareVersion Field

This field indicates the version of the image being provided to the OTA Requestor by the OTA Provider.

Beware, this field is conditionally present based on the conformance listed in [Section 11.20.6.5.2, “QueryImageResponse Command”](#).

See [Section 11.20.3.2, “Querying the OTA Provider”](#) for additional details about the flow and acceptable values.

SoftwareVersionString Field

This field provides a string version of the image being provided to the OTA Requestor by the OTA

Provider.

Beware, this field is conditionally present based on the conformance listed in [Section 11.20.6.5.2, “QueryImageResponse Command”](#).

See [Section 11.20.3.2, “Querying the OTA Provider”](#) for additional details about the flow and acceptable values.

UpdateToken Field

Beware, this field is conditionally present based on the conformance listed in [Section 11.20.6.5.2, “QueryImageResponse Command”](#).

See [Section 11.20.3.6.1, “UpdateToken usage”](#) for additional details about the generation and usage of UpdateToken.

UserConsentNeeded Field

This field, if present, SHALL only be interpreted if the OTA Requestor had previously indicated a value of True in the [RequestorCanConsent](#) field of the QueryImageRequest. This field, when present and set to True, SHALL indicate that a capable OTA Requestor must obtain user-visible consent prior to downloading the OTA Software Image.

See [Section 11.20.3.4, “Obtaining user consent for updating software”](#) for application details about usage.

MetadataForRequestor Field

This optional field, if present, SHALL consist of a [top-level anonymous list](#); each list element SHALL have a [profile-specific](#) tag encoded in [fully-qualified](#) form. Each list element SHALL contain a manufacturer-specific payload, which the OTA Provider wants to expose to the receiving OTA Requestor. This payload MAY be used for any purpose and SHOULD be as small as practical.

The presence of this field SHALL NOT be required for correct operation of any OTA Provider compliant with this Cluster specification.

The data for this field does not exist in any [Distributed Compliance Ledger](#) record and SHOULD only be emitted by an OTA Provider with this additional knowledge if it has knowledge that the receiving OTA Requestor MAY be able to use it.

ApplyUpdateRequest Command

This command requests the specified version be installed on the device.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	UpdateToken	octstr	8 to 32			M
1	NewVersion	uint32	all			M

UpdateToken Field

This field SHALL contain the UpdateToken as specified in [Section 11.20.3.6.1, “UpdateToken usage”](#). This field MAY be used by the OTA Provider to track minimal lifecycle state to allow finer-grained scheduling of the application of Software Images by OTA Requestors.

NewVersion Field

The NewVersion field included in the request payload SHALL provide the SoftwareVersion value of the new Software Image which the OTA Requestor is ready to start applying. The OTA Provider MAY use this new version to track or record Software Image application by OTA Requestors.

When Generated

The ApplyUpdateRequest Command SHALL be invoked by an OTA Requestor once it is ready to apply a previously downloaded Software Image.

Effect on Receipt

Upon receipt of this command the OTA Provider SHALL respond with an Action field consistent with the next action the OTA Requestor should take, including any possible time delay.

The OTA Provider SHALL NOT refer to previously stored state about any download progress to reply. If any state keeping is done by the OTA Provider, it SHALL only relate to the UpdateToken and the history of prior ApplyUpdateRequest commands.

See [Section 11.20.3.6, “Applying a software update”](#) for a description of the flow in response to an OTA Provider receiving an invocation of this command.

Handling Error Cases

See [Section 11.20.3.6, “Applying a software update”](#) for all error-handling information.

ApplyUpdateResponse Command

This command is sent in response to the [ApplyUpdateRequest](#) command.

ID	Name	Type	Constraint	Quality	Fallback	Confor- mance
0	Action	ApplyUp- dateActio- nEnum	all			M
1	DelayedAc- tionTime	uint32	all			M

Action Field

The Action field SHALL express the action that the OTA Provider requests from the OTA Requestor. See [Section 11.20.3.6, “Applying a software update”](#) for a description of the Action values provided in response to an OTA Provider receiving an invocation of this command.

DelayedActionTime Field

The minimum time period the OTA Requestor SHALL wait before executing the Action, in seconds from receipt.

If this field has a value higher than 86400 seconds (24 hours), then the OTA Requestor MAY assume a value of 86400, in order to reduce undue Software Image application delays.

NotifyUpdateApplied Command

This command tells the Provider that the specified update has been applied.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	UpdateToken	octstr	8 to 32			M
1	SoftwareVersion	uint32	all			M

UpdateToken Field

This field SHALL contain the UpdateToken as specified in [Section 11.20.3.6.1, “UpdateToken usage”](#).

SoftwareVersion Field

The SoftwareVersion included in the request payload SHALL provide the same value as the [SoftwareVersion](#) attribute in the invoking OTA Requestor’s [Basic Information Cluster](#), and SHOULD be consistent with the value representing a new version running on the Node invoking the command.

When Generated

The NotifyUpdateApplied command SHOULD be invoked in the following two circumstances:

1. An OTA Requestor has just successfully applied a Software Image it had obtained from a previous [QueryImageResponse](#).
2. An OTA Requestor has just successfully applied a Software Image it had obtained through means different than those of this Cluster.

An OTA Provider MAY use the state of invocation of this command to help track the progress of update for OTA Requestors it knows require a new OTA Software Image. However, due to the possibility that an OTA Requestor MAY never come back (e.g. device removed from Fabric altogether, or a critical malfunction), an OTA Provider SHALL NOT expect every OTA Requestor to invoke this command for correct operation of the OTA Provider.

This command SHALL be considered optional and SHALL NOT result in reduced availability of the OTA Provider functionality if OTA Requestors never invoke this command.

Effect on Receipt

An OTA Provider receiving an invocation of this command MAY log it internally.

On receiving this command, an OTA Provider MAY use the information to update its bookkeeping of cached Software Images, or use it for other similar administrative purposes.

11.20.7. OTA Software Update Requestor Cluster

This cluster implements the Requestor role in the OTA process.

11.20.7.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.20.7.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	OTAR

11.20.7.3. Cluster ID

ID	Name
0x002A	OTA Software Update Requestor

11.20.7.4. Data Types

AnnouncementReasonEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	SimpleAnnouncement	An OTA Provider is announcing its presence.	M
1	UpdateAvailable	An OTA Provider is announcing, either to a single Node or to a group of Nodes, that a new Software Image MAY be available.	M

Value	Name	Summary	Conformance
2	UrgentUpdateAvailable	An OTA Provider is announcing, either to a single Node or to a group of Nodes, that a new Software Image MAY be available, which contains an update that needs to be applied urgently.	M

SimpleAnnouncement Value

An OTA Provider is announcing its presence, but there is no implication that an OTA Requestor would have a new Software Image available if it queried immediately.

UpdateAvailable Value

An OTA Provider is announcing, either to a single Node or to a group of Nodes, that a new Software Image MAY be available. The details may only be obtained by executing a [OTA Software Update Query](#) procedure. A receiving OTA Requestor SHOULD only query the indicated OTA Provider at the ProviderLocation at its next upcoming OTA Provider query.

UrgentUpdateAvailable Value

An OTA Provider is announcing, either to a single Node or to a group of Nodes, that a new Software Image MAY be available, which contains an update that needs to be applied urgently. The details may only be obtained by executing a [OTA Software Update Query](#) procedure. A receiving OTA Requestor SHOULD query the indicated OTA Provider at the ProviderLocation after a random jitter delay between 1 and 600 seconds. This particular reason SHOULD only be employed when an urgent update is available, such as an important security update, or just after initial commissioning of a device, to assist OTA Requestors in more rapidly obtaining updated software.

UpdateStateEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unknown	Current state is not yet determined.	M
1	Idle	Indicate a Node not yet in the process of software update.	M
2	Querying	Indicate a Node in the process of querying an OTA Provider.	M

Value	Name	Summary	Conformance
3	DelayedOnQuery	Indicate a Node waiting after a Busy response.	M
4	Downloading	Indicate a Node currently in the process of downloading a software update.	M
5	Applying	Indicate a Node currently in the process of verifying and applying a software update.	M
6	DelayedOnApply	Indicate a Node waiting caused by AwaitNextAction response.	M
7	RollingBack	Indicate a Node in the process of recovering to a previous version.	M
8	DelayedOnUserConsent	Indicate a Node is capable of user consent.	M

Unknown Value

This value SHALL indicate that the current state is not yet determined. Nodes SHOULD attempt a better state reporting.

Idle Value

This value SHALL indicate a Node not yet in the process of software update, for example because it is awaiting the moment when a query will be made.

Querying Value

This value SHALL indicate a Node in the process of querying an OTA Provider with [QueryImage command](#), including during the process of awaiting a response to that command.

DelayedOnQuery Value

This value SHALL indicate a Node waiting because it received a prior [QueryImageResponse](#) with a [Status](#) field indicating Busy.

Downloading Value

This value SHALL indicate a Node currently in the process of downloading a software update.

Applying Value

This value SHALL indicate a Node currently in the process of verifying and applying a software

update.

DelayedOnApply Value

This value SHALL indicate a Node waiting because it received a prior [ApplyUpdateResponse](#) with an [Action](#) field set to AwaitNextAction.

RollingBack Value

This value SHALL indicate a Node in the process of recovering to a previous version from a new version that was applied, but that could not remain in force, for reasons such as invalid data detected on boot, or significant runtime issues such as reboot loops. Eventually, the next state seen SHOULD be Unknown or Idle.

DelayedOnConsent Value

This value SHALL indicate a Node is capable of [obtaining user consent](#) through its own means, but is currently awaiting that consent after having determined from a prior [QueryImageResponse](#) that an update was available.

ChangeReasonEnum Type

The ChangeReasonEnum Data Type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Unknown	The reason for a state change is unknown.	M
1	Success	The reason for a state change is the success of a prior operation.	M
2	Failure	The reason for a state change is the failure of a prior operation.	M
3	TimeOut	The reason for a state change is a time-out.	M
4	DelayByProvider	The reason for a state change is a request by the OTA Provider to wait.	O

Unknown Value

This value SHALL indicate that the reason for a state change is unknown.

Success Value

This value SHALL indicate that the reason for a state change is the success of a prior operation.

Failure Value

This value SHALL indicate that the reason for a state change is the failure of a prior operation.

TimeOut Value

This value SHALL indicate that the reason for a state change is a time-out condition as determined by the OTA Requestor.

DelayByProvider Value

This value SHALL indicate that the reason for a state change is a request by the OTA Provider to await for a delay.

ProviderLocation Type

This structure encodes a [fabric-scoped](#) location of an OTA provider on a given fabric.

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	ProviderNodeID	node-id					M
2	Endpoint	endpoint-no					M

ProviderNodeID Field

This field SHALL contain the [Node ID](#) of the OTA Provider to contact within the Fabric identified by the FabricIndex.

Endpoint Field

This field SHALL contain the endpoint number which has the OTA Provider device type and OTA Software Update Provider cluster server on the ProviderNodeID. This is provided to avoid having to do discovery of the location of that endpoint by walking over all endpoints and checking their [Descriptor Cluster](#).

11.20.7.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	DefaultOTAProviders	list[ProviderLocation]	desc	N	[]	RW F VA	M
0x0001	UpdatePossible	bool	all		True	R V	M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0002	UpdateState	UpdateStateEnum	all		Unknown	R V	M
0x0003	UpdateStateProgress	uint8	0 to 100	X	null	R V	M

DefaultOTAProviders Attribute

This field is a list of [ProviderLocation](#) whose entries SHALL be set by Administrators, either during Commissioning or at a later time, to set the [ProviderLocation](#) for the default OTA Provider Node to use for software updates on a given Fabric.

There SHALL NOT be more than one entry per Fabric. On a list update that would introduce more than one entry per fabric, the write SHALL fail with CONSTRAINT_ERROR status code.

Provider Locations obtained using the [AnnounceOTAProvider](#) command SHALL NOT overwrite values set in the DefaultOTAProviders attribute.

UpdatePossible Attribute

This field SHALL be set to True if the OTA Requestor is currently able to be updated. Otherwise, it SHALL be set to False in case of any condition preventing update being possible, such as insufficient capacity of an internal battery. This field is merely informational for diagnostics purposes and SHALL NOT affect the responses provided by an OTA Provider to an OTA Requestor.

UpdateState Attribute

This field SHALL reflect the current state of the OTA Requestor with regards to obtaining software updates. See [Section 11.20.7.4.2, “UpdateStateEnum Type”](#) for possible values.

This field SHOULD be updated in a timely manner whenever OTA Requestor internal state updates.

UpdateStateProgress Attribute

This field SHALL reflect the percentage value of progress, relative to the current [UpdateState](#), if applicable to the state.

The value of this field SHALL be null if a progress indication does not apply to the current state.

A value of 0 SHALL indicate that the beginning has occurred. A value of 100 SHALL indicate completion.

This field MAY be updated infrequently. Some care SHOULD be taken by Nodes to avoid over-reporting progress when this attribute is part of a subscription.

11.20.7.6. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	AnnounceOTAProvider	client ⇒ server	Y	A	O

AnnounceOTAProvider Command

This command MAY be invoked by Administrators to announce the presence of a particular OTA Provider.

This command SHALL be [scoped](#) to the [accessing fabric](#).

If the [accessing fabric index](#) is 0, this command SHALL fail with an UNSUPPORTED_ACCESS status code.

Access Modifier: Fabric Scoped						
ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ProviderNodeID	node-id				M
1	VendorID	vendor-id				M
2	AnnouncementReason	AnnouncementReasonEnum				M
3	MetadataForNode	octstr	max 512			O
4	Endpoint	endpoint-no				M

ProviderNodeID Field

This field SHALL contain the [Node ID](#) of a Node implementing the OTA Provider cluster server, on the [accessing fabric](#).

VendorID Field

This field SHALL contain the assigned Vendor ID of the Node invoking this command, as it would appear in that Node's [Basic Information Cluster VendorID](#) attribute.

AnnouncementReason Field

This field SHALL contain a value expressing the reason for the announcement.

MetadataForNode Field

This optional field, if present, SHALL consist of a [top-level anonymous list](#); each list element SHALL have a [profile-specific](#) tag encoded in [fully-qualified](#) form. Each list element SHALL contain a manufacturer-specific payload, which the Node invoking this command wants to expose to the receiving

Node. This payload MAY be used for any purpose and SHOULD be as small as practical, especially if invoked to groups, in order to reduce networking burden of these payloads.

This field SHOULD only be included if the sending OTA Provider has knowledge that some recipient can make use of it.

Endpoint Field

This field SHALL contain the endpoint number which has the OTA Provider device type and OTA Software Update Provider cluster server on the ProviderNodeID. This is provided to avoid having to do discovery of the location of that endpoint by walking over all endpoints and checking their [Descriptor Cluster](#).

When Generated

An OTA Provider MAY invoke this command directly to an OTA Requestor, to announce its presence as an OTA Provider on the Fabric.

These announcements, if made, SHOULD be made at most once every 24 hours for any given target Node, to assist OTA Requestors in discovering available OTA Provider resources, unless the AnnouncementReason is UrgentUpdateAvailable, in which case this command MAY be more frequent.

Any invocation SHALL be made with a delay of at least 1 second between invocations from a given OTA Provider, to reduce burden on the networking infrastructure and affect a form of serialized jitter. It is RECOMMENDED to offset the first announcement of a round (i.e. new set of announcements after a previous complete set) by a random delay time with a distribution span of ≥ 60 seconds to jitter announcement schedules over time.

Effect on Receipt

On receipt of this command, an OTA Requestor SHOULD consider the new ProviderNodeID and AnnouncementReason to possibly query for new software sooner than it would have with its default behavior.

The OTA Requestor SHOULD NOT update entries in the [DefaultOTAProviders](#) list based on announcements.

The receiving Node MAY ignore the content of the announcement if it is unable or unwilling to further query OTA Providers temporarily, or if its provider list is full. If the announcement is ignored, the response SHOULD be SUCCESS.

Depending on the value of the AnnouncementReason field, the OTA Requestor MAY have to query the OTA Provider. See [Section 11.20.7.6.1.3, “AnnouncementReason Field”](#) for the different values and their meaning.

If present, the [MetadataForNode](#) field's MAY be used by a receiving OTA Requestor in any way it deems satisfactory. The MetadataForNode field SHOULD be empty under most normal operational circumstance, but can be useful in environments such as field trials or integration test environments to hint at additional capabilities which OTA Requestors MAY use in a particular Vendor-spe-

cific context.

11.20.7.7. Events

ID	Name	Priority	Access	Conformance
0x00	StateTransition	INFO	V	M
0x01	VersionApplied	CRITICAL	V	M
0x02	DownloadError	INFO	V	M

StateTransition Event

This event SHALL be generated when a change of the [UpdateState](#) attribute occurs due to an OTA Requestor moving through the states necessary to query for updates.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	PreviousState	UpdateStateEnum			Unknown	M
1	NewState	UpdateStateEnum				M
2	Reason	ChangeReasonEnum				M
3	TargetSoftwareVersion	uint32		X	null	M

PreviousState Field

This field SHALL be set to the state that preceded the transition causing this event to be generated, if such a state existed. If no previous state exists, the value SHALL be Unknown.

NewState Field

This field SHALL be set to the state now in effect through the transition causing this event to be generated.

Reason Field

This field SHALL be set to the reason why this event was generated.

TargetSoftwareVersion Field

This field SHALL be set to the target SoftwareVersion which is the subject of the operation, whenever the NewState is Downloading, Applying or RollingBack. Otherwise TargetSoftwareVersion SHALL be null.

VersionApplied Event

This event SHALL be generated whenever a new version starts executing after being applied due to a software update. This event SHOULD be generated even if a software update was done using means outside of this cluster.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Software-Version	uint32				M
1	ProductID	uint16				M

SoftwareVersion Field

This field SHALL be set to the same value as the one available in the [Software Version attribute](#) of the [Basic Information Cluster](#) for the newly executing version.

ProductID Field

This field SHALL be set to the [ProductID](#) applying to the executing version, as reflected by the [Basic Information Cluster](#). This can be used to detect a product updating its definition due to a large-scale functional update that may impact aspects of the product reflected in the [DeviceModel schema](#) of the Distributed Compliance Ledger.

DownloadError Event

This event SHALL be generated whenever an error occurs during OTA Requestor download operation.

The data of this event SHALL contain the following information:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Software-Version	uint32				M
1	BytesDownloaded	uint64				M
2	ProgressPercent	uint8	0 to 100	X	null	M
3	Platform-Code	int64		X	null	M

SoftwareVersion Field

This field SHALL be set to the value of the SoftwareVersion being downloaded, matching the [SoftwareVersion](#) field of the QueryImageResponse that caused the failing download to take place.

BytesDownloaded Field

This field SHALL be set to the number of bytes that have been downloaded during the failing transfer that caused this event to be generated.

ProgressPercent Field

This field SHALL be set to the nearest integer percent value reflecting how far within the transfer the failure occurred during the failing transfer that caused this event to be generated, unless the total length of the transfer is unknown, in which case it SHALL be null.

PlatformCode Field

This field SHOULD be set to some internal product-specific error code, closest in temporal/functional proximity to the failure that caused this event to be generated. Otherwise, it SHALL be null. This event field MAY be used for debugging purposes and no uniform definition exists related to its meaning.

11.21. Over-the-Air (OTA) Software Update File Format

11.21.1. Scope & Purpose

The majority of devices will undergo an over-the-air (OTA) software update at some point during their operational lifecycle. It cannot be assumed that the Node responsible for serving OTA updates (OTA Provider) has any specific knowledge about the internals of OTA Requestor Nodes that are receiving OTA updates. This section provides a standardized header that SHALL be included in all OTA Software Images, in order to provide the necessary information for OTA Providers to validate images available for a given OTA Requestor.

It should be noted that while this specification standardizes an OTA Software Image header that SHALL be used by all OTA Software Images, this specification does not further attempt to standardize the remaining contents of those files.

11.21.2. General Structure

The OTA Software Image file format is composed of a header followed by an opaque body. The header describes general information about the file such as software version, and the [Vendor ID](#) and [Product ID](#) for which the image applies, see [Section 11.20.3.3.2, “Conceptual algorithm for matching OTA Software Images applicable to a query”](#)).

OTA Software Image files SHALL use a fixed encoding. Individual fields of the OTA Software Image file may be comprised of more complex data types that utilize other encoding schemes.

The fields that comprise an OTA Software Image file, listed in the sequential order in which they

SHALL appear, are provided below.

Table 109. OTA Software Image File Layout

Name	Type
FileIdentifier	uint32
TotalSize	uint64
HeaderSize	uint32
Header	see Appendix A, Tag-length-value (TLV) Encoding Format
Payload	N/A

11.21.2.1. FileIdentifier field

The FileIdentifier field is a fixed-width, little-endian-encoded, unsigned 32-bit value that SHALL be included at the beginning of all OTA software image files in order to quickly identify and distinguish the file as being of that format, without having to examine the contents of the whole file. This helps distinguishing the file from other file types in storage. The fixed constant value is defined to be 0x1BEEF11E.

11.21.2.2. TotalSize field

The TotalSize field is a fixed-width, little-endian-encoded, unsigned 64-bit value that SHALL indicate the total size, in bytes, of the entire file, including all fields and Payload. This field SHALL match the total stored size of the file. It SHALL match the sum of:

- The sizes of the FileIdentifier, TotalSize, and HeaderSize fields.
- The value of the HeaderSize field of this header.
- The value of the PayloadSize field of the Header subfield.

For example, given:

- Size of total Header structure is 105 bytes, reflected as HeaderSize = 105
- Payload size is 128KiB = 128 * 1024 bytes = 131072 bytes, reflected as Header.PayloadSize = 131072

Then the TotalSize would be the sum of:

- the size of the FileIdentifier, TotalSize, and HeaderSize fields: 4 + 8 + 4 = 16 bytes
- the HeaderSize value = 105 bytes
- the Header.PayloadSize value = 131072 bytes

This would give a total of 16 + 105 + 131072 = 131193 bytes. The overall file SHALL have this size and no more, and the TotalSize field SHALL contain that value.

11.21.2.3. HeaderSize field

The HeaderSize field is fixed-width, little-endian-encoded, unsigned 32-bit value that SHALL indicate the total size, in bytes, of the TLV-encoded Header field.

11.21.2.4. Header field

The Header is a [TLV](#) structure, encoded with anonymous outer tag, with the following format:

```
ota-image-header-struct => STRUCTURE [ tag-order ]
{
  VendorID           [0] : UNSIGNED INTEGER [ range 16bits ],
  ProductID          [1] : UNSIGNED INTEGER [ range 16bits ],
  SoftwareVersion     [2] : UNSIGNED INTEGER [ range 32bits ],
  SoftwareVersionString [3] : STRING [ length 1..64 ],
  PayloadSize        [4] : UNSIGNED INTEGER [ range 64bits ],
  MinApplicableSoftwareVersion [5, optional] : UNSIGNED INTEGER [ range 32bits ],
  MaxApplicableSoftwareVersion [6, optional] : UNSIGNED INTEGER [ range 32bits ],
  ReleaseNotesURL     [7, optional] : STRING [ length 1..256 ],
  ImageDigestType     [8] : UNSIGNED INTEGER [ range 8bits ],
  ImageDigest         [9] : OCTET STRING [ length 0..64 ]
}
```

VendorID field

The VendorID field SHALL be used by an OTA Provider to determine if a Node is the intended recipient of the OTA software update file by checking that the VendorID field in the OTA software update file matches the [VendorID](#) received in the [QueryImage](#) command from the OTA Requestor. This VendorID field MAY be zero, in which case this OTA software update file MAY apply to more than one vendor.

ProductID field

The ProductID field MAY be used by an OTA Provider to determine if a Node is the intended recipient of the OTA software update file by checking that the ProductID field in the OTA software update file matches the [ProductID](#) received in the [QueryImage](#) command from the OTA Requestor. This ProductID field MAY be zero, in which case this OTA software update file MAY apply to more than one product.

SoftwareVersion field

The SoftwareVersion field SHALL contain the version for the software contained within the file. A larger value of SoftwareVersion is newer than a lower value. The SoftwareVersion value SHOULD NOT be displayed to an end-user.

For a given version, this SoftwareVersion field SHALL match what the Node will report in its [SoftwareVersion](#) attribute in the [Basic Information](#) cluster, once executing the version.

SoftwareVersionString field

The SoftwareVersionString field SHALL contain a human readable (displayable) representation of the version for the software contained within the file. The SoftwareVersionString value SHALL NOT be used by an OTA Provider to determine if the OTA software update file contains a newer image than what is currently running on a Node. The SoftwareVersionString value SHOULD be displayed to an end-user when communicating an identification for the software version.

Format constraints for this field SHALL match the constraints of the [SoftwareVersionString](#) attribute in the [Basic Information](#) cluster.

For a given version, this SoftwareVersionString field SHALL match what the Node will report in its [SoftwareVersionString](#) attribute in the [Basic Information](#) cluster, once executing the version.

PayloadSize field

The PayloadSize field SHALL indicate the total size, in bytes, of the payload contained within this OTA software update file, beyond the header. The length of all data beyond the terminating byte of the header structure SHALL be equal to this field's value.

MinApplicableSoftwareVersion field

The MinApplicableSoftwareVersion field, if present, SHALL be used by an OTA Provider to determine if the OTA Software Image is suitable for the Node, by checking that the MinApplicableSoftwareVersion field in the OTA software update file is less than or equal to the [SoftwareVersion](#) received in the [QueryImage](#) command from the OTA Requestor.

MaxApplicableSoftwareVersion field

The MaxApplicableSoftwareVersion field, if present, SHALL be used by an OTA Provider to determine if the OTA Software Image is suitable for the Node, by checking that the MaxApplicableSoftwareVersion field in the OTA software update file is greater than or equal to the [SoftwareVersion](#) received in the [QueryImage](#) command from the OTA Requestor.

ReleaseNotesUrl field

The ReleaseNotesUrl field, if present, SHOULD specify a link to a product specific web page that contains release notes for the OTA software update file. The specified URL SHOULD resolve to a maintained web page available at that URL for the lifetime of the software version's availability. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme. The maximum length of this field is 256 ASCII characters.

ImageDigestType field

The ImageDigestTypeField SHALL contain the algorithm used to compute the [ImageDigest](#) field.

The value of this field SHALL be a supported numerical identifier value from the [IANA Named Information Hash Algorithm Registry](#) [<https://www.iana.org/assignments/named-information/named-information.xhtml#hash-alg>] established as part of [RFC 6920](#). For example, a value of **1** would match the [sha-256](#) identifier, which maps to the SHA-256 digest algorithm per Section 6.2 of [FIPS 180-4](#)

It is RECOMMENDED that a digest algorithm be chosen that has a minimum digest length of 256 bits, such as **sha-256** (ID 1 in the registry).

ImageDigest field

The ImageDigestField SHALL contain the digest of the entire payload of length **PayloadSize** that follows the header. The digest SHALL be computed using the algorithm indicated in the **ImageDigest-Type** field. This digest SHOULD be used by OTA Providers to ensure they have obtained the entire image expected, and that the contents matches the expectations.

11.21.3. Security considerations

The OTA Software Image file format does not specify the mechanisms that an OTA Requestor should use to validate that a software image is valid for itself. Considerations relative to OTA Software Image Signing are presented in [Section 11.20.4.2, “Image Verification”](#).

11.22. Bulk Data Exchange Protocol (BDX)

11.22.1. Overview

Nodes need the ability to transfer "files" between nodes. For example, uploading sensor data or diagnostics data to another node or downloading software update images from a software update server both require such a protocol.

This document defines a Bulk Data Exchange (BDX) protocol, where files are modeled as collections of bytes with some attached metadata. For the purposes of this protocol, files are opaque, and no attempt is made to specify their format. However, the protocol allows for extensible metadata so that higher-level applications can participate in the decision whether to proceed with a requested file transfer.

The Bulk Data Exchange (BDX) protocol has some semantic elements influenced by the Trivial File Transfer Protocol (TFTP) [RFC 1350](#).

One major difference is that TFTP is defined to run over UDP only, while Bulk Data Transfers can proceed over various reliable transports including both TCP and UDP, using the Message Reliability Protocol (see [Section 4.12, “Message Reliability Protocol \(MRP\)”](#)). Therefore, **BlockAck** and **BlockQuery** fulfill the role of application-layer control flow and acknowledgement rather than providing a reliability and retransmission mechanism. The availability of application-level flow control enables highly power-constrained nodes to pace transfers in a way that respects their power limitations.

11.22.2. Terminology

11.22.2.1. BDX Sender

The node that has bulk data to send to another node.

11.22.2.2. BDX Receiver

The node that receives bulk data from a Sender.

11.22.2.3. BDX Initiator

The node that initiates a bulk data transfer. The Initiator of a data transfer can either be the Sender (for "upload", which starts with a **SendInit**) or the Receiver (for "download", which starts with a **ReceiveInit**).

11.22.2.4. BDX Responder

The node that responds to the initiator by either accepting or rejecting the proposed bulk data transfer and choosing parameters of the transfer compatible with those proposed by the Initiator. It can also either be the Sender (for "download", which is when the Responder receives a **ReceiveInit**) or the Receiver (for "upload", which is when the Responder receives a **SendInit**).

11.22.2.5. BDX Synchronous / Asynchronous Modes

Bulk data transfers can operate in synchronous ("driven") or asynchronous mode. When operating in synchronous mode, one party (the Driver) is responsible for controlling the rate at which the transfer proceeds, and each message in the bulk data transfer protocol SHALL be acknowledged before the next message will be sent.

In asynchronous mode, there is no driver and successive messages are freely sent without waiting for **BlockAck** responses from the Receiver. In asynchronous mode, flow control is provided by the underlying transport (e.g. Matter over TCP).

NOTE	Support for the asynchronous mode is provisional.
-------------	---

11.22.2.6. BDX Driver

The node (either Sender or Receiver) that controls the rate at which a synchronous data transfer proceeds by sending **Block** or **BlockQuery** messages to advance the transfer. This facility allows for **intermittently connected devices** operating (e.g. due to battery constraints) to complete a bulk data transfer without requiring them to stay active continuously during the transfer. In every synchronous bulk data transfer, exactly one device acts as Driver, and the transfer consists of a series of request/responses, each one initiated by the Driver.

BDX utilizes features of the underlying message layer to provide reliability and at-most-once delivery semantics. If the transport fails to deliver the message within the specified parameters, the BDX session SHOULD be aborted. For synchronous transfers, if the Driver fails to receive the response to any given request that is not received within a particular application-determined time, it SHOULD abort the session.

11.22.2.7. BDX Follower

The node (either Sender or Receiver) that "follows" the driver in the protocol flow. The protocol defines the followers as devices that can never be sleepy. Follower either receives a **BlockQuery** to send the data upon request from the driver or receives a new **Block** message and acknowledges it with a **BlockAck** message (in synchronous mode).

11.22.2.8. BDX Session

A bulk data transfer Session is a series of messages passed between a Sender and Receiver that begins with the Initiator starting the transfer negotiation, and ends with a **BlockAckEOF** from the Receiver which indicates receipt of all transmitted data and ends the session. All messages in a session SHALL be sent within the scope of a single **Exchange**. Only one bulk data transfer session can be in progress at any time during a **Exchange**.

11.22.3. Protocol Opcodes and Status Report Values

11.22.3.1. BDX Protocol Messages

Each message in the BDX protocol is mapped to a unique **Protocol Opcode**, namespaced under the **PROTOCOL_ID_BDX** Protocol ID:

- **Vendor ID** = 0x0000 (Matter Common)
- **Protocol ID** = **PROTOCOL_ID_BDX**

Table 110. BDX Protocol Opcodes

Protocol Opcode	Message
0x01	SendInit
0x02	SendAccept
0x03	Reserved for future use
0x04	ReceiveInit
0x05	ReceiveAccept
0x06 ... 0x0F	Reserved for future use
0x10	BlockQuery
0x11	Block
0x12	BlockEOF
0x13	BlockAck
0x14	BlockAckEOF
0x15	BlockQueryWithSkip

11.22.3.2. BDX Status Codes

The list of status codes used in **StatusReport** messages to signify a reason for failing or rejecting a transfer is provided in **Table 111, “BDX Status reports”**. For **StatusReport** messages, the protocol needs to be defined as the BDX protocol, i.e. **StatusReport(GeneralCode: FAILURE, ProtocolId: {VendorID=0x0000, ProtocolId=BDX}, ProtocolCode: <value>)**.

Table 111. BDX Status reports

Status code value	Error	Description
0x0012	LENGTH_TOO_LARGE	Definite length too large to support. For example, trying to SendInit with too large of a file.
0x0013	LENGTH_TOO_SHORT	Definite length proposed for transfer is too short for the context based on the responder's knowledge of expected size.
0x0014	LENGTH_MISMATCH	Pre-negotiated size of transfer was not fulfilled prior to BlockAckEOF.
0x0015	LENGTH_REQUIRED	Responder can only support proposed transfer if definite length is provided.
0x0016	BAD_MESSAGE_CONTENTS	Received a malformed protocol message.
0x0017	BAD_BLOCK_COUNTER	Received block counter out of order from expectation.
0x0018	UNEXPECTED_MESSAGE	Received a well-formed message that was contextually inappropriate for the current state of the transfer.
0x0019	RESPONDER_BUSY	Responder is too busy to proceed with a new transfer at this moment.
0x001F	TRANSFER_FAILED_UNKNOWN_ERROR	Other error occurred, such as perhaps an input/output error occurring at one of the peers.
0x0050	TRANSFER_METHOD_NOT_SUPPORTED	Received a message that mismatches the current transfer mode.
0x0051	FILE_DESIGNATOR_UNKNOWN	Attempted to request a file whose designator is unknown to the responder.
0x0052	START_OFFSET_NOT_SUPPORTED	Proposed transfer with explicit start offset is not supported in current context.
0x0053	VERSION_NOT_SUPPORTED	Could not find a common supported version between initiator and responder.
0x005F	UNKNOWN	Other unexpected error.

The following **StatusReport** message **ProtocolCode** values MAY occur at any time in response to any BDX message:

- **UNEXPECTED_MESSAGE**: When a peer receives a well-formed message that was contextually inappropriate for the current state of the transfer. For example, receiving a **Block** message before a **SendAccept** message, or other logical inconsistencies.
- **BAD_MESSAGE_CONTENTS**: When a peer receives a malformed protocol message.
- **TRANSFER_FAILED_UNKNOWN_ERROR**: Other error occurred, such as perhaps an input/output error occurring at either peer.

- **UNKNOWN**: Internal error beyond the usual error handling.

If any such **StatusReport** message is received, or any other unexpected **StatusReport** is received, the receiving peer SHALL terminate its processing of the transfer and invalidate the exchange.

11.22.4. Security and Transport Constraints

In order to maintain data-in-transit confidentiality, and ensure authenticated message flows, the BDX protocol SHALL only be executed over **PASE** or **CASE** encrypted session. This is enforced by the fact that the only messages allowed to be transmitted without message security are the actual PASE and CASE session establishment messages.

The BDX protocol MAY be carried over any supported Matter messaging transport, such as BTP, **PAFTP**, **NTL**, TCP or MRP, as long as the messages appear in a PASE or CASE session.

Furthermore, the BDX protocol relies on transport-level reliability. Therefore, BDX SHALL always be used over reliable transports. For example, usage with Matter messaging over UDP without MRP reliability, that is, without using the **R Flag** in the **Exchange Flags**, would prevent the necessary reliability.

NOTE

Since using BDX is likely to be used to transfer a considerable amount of data which will take some time, it might not be realistic to use it over **NTL** since it requires steady close proximity between both ends.

11.22.5. Transfer Management Messages

11.22.5.1. **SendInit** and **ReceiveInit** Messages

A **SendInit** message is sent by an Initiator to propose a BDX Transfer session where the Initiator wants to be a Sender and deliver information to another node.

A **ReceiveInit** message is sent by an Initiator to propose a BDX Transfer session where the Initiator wants to be a Receiver and obtain information from another node.

Any BDX transfer exchange begins with one of these two messages.

Both **SendInit** and **ReceiveInit** initiate a negotiation for a number of parameters:

- the Initiator (sender of the message) proposes the transfer parameters and the block size;
- the Responder (receiver of the message) responds with a set of parameters compatible with the Initiator's proposal, according to the Responder's subset of capabilities in common with the Initiator.

The Responder SHALL indicate all the supported modes of operation applicable to the session by replying, upon success, with **SendAccept** (in response to **SendInit**) or **ReceiveAccept** (in response to **ReceiveInit**).

The parameters in the **SendAccept/ReceiveAccept** message SHALL be used in the transfer. If those parameters are unacceptable to the Initiator, it SHALL abort the transfer with an appropriate error

per [Table 111, “BDX Status reports”](#).

Possible replies for **SendInit** (See [Section 11.22.7, “Synchronous Transfers Message Flows”](#) for examples):

- Success: **SendAccept**, if the transfer is accepted by the Responder
- Failure: **StatusReport**(GeneralCode: **FAILURE**, ProtocolId: **BDX**, ProtocolCode: **<error-specific>**), if the transfer is rejected by the Responder, or anything that prevents the Responder from being able to proceed with the transfer as requested.
 - **FILE_DESIGNATOR_UNKNOWN**: The **file designator** field was present and contained a file designator not supported by the responder.
 - **START_OFFSET_NOT_SUPPORTED**: The **start offset** field contained an invalid start offset, or presence of start offset indicated by **RC[STARTOFS]** is not supported by the responder.
 - **LENGTH_TOO_LARGE**: The **definite length** field was present, but too large for the responder.
 - **LENGTH_TOO_SHORT**: The **definite length** field was present, but contextually too short based on the responder’s knowledge of expected size.
 - **LENGTH_REQUIRED**: The responder can only support the proposed transfer if the **definite length** field is provided.
 - **TRANSFER_METHOD_NOT_SUPPORTED**: The responder does not support the proposed transfer control method.
 - **RESPONDER_BUSY**: The responder is too busy to process another transfer. An initiator SHOULD wait at least 60 seconds before attempting to initiate a new **SendInit** with this responder.

Possible replies for **ReceiveInit** (See [Section 11.22.7, “Synchronous Transfers Message Flows”](#) for examples):

- Success: **ReceiveAccept**, if the transfer is accepted by the Responder
- Failure: **StatusReport**(GeneralCode: **FAILURE**, ProtocolId: **PROTOCOL_ID_BDX**, ProtocolCode: **<error-specific>**), if the transfer is rejected by the Responder, or anything that prevents the Responder from being able to proceed with the transfer as requested:
 - **FILE_DESIGNATOR_UNKNOWN**: The **file designator** field was present and contained a file designator not found or supported by the responder.
 - **START_OFFSET_NOT_SUPPORTED**: The **start offset** field contained an invalid start offset, or presence of start offset indicated by **RC[STARTOFS]** is not supported by the responder.
 - **LENGTH_TOO_LARGE**: The **definite length** field was present, but too large for the responder.
 - **LENGTH_REQUIRED**: The responder can only support the proposed transfer if the **definite length** field is provided.
 - **TRANSFER_METHOD_NOT_SUPPORTED**: The responder does not support the proposed transfer control method.
 - **RESPONDER_BUSY**: The responder is too busy to process another transfer. An initiator SHOULD wait at least 60 seconds before attempting to initiate a new **ReceiveInit** with this responder.

The format of the **SendInit** and **ReceiveInit** messages is enumerated in Table 112, “**SendInit/ReceiveInit message fields**”.

Table 112. *SendInit/ReceiveInit message fields*

Field	Size	Notes
Message ID: SendInit (0x01), ReceiveInit (0x04)		
Proposed Transfer Control (PTC)	1 octet	
Range Control (RC)	1 octet	
Proposed Max Block Size (PMBS)	2 octets	Unsigned little-endian
Start Offset (STARTOFS)	4/8 octets	Optional. • Present if RC[STARTOFS] = 1 • Size = 4 if RC[WIDERANGE] = 0 • Size = 8 if RC[WIDERANGE] = 1
Proposed Max Length (LEN)	4/8 octets	Optional. • Present if RC[DEFLN] = 1 • Size = 4 if RC[WIDERANGE] = 0 • Size = 8 if RC[WIDERANGE] = 1
File Designator Length (FDL)	2 octets	Unsigned little-endian
File Designator (FD)	variable length	
Metadata (MDATA)	variable	Optional, TLV.

The following subsections describe the fields composing the **SendInit** and **ReceiveInit** messages.

Proposed Transfer Control (**PTC**)

The Proposed Transfer Control (**PTC**) field specifies, as subfields, the highest version of the protocol and the transfer modes supported by the Initiator of the transfer. All PTC subfields are Initiator-proposed. It is up to the Responder to decide what version and modes it will use. The first version, as of Matter specification 1.0 is BDX Version 0.

At least one of the **PTC**[**RECEIVER_DRIVE**] or **PTC**[**SENDER_DRIVE**] field bits SHALL be set in order for the Responder to set the final transfer control. If neither **PTC**[**RECEIVER_DRIVE**] or **PTC**[**SENDER_DRIVE**] is set, the transfer SHALL be rejected by the Responder.

The **PTC[ASYNC]** bit is an optimization of the transfer and is subject to negotiation with the Responder. In general if the Initiator proposes an ASYNC transfer, it SHOULD also be prepared to accept a synchronous transfer, and SHOULD at least list one of **PTC[RECEIVER_DRIVE]** or **PTC[SENDER_DRIVE]** in the request.

Multiple transfer modes can be specified, which signifies that the Initiator can support any of those transfer modes. For example, if **PTC[ASYNC]**, **PTC[RECEIVER_DRIVE]** and **PTC[SENDER_DRIVE]** bits are set on a SendInit, it indicates that the Initiator supports 3 distinct transfer modes: synchronous Sender drive, synchronous Receiver drive and asynchronous (drive is implied). Only one transfer mode SHALL be used in the actual transfer. The Responder SHALL choose which one to use. If the Responder supports both Sender and Receiver drive, it SHOULD prefer to use Sender drive to retain the request/response semantics.

NOTE Support for the asynchronous mode is provisional and SHALL not be chosen by the Responder.

Table 113. SendInit/ReceiveInit Proposed Transfer Control (PTC) field structure

bit 7	6	5	4	3	2	1	0
RFU	ASYNC	RECEIVER_DRIVE	SENDER_DRIVE	VERSION			

The fields for **PTC** are:

SendInit/ReceiveInit **PTC[VERSION]: proposed protocol version**

- Width: 4 bits
- Values: 0x00-0x0F proposed protocol version

SendInit/ReceiveInit **PTC[SENDER_DRIVE]: sender drive supported**

- Width: 1 bit
- Values:
 - 0 if Sender drive is not supported
 - 1 if Sender drive is supported by Initiator

SendInit/ReceiveInit **PTC[RECEIVER_DRIVE]: receiver drive supported**

- Width: 1 bit
- Values:
 - 0 if Receiver drive is not supported
 - 1 if Receiver drive is supported by Initiator

SendInit/ReceiveInit **PTC[ASYNC]: asynchronous mode supported**

- Width: 1 bit

- Values:
 - 0 if asynchronous mode is not supported
 - 1 if asynchronous mode is supported by Initiator
 - NOTE: Synchronous mode is always implicitly supported.

Range Control (RC)

The Range Control (RC) field specifies parameters related to offset and length of the transfer:

- whether the transfer has a definite length (DEFLEN);
- whether an offset is present (STARTOFS);
- the width of the length and offset fields (WIDERANGE).

Table 114. SendInit/ReceiveInit Range Control (RC) field structure

bit 7	6	5	4	3	2	1	0
RFU			WIDERA NGE	RFU		STARTOF S	DEFLEN

The fields for RC are:

SendInit/ReceiveInit RC[DEFLEN]: definite length present

- Width: 1 bit
 - Values:
 - 0 if no length is present (indefinite length)
 - 1 if a definite length is present (see [Section 11.22.5.1.5](#), “Length (DEFLEN)”)

SendInit/ReceiveInit RC[STARTOFS]: start offset present

- Width: 1 bit
 - Values:
 - 0 if a start offset is not present
 - 1 if a start offset is present

SendInit/ReceiveInit RC[WIDERANGE]: wide (64-bit) range enable for values

- Width: 1 bit
 - Values:
 - 0 to indicate that offset (STARTFOFS) and length (DEFLEN) are 4 octets (32-bit) little-endian unsigned quantities.
 - 1 to indicate that offset (STARTFOFS) and length (DEFLEN) are 8 octets (64-bit) little-endian unsigned quantities.

Proposed Max Block Size (PMBS)

The Proposed Max Block Size (PMBS) field specifies the maximum data size (in bytes) of the block that the Initiator supports, exclusive of block header fields, such as a block counter.

Start Offset (STARTOFS)

The Start Offset (STARTOFS) field is an optional 32-bit/64-bit length that specifies the offset in bytes from start of the file from which the Sender would start the transfer. This allows large file transfers to be segmented into multiple bulk transfer sessions. The RC[STARTOFS] bit indicates whether this start offset is present in the message. The RC[WIDERANGE] bit defines the size of start offset quantity (32-bit or 64-bit). Receivers are not required to accept non-zero start offset transfers. Devices SHOULD make every attempt to support non-zero start offset. If a Responder cannot accept a given start offset, it SHALL reject the SendInit with a StatusReport(GeneralCode: FAILURE, ProtocolId: BDX, ProtocolCode: START_OFFSET_NOT_SUPPORTED). See Table 111, “BDX Status reports”.

Length (DEFLN)

The optional length (DEFLN) field specifies a predetermined definite length for the transfer. For a SendInit message, this field defines the number of bytes the Sender commits to sending to the Receiver. For a ReceiveInit message, this field defines the maximum number of bytes that the Receiver wishes to receive from the Sender. The SendAccept or ReceiveAccept response will then contain the expected length for the transfer. A Receiver receiving a premature BlockEOF would have to consider the transfer failed. A length of 0 or a missing length field signifies an indefinite length. The RC(DEFLN) bit indicates the presence of this field. The RC[WIDERANGE] bit indicates the width of this field.

File Designator Length (FDL)

The File Designator Length (FDL) field is a 16-bit unsigned little-endian field over 2 octets that specifies the length of the upcoming file designator (FD) field, which has variable length.

File Designator (FD)

The File Designator (FD) field is a variable-length identifier chosen by the Initiator to identify the payload to be transferred. In some applications, this identifier will need to be negotiated in advance, so that the Responder will know how to handle the file designator. In other applications, the file designator and optional metadata will be sufficient for the Responder to determine whether to accept the file transfer and how to handle the data, allowing the transfer to proceed without any additional message exchanges.

The length of this field in bytes is provided by the previous File Designator Length (FDL) field.

Metadata (MDATA)

The Metadata (MDATA) field is an optional field, and allows the Initiator to send additional application-specific information about the file to be transferred. The contents of this field are not specified here; applications can use it to avoid needing a separate round-trip negotiation of the file designator, as described above. The TLV metadata consumes the rest of the payload for the SendInit/ReceiveInit message, after all previous fields.

11.22.5.2. SendAccept Message

A **SendAccept** message is sent by the Responder as a response to **SendInit** in order to accept a BDX Transfer session where the Initiator wants to be a Sender and deliver information (upload) to the Responder. The final transfer parameters used are decided by the Responder, given the Initiator proposals in the **SendInit** message.

Responds to (See [Section 11.22.7, “Synchronous Transfers Message Flows”](#) for examples):

- **SendInit** - to accept the proposed transfer.

Possible replies (See [Section 11.22.7, “Synchronous Transfers Message Flows”](#) for examples):

- **Block** - if the Initiator accepts the parameters from the **SendAccept** message, and the transfer method is Sender drive.
- **StatusReport** - If there was another error the Initiator encountered.

Possible follow-ups (See [Section 11.22.7, “Synchronous Transfers Message Flows”](#) for examples):

- **BlockQuery** - if the proposed method is Receiver drive. If the Initiator does not accept the parameters, it SHOULD ignore this and send a **StatusReport(GeneralCode: FAILURE, ProtocolId: BDX, ProtocolCode: TRANSFER_METHOD_NOT_SUPPORTED)** to end the transfer.

11.22.5.3. ReceiveAccept Message

A **ReceiveAccept** message is sent by the Responder as a response to **ReceiveInit** in order to accept a BDX Transfer session where the Initiator wants to be a Receiver and receive information (download) from the Responder. The final transfer parameters used are decided by the Responder, given the Initiator proposals in the **ReceiveInit** message.

Responds to (See [Section 11.22.7, “Synchronous Transfers Message Flows”](#) for examples):

- **ReceiveInit** - to accept the proposed transfer.

Possible replies (See Message Flows for examples):

- **BlockQuery** - if the Initiator accepts the parameters from the **ReceiveAccept** message, and the transfer method is Receiver drive.
- **StatusReport** - If there was another error the Initiator encountered.

Possible follow-ups (See [Section 11.22.7, “Synchronous Transfers Message Flows”](#) for examples):

- **Block** - if the proposed method is Sender drive. If the Initiator does not accept the parameters, it SHOULD ignore this and send a **StatusReport(GeneralCode: FAILURE, ProtocolId: BDX, ProtocolCode: TRANSFER_METHOD_NOT_SUPPORTED)** to end the transfer.

11.22.5.4. SendAccept and ReceiveAccept message fields

The format of the **SendAccept** message is enumerated in [Table 115, “SendAccept message fields”](#).

The format of the **ReceiveAccept** message is enumerated in [Table 116, “ReceiveAccept message](#)

fields”.

Since the meaning of many fields overlap in these messages, they are described only once following the ReceiveAccept message fields.

Table 115. SendAccept message fields

Field	Size	Notes
Message ID: SendAccept(0x02)		
Transfer Control (TC)	1 octet	Must specify at least one of: • TC[RECEIVER_DRIVE] • TC[SENDER_DRIVE] and optionally: • TC[ASYNC] (NOTE: Support for the asynchronous mode is provisional)
Max Block Size (MBS)	2 octets	Unsigned little-endian. Must be $\leq$ PMBS .
Metadata (MDATA)	variable	Optional, TLV.

Table 116. ReceiveAccept message fields

Field	Size	Notes
Message ID: ReceiveAccept(0x05)		
Transfer Control (TC)	1 octet	Must specify exactly one of: • TC[ASYNC] • TC[RECEIVER_DRIVE] • TC[SENDER_DRIVE]
Range Control (RC)	1 octet	
Max Block Size (MBS)	2 octets	Unsigned little-endian. Must be $\leq$ PMBS .
Length (LEN)	4/8 octets	Optional. • Present if RC[DEFLLEN] = 1 • Size = 4 if RC[WIDERANGE] = 0 • Size = 8 if RC[WIDERANGE] = 1
Metadata (MDATA)	variable	Optional, TLV.

The following subsections describe the fields composing the **SendAccept** and **ReceiveAccept** mes-

sages.

Transfer control (TC)

The Transfer Control (TC) field specifies, as subfields, the transfer mode and protocol version.

For the transfer mode (TC[SENDER_DRIVE], TC[RECEIVER_DRIVE]), exactly one mode SHALL be chosen for this transfer, which SHALL be one of the original proposed transfer methods sent by the Initiator.

The version SHALL be the newest version that is supported by the Responder and is not newer than the proposed version sent by the Initiator. If the Responder cannot support a version equal or older to the proposed version, the Responder SHALL send a StatusReport(GeneralCode: FAILURE, ProtocolId: BDX_, ProtocolCode: VERSION_NOT_SUPPORTED) to end the transfer.

The Responder MAY reject a proposed asynchronous transfer (PTC[ASYNC] = 1) by sending a Status-Report(GeneralCode: FAILURE, ProtocolId: BDX, ProtocolCode: TRANSFER_METHOD_NOT_SUPPORTED) to end the transfer. If the Initiator proposed both the PTC[RECEIVER_DRIVE] and PTC[SENDER_DRIVE], the Responder SHALL select exactly one of those options. In that case, to retain the request/response semantics, the Responder SHALL default to TC[SENDER_DRIVE].

Table 117. SendAccept/ReceiveAccept Transfer Control (TC) field structure

bit 7	6	5	4	3	2	1	0
RFU	ASYNC	RECEIVER_DRIVE	SENDER_DRIVE	VERSION			

The fields for TC are:

SendAccept/ReceiveAccept TC[VERSION]: protocol version

- Width: 4 bits
- Values: 0x00-0x0F accepted protocol version

SendAccept/SendAccept TC[SENDER_DRIVE]: sender drive enabled

- Width: 1 bit
- Values:
 - 0 if Sender drive was not chosen
 - 1 if Sender drive was chosen (if this is set, TC[RECEIVER_DRIVE] SHALL be 0)

SendAccept/ReceiveAccept TC[RECEIVER_DRIVE]: receiver drive enabled

- Width: 1 bit
- Values:
 - 0 if Receiver drive was not chosen
 - 1 if Receiver drive was chosen (if this is set, TC[SENDER_DRIVE] SHALL be 0)

SendAccept/ReceiveAccept TC[ASYNC]: asynchronous mode enabled

- Width: 1 bit
- Values:
 - 0 if synchronous mode was chosen for the transfer.
 - 1 if asynchronous mode was chosen for the transfer.
 - NOTE: Synchronous mode is always implicitly supported.
 - NOTE: Support for the asynchronous mode is provisional and SHALL not be chosen by the Responder.

ReceiveAccept Range Control (RC)

NOTE This field is only present in **ReceiveAccept** messages.

The Range Control (RC) field specifies parameters related to offset and length of the transfer:

- whether the transfer has a definite length (**DEFLEN**);
- the width of the length and offset fields (**WIDERANGE**).

Table 118. ReceiveAccept Range Control (RC) field structure

bit 7	6	5	4	3	2	1	0
RFU			WIDERA NGE	RFU			DEFLEN

The fields for **RC** are:

ReceiveAccept RC[DEFLEN]: definite length present

- Width: 1 bit
 - Values:
 - 0 if no length is present (indefinite length)
 - 1 if a definite length is present (see [Section 11.22.5.4.4](#), “Length (LEN)”)

ReceiveAccept RC[WIDERANGE]: wide (64-bit) range enable for values

- Width: 1 bit
 - Values:
 - 0 to indicate that length **LEN** is a 4 octets (32-bit) little-endian unsigned quantity.
 - 1 to indicate that length **LEN** is a 8 octets (64-bit) little-endian unsigned quantity.

Max Block Size (MBS)

The Max Block Size (**MBS**) field specifies the maximum size, in bytes, of each block for this transfer. The maximum whole block payload size will be the sum of the Max Block Size and the size of the block parameters (such as the block counter). This value SHALL be less than or equal to the [pro-](#)

posed max block size (PMBS).

Length (LEN)

NOTE This field is only present in **ReceiveAccept** messages, and only if **RC[DEFLEN] = 1**.

The Length (LEN) field specifies the length of the transfer. If the Initiator indicated a definite length, this length SHALL either be:

- equal to the proposed definite length, if the remaining data in the file beyond the Start Offset (**STARTOFS**) is larger or equal to the proposed length;
- smaller than the proposed definite length, if the remaining data in the file beyond the Start Offset (**STARTOFS**) is smaller than the proposed length.

If this field is present, and the Initiator indicated an indefinite length (i.e. **RC[DEFLEN] = 0** in **ReceiveInit**), this Length field SHALL be equal to the size of the file remaining (if known), or 0, for indefinite. The absence of the Length (LEN) field implies indefinite length.

Metadata (MDATA)

The Metadata (MDATA) field is optional and allows the Responder to provide application-specific metadata about the transfer in **TLV format**. The TLV metadata consumes the rest of the payload for the **SendAccept/ReceiveAccept** message, after all previous fields.

11.22.6. Data Transfer Messages

11.22.6.1. Block Ordering Rules

The following behavior applies to all messages containing a Block Counter field.

Queries (**BlockQuery** message) SHALL be made in ascending and sequential Block Counter order. If the arriving Block Counter at the recipient is not exactly equal to the previous **BlockQuery**'s Block Counter, plus 1 (i.e. next block is being directly asked), a block counter SHALL be considered out-of-order by the recipient.

Blocks (**Block**, **BlockEOF** message) SHALL be sent in ascending and sequential Block Counter order. If the arriving Block Counter at the recipient is not exactly equal to current expected Block Counter, the block counter SHALL be considered out-of-order by the recipient.

Block acknowledgements (**BlockAck**, **BlockAckEOF** messages) SHALL be sent with the same Block Counter as the previously received **Block/BlockEOF**, since they convey a direct acknowledgement. If the arriving Block Counter at the recipient is not exactly equal to the last sent Block Counter, the Block Counter SHALL be considered out-of-order by the recipient.

On any out-of-order block counter condition described above, the recipient of the out-of-order message SHALL send a **StatusReport(GeneralCode: FAILURE, ProtocolId: BDX, ProtocolCode: BAD_BLOCK_COUNTER)** and abort the transfer. On receiving a **StatusReport(GeneralCode: FAILURE, ProtocolId: BDX, ProtocolCode: BAD_BLOCK_COUNTER)**, the recipient SHALL abort the transfer.

Since Block Counter fields are always 32-bit unsigned integer values, but file sizes may be specified

over 64-bit lengths to support very large transfers, the ordering for a "next block counter" SHALL use modulo 2^{32} integer arithmetic. Therefore, if the last Block Counter was 0xFFFF_FFFF, the next expected Block Counter would be 0x0000_0000.

NOTE

Since the data length of blocks does not need to exactly match the Max Block Size in **Block/BlockEOF** messages, it is application-dependent whether Block Counters can be used to statelessly track the offset within a span of initiated transfer. That is, the relationship: $\text{current_offset} = \text{start_offset} + (\text{max_block_size} * \text{block_counter})$ is true **only** if it is contextually known that this usage applies for a given transfer application. This relationship, is SHALL NOT be assumed, since it may not apply, such as when variable-sized blocks are being sent to optimize a given data transfer flow.

11.22.6.2. BlockQuery Message

The **BlockQuery** message is sent by the driving Receiver to the follower to request the next block of data. This message implies a **BlockAck** of the previous block if no **BlockAck** was explicitly sent. The block counter SHOULD start at 0 at the start of the transfer, and increment by one for each subsequent block.

In asynchronous transfers, no **BlockQuery** messages are sent.

The fields of the BlockQuery message are:

Table 119. BlockQuery message fields

Field	Size	Notes
Message ID: BlockQuery(0x10)		
BlockCounter	4 octets	Unsigned little-endian

11.22.6.3. BlockQueryWithSkip Message

The **BlockQueryWithSkip** message MAY be sent by a driving Receiver to the follower to request the next block of data, after skipping an amount forward within the stream. This message implies a **BlockAck** of the previous block if no **BlockAck** was explicitly sent.

This message is semantically equivalent to a **BlockQuery**, but with the following additions:

- Before the next Block is sent by the Sender, the cursor within the underlying data transferred by the Sender SHALL be advanced by **BytesToSkip** bytes.
- If, after skipping **BytesToSkip** bytes, the cursor reaches the end of the file, or beyond, then the next message from the Sender SHALL be a **BlockEOF** with empty contents. In other words, there SHALL be no error indicated when receiving a request to skip past the end of the transferable data.
- The amount of **BytesToSkip** MAY be a size that does not match the current transfer’s maximum block size.

This message enables seeking forward in BDX transfers without requiring the termination of a

transfer followed by the re-opening of a new one with a different **STARTOFS** field specified.

In asynchronous transfers, no **BlockQueryWithSkip** messages are sent.

The fields of the **BlockQueryWithSkip** message are:

Table 120. *BlockQueryWithSkip* message fields

Field	Size	Notes
Message ID: BlockQueryWithSkip(0x15)		
BlockCounter	4 octets	Unsigned little-endian
BytesToSkip	8 octets	Unsigned little-endian

11.22.6.4. Block Message

The **Block** message is the container for the actual bulk transfer data.

Blocks SHALL be sent ascending and sequential block counter order (see [Section 11.22.6.1](#), “**Block Ordering Rules**”). Block data payload length does not need to exactly match the Max Block Size for the transfer.

The fields of the **Block** message are:

Table 121. *Block* message fields

Field	Size	Notes
Message ID: Block(0x11)		
Block Counter	4 octets	Unsigned little-endian
Data	Variable Length	- The data field’s length is that of the remainder of the message payload after the Block Counter field, since Matter messages have definite length. - The length SHALL be in the range $[0 < \text{Length} \leq \text{Max Block Size}]$, where Max Block Size is the negotiated Max Block Size matching the SendAccept / ReceiveAccept message that initiated the transfer.

11.22.6.5. BlockEOF Message

The **BlockEOF** message represents the final block in a data transfer.

Note that, unlike a **Block**, **BlockEOF** MAY have a data length of zero. If the entire transfer fits within the negotiated block size, the **BlockEOF** SHALL be the one and only message sent in the exchange and no **Block** messages will be sent. In that trivial case, the Block Counter would be 0 in the **BlockEOF**.

On receipt of this message, the recipient SHALL verify that the pre-negotiated file size was trans-

ferred, if a definite size had been given. If the Receiver finds a discrepancy between the pre-negotiated size of the file and the amount of data that the Sender has sent, then it MAY consider the transfer failed. In that case, the Receiver MAY respond with a **StatusReport(GeneralCode: FAILURE, ProtocolId: BDX, ProtocolCode: LENGTH_MISMATCH)** message.

Blocks SHALL be sent ascending and sequential block counter order (see [Section 11.22.6.1, “Block Ordering Rules”](#)). Block data payload length does not need to exactly match the Max Block Size for the transfer.

Table 122. BlockEOF message fields

Field	Size	Notes
Message ID: BlockEOF(0x12)		
Block Counter	4 octets	Unsigned little-endian
Data	Variable Length	- The data field’s length is that of the remainder of the message payload after the Block Counter field, since Matter messages have definite length. - The length SHALL be in the range $[0 \leq \text{Length} \leq \text{Max Block Size}]$, where Max Block Size is the negotiated Max Block Size matching the SendAccept / ReceiveAccept message that initiated the transfer. In contrast to the Block message, a length of 0 is permissible to indicate an empty file.

11.22.6.6. BlockAck Message

The **BlockAck** message is an application-level acknowledgement that a **Block** was received, and not necessarily that the Block’s data was stored.

If the Sender is driving in a synchronous transfer, the Receiver SHALL send a **BlockAck** in response to each block of data received. If the Receiver is driving in a synchronous transfer, the Receiver MAY send a **BlockAck** after receipt of a **Block**, and before its next **BlockQuery**. For example, the Receiver SHOULD send a **BlockAck** if it knows it will be going to sleep for some time before the next **BlockQuery**, so that the Sender can free resources associated with the last block.

In asynchronous transfers, no **BlockAck** messages are sent.

The Block Counter in a **BlockAck** SHALL correspond to the Block Counter which was embedded in the **Block** being acknowledged (see [Section 11.22.6.1, “Block Ordering Rules”](#)).

Table 123. BlockAck message fields

Field	Size	Notes
Message ID: BlockAck(0x13)		
BlockCounter	4 octets	Unsigned little-endian

11.22.6.7. BlockAckEOF Message

The **BlockAckEOF** message is sent in response to **BlockEOF** to indicate that the Receiver has received all data.

This message SHALL be sent in both synchronous and asynchronous transfers. This signals the end of the ongoing bulk data transfer session and a successful transfer of the file. The table below lists the contents of the message

The Block Counter in a **BlockAckEOF** SHALL correspond to the Block Counter which was embedded in the **BlockEOF** being acknowledged (see [Section 11.22.6.1, “Block Ordering Rules”](#)).

Table 124. BlockAckEOF message fields

Field	Size	Notes
Message ID: BlockAckEOF(0x14)		
BlockCounter	4 octets	Unsigned little-endian

11.22.7. Synchronous Transfers Message Flows




Figure 99. Version negotiation: both participants support V1 of the protocol




Figure 100. Version negotiation: Initiator supports V2 of the protocol, while Responder supports V1




Figure 101. Synchronous file sending from sleepy device acting as driving Sender


Synchronous file sending to sleepy device acting as driving Receiver

Figure 102. Synchronous file sending to sleepy device acting as driving Receiver




Figure 103. Synchronous file receiving by sleepy device acting as driving Receiver


Synchronous file receiving by device acting as driving Receiver, including BlockQueryWithSkip

Figure 104. Synchronous file receiving by device acting as driving Receiver, including BlockQueryWithSkip




Figure 105. Synchronous file receive from sleepy device acting as driving Sender

In the following flow, the Initiator wants to continue downloading (receiving) a file after the first 200 bytes were received in a previous transfer. The entire remaining contents is desired until the end of the file. Therefore, the proposed start offset (**STARTOFS**) is set to the offset of the first byte desired in the first block transferred. The proposed length (**PLEN**) is set to 0 (or omitted) to announce desire to receive as much as is available from the Sender.

During negotiation, the length (**LEN**) field of **ReceiveAccept** is set to the known remaining file size by the Sender. This could have also been omitted in case the Sender could not or would not state the maximal amount of data to read. Observe that block numbering begins at 0 for every transfer, even if the start offset is not 0. Block indices are always 0-based for every transfer.


Re-started range-based receive from sleepy device acting as driving Sender

Figure 106. Re-started receive from sleepy device acting as driving Sender

In the following flow, the Initiator wants to download (receive) only a section of a file, starting at offset 200, and of length 1000.

During negotiation, the length (LEN) field of `ReceiveAccept` is set to match the proposed desired region length by the Sender. The range is fully-specified by the `\[startOffset .. startOffset +`

length] span.




Figure 107. Range-based receive from sleepy device acting as driving Sender

11.22.8. Asynchronous Transfers Message Flows


Asynchronous file sending

Figure 108. Asynchronous file sending




Figure 109. Asynchronous file receiving

11.23. Distributed Compliance Ledger

11.23.1. Scope & Purpose

The Connectivity Standards Alliance’s DCL (Distributed Compliance Ledger) is a distributed store of information used for tracking certification status and Vendor maintained information such as, but not limited to, product name, product description, and firmware URL. This information is cryptographically secured by digital signatures and is made available via Connectivity Standards Alliance-approved synchronized servers or nodes that are geographically distributed.

The Policies, Procedure and Governance of DCL is maintained by Board approved committees that

comprise **Test Certification Oversight Committee** (TCOC) and **Security Advisory Group** (SEC AG).

Distributed Compliance Ledger [<https://github.com/zigbee-alliance/distributed-compliance-ledger>] is an open-source project designed to:

- Act as a data store for device models' information and their compliance status.
- Act as a secure distribution point of device meta data as made available by vendors.
- Act as a secure distribution point of **Product Attestation Authorities certificates**.
- Act as a secure distribution point of **Operational Root CA Certificates (RCAC)**.

The DCL is owned and hosted by members of the Connectivity Standards Alliance in a way that,

- Write access to Ledger is restricted
 - Vendors role can add new device models that belong to the VendorID that is associated with the public key of that vendor. VendorID is associated to the vendor public key during vendor account creation process.
 - Vendor role can update a subset of existing device model information, such as product name, product description, firmware and hardware info. Updates are only allowed if the device is associated with the same vendor account.
 - A TestHouse role can write the test status for each device to the Ledger.
 - A ZBCertificationCenter role can write the confirmation of the Compliance or revoke the Compliance status of a device model to the Ledger.
- Read access from DCL is public
 - Read DeviceModel info, including firmware and hardware versions from the DCL.
 - Read the Device compliance state from the Ledger.
 - Read the **Product Attestation Authorities certificates**.
 - Read the **Operational Root CA Certificates (RCAC)**.

The DCL is cryptographically secure, machine-readable, and distributed. More details about the Ledger can be found [here](https://github.com/zigbee-alliance/distributed-compliance-ledger/blob/master/docs/DCL-Overview.pdf) [<https://github.com/zigbee-alliance/distributed-compliance-ledger/blob/master/docs/DCL-Overview.pdf>].

While the **DCL repository** [<https://github.com/zigbee-alliance/distributed-compliance-ledger/>] functionality MAY contain more features which MAY evolve independently from those described in this section, this section SHALL be the normative source of truth for usage of the DCL within this specification.

The DCL best practice guidelines are available in [Section 13.6.10, “Distributed Compliance Ledger”](#).

11.23.2. Schemas

Ledger data is available in following representational schemas.

- **Vendor Schema**
 - Provide general information about a Vendor such as Company legal name, Preferred brand name associated with VendorID, Landing page URL for vendor, etc.

- **Product Attestation Authority and Intermediate Certificate Schema**
 - Provide a list of **Product Attestation Authorities Certificates** for the approved PAAs.
 - Optionally, provide a list of **Product Attestation Authorities Intermediate Certificates** that have been added into this schema.
- **Operational Trust Anchors Schema**
 - Provide a list of **Operational Root CA Certificates (RCAC)** have been added into this schema.
 - Provide a list of **Operational Intermediate CA Certificates (ICAC)** have been added into this schema.
- **DeviceModel Schema**
 - Provide general information about a device, and the information is shared across all software versions.
 - e.g ProductName, ProductLabel, PartNumber, Commissioning info, etc.
- **DeviceSoftwareVersionModel Schema**
 - Provide software version specific information.
 - e.g Release Notes URL, FirmwareInformation, OTA Software Image URL (OtaUrl), etc.
- **Compliance / Compliance test result Schema**
 - Provide compliance and test result data about a software version.
- **Device Attestation PKI Revocation Distribution Points**
 - Provide the distribution points (URLs) for Device Attestation PKI Revocation

11.23.2.1. SchemaVersion

Each Schema SHALL have a SchemaVersion attribute to indicate the version of the schema. This helps provide backwards and forwards compatibility when adding new fields to or removing fields from an existing schema.

SchemaVersion is a monotonically increasing positive integer indicating the latest available version of the schema in this specification.

When the initial version of any schema is published, the value of this field SHALL be 0. For schemas that have been published prior to the addition of SchemaVersion field, its value SHALL default to 0 for the purpose of interoperability.

When a new revision of any schema is published, this value SHALL monotonically increase by 1.

HardwareVersion and HardwareVersionString are NOT part of the schemas. If there is different software for a VendorID/ProductID/HardwareVersion, then it SHALL be a new VendorID/ProductID.

11.23.3. Vendor Schema

Vendor Schema provides contact information associated with the vendor for a given VendorID. Information in the Vendor schema is populated by the respective vendors as part of creating a vendor account on the DCL.

Name	Type	Constraint	Conformance	Mutable
VendorID	vendor-id	all	M	No
VendorName	string	max 128	M	Yes
CompanyLegal-Name	string	max 256	M	Yes
CompanyPre-ferredName	string	max 256	O	Yes
Creator	string	max 90	M	No
VendorLanding-PageUrl	string	max 256	O	Yes
SchemaVersion	uint32	all	M	Yes

11.23.3.1. VendorID

This field SHALL uniquely identify this Vendor Schema entry and it SHALL match the Vendor's assigned [Vendor ID](#).

11.23.3.2. VendorName

This field SHALL provide a human readable (displayable) name for the product manufacturer associated with this record.

11.23.3.3. CompanyLegalName

The CompanyLegalName field SHALL specify the legal name for the product manufacturer.

11.23.3.4. CompanyPreferredName

The CompanyPreferredName field, if provided, SHALL specify the Preferred Name that MAY be used for display purposes instead of the CompanyLegalName.

11.23.3.5. Creator

This field uniquely identifies the DCL key that was used to register the vendor record in DCL, pursuant to DCL policies.

11.23.3.6. VendorLandingPageUrl

The VendorLandingPageUrl field (when provided) SHALL contain a link to a maintained web page containing more information about the device manufacturer, such as contact information, address, etc. During the lifetime of the products of this manufacturer, the specified URL SHOULD resolve to a maintained web page. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme. The maximum length of this field is 256 ASCII characters.

11.23.3.7. SchemaVersion

The SchemaVersion field value history for this schema is provided below:

Version	Description
0	Initial Release

11.23.4. Product Attestation Authority and Intermediate Certificate Schema

This schema allows approved [Product Attestation Authorities Certificates](#) to be uploaded and made available via DCL queries. Information in this schema is populated by the respective vendors who have approved [Product Attestation Authorities Certificates](#). It allows vendors to upload [Product Attestation Intermediate certificate](#) under following conditions:

- A corresponding [Product Attestation Authorities Certificates](#) exists in DCL that has signed the [Product Attestation Intermediate certificate](#), and
- Both [Product Attestation Authorities Certificates](#) and [Product Attestation Intermediate certificate](#) are associated with the same [Vendor ID](#).

Note that Matter Operational Root CA Certificates always appear in [Matter TLV Certificate](#) format when they appear in cluster data fields. Careful conversions from PEM to Matter TLV Certificate format SHOULD be done whenever a comparison is made.

Name	Type	Constraint	Conformance	Mutable
PEMCert	string	max 102400	M	No
SerialNumber	octstr	20	M	No
Issuer	string	max 1024	O	No
AuthorityKeyID	string	max 256	O	No
RootSubject	string	max 1024	O	No
RootSubjectKeyID	string	max 256	O	No
IsRoot	boolean	all	M	No
Owner	string	max 90	M	No
Subject	string	max 1024	M	No
SubjectAsText	string	max 1024	M	No
SubjectKeyID	string	max 256	M	No
Approvals	list[Grant]	max 12	O	No
Rejects	list[Grant]	max 12	O	No
VID	vendor-id	all	M	No
SchemaVersion	uint32	all	M	Yes

11.23.4.1. PEMCert

This field uniquely identifies a certificate and SHALL contain the body of a certificate that has been added in the DCL. It SHALL be encoded in PEM format. The certificate SHALL respect the format constraints provided for that certificate type.

11.23.4.2. SerialNumber

The field SHALL be used to identify the serial number field in the Matter certificate structure. A Matter certificate follows the same limitation on admissible serial numbers as in [RFC 5280], i.e., that implementations SHALL admit serial numbers up to 20 octets in length, and certificate authorities SHALL NOT use serial numbers longer than 20 octets in length.

11.23.4.3. Issuer

The field SHALL be used to identify the Certificate Authority that issues the certificate. This field SHALL be set only for a [PAI Certificate](#).

11.23.4.4. AuthorityKeyID

The authority key identifier extension provides a means of identifying the public key corresponding to the private key used to sign a Matter certificate. This field SHALL be set only for a [PAI Certificate](#).

11.23.4.5. RootSubject

This field SHALL contain the PAA certificate's **Subject** field, as defined in PAA in [PAA Certificate](#). This is OPTIONAL for [PAA Certificates](#). This is encoded as defined in [Section 6.1, "Certificate Common Conventions"](#).

11.23.4.6. RootSubjectKeyID

This field SHALL uniquely identify the PAA certificate's **Subject Key Identifier** mandatory extension. It is defined in [PAA Certificate](#) and [Operational Root CA Certificates \(RCAC\)](#). This is OPTIONAL for [PAA Certificates](#). This is encoded as defined in [Section 6.1, "Certificate Common Conventions"](#).

11.23.4.7. Owner

This field uniquely identifies the DCL key that was used to register the certificate in DCL, pursuant to DCL policies.

11.23.4.8. IsRoot

This field SHALL signify whether the associated certificate is [PAA Certificate](#).

11.23.4.9. Subject

This field SHALL contain the certificate's **Subject** field. This is OPTIONAL for [PAA Certificates](#). This is encoded as defined in [Section 6.1, "Certificate Common Conventions"](#).

11.23.4.10. SubjectAsText

This field SHALL contain the certificate's **Subject** field, encoded as a human-readable string.

11.23.4.11. SubjectKeyID

This field SHALL uniquely identify the PAA certificate's **Subject Key Identifier** mandatory extension.

sion. This is encoded as defined in [Section 6.1.2, “Key Identifier Extension Constraints”](#).

11.23.4.12. Approvals

This field SHALL contain a list of [Grants](#) representing approvals for the [PAA Certificate](#) admission into DCL. This field SHALL be set only for a [PAA Certificate](#).

11.23.4.13. Rejects

This field SHALL contain a list of [Grants](#) representing rejections for the [PAA Certificate](#) admission into DCL. This field SHALL be set only for a [PAA Certificate](#)

11.23.4.14. VID

This field SHALL uniquely identify this Vendor Schema entry and it SHALL match the Vendor’s assigned [Vendor ID](#).

11.23.4.15. Grant

The Grant schema is used to store information about an approval or rejection.

Name	Type	Constraint	Conformance	Mutable
Address	string	max 90	M	No
Time	epoch-s	all	M	No
Info	string	max 4096	O	No
SchemaVersion	uint32	all	M	Yes

Address

This field SHALL contain the DCL Key address of the entity granting the approval or rejection.

Time

This field SHALL contain the timestamp of the grant and follow the syntax as specified in [epoch-s](#).

Info

This field, when present, SHALL contain additional information or reason for the grant.

SchemaVersion

The SchemaVersion field value history for this schema is provided below:

Version	Description
0	Initial Release

11.23.4.16. SchemaVersion

The SchemaVersion field value history for this schema is provided below:

Version	Description
0	Initial Release

11.23.5. Operational Trust Anchors Schema

This schema allows [Matter Certificates](#) to be uploaded and made available via DCL queries. Information in this schema is populated by the respective vendors who have [Matter Certificates](#). It also allows vendors to upload [Matter Intermediate CA \(ICA\) Certificate](#) under following conditions:

- A corresponding [Operational Root CA Certificates \(RCAC\)](#) exists in DCL that has signed the [Matter Intermediate CA \(ICA\) Certificate](#), and
- Both [Operational Root CA Certificates \(RCAC\)](#) and [ICAC](#) are associated with the same [Vendor ID](#).

Note that Matter Operational Root CA Certificates always appear in [Matter TLV Certificate](#) format when they appear in cluster data fields. Careful conversions from PEM to Matter TLV Certificate format SHOULD be done whenever comparison are made. Each certificate from the this schema SHALL respect the format constraints provided in [Section 6.4.5.3, “Trusted Root CA Certificates”](#) and [Section 6.4.5.2, “Intermediate CA \(ICA\) Certificate”](#).

Name	Type	Constraint	Conformance	Mutable
PEMCert	string	max 102400	M	No
SerialNumber	octstr	20	M	No
Issuer	string	max 1024	O	No
AuthorityKeyID	string	max 256	O	No
RootSubject	string	max 1024	O	No
RootSubjectKeyID	string	max 256	O	No
IsRoot	boolean	all	M	No
CertificationType	CertificationType-Enum	all	M	No
Owner	string	max 90	M	No
Subject	string	max 1024	M	No
SubjectAsText	string	max 1024	M	No
SubjectKeyID	string	max 256	M	No
VID	vendor-id	all	M	No
SchemaVersion	uint32	all	M	Yes

11.23.5.1. PEMCert

This field uniquely identifies a certificate and SHALL contain the body of a certificate that has been added in the DCL. It SHALL be encoded in PEM format. The certificate SHALL respect the format constraints provided for that certificate type.

11.23.5.2. SerialNumber

The field SHALL be used to identify the serial number field in the Matter certificate structure. A Matter certificate follows the same limitation on admissible serial numbers as in [RFC 5280], i.e., that implementations SHALL admit serial numbers up to 20 octets in length, and certificate authorities SHALL NOT use serial numbers longer than 20 octets in length.

11.23.5.3. Issuer

The field SHALL be used to identify the Certificate Authority that issues the certificate. This field SHALL be set only for a [Operational Intermediate CA Certificates\(ICAC\)](#).

11.23.5.4. AuthorityKeyID

The authority key identifier extension provides a means of identifying the public key corresponding to the private key used to sign a Matter certificate. This field SHALL be set only for a [Operational Intermediate CA Certificates\(ICAC\)](#).

11.23.5.5. RootSubject

This field SHALL contain the PAA certificate's **Subject** field, as defined [Operational Root CA Certificates \(RCAC\)](#). This is OPTIONAL for [Operational Root CA Certificates \(RCAC\)](#). This is encoded as defined in [Section 6.1, "Certificate Common Conventions"](#).

11.23.5.6. RootSubjectKeyID

This field SHALL uniquely identify the PAA certificate's **Subject Key Identifier** mandatory extension. It is defined in [PAA Certificate](#) and [Operational Root CA Certificates \(RCAC\)](#). This is OPTIONAL for [Operational Root CA Certificates \(RCAC\)](#). This is encoded as defined in [Section 6.1, "Certificate Common Conventions"](#).

11.23.5.7. Owner

This field uniquely identifies the DCL key that was used to register the certificate in DCL, pursuant to DCL policies.

11.23.5.8. IsRoot

This field SHALL indicate whether the associated certificate is a [Operational Root CA Certificate \(RCAC\)](#).

11.23.5.9. CertificationType

This field SHALL indicate the type of the certificate. This field SHALL have a value from [CertificationTypeEnum](#) reflecting the type of the certificate.

CertificationTypeEnum Type

This data type is derived from enum8 and has its values listed below.

Value	Name	Summary	Conformance	Description
0	DeviceAttestationPKI	DeviceAttestationPKI	M	used for Device Attestation PKI (PAA/PAI)
1	OperationalPKI	OperationalPKI	M	used for Operational PKI (RCAC/ICAC)
2	VIDSignerPKI	VIDSignerPKI	M	used for Vendor ID Verification Signer Certificate (VVSC)

11.23.5.10. Subject

This field SHALL contain the certificate's **Subject** field. This is encoded as defined in [Section 6.1, "Certificate Common Conventions"](#).

11.23.5.11. SubjectAsText

This field SHALL contain the certificate's **Subject** field, encoded as a human-readable string.

11.23.5.12. SubjectKeyID

This field SHALL uniquely identify the PAA certificate's **Subject Key Identifier** mandatory extension. This is encoded as defined in [Section 6.1.2, "Key Identifier Extension Constraints"](#).

11.23.5.13. VID

This field SHALL uniquely identify this Vendor Schema entry and it SHALL match the Vendor's assigned [Vendor ID](#).

11.23.5.14. SchemaVersion

The SchemaVersion field value history for this schema is provided below:

Version	Description
0	Initial Release
1	Introduction of IsVidVerificationSigner

11.23.6. DeviceModel Schema

A unique combination of the [VendorID](#) and [ProductID](#) is used to identify a DeviceModel. The record schema available to vendors to provide general information shared across all software versions of a given product is presented below.

Name	Type	Constraint	Conformance	Mutable
VendorID	vendor-id	all	M	No

Name	Type	Constraint	Conformance	Mutable
ProductID	uint16	all	M	No
DeviceTypeID	devtype-id	all	M	No
ProductName	string	max 128	M	Yes
ProductLabel	string	max 256	M	Yes
PartNumber	string	max 32	M	Yes
DiscoveryCapabilitiesBitmask	map8	0 to 62	desc	No
Commissioning-CustomFlow	uint8	0 to 2	M	No
Commissioning-CustomFlowUrl	string	max 256	desc	Yes
Commissioning-ModeInitial-StepsHint	map32	desc	M	Yes
Commissioning-ModeInitial-StepsInstruction	string	max 1024	desc	Yes
Commissioning-ModeSecondaryStepsHint	map32	desc	M	Yes
Commissioning-ModeSecondaryStepsInstruction	string	max 1024	desc	Yes
Commissioning-FallbackUrl	string	max 256	desc	Yes
Creator	string	max 90	M	No
UserManualUrl	string	max 256	O	Yes
SupportUrl	string	max 256	O	Yes
ProductURL	string	max 256	O	Yes
LsfUrl	string	max 256	O	Yes
LsfRevision	uint16	all	desc	Yes
SchemaVersion	uint32	all	M	Yes
EnhancedSetupFlowOptions	map16	desc	O	Yes
EnhancedSetupFlowTCUrl	string	max 256	desc	Yes

Name	Type	Constraint	Conformance	Mutable
EnhancedSetupFlowTCRevision	uint16	all	desc	Yes
EnhancedSetupFlowTCDigest	string	max 128	desc	Yes
EnhancedSetupFlowTCFileSize	uint32	desc	desc	Yes
EnhancedSetupFlowMaintenanceUrl	string	desc	desc	Yes
IcdUserActive-ModeTriggerHint	map32	desc	desc	Yes
IcdUserActive-ModeTriggerInstruction	string	desc	desc	Yes
FactoryReset-StepsHint	map32	desc	O	Yes
FactoryReset-StepsInstruction	string	max 1024	desc	Yes

11.23.6.1. VendorID

This field SHALL identify the vendor of the product by its [Vendor ID](#) and SHALL match the [VendorID](#) field in the [Basic Information Cluster](#) of a device running the software referenced by this [DeviceModel/DeviceSoftwareVersionModel](#) record.

11.23.6.2. ProductID

This field SHALL identify the [Product ID](#) of the product instance to which a certification declaration, and thus a DCL entry, applies. This field SHALL match the [ProductID](#) field in the [Basic Information Cluster](#) of a device running the software referenced by this [DeviceModel/DeviceSoftwareVersionModel](#) record.

11.23.6.3. DeviceTypeID

DeviceTypeID is the [Primary Device Type](#) identifier for the device. For example, DeviceTypeID is 10 (0x000A), which is the device type identifier for a Door Lock.

11.23.6.4. ProductName

This field SHOULD match the [ProductName](#) field in the [Basic Information Cluster](#) of a device running the software referenced by this DeviceModel record.

11.23.6.5. ProductLabel

This field SHOULD match the **ProductLabel** field in the **Basic Information Cluster** of a device running the software referenced by this DeviceModel record.

11.23.6.6. PartNumber

This field SHALL match the **PartNumber** field in the **Basic Information Cluster** of a device running the software referenced by this DeviceModel record.

Multiple products (and hence PartNumbers) can share a ProductID. For instance, there may be different packaging (with different PartNumbers) for different regions; also different colors of a product might share the ProductID but may have a different PartNumber. In such cases, the choice of a single PartNumber out of the available set of PartNumbers using this ProductID SHALL be used to populate this field.

11.23.6.7. DiscoveryCapabilitiesBitmask

This field SHALL identify the device's available technologies for device discovery which SHALL be encoded as defined in **Table 60, "Discovery Capabilities Bitmask"**. This field SHALL match the **Table 60, "Discovery Capabilities Bitmask"** provided in the **Section 5.1.3, "QR Code"** (if a QR-code or NFC Tag is provided on or with the product). This field SHALL be populated if the CommissioningFallbackUrl is populated, and SHOULD be populated for all other products.

11.23.6.8. CommissioningCustomFlow

This field SHALL identify the **device's commissioning flow** with encoding as described in **Custom Flow**.

11.23.6.9. CommissioningCustomFlowUrl

This field SHALL identify a vendor-specific commissioning URL for the device model when the CommissioningCustomFlow field is set to '2', and MAY be set for other values of CommissioningCustomFlow. See **Section 5.7.3, "Custom Commissioning Flow"** for how this URL is used. During the lifetime of the product, the specified URL SHOULD resolve to a maintained web page. The syntax of this field SHALL follow the syntax as specified in **RFC 1738** and SHALL use the **https** scheme. The maximum length of this field is 256 ASCII characters.

11.23.6.10. CommissioningModeInitialStepsHint

This field SHALL identify a hint for the steps that MAY be used to put a device that has not yet been commissioned into commissioning mode without factory resetting it. This field is a bitmap with values defined in the **Pairing Hint Table**. For example, a value of 1 (bit 0 is set) indicates that a device that has not yet been commissioned will enter Commissioning Mode upon a power cycle.

Devices that implement **Extended Discovery** SHALL reflect this value in the **Pairing Hint** field of **Commissionable Node Discovery** when they have not yet been commissioned.

11.23.6.11. CommissioningModeInitialStepsInstruction

This field SHALL be populated with the appropriate **Pairing Instruction** for those values of CommissioningModeInitialStepsHint, for which the **Pairing Hint Table** indicates a Pairing Instruction (PI)

dependency.

Devices that implement [Extended Discovery](#) SHALL reflect this value in the [Pairing Instruction](#) field of [Commissionable Node Discovery](#) when they have not yet been commissioned.

11.23.6.12. CommissioningModeSecondaryStepsHint

This field SHALL identify a hint for the steps that MAY be used to put a device that has already been commissioned into commissioning mode without factory resetting it. This field is a bitmap with values defined in the [Pairing Hint Table](#). At least bit 2 SHALL be set, to indicate that a current Administrator can be used to put a device that has already been commissioned into commissioning mode (see [Section 5.6.3, “Enhanced Commissioning Method \(ECM\)”](#)). Devices which implement *additional* mechanisms to put a device that has already been commissioned into commissioning mode without factory resetting it SHALL reflect these mechanism by setting the corresponding bits in this field. Devices that implement [Extended Discovery](#) SHALL reflect this value in the [Pairing Hint](#) field of [Commissionable Node Discovery](#) when they have already been commissioned.

11.23.6.13. CommissioningModeSecondaryStepsInstruction

This field SHALL be populated with the appropriate [Pairing Instruction](#) for those values of CommissioningModeSecondaryStepsHint, for which the [Pairing Hint Table](#) indicates a Pairing Instruction (PI) dependency.

Devices that implement [Extended Discovery](#) SHALL reflect this value in the [Pairing Instruction](#) field of [Commissionable Node Discovery](#) when they have already been commissioned.

11.23.6.14. CommissioningFallbackUrl

This field SHALL identify a vendor-specific commissioning-fallback URL for the device model, which can be used by a Commissioner in case commissioning fails to direct the user to a manufacturer-provided mechanism to provide resolution to commissioning issues. See [Section 5.7.5, “Commissioning Fallback Mechanism”](#) for how this URL is used.

During the lifetime of the product, the specified URL SHOULD resolve to a maintained web page. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme. The maximum length of this field is 256 ASCII characters.

11.23.6.15. Creator

This field uniquely identifies the DCL key that was used to register the device model record in DCL, pursuant to DCL policies.

11.23.6.16. UserManualUrl

This field (when provided) SHALL identify a product-specific web page containing a user manual for the device model. During the lifetime of the product, the specified URL SHOULD resolve to a maintained web page. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme. The maximum length of this field is 256 ASCII characters.

11.23.6.17. SupportUrl

This field (when provided) SHALL identify a product specific support web page. During the lifetime

of the product, the specified URL SHOULD resolve to a maintained web page. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme. The maximum length of this field is 256 ASCII characters.

11.23.6.18. ProductURL

This field (when provided) SHALL identify a link to a product specific web page. This field SHALL match the [ProductURL](#) field in the [Basic Information Cluster](#) of a device running the software referenced by this DeviceModel record. During the lifetime of the product, the specified URL SHOULD resolve to a maintained web page. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme. The maximum length of this field is 256 ASCII characters.

11.23.6.19. LsfUrl

This field (when provided) SHALL identify a link to the Localized String File of this product. This field SHALL NOT have a localized string identifier. During the lifetime of the product, the specified URL SHOULD resolve to a maintained Localized String File. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme. The maximum length of this field is 256 ASCII characters.

11.23.6.20. LsfRevision

LsfRevision is a monotonically increasing positive integer indicating the latest available version of Localized String File. Any client can use this field to check whether it has the latest version of the Localized String File cached. When the first version of the Localized String File is published, the value of this field SHOULD be 1. When a new revision of the Localized String File is published, this value SHALL monotonically increase by 1. When a client of this information finds this to be greater than its currently stored LSF revision number, it SHOULD download the latest version of the LSF from the [LsfUrl](#), and update its local value of this field.

This field SHALL be provided if and only if when LsfUrl is provided.

11.23.6.21. SchemaVersion

The SchemaVersion field value history for this schema is provided below:

Version	Description
0	Initial Release

11.23.6.22. EnhancedSetupFlowOptions

This field SHALL identify the configuration options for the Enhanced Setup Flow. This field is a bitmap with values defined in [Enhanced Setup Flow Options Table](#).

11.23.6.23. EnhancedSetupFlowTCUrl

This field (when provided) SHALL identify a link to the Enhanced Setup Flow Terms and Condition File for this product. During the lifetime of the product, the specified URL SHOULD resolve to a

maintained Terms and Conditions File. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#). The maximum length of this field is 256 ASCII characters. All URLs SHALL use the [https](#) scheme. This field SHALL be present if and only if the EnhancedSetupFlowOptions field has bit 0 set.

11.23.6.24. EnhancedSetupFlowTCRevision

This field (when provided) is an increasing positive integer indicating the latest available version of the Enhanced Setup Flow Terms and Conditions file. When a new revision of the File is published, this value SHALL increase (and SHOULD increase by 1). This field SHALL be present if and only if the EnhancedSetupFlowOptions field has bit 0 set.

11.23.6.25. EnhancedSetupFlowTCDigest

This field (when provided) SHALL contain the digest of the entire contents of the associated file downloaded from the EnhancedSetupFlowTCUrl field, encoded in base64 string representation and SHALL be used to ensure the contents of the downloaded file are authentic. The digest SHALL have been computed using the SHA-256 digest algorithm. This field SHALL be present if and only if the EnhancedSetupFlowOptions field has bit 0 set.

11.23.6.26. EnhancedSetupFlowTCFileSize

This field (when provided) SHALL indicate the total size of the Enhanced Setup Flow Terms and Conditions file in bytes, and SHALL be used to ensure the downloaded file size is within the bounds of EnhancedSetupFlowTCFileSize. This field SHALL be provided if and only if the EnhancedSetupFlowTCUrl field is present.

11.23.6.27. EnhancedSetupFlowMaintenanceUrl

This field (when provided) SHALL identify a link to a vendor-specific URL which SHALL provide a manufacturer specific means to resolve any functionality limitations indicated by the TERMS_AND_CONDITIONS_CHANGED status code. This field SHALL be present if the EnhancedSetupFlowOptions field has bit 0 set. During the lifetime of the product, the specified URL SHOULD resolve to a maintained web page. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#). The maximum length of this field is 256 ASCII characters. All URLs SHALL use the [https](#) scheme.

11.23.6.28. IcdUserActiveModeTriggerHint

This field (when provided) is applicable to an [ICD](#) that supports the [UserActiveModeTrigger](#) feature. This field SHALL indicate which user action(s) will trigger the ICD to switch to Active mode. This field SHALL follow the requirements specified in [UserActiveModeTriggerHint](#).

11.23.6.29. IcdUserActiveModeTriggerInstruction

This field (when provided) is applicable to an [ICD](#) that supports the [UserActiveModeTrigger](#) feature. The meaning of this field is dependent upon the [UserActiveModeTriggerHint](#) field value, and the conformance is indicated in the "dependency" column in [UserActiveModeTriggerHint](#) table. This field SHALL follow the requirements specified in [UserActiveModeTriggerInstruction](#).

11.23.6.30. FactoryResetStepsHint

This field (when provided) SHALL identify a hint for the steps that MAY be used to factory reset a device. This field is a bitmap with values defined in the [Pairing/Reset Hint Table](#). For example, a value of 64 (bit 6 is set) indicates that a device will be factory reset when the Reset Button is pressed.

11.23.6.31. FactoryResetStepsInstruction

This field SHALL be populated with the appropriate factory reset instruction for those values of FactoryResetStepsHint, for which the [Pairing/Reset Hint Table](#) indicates a dependency in the Instruction Dependency column.

11.23.7. DeviceSoftwareVersionModel Schema

A unique combination of the VendorID, ProductID and SoftwareVersion is used to identify a Device-SoftwareVersionModel. The record schema available to vendors to provide information specific to a particular software version for a given product is presented below.

Name - (Matching Basic Information Cluster Field)	Type	Constraint	Conformance	Mutable
VendorID	vendor-id	all	M	No
ProductID	uint16	all	M	No
SoftwareVersion	uint32	all	M	No
SoftwareVersion-String	string	1 to 64	M	No
CDVersionNumber	uint16	all	M	No
Creator	string	max 90	M	No
FirmwareInformation	string	max 512	O	No
SoftwareVersion-Valid	boolean	all	M	Yes
OtaUrl	string	max 256	desc	Yes
OtaFileSize	uint64	all	OtaUrl	No
OtaChecksum	string	max 64	OtaUrl	No
OtaChecksumType	uint16	desc	OtaUrl	No
MinApplicable-SoftwareVersion	uint32	all	M	Yes
MaxApplicable-SoftwareVersion	uint32	all	M	Yes
ReleaseNotesUrl	string	max 256	O	Yes

Name - (Matching Basic Information Cluster Field)	Type	Constraint	Conformance	Mutable
SchemaVersion	uint32	all	M	Yes
SpecificationVersion	uint32	all	M	No

11.23.7.1. VendorID

See [Section 11.23.6.1, “VendorID”](#).

11.23.7.2. ProductID

See [Section 11.23.6.2, “ProductID”](#).

11.23.7.3. SoftwareVersion

SoftwareVersion SHALL identify the software version number for the device model consistent with the value found in [Section 11.21.2.4.3, “SoftwareVersion field”](#). The SoftwareVersionNumber value SHOULD NOT be displayed to an end-user. SoftwareVersion is not editable, but it would be possible to create a new device model for the same VendorID/ProductID for different versions. Both SoftwareVersion and SoftwareVersionString SHALL be included. This field SHALL match the **Software-Version** attribute in the [Basic Information Cluster](#) of a device running the software certified by this DeviceModel record.

11.23.7.4. SoftwareVersionString

This field SHALL match the **Software Version String** attribute in the [Basic Information Cluster](#) of a device running the software referenced by this DeviceModel record, including format constraints on that field.

11.23.7.5. CDVersionNumber

CDVersionNumber SHALL identify the CD Version Number of the Certification that applies to this Software Image. The **CDVersionNumber** maps to **version_number** defined in [Certification Elements TLV structure](#).

11.23.7.6. Creator

This field uniquely identifies the DCL key that was used to register the device software version model record in DCL, pursuant to DCL policies.

11.23.7.7. FirmwareInformation

The FirmwareInformation field, if present, SHALL match the **firmware_information** field in [attestation-elements](#) field included in the Device Attestation response when this Software Image boots on the device. It is an OPTIONAL field that MAY be present only for devices that meet the requirements listed in [Section 6.3.2, “Firmware Information”](#).

11.23.7.8. SoftwareVersionValid

This field SHALL indicate whether this SoftwareVersion is valid. When creating an entry for a new SoftwareVersion, this typically is set to True. When the manufacturer later finds out there is some reason that this version should no longer be used (e.g. due to some bugs), the field SHALL be updated to False.

NOTE

This mechanism is for "withdrawal" of a SoftwareVersion by the *manufacturer*, not to be confused with *certification revocation* by the Connectivity Standards Alliance (see [Section 11.23.8.1, "SoftwareVersionCertificationStatus"](#)).

11.23.7.9. OTA

OtaUrl

OtaUrl SHALL identify the URL where to obtain the OTA image. The OtaUrl field SHALL be populated unless the device manufacturer provides alternative means to update the product's firmware. The specified URL SHOULD be available for the relevant lifetime of the corresponding software. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme. The maximum length of this field is 256 ASCII characters.

OtaFileSize

OtaFileSize is the total size of the OTA software image in bytes. This field SHALL be provided if the OtaUrl field is populated.

OtaChecksum

OtaChecksum SHALL contain the digest of the entire contents of the associated OTA Software Update Image under the [OtaUrl](#) field, encoded in base64 string representation. The digest SHALL have been computed using the algorithm specified in [OtaChecksumType](#). This field SHALL be provided if the OtaUrl field is populated.

OtaChecksumType

OtaChecksumType SHALL identify the type of OtaChecksum. This field SHALL be provided if the OtaUrl field is populated.

The value of this field SHALL be a supported numerical identifier value from the [IANA Named Information Hash Algorithm Registry](#) [<https://www.iana.org/assignments/named-information/named-information.xhtml#hash-alg>] established as part of [RFC 6920](#). For example, a value of [1](#) would match the [sha-256](#) identifier, which maps to the SHA-256 digest algorithm per Section 6.2 of [FIPS 180-4](#).

The digest algorithm chosen SHALL have a minimum digest length of 256 bits, such as [sha-256](#) (ID 1 in the registry).

To increase interoperability, OtaChecksumType SHALL be within the list of [1, 7, 8, 10, 11, 12].

11.23.7.10. MinApplicableSoftwareVersion

MinApplicableSoftwareVersion SHALL be equal to the lowest SoftwareVersion for which this image

can be applied. Also see [Section 11.21.2.4.6, “MinApplicableSoftwareVersion field”](#).

11.23.7.11. MaxApplicableSoftwareVersion

MaxApplicableSoftwareVersion SHALL be equal to the highest SoftwareVersion for which this image can be applied. Also see [Section 11.21.2.4.7, “MaxApplicableSoftwareVersion field”](#).

11.23.7.12. ReleaseNotesUrl

ReleaseNotesUrl SHALL identify a product specific web page that contains release notes for the device model software. The specified URL SHOULD resolve to a maintained web page available for the lifetime of the corresponding software being relevant. The syntax of this field SHALL follow the syntax as specified in [RFC 1738](#) and SHALL use the [https](#) scheme. The maximum length of this field is 256 ASCII characters.

11.23.7.13. SchemaVersion

The SchemaVersion field value history for this schema is provided below:

Version	Description
0	Initial Release

11.23.7.14. SpecificationVersion

SpecificationVersion SHALL identify the specification version applicable to the device model. This field SHALL match the [SpecificationVersion](#) attribute in the [Basic Information Cluster](#) of a device running the software certified by this DeviceModel record.

11.23.8. DeviceSoftwareCompliance / Compliance test result Schema

A unique combination of the VendorID, ProductID and SoftwareVersion is used to identify a DeviceSoftwareCompliance record. This record schema is available to the Connectivity Standards Alliance to provide information specific to certification of a particular software version for a given product is presented below. Note that this schema is writable and mutable by the Connectivity Standards Alliance, not by the manufacturer.

Name - (Matching Basic Information Cluster Field)	Type	Constraint	Conformance	Mutable
VendorID	vendor-id	all	M	No
ProductID	uint16	all	M	No
SoftwareVersion	uint32	all	M	No
SoftwareVersion-String	string	1 to 64	M	No
CDVersionNumber	uint16	all	M	Yes

Name - (Matching Basic Information Cluster Field)	Type	Constraint	Conformance	Mutable
SoftwareVersion-CertificationStatus	SoftwareVersion-CertificationStatusEnum	all	M	Yes
CDCertificateID	string	19	M	Yes
Owner	string	max 90	M	No
SchemaVersion	uint32	all	M	Yes

For the description of the first five fields, see the corresponding fields in [DeviceSoftwareVersion-Model Schema](#).

11.23.8.1. SoftwareVersionCertificationStatus

This field SHALL have a value from [SoftwareVersionCertificationStatusEnum](#) reflecting the current certification status of this SoftwareVersion.

11.23.8.2. SoftwareVersionCertificationStatusEnum Type

This data type is derived from enum8 and has its values listed below.

Value	Name	Summary	Conformance	Description
0	dev-test	dev-test	M	used for development and test purposes (These will typically not be placed in DCL)
1	provisional	provisional	M	used for a SoftwareVersion when going into certification testing (These might or might not be placed in DCL, depending on the Connectivity Standards Alliance policy and procedures)
2	certified	certified	M	used for a SoftwareVersion which has been certified

Value	Name	Summary	Conformance	Description
3	revoked	revoked	M	used for a SoftwareVersion which has been revoked

The values 0 through 2 SHALL correspond to the values 0 through 2 used in [certification_type](#) in the [Certification Declaration](#).

11.23.8.3. CDCertificateID

This field SHALL have the Connectivity Standards Alliance certification's certificate ID for the Certification that applies to this record. The value of this field is used in the [Certification Declaration](#)'s [certificate_id](#) field (see [Certification Elements TLV structure](#)) for products using the [VendorID](#), [ProductID](#) and [SoftwareVersion](#) in this schema entry.

11.23.8.4. Owner

This field uniquely identifies the DCL key that was used to register the device software compliance record in DCL, pursuant to DCL policies.

11.23.8.5. SchemaVersion

The SchemaVersion field value history for this schema is provided below:

Version	Description
0	Initial Release

11.23.9. Device Attestation PKI Revocation Distribution Points Schema

The Device Attestation PKI revocation distribution points are used to identify the URL where PAAs and PAIs provide revocation information for PAIs and DACs, respectively.

See [Section 6.2, "Device Attestation"](#) for more details on the PKI elements discussed in this subsection (e.g. PAA, PAI, DAC).

The schema is presented in the table below:

Name	Type	Constraint	Conformance	Mutable
VendorID	vendor-id	all	M	No
ProductID	uint16	all	desc	No
IsPAA	bool	all	M	No
Label	string	max 64	M	No
CRLSignerDelegator	string	max 2048	desc	Yes

Name	Type	Constraint	Conformance	Mutable
CRLSignerCertificate	string	max 2048	M	Yes
IssuerSubjectKeyID	string	max 64	M	No
DataUrl	string	max 256	M	Yes
DataFileSize	uint64	desc	O	Yes
DataDigest	string	max 128	desc	Yes
DataDigestType	uint32	all	desc	Yes
RevocationType	uint32	desc	M	No
SchemaVersion	uint32	all	M	Yes

11.23.9.1. VendorID

This field SHALL indicate the [VendorID](#) associated with the PAA or PAI whose revocation information is provided. For a non-vendor-scoped PAA, this SHALL be the VendorID associated with the organization operating the PAA. For vendor-scoped PAA and for PAIs, this field SHALL contain the VendorID as found in the CRLSignerCertificate (see [Section 6.2.2.2, “Encoding of Vendor ID and Product ID in subject and issuer fields”](#) for encoding).

11.23.9.2. ProductID

This field is only required when a PAI is making the distribution point available. This field SHALL only be provided if the IsPAA field is false and if the CRLSignerCertificate field has a ProductID in its subject (see [Section 6.2.2.2, “Encoding of Vendor ID and Product ID in subject and issuer fields”](#) for encoding).

11.23.9.3. IsPAA

This field SHALL be set to true if the revocation information distribution point relates to a PAA, otherwise it SHALL be set to false (i.e. it relates to a PAI, not a PAA)."

11.23.9.4. Label

This field contains a label to disambiguate multiple revocation information partitions of a particular issuer. Uniqueness within the Device Attestation PKI Revocation Distribution Points schema SHALL be enforced against the tuple containing all of:

- VendorID
- Label
- IssuerSubjectKeyID

Therefore, there MAY be multiple entries for the same VendorID and IssuerSubjectKeyID in case partitioning is done, which are disambiguated by the Label.

Enforcement of uniqueness constraints SHALL be done by the Distributed Compliance Ledger's block transaction processing and SHALL also be done by clients making use of the information from this schema.

11.23.9.5. CRLSignerDelegator

This field SHALL be present and non-empty if all of the following are true:

- Certificate in CRLSignerCertificate field is not self-signed.
- IsPAA is false.
- CRLSignerCertificate is not a PAI.

When present, this field SHALL contain the issuer certificate which signed the CRLSignerCertificate, encoded in X.509v3 PEM format.

Additional constraints related to the value of this field are specified in [Section 11.23.9.6, "CRLSignerCertificate"](#).

11.23.9.6. CRLSignerCertificate

This field SHALL contain the issuer certificate who signed the revocation information that is provided in the distribution point entry, encoded in X.509v3 PEM format.

If the RevocationType is **1** (X.509v3 CRL):

- The certificate's subject public key SHALL be used to validate the **signatureValue** of the **Certifi-cateList** object in the CRL (see [RFC 5280](#) sec 5.1.1.3). These checks SHALL be done by the client making use of the information, since they cannot be done by the Distributed Compliance Ledger implementation which cannot rely on out-of-ledger records or files.
- The certificate's **SubjectKeyIdentifier** mandatory extension SHALL match the **Authority Key Identifier** found in the file at the **DataUrl**.
- The certificate SHALL be checked for correct chaining by the dynamic validation within the Distributed Compliance Ledger's block transaction processing:
 - If the certificate is self-signed, it SHALL be considered as a PAA certificate, the **IsPAA** field SHALL be true, and the certificate SHALL be found in the list of approved PAA certificates currently in force in the Distributed Compliance Ledger using exact byte equality.
 - If the certificate is not self-signed and the **IsPAA** field is set to true, it SHALL be considered as a CRL signer certificate delegated by a PAA. The certificate format SHALL be as specified in [Section 11.23.9.6.1, "PAA-delegated CRLSignerCertificate format"](#). The certificate SHALL be validated using the set of approved PAA certificates currently in force in the Distributed Compliance Ledger and the **IssuerSubjectKeyID** field SHALL match the **SubjectKeyIdentifier** extension value in the matching PAA. Commissioners SHALL verify that the PAA that issued the CRLSignerCertificate is the same PAA that issued the PAI being checked for revocation.
 - If the certificate is not self-signed and the **IsPAA** field is set to false, it SHALL be considered as a PAI certificate or a CRL signer certificate delegated by a PAI. The format of a CRL signer certificate delegated by a PAI SHALL be as specified in [Section 11.23.9.6.2, "PAI-delegated CRLSignerCertificate format"](#). If the CRL signer certificate is not a PAI certificate but is dele-

gated by a PAI, then the PAI certificate SHALL be present in the CRLSignerDelegator field so that the full certificate chain can be verified. The certification path containing the certificate SHALL be validated using the set of approved PAA certificates currently in force in the Distributed Compliance Ledger as the trust store. Commissioners SHALL verify that either the CRLSignerCertificate or the CRLSignerDelegator is the same PAI that issued the DAC being checked for revocation.

PAA-delegated CRLSignerCertificate format

In case the CRLSignerCertificate is a delegate of a PAA for signing CRL for revoking PAI certificates, then the certificate SHALL follow the following constraints layered on top of the encoding specified by [RFC 5280](#) within the [TBSCertificate](#) structure:

1. The [version](#) field SHALL be set to [2](#) to indicate v3 certificate.
2. The [signature](#) field SHALL contain the identifier for signatureAlgorithm [ecdsa-with-SHA256](#).
3. The [issuer](#) field SHALL be a sequence of [RelativeDistinguishedName](#) s.
4. The [issuer](#) field SHALL match the [subject](#) field of the parent PAA certificate.
5. The [subject](#) field SHALL be a sequence of [RelativeDistinguishedName](#) s.
6. A ProductID value SHALL NOT be present in either the [subject](#) or [issuer](#) fields.
7. The algorithm field in [subjectPublicKeyInfo](#) field SHALL be the object identifier for [prime256v1](#).
8. The certificate SHALL carry the following Extensions:
 - a. Basic Constraint extension SHALL be marked [critical](#) and have the [cA](#) field set to FALSE.
 - b. Key Usage extension SHALL be marked [critical](#).
 - i. The [cRLSign](#) bits SHALL be set in the [KeyUsage](#) bitstring
 - ii. The [digitalSignature](#) bit MAY be set in the [KeyUsage](#) bitstring
 - iii. Other bits SHALL NOT be set
 - c. Subject Key Identifier
9. The certificate MAY also carry the following additional Extensions:
 - a. Extended Key Usage
 - b. Authority Key Identifier
 - c. Any other extension allowed in [RFC 5280](#) where inclusion does not violate size limitations. These extensions insofar not defined in this specification SHALL be ignored, but MAY be present to allow flexibility in CA operation.

PAI-delegated CRLSignerCertificate format

In case the CRLSignerCertificate is a delegate of a PAI for signing CRL for revoking DAC certificates, then the certificate SHALL follow the following constraints layered on top of the encoding specified by [RFC 5280](#) within the [TBSCertificate](#) structure:

1. The [version](#) field SHALL be set to [2](#) to indicate v3 certificate.
2. The [signature](#) field SHALL contain the identifier for signatureAlgorithm [ecdsa-with-SHA256](#).

3. The **issuer** field SHALL be a sequence of **RelativeDistinguishedName** s.
4. The **issuer** field SHALL match the **subject** field of the parent PAI certificate.
5. The **subject** field SHALL be a sequence of **RelativeDistinguishedName** s.
6. The algorithm field in **subjectPublicKeyInfo** field SHALL be the object identifier for **prime256v1**.
7. The certificate SHALL carry the following Extensions:
 - a. Basic Constraint extension SHALL be marked **critical** and have the **cA** field set to FALSE.
 - b. Key Usage extension SHALL be marked **critical**
 - i. The **cRLSign** bit SHALL be set in the **KeyUsage** bitstring
 - ii. The **digitalSignature** bit MAY be set in the **KeyUsage** bitstring
 - iii. Other bits SHALL NOT be set
 - c. Authority Key Identifier
 - d. Subject Key Identifier
8. The certificate MAY also carry the following additional Extensions:
 - a. Extended Key Usage
 - b. Any other extension allowed in **RFC 5280** where inclusion does not violate size limitations. These extensions insofar not defined in this specification SHALL be ignored, but MAY be present to allow flexibility in CA operation.

11.23.9.7. IssuerSubjectKeyID

This field SHALL uniquely identify the PAA or PAI for which this revocation distribution point is provided, via the certificate's **SubjectKeyIdentifier** mandatory extension. This field is provided to assist queries without requiring additional certificate parsing. This field SHALL provide the subject key identifier as an even number of uppercase hexadecimal characters (**[0-9A-F]**), with no white-space and no non-hexadecimal characters.

For example, subject key ID **A3:03:13:6D:54:A8:4B:E2:4C:48:87:B3:41:06:6D:C2:70:96:2F:99** (as it would appear in **openssl x509** output, for human consumption) would be recorded as **A303136D54A84BE24C4887B341066DC270962F99**.

When processing revocation information during the device **Device Attestation Procedure**, clients SHALL only use entries whose **IssuerSubjectKeyID** matches a candidate certificate's Authority Key Identifier extension.

11.23.9.8. DataUrl

This field SHALL indicate the URL where to obtain the information in the format indicated by the **RevocationType** field. The syntax of this field SHALL follow the syntax as specified in **RFC 1738**. The maximum length of this field is 256 ASCII characters. All URLs SHALL use either the **http** or **https** scheme.

For entries of **RevocationType** 1 (**RFC 5280** CRL):

- The content located at the **DataUrl** SHALL be a **RFC 5280** Certificate Revocation List document

(see section 5 in the RFC) encoded in DER format.

- Note that conformance to the [RFC 5280](#) CRL profile requires at least the [AuthorityKeyIdentifier](#) and [CRLNumber](#) extensions to be present in the document.
- If multiple entries exist in the schema, which match in VendorID and IssuerSubjectKeyID, then:
 - All entries SHALL have a different Label.
 - All entries SHALL have a different DataUrl.
 - The content located at the [DataURL](#) SHALL have the Issuing Distribution Point critical CRL extension present (see [RFC 5280](#) section 5.2.5).
 - DCL software SHALL validate the uniqueness of Label and DataUrl of each entry.
 - Note that the content located at the [DataURL](#) will never be accessed or validated by DCL software which cannot rely on out-of-ledger records or files.
 - Based on the Issuing Distribution Point critical CRL extension (see [RFC 5280](#) section 5.2.5), the following validation SHALL be applied:
 - The [distributionPoint](#) field of the extension SHALL have a single [GeneralName](#) of type [uniformResourceIdentifier](#).
 - The [uniformResourceIdentifier](#) SHALL match exactly, byte-for-byte, the value found in the [DataUrl](#) field of this schema entry.
 - If the value of the [CRLSignerCertificate](#) is a CRL signer certificate delegated by a PAA or delegated by a PAI, then the [indirectCRL](#) field of the extension SHALL be set to true. Otherwise, the [indirectCRL](#) field of the extension SHALL be set to false.

11.23.9.9. DataFileSize

This field, if present, SHALL indicate the total size in bytes of the file found at the DataUrl. This field SHALL be omitted if the RevocationType is 1, which refers to a type having built-in signatures.

11.23.9.10. DataDigest

This field, if present, SHALL contain the digest of the entire contents of the associated file downloaded from the [DataUrl](#) field, encoded in base64 string representation. The digest SHALL have been computed using the algorithm specified in [DataDigestType](#). This field SHALL be present if and only if the DataFileSize field is present.

11.23.9.11. DataDigestType

This field, if present, SHALL indicate the type of digest used in the DataDigest field. This field SHALL be provided if and only if the DataDigest field is present.

The value of this field SHALL be a supported numerical identifier value from the [IANA Named Information Hash Algorithm Registry](#) [<https://www.iana.org/assignments/named-information/named-information.xhtml#hash-alg>] established as part of [RFC 6920](#). For example, a value of **1** would match the [sha-256](#) identifier, which maps to the SHA-256 digest algorithm per Section 6.2 of [FIPS 180-4](#).

The digest algorithm chosen SHALL have a minimum digest length of 256 bits, such as [sha-256](#) (ID 1 in the registry).

To increase interoperability, `DataDigestType`, if present, SHALL be within the list of [1, 7, 8, 10, 11, 12].

11.23.9.12. RevocationType

This field SHALL identify the type of file found at the `DataUrl` for this entry. Values supported are:

- 0: Undefined. This value SHALL NOT appear
- 1: [RFC 5280](#) Certificate Revocation List (CRL)
- Other: reserved for future use.

11.23.9.13. SchemaVersion

The `SchemaVersion` field value history for this schema is provided below:

Version	Description
0	Initial Release

11.23.10. APIs / CLI

The Ledger comes with a set of secure APIs and CLI. More details available [here](https://github.com/zigbee-alliance/distributed-compliance-ledger/blob/master/docs/transactions.md) [https://github.com/zigbee-alliance/distributed-compliance-ledger/blob/master/docs/transactions.md].

11.24. Joint Fabric Datastore Cluster

The Joint Fabric Datastore Cluster is a cluster that provides a mechanism for the [Joint Fabric](#) Administrators to manage the set of Nodes, Groups, and Group membership among Nodes in the [Joint Fabric](#).

When an Ecosystem Administrator Node is commissioned onto the [Joint Fabric](#), the Ecosystem Administrator Node has no knowledge of what Nodes and Groups are present, or what set-up information related to the [Joint Fabric](#) is provided by the user. To address lack of knowledge, the Joint Fabric Datastore provides the information required for all Ecosystem Administrators to maintain a consistent view of the [Joint Fabric](#) including Nodes, Groups, settings and privileges.

The Joint Fabric Datastore contains the access control configuration for the Joint Fabric - the groups, the group keys, the group membership (expressed in terms of binding and ACL entries), as well as the CAT value and version for each group. This section describes how all changes to the Joint Fabric access control configuration are made via the Joint Fabric Datastore: a change is first made to the Datastore where the impacted configuration of individual nodes is marked as pending; the Datastore is then responsible for propagating the change to all impacted nodes and then updating its per-node (and sometimes per-endpoint) pending state to committed.

The Joint Fabric Datastore cluster server SHALL only be accessible on a Node which is acting as the Joint Fabric Anchor Administrator. When not acting as the Joint Fabric Anchor Administrator, the Joint Fabric Datastore cluster SHALL NOT be accessible.

The **Admin** level of access to the Joint Fabric Datastore cluster server SHALL be limited to JF Adminis-

trator Nodes identified using the [Administrator CAT](#).

11.24.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.24.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	JFDS

11.24.3. Cluster ID

ID	Name
0x0752	Joint Fabric Datastore Cluster

11.24.4. Usage Requirements

This section describes the required flows for Joint Fabric Administrators to perform common actions related to adding and removing Nodes to a Joint Fabric, and managing Groups and Group Membership on the Joint Fabric.

Many flows performed by a Joint Fabric Administrator require commands to be sent to the Datastore in order to keep the Datastore up-to-date with regard to Nodes and Groups on the Fabric. While a Joint Fabric Administrator SHALL update the Datastore with changes made to Nodes and Groups, it MAY complete these configuration and commissioning tasks when the Datastore is unreachable. If attempts to update the Datastore fail due to temporary conditions (reachability issues, BUSY response error, RESOURCE_EXHAUSTED error, etc.), the Commissioner SHALL periodically retry the attempted commands until successful.

The RESOURCE_EXHAUSTED error MAY be used by the Datastore to indicate that a storage capacity limit of the Datastore has been reached and SHALL notify the user of this condition using proprietary means outside of this specification (ex. email, screen-based notification, etc.), and SHALL provide a means for the user to purge information pertaining to unused Nodes.

11.24.4.1. Adding a Node to the Joint Fabric

During commissioning of a Node onto a Joint fabric, the Fabric Administrator performing the commissioning first populates the Datastore in order to ensure the new Node can be discovered by any client on the fabric for which the user has granted access. Specifically, before completing the commissioning process, the Commissioner SHALL invoke the [AddPendingNode](#) command in order to add a node entry to the Datastore (which will be marked with “pending” status since commissioning of the device has not yet completed), and follow the rules defined in [Section 11.24.4.9](#),

“Adding/Removing access privileges for a Controller”, to ensure access privileges are managed consistently across the Joint Fabric. After completion of commissioning and associated configuration, the Commissioner SHALL invoke the [RefreshNode](#) command in order to indicate completion of commissioning and to cause the Datastore to update the [DatastoreNodeInformationEntryStruct](#) with state read from the Node and mark the entry as "committed". By performing this work in two steps (first pending status, then committed status), the design can prevent error scenarios where a node is brought onto a fabric without appearing in the Datastore.

Commissioning Flow - Datastore Impact

1. Regular Flow: Commissioner obtains Onboarding Payload, discovers device, completes PASE portion of commissioning.
2. Datastore Step: Commissioner SHALL add a [DatastoreNodeInformationEntryStruct](#) to the Datastore (see [Section 11.24.7.10, “AddPendingNode Command”](#)).
3. Regular Flow: Commissioner finalizes commissioning via CASE.

If the finalize commissioning step fails, the Commissioner removes the pending node from the Datastore (see [Section 11.24.7.13, “RemoveNode Command”](#)).

If the finalize commissioning step succeeds, the Commissioner SHALL indicate completion by calling the [RefreshNode](#) command.

11.24.4.2. Removing a Node from the Joint Fabric

When removing a Node from the [Joint Fabric](#) the following steps SHALL be followed:

1. The [Joint Fabric](#) Administrator SHALL first remove access privileges (see [Section 11.24.4.9, “Adding/Removing access privileges for a Controller”](#)) for the Node.
2. The [DatastoreNodeInformationEntryStruct](#) SHALL be removed from the Joint Fabric Datastore using the [RemoveNode](#) command.

11.24.4.3. Pending Data Cleanup

The Datastore SHALL periodically review its data in a [Pending](#) and [PendingDeletion](#) state and attempt to reach the corresponding Node in order to apply these updates. Examples of configuration that may be in such a state include:

1. Node GroupKey additions and removals.
2. Node Endpoint Group membership additions and removals.
3. Node ACL additions, updates and removals for groups, including changes relating to CAT version updates.
4. Node Binding additions and removals for groups.

See [Section 11.24.7.11, “RefreshNode Command”](#) command for a detailed description of each step. An Administrator MAY invoke this command at any time to request that the Datastore perform this operation immediately for a given Node.

When a new node entry is left in the pending state, a [Joint Fabric](#) Administrator may not be able to

determine whether commissioning completed or failed. In this scenario, the [Joint Fabric](#) Administrator must continue to monitor and attempt maintenance duties upon the node until the node entry is removed by another [Joint Fabric](#) Administrator or by the user.

11.24.4.4. Adding a Node to a Group

When adding a Node to a Group, if the GroupKeySetID is set, then the [Joint Fabric](#) Administrator SHALL use the [AddGroupIDToEndpointForNode](#) command to update the Node with the KeySet and Group relationship required for the new Group membership.

When the Group membership is intended for control of the Node, the [Joint Fabric](#) Administrator SHALL use the [AddACLToNode](#) command to update the Node with the appropriate ACL.

When the Group membership is intended for control by the Node, the [Joint Fabric](#) Administrator SHALL use the [AddBindingToEndpointForNode](#) command to update the Node with the appropriate Binding.

11.24.4.5. Removing a Node from a Group

When the Group membership was used for control of the Node, the [Joint Fabric](#) Administrator SHALL use the [RemoveACLFromNode](#) command to remove the ACL from the Node.

When the Group membership was used for control by the Node, the [Joint Fabric](#) Administrator SHALL use the [RemoveBindingFromEndpointForNode](#) command to remove the Binding from the Node.

When removing a Node from a Group, if the GroupKeySetID is set, then the [Joint Fabric](#) Administrator SHALL use the [RemoveGroupIDFromEndpointForNode](#) command to ensure the KeySet and Group relationship is removed from the Node and the Datastore reflects the new Group membership.

For groups with access control considerations, the [Joint Fabric](#) Administrator SHALL use the [UpdateGroup](#) command to update the GroupKeySetID (if set) and increment the GroupCATVersion. If the Node is a Controller, then the [Joint Fabric](#) Administrator SHALL also update the GroupCATVersion of the Group.

11.24.4.6. Making a Node a Group Controller

In this example of an On/Off Light Switch controlling a group, Node N's Endpoint E (On/Off Light Switch) should control Group G with Group Key Set K:

1. Make sure Node N is in Group G by calling [AddGroupIDToEndpointForNode](#) command with arguments N, E, and G.
2. Add a binding to Endpoint E on Node N by calling [AddBindingToEndpointForNode](#) command with arguments N, E, and the desired Binding.

11.24.4.7. Adding a Joint Fabric Administrator Node to a Fabric

When a Node is a Joint Fabric Administrator, its Group membership does not need to be managed by the Datastore since it has the ability to monitor the Datastore and issue operational certs for its

own Administrator Nodes at any time. As a result, rather than adding Joint Fabric Administrator Nodes to the regular NodeList attribute via the [AddPendingNode](#) command, these Nodes SHALL be added to the AdminList via the [AddAdmin](#) command.

It's important to explicitly track the vendor ID of an [Joint Fabric](#) Administrator so that the [Joint Fabric](#) Administrator controlling the Root Operational Cert for the fabric can sign intermediate certificates which can be used by an [Joint Fabric](#) Administrator to issue operational certificates.

11.24.4.8. Adding/Updating/Removing a Group

When adding a group, a client must create the [DatastoreGroupInformationEntryStruct](#) entry in the GroupList attribute before adding membership associations between the Node and the Group.

Adding a Group

Add Group Key Set Entry (when applicable) using the [AddKeySet](#) command, followed by [DatastoreGroupInformationEntryStruct](#) entry using the [AddGroup](#) command.

Updating a Group

If the GroupKeySetID is changing, update the [DatastoreGroupKeySetStruct](#) of the GroupKeySetList attribute using the [UpdateKeySet](#) command. Update the [DatastoreGroupInformationEntryStruct](#) entry of the GroupList attribute using the [UpdateGroup](#) command. KeySet and Group updates performed using these commands will cause the Datastore to attempt to update all Node members of the group with these changes. Any Node changes that do not complete (for example, if the Node is unreachable) will result in [Pending](#) status on the data in the [DatastoreNodeInformationEntryStruct](#).

Removing a Group

Remove the [DatastoreGroupInformationEntryStruct](#) entry of the GroupList attribute using the [UpdateGroup](#) command. This command will fail if the Group is in use (and not with [PendingDelete](#) Status) by any Nodes.

If the GroupKeySetID is set, remove the [DatastoreGroupKeySetStruct](#) of the GroupKeySetList attribute using the [RemoveKeySet](#) command. This command will fail if this KeySet is in use by any Groups.

11.24.4.9. Adding/Removing access privileges for a Controller

In order to ensure a consistent approach to access control by Administrators on a [Joint Fabric](#), [Joint Fabric](#) Administrators SHALL use groups in the Datastore for managing access to Nodes. One goal of this approach is to avoid creating ACL entries on Nodes which reference a single NodeID, which does not scale as the number of Controller NodeIDs increases. Instead, ACLs can reference the CAT tag of the group or the GroupID of the group, depending upon the type of access intended. Each group in the Datastore may contain a CAT tag for use in unicast access to group members, and may contain a KeySet for use in broadcast access to group members.

In order to add a new access privilege to NodeA by ControllerB, the [Joint Fabric](#) Administrator SHALL identify a group to use for this grant (using [Section 11.24.7.4, "AddGroup Command"](#) to create a new one if needed), and call [Section 11.24.7.15, "AddGroupIDToEndpointForNode Command"](#)

and [Section 11.24.7.19, “AddACLToNode Command”](#) in order to configure NodeA to allow access by the group. Next, the Administrator SHALL provide a NOC to the controller containing the CAT tag specified for the given group.

In order to remove an access privilege grant for ControllerB on NodeA, the [Joint Fabric Administrator](#) SHALL revoke access to the group by ControllerB (see [Section 11.24.4.5, “Removing a Node from a Group”](#)) and in doing so, ensure that the CAT version for the [DatastoreGroupInformationEntryStruct](#) entry for the given GroupID in the GroupList attribute is incremented.

11.24.5. Data Types

11.24.5.1. DatastoreStateEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Pending	Target device operation is pending	M
1	Committed	Target device operation has been committed	M
2	DeletePending	Target device delete operation is pending	M
3	CommitFailed	Target device operation has failed	M

11.24.5.2. DatastoreStatusEntryStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	State	DatastoreStateEnum	all			R V	M
1	Update-Time-stamp	epoch-s	all			R V	M
2	Failure-Code	Status Code	all			R V	M

State Field

This field SHALL contain the current state of the target device operation.

UpdateTimestamp Field

This field SHALL contain the timestamp of the last update.

FailureCode Field

This field SHALL contain the Status Code of the last failed operation where the State field is set to **CommitFailure**.

11.24.5.3. DatastoreNodeKeySetEntryStruct Type

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
0	NodeID	node-id	all			R V	M
1	GroupKeySetID	uint16	all			R V	M
2	StatusEntry	DatastoreStatusEntryStruct				R V	M

NodeID Field

The unique identifier for the node.

GroupKeySet Field

Group Key Set structure containing the list of epoch keys.

StatusEntry Field

Indicates whether entry in this list is pending, committed, delete-pending, or commit-failed.

11.24.5.4. DatastoreAccessControlEntryPrivilegeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
1	View	Can read and observe all (except Access Control Cluster)	M
3	Operate	View privileges, and can perform the primary function of this Node (except Access Control Cluster)	M
4	Manage	Operate privileges, and can modify persistent configuration of this Node (except Access Control Cluster)	M

Value	Name	Summary	Conformance
5	Administer	Manage privileges, and can observe and modify the Access Control Cluster	M

11.24.5.5. DatastoreGroupInformationEntryStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	GroupID	uint64	all			R V	M
1	Friendly-Name	string	max 32			R V	M
2	GroupKey-SetID	uint16	1 to 65534	X		R V	M
3	GroupCAT	uint16	desc	X		R V	M
4	Group-CATVersion	uint16	1 to 65534	X		R V	M
5	GroupPermission	Datastore-AccessControlEntryPrivilegeEnum	all			R V	M

GroupID Field

The unique identifier for the group.

FriendlyName Field

The friendly name for the group.

GroupKeySetID Field

The unique identifier for the group key set.

This value MAY be null when multicast communication is not used for the group. When GroupPermission is Admin or Manage, this value SHALL be null.

A value of 0 is not allowed since this value is reserved for IPK and the group entry for this value is not managed by the Datastore.

GroupCAT Field

CAT value for this group. This is used for control of individual members of a group (non-broadcast commands).

Allowable values include the range 0x0000 to 0xEFFF, and the [Administrator CAT](#) and [Anchor CAT](#) values.

This value MAY be null when unicast communication is not used for the group.

GroupCATVersion Field

Current version number for this CAT.

This value SHALL be null when GroupCAT value is null.

GroupPermission Field

The permission level associated with ACL entries for this group. There should be only one Administrator group per fabric, and at most one Manage group per Ecosystem (Vendor Entry).

11.24.5.6. DatastoreBindingTargetStruct Type

The DatastoreBindingTargetStruct represents a Binding on a specific Node (identified by the DatastoreEndpointBindingEntryStruct) which is managed by the Datastore. Only bindings on a specific Node that are fabric-scoped to the Joint Fabric are managed by the Datastore. As a result, references to nodes and groups are specific to the Joint Fabric.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	Node	node-id	all				Endpoint
2	Group	group-id	min 1				!Endpoint
3	Endpoint	endpoint-no	all				!Group
4	Cluster	cluster-id	all				O

Node Field

This field is the binding's remote target node ID. If the Endpoint field is present, this field SHALL be present.

Group Field

This field is the binding's target group ID that represents remote endpoints. If the Endpoint field is present, this field SHALL NOT be present.

Endpoint Field

This field is the binding's remote endpoint that the local endpoint is bound to. If the Group field is present, this field SHALL NOT be present.

Cluster Field

This field is the binding's cluster ID (client & server) on the local and target endpoint(s). If this field is present, the client cluster SHALL also exist on this endpoint (with this Binding cluster). If this

field is present, the target SHALL be this cluster on the target endpoint(s).

11.24.5.7. DatastoreEndpointBindingEntryStruct Type

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
0	NodeID	node-id	all			R V	M
1	EndpointID	endpoint-no	all			R V	M
2	ListID	uint16	all			R V	M
3	Binding	Datastore-Binding-Target-Struct				R V	M
4	StatusEntry	DatastoreStatusEntryStruct				R V	M

NodeID Field

The unique identifier for the node.

EndpointID Field

The unique identifier for the endpoint.

ListID Field

The unique identifier for the entry in the Datastore's EndpointBindingList attribute, which is a list of DatastoreEndpointBindingEntryStruct.

This field is used to uniquely identify an entry in the EndpointBindingList attribute for the purpose of deletion (RemoveBindingFromEndpointForNode Command).

Binding Field

The binding target structure.

StatusEntry Field

Indicates whether entry in this list is pending, committed, delete-pending, or commit-failed.

11.24.5.8. DatastoreEndpointGroupIDEntryStruct Type

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
0	NodeID	node-id	all			R V	M

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
1	EndpointID	endpoint-no	all			R V	M
2	GroupID	group-id	all			R V	M
3	StatusEntry	DatastoreStatusEntryStruct				R V	M

NodeID Field

The unique identifier for the node.

EndpointID Field

The unique identifier for the endpoint.

GroupID Field

The unique identifier for the group.

StatusEntry Field

Indicates whether entry in this list is pending, committed, delete-pending, or commit-failed.

11.24.5.9. DatastoreEndpointEntryStruct Type

The DatastoreEndpointEntryStruct represents an Endpoint on a specific Node which is managed by the Datastore. Only Nodes on the Joint Fabric are managed by the Datastore. As a result, references to NodeID are specific to the Joint Fabric.

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
0	EndpointID	endpoint-no	all			R V	M
1	NodeID	node-id	all			R V	M
2	Friendly-Name	string	max 32			R V	M

EndpointID Field

The unique identifier for the endpoint.

NodeID Field

The unique identifier for the node.

FriendlyName Field

This field SHALL indicate a user-assigned label for this endpoint, as captured by a [Joint Fabric Administrator](#)'s user interface. By maintaining this value in the Joint Fabric Datastore, all [Joint Fabric](#) Administrators can keep these values synchronized. This is particularly useful for complex multi-application endpoint devices (such as appliances) and for endpoints exposed via a Bridge, where individual endpoints might have custom names assigned by the user. For basic devices, only the node-level FriendlyName might be used.

Administrators MAY keep this field in sync with the NodeLabel field from the Basic Information or Bridged Basic Information clusters.

11.24.5.10. DatastoreAccessControlEntryAuthModeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
1	PASE	Passcode authenticated session	M
2	CASE	Certificate authenticated session	M
3	Group	Group authenticated session	M

11.24.5.11. DatastoreAccessControlTargetStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	Cluster	cluster-id	all	X			M
1	Endpoint	endpoint-no	all	X			M
2	Device-Type	devtype-id	all	X			M

11.24.5.12. DatastoreAccessControlEntryStruct Type

The DatastoreAccessControlEntryStruct represents an ACL on a specific Node (identified by the DatastoreACLEntryStruct) which is managed by the Datastore. Only ACLs on a specific Node that are fabric-scoped to the Joint Fabric are managed by the Datastore. As a result, references to nodes and groups are specific to the Joint Fabric.

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
1	Privilege	DatastoreAccessControlEntryPrivilegeEnum	all				M
2	AuthMode	DatastoreAccessControlEntryAuthModeEnum	all				M
3	Subjects	list[subject-id]	max SubjectsPerAccessControlEntry	X			M
4	Targets	list[DatastoreAccessControlTargetStruct]	max TargetsPerAccessControlEntry	X			M

11.24.5.13. DatastoreACLEntryStruct Type

The DatastoreACLEntryStruct is a holder for an ACL (DatastoreAccessControlEntryStruct) on a specific Node which is managed by the Datastore. Only ACLs on a specific Node that are fabric-scoped to the Joint Fabric are managed by the Datastore. As a result, references to nodes and groups are specific to the Joint Fabric.

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
0	NodeID	node-id	all			R V	M
1	ListID	uint16	all			R V	M
2	ACLEntry	DatastoreAccessControlEntryStruct				R V	M
3	StatusEntry	DatastoreStatusEntryStruct				R V	M

NodeID Field

The unique identifier for the node.

ListID Field

The unique identifier for the ACL entry in the Datastore’s list of DatastoreACLEntry.

ACLEntry Field

The Access Control Entry structure.

StatusEntry Field

Indicates whether entry in this list is pending, committed, delete-pending, or commit-failed.

11.24.5.14. DatastoreNodeInformationEntryStruct Type

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
0	NodeID	node-id	all			R V	M
1	Friendly-Name	string	max 32			R V	M
2	CommissioningStatusEntry	DatastoreStatusEntryStruct				R V	M

NodeID Field

The unique identifier for the node.

FriendlyName Field

This field SHALL contain a user-assigned label for this node, as captured by a [Joint Fabric](#) Administrator’s user interface. By maintaining this value in the Joint Fabric Datastore, all [Joint Fabric](#) Administrators can keep these values synchronized. This value is not propagated to the node itself.

Administrators MAY keep this field in sync with the NodeLabel field from the Basic Information or Bridged Basic Information clusters.

CommissioningStatusEntry Field

Set to **Pending** prior to completing commissioning, set to **Committed** after commissioning complete is successful, or set to **CommitFailed** if commissioning failed with the FailureCode Field set to the error.

11.24.5.15. DatastoreAdministratorInformationEntryStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	NodeID	node-id	all				M
1	Friendly-Name	string	max 32				M
2	VendorID	vendor-id	all				M
3	ICAC	octstr	max 400				M

NodeID Field

The unique identifier for the node.

FriendlyName Field

Friendly name for this node which is not propagated to nodes.

VendorID Field

The Vendor ID for the node.

ICAC Field

The ICAC used to issue the NOC.

11.24.5.16. DatastoreGroupKeySecurityPolicyEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	TrustFirst	Message counter synchronization using trust-first	M

11.24.5.17. DatastoreGroupKeyMulticastPolicyEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	PerGroupID	Indicates filtering of multicast messages for a specific Group ID	M
1	AllNodes	Indicates not filtering of multicast messages	M

11.24.5.18. DatastoreGroupKeySetStruct Type

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	GroupKey-SetID	uint16	all				M
1	GroupKey-Security-Policy	Datastore-GroupKey-Security-PolicyEnum	all				M
2	EpochKey0	octstr	16	X			M
3	EpochStartTime0	epoch-us	all	X			M
4	EpochKey1	octstr	16	X			M
5	EpochStartTime1	epoch-us	all	X			M
6	EpochKey2	octstr	16	X			M
7	EpochStartTime2	epoch-us	all	X			M
8	GroupKey-Multicast-Policy	Datastore-GroupKey-Multicast-PolicyEnum	all				D

11.24.6. Attributes

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
0x00	Anchor-RootCA	octstr	max 400	N		A R	M
0x01	AnchorNodeID	node-id	all	N		A R	M
0x02	Anchor-VendorID	vendor-id	all	N		A R	M
0x03	Friendly-Name	string	max 32	N		A R	M

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
0x04	GroupKey-SetList	list[DatatstoreGroupKeySetStruct]	1 to 254	N		A R	M
0x05	GroupList	list[DatatstoreGroupInformationEntryStruct]	1 to 254	N		A R	M
0x06	NodeList	list[DatatstoreNodeInformationEntryStruct]	max 1024	N		A R	M
0x07	AdminList	list[DatatstoreAdministratorInformationEntryStruct]	1 to 32	N		A R	M
0x08	Status	DatatstoreStatusEntryStruct	all	N		A R	M
0x09	Endpoint-GroupIDList	list[DatatstoreEndpointGroupIDEntryStruct]	max 8192	N		A R	M
0x0A	Endpoint-BindingList	list[DatatstoreEndpointBindingEntryStruct]	max 8192	N		A R	M
0x0B	NodeKey-SetList	list[DatatstoreNodeKeySetEntryStruct]	max 8192	N		A R	M

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
0x0C	NodeACLList	list[DatastoreACLEntryStruct]	max 8192	N		A R	M
0x0D	NodeEndpointList	list[DatastoreEndpointEntryStruct]	max 8192	N		A R	M

11.24.6.1. AnchorRootCA Attribute

This SHALL indicate the Anchor Root CA used to sign all [NOC](#) Issuers in the [Joint Fabric](#) for the accessing fabric. A null value indicates that the [Joint Fabric](#) is not yet formed.

11.24.6.2. AnchorNodeID Attribute

This SHALL indicate the Node identifier of the [Joint Fabric](#) Anchor Root CA for the accessing fabric.

11.24.6.3. AnchorVendorID Attribute

This SHALL indicate the Vendor identifier of the [Joint Fabric](#) Anchor Root CA for the accessing fabric.

11.24.6.4. FriendlyName Attribute

Friendly name for the accessing fabric.

11.24.6.5. GroupKeySetList Attribute

This SHALL indicate the list of [DatastoreGroupKeySetStruct](#) used in the [Joint Fabric](#) for the accessing fabric.

This attribute SHALL contain at least one entry, the [IPK](#), which has GroupKeySetID of 0.

11.24.6.6. GroupList Attribute

This SHALL indicate the list of groups in the [Joint Fabric](#) for the accessing fabric.

This list SHALL include, at a minimum, one group with GroupCAT value set to [Administrator CAT](#) and one group with GroupCAT value set to [Anchor CAT](#).

11.24.6.7. NodeList Attribute

This SHALL indicate the list of nodes in the [Joint Fabric](#) for the accessing fabric.

11.24.6.8. AdminList Attribute

This SHALL indicate the list of administrators in the [Joint Fabric](#) for the accessing fabric.

Only one Administrator may serve as the Anchor Root CA and Anchor Fabric Administrator and SHALL have index value 0. All other [Joint Fabric](#) Administrators SHALL be referenced at index 1 or greater.

An empty list indicates that the [Joint Fabric](#) is not yet formed.

11.24.6.9. Status Attribute

This SHALL indicate the current state of the [Joint Fabric](#) Datastore Cluster for the accessing fabric.

The value SHALL be one of the following states:

- **Committed** - indicates the DataStore is ready for use.
- **Pending** - indicates that the DataStore is not yet ready for use.
- **DeletePending** - indicates that the DataStore is in the process of being transferred to another [Joint Fabric](#) Anchor Administrator.

11.24.6.10. EndpointGroupIDList Attribute

This SHALL indicate the group membership of endpoints in the accessing fabric.

11.24.6.11. EndpointBindingList Attribute

This SHALL indicate the binding list for endpoints in the accessing fabric.

11.24.6.12. NodeKeySetList Attribute

This SHALL indicate the KeySet entries for nodes in the accessing fabric.

11.24.6.13. NodeACLList Attribute

This SHALL indicate the ACL entries for nodes in the accessing fabric.

11.24.6.14. NodeEndpointList Attribute

This SHALL indicate the Endpoint entries for nodes in the accessing fabric.

11.24.7. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	AddKeySet	Client ⇒ Server	Y	A	M
0x01	UpdateKeySet	Client ⇒ Server	Y	A	M
0x02	RemoveKey-Set	Client ⇒ Server	Y	A	M
0x03	AddGroup	Client ⇒ Server	Y	A	M
0x04	UpdateGroup	Client ⇒ Server	Y	A	M
0x05	RemoveGroup	Client ⇒ Server	Y	A	M

ID	Name	Direction	Response	Access	Conformance
0x06	AddAdmin	Client ⇒ Server	Y	A	M
0x07	UpdateAdmin	Client ⇒ Server	Y	A	M
0x08	RemoveAdmin	Client ⇒ Server	Y	A	M
0x09	AddPendingNode	Client ⇒ Server	Y	A	M
0x0A	RefreshNode	Client ⇒ Server	Y	A	M
0x0B	UpdateNode	Client ⇒ Server	Y	A	M
0x0C	RemoveNode	Client ⇒ Server	Y	A	M
0x0D	UpdateEndpointForNode	Client ⇒ Server	Y	A	M
0x0E	AddGroupID-ToEndpoint-ForNode	Client ⇒ Server	Y	A	M
0x0F	Remove-GroupID-FromEndpointForNode	Client ⇒ Server	Y	A	M
0x10	AddBinding-ToEndpoint-ForNode	Client ⇒ Server	Y	A	M
0x11	RemoveBindingFromEndpointForNode	Client ⇒ Server	Y	A	M
0x12	AddACLToNode	Client ⇒ Server	Y	A	M
0x13	RemoveACL-FromNode	Client ⇒ Server	Y	A	M

11.24.7.1. AddKeySet Command

Upon receipt, this SHALL add a KeySet to the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupKey-Set	Datastore-GroupKey-SetStruct				M

GroupKeySet Field

This field SHALL indicate the KeySet to be added to the [Joint Fabric](#) Datastore Cluster.

The Datastore SHALL:

- 1. Ensure there are no KeySets in the KeySetList attribute with the given GroupKeySetID.
- 2. If a match is found, then this command SHALL fail with a CONSTRAINT_ERROR status code.
- 3. Add the [Epoch Key Entry](#) for the KeySet to the KeySetList attribute.

11.24.7.2. UpdateKeySet Command

Upon receipt, this SHALL update a KeySet in the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupKey-Set	Datastore-GroupKey-SetStruct				M

GroupKeySet Field

This field SHALL indicate the KeySet to be updated in the [Joint Fabric](#) Datastore Cluster.

The Datastore SHALL:

- 1. Find the Epoch Key Entry for the KeySet in the KeySetList attribute with the given GroupKeySetID, and update any changed fields.
- 2. If entry is not found, then this command SHALL fail with a NOT_FOUND status code.
- 3. If any fields are changed as a result of this command:
 - a. Iterate through each [DatastoreNodeInformationEntryStruct](#):
 - i. If the NodeKeySetList contains an entry with the given GroupKeySetID:
 - A. Update the Status on the given [DatastoreNodeKeySetEntryStruct](#) to **Pending**.
 - B. Update the GroupKeySet on the given Node with the new values.
 - I. If successful, update the Status on this [DatastoreNodeKeySetEntryStruct](#) to **Committed**.
 - II. If not successful, update the State field of the StatusEntry on this [DatastoreNodeKeySetEntryStruct](#) to **CommitFailed** and FailureCode code to the returned error. The pending change SHALL be applied in a subsequent [Node Refresh](#).

11.24.7.3. RemoveKeySet Command

Upon receipt, this SHALL remove a KeySet from the [Joint Fabric](#) Datastore Cluster of the accessing

fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupKey-SetID	uint16	all			M

GroupKeySetID Field

This field SHALL indicate the unique identifier for the KeySet to be removed from the [Joint Fabric](#) Datastore Cluster.

Attempt to remove the [IPK](#), which has GroupKeySetID of 0, SHALL fail with response CONSTRAINT_ERROR.

The Datastore SHALL:

1. If entry is not found, then this command SHALL fail with a NOT_FOUND status code.
2. Ensure there are no Nodes using this KeySet. To do this:
 - a. Iterate through each [DatastoreNodeInformationEntryStruct](#):
 - i. If the NodeKeySetList list contains an entry with the given GroupKeySetID, and the entry does NOT have Status [DeletePending](#), then this command SHALL fail with a CONSTRAINT_ERROR status code.
3. Remove the [DatastoreGroupKeySetStruct](#) for the given GroupKeySetID from the GroupKeySetList attribute.

11.24.7.4. AddGroup Command

Upon receipt, this SHALL add a group to the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupID	group-id	all			M
1	Friendly-Name	string	max 32			M
2	GroupKey-SetID	uint16	1 to 65534	X		M
3	GroupCAT	uint16	desc	X		M
4	Group-CATVersion	uint16	1 to 65534	X		M

ID	Name	Type	Constraint	Quality	Fallback	Conformance
5	GroupPermission	DatastoreAccessControlEntryPrivilegeEnum				M

GroupID Field

This field SHALL indicate the unique identifier for the group to be added to the [Joint Fabric](#) Datastore Cluster.

FriendlyName Field

This field SHALL indicate the friendly name for the group.

GroupKeySetID Field

This field SHALL indicate the unique identifier for the group key set.

GroupCAT Field

This field SHALL indicate the CAT value for this group.

GroupCAT values SHALL fall within the range 1 to 65534. Attempts to add a group with a GroupCAT value of [Administrator CAT](#) or [Anchor CAT](#) SHALL fail with CONSTRAINT_ERROR.

GroupCATVersion Field

This field SHALL indicate the current version number for this CAT.

GroupPermission Field

This field SHALL indicate the permission level associated with ACL entries for this group.

The Datastore SHALL:

1. Ensure there are no Groups in the GroupList attribute with the given GroupID. If a match is found, then this command SHALL fail with a CONSTRAINT_ERROR status code.
2. Add the [DatastoreGroupInformationEntryStruct](#) for the Group with the given GroupID to the GroupList attribute.

11.24.7.5. UpdateGroup Command

Upon receipt, this SHALL update a group in the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupID	group-id	all			M
1	Friendly-Name	string	max 32	X		M
2	GroupKey-SetID	uint16	1 to 65535	X		M
3	GroupCAT	uint16	desc	X		M
4	Group-CATVersion	uint16	1 to 65535	X		M
5	GroupPermission	DatastoreAccessControlEntryPrivilegeEnum		X		M

GroupID Field

This field SHALL indicate the group to be updated in the [Joint Fabric](#) Datastore Cluster.

FriendlyName Field

This field SHALL indicate the friendly name for the group. NULL values will be ignored (not updated).

GroupKeySetID Field

This field SHALL indicate the unique identifier for the group key set. NULL values will be ignored (not updated).

GroupCAT Field

This field SHALL indicate the CAT value for this group. NULL values will be ignored (not updated).

GroupCAT values SHALL fall within the range 1 to 65534. Attempts to update the GroupCAT on an existing group which has a GroupCAT value of [Administrator CAT](#) or [Anchor CAT](#) SHALL fail with CONSTRAINT_ERROR.

Attempts to set the GroupCAT to [Administrator CAT](#) or [Anchor CAT](#) SHALL fail with CONSTRAINT_ERROR.

GroupCATVersion Field

This field SHALL indicate the current version number for this CAT. NULL values will be ignored (not updated).

GroupPermission Field

This field SHALL indicate the permission level associated with ACL entries for this group. NULL val-

ues will be ignored (not updated).

The Datastore SHALL:

1. If entry is not found, then this command SHALL fail with a NOT_FOUND status code.
2. Update the [DatastoreGroupInformationEntryStruct](#) for the Group with the given GroupID to match the non-NULL fields passed in.
3. If any fields are changed as a result of this command:
 - a. Iterate through each [DatastoreNodeInformationEntryStruct](#):
 - i. If the GroupKeySetID changed:
 - I. Add a [DatastoreNodeKeySetEntryStruct](#) with the new GroupKeySetID, and Status set to **Pending**.
 - II. Add this KeySet to the Node.
 1. If successful, Set the Status to **Committed** for this entry in the NodeKeySetList.
 2. If not successful, Set the Status to **CommitFailed** and the FailureCode to the returned error. The pending change SHALL be applied in a subsequent [Node Refresh](#).
 - A. If the NodeKeySetList list contains an entry with the previous GroupKeySetID:
 - III. Set the Status set to **DeletePending**.
 - IV. Remove this KeySet from the Node.
 1. If successful, Remove this entry from the NodeKeySetList.
 2. If not successful, the pending change SHALL be applied in a subsequent [Node Refresh](#).
 - ii. If the GroupCAT, GroupCATVersion or GroupPermission changed:
 - A. If the ACLList contains an entry for this Group, update the ACL List Entry in the Datastore with the new values and Status **Pending**, update the ACL attribute on the given Node with the new values. If the update succeeds, set the Status to **Committed** on the ACLList Entry in the Datastore.
 - iii. If the FriendlyName changed:
 - A. Iterate through each [DatastoreEndpointGroupIDEntryStruct](#) in the Endpoint-GroupIDList attribute:
 - I. If the [DatastoreEndpointGroupIDEntryStruct](#) contains an entry with the given GroupID:
 1. Update the [DatastoreEndpointGroupIDEntryStruct](#) Entry in the Datastore with the new values and Status **Pending**
 2. Update the Groups on the given Node with the new values.
 1. If the update succeeds, set the Status to **Committed** on the GroupIDList Entry in the Datastore.
 2. If not successful, the pending change SHALL be applied in a subsequent [Node Refresh](#).

11.24.7.6. RemoveGroup Command

Upon receipt, this SHALL remove a group from the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupID	group-id	all			M

GroupID Field

This field SHALL indicate the unique identifier for the group to be removed from the [Joint Fabric](#) Datastore Cluster.

Attempts to remove a group with GroupCAT value set to [Administrator CAT](#) or [Anchor CAT](#) SHALL fail with CONSTRAINT_ERROR.

The Datastore SHALL:

1. If entry is not found, then this command SHALL fail with a NOT_FOUND status code.
2. Ensure there are no Nodes in this group. To do this:
 - a. Iterate through each [DatastoreNodeInformationEntryStruct](#):
 - i. If the GroupIDList contains an entry with the given GroupID, and the entry does NOT have Status **DeletePending**, then this command SHALL fail with a CONSTRAINT_ERROR status code.
3. Remove the [DatastoreGroupInformationEntryStruct](#) for the Group with the given GroupID from the GroupList attribute.

11.24.7.7. AddAdmin Command

Upon receipt, this SHALL add an admin to the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M
1	Friendly-Name	string	max 32			M
2	VendorID	vendor-id	all			M
3	ICAC	octstr	max 400			M

NodeID Field

This field SHALL indicate the unique identifier for the admin node.

FriendlyName Field

This field SHALL indicate the friendly name for the admin node.

VendorID Field

This field SHALL indicate the Vendor ID for the admin node.

ICAC Field

This field SHALL indicate the ICAC used to issue the NOC.

The Datastore SHALL:

1. Ensure there are no Admins in the AdminList attribute with the given NodeID. If a match is found, then this command SHALL fail with a CONSTRAINT_ERROR status code.
2. Add the [DatastoreAdministratorInformationEntryStruct](#) for the Admin with the given NodeID to the AdminList attribute.

11.24.7.8. UpdateAdmin Command

Upon receipt, this SHALL update an admin entry in the AdminList attribute.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M
1	Friendly-Name	string	max 32			O
2	ICAC	octstr	max 400			O

NodeID Field

This field SHALL indicate the admin to be updated in the [Joint Fabric](#) Datastore Cluster.

FriendlyName Field

This field SHALL indicate the friendly name for the admin node.

ICAC Field

This field SHALL indicate the ICAC used to issue the NOC.

The Datastore SHALL:

1. Find the entry in the AdminList attribute with the given NodeID, and update any changed fields.
2. If entry is not found, then this command SHALL fail with a NOT_FOUND status code.

11.24.7.9. RemoveAdmin Command

Upon receipt, this SHALL remove an admin from the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M

NodeID Field

This field SHALL indicate the unique identifier for the admin to be removed from the [Joint Fabric](#) Datastore Cluster.

The Datastore SHALL:

1. If an entry in the AdminList attribute with the given NodeID is not found, then this command SHALL fail with a NOT_FOUND status code.
2. Remove the [DatastoreAdministratorInformationEntryStruct](#) for the Admin with the given NodeID from the AdminList attribute.

11.24.7.10. AddPendingNode Command

Upon receipt, this SHALL add a node to the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M
1	Friendly-Name	string	max 32			M

NodeID Field

This field SHALL indicate the node to be added to the [Joint Fabric](#) Datastore Cluster.

FriendlyName Field

This field SHALL indicate the friendly name for the node.

If a [DatastoreNodeInformationEntryStruct](#) exists for the given NodeID, then this command SHALL fail with a INVALID_CONSTRAINT status code.

The Datastore SHALL:

1. Add an entry for the given NodeID to the NodeList attribute and set the status for the [DatastoreNodeInformationEntryStruct](#) to **Pending**.

11.24.7.11. RefreshNode Command

Upon receipt, this SHALL request that Datastore information relating to a Node of the accessing fabric is refreshed.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M

NodeID Field

This field SHALL indicate the node for which Datastore information should be refreshed.

The Datastore SHALL:

1. Confirm that a [DatastoreNodeInformationEntryStruct](#) exists in the NodeList attribute for the given NodeID, and if not, then this command SHALL fail with a NOT_FOUND status code.
2. Update the CommissioningStatusEntry for the [DatastoreNodeInformationEntryStruct](#) to **Pending**.
3. Ensure the Endpoint List for the [DatastoreNodeInformationEntryStruct](#) with the given NodeID matches Endpoint list on the given Node. This involves the following steps:
 - a. Read the PartsList of the Descriptor cluster from the Node.
 - b. For each [DatastoreEndpointEntryStruct](#) in the NodeEndpointList attribute with the given NodeID that does not match an Endpoint ID in the PartsList, remove the [DatastoreEndpointEntryStruct](#).
 - c. For each [DatastoreEndpointEntryStruct](#) in the NodeEndpointList attribute with the given NodeID that matches an Endpoint ID in the PartsList:
 - i. Check that each entry in Node's Group List occurs in the EndpointGroupIDList attribute.
 - A. Add any missing entries to the EndpointGroupIDList.
 - B. For any entries in the EndpointGroupIDList attribute with the given NodeId and EndpointId with Status of **Pending**:
 - I. Add the corresponding change to the Node's Group List.
 1. If successful, mark the Status to **Committed**.
 2. If not successful, update the Status to **CommitFailed** and the FailureCode to the returned error. The error SHALL be handled in a subsequent [Node Refresh](#).
 - C. For any entries in the EndpointGroupIDList attribute with the given NodeID and EndpointID with Status of **DeletePending**:
 1. If successful, remove the corresponding entry from the Node's Group List.
 2. If not successful, update the Status to **CommitFailed** and the FailureCode to the returned error. The error SHALL be handled in a subsequent [Node Refresh](#).
 - D. For any entries in the EndpointGroupIDList attribute with the given NodeID and EndpointID with Status of **CommitFailure**:

- I. A **CommitFailure** with an unrecoverable FailureCode SHALL be handled by removing the entry from the GroupIDList.
 - II. A **CommitFailure** with a recoverable FailureCode (i.e. **TIMEOUT**, **BUSY**) SHALL be handle in a subsequent **Node Refresh**.
 - ii. Check that each entry in Node's Binding List occurs in the EndpointBindingList attribute with the given NodeId and EndpointId.
 - A. Add any missing entries to the EndpointBindingList attribute.
 - B. For any entries in the EndpointBindingList attribute with the given NodeID and EndpointID with Status of **Pending**:
 - I. Add the corresponding change to the Node's Binding List.
 1. If successful, mark the Status to **Committed**.
 2. If not successful, update the Status to **CommitFailed** and the FailureCode to the returned error. The error SHALL be handled in a subsequent **Node Refresh**.
 - C. For any entries in the EndpointBindingList attribute with the given NodeID and EndpointID with Status of **DeletePending**:
 1. If successful, remove the corresponding entry from the Node's BindingList.
 2. If not successful, update the Status to **CommitFailed** and the FailureCode to the returned error. The error SHALL be handled in a subsequent **Node Refresh**.
 - D. For any entries in the EndpointBindingList attribute with the given NodeID and EndpointID with Status of **CommitFailure**:
 - I. A **CommitFailure** with an unrecoverable FailureCode SHALL be handled by removing the entry from the BindingList.
 - II. A **CommitFailure** with a recoverable FailureCode (i.e. **TIMEOUT**, **BUSY**) SHALL be handle in a subsequent **Node Refresh**.
 4. Ensure the GroupKeySetList entries with the given NodeID match the Group Keys on the given Node. This involves the following steps:
 - a. Read the Group Keys from the Node.
 - b. For each **DatastoreGroupKeySetStruct** in the GroupKeySetList attribute for the given NodeID with a **Pending** Status:
 - i. Add the corresponding **DatastoreGroupKeySetStruct** to the Node's Group Key list.
 - A. If successful, mark the Status to **Committed**.
 - B. If not successful, update the Status to **CommitFailed** and the FailureCode to the returned error. The error SHALL be handled in a subsequent **Node Refresh**.
 - c. For each **DatastoreGroupKeySetStruct** in the GroupKeySetList attribute for the given NodeID with a **CommitFailure** Status:
 - i. A **CommitFailure** with an unrecoverable FailureCode SHALL be handled by removing the entry from the GroupKeySetList.
 - ii. A **CommitFailure** with a recoverable FailureCode (i.e. **TIMEOUT**, **BUSY**) SHALL be handle in a subsequent **Node Refresh**.

- d. All remaining entries in the GroupKeySetList attribute for the given NodeId should be replaced by the remaining entries on the Node.
5. Ensure the NodeACLList attribute for the given NodeID matches the ACL attribute on the given Node. This involves the following steps:
 - a. Read the ACL attribute on the Node.
 - b. For each DatastoreACLEntryStruct in the ACLList attribute with the given NodeID with a Pending Status:
 - i. Add the corresponding DatastoreACLEntryStruct to the Node's ACL attribute.
 - A. If successful, mark the Status to Committed.
 - B. If not successful, update the Status to CommitFailed and the FailureCode to the returned error. The error SHALL be handled in a subsequent Node Refresh.
 - c. For each DatastoreACLEntryStruct in the ACLList attribute with the given NodeID with a CommitFailure Status:
 - i. A CommitFailure with an unrecoverable FailureCode (i.e. RESOURCE_EXHAUSTED, CONSTRAINT_ERROR) SHALL be handled by removing the entry from the ACLList.
 - ii. A CommitFailure with a recoverable FailureCode (i.e. TIMEOUT, BUSY) SHALL be handle in a subsequent Node Refresh.
 - d. All remaining entries in the ACLList should be replaced by the remaining entries on the Node.
6. Update the CommissioningStatusEntry for the DatastoreNodeInformationEntryStruct to Committed.

11.24.7.12. UpdateNode Command

Upon receipt, this SHALL update the friendly name for a node in the Joint Fabric Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M
1	Friendly-Name	string	max 32			M

NodeID Field

This field SHALL indicate the node to be updated in the Joint Fabric Datastore Cluster.

FriendlyName Field

This field SHALL indicate the friendly name for the node.

The Datastore SHALL:

1. If a [DatastoreNodeInformationEntryStruct](#) does not exist for the given NodeID in the NodeList attribute, then this command SHALL fail with a NOT_FOUND status code.
2. Update the FriendlyName field of the [DatastoreNodeInformationEntryStruct](#) with the given NodeID.

11.24.7.13. RemoveNode Command

Upon receipt, this SHALL remove a node from the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M

NodeID Field

This field SHALL indicate the unique identifier for the node to be removed from the [Joint Fabric](#) Datastore Cluster.

The Datastore SHALL:

1. If a [DatastoreNodeInformationEntryStruct](#) does not exist for the given NodeID in the NodeList attribute, then this command SHALL fail with a NOT_FOUND status code.
2. Remove the [DatastoreNodeInformationEntryStruct](#) with the given NodeID from the NodeList attribute.

11.24.7.14. UpdateEndpointForNode Command

Upon receipt, this SHALL update the state of an endpoint for a node in the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	EndpointID	endpoint-no	all			M
1	NodeID	node-id	all			M
2	Friendly-Name	string	max 32			M

EndpointID Field

This field SHALL indicate the unique identifier for the endpoint to be updated in the [Joint Fabric](#) Datastore Cluster.

NodeID Field

This field SHALL indicate the unique identifier for the node to which the endpoint belongs.

FriendlyName Field

This field SHALL indicate the friendly name for the endpoint.

The Datastore SHALL:

1. If an [DatastoreNodeInformationEntryStruct](#) does not exist for the given NodeID and EndpointID in the NodeEndpointList attribute, then this command SHALL fail with a NOT_FOUND status code.
2. Update the FriendlyName field of the [DatastoreNodeInformationEntryStruct](#) with the given NodeID and EndpointID.

11.24.7.15. AddGroupIDToEndpointForNode Command

Upon receipt, this SHALL add a Group ID to an endpoint for a node in the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M
1	EndpointID	endpoint-no	all			M
2	GroupID	group-id	all			M

NodeID Field

This field SHALL indicate the unique identifier for the node to which the endpoint belongs.

EndpointID Field

This field SHALL indicate the unique identifier for the endpoint to be updated in the [Joint Fabric](#) Datastore Cluster.

GroupID Field

This field SHALL indicate the unique identifier for the group to be added to the endpoint.

The Datastore SHALL:

1. Confirm that an [DatastoreNodeInformationEntryStruct](#) exists for the given NodeID and EndpointID in the NodeEndpointList attribute, and if not, then this command SHALL fail with a NOT_FOUND status code.
2. Ensure the Group Key List for the [DatastoreNodeInformationEntryStruct](#) with the given NodeID includes the KeySet for the given Group ID. If it does not:

- a. Add an entry for the KeySet of the given Group ID to the Group Key List for the Node. The new entry’s status SHALL be set to Pending.

b. Add a Group Key Entry for this KeySet to the given Node ID.

i. If this succeeds, update the new KeySet entry in the Datastore to Committed.

ii. If not successful, the pending change SHALL be applied in a subsequent Node Refresh.
3. Ensure the Group List for the DatastoreNodeInformationEntryStruct with the given NodeID and EndpointID includes an entry for the given Group. If it does not:

a. Add a Group entry for the given Group ID to the Group List for the Endpoint and Node. The new entry’s status SHALL be set to Pending.

b. Add this Group entry to the given Endpoint ID on the given Node ID.

i. If this succeeds, update the new Group entry in the Datastore to Committed.

ii. If not successful, update the Status to CommitFailed and the FailureCode to the returned error. The error SHALL be handled in a subsequent Node Refresh.

11.24.7.16. RemoveGroupIDFromEndpointForNode Command

Upon receipt, this SHALL remove a Group ID from an endpoint for a node in the Joint Fabric Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M
1	EndpointID	endpoint-no	all			M
2	GroupID	group-id	all			M

NodeID Field

This field SHALL indicate the unique identifier for the node to which the endpoint belongs.

EndpointID Field

This field SHALL indicate the unique identifier for the endpoint to be updated in the Joint Fabric Datastore Cluster.

GroupID Field

This field SHALL indicate the unique identifier for the group to be removed from the endpoint.

The Datastore SHALL:

1. Confirm that an DatastoreNodeInformationEntryStruct exists for the given NodeID and EndpointID in the NodeEndpointList attribute, and if not, then this command SHALL fail with a NOT_FOUND status code.

2. Ensure the EndpointGroupIDList entries with the given NodeID and EndpointID does not include an entry for the given Group. If it does:
 - a. Update the status to **DeletePending** of the Group entry for the given Group ID in the Group List.
 - b. Remove this Group entry for the given Endpoint ID on the given Node ID.
 - i. If this succeeds, remove the Group entry for the given Group ID in the Group List for this NodeID and EndpointID in the Datastore.
 - ii. If not successful, the pending change SHALL be applied in a subsequent [Node Refresh](#).
3. Ensure the Group Key List for the [DatastoreNodeInformationEntryStruct](#) with the given NodeID does not include the KeySet for the given Group ID. If it does:
 - a. Update the status to **DeletePending** for the entry for the KeySet of the given Group ID in the Node Group Key List.
 - b. Remove the Group Key Entry for this KeySet from the given Node ID.
 - i. If this succeeds, remove the KeySet entry for the given Node ID.
 - ii. If not successful, update the Status to **CommitFailed** and the FailureCode to the returned error. The error SHALL be handled in a subsequent [Node Refresh](#).

11.24.7.17. AddBindingToEndpointForNode Command

Upon receipt, this SHALL add a binding to an endpoint for a node in the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M
1	EndpointID	endpoint-no	all			M
2	Binding	Datastore-BindingTargetStruct				M

NodeID Field

This field SHALL indicate the unique identifier for the node to which the endpoint belongs.

EndpointID Field

This field SHALL indicate the unique identifier for the endpoint to be updated in the [Joint Fabric](#) Datastore Cluster.

Binding Field

This field SHALL indicate the binding to be added to the endpoint.

The Datastore SHALL:

1. Confirm that an [DatastoreNodeInformationEntryStruct](#) exists for the given NodeID and EndpointID in the NodeEndpointList attribute, and if not, then this command SHALL fail with a NOT_FOUND status code.
2. Ensure the Binding List for the [DatastoreNodeInformationEntryStruct](#) with the given NodeID includes the given Binding. If it does not:
 - a. Add the [DatastoreEndpointBindingEntryStruct](#) entry to the EndpointBindingList attribute for the given NodeID and EndpointID. The new entry’s status SHALL be set to **Pending**.
 - b. Add this Binding to the given Node ID.
 - i. If this succeeds, update the new Binding in the Datastore to **Committed**.
 - ii. If not successful, update the Status to **CommitFailed** and the FailureCode to the returned error. The error SHALL be handled in a subsequent [Node Refresh](#).

11.24.7.18. RemoveBindingFromEndpointForNode Command

Upon receipt, this SHALL remove a binding from an endpoint for a node in the [Joint Fabric](#) Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ListID	uint16	all			M
1	EndpointID	endpoint-no	all			M
2	NodeID	node-id	all			M

ListID Field

This field SHALL indicate the unique identifier for the binding entry in the Datastore’s Endpoint-BindingList attribute to be removed from the endpoint.

EndpointID Field

This field SHALL indicate the unique identifier for the endpoint to be updated in the [Joint Fabric](#) Datastore Cluster.

NodeID Field

This field SHALL indicate the unique identifier for the node to which the endpoint belongs.

The Datastore SHALL:

1. Confirm that an [DatastoreNodeInformationEntryStruct](#) exists for the given NodeID and EndpointID in the EndpointBindingList attribute, and if not, then this command SHALL fail with a NOT_FOUND status code.
2. Ensure the EndpointBindingList entries with the given NodeID does not include an entry with

the given ListID. If it does:

- a. Update the status to **DeletePending** for the given Binding in the Binding List.
- b. Remove this Binding from the given Node ID.
 - i. If this succeeds, remove the given Binding from the Binding List.
 - ii. If not successful, update the Status to **CommitFailed** and the FailureCode to the returned error. The error SHALL be handled in a subsequent **Node Refresh**.

11.24.7.19. AddACLToNode Command

Upon receipt, this SHALL add an ACL to a node in the **Joint Fabric** Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	NodeID	node-id	all			M
1	ACLEntry	DatastoreAccessControlEntryStruct				M

NodeID Field

This field SHALL indicate the unique identifier for the node to which the ACL is to be added.

ACLEntry Field

This field SHALL indicate the ACL to be added to the **Joint Fabric** Datastore Cluster.

The Datastore SHALL:

- 1. Confirm that a [DatastoreNodeInformationEntryStruct](#) exists for the given NodeID in the NodeList attribute, and if not, then this command SHALL fail with a NOT_FOUND status code.
- 2. Ensure the ACL List for the given NodeID includes the given ACLEntry. If it does not:
 - a. Add the ACLEntry to the ACL List for the given NodeID. The new entry’s status SHALL be set to **Pending**.
 - b. Add this ACLEntry to the given Node ID.
 - i. If this succeeds, update the new ACLEntry in the Datastore to **Committed**.
 - ii. If not successful, update the Status to **CommitFailed** and the FailureCode to the returned error. The error SHALL be handled in a subsequent **Node Refresh**.

11.24.7.20. RemoveACLFromNode Command

Upon receipt, this SHALL remove an ACL from a node in the **Joint Fabric** Datastore Cluster of the accessing fabric.

The data for this command is as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	ListID	uint16	all			M
1	NodeID	node-id	all			M

ListID Field

This field SHALL indicate the unique identifier for the [DatastoreACLEntryStruct](#) to be removed from the Datastore's list of DatastoreACLEntry.

NodeID Field

This field SHALL indicate the unique identifier for the node from which the ACL is to be removed.

The Datastore SHALL:

1. Confirm that a [DatastoreNodeInformationEntryStruct](#) exists for the given NodeID in the NodeList attribute, and if not, then this command SHALL fail with a NOT_FOUND status code.
2. Ensure the ACL List for the given NodeID does not include the given ACLEntry. If it does:
 - a. Update the status to [DeletePending](#) for the given ACLEntry in the ACL List.
 - b. Remove this ACLEntry from the given Node ID.
 - i. If this succeeds, remove the given ACLEntry from the Node ACL List.
 - ii. If not successful, update the Status to [CommitFailed](#) and the FailureCode to the returned error. The error SHALL be handled in a subsequent [Node Refresh](#).

11.25. Joint Fabric Administrator Cluster

An instance of the Joint Fabric Administrator Cluster only applies to Joint Fabric Administrator nodes fulfilling the role of Anchor CA.

11.25.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.25.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	JFADMIN

11.25.3. Cluster ID

ID	Name
0x0753	Joint Fabric Administrator

11.25.4. Data Types

11.25.4.1. ICACResponseStatusEnum Type

This data type is derived from [enum8](#).

This enumeration is used by the [AddICAC](#) command to convey the outcome of this cluster's operations.

Value	Name	Summary	Conformance
0	OK	No error	M
1	InvalidPublicKey	Public Key in the ICAC is invalid	M
2	InvalidICAC	ICAC chain validation failed / ICAC DN Encoding rules verification failed	M

11.25.4.2. TransferAnchorResponseStatusEnum Type

This data type is derived from [enum8](#).

This enumeration is used by the [TransferAnchorResponse](#) command to convey the detailed outcome of this cluster's [TransferAnchorRequest](#) command.

Value	Name	Summary	Conformance
0	OK	No error	M
1	TransferAnchorStatusDatastoreBusy	Anchor Transfer was not started due to ongoing Datastore operations	M
2	TransferAnchorStatusNoUserConsent	User has not consented for Anchor Transfer	M

11.25.4.3. ICACCSRResponseStatusCodeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	OK	No error	M

Value	Name	Summary	Conformance
1	Busy	Could not be completed because another commissioning is in progress	M
2	PAKEParameterError	Provided PAKE parameters were incorrectly formatted or otherwise invalid	M
3	WindowNotOpen	The Joint Commissioning Method commissioning window is not currently open	M
4	VIDNotVerified	ICACSRRequest command has been invoked by a peer against which Fabric Table VID Verification hasn't been executed	M
5	InvalidAdministratorFabricIndex	OpenJointCommissioningWindow command has been invoked but the AdministratorFabricIndex field has the value of null	M

11.25.5. Attributes

ID	Name	Type	Constraint	Quality	Default	Access	Conformance
0x00	AdministratorFabricIndex	fabric-idx	1 to 254	X		A	M

11.25.5.1. AdministratorFabricIndex

The `AdministratorFabricIndex` attribute SHALL indicate the `FabricIndex` from the Endpoint 0's Operational Cluster **Fabrics** attribute (i.e. the Fabric Table) which is associated with the [JointFabric](#). This field SHALL have the value of null if there is no fabric associated with the [JointFabric](#).

11.25.6. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	ICACCSRRequest	client ⇒ server	ICACCSRResponse	A	M
0x01	ICACCSRResponse	client ⇐ server	N		M
0x02	AddICAC	client ⇒ server	ICACResponse	A	M
0x03	ICACResponse	client ⇐ server	N		M
0x04	OpenJoint-CommissioningWindow	client ⇒ server	Y	A	M
0x05	TransferAnchorRequest	client ⇒ server	TransferAnchorResponse	A	M
0x06	TransferAnchorResponse	client ⇐ server	N	A	M
0x07	TransferAnchorComplete	client ⇒ server	Y	A	M
0x08	Announce-JointFabricAdministrator	client ⇒ server	Y	A	M

11.25.6.1. ICACCSRRequest Command

This command SHALL be generated during [Joint Commissioning Method](#) and subsequently be responded in the form of an [ICACCSRResponse](#) command.

Effect on Receipt

This command SHALL be received over a [CASE](#) session otherwise it SHALL fail with an INVALID_COMMAND status code.

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

If the [FabricFabric Table Vendor ID Verification Procedure](#) has not been executed against the initiator of this command, the command SHALL fail with a [JfVidNotVerified](#) status code SHALL be sent back to the initiator.

If the [AdministratorFabricIndex](#) attribute has the value of null, the command SHALL fail with a [InvalidAdministratorFabricIndex](#) status code sent back to the initiator.

If a prior AddICAC command was successfully executed within the fail-safe timer period, then this command SHALL fail with a CONSTRAINT_ERROR status code sent back to the initiator.

11.25.6.2. ICACSRResponse Command

This command SHALL be generated in response to a [ICACSRRequest](#) command.

Check [ICAC Cross Signing](#) for details about the generation of the [ICACSR](#).

ID	Name	Type	Constraint	Quality	Default	Conformance
0	StatusCode	ICACSRResponseStatusCodeEnum	all			M
1	ICACSR	octstr	max 600			StatusCode == OK, O

ICACSR field

This field SHALL be a DER-encoded octet string of a properly encoded [PKCS #10](#) Certificate Signing Request (CSR).

11.25.6.3. AddICAC Command

This command SHALL be generated and executed during [Joint Commissioning Method](#) and subsequently be responded in the form of an [ICACResponse](#) command.

A Commissioner or Administrator SHALL issue this command after issuing the [ICACSRRequest](#) command and receiving its response.

A Commissioner or Administrator SHALL issue this command after performing the Attestation Procedure, Fabric Table VID Verification and after validating that the peer is authorized to act as an Administrator in its own Fabric.

Check [ICA Cross Signing](#) for details about the generation of [ICACValue](#).

ID	Name	Type	Constraint	Quality	Default	Conformance
0	ICACValue	octstr	max 400			M

ICACValue Field

This field SHALL contain an [ICAC](#) encoded using [Matter Certificate Encoding](#).

Effect on Receipt

This command SHALL be received over a [CASE](#) session otherwise it SHALL fail with an INVALID_COMMAND status code.

If this command is received without an armed fail-safe context (see [Section 11.10.7.2, “ArmFail-Safe”](#)), then this command SHALL fail with a FAILSAFE_REQUIRED status code sent back to the initiator.

If a prior AddICAC command was successfully executed within the fail-safe timer period, then this command SHALL fail with a CONSTRAINT_ERROR status code sent back to the initiator.

Upon receipt, the ICACValue SHALL be validated in the following ways:

- 1. Verify the ICAC using `Crypto_VerifyChain(certificates = [ICACValue, RootCACertificate])` where RootCACertificate is the associated RCAC of the `accessing fabric`. If this check fails, the error status SHALL be InvalidICAC.
- 2. The public key of the ICAC SHALL match the public key present in the last `ICACSRResponse` provided to the Administrator that sent the `AddICAC` command. If this check fails, the error status SHALL be InvalidPublicKey.
- 3. The `DN Encoding Rules` SHALL be validated for the ICAC. If this check fails, the error status SHALL be InvalidICAC.

If any of the above validation checks fail, the server SHALL immediately respond to the client with an `ICACResponse`. The `StatusCode` field of the `ICACResponse` SHALL be set to the error status value specified in the above validation checks.

If all the checks succeed, then the ICACValue SHALL be used as described in the `Joint Commissioning Method`.

11.25.6.4. ICACResponse Command

This command SHALL be generated in response to the `AddICAC` command.

ID	Name	Type	Constraint	Quality	Default	Conformance
0	StatusCode	ICACResponseStatusEnum				M

StatusCode Field

This field SHALL contain an `ICACResponseStatusEnum` value representing the status of the `AddICAC` operation.

11.25.6.5. OpenJointCommissioningWindow Command

NOTE

This is an alias onto the `OpenCommissioningWindow` command within the Joint Fabric Administrator Cluster. Refer to the `OpenCommissioningWindow` command for a description of the command behavior and parameters.

This command SHALL fail with a `InvalidAdministratorFabricIndex` status code sent back to the initiator if the `AdministratorFabricIndex` attribute has the value of null.

The parameters for OpenJointCommissioningWindow command are as follows:

ID	Name	Type	Constraint	Quality	Default	Conformance
0	Commissioning-Timeout	uint16	desc			M
1	PAKEPass-codeVerifier	octstr	97			M
2	Discriminator	uint16	max 4095			M
3	Iterations	uint32	1000 to 100000			M
4	Salt	octstr	16 to 32			M

11.25.6.6. TransferAnchorRequest Command

This command SHALL be sent by a candidate Joint Fabric Anchor Administrator to the current Joint Fabric Anchor Administrator to request transfer of the Anchor Fabric.

11.25.6.7. TransferAnchorResponse Command

This command SHALL be generated in response to the [Transfer Anchor Request](#) command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	StatusCode	TransferAnchorResponseStatusEnum				M

11.25.6.8. TransferAnchorComplete Command

This command SHALL indicate the completion of the transfer of the Anchor Fabric to another Joint Fabric Ecosystem Administrator.

11.25.6.9. AnnounceJointFabricAdministrator Command

This command SHALL be used for communicating to client the endpoint that holds the Joint Fabric Administrator Cluster.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	EndpointID	endpoint-no	all			M

This field SHALL contain the unique identifier for the endpoint that holds the Joint Fabric Administrator Cluster.

11.26. Commissioner Control Cluster

The Commissioner Control Cluster supports the ability for clients to request the commissioning of themselves or other nodes onto a fabric which the cluster server can commission onto. An example use case is ecosystem to ecosystem [Fabric Synchronization setup](#).

The generalized flow supported by the Commissioner Control Cluster can be seen in the following diagram.




Figure 110. Commissioner Control Cluster - General Flow

11.26.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial revision

11.26.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	CCTRL

11.26.3. Cluster ID

ID	Name
0x0751	Commissioner Control

11.26.4. Data Types

11.26.4.1. SupportedDeviceCategoryBitmap Type

This data type is derived from map32.

Bit	Name	Summary	Conformance
0	FabricSynchronization	Aggregators which support Fabric Synchronization may be commissioned.	O

FabricSynchronization Bit

The FabricSynchronization bit SHALL be set to 1 if and only if the server supports commissioning nodes that support [Fabric Synchronization](#).

11.26.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Supported-DeviceCategories	Supported-DeviceCategory-Bitmap	all		0	R M	M

11.26.5.1. SupportedDeviceCategories Attribute

This attribute SHALL indicate the device categories specified in SupportedDeviceCategoryBitmap that are supported by this Commissioner Control Cluster server.

A client SHALL NOT send the RequestCommissioningApproval command if the intended node to be commissioned does not conform to any of the values specified in SupportedDeviceCategories.

11.26.6. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	RequestCommissioningApproval	client ⇒ server	Y	M	M
0x01	CommissionNode	client ⇒ server	ReverseOpen-CommissioningWindow	M	M
0x02	ReverseOpenCommissioningWindow	server ⇒ client	N		M

11.26.6.1. RequestCommissioningApproval Command

This command is sent by a client to request approval for a future CommissionNode call. This is required to be a separate step in order to provide the server time for interacting with a user before informing the client that the CommissionNode operation may be successful.

If the command is not executed via a CASE session, the command SHALL fail with a status code of UNSUPPORTED_ACCESS.

The server MAY request approval from the user, but it is not required.

The server SHALL always return SUCCESS to a correctly formatted RequestCommissioningApproval command, and then generate a CommissioningRequestResult event associated with the command's accessing fabric once the result is ready.

Clients SHOULD avoid using the same RequestID. If the RequestID and client NodeID of a RequestCommissioningApproval match a previously received RequestCommissioningApproval and the server has not returned an error or completed commissioning of a device for the prior request, then the server SHOULD return FAILURE.

The parameters for RequestCommissioningApproval command are as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	RequestID	uint64				M
1	VendorID	vendor-id				M
2	ProductID	uint16				M
3	Label	string	max 64			O

11.26.6.2. RequestID

This SHALL be set by the client to be used again in future interactions.

11.26.6.3. VendorID / ProductID fields

These fields SHALL contain the VendorID and ProductID which match the values in the Device Attestation Certificate of the device that the client would like to have commissioned in the subsequent ReverseOpenCommissioningWindow command.

11.26.6.4. Label

The Label field SHOULD be set to a user readable name for the device that the client would like to have commissioned in the subsequent ReverseOpenCommissioningWindow command. This may be an ecosystem defined label or user provided name.

For example, in [FabricSynchronization](#) between two ecosystems (E1 as initial commissioner and E2 as initial commissionee), the label here represents the node on E1 that will be commissioned by E2 in the reverse commissioning operation after E1 has initially commissioned E2. This enables E2 to provide the user with a label for the node on E1 prior to completing the reverse commissioning

operation.

11.26.6.5. CommissionNode Command

This command is sent by a client to request that the server begins commissioning a previously approved request.

The server SHALL return FAILURE if the CommissionNode command is not sent from the same NodeID and on the same fabric as the RequestCommissioningApproval or if the provided RequestID to CommissionNode does not match the value provided to RequestCommissioningApproval.

If the command is not executed via a CASE session, the command SHALL fail with a status code of UNSUPPORTED_ACCESS.

Upon receipt, the server SHALL respond with ReverseOpenCommissioningWindow if CommissioningRequestResult was generated with StatusCode of SUCCESS for the matching RequestID field and NodeID of the client.

The server SHALL return FAILURE if the CommissionNode command is received after the server has already responded to a client with ReverseOpenCommissioningWindow for a matching RequestID field and NodeID of the client unless the client has sent another RequestCommissioningApproval and received an additional CommissioningRequestResult.

The parameters for CommissionNode command are as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	RequestID	uint64				M
1	Response-TimeoutSeconds	uint16	30 to 120		30	M

11.26.6.6. RequestID Field

This SHALL be set by the client to be the same value provided to RequestCommissioningApproval.

11.26.6.7. ResponseTimeoutSeconds Field

Timeout in seconds that the client SHALL wait to receive a response before considering the responses invalid.

11.26.6.8. ReverseOpenCommissioningWindow Command

When received within the timeout specified by ResponseTimeoutSeconds in the CommissionNode command, the client SHALL open a commissioning window on a node which matches the VendorID and ProductID provided in the associated RequestCommissioningApproval command.

When commissioning this node, the server SHALL check that the VendorID and ProductID fields provided in the RequestCommissioningApproval command match the VendorID and ProductID

attributes of the Basic Information Cluster which have already been verified during the [Device Attestation Procedure](#). If they do not match, the server SHALL NOT complete commissioning and SHOULD indicate an error to the user.

NOTE

This is an alias onto the [OpenCommissioningWindow](#) command within the Administrator Commissioning Cluster. Refer to the [OpenCommissioningWindow](#) command for a description of the command behavior and parameters.

The parameters for ReverseOpenCommissioningWindow command are as follows:

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Commissioning-Timeout	uint16	desc			M
1	PAKEPass-codeVerifier	octstr	97			M
2	Discriminator	uint16	max 4095			M
3	Iterations	uint32	1000 to 100000			M
4	Salt	octstr	16 to 32			M

11.26.7. Events

ID	Name	Priority	Quality	Access	Conformance
0x00	CommissioningRequestResult	INFO		M S	M

11.26.7.1. CommissioningRequestResult Event

This event SHALL be generated by the server following a RequestCommissioningApproval command which the server responded to with SUCCESS.

NOTE

The approval is valid for a period determined by the manufacturer and characteristics of the node presenting the Commissioner Control Cluster. Clients SHOULD send the CommissionNode command immediately upon receiving a CommissioningRequestResult event.

Access Modifier: Fabric-Sensitive

ID	Name	Type	Constraint	Quality	Fallback	Conformance
----	------	------	------------	---------	----------	-------------

Access Modifier: Fabric-Sensitive						
0	RequestID	uint64	all			M
1	ClientNodeID	node-id	all			M
2	StatusCode	status	desc			M

11.26.7.2. RequestID / ClientNodeID Fields

The RequestID SHALL match the RequestID provided to RequestCommissioningApproval and the ClientNodeID shall match the NodeID of the client which generated the RequestCommissioningApproval command.

11.26.7.3. StatusCode Field

The server SHALL set this field to SUCCESS if the server is ready to begin commissioning the requested device, TIMEOUT if the server timed out due to user inaction or FAILURE for any other error.

11.27. Groupcast Cluster

The Groupcast Cluster defines a unified, simpler, and more scalable mechanism for configuring, managing, and using groups in Matter. It is designed to replace the legacy Groups cluster, addressing long-standing limitations that have made group configuration complex and inconsistently supported across devices and ecosystems.

The Groupcast Cluster consolidates group-related functionality that was previously distributed across multiple clusters into a single, coherent model for membership, addressing, key management, and ACL permissions. It introduces the usage of a unique, IANA-defined multicast address, usable by all groups, improving support for devices with limited multicast address capacity. The cluster also defines a streamlined key management and access control model that reduces storage requirements and simplifies controller logic, while remaining compatible with existing security concepts.

In addition, the Groupcast Cluster explicitly supports specialization of device roles, allowing endpoints to operate as group senders, receivers, or both. This enables low-power and resource-constrained devices to participate in group communication without incurring unnecessary state or processing overhead. Together, these design elements provide a clearer and more interoperable basis for group functionality and form the replacement for the legacy Groups cluster.

NOTE Support for Groupcast cluster is provisional.

11.27.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial release, replacing functionality previously in Groups and Group Key Management clusters and addressing issues with application of those clusters (TCR 11433)

11.27.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	GC

11.27.3. Cluster ID

Cluster ID	Cluster Name	Conformance
0x0065	Groupcast	P

11.27.4. Features

This cluster SHALL support the FeatureMap bitmap attribute as defined below.

Bit	Code	Feature	Conformance	Summary
0	LN	Listener	P, O.a+	Supports joining a multicast group of nodes as a listener.
1	SD	Sender	P, O.a+	Supports sending multicast message to a targeted group of nodes.
2	PGA	PerGroup	P, O	Supports Per-Group multicast addresses.

11.27.4.1. Listener Feature

This feature indicates that the device can join one or more Groupcast groups and receive multicast messages targeted to those groups.

11.27.4.2. Sender Feature

This feature indicates the ability to send multicast messages to one or more targeted groups of nodes to which it belongs. Being a sender does not imply the ability to listen to messages sent to those multicast addresses.

PerGroup Feature

The PerGroup feature indicates that the node MAY be configured to use [Multicast Address](#) based on the Group ID instead of the [IANA Assigned IPv6 Multicast Address](#). This feature enables the formation of groups that are interoperable between pre-Groupcast nodes and post-Groupcast nodes.

11.27.5. Data Types

11.27.5.1. MulticastAddrPolicyEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	IanaAddr	Indicates group uses the IANA-assigned multicast address (default)	M
1	PerGroup	Group uses multicast address scoped to Fabric ID and Group ID	PGA

IanaAddr Value

The [Multicast Address](#) of the group SHALL be the [IANA Assigned IPv6 Multicast Address](#).

The **IanaAddr** method minimizes the number of multicast addresses to which a node must subscribe, reducing the state it must maintain. This approach comes at the cost of receiving multicast traffic for all groups that use this policy, potentially from any fabric. The node will then filter this traffic at the message layer by attempting decryption with its available group keys. Due to various network infrastructure scalability limits, such as the maximum number of [MPL](#) registrations a Border Router can support, a controller SHOULD default to configure all groups to use the AllNodes address.

PerGroup Value

The 16-bit Group Identifier of the [Multicast Address](#) SHALL be the Group ID of the group.

The **PerGroupID** multicast address policy maximizes network-layer filtering, ensuring that a node only receives multicast traffic for groups to which it has subscribed. This method requires a node to join a unique multicast group for each (fabric, group) combination that uses this policy. This method SHOULD be reserved for the following purposes: . Use by groups with listener members that require extra filtering at the network layer, such as low power devices. . Use by groups with members that include legacy listeners running a version 2 or earlier of this cluster. Controllers SHOULD limit the number of groups that use this policy within a network environment to account for infrastructure limits mentioned above.

11.27.5.2. GroupcastTestingEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	DisableTesting	Disable all test event generation	M
1	EnableListenerTesting	Enable Listener test event generation	LN
2	EnableSenderTesting	Enable Sender test event generation	SD

See the [GroupcastTesting](#) command for a description of these operations.

11.27.5.3. GroupcastTestResultEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary	Conformance
0	Success	There was no failure	M
1	GeneralError	Other error during processing not otherwise covered	M
2	MessageReplay	Message counter was $\leftarrow$ counter from last received	M
3	FailedAuth	Message authentication failed	M
4	NoAvailableKey	No key was found to process the message	M
5	SendFailure	Issue with attempting to send a multicast IPv6 message	M

11.27.5.4. MembershipStruct Type

Access Modifier: Fabric-Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	GroupID	group-id	min 1				M
1	Endpoints	list[end-point-no]	max 255				LN
2	KeySetID	uint16	min 1			S	M
3	HasAuxiliaryACL	bool	all				LN

Access Modifier: Fabric-Scoped							
4	McastAddrPolicy	MulticastAddrPolicyEnum	all		0		M

GroupID Field

This field SHALL indicate an identifier for the multicast group within the fabric. [Group Identifier](#) is a 16-bit value that SHALL be assigned by the administrator when the group is created.

Endpoints Field

This field SHALL indicate a list of endpoint numbers that are members of the multicast group. This field is only relevant for Listeners. The content of this field SHALL be considered in [Request Path Expansion](#) when processing the relevant Interaction for the received Groupcast message for the associated GroupID. When the device is acting as a Listener as member of the group, the Endpoints field SHALL list at least one endpoint and SHALL NOT contain more than 255 endpoints.

If both the Listener and Sender bits are set in the FeatureMap attribute, then the list MAY be either empty or missing, until such time that the JoinGroup command is done with a non-empty Endpoints field. With these FeatureMap bits both set, the empty or missing list indicates that the group is joined for sending, but since there are not endpoints available in the group membership, no group messages will be possible to receive.

KeySetID Field

This field SHALL indicate the value used to identify the group operational key used for this group.

This field maps directly to the GroupKeySetID as defined in the [Group Key Management](#) cluster. This field is provided by the administrator with the provided Key within the [JoinGroup](#) or [UpdateGroupKey](#) commands. This field SHALL be used by administrators of the fabric to determine if the server has the correct group operational key for their group. If not, the administrator MAY update the GroupKey using the [UpdateGroupKey](#) command to recover communication.

HasAuxiliaryACL Field

This field SHALL indicate if AuxiliaryACL entries are to be generated. The field SHALL be set to true if the group membership was configured to cause AuxiliaryACL entries to be generated, or false otherwise. See [Groupcast Auxiliary ACL Handling](#) for the handling of this field.

See also the [ConfigureAuxiliaryACL](#) command.

McastAddrPolicy Field

This field SHALL indicate how the [IPv6 Multicast Address](#) SHALL be constructed for this group.

11.27.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Membership	list[MembershipStruct]	desc	N		R V F	P, M
0x0001	MaxMembershipCount	uint16	min 10	F		R V	P, M
0x0002	MaxMcastAddrCount	uint16	min 1	F		R V	P, M
0x0003	UsedMcastAddrCount	uint16	all	F		R V	P, M
0x0004	FabricUnderTest	fabric-idx	all			R V	P, M

11.27.6.1. Membership Attribute

This attribute SHALL indicate the list of groups memberships currently active on the node.

Each entry specifies the group ID and the list of the endpoints that participate in that group, as well as policy for that group.

For any single fabric, the server SHALL limit the total number of GroupIDs used across all entries in the Membership attribute to no more than half (rounded down) of the [MaxMembershipCount](#) value.

If processing a command would cause this per-fabric limit to be exceeded, then the resource exhaustion behavior of the associated command SHALL apply.

Since every [MembershipStruct](#) entry only supports up to 255 endpoints in the case where the Listener feature is enabled in the FeatureMap, there MAY be situations where the server has to employ more than one entry to represent the full list of joined endpoints in the list. In those cases:

- The entries associated with such a group SHALL be consecutive.
- These entries SHALL be identical in all fields except the Endpoints field.
- These entries SHALL NOT have any intersection of EndpointID values between their Endpoints field.
- The maximum number of entries over which a group membership is split SHALL NOT be more than the minimum number of entries needed to fully encode the Endpoints list while respecting that no entry has more than 255 Endpoints listed.
 - Therefore, the maximum number of entries for that group SHALL be $\text{ceil}(\text{length}(\text{Endpoints}) / 255)$.

The actual number of entries in the list across all fabrics SHALL be at most $\text{ceil}(\text{num_end-}$

`points_in_node / 255) * MaxMembershipCount.`

For example, if a Group 123 is joined on endpoints 1 through 300 (i.e. a large group on a large bridge), the following Membership list contents would both be valid examples, and there MAY be more valid entries matching the rules.

```
# Example 1: Split between 1..256 and 257..300
[
  # Other unrelated group 17
  {
    GroupID: 17,
    Endpoints: [4, 5],
    KeySetID: 1,
    HasAuxiliaryACL: true,
    McastAddrPolicy: IanaAddr,
    FabricIndex: 1
  },
  # First entry for group 123
  {
    GroupID: 123,
    Endpoints: [1, 2, 3, ..., 256], # 255 entries
    KeySetID: 1,
    HasAuxiliaryACL: true,
    McastAddrPolicy: IanaAddr,
    FabricIndex: 1
  },
  # Second "overflow" entry for group 123, consecutive to previous.
  {
    GroupID: 123,
    Endpoints: [257, 258, 259, ..., 300], # 44 entries
    KeySetID: 1,
    HasAuxiliaryACL: true,
    McastAddrPolicy: IanaAddr,
    FabricIndex: 1
  }
]

# Example 2: Split between 1..150 and 151..300
[
  # First entry for group 123
  {
    GroupID: 123,
    Endpoints: [1, 2, 3, ..., 150], # 150 endpoints
```

```

    KeySetID: 1,
    HasAuxiliaryACL: true,
    McastAddrPolicy: IanaAddr,
    FabricIndex: 1
  },
  # Second "overflow" entry for group 123, consecutive to previous.
  {
    GroupID: 123,
    Endpoints: [151, 152, 153, ..., 300], # 150 endpoints also
    KeySetID: 1,
    HasAuxiliaryACL: true,
    McastAddrPolicy: IanaAddr,
    FabricIndex: 1
  },
  # Other unrelated group 54. Note the group order between GroupIDs is not guaranteed.
  {
    GroupID: 54,
    Endpoints: [19, 29],
    KeySetID: 1,
    HasAuxiliaryACL: true,
    McastAddrPolicy: IanaAddr,
    FabricIndex: 1
  }
]

```

The following examples would be ILLEGAL:

```

# Example 3: Illegal split of large group entries across non-consecutive list entries
[
  # First entry for group 123
  {
    GroupID: 123,
    Endpoints: [1, 2, 3, ..., 150], # 150 endpoints
    KeySetID: 1,
    HasAuxiliaryACL: true,
    McastAddrPolicy: IanaAddr,
    FabricIndex: 1
  },
  # Other unrelated group 54.
  {
    GroupID: 54,
    Endpoints: [19, 29],

```

```
    KeySetID: 1,
    HasAuxiliaryACL: true,
    McastAddrPolicy: IanaAddr,
    FabricIndex: 1
  },
# ILLEGAL: Second "overflow" entry for group 123, due to not being
#           consecutive to previous entry for group 123.
{
  GroupID: 123,
  Endpoints: [151, 152, 153, ..., 300], # 150 endpoints also
  KeySetID: 1,
  HasAuxiliaryACL: true,
  McastAddrPolicy: IanaAddr,
  FabricIndex: 1
}
]

# Example 4: Illegal split of large group entries across too many list entries
[
# First entry for group 123
{
  GroupID: 123,
  Endpoints: [1, 2, 3, ..., 256], # 255 endpoints
  KeySetID: 1,
  HasAuxiliaryACL: true,
  McastAddrPolicy: IanaAddr,
  FabricIndex: 1
},
# Second entry for group 123
{
  GroupID: 123,
  Endpoints: [257], # 1 endpoint
  KeySetID: 1,
  HasAuxiliaryACL: true,
  McastAddrPolicy: IanaAddr,
  FabricIndex: 1
},
# Third entry for group 123
# ILLEGAL: Used 3 list entries when  $\text{ceil}(300 / 255) == 2$ .
{
  GroupID: 123,
  Endpoints: [258, 259, ..., 300], # 43 endpoints
```

```
    KeySetID: 1,  
    HasAuxiliaryACL: true,  
    McastAddrPolicy: IanaAddr,  
    FabricIndex: 1  
  }  
]
```

11.27.6.2. MaxMembershipCount Attribute

This attribute SHALL indicate the maximum number of Groups which can be joined and appear in entries of the [Membership](#) attribute.

This attribute does not specify the maximum number of list entries in the Membership attribute list, but rather the maximum number of different GroupID values which can appear across all entries of the list.

11.27.6.3. MaxMcastAddrCount Attribute

This attribute SHALL indicate the maximum number of unique multicast addresses the node can support. The value of this attribute SHALL be at least 1, to support the [IANA Assigned IPv6 Multicast Address](#).

When the PerGroup feature is supported, this attribute SHOULD be equal to [MaxMembershipCount](#). However, the value MAY be less than [MaxMembershipCount](#). When the PerGroup feature is supported, the value of this attribute SHALL be at least 4.

11.27.6.4. UsedMcastAddrCount Attribute

This attribute SHALL indicate the number of unique multicast addresses currently in use by the Groupcast cluster. This count SHALL include the [IANA Assigned IPv6 Multicast Address](#) if at least one group is configured to use the [IanaAddr](#) policy. This count SHALL include one unique multicast address for each group configured to use the [PerGroup](#) policy. The value of this attribute SHALL NOT exceed [MaxMcastAddrCount](#).

11.27.6.5. FabricUnderTest Attribute

This attribute SHALL indicate the FabricIndex of the fabric currently testing the Groupcast feature with the [GroupcastTesting](#) command.

The last accessing fabric to have invoked a successful TestGroupcast command that caused testing mode to be engaged SHALL be the value indicated.

Otherwise, if Groupcast testing is currently disabled, this attribute SHALL have a value of zero, which is the invalid FabricIndex value.

This attribute SHALL be set to zero when the server initializes.

11.27.7. Commands

ID	Name	Direction	Response	Access	Conformance
0x00	JoinGroup	client ⇒ server	Y	M F	P, M
0x01	LeaveGroup	client ⇒ server	LeaveGroupResponse	M F	P, M
0x02	Leave-GroupResponse	client ⇐ server	N	F	P, M
0x03	UpdateGroup-Key	client ⇒ server	Y	M F	P, M
0x04	ConfigureAuxiliaryACL	client ⇒ server	Y	A F	P, LN
0x05	Group-castTesting	client ⇒ server	Y	A F	P, M

11.27.7.1. JoinGroup Command

This command SHALL be used to instruct the server to join a multicast group. It provides a comprehensive way to add all information required to create a new group in one command. This command SHALL be used to create a new GroupID or to add specified endpoints to an existing GroupID. This command MAY also be used to change the OperationalGroupKey associated with an existing GroupID, but if this is the only operation desired, the [UpdateGroupKey](#) command SHOULD be used instead. This command SHALL be used following the rules defined in [Groupcast Key Management](#) section.

This command SHALL have the following data fields subject to the listed conformance.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupID	group-id	min 1			M
1	Endpoints	list[end-point-no]	desc			M
2	KeySetID	uint16	min 1			M
3	Key	octstr	16			O
4	UseAuxiliaryACL	bool	all			[LN]
5	ReplaceEndpoints	bool	all			[LN]
6	McastAddr-Policy	MulticastAddrPolicyEnum	all			O

GroupID Field

This field SHALL indicate the destination group to which the server is to be added.

Endpoints Field

This field SHALL indicate the list of endpoints that SHALL receive this multicast group's messages. When joining a group as a Sender only, this field SHALL be empty. For Listeners, this field SHALL list at least 1 endpoint and up to 20 endpoints. The RootEndpoint (Endpoint 0) SHALL never be included in the Endpoints list as groupcast commands SHALL not interact with clusters on the RootEndpoint.

KeySetID Field

This field SHALL indicate the value used to identify the OperationalGroupKey used for this group.

This field is generated by the administrator in any manner it deems appropriate, ensuring that there is no collision within the fabric. Some examples may include but are not limited to:

- Truncated Key Hash
- Randomly Generated Value
- Key creation count

NOTE

This field represents maps to the GroupKeySetID as defined in the [Group Key Management](#) cluster.

Key Field

This field SHALL indicate the value used to derive the OperationalGroupKey.

This key SHALL be new for the group, with a unique KeySetID that does not exist on the node yet. When a new key is added using this field, a [GroupKeySetStruct](#) SHALL be added automatically to the [Group Key Management](#) cluster with the following entries:

- GroupKeySetID = KeySetID field
- GroupKeySecurityPolicy = TrustFirst
- EpochKey0 = Key
- EpochStartTime0 = 1
- EpochKey1 = null
- EpochStartTime1 = null
- EpochKey2 = null
- EpochStartTime2 = null The derived OperationalGroupKey SHALL be persistently stored for the lifetime of the group; however, the InputKey itself SHALL NOT be stored. This field MAY be omitted to join endpoints to an existing GroupID or to associate the group to a different, pre-existing KeySetID.

UseAuxiliaryACL Field

When this field is present and set to true, necessary [AuxiliaryACL](#) entries SHALL be generated in the [Access Control Cluster](#) referring to each endpoint in the group identified by the given GroupID. When this field is present and set to false, all previously auto-generated AuxiliaryACL entries for the endpoints in this group SHALL be removed. If this field is omitted, there SHALL NOT be any change. See the Effect on Receipt section of this command for more detail.

ReplaceEndpoints Field

This field SHALL indicate whether the Endpoints list of the membership is replaced by the command's Endpoints field, or whether those endpoints are added to the existing membership. See [Section 11.27.7.1.7, "Effect on Receipt"](#).

Effect on Receipt

On receipt of this command, the server SHALL perform the following procedure:

1. If there is no available space in the server's Membership attribute to add a new entry or to extend the Endpoints list of an existing group, the server SHALL stop command processing and SHALL return with status RESOURCE_EXHAUSTED.
2. If the Endpoints field is empty:
 - a. If Sender feature is not supported by the server, the server SHALL stop command processing and SHALL return with status CONSTRAINT_ERROR.
3. If the Endpoints field is not empty:
 - a. If Listener feature is not supported by the server, the server SHALL stop command processing and SHALL return with status CONSTRAINT_ERROR.
 - b. The server SHALL verify that all endpoints listed in the Endpoints field are valid.
 - c. If any endpoint is invalid or is the RootEndpoint (Endpoint 0), the server SHALL stop command processing and SHALL return with status UNSUPPORTED_ENDPOINT.
 - d. If the UseAuxiliaryACL field is provided and the command was invoked by a client that does not have at least Administer privilege granted, then the server SHALL stop command processing and SHALL return with status UNSUPPORTED_ACCESS.
4. If the Key field is provided:
 - a. The server SHALL verify that the KeySetID field does not already map to an existing OperationalGroupKey of the fabric. Otherwise, the server SHALL stop command processing and SHALL fail with a status of ALREADY_EXISTS.
 - b. The server SHALL use the 128-bit Key as the InputKey to derive the OperationalGroupKey for the specified GroupID (per [Operational Group Key Derivation](#)).
 - c. The server SHALL store the derived OperationalGroupKey to non-volatile memory. The given InputKey itself SHALL NOT be stored.
 - d. If the operational group key storage fails, the server SHALL stop command processing and SHALL return with status RESOURCE_EXHAUSTED if the failure was due to insufficient storage space, or with status FAILURE for any other internal errors.

5. If the Key field is omitted:
 - a. The server SHALL verify that the KeySetID field maps to an existing OperationalGroupKey of the fabric. Otherwise, the server SHALL stop command processing and SHALL fail with a status of NOT_FOUND.
6. The server SHALL associate the given endpoints to the group indicated by the GroupID, and update its Membership attribute to reflect these changes.
 - a. If ReplaceEndpoints is provided and its value is true, the provided endpoints list SHALL overwrite the existing Endpoints list of the Membership entry identified by the GroupID, omitting duplicates.
 - b. Else, the provided endpoints list SHALL be appended to the existing Endpoints list of the Membership entry identified by the GroupID, omitting duplicates.
 - c. For a new GroupID, the ReplaceEndpoints field's value has no effect since the initial list effectively "replaces" any existing membership.
7. If the UseAuxiliaryACL field is present, then the Effect on Receipt of the ConfigureAuxiliaryACL command SHALL be applied with the given GroupID and UseAuxiliaryACL field values, as if the ConfigureAuxiliaryACL command was invoked with those fields. Any failure to execute the equivalent of ConfigureAuxiliaryACL SHALL be ignored, leaving the effect of prior successful steps unchanged.
8. The server SHALL then return a SUCCESS status.

McastAddrPolicy Field

This field SHALL indicate the [IPv6 Multicast Address](#) that SHALL be used by the group. If omitted, the group SHALL use the [IANA-assigned Multicast Address](#).

11.27.7.2. LeaveGroup Command

This command SHALL allow a maintainer to request that the server withdraws itself or specific endpoints from a specific group or from all groups of this client's fabric.

This command SHALL have the following data fields subject to the listed conformance.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupID	group-id	all			M
1	Endpoints	list[end-point-no]	1 to 20			O

GroupID Field

This field SHALL indicate the group affected by this command. The [unspecified GroupID 0](#) SHALL indicate that all groups on the node for this fabric are affected by the command.

Endpoints Field

This field SHALL indicate a list of endpoints that SHALL leave the group identified by the GroupID

field. See 4. Omitting this field indicates that the server SHALL remove itself completely from the group identified by the GroupID field. See 5.

Effect on Receipt

On receipt of this command, the server SHALL follow these steps to withdraw itself or the specified endpoints from the group identified by the GroupID field.

1. If the GroupID is 0x00, the server SHALL confirm, by parsing its Membership list, that it is a member of at least one group associated with the client's fabric.
 - a. If the device is not a member of any group, the process SHALL stop and return the status NOT_FOUND.
 - b. Else, the server SHALL perform steps 2 through 5 as defined for each Membership entry of the fabric.
2. The server SHALL confirm, by parsing its Membership list, that it is a member of the group identified by the GroupID field.
3. If the GroupID is not found in the Membership list, the process SHALL stop and return the status NOT_FOUND.
4. If the Endpoints field is provided:
 - a. The server SHALL remove all endpoints specified in the Endpoints field from its Membership list entry matching the GroupID field.
 - b. If a listed endpoint is not a member of the group, the server SHALL ignore it and continue processing the remaining endpoints. the ignored endpoints SHALL be excluded from the [LeaveGroupResponse](#).
 - c. If a Membership entry is left with no endpoints, and the device is a Listener only.
 - i. The server SHALL remove the Membership attribute list entry associated with the received GroupID field.
 - ii. The server SHALL ensure that no AuxiliaryACL entries (see [Groupcast Auxiliary ACL Handling](#)) remain which refer to the GroupID provided.
 - iii. The server SHALL also delete all operational group keys of the removed group.
5. Else:
 - a. The server SHALL remove the Membership attribute list entry associated with the received GroupID field.
 - b. The server SHALL also delete all operational group keys of the removed group.
6. The server SHALL then return the list of endpoint numbers that were removed from the group through the [LeaveGroupResponse](#) command.

11.27.7.3. LeaveGroupResponse Command

This command SHALL allow the server to inform the client about the result of the LeaveGroup command.

This command SHALL have the following data fields subject to the listed conformance.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupID	group-id	all			M
1	Endpoints	list[end-point-no]	max 20			M

GroupID Field

This field SHALL indicate the group affected by the LeaveGroup command. The value of this field SHALL be the same value as the GroupID field in the LeaveGroup command triggering the response.

Endpoints Field

This field SHALL indicate a list of endpoints that were removed from the group identified by the GroupID field. Endpoints listed in the Endpoints field of the [LeaveGroup](#) command that were not members of the affected group SHALL be excluded from this response.

This field SHALL be an empty list if any of the following conditions is true: * The Endpoints list size exceeds the maximum constraint. * The GroupID field is 0, indicating multiple groups were affected, resulting in ambiguous Endpoints list content.

11.27.7.4. UpdateGroupKey Command

This command SHALL allow a fabric administrator to update the OperationalGroupKey associated with the existing group identified by GroupID, which is already joined. This command SHALL be used following the rules defined in [Groupcast Key Management](#) section.

This command SHALL have the following data fields subject to the listed conformance.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupID	group-id	min 1			M
1	KeySetID	uint16	min 1			M
2	Key	octstr	16			O

GroupID Field

This field SHALL indicate the destination group for which the operational key is being updated.

KeySetID Field

This field SHALL indicate the value used to identify the operational group key used for this group.

This field is generated by the administrator in any manner it deems appropriate, ensuring that there is no collision within the fabric. Some examples may include but are not limited to:

- Truncated Key Hash with deduplication such as Cuckoo Hashing

- Randomly Generated Value
- Key creation count

Key Field

This field SHALL indicate the 128-bit InputKey used to derive the OperationalGroupKey for the specified GroupID (per [Operational Group Key Derivation](#)). The derived OperationalGroupKey SHALL be persistently stored for the lifetime of the group; however, the InputKey itself SHALL NOT be stored.

Effect on Receipt

On receipt of this command, the server SHALL perform the following steps:

1. The server SHALL verify that it is a member of the group identified by the GroupID field, if the GroupID is not found in the Membership list, the process SHALL stop and return the status NOT_FOUND.
2. If the Key field is provided:
 - a. The server SHALL verify that the KeySetID field does not already map to an existing OperationalGroupKey for the fabric. Otherwise, the server SHALL stop command processing and SHALL fail with a status of ALREADY_EXISTS.
 - b. The server SHALL use the 128-bit Key as the InputKey used to derive the OperationalGroupKey for the specified GroupID (per [Operational Group Key Derivation](#)) and store persistently the obtained OperationalGroupKey.
 - c. If the operational group key storage fails, the server SHALL stop command processing and SHALL return with status RESOURCE_EXHAUSTED if the failure was due to insufficient storage space, or with status FAILURE for any other internal errors.
3. If the Key field is omitted:
 - a. The server SHALL verify that the KeySetID field maps to an existing OperationalGroupKey of the fabric. Otherwise, the server SHALL stop command processing and SHALL fail with a status of NOT_FOUND.
4. The server SHALL then return a SUCCESS status.

11.27.7.5. ConfigureAuxiliaryACL Command

This command SHALL allow an Administrator to enable or disable the generation of [AuxiliaryACL](#) entries in the [Access Control Cluster](#) based on the groups joined (see [Groupcast Auxiliary ACL Handling](#)).

This command SHALL have the following data fields subject to the listed conformance.

Field ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	GroupID	group-id	all			M

Field ID	Name	Type	Constraint	Quality	Fallback	Conformance
1	UseAuxiliaryACL	bool	all			M

GroupID field

This field SHALL indicate the group for which Auxiliary ACL entry generation state will be enabled/disabled.

UseAuxiliaryACL field

This field SHALL indicate if Auxiliary ACL entries SHALL be generated or removed.

If the field is set to true, necessary [AuxiliaryACL](#) entries SHALL be generated in the [Access Control Cluster](#) referring to each endpoint in the group identified by the given GroupID. Otherwise, if the field is set to false, all previously auto-generated AuxiliaryACL entries for the endpoints in this group SHALL be removed.

Effect on Receipt

On receipt of this command, the server SHALL perform the following steps:

1. The server SHALL verify that it is a member of the group identified by the GroupID field, if the GroupID is not found in the Membership list, processing SHALL stop and a status of NOT_FOUND SHALL be returned to the client.
2. If the UseAuxiliaryACL field is set to true, then the AuxiliaryACL entries for the group SHALL be generated according to [Groupcast Auxiliary ACL Handling](#). This MAY lead to no changes, depending on existing configuration of the AuxiliaryACL.
3. Otherwise if the UseAuxiliaryACL field is set to false, then the AuxiliaryACL entries for the group, if any already existed, SHALL be removed and the AuxiliaryACL attribute of the Access Control Cluster SHALL NOT contain any entries with AuxiliaryType set to Groupcast where the subject is set to this command's GroupID field.
4. For each [MembershipStruct](#) entry in the [Membership](#) attribute associated with the given GroupID, the HasAuxiliaryACL field SHALL be set to the value of this command's UseAuxiliaryACL field.
5. The server SHALL then return a SUCCESS status.

11.27.7.6. GroupcastTesting Command

This command SHALL allow an Administrator to configure test modes that allow validation of Groupcast communication.

This command SHALL have the following data fields subject to the listed conformance.

Field ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	TestOperation	Group-castTestingEnum	all			M
1	DurationSeconds	uint16	10 to 1200		60	O

TestOperation field

This field SHALL identify the type of TestOperation to execute. See [Effect on Receipt](#).

DurationSeconds field

This field SHALL indicate the time duration in seconds for the test operation until all testing operations reverts to a disabled state.

This field is ignored when the TestOperation field has a value of DisableTesting.

Effect on Receipt

The GroupcastTesting functionality allows for certification and production validation of groupcast communication.

Testing events SHALL only be reported for the last fabric for which testing was enabled.

The FabricIndex of the fabric currently "owning" the testing is indicated in the [FabricUnderTest](#), until testing is complete, at which point the [FabricUnderTest](#) returns to value 0.

On receipt of this command, the server SHALL perform the following steps:

- If TestOperation is EnableListenerTesting, then, for a period indicated in the field DurationSeconds, or until a GroupCastTesting command is received again from the same accessing fabric, the following process should be started in the Server, prior to responding immediately with a status of SUCCESS:
 1. Every successfully processed InvokeRequest received over Matter Groupcast communication and expanded to a given destination concrete path (1 per endpoint listed in the group membership, if the cluster was present and element was present) associated with the fabric indicated in the [FabricUnderTest](#) due to group memberships SHALL cause the generation of a [GroupcastTesting](#):
 - Set the SourceIpAddress set to the sender's IPv6 address from the UDP datagram.
 - Set the DestinationIpAddress set to the receiver's IPv6 group address from the UDP datagram.
 - Set the GroupID field set to the initial GroupID that led to InvokeRequest processing.
 - Set the EndpointID field set to the concrete path's endpoint ID.
 - Set the ClusterID field set to the concrete path's cluster ID.

- Set the ElementID field set to the concrete path's command ID.
 - Set the AccessAllowed field to true if the client was allowed by Access Control to invoke the command, and false otherwise.
 - Set the GroupcastTestResult field to Success.
2. Every incorrectly processed Groupcast message that failed decoding or decryption SHALL cause the generation of a [GroupcastTesting](#):
- Set the SourceIpAddress to the sender's IPv6 address from the UDP datagram.
 - Set the DestinationIpAddress to the receiver's IPv6 group address from the UDP datagram.
 - If the message was properly authenticated for the fabric associated with the [FabricUnderTest](#), set the GroupID field to the GroupID from the request.
 - If the message was improperly authenticated, set the GroupcastTestResult field to FailedAuth;
 - Else if the message was considered a duplicate of a previously received message according to processing of the [Global Group Encrypted Message Counters](#), set the GroupcastTestResult field to MessageReplay;
 - Else if no key was even found to attempt decrypting the message, set the GroupcastTestResult field to NoAvailableKey;
 - Else set the GroupcastTestResult field to GeneralError.
- Else if TestOperation is EnableSenderTesting, then, for a period indicated in the field DurationSeconds, or until a GroupCastTesting command is received again from the same accessing fabric, the following process should be started in the Server, prior to responding immediately with a status of SUCCESS:
 - Every successfully prepared Groupcast message initiated by the fabric indicated in the [FabricUnderTest](#) due to group memberships SHALL cause the generation of a [GroupcastTesting](#):
 - Set the SourceIpAddress to the sender's IPv6 address from the UDP datagram.
 - Set the DestinationIpAddress to the destination IPv6 group address from the UDP datagram.
 - Set the GroupID field to the GroupID associated with the message.
 - If the message failed to transmit after preparation, set the GroupcastTestResult to Send-Failure, otherwise to Success.
 - Else if TestOperation is DisableTesting, immediately terminate any ongoing Groupcast test processing, setting the [FabricUnderTest](#) to 0, and respond immediately with a status of Success.

11.27.8. Events

ID	Name	Priority	Quality	Access	Conformance
0x00	Group-castTesting	INFO		A S	P, M

11.27.8.1. GroupcastTesting Event

This event SHALL be generated during Groupcast testing processing after invocation of the [GroupcastTesting](#) command, under the conditions stated in that command's [Effect on Receipt](#) section. Some of the fields may be present or absent depending of the test mode involved.

This event SHALL contain the following fields:

Access Modifier: Fabric-Sensitive						
ID	Name	Type	Constraint	Quality	Fallback	Conformance
0x00	SourceIpAddress	ipv6adr	all			O
0x01	DestinationIpAddress	ipv6adr	all			O
0x02	GroupID	group-id	all			O
0x03	EndpointID	endpoint-id	all			O
0x04	ClusterID	cluster-id	all			O
0x05	ElementID	uint32	all			O
0x06	AccessAllowed	bool	all			O
0x07	GroupcastTestResult	GroupcastTestResultEnum	all			M

SourceIpAddress Field

This field, if present, SHALL be set to the source IPv6 address obtained from the UDP datagram of the groupcast message.

DestinationIpAddress Field

This field, if present, SHALL be set to the destination IPv6 group address obtained from the UDP datagram of a groupcast message.

GroupID Field

This field, if present, SHALL be set to the GroupID associated with the groupcast message. This represents the initial Group ID that led to InvokeRequest processing, the Group ID from the request, or the Group ID used for message transmission.

EndpointID Field

This field, if present, SHALL be set to the concrete path's endpoint ID derived from the processed request.

ClusterID Field

This field, if present, SHALL be set to the concrete path's cluster ID derived from the processed request.

ElementID Field

This field, if present, SHALL be set to the concrete path's element ID (Command ID or Attribute ID) derived from the processed request.

AccessAllowed Field

This field, if present, SHALL indicate whether the Groupcast sender was allowed to invoke the command by Access Control, in this Node. The value SHALL be set to true if the request was allowed, and false otherwise.

GroupcastTestResult Field

This field SHALL indicate the outcome of the groupcast operation or the specific error encountered. The value SHALL be one of the values defined in [GroupcastTestResultEnum](#).

11.27.9. Groupcast Key Management

The Groupcast relies on the [Group Key Management Cluster](#) to provide a complete key management solution for group communication. The KeySetID field in the Groupcast cluster refers to the GroupKeySetId in the Group Key Management cluster. Administrators and controllers MAY use either cluster to manage group keys, with both clusters providing visibility and manipulation of the KeySets for the fabric.

In the Groupcast cluster, the operational keys are managed through the [JoinGroup](#) and [UpdateGroupKey](#) commands. The [JoinGroup](#) and the [UpdateGroupKey](#) commands MAY create a new KeySet (by KeySetID), containing a single key, and assign or reassign an existing KeySet to a Group.

However, the Groupcast cluster does NOT support modifying the operational key of an existing GroupKeySet. If an administrator needs to provide a new key to an existing GroupKeySet, this operation SHALL be done through the [Group Key Management Cluster](#), using the appropriate Group Key Management command ([KeySetWrite Command](#)).

Note that each operational group key is identified by a unique KeySetID within the fabric. Multiple groups in the same fabric may share the same operational group KeySet. A Group Key Set created through the Groupcast cluster SHALL be visible through the [Group Key Management Cluster](#)'s GroupKeyMap attribute. A Group Key Set created through the [Group Key Management Cluster](#) MAY be used in Groupcast commands to assign the existing KeySetID to a group by GroupID.

11.27.10. Groupcast Auxiliary ACL Handling

The generation rules in this section SHALL be used to manage [AuxiliaryACL](#) entries in the [Access Control Cluster](#) for all LeaveGroup commands, and for JoinGroup and ConfigureAuxiliaryACL commands as determined by their UseAuxiliaryACL field.

For every [MembershipStruct](#) entry in the [Membership](#) attribute that has the HasAuxiliaryACL field

set to true, there SHALL exist AuxiliaryACL entries where all the following are true:

- the [AccessControlEntryStruct](#)'s FabricIndex field matches the [MembershipStruct](#)'s FabricIndex;
- the [AccessControlEntryStruct](#)'s AuxiliaryType field is set to Groupcast;
- the [AccessControlEntryStruct](#)'s Privilege field is set to Operate;
- the [AccessControlEntryStruct](#)'s AuthMode field is set to Group;
- the [AccessControlEntryStruct](#)'s Subjects field contains the [MembershipStruct](#)'s GroupID;
- the [AccessControlEntryStruct](#)'s Targets field contains some of the [MembershipStruct](#)'s Endpoints.

For each [MembershipStruct](#) with its HasAuxiliaryACL field set to true, every endpoint in that struct SHALL have its access granted by an entry in the [AuxiliaryACL](#) attribute. This requirement is satisfied if every logical (FabricIndex, GroupID, EndpointID) combination derived from these membership entries is covered by at least one AuxiliaryACL entry that links the GroupID (in its Subjects field) to the EndpointID (in its Targets field).

Since the [AccessControlEntryStruct](#) has limits in the number of Targets listed (see [TargetsPerAccessControlEntry](#)) per entry, the number of AuxiliaryACL entries needed to express the access granted to members of a group MAY be more than one. The policy SHALL be considered properly applied if, when looking at every AuxiliaryACL entry matching the FabricIndex associated with the [MembershipStruct](#), the union of all such entries, when expanded to equivalent entries having a single subject and a single target, is the same as if one AuxiliaryACL entry was created for each individual Endpoint referenced in the ref_GroupcastMembershipStruct.

For example, using FabricIndex 1 and the following [Membership](#) attribute contents scoped to that fabric:

```
GroupCastMembership := [
  {FabricIndex: 1, GroupID: 0x12AB, Endpoints: [1,2,5,6,7], KeySetID: 1,
  HasAuxiliaryACL: true},
  {FabricIndex: 1, GroupID: 0x34CD, Endpoints: [2,8], KeySetID: 2, HasAuxiliaryACL:
  true},
  {FabricIndex: 1, GroupID: 0x55AA, Endpoints: [2], KeySetID: 1, HasAuxiliaryACL:
  false}
]
```

The following [AuxiliaryACL](#) attribute contents would express most directly the application of the above rules:

```
AuxiliaryACL := [
  # Entries for GroupID 0x12AB, due to HasAuxiliaryACL set to true
  { FabricIndex: 1, Privilege: Operate, AuthMode: Group,
    Subjects: [0x12AB], Targets: [{Endpoint: 1}], AuxiliaryType: Groupcast },
  { FabricIndex: 1, Privilege: Operate, AuthMode: Group,
    Subjects: [0x12AB], Targets: [{Endpoint: 2}], AuxiliaryType: Groupcast },
```

```

{ FabricIndex: 1, Privilege: Operate, AuthMode: Group,
  Subjects: [0x12AB], Targets: [{Endpoint: 5}], AuxiliaryType: Groupcast },
{ FabricIndex: 1, Privilege: Operate, AuthMode: Group,
  Subjects: [0x12AB], Targets: [{Endpoint: 6}], AuxiliaryType: Groupcast },
{ FabricIndex: 1, Privilege: Operate, AuthMode: Group,
  Subjects: [0x12AB], Targets: [{Endpoint: 7}], AuxiliaryType: Groupcast },

# Entries for GroupID 0x34CD, due to HasAuxiliaryACL set to true
{ FabricIndex: 1, Privilege: Operate, AuthMode: Group,
  Subjects: [0x34CD], Targets: [{Endpoint: 2}], AuxiliaryType: Groupcast },
{ FabricIndex: 1, Privilege: Operate, AuthMode: Group,
  Subjects: [0x34CD], Targets: [{Endpoint: 8}], AuxiliaryType: Groupcast },

# NO entries for GroupID 0x55AA, due to HasAuxiliaryACL set to false
]

```

Similarly, another legal encoding of the AuxiliaryACL attribute, where commonality of Subjects and Targets are taken in consideration, and where a [TargetsPerAccessControlEntry](#) value of 3 is taken in consideration, could be:

```

AuxiliaryACL := [
  # Entries for Endpoint 2 shared between two groups:
  { FabricIndex: 1, Privilege: Operate, AuthMode: Group,
    Subjects: [0x12AB, 0x34CD], Targets: [{Endpoint: 2}], AuxiliaryType: Groupcast },

  # Remaining entries for GroupID 0x12AB, due to HasAuxiliaryACL set to true, using
  # as many targets as is possible to reduce number of entries:
  { FabricIndex: 1, Privilege: Operate, AuthMode: Group,
    Subjects: [0x12AB], Targets: [{Endpoint: 1}, {Endpoint: 5}, {Endpoint: 6}],
    AuxiliaryType: Groupcast },
  { FabricIndex: 1, Privilege: Operate, AuthMode: Group,
    Subjects: [0x12AB], Targets: [{Endpoint: 7}], AuxiliaryType: Groupcast },

  # Entry for Endpoint 8 with GroupID 0x34CD, due to HasAuxiliaryACL set to true:
  { FabricIndex: 1, Privilege: Operate, AuthMode: Group,
    Subjects: [0x34CD], Targets: [{Endpoint: 8}], AuxiliaryType: Groupcast },

  # NO entries for GroupID 0x55AA, due to HasAuxiliaryACL set to false
]

```

As can be seen above, there might be multiple different encodings of the [AuxiliaryACL](#) attribute which grant the Operate privilege to the necessary (FabricIndex, GroupID, EndpointID) tuples.

When determining the privilege granted to group senders, it is important to expand the list into access control tuples of (FabricIndex, Single Subject, Single Target, Privilege) to determine equivalence.

Data-dependent commands like the Groups cluster's RemoveAllGroups command cause an impact on data in the node's [Groupcast Session Context](#). Therefore, implementations SHALL ensure that removal of groups through any means (including outside the Groupcast cluster's commands and attributes) cause the requisite updates of AuxiliaryAclEntries if the entry would have disappeared from the Membership attribute.

To summarize, when the HasAuxiliaryACL field is set in a [MembershipStruct](#), all Senders on that Group will have at least Operate privilege on all Endpoints listed for that membership, automatically.

WARNING

Whenever the [SoftwareVersion](#) or [ConfigurationVersion](#) attributes are modified in the [Basic Information cluster](#) of a Node's Root Node endpoint, Administrators SHOULD evaluate whether the current UseAuxiliaryACL settings remain appropriate for the fabric's access policy. Because the AuxiliaryACL policy grants broad "operate" privilege at the endpoint level, the introduction of new clusters in a newer version MAY inadvertently extend access permissions beyond the original intent. This evaluation SHOULD be performed during session reestablishment to a Node or via monitoring of attribute reports from a regular subscription.

NOTE

Use of the HasAuxiliaryACL field is an optional administrative choice designed to simplify and scale access control for common group scenarios. An Administrator MAY always bypass this automated feature and instead configure granular permissions manually using the [ACL](#) attribute. Any AuxiliaryACL entries auto-generated by this feature MAY be revoked at any time by invoking the [ConfigureAuxiliaryACL](#) command with the UseAuxiliaryACL field set to false.

Chapter 12. Multiple Fabrics

12.1. Introduction

The Multiple Fabric feature allows a Node to be commissioned to multiple separately-administered Fabrics. With this feature a current Administrator can (with user consent) allow the Commissioner for another fabric to commission that Node within its Fabric. The new Commissioner SHALL have their own [Node Operational Certificate \(NOC\)](#) issued by its Trusted Root Certificate Authority (TRCA). Once commissioning is completed and the Node is properly configured, Administrators on the newly joined Fabric have access to the Node and can perform all administrative tasks.

A Fabric is anchored by its Trusted Root Certificate Authority (TRCA). A TRCA MAY delegate to one or more Intermediate Certificate Authorities (ICA) which issue [NOCs](#). Multiple vendors or companies can use the same CA hierarchy in which case they will be governed under the same Trusted Root Certificate Authority.

12.2. Joint Fabric

12.2.1. Introduction

When multiple vendors or companies use the same CA hierarchy, governed under the same Trusted Root Certificate Authority, they create a Joint Fabric using the [Joint Commissioning Method](#) flow. This flow enables a set of one or more companies to use a single fabric anchored by the Trusted Root Certificate Authority (TRCA). In the context of [JCM](#), the fabric that is anchored by this TRCA is known as the Anchor Fabric.

The Joint Fabric architecture allows any Ecosystem Administrator to administer devices on the Joint Fabric, where the Joint Fabric Administrators [Node Operational Certificate \(NOC\)](#) and devices being administered have ACL entries containing the [Administrator CAT](#), regardless of what Ecosystem Commissioner/Administrator commissions a device onto the Joint Fabric. As a result, the other Ecosystem Administrator(s) will be notified, via [Joint Fabric Datastore Cluster](#) Joint Fabric.

Similarly, any Ecosystem Administrator participating in the Joint Fabric can remove a device from the Joint Fabric and all other Ecosystem Administrator will be notified, via Joint Fabric [NodeList Attribute](#) subscriptions, of its removal and will no longer be able to communicate with the device. Impact of a device's removal can be mitigated by the fact that any bindings or subscriptions to the device from other devices in the Joint Fabric are known within the context of the [Joint Fabric Datastore Cluster](#) and can be in turn removed as part of the removal process or the removal process can be aborted.

12.2.2. Node ID Generation

Any newly-allocated Node ID SHALL:

- be greater than 0x0000_0000_0000_0000, but less than 0xFFFF_FFEF_FFFF_FFFF, representing a value within the Operational NodeID range (see [Table 4, “Node Identifier Allocations”](#));
- be checked to ensure its uniqueness in the [NodeList](#) attribute of the [Section 11.24, “Joint Fabric](#)

Datastore Cluster”

The Node ID SHALL be regenerated if these constraints are not met.

It is RECOMMENDED to use random allocation within the valid range to avoid having to regenerate the Node ID.

12.2.3. Anchor ICAC requirements

The Anchor ICAC SHALL be the ICAC corresponding to the Anchor Administrator. The Anchor ICAC SHALL contain the reserved **org-unit-name** attribute from the Table 87, “Standard DN Object Identifiers” with value **jf-anchor-icac** in its Subject DN. The Anchor ICAC SHALL be issued only by the Anchor CA to an Anchor Administrator.

12.2.4. Joint Fabric ACL Architecture

A set of **CASE Authenticated Tags** defined to be used in the context of Joint Fabric. In combination with the **Anchor ICAC**, these CATs can provide a greater level of security.

12.2.4.1. Administrator CAT

The Administrator CAT is meant to simplify management of the group of Administrator Nodes as opposed to managing a list of subject Node IDs. All devices participating in Section 12.2, “Joint Fabric” SHALL contain an ACL entry granting **Administer** privilege to **CaseSubjectAdmin** set to the Administrator CAT. During commissioning of any Node onto the Section 12.2, “Joint Fabric” the **CaseAdmin-Subject** field SHALL be set to the Administrator CAT upon invoking the **AddNOC** command.

Any Node advertising as a Section 12.2, “Joint Fabric” Administrator SHALL contain the Administrator CAT in its NOC. A NOC containing the Administrator CAT MAY be issued by any Section 12.2, “Joint Fabric” Administrator.

Any client that discovers an Administrator Node with DNS-SD and connects to the Node via CASE SHALL check if the Administrator CAT is present in the NOC using the **Peer CASE Authenticated Tags** before taking any action on the Node. This SHALL be required in order to verify that the Node is authorized to act as an Administrator in the Section 12.2, “Joint Fabric”.

If the Administrator CAT is already in use in a fabric that wants to participate in the Section 12.2, “Joint Fabric” then its Administrator first needs to take the required steps for updating it with a non-overlapping Administrator CAT. This may involve updating ACLs and **NOCs** on the Nodes wishing to participate in the Section 12.2, “Joint Fabric”.

Concern for its use in gaining unauthorized **Administer** access is mitigated by the fact that the Administrator CAT is a **special subject distinguished name** within the **NOC** and that the **NOC** SHALL be issued only by an Administrator.

User initiated and granted revocation of an Administrator to administer nodes SHALL be achieved by updating the Administrator CAT. The Joint Fabric Anchor Administrator SHALL increment the version number of the Administrator CAT to a value higher than its current value (e.g., from 0x0000 to 0x0001), update the existing credentials (**NOC**) for all Administrator Nodes that are NOT being revoked with the new version of the Administrator CAT, and update the ACL entry of all Nodes

whose subject list contains the prior version of the Administrator CAT with the new version of the Administrator CAT. This has the effect of revoking the **Administer** access of any Administrator Nodes that did not receive updated credentials (**NOC**) with the new version of the Administrator CAT. Completing this operation requires visiting all the nodes in the [Section 12.2, “Joint Fabric”](#), a task which might take a long time to complete or might never complete if some Nodes are permanently offline or otherwise unreachable. However, any Nodes that are permanently offline are probably not at risk due to the no longer trusted Administrator Node because they are inaccessible to it.

12.2.4.2. Anchor CAT

The Anchor CAT is meant to restrict **Administer** privilege to [Section 12.2, “Joint Fabric”](#) Administrator nodes, thereby granting **Administer** access to only the [Section 12.2, “Joint Fabric”](#) Anchor Node with a **NOC** containing the Anchor CAT. All Administrator devices participating in [Section 12.2, “Joint Fabric”](#) SHALL contain an ACL entry granting **Administer** privilege with the **CaseSubjectAdmin** set to the Anchor CAT. During commissioning of any Administrator Node onto the [Section 12.2, “Joint Fabric”](#) the **CaseAdminSubject** field SHALL be set to the Anchor CAT upon invoking the **AddNOC** command. Any Node advertising as a [Section 12.2, “Joint Fabric”](#) Anchor or Datastore SHALL contain the Anchor CAT in its **NOC**. A **NOC** containing the Anchor CAT SHALL be issued only by the [Section 12.2, “Joint Fabric”](#) Anchor ICAC.

Any client that discovers an Anchor Node with DNS-SD and connects to the Node via CASE SHALL check if the Anchor CAT is present in the **NOC** using the **Peer CASE Authenticated Tags** and that the **NOC** chains up to the [Anchor ICAC](#) before taking any [Section 12.2, “Joint Fabric”](#) actions on the Node. This SHALL be required in order to verify that the Node is authorized to act as an Anchor of the [Section 12.2, “Joint Fabric”](#).

If the Anchor CAT is already in use in a fabric that wants to participate in the [Section 12.2, “Joint Fabric”](#) then its Administrator first needs to take the required steps for updating it with a non-overlapping Anchor CAT. This may involve updating ACLs and **NOCs** on the Nodes wishing to participate in the [Section 12.2, “Joint Fabric”](#).

12.2.5. Joint Commissioning Method (JCM)

This method SHALL be implemented for Commissioners and Administrators that support Joint Fabric.

Given a pre-existing Matter fabric, Fabric A, managed by Ecosystem A with its own administrator and commissioner nodes, a Joint Fabric can be created with another Ecosystem B. This allows all devices of Ecosystem A and Ecosystem B to communicate using a single operational root of trust. Either ecosystem can add devices to the new joint fabric. Joint Commissioning Management (JCM) is the process through which an administrator of Ecosystem B receives a **cross-signed ICAC** from the Anchor Administrator of Ecosystem A. While Ecosystem A MAY contain multiple administrator nodes, only the Anchor Administrator SHALL be able to execute JCM.

The joint operation ensures that ICA's from both Fabric A and Fabric B are signed by the Anchor CA (the Root CA of the fabric selected to become the Anchor Fabric), establishing a common trust between all devices of the original Fabric A and Fabric B and all newly added devices to the “Joint Fabric”.

The Joint Fabric contains multiple ecosystems whose ICACs are cross-signed by the Anchor CA. Therefore, all ecosystems participating in the [Section 12.2, “Joint Fabric”](#) SHALL use an ICAC to sign NOCs. Signing NOCs with a Root CA different than the Anchor CA is not permitted in the Joint Fabric.

In the figure below, Ecosystem A with Fabric A acts as the Anchor Fabric and Ecosystem B with Fabric B as joining fabric to create a Joint Fabric between the two Ecosystems. Here are the JCM steps:

1. Ecosystem B Administrator SHALL become commissionable for JCM. This SHALL happen using one of the two methods below:
 - a. Another Node which is an Administrator in Ecosystem B instructs the Ecosystem B Administrator, over a [CASE](#) session inside Fabric B, to go into [Section 11.25.6.5, “OpenJointCommissioningWindow”](#)
 - b. Ecosystem B Administrator puts itself in commissionable mode using implementation specific means for calling [Section 11.25.6.5, “OpenJointCommissioningWindow”](#)
2. A new **RANDOM** passcode SHALL be computed and [Section 11.25.6.5, “OpenJointCommissioningWindow”](#) SHALL be called using the corresponding [PAKE](#) passcode verifier. The new passcode and the discriminator SHALL be presented to Ecosystem A Commissioner using the same rules described by ECM in [Section 5.6.3.1, “Presentation of Onboarding Payload for ECM”](#).
3. On Ecosystem A side, User SHALL be made aware that it has the option of using the [Section 5.1, “Onboarding Payload”](#) in order to start the JCM steps. Once this option is selected, Ecosystem A Commissioner SHALL start execution of the steps outlined in [Figure 47, “Commissioning flow diagram”](#) with Ecosystem B Administrator.
 - a. Commissioning steps proceed through the successful end of Device Attestation Procedure (that would be initial phase step [10](#)), with standard error handling behavior unchanged.
4. A new step is introduced for verifying trust against Ecosystem B Administrator:
 - a. Ecosystem A Commissioner SHALL search all endpoints' Descriptor cluster of Ecosystem B Administrator for the Joint Fabric Administrator device type, which requires the [Joint Fabric Administrator](#) cluster.
 - i. if **Joint Fabric Administrator** cluster is not found, process SHALL be terminated and the User SHALL be informed that Ecosystem B lacks the expected functionality
 - ii. otherwise the endpoint that contains the [Joint Fabric Administrator](#) cluster SHALL be saved as **JointEndPointB**
 - b. Ecosystem A Commissioner SHALL check that the NOC used by the Fabric indicated by the **AdministratorFabricIndex** of the **Joint Fabric Administrator** Cluster on **JointEndPointB** contains an Administrator CAT.
 - i. if verification fails, process SHALL be terminated and the User SHALL be informed that Ecosystem B Administrator is not trusted
 - c. Ecosystem A Commissioner SHALL execute [Section 6.4.10, “Fabric Table Vendor ID Verification Procedure”](#) against the Fabric indicated by the **AdministratorFabricIndex** of the **Joint Fabric Administrator** Cluster on **JointEndPointB** (i.e. [Section 7.13.6, “FabricIndex”](#) corresponding to Fabric B)
 - i. if verification fails, process SHALL be terminated and the User SHALL be informed that

Ecosystem B Administrator is not trusted

- d. On Ecosystem A side, User SHALL be asked if they agree to onboard the Fabric indicated by **AdministratorFabricIndex**. At least the name associated with the **VendorID** of the onboarded Fabric SHALL be presented to the User.
 - e. Ecosystem A Commissioner SHALL save the **SubjectPublicKey** of the ICAC belonging to the Fabric indicated by **AdministratorFabricIndex** as **TrustedIcacPublicKeyB**
5. **Figure 47, “Commissioning flow diagram”** SHALL continue with the original steps between initial phase step 11 and setup phase step 3:
- a. during the **AddNOC** step, Ecosystem B Administrator SHALL receive a NOC with Administrator CAT inside Fabric A. The version of the Administrator CAT SHALL be the one indicated by the corresponding entry of the **GroupList** attribute hold by the Joint Fabric Datastore of Fabric B.
 - i. If the NOC doesn’t contain an Administrator CAT then this command SHALL process an error by responding with a **Section 11.18.6.10, “NOCResponse”** with a **StatusCode** of **InvalidNOC** as described in **Section 11.18.6.7.2, “Handling Errors”**. The process SHALL be terminated and the User SHALL be informed.
 - b. during the **AddNOC** step, the **CaseAdminSubject** field SHALL contain an Anchor CAT that grants Administer privileges only to Ecosystem A Anchor Administrator over Ecosystem B Administrator in Fabric A. The version of the Anchor CAT SHALL be the one indicated by the corresponding entry of the **GroupList** attribute hold by the Joint Fabric Datastore of Fabric B.
 - i. If the **CaseAdminSubject** doesn’t contain an Anchor CAT then this command SHALL process an error by responding with a **StatusCode** of **InvalidAdminSubject** as described in **Section 11.18.6.7.2, “Handling Errors”**. The process SHALL be terminated and the User SHALL be informed.
6. A new step is introduced for verifying trust against Ecosystem A Administrator:
- a. Ecosystem A Administrator SHALL invoke the **Section 11.25.6.9, “AnnounceJointFabricAdministrator”** command of the **Joint Fabric Administrator** cluster belonging to **JointEndPointB** on Ecosystem B Administrator. The **EndpointID** parameter SHALL be set to the value of the endpoint that holds the **Joint Fabric Administrator** on Ecosystem B Administrator.
 - i. Ecosystem B Administrator SHALL save the value of the **EndpointID** as **JointEndPointA**
 - b. Ecosystem B Administrator SHALL read the **Section 11.25.5.1, “AdministratorFabricIndex”** attribute of the **Joint Fabric Administrator** cluster belonging to **JointEndPointA** on Ecosystem A Administrator and executes **Section 6.4.10, “Fabric Table Vendor ID Verification Procedure”** against the Fabric indicated by **AdministratorFabricIndex** (i.e. **Section 7.13.6, “FabricIndex”** corresponding to Fabric A). Ecosystem B Administrator SHALL check that the **RootPublicKey** and **FabricID** of the **accessing fabric** (found in the **FabricDescriptorStruct**) match the **RootPublicKey** and **FabricID** of the Fabric indicated by **AdministratorFabricIndex**.
 - i. if verification fails, the process SHALL be terminated and the User SHALL be informed that Ecosystem B Administrator is not trusted
7. Ecosystem A Administrator SHALL follow the **ICAC cross-signing** steps using **TrustedIcacPublicKeyB** as input parameter.
- a. in case of error, the process SHALL be terminated and the User SHALL be informed

8. Ecosystem B Administrator SHALL save the cross-signed ICAC (see [Section 11.25.6.3.1, “ICAC-Value”](#) parameter of the [Section 11.25.6.3, “AddICAC”](#) command) as **JFCrossSignedICAC**
9. Ecosystem B Administrator SHALL set the [Section 11.25.5.1, “AdministratorFabricIndex”](#) attribute of its own Joint Fabric Administrator cluster to the index that Fabric A has inside the **FabricDescriptorStruct**
10. Setup phase step 4 (**CommissioningComplete**) of [Section 5.5, “Commissioning Flows”](#) SHALL be executed.
11. User MAY be notified that Ecosystem B Administrator is now onboarded as an Administrator in the Joint Fabric formed with Ecosystem A. At least the VendorID of Ecosystem A SHOULD be provided in this notification and the User MAY have the option of leaving the Joint Fabric (e.g.: User doesn't want to join Ecosystem B with A's Vendor ID).

Ecosystem B Administrator MAY now use the below commands from [Operational Credentials cluster](#) to update each Node in the Fabric B to join Fabric A:

1. Update [Section 11.18.5.5, “TrustedRootCertificates”](#) by using [Section 11.18.6.13, “AddTrusted-RootCertificate”](#) with the **RootCACertificate** parameter set to the RCAC of Fabric A
2. Update Node Operational Credentials using **AddNOC**:
 - a. **NOCValue** SHALL contain a NOC issued by the **JFCrossSignedICAC**. Subject DN of the **NOCValue** SHALL encode a **matter-fabric-id** attribute whose value SHALL be identical with the value of the **matter-fabric-id** attribute from **JFCrossSignedICAC**
 - b. **ICACValue** parameter SHALL be set to the value of **JFCrossSignedICAC**
 - c. **IPKValue** parameter SHALL be set to the value of the IPK found in the **GroupKeySetList** hold by the Joint Fabric Datastore of Fabric A
 - d. **CaseAdminSubject** SHALL contain an Administrator CAT whose version is indicated by the corresponding entry from **GroupList** attribute hold by the Joint Fabric Datastore of Fabric B.
 - e. **AdminVendorId** field SHALL be set to the value of the AdminVendorIdValue found in the local Fabric Table under the **AdministratorFabricIndex**.
3. Remove the old Trusted Root CA Certificate and the associated Fabric B using the [Section 11.18.6.12, “RemoveFabric”](#) command

User Consent MAY be needed each time a new Node/subset of Nodes is shared from Fabric B to Fabric A. While moving Nodes from Fabric B to Fabric A, Ecosystem B SHALL invoke the required commands on the Joint Fabric Datastore cluster owned by the Anchor Administrator of Fabric B in order to remove those Nodes and on the Joint Fabric Datastore cluster owned by Ecosystem A Anchor Administrator in order to add those Nodes.




Figure 111. Joint Commissioning Steps

12.2.5.1. Scope of User Consent

Before commissioning a Joint Fabric Administrator, the user SHALL be asked for consent to enable Joint Fabric functionality between ecosystems. Each ecosystem joining the Joint Fabric SHALL independently ask the user for consent before initiating JCM.

12.2.5.2. Discovery

The user SHALL be able to enable **JCM** through an appropriate interface of the devices on that ecosystem. For example, a mobile application, a web configuration, or an on-device interface.

12.2.5.3. Vendor ID Validation

The Joint Fabric requires that the fabrics participating in a Joint Fabric perform mutual Vendor ID validation so that they can be sure of the authenticity of the Vendor ID being asserted. Vendor ID validation SHALL be achieved by using the [Section 6.4.10, “Fabric Table Vendor ID Verification Procedure”](#).

12.2.5.4. ICAC Cross Signing

A joining device, acting also as an Administrator in Fabric B, SHALL obtain the ability to issue [NOCs](#) chaining up to the Anchor CA of Fabric A once it has become an Administrator in Fabric A by following the [Joint Commissioning Method](#). For simplicity, the joining device will be called joining Administrator.

To obtain this ability it SHALL receive an [ICAC](#) issued by the Anchor Administrator from Fabric A. Cross-signing means that the SubjectPublicKey of the [ICAC](#) chaining up to the Anchor CA of Fabric A is identical to the SubjectPublicKey of the [ICAC](#) associated with the Fabric indicated by [Section 11.25.5.1, “AdministratorFabricIndex”](#) of the joining Administrator (i.e. Fabric B).

A pre-requisite for the ICA Cross Signing process is the execution of [FabricFabric Table Vendor ID Verification Procedure](#) against the Fabric indicated by [Section 11.25.5.1, “AdministratorFabricIndex”](#) of the joining Administrator. Also the public ICAC of that Fabric SHALL be passed as input parameter to this procedure as [trusted SubjectPublicKey](#) and it SHALL be encoded using the specific rules for [ec-pub-key](#), [pub-key-algo](#) and [ec-curve-id](#).

The ICAC Cross Signing process is started by the Anchor Administrator from Fabric A by sending an [Section 11.25.6.1, “ICACCSRRequest”](#) command to the joining Administrator. The joining Administrator SHALL create a Certificate Signing Request (ICA CSR) as described in [PKCS #10](#):

1. ICAC CSR SHALL include a signature (see [RFC 2986](#) section 4.2, SignatureAlgorithm) generated using the Private Key of its pre-installed [ICAC](#).
2. The Public Key associated with the Private Key used to sign the ICA CSR SHALL appear in the SubjectPublicKey of the ICA CSR.
3. Minimum attributes that SHALL be present are shown in the image below.
4. Once the ICAC CSR is created it SHALL be sent in a DER-encoded string inside the [ICACCSR](#) parameter of the [Section 11.25.6.2, “ICACCSRResponse”](#) command.

ICACSR

Certificate Request:

Data:

Version: 1 (0x0)

Subject:

Subject Public Key Info:

Public Key Algorithm: id-ecPublicKey

Public-Key: (256 bit)

pub:

04:12:3b:90:f5:.....

ASN1 OID: prime256v1

NIST CURVE: P-256

Attributes:

Requested Extensions:

Signature Algorithm: ecdsa-with-SHA256

30:46:02:21:00:95:ff:.....

The ICA cross signing requires the Anchor Administrator to receive and validate the Certificate Signing Request. The validation steps are as follows:

1. ICAC CSR follows the encoding and rules from [PKCS #10](#).
2. Signature SHALL be validated using the [trusted SubjectPublicKey](#).
3. Convert the Subject Public Key Info of the ICA CSR to the equivalent Matter certificate TLV values by using the specific rules for [ec-pub-key](#), [pub-key-algo](#) and [ec-curve-id](#).
 - a. Values for [ec-pub-key](#), [pub-key-algo](#), [ec-curve-id](#) SHALL match with the ones of the [trusted SubjectPublicKey](#)
4. Upon success, ICA's certificate SHALL be signed by the root CA of the Anchor Administrator, such that the ICAC chains to the [RCAC](#) of the Anchor Administrator.
 - a. The subject DN SHALL encode a [matter-fabric-id](#) attribute. The attribute's value SHALL be identical to the Fabric ID of the Anchor Fabric.
5. ICAC SHALL be returned inside the [ICAC](#) field of the [Section 11.25.6.3, "AddICAC"](#) command. If any of the above validation checks fail, the server SHALL immediately terminate the procedure and inform the User.

12.2.6. Anchor Administrator Selection

Anchor Administrator Selection is a process where the Joint Fabric Administrator is selected by the user. Given a Joint Fabric Administrator holds additional responsibilities such as [Section 12.2.5.4, "ICAC Cross Signing"](#), hosting the [Joint Fabric Datastore](#) and being the only entity allowed to issue [Joint Fabric Datastore](#) updates, assignment of this role SHALL be based on the intent and consent of the user. This role SHALL be assigned during the creation of the Joint Fabric using the [Section 12.2.5, "Joint Commissioning Method \(JCM\)"](#) or MAY be transferred to another Joint Fabric Administrator using the [Transfer Anchor](#) procedure, both requiring [Section 11.20.3.4, "Obtaining user consent for updating software"](#).

The [Transfer Anchor](#) procedure allows transferring the Anchor Administrator role from one Joint Fabric Administrator to another Joint Fabric Administrator. This process is initiated when the current Anchor Administrator needs to be replaced by another Joint Fabric Administrator and is controlled by the user.

User SHALL provide consent when an Anchor Administrator is removed and a new Anchor CA is chosen. In the example below, the Anchor Administrator is removed from the Joint Fabric and the user has selected a candidate Administrator to be promoted to Anchor Administrator. For clarity, Joint Fabric Anchor Administrator **A** is being replaced with the candidate Joint Fabric Administrator **B**.

User consent SHALL be mutual:

- on Administrator **A** side, User provides consent that transfer of the Anchor role to Administrator **B** is allowed.
- on Administrator **B** side, User provides consent that receipt of the Anchor role from Administra-

tor **A** is allowed.

Flow for Anchor Transfer is as follows:

1. Administrator B SHALL:
 - a. send [Section 11.25.6.6, “TransferAnchorRequest”](#) to Administrator A to set itself as Joint Fabric Anchor Administrator.
 - b. obtain user consent prior to sending the [Section 11.25.6.6, “TransferAnchorRequest”](#) command.
2. Administrator A SHALL:
 - a. check that the NOC used by Administrator B during the CASE session contains an Administrator CAT.
 - b. check that user provided consent that allows the transfer of the Anchor role to a different Administrator. If not, then [Section 11.25.6.7, “TransferAnchorResponse”](#) command with StatusCode set to [TransferAnchorStatusNoUserConsent](#) SHALL be sent and the procedure stopped here.
 - c. check all the [Datastore](#) entries of type [DatastoreStatusEntryStruct](#). If any of these entries has a value that equals to [Pending](#) or to [DeletePending](#) then [Section 11.25.6.7, “TransferAnchorResponse”](#) command with StatusCode set to [TransferAnchorStatusDatastoreBusy](#) SHALL be sent and the procedure stopped here.
 - d. put Joint Fabric Datastore in read only state by setting the Datastore [StatusEntry](#) to [DeletePending](#).
 - e. stop DNS-SD advertising of the [Administrator](#), [Anchor](#) and [Datastore](#) capability inside the [JF TXT](#) key.
3. All other Joint Fabric Administrators SHALL:
 - a. stop commissioning of any new devices into the Joint Fabric once it detects that Datastore [StatusEntry](#) equals [DeletePending](#).
4. Administrator B SHALL:
 - a. copy (through attribute read, BDX) Joint Fabric Datastore from Administrator A.
 - b. set the Datastore [StatusEntry](#) to [Committed](#).
 - c. set the value of the Datastore [AnchorNodeId](#) attribute to the value of its Node ID.
 - d. increase [Section 12.2.4.1, “Administrator CAT”](#) version.
 - e. use following commands from [Operational Credentials cluster](#) on the other Joint Fabric Administrators:
 - i. update [Section 11.18.5.5, “TrustedRootCertificates”](#) by using [Section 11.18.6.13, “AddTrustedRootCertificate”](#).
 - ii. update Node Operational Credentials using [AddNOC](#). [NOCs](#) SHALL contain an updated version of the Administrator CAT.
 - f. issue itself a new NOC with the updated version of the Anchor CAT.
 - g. set itself as the new Datastore provider and Anchor Administrator by advertising the [Admin-](#)

istrator, Anchor and Datastore capabilities inside the JF TXT key.

- h. send Section 11.25.6.8, “TransferAnchorComplete” to Administrator A to announce that transition to Anchor Administrator is complete.
5. All other Joint Fabric Administrators SHALL:
 - a. subscribe to the new Datastore provider (Administrator B) having discovered the new Datstore capability via Service Discovery of the JF TXT key.
 - b. check that the NOC used by the new Datastore provider (Administrator B) during the first CASE session contains an Section 12.2.4.2, “Anchor CAT” and that the NOC chains up to the Anchor ICAC.
 - c. request ICA Cross Signing from the new Joint Fabric Anchor Administrator discovered via Service Discovery (Administrator and Anchor flags are set in JF TXT key).
 - d. remove devices from the fabric governed by the old Anchor CA using the Section 11.18.6.12, “RemoveFabric” command.
 - e. start commissioning devices onto the Joint Fabric using the new ICAC for NOC issuance.

12.2.7. Administrator Removal

Since Section 12.2, “Joint Fabric” has multiple Intermediate NOC Certificate Authorities trusted by Joint Fabric Anchor CA, the following steps SHALL be taken to remove an Intermediate NOC Certificate Authority from a Section 12.2, “Joint Fabric”.

The Section 11.18.6.12, “RemoveFabric” section outlines a Warning that SHALL apply here for removing a Joint Fabric Administrator.

1. Joint Fabric Anchor Administrator SHALL send a Section 11.18.6.12, “RemoveFabric” command to a Joint Fabric Admin that user has consented to remove.
2. Joint Fabric Administrator SHALL remove the outgoing Administrator from the Joint Fabric Datastore.
3. The outgoing Administrator SHALL remove the given Fabric and **delete all associated fabric-scoped data**.
4. Joint Fabric Anchor Administrator SHALL increase Section 12.2.4.1, “Administrator CAT” version and issue itself a new NOC.
5. Administrator B SHALL set new NOC for all the Joint Fabric Administrators.

12.2.7.1. Security Consideration

Matter does not currently include any method for a Trusted Root Certificate to revoke an ICAC previously issued. Thus, to ensure proper fail proof removal of a Joint Fabric Administrator from a Joint Fabric, the Anchor Administrator SHOULD trigger a transition to a new Trusted Root Certificate as described in the Section 12.2.6, “Anchor Administrator Selection” section. In this case, the new Anchor can be run by the same ecosystem as the old Anchor but the new Trusted Root Certificate will not issue an ICAC to the Joint Fabric Administrator that is to be removed.

12.3. User Consent

A user who wishes to have an already commissioned Node join another Fabric (and therefore another Security Domain) provides consent by instructing an existing Administrator, which SHALL put the Node into commissioning mode by using steps outlined in [Section 5.6.4, “Open Commissioning Window”](#). Administrators SHALL provide a mechanism for the user to thus instruct them.

12.4. Administrator-Assisted Commissioning Method

Administrators SHALL support opening a commissioning window on a Node using the mandatory method described in [Section 5.6.3, “Enhanced Commissioning Method \(ECM\)”](#). All Nodes SHALL support having a commissioning window opened using the mandatory method described in [Section 5.6.3, “Enhanced Commissioning Method \(ECM\)”](#).

An Administrator MAY open a commissioning window on a Node using the optional method described in [Section 5.6.2, “Basic Commissioning Method \(BCM\)”](#), if the Node supports the method.

12.5. Node Behavior

The Node SHALL host an [Section 11.19, “Administrator Commissioning Cluster”](#). The Cluster exposes commands which enable the entry into commissioning mode for a prescribed time, and which SHALL be invoked over a [CASE](#) secure channel. See [Section 11.19.8.1, “OpenCommissioningWindow”](#) and [Section 11.19.8.2, “OpenBasicCommissioningWindow”](#). During such a commissioning window, the Node SHALL maintain its existing configuration, such as its operational network connection and identities, and SHOULD allow normal interactions from other Nodes.

12.6. Fabric Synchronization

12.6.1. Introduction

The Fabric Synchronization feature enables commissioning of devices from one fabric to another without requiring user intervention for every device. It defines mechanisms that can be used by multiple ecosystems/controllers to communicate with one another to simplify the experience for users.

The following diagram shows how two ecosystems/controllers would communicate after a single setup sequence.




Figure 112. Example of two synchronized fabrics

12.6.2. Terminology

A component within an ecosystem which supports Fabric Synchronization will be referred to as a **Fabric Synchronizing Administrator**. This includes the Aggregator and Fabric Administrator node(s).

A device which is being synchronized between ecosystems will be referred to as a **Synchronized Device**.

12.6.3. Fabric Synchronization Composition

Fabric Synchronization is present when an Aggregator satisfies the FabricSynchronization condition or one or more Bridged Nodes satisfy the FabricSynchronizedNode condition. (See Device Library, *Aggregator* and *Bridged Node*.)




Figure 113. Composition of a device supporting Fabric Synchronization.

An Aggregator supporting Fabric Synchronization SHALL be composed of the following components.

12.6.3.1. Matter Fabric Administrator Node

The device providing the Aggregator SHALL be able to commission nodes on its fabric.

12.6.3.2. Aggregator Node

When Fabric Synchronization is supported, the Aggregator with FabricSynchronization condition (see Device Library, *Aggregator*) SHALL be met on an endpoint with the following endpoints in the Descriptor cluster PartsList.

Commissioner Control Cluster

The Commissioner Control Cluster enables another device which supports Fabric Synchronization to set up a bidirectional synchronization relationship without the user having to scan a QR code or enter a Manual Pairing Code (as used in [ECM](#) flow) for a device.

Fabric Synchronization SHALL be supported when the SupportedDeviceCategories attribute in the Commissioner Control Cluster has the FabricSynchronization bit set.

Bridged Node Endpoint(s)

Every Bridged Node endpoint presented that conforms to the FabricSynchronizedNode condition (see Device Library, *Bridged Node*) represents a device that can be synchronized (A Synchronized Device).

The Bridged Node with FabricSynchronizedNode condition enables the synchronization of devices between fabrics by the following:

- The Bridged Node SHALL include the Ecosystem Information Cluster. (Enables discovery and directory synchronization.)
- The Bridged Node SHALL include the Administrator Commissioning Cluster when the user consents to share a device. (Enables synchronization of devices between fabrics.)
- The Bridged Node SHOULD support the BridgedICDSupport feature in the Bridged Device Basic Information Cluster if the Synchronized Device is an Intermittently Connected Device (ICD). (Enables communication with ICDs.)

12.6.4. Preventing Device Duplication

A Fabric Synchronizing Administrator SHOULD NOT commission devices onto the same fabric that they are already on. To avoid this, the Fabric Synchronizing Administrator SHOULD examine the UniqueID of a potential Synchronized Device's Bridged Device Basic Information Cluster and SHOULD examine the VendorID and ProductID fields if they are present in the Bridged Device Basic Information Cluster. If all of the provided values for UniqueID, ProductID, and VendorID match a known device that is already on the Fabric Synchronizing Administrator's fabric, then it SHOULD NOT attempt to commission the device.

12.6.4.1. Missing UniqueID

The UniqueID field in the Basic Information Cluster became mandatory after the initial release of this specification. Therefore, the following must be performed to support devices without

UniqueIDs.

When a Fabric Synchronizing Administrator commissions a Synchronized Device, it SHALL persist and maintain an association with the UniqueID in the Bridged Device Basic Information Cluster exposed by another Fabric Synchronizing Administrator.

If a Fabric Synchronizing Administrator exposes a Synchronized Device which does not have a UniqueID in its Basic Information Cluster, then the Fabric Synchronizing Administrator SHALL generate and persist a new UniqueID to be used in the Bridged Device Basic Information Cluster.

12.6.4.2. Unifying Generated UniqueID

When a Fabric Synchronizing Administrator establishes a PASE session to a Synchronized Device for the purposes of commissioning, the Fabric Synchronizing Administrator SHOULD verify that the device is not already present on the intended fabric as follows:

- The Fabric Synchronizing Administrator MAY check if the UniqueID is present in the Basic Information Cluster. If the UniqueID is present, the Fabric Synchronizing Administrator can skip the below check.
- If the UniqueID is not present or not checked, the Fabric Synchronizing Administrator SHOULD check if the intended fabric is already present in the Fabric Table.
 - If it is present, the Fabric Synchronizing Administrator SHOULD NOT complete commissioning and SHOULD avoid attempting to commission the device (or establish PASE sessions) in the future by persisting the UniqueID exposed by the other Fabric Synchronizing Administrator's Bridged Device Basic Information Cluster. (If the Fabric Synchronizing Administrator exposes the device through a Bridged Node endpoint, then the Fabric Synchronizing Administrator SHOULD expose the UniqueID through its Bridged Device Basic Information Cluster.)
 - If it is not present, the Fabric Synchronizing Administrator MAY continue the commissioning process.

To avoid the potential for getting stuck in a loop of checking the device and updating UniqueIDs in the Bridged Device Basic Information Cluster, a caching and tie-breaker policy is required. If a Fabric Synchronizing Administrator updated its Bridged Device Basic Information Cluster as a result of a duplication detected when checking the Fabric Table during a PASE session, then it SHOULD perform the following:

- The Fabric Synchronizing Administrator SHOULD create a cache of prior known UniqueIDs scoped to the NodeID of the Synchronized Device. The cache SHOULD have space for at least 5 entries per NodeID.
- If a Fabric Synchronizing Administrator (hereafter denoted A) receives a UniqueID from another Fabric Synchronizing Administrator's (hereafter denoted B) Bridged Device Basic Information Cluster that matches an entry in the cache, but is not the entry currently presented in the Bridged Device Basic Information Cluster of the client Fabric Synchronizing Administrator (A), then the Fabric Synchronizing Administrator (A) SHOULD set the UniqueID in its Bridged Device Basic Information Cluster to the value stored in the cache which is lexicographically smaller than all other entries.

12.6.5. Changes to device and locations of Synchronized Devices

A fabric administrator typically has information on topology or logical grouping of the Synchronized Devices, which can be of use to the user when shared with other administrators.

For example, consider a fabric with 50 lights. If the locations (such as rooms) and device naming are not present, the user would just see a list of 50 lights on their synchronized controller and would not know which of those lights would be in which location.

Typically, the user has some means (e.g. a Manufacturer-provided app) to assign names to the Synchronized Devices, or names could be assigned automatically by the administrator. Sometimes these names are written to the Basic Information Cluster on the device, while at other times the administrator does not write these names to the Basic Information Cluster.

A Fabric Synchronizing Administrator SHOULD expose such names in the Ecosystem Information Cluster on the associated endpoint. A Fabric Synchronizing Administrator MAY expose such names in the Basic Information Cluster on the associated endpoint.

If a Fabric Synchronizing Administrator exposes such names in the Basic Information Cluster for a Synchronized Device, then the same associated names SHALL be exposed in the Ecosystem Information Cluster.

Similarly, the user typically has some means to assign and change the location (e.g. room/zone) names. The grouping MAY also be applied automatically by the administrator.

A Fabric Synchronizing Administrator SHOULD expose such grouping using the Ecosystem Information Cluster as described above.

Nodes that wish to be notified of a change in such a name or location SHOULD monitor changes of the Ecosystem Information Cluster.

A Fabric Synchronizing Administrator MAY make it possible (e.g. through a Manufacturer's app) for its users to restrict whether all or some of the Ecosystem Information Cluster is exposed to the Fabric.

12.6.6. Changes to the set of Synchronized Devices

Matter Devices can be added to or removed from the set of Synchronized Devices through Administrator-specific means. For example, the user can use a Manufacturer-provided app to disable synchronization of specific devices. When an update to the set of synchronized Devices occurs, the Fabric Synchronizing Administrator SHALL:

- Update the PartsList attribute on the Descriptor clusters of the [Root Node Endpoint](#) and of the endpoint which holds the Aggregator device type.
- Update the exposed endpoints and their descriptors according to the new set of Synchronized Devices

Nodes that wish to be notified of added/removed devices SHOULD monitor changes of the PartsList attribute in the Descriptor cluster on the [Root Node Endpoint](#) and the endpoint which holds the Aggregator device type.

Allocation of endpoints for Synchronized Devices SHALL be performed as described in [Dynamic Endpoint allocation](#).

12.6.7. Fabric Synchronized Relationships

Each Client ecosystem to a Fabric Synchronizing Administrator (possibly multiple unique clients per ecosystem) will use the client ecosystem's fabric to commission the Aggregator. The client ecosystems will each have their own dedicated, isolated fabric that is separate from the fabric used by the Aggregator to interact with the Matter devices.

This relationship can be seen in the [overview figure](#). In that figure, Ecosystem E1 can directly access all devices and the Aggregator of Ecosystem E2 from Ecosystem E1's fabric. This is done without requiring Ecosystem E2 to have any access granted on Ecosystem E1's fabric.

12.6.8. Setup flow for Fabric Synchronization

This section will cover generic setup sequences, which ecosystems MAY extend to provide the desired level of configuration options.

12.6.8.1. Fabric Synchronization Mutual Authentication

As a precondition to enabling the Fabric Synchronization feature, ecosystems MAY wait until the complete bi-directional commissioning of both ecosystems has been completed before exposing any Bridged Nodes. At this point, both ecosystems have received the device attestation from the other.

12.6.8.2. Scope of User Consent

Before or after commissioning a Fabric Synchronizing Administrator, the user SHALL be asked for consent to enable Fabric Synchronization functionality between ecosystems. This matches the existing single-device Administrator-Assisted Commissioning consent model that requires user consent when Matter devices are shared.

Each ecosystem SHALL independently ask the user for consent. This can be done before or after commissioning the device.

Similar to the Multi-Admin flow, after the share operation has been completed, the second Admin can now perform additional sharing operations with the device and/or data according to consent the user has given the second ecosystem.

12.6.8.3. Starting Condition

For the purpose of demonstrating the setup process, consider two ecosystems that provide Fabric Synchronization. These may have existing Matter devices commissioned on their independent fabrics.




Figure 114. Fabric Synchronization Setup - Example Starting Condition.

12.6.8.4. Enabling Fabric Synchronization Between Two Ecosystems

User Action Summary (Informational)

To initiate the setup of Fabric Synchronization, the manufacturer-specific setup SHOULD include the following steps:

- 1. Scan a QR code or enter the Manual Pairing Code of the Commissionee Fabric Synchronizing Administrator. This is similar to the current Administrator-Assisted Commissioning of Matter devices.
- 2. Consent and configure relationships on both ecosystems. (The Commissionee ecosystem MAY provide a configuration step prior to providing the QR code or Manual Pairing Code.)

If full bi-directional Fabric Synchronization is enabled, the end result of the setup process will resemble the Introductory Example.

Initiating Discovery

The user SHALL be able to enable the administrator-assisted commissioning of an ecosystem’s Fabric Synchronization feature through an appropriate interface of the devices on that ecosystem. For example, a mobile application, a web configuration, or an on-device interface.

The Fabric Synchronizing Administrator that advertises its presence over DNS-SD will be designated the Commissionee. The Fabric Synchronizing Administrator that performs service discovery will be designated the Commissioner.

The Commissionee SHALL provide the user with Manual Pairing Code and MAY provide the user with a QR code to initiate commissioning of the Commissionee by the Commissioner.

Discovery

The Commissionee SHALL advertise its presence over DNS-SD (see Section 5.4.2.7, “Using Existing IP-bearing Network” and Section 4.3.1, “Commissionable Node Discovery”).

A Commissioner MAY discover the Commissionee device and provide the user with a notification prior to additional user action.




Figure 115. Fabric Synchronization Setup - Commissioner Relationship Established.

Forward Commissioning

The user SHALL then be able to initiate commissioning on another administrator with the Fabric Synchronization feature using the provided QR code or manual pairing code from the Commissioner.

The Commissioner SHALL commission the Commissionee using the steps outlined in the [Concurrent connection commissioning flow](#).

After commissioning is complete, the Commissionee Fabric Synchronizing Administrator’s Aggregator will now be accessible on the Fabric administrated by the Commissioner.




Figure 116. Fabric Synchronization Setup - Forward Commissioning Complete.

Reverse Commissioning

After the Commissioner has finished commissioning the Commissionee, the Commissioner SHOULD initiate Reverse Commissioning using the Commissioner Control Cluster.

NOTE

A Commissionee Fabric Synchronizing Administrator MAY choose to require reverse commissioning before enabling Fabric Synchronization.




Figure 117. Fabric Synchronization Setup - Reverse Commissioning Complete.

Fabric Synchronization Configuration

After or before commissioning (and optionally Reverse Commissioning,) the device-appropriate interfaces for the Fabric Synchronization feature SHALL ask the user for consent to synchronize devices between fabrics according to [Scope Of User Consent](#).

The following are example configurable options:

- The Fabric Synchronization feature MAY ask the user for consent for all Synchronized Devices or consent for smaller subsets independently.
- The Fabric Synchronization feature MAY ask the user for consent to perform this operation automatically when new Synchronized Devices are commissioned.
- The Fabric Synchronization feature MAY ask the user for consent to share Synchronized Device metadata such as device names and locations.

Fabric Synchronization

Once commissioning and configuration are complete, the ecosystems are now ready to synchronize their fabrics.

The Fabric Synchronizing Administrator SHALL commission the synchronized devices as configured by the user.

Commissioning of Matter devices SHALL be performed by:

1. Discover available Synchronized Devices to commission by identifying endpoints specified in the PartsList of the Aggregator within the Fabric Synchronizing Administrator.
2. Identify if the Synchronized Device supports the BridgedICDSupport feature in the Bridged Device Basic Information Cluster by presence of the BridgedICDSupport feature.
 - a. If the BridgedICDSupport feature is present, the client MAY use the BridgedICDSupport feature to ensure the device is active.
3. Initiate commissioning by sending an OpenCommissioningWindow command to the Administrator Commissioning Cluster exposed on the endpoint with the Bridged Node device type speci-

fied in the PartsList and the Administrator Commissioning Cluster present.

4. Complete commissioning using the Enhanced Commissioning Method.




Figure 118. Fabric Synchronization Setup - Setup Complete.

Device Interaction Summary

The below sequence diagram summarizes the interactions between the Fabric Synchronizing Administrator described in the steps above.




Figure 119. Fabric Synchronization Setup - Example Complete Flow.

Chapter 13. Security Requirements

13.1. Overview

Each Matter security and privacy requirement references the underlying countermeasure (CM) and threat (T) in the Threat Model that motivated the need for the requirement. The requirements are grouped by topic. Unless stated otherwise, these requirements typically address all Devices and Nodes (i.e. all implementations of Matter functionality). Some requirements are called out specifically for a particular group of implementations, or Roles.

13.2. Device vs. Node

For some requirements, we need to differentiate between a Node (the entity which supports the Matter protocol stack) and a Device (a piece of equipment containing one or more Nodes).

- If the Node contains all of the Matter functionality, nevertheless it will probably rely on some security features of the Device.
- If the Node uses Matter functionality provided by the Device, the requirements obviously hold for both the Node and the Device.

13.3. Commissioning

- a. Nodes SHALL stop both commissioning and unsecured rendezvous actions after a specified time period from the beginning of the commissioning mode when commissioning has not been concluded, unless allowed for other purposes such as [Commissionable Node Discovery](#). The time period for commissioning and unsecured rendezvous announcements SHALL follow requirements as specified in [Section 5.5, “Commissioning Flows”](#) and [Section 5.4.2.3, “Announcement Duration”](#) respectively. [CM8 for T5, T7]
- b. Nodes SHALL utilize multiple hash iterations in PBKDF, as required by definitions in [Section 3.9, “Password-Based Key Derivation Function \(PBKDF\)”](#). Nodes SHALL validate the bounds of the iteration count for PBKDF, to respect the minimum and maximum values stated in the cryptography primitives section (see [Section 3.9, “Password-Based Key Derivation Function \(PBKDF\)”](#)). [CM99 in T102]
- c. Nodes SHALL exit commissioning mode after 20 failed commission attempts (see [Section 5.5, “Commissioning Flows”](#)). [CM100 for T101, T112]
- d. Devices SHALL include a Device Attestation Certificate and private key, unique to that Device. [CM23 for T22, T24, T34, T86]
- e. If an [NFC Tag](#) is used to convey the Onboarding Payload from a device to a Commissioner, the NFC Tag SHALL only allow the alteration of the Onboarding Payload and the reading thereof in ways and in device states which attackers cannot exploit to surreptitiously commission the device. For example, the NFC Tag Onboarding Payload is protected from modification outside of manufacturing, or the NFC tag read-out is impossible when the device is still in its original packaging. [CM4 for T90]
- f. The Commissioner or Administrator completing the commissioning, after the first phase has

been executed using NFC-based commissioning, SHALL actively inform a user about devices not appearing on the operational network in the second phase, allowing a user to detect potential re-commissioning over NFC by unauthorized users. [CM276 for T255]

- g. After initial Commissioning of a Device, subsequent Commissioning SHALL only be triggered by an Administrator or an equivalent entity where the user gives administrative consent. [CM2 for T1]

13.4. Factory Reset

- a. Devices and Nodes SHALL have a factory reset capability. [CM15 for T16, T17, T79, T82]
- b. Factory reset SHALL remove from the Node all security- and privacy-related data and key material created during or after commissioning except data explicitly required to persist across resets. [CM35 for T16, T17, T79, T82]

See the following sections for detailed factory reset requirements.

- [Section 6.4.3, “Node Operational Identifier Composition”](#)
- [Section 6.4.11, “Security Considerations”](#)
- [Section 6.6.3, “Access Control List Examples”](#)
- [Section 7.12.1, “Persistence”](#)
- [Section 7.14, “Event”](#)
- [Section 11.12, “General Diagnostics Cluster”](#)
- [Section 11.20.4.2, “Image Verification”](#)
- [NOCs](#)
- [Fabrics](#)
- [CommissionedFabrics](#)
- [TrustedRootCertificates](#)

13.5. Firmware

- a. Nodes SHALL support OTA firmware updates, either using Matter-provided means (see [Section 11.20, “Over-the-Air \(OTA\) Software Update”](#)) or proprietary means. [CM58 for T59, T226]
- b. Nodes SHALL validate the authenticity and integrity of the firmware prior to installation, such as through cryptographic means (see [Section 11.20.4.2, “Image Verification”](#)). [CM58 for T59, T226]
- c. Nodes SHOULD validate configuration and input for length, and acceptable values and ranges before applying them. This validation is dependent on the configuration or input being applied (see [Section 9.10, “Access Control Cluster”](#)). Configuration and input validation is explicitly defined in relevant sections of the specification. [CM46 for T55]

13.6. Security Best Practices

This section describes best practices that Matter implementors SHOULD implement. Matter implementors SHALL indicate whether they comply or not to the best practices. Matter certification will not itself test that these requirements have been met.

13.6.1. Cryptography

- a. Devices and Nodes SHOULD include protection (if it exists) against known remote attacks that can be used to extract or infer cryptographic key material. [CM107 for T94]
- b. Devices SHOULD protect the confidentiality of attestation (DAC) private keys. The level and nature of protection for these keys may vary depending on the nature of the Device. [CM77 for T22]
- c. Nodes SHOULD protect the confidentiality of Node Operational Private Keys. The level and nature of protection for these keys may vary depending on the nature of the Nodes. [CM87 for T87, T110, T120]
- d. For operational PKI CAs, cryptographic operational private keys SHOULD be stored in secure sections of memory. An example could be a Trusted Execution Environment (TEE) or a dedicated cryptographic module [CM287 for T193, T194]
- e. Cryptographic keys SHALL be randomly chosen using a cryptographically secure random number generator in accordance with algorithms defined in [Deterministic Random Bit Generator \(DRBG\)](#). [CM39 for T39]
- f. Devices SHALL use non-repeating initialization vectors for a given session key. [CM78 for T81]

13.6.2. Commissioning and Administration

- a. Manufacturers SHOULD control the number of DACs issued under their Vendor ID. [CM24 for T23, T34]
- b. Where practical, the Onboarding Payload SHOULD be protected from capture from passive visual or photographable methods, regardless of mode of Payload (QR Code, Manual Pairing Code, etc.), e.g., hidden with a flap. [CM89 for T90]
- c. Uncommissioned Devices SHOULD only be available to be commissioned with some form of physical proximity user interaction (e.g. power cycle or button press). [CM3 for T15, T90, T92]
- d. For Devices subject to physical tampering (e.g. doorbell, camera, door lock, devices designed for outdoor use cases), the physical interaction to initiate commissioning and/or the Onboarding Payload (QR code, NFC Tag or Manual Pairing Code) SHOULD NOT be accessible to a physical attacker. E.g. Onboarding Payload is removable or not on the device, the button used to initiate commissioning for the lock is inside the house. [CM4 for T3]
- e. A Commissioner or Administrator SHOULD only add Root Certificates that it trusts to a Node. [CM36 for T153]
- f. A device manufacturer SHOULD implement [Basic Commissioning Method](#) only for devices that adequately secure the [Passcode](#). [CM154 for T173]
- g. Commissioners and Administrators SHOULD carefully control which Nodes get Administer,

Manage, or Operate privilege, especially for safety-critical systems like Door Locks and Smoke CO Alarms. [CM244 for T17, T20, T59, T94, T230]

- h. The NFC-based commissioning session started by a non-authorized user (for instance in a shop) SHALL be prevented. Different solutions exist depending on the form factor of the Matter device. For example, packaging can prevent access to the NFC interface. [CM135 for T131]
- i. NFC-based commissioning requires that commissioning assets are processed by the NFC powered part of the Commissionee Device. As there is potentially a higher risk for certain attacks on the NFC implementation (e.g. if devices are left unattended before completing their installation), depending on the value of the use case, the implementation SHOULD be protected against local logical and physical attacks. [CM62 for T60]
- j. For Locks and Barrier Devices that support [Standard commissioning flow](#), the QR code and power source SHALL be located on the secure side of the property. [CM280 for T260]
- k. For Locks and Barrier Devices that support [Standard commissioning flow](#), the Device documentation SHOULD include a warning to the user about the window of vulnerability upon initial power up during which the Device could be brought onto another network. [CM282 for T260]

13.6.3. Firmware

- a. Vendors of Matter implementations (including their suppliers of Matter functionality) SHOULD have a public vulnerability reporting mechanism and policy and actively monitor, identify and rectify in a timely manner security vulnerabilities throughout the publicly stated security life-cycle policy of the product. Typical responsible disclosure guidelines allow vendors from 60 to 120 days to patch a vulnerability, but the implementation of such a program is at each vendor's discretion. [CM59 for T9, T226]
- b. Devices SHOULD have a verified boot based in an immutable root of trust to verify the authenticity of firmware. Commissioners SHOULD only be loaded on a computing platform that has such a verified boot. If devices cannot support verified boot, devices SHOULD perform verification on any firmware update before applying the new firmware. [CM22 for T16, T20, T226]
- c. The private part of the code signing key (used to sign firmware) SHOULD be strongly protected against disclosure or misuse. For example, it could be stored in an HSM on a secure server outside the factory with very restricted access to only a small number of Device Manufacturer employees. [CM28 for T72]

13.6.4. Manufacturing

- a. Fusing of Devices in production SHOULD be done to limit unintended access to hardware components. For example, vendors may disable debug interfaces, and program trust anchors for secure boot. There are multiple options to secure fusing on the factory floor (e.g., physically securing the fusing station, pre-fusing the silicon, etc). [CM113 for T117]

13.6.5. Resiliency

- a. Matter implementations SHOULD implement resiliency features (e.g., responding to secure boot failures with recovery or error signaling mode) to detect and handle compromise. [CM57 for T59, T226]

13.6.6. Battery Powered Devices

- a. Battery powered Devices SHOULD respond to excessive queries by rate limiting (even limiting the rate to zero if desired). [CM51 for T52, T53]

13.6.7. Tamper Resistance

- a. Protection against physical attacks (especially those that impact cybersecurity) MAY be needed for some Devices, as determined by the manufacturer. [CM62 for T60]

13.6.8. Product Security

- a. Manufacturers SHOULD consider implementing Product Security measures such as those called for in the Alliance's [IoT Device Security Specification](#), based on considerations of device type, use cases, deployment region, regulations, etc. The Alliance's Product Security certification program can be used but this is not required for Matter. [CM272 for T5, T9, T87, T226]

13.6.9. Bridging

- a. Admins SHOULD only grant privileges to a Bridge or Bridged Device (sending commands from that Bridged Device towards a Matter node) that the User is comfortable implicitly granting to all other Bridged Devices exposed by that Bridge. Background: Matter's ACL mechanism does not provide a way to grant privileges to only a single endpoint (Bridged Device) from a node (the Bridge). [CM149 for T162, T165, and T167]

13.6.10. Distributed Compliance Ledger

- a. Vendors SHOULD avail themselves of the DCL store-and-forward functionality so that they can control posting of data about their products to the DCL. [CM160 for T170]
- b. Access to Validator Nodes SHOULD be restricted (e.g., with VPN that only permits Validator Nodes, Observer Nodes, and authenticated clients with write access). [CM163 for T169, T177, T180, T183]
- c. Vendors SHOULD run and use their own Observer Nodes and restrict access to it to make sure that it stays available to the vendors' DCL clients. [CM166 for T180, T182]
- d. Vendors SHOULD protect DCL private keys in Hardware Security Module (HSM) equipped servers. [CM167 for T168]
- e. All parameters passed in transactions and queries to the DCL SHALL pass input validation checks done by the implementation of the DCL. [CM169 for T185]

13.6.11. Safety

- a. Safety-relevant Devices SHOULD revert to a safe state before performing operations that will result in an outage (e.g., restart or firmware update) and, if possible, when an outage is detected. Safety-relevant Devices are Devices whose operation may have a safety or physical impact. The definitions of "safety-relevant" and "safe state" may vary depending on the kind of Device and on that Device's usage. [CM263 for T249]
- b. Vendors SHOULD implement controls to ensure that safety-relevant Devices remain safe and

can enter a failsafe state. Any settings that could lead to a safety or physical impact and which can be modified via Matter clusters, SHOULD be vetted and ignored if outside safe limits. For example, maximum temperature of Water Heaters should remain at factory setting and never change during the lifetime of the device. [CM269 for T55]

- c. Energy Management SHOULD NOT make automated decisions based on unauthenticated data. When using such data, it should treat it as advisory only. Examples include weather data that is pulled in from online resources or tariff and grid event data provided by Utilities. [CM254 to T239]

13.7. Threats and Countermeasures

This section lists identified threats to Matter and countermeasures to mitigate those threats. This section is meant to be informational and not as normative requirements.

Table 125, “Threats” describes the threats, the agent involved in the threat as well as evaluates the consequences. This includes the severity, impact and likelihood of the threat being exploited.

Table 125. Threats

Threat Description			Threat Evaluation		Counter-measure
ID	Description	Threat Agent	Impact/Consequence	Severity	ID
T1	Commission an already Commissioned Node for control that may be difficult to detect (e.g., IP Camera to stream video)	Malicious house guest (brief physical access).	Silent control of Node and access to sensitive Node data (e.g. IP Camera traffic). If only 1 commissioning is allowed, user would detect issue later if/when they try to use Node.	High	CM2
T3	Reset Device and Commission for silent control (e.g. IP Camera to stream video).	Malicious house guest (brief physical access).	Silent control of Node and access to sensitive Node data (e.g. IP Camera traffic). If only 1 commissioning is allowed, user would detect issue later if/when they try to use Node.	High	CM4
T5	IoT device advertises information that can be used to identify vulnerabilities.	Malicious device or person with local network access.	Use information about the Device to exploit a (known) vulnerability.	High	CM6, CM7, CM8

Threat Description			Threat Evaluation		Counter-measure
T7	IoT device advertises information that can be used to identify, profile, or target a home or a user.	Malicious device or person with local network access.	Use information about available accessories to target a given home or user (e.g. to rob).	Medium	CM6, CM7, CM8, CM183
T9	Exploit vulnerability in the Device to gain arbitrary control over Device.	Any.	Unexpected control of Device to gain access to home data, instill fear, etc.	High	CM21, CM22, CM57, CM58, CM59, CM110, CM112, CM259
T15	Commission an uncommissioned Node without physical access to Device	Malicious neighbor or other nearby active attacker	Silent control of Node and access to sensitive Node data (e.g. IP Camera traffic).	High	CM3
T16	Seller of an Device (most likely a used one) intentionally leaves malicious software or configuration on the Device to compromise future traffic.	Malicious Device seller.	Control or access sensitive data of Device.	Medium	CM15, CM16, CM17, CM21, CM22, CM35
T17	Device buyer or trash picker dumps memory to find previous owner's Device keys, group keys, and network credentials.	Malicious Device buyer or trash picker.	Access to sensitive data and ability to inject trusted data or even commands.	Medium	CM15, CM16, CM17, CM35, CM244
T20	Firmware (any software on Device that can be modified) is modified by attacker in factory (or remotely through factory)	Worker at factory or programming location or remote attacker	1. Local network behavior issues 2. Infected Nodes leading to data breach, malfunction, denial of service, or attacks by this Node on other Nodes	Medium	CM21, CM22, CM244

Threat Description			Threat Evaluation		Counter-measure
T22	Cloned Device produced (with identical credentials to a proper Device).	Anyone with physical access to a Device from which they can extract Device Attestation credentials.	1. Brand damage if Device is of lower quality. 2. Loss of revenue. 3. Lack of function or interoperability if Device does not function properly. 4. Lack of security if Device does not have proper security measures. 5. Lack of support from manufacturer for cloned Devices.	Medium	CM23, CM77
T23	Counterfeit Device produced (with unique but apparently authorized credentials)	Worker at factory or programming location	1. Brand damage if Device is of lower quality. 2. Loss of revenue. 3. Lack of function or interoperability if Device does not function properly. 4. Lack of security if Device does not have proper security measures. 5. Lack of support from manufacturer for cloned Devices.	Medium	CM24
T24	Factory provisioned keys captured in factory, transit, or store (e.g., with fault injection or other tampering).	1. Worker in supply chain. 2. Attacker who goes to retail store	Control of Device and access to sensitive Device data (e.g. IP Camera traffic).	Medium	CM23, CM113
T34	Cloned Device enters the network.	Manufacturer selling cloned Devices.	Loss of revenue or brand issues for original manufacturer.	Low	CM23, CM24

Threat Description			Threat Evaluation		Counter-measure
T39	Guessing security keys via brute force attack.	Attacker within radio range, captures encrypted network packets.	Control or access sensitive data of Device. Even control entire network if the private key for the Operational Certificate of a Controller can be guessed.	High	CM39
T52	Malicious Device off the network causes battery powered Device to wake too often.	Attacker using a Device on the network.	Shortened battery life of nearby Devices.	Medium	CM44, CM51
T53	Malicious Device off the network interrupts battery powered messages too often and greatly reduces battery life.	Attacker using a Device on the network.	Shortened battery life of nearby Devices.	Medium	CM44, CM51
T55	Device reconfigured improperly.	Attacker using a Device on the network.	Device could be configured such that it does not properly behave.	Medium	CM45, CM46, CM47, CM269
T59	Maliciously crafted message exploits Device vulnerability, causing Device compromise.	Attacker using a Device on the network.	Trusted Device could be hijacked.	High	CM57, CM58, CM244
T60	Physical tampering with Device permits compromise.	Attacker with physical access to Device.	Trusted Device could be hijacked.	Medium	CM62
T72	Code signing key copied, permitting signing of unauthorized code.	Attacker gets administrative access to code signing server.	1. Bad images could be installed on many Devices 2. Devices could be destroyed or hijacked, even invisibly and irreparably.	High	CM28

Threat Description			Threat Evaluation		Counter-measure
T79	Device marked for destruction reused in network.	Installer or return agent.	Damaged or obsolete Devices may re-enter the network.	Low	CM15, CM16, CM20, CM35
T81	Attacker uses predictable Initialization Vectors to derive crypto keys.	Attacker able to observe network traffic from the Device.	Attacker discovery of Device crypto keys and other secrets, which can lead to control of other Devices if the Device has such privileges.	High	CM78
T82	Device buyer dumps memory to find previous owner's user data.	Malicious Device buyer or dumpster diver.	User data may be leaked.	Medium	CM15, CM35
T86	Leak certificate or Device identity private key to appear as valid Device.	Physical or locally compromised attacker.	Device appears as valid Device.	Low	CM23
T87	Malicious Device or party poses as valid Matter Node using operational credentials	Attacker on the local network or remotely controlling a Node on the local network	Group keys and sensitive data revealed to an invalid Node	Medium	CM87
T90	Long range camera captures QR code at Commissioning time or otherwise.	Malicious neighbor, robber, or other nearby active attacker.	Attacker manages to connect Device to their gateway or account.	Medium	CM3, CM89
T92	Microphone in the house can capture person speaking the Manual Pairing Code and use that to MITM Commissioning.	Malicious neighbor, robber, or other nearby active attacker	Attacker manages to connect the Node to their gateway or account	Medium	CM3

Threat Description			Threat Evaluation		Counter-measure
T94	Remote attack used to extract keys and other secrets.	Attacker able to access the Device remotely or over local network.	Attacker discovery of Device crypto keys and other secrets, which can lead to control of other Devices if the Device has such privileges.	High	CM107, CM244
T101	Malicious Device or person with local network access attempts to guess Passcode via online brute force attack.	Attacker on the local network.	Control of Device and access to sensitive Device data (e.g. IP Camera traffic).	High	CM5, CM100
T102	Malicious Device or person with knowledge of passcode verifier uses offline brute force attack to derive Passcode.	Attacker with remote access.	Control of Device and access to sensitive Device data (e.g. IP Camera traffic).	High	CM5, CM99
T110	Malicious device or party poses as valid Matter Administrative Node using operational credentials	Attacker on the local network or remotely controlling a Node on the local network	Control of Node and access to sensitive Node data (e.g., IP camera traffic)	High	CM87
T112	Malicious Device(s) or person(s) with local network access attempts to guess Passcode of many Devices via online brute force attack.	Attackers on the local network.	Control of Device and access to sensitive Device data (e.g. IP Camera traffic).	Medium	CM5, CM100
T117	Incorrect fusing of production Devices.	Contract manufacturer, factory worker.	Device assets are vulnerable, security policies including secure boot might not be enforceable, etc.	High	CM113

Threat Description			Threat Evaluation		Counter-measure
T120	Data from Matter Nodes is shared with non-Matter or unauthorized entities (e.g. data leakage to adjacent non-Matter Nodes)	Any adversarial process running in the node that has enough privileges to modify ACLs. Secure boot is not sufficient protection if the device boots rarely and the malicious process was spawned post-boot up.	Matter data may be used in inappropriate or unauthorized ways potentially harming the owner.	High	CM87
T131	Attacker reads the NFC code of a device while it is being transported or at a mailbox	Attacker with close proximity to device (10cm with standard hw and ~40cm with specialized equipment)	Attacker manages to connect Device to their gateway or account; attacker harvests VID and PID for large number of NFC devices.	High	CM135
T153	Commissioner/Administrator adds untrustworthy Root CA to Node.	Malicious, compromised, or poorly advised Commissioner/Administrator.	Attacker can create NOCs that enable impersonation and MITM of Secure Channels.	High	CM36
T162	Compromise of Bridged Device. Bridged device can have lower security than Matter device. (Applicable to bridged devices controlling Matter device or exposing data to Matter devices. Not applicable to bridged devices being controlled from Matter devices.)	Attacker who can compromise a bridged device.	Control of Matter-side Devices and access to sensitive Device data. Control Camera, unlock door.	High	CM149

Threat Description			Threat Evaluation		Counter-measure
T165	Privilege Escalation. Bridged device is able to impersonate higher privileged bridged device. Unlock door on presence. (Applicable to bridged devices controlling Matter device or exposing data to Matter devices. Not applicable to bridged devices being controlled from Matter devices.)	Attacker that controls Bridged Device can use that device to impersonate other Bridged Devices (potentially ones that have more access to the Matter network.)	Control of all interactions between Matter and bridged devices. This can include control of Matter devices.	High	CM149
T167	Attacker with privileges on Bridged Ecosystem	Attacker has no privileges on Matter fabrics but does have privileges on Bridged Ecosystems.	Control of all interactions between Matter and bridged devices. This can include control of Matter devices.	High	CM149
T168	DCL Private Key Exfiltration.	Attacker obtains one or more private keys of a company with DCL writer privilege.	Attacker can add to the block chain on behalf of the company. Can change the OtaUrl to point to old and known-vulnerable firmware or prevent an update from being installed.	High	CM167
T169	DoS/DDoS of Validators Nodes.	Attack can direct a DoS attack or resource exhaustion attack against validators. Attacker only needs to DoS 1/3+1 of validators to DoS consensus.	New blocks cannot be added.	High	CM163

Threat Description			Threat Evaluation		Counter-measure
T170	Unintended or premature exposure of information.	Company or certification lab posts device details to the chain and it is validated.	Immutability of block chain means the information is permanently on the chain.	High	CM160
T173	Malicious Device or person with local network access and knowledge of the pass-code attempts to pair with a commissioned device when someone else opens the commissioning window using Section 5.6.2, “Basic Commissioning Method (BCM)” and the device’s Passcode.	Attacker on the local network	Control of Device and access to sensitive Device data (e.g. IP Camera traffic)	Medium	CM41, CM152, CM154
T174	Malicious Device gains knowledge of the Passcode on an uncommissioned Device, commissions the Device, and then puts it back into commissioning mode with the same Passcode using Section 5.6.2, “Basic Commissioning Method (BCM)” or Section 5.6.3, “Enhanced Commissioning Method (ECM)” to avoid detection.	Attacker on the local network	Control of Device and access to sensitive Device data (e.g. IP Camera traffic)	Medium	CM41, CM152

Threat Description			Threat Evaluation		Counter-measure
T175	Malicious Device with knowledge of the Passcode commissions an uncommissioned Device and then puts it back into commissioning mode with the same Passcode using Section 5.6.2, “Basic Commissioning Method (BCM)” or Section 5.6.3, “Enhanced Commissioning Method (ECM)” to avoid detection.	Attacker on the local network	Control of Device and access to sensitive Device data (e.g. IP Camera traffic)	Medium	CM41, CM152
T177	Attacker exploits a vulnerability that is common to most or all of the Validator Node software.	Attacker with some sort of DCL access (maybe just read, which is open to all).	Many Validator Nodes misbehave (e.g., approving adding or revoking a PAA or changing an OtaURL).	High	CM163
T180	Attacker accesses Observer Node or Validator Node with unauthenticated READ CLI protocol, mounts a DoS or DDoS attack.	Attacker that can send network messages to a DCL Observer Node or Validator Node.	DCL capacity diminished or eliminated. Unable to communicate important things like revocation of PAA.	High	CM163, CM166
T182	DoS/DDoS of Observers Nodes.	Attack can direct a DoS attack or resource exhaustion attack against Observer Nodes. If enough Observer Nodes are impacted by a DoS attack, the DCL may become unavailable.	DCL unavailable	High	CM166

Threat Description			Threat Evaluation		Counter-measure
T183	DoS on Trustees' approval process.	Submit many PROPOSE_ADD_ACCOUNT requests. The Trustees can be overwhelmed with illegitimate requests. Requires compromise of a Trustee, although replay is possible.	Trustees overloaded	Medium	CM163
T185	DCL Denial of Service. Attacker writes a value to the ledger that is very large or out of bounds.	Authorized attacker sends a write request with very large parameter payload.	Very large ledger blocks added to ledger. This could cause validation problems. DOS on Observer Nodes if response is very large.	High	CM169
T193	Compromised Operational Root CA	Attacker that is able to gain control of the Operational Root CA or obtain its private key.	Introduce an unauthorized or untrustworthy device into one or more homes, enable impersonation within a home by issuing intermediate operational certs, that can then be used to issue operational device certs.	High	CM181, CM287
T194	Compromised Operational Intermediate CA	Attacker that is able to gain control of the Operational Intermediate CA or obtain its private key.	Introduce an unauthorized or untrustworthy device into one or more homes, enable impersonation within a home by issuing certs with duplicate subject names and different private keys, or surveil device introduction into homes by monitoring operational certificate issuance.	High	CM287

Threat Description			Threat Evaluation		Counter-measure
T226	False information is provided by a compromised, vulnerable, or malicious device.	Attacker that is able to compromise an existing vulnerable device or introduce a new malicious device	Impacts may include triggering an action (e.g., unlocking a door due to a false report of a person detected in the home) or tricking the user into thinking configuring the wrong device (e.g., changing the NodeLabel of a device to impersonate another device).	High	CM21, CM22, CM57, CM58, CM59
T230	Spurious commands are sent to a Device, leading the consumer to disable it.	Attacker authorized by ACL to send commands to Device.	Consumer frustration with and distrust of Device leading consumer to disable Device	High	CM244
T231	An attacker remotely starts a household appliance that belongs to a device type which may cause harm without attended confirmation.	Attacker that is able to access the device remotely	Cause a flood or a fire or electric overload in a home	High	CM245
T232	CA or other parties failing to meet obligations laid out in the Certificate Policy	Attacker able to exploit undressed Certificate Policy rules and/or absence of technical and operational controls associated with the identity information, key generation, distribution, revocation, and administration of certificates	Loss of relying party trust	High	CM24, CM101, CM117

Threat Description			Threat Evaluation		Counter-measure
T233	Malicious PAA/PAI/DAC revocation request initiated by DCL observer node	Attacker able to compromise or control a DCL observer node able to corrupt revocation information to clients	Illegitimate revocation information shared to clients	High	CM199
T234	Malicious proposal and revocation of PAA/PAI/DAC using compromised Vendor's DCL Account and its associated Key	Attacker able to compromise or control a DCL account and its associated key (Vendor, Trustee, NodeAdmin, Cert-Center) able to propose and/or revoke certificates	Denial of Service	High	CM76, CM167
T239	Ability to influence many devices with large electric load in a small geographic area could permit an attack to increase or decrease load simultaneously and impact the power grid	Influence Energy Management decision making	Damage to power grid	High	CM254
T240	Ability to perform phishing exploits using a compromised Commissionee Device or Video content sharing App	Client uses Screen Sharing and Messaging feature to trick the customer into thinking a message comes from a different source. User responds thinking this is the TV or Content App asking	Loss of relying party trust. Display of false information, including malicious instructions	High	CM248

Threat Description			Threat Evaluation		Counter-measure
T241	Exposure of data in transit using MITM, remote code execution and adversary lateral movement exploiting WebRTC communication protocol	Some clusters utilize WebRTC. By default, the WebRTC communication is potentially vulnerable to eavesdropping and MITM	Denial of Service	High	CM249
T242	Lack of traceability and identification of Commissionees (Video content sharing apps) temporarily using the TV for Video content-sharing	Airbnb/Hoteling scenario - Renter commissions a device as a message source using the on-screen flow. Homeowner returns and does not know which actual device is sending messages to their TV. Maybe the device is hidden in a closet	Loss of relying party trust. Display of false information, including malicious instructions	High	CM248, CM250
T243	Lack of mechanisms to establish parental control requirements as per prevailing standards and regional regulations	Customer uses Content Control feature to set Parental Controls on TV to limit access to content with rating N, child opens an app, app allows child to access content with rating N	Regulatory and compliance risk - Loss of customer trust or brand issues for original manufacturer	High	CM251

Threat Description			Threat Evaluation		Counter-measure
T244	Unknown commissioned node acting maliciously (ex. Video content sharing app from a client device)	Unrecognized NodeId. Adversary hijacked client or phone app sends commands or eavesdrops on network traffic - Ex: a UDC/CDC message indicating PIN code dialog is requested. Attacker leverages a hijacked camera to see PIN code displayed on TV screen	Loss of relying party trust. Display of false information, including malicious instructions	High	CM248, CM250, CM253
T245	Adversary exploits TV static passcode / PIN feature to request that the TV generate and display a Passcode	TV reuses the same Passcode code every time (no dynamic passcode). Vulnerable to brute-forcing SPAKE2 (during PASE)	Denial of Service	High	CM252
T246	Adversary eavesdrops on network traffic and sees/sends UDC/CDC messages exploiting a device vulnerability that allows adversary to use information for discovering available apps and push prompts on the TV	Attacker uses UDC to mount a DoS by continually sending messages to the TV to request commissioning, allows UDC flow to partially complete (performs commissionable node discovery) and times out, effectively blocking any other client from using this feature	Denial of Service	High	CM253
T249	Attacker triggers Device outage while Device is in an unsafe state	Attacker able to cause Device outage (e.g., through Denial of Service)	Attacker can keep Device in unsafe state indefinitely, potentially leading to damage	Medium	CM263

Threat Description			Threat Evaluation		Counter-measure
T255	An attacker re-starts the commissioning process after the first phase over the NFC Transport Layer has already been completed by the user (or installer). Then the Device may be commissioned on the attackers' Matter network.	Attacker with close proximity to device (1m to 10m with specialized equipment)	Attacker manages to connect Device to their gateway or account	Medium	CM276
T260	A malicious individual causes a distraction during initial installation, in order to steal the QR Code or Pass code, in an attempt to onboard device to their network during initial commissioning.	Malicious neighbor or other nearby active attacker	An individual maliciously connects the Device to their network, gaining control and unlocking it to physically access property.	Medium	CM280, CM282

Table 126, “Countermeasures” describes the various countermeasures to the threats listed in Table 125, “Threats”.

Table 126. Countermeasures

ID	Description
CM2	After initial Commissioning of a Device, subsequent Commissioning can only be triggered by an Administrator or an equivalent entity where the user gives administrative consent.
CM3	Commissioning is started with some form of physical user interaction (e.g. power cycle , NFC tap or button press).
CM4	For Devices subject to physical tampering (e.g. doorbell, camera, door lock, devices designed for outdoor use cases), the physical interaction to initiate commissioning and/or the Onboarding Payload is not accessible to a physical attacker. E.g. setup Passcode is removable or not on the device, the button for the lock is inside the house.

ID	Description
CM5	All Devices include a randomly generated setup passcode and a corresponding passcode verifier derived from the setup passcode via a PBKDF. The passcode includes at least 27 bits of entropy compliant with a recognized standard (e.g., NIST SP 800-90B).
CM6	Unsecured rendezvous are enabled by a user action and, upon time-out or commissioning failure, will cause deletion of any state information. Examples of "user actions" are pressing a physical button, power-cycling a Device, and leveraging a previously commissioned account.
CM7	Minimize OS and other version information advertised during discovery.
CM8	Both commissioning and unsecured rendezvous actions time-out after at most 15 minutes from the beginning of the commissioning mode when commissioning has not been concluded.
CM15	Devices have a physical button or trigger for factory reset.
CM16	Devices rotate keys at specified triggers (e.g. Factory Device Reset).
CM17	Devices implement Perfect Forward secrecy key agreement protocols that give assurances that session keys will not be compromised even if long-term secrets used in the session key exchange are compromised.
CM20	Revoke Device credentials and privileges when the Device is removed from the home.
CM21	Devices have cryptographically signed firmware, including all firmware and software on the Device.
CM22	Devices have a verified boot based in an immutable root of trust to verify the authenticity of firmware.
CM23	All Devices include a Device Attestation Certificate and private key, unique to that Device.
CM24	Manufacturers control the number of DACs issued under their Vendor ID.
CM28	Private part of code signing key is strongly protected against disclosure or misuse. For example, it could be stored in an HSM on a secure server outside the factory with very restricted access to only a small number of Device Manufacturer employees.
CM35	Factory reset removes all local data and key material created during or after commissioning except data explicitly required to persist across resets.
CM36	A Commissioner / Administrator only adds Root Certificates that it trusts to a node.
CM39	Cryptographic keys are randomly chosen using the strong entropy separately required and the cryptographic algorithms and key lengths specified by Matter.
CM41	Administrators can view the set of Fabrics on each Device (i.e., Attributes for the Node Operational Credentials Cluster).

ID	Description
CM44	Administrator responds to reported or detected attacks and malfunctions (e.g., by adding Devices to a denylist, notifying the user, changing group keys, or hopping channels).
CM45	Configuration for secure channel protocol is carefully negotiated and validated by both parties.
CM46	Devices validate configuration and input changes for length, character type, and acceptable values and ranges before applying them. This validation is dependent on the configuration or input being applied (e.g. ACL entries). Configuration and input validation is explicitly defined in relevant sections of the specification.
CM47	Device management service uses a secure communication mechanism for reconfiguration.
CM51	Battery powered Devices respond to excessive queries by rate limiting (even limiting the rate to zero if desired).
CM57	Devices implement resiliency features (e.g., responding to secure boot failures with recovery or error signaling mode) to detect and handle compromise.
CM58	Devices support OTA firmware updates. Devices validate the authenticity and integrity of the firmware prior to installation.
CM59	Manufacturers monitor newly discovered vulnerabilities and provide software updates to address them.
CM62	Protection against physical attacks (especially those that impact cybersecurity) is needed for some Devices, as determined by the manufacturer.
CM76	All DCL writers constantly watch for DCL changes that relate to them and complain about them if they are unauthorized.
CM77	All Devices protect the confidentiality of attestation (DAC) private keys. The level and nature of protection for these keys may vary depending on the nature of the Device.
CM78	Devices use random initialization vectors.
CM87	All Nodes protect the confidentiality of operational credential private keys. The level and nature of protection for these keys may vary depending on the nature of the Nodes.
CM89	Where practical, the Onboarding Payload is protected from capture from passive visual or photographable methods, regardless of mode of Payload (QR Code, Manual Pairing Code, etc.), e.g., hidden with a flap.
CM99	Devices utilize multiple hashes in PBKDF.
CM100	Device exits commissioning mode after 20 failed commissioning attempts.
CM101	Use PKI best practices for all CAs, including secure backup of private key, use of an HSM, four eyes principle, etc. These are described in the PKI CP.
CM107	Devices include protection (if it exists) against known remote attacks that can be used to extract or infer cryptographic key material.

ID	Description
CM110	Devices support remote Firmware attestation with trusted root of measurement to allow local and cloud agents to validate the cryptographic identity of the Device firmware and refuse to perform actions on behalf of or by relying to information acquired by known compromised versions. Firmware attestation is performed at initial commissioning and when revocation lists get updated.
CM112	ACLs restrict access granted to Nodes.
CM113	Fusing of Production Devices is done correctly. For example, disabling debug interfaces, and programming trust anchors for secure boot. There are multiple options to secure fusing on the factory floor (e.g., physically securing the fusing station, pre-fusing the silicon, etc).
CM117	Design, implement, and operate the PAI and associated infrastructure so that unauthorized Devices cannot be created with proper credentials.
CM132	Use passive mode NFC and lowest power on the initiator if possible — this reduces the range of NFC making eavesdropping more difficult. Doing so, however, this will reduce the distance at which the NFC interaction will be effective and, possibly, the device setup experience.
CM133	If possible, use RSSI on at least one of the two NFC endpoints to disregard high power radiation from attackers. Attackers cannot always control the strength of the field near the legitimate endpoints if they are far away and excess power can be detected.
CM135	Packaging protects against RF devices reading passive NFC tags.
CM136	Position NFC tags to be readable only from a certain direction.
CM139	All cryptographic keys and secrets are protected against unauthorized modifications on the device, at rest and during compute. This could include protections against hardware attacks or logical attacks.
CM149	Administrators only grant privileges to a Bridge if the Administrator is comfortable granting those same privileges to all Bridged Devices behind that Bridge.
CM152	Device manufacturers provide a way to secure a static Passcode after initial commissioning so that it is not available for unauthorized agents.
CM154	A device manufacturer implements Section 5.6.2, “Basic Commissioning Method (BCM)” only for devices that adequately secure the Passcode .
CM160	Vendors sign off on some other entity posting data about their products to the DCL.
CM163	Tightly restrict access to Validator Nodes (e.g., with VPN that only permits Validator Nodes, Observer Nodes, and authenticated clients with write access).
CM166	Matter vendors run and use their own Observer Node and restrict access to it to make sure that it stays available to that company’s DCL clients.
CM167	Matter vendors protect DCL private keys in HSM equipped servers.
CM169	All parameters passed in transactions and queries to the DCL pass input validation checks.

ID	Description
CM181	Include the ability to update the list of trusted Operational Root CAs (i.e., by removing a Fabric)
CM183	VID & PID are not advertised before Commissioning.
CM199	Light client cryptographically verifies responses received from nodes.
CM244	Commissioners and Administrators carefully control which Nodes get Administer, Manage, or Operate privilege, especially for safety-critical systems like Door Locks and Smoke CO Alarms.
CM245	Device manufacturers consider carefully whether their products that support remotely starting its operation could do harm to the end user; if so, require some additional operations (e.g. attended confirmation for remote start).
CM248	Devices with content launching features require source of content origin and perform verification of origin information and display human-friendly information before sharing Video content and/or commands to perform sharing of content
CM249	The Screen device and content sharing client support implementing the minimum required WebRTC security requirements as stated in w3c and webrtc-security WebRTC security requirements (see https://www.w3.org/TR/webrtc/#privacy-and-security-considerations and https://webrtc-security.github.io)
CM250	The commissioner require utilizing ACL or similar access control mechanisms for uniquely identifying and controlling access by Video content-sharing clients so that access control can be managed and revoked without impacting the primary device admin account
CM251	Devices supporting content control functionality will enforce the required parental controls specific to the content being shared and inform the user of limitations of this control, for example, when these settings do not apply to content provided by Content Apps on the TV
CM252	The Screen device provide access using dynamic passcode or user-defined passphrase (similar to other casting protocols) with options to reset frequently
CM253	Detect and surface to the user suspicious behavior by unauthenticated network clients and provide the user with the option to disable any information or functionality exposed to unauthenticated clients
CM254	Treat unauthenticated data as advisory only
CM259	Include the ability to revoke an Operational Certificate
CM263	Device reverts to safe state before performing unsafe operations and in case of outage
CM269	Device vets settings that could lead to a safety or physical impact and ignores if outside safe limits.
CM272	Products implement widely agreed Product Security best practices.

ID	Description
CM276	The Commissioner or Administrator completing the commissioning, after the first phase has been executed using NFC-based commissioning, SHALL actively inform a user about devices not appearing on the operational network in the second phase, allowing a user to detect potential re-commissioning over NFC by unauthorized users.
CM280	For physical security devices like door locks, barriers, and closures, if the device supports standard commissioning then the QR code and power source are located on the secure side of the property.
CM282	For physical security devices like door locks, barriers, and closures, if the device supports standard commissioning then the device documentation includes a warning to the user about the window of vulnerability after initial power up.
CM287	For operational PKI CAs, cryptographic operational private keys are stored in secure sections of memory. An example could be a Trusted Execution Environment (TEE) or a dedicated cryptographic module.

Chapter 14. Transport Layer Security

14.1. Introduction

Transport Layer Security (TLS) is a cryptographic protocol designed to secure data transmission over computer networks and is widely used to ensure data privacy, integrity and authenticity. Matter TLS compatibility is intended to provide a mechanism for Application layers to use TLS to communicate securely with devices or servers outside the Matter [Fabric](#).

Since TLS employs digital certificates to authenticate both the identity of servers and clients, this chapter discusses the provisioning of nodes with X.509 digital certificates, as well as the provisioning of nodes with other TLS endpoint information needed to facilitate TLS communication.

Nodes MAY choose to use TLS in other, non-Matter, contexts. Therefore in this chapter all references to Nodes supporting or using TLS, and the mechanisms defined to manage TLS connections, pertain only to the use of the TLS within Matter features that depend on it.

14.2. Common Conventions

The following conventions are used.

14.2.1. Transport Layer Security (TLS)

Nodes supporting TLS SHALL implement Versions 1.3 (as defined in [RFC 8446](#)) or later, and SHALL NOT support any earlier version. Matter's use of TLS SHALL support only key exchange/cipher suites which provide Perfect Forward Security.

Since Matter TLS connections by their nature connect to services outside the Matter Fabric, TLS connections MAY utilize IPv4 transport. Ecosystems implementing Matter linked TLS services SHOULD therefore ensure their TLS services have both IPv6 and IPv4 connectivity.

14.2.2. TLS Certificates

TLS certificates used within this chapter SHALL be based on X.509v3-compliant certificates as defined in [RFC 5280](#). TLS certificates used in this chapter SHALL be encoded using the Distinguished Encoding Rules (DER) format as defined in [X.690](#).

14.2.2.1. X.509 Certificate Fingerprint

An X.509 certificate fingerprint is created by applying a cryptographic hash function to the X.509 certificate's binary data, enabling easy comparison and verification. X.509 certificate fingerprints used in this chapter SHALL be calculated by using the SHA-256 hash algorithm as defined in [RFC 6234](#), using the DER (binary) encoding format of the certificate.

14.2.2.2. Certificate Authority (CA)

A Certificate Authority (CA) is a trusted entity that issues digital certificates. These digital certificates are used in Public Key Infrastructure (PKI) to facilitate secure communication.

14.2.2.3. TLS Root CA Certificate (TLSRCAC)

A TLS Root CA Certificate (TLSRCAC) is a special type of digital certificate that plays a fundamental role in the Public Key Infrastructure (PKI). It is the topmost certificate in the hierarchical structure of the certificate chain and holds the highest level of trust. TLS clients use TLSRCACs to authenticate a server during the handshake process, using the "Chain of Trust" described in [RFC 4158](#).

14.2.2.4. TLS Client Certificate

A TLS Client Certificate is a type of digital certificate used to authenticate clients during the TLS handshake. Use of client certificates in TLS is covered in [RFC 8446](#) (or later).

14.2.2.5. Certificate Signing Request (CSR)

A Certificate Signing Request (CSR) is used by an entity to request a CA to issue a digital certificate to that entity. The CSR contains information about the entity's public key and MAY contain other relevant information. All CSRs in this chapter SHALL be created following the format and procedure in [PKCS #10](#).

14.2.3. TLS Endpoint

A TLS Endpoint is a specific network address or hostname to which communication using TLS can be established as defined in [RFC 8446](#) (or later), and thus represents a potential connection from the Node to another device. A TLS Endpoint can be either a TLS Client, which is the TLS Endpoint initiating the connection, or a TLS Server, which is the TLS Endpoint serving the requests from the TLS Client. In this chapter the Node is presumed to be the TLS Client, and the TLS Server is some other TLS Endpoint that the Node wishes to connect to.

In the context of this chapter, a TLS Endpoint is used to refer to the combination of a TLS Hostname and a TLS Port Number, along with certain other details, representing a TLS Server.

TLS Endpoints are managed by the [TLS Client Management Cluster](#).

14.2.3.1. TLS Hostname

A TLS Hostname SHALL be a domain name or IP Address. For cloud based TLS servers this MAY be a public resolvable DNS name. For local TLS servers on the same broadcast domain or locally reachable IP Addresses, this MAY be a multicast DNS based .local name, or other resolvable private name.

14.2.3.2. TLS Port Number

A TLS Port Number is the network port designated for TLS connections. If a Port Number is omitted, the value of 443 SHALL be assumed.

14.2.4. TLS Connection

A TLS Connection is defined as a connection between two TLS Endpoints, one of which is the TLS client side and the other as the TLS server side. Establishment and usage of all TLS Connections uses the standard procedure defined in [RFC 8446](#) (or later). A TLS Client Connection is defined as a TLS Connection where the Node is the TLS client side.

Only TLS Client Connections SHALL be allowed using the clusters and commands in this chapter.

14.3. Description

TLS compatibility will be enabled by the following features.

14.3.1. TLS Certificate Management

This cluster allows for the provisioning of X.509 Root Certificates and the creation of and management of X.509 Client Certificates as used with the TLS protocol implementation.

Nodes wishing to use TLS as part of Matter SHALL support the [TLS Certificate Management Cluster](#).

Since standard Web PKI and TLS policy performs time and date validation of all X.509 certificates in the chain, Nodes supporting TLS Certificate Management SHALL support the [Time Synchronization](#) cluster with either the [NTPClient](#) or [TimeSyncClient](#) feature enabled. See [Section 11.17.11, “Time Synchronization at Commissioning”](#) and [Section 11.17.12, “Time Synchronization during operation”](#) for additional details on how time information gets set and updated.

14.3.1.1. TLS Root CA Certificate Management

Authentication of a server ensures that the client is connecting to the correct server and not to an impostor. TLS uses Root CA Certificates to authenticate a server during the handshake process.

Nodes wishing to provide Root CA Certificate Management SHALL support the [TLS Certificate Management Cluster](#). In addition, Nodes SHALL provide the following.

Certificate Authority ID (CAID)

A Certificate Authority ID ([CAID](#)) is used to uniquely identify a provisioned [TLSRCAC](#) associated with a Fabric on a Node.

Nodes SHALL provide a mechanism where the CAID can be used to get, update, or remove the [TLSRCAC](#) associated with the CAID (see [TLS Certificate Management Commands](#)).

14.3.1.2. TLS Client Certificate Management

Some TLS Servers MAY choose to implement Client Certificate Authentication. To facilitate this, Nodes MAY choose to provision Client Certificate Details ([TLSClientCertificateDetailStruct](#)) which are associated with a Fabric for use in TLS communication. Client Certificate Details (CCD) is used to refer to the combination of client certificate, private key and chain of intermediate certificates that is needed to complete the TLS handshake during Client Authentication.

Nodes wishing to provide Client Certificate Management SHALL support the [TLS Certificate Management Cluster](#).

Client Certificate Details ID (CCDID)

A Client Certificate Details ID ([CCDID](#)) is used to uniquely identify Client Certificate Details associated with a Fabric on a Node.

A Node SHALL only store a single Client Certificate Details for a particular CCDID. Nodes SHALL provide a mechanism where the CCDID can be used to retrieve, update or remove the CCD associated with the CCDID (see [TLS Client Management Commands](#)).

TLS Client Certificate Signing Request (CSR) Procedure

TLS client certificates can be generated, installed, and rotated using the [ClientCSR](#) and [ClientCSRResponse](#) commands.

The TLS client certificate issuance process uses a freshness check by having the client send a nonce value that must be signed in addition to the signature returned inside the X.509 certificate signing request.

TLS Client CSR Generation, Validation, and Provisioning process

1. The client SHALL generate a random 32 byte nonce using [Crypto_DRBG\(\)](#).
2. The client SHALL send the nonce to the Node and request a CSR using the [ClientCSR](#) command. If rotating, an existing CCDID SHALL be provided.
3. The client SHALL verify the inner signature in the [PKCS #10](#) of the [CSR](#) field in the [ClientCSRResponse](#), per the definition of CSR signatures in [PKCS #10](#).
4. The client SHALL verify [NonceSignature](#) field on the [ClientCSRResponse](#) is valid against the nonce value it provided in the corresponding [ClientCSR](#) command.
5. The client SHALL verify that the public key used for the [NonceSignature](#) is the same public key as the inner signature in the [PKCS #10](#) of the [CSR](#) field.
6. The client performs a procedure not covered in this specification to create and sign a TLS Client Certificate from the returned [CSR](#) field.
7. The client SHALL provision the resulting TLS Client Certificate and any intermediate certificate chain into the NODE using the [ProvisionClientCertificate](#) command using the [CCDID](#) value returned in the [ClientCSRResponse](#).

14.3.2. TLS Client Management

Nodes which have application or utility clusters that use TLS client connections SHALL support the following:

1. [TLS Client Management Cluster](#)
2. [TLS Certificate Management Cluster](#)
3. [Time Synchronization Cluster](#)

This is intended to allow provisioning on the Node of sufficient information to facilitate subsequent TLS client connections and validate the server presented certificates used in those connections.

14.3.2.1. TLS Endpoints

TLS endpoints SHALL be provisioned using the [ProvisionEndpoint](#) command and removed when no longer needed using the [RemoveEndpoint](#) command. In addition, Nodes SHALL provide the following:

14.3.3. TLS Minimum Requirements

This section describes the minimum requirements when supporting the various TLS clusters.

14.3.3.1. Minimum storage for TLS Certificates

Nodes SHALL provide enough storage space for at least 5 Root certificates per Fabric on a Node. Certain device types MAY choose to increase this minimum.

Nodes SHALL provide enough storage space for at least 2 Client Certificate Details per Fabric on a Node. Certain device types MAY choose to increase this minimum.

14.3.3.2. Minimum storage for TLS Endpoints

Nodes SHALL provide enough storage to provision at least 2 TLS Endpoints per Fabric. Certain device types MAY choose to increase this minimum. In addition, Nodes SHALL be capable of at least 2 concurrent TLS client connections per Fabric using the provisioned TLS Endpoints.

14.4. TLS Certificate Management Cluster

This cluster is used to manage TLS CA Root and Client Certificates on a Node, which are then used by other clusters to provision and manage their usage of TLS.

Commands in this cluster uniformly use the [Large Message](#) qualifier, even when the command doesn't require it, to reduce the testing matrix.

This cluster SHALL be present on the [root node endpoint](#) when required by a device type, MAY be present on that endpoint otherwise, and SHALL NOT be present on any other Endpoint of any Node.

14.4.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial Release

14.4.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	TLSCERT

14.4.3. Cluster ID

ID	Name
0x0801	TLS Certificate Management

14.4.4. Data Types

14.4.4.1. TLSCAID Type

This data type is derived from uint16 and represents a provisioned certificate authority certificate with valid values from 0 to 65534. This value SHOULD start at 0 and monotonically increase by 1 with each new allocation provisioned by the Node. A value incremented past 65534 SHOULD wrap to 0. The Node SHALL verify that a new value does not match any other value for this type. If such a match is found, the value SHALL be changed until a unique value is found.

14.4.4.2. TLSCCDID Type

This data type is derived from uint16 and represents a provisioned client certificate with valid values from 0 to 65534. This value SHOULD start at 0 and monotonically increase by 1 with each new allocation provisioned by the Node. A value incremented past 65534 SHOULD wrap to 0. The Node SHALL verify that a new value does not match any other value for this type. If such a match is found, the value SHALL be changed until a unique value is found.

14.4.4.3. TLSCertStruct Type

This encodes the mapping between a [TLSCAID](#) and the associated root certificate.

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	CAID	TLSCAID	0 to 65534				M
1	Certificate	octstr	max 3000				O

CAID Field

This field SHALL be a [TLSCAID](#) representing the unique Certificate Authority ID.

Certificate Field

This field SHALL be an octet string that represents a certificate encoded using DER encoding.

When this field exists and is read over a [Large Message](#) capable transport, it SHALL be included. When this field exists and is read over a non [Large Message](#) capable transport, it SHALL NOT be included. To get the full details of a certificate use the [FindRootCertificate](#) command.

14.4.4.4. TLSClientCertificateDetailStruct Type

This encodes a TLS Client Certificate and corresponding ICAC chain.

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	CCDID	TLSCCDID	0 to 65534				M

Access Modifier: Fabric Scoped							
1	ClientCertificate	octstr	max 3000	X			O
2	IntermediateCertificates	list[octstr]	max 10 [max 3000]				O

CCDID Field

This field SHALL be a [TLSCCID](#) representing the unique Client Certificate ID.

ClientCertificate Field

This field SHALL be an octet string that represents a TLS Client Certificate encoded using DER encoding.

When this field exists and is read over a [Large Message](#) capable transport, it SHALL be included. When this field exists, is non-NULL, and is read over a non [Large Message](#) capable transport, it SHALL NOT be included. To get the full details of a certificate use the [FindClientCertificate](#) command.

A NULL value indicates that the [TLS Client Certificate Signing Request \(CSR\) Procedure](#) has not yet completed.

IntermediateCertificates Field

This field SHALL be a list of octet strings representing one or more ICACs (also encoded using DER) that form a Certificate Chain up to, but not including, the [TLSRCAC](#).

When this field exists and is read over a [Large Message](#) capable transport, it SHALL be included. When this field exists, is non-empty, and is read over a non [Large Message](#) capable transport, it SHALL NOT be included. To get the full details of a certificate use the [FindClientCertificate](#) command.

An empty value means that no intermediate certificates are needed for the TLS Server to validate the [ClientCertificate](#).

14.4.5. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	Max-RootCertificates	uint8	5 to 254	F		R V	M
0x0001	Provisioned-RootCertificates	list[TLSCertStruct]	max Max-RootCertificates	N		R V S	M

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0002	Max-ClientCertificates	uint8	2 to 254	F		R V	M
0x0003	Provisioned-ClientCertificates	list[TLSClientCertificateDetailStruct]	max Max-ClientCertificates	N		R V S	M

14.4.5.1. MaxRootCertificates Attribute

This attribute SHALL contain the maximum number of per fabric [TLSRCACs](#) that can be installed on this Node.

14.4.5.2. ProvisionedRootCertificates Attribute

This attribute SHALL be a list of all provisioned [TLSCertStruct](#) that are currently installed on this Node. When this attribute is read over a non [Large Message](#) capable transport, the Certificate field SHALL NOT be included. To get the full details of a certificate use the [FindRootCertificate](#) command.

14.4.5.3. MaxClientCertificates Attribute

This attribute SHALL contain the maximum number of per fabric [Client Certificates](#) that can be installed on this Node.

14.4.5.4. ProvisionedClientCertificates Attribute

This attribute SHALL be a list of all provisioned [TLSCCDID](#) that are currently installed on this Node. When this attribute is read over a non [Large Message](#) capable transport, the ClientCertificate and IntermediateCertificates fields SHALL NOT be included. To get the full details of a client certificate use the [FindClientCertificate](#) command.

14.4.6. Commands

ID	Name	Direction	Response	Access	Quality	Conformance
0x00	Provision-RootCertificate	client ⇒ server	Provision-RootCertificateResponse	A F	L	M
0x01	Provision-RootCertificateResponse	client ⇐ server	N		L	M

ID	Name	Direction	Response	Access	Quality	Conformance
0x02	Find-RootCertificate	client ⇒ server	FindRootCertificateResponse	O F	L	M
0x03	Find-RootCertificateResponse	client ⇐ server	N		L	M
0x04	Lookup-RootCertificate	client ⇒ server	Lookup-RootCertificateResponse	O F	L	M
0x05	Lookup-RootCertificateResponse	client ⇐ server	N		L	M
0x06	Remove-RootCertificate	client ⇒ server	Y	A F	L	M
0x07	ClientCSR	client ⇒ server	ClientCSRResponse	A F	L	M
0x08	ClientCSR-Response	client ⇐ server	N		L	M
0x09	Provision-ClientCertificate	client ⇒ server	Y	A F	L	M
0x0A	Find-ClientCertificate	client ⇒ server	Find-ClientCertificateResponse	O F	L	M
0x0B	Find-ClientCertificateResponse	client ⇐ server	N		L	M
0x0C	Lookup-ClientCertificate	client ⇒ server	Lookup-ClientCertificateResponse	O F	L	M
0x0D	Lookup-ClientCertificateResponse	client ⇐ server	N		L	M

ID	Name	Direction	Response	Access	Quality	Conformance
0x0E	Remove-ClientCertificate	client ⇒ server	Y	A F	L	M

14.4.6.1. ProvisionRootCertificate Command

This command SHALL provision a newly provided certificate, or rotate an existing one, based on the contents of the CAID field.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Certificate	octstr	max 3000			M
1	CAID	TLSCAID	all	X		M

Certificate Field

This field SHALL be an octet string that represents a certificate encoded using DER encoding.

CAID Field

This field SHALL be a [TLSCAID](#) representing the unique Certificate Authority ID. A null requests a new certificate to be added, and a non-null allows for updating / rotating an existing certificate.

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the [UTCTime](#) attribute of the [Time Synchronization](#) cluster is null:
 - Fail the command with the status code INVALID_IN_STATE, and end processing with no other side effects.
- If the passed in *Certificate* is an invalid [TLS Certificate](#):
 - Fail the command with the status code DYNAMIC_CONSTRAINT_ERROR, and end processing with no other side effects.
- If any existing entry for Certificate is found in [ProvisionedRootCertificates](#) which has both a matching Fingerprint and an [associated fabric](#) which matches the [accessing fabric](#):
 - Fail the command with the status code ALREADY_EXISTS, and end processing with no other side effects.
- If the passed in *CAID* is null:
 - If the count of entries in the [ProvisionedRootCertificates](#) list where the [associated fabric](#) matches the [accessing fabric](#), is equal to the [MaxRootCertificates](#) value:
 - Fail the command with the status code RESOURCE_EXHAUSTED, and end processing with no other side effects.

- Generate a new [TLSCAID](#)
- Create and populate a [TLSCertStruct](#) with the generated TLSCAID and the passed in *Certificate* field, associated with the [accessing fabric](#)
- Add the resulting [TLSCertStruct](#) to the [ProvisionedRootCertificates](#) list.
- Else if the passed in *CAID* is not null:
 - If there is no matching entry found for the passed in *CAID* in the [ProvisionedRootCertificates](#) list:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
 - If the [associated fabric](#) of that entry does not equal the [accessing fabric](#):
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
 - Update the [Certificate Field](#) field of that entry with the passed in *Certificate* field.
- Return the TLSCAID as the [CAID](#) field in the corresponding [ProvisionRootCertificateResponse](#) command.

Note when using this command for certificate rotation, the updated certificate will only be used for new underlying TLS connections established after this call.

14.4.6.2. ProvisionRootCertificateResponse Command

This command SHALL be generated in response to a ProvisionRootCertificate command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CAID	TLSCAID	0 to 65534			M

CAID Field

This field SHALL be a [TLSCAID](#) representing the unique Certificate Authority ID.

14.4.6.3. FindRootCertificate Command

This command SHALL return the specified TLS root certificate, or all provisioned TLS root certificates for the [accessing fabric](#), based on the contents of the CAID field.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CAID	TLSCAID	all	X		M

CAID Field

This field SHALL be a [TLSCAID](#) representing the unique Certificate Authority ID to return, or null to return all provisioned root certificates.

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the [ProvisionedRootCertificates](#) list is empty:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the passed in *CAID* is null:
 - Create an empty list of [TLSCertStruct](#).
 - For each entry in [ProvisionedRootCertificates](#):
 - If the [associated fabric](#) of the entry matches the [accessing fabric](#):
 - Add a populated [TLSCertStruct](#) entry for the CAID to the resulting list.
 - If the resulting list has no entries:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- Else if the passed in *CAID* is not null:
 - If there is no entry in the [ProvisionedRootCertificates](#) list that has a [CAID Field](#) matching the passed in *CAID*:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
 - If the [associated fabric](#) of the [TLSCertStruct](#) for that entry does not equal the [accessing fabric](#):
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
 - Create a list of one [TLSCertStruct](#) and populate with the values from that entry.
- Return the resulting list in the corresponding [FindRootCertificateResponse](#) command.

14.4.6.4. FindRootCertificateResponse Command

This command SHALL be generated in response to a FindRootCertificate command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CertificateDetails	list[TLSCertStruct]	1 to Max-RootCertificates			M

CertificateDetails Field

This field SHALL be a list of [TLSCertStructs](#) containing a minimum of one [TLSCertStruct](#).

14.4.6.5. LookupRootCertificate Command

This command SHALL return the [CAID](#) for the passed in fingerprint.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Fingerprint	octstr	max 64			M

Fingerprint Field

This field SHALL be an octet string that represents the [certificate fingerprint](#).

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the [ProvisionedRootCertificates](#) list is empty:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If there is no entry in the [ProvisionedRootCertificates](#) list that has a matching [Fingerprint](#), or the [associated fabric](#) of that entry does not equal the [accessing fabric](#):
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- Return the [CAID](#) of that entry, in the corresponding [LookupRootCertificateResponse](#) command.

14.4.6.6. LookupRootCertificateResponse Command

This command SHALL be generated in response to a LookupRootCertificate command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CAID	TLSCAID	0 to 65534			M

CAID Field

This field SHALL be a [TLSCAID](#) representing the unique Certificate Authority ID.

14.4.6.7. RemoveRootCertificate Command

This command SHALL be generated to request the server removes the certificate provisioned to the provided Certificate Authority ID.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CAID	TLSCAID	0 to 65534			M

CAID Field

This field SHALL be a [TLSCAID](#) representing the unique Certificate Authority ID.

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the [ProvisionedRootCertificates](#) list is empty:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If there is no entry in the [ProvisionedRootCertificates](#) list that has a [CAID Field](#) matching the passed in *CAID*:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the [associated fabric](#) of the [TLSCertStruct](#) for that entry does not equal the [accessing fabric](#):
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the passed in *CAID* equals the [CAID](#) of any entry in the [ProvisionedEndpoints](#) list in the [TLS Client Management Cluster](#):
 - Fail the command with the status code INVALID_IN_STATE, and end processing with no other side effects.
- Remove the entry for the passed in *CAID* from the [ProvisionedRootCertificates](#) list.

14.4.6.8. ClientCSR Command

This command SHALL be generated to request the Node generates a certificate signing request for a new TLS key pair or use an existing CCDID for certificate rotation.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Nonce	octstr	32			M
1	CCDID	TLSCCDID	0 to 65534	X		M

Nonce Field

This field SHALL be an octet string that represents the nonce to be signed by the private key used in the CSR, with the resulting signature returned in the [NonceSignature](#) field of [ClientCSRResponse](#).

CCDID Field

This field SHALL be a [TLSCCDID](#) representing the unique Client Certificate Details ID. If NULL, a new key pair and CCDID will be generated. If non-NULL, the existing key-pair for the CCDID will be used.

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the passed in *CCDID* is NULL:
 - If the count of entries in the [ProvisionedClientCertificates](#) list where the [associated fabric](#) matches the [accessing fabric](#), is equal to the [MaxClientCertificates](#) value:
 - Fail the command with the status code `RESOURCE_EXHAUSTED`, and end processing with no other side effects.
 - Generate a new key pair using [Crypto_GenerateKeypair](#).
 - If a key collision is detected against any other TLS key pair or Operational credential key pair:
 - Discard the new key pair.
 - Fail the command with the status code `DYNAMIC_CONSTRAINT_ERROR`, and end processing with no other side effects.
 - Generate a new [TLSCCDID](#) value.
 - Create a new [TLSClientCertificateDetailStruct](#) associated with the [accessing fabric](#).
 - Set the [CCDID](#) field with the newly created [TLSCCDID](#) value, and associate the key pair with it.
 - Set the [ClientCertificate](#) and [IntermediateCertificates](#) fields to NULL.
 - Add the [TLSClientCertificateDetailStruct](#) to the [ProvisionedClientCertificates](#) list.
- Else if the passed in *CCDID* is not NULL:
 - If there is no entry in the [ProvisionedClientCertificates](#) list that has a matching [CCDID](#) to the passed in *CCDID*:
 - Fail the command with the status code `NOT_FOUND`, and end processing with no other side effects.
 - If the [associated fabric](#) of that entry does not equal the [accessing fabric](#):
 - Fail the command with the status code `NOT_FOUND`, and end processing with no other side effects.
- Generate a `tls_csr` using the TLS key pair by following the format and procedure in [PKCS #10](#), which includes a signature using the private key (see [RFC 2986](#) section 4.2) associated with the public key which is the `subjectPublicKey` field of the CSR. The CSR subject MAY be any value and the device SHOULD NOT expect the final certificate to contain any of the CSR subject DN attributes.
- Compute an [ec-signature](#) using [Crypto_Sign\(\)](#) of the passed in *Nonce*, and encode the result as an octet string into `tls_nonce_signature`.

```
tls_nonce_signature = Crypto_Sign(  
    message = Nonce,  
    privateKey = TLS Private Key
```

)

- Return the CCDID as [CCDID](#), the DER-encoded `tls_csr` as [CSR](#), and `tls_nonce_signature` as [NonceSignature](#), in the corresponding [ClientCSRResponse](#) command.

14.4.6.9. ClientCSRResponse Command

This command SHALL be generated in response to a [ClientCSR](#) command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CCDID	TLSCCDID	0 to 65534			M
1	CSR	octstr	max 3000			M
2	NonceSignature	octstr	max 128			M

CCDID Field

This field SHALL be a [TLSCCDID](#) representing the unique Client Certificate Details ID.

CSR Field

This field SHALL be a DER-encoded octet string of a [PKCS #10](#) format Certificate Signing Request.

NonceSignature Field

This field SHALL be an octet string of the [ec-signature](#) of the [Nonce](#) field from the corresponding [ClientCSR](#) command.

14.4.6.10. ProvisionClientCertificate Command

This command SHALL be generated to request the Node provisions newly provided Client Certificate Details, or rotate an existing client certificate.

This command is typically invoked after having created a new client certificate using the CSR requested in [ClientCSR](#), with the [TLSCCDID](#) returned by [ClientCSRResponse](#).

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CCDID	TLSCCDID	0 to 65534			M
1	ClientCertificate	octstr	max 3000			M
2	IntermediateCertificates	list[octstr]	0 to 10 [max 3000]			M

CCDID Field

This field SHALL be a [TLSCCDID](#) representing the unique Client Certificate Details ID.

ClientCertificate Field

This field SHALL be an octet string that represents a TLS Client Certificate encoded using DER encoding.

IntermediateCertificates Field

This field SHALL be a list of octet strings representing one or more ICACs (also encoded using DER) that form a Certificate Chain up to, but not including, the [TLSRCAC](#). An empty value means no intermediate certificates are needed.

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the [ProvisionedClientCertificates](#) list is empty:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If there is an existing entry for the passed in *ClientCertificate* in the [ProvisionedClientCertificates](#) list, which has both a matching [Fingerprint](#) and an [associated fabric](#) that equals the [accessing fabric](#):
 - Fail the command with the status code ALREADY_EXISTS, and end processing with no other side effects.
- If there is no entry in the [ProvisionedClientCertificates](#) list that has a [CCDID](#) matching the passed in *CCDID*:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the [associated fabric](#) for that entry does not equal the [accessing fabric](#):
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If there is any invalid [TLS Certificate](#) in the passed in *ClientCertificate* or *IntermediateCertificates*:
 - Fail the command with the status code DYNAMIC_CONSTRAINT_ERROR, and end processing with no other side effects.
- If the public key of the passed in *ClientCertificate* does not correspond to the private key of the matching entry:
 - Fail the command with the status code DYNAMIC_CONSTRAINT_ERROR, and end processing with no other side effects.
- Update the [ClientCertificate](#) and [IntermediateCertificates](#) fields of that entry to the passed in *ClientCertificate* and *IntermediateCertificates*.
- Return SUCCESS.

Note: When using this command for client certificate rotation, only new underlying TLS connections (established after this finishes processing), will use the updated Certificate.

14.4.6.11. FindClientCertificate Command

This command SHALL return the [TLSClientCertificateDetailStruct](#) for the passed in CCDID, or all TLS client certificates for the [accessing fabric](#), based on the contents of the CCDID field.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CCDID	TLSCCDID	all	X		M

CCDID Field

This field SHALL be a [TLSCCDID](#) representing the unique Client Certificate Details ID.

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the [ProvisionedClientCertificates](#) list is empty:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the passed in *CCDID* is null:
 - Create a list of [TLSClientCertificateDetailStruct](#)
 - For each entry in [ProvisionedClientCertificates](#):
 - If the entry’s [associated fabric](#) matches the [accessing fabric](#):
 - Add a populated [TLSClientCertificateDetailStruct](#) entry for the passed in *CCDID* to the resulting list.
 - If the resulting list has no entries:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- Else if the passed in *CCDID* is not null:
 - If there is no entry in the [ProvisionedClientCertificates](#) list that has a [CCDID](#) matching the passed in *CCDID*:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
 - If the [associated fabric](#) of that entry does not equal the [accessing fabric](#):
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
 - Create a list of one [TLSClientCertificateDetailStruct](#) and populate with the values from that entry for the requested *CCDID*.
- Return the list as the [CertificateDetails](#) field, in the corresponding [FindClientCertificateResponse](#)

command.

Note: If an entry in the returned list has an empty [ClientCertificate](#) field, it means the [ClientCSR](#) command was invoked, but the corresponding [ProvisionClientCertificate](#) has not been invoked yet.

14.4.6.12. FindClientCertificateResponse Command

This command SHALL be generated in response to a FindClientCertificate command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CertificateDetails	list[TLSClientCertificateDetailStruct]	1 to Max-ClientCertificates			M

CertificateDetails Field

This field SHALL be a list of [TLSClientCertificateDetailStruct](#) containing a minimum of one entry.

14.4.6.13. LookupClientCertificate Command

This command SHALL return the CCDID for the passed in Fingerprint.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Fingerprint	octstr	max 64			M

Fingerprint Field

This field SHALL be an octet string that represents the [certificate fingerprint](#).

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the [ProvisionedClientCertificates](#) list is empty:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If there is no entry in the [ProvisionedClientCertificates](#) list that has a [Fingerprint](#) matching the passed in *Fingerprint*:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the [associated fabric](#) of that entry does not equal the [accessing fabric](#):
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- Return the [CCDID](#) field of the matching entry, as the [CCDID](#) field in the corresponding [Lookup-](#)

[ClientCertificateResponse](#) command.

14.4.6.14. LookupClientCertificateResponse Command

This command SHALL be generated in response to a LookupClientCertificate command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CCDID	TLSCCDID	0 to 65534			M

CCDID Field

This field SHALL be a [TLSCCDID](#) representing the unique Client Certificate Details ID.

14.4.6.15. RemoveClientCertificate Command

This command SHALL be used to request the Node removes all stored information for the provided CCDID.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	CCDID	TLSCCDID	0 to 65534			M

CCDID Field

This field SHALL be a [TLSCCDID](#) representing the unique Client Certificate Details ID.

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the [ProvisionedClientCertificates](#) list is empty:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If there is no entry in the [ProvisionedClientCertificates](#) list that has a [CCDID](#) matching the passed in *CCDID*:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the [associated fabric](#) of that entry does not equal the [accessing fabric](#):
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the passed in *CCDID* equals the [CCDID](#) of any entry in the [ProvisionedEndpoints](#) list in the [TLS Client Management Cluster](#):
 - Fail the command with the status code INVALID_IN_STATE, and end processing with no other side effects.
- Remove the entry for the passed in *CCDID* from the [ProvisionedClientCertificates](#) list.

- Remove the TLS Key Pair belonging to the passed in *CCDID*.

14.5. TLS Client Management Cluster

This Cluster is used to provision TLS Endpoints with enough information to facilitate subsequent connection.

Commands in this cluster uniformly use the [Large Message](#) qualifier, even when the command doesn't require it, to reduce the testing matrix.

This cluster SHALL be present on the [root node endpoint](#) when required by a device type, MAY be present on that endpoint otherwise, and SHALL NOT be present on any other Endpoint of any Node.

14.5.1. Revision History

The global ClusterRevision attribute value SHALL be the highest revision number in the table below.

Revision	Description
1	Initial Release

14.5.2. Classification

Hierarchy	Role	Scope	PICS Code
Base	Utility	Node	TLSCIENT

14.5.3. Cluster ID

ID	Name
0x0802	TLS Client Management

14.5.4. Data Types

14.5.4.1. TLSEndpointID Type

This data type is derived from uint16 and represents a provisioned endpoint with valid values from 0 to 65534. This value SHOULD start at 0 and monotonically increase by 1 with each new allocation provisioned by the Node. A value incremented past 65534 SHOULD wrap to 0. The Node SHALL verify that a new value does not match any other value for this type. If such a match is found, the value SHALL be changed until a unique value is found.

14.5.4.2. TLSEndpointStruct Type

This struct encodes details about a TLS Endpoint.

Access Modifier: Fabric Scoped							
ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0	End-pointID	TLSEnd-pointID	0 to 65534				M
1	Hostname	octstr	4 to 253				M
2	Port	uint16	1 to 65535				M
3	CAID	TLSCAID	0 to 65534				M
4	CCDID	TLSCCDID	all	X			M
5	Reference-Count	uint8	all				M

EndpointID Field

This field SHALL represent the unique TLS Endpoint ID.

Hostname Field

This field SHALL represent a TLS Hostname.

Port Field

This field SHALL represent a TLS Port Number.

CAID Field

This field SHALL be a [TLSCAID](#) representing the associated Certificate Authority ID.

CCDID Field

This field SHALL be a [TLSCCDID](#) representing the associated Client Certificate Details ID. A NULL value means no client certificate is used with this endpoint.

ReferenceCount Field

This field SHALL indicate a reference count of the number of entities currently using this TLS Endpoint. The node SHALL recompute this field to reflect the correct value at runtime (e.g., when restored from a persisted value after a reboot).

14.5.5. Status Codes

14.5.5.1. StatusCodeEnum Type

This data type is derived from [enum8](#).

Value	Name	Summary
0x02	EndpointAlreadyInstalled	The endpoint is already installed.
0x03	RootCertificateNotFound	No root certificate exists for this CAID.
0x04	ClientCertificateNotFound	No client certificate exists for this CCDID.
0x05	EndpointInUse	The endpoint is in use and cannot be removed.
0x06	InvalidTime	Time sync has not yet occurred.

14.5.6. Attributes

ID	Name	Type	Constraint	Quality	Fallback	Access	Conformance
0x0000	MaxProvisioned	uint8	5 to 254	F		R V	M
0x0001	ProvisionedEndpoints	list[TLSEndpointStruct]	desc	N		R V S	M

14.5.6.1. MaxProvisioned Attribute

This attribute SHALL indicate the maximum number of per fabric [TLSEndpoints](#) that can be installed on this Node.

14.5.6.2. ProvisionedEndpoints Attribute

This attribute SHALL indicate a list of currently provisioned TLS Endpoints on this Node. The maximum length of this list when read will be the value of [MaxProvisioned](#).

14.5.7. Commands

ID	Name	Direction	Response	Access	Quality	Conformance
0x00	ProvisionEndpoint	client ⇒ server	ProvisionEndpointResponse	A F	L	M
0x01	ProvisionEndpointResponse	client ⇐ server	N		L	M

ID	Name	Direction	Response	Access	Quality	Conformance
0x02	FindEndpoint	client ⇒ server	FindEndpointResponse	O F	L	M
0x03	FindEndpointResponse	client ← server	N		L	M
0x04	RemoveEndpoint	client ⇒ server	N	A F	L	M

14.5.7.1. ProvisionEndpoint Command

This command is used to provision a TLS Endpoint for the provided Hostname / Port combination.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0x00	Hostname	octstr	4 to 253			M
0x01	Port	uint16	1 to 65535			M
0x02	CAID	TLSCAID	0 to 65534			M
0x03	CCDID	TLSCCDID	all	X		M
0x04	EndpointID	TLSEndpointID	all	X		M

Hostname Field

This field SHALL represent a TLS Hostname.

Port Field

This field SHALL represent a TLS Port Number.

CAID Field

This field SHALL be the [TLSCAID](#) used to associate the [TLSRCAC](#) with this endpoint.

CCDID Field

This field SHALL be the [TLSCCDID](#) used to associate the Client Certificate Details with this endpoint. A NULL value means no client certificate is associated.

EndpointID Field

This field SHALL represent the unique TLS Endpoint. A NULL value causes a new endpoint to be created and a non-NULL value allows for updating an existing endpoint.

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the **UTCTime** attribute of the **Time Synchronization** cluster is null:
 - Fail the command with the cluster-specific status code of **InvalidTime**, and end processing with no other side effects.
- If there is no entry for the passed in **CAID** in the **ProvisionedRootCertificates** list:
 - Fail the command with the cluster-specific status code of **RootCertificateNotFound**, and end processing with no other side effects.
- If the **associated fabric** for that matching entry, does not equal the **accessing fabric**:
 - Fail the command with the cluster-specific status code of **RootCertificateNotFound**, and end processing with no other side effects.
- If the passed in **CCDID** is not NULL, and the **associated fabric** for the passed in **CCDID** entry in the **ProvisionedClientCertificates** list does not equal the **accessing fabric**:
 - Fail the command with the cluster-specific status code of **ClientCertificateNotFound**, and end processing with no other side effects.
- If the passed in **EndpointID** is NULL:
 - If the length of **ProvisionedEndpoints** is equal to the **MaxProvisioned** value:
 - Fail the command with the status code **RESOURCE_EXHAUSTED**, and end processing with no other side effects.
 - If there is an existing entry for the **Hostname** / **Port** combination in the **ProvisionedEndpoints** list, where the **associated fabric** equals the **accessing fabric**:
 - Fail the command with the cluster-specific status code of **EndpointAlreadyInstalled**, and end processing with no other side effects.
 - Generate a new **TLSEndpointID**
 - Create and populate a **TLSEndpointStruct** with the passed in values, associated with the **accessing fabric**
 - Set the **EndpointID** field to the newly generated **TLSEndpointID**.
 - Set the **ReferenceCount** field to 0.
 - Add the resulting **TLSEndpointStruct** to the **ProvisionedEndpoints** list and store the results.
- Else if the passed in **EndpointID** is not NULL:
 - If there is no entry found for the passed in **EndpointID** in the **ProvisionedEndpoints** list:
 - Fail the command with the status code **NOT_FOUND**, and end processing with no other side effects.
 - If the **associated fabric** for that matching entry does not equal the **accessing fabric**:
 - Fail the command with the status code **NOT_FOUND**, and end processing with no other side effects.
 - If another entry exists for the passed in **Hostname** / **Port** combination in the **ProvisionedEndpoints** list, where the **associated fabric** equals the **accessing fabric**:

- Fail the command with the cluster-specific status code of EndpointAlreadyInstalled, and end processing with no other side effects.
- Update the fields of that matching entry with the passed in values and store the results.
- Return the TLSEndpointID as the EndpointID field in the corresponding ProvisionEndpointResponse command.

General Notes When the Node is making a TLS connection to this TLS Endpoint, the TLSRCAC represented by the CAID SHALL be used to authenticate the TLS Endpoint.

When the Node is making a TLS connection that has a non-NULL CCDID, the Client Certificate Details represented by the CCDID SHALL be used during client authentication in the TLS Handshake. In addition, a Node SHALL fail to connect to the TLS server, if that TLS Server did not require TLS Client Authentication for the connection, when a CCDID is provisioned on a TLS Endpoint.

14.5.7.2. ProvisionEndpointResponse Command

This command is used to report the result of the ProvisionEndpoint command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	EndpointID	TLSEndpointID	0 to 65534			M

EndpointID Field

This field SHALL be the TLS Endpoint ID created or updated by the ProvisionEndpoint command.

14.5.7.3. FindEndpoint Command

This command is used to find a TLS Endpoint by its ID.

This command SHALL return the TLSEndpointStruct for the passed in EndpointID.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	EndpointID	TLSEndpointID	0 to 65534			M

EndpointID Field

This field SHALL be the TLS Endpoint ID being looked up.

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the ProvisionedEndpoints list is empty:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side

effects.

- If there is no matching entry found for the passed in *EndpointID* in the [ProvisionedEndpoints](#) list:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the [associated fabric](#) for that matching entry, does not equal the [accessing fabric](#):
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- Return that entry as the [Endpoint](#) field in the corresponding [FindEndpointResponse](#) command.

14.5.7.4. FindEndpointResponse Command

This command is used to report the result of the FindEndpoint command.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	Endpoint	TLSEndpointStruct	0 to 65534			M

Endpoint Field

The field SHALL be a [TLSEndpointStruct](#) containing the requested entry.

14.5.7.5. RemoveEndpoint Command

This command is used to remove a TLS Endpoint by its ID.

This command SHALL be generated to request the Node remove any TLS Endpoint.

ID	Name	Type	Constraint	Quality	Fallback	Conformance
0	EndpointID	TLSEndpointID	0 to 65534			M

EndpointID Field

This field SHALL represent the unique [TLSEndpointID](#) of the TLS Endpoint to remove.

Effect on Receipt

The following process SHALL be followed when the server receives this command:

- If the [ProvisionedEndpoints](#) list is empty:
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If there is no matching entry found for the passed in *EndpointID* in the [ProvisionedEndpoints](#)

list:

- Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the [associated fabric](#) for that matching entry, does not equal the [accessing fabric](#):
 - Fail the command with the status code NOT_FOUND, and end processing with no other side effects.
- If the [ReferenceCount](#) of that matching entry is greater than 0:
 - Fail the command with the status code INVALID_IN_STATE, and end processing with no other side effects.
- Remove that entry from the [ProvisionedEndpoints](#) list and store the results.

Appendix A: Tag-length-value (TLV) Encoding Format

A.1. Scope & Purpose

The Matter TLV (*Tag-Length-Value*) format is a generalized encoding method for simple structured data, used throughout this specification.

Values in the Matter TLV format are encoded as *TLV elements*. Each TLV element has a type. Element types fall into two categories: *primitive types* and *container types*. Primitive types convey fundamental data values such as integers and strings. Container types convey collections of elements that themselves are either primitives or containers. The Matter TLV format supports three different container types: structures, arrays and lists.

All valid *TLV encodings* consist of a single top-level element. This value can be either a primitive type or a container type.

A.2. Tags

A TLV element includes an optional numeric tag that identifies its purpose. A TLV element without a tag is called an *anonymous* element. For elements with tags, two categories of tags are defined: *profile-specific* and *context-specific*.

A.2.1. Profile-Specific Tags

Profile-specific tags identify elements globally. A profile-specific tag is a 64-bit number composed of the following fields:

- 16-bit [Vendor ID](#)
- 16-bit profile number
- 32-bit tag number

Profile-specific tags are defined either by Matter or by vendors. Additionally the Matter Common Profile includes a set of predefined profile-specific tags that can be used across organizations.

A.2.2. Context-Specific Tags

Context-specific tags identify elements within the context of a containing structure element. A context-specific tag consists of a single 8-bit tag number. The meaning of a context-specific tag derives from the structure it resides in, implying that the same tag number may have different meanings in the context of different structures. Effectively, the interpretation of a context-specific tag depends on the tag attached to the containing element. Because structures themselves can be assigned context-specific tags, the interpretation of a context-specific tag may ultimately depend on a nested chain of such tags.

Context-specific tags can only be assigned to elements that are immediately within a structure. This

implies that an element with a context-specific tag cannot appear as the outermost element of a TLV encoding.

A.2.3. Anonymous Tags

A special "anonymous tag" is used to denote TLV elements that lack a tag value. Such a TLV element is referred to as an anonymous element.

A.2.4. Canonical Ordering of Tags

Where a distinguished ordering of tags is required (e.g. for the purposes of generating a hash or cryptographic signature of elements within a structure), the following ordering rules SHALL be used:

- Anonymous tags SHALL be ordered before all other tags.
- Context-specific tags SHALL be ordered before profile-specific tags.
- Context-specific tags with numerically lower tag values SHALL be ordered before those with higher tag values.
- Profile-specific tags with numerically lower Vendor IDs SHALL be ordered before those with higher Vendor IDs.
- Profile-specific tags with the same Vendor ID, but numerically lower profile numbers SHALL be ordered before those with higher profile numbers.
- Profile-specific tags with the same Vendor ID and the same profile numbers but numerically lower tag numbers SHALL be ordered before those with higher tag numbers.

The ordering rules SHALL apply to elements at the same level within a container.

A.3. Lengths

Depending on its type, a TLV element may contain a length field that gives the length, in octets, of the element's value field. A length field is only present for string types (character and octet strings). Other element types either have a predetermined length or are encoded with a marker that identifies their end.

A.4. Primitive Types

The Matter TLV format supports the following primitive types:

- Signed integers
- Unsigned integers
- UTF-8 Strings
- Octet Strings
- Single or double-precision floating point numbers (following [IEEE 754-2019](#))
- Booleans

- Nulls

Of the primitive types, integers, floating point numbers, booleans and nulls have a predetermined length specified by their type. Octet strings and UTF-8 strings include a length field that gives their lengths in octets.

A.5. Container Types

The Matter TLV format supports the following container types:

- Structures
- Arrays
- Lists

Each of the container types is a form of element collection that can contain primitive types and/or other container types. The elements appearing immediately within a container type are called its *members*. A container type can contain any number of member elements, including none. Container types can be nested to any depth and in any combination. The end of a container type is denoted by a special element called the ‘end-of-container’ element. Although encoded as a member, conceptually the end-of-container element is not included in the members of the containing type.

A.5.1. Structures

A structure is a collection of member elements that each have a distinct meaning. All member elements within a structure SHALL have a unique tag as compared to the other members of the structure. Member elements without tags (anonymous elements) are not allowed in structures. The encoded ordering of members in a structure may or may not be important depending on the intent of the sender or the expectations of the receiver. For example, in some situations, senders and receivers may agree on a particular ordering of elements to make encoding and decoding easier.

Where a distinguished ordering of members is required (for example, for the purposes of generating a hash or cryptographic signature of the structure), the members of the structure SHALL be encoded as specified in [Section A.2.4, “Canonical Ordering of Tags”](#).

A.5.2. Arrays

An array is an ordered collection of member elements that either do not have distinct meanings, or whose meanings are implied by their encoded positions in the array. An array can contain any type of element, including other arrays. All member elements of an array SHALL be anonymous elements – that is, they SHALL be encoded with an anonymous tag.

A.5.3. Lists

A list is an ordered collection of member elements, each of which may be encoded with a tag. The meanings of member elements in a list are denoted by their position within the list in conjunction with any associated tag value they may have.

A list can contain any type of element, including other lists. The members of a list may be encoded

with any form of tag, including an anonymous tag. The tags within a list do not need to be unique with respect to other members of the list.

A.6. Element Encoding

A TLV element is encoded a single control octet, followed by a sequence of tag, length and value octets. Depending on the nature of the element, any of the tag, length or value fields may be omitted.

Control Octet	Tag	Length	Value
1 octet	0 to 8 octets	0 to 8 octets	Variable

A.7. Control Octet Encoding

The control octet specifies the type of a TLV element and how its tag, length and value fields are encoded. The control octet consists of two subfields: an *element type field* which occupies the lower 5 bits, and a *tag control field* which occupies the upper 3 bits.

A.7.1. Element Type Field

The element type field encodes the element's type as well as how the corresponding length and value fields are encoded. In the case of Booleans and the **Null** value, the element type field also encodes the value itself.

Control Octet								Description
Tag Control			Element Type					
7	6	5	4	3	2	1	0	
x	x	x	0	0	0	0	0	Signed Integer, 1-octet value
x	x	x	0	0	0	0	1	Signed Integer, 2-octet value
x	x	x	0	0	0	1	0	Signed Integer, 4-octet value
x	x	x	0	0	0	1	1	Signed Integer, 8-octet value
x	x	x	0	0	1	0	0	Unsigned Integer, 1-octet value
x	x	x	0	0	1	0	1	Unsigned Integer, 2-octet value
x	x	x	0	0	1	1	0	Unsigned Integer, 4-octet value
x	x	x	0	0	1	1	1	Unsigned Integer, 8-octet value
x	x	x	0	1	0	0	0	Boolean False
x	x	x	0	1	0	0	1	Boolean True
x	x	x	0	1	0	1	0	Floating Point Number, 4-octet value
x	x	x	0	1	0	1	1	Floating Point Number, 8-octet value
x	x	x	0	1	1	0	0	UTF-8 String, 1-octet length

Control Octet								Description
x	x	x	0	1	1	0	1	UTF-8 String, 2-octet length
x	x	x	0	1	1	1	0	UTF-8 String, 4-octet length
x	x	x	0	1	1	1	1	UTF-8 String, 8-octet length
x	x	x	1	0	0	0	0	Octet String, 1-octet length
x	x	x	1	0	0	0	1	Octet String, 2-octet length
x	x	x	1	0	0	1	0	Octet String, 4-octet length
x	x	x	1	0	0	1	1	Octet String, 8-octet length
x	x	x	1	0	1	0	0	Null
x	x	x	1	0	1	0	1	Structure
x	x	x	1	0	1	1	0	Array
x	x	x	1	0	1	1	1	List
0	0	0	1	1	0	0	0	End of Container
x	x	x	1	1	0	0	0	Reserved (where xxx are not 000)
x	x	x	1	1	0	0	1	Reserved
x	x	x	1	1	0	1	0	Reserved
x	x	x	1	1	0	1	1	Reserved
x	x	x	1	1	1	0	0	Reserved
x	x	x	1	1	1	0	1	Reserved
x	x	x	1	1	1	1	0	Reserved
x	x	x	1	1	1	1	1	Reserved

For both signed and unsigned integer types the bottom two bits of the element type field signal the width of the corresponding field as follows:

- 00 — 1 octet
- 01 — 2 octets
- 10 — 4 octets
- 11 — 8 octets

For UTF-8 and octet string types the bottom two bits of the element type field signal the width of the length field as follows:

- 00 — 1 octet
- 01 — 2 octets
- 10 — 4 octets
- 11 — 8 octets

For end of container element type the tag control bits are set to 0. Any other combination of the tag control bits for this element type only is reserved. See [Section A.10, “End of Container Encoding”](#).

A.7.2. Tag Control Field

The tag control field identifies the form of tag assigned to the element (including none) as well as the encoding of the tag octets.

Control Octet								Description
Tag Control			Element Type					
7	6	5	4	3	2	1	0	
0	0	0	x	x	x	x	x	Anonymous Tag Form, 0 octets
0	0	1	x	x	x	x	x	Context-specific Tag Form, 1 octet
0	1	0	x	x	x	x	x	Common Profile Tag Form, 2 octets
0	1	1	x	x	x	x	x	Common Profile Tag Form, 4 octets
1	0	0	x	x	x	x	x	Implicit Profile Tag Form, 2 octets
1	0	1	x	x	x	x	x	Implicit Profile Tag Form, 4 octets
1	1	0	x	x	x	x	x	Fully-qualified Tag Form, 6 octets
1	1	1	x	x	x	x	x	Fully-qualified Tag Form, 8 octets

A.8. Tag Encoding

Tags are encoded in 0, 1, 2, 4, 6 or 8 octet widths as specified by the tag control field. Tags consist of up to three numeric fields: a *Vendor ID field*, a *profile number field*, and a *tag number field*. All fields are encoded in little-endian order. The tag fields are ordered as follows:

Vendor ID	Profile Number	Tag Number
0 or 2 octets	0 or 2 octets	1, 2, or 4 octets

A.8.1. Fully-Qualified Tag Form

A profile-specific tag can be encoded in *fully-qualified tag form*, where the encoding includes all three tag components (Vendor ID, profile number and tag number). Two variants of this form are supported, one with a 16-bit tag number and one with a 32-bit tag number. The 16-bit variant SHALL be used with tag numbers < 65536, while the 32-bit variant SHALL be used with tag numbers >= 65536.

Tag Control	Vendor ID Size	Profile Number Size	Tag Number Size	
110b	2 octets	2 octets	2 octets	For tag numbers < 65536

Tag Control	Vendor ID Size	Profile Number Size	Tag Number Size	
111b	2 octets	2 octets	4 octets	For tag numbers >= 65536

A.8.2. Implicit Profile Tag Form

A profile-specific tag can also be encoded in *implicit profile tag form*, where the encoding includes only the tag number, and the Vendor ID and profile number are inferred from the protocol context in which the TLV encoding is communicated. This form also has two variants based on the magnitude of the tag number.

Tag Control	Tag Number Size	
100b	2 octets	For tag numbers < 65536
101b	4 octets	For tag numbers >= 65536

A.8.3. Common Profile Tag Form

A special encoding exists for profile-specific tags that are defined by the Matter Common Profile. These are encoded in the same manner as implicit profile tags except that they are identified as common profile tags, rather than implicit profile tags in the tag control field.

Tag Control	Tag Number Size	
010b	2 octets	For tag numbers < 65536
011b	4 octets	For tag numbers >= 65536

A.8.4. Context-Specific Tag Form

Context-specific tags are encoded as a single octet conveying the tag number.

Tag Control	Tag Number Size	
001b	1 octets	All tag numbers 0 - 255

A.8.5. Anonymous Tag Form

Anonymous elements do not encode any tag octets.

Tag Control	Tag Size	
000b	0 octets	No data encoded

A.9. Length Encoding

Length fields are encoded in 0, 1, 2, 4 or 8 octet widths, as specified by the element type field. Length fields of more than one octet are encoded in little-endian order. The choice of width for the

length field is up to the discretion of the sender, implying that a sender can choose to send more length octets than strictly necessary to encode the value.

A.10. End of Container Encoding

The end of a container type is marked with a special element called the end-of-container element. The end-of-container element is encoded as a single control octet with the value 18h. The tag control bits within the control octet SHALL be set to zero, implying that end-of-container element can never have a tag.

Control Octet
1 octet

A.11. Value Encodings

A.11.1. Integers

An integer element is encoded as follows:

Control Octet	Tag	Value
1 octet	0 to 8 octets	1, 2, 4 or 8 octets

The number of octets in the value field is indicated by the element type field within the control octet. The choice of value octet count is at the sender’s discretion, implying that a sender is free to send more octets than strictly necessary to encode the value. Within the value octets, the integer value is encoded in little-endian format (two’s complement format for signed integers).

A.11.2. UTF-8 and Octet Strings

UTF-8 and octet strings are encoded as follows:

Control Octet	Tag	Length	Value
1 octet	0 to 8 octets	1 to 8 octets	0 to 2 ⁶⁴ -1 octets

The length field of a UTF-8 or octet string encodes the number of octets (not characters) present in the value field. The number of octets in the length field is implied by the type specified in the element type field (within the control octet).

For octet strings, the value can be any arbitrary sequence of octets. For UTF-8 strings, the value octets SHALL encode a valid UTF-8 character (code points) sequence. Senders SHALL NOT include a terminating null character to mark the end of a string.

A.11.3. Booleans

Boolean elements are encoded as follows:

Control Octet	Tag
1 octet	0 to 8 octets

The value of a Boolean element (true or false) is implied by the type indicated in the element type field.

A.11.4. Arrays, Structures and Lists

Array, structure and list elements are encoded as follows:

Control Octet	Tag	Value	End-of-Container
1 octet	0 to 8 octets	<i>Variable</i>	1 octet

The value field of an array/structure/list element is a sequence of encoded TLV elements that constitute the members of the element, followed by an end-of-container element. The end-of-container element SHALL always be present, even in cases where the end of the array/structure/list element could be inferred by other means (e.g. the length of the packet containing the TLV encoding).

A.11.5. Floating Point Numbers

A floating point number is encoded as follows:

Control Octet	Tag	Value
1 octet	0 to 8 octets	4 or 8 octets

The value field of a floating point element contains an [IEEE 754-2019](#) single or double precision floating point number encoded in little-endian format (specifically, the reverse of the order described in External Data Representation, [RFC 4506](#)). The choice of precision is implied by the type specified in the element type field (within the control octet). The sender is free to choose either precision at their discretion.

A.11.6. Nulls

A Null value is encoded as follows:

Control Octet	Tag
1 octet	0 to 8 octets

A.12. TLV Encoding Examples

In order to better ground the TLV concepts, this subsection provides a set of sample encodings. In the tables below, type and values column uses a decimal representation for all number whereas the encoding is represented with hexadecimal numbers.

[Table 127, “Sample encoding of primitive types”](#) shows sample encodings for primitive types. All examples in the table below are encoded as anonymous elements.

Table 127. Sample encoding of primitive types

Type and Value	Encoding (hex)
Boolean false	08
Boolean true	09
Signed Integer, 1-octet, value 42	00 2a
Signed Integer, 1-octet, value -17	00 ef
Unsigned Integer, 1-octet, value 42U	04 2a
Signed Integer, 2-octet, value 42	01 2a 00
Signed Integer, 4-octet, value -170000	02 f0 67 fd ff
Signed Integer, 8-octet, value 40000000000	03 00 90 2f 50 09 00 00 00
UTF-8 String, 1-octet length, "Hello!"	0c 06 48 65 6c 6c 6f 21
UTF-8 String, 1-octet length, "Tschüs"	0c 07 54 73 63 68 c3 bc 73
Octet String, 1-octet length, octets 00 01 02 03 04	10 05 00 01 02 03 04
Null	14
Single precision floating point 0.0	0a 00 00 00 00
Single precision floating point (1.0 / 3.0)	0a ab aa aa 3e
Single precision floating point 17.9	0a 33 33 8f 41
Single precision floating point infinity (∞)	0a 00 00 80 7f
Single precision floating point negative infinity ($-\infty$)	0a 00 00 80 ff
Double precision floating point 0.0	0b 00 00 00 00 00 00 00 00
Double precision floating point (1.0 / 3.0)	0b 55 55 55 55 55 55 d5 3f
Double precision floating point 17.9	0b 66 66 66 66 66 e6 31 40
Double precision floating point infinity (∞)	0b 00 00 00 00 00 00 f0 7f
Double precision floating point negative infinity ($-\infty$)	0b 00 00 00 00 00 00 f0 ff

Table 128, “Sample encoding of containers” shows sample encodings for container types. In each of the examples below, the outermost container is encoded as an anonymous element.

Table 128. Sample encoding of containers

Type and Value	Encoding (hex)
Empty Structure, {}	15 18
Empty Array, []	16 18
Empty List, []	17 18
Structure, two context specific tags, Signed Integer, 1 octet values, {0 = 42, 1 = -17}	15 20 00 2a 20 01 ef 18

Type and Value	Encoding (hex)
Array, Signed Integer, 1-octet values, [0, 1, 2, 3, 4]	16 00 00 00 01 00 02 00 03 00 04 18
List, mix of anonymous and context tags, Signed Integer, 1 octet values, [[1, 0 = 42, 2, 3, 0 = -17]]	17 00 01 20 00 2a 00 02 00 03 20 00 ef 18
Array, mix of element types, [42, -170000, {}, 17.9, "Hello!"]	16 00 2a 02 f0 67 fd ff 15 18 0a 33 33 8f 41 0c 06 48 65 6c 6c 6f 21 18

Table 129, “Sample encoding of different tag types” shows sample encoding of a value with different associated tags, using Vendor ID = : 65522 (0xFFFF2), one of the [Vendor IDs allocated for testing purposes](#).

Table 129. Sample encoding of different tag types

Type and Value	Encoding (hex)
Anonymous tag, Unsigned Integer, 1-octet value, 42U	04 2a
Context tag 1, Unsigned Integer, 1-octet value, 1 = 42U	24 01 2a
Common profile tag 1, Unsigned Integer, 1-octet value, Matter::1 = 42U	44 01 00 2a
Common profile tag 100000, Unsigned Integer, 1-octet value, Matter::100000 = 42U	64 a0 86 01 00 2a
Fully qualified tag, Vendor ID 0xFFFF1/65521, profile number 0xDEED/57069, 2-octet tag 1, Unsigned Integer, 1-octet value 42, 65521::57069:1 = 42U	c4 f1 ff ed de 01 00 2a
Fully qualified tag, Vendor ID 0xFFFF1/65521, profile number 0xDEED/57069, 4-octet tag 0xAA55FEED/2857762541, Unsigned Integer, 1-octet value 42, 65521::57069:2857762541 = 42U	e4 f1 ff ed de ed fe 55 aa 2a
Structure with the fully qualified tag, Vendor ID 0xFFFF1/65521, profile number 0xDEED/57069, 2-octet tag 1. The structure contains a single element labeled using a fully qualified tag under the same profile, with 2-octet tag 0xAA55/43605. 65521::57069:1 = {65521::57069:43605 = 42U}	d5 f1 ff ed de 01 00 c4 f1 ff ed de 55 aa 2a 18

Appendix B: Tag-length-value (TLV) Schema Definitions

B.1. Introduction

A TLV Schema provides a simple textual description of the structure of data encoded in the [Matter TLV](#) format. A single TLV Schema MAY define the structure of multiple different TLV-encoded payloads. This section describes the syntax one can use to define a TLV Schema.

B.1.1. Basic Structure

A TLV Schema takes the form of a series of definitions. Each definition describes some construct, such as a data type. Each definition has an associated human readable name separated from the definition with a $\Rightarrow$ symbol. As a mnemonic device, it is useful to read the $\Rightarrow$ symbol as “is a”. For example, the following definition defines a data type that MAY be used to represent a sensor sample:

Example

```
/** Sensor sample structure */
sensor-sample => STRUCTURE
{
    timestamp [1] : UNSIGNED INTEGER [ range 32-bits ],
    value [2] : FLOAT64,
}
```

This example would be read as “*sensor-sample* is a structure containing a *timestamp* and *value*”.

A TLV Schema MAY contain multiple definitions. The order of definitions within a TLV Schema is unimportant.

B.1.2. Keywords

TLV Schemas employ various keywords when describing a construct. These keywords (e.g. **STRUCTURE**, **SIGNED INTEGER**, and **range**) are an inherent part of the schema language. Keywords in TLV Schemas are always case-insensitive. However, by convention, keywords associated with types and other high-level constructs are capitalized for emphasis in text-only contexts.

B.1.3. Naming

Each definition in a TLV Schema assigns a human-readable name to the construct being defined. This name serves both as a descriptive title as well as a means to refer to the construct from elsewhere in the schema.

Names in TLV Schemas are limited to ASCII alphanumeric characters, plus dash (-) and underscore (_). Additionally, all names SHALL begin with either an alphabetic character or an underscore. In

general, any name conforming to these rules MAY be used, as long as it does not collide with a [key-word](#) used by the schema language.

B.1.4. Namespaces

The name assigned to a schema construct SHALL be unique relative to all other named constructs in the same scope. To facilitate this, TLV Schemas support a namespacing mechanism similar to that provided in languages like C++.

The names of constructs defined within a [namespace](#) definition are only required to be unique within the given namespace. Namespaces themselves MAY be nested to any depth.

Constructs defined in other namespaces MAY be referenced using a name that gives the enclosing namespaces, plus the construct name, each separated by dots (.). Such a multi-part name is called a *scoped name*. For example:

Namespaces Example

```
namespace a
{
    x => STRING,
    other-x => b.x
}

namespace b
{
    x => SIGNED INTEGER
}
```

See [namespace-def](#) for further details.

B.1.5. Qualifiers

Constructs within a TLV Schema MAY be annotated with additional information using a *qualifier*. Qualifiers appear within square brackets ([...]) immediately following the construct they affect. In most cases the use of qualifiers is optional, but there are some situations where the schema syntax requires a qualifier.

Often qualifiers are used to place restrictions on the form or range of values that a construct can assume. For example a [length](#) qualifier MAY be used to constrain the length of a STRING type:

```
international-standard-book-number => STRING [length 13]
```

Multiple qualifiers MAY appear within the square brackets, and SHALL be separated by commas.

See [Section B.5, “Qualifiers”](#) for further details.

B.1.6. Tagging

In a TLV Schema, [tag numbers](#) appear as [qualifiers](#) attached to a particular named construct, such as a field within a structure. This association reflects the tag's role as an alias for the textual name in the TLV encoding. The syntax for tag qualifiers is defined in [tag](#). For example:

Tagging Example

```
certificate [my-protocol:1] => STRUCTURE
{
    serial-num [1] : OCTET STRING,
    ...
}
```

B.2. Definitions

A Matter TLV Schema consists of a set of one or more definitions. The definitions that MAY appear within a schema are:

- [type-def](#)
- [field-group-def](#)
- [namespace-def](#)
- [protocol-def](#)
- [vendor-def](#)

B.2.1. Type Definition (type-def)

Type Definition Syntax

```
type-name [ qualifier ] => type-or-ref
```

```
type-or-ref:
```

```
    type
```

```
    type-ref
```

```
type:
```

```
    ANY
```

```
    ARRAY
```

```
    ARRAY OF
```

```
    BOOLEAN
```

```
    CHOICE OF
```

```
    FLOAT32
```

```
    FLOAT64
```

```
    LIST
```


LIST OF
 NULL
 OCTET STRING
 SIGNED INTEGER
 STRING
 STRUCTURE
 UNSIGNED INTEGER

type-ref:
 type-name
 scoped-type-name

qualifier:
 tag

A type definition associates a name (**type-name**) with a schema construct representing a TLV type or **pseudo-type**. The given name serves as a descriptive title for the type, as well as a means to refer to the type from elsewhere in the schema.

Type definitions (**type-def**) are often used to describe TLV types that appear directly in some form of communication. For example, a type definition MAY define the structure of data carried within the payload of a message. Some type definitions may be used to define general purpose TLV constructs which are then employed in the definitions of other types.

The type (**type-name**) associated with a type definition MAY be any one of the available TLV types or **pseudo-type**. Alternatively, a type definition MAY contain a scoped type (**scoped-type-name**) referring to another type definition appearing elsewhere in the schema. This form is referred to as a type reference (**type-ref**). The ordering of type definitions and type references within a schema is unimportant, implying that a type reference MAY refer to a type that is defined later in the schema.

A **tag** qualifier MAY be applied to the name within a type definition to associate a default tag with that name. The default tag will be used in an encoding of the type whenever an explicit tag has not been given.

B.2.2. FIELD GROUP Definition (field-group-def)

FIELD GROUP Definition Syntax

```
field-group-name => FIELD GROUP { field-group-members }
```

field-group-members:

```
    field-group-member
    field-group-members, field-group-member
```

field-group-member:

```
    identifier [ id-qualifier ] : type-or-ref    // field definition
```

```
includes field-group-name
```

```
// field group include
```

```
id-qualifier:
```

```
  tag
  optional
```

```
// SHALL be present
```

FIELD GROUP declares a collection of fields that MAY be included in a TLV **Structure**. A **FIELD GROUP** is never directly encoded in a TLV encoding. A **FIELD GROUP** is used with **includes** statement to define common patterns of fields such that they MAY be reused across different **STRUCTURE** definitions.

A **FIELD GROUP** definition (**field-group-def**) contains a list of field definitions, each of which gives the type of the field, its tag, and an associated textual name. The field type MAY be either a fundamental type, a **CHOICE OF** pseudo-type, an **ANY** pseudo-type, or a reference to one of these types defined outside the **FIELD GROUP** definition.

A **FIELD GROUP** definition MAY also contain one or more **includes** statements. Each such statement identifies another **FIELD GROUP** whose fields are to be included within the referencing **FIELD GROUP**. Such nested inclusion MAY be specified to any depth.

The rules governing the names and tags associated with fields within a **FIELD GROUP** are the same as those defined for **STRUCTURE**.

FIELD GROUP Definition Examples

```
common-sensor-sample-fields => FIELD GROUP
{
    timestamp [1]          : UNSIGNED INTEGER [ range 32-bits ],
}

temperature-sensor-sample => STRUCTURE [tag-order]
{
    includes common-sensor-sample-fields,

    value [2]              : FLOAT64,
}

humidity-sensor-sample => STRUCTURE [tag-order]
{
    includes common-sensor-sample-fields,

    value [2]              : UNSIGNED INTEGER [ range 64-bits ],
}
```

B.2.3. Namespace Definition (namespace-def)

Namespace Definition Syntax

```
namespace ns-name { ns-scoped-defs }
```

ns-name:

name

scoped-name

ns-scoped-defs:

ns-scoped-def

ns-scoped-defs, ns-scoped-def

ns-scoped-def:

type-def

field-group-def

protocol-def

namespace-def

namespace introduces a new naming scope. Definitions that appear within the braces of a namespace definition are scoped to that namespace, such that their names need only be unique within the bounds of the enclosing scope. The namespace scoped definitions SHALL be separated by commas.

In general, four forms of definitions MAY appear within a namespace: type definitions (**type-def**), **FIELD GROUP** definitions (**field-group-def**), protocol definitions (**protocol-def**) and further namespace definitions (**namespace-def**). Namespace definitions MAY be nested to any level. Protocol definitions, however, are restricted such that they SHALL NOT be nested. Thus a namespace can only contain a protocol definition if the namespace itself is not located, at any level, within another protocol definition.

The name used in a namespace definition MAY be either a simple **name**, such as **a**, or a **scoped-name**, such as **a.b.c**. When a **scoped-name** is used, the effect is exactly as if multiple nested namespaces had been declared, each named after a part of the scoped name.

It is legal to have multiple namespace definitions, each with the same name, defined within the same scope. The effect is as if there were only a single namespace definition containing a union of the enclosed definitions. Thus, a namespace definition with the same name as a preceding definition MAY be seen as a kind of continuation of the earlier one.

Namespace Definition Examples

```
namespace abc
{
    property => FLOAT32 [ range 0..50 ],
```

```

point => STRUCTURE
{
    day [0] : UNSIGNED INTEGER,
    prop [1] : property
}

namespace matter.protocols.aaa
{
    config => STRUCTURE
    {
        points [0] : ARRAY OF abc.point,
        ...
    }
}

```

B.2.4. PROTOCOL Definition (protocol-def)

PROTOCOL Definition Syntax

```

name => PROTOCOL [ qualifier ] { protocol-scoped-defs }

qualifier:
    id

protocol-scoped-defs:
    protocol-scoped-def
    protocol-scoped-defs, protocol-scoped-def

protocol-scoped-def:
    type-def
    field-group-def
    namespace-def

```

PROTOCOL defines a Matter protocol. A Matter protocol is a group of logically related Matter TLV constructs that together serve a common purpose.

Similar to a **namespace** definition, a **PROTOCOL** definition introduces a new naming scope in which further definitions may appear. The names of definitions appearing within the braces of a **PROTOCOL** are scoped in exactly the same way as if they had appeared within a **namespace** definition. Likewise, constructs outside the **PROTOCOL** definition MAY refer to definitions within the protocol by using a scoped name that includes the protocol name. The **PROTOCOL** scoped definitions SHALL be separated by commas.

PROTOCOL definitions MAY appear at the global naming scope, or within a namespace definition. However, **PROTOCOL** definitions SHALL NOT be nested within other **PROTOCOL** definitions at any depth.

Every **PROTOCOL** definition SHALL include an **id** qualifier giving the **id** of the protocol, that uniquely identifies the protocol among all other protocols. The **id** given in a **PROTOCOL** definition SHALL be unique relative to all other **PROTOCOL** definitions in a schema. However, it is legal to have multiple **PROTOCOL** definitions with the same protocol id, provided that they also have the same name and appear within the same naming scope. The effect of this is as if there were only a single **PROTOCOL** definition containing a union of the enclosed definitions. This makes it possible to break up a **PROTOCOL** definition across multiple schema files.

PROTOCOL Definition Example

```
namespace some.names {  
  
    my-protocol => PROTOCOL [ VENDOR:0x0008 ]  
    {  
        laser-transducer-metadata [1] => STRUCTURE  
        {  
            serial-num [1] : OCTET STRING,  
            ...  
        },  
  
        optics-array [2] => ARRAY OF optical-specification  
    }  
  
}
```

B.2.5. VENDOR Definition (vendor-def)

VENDOR Definition Syntax

```
name => VENDOR [ qualifier ]  
  
qualifier:  
    id
```

VENDOR associates a name with a Vendor ID. **VENDOR** definition SHALL include an **id** qualifier giving the **id** of the vendor.

In a TLV Schema that includes a **VENDOR** definition, the vendor name MAY be used elsewhere in the schema as a stand-in for the associated **Vendor ID**. One such place where a vendor name may appear is within the id qualifier of a **PROTOCOL** definition.

VENDOR definitions MAY only appear at the global name scope, implying they SHALL NOT be placed within the body of a **namespace** or **PROTOCOL** definition.

Both the name and id value used in a **VENDOR** definition SHALL be unique across all such definitions. However, for convenience, a **VENDOR** definition MAY be repeated provided that the name and id are the same.

The Matter vendor (0x00000) is implicitly defined in all schemas, although it MAY be explicitly defined as well:

VENDOR Definition Example

```
Matter => VENDOR [ 0x0000 ]
```

B.3. Types

The TLV format supports 10 fundamental types: integers (signed and unsigned), floats, booleans, UTF-8 strings, octet strings, structures, arrays, lists and nulls. Accordingly, a TLV Schema MAY use one of the following type constructs to constrain an encoding to be one of these fundamental types.

B.3.1. ARRAY / ARRAY OF

ARRAY Syntax

```

ARRAY [ qualifier ] OF type-or-ref           // uniform array
ARRAY [ qualifier ] { type-pattern }         // pattern array

qualifier (optional):
    length
    nullable

type-pattern:
    type-pattern-item
    type-pattern, type-pattern-item

type-pattern-item:
    type-or-ref quantifier                   // unnamed item
    identifier : type-or-ref quantifier      // named item

quantifier (optional):
    *                                       // zero or more
    +                                       // one or more
    { count }                               // exactly count
    { min..max }                             // between min and max
    { min.. }                               // at least min

```

ARRAY and **ARRAY OF** declare an element that is encoded as a TLV [Array](#).

ARRAY OF declares an array where all the items in the array are of the same fundamental type, or taken from the same set of possible types. This form of array is called a *uniform array*, and is generally used to represent ordered collections of values.

ARRAY declares an array where the types of the array items follow a particular pattern. In this form, known as a *pattern array*; the allowed type for an item depends on its position in the array. The overall pattern of types allowed in the array is declared using a schema construct called a linear type pattern, which is similar to a regular expression (see below). Pattern arrays are typically used to represent vectors, tuples or paths.

A **length** qualifier on an array MAY be used to constraint the minimum and maximum number of items in the array. For a pattern array, the given length constraint SHALL be consistent with (i.e. fall within) the minimum and maximum number of items implied by the type pattern. In cases where the **length** qualifier places a narrower constraint on the length of an array than that implied by the type pattern, the **length** qualifier constraint takes precedence.

A **nullable** qualifier MAY be used to indicate that a TLV **Null** MAY be encoded in place of the **ARRAY** or **ARRAY OF**. Note that an array that has been replaced by a Null is distinct in terms of its encoding from an array that has no items.

B.3.1.1. Linear Type Patterns

A linear type pattern describes the sequence of TLV types that MAY appear in a TLV Array or List element. In its simplest form, a linear type pattern is a list of type definitions, or references to defined types, where each item constrains the TLV type that appears at the corresponding position in the collection. The type pattern is always anchored at the start of the collection, with the first type constraining the first item in the collection. Any type or pseudo-type MAY appear within a linear type pattern.

More complex type patterns can be created by using a quantifier. Quantifiers appear after a type in a type pattern and specify the number of times the associated type MAY appear at that position in the collection. Quantifiers borrow common regular expression notation to denote repetition, with ***** meaning zero or more, **+** meaning one or more, and **{ }** expressing specific counts. Using quantifiers, one can express complex sequences of types, including some that require arbitrary look-ahead to match.

B.3.1.2. Item Names

Items or groups of items in a pattern array MAY be given textual names. These names do not affect the encoding of the array, but serve as user documentation, or as input to code generation tools. Item names within a pattern array SHALL be unique.

Per the rules for encoding TLV arrays, array items SHALL NOT have tags. Thus the tag qualifier SHALL NOT be applied to an item name with a pattern array.

ARRAY Example

```
supported-country-codes => ARRAY [ length 0..10 ] OF STRING [ length 2 ]
```

```

weather-tuple => ARRAY
{
    timestamp          : UNSIGNED INTEGER [ range 32-bits ],
    temperature        : FLOAT64,
    relative-humidity  : UNSIGNED INTEGER [ range 0..100 ],
    precipitation       : UNSIGNED INTEGER [ range 0..100 ],
}

named-vector => ARRAY
{
    name                : STRING,
                      FLOAT64 *,
}

```

B.3.2. BOOLEAN

BOOLEAN Syntax

```

BOOLEAN [ qualifier ]

qualifier (optional):
    nullable

```

BOOLEAN declares an element that SHALL be encoded as a TLV **Boolean**.

If the **nullable** qualifier is given, a TLV **Null** MAY be encoded in its place.

BOOLEAN Example

```

pathlight-enabled => BOOLEAN

```

B.3.3. FLOAT32 / FLOAT64

FLOAT Syntax

```

FLOAT32 [ qualifier ]
FLOAT64 [ qualifier ]

qualifier (optional):
    range
    nullable

```

FLOAT32 declares an element that SHALL be encoded as a TLV **floating point number** with the **element type** indicating a 4-octet **IEEE 754-2019** single-precision value. Correspondingly, **FLOAT64** declares a TLV element that SHALL be encoded as a TLV **floating point number** with the **element**

type indicating an 8-octet [IEEE 754-2019](#) double-precision value.

If the [nullable](#) qualifier is given, a TLV **Null** MAY be encoded in place of the number.

The allowed range of values can be constrained using the [range](#) qualifier. If omitted, the value is constrained by what the relevant TLV type can represent.

Float Example

```
set-value => FLOAT32 [ range 0..50 ]
```

B.3.4. SIGNED INTEGER / UNSIGNED INTEGER

Integer Syntax

```
SIGNED INTEGER [ qualifier ] { enum }  
UNSIGNED INTEGER [ qualifier ] { enum }
```

qualifier (optional):

[range](#)

[nullable](#)

enum:

identifier = int-value

SIGNED INTEGER declares an element that SHALL be encoded as a TLV [integer](#) with the [element type](#) indicating the integer is signed. Correspondingly, **UNSIGNED INTEGER** declares a TLV element that SHALL be encoded as a TLV [integer](#) with the [element type](#) indicating the integer is unsigned.

If the [nullable](#) qualifier is given, a TLV **Null** MAY be encoded in place of the integer.

The allowed range of values MAY be constrained using the [range](#) qualifier. If omitted, the value is constrained by what the relevant TLV type can represent.

SIGNED INTEGER and **UNSIGNED INTEGER** definitions MAY include a set of enumerated values ([enum](#)), each of which associates a textual name ([identifier](#)) with a constant integer value ([int-value](#)). Each value SHALL conform to the allowed range of values for the **SIGNED INTEGER** definition as given by its sign and any [range](#) qualifier. The presence of enumerated values SHALL NOT restrict senders to only encoding those values. Rather, enumerations merely give symbolic names to particular noteworthy values.

Integer Examples

```
sensor-value => SIGNED INTEGER [ range -100..100 ]
```

```
counter => UNSIGNED INTEGER [ range 32-bits ]
```

B.3.5. LIST / LIST OF

LIST Syntax

```

LIST [ list-qualifier ] OF type-or-ref           // uniform list
LIST [ list-qualifier ] { type-pattern }         // pattern list

list-qualifier (optional):
    length
    nullable

type-pattern:
    type-or-ref quantifier                       // unnamed item
    identifier [ qualifier ] : type-or-ref quantifier // named item

quantifier (optional):
    *                                           // zero or more
    +                                           // one or more
    { count }                                   // exactly count
    { min..max }                               // between min and max
    { min.. }                                   // at least min

qualifier (optional):
    tag

```

LIST and **LIST OF** declare an element that is encoded as a TLV [List](#). **LIST** and **LIST OF** declare the same fundamental type, but differ based on how the allowed types of their items are expressed.

LIST OF declares a list where all the items in the list are of the same fundamental type, or taken from the same set of possible types. This form of list is called a *uniform list*. Uniform lists are generally used to represent ordered collections of values where the tags differentiate the semantic meaning of the value.

LIST declares a *pattern list* where the types of the items in the list follow a particular pattern. In this form, the allowed type(s) for an item depends on its position in the array. Pattern lists are typically used to represent path-like constructs.

The overall pattern of types allowed in a pattern list is declared using a schema construct called a linear type pattern. The syntax and interpretation of linear type patterns for pattern lists are the same as those for pattern arrays (see [Section B.3.1.1, “Linear Type Patterns”](#)).

The [length](#) qualifier MAY be used to constraint the minimum and maximum number of items in the list. For a pattern list, the given length constraint SHALL be consistent with (i.e. fall within) the minimum and maximum number of items implied by the type pattern. In cases where the [length](#) qualifier places a narrower constraint on the length of a list than that implied by the type pattern, the [length](#) qualifier constraint takes precedence.

A **nullable** qualifier MAY be used to indicate that a TLV **Null** MAY be encoded in place of the **LIST** or **LIST OF**. Note that a list that has been replaced by a Null is distinct (in terms of its encoding) from a list that has no items.

B.3.5.1. Item Names

As with the **ARRAY** type, items or groups of items in a pattern list MAY be given textual names to distinguish their purposes. Item names within a pattern list SHALL be unique.

B.3.5.2. Item Tags

Items within a pattern list can have a **tag** qualifier that specifies a particular tag value that SHALL be encoded with the item. The specific tag can be protocol-specific or context-specific, or the **anonymous** tag. The assigned tag values are not required to be unique among the items in a pattern list.

When no explicit tag qualifier is given (which is always the case for uniform lists) the items in a list automatically assume the default tag of their underlying types, if such a tag is provided. This can occur in two situations: 1) when the underlying type is a reference to a type definition that declares a default tag, and 2) when the underlying type is a **CHOICE OF** whose alternates declare default tags. See [default tag](#) for further information.

If no tag qualifier is given, and no default tag is available, an encoder is allowed to encode list items with any tag of their choosing.

B.3.6. OCTET STRING

OCTET STRING Syntax

```
OCTET STRING [ qualifier ]
```

qualifier (optional):

[length](#)
[nullable](#)

OCTET STRING declares an element that is encoded as a TLV **Octet String**, and in particular with the **element type** indicating it's an Octet String.

The minimum and maximum number of bytes can be constrained using the [length](#) qualifier.

OCTET STRING Example

```
address => OCTET STRING [ length 8 ]
```

B.3.7. NULL

NULL Syntax

```
NULL
```

NULL declares an element that SHALL be encoded as a TLV **Null**. There are no qualifiers that can be associated with a **NULL** type.

NULL Example

```
serial-num => CHOICE OF { STRING, UNSIGNED INTEGER, NULL }
```

B.3.8. STRING

STRING Syntax

```
STRING [ qualifier ]
```

qualifier (optional):

```
length
nullable
```

STRING declares an element that is encoded as a TLV **UTF-8 String**, and in particular with the **element type** indicating it's a UTF-8 String.

If the **nullable** qualifier is given, a TLV **Null** MAY be encoded in place of the string.

The minimum and maximum length of the string can be constrained using the **length** qualifier.

STRING Example

```
name-field => STRING [ length 0..32 ]
```

B.3.9. STRUCTURE

STRUCTURE Syntax

```
STRUCTURE [ structure-qualifier ] { structure-fields }
```

structure-qualifier (optional):

```
extensible
order
nullable
```

structure-fields:

```
structure-field
structure-fields, structure-field
```

structure-field:

```
identifier [ id-qualifier ] : type-or-ref    // field definition
includes field-group-name                    // field group include
```

id-qualifier:

tag
optional

STRUCTURE declares an element that is encoded as a TLV **Structure**. The **STRUCTURE** fields SHALL be separated by commas.

A **STRUCTURE** definition declares the list of fields that MAY appear within the corresponding TLV **Structure**. Each field definition gives the type of the field, its tag, and an associated textual name. The field type MAY be either a fundamental type, a **CHOICE OF** pseudo-type, an **ANY** pseudo-type, or a reference to one of these types defined outside the **STRUCTURE** definition.

A **STRUCTURE** definition MAY also contain one or more **includes** statements. Each such statement identifies a **FIELD GROUP** definition whose fields are to be included within the TLV **Structure** as if they had been declared within the **STRUCTURE** definition itself (see **Includes FIELD GROUP** below).

An **extensible** qualifier MAY be used to declare that a structure can be extended at encoding time by the inclusion of fields not listed in the **STRUCTURE** definition.

The **order** qualifiers (**any-order**, **schema-order** and **tag-order**) MAY be used to specify a particular order for the encoding of fields within a TLV **Structure**.

A **nullable** qualifier MAY be used to indicate that a TLV **Null** MAY be encoded in place of the **STRUCTURE**.

B.3.9.1. Fields

Fields within a **STRUCTURE** are assigned textual names to distinguish them from one another. Each such name SHALL be distinct from all other field names defined within the **STRUCTURE** or included via a **includes** statement. Fields names do not affect the encoding of the resultant TLV, but MAY serve as either user documentation or input to code generation tools.

Per the rules of TLV, all fields within a TLV **Structure** SHALL be encoded with a distinct TLV tag. Field tags are declared by placing a **tag** qualifier on the field name. Both protocol-specific and context-specific tags are allowed on the fields in a **STRUCTURE** definition.

For a given field if the tag qualifier is missing then the underlying type SHALL provide a **default tag**. This can occur in two situations:

1. the underlying type is a reference to a type definition that provides a default tag
2. the underlying type is a **CHOICE OF** pseudo-type whose alternates provide default tags.

The tags associated with **includes** fields are inherited from the target **FIELD GROUP** definition.

All tags associated with the fields of a TLV **Structure** SHALL be unique. This is true not only for tags declared directly within the **STRUCTURE** definition, but also for any tags associated with fields that are incorporated via an **includes** statement.

The **anonymous** tag SHALL NOT be used as the tag for a field within a **STRUCTURE** definition.

The **optional** qualifier MAY be used to declare a field which can be omitted from the structure encoding under some circumstances.

B.3.9.2. CHOICE OF Fields

A field within a **STRUCTURE** definition MAY be defined to be a **CHOICE OF** (either directly within the **STRUCTURE** definition or via a type reference). Over the wire, such a field is encoded as one of the alternate types given in the **CHOICE OF** definition. For example, the **user-id** field in the following **STRUCTURE** MAY be encoded as either a TLV **UTF-8 String** or an **Unsigned Integer**.

STRUCTURE with CHOICE OF Field Example

```
user-information => STRUCTURE [ extensible ]
{
  user-id [1] : CHOICE OF
  {
    UNSIGNED INTEGER,
    STRING
  }
}
```

If a **tag** qualifier is given for a **CHOICE OF** field (e.g. **[1]** as shown above), that tag SHALL be used in the encoding of the field for all possible alternates. On the other hand, if a **tag** qualifier is *not* given, then the default tag associated with the selected **CHOICE OF** alternate SHALL be used in the encoding. For example, in the following structure, a context-tag of **1** will be encoded if the **user-id** field is an Unsigned Integer, or **2** if the field is a String.

STRUCTURE with CHOICE OF Field with Default Tag Example

```
user-information => STRUCTURE [ extensible ]
{
  user-id : CHOICE OF
  {
    id [1] : UNSIGNED INTEGER,
    name [2] : STRING
  }
}
```

Note that, in all cases, the tag or tags associated with a **CHOICE OF** field SHALL be unique within the context of the containing **STRUCTURE**.

B.3.9.3. Includes FIELD GROUP

A **includes** statement MAY be used within a **STRUCTURE** definition to incorporate the fields of a **FIELD GROUP** defined outside the **STRUCTURE**. The fields of the **FIELD GROUP** are included in the **STRUCTURE** as if they had been listed within the **STRUCTURE** definition itself.

A particular **FIELD GROUP** SHALL NOT be included more than once within a given **STRUCTURE**.

The names assigned to fields within an included **FIELD GROUP** SHALL be distinct with respect to all other fields contained within the enclosing **STRUCTURE**, whether defined directly within the **STRUCTURE** itself, or included from another **FIELD GROUP**.

Likewise, tags assigned to fields within an included **FIELD GROUP** SHALL be distinct with respect to all other fields within the enclosing **STRUCTURE**.

B.4. Pseudo-Types

Pseudo-types are type-like constructs that provide flexibility in schema definitions. Some pseudo-types, like **CHOICE OF** and **ANY**, allow for variance in the fundamental TLV types that may appear in an encoding. Others make it easier to reuse schema constructs in multiple contexts.

B.4.1. ANY

ANY Syntax

```
ANY
```

ANY declares an element that can be encoded as any fundamental TLV type. Note that **ANY** is not a fundamental TLV type itself, but rather a pseudo-type that identifies a range of possible encodings. An **ANY** type serves a shorthand for (and is exactly equivalent to) a **CHOICE OF** all possible fundamental types.

There are no qualifiers that can be associated with an **ANY** type.

ANY Example

```
app-defined-metadata => ANY
```

B.4.2. CHOICE OF

CHOICE OF Syntax

```
CHOICE [ qualifier ] OF { alternates }
```

```
qualifier (optional):
```

```
    nullable
```

```
alternates:
```

```
    alternate
```

```
    alternates, alternate
```

```
alternate:
```

```
    type-or-ref
```

```
// unnamed alternate
```

```
identifier [ id-qualifier ] : type-or-ref    // named alternate
```

```
id-qualifier:
  tag
```

CHOICE OF declares an element that MAY be any of a set of TLV types. **CHOICE OF** is considered a **pseudo-type**, rather than a fundamental type, in that the **CHOICE OF** itself doesn't have a representation in the final TLV encoding.

The allowed TLV types for a **CHOICE OF**, known as **alternates**, are given in the body of the definition. An **alternate** MAY be any of the fundamental TLV types, an **ANY** pseudo-type, or another **CHOICE OF** definition (more on this below). Additionally, an alternate MAY be a type reference (in the form of a scoped type name) referring to a type defined outside of the **CHOICE OF** definition.

A **nullable** qualifier MAY be used to indicate that a TLV **Null** can be encoded in place of the **CHOICE OF**. This is exactly the same as if NULL had been listed as one of the alternates.

B.4.2.1. Alternate Names and Tags

Alternates MAY be assigned textual names to distinguish them from one another. Each such name SHALL be unique within the particular **CHOICE OF** definition. Alternate names do not affect the encoding of the resultant TLV. Rather, alternate names serve as user documentation, or as input to code generation tools.

Named **CHOICE OF** alternates MAY include a **tag** qualifier assigning a particular **tag** value to the **alternate**. When qualified in this way, the given tag value serves as a default tag for the alternate whenever the **CHOICE OF** appears in a context that doesn't otherwise specify a tag. The tags assigned within a **CHOICE OF** do not need to be unique, although see the discussion of **Ambiguous Alternates** below.

Both protocol-specific and context-specific tags are allowed on the alternates of a **CHOICE OF** definition.

B.4.2.2. Nested CHOICE OF and CHOICE OF Merging

It is legal for an alternate within a **CHOICE OF** to be another **CHOICE OF** definition, or a type reference to such. In this case, the effect is exactly as if the alternates of the inner **CHOICE OF** definition had been declared directly with the outer definition. This merging of **CHOICE OF** alternates occurs to any level of nesting, and MAY be used as a means of declaring multiple **CHOICE OF** that are supersets of other **CHOICE OF**.

When alternates are merged, their names are preserved. In cases where the same name appears in nested **CHOICE OF** definitions, the name of the outer alternate is prepended to that of the inner alternate, separated by a dot, to form a unique name for the merged alternate. In these cases, the outer alternate SHALL have a name in the schema, to ensure uniqueness.

An example of invalid **CHOICE OF** syntax, which results in a name conflict when alternates are merged:

CHOICE OF Invalid Alternates Merge Example

```
CHOICE OF {  
  CHOICE OF {  
    foo: STRING,  
    bar: UNSIGNED INTEGER  
  },  
  CHOICE OF {  
    foo: BOOLEAN,  
    bar: FLOAT64  
  }  
}
```

The example below shows how a valid schema should look to avoid conflict:

CHOICE OF Valid Alternates Merge Example

```
CHOICE OF {  
  alt1: CHOICE OF {  
    foo: STRING,  
    bar: UNSIGNED INTEGER  
  },  
  alt2: CHOICE OF {  
    foo: BOOLEAN,  
    bar: FLOAT64  
  }  
}
```

B.4.2.3. Ambiguous Alternates

A **CHOICE OF** MAY contain multiple alternates having the same fundamental TLV type (e.g. two alternates that are both **SIGNED INTEGER**). If these alternates are also encoded using the same tag, their encoded forms are effectively indistinguishable from one another. Such alternates are referred to as ambiguous alternates.

Ambiguous alternates MAY occur due to the merging of nested **CHOICE OF** definitions (see above). They MAY also arise in cases where the tags associated with the alternates are overridden by a tag qualifier in an outer context; e.g. when a **STRUCTURE** incorporates a **CHOICE OF** field that has a specific tag qualifier assigned to the field.

Ambiguous alternates are legal in TLV Schemas. However, care SHALL be taken when introducing ambiguous alternates to ensure that a decoder can correctly interpret the resulting encoding. This can be achieved, for example, by signaling the appropriate interpretation via a data value (e.g. an enumerated integer) contained elsewhere in the encoding.

B.5. Qualifiers

Qualifiers are annotations that provide additional information regarding the use or interpretation of a schema construct. Often qualifiers are used to place restrictions on the form or range of values that the construct can assume.

B.5.1. any-order / schema-order / tag-order

Order Qualifiers Syntax

```
STRUCTURE [ any-order ]  
STRUCTURE [ schema-order ]  
STRUCTURE [ tag-order ]
```

The **any-order**, **schema-order** and **tag-order** qualifiers MAY be used to specify a particular order for the encoding of fields within a **STRUCTURE**.

The **any-order** qualifier specifies that the encoder of a TLV structure is free to encode the fields of the structure in any desired order.

The **schema-order** qualifier specifies that the fields of a structure SHALL be encoded in the order given within the associated **STRUCTURE** definition. If the **STRUCTURE** definition contains one or more **includes** statements, the fields of the referenced **FIELD GROUPs** SHALL be encoded in the order given in the respective **FIELD GROUP** definition, and at the position of the **includes** statement relative to other fields within the **STRUCTURE**.

The **tag-order** qualifier specifies that the fields of a structure SHALL be encoded in the order specified by their tags, as defined in [Section A.2.4, “Canonical Ordering of Tags”](#).

Only a single ordering qualifier MAY be applied to a given **STRUCTURE** type.

In the absence of an order qualifier, fields within TLV structure MAY generally be encoded in any order. However, the author of a **STRUCTURE** definition MAY choose to impose custom ordering constraints on some or all of the fields if so desired. Such constraints SHALL be clearly described in the prose documentation for the schema.

B.5.2. extensible

extensible Qualifier Syntax

```
STRUCTURE [ extensible ]
```

The **extensible** qualifier is only allowed on **STRUCTURE** types, and declares that the structure MAY be extended by the inclusion of fields not listed in its definition. When a structure is extended in this way, any new fields that are included SHALL use tags that are distinct from any of those associated with defined or included fields.

Absent the **extensible** qualifier, a structure encoding SHALL NOT include fields beyond those given

in the **STRUCTURE** definition.

extensible Qualifier Example

```
user-information => STRUCTURE [ extensible ]
{
    user-id [1]      : UNSIGNED INTEGER,
    first-name [2]   : STRING,
    last-name [3]    : STRING,
    email-address [4] : STRING,
}
```

B.5.3. id

id Qualifier Syntax

```
vendor-name => VENDOR [ id uint-value ]

protocol-name => PROTOCOL [ id uint-value ]

protocol-name => PROTOCOL [ id uint-value:uint-value ]

protocol-name => PROTOCOL [ id vendor-name:uint-value ]
```

The **id** qualifier is used to specify an identifying number associated with a **VENDOR** or **PROTOCOL** definition.

When applied to a **VENDOR** definition, the **id** value is a 16-bit unsigned integer specifying the **Protocol Vendor ID**, which uniquely identify an organization or company. **VENDOR** ids are used to scope other identifiers (e.g. **PROTOCOL** ids) such that organizations can independently mint these identifiers without fear of collision.

When applied to a **PROTOCOL** definition, the **id** value MAY take three forms:

- 32-bit unsigned integer, which is composed of a **Protocol Vendor ID** in the high 16-bits and a **protocol id** in the low 16-bits
- two 16-bit unsigned integers (separated by a colon) specifying the **Protocol Vendor ID** and **protocol id**
- **vendor-name** and 16-bit **protocol id** (separated by a colon). The **vendor-name** definition SHALL exist elsewhere in the schema

id Qualifier Examples

```
MATTER-VENDOR-AB => VENDOR [ 0x00AB ]

// Equivalent definitions of the protocol introduced by MATTER-VENDOR-AB
```

```

vendor-ab-prot8 => PROTOCOL [ 0x00AB0008 ]
vendor-ab-prot8 => PROTOCOL [ 0x00AB:0x0008 ]
vendor-ab-prot8 => PROTOCOL [ MATTER-VENDOR-AB:8 ]

```

B.5.4. length

length Qualifier Syntax

```

type [ length count ]           // exactly count
type [ length min..max ]       // between min and max (inclusive)
type [ length min.. ]          // at least min

```

The **length** qualifier MAY be used to constrain the number of elements in a collection type, such as an **ARRAY** or **LIST**, or the number of bytes in a **STRING** or **OCTET STRING** type.

B.5.5. nullable

nullable Qualifier Syntax

```

type [ nullable ]

```

The **nullable** qualifier is used with **ARRAY**, **LIST**, **STRUCTURE**, **STRING**, **OCTET STRING**, **BOOLEAN**, **SIGNED INTEGER**, **UNSIGNED INTEGER**, **FLOAT32**, **FLOAT64** types. The **nullable** qualifier declares that a TLV **Null** MAY be substituted for a value of the specified type at a particular point in an encoding. For example, in the following **sensor-sample** structure, a null value MAY be encoded for the **value** field (e.g. in the case the sensor was off-line at the sample time):

nullable Qualifier Example

```

sensor-sample => STRUCTURE
{
    timestamp [1] : UNSIGNED INTEGER,
    value [2] : FLOAT64 [ nullable ],
}

```

Applying a **nullable** qualifier to a type is exactly the same as defining a **CHOICE OF** type with **alternates** for the primary and **NULL**. For example, the sensor sample structure could also be defined as follows:

nullable Qualifier Example

```

sensor-sample => STRUCTURE
{
    timestamp [1] : UNSIGNED INTEGER,
    value [2] : CHOICE OF
    {

```

```

        FLOAT64,
        NULL
    }
}

```

B.5.6. optional

optional Qualifier Syntax

```

...
field-name [ optional ] : type-or-ref,
...

```

The **optional** qualifier declares that a field within a **STRUCTURE** or **FIELD GROUP** is optional, and MAY be omitted by an encoder. The **optional** qualifier MAY only appear on the name portion of a field definition within either a **STRUCTURE** or **FIELD GROUP**.

Note that an optional field is distinct, both semantically and in terms of encoding, from a field whose type has been declared **nullable**. In the former case the field MAY be omitted from the encoding altogether. In the latter case the field SHALL appear within the encoding, however its value MAY be encoded as a TLV **Null**. It is legal to declare a field that is both **optional** and **nullable**.

The conditions under which an optional field can be omitted depend on the semantics of the structure. In some cases, fields MAY be omitted entirely at the discretion of the sender. In other cases, omission of a field MAY be contingent on the value present in another field. In all cases, prose documentation associated with the field definition SHALL make clear the rules for when the field may be omitted.

Optional fields are allowed within **FIELD GROUP** and retain their optionality when included within **STRUCTURE**.

optional Qualifier Example

```

user-information => STRUCTURE [ extensible ]
{
    user-id [1]                : UNSIGNED INTEGER,
    first-name [2]             : STRING,
    middle-name [3, optional]  : STRING,           // MIGHT be omitted
    last-name [4]              : STRING,
    email-address [5, optional] : STRING,          // MIGHT be omitted
}

```

B.5.7. range

range Qualifier Syntax

```
integer-type [ range min..max ]           // explicit constraint (inclusive)
integer-type [ range 8-bits ]             // width constraint
integer-type [ range 16-bits ]
integer-type [ range 32-bits ]
integer-type [ range 64-bits ]
```

The **range** qualifier MAY be used to constrain the range of values for a numeric type such as **SIGNED INTEGER**, **UNSIGNED INTEGER**, **FLOAT32**, or **FLOAT64**. Two forms are supported: *explicit constraints* and *width constraints*. Only one form MAY be applied to a given type.

An *explicit constraint* gives specific minimum and maximum (inclusive) values for the type. These MAY be any value that is legal for the underlying type.

A *width constraint* constrains the value to fit within a specific number of bytes. Any of the width constraints (8-bits, 16-bits, 32-bits or 64-bits) MAY be applied to **SIGNED INTEGER** and **UNSIGNED INTEGER** types, where 8-bits, 16-bits, 32-bits and 64-bits constraints correspond to 1-octet, 2-octet, 4-octet and 8-octet **element type** respectively; only 32-bits constraint MAY be applied to **FLOAT32** type and only 64-bits constraint MAY be applied to **FLOAT64** type.

Note that a width constraint **range** qualifier does not obligate an encoder to always encode the specified number of bits. Per the TLV encoding rules, senders are always free to encode integer and floating point values in any encoding size, bigger or smaller, that will accommodate the value.

range Qualifier Example

```
system-status-event => STRUCTURE
{
    timestamp [1]           : UNSIGNED INTEGER [ range 32-bits ],
    num-processes [2]       : UNSIGNED INTEGER [ range 16-bits ],
    percent-busy [3]        : UNSIGNED INTEGER [ range 0..100 ],
}
```

B.5.8. tag*tag Qualifier Syntax*

```
identifier [ tag-num ]           // context-specific tag
identifier [ protocol-id:tag-num ] // protocol-specific tag
identifier [ protocol-name:tag-num ] // protocol-specific tag
identifier [ *:tag-num ]         // protocol-specific tag (cur. protocol)
identifier [ anonymous ]         // no tag
```

The **tag** qualifier is allowed on type names, field names within a **STRUCTURE** (**STRUCTURE Fields**) or **FIELD GROUP**, item names within a **LIST** (**LIST Item Tags**), alternate names within a **CHOICE OF** (**CHOICE**

OF Fields).

The **tag** qualifier specifies a numeric tag value to be used when encoding a particular value. For brevity, the **tag** keyword SHALL be omitted when specifying a tag qualifier. As a special case, the keyword **anonymous** MAY be used to signal a value that SHALL be encoded without a tag.

Matter TLV supports two forms of tags: **Protocol-Specific Tags** and **Context-Specific Tags**. A protocol-specific tag is a colon-separated tuple containing a **protocol-id** and a **tag-num**. Protocol ids MAY also be specified indirectly, by giving the name of a **PROTOCOL** definition (**protocol-name**) located elsewhere in the schema. An asterisk (*) MAY be used as a shorthand to refer to the id of the **PROTOCOL** definition in which the tag qualifier appears. This protocol is referred to as the *current protocol*.

B.5.8.1. Explicit Tags

A **tag** qualifier that appears on a field within a **STRUCTURE** or **FIELD GROUP**, or on an item within a **LIST**, specifies the exact tag to be used when encoding the associated field/item. Such a tag is called an explicit tag, and MAY be either a **context-specific**, **protocol-specific** or **anonymous** (for **LIST**) tag.

If a field or item lacks a **tag** qualifier, then the encoding will use a **default tag** associated with the underlying field type, if such a tag has been specified.

B.5.8.2. Default Tags

A **tag** qualifier that appears on a type definition, or on an alternate within a **CHOICE OF**, serves as a default tag. A default tag is used to encode a value when an explicit tag has not been given in the schema.

For example, a field within a **STRUCTURE** that refers to a type with a default tag will use that tag if no **tag** qualifier has been specified on the field itself. Similarly, **tag** qualifiers that appear on the alternates of a **CHOICE OF** serve as default tags to be used when no other tag has been specified.

Both **context-specific** and **protocol-specific** tags MAY be used as default tags. 'anonymous' tag SHALL NOT be used as default tag.

Default Tag Qualifier Example

```
vendor-ab-prot8 => PROTOCOL [ id 0x00AB0008 ]
{
  ec-pub-key [0x00AB0008:1] => OCTET STRING,    // default protocol-specific tag using
                                              // a numeric protocol id

  ec-priv-key [vendor-ab-prot8:2] => STRUCTURE // default protocol-specific
                                              // tag using a name
  {
    priv-key [1]      : OCTET STRING,          // explicit context-specific tag

    pub-key [2, optional] : ec-pub-key         // explicit tag overrides default tag on
                                              // ec-pub-key
```

```

    curve          : CHOICE OF          // tag depends on choice of id or name
    {
        id [3]      : UNSIGNED INTEGER,  // default tag if id chosen
        name [4]    : STRING              // default tag if name chosen
    }
},

ecdsa-sig [*:3] => STRUCTURE           // shorthand for vendor-ab-prot8:3
{
    r [1]           : OCTET STRING,
    s [2]           : OCTET STRING
}
} // end vendor-ab-prot8 PROTOCOL

```

B.5.9. Documentation and Comments

TLV Schemas MAY include inline annotations that support the automatic generation of reference documentation and the production of documented code. TLV Schemas follow the Javadoc style of annotation wherein documentation is wrapped in the special multi-line comment markers **/**** and ***/**.

Documentation and Comments Example

```

/** Sensor sample structure */
sensor-sample => STRUCTURE
{
    timestamp [1] : UNSIGNED INTEGER,
    value [2] : FLOAT64,
}

```

In certain cases, documentation MAY also be placed after a construct, using **/**<** and ***/**.

Documentation and Comments for a Construct Example

```

/** Sensor sample structure */
sensor-sample => STRUCTURE
{
    timestamp [1] : UNSIGNED INTEGER, /**< Unix timestamp */
    value [2] : FLOAT64,              /**< Sensor value */
}

```

Postfix annotations are allowed on **STRUCTURE** and **FIELD GROUP** members, **ARRAY** and **LIST** items, **CHOICE OF** alternates, **SIGNED INTEGER** and **UNSIGNED INTEGER** enumerated values.

Non-documentation comments follow the standard C++ commenting style.

Documentation and Comments C++ Style Example

```
user-information => STRUCTURE [ extensible ]
{
    user-id [1]      : UNSIGNED INTEGER, // 0 = Unknown user id
    first-name [2]   : STRING,
    last-name [3]    : STRING,
    email-address [4] : STRING,

    /* TODO: additional fields to be added later */
}
```

Appendix C: Tag-length-value (TLV) Payload Text Representation Format

C.1. Introduction

This section describes a means by which to depict TLV payloads in a more user-friendly, textual representation.

C.2. Format Specification

C.2.1. Tag/Value

TLV elements are tag/value pairs. As such, their general textual representation is as follows:

```
tag = value
```

C.2.2. Context-Specific Tags

The basic representation of a context-specific tag is a single scalar number.

TLV entries using context-specific tags MAY use the basic representation alone:

```
2 = "hello"
```

If the tag has a name from an associated schema, it MAY be represented using that name. The basic representation MAY also be appended in parentheses ("(", ")"):

```
name (2) = "hello"
```

C.2.3. Protocol-Specific Tags

The basic representation of a protocol-specific tag SHALL be fully-qualified with "::" separating the vendor id and the protocol number and ":" separating the protocol number and tag number. The vendor id, protocol number and tag number are each represented using a single scalar number represented in hexadecimal notation.

```
0x0000::0x0000:0x01 = 10
```

If the tag has a name from an associated schema, it MAY be represented using that name. The basic representation MAY also be appended in parentheses ("(", ")"):

```
SmartSensorsCompany::SensingProtocol:Extension (0x00ef::0x00aa:0x01) = 10
```

C.2.4. Anonymous Tags

TLV entries using anonymous tags SHALL display the value alone:

```
"hello"
```

C.2.5. Primitive Types

Signed Integer:

```
duration = 20
```

Unsigned Integer:

```
duration = 20U
```

If the value is a defined constant, or enumerated value, then the string literal MAY be provided as well:

```
mode = FAST (20U)
```

UTF-8 string:

```
name = "Jonah"
```

Octet String (listed as 8-bit hex digits):

```
data = 2f 2a fd 11 33 e2 ...
```

Floats:

```
temp = 20.234
```

Booleans:

```
isOn = false
```

```
isOn = true
```

Null:

```
temp = null
```

C.2.6. Complex Types: Structure

Braces (`{ ... }`) SHALL be used to convey the start and end of structure scope, with the members separated by commas (`,`):

```
user-record = {  
    name = "Jonah",  
    pin = 1122  
}
```

C.2.7. Complex Types: Arrays

Square brackets (`[ ... ]`) SHALL be used to convey array scope, with elements in the array separated by commas (`,`). Since elements in the array are required to be anonymous, each element SHALL display the value alone:

```
temp-samples = [20, 30, 40]
```

C.2.8. Complex Types: List

Double square brackets (`[ [ ... ] ]`) SHALL be used to convey list scope, with elements in the list separated by commas (`,`). Since a diversity of tag types can be used in a list (including duplicates), the tags SHALL always be present and explicitly stated:

```
AttributePath = [ [ EndpointId = 20, ClusterId = 40 ] ]
```

C.3. Examples

C.3.1. TLV Schema

This is a sample TLV schema that will be used to define example TLV payloads.

```
temp-sample => STRUCTURE  
{  
    timestamp [1] : UNSIGNED INTEGER [ range 32-bits ],  
    temperature [2] : FLOAT64,  
}
```

```

accel-sample => STRUCTURE
{
    x [1] : SIGNED INTEGER [ range 16-bits ],
    y [2] : SIGNED INTEGER [ range 16-bits ],
    z [3] : SIGNED INTEGER [ range 16-bits ]
}

temp-features-enum => UNSIGNED INTEGER [ range 0...3 ]
{
    HAS_TEMP_COMPENSATION = 1,
    SUPPORTS_THRESHOLD_TRIGGERS = 2
}

accel-features-enum => UNSIGNED INTEGER [ range 0...3 ]
{
    SUPPORTS_HIGH_SAMPLING = 1,
    SUPPORTS_THRESHOLD_TRIGGERS = 2
}

sensor-state => STRUCTURE
{
    temperature-samples [1] : ARRAY OF temperature-sample,
    accel-samples [2] : ARRAY OF accel-sample,
    manufacturer-name [3]: STRING,

    // List of lists. If present, one or more of the feature lists will be present.
    feature-map [4,optional] : LIST {temp-features[1]: ARRAY OF temp-features-enum,
    accel-features[2]: ARRAY OF accel-features-enum},

    supports-idle [5] : BOOLEAN,

    num-power-modes[6]: UNSIGNED INTEGER [ range 8-bits ],

    supported-extensions[7] : LIST OF STRING
}

```

C.3.2. TLV Payloads

C.3.2.1. Temperature Sample

```

temperature-sample-example =
{

```

```
timestamp (1) = 2023423U,  
temperature (2) = 72.0  
}
```

C.3.2.2. Accelerometer Sample

```
accelerometer-sample-example =  
{  
    x (1) = 10,  
    y (2) = 20,  
    z (3) = 30  
}
```

C.3.2.3. Sensor State

```
sensor-state-example =  
{  
    temperature-samples (1) =  
    [  
        {  
            timestamp (1) = 2023423U,  
            temperature (2) = 72.0  
        },  
        {  
            timestamp (1) = 2023U,  
            temperature (2) = 69.2  
        },  
    ],  
    accel-samples (2) =  
    [  
        {  
            x (1) = 10,  
            y (2) = 20,  
            z (3) = 30,  
        },  
        {  
            x (1) = 1,  
            y (2) = 2,  
            z (3) = 3,  
        },  
    ],  
}
```

```
manufacturer-name (3) = "SmartSensors Ltd",

feature-map (4) =
[[
    temp-features (1) =
    [
        HAS_TEMP_COMPENSATION (1),
        SUPPORTS_THRESHOLD_TRIGGERS (2)
    ],

    accel-features (2) =
    [
        SUPPORTS_THRESHOLD_TRIGGERS (2)
    ]
]],

supports-idle (5) = false,

num-power-modes (6) = 2U,

supported-extensions (7) =
[[
    SMARTSENSORS::SensingProtocol:Extension (0x00ef::0x00aa:0x01) =
    "SUPPORTS_SMART_AVERAGING"
]]
}
```

Appendix D: Status Report Messages

D.1. Overview

The **StatusReport** is a core message that encapsulates the result of an operation which a responder sends as a reply for requests sent from an initiator, using a common message of the **Secure Channel Protocol** (Protocol ID = **PROTOCOL_ID_SECURE_CHANNEL**).

This section details the standard Status Report message format encoding as Matter Message Format payloads.

D.2. Status Report elements

A Status Report message describes a protocol-specific operation result or status.

The Status Report message SHALL have the following message header values (some of which may be omitted within protocol messages, as per header flag rules), no matter which protocol actually generated the status report:

- A **Protocol Vendor ID** set to 0 (Matter common Vendor ID)
- A **Protocol ID** set to 0x0000 (**PROTOCOL_ID_SECURE_CHANNEL**)
- A **Protocol Opcode** set to 0x40 (**StatusReport**)

The report message’s **Application Payload** SHALL consist of:

- A mandatory *GENERAL CODE* field, providing a general description of the status being reported.
- A mandatory *PROTOCOL SPECIFIC STATUS* field, providing additional details
- An **optional** protocol-specific data section that MAY include any additional information that a protocol requires
 - Individual protocols define the contents of this data section and how it is handled

D.3. Message Format

Length	Octets 0	1
<i>Structure, little-endian</i>		
2 octets	GeneralCode	
4 octets	ProtocolId of Protocol-Specific Status	
	...	
2 octets	ProtocolCode of Protocol-Specific Status	
Variable	Optional ProtocolData for protocol-specific additional details, MAY be empty	

D.3.1. General status codes (**GeneralCode**)

General status codes conveyed in the **GeneralCode** field are uniform codes that convey both success and failures.

The following general status codes are defined:

Code	Numeric Value	Description
SUCCESS	0	Operation completed successfully.
FAILURE	1	Generic failure, additional details may be included in the <i>protocol specific status</i> .
BAD_PRECONDITION	2	Operation was rejected by the system because the system is in an invalid state.
OUT_OF_RANGE	3	A value was out of a required range
BAD_REQUEST	4	A request was unrecognized or malformed
UNSUPPORTED	5	An unrecognized or unsupported request was received
UNEXPECTED	6	A request was not expected at this time
RESOURCE_EXHAUSTED	7	Insufficient resources to process the given request
BUSY	8	Device is busy and cannot handle this request at this time
TIMEOUT	9	A timeout occurred
CONTINUE	10	Context-specific signal to proceed
ABORTED	11	Failure, may be due to a concurrency error.
INVALID_ARGUMENT	12	An invalid/unsupported argument was provided
NOT_FOUND	13	Some requested entity was not found
ALREADY_EXISTS	14	The sender attempted to create something that already exists
PERMISSION_DENIED	15	The sender does not have sufficient permissions to execute the requested operations.
DATA_LOSS	16	Unrecoverable data loss or corruption has occurred.
MESSAGE_TOO_LARGE	17	Message size is larger than the recipient can handle.

If none of the specific codes above fits for application usage, a protocol SHALL use *FAILURE* and provide more information encoded in the **ProtocolId** and **ProtocolCode** subfields.

D.3.2. Protocol-specific codes (**ProtocolId** and **ProtocolCode**)

The protocol-specific portion of StatusReport messages is composed of a fully-qualified **ProtocolId** which qualifies the subsequent **ProtocolCode** space.

The **ProtocolId** is encoded as a 32 bit value of **Protocol Vendor ID** (upper 16 bits) and **Protocol ID**

under that Protocol Vendor ID (lower 16 bits), similarly to how message [Protocol ID](#) and Protocol Vendor ID are encoded in the [Protocol Header](#).

The following rules apply to the encoding of the [ProtocolCode](#) protocol-specific field:

- [ProtocolCode](#) value 0x0000 SHALL be reserved for use as success placeholder when either a [GeneralCode](#) of *SUCCESS* (0) or *CONTINUE* (10) are present.
- [ProtocolCode](#) value 0xFFFF SHALL be reserved to indicate that no additional protocol-specific status code is available.
- When the [GeneralCode](#) is *FAILURE*, the [ProtocolCode](#) value of 0xFFFF SHOULD NOT be used, since the conveyance of specific error codes assists in troubleshooting.
- [ProtocolCode](#) values 0x0001 through 0xFFFE SHALL be used to convey protocol-specific status indications.

Since protocol-specific status reports are meant to convey more information than generic codes, it is RECOMMENDED to always use a specific [ProtocolCode](#) value, rather than 0xFFFF, unless there are no additional details to convey.

D.3.3. Protocol-specific data ([ProtocolData](#))

The [ProtocolData](#) portion of the StatusReport message is composed of all data beyond the [ProtocolCode](#) field. If a StatusReport message of size *N* octets is received, the first 8 octets of payload encode the [GeneralCode](#), [ProtocolId](#) and [ProtocolCode](#), while the remaining *N* - 8 bytes represent the protocol-specific [ProtocolData](#).

Encoding of the [ProtocolData](#) portion of the payload depends on the [ProtocolId](#) and potentially [ProtocolCode](#). To decode this data, the [ProtocolId](#) has to be examined and decoding SHALL be done according to that protocol specification. For example:

- A vendor-specific protocol would encode additional custom error metadata in the [ProtocolData](#).
- The [Bulk transfer \(BDX\)](#) protocol does not require additional error information and will always have [ProtocolData](#) empty.

D.4. Presenting [StatusReport](#) messages in protocol specifications

In order to simplify referring to StatusReport messages, the following mnemonic encoding will be used in the descriptive text for a given protocol.

References to [StatusReport](#) messages take one of the following forms:

- No [ProtocolData](#) present:
 - [StatusReport](#)([GeneralCode](#): <value>, [ProtocolId](#): <value>, [ProtocolCode](#): <value>)
 - Example 1: [StatusReport](#)([GeneralCode](#): *FAILURE*, [ProtocolId](#): *BDX*, [ProtocolCode](#): *START_OFFSET_NOT_SUPPORTED*)
 - Encodes as: 01 00 02 00 00 00 52 00

- Example 2: `StatusReport(GeneralCode: SUCCESS, ProtocolId: {VendorID=0xFFFF1, ProtocolId=0xAABB}, ProtocolCode: 0)`
 - Encodes as: 00 00 BB AA F1 FF 00 00
- Additional `ProtocolData` present:
 - `StatusReport(GeneralCode: <value>, ProtocolId: <value>, ProtocolCode: <value>, ProtocolData: <value>)`
 - Example: `StatusReport(GeneralCode: FAILURE, ProtocolId: {VendorID=0xFFFF1, ProtocolId=0xAABB}, ProtocolCode: 9921, ProtocolData: [0x55, 0x66, 0xEE, 0xFF])`
 - Encodes as: 01 00 BB AA F1 FF C1 26 55 66 EE FF

Appendix E: Matter-Specific ASN.1 Object Identifiers (OIDs)

Matter defines custom ASN.1 OID values, which are listed in the table below under the 1.3.6.1.4.1.37244 private arc. These OID values are assigned by the Connectivity Standards Alliance for use with Matter.

Table 130. ASN.1 Matter-Specific Object Identifiers

Dot Notation	ASN.1 Notation	Description
1.3.6.1.4.1.37244.1.1	iso(1) org(3) dod(6) internet(1) private(4) enterprise(1) zigbee(37244) matter-op-cert(1) node-id(1)	Matter Operational Certificate DN attribute for node identifier
1.3.6.1.4.1.37244.1.2	iso(1) org(3) dod(6) internet(1) private(4) enterprise(1) zigbee(37244) matter-op-cert(1) firmware-signing-id(2)	Matter Operational Certificate DN attribute for firmware signing identifier
1.3.6.1.4.1.37244.1.3	iso(1) org(3) dod(6) internet(1) private(4) enterprise(1) zigbee(37244) matter-op-cert(1) ica-id(3)	Matter Operational Certificate DN attribute for Intermediate CA (ICA) identifier
1.3.6.1.4.1.37244.1.4	iso(1) org(3) dod(6) internet(1) private(4) enterprise(1) zigbee(37244) matter-op-cert(1) root-ca-id(4)	Matter Operational Certificate DN attribute for Root Certificate Authority (CA) identifier
1.3.6.1.4.1.37244.1.5	iso(1) org(3) dod(6) internet(1) private(4) enterprise(1) zigbee(37244) matter-op-cert(1) fabric-id(5)	Matter Operational Certificate DN attribute for fabric identifier
1.3.6.1.4.1.37244.1.6	iso(1) org(3) dod(6) internet(1) private(4) enterprise(1) zigbee(37244) matter-op-cert(1) case-authenticated-tag(6)	Matter Operational Certificate DN attribute for CASE Authenticated Tag
1.3.6.1.4.1.37244.1.7	iso(1) org(3) dod(6) internet(1) private(4) enterprise(1) zigbee(37244) matter-op-cert(1) vendor-verification-signer-id(7)	Matter Operational Certificate DN attribute for vendor verification signer identifier
1.3.6.1.4.1.37244.2.1	iso(1) org(3) dod(6) internet(1) private(4) enterprise(1) zigbee(37244) matter-att-cert(2) vid(1)	Matter Device Attestation Certificate DN attribute for the Vendor ID (VID)
1.3.6.1.4.1.37244.2.2	iso(1) org(3) dod(6) internet(1) private(4) enterprise(1) zigbee(37244) matter-att-cert(2) pid(2)	Matter Device Attestation Certificate DN attribute for the Product ID (PID)

Appendix F: Cryptographic test vectors for some procedures

F.1. Certification Declaration CMS test vector

This subsection contains worked examples of encoding a [Certification Declaration](#), which is conveyed by the [Attestation Information](#) payload during the [Device Attestation Procedure](#).

The Connectivity Standards Alliance CD signing certificate and associated private key which are provided in the vectors are only for exemplary purposes and are not official CD signing material.

The first example [Certification Declaration](#) has the following qualities:

- Both `dac_origin_vendor_id` and `dac_origin_product_id` are absent
- The `product_id_array` contains a single PID

The content of this first example is shown below:

```
===== Algorithm inputs =====
-> format_version = 1
-> vendor_id = 0xFFFF1
-> product_id_array = [ 0x8000 ]
-> device_type_id = 0x1234
-> certificate_id = "ZIG20141ZB330001-24"
-> security_level = 0
-> security_information = 0
-> version_number = 0x2694
-> certification_type = 0
-> dac_origin_vendor_id is not present
-> dac_origin_product_id is not present
-> authorized_paa_list is not present

-> Sample Connectivity Standards Alliance CD Signing Certificate:
-----BEGIN CERTIFICATE-----
MIIBszCCAVqgAwIBAgIIRdrzneR6oI8wCgYIKoZIzj0EAwIwKzEpMCcGA1UEAwg
TWf0dGvYIFRlc3QgQ0QgU2lnbmLuZyBBdXR0b3JpdHkwIBcNMjEwNjI4MTQyMzQz
WhgPOTk5OTEyMzEyMzU5NTlaMCsxKTAnBgNVBAMMIE1hdHRlcjBUZXNOIENEIFNp
Z25pbmcgQXV0aG9yaXR5MFkwEwYHKoZIzj0CAQYIKoZIzj0DAQcDQgAEpDmJIkUr
VcrzicJb0bykZWLSzLk0iGgkmtHHRlMBTL+V1oeWXgNrUhxRA35rj03vyh60QEzP
T6CIgu7WUZ3suqNmMGQwEgYDVROTAQH/BAGwBgEB/wIBATA0BgNVHQ8BAf8EBAMC
AQYwHQYDVRO0BBYEFGL6gjNZrPqplj4c+hQK3fUE83FgMB8GA1UdIwQYMBaAFGL6
gjNZrPqplj4c+hQK3fUE83FgMAoGCCqGSM49BAMCA0cAMEQCICxUXOTkV9im8NnZ
u+vW70Hd/n+MbZps83UyH8b6xx0EAiBUB3jodDlyUn7t669YaGIgtUB48s10Yqdq
```

```
58u5L/VMiw==
```

```
-----END CERTIFICATE-----
```

```
-> Sample Connectivity Standards Alliance CD Signing Private Key:
```

```
-----BEGIN EC PRIVATE KEY-----
```

```
MHcCAQEEIK7zSEEW6UgexXvgRy30G/SZBk5QJK2GnspeiJgC1IB1oAoGCCqGSM49
```

```
AwEHoUQDQgAEPDmJIkUrVcrzicJb0bykZWLSzLk0iGkkmthHRLMBTL+V1oeWXgNr
```

```
UhxRA35rj03vyh60QEzPT6CIgu7WUZ3sug==
```

```
-----END EC PRIVATE KEY-----
```

```
===== Intermediate outputs =====
```

```
-> Encoded TLV of sample Certification Declaration (54 bytes):
```

```
00000000 15 24 00 01 25 01 f1 ff 36 02 05 00 80 18 25 03 |.$.%....6.....%|
00000010 34 12 2c 04 13 5a 49 47 32 30 31 34 31 5a 42 33 |4.,...ZIG20141ZB3|
00000020 33 30 30 30 31 2d 32 34 24 05 00 24 06 00 25 07 |30001-24$.$.%.|
00000030 94 26 24 08 00 18 |.&$....|
00000036
```

```
===== Algorithm outputs =====
```

```
-> Encoded CMS SignedData of Certification Declaration (235 bytes):
```

```
00000000 30 81 e8 06 09 2a 86 48 86 f7 0d 01 07 02 a0 81 |0....*.H.....|
00000010 da 30 81 d7 02 01 03 31 0d 30 0b 06 09 60 86 48 |.0.....1.0...`.H|
00000020 01 65 03 04 02 01 30 45 06 09 2a 86 48 86 f7 0d |.e....0E...*.H...|
00000030 01 07 01 a0 38 04 36 15 24 00 01 25 01 f1 ff 36 |...8.6$.%....6|
00000040 02 05 00 80 18 25 03 34 12 2c 04 13 5a 49 47 32 |...%.4.,...ZIG2|
00000050 30 31 34 31 5a 42 33 33 30 30 30 31 2d 32 34 24 |0141ZB330001-24$|
00000060 05 00 24 06 00 25 07 94 26 24 08 00 18 31 7c 30 |..$.%.&$....1|0|
00000070 7a 02 01 03 80 14 62 fa 82 33 59 ac fa a9 96 3e |z....b..3Y....>|
00000080 1c fa 14 0a dd f5 04 f3 71 60 30 0b 06 09 60 86 |.....q`0...`.|
00000090 48 01 65 03 04 02 01 30 0a 06 08 2a 86 48 ce 3d |H.e....0...*.H.=|
000000a0 04 03 02 04 46 30 44 02 20 43 a6 3f 2b 94 3d f3 |....FOD. C.?+.=.|
000000b0 3c 38 b3 e0 2f ca a7 5f e3 53 2a eb bf 5e 63 f5 |<8../...S*..^c.|
000000c0 bb db c0 b1 f0 1d 3c 4f 60 02 20 4c 1a bf 5f 18 |.....<0`. L..._|
000000d0 07 b8 18 94 b1 57 6c 47 e4 72 4e 4d 96 6c 61 2e |....WlG.rNM.la.|
000000e0 d3 fa 25 c1 18 c3 f2 b3 f9 03 69 |..%......i|
000000eb
```

The second example [Certification Declaration](#) has the following qualities:

- Both [dac_origin_vendor_id](#) and [dac_origin_product_id](#) are present
- The [product_id_array](#) contains a two PIDs (0x8001, 0x8002)
- It uses the [authorized_paa_list](#) to indicate the Subject Key Identifier (SKI) extension value of the expected PAA in the certificate chain of the Device Attestation Certificate for a product carrying

this Certification Declaration

The content of this second example is shown below:

```

===== Algorithm inputs =====
-> format_version = 1
-> vendor_id = 0xFFF2
-> product_id_array = [ 0x8001, 0x8002 ]
-> device_type_id = 0x1234
-> certificate_id = "ZIG20142ZB330002-24"
-> security_level = 0
-> security_information = 0
-> version_number = 0x2694
-> certification_type = 0
-> dac_origin_vendor_id = 0xFFF1
-> dac_origin_product_id = 0x8000
-> authorized_paa_list = [ 78:5c:e7:05:b8:6b:8f:4e:6f:c7:93:aa:60:cb:43:ea:69:68:82:d5
]

-> Sample Connectivity Standards Alliance CD Signing Certificate:
-----BEGIN CERTIFICATE-----
MIIBszCCAVqgAwIBAgIIRdrzneR6oI8wCgYIKoZIzj0EAwIwKzEpMCcGA1UEAwg
TWf0dGVyIFRlc3QgQ0QGU2lnbmLuZyBBdXRob3JpdHkwIBcNMjE4MTQyMzQz
WhgPOTk5OTEyMzEyMzU5NTlaMCsxKTAnBgNVBAMMIE1hdHRlc iBUZXNOIENEIFNp
Z25pbmcgQXV0aG9yaXR5MFkwEwYHKoZIzj0CAQYIKoZIzj0DAQcDQgAEpDmJIkUr
VcrzicJb0bykZWLSzLk0iGkkmthHRLMBTL+V1oeWXgNrUhxRA35rj03vyh60QEZp
T6CIgu7WUZ3suqNmMGQwEgYDVROTAQH/BAGwBgEB/wIBATAOBgNVHQ8BAf8EBAMC
AQYwHQYDVRO0BBYEFGL6gjNZrPqplj4c+hQK3fUE83FgMB8GA1UdIwQYMBaAFGL6
gjNZrPqplj4c+hQK3fUE83FgMAoGCCqGSM49BAMCA0cAMEQCICxUX0TkV9im8NnZ
u+vW70Hd/n+MbZps83UyH8b6xx0EAiBUB3jodDlyUn7t669YaGIgtUB48s10Yqdq
58u5L/VMiw==
-----END CERTIFICATE-----

-> Sample Connectivity Standards Alliance CD Signing Private Key:
-----BEGIN EC PRIVATE KEY-----
MHcCAQEEIK7zSEEW6UgexXvgRy30G/SZBk5QJK2GnspeiJgC1IB1oAoGCCqGSM49
AwEHoUQDQgAEpDmJIkUrVcrzicJb0bykZWLSzLk0iGkkmthHRLMBTL+V1oeWXgNr
UhxRA35rj03vyh60QEZpT6CIgu7WUZ3sug==
-----END EC PRIVATE KEY-----

===== Intermediate outputs =====
-> Encoded TLV of sample Certification Declaration (90 bytes):
00000000 15 24 00 01 25 01 f2 ff 36 02 05 01 80 05 02 80 |.$.%....6.....|

```



```

00000010 18 25 03 34 12 2c 04 13 5a 49 47 32 30 31 34 32 |.%.4.,...ZIG20142|
00000020 5a 42 33 33 30 30 30 32 2d 32 34 24 05 00 24 06 |ZB330002-24$..$.|
00000030 00 25 07 94 26 24 08 00 25 09 f1 ff 25 0a 00 80 |.%.&$..%...%...|
00000040 36 0b 10 14 78 5c e7 05 b8 6b 8f 4e 6f c7 93 aa |6...x\...k.No...|
00000050 60 cb 43 ea 69 68 82 d5 18 18 |`.C.ih....|
0000005a

```

===== Algorithm outputs =====

-> Encoded CMS SignedData of Certification Declaration (273 bytes):

```

00000000 30 82 01 0d 06 09 2a 86 48 86 f7 0d 01 07 02 a0 |0.....*.H.....|
00000010 81 ff 30 81 fc 02 01 03 31 0d 30 0b 06 09 60 86 |..0.....1.0...`|
00000020 48 01 65 03 04 02 01 30 69 06 09 2a 86 48 86 f7 |H.e....0i...*.H..|
00000030 0d 01 07 01 a0 5c 04 5a 15 24 00 01 25 01 f2 ff |.....\Z.$..%...|
00000040 36 02 05 01 80 05 02 80 18 25 03 34 12 2c 04 13 |6.....%.4.,...|
00000050 5a 49 47 32 30 31 34 32 5a 42 33 33 30 30 30 32 |ZIG20142ZB330002|
00000060 2d 32 34 24 05 00 24 06 00 25 07 94 26 24 08 00 |-24$..$.%..&$..|
00000070 25 09 f1 ff 25 0a 00 80 36 0b 10 14 78 5c e7 05 |%...%...6...x\..|
00000080 b8 6b 8f 4e 6f c7 93 aa 60 cb 43 ea 69 68 82 d5 |.k.No...`.C.ih..|
00000090 18 18 31 7d 30 7b 02 01 03 80 14 62 fa 82 33 59 |..1}0{.....b..3Y|
000000a0 ac fa a9 96 3e 1c fa 14 0a dd f5 04 f3 71 60 30 |....>.....q`0|
000000b0 0b 06 09 60 86 48 01 65 03 04 02 01 30 0a 06 08 |...`.H.e....0...|
000000c0 2a 86 48 ce 3d 04 03 02 04 47 30 45 02 20 4a e9 |*.H.=....G0E. J.|
000000d0 c9 b7 f8 aa 68 61 0a dd 84 e4 12 91 fc 8f 4d c5 |....ha.....M.|
000000e0 33 fc a2 9d c1 ff f2 25 3c 09 cd 32 f7 75 02 21 |3.....%<..2.u.!!|
000000f0 00 9c 0a 5f de f9 e0 08 d1 cc 8b b7 c3 95 9c db |..._.....|
00000100 65 c4 61 25 cb 72 95 08 1e 47 b5 c1 31 e4 d1 f4 |e.a%.r...G..1...|
00000110 8c |.|
00000111

```

F.2. Device Attestation Response test vector

This subsection contains a worked example of the [Attestation Information](#) to be generated in the [AttestationResponse Command](#) when executing the [Device Attestation Procedure](#).

The Device Attestation key pair shown is an example, not to be reused in implementations.

NOTE

This test vector does NOT contain the optional [Firmware Information](#) payload. It is omitted.

===== Algorithm inputs =====

-> AttestationNonce (example):

```

e0:42:1b:91:c6:fd:cd:b4:0e:2a:4d:2c:f3:1d:b2:b4:e1:8b:41:1b:1d:3a:d4:d1:2a:9d:90:aa:8e
:52:fa:e2

```

-> Attestation challenge (example): 7a:49:53:05:d0:77:79:a4:94:dd:39:a0:85:1b:66:0d

-> Device attestation private key (example):

38:f3:e0:a1:f1:45:ba:1b:f3:e4:4b:55:2d:ef:65:27:3d:1d:8e:27:6a:a3:14:ac:74:2e:b1:28:93
:3b:a6:4b

-----BEGIN EC PRIVATE KEY-----

MHcCAQEEIDjz4KHxRbob8+RLVS3vZSc9HY4naqMURHQusSiT06ZLoAoGCCqGSM49
AwEHoUQDQgAEzLz477BdTu55DQpx1cARu3RyQNuiffiEXTpjSwr2ZRYzBjqASy/4
XcqyAZoKtvVZV3X+jYX716B8joN9pNWouQ==

-----END EC PRIVATE KEY-----

-> Device attestation public key (example):

04:ce:5c:f8:ef:b0:5d:4e:ee:79:0d:0a:71:d5:c0:11:bb:74:72:40:db:a2:14:58:84:5d:33:e3:4b
:0a:f6:65:16:33:06:3a:80:4b:2f:f8:5d:ca:b2:01:9a:0a:b6:f5:59:57:75:fe:8d:85:fb:d7:a0:7
c:8e:83:7d:a4:d5:a8:b9

-----BEGIN PUBLIC KEY-----

MFkwEwYHKoZIzj0CAQYIKoZIzj0DAQcDQgAEzLz477BdTu55DQpx1cARu3RyQNuiffiEXTpjSwr2ZRYzBjqASy/4XcqyAZoKtvVZV3X+jYX716B8joN9pNWouQ==

-----END PUBLIC KEY-----

-> Desired timestamp: 2021-06-15T20:15:57Z

-> Desired timestamp in epoch-s: 677103357

-> vendor specific [0xfff1:0x3e:0x1] =

73:61:6d:70:6c:65:5f:76:65:6e:64:6f:72:5f:72:65:73:65:72:76:65:64:31

("sample_vendor_reserved1")

-> vendor specific [0xfff1:0x3e:0x3] =

76:65:6e:64:6f:72:5f:72:65:73:65:72:76:65:64:33:5f:65:78:61:6d:70:6c:65

("vendor_reserved3_example")

===== Intermediate outputs =====

-> attestation_elements_message:

```
00000000 15 31 01 11 01 30 82 01 0d 06 09 2a 86 48 86 f7 |.1...0.....*.H..|
00000010 0d 01 07 02 a0 81 ff 30 81 fc 02 01 03 31 0d 30 |.....0.....1.0|
00000020 0b 06 09 60 86 48 01 65 03 04 02 01 30 69 06 09 |...`.H.e....0i..|
00000030 2a 86 48 86 f7 0d 01 07 01 a0 5c 04 5a 15 24 00 |*.H.....\Z.$.|
00000040 01 25 01 f2 ff 36 02 05 01 80 05 02 80 18 25 03 |.%.6.....%.|
00000050 34 12 2c 04 13 5a 49 47 32 30 31 34 32 5a 42 33 |4.,...ZIG20142ZB3|
00000060 33 30 30 30 32 2d 32 34 24 05 00 24 06 00 25 07 |30002-24$.$.%.|
00000070 94 26 24 08 00 25 09 f1 ff 25 0a 00 80 36 0b 10 |.&$.%.%.6..|
00000080 14 78 5c e7 05 b8 6b 8f 4e 6f c7 93 aa 60 cb 43 |.x\...k.No...`.C|
00000090 ea 69 68 82 d5 18 18 31 7d 30 7b 02 01 03 80 14 |.ih....1}0{.....|
000000a0 62 fa 82 33 59 ac fa a9 96 3e 1c fa 14 0a dd f5 |b..3Y....>.....|
000000b0 04 f3 71 60 30 0b 06 09 60 86 48 01 65 03 04 02 |..q`0...`.H.e...|
000000c0 01 30 0a 06 08 2a 86 48 ce 3d 04 03 02 04 47 30 |.0...*.H.=....G0|
```

```

000000d0 45 02 20 4a e9 c9 b7 f8 aa 68 61 0a dd 84 e4 12 |E. J.....ha.....|
000000e0 91 fc 8f 4d c5 33 fc a2 9d c1 ff f2 25 3c 09 cd |...M.3.....%<..|
000000f0 32 f7 75 02 21 00 9c 0a 5f de f9 e0 08 d1 cc 8b |2.u.!..._.....|
00000100 b7 c3 95 9c db 65 c4 61 25 cb 72 95 08 1e 47 b5 |.....e.a%.r...G.|
00000110 c1 31 e4 d1 f4 8c 30 02 20 e0 42 1b 91 c6 fd cd |.1....0. .B.....|
00000120 b4 0e 2a 4d 2c f3 1d b2 b4 e1 8b 41 1b 1d 3a d4 |..*M,.....A...|
00000130 d1 2a 9d 90 aa 8e 52 fa e2 26 03 fd c6 5b 28 d0 |.*....R.&...[(.|
00000140 f1 ff 3e 00 01 00 17 73 61 6d 70 6c 65 5f 76 65 |..>....sample_ve|
00000150 6e 64 6f 72 5f 72 65 73 65 72 76 65 64 31 d0 f1 |ndor_reserved1..|
00000160 ff 3e 00 03 00 18 76 65 6e 64 6f 72 5f 72 65 73 |.>....vendor_res|
00000170 65 72 76 65 64 33 5f 65 78 61 6d 70 6c 65 18 |erved3_example.|
0000017f

```

-> attestation_tbs := attestation_elements_message || attestation_challenge

-> attestation_tbs (NOT sent over the wire):

```

00000000 15 31 01 11 01 30 82 01 0d 06 09 2a 86 48 86 f7 |.1...0.....*.H..|
00000010 0d 01 07 02 a0 81 ff 30 81 fc 02 01 03 31 0d 30 |.....0.....1.0|
00000020 0b 06 09 60 86 48 01 65 03 04 02 01 30 69 06 09 |...`.H.e....0i..|
00000030 2a 86 48 86 f7 0d 01 07 01 a0 5c 04 5a 15 24 00 |*.H.....\Z.$.|
00000040 01 25 01 f2 ff 36 02 05 01 80 05 02 80 18 25 03 |.%...6.....%.|
00000050 34 12 2c 04 13 5a 49 47 32 30 31 34 32 5a 42 33 |4.,..ZIG20142ZB3|
00000060 33 30 30 30 32 2d 32 34 24 05 00 24 06 00 25 07 |30002-24$..$.%.|
00000070 94 26 24 08 00 25 09 f1 ff 25 0a 00 80 36 0b 10 |.&$%.%...%.6..|
00000080 14 78 5c e7 05 b8 6b 8f 4e 6f c7 93 aa 60 cb 43 |.x\...k.No...`.C|
00000090 ea 69 68 82 d5 18 18 31 7d 30 7b 02 01 03 80 14 |.ih....1}0{.....|
000000a0 62 fa 82 33 59 ac fa a9 96 3e 1c fa 14 0a dd f5 |b..3Y....>.....|
000000b0 04 f3 71 60 30 0b 06 09 60 86 48 01 65 03 04 02 |..q`0...`.H.e...|
000000c0 01 30 0a 06 08 2a 86 48 ce 3d 04 03 02 04 47 30 |.0...*.H.=....G0|
000000d0 45 02 20 4a e9 c9 b7 f8 aa 68 61 0a dd 84 e4 12 |E. J.....ha.....|
000000e0 91 fc 8f 4d c5 33 fc a2 9d c1 ff f2 25 3c 09 cd |...M.3.....%<..|
000000f0 32 f7 75 02 21 00 9c 0a 5f de f9 e0 08 d1 cc 8b |2.u.!..._.....|
00000100 b7 c3 95 9c db 65 c4 61 25 cb 72 95 08 1e 47 b5 |.....e.a%.r...G.|
00000110 c1 31 e4 d1 f4 8c 30 02 20 e0 42 1b 91 c6 fd cd |.1....0. .B.....|
00000120 b4 0e 2a 4d 2c f3 1d b2 b4 e1 8b 41 1b 1d 3a d4 |..*M,.....A...|
00000130 d1 2a 9d 90 aa 8e 52 fa e2 26 03 fd c6 5b 28 d0 |.*....R.&...[(.|
00000140 f1 ff 3e 00 01 00 17 73 61 6d 70 6c 65 5f 76 65 |..>....sample_ve|
00000150 6e 64 6f 72 5f 72 65 73 65 72 76 65 64 31 d0 f1 |ndor_reserved1..|
00000160 ff 3e 00 03 00 18 76 65 6e 64 6f 72 5f 72 65 73 |.>....vendor_res|
00000170 65 72 76 65 64 33 5f 65 78 61 6d 70 6c 65 18 7a |erved3_example.z|
00000180 49 53 05 d0 77 79 a4 94 dd 39 a0 85 1b 66 0d |IS..wy...9...f.|
0000018f

```

-> SHA-256 of attestation_tbs used for signature (NOT sent over the wire):

```
1d:f1:05:b1:30:84:c3:cc:13:19:9e:df:07:b8:76:9e:be:2e:26:0d:84:8f:27:a6:ca:b6:6d:d9:a5
:8c:ea:b1
```

-> Fixed K for sample signature of attestation_tbs:

```
c5:35:83:3f:47:86:4f:cb:d8:b5:e3:2e:fb:a8:84:35:c0:fb:0c:9f:db:0f:00:34:98:0a:41:84:cc
:f0:52:4d
```

-> Attestation signature:

```
79:82:53:5d:24:cf:e1:4a:71:ab:04:24:cf:0b:ac:f1:e3:45:48:7e:d5:0f:1a:c0:bc:25:9e:cc:fb
:39:08:1e:d7:a7:52:18:8d:9f:76:f9:06:37:03:eb:24:0f:9c:d1:4b:0a:43:e7:41:fe:60:ef:2a:8
1:63:5a:ea:5b:48:4d
```

==== Algorithm outputs ====

-> AttestationElements field of AttestationResponse (len 383 bytes):

```
00000000 15 31 01 11 01 30 82 01 0d 06 09 2a 86 48 86 f7 |.1...0.....*.H..|
00000010 0d 01 07 02 a0 81 ff 30 81 fc 02 01 03 31 0d 30 |.....0.....1.0|
00000020 0b 06 09 60 86 48 01 65 03 04 02 01 30 69 06 09 |...`.H.e....0i..|
00000030 2a 86 48 86 f7 0d 01 07 01 a0 5c 04 5a 15 24 00 |*.H.....\Z.$.|
00000040 01 25 01 f2 ff 36 02 05 01 80 05 02 80 18 25 03 |.%.6.....%|.
00000050 34 12 2c 04 13 5a 49 47 32 30 31 34 32 5a 42 33 |4.,...ZIG20142ZB3|
00000060 33 30 30 30 32 2d 32 34 24 05 00 24 06 00 25 07 |30002-24$..$.%|.
00000070 94 26 24 08 00 25 09 f1 ff 25 0a 00 80 36 0b 10 |.&$.%...%...6..|
00000080 14 78 5c e7 05 b8 6b 8f 4e 6f c7 93 aa 60 cb 43 |.x\...k.No...`.C|
00000090 ea 69 68 82 d5 18 18 31 7d 30 7b 02 01 03 80 14 |.ih....1}0{.....|
000000a0 62 fa 82 33 59 ac fa a9 96 3e 1c fa 14 0a dd f5 |b..3Y....>.....|
000000b0 04 f3 71 60 30 0b 06 09 60 86 48 01 65 03 04 02 |..q`0...`.H.e...|
000000c0 01 30 0a 06 08 2a 86 48 ce 3d 04 03 02 04 47 30 |.0...*.H.=....G0|
000000d0 45 02 20 4a e9 c9 b7 f8 aa 68 61 0a dd 84 e4 12 |E. J.....ha.....|
000000e0 91 fc 8f 4d c5 33 fc a2 9d c1 ff f2 25 3c 09 cd |...M.3.....%<..|
000000f0 32 f7 75 02 21 00 9c 0a 5f de f9 e0 08 d1 cc 8b |2.u.!..._.....|
00000100 b7 c3 95 9c db 65 c4 61 25 cb 72 95 08 1e 47 b5 |.....e.a%.r...G.|
00000110 c1 31 e4 d1 f4 8c 30 02 20 e0 42 1b 91 c6 fd cd |.1...0. .B.....|
00000120 b4 0e 2a 4d 2c f3 1d b2 b4 e1 8b 41 1b 1d 3a d4 |..*M,.....A...|
00000130 d1 2a 9d 90 aa 8e 52 fa e2 26 03 fd c6 5b 28 d0 |.*....R.&...[(.|
00000140 f1 ff 3e 00 01 00 17 73 61 6d 70 6c 65 5f 76 65 |..>....sample_ve|
00000150 6e 64 6f 72 5f 72 65 73 65 72 76 65 64 31 d0 f1 |ndor_reserved1..|
00000160 ff 3e 00 03 00 18 76 65 6e 64 6f 72 5f 72 65 73 |.>....vendor_res|
00000170 65 72 76 65 64 33 5f 65 78 61 6d 70 6c 65 18 |erved3_example.|
0000017f
```

-> AttestationSignature field of AttestationResponse (len 64 bytes):

```
00000000 79 82 53 5d 24 cf e1 4a 71 ab 04 24 cf 0b ac f1 |y.S]$.Jq.$.$.|
00000010 e3 45 48 7e d5 0f 1a c0 bc 25 9e cc fb 39 08 1e |.EH~.....%...9..|
00000020 d7 a7 52 18 8d 9f 76 f9 06 37 03 eb 24 0f 9c d1 |..R...v...7..$.|
```

```
00000030 4b 0a 43 e7 41 fe 60 ef 2a 81 63 5a ea 5b 48 4d |K.C.A.`*.cZ.[HM|
00000040
```

F.3. Node Operational CSR Response test vector

This subsection contains a worked example of the [NOCSR Information](#) to be generated in the [CSR-Response Command](#) when executing the [Node Operational CSR Procedure](#).

The CSR shown is valid for the provided Node Operational public key.

The Device Attestation key pair shown is an example, not to be reused in implementations.

```
===== Algorithm inputs =====
```

```
-> CSRNonce:
```

```
81:4a:4d:4c:1c:4a:8e:bb:ea:db:0a:e2:82:f9:91:eb:13:ac:5f:9f:ce:94:30:93:19:aa:94:09:6c
:8c:d4:b8
```

```
-> Attestation challenge (example): 7a:49:53:05:d0:77:79:a4:94:dd:39:a0:85:1b:66:0d
```

```
-> Device attestation private key (example):
```

```
38:f3:e0:a1:f1:45:ba:1b:f3:e4:4b:55:2d:ef:65:27:3d:1d:8e:27:6a:a3:14:ac:74:2e:b1:28:93
:3b:a6:4b
```

```
-----BEGIN EC PRIVATE KEY-----
```

```
MHcCAQEEIDjz4KHxRbob8+RLVS3vZSc9HY4naqMURHQusSiT06ZLoAoGCCqGSM49
```

```
AwEHoUQDQgAEzLz477BdTu55DQpx1cARu3RyQNu iFFiEXTPjSwr2ZRYzBjqASy/4
```

```
XcqyAZoKtvVZV3X+jYX716B8joN9pNWouQ==
```

```
-----END EC PRIVATE KEY-----
```

```
-> Device attestation public key (example):
```

```
04:ce:5c:f8:ef:b0:5d:4e:ee:79:0d:0a:71:d5:c0:11:bb:74:72:40:db:a2:14:58:84:5d:33:e3:4b
:0a:f6:65:16:33:06:3a:80:4b:2f:f8:5d:ca:b2:01:9a:0a:b6:f5:59:57:75:fe:8d:85:fb:d7:a0:7
c:8e:83:7d:a4:d5:a8:b9
```

```
-----BEGIN PUBLIC KEY-----
```

```
MFkwEwYHKoZIzj0CAQYIKoZIzj0DAQcDQgAEzLz477BdTu55DQpx1cARu3RyQNu i
```

```
FFiEXTPjSwr2ZRYzBjqASy/4XcqyAZoKtvVZV3X+jYX716B8joN9pNWouQ==
```

```
-----END PUBLIC KEY-----
```

```
===== Intermediate outputs =====
```

```
-> Candidate Operational Private Key:
```

```
1c:18:82:e8:7f:80:d8:1a:25:9a:62:b6:ea:02:db:08:17:e2:10:68:46:84:2b:eb:3a:ab:c2:53:86
:a9:1e:89
```

```
-----BEGIN EC PRIVATE KEY-----
```

```
MHcCAQEEIBwYguh/gNgajZpituoC2wgX4hBoRoQr6zqrwLOGqR6JoAoGCCqGSM49
```

```
AwEHoUQDQgAEXKJ542aCwtRs59TPiWeEZwi1ufhbnNr9jKiFJhLLDwx6cTF0yNyc
```

```
LjTd7v7p9j80i9faz802pFMqrdiallHNbg==
```

```
-----END EC PRIVATE KEY-----
```

-> Candidate Operational Public Key:

04:5c:a2:79:e3:66:82:c2:d4:6c:e7:d4:cf:89:67:84:67:08:b5:b9:f8:5b:9c:da:fd:8c:a8:85:26
:12:cb:0f:0c:7a:71:31:4e:c8:dc:9c:96:34:dd:ee:fe:e9:f6:3f:0e:8b:d7:da:cf:c3:b6:a4:53:2
a:ad:d8:9a:96:51:cd:6e

-----BEGIN PUBLIC KEY-----

MFkwEwYHKoZIzj0CAQYIKoZIzj0DAQcDQgAEEXKJ542aCwtRs59TPiWeEZwiufhb
nNr9jKiFJhLLDwx6cTF0yNycljTd7v7p9j80i9faz802pFMqrdiallHNbg==

-----END PUBLIC KEY-----

Certificate Request:

Data:

Version: 1 (0x0)

Subject: 0 = CSA

Subject Public Key Info:

Public Key Algorithm: id-ecPublicKey

Public-Key: (256 bit)

pub:

04:5c:a2:79:e3:66:82:c2:d4:6c:e7:d4:cf:89:67:
84:67:08:b5:b9:f8:5b:9c:da:fd:8c:a8:85:26:12:
cb:0f:0c:7a:71:31:4e:c8:dc:9c:96:34:dd:ee:fe:
e9:f6:3f:0e:8b:d7:da:cf:c3:b6:a4:53:2a:ad:d8:
9a:96:51:cd:6e

ASN1 OID: prime256v1

NIST CURVE: P-256

Attributes:

Requested Extensions:

Signature Algorithm: ecdsa-with-SHA256

30:45:02:20:0e:67:5e:e1:b3:bb:fe:15:2a:17:4a:f5:35:e2:
2d:55:ce:10:c1:50:ca:c0:1b:31:18:de:05:e8:fd:9f:10:48:
02:21:00:d8:8c:57:cc:6e:74:f0:e5:48:8a:26:16:7a:07:fd:
6d:be:f1:aa:ad:72:1c:58:0b:6e:ae:21:be:5e:6d:0c:72

-> CSR bytes DER:

```
00000000 30 81 da 30 81 81 02 01 00 30 0e 31 0c 30 0a 06 |0..0....0.1.0..|
00000010 03 55 04 0a 0c 03 43 53 41 30 59 30 13 06 07 2a |.U....CSA0Y0...*|
00000020 86 48 ce 3d 02 01 06 08 2a 86 48 ce 3d 03 01 07 |.H.=....*.H.=...|
00000030 03 42 00 04 5c a2 79 e3 66 82 c2 d4 6c e7 d4 cf |.B..\.y.f...l...|
00000040 89 67 84 67 08 b5 b9 f8 5b 9c da fd 8c a8 85 26 |.g.g....[.....&|
00000050 12 cb 0f 0c 7a 71 31 4e c8 dc 9c 96 34 dd ee fe |....zq1N....4...|
00000060 e9 f6 3f 0e 8b d7 da cf c3 b6 a4 53 2a ad d8 9a |..?.....S*...|
00000070 96 51 cd 6e a0 11 30 0f 06 09 2a 86 48 86 f7 0d |.Q.n..0...*.H...|
00000080 01 09 0e 31 02 30 00 30 0a 06 08 2a 86 48 ce 3d |...1.0.0...*.H.=|
00000090 04 03 02 03 48 00 30 45 02 20 0e 67 5e e1 b3 bb |....H.OE. .g^...|
```

```

000000a0 fe 15 2a 17 4a f5 35 e2 2d 55 ce 10 c1 50 ca c0 |..*.J.5.-U...P..|
000000b0 1b 31 18 de 05 e8 fd 9f 10 48 02 21 00 d8 8c 57 |.1.....H.!...W|
000000c0 cc 6e 74 f0 e5 48 8a 26 16 7a 07 fd 6d be f1 aa |.nt..H.&.z...m...|
000000d0 ad 72 1c 58 0b 6e ae 21 be 5e 6d 0c 72 |.r.X.n.!.^m.r|
000000dd

```

-> Sample vendor_reserved1:

73:61:6d:70:6c:65:5f:76:65:6e:64:6f:72:5f:72:65:73:65:72:76:65:64:31

-> Sample vendor_reserved3:

76:65:6e:64:6f:72:5f:72:65:73:65:72:76:65:64:33:5f:65:78:61:6d:70:6c:65

-> nocsr_elements_message:

```

00000000 15 30 01 dd 30 81 da 30 81 81 02 01 00 30 0e 31 |.0..0..0.....0.1|
00000010 0c 30 0a 06 03 55 04 0a 0c 03 43 53 41 30 59 30 |.0...U....CSA0Y0|
00000020 13 06 07 2a 86 48 ce 3d 02 01 06 08 2a 86 48 ce |...*.H.=....*.H.|
00000030 3d 03 01 07 03 42 00 04 5c a2 79 e3 66 82 c2 d4 |=....B..\y.f...|
00000040 6c e7 d4 cf 89 67 84 67 08 b5 b9 f8 5b 9c da fd |l....g.g....[...|
00000050 8c a8 85 26 12 cb 0f 0c 7a 71 31 4e c8 dc 9c 96 |...&....zq1N....|
00000060 34 dd ee fe e9 f6 3f 0e 8b d7 da cf c3 b6 a4 53 |4.....?.....S|
00000070 2a ad d8 9a 96 51 cd 6e a0 11 30 0f 06 09 2a 86 |*....Q.n..0...*.|
00000080 48 86 f7 0d 01 09 0e 31 02 30 00 30 0a 06 08 2a |H.....1.0.0...*.|
00000090 86 48 ce 3d 04 03 02 03 48 00 30 45 02 20 0e 67 |.H.=....H.OE. .g|
000000a0 5e e1 b3 bb fe 15 2a 17 4a f5 35 e2 2d 55 ce 10 |^.....*.J.5.-U..|
000000b0 c1 50 ca c0 1b 31 18 de 05 e8 fd 9f 10 48 02 21 |.P...1.....H.!|
000000c0 00 d8 8c 57 cc 6e 74 f0 e5 48 8a 26 16 7a 07 fd |...W.nt..H.&.z..|
000000d0 6d be f1 aa ad 72 1c 58 0b 6e ae 21 be 5e 6d 0c |m....r.X.n.!.^m.|
000000e0 72 30 02 20 81 4a 4d 4c 1c 4a 8e bb ea db 0a e2 |r0. .JML.J.....|
000000f0 82 f9 91 eb 13 ac 5f 9f ce 94 30 93 19 aa 94 09 |....._...0.....|
00000100 6c 8c d4 b8 30 03 17 73 61 6d 70 6c 65 5f 76 65 |l...0..sample_ve|
00000110 6e 64 6f 72 5f 72 65 73 65 72 76 65 64 31 30 05 |ndor_reserved10.|
00000120 18 76 65 6e 64 6f 72 5f 72 65 73 65 72 76 65 64 |.vendor_reserved|
00000130 33 5f 65 78 61 6d 70 6c 65 18 |3_example.|
0000013a

```

-> nocsr_tbs (NOT sent over the wire):

```

00000000 15 30 01 dd 30 81 da 30 81 81 02 01 00 30 0e 31 |.0..0..0.....0.1|
00000010 0c 30 0a 06 03 55 04 0a 0c 03 43 53 41 30 59 30 |.0...U....CSA0Y0|
00000020 13 06 07 2a 86 48 ce 3d 02 01 06 08 2a 86 48 ce |...*.H.=....*.H.|
00000030 3d 03 01 07 03 42 00 04 5c a2 79 e3 66 82 c2 d4 |=....B..\y.f...|
00000040 6c e7 d4 cf 89 67 84 67 08 b5 b9 f8 5b 9c da fd |l....g.g....[...|
00000050 8c a8 85 26 12 cb 0f 0c 7a 71 31 4e c8 dc 9c 96 |...&....zq1N....|
00000060 34 dd ee fe e9 f6 3f 0e 8b d7 da cf c3 b6 a4 53 |4.....?.....S|
00000070 2a ad d8 9a 96 51 cd 6e a0 11 30 0f 06 09 2a 86 |*....Q.n..0...*.|
00000080 48 86 f7 0d 01 09 0e 31 02 30 00 30 0a 06 08 2a |H.....1.0.0...*.|

```



```

00000090 86 48 ce 3d 04 03 02 03 48 00 30 45 02 20 0e 67 |.H.=...H.OE. .g|
000000a0 5e e1 b3 bb fe 15 2a 17 4a f5 35 e2 2d 55 ce 10 |^.....*.J.5.-U..|
000000b0 c1 50 ca c0 1b 31 18 de 05 e8 fd 9f 10 48 02 21 |.P...1.....H.!!|
000000c0 00 d8 8c 57 cc 6e 74 f0 e5 48 8a 26 16 7a 07 fd |...W.nt..H.&.z..|
000000d0 6d be f1 aa ad 72 1c 58 0b 6e ae 21 be 5e 6d 0c |m....r.X.n.!.^m.|
000000e0 72 30 02 20 81 4a 4d 4c 1c 4a 8e bb ea db 0a e2 |r0. .JML.J.....|
000000f0 82 f9 91 eb 13 ac 5f 9f ce 94 30 93 19 aa 94 09 |....._...0.....|
00000100 6c 8c d4 b8 30 03 17 73 61 6d 70 6c 65 5f 76 65 |l...0..sample_ve|
00000110 6e 64 6f 72 5f 72 65 73 65 72 76 65 64 31 30 05 |ndor_reserved10.|
00000120 18 76 65 6e 64 6f 72 5f 72 65 73 65 72 76 65 64 |.vendor_reserved|
00000130 33 5f 65 78 61 6d 70 6c 65 18 7a 49 53 05 d0 77 |3_example.zIS..w|
00000140 79 a4 94 dd 39 a0 85 1b 66 0d |y...9...f.|
0000014a

```

-> SHA-256 of nocsr_tbs used for signature (NOT sent over the wire):

```
e2:62:65:69:65:2b:49:e1:5b:6e:d5:b2:42:92:bf:28:e8:e0:e9:5d:e4:25:14:e1:03:a4:30:30:18
:16:cf:3f
```

-> Fixed K for sample signature of nocsr_tbs:

```
a9:c0:d7:f2:b5:1f:51:e3:75:05:3d:c7:0e:53:f5:4e:b1:86:59:c7:d2:99:47:94:f6:8d:b5:08:bb
:53:05:5f
```

-> Attestation signature:

```
87:8e:46:cf:fa:83:c8:32:96:eb:27:2e:bc:37:1c:1f:ef:ee:6d:69:54:f3:78:9f:d3:d2:27:e1:64
:13:d3:d4:75:a6:2f:d0:12:b9:19:d9:95:8b:c7:3d:7c:63:b3:cc:1e:f2:b6:2c:18:e0:cc:10:2e:d
1:ba:4d:ac:85:fe:ea
```

===== Algorithm outputs =====

-> NOCSRElements field of CSRResponse (len 314 bytes):

```

00000000 15 30 01 dd 30 81 da 30 81 81 02 01 00 30 0e 31 |.0..0..0.....0.1|
00000010 0c 30 0a 06 03 55 04 0a 0c 03 43 53 41 30 59 30 |.0...U....CSA0Y0|
00000020 13 06 07 2a 86 48 ce 3d 02 01 06 08 2a 86 48 ce |...*.H.=...*.H.|
00000030 3d 03 01 07 03 42 00 04 5c a2 79 e3 66 82 c2 d4 |=....B..\.y.f...|
00000040 6c e7 d4 cf 89 67 84 67 08 b5 b9 f8 5b 9c da fd |l....g.g....[...|
00000050 8c a8 85 26 12 cb 0f 0c 7a 71 31 4e c8 dc 9c 96 |...&....zq1N....|
00000060 34 dd ee fe e9 f6 3f 0e 8b d7 da cf c3 b6 a4 53 |4.....?.....S|
00000070 2a ad d8 9a 96 51 cd 6e a0 11 30 0f 06 09 2a 86 |*....Q.n..0...*.|
00000080 48 86 f7 0d 01 09 0e 31 02 30 00 30 0a 06 08 2a |H.....1.0.0...*|
00000090 86 48 ce 3d 04 03 02 03 48 00 30 45 02 20 0e 67 |.H.=...H.OE. .g|
000000a0 5e e1 b3 bb fe 15 2a 17 4a f5 35 e2 2d 55 ce 10 |^.....*.J.5.-U..|
000000b0 c1 50 ca c0 1b 31 18 de 05 e8 fd 9f 10 48 02 21 |.P...1.....H.!!|
000000c0 00 d8 8c 57 cc 6e 74 f0 e5 48 8a 26 16 7a 07 fd |...W.nt..H.&.z..|
000000d0 6d be f1 aa ad 72 1c 58 0b 6e ae 21 be 5e 6d 0c |m....r.X.n.!.^m.|
000000e0 72 30 02 20 81 4a 4d 4c 1c 4a 8e bb ea db 0a e2 |r0. .JML.J.....|

```



```

000000f0 82 f9 91 eb 13 ac 5f 9f ce 94 30 93 19 aa 94 09 | ....._...0.....|
00000100 6c 8c d4 b8 30 03 17 73 61 6d 70 6c 65 5f 76 65 | l...0..sample_ve|
00000110 6e 64 6f 72 5f 72 65 73 65 72 76 65 64 31 30 05 | ndor_reserved10.|
00000120 18 76 65 6e 64 6f 72 5f 72 65 73 65 72 76 65 64 | .vendor_reserved|
00000130 33 5f 65 78 61 6d 70 6c 65 18 |3_example.|
0000013a
-> AttestationSignature field of CSRResponse (len 64 bytes):
00000000 87 8e 46 cf fa 83 c8 32 96 eb 27 2e bc 37 1c 1f | ..F....2..'..7..|
00000010 ef ee 6d 69 54 f3 78 9f d3 d2 27 e1 64 13 d3 d4 | ..miT.x...'..d...|
00000020 75 a6 2f d0 12 b9 19 d9 95 8b c7 3d 7c 63 b3 cc | u./.....=|c..|
00000030 1e f2 b6 2c 18 e0 cc 10 2e d1 ba 4d ac 85 fe ea | ...,.....M....|
00000040

```

F.4. Check-In Protocol test vectors

This subsection contains worked examples of the [Check-In Protocol encryption](#) and [decryption](#).

Test 1

===== Encryption algorithm inputs =====

```

-> Symmetric Key: d9:0e:13:18:0d:00:ba:ad:d2:0c:f5:ed:49:13:d3:ff
-> Server counter: 0c:00:00:00
-> Application data : ''

```

===== Intermediate outputs =====

```

-> generated Nonce: 45:80:d2:c6:f1:31:0d:c4:eb:64:f1:f8:e8
-> Ciphertext: bd:c2:1f:b5
-> Tag: 19:5d:74:7d:d2:87:9b:2b:0d:43:ce:5b:1c:56:50:78

```

===== Encryption algorithm outputs =====

```

-> Check-in message payload :
45:80:d2:c6:f1:31:0d:c4:eb:64:f1:f8:e8:bd:c2:1f:b5:19:5d:74:7d:d2:87:9b:2b:0d:43:ce:5b:1c:56:50:78

```

===== Decryption algorithm inputs =====

```

-> Check-in message payload :
45:80:d2:c6:f1:31:0d:c4:eb:64:f1:f8:e8:bd:c2:1f:b5:19:5d:74:7d:d2:87:9b:2b:0d:43:ce:5b:1c:56:50:78
-> Symmetric Key: d9:0e:13:18:0d:00:ba:ad:d2:0c:f5:ed:49:13:d3:ff
-> Client counter: 0b:00:00:00

```

===== Decryption algorithm outputs =====

```

-> Received counter: 0c:00:00:00

```

```
-> Application data : ''  
-> Valid Check-In message
```

Test 2

===== Encryption algorithm inputs =====

```
-> Symmetric Key: 18:fd:bc:ea:ef:01:95:5b:0e:c8:75:ed:a3:ae:6e:e8  
-> Server counter: 0f:00:00:00  
-> Application data : 'This'
```

===== Intermediate outputs =====

```
-> generated Nonce: 9b:02:ed:21:ee:0c:7b:49:19:85:50:2e:37  
-> Ciphertext: 2d:bd:7b:3f:8b:4f:8e:3c  
-> Tag: 5a:d9:94:19:38:9f:41:a8:d6:09:93:8c:67:a8:6d:65
```

===== Encryption algorithm outputs =====

```
-> Check-in message payload :  
9b:02:ed:21:ee:0c:7b:49:19:85:50:2e:37:2d:bd:7b:3f:8b:4f:8e:3c:5a:d9:94:19:38:9f:41:a8  
:d6:09:93:8c:67:a8:6d:65
```

===== Decryption algorithm inputs =====

```
-> Check-in message payload :  
9b:02:ed:21:ee:0c:7b:49:19:85:50:2e:37:2d:bd:7b:3f:8b:4f:8e:3c:5a:d9:94:19:38:9f:41:a8  
:d6:09:93:8c:67:a8:6d:65  
-> Symmetric Key: 18:fd:bc:ea:ef:01:95:5b:0e:c8:75:ed:a3:ae:6e:e8  
-> Client counter: 0b:00:00:00
```

===== Decryption algorithm outputs =====

```
-> Received counter: 0f:00:00:00  
-> Application data : 'This'  
-> Valid Check-In message
```

Test 3

===== Encryption algorithm inputs =====

```
-> Symmetric Key: d9:0e:13:18:0d:00:ba:ad:d2:0c:f5:ed:49:13:d3:ff  
-> Server counter: 0b:00:00:00  
-> Application data : 'This is a'
```

===== Intermediate outputs =====

```
-> generated Nonce: aa:84:bc:60:88:6a:63:a8:47:5d:5d:be:b5  
-> Ciphertext: 6d:63:5f:a9:52:85:ae:33:62:66:13:c7:63
```

-> Tag: 6c:e3:e3:b2:a8:b1:3a:8c:89:be:f7:68:91:e8:e2:96

===== Encryption algorithm outputs =====

-> Check-in message payload :

aa:84:bc:60:88:6a:63:a8:47:5d:5d:be:b5:6d:63:5f:a9:52:85:ae:33:62:66:13:c7:63:6c:e3:e3
:b2:a8:b1:3a:8c:89:be:f7:68:91:e8:e2:96

===== Decryption algorithm inputs =====

-> Check-in message payload :

aa:84:bc:60:88:6a:63:a8:47:5d:5d:be:b5:6d:63:5f:a9:52:85:ae:33:62:66:13:c7:63:6c:e3:e3
:b2:a8:b1:3a:8c:89:be:f7:68:91:e8:e2:96

-> Symmetric Key: d9:0e:13:18:0d:00:ba:ad:d2:0c:f5:ed:49:13:d3:ff

-> Client counter: 0b:00:00:00

===== Decryption algorithm outputs =====

-> Received counter: 0b:00:00:00

-> Application data : 'This is a'

-> Invalid Check-In message - Received counter has already been used

Test 4

===== Encryption algorithm inputs =====

-> Symmetric Key: ca:67:d4:1f:f7:11:29:10:fd:d1:8a:1b:f9:9e:a9:74

-> Server counter: 0b:00:00:00

-> Application data : 'This is a longer'

===== Intermediate outputs =====

-> generated Nonce: 7a:97:72:24:3c:97:c8:7d:5f:3a:31:c4:e6

-> Ciphertext: db:bc:1a:a5:66:c4:43:c2:05:86:06:6b:42:7b:fc:aa:ad:78:da:4a

-> Tag: 10:5a:13:42:ad:bf:3f:47:98:cd:81:b9:ef:97:bb:b7

===== Encryption algorithm outputs =====

-> Check-in message payload :

7a:97:72:24:3c:97:c8:7d:5f:3a:31:c4:e6:db:bc:1a:a5:66:c4:43:c2:05:86:06:6b:42:7b:fc:aa
:ad:78:da:4a:10:5a:13:42:ad:bf:3f:47:98:cd:81:b9:ef:97:bb:b7

===== Decryption algorithm inputs =====

-> Check-in message payload

:7a:97:72:24:3c:97:c8:7d:5f:3a:31:c4:e6:db:bc:1a:a5:66:c4:43:c2:05:86:06:6b:42:7b:fc:a
a:ad:78:da:4a:10:5a:13:42:ad:bf:3f:47:98:cd:81:b9:ef:97:bb:b7

-> Symmetric Key: ca:67:d4:1f:f7:11:29:10:fd:d1:8a:1b:f9:9e:a9:74

-> Client counter: 0f:00:00:00

```
===== Decryption algorithm outputs =====
-> Received counter: 0b:00:00:00
-> Application data : 'This is a longer'
-> Invalid Check-In message - Received counter has already been used
```

Test 5

This test only had decryption processing because it tests the nonce validation that only applies to the decryption process.

```
===== Decryption algorithm inputs =====
-> Check-in message payload :
f9:34:67:6e:a6:e0:70:7b:7a:d7:81:4f:f8:2e:5b:18:d1:9a:23:b2:e4:fa:df:82:92:53:51:7f:f3
:c9:1d:8d:47:84:31:5a:1e:32:08:b8:ec:f6:11:8b:02:1a:5a:4c:d4:e9:d4:13:8d:ff:29:71
-> Symmetric Key: ca:67:d4:1f:f7:11:29:10:fd:d1:8a:1b:f9:9e:a9:74
-> Client counter: 0b:00:00:00

===== Decryption algorithm outputs =====
-> Failed to decrypt Check-In message.
```

Test 6

This test only had decryption processing because it tests the nonce validation that only applies to the decryption process.

```
===== Decryption algorithm inputs =====
-> Check-in message payload :
06:34:67:6e:a6:e0:70:7b:7a:d7:81:4f:f8:29:5b:18:d1:9a:23:b2:e4:fa:df:82:92:53:51:7f:f3
:c9:1d:8d:47:84:2e:41:02:3c:03:ad:66:ac:4d:ca:72:47:e0:e4:c6:6b:d9:d3:99:13:e2:3d:82:3
2:b9:61:fa:92:26
-> Symmetric Key: ca:67:d4:1f:f7:11:29:10:fd:d1:8a:1b:f9:9e:a9:74
-> Client counter: 0f:00:00:00

===== Decryption algorithm outputs =====
-> Received counter: 0b:00:00:00
-> Application data : 'This is a longer longer string'
-> Invalid Check-In message - Nonce and received counter do not match.
```

F.5. Fabric Table Vendor ID Verification Procedure Test Vector

This subsection contains a worked example of the [Fabric Table Vendor ID Verification Procedure](#).

All key pairs used are examples, not to be reused in implementations.

***** Generation of VID verification statement *****

===== Algorithm inputs =====

Fabric Root Public Key (SEC1 uncompressed point form):

-> root_public_key_bytes:

```
00000000 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df d5 58 |...>.....o.....X|
00000010 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 66 d7 |...=.!bM.....f.|
00000020 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 77 19 d9 |;..H.....e:...w..|
00000030 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f 7c bc 6a |..t}_}....~.0|.j|
00000040 de |.|
00000041
```

Fabric ID: (0x1122334455667788)

Vendor ID: (0xFFf1)

VID Verifier Signer private key (raw point):

->

9f:5e:44:86:a0:7d:56:f3:8e:63:bb:27:6b:ef:91:aa:dd:d8:9d:ce:b1:56:98:26:bc:27:ba:47:f7:97:19:ee

VID Verifier Signer uncompressed public key:

->

04:9e:c0:a3:6c:b9:3c:01:35:79:f2:9e:d3:60:4e:a9:8e:14:40:99:36:c5:59:7a:50:14:d9:6e:1d:77:88:cb:55:19:f8:d8:80:6f:e8:96:40:f8:a8:75:9d:4d:35:9b:78:30:d4:4a:a2:44:e7:b8:71:dc:71:3b:b5:59:12:16:7d

VID Verifier Signer subject key ID:

-> fa:c0:e7:d0:08:94:73:f1:5c:2d:fe:fe:54:6f:ff:a0:b0:0a:8b:b2

===== Algorithm Outputs =====

Fabric ID as 64-bits big-endian bytes: 11:22:33:44:55:66:77:88

Vendor ID as 16-bits big-endian bytes: ff:f1

vendor_fabric_binding_message := fabric_binding_version (0x01) || root_public_key || fabric_id || vendor_id

-> vendor_fabric_binding_message (len 76 bytes):

```
00000000 01 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df d5 |....>.....o.....|
00000010 58 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 66 |X...=.!bM.....f|
00000020 d7 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 77 19 |.;..H.....e:...w.|
00000030 d9 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f 7c bc |...t}_}....~.0|.j|
00000040 6a de 11 22 33 44 55 66 77 88 ff f1 |j.."3DUfw...|
0000004c
```

VID Verification Statement to-be-signed := version || vendor_fabric_binding_message || vid_signer_skid

-> vid_verification_statement_tbs (len 97 bytes):

```
00000000 21 01 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df |!....>.....o....|
00000010 d5 58 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 |.X...=!.bM.....|
00000020 66 d7 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 77 |f.;.H.....e:..w|
00000030 19 d9 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f 7c |....t}_}....~.0||
00000040 bc 6a de 11 22 33 44 55 66 77 88 ff f1 fa c0 e7 |.j.."3DUfw.....|
00000050 d0 08 94 73 f1 5c 2d fe fe 54 6f ff a0 b0 0a 8b |...s.\-..To.....|
00000060 b2 |.|
00000061
```

--> vid_verification_signer uncompressed public key:

```
00000000 04 9e c0 a3 6c b9 3c 01 35 79 f2 9e d3 60 4e a9 |....l.<.5y...`N.|
00000010 8e 14 40 99 36 c5 59 7a 50 14 d9 6e 1d 77 88 cb |..@.6.YzP...n.w..|
00000020 55 19 f8 d8 80 6f e8 96 40 f8 a8 75 9d 4d 35 9b |U....o..@..u.M5.|
00000030 78 30 d4 4a a2 44 e7 b8 71 dc 71 3b b5 59 12 16 |x0.J.D..q.q;.Y..|
00000040 7d |}|
00000041
```

--> vid_verification_signer private key:

```
9f:5e:44:86:a0:7d:56:f3:8e:63:bb:27:6b:ef:91:aa:dd:d8:9d:ce:b1:56:98:26:bc:27:ba:47:f7
:97:19:ee
```

--> Signature:

```
00000000 32 32 84 af 77 38 f9 b3 21 d0 0b bf 2b ff 94 2c |22..w8..!...+..|
00000010 82 15 c4 98 df dd b1 b6 6b 34 fb d7 db 31 8a d7 |.....k4...1..|
00000020 54 82 9c d5 ca 1d 9c f5 3b eb 2c 93 57 1a f3 52 |T.....;.,.W..R|
00000030 57 ee ad 67 14 f6 8e 91 0f 7b b7 ac 9a 90 9a 72 |W..g.....{.....r|
00000040
```

VID Verification Statement := statement_version || vid_verification_signer_skid || signature

--> vid_verification_statement (len 85 bytes):

```
00000000 21 fa c0 e7 d0 08 94 73 f1 5c 2d fe fe 54 6f ff |!.....s.\-..To.|
00000010 a0 b0 0a 8b b2 32 32 84 af 77 38 f9 b3 21 d0 0b |....22..w8..!...|
00000020 bf 2b ff 94 2c 82 15 c4 98 df dd b1 b6 6b 34 fb |.+..,.....k4..|
00000030 d7 db 31 8a d7 54 82 9c d5 ca 1d 9c f5 3b eb 2c |..1..T.....;.,|
00000040 93 57 1a f3 52 57 ee ad 67 14 f6 8e 91 0f 7b b7 |.W..RW..g.....{.|
00000050 ac 9a 90 9a 72 |....r|
00000055
```

===== Global roots case =====

-> VID verification statement implied by RCAC being a global root

- Process SignVendorIdVerificationRequest, without VID verification statement

***** Handling of SignVendorIdVerificationRequest (with device) *****

===== Algorithm inputs =====

SignVendorIdVerificationRequest's FabricIndex field: 0x05

SignVendorIdVerificationRequest's NonceChallenge field:

a1:a2:a3:a4:a5:a6:a7:a8:a9:aa:ab:ac:ad:ae:af:b0:b1:b2:b3:b4:b5:b6:b7:b8:b9:ba:bb:bc:bd
:be:bf:c0

AttestationChallenge from secure session:

90:4b:82:e6:6e:98:57:85:b4:a3:ff:18:c2:59:47:f3

NOC private key (raw point):

->

65:65:6f:be:57:31:c2:0c:c6:80:7a:86:c1:b8:ba:f4:2f:ab:a8:89:53:d2:f6:6a:44:b0:8a:67:0a
:65:4e:3f

Fabric Root Public Key (SEC1 uncompressed point form):

-> root_public_key_bytes:

```
00000000 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df d5 58 |...>.....o.....X|
00000010 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 66 d7 |...=.!bM.....f.|
00000020 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 77 19 d9 |;..H.....e:..w..|
00000030 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f 7c bc 6a |..t}_}....~.0|.j|
00000040 de                                     |.|
00000041
```

Fabric ID from fabric table: (0x1122334455667788)

Vendor ID from fabric table: (0xFFFF1)

VIDVerificationStatement: <absent>

===== Algorithm Outputs =====

- Locally generate vendor_fabric_binding_message

Fabric ID as 64-bits big-endian bytes: 11:22:33:44:55:66:77:88

Vendor ID as 16-bits big-endian bytes: ff:f1

vendor_fabric_binding_message := fabric_binding_version (0x01) || root_public_key ||
fabric_id || vendor_id

-> vendor_fabric_binding_message (len 76 bytes):

```
00000000 01 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df d5 |....>.....o.....|
00000010 58 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 66 |X...=.!bM.....f|
00000020 d7 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 77 19 |.;..H.....e:..w.|
00000030 d9 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f 7c bc |...t}_}....~.0|.j|
00000040 6a de 11 22 33 44 55 66 77 88 ff f1          |j.."3DUfw...|
0000004c
```

```
- Generate vendor_id_verification_tbs
-> vendor_id_verification_tbs := fabric_binding_version || client_challenge ||
attestation_challenge ||
                                fabric_index || vendor_fabric_binding_message ||
                                <vid_verification_statement?>
```

vendor_id_verification_tbs (len 126 bytes):

```
00000000 01 a1 a2 a3 a4 a5 a6 a7 a8 a9 aa ab ac ad ae af |.....|
00000010 b0 b1 b2 b3 b4 b5 b6 b7 b8 b9 ba bb bc bd be bf |.....|
00000020 c0 90 4b 82 e6 6e 98 57 85 b4 a3 ff 18 c2 59 47 |..K..n.W.....YG|
00000030 f3 05 01 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff |.....>.....0...|
00000040 df d5 58 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 |..X...=.!bM....|
00000050 07 66 d7 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 |.f.;..H.....e:..|
00000060 77 19 d9 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f |w....t}_}....~.0|
00000070 7c bc 6a de 11 22 33 44 55 66 77 88 ff f1      ||.j.."3DUfw...|
0000007e
```

- Sign vendor_id_verification_tbs with NOC private key

-> signature:

```
00000000 e2 af 89 2d 77 28 4b 86 c4 f0 ec 50 22 83 69 74 |...-w(K....P".it|
00000010 89 1e d4 a6 3b ad a3 48 10 9a 85 ad b6 a5 c2 28 |....;..H.....(|
00000020 03 ec 97 e4 86 83 59 24 b3 21 a8 25 f8 a0 86 84 |.....Y$.!%....|
00000030 3f 4c 9a ee c7 ff 05 5d ad e5 52 16 57 e4 59 9a |?L.....]..R.W.Y.|
00000040
```

Output SignVendorIdVerificationResponse:

-> FabricIndex: 0x05

-> FabricBindingVersion: 0x01

-> Signature:

```
00000000 e2 af 89 2d 77 28 4b 86 c4 f0 ec 50 22 83 69 74 |...-w(K....P".it|
00000010 89 1e d4 a6 3b ad a3 48 10 9a 85 ad b6 a5 c2 28 |....;..H.....(|
00000020 03 ec 97 e4 86 83 59 24 b3 21 a8 25 f8 a0 86 84 |.....Y$.!%....|
00000030 3f 4c 9a ee c7 ff 05 5d ad e5 52 16 57 e4 59 9a |?L.....]..R.W.Y.|
00000040
```

- (alternate): Verify challenge signature, NOC chain and globally trusted RCAC

***** Handling of SignVendorIdVerificationResponse (with device) *****

===== Algorithm inputs =====

AttestationChallenge from secure session:

90:4b:82:e6:6e:98:57:85:b4:a3:ff:18:c2:59:47:f3

Previously-sent nonce challenge in request:

a1:a2:a3:a4:a5:a6:a7:a8:a9:aa:ab:ac:ad:ae:af:b0:b1:b2:b3:b4:b5:b6:b7:b8:b9:ba:bb:bc:bd
:be:bf:c0

SignVendorIdVerificationResponse's FabricIndex field: 0x05

SignVendorIdVerificationResponse's Signature Version field: 0x01

SignVendorIdVerificationResponse's Signature field:

```
00000000 e2 af 89 2d 77 28 4b 86 c4 f0 ec 50 22 83 69 74 |...-w(K....P".it|
00000010 89 1e d4 a6 3b ad a3 48 10 9a 85 ad b6 a5 c2 28 |....;..H.....(|
00000020 03 ec 97 e4 86 83 59 24 b3 21 a8 25 f8 a0 86 84 |.....Y$.!.%....|
00000030 3f 4c 9a ee c7 ff 05 5d ad e5 52 16 57 e4 59 9a |?L.....]..R.W.Y.|
00000040
```

NOC cert:

-----BEGIN CERTIFICATE-----

```
MIIB2DCCAX+gAwIBAgIBATAKBggqhkJOPQQDAjAiMSAwHgYKKwYBBAGConwBAwwQ
MDAwMDAwMDAwMDAwMDAwMjAeFw0yMTAxMDEwMDAwMDBaFw0zMDEyMzAwMDAwMDBa
MEQxIDAeBgorBgEEAYKifAEFDBAwMDAwMDAwMDAwMDAwMDAxMSAwHgYKKwYBBAGC
onwBAQwQMDAwMDAwMDAwMjM0NDMyMTBZMBMGByqGSM49AgEGCCqGSM49AwEHA0IA
BA57wmb5bT+IcT0mAhIV8KTAaifJxdAmQbaGQgKbYMECC4PxM2oJ0DrLI7VIi66
2UtMcu8lU/pUyX/YSzt25LWjgYmWgYAwDAYDVR0TAQH/BAIwADA0BgNVHQ8BAf8E
BAMCB4AwIAYDVR0LAQH/BBYwFAYIKwYBBQUHAWIGCCsGAQUFBwMBMBOGA1UdDgQW
BBTUXt25i0KFMXNdC2izZ9XyUn92QTAfBgNVHSMEGDAwBS3DcaCNnwGPhP01E+T
dAiYFPsI2TAKBggqhkJOPQQDAgNHADBEAIB8FazXDfKZesqUvL0RAr5qIJfB6F/R
QgpwbLq3785IVgIgeEnxza8qMpNtGhzgfhWcpsDi/DStQapeQ1ISQRDiOT4=
```

-----END CERTIFICATE-----

ICAC cert:

-----BEGIN CERTIFICATE-----

```
MIIBLTCCATygAwIBAgIBATAKBggqhkJOPQQDAjAiMSAwHgYKKwYBBAGConwBBawQ
MDAwMDAwMDAwMDAwMDAwMTAeFw0yMTAxMDEwMDAwMDBaFw0zMDEyMzAwMDAwMDBa
MCIXIDAeBgorBgEEAYKifAEEDBAwMDAwMDAwMDAwMDAwMDAyMFkwEwYHKoZIzjOC
AQYIKoZIzjODAQcDQgAEAQkkmAIPq0p7LLkgYMUdZfWgjIcphuyDuqpETefdmAM
XJJbzRc5ipLD5fSQegPAKJdRKyn9oVeu71vYorDP16NjMGEwDwYDVR0TAQH/BAUw
AwEB/zAOBgNVHQ8BAf8EBAMCAQYwHQYDVR00BBYEFLcNxoI2fCA+E/TUT5N0AjIU
+yLZMB8GA1UdIwQYMBaAFBkNojfgcCqWP7LvV3UiFbQAMrRMAoGCCqGSM49BAMC
A0cAMEQCIH5Wh/Ne804aaFs8wR+QT4L508z0E8vm0DwqOS9YQzfGaiBN/iku1zBR
GmBmUt0zWm/whEgIzCd8tX0kTaeZk463dg==
```

-----END CERTIFICATE-----

RCAC cert:

-----BEGIN CERTIFICATE-----

```
MIIBLzCCATygAwIBAgIBATAKBggqhkJOPQQAjAiMSAwHgYKKwYBBAGConwBBAwQ
MDAwMDAwMDAwMDAwMDAwMTAeFw0yMTAxMDEwMDAwMDBaFw0zMDEyMzAwMDAwMDBa
MCIXIDAeBgorBgEEAYKifAEEDBAwMDAwMDAwMDAwMDAwMDAxMFkwEwYHKOZIzjOC
AQYIKoZIzjODAQcDQgAE1Z8+sLnCFqRvBAR/39VYF8qKPckhqmJNpRaLqQdm1zvw
jkgMOPXMLGU6ktd3GdmwBHR9X32jne22fsVPfLxq3qNjMGEwDwYDVR0TAQH/BAUw
AwEB/zA0BgNVHQ8BAf8EBAMCAQYwHQYDVR00BBYEFBkNojfgeCqcWP7LvV3Ui fbQ
AMrMB8GA1UdIwQYMBaAFBkNojfgeCqcWP7LvV3Ui fbQAMrRMAoGCCqGSM49BAMC
A0kAMEYCIQDD0q80q3nQa9/iT4hK7tBp4QUtLLvHGCDnSwDW7s0aWQIhANwRcs4M
BcDq92J61vDDxosEg8m4LZkknPk+pirVnH79
```

-----END CERTIFICATE-----

Fabric Root Public Key (SEC1 uncompressed point form) from RCAC:

-> root_public_key_bytes:

```
00000000 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df d5 58 |...>.....o.....X|
00000010 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 66 d7 |...=!.bM.....f.|
00000020 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 77 19 d9 |;.H.....e:...w..|
00000030 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f 7c bc 6a |..t}_}....~.0|.j|
00000040 de |.|
00000041
```

Fabric ID from fabric table: (0x1122334455667788)

Vendor ID from fabric table: (0xFFFF1)

VIDVerificationStatement: <absent>

===== Algorithm Outputs =====

- Validate NOC chain certification path

--> Passed

- Validate response's signature version and fabric index match expectations

--> Passed: signature version 0x01 matches expected value 0x01

--> Passed: response fabric index 0x05 matches expected value 0x05

- Locally generate vendor_fabric_binding_message

Fabric ID as 64-bits big-endian bytes: 11:22:33:44:55:66:77:88

Vendor ID as 16-bits big-endian bytes: ff:f1

vendor_fabric_binding_message := fabric_binding_version (0x01) || root_public_key ||
fabric_id || vendor_id

-> vendor_fabric_binding_message (len 76 bytes):

```
00000000 01 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df d5 |...>.....o.....|
00000010 58 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 66 |X...=!.bM.....f|
```

```

00000020  d7 3b f0 8e 48 0c d0 f5  cc 94 65 3a 92 d7 77 19 |.;...H.....e:...w.|
00000030  d9 b0 04 74 7d 5f 7d a3  9d ed b6 7e c5 4f 7c bc |...t}_}....~.0|. |
00000040  6a de 11 22 33 44 55 66  77 88 ff f1              |j.."3DUfw...|
0000004c

```

```

- Generate vendor_id_verification_tbs
-> vendor_id_verification_tbs := fabric_binding_version || client_challenge ||
attestation_challenge ||
                                fabric_index || vendor_fabric_binding_message ||
                                <vid_verification_statement?>

```

vendor_id_verification_tbs (len 126 bytes):

```

00000000  01 a1 a2 a3 a4 a5 a6 a7  a8 a9 aa ab ac ad ae af |.....|
00000010  b0 b1 b2 b3 b4 b5 b6 b7  b8 b9 ba bb bc bd be bf |.....|
00000020  c0 90 4b 82 e6 6e 98 57  85 b4 a3 ff 18 c2 59 47 |..K..n.W.....YG|
00000030  f3 05 01 04 d5 9f 3e b0  b9 c2 16 a4 6f 04 0a ff |.....>.....o...|
00000040  df d5 58 17 ca 8a 3d c9  21 aa 62 4d a5 16 8b a9 |..X...=..!..bM....|
00000050  07 66 d7 3b f0 8e 48 0c  d0 f5 cc 94 65 3a 92 d7 |.f.;...H.....e:...|
00000060  77 19 d9 b0 04 74 7d 5f  7d a3 9d ed b6 7e c5 4f |w....t}_}....~.0|
00000070  7c bc 6a de 11 22 33 44  55 66 77 88 ff f1      ||.j.."3DUfw...|
0000007e

```

```

- Verify NOC signature from SignVendorIdVerificationResponse.signature over
vendor_id_verification_tbs
--> Passed: signature matches over locally re-generated vendor_id_verification_tbs

- VID Verification Statement absent: skip to find RCAC in DCL trust store
- Find RCAC in DCL trust store
--> Extracting subject key ID from RCAC
--> RCAC Subject Key ID: 19:0d:a2:37:e0:78:2a:9c:58:fe:cb:bd:5d:d4:89:f6:d0:00:ca:d1
--> Found matching RCAC in truststore by Subject Key ID:
-----BEGIN CERTIFICATE-----
MIIBlzCCATygAwIBAgIBATAKBggqhkJOPQQDAjAiMSAwHgYKKwYBBAGConwBBAwQ
MDAwMDAwMDAwMDAwMDAwMTAeFw0yMTAxMDEwMDAwMDBaFw0zMDEyMzAwMDAwMDBa
MCIXIDAeBgorBgEEAYKifAEEDBAwMDAwMDAwMDAwMDAwMDAxMFkwEwYHKoZIzj0C
AQYIKoZIzj0DAQcDQgAE1Z8+sLnCFqRvBAr/39VYF8qKPckhqmJNpRaLqQdm1zvW
jkgMOPXMLGU6ktd3GdmwBHR9X32jne22fsVPfLxq3qNjMGEwDwYDVR0TAQH/BAUw
AwEB/zA0BgNVHQ8BAf8EBAMCAQYwHQYDVR00BBYEFBkNojfgeCqcWP7LvV3Ui fbQ
AMrRMB8GA1UdIwQYMBaAFBkNojfgeCqcWP7LvV3Ui fbQAMrRMAoGCCqGSM49BAMC
A0kAMEYCIQDD0q80q3nQa9/iT4hK7tBp4QuTLLvHGCDnSwDW7s0aWQIhANwRcs4M
BcDq92J61vDDxosEg8m4LZkknPk+pirVnH79
-----END CERTIFICATE-----

```

- Re-validate NOC chain with RCAC replaced with the one found in DCL

--> Validate NOC against the known root successfully!

**** All steps completed, the fabric 0x1122334455667788 under root

04:d5:9f:3e:b0:b9:c2:16:a4:6f:04:0a:ff:df:d5:58:17:ca:8a:3d:c9:21:aa:62:4d:a5:16:8b:a9:07:66:d7:3b:f0:8e:48:0c:d0:f5:cc:94:65:3a:92:d7:77:19:d9:b0:04:74:7d:5f:7d:a3:9d:ed:b6:7e:c5:4f:7c:bc:6a:de is associated with Vendor ID 0xFFFF1!

Successfully validated VID without VID verification statement.

===== VID verification statement case =====

- (alternate): Process SignVendorIdVerificationRequest with VID verification statement

***** Handling of SignVendorIdVerificationRequest (with device) *****

===== Algorithm inputs =====

SignVendorIdVerificationRequest's FabricIndex field: 0x05

SignVendorIdVerificationRequest's NonceChallenge field:

c1:c2:c3:c4:c5:c6:c7:c8:c9:ca:cb:cc:cd:ce:cf:d0:d1:d2:d3:d4:d5:d6:d7:d8:d9:da:db:dc:dd:de:df:e0

AttestationChallenge from secure session:

14:5f:2e:f3:9e:3e:4f:fc:f5:64:9c:09:09:cb:c8:e4

NOC private key (raw point):

->

65:65:6f:be:57:31:c2:0c:c6:80:7a:86:c1:b8:ba:f4:2f:ab:a8:89:53:d2:f6:6a:44:b0:8a:67:0a:65:4e:3f

Fabric Root Public Key (SEC1 uncompressed point form):

-> root_public_key_bytes:

```
00000000 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df d5 58 |...>....o....X|
00000010 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 66 d7 |...=.!bM....f.|
00000020 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 77 19 d9 |;.H....e:..w..|
00000030 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f 7c bc 6a |..t}_}....~.0|.j|
00000040 de |.|
00000041
```

Fabric ID from fabric table: (0x1122334455667788)

Vendor ID from fabric table: (0xFFFF1)

VIDVerificationStatement:

```
00000000 21 fa c0 e7 d0 08 94 73 f1 5c 2d fe fe 54 6f ff |!.....s.\-...To.|
00000010 a0 b0 0a 8b b2 32 32 84 af 77 38 f9 b3 21 d0 0b |....22..w8...!..|
00000020 bf 2b ff 94 2c 82 15 c4 98 df dd b1 b6 6b 34 fb |.+..,.....k4.|
00000030 d7 db 31 8a d7 54 82 9c d5 ca 1d 9c f5 3b eb 2c |..1..T.....;.,|
```

```

00000040 93 57 1a f3 52 57 ee ad 67 14 f6 8e 91 0f 7b b7 |.W..RW..g.....{.|
00000050 ac 9a 90 9a 72                                |....r|
00000055

```

==== Algorithm Outputs =====

- Locally generate vendor_fabric_binding_message

Fabric ID as 64-bits big-endian bytes: 11:22:33:44:55:66:77:88

Vendor ID as 16-bits big-endian bytes: ff:f1

```

vendor_fabric_binding_message := fabric_binding_version (0x01) || root_public_key ||
fabric_id || vendor_id

```

-> vendor_fabric_binding_message (len 76 bytes):

```

00000000 01 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df d5 |....>.....o.....|
00000010 58 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 66 |X...=!.bM.....f|
00000020 d7 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 77 19 |.;..H.....e:..w.|
00000030 d9 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f 7c bc |...t}_}....~.0|. |
00000040 6a de 11 22 33 44 55 66 77 88 ff f1            |j.."3DUfw...|
0000004c

```

- Generate vendor_id_verification_tbs

```

-> vendor_id_verification_tbs := fabric_binding_version || client_challenge ||
attestation_challenge ||

```

```

fabric_index || vendor_fabric_binding_message ||
<vid_verification_statement?>

```

vendor_id_verification_tbs (len 211 bytes):

```

00000000 01 c1 c2 c3 c4 c5 c6 c7 c8 c9 ca cb cc cd ce cf |.....|
00000010 d0 d1 d2 d3 d4 d5 d6 d7 d8 d9 da db dc dd de df |.....|
00000020 e0 14 5f 2e f3 9e 3e 4f fc f5 64 9c 09 09 cb c8 |.._...>0..d.....|
00000030 e4 05 01 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff |.....>.....o...|
00000040 df d5 58 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 |..X...=!.bM....|
00000050 07 66 d7 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 |.f.;..H.....e:..|
00000060 77 19 d9 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f |w....t}_}....~.0|
00000070 7c bc 6a de 11 22 33 44 55 66 77 88 ff f1 21 fa ||.j.."3DUfw...!.|
00000080 c0 e7 d0 08 94 73 f1 5c 2d fe fe 54 6f ff a0 b0 |.....s.\-..To...|
00000090 0a 8b b2 32 32 84 af 77 38 f9 b3 21 d0 0b bf 2b |...22..w8..!...+|
000000a0 ff 94 2c 82 15 c4 98 df dd b1 b6 6b 34 fb d7 db |..,.....k4...|
000000b0 31 8a d7 54 82 9c d5 ca 1d 9c f5 3b eb 2c 93 57 |1..T.....;..W|
000000c0 1a f3 52 57 ee ad 67 14 f6 8e 91 0f 7b b7 ac 9a |..RW..g.....{...|
000000d0 90 9a 72                                |..r|
000000d3

```

- Sign vendor_id_verification_tbs with NOC private key

-> signature:

```

00000000 e2 af 89 2d 77 28 4b 86 c4 f0 ec 50 22 83 69 74 |...-w(K....P".it|
00000010 89 1e d4 a6 3b ad a3 48 10 9a 85 ad b6 a5 c2 28 |....;.H.....(|
00000020 ab fb b9 d1 a8 35 eb dc 7f 51 0b 8c d7 53 2c 8f |.....5...Q...S,.|
00000030 35 b3 c1 92 73 75 19 fc 14 f3 8e b4 42 68 d1 ea |5...su.....Bh..|
00000040

```

Output SignVendorIdVerificationResponse:

-> FabricIndex: 0x05

-> FabricBindingVersion: 0x01

-> Signature:

```

00000000 e2 af 89 2d 77 28 4b 86 c4 f0 ec 50 22 83 69 74 |...-w(K....P".it|
00000010 89 1e d4 a6 3b ad a3 48 10 9a 85 ad b6 a5 c2 28 |....;.H.....(|
00000020 ab fb b9 d1 a8 35 eb dc 7f 51 0b 8c d7 53 2c 8f |.....5...Q...S,.|
00000030 35 b3 c1 92 73 75 19 fc 14 f3 8e b4 42 68 d1 ea |5...su.....Bh..|
00000040

```

- (alternate): Verify challenge signature, NOC chain and VID verification statement

***** Handling of SignVendorIdVerificationResponse (with device) *****

===== Algorithm inputs =====

AttestationChallenge from secure session:

14:5f:2e:f3:9e:3e:4f:fc:f5:64:9c:09:09:cb:c8:e4

Previously-sent nonce challenge in request:

c1:c2:c3:c4:c5:c6:c7:c8:c9:ca:cb:cc:cd:ce:cf:d0:d1:d2:d3:d4:d5:d6:d7:d8:d9:da:db:dc:dd
:de:df:e0

SignVendorIdVerificationResponse's FabricIndex field: 0x05

SignVendorIdVerificationResponse's Signature Version field: 0x01

SignVendorIdVerificationResponse's Signature field:

```

00000000 e2 af 89 2d 77 28 4b 86 c4 f0 ec 50 22 83 69 74 |...-w(K....P".it|
00000010 89 1e d4 a6 3b ad a3 48 10 9a 85 ad b6 a5 c2 28 |....;.H.....(|
00000020 ab fb b9 d1 a8 35 eb dc 7f 51 0b 8c d7 53 2c 8f |.....5...Q...S,.|
00000030 35 b3 c1 92 73 75 19 fc 14 f3 8e b4 42 68 d1 ea |5...su.....Bh..|
00000040

```

NOC cert:

-----BEGIN CERTIFICATE-----

```

MIIB2DCCAX+gAwIBAgIBATAKBggqhkJOPQQDAjAiMSAwHgYKKwYBBAGConwBAwwQ
MDAwMDAwMDAwMDAwMDAwMjAeFw0yMTAxMDEwMDAwMDAwMDAwMDAwMDAwMDAwMDAw
MEQxIDAeBgorBgEEAYKIFAEFDBAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAw
onwBAQwQMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAwMDAw
BA57wmb5bT+IcT0mAhIV8KTAaIFJxdIAmQbaGQgKbYMECC4PxM2oJ0DrLI7VIi66

```

```

2UtMcu8lU/pUyX/YSzt25LWjgYMwgYAwDAYDVR0TAQH/BAIwADA0BgNVHQ8BAf8E
BAMCB4AwIAYDVR0LAQH/BBYwFAYIKwYBBQUHAWIGCCsGAQUFBwMBMB0GA1UdDgQW
BBTUXt25i0KFMXNdc2izZ9XyUn92QTAFBgNVHSMEGDAWgBS3DcaCNnwgPhP01E+T
dAIyFPsi2TAKBggqhkJOPQDAGNHADBEAiB8FazXDfKZesqUvL0RAR5qIJfB6F/R
QgpwbLq3785IVgIgeEnxza8qMpNtGhzgfhWcpsDi/DStQapeQ1ISQRDiOT4=
-----END CERTIFICATE-----

```

ICAC cert:

```

-----BEGIN CERTIFICATE-----
MIIBLTCCATygAwIBAgIBATAKBggqhkJOPQDAjAiMSAwHgYKKwYBBAGConwBBawQ
MDAwMDAwMDAwMDAwMDAwMTAeFw0yMTAxMDEwMDAwMDBaFw0zMDEyMzAwMDAwMDBa
MCIXIDAeBgorBgEEAYKifAEEDBAwMDAwMDAwMDAwMDAwMDAyMFkwEwYHKoZIzj0C
AQYIKoZIzj0DAQcDQgAEAQkkmAiPq0p7LLkgYMUSDzfWgjIcphuyDuqpETefdMAM
XJJbzRc5ipLD5fSQegPAKJdRKyn9oVeu71vYorDP16NjMGEwDwYDVR0TAQH/BAUw
AwEB/zA0BgNVHQ8BAf8EBAMCAQYwHQYDVR00BBYEFLnXoI2fCA+E/TUT5N0AjIU
+yLZMB8GA1UdIwQYMBaAFBkNoj fgeCqcWP7LvV3Ui fbQAMrRMAoGCCqGSM49BAMC
A0cAMEQCIIH5Wh/Ne804aaFs8wR+QT4L508z0E8vm0DwqOS9YQzfGaiBN/iku1zBR
GmBmUt0zWm/whEgIzCd8tX0kTaeZk463dg==
-----END CERTIFICATE-----

```

RCAC cert:

```

-----BEGIN CERTIFICATE-----
MIIBLzCCATygAwIBAgIBATAKBggqhkJOPQDAjAiMSAwHgYKKwYBBAGConwBBawQ
MDAwMDAwMDAwMDAwMDAwMTAeFw0yMTAxMDEwMDAwMDBaFw0zMDEyMzAwMDAwMDBa
MCIXIDAeBgorBgEEAYKifAEEDBAwMDAwMDAwMDAwMDAwMDAxMFkwEwYHKoZIzj0C
AQYIKoZIzj0DAQcDQgAEIz8+sLnCFqRvBAR/39VYF8qKPckhqmJNpRaLqQdm1zvW
jkgMOPXMLGU6ktd3GdmwBHR9X32jne22fsVPfLxq3qNjMGEwDwYDVR0TAQH/BAUw
AwEB/zA0BgNVHQ8BAf8EBAMCAQYwHQYDVR00BBYEFBkNoj fgeCqcWP7LvV3Ui fbQ
AMrRMB8GA1UdIwQYMBaAFBkNoj fgeCqcWP7LvV3Ui fbQAMrRMAoGCCqGSM49BAMC
A0kAMEYCIQDD0q80q3nQa9/iT4hK7tBp4QUtLLvHGCDnSwDW7s0aWQIhANwRcs4M
BcDq92J61vDDxosEg8m4LZkknPk+pirVnH79
-----END CERTIFICATE-----

```

Fabric Root Public Key (SEC1 uncompressed point form) from RCAC:

-> root_public_key_bytes:

```

00000000 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df d5 58 |...>.....o.....X|
00000010 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 66 d7 |...=.!.bM.....f.|

```



```

00000020 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 77 19 d9 |;..H.....e:...w..|
00000030 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f 7c bc 6a |..t}_}....~.0|.j|
00000040 de |.|
00000041

```

Fabric ID from fabric table: (0x1122334455667788)

Vendor ID from fabric table: (0xFFFF1)

VIDVerificationStatement:

```

00000000 21 fa c0 e7 d0 08 94 73 f1 5c 2d fe fe 54 6f ff |!.....s.\-...To.|
00000010 a0 b0 0a 8b b2 32 32 84 af 77 38 f9 b3 21 d0 0b |....22..w8...!..|
00000020 bf 2b ff 94 2c 82 15 c4 98 df dd b1 b6 6b 34 fb |.+...k4.|
00000030 d7 db 31 8a d7 54 82 9c d5 ca 1d 9c f5 3b eb 2c |..1..T.....;.,|
00000040 93 57 1a f3 52 57 ee ad 67 14 f6 8e 91 0f 7b b7 |.W..RW..g.....{.|
00000050 ac 9a 90 9a 72 |....r|
00000055

```

===== Algorithm Outputs =====

- Validate NOC chain certification path

--> Passed

- Validate response's signature version and fabric index match expectations

--> Passed: signature version 0x01 matches expected value 0x01

--> Passed: response fabric index 0x05 matches expected value 0x05

- Locally generate vendor_fabric_binding_message

Fabric ID as 64-bits big-endian bytes: 11:22:33:44:55:66:77:88

Vendor ID as 16-bits big-endian bytes: ff:f1

```

vendor_fabric_binding_message := fabric_binding_version (0x01) || root_public_key ||
fabric_id || vendor_id

```

-> vendor_fabric_binding_message (len 76 bytes):

```

00000000 01 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff df d5 |....>.....o.....|
00000010 58 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 07 66 |X...=.!bM.....f|
00000020 d7 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 77 19 |;..H.....e:...w..|
00000030 d9 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f 7c bc |...t}_}....~.0|.j|
00000040 6a de 11 22 33 44 55 66 77 88 ff f1 |j.."3DUfw...|
0000004c

```

- Generate vendor_id_verification_tbs

```

-> vendor_id_verification_tbs := fabric_binding_version || client_challenge ||
attestation_challenge ||

```

```

fabric_index || vendor_fabric_binding_message ||
<vid_verification_statement?>

```

vendor_id_verification_tbs (len 211 bytes):


```

00000000 01 c1 c2 c3 c4 c5 c6 c7 c8 c9 ca cb cc cd ce cf |.....|
00000010 d0 d1 d2 d3 d4 d5 d6 d7 d8 d9 da db dc dd de df |.....|
00000020 e0 14 5f 2e f3 9e 3e 4f fc f5 64 9c 09 09 cb c8 |...>0..d....|
00000030 e4 05 01 04 d5 9f 3e b0 b9 c2 16 a4 6f 04 0a ff |.....>.....0...|
00000040 df d5 58 17 ca 8a 3d c9 21 aa 62 4d a5 16 8b a9 |..X...=!.bM....|
00000050 07 66 d7 3b f0 8e 48 0c d0 f5 cc 94 65 3a 92 d7 |.f.;..H.....e:..|
00000060 77 19 d9 b0 04 74 7d 5f 7d a3 9d ed b6 7e c5 4f |w....t}_}....~.0|
00000070 7c bc 6a de 11 22 33 44 55 66 77 88 ff f1 21 fa ||.j.."3DUfw...!.|
00000080 c0 e7 d0 08 94 73 f1 5c 2d fe fe 54 6f ff a0 b0 |....s.\-.To...|
00000090 0a 8b b2 32 32 84 af 77 38 f9 b3 21 d0 0b bf 2b |...22..w8...!...+|
000000a0 ff 94 2c 82 15 c4 98 df dd b1 b6 6b 34 fb d7 db |..,.....k4...|
000000b0 31 8a d7 54 82 9c d5 ca 1d 9c f5 3b eb 2c 93 57 |1..T.....;.,.W|
000000c0 1a f3 52 57 ee ad 67 14 f6 8e 91 0f 7b b7 ac 9a |..RW..g.....{...|
000000d0 90 9a 72 |..r|
000000d3

```

- Verify NOC signature from SignVendorIdVerificationResponse.signature over vendor_id_verification_tbs

--> Passed: signature matches over locally re-generated vendor_id_verification_tbs

- VID Verification Statement present: find VID verification signer cert

--> Extracting subject key ID from vid_verification_statement

--> Fabric binding version 0x01 matches expected version 0x01

--> Statement version 0x21 matches expected statement version 0x21

--> VID Verification Signer Subject Key ID:

fa:c0:e7:d0:08:94:73:f1:5c:2d:fe:fe:54:6f:ff:a0:b0:0a:8b:b2

--> Found VID Verification Signer cert:

-----BEGIN CERTIFICATE-----

```

MIICGjCCAcCgAwIBAgIUbvVpiZNHsn7VZ1Jh/3TXFzVTSpMwCgYIKoZIzj0EAwIw
YTERMA8GA1UECgwIQWntZSBMTkxkYjAoBgNVBAMMIVZJRCBGRkYxIFZlbmRvcjB2
ZXJpZmLjYXRpb24gcm9vdDEgMB4GCisGAQQBgqJ8AQQMEExMjIzMzQ0Q0UFCQkND
REQwHhcNMjUwNTMwMTkzMjUxWWhcNMzUwNTI4MTkzMjUxWjBjLMREwDwYDVQQKDAhB
Y21lIEExMQZuEuMCwGA1UEAwwLVkLEIEZGRjEgVmVuZG9yIHZlcmlmaWNhdGlubiBz
aWduZXIEMTEgMB4GCisGAQQBgqJ8AQCMEIJCruVEREVMDAxMTIyMzUwMTkzMjUxWjBj
hkjOPQIBBgqhkhjOPQMBBwNCAASewKNsuTwBNXnyntNgTqm0FECZNSVZeLAU2W4d
d4jLVrN42IBv6JZA+Kh1nU01m3gw1EqiR0e4cdxx07VZEhZ9o1IwUDA0BgNVHQ8B
Af8EBAMCB4AwHQYDVR00BBYEFPrA59AI LHPxXC3+/LRv/6CwCouyMB8GA1UdIwQY
MBaAFBBSQdgvXtuXpsRezyXJkYujUFABMAoGCCqGSM49BAMCA0gAMEUCIH416CEz
dF1DQJnQKjtwHqPlYZ5jIvqFIBqHcnVbKssWAiEAqoNk7pNjg0AsW5E+T0Dqx9aS
YWY4sBvL3cod7wE7/D4=

```

-----END CERTIFICATE-----

```
--> VID Verification Signer certificate path was validated successfully with DCL-
based trust store.
```

```
- Locally regenerate vendor_id_verification_tbs from vendor_fabric_binding_message of
prior step
```

```
VID Verification Statement to-be-signed := version || vendor_fabric_binding_message ||
vid_signer_skid
```

```
-> vid_verification_statement_tbs (len 97 bytes):
```

```
00000000  21 01 04 d5 9f 3e b0 b9  c2 16 a4 6f 04 0a ff df  |!....>.....o....|
00000010  d5 58 17 ca 8a 3d c9 21  aa 62 4d a5 16 8b a9 07  |.X...=!.bM.....|
00000020  66 d7 3b f0 8e 48 0c d0  f5 cc 94 65 3a 92 d7 77  |f.;..H.....e:..w|
00000030  19 d9 b0 04 74 7d 5f 7d  a3 9d ed b6 7e c5 4f 7c  |...t}_}....~.0||
00000040  bc 6a de 11 22 33 44 55  66 77 88 ff f1 fa c0 e7  |.j.."3DUfw.....|
00000050  d0 08 94 73 f1 5c 2d fe  fe 54 6f ff a0 b0 0a 8b  |...s.\-..To.....|
00000060  b2                                |.|
00000061
```

```
- Validate vid_verification_statement signature
```

```
--> vid_verification_signer_public_key_bytes:
```

```
00000000  04 9e c0 a3 6c b9 3c 01  35 79 f2 9e d3 60 4e a9  |....l.<.5y...`N.|
00000010  8e 14 40 99 36 c5 59 7a  50 14 d9 6e 1d 77 88 cb  |..@.6.YzP..n.w..|
00000020  55 19 f8 d8 80 6f e8 96  40 f8 a8 75 9d 4d 35 9b  |U....o..@..u.M5.|
00000030  78 30 d4 4a a2 44 e7 b8  71 dc 71 3b b5 59 12 16  |x0.J.D..q;q;.Y..|
00000040  7d                                |}|
00000041
```

```
--> vid_verification_statement signature validated!
```

```
**** All steps completed, the fabric 0x1122334455667788 under root
```

```
04:d5:9f:3e:b0:b9:c2:16:a4:6f:04:0a:ff:df:d5:58:17:ca:8a:3d:c9:21:aa:62:4d:a5:16:8b:a9
:07:66:d7:3b:f0:8e:48:0c:d0:f5:cc:94:65:3a:92:d7:77:19:d9:b0:04:74:7d:5f:7d:a3:9d:ed:b
```

```
6:7e:c5:4f:7c:bc:6a:de is associated with Vendor ID 0xFFF1!
```

```
Successfully validated VID with VID verification statement.
```


Appendix G: Minimal Resource Requirements

This is a list of various resources required by a Node implementation, along with references to where the minimal requirements for each resource type are defined.

Resource	Minimal requirement definition
Fabric	Section 11.18.5, “Attributes”
Operational Certificate	Section 11.18.5, “Attributes”
ACL Entry	ACL
ACL Entry subject	AccessControlEntryStruct
ACL Entry target	AccessControlEntryStruct
Scene	"Maximum Number of Scenes" in Scenes Management cluster in the Cluster Library
Binding	Section 9.6.1, “Binding Mutation”
Group	Section 2.11.1.2, “Group Limits”
Group key	Section 2.11.1.2, “Group Limits”
Group peer state entries	Section 4.18.2, “Group Peer State”
Read path	Section 2.11.2.1, “Read Interaction Limits”
Subscription	Section 2.11.2.2, “Subscribe Interaction Limits”
Subscription path	Section 2.11.2.2, “Subscribe Interaction Limits”
IPv6 multicast group	Section 2.11.1.2, “Group Limits”
IPv6 Prefix	Section 4.2.2, “Matter Node Behavior”
IPv6 route	Section 4.2.2, “Matter Node Behavior”
IPv6 neighbor cache entry	Section 4.2.2, “Matter Node Behavior”
CASE session	Section 4.14.2.8, “Minimal Number of CASE Sessions”
ICD Check-In registered clients	ClientsSupportedPerFabric
TLS	Section 14.3.3, “TLS Minimum Requirements”